Bytes to Threats: Byte-Sequence Transformer For Malware Detection

1st Abdelrahman Yassin Eljamal
Department of Computer Science
University of Rhode Island
Kingston, USA
abdelrahman.eljamal@uri.edu

2nd Abdeltawab Hendawi
Department of Computer Science
University of Rhode Island
Kingston, USA
hendawi@uri.edu

arXiv:2610.18787v1 [cs.CR] 13 Oct 2026

***Abstract*—Existing malware detection approaches heavily rely on handcrafted features derived from static or dynamic analysis of the executable files. These methods require specialized domain expertise, dedicated analysis tools, and time-consuming analysis, which limits their scalability and robustness in the face of the increasing complexity and diversity of modern malware. Consequently, there is a growing interest in approaches that operate directly on raw bytes, which contain rich semantic and structural information. However, capturing long-range dependencies within executable files remains challenging for conventional static and CNN-based methods. This paper presents MalwareFormer, a transformer-based, byte-level malware detection framework that operates directly on raw bytes, without requiring domain knowledge or feature extraction. MalwareFormer achieves an accuracy of 98.35%, a balanced accuracy of 98.32%, a macro-F1 of 98.32%, and a false positive rate of 1.94%, outperforming both baselines across all threshold-dependent classification metrics and demonstrating that transformers provide a powerful and scalable approach for byte-level malware detection.**

***Index Terms*—Malware detection, Byte-level analysis, Transformer models, Deep learning, Static analysis, Dynamic analysis, Malware classification**

I. Introduction

The rapid growth of malware variants poses a significant and escalating economic threat. Recent estimates project that the global cost of cybercrime will reach approximately $10.5 trillion annually by 2025, underscoring its scale as one of the largest economic threats worldwide. At the organizational level, the average cost of a single data breach is estimated at $4.44 million globally. Ransomware and phishing attacks remain dominant, accounting for approximately 44% and 16% of reported breaches, respectively [1]. In the past year alone, cyber incidents led to substantial financial losses, including an estimated $3.09 billion impact on UnitedHealth and a $50 million ransom demand targeting CDK Global [2]. According to Kaspersky [3], their detection systems identified an average of 467,000 malicious files per day in 2024, with Trojan-based attacks showing a notable increase compared to 2023. Despite continuous advancements, anti-virus technologies are struggling to keep pace with the rapid evolution and sophistication of modern malware, highlighting the need for more adaptive detection methods.

Existing machine learning approaches rely heavily on handcrafted features extracted from executable files, including header information, imported libraries, and opcode statistics [4], [5]. While effective, these methods require significant domain expertise and extensive effort to extract such features. For instance, dynamic analysis captures behavioral and runtime features by executing malware in a controlled environment, such as a virtual machine. This allows the collection of system call traces, registry changes, file and network activity, and process behaviors, but the approach is slow, resource-intensive, and difficult to scale. Static analysis, in contrast, operates directly on the file without execution, making it faster and less resource-demanding. However, its effectiveness remains limited by the handcrafted features and specialized tools required for tasks such as extracting PE header information, imported libraries, opcode sequences, strings, and metadata, as well as performing reverse engineering and parsing of embedded resources, all of which necessitate expert knowledge.

Deep learning approaches have emerged as a promising alternative by enabling direct processing of raw bytes without feature engineering. CNN-based models such as MalConv [6] demonstrated the feasibility of learning patterns from entire executable files; however, their reliance on local convolutional filters limits their ability to capture long-range dependencies inherent in executable files. Transformer architectures, with their global self-attention mechanism, offer a better framework for modeling byte sequences.

To address these limitations, this paper proposes **MalwareFormer**, a transformer-based, byte-level malware detection framework designed to operate directly on raw bytes without requiring hand-crafted features or domain knowledge. At its core, **MalwareFormer** leverages a transformer model, ByteFormer [7], which is fine-tuned within our framework to efficiently capture long-range dependencies in malware byte sequences. We conducted extensive experiments to evaluate the model's effectiveness, comparing training from scratch with fine-tuning and analyzing their impact on convergence speed and final performance. We further examine how

© 2026 IEEE. Personal use of this material is permitted. Permission from IEEE must be obtained for all other uses, in any current or future media, including reprinting/republishing this material for advertising or promotional purposes, creating new collective works, for resale or redistribution to servers or lists, or reuse of any copyrighted component of this work in other works.
DOI: 10.1109/ICAD69378.2026.11608884

varying input sequence lengths influence model accuracy and stability. To contextualize its performance, we compare **MalwareFormer** against two strong CNN-based baselines, MalConv2 [8] and AvastNet [9], using a 50K-byte input size. **MalwareFormer** achieves an accuracy of **98.35%**, a balanced accuracy of **98.32%**, a macro-F1 of **98.32%**, and a false positive rate of **1.94%**, outperforming both baselines across all threshold-dependent classification metrics. MalConv attains 95.71% accuracy, a macro-F1 of 95.64%, and a false positive rate of 3.56%, while AvastNet achieves 97.98% accuracy, a macro-F1 of 97.94%, and a false positive rate of 2.46%. Notably, while AvastNet attains a marginally higher ROC-AUC (99.68% vs. 99.61%), MalwareFormer yields a better trade-off between detection rate and false alarm rate at the operating threshold used in practice. These results demonstrate that transformer-based byte modeling enables **MalwareFormer** to capture long-range structural dependencies within executables that CNN-based approaches struggle to model, translating into more reliable detection under realistic deployment conditions.

This paper investigates the capabilities of transformer-based, byte-level malware detection by addressing the following research questions. **RQ1:** *How effective is a transformer-based model compared to established CNN-based baselines in classifying malware?* **RQ2:** *How does fine-tuning pretrained representations impact training convergence and detection performance?* **RQ3:** *Can accurate detection be achieved without processing the entire executable file?* **RQ4:** *How does entropy affect classification accuracy, and what asymmetric failure modes emerge in high-entropy regions?*

Our key contributions are summarized as follows:

- We present **MalwareFormer**, a byte-level transformer framework for malware detection that incorporates dataset integrity verification, label consistency validation, byte-sequence extraction, and model fine-tuning, achieving **98.35%** accuracy and a false positive rate of **1.94%** without requiring handcrafted features or domain expertise.
- Investigate the transferability of pretrained ByteFormer representations to the malware domain by fine-tuning the model and analyzing their impact on training convergence and detection performance across varying input sequence lengths. To the best of our knowledge, this is the first work to utilize the ByteFormer model for malware detection.
- Compare this transformer-based model with established CNN-based baselines, providing empirical evidence of its effectiveness for static malware detection.
- We conduct an entropy-stratified analysis of model predictions, revealing an asymmetric failure mode in high-entropy regions (entropy 7–8) where FPR reaches 14.24% while FNR drops to 0.08%, exposing a fundamental limitation of byte-level detection in entropy-ambiguous regimes and motivating future entropy-aware approaches.

The rest of our paper is organized as follows. Section II reviews related work. Section III presents the **MalwareFormer** framework, detailing its transformer-based, byte-level architecture and the dataset employed in our experiments. Section IV provides the experimental setup. Section V presents the experimental results. Section VI provides an analysis and discussion of the results. Finally, Section VII concludes our work.

II. Related Work

This section reviews prior research on malware detection using raw bytes as input, beginning with convolutional neural network (CNN)–based approaches and progressing to more recent Transformer-based architectures.

A. CNN Based Models

Early work demonstrated that deep learning models can effectively analyze malware directly from raw executable bytes, eliminating the reliance on manual feature engineering. Reference [6] pioneered this direction with MalConv, a convolutional neural network designed specifically for Windows PE files. MalConv processes executable files as byte sequences using a gated convolutional layer followed by global max pooling and a fully connected classifier. Their results showed that meaningful representations can be learned directly from raw bytes, capturing global structural patterns present within the entire executable file. Building upon this idea, reference [9] investigated deeper convolutional architectures trained exclusively on raw executable files. Their study confirmed that CNNs can generalize well without hand-crafted input features and systematically analyzed architectural factors that influence detection performance. These findings strengthened the evidence that binary-level features can be learned directly from byte sequences. A major challenge in byte-level malware analysis is the extreme input length, which can exceed several million bytes. This limitation was addressed by proposing memory-efficient sequence classification methods with constant memory complexity, enabling processing of very long byte sequences [8]. Their approach increased the scalability of raw-byte malware detection and facilitated practical large-scale deployment. Other works focus on non-executable documents such as PDF files. Reference [10] models PDF files as raw byte streams and applies convolutional neural networks to learn discriminative byte-level representations without relying on explicit file parsing. Their approach captures spatial and structural patterns within byte sequences to identify malicious content embedded in non-executable documents. In contrast, byte-stream–based detection extended to the file level by aggregating stream-level predictions across an entire document [11]. This aggregation strategy enables more robust file-level malware detection, improving detection performance while maintaining resilience against obfuscation techniques in document-based malware.

Hybrid approaches have also been explored. In [12], they combined CNN layers (for local feature extraction) with LSTM layers (for long-range dependencies), demonstrating improved performance by modeling both spatial and sequential patterns in byte streams.

B. Transformer Based Models

Since the introduction of the Transformer architecture [13], Transformer-based models have achieved state-of-the-art performance across a wide range of domains beyond natural language processing, including malware detection. Reference [14] provided a comprehensive survey of Transformer-based approaches for malware classification, noting that the majority of existing methods rely on engineered or extracted features rather than operating directly on raw binary data.

Several studies apply Transformers to sequential representations of program content, particularly assembly code. Works such as [15]–[17] model assembly instructions as token sequences, enabling Transformers to capture long-range dependencies and instruction-level semantics for malware detection.

Beyond linear sequences, In [18], the author introduced a Transformer architecture designed to model function-call graphs, generating graph-based embeddings that capture semantic relationships in program behavior. This approach demonstrates the flexibility of Transformers in handling structured representations beyond flat sequences.

Other researchers have explored transforming binary data into image-like representations suitable for vision Transformers. Seneviratne *et al.* [19] and subsequent works [20], [21] demonstrate that such visual encodings can reveal spatial patterns in malware binaries that are difficult to capture using traditional static analysis.

In addition to code and binary representations, Transformer-based models have been applied to higher-level feature types, including manifest information [22]–[24] and network behavior data [25]–[27]. These approaches further demonstrate the versatility of Transformers in modeling heterogeneous sources of malware-related information.

1) Byte Only Transformer Models: Despite substantial progress, most existing Transformer-based malware detection methods continue to rely on extracted or handcrafted features (assembly instructions, graphs, images, etc.). Reference [28] introduced ByteEnc, a byte-level pre-trained Transformer encoder designed for non-executable document formats such as PDF, HWP, and MS Office files. The work proposed two novel pre-training objectives, Real-or-Fake Prediction (RFP) and Same-Sequence Prediction (SSP), combined with masked language modeling (MLM) and next-sentence prediction (NSP). It further demonstrated that relative positional encodings and byte-level unigram tokenization improve performance on byte-level understanding tasks.

Our work focuses exclusively on raw byte input, employing a transformer-based architecture within the proposed **MalwareFormer** framework to capture long-range dependencies in executable files.

III. Methodology

Fig. 1 illustrates our proposed framework, **MalwareFormer**. The process begins with an integrity check of the dataset files, followed by label consistency validation through cross-verification of samples using VirusTotal [29]. A diversity check is then conducted to ensure variability among samples in terms of file size and entropy. In the second phase, raw byte sequences are fed into the model described in the Model Overview subsection (III-A). Finally, the model is evaluated using accuracy, balanced accuracy, macro-F1, false positive rate (FPR), false negative rate (FNR), and ROC-AUC. We further conduct an entropy-stratified analysis of model predictions across six entropy bins to identify entropy-induced false positive bias and characterize regimes where byte-level detection is inherently limited.

A. Model Overview

We follow the Byteformer model [7]. It is a transformer-based encoder model designed to capture contextual representations directly from byte-only input sequences without requiring any modality-specific preprocessing. The model consists of a learnable positional embedding table, an initial convolutional token-reduction module that downsamples the raw byte stream, and a stack of 12 windowed transformer layers. Within these 12 layers, 6 token merging operations are distributed at layers (0, 1, 3, 5, 7, 9) to hierarchically reduce the sequence length and aggregate local byte-level patterns. Each transformer layer applies multi-head self-attention with windowed attention, followed by a feedforward network. After the final transformer block, a global mean pooling operation is applied to the token representations before feeding them into a linear classification head.

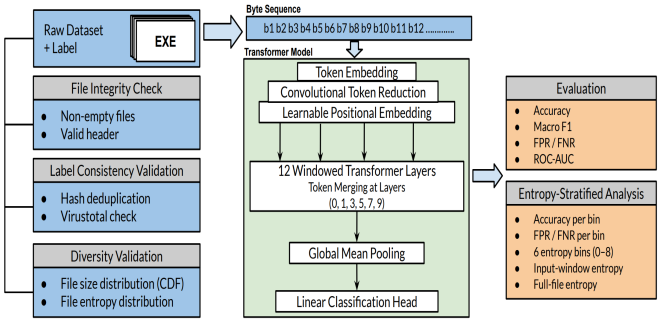

Fig. 1. MalwareFormer Framework.

B. Dataset

For this task, we used "PE Malware Machine Learning Dataset" [30], which is collected from multiple public repositories, including VirusShare, MalShare, and TheZoo, and contains 201,549 executable files, 86,812 benign, and 114,737 malicious files.

Fig. 2 illustrates the cumulative distribution function (CDF) of file sizes for benign and malicious executable files. The two curves are highly overlapped, indicating that benign and malicious samples share very similar size distributions overall. Benign executable files predominantly start at file sizes above 10^3 bytes, while malicious samples include smaller binaries, with the distribution extending below 10^3 bytes. In particular, malware files exhibit finer-grained steps at smaller sizes, with

approximately five distinct steps down to 512 bytes and a non-negligible cumulative probability above 10^{-3}. Despite this difference in the lower-size regime, the two curves converge quickly as file size increases, remaining close across the majority of the distribution. This substantial overlap suggests that file size alone is insufficient to reliably distinguish benign and malicious executable files, motivating the need for content-aware byte-level modeling approaches.

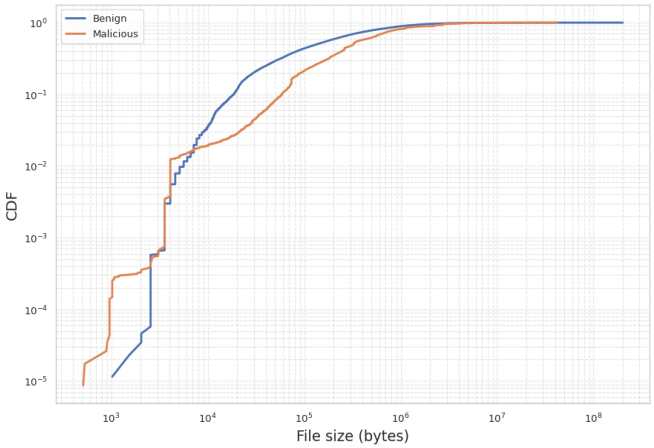

Fig. 2. Cumulative Distribution of File Sizes

Entropy is employed as a descriptive metric to characterize the statistical properties of the dataset. File entropy is computed using Shannon entropy with probabilities estimated via the maximum likelihood estimator (MLE) over byte-level distributions, which provides insight into the degree of randomness present in benign and malicious files; therefore, entropy is used as an indicator for high randomness, packing, compression, and encryption in executable files [31], [32].

Fig. 3 shows a peak in the benign distribution around entropy of 5.8 to 6.3, which corresponds to typical compiled executable files that contain a mixture of machine code, static data, and limited compression. In contrast, malware distribution is more dispersed, indicating greater structural variability in malware binaries. Another notable feature is the presence of localized spikes in the malicious entropy curve, especially in the high-entropy tail. These spikes suggest that malware does not form a smooth statistical population but instead consists of multiple sub-populations, likely corresponding to different malware families or packing strategies. Each family may rely on a specific packer or encryption scheme, resulting in characteristic entropy values.

IV. Experimental Setup

We conducted multiple experiments to study its behavior under different training strategies, including fine-tuning from a model pre-trained on image bytes and training from scratch on executable files. In addition, we evaluated the impact of different input lengths by feeding the model the first 50K, 100K, 150K, and 200K bytes of each executable. This analysis is motivated by the large variation in file sizes within our

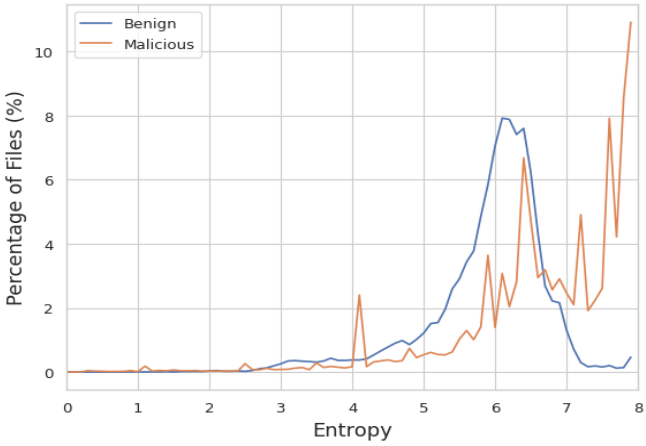

Fig. 3. Percentage of Files as a Function of Entropy (0–8)

dataset, which ranges from approximately 500 bytes to more than 2 million bytes for larger files.

All experiments were conducted within the CoreNet framework [33], which natively supports ByteFormer. To ensure a fair and reproducible comparison, we integrated the additional baseline models (MalConv and AvastNet) into the same framework and trained each model using the hyperparameter configuration reported in its original publication whenever possible. For consistency, all models were trained for 30 epochs.

For ByteFormer, training was performed using the AdamW optimizer with a weight decay of 0.05, excluding batch normalization layers and bias parameters from weight decay. The optimizer parameters were set to $\beta_1 = 0.9$ and $\beta_2 = 0.999$. A cosine learning rate scheduler was used with a linear warmup of 1,000 iterations, using an initial learning rate of 1×10^{-6}, a peak learning rate of 3×10^{-5}, and a minimum learning rate of 1×10^{-6}. The scheduler was configured to update on an epoch basis.

For MalConv and AvastNet, we followed the training configuration reported by the original authors. Both models were trained with Adam optimizer ($\beta_1 = 0.9$, $\beta_2 = 0.999$) with a fixed learning rate of 1×10^{-3} and no weight decay. Training was conducted for 30 epochs with a batch size of 128. The input size was fixed at 50,000 bytes, and the cross-entropy loss was used for binary classification.

V. Results

Fig. 4 compares the training and validation loss curves for fine-tuning and training **MalwareFormer** from scratch. The fine-tuned model, which uses pre-trained weights, converges significantly faster, with the validation loss decreasing rapidly and falling below 0.1 by approximately epoch 11 of the 30-epoch training schedule. In contrast, the model trained from scratch over 20 epochs maintains a higher validation loss throughout the same period.

Although the fine-tuned model starts with a higher training loss in the first epoch, it drops sharply after the initial epochs,

reaching a lower training loss than the model trained from scratch by the second epoch. These observations highlight the benefits of leveraging pre-trained representations to accelerate convergence and improve model stability.

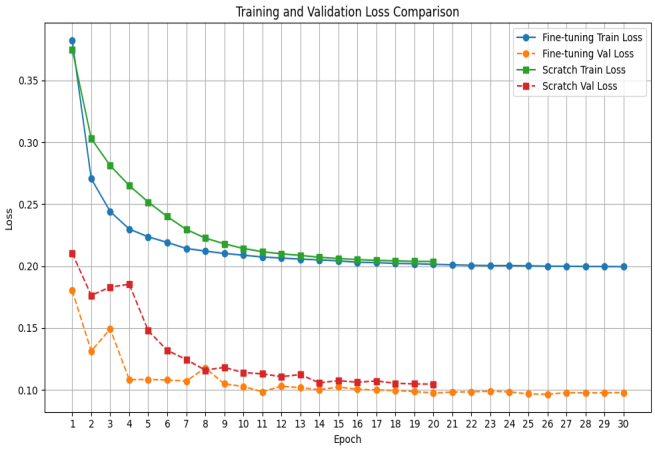

Fig. 4. Training Loss Across Epochs: Fine-Tuning vs. Training from Scratch

A. Effect of Input Sequence Length

Table I reports the effect of input sequence length on classification performance for both fine-tuning and training from scratch. For the fine-tuned model, increasing the input size from 50K to 100K bytes results in consistent improvements in both accuracy and AUC. Further increasing the input size to 150K or 200K bytes yields only marginal gains.

Overall, these results indicate that **MalwareFormer** achieves stable performance even with truncated byte sequences, while larger input sizes can provide modest additional benefits.

TABLE I
MODEL RESULTS WITH DIFFERENT INPUT SIZES.

Input Size	Acc. (%)	ROC-AUC. (%)
Fine-Tuning		
50K	98.13	99.17
100K	98.45	99.21
150K	98.35	99.45
200K	98.48	99.42

B. Comparison with CNN-Based Malware Detectors

To evaluate the effectiveness of MalwareFormer, we compare its performance against two widely used CNN-based malware detection models, MalConv [8] and AvastNet [9], using a 50K-byte input size across all models. As shown in Table II, **MalwareFormer** achieves the highest overall classification performance, with an accuracy of **98.35%**, a balanced accuracy of **98.32%**, a macro-F1 of **98.32%**, and a false positive rate (FPR) of **1.94%**, the lowest among all evaluated models. Although AvastNet attains a marginally higher ROC-AUC (99.68% vs. 99.61%), MalwareFormer consistently outperforms both CNN-based baselines across all threshold-dependent metrics, which are more directly relevant to practical deployment where false alarm rates and missed detections carry real operational costs.

On the shared test set of 30,233 samples (17,211 malware, 13,022 benign), MalwareFormer correctly detects 16,966 malware samples while missing only 245 (FNR: 1.42%), and correctly classifies 12,769 benign files with just 253 false alarms (FPR: 1.94%). Compared to AvastNet, this represents 46 fewer missed malware detections and 67 fewer false positives on benign files, demonstrating improved reliability under realistic class distributions.

TABLE II
PERFORMANCE COMPARISON OF MALWAREFORMER, MALCONV, AND AVASTNET (50K-BYTE INPUT). FPR = FALSE POSITIVE RATE; FNR = FALSE NEGATIVE RATE. BEST RESULTS PER METRIC IN **BOLD**.

Model	Acc. (%)	Bal. Acc. (%)	Macro F1 (%)	FPR (%)	FNR (%)	ROC-AUC (%)
MalwareFormer	**98.35**	**98.32**	**98.32**	**1.94**	**1.42**	99.61
AvastNet	97.98	97.93	97.94	2.46	1.69	**99.68**
MalConv	95.71	95.80	95.64	3.56	4.84	99.14

VI. DISCUSSION

The results demonstrate that pre-training on byte-level data provides a strong initialization for downstream malware classification. The faster convergence of the fine-tuned model suggests that the pre-trained weights capture generalizable byte-level patterns that transfer effectively to executable files.

The initially higher training loss observed during the first epoch of fine-tuning is consistent with the semantic differences between image bytes and executable bytes. Despite this mismatch, the model rapidly adapts, achieving lower training loss than the scratch model by epoch two.

The pre-trained model was trained on byte sequences extracted from TIFF files. Although TIFF is an image format, it is also a structured binary container composed of headers, metadata, compressed blocks, and padding. Because **MalwareFormer** operates directly on raw bytes without visual inductive biases, it learns representations that capture the statistical and structural properties of binary streams rather than image semantics. These learned patterns likely facilitate transfer to executable analysis, helping explain the improved convergence and generalization behavior observed during fine-tuning.

To further characterize model behaviour, we stratify classification performance by input-window entropy, examining how byte randomness affects detection accuracy, FPR, and FNR across six entropy bins. As shown in Table III, the model maintains balanced and stable error rates across the moderate entropy range (0–7), with both FPR and FNR remaining below 3.5%. However, in the highest entropy bin (7–8), which contains 5,011 samples predominantly corresponding to packed or encrypted content, the model's FPR spikes to **14.24%** while FNR drops to near zero (0.08%). This reveals that the model has learned to strongly associate extreme entropy with malware, correctly capturing a common obfuscation pattern but at the cost of over-flagging high-entropy benign files such as compressed archives or resource-heavy executables.

Notably, overall accuracy in this bin remains high (99.06%), illustrating how aggregate accuracy can mask a severe class-asymmetric failure that would be operationally significant in deployment. The same asymmetry is observed when stratifying by full-file entropy (FPR: 11.63%, FNR: 0.67% in the 7–8 bin), confirming that this behaviour reflects a genuine bias in the learned representations rather than an artifact of input truncation. These findings suggest that entropy-aware training strategies or a dedicated high-entropy classifier could meaningfully reduce false positives in this regime.

TABLE III
CLASSIFICATION PERFORMANCE OF **MALWAREFORMER** STRATIFIED BY INPUT-WINDOW ENTROPY BIN. ACCURACY, FPR, AND FNR ARE REPORTED PER BIN. THE 7–8 BIN REVEALS A SEVERE ASYMMETRIC FAILURE: NEAR-ZERO FNR BUT A SUBSTANTIALLY ELEVATED FPR, INDICATING SYSTEMATIC OVER-FLAGGING OF HIGH-ENTROPY BENIGN FILES.

Entropy bin	Samples	Accuracy (%)	FPR (%)	FNR (%)
0–2	86	98.84	0.00	1.82
2–4	1,279	98.59	0.97	2.54
4–5	3,165	98.64	1.17	1.54
5–6	6,044	97.96	1.34	3.30
6–7	14,648	98.18	2.07	1.62
7–8	5,011	99.06	**14.24**	**0.08**

Comparison with CNN-based malware detectors highlights the advantages of transformer-based architectures for byte sequence modeling. CNNs use fixed receptive fields and local convolutional operations, which limit their ability to capture long-range dependencies. In contrast, **MalwareFormer** combines token reduction with windowed self-attention, enabling aggregation of both local and global byte-level structures. This architectural expressiveness likely contributes to its superior performance across all threshold-dependent metrics, even under truncated input constraints.

VII. CONCLUSION

In this work, we presented **MalwareFormer**, a transformer-based, byte-level malware detection framework that operates directly on raw executable bytes without requiring handcrafted features or domain expertise. Leveraging transformer self-attention, the model captures long-range byte dependencies and consistently outperforms CNN-based baselines across all threshold-dependent metrics, achieving **98.35%** accuracy, a macro-F1 of **98.32%**, and a false positive rate of **1.94%** on 50K-byte inputs.

We further analyzed the impact of input sequence length and found that while larger byte sequences can provide modest gains, the framework shows stable performance even with truncated inputs. An entropy-stratified analysis reveals a nuanced failure mode: while error rates remain balanced across moderate entropy ranges, the model exhibits an entropy-induced false positive bias in the highest entropy bin (7–8), near-zero FNR (0.08%) but an elevated FPR (14.24%), over-flagging compressed or encrypted benign files due to an overgeneralized association between extreme entropy and malware. Future work will explore entropy-aware training and hybrid approaches to address this bias.

Overall, **MalwareFormer** demonstrates that transformer architectures provide a powerful, scalable, and practical solution for byte-level malware analysis, eliminating the need for handcrafted features.

REFERENCES

[1] M. Khalil, "Cybersecurity statistics 2025: Breach costs, ransomware & ai threats," https://deepstrike.io/blog/cybersecurity-statistics-2025-threats-trends-challenges.

[2] sygnia, "Ransomware attacks in 2024: the most devastating year yet?" https://www.sygnia.co/blog/ransomware-attacks-2024/, 4 2025.

[3] kaspersky, "The cyber surge: Kaspersky detected 467,000 malicious files daily in 2024," https://www.kaspersky.com/about/press-releases/the-cyber-surge-kaspersky-detected-467000-malicious-files-daily-in-2024, Dec. 2024.

[4] G. M. and S. C. Sethuraman, "A comprehensive survey on deep learning based malware detection techniques," *Comput. Sci. Rev.*, vol. 47, p. 100529, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:255020688

[5] E. Martins, R. Sant'Ana, J. R. B. Higuera, J. A. S. Montalvo, J. B. Higuera, and D. P. Castillo, "Semantic malware classification using artificial intelligence techniques," *Computer Modeling in Engineering & Sciences*, 2025. [Online]. Available: https://api.semanticscholar.org/CorpusID:276758036

[6] E. Raff, J. Barker, J. Sylvester, R. Brandon, B. Catanzaro, and C. K. Nicholas, "Malware detection by eating a whole exe," in *AAAI Workshops*, 2017. [Online]. Available: https://api.semanticscholar.org/CorpusID:33641567

[7] M. Horton, S. Mehta, A. Farhadi, and M. Rastegari, "Bytes are all you need: Transformers operating directly on file bytes," *Trans. Mach. Learn. Res.*, vol. 2024, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:258999744

[8] E. Raff, W. Fleshman, R. Zak, H. S. Anderson, B. Filar, and M. McLean, "Classifying sequences of extreme length with constant memory applied to malware detection," *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 11, pp. 9386–9394, May 2021. [Online]. Available: https://ojs.aaai.org/index.php/AAAI/article/view/17131

[9] M. Krčál, O. Švec, M. Bálek, and O. Jašek, "Deep convolutional malware classifiers can learn from raw executables and labels only," 2018. [Online]. Available: https://openreview.net/forum?id=HkHrmM1PM

[10] Y.-S. Jeong, J. myung Woo, and A. R. Kang, "Malware detection on byte streams of pdf files using convolutional neural networks," *Secur. Commun. Networks*, vol. 2019, pp. 8 485 365:1–8 485 365:9, 2019. [Online]. Available: https://api.semanticscholar.org/CorpusID:109929301

[11] Y.-S. Jeong, M. E. Mswahili, and A. R. Kang, "File-level malware detection using byte streams," *Scientific Reports*, vol. 13, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:259024861

[12] Q. Le, O. Boydell, B. M. Namee, and M. Scanlon, "Deep learning at the shallow end: Malware classification for non-domain experts," *Digit. Investig.*, vol. 26 Supplement, pp. S118–S126, 2018. [Online]. Available: https://api.semanticscholar.org/CorpusID:49908612

[13] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, "Attention is all you need," in *Neural Information Processing Systems*, 2017. [Online]. Available: https://api.semanticscholar.org/CorpusID:13756489

[14] P. Kunwar, K. Aryal, M. Gupta, M. Abdelsalam, and E. Bertino, "Sok: Leveraging transformers for malware analysis," *IEEE Transactions on Dependable and Secure Computing*, vol. 22, pp. 5888–5905, 2024. [Online]. Available: https://api.semanticscholar.org/CorpusID:270062599

[15] L. T. Li, S. H. H. Ding, and P. Charland, "Gental: Generative denoising skip-gram transformer for unsupervised binary code similarity detection," *2023 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–8, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:260385610

[16] M. Q. Li, B. C. M. Fung, P. Charland, and S. H. H. Ding, "I-mad: Interpretable malware detector using galaxy transformer," *Comput. Secur.*, vol. 108, p. 102371, 2019. [Online]. Available: https://api.semanticscholar.org/CorpusID:235490302

[17] D. Demirci, N. Sahin, M. Sirlanci, and C. Acartürk, "Static malware detection using stacked bilstm and gpt-2," *IEEE Access*, vol. PP, pp. 1–1, 2022. [Online]. Available: https://api.semanticscholar.org/CorpusID:249215209

[18] Z. He, Y. Sun, J. Liu, and R. Li, "Adatrans: Feature-wise and sample-wise adaptive transfer learning for high-dimensional regression," *ArXiv*, vol. abs/2403.13565, 2024. [Online]. Available: https://api.semanticscholar.org/CorpusID:268536703

[19] S. Seneviratne, R. Shariffdeen, S. Rasnayaka, and N. Kasthuriarachchi, "Self-supervised vision transformers for malware detection," *IEEE Access*, vol. 10, pp. 103 121–103 135, 2022. [Online]. Available: https://api.semanticscholar.org/CorpusID:251564607

[20] A. Ravi, V. Chaturvedi, and M. Shafique, "Vit4mal: Lightweight vision transformer for malware detection on edge devices," *ACM Transactions on Embedded Computing Systems*, vol. 22, pp. 1 – 26, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:261665937

[21] K.-W. Park and S.-B. Cho, "A vision transformer enhanced with patch encoding for malware classification," in *Ideal*, 2022. [Online]. Available: https://api.semanticscholar.org/CorpusID:253762862

[22] A. Rahali and M. A. Akhloufi, "Malbert: Malware detection using bidirectional encoder representations from transformers*," *2021 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, pp. 3226–3231, 2021. [Online]. Available: https://api.semanticscholar.org/CorpusID:245802628

[23] Y. Fan, M. Ju, S. Hou, Y. Ye, W. Wan, K. Wang, Y. Mei, and Q. Xiong, "Heterogeneous temporal graph transformer: An intelligent system for evolving android malware detection," *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, 2021. [Online]. Available: https://api.semanticscholar.org/CorpusID:236980156

[24] A. Rahali and M. A. Akhloufi, "Malbertv2: Code aware bert-based model for malware identification," *Big Data Cogn. Comput.*, vol. 7, p. 60, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:257787274

[25] W. Wangwang, Z. Yunchun, L. Chengjie, S. Xuchenming, Z. Yuting, and Z. Xin, "Network traffic oriented malware detection in iot (internet-of-things)," *2021 International Conference on Networking and Network Applications (NaNA)*, pp. 301–307, 2021. [Online]. Available: https://api.semanticscholar.org/CorpusID:245015225

[26] F. Ullah, A. Alsirhani, M. M. Alshahrani, A. Alomari, H. Naeem, and S. A. Shah, "Explainable malware detection system using transformers-based transfer learning and multi-model visual representation," *Sensors (Basel, Switzerland)*, vol. 22, 2022. [Online]. Available: https://api.semanticscholar.org/CorpusID:252250027

[27] D. Trizna, L. Demetrio, B. Biggio, and F. Roli, "Nebula: Self-attention for dynamic malware analysis," *IEEE Transactions on Information Forensics and Security*, vol. 19, pp. 6155–6167, 2023. [Online]. Available: https://api.semanticscholar.org/CorpusID:264172387

[28] E.-J. Kim, Y.-K. Lee, S.-M. Lee, J.-N. Kim, A. R. Kang, M.-S. Kim, and Y.-S. Jeong, "Malware detection using pre-trained transformer encoder with byte sequences," *PLOS One*, vol. 20, 2025. [Online]. Available: https://api.semanticscholar.org/CorpusID:282059052

[29] VirusTotal, "Virustotal," https://www.virustotal.com/.

[30] M. Lester, "PE malware machine learning dataset," https://practicalsecurityanalytics.com/pe-malware-machine-learning-dataset/, 2021, [Online].

[31] F. de Gaspari, D. Hitaj, G. Pagnotta, L. D. Carli, and L. V. Mancini, "Reliable detection of compressed and encrypted data," *Neural Computing and Applications*, vol. 34, pp. 20 379 – 20 393, 2021. [Online]. Available: https://api.semanticscholar.org/CorpusID:232427932

[32] R. Lyda and J. Hamrock, "Using entropy analysis to find encrypted and packed malware," *IEEE Security & Privacy*, vol. 5, no. 2, pp. 40–45, 2007.

[33] apple, "GitHub - apple/corenet: CoreNet: A library for training deep neural networks." [Online]. Available: https://github.com/apple/corenet