Configuring RAG Systems: What Breaks, Where Originates, and How to Fix

LIANGYU CHEN, University of Electronic Science and Technology of China, China
YOUPENG MA, University of Electronic Science and Technology of China, China
YULONG YE, University of Birmingham, United Kingdom
ZEZHEN XIANG, University of Electronic Science and Technology of China, China
TAO CHEN*, University of Birmingham, United Kingdom

arXiv:2610.18567v1 [cs.AI] 13 Oct 2026

Retrieval-Augmented Generation (RAG) has emerged as a dominant paradigm for improving Large Language Model (LLM) by intertwining it with other components, such as a vector/graph database. Therefore, the resulting RAG systems are often highly configurable, under which a non-trivial amount of failures are related to the configurations. However, there remains a lack of systematic understanding regarding the configuration issues for RAG systems. To bridge this gap, in this paper, we conduct the first large-scale empirical study of 654 real-world configuration issues collected from 10 diverse RAG systems. Drawing on this, we provide a complete categorization of the symptoms, root causes, relevant stages, and fix patterns of configuration issues for RAG systems, as well as several new observations. Our findings not only reveal the characteristics of configuration issues, but also strongly link them to the unique pipeline stages of RAG, e.g., we found that most symptoms can be observed at the early or late stage, while the stage of the root cause might not be at the same stage as where the symptom is identified. We also provide nine actionable insights to the practitioners on testing, localizing, and fixing the configuration issues for RAG systems.

CCS Concepts: • **Software and its engineering → Empirical software validation**; **Software configuration management and version control systems**.

Additional Key Words and Phrases: Retrieval-augmented generation, large language models, configuration bugs, misconfiguration

1 Introduction

Large Language Models (LLMs) have witnessed unprecedented growth and evolution in recent years, demonstrating remarkable capabilities across a wide spectrum of tasks [6, 37, 58]. To mitigate the issue of LLM hallucination [27, 35, 40] and outdated knowledge [26, 44], Retrieval-Augmented Generation (RAG) system has emerged as a dominant paradigm [21, 31]. Typically, as a new type of software systems, the RAG systems augment LLMs with external knowledge sources, such as documents, vector databases, and retrieval pipelines, forming the backbone of many question answering systems, developer assistants, and knowledge-based applications [2, 17, 18].

To ensure flexibility, RAG systems are designed to be highly configurable [3, 5, 19, 41, 42]—a key property of many types of software systems. For example, changing the retrieval stride length will significantly affect the quality of the generated text in RAG systems [32]. However, configuration has long been recognized as a critical source of reliability issues in traditional configurable software systems [16, 22–25, 33, 36, 45, 50, 51], RAG systems are of no exception. Yet, the unique complex stages in the pipeline of RAG systems (e.g., indexing, retrieval, and generation) make their configuration even more error-prone [3, 5, 41]. For example, it has been shown that configuration issues of RAG systems (including both configuration bugs and misconfiguration) can

*Corresponding author.

Authors' Contact Information: Liangyu Chen, University of Electronic Science and Technology of China, Chengdu, China, cly009@std.uestc.edu.cn; Youpeng Ma, University of Electronic Science and Technology of China, Chengdu, China, myp@std.uestc.edu.cn; Yulong Ye, yxy382@student.bham.ac.uk, University of Birmingham, Birmingham, United Kingdom; Zezhen Xiang, zz.xiang.work@gmail.com, University of Electronic Science and Technology of China, Chengdu, China; Tao Chen, University of Birmingham, Birmingham, United Kingdom, t.chen@bham.ac.uk.

cause a latency performance degradation of 10×, leading to a two orders of magnitude cost [5]. Another example showcases the complex interaction of configuration options among the RAG stages: a subtle change in the indexing strategy can unexpectedly invalidate assumptions in the retrieval or generation stages, causing significant errors to propagate silently across the pipeline, and hence a system failure. Therefore, understanding the configuration issues of RAG systems and how they have been fixed is in high demand.

Despite the importance, our current knowledge of the configuration issues within RAG systems remains rather limited. While there are prior empirical studies of RAG systems, they focus on two different aspects: either investigates how configuration choices influence system performance [3, 5, 41], typically from a configuration tuning and hyperparameter optimization perspective; or examines failure cases and practical challenges in RAG systems through empirical analyses of observed bugs in general, which are not configuration-related [1, 4, 30]. There has been no systematic study of the configuration issues as the root of RAG systems' failures, particularly with respect to the RAG pipeline.

To bridge this gap, in this paper, we conduct a large-scale empirical study on RAG configuration issues, covering 654 issues and 345 pull requests from 10 RAG systems on GitHub over six person-months. We make several contributions:

- **A Categorization:** We summarize a comprehensive categorization for configuration issues in RAG systems with respect to its pipeline stages, including five types of symptoms, nine root causes, and six fix patterns.
- **Empirical Findings:** Our empirical study reveals ten previously unknown findings, the most notable ones are:
  - Core pipeline fault, in which failures occurred at a pipeline stage with exceptions, is the predominant symptom of configuration issues in RAG systems.
  - For configuration bugs, invocation of external libraries/systems is the most common cause; while resource provision that does not fulfil the needs is the most common reason for user misconfiguration.
  - Most of the symptoms occurred at the initialization, indexing and generation stage of RAG systems; while root causes mostly lie in the early stages.
  - The stage where the symptom is observed might not necessarily correspond to the stage of the root cause—as observed in one third of the cases.
  - When the symptom and root cause deviate at more distant stages (e.g., from initialization to the last stage of generation), the resolving cost (time) of real configuration bugs can increase by up to 2.83×.
  - Patching the logic and algorithm is the most common fix pattern for configuration bugs.
- **Insights:** Based on these results, we distil several insights about configuration issues for RAG systems on their testing/detection, localization, and fixing.

All the data and code can be accessed at our repository: https://github.com/ideas-labo/rag-config-bugs.

This paper is organized as follows. Section 2 introduces the background. Section 3 presents our methodology. Section 4 delineates the categorization of configuration-related issues and the findings. Section 5 discusses actionable insights. Section 6 addresses threats to validity. Section 7 reviews related work. Finally, Section 8 concludes the paper.

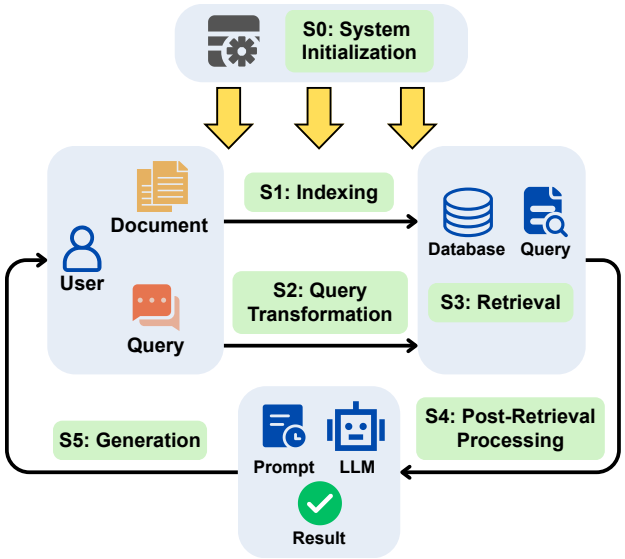

Fig. 1. Workflow of RAG with different pipeline stages.

2 Background

2.1 RAG Pipeline

As shown in Figure 1, RAG systems operate in their unique pipeline to work with LLM. Specifically, given a task (e.g., question answering), there are several key stages [18]:

- **S0: Initialization** prepares, validates the environment, and fulfils the resources needed to boot the system.
- **S1: Indexing** transforms raw data into a searchable knowledge base via chunking and embedding.
- **S2: Query Transformation** rewrites or decomposes user inputs to improve retrieval capability.
- **S3: Retrieval** fetches contexts from vector/graph database.
- **S4: Post-Retrieval Processing** refines results through re-ranking and compression.
- **S5: Generation** synthesizes the final prompt with contextual information for the LLM(s) to process.

Each stage comes with its own configuration options. The final output serves as the answer to different tasks.

2.2 Configuration Issues for RAG Systems

In this paper, we define a *configuration issue* in RAG systems as a defect that satisfies two conditions:

- The issue must manifest as an observable system-level problem: functional defects, e.g., execution failures, crashes, or core pipeline breakdowns, and performance-related degradations, e.g., latency or answer quality.
- It must be directly related to configuration.

From the example in Table 1, we see that the reported issue exhibits a system-level problem and is relevant to the configuration of the RAG system. Clearly, it indicates an observable system-level problem, i.e., a degradation in document chunking quality during indexing, with clear statements confirming that the problem is directly related to the configuration option `chunk_token_num`.

Table 1. A real-world configuration issue from RAG system RAGFlow.

ID #7790**:** Potential Issue in `naive_merge_docx` Function
Description (excerpt from the issue report): ...This can result in the first chunk being larger than intended, as there's no proper size control for the initial chunk... ...The function initializes `tk_nums = [0]` and uses `tk_nums[-1] > chunk_token_num` as a condition for creating new chunks... Since `0 <` chunk_token_num (default 128), the first section is always added to the first chunk...

In general, configuration issues can be broadly distinguished by two types depending on their root causes [25, 51], which are both considered in this work:

- **Configuration bugs:** Here, the issue is caused by system defects, where the system fails to correctly handle specific configuration combinations due to flawed adaptation logic, incorrect assumptions, or incomplete validation.
- **Misconfigurations:** The issues are mainly caused by incompatible or improper configuration settings (e.g., model choices and retrieval strategies) made by the users.

3 Empirical Methodology

Figure 2 shows our methodology. As can be seen, we select representative RAG systems from GitHub, followed by rigorous extraction and analysis of the issues. These then allow us to construct a categorization of configuration issues for RAG systems, together with a rich set of findings. All of those are systematically validated for ensuring inter-rater reliability.

3.1 Research Questions

We seek to answer the following research questions (RQs):

- **RQ1:** What common symptoms do configuration issues exhibit in RAG systems?
- **RQ2:** What commonly causes the configuration issues?
- **RQ3:** How do the deviated stages between symptom and root cause affect the resolution of configuration issues?
- **RQ4:** How are the actual configuration bugs often fixed?

RQ1 and **RQ2** summarize the general characteristics of configuration issues for RAG systems. **RQ3** is derived from the fact that the stage of the exhibited symptom of a configuration issue might deviate from that of its root cause (e.g., a symptom from the *indexing* stage might be caused by bugs in the *generation* stage), hence we aim to study how such a phenomenon might affect its bug fixing. **RQ4** investigates the patterns of actual configuration bug fixing, if any.

Notably, we centre around the above aspects because RAG developers typically reasoned about configuration issues from three complementary pillars: (1) how the issue manifests in practice; (2) where the issue originates; and (3) how the issue is eventually resolved. Thus, understanding them enables us to analyze issue propagation across the RAG pipeline and derive actionable guidance for debugging and maintenance.

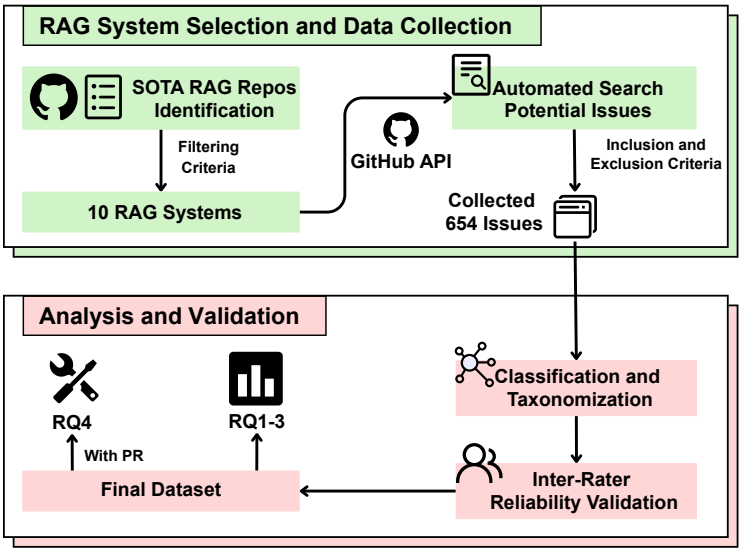

Fig. 2. The empirical study methodology.

3.2 RAG System Selection

To identify the RAG system, we use several keywords to search over GitHub, i.e., *"RAG"* or *"Retrieval-Augmented Generation"*. This returns more than a thousand projects. We then filter the systems using the criteria below:

- **Configuration Accessibility**: The repository contains user-adjustable configuration artefacts, such as configuration files, environment variables, or documented options.
- **Popularity**: To ensure sufficient issue discussion and relevance, we remove repositories with less than 2,000 stars[1]—a proxy for both community adoption and the likelihood of reporting configuration-related issues.
- **Issue Richness**: Repositories with fewer than 20 issues are removed, as there is little information to be analyzed.

After applying the above criteria, we identify several tens of repositories of the RAG systems[2]. To ensure both representativeness and realism, we select the top 10 repositories ranked by the number of GitHub stars among those eligible.

As can be seen from Table 2, the RAG systems are of distinct domains, including general end-to-end RAG platforms[3] (e.g., FastGPT), research prototypes (e.g. HippoRAG), and domain specific RAG system (e.g. DocsGPT). Further, the selected RAG systems cover the major forms of popular open-source RAG software, including end-to-end platforms, developer-oriented frameworks, research prototypes, and user-facing or domain-specific applications. They also incorporate heterogeneous RAG components and external services, such as LLMs, embedding models, vector databases, retrievers, rerankers, and data sources. This diversity provides broad coverage of the architectures, application scenarios, and configuration surfaces found in popular and active open-source RAG systems, making the selected repositories representative within the scope targeted by our study.

[1]This was chosen pragmatically as a general measurement of popularity.
[2]As of 1st Sep, 2025.
[3]Some of these can be used for building other types of LLM systems, but we explicitly extract the RAG-related issues only.

Table 2. Statistics of the selected RAG systems.

RAG System	Stars	Retrieved Issues		Filtered Issues	
		Total	Closed	Total	with PR
RAGFlow	61.1k	1,417	685	217	161
DB-GPT	17.0k	812	559	120	76
FastGPT	25.3k	494	412	74	19
GraphRAG	26.9k	195	152	64	28
LightRAG	18.9k	214	197	63	22
Obsidian-copilot	5.0k	564	353	35	15
Private-GPT	56.4k	255	225	31	3
Quivr	38.0k	283	282	22	8
DocsGPT	16.9k	332	318	21	10
HippoRAG	2.6k	86	75	7	3
Total	–	4,652	3,258	654	345

3.3 Data Collection and Filtering

We collected issues from the repositories of the 10 RAG systems (till 1st August 2025) using a two-step process: automated retrieval and then manual validation by the authors.

3.3.1 Automatic Search. To comprehensively identify configuration issues under a repository of RAG system, we constructed a set of search queries by combining two primary dimensions: commonly used configuration terms in software systems (e.g., `config`, `setting`, `parameter`, `env`) and RAG-specific terms corresponding to core pipeline components (e.g., `embedding`, `vectorstore`, `retriever`, `prompt`). Each pair of the above terms is used in isolation; the resulting issues are then merged with redundancy removal.

We tailored our search scope based on the volume of issues and the effectiveness of labels in each repository:

- **Labeled Search**: For projects with a large volume of issues and well-maintained labeling systems (e.g., LightRAG), we combine the keywords with the project-specific label (e.g., `keyword AND label:bug`) to locate related issues while discarding unrelated discussions.
- **Keyword-Only Search**: For projects that did not maintain a meaningful labeling protocol (or those returned fewer than 20 issues), we conducted keyword-based searches paired with manual screening.

As from Table 2, this process yields a preliminary dataset of 4,652 candidate issues. From this pool, we focus on the 3,258 closed issues due to the needs for **RQ3** and **RQ4**.

3.3.2 Manual Filtering. Since the retrieved issues might contain noise, e.g., some of them could be usage questions or feature requests. To mitigate this, we conducted a rigorous manual inspection to identify genuine configuration-related issues. We say that an issue should be included in the final dataset if and only if it meets all the criteria below:

- **Validity**: The issue must describe a confirmed system defect (e.g., crash, functional error, performance degradation) or user misconfiguration (e.g., parameter/option misunderstanding, usage errors, or documentation-induced mistakes), excluding feature requests and inquiries.

- **Relevance**: The issue's trigger, manifestation, or expected resolution must be logically attributable to at least one user-adjustable configuration option of RAG, such as those specified in configuration files (e.g., `YAML` or `TOML`), environment variables, or API arguments.
- **Traceability**: Only closed issues are considered, and sufficient evidence (e.g., linked pull requests, reproduction steps, or conclusive developer discussions) has to be available to reasonably infer the root cause of the issues.

Applying these criteria resulted in 654 configuration issues for RAG systems in total, which we use to address **RQ1**–**RQ3**. Among them, 345 issues explicitly link to pull requests (PR), enabling a deeper analysis of code-level fix patterns for **RQ4**.

3.4 Analysis and Validation Procedure

To systematically analyze configuration issues in RAG systems, we carefully read the issues' description, comments/discussions, and pull requests (if any), which provide rich information on the symptoms (**RQ1**), root causes (**RQ2**), and fix patterns (**RQ4**), following iterative open-coding [15, 38].

3.4.1 Iterative Categorization. Four authors, who are experienced researchers of RAG systems, lead the categorization. The analysis starts with an exploratory open-coding phase. We randomly sample an initial 10% issues (stratified across systems) in an open-coded manner without predefined categories. Based on recurring patterns, we propose preliminary categories (and sub-categories as needed) for **RQ1**, **RQ2**, and **RQ4**. In the first iteration, the full dataset is then labeled, with iterative refinements (merging, splitting, redefining categories); we stop at the level of categories where the next level of sub-categories resulted in highly biased labeling—over 90% of the issues are condensed to one sub-category—or no finer-grained common phenomenon across the RAG systems can be observed. All authors need to agree upon a final categorization.

3.4.2 Inter-Rater Reliability Validation. With the final categories, each author then independently conducts the second iteration, aiming to correct the labels of issues using those categories. Note that in some cases, an issue might involve multiple categories/sub-categories, in which case we label it with the most fundamental/crucial one, as agreed by all raters. We also seek to assign the issues to the relevant RAG stages for **RQ3**. To measure inter-rater reliability, we use Cohen's κ [29]. Through multiple iterations of relabeling, agreeing, and consulting external experts, we seek to reach $\kappa \geq 0.7$ for all levels of the categories and stages[4] [14, 46].

As summarized in Table 3, we achieve $\kappa \geq 0.7$ at all levels, including all the categories and identification of the RAG stage, with three iterations. In particular, we have perfect agreement for the main category of symptoms and root causes.

3.4.3 Unweighted vs. Weighted Percentage. To provide overall statistics, we report on both the unweighted and weighted proportions of issues per category/sub-category over all systems, i.e., the former is robust to the diverse sample sizes of the RAG systems, while the latter is sensitive to the sizes.

4 Results

4.1 Categorization of RAG Configuration Issues

Unlike others, RAG systems operate in a multi-stage semantic pipeline that integrates heterogeneous components, which have been captured by our categorization in Figure 3. It is worth noting that our categorization contains both general software-failure manifestations and RAG-specific

[4] 0.7 is a pragmatic "good" agreement in software engineering [14, 46].

Table 3. Inter-rater agreements (Cohen's κ) recorded.

	2nd Iteration	3rd Iteration	4th Iteration
Symptom category	0.78	1.00	1.00
Symptom sub-category	0.61	0.73	0.73
Symptom stage	0.89	0.89	0.89
Root cause category	0.68	1.00	1.00
Root cause sub-category	0.53	0.69	0.80
Root cause stage	0.62	0.74	0.74
Fix pattern	0.63	0.82	0.82

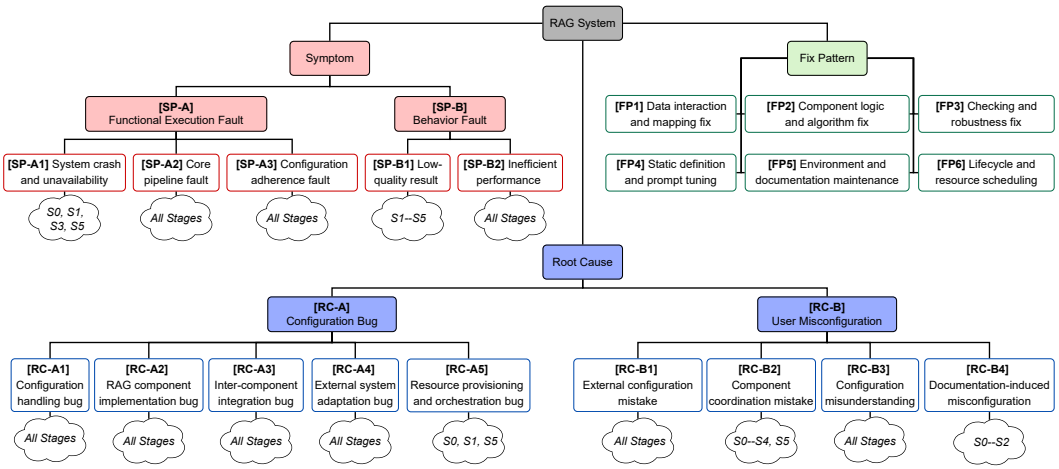

Fig. 3. Categorization for configuration issues of RAG systems.

refinements. High-level symptoms such as crashes or performance degradation may also occur in conventional software systems. However, their finer-grained manifestations in our dataset are tied to the functional components of RAG systems. In particular, *core pipeline faults* are distinguished according to the RAG stage in which the affected functionality resides, while *external system adaptation bugs* capture incompatibilities involving LLMs, embedding models, vector databases, external data sources or applications, and rerankers. These component-specific categories are further connected to the symptom and root cause stages, providing a stage-aware characterization that is not captured by generic configuration bug categories.

4.1.1 Symptoms. The Symptom dimension characterizes how configuration issues become externally observable from a user perspective. We divide the observable symptoms into two categories, each of which has different sub-categories.

Functional Execution Fault. This corresponds to cases where the system or a specific pipeline stage fails to fulfil the intended needs, including three sub-categories:

- **System crash and unavailability** mean the system terminates unexpectedly without explicit exceptions or hangs indefinitely. The RAG system becomes entirely non-responsive without reasons, preventing any further interaction: e.g., issue #822 for Obsidian-copilot.

- **Core pipeline fault** denotes that the system remains functional, but fails to fully complete a specific stage of the pipeline with explicit exceptions, causing the system to stop: e.g., issue #1642 for GraphRAG.
- **Configuration adherence fault** means that the system has good quality/performance, but violates the user's configuration intent: e.g., issue #2636 for DB-GPT.

Behavior Fault. Here, the RAG system suffers severe behavior-level degradations, where the system delivers **low-quality results**, i.e., the system does generate answers, but they are factually incorrect/incomplete; or exhibits **poor performance**: a good quality outcome with abnormal resource consumption or excessive latency. For example, issue #1321 for LightRAG and issue #7518 for RAGFlow.

4.1.2 Root Cause. Our results indicate that root causes of the configuration issues for RAG systems can be divided into two categories, each of which contains several sub-categories.

Configuration Bug. The configuration issues whose root causes lie in the system's code. There are five sub-categories:

- **Configuration handling bug** arises when configuration options are incorrectly parsed, transmitted, or bound to execution logic, e.g., issue #1313 for DocsGPT.
- **RAG component implementation bug** originates from flawed internal logic at core RAG data processing (such as chunking, retrieval strategy, graph clustering, etc.) related to configuration, e.g., issue #436 for RAGFlow.
- **Inter-component integration bug** occurs at component boundaries, where mismatches between upstream outputs and downstream input assumptions violate implicit data contracts within the RAG pipeline under specific configurations. For example, issue #1642 for GraphRAG.
- **External system adaptation bug** stems from incompatible configuration between the RAG system and external artefacts such as LLMs, databases, embedding models, or re-rankers: issue #990 for Private-GPT.
- **Resource provisioning and orchestration bug** reflects failures in allocating, scheduling, or utilizing underlying computational resources for fulfilling a given configuration. For example, issue #1288 for RAGFlow.

User Misconfiguration. Unlike configuration bugs, user misconfiguration mainly originates from users' misunderstanding of the RAG system's configuration options:

- **External configuration mistake** occurs when users select or configure external artefacts—such as LLMs or databases—that are incapable of supporting the intended workload. For example, issue #1567 for LightRAG.
- **Component coordination mistake** arises when users set the configuration to combine incompatible internal RAG components, violating functional dependencies between pipeline stages. For example, issue #614 for FastGPT.
- **Configuration misunderstanding** reflects misunderstandings of configuration semantics, constraints, or deprecated options, leading to behaviors that diverge from user intents despite technically valid inputs. For example, issue #144 for Obsidian-copilot.
- **Documentation-induced misconfiguration** captures cases where users are misled by incorrect, outdated, or ambiguous documentation, resulting in configurations that follow official guidance but fail in practice. For example, issue #8556 for RAGFlow.

4.1.3 Stage Labeling. We also study how configuration issues are localized within and propagated through the RAG pipeline with respect to the symptoms and root cause[5]:

- **Symptom Stage:** The stage where the abnormal behavior first becomes observable to the user. Crucially, this measures failure observability rather than the exact moment an internal pipeline data flow might first deviate silently.
- **Root Cause Stage:** This pinpoints where the configuration issue originates. For configuration bugs, it is the stage of the flawed code logic; for user misconfiguration, it is the stage where users set improper configuration options. We infer the root-cause stage by jointly analyzing issue content and developer discussions. For issues with pull requests, we examine the modified code and the comments to identify the originating pipeline stage.

We found that some symptoms/root causes are linked to all stages, while others are only related to certain stages.

4.1.4 Fix Pattern. We extract six fix patterns for configuration issues in RAG systems based on 345 issues with pull requests:

- **Data interaction and mapping:** This pattern involves ensuring configuration-related data mapping and interface alignment between disparate RAG components or external APIs (e.g., key renaming, context propagation).
- **Component logic and algorithm:** This involves modifying the core business logic or internal algorithms (e.g., rewriting chunking flows) to better handle configurations.
- **Checking and robustness:** This pattern focuses on increasing system stability by adding error checking/warning code, e.g., null-pointer checkers, data type validation, and improved error-handling logic to prevent crashes under invalid or incompatible configurations.
- **Static definition and prompt tuning:** These fixes involve tuning static constants, default values, or prompt templates. Developers often use this to refine hard-coded defaults or update prompt constraints within the *generation* stage to align with the intents of the configurations.
- **Environment and documentation maintenance:** This covers configuration-related but non-code interventions, such as fixing external dependencies, updating deployment (e.g., Docker), or correcting misleading documentation to ensure correct RAG system deployment.
- **Lifecycle and resource scheduling** This involves governing configuration-related resources for RAG's lifecycle, including concurrency and consistency. Typical fixes include singleton enforcement, implementing cache invalidation, and optimizing pooling for external databases.

4.2 RQ1: Symptoms of Configuration Issues

4.2.1 Commonality of Symptoms. As from Table 4, the observed *functional execution fault* (86%/88% for SP-A1, SP-A2, and SP-A3) is predominantly more common than *behavior fault* (14%/12% for SP-B1 and SP-B2) on both weighted and unweighted percentages. These results indicate that configuration issues in RAG systems primarily act as execution blockers rather than behaviors influencers. Within the former, for all RAG systems, the *core pipeline fault* is the most widespread exhibited issue, consisting of 64%, since the related options, if contain issues, often cause severe results such as crashes; while for the latter, *quality fault* is more commonly observed, as it is often more critical than performance.

There are also discrepancies among the systems, e.g., Quivr—external tools-heavy RAG systems—show particularly high proportions of core pipeline faults, as their orchestration logic is strongly coupled with configuration. FastGPT, a lightweight RAG system, exhibits more behavior faults.

[5]We found that most symptoms/root causes can correspond to one stage only. For the rare cases of symptoms/root causes that are related to multiple stages, we use the one that is the most relevant, as agreed by all authors.

Table 4. Symptoms proportion by systems/categories—% per system (count). The highest % is highlighted in blue.

System	SP-A1	SP-A2	SP-A3	SP-B1	SP-B2	Total
RAGFlow	8% (17)	63% (137)	17% (37)	9% (19)	3% (7)	217
DB-GPT	11% (13)	66% (79)	9% (11)	13% (16)	1% (1)	120
FastGPT	4% (3)	55% (41)	18% (13)	20% (15)	3% (2)	74
GraphRAG	6% (4)	59% (38)	23% (15)	9% (6)	2% (1)	64
LightRAG	11% (7)	63% (40)	6% (4)	14% (9)	5% (3)	63
Obsidian-copilot	6% (2)	77% (27)	6% (2)	11% (4)	0% (0)	35
Private-GPT	6% (2)	74% (23)	3% (1)	6% (2)	10% (3)	31
Quivr	5% (1)	86% (19)	5% (1)	5% (1)	0% (0)	22
DocsGPT	14% (3)	62% (13)	14% (3)	0% (0)	10% (2)	21
HippoRAG	29% (2)	57% (4)	14% (1)	0% (0)	0% (0)	7
Weighted	8% (54)	64% (421)	14% (88)	11% (72)	3% (19)	654
Unweighted	10% (54)	66% (421)	12% (88)	9% (72)	3% (19)	654

Table 5. Overall symptoms proportion over pipeline stages. The highest is highlighted in blue.

Pipeline Stage	Count	Weighted	Unweighted
S0: Initialization	153	24%	22%
S1: Indexing	208	32%	30%
S2: Query Transformation	23	3%	3%
S3: Retrieval	70	11%	11%
S4: Post-Retrieval Processing	23	3%	2%
S5: Generation	177	27%	32%

Finding 1

Configuration issues in RAG systems are more predominantly manifested as functional execution faults (86%/88%) rather than behavior faults (14%/12%). Faults related to the core pipeline and the answers' quality are the most common.

4.2.2 Symptoms-Stage Relations. Table 5 reveals two critical observations on the symptoms-stage relations: (1) the stage *initialization* accounts for 24%/22% of the issues, meaning that users frequently fail to even start the system. (2) The issues are more associated with either the early or late stage. That is, we see that stage *indexing* and *generation* own 32%/30% and 27%/32% of the issues, respectively, which are considerably higher than any other intermediate stages.

Finding 2

Configuration issues in RAG systems predominantly surface at three critical stages: the *initialization* (24%/22%), *indexing* (32%/30%) and *generation* (27%/32%).

To take a closer look, as shown in Figure 4, there is a distinct temporal separation in symptom manifestation. 54%-92%/36%-96% of the *functional execution fault* occurs at the first two stages, often blocking further execution. In contrast, the *behavior fault* is related to the late stages, with 42%-64%/24%-59% at the *generation* stage, manifesting after the full pipeline has executed. This is

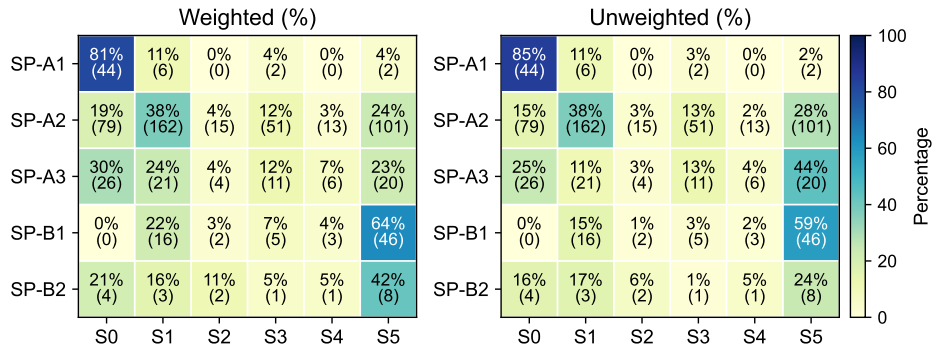

Fig. 4. Detailed symptoms % per category over stages.

Table 6. Root cause proportion by systems/categories. The same format as Table 4.

System	RC-A1	RC-A2	RC-A3	RC-A4	RC-A5	RC-B1	RC-B2	RC-B3	RC-B4	Total
RAGFlow	18% (38)	28% (61)	10% (21)	26% (56)	3% (7)	5% (10)	1% (3)	7% (16)	2% (5)	217
DB-GPT	12% (14)	13% (16)	13% (16)	32% (38)	12% (14)	8% (10)	2% (2)	8% (9)	1% (1)	120
FastGPT	5% (4)	8% (6)	0% (0)	28% (21)	3% (2)	38% (28)	4% (3)	14% (10)	0% (0)	74
GraphRAG	33% (21)	19% (12)	9% (6)	19% (12)	2% (1)	6% (4)	3% (2)	9% (6)	0% (0)	64
LightRAG	6% (4)	17% (11)	6% (4)	40% (25)	0% (0)	14% (9)	0% (0)	16% (10)	0% (0)	63
Obsidian-copilot	29% (10)	3% (1)	0% (0)	17% (6)	6% (2)	34% (12)	3% (1)	9% (3)	0% (0)	35
Private-GPT	6% (2)	0% (0)	10% (3)	23% (7)	10% (3)	42% (13)	3% (1)	6% (2)	0% (0)	31
Quivr	9% (2)	9% (2)	0% (0)	59% (13)	0% (0)	18% (4)	0% (0)	5% (1)	0% (0)	22
DocsGPT	24% (5)	0% (0)	0% (0)	24% (5)	10% (2)	29% (6)	5% (1)	10% (2)	0% (0)	21
HippoRAG	14% (1)	14% (1)	0% (0)	14% (1)	0% (0)	43% (3)	0% (0)	14% (1)	0% (0)	7
Weighted	15% (101)	17% (110)	8% (50)	28% (184)	5% (31)	15% (99)	2% (13)	9% (60)	1% (6)	654
Unweighted	16% (101)	11% (110)	5% (50)	28% (184)	4% (31)	24% (99)	2% (13)	10% (60)	0% (6)	654

because the early stages can easily exhibit errors related to preparations, while the late stage is often where the problematic answers are observed.

Finding 3

Functional execution fault often occurs much earlier than the *behavior fault* at the pipeline stages of RAG systems.

4.3 RQ2: Root Causes of Configuration Issues

4.3.1 Commonality of Root Causes. As shown in Table 6, *configuration bug* (RC-A1 to RC-A5) dominates the root causes, accounting for 73%/64% of the weighted/unweighted percentage, while *user misconfiguration* accounts for the remaining 27%/36%. Notably, *external system adaptation bug* (RC-A4) emerges as the most prevalent root cause of configuration bugs. This reflects the fundamental nature of RAG systems as integration-centric architectures: they act as glue code that orchestrates multiple heterogeneous components with evolving interfaces and implicit assumptions. As a result, configuration bugs frequently surface at those boundaries.

As for *user misconfiguration*, the *external configuration mistake* (RC-B1) is the most common cause. This suggests that user mistakes are most likely to arise from configuring external artefacts that do not meet the RAG system's needs.

Table 7. Overall root cause proportion over pipeline stages. The highest is highlighted in blue.

Pipeline Stage	Count	Weighted	Unweighted
S0: Initialization	255	39%	40%
S1: Indexing	177	27%	25%
S2: Query Transformation	27	4%	3%
S3: Retrieval	43	7%	5%
S4: Post-Retrieval Processing	23	3%	2%
S5: Generation	129	20%	24%

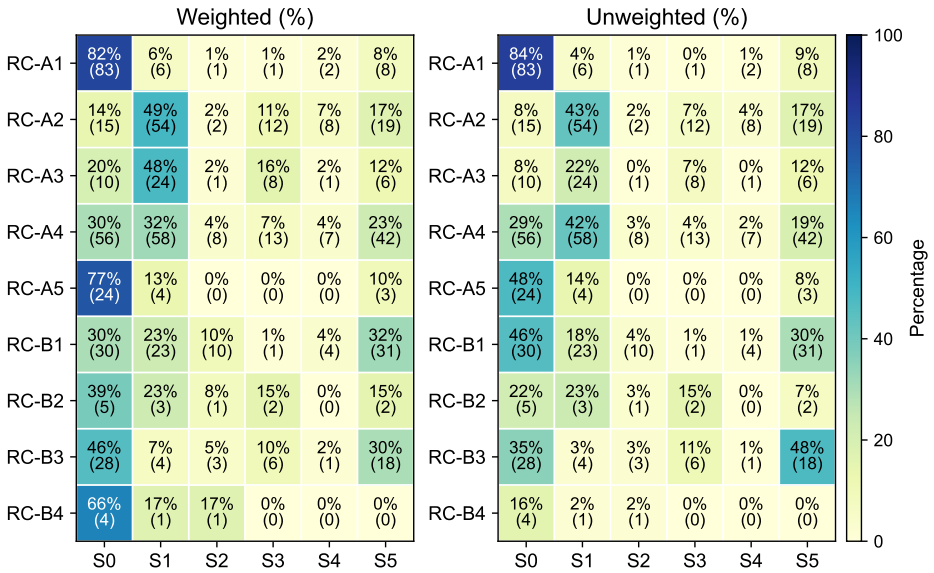

Fig. 5. Detailed root cause % per category over stages.

There is also a notable system discrepancy: highly-integrated and cloud-native assistants, e.g., Quivr, predominantly suffer from external system adaptation bugs due to their extensive reliance on third-party APIs, whereas local-deployment or lightweight systems, e.g., Private-GPT and FastGPT exhibit a significantly higher proportion of external configuration mistakes (38%–42%) as users frequently struggle with local environment and model provisioning.

Finding 4

Configuration bug (73%/64%) is the primary cause of configuration issues in RAG systems against *user misconfiguration* (27%/36%). Notably, *external system adaptation bug* is the most predominant configuration bug, while the most common user mistakes are from configuring external artefacts.

4.3.2 Root Cause-Stage Relation. We further analyze at which RAG stage the root causes are identified. As shown in Table 7, the majority of root causes lie in the early stage (*initialization* and *indexing*), which together account for over 66%/65% of the cases. The same can be observed over the sub-categories (Figure 5). This highlights that many configuration issues are rooted in system setup and data preparation.

Table 8. Proportion of 607 error-prone configuration option over stages. The highest is highlighted in blue.

Pipeline Stage	Count	Weighted	Unweighted
S0: Initialization	131	21%	22%
S1: Indexing	239	39%	38%
S2: Query Transformation	25	4%	3%
S3: Retrieval	40	7%	7%
S4: Post-Retrieval Processing	34	6%	4%
S5: Generation	138	23%	26%

Finding 5

All types of root causes of configuration issues in RAG are often related to early RAG pipeline stages (66%/65% cases).

4.3.3 Configuration Options-Stage Relations. To understand the granularity of configuration risks, we extracted the specific configuration options of root causes in the issues. Notblity, to prevent repeatedly reported problems from disproportionately increasing the frequency of a configuration option, we deduplicated the option-level data within each repository. Specifically, when multiple issues from the same repository involved the same configuration option, that option was counted only once. The same option name appearing in different repositories was retained separately because its implementation, semantics, and affected component may differ across systems. Consequently, the resulting 607 repository–option pairs characterize the breadth of distinct issue-linked configuration options rather than the raw frequency of issue reports.

In total, we identify 607 unique options across the 654 issues (an issue may involve multiple options). As in Table 8, we say that:

Finding 6

The complexity of RAG configuration is unevenly distributed: the *indexing* stage involves the highest volume of error-prone options (39%/38%), such as embedding settings.

4.4 RQ3: Stage Gap and Resolving Costs

RQ3 investigates how costly it is to resolve the configuration issues of RAG systems in practice, depending on the deviation between the stage of symptoms and root causes. To that end, we compute the stage gap $\mathcal{G}$ of each issue:

$$\mathcal{G} = |\mathcal{S} - \mathcal{R}| \tag{1}$$

whereby $\mathcal{S}$ and $\mathcal{R}$ denote the stage level for symptom and root cause, respectively. Clearly, a $\mathcal{G} = 0$ indicates the issue manifests at the same stage where it was introduced. A $\mathcal{G} > 0$ indicates cross-stage propagation, where the configuration issues are propagated through other stages before revealing a symptom. Note that the cause might be related to a stage earlier or later than the stage of the symptom.

For example, the issue #144 for Obsidian-copilot is a misconfiguration issue. Here, the user configured the RAG system to use the *"Active Note"* context mode, mistakenly believing it would feed the entire document as full context. However, the system's underlying logic uses similarity-based retrieval. As such, the stage of root cause is *retrieval*, as the user's misunderstanding led to selecting a retrieval strategy unsuitable for the task. Yet, the symptom only became observable at the *generation* stage, when the LLM produced a low-quality response. This results in a $\mathcal{G} = |5 - 3| = 2$.

Table 9. Stage gap proportion across root cause categories; W and U denote weighted and unweighted percentage, respectively. The highest is highlighted in blue.

Stage Gap	Overall [Count (W/U%)]	RC-A [Count (W/U%)]	RC-B [Count (W/U%)]
Gap = 0	440 (67%/64%)	332 (70%/67%)	108 (61%/57%)
Gap = 1	93 (15%/13%)	61 (13%/8%)	32 (18%/21%)
Gap = 2	27 (4%/5%)	19 (4%/3%)	8 (5%/5%)
Gap = 3	19 (3%/5%)	17 (3%/9%)	2 (1%/2%)
Gap = 4	15 (2%/2%)	9 (2%/2%)	6 (3%/2%)
Gap = 5	60 (9%/11%)	38 (8%/11%)	22 (12%/13%)
Total Gap > 0	214 (33%/36%)	144 (30%/33%)	70 (39%/43%)

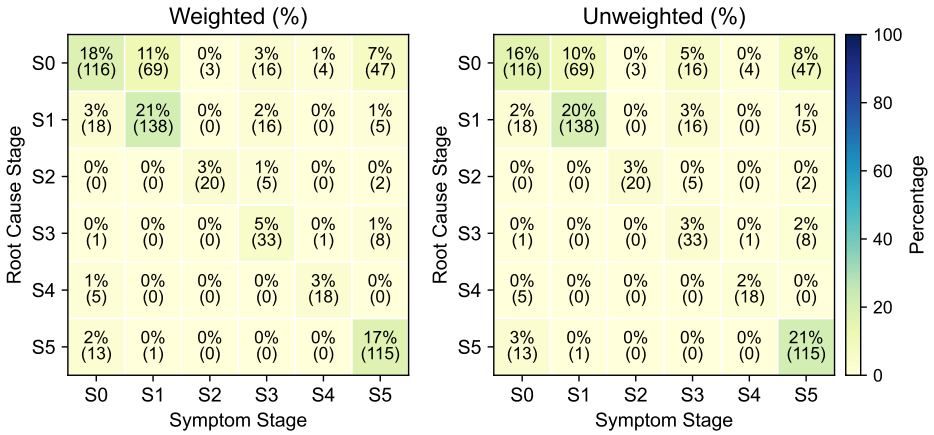

Fig. 6. Stage matrix for symptom and root cause.

4.4.1 Prevalence and Patterns of Stage Gap. As summarized in Table 9, 33%/36% (weighted/unweighted) of the configuration issues exhibit a non-zero stage gap between root cause and symptom. This indicates that nearly one-third of configuration issues in RAG systems cannot be diagnosed by inspecting the stage where the failure is observed alone. Notably, this gap is present in both primary categories of the root cause, suggesting that cross-stage reasoning is a challenge in debugging the configuration issues for RAG systems.

Figure 6 reveals that for those configuration issues with deviated stages, most commonly the symptoms can be at either *indexing* (S1) or *generation* (S5) while the root cause is associated with *initialization* (S0) stage—the configuration issues introduced during system *initialization* only become apparent as abnormalities propagate to the *indexing* and *generation*.

Finding 7

Nearly one-third of RAG configuration issues exhibit a mismatch between the stages of symptom and root cause. This often happens when root problems at the *initialization* cause the symptoms at either *indexing* or *generation* stage.

Table 10. Median resolving cost over stage gap.

Stage gap	0	1	2	3	4	5
Median cost (h)	66.0	35.1	102.4	357.2	47.1	167.2

4.4.2 Impacts on Resolving Costs. To understand the impacts of the stage gap, we quantify the duration required to resolve each configuration issue. Specifically, for each issue, the start time is defined as the issue creation timestamp (`created_at`). The end time is determined depending on two cases:

- if the issue is relevant to a fixing pull request, we use the pull request merge time (`merged_at`);
- otherwise, we use the issue closing time (`closed_at`).

The resolving cost is calculated as the elapsed time between the start and end timestamps. From Table 10, we observe that there seems to be a complex relationship between the stage gap and the resolving cost. We hence take two further analyses.

Correlation Analysis. Our initial hypothesis is that there could be a monotonic relationship between those two. To confirm this, we apply Spearman correlation r on the stage gap ($\mathcal{G} \in [0, 5]$) to the resolving cost for all configuration issues. The result indicates that they have a trivial correlation ($r \approx 0.068$, p-value ≈ 0.08). As such, a larger stage gap might not mean that the configuration issue takes longer to resolve.

Grouping Analysis. Based on the correlation results, we revised our hypothesis: there might exist a *"complexity threshold"* where cross-stage debugging of configuration issues becomes significantly harder. To that end, we apply the Mann-Whitney U test to compare the resolving costs between two groups of issues separated by different thresholds, and we found that only when $\mathcal{G} = 2$ yields statistically significant differences between the groups in the resolving cost: we evaluated each feasible cutoff $c \in \{1, 2, 3, 4, 5\}$ by comparing issues with $\mathcal{G} < c$ against those with $\mathcal{G} \geq c$. The cutoffs $c = 1$, $c = 4$, and $c = 5$ did not produce statistically significant differences in resolving costs, whereas both $c = 2$ and $c = 3$ produced significant differences. We selected $c = 2$ for the reported grouping because, compared with $c = 3$, it yielded a smaller overall p-value (0.0025 versus 0.0059), a larger cost multiplier for configuration bugs (2.83× versus 2.77×), and a more balanced high-gap sample ($N = 121$ versus $N = 94$). Therefore, we dichotomize the configuration issues for RAG systems as:

- **Low gap ($\mathcal{G} \in [0, 1]$)**: Issues that are localized or manifest in the same or immediate adjacent stage.
- **High gap ($\mathcal{G} \in [2, 5]$)**: Issues that propagate more stages.

As shown in Table 11, we observe:

- **Overall impact**: For all configuration issues, those with the high gap have a median resolving cost of 160.6 hours, which is 2.48× higher than that of those within a low gap (64.7 hours) with statistical significance ($p < 0.01$).
- **Configuration bugs (RC-A) as root causes**: A deeper breakdown reveals that the large resolving cost difference is primarily driven by configuration bugs: the cost multiplier rises to 2.83× ($p < 0.01$) in the high gap issues. This implies that when internal code logic fails and propagates downstream, developers face substantial difficulty in tracing the RAG pipeline to the root causes.
- **Misconfiguration (RC-B) as root causes**: For user misconfigurations, although the median cost increases (1.87×) in the high gap group, the difference is not statistically significant. This

Table 11. Statistical difference of resolving costs (hours) on RAG configuration issues with low and high gap.

Group	Low Gap		High Gap		U Test	
	Count	Median	Count	Median	Increase	p-value
Overall	533	64.7	121	160.6	2.48×	0.003
RC-A: bugs	393	58.5	83	165.2	2.83×	0.002
RC-B: misconfig.	140	78.3	38	146.1	1.87×	0.273

confirms that while user misconfiguration issues exhibit large stage gaps, they are "cheaper" to resolve (often a single option update).

Therefore, we conclude:

Finding 8

There is no monotonic correlation between stage gap size and resolving cost. Instead, we observe a *"threshold effect"*: costs spike significantly when issues have $\mathcal{G} \geq 2$, leading to largely deviated costs mainly due to configuration bugs.

4.5 RQ4: Fix Pattern

We now focus on studying the 345 configuration issues for RAG systems, which are linked with pull requests (including both configuration bugs and user misconfiguration). As shown in Table 12, we see that for both weighted and unweighted percentages, the *component logic and algorithm* (FP2, 40%/43%) and *data interaction and mapping* (FP1, 28%/26%) are among the most common fix patterns. However, they are not necessarily those that require the greatest effort in terms of fixing. Instead, patterns that involve deeper system-level intervention, such as *lifecycle and resource scheduling* (FP6) exhibit noticeably higher median fix durations.

Specifically, graph-based and multi-source integrated systems, e.g., GraphRAG and Quivr require *data interaction and mapping* fix to resolve complex data structure mismatches, whereas other systems heavily rely on *component logic and algorithm fix* to repair their routing and execution flows.

From Figure 7, we see several patterns between root causes and fix patterns, where the configuration bugs are often linked to *data interaction and mapping* fixes (FP1) and *component logic and algorithm fixes* (FP2). Interestingly, misconfiguration can also lead to fixes, mainly on the improved error checking/handling (FP3) and document/manual updates (FP5).

Finding 9

The *component logic and algorithm* (40%/43%) and *data interaction and mapping* (28%/26%) are the most common patterns. Yet, those patterns that require deeper system-level interaction often require a higher cost to fix.

Finding 10

Different root causes indeed exhibit distinct repair tendencies: configuration bugs are often related to data mapping and core logic fixes, while the misconfiguration can lead to improved error checking and document updates.

5 Insights for Configuration of RAG Systems

Here, we discuss the insights derived from our empirical observations:

Table 12. Fix pattern proportion by systems/categories. The same format as Table 4.

System	FP1	FP2	FP3	FP4	FP5	FP6	Total
RAGFlow	25% (40)	40% (64)	21% (34)	6% (9)	5% (8)	4% (6)	161
DB-GPT	26% (20)	39% (30)	8% (6)	8% (6)	7% (5)	12% (9)	76
GraphRAG	50% (14)	32% (9)	11% (3)	4% (1)	0% (0)	4% (1)	28
LightRAG	32% (7)	36% (8)	5% (1)	18% (4)	0% (0)	9% (2)	22
FastGPT	26% (5)	47% (9)	5% (1)	11% (2)	0% (0)	11% (2)	19
Obsidian-Copilot	7% (1)	60% (9)	13% (2)	20% (3)	0% (0)	0% (0)	15
DocsGPT	40% (4)	20% (2)	10% (1)	0% (0)	20% (2)	10% (1)	10
Quivr	50% (4)	25% (2)	0% (0)	0% (0)	12% (1)	12% (1)	8
HippoRAG	0% (0)	67% (2)	33% (1)	0% (0)	0% (0)	0% (0)	3
Private-GPT	0% (0)	67% (2)	0% (0)	0% (0)	0% (0)	33% (1)	3
Weighted	28% (95)	40% (137)	14% (49)	7% (25)	5% (16)	6% (23)	345
Unweighted	26% (95)	43% (137)	11% (49)	7% (25)	4% (16)	9% (23)	345
Median cost (h)	31.7	58.6	41.0	30.4	63.1	103.2	–

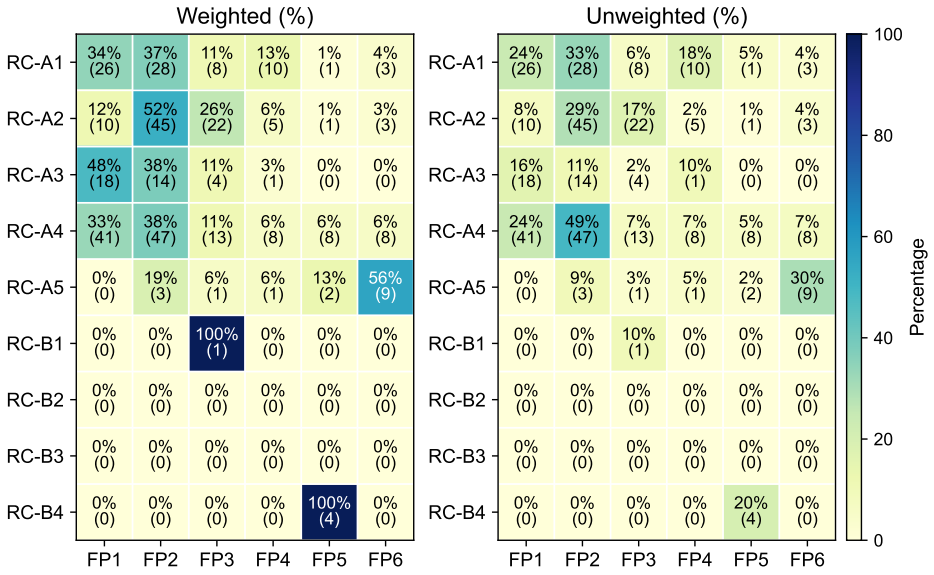

Fig. 7. Distribution of fix patterns across different root causes.

5.1 For RAG System Developers

Insight 1: Prioritize detection on stages other than query or retrieval-related ones.

Most configuration issues in our dataset manifest as functional execution faults rather than only as degraded answer quality. Core pipeline faults are the most common symptom category, and symptoms concentrate at initialization, indexing, and generation. In contrast, query transformation, retrieval, and post-retrieval processing account for a smaller share of directly observable symptoms.

Guidance. Developers should treat configuration validation as a pipeline-level reliability mechanism rather than as a retrieval-only quality check. A practical validation suite should include at least three layers:

- **Startup validation:** check whether required environment variables, model credentials, service endpoints, and database connections are available before the RAG service accepts user requests.
- **Indexing validation:** run a small indexing job with a known document and verify that chunking, embedding, metadata persistence, and vector-store insertion all complete successfully.
- **Generation validation:** run a short end-to-end query over a controlled fixture and verify that the configured prompt template, LLM endpoint, context window, and output parser are mutually compatible.

Retrieval metrics such as recall or ranking quality remain important, but they should not be the only oracle. Many RAG configuration failures prevent the system from reaching the retrieval-quality evaluation stage at all.

Example. Suppose a RAG framework allows users to configure `embedding_model`, `chunk_size`, `vector_store`, and `llm_provider`. A developer can add a `rag doctor` command that performs the following checks before deployment: instantiate the embedding model, index one sample document, retrieve one chunk, and send a minimal prompt to the configured LLM. If the embedding dimension does not match the vector database collection, the command should fail during indexing with an explicit diagnostic, instead of letting the mismatch later appear as a retrieval or generation failure.

Insight 2: Focus debugging efforts on early-stage configurations.

Most root causes originate from the initialization and indexing stages. This means that downstream failures are often caused by earlier configuration handling, resource binding, model loading, document processing, or index construction decisions.

Guidance. When a configuration issue is reported, developers should first reconstruct the early-stage configuration path:

(1) Identify which configuration source is authoritative, such as a YAML file, environment variable, command-line argument, UI setting, or database record.
(2) Confirm how the option is parsed, normalized, and propagated into the initialization or indexing component.
(3) Check whether the option silently falls back to a default value.
(4) Verify that derived artifacts, such as embedding dimensions, chunk metadata, index schema, and collection names, match the configured values.

This debugging order is especially useful when the symptom appears later in the pipeline. A generation-stage error may still be caused by an incorrectly constructed index or by a missing initialization-time dependency.

Example. A user reports that a generated answer ignores uploaded documents. Rather than starting from prompt tuning, the maintainer first checks whether the configured document parser and embedding model were used during indexing. The investigation reveals that the system accepted a custom embedding model in the UI, but the indexing worker still used the default model because the configuration was not propagated to the background job. The fix is to pass the selected model ID into the indexing job and record the model metadata with the generated index.

Insight 3: Cross-stage bug localization is necessary.

Nearly one-third of configuration issues exhibit a non-zero stage gap between the root-cause stage and the symptom stage. Therefore, inspecting only the stage where the failure becomes visible is often insufficient.

Guidance. Developers should provide observability that connects configuration values to pipeline artifacts across stages. Useful mechanisms include:

- a configuration provenance record that shows where each effective option came from;
- stage-level logs that include configuration fingerprints rather than only raw errors;
- metadata attached to indexed chunks, retrieved passages, and generated answers;
- an issue template that asks reporters for both runtime symptoms and upstream setup and indexing details.

The goal is to make it possible to trace how a configuration value affects later-stage behavior. Without this trace, maintainers may overfit the fix to the visible symptom and leave the upstream cause unchanged.

Example. In a note-taking RAG application, a user chooses an "active note" mode expecting the whole note to be passed to the LLM. The system actually interprets this mode as a retrieval strategy over the note, and the symptom appears only when the LLM answers that it lacks enough context. The visible failure is in generation, but the root cause is a retrieval-stage configuration misunderstanding. A cross-stage trace should show: selected context mode → retrieval strategy → retrieved passages → final prompt. This trace would guide both documentation and UI improvements.

Insight 4: Prioritize the fix of issues with large stage gaps.

Issues whose root causes and symptoms are separated by multiple pipeline stages incur substantially higher resolving costs. In our analysis, high-gap issues ($\mathcal{G} \geq 2$) have a higher median resolving cost than low-gap issues, and the effect is particularly strong for system implementation defects.

Guidance. Maintainers should add stage-gap information to issue triage. Severity should still matter, but a non-crashing issue with a large stage gap may consume more debugging time than a localized crash with an obvious cause. A practical triage workflow can classify each report as:

- **localized:** symptom and suspected root cause are in the same or adjacent stages;
- **propagated:** symptom appears two or more stages away from the suspected root cause;
- **unknown:** insufficient evidence, requiring reproduction logs or configuration provenance.

Propagated issues should receive earlier maintainer attention because they often require understanding multiple components and data contracts.

Example. A generation-stage formatting failure may appear minor because the system still returns an answer. However, if the actual cause is that indexing stored metadata under the wrong key, then the issue spans indexing, retrieval, and generation. The maintainer should prioritize it over a localized prompt typo because the metadata mismatch can affect multiple retrieval strategies and future indexed data. A robust fix should migrate or normalize metadata at the indexing boundary, not merely patch the final prompt.

Insight 5: Specific tools are required for configuration repair.

The most common fixes involve component logic and algorithms, followed by data interaction and mapping. Therefore, RAG configuration repair requires tools that understand component interactions, external dependencies, and pipeline data contracts, rather than only tools that edit configuration files.

Guidance. Repair support should be configuration-aware and component-aware. Useful tool features include:

- checking whether configuration keys are consistently mapped across UI, API, worker, and storage layers;
- detecting incompatible component combinations, such as embedding models and existing vector-store schemas;
- generating regression tests from real configuration issues;
- suggesting code-level repairs when the system mishandles valid configurations.

This is different from a generic linter. Many observed repairs modify internal logic, interface mapping, or lifecycle management, which means the correct repair often belongs in framework code rather than in the user's configuration file.

Example. A vector database adapter changes its expected field name from `metadata` to `payload`. Existing configurations still point to the same database, but retrieval fails because the RAG system maps chunk metadata to the old field. A configuration repair assistant should detect the adapter version, compare the expected schema with the configured mapping, and recommend an interface-level fix. The repair may update the adapter code and add a compatibility test for both field names.

5.2 For RAG Systems Users

Insight 6: Validate external artefacts before deployment.

External configuration mistakes are the most common type of user misconfiguration. These mistakes involve external artefacts such as LLMs, embedding models, vector databases, rerankers, local model files, credentials, and service endpoints.

Guidance. Users should validate external artefacts before relying on a full RAG deployment. A lightweight pre-deployment checklist should include:

- the selected LLM endpoint is reachable and supports the configured context length, model name, authentication mode, and streaming option;
- the embedding model is available and its output dimension matches the vector database collection;

- the vector database is reachable, writable, and uses the expected metric, schema, and collection name;
- any re-ranker, parser, or external application integration uses a compatible API version.

The key is to test each external artefact independently before debugging the entire RAG pipeline.

Example. A user switches from one embedding provider to another because the new model is cheaper. Before reusing the old vector collection, the user should check the new model's embedding dimension. If the old collection was built with 1536-dimensional vectors and the new model returns 1024-dimensional vectors, the correct action is to rebuild the index or create a separate collection. Otherwise, the system may fail during insertion or silently retrieve poor matches.

Insight 7: Pay close attention to configuration semantics and documentation.

Some user misconfigurations arise from misunderstanding option semantics, constraints, deprecated parameters, or misleading documentation. A configuration can be syntactically valid but still conflict with the intended RAG behavior.

Guidance. Users should not treat defaults and option names as complete specifications. Before production deployment, they should inspect the semantic effects of critical options, especially those controlling context construction, chunking, retrieval mode, model selection, prompt templates, and resource limits. When documentation is ambiguous, users should run a small controlled example and inspect intermediate artifacts such as chunks, retrieved passages, and final prompts.

Example. A configuration option named `context_mode=active` may sound as if the system passes the entire active document to the LLM. In a particular implementation, it may instead mean that the active document is used as a search scope for similarity-based retrieval. A user preparing a legal or academic assistant should verify the final prompt for a test query. If only a few retrieved snippets appear, the user may need a different mode, a larger `top_k`, or a custom prompt assembly strategy.

5.3 For RAG Researchers

Insight 8: Explore compatibility-sensitive configuration assistants.

External system adaptation bugs are the most prevalent system-side root cause. This suggests that RAG systems need support for handling evolving APIs, changing model capabilities, incompatible vector database schemas, and heterogeneous component assumptions.

Guidance. Researchers can design assistants that reason about compatibility as a first-class property. Promising directions include:

- mining dependency versions, model cards, API specifications, and adapter code to infer compatibility constraints;
- generating warnings when a selected LLM, embedding model, vector database, or re-ranker violates known constraints;
- learning from historical issue-fix pairs to identify fragile configuration combinations;
- continuously checking compatibility as external services evolve.

Such assistants should explain why a configuration is risky, not merely report that it is invalid. Many failures are caused by implicit assumptions that are not encoded in a single configuration schema.

Example. A research prototype could monitor a RAG project's dependency file, environment variables, and adapter implementations. When the vector database client is upgraded, the assistant checks whether the configured collection schema and metadata mapping still match the adapter's expected API. If not, it suggests a migration plan and generates a regression test that indexes, retrieves, and generates over a small fixture.

Insight 9: Investigate pipeline-aware localization and repair techniques.

Cross-stage issue propagation and the high resolving cost of large-gap issues indicate that RAG debugging tools should model the pipeline structure explicitly. Stage-agnostic bug localization may miss upstream causes that only become visible after several transformations.

Guidance. Future localization and repair techniques should represent RAG execution as a chain of configuration-dependent transformations: initialization, indexing, query transformation, retrieval, post-retrieval processing, and generation. A pipeline-aware tool can then reason over both control flow and data flow:

- which configuration value affects which component;
- which intermediate artifact was produced under which configuration;
- where the first divergence from an expected artifact occurs;
- whether the repair should update user documentation, configuration validation, adapter mapping, or core pipeline logic.

Example. Consider an issue where the final answer cites irrelevant documents. A pipeline-aware localizer should compare the execution trace against a known-good fixture: initialization loads the expected model, indexing creates chunks with correct metadata, query transformation preserves user intent, retrieval returns the expected candidates, post-retrieval filtering keeps valid passages, and generation uses the intended prompt template. If the first divergence appears in post-retrieval filtering because a configured metadata key was renamed during indexing, the repair should target metadata propagation rather than the LLM prompt.

6 Threats to Validity

To mitigate threats to **internal validity**, all authors are involved in the manual analysis for reaching agreements. We also calculate the Cohen's κ coefficient to evaluate inter-rater reliability iteratively—we stop only when there is $\kappa \geq 0.7$ for all categories. Our categorization is also supported by discussion and agreement among all authors, together with external experts when needed. Yet, we agree that we might not be able to completely rule out subjectivity.

External validity concerns the generalizability. Although we analyzed 654 issues from 10 popular RAG systems, which serve as good representations, our results may not fully generalize to all RAG systems, particularly proprietary systems. Our selection criteria favor popular and active open-source repositories with sufficiently rich issue histories. Therefore, the findings may not generalize to proprietary or long-tail RAG systems, newly emerging projects with fewer stars, or systems whose configuration problems are not publicly reported. Moreover, we did not explicitly select or analyze systems according to agentic or cyclic execution patterns. Caution is therefore required when extending our findings to newly emerging types of RAG systems involving planning, tool use, memory, iterative retrieval, or multi-agent orchestration.

Construct validity relates to the metrics/procedure used in the study. We choose Cohen's κ as it has been widely used for empirical studies in software engineering [14, 46]; the same for Spearman correlation and Mann-Whitney U test. Further, we use elapsed calendar time as

an observable proxy for resolving cost rather than a direct measurement of developers' active debugging effort. The interval may also include issue triage, waiting for maintainer responses, implementation scheduling, code review, and periods of inactivity. It might be affected by repository activity, maintainer availability, issue priority and severity, reporting quality, and project-specific development practices.

7 Related Work

Empirical studies on AI systems. A growing body of empirical research investigates reliability challenges in AI-driven software systems [13, 34, 43, 47, 52, 57]. Long and Chen [34] investigate the reporting and resolution of performance and accuracy bugs in deep learning frameworks. Shao et al. [43] study LLM-integrated software applications and systematically analyzed their integration defects, while Xue et al. [52] examine the symptoms and root causes of bugs in LLM agent orchestration frameworks. While these studies provide valuable empirical insights into the reliability and development challenges of AI systems, they primarily focus on deep learning frameworks and LLM-based applications, leaving the configuration of RAG systems largely unexplored. There are studies that systematically characterized the RAG landscape, detailing its architectural evolution and core paradigms [2, 17, 18, 39, 42]. Among them, Gao et al. [18] provide a survey of RAG, categorizing its evolution into different types. Fan et al. [17] systematically review mainstream relevant work by RAG architectures, training strategies, and application areas. Şakar and Emekci [42] address the optimization of RAG processes by further exploring various methodologies. However, those works do not focus on real-world bug patterns or configuration-induced issues of RAG systems. Some other efforts are specific to RAG deployments [1, 4, 30]. Agrawal et al. [1] identify failure points in knowledge graph-based RAG systems, while Leung et al. [30] develop a taxonomy of RAG errors and propose an automated evaluation approach. Barnett et al. [4] identify seven failure points in RAG engineering through three case studies. However, these studies primarily focus on RAG failure modes and output errors, without systematically investigating configuration-related issues and their relationships with the RAG pipeline.

Unlike these studies, our work conducts a systematic empirical investigation of real-world configuration issues in RAG systems, characterizing their symptoms, root causes, and fix patterns. In particular, we examine how configuration issues manifest and originate across different RAG pipeline stages, revealing the discrepancies between symptom and root-cause locations and their implications for issue resolution which are previously unknown.

Empirical studies on configurable systems. Configuration-related errors have long been recognized as a key failure in general software systems. Classic studies [16, 33] analyze real-world misconfigurations to understand their root causes and impact on system availability. Design-centric research [50, 51] explore the challenges of configuration design and usage, highlighting how inadequate configuration handling and excessive configuration options complicate system management and increase the likelihood of user errors. Performance-focused studies [10, 11, 22, 24, 25, 28, 55] characterize configuration-related performance bugs where internal code defects trigger performance degradation. Other studies [22, 28] empirically characterize real-world performance bugs, including the challenges of testing and debugging them in highly configurable systems. Subsequent research [24, 25] specifically investigates configuration-related performance bugs and develops detection techniques based on unexpected performance behaviors under different configuration settings.

Being a new type of system, RAG systems introduce distinct configuration challenges that may limit the direct applicability of findings from traditional configurable systems, for two main reasons:

- **Heterogeneous component integration.** RAG systems are integration-driven applications that depend on diverse external components, such as LLM APIs, embedding models, and vector databases. Their configurations must accommodate evolving interfaces, component dependencies, and compatibility requirements, which forms special categories for RAG systems such as *inter-component implementation bug*.
- **Multi-stage pipeline.** RAG systems operate in their unique multi-stage pipelines, where configuration options influence inter-stage data flow and semantic processing logic. Configuration choices made at one stage can influence downstream data processing and system behavior, potentially causing failures to manifest at stages different from where their root causes originate.

Those characteristics are what make our configuration study unique, as we unveil new, previously unknown findings/insights specifically on RAG systems.

Configuration testing, tuning, and modeling. Configurable systems have been widely researched, including performance modeling [20, 45, 48] for learning the effects of configuration options and their interactions on system performance under varying conditions; configuration tuning [8, 9, 12, 49, 53, 54] for efficiently searching for configurations that optimize system performance; and configuration testing [36] for efficiently detecting configuration-related performance bugs. Additionally, a major line of work focuses on the automated diagnosis of configuration errors and detection of performance bugs [23, 36, 56], leveraging program analysis, historical bug reports, and testing prioritization techniques to identify and diagnose defects. Other studies investigate automated configuration optimization and adaptation for RAG systems [3, 5, 7, 41], aiming to improve answer quality, latency, and cost efficiency through parameter optimization, configuration selection, and workload-aware adaptation. Our study is complementary to those by revealing the unique characteristics of configuration issues for RAG systems.

8 Conclusion

In the paper, we conduct an empirical study consisting of 10 RAG systems and 654 configuration issues on GitHub. From this, we summarize a categorization that documents symptoms, root causes, relevant stages, and fix patterns for configuration issues of RAG systems. In particular, we reveal ten previously unknown findings. Drawing on those, we conclude several actionable insights for future configuration engineering of RAG systems. These results highlight that configuration reliability in RAG systems is not merely an option-selection concern, but a stage-dependent engineering challenge rooted in their integration complexity. We hope this work lays the foundation for more principled debugging and design practices of the next-generation RAG systems.

References

[1] Garima Agrawal, Tharindu Kumarage, Zeyad Alghamdi, and Huan Liu. 2024. Mindful-RAG: A Study of Points of Failure in Retrieval Augmented Generation. In *2nd International Conference on Foundation and Large Language Models, FLLM 2024, Dubai, United Arab Emirates, November 26-29, 2024*. IEEE, 607–611. doi:10.1109/FLLM63129.2024.10852457

[2] Muhammad Arslan, Hussam Ghanem, Saba Munawar, and Christophe Cruz. 2024. A Survey on RAG with LLMs. *Procedia computer science* 246 (2024), 3781–3790.

[3] Matthew Barker, Andrew Bell, Evan Thomas, James Carr, Thomas Andrews, and Umang Bhatt. 2025. Faster, cheaper, better: Multi-objective hyperparameter optimization for LLM and RAG systems. *arXiv preprint arXiv:2502.18635* (2025).

[4] Scott Barnett, Stefanus Kurniawan, Srikanth Thudumu, Zach Brannelly, and Mohamed Abdelrazek. 2024. Seven Failure Points When Engineering a Retrieval Augmented Generation System. In *Proceedings of the IEEE/ACM 3rd International Conference on AI Engineering - Software Engineering for AI, CAIN 2024, Lisbon, Portugal, April 14-15, 2024*, Jane Cleland-Huang, Jan Bosch, Henry Muccini, and Grace A. Lewis (Eds.). ACM, 194–199. doi:10.1145/3644815.3644945

[5] Archisman Bhowmick and Rekha Singhal. 2025. OptiRAG: Optimized RAG Configuration Under Performance Constraints. In *2025 5th International Conference on AI-ML-Systems (AIMLSystems)*. 318–323. doi:10.1109/AIMLSystems67835.2025.11330856

[6] Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, et al. 2024. A survey on evaluation of large language models. *ACM transactions on intelligent systems and technology* 15, 3 (2024), 1–45.

[7] Pengzhou Chen and Tao Chen. 2026. CDS4RAG: Cyclic Dual-Sequential Hyperparameter Optimization for RAG. In *Proceedings of the Thirty-Fifth International Joint Conference on Artificial Intelligence and the 29th European Conference on Artificial Intelligence, IJCAI/ECAI 2026, Bremen, Germany, 15-21 August 2026*. ijcai.org, 4188–4196. doi:10.24963/IJCAI.2026/466

[8] Pengzhou Chen and Tao Chen. 2026. PromiseTune: Unveiling Causally Promising and Explainable Configuration Tuning. In *Proceedings of the 2026 IEEE/ACM 48th International Conference on Software Engineering* (Rio de Janeiro, Brazil) *(ICSE '26)*. Association for Computing Machinery, New York, NY, USA, 14–26. doi:10.1145/3744916.3764552

[9] Pengzhou Chen, Tao Chen, and Miqing Li. 2024. MMO: Meta Multi-Objectivization for Software Configuration Tuning. *IEEE Trans. Software Eng.* 50, 6 (2024), 1478–1504. doi:10.1109/TSE.2024.3388910

[10] Pengzhou Chen, Jingzhi Gong, and Tao Chen. 2025. Accuracy Can Lie: On the Impact of Surrogate Model in Configuration Tuning. *IEEE Trans. Software Eng.* 51, 2 (2025), 548–580. doi:10.1109/TSE.2025.3525955

[11] Pengzhou Chen, Hongyuan Liang, and Tao Chen. 2026. Evaluating Useful Surrogate Models for Configuration Tuning Beyond Accuracy: A Fitness Landscape Analysis Perspective. *ACM Trans. Softw. Eng. Methodol.* (Aug. 2026). doi:10.1145/3844143 Just Accepted.

[12] Tao Chen and Miqing Li. 2021. Multi-objectivizing software configuration tuning. In *Proceedings of the 29th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE '21)*. ACM, 453–465. doi:10.1145/3468264.3468555

[13] Xiang Chen, Chaoyang Gao, Chunyang Chen, Guangbei Zhang, and Yong Liu. 2025. An Empirical Study on Challenges for LLM Application Developers. *ACM Trans. Softw. Eng. Methodol.* 34, 7, Article 205 (Aug. 2025), 37 pages. doi:10.1145/3715007

[14] Jithin Cheriyan, Bastin Tony Roy Savarimuthu, and Stephen Cranefield. 2021. Towards offensive language detection and reduction in four software engineering communities. In *Proceedings of the 25th International Conference on Evaluation and Assessment in Software Engineering*. 254–259.

[15] Daniela S. Cruzes and Tore Dyba. 2011. Recommended Steps for Thematic Synthesis in Software Engineering. In *Proceedings of the 2011 International Symposium on Empirical Software Engineering and Measurement (ESEM '11)*. IEEE Computer Society, USA, 275–284. doi:10.1109/ESEM.2011.36

[16] Constanze Dietrich, Katharina Krombholz, Kevin Borgolte, and Tobias Fiebig. 2018. Investigating System Operators' Perspective on Security Misconfigurations. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security* (Toronto, Canada) *(CCS '18)*. Association for Computing Machinery, New York, NY, USA, 1272–1289. doi:10.1145/3243734.3243794

[17] Wenqi Fan, Yujuan Ding, Liangbo Ning, Shijie Wang, Hengyun Li, Dawei Yin, Tat-Seng Chua, and Qing Li. 2024. A survey on rag meeting llms: Towards retrieval-augmented large language models. In *Proceedings of the 30th ACM SIGKDD conference on knowledge discovery and data mining*. 6491–6501.

[18] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Qianyu Guo, Meng Wang, and Haofen Wang. 2023. Retrieval-Augmented Generation for Large Language Models: A Survey. *CoRR* abs/2312.10997 (2023). arXiv:2312.10997 doi:10.48550/ARXIV.2312.10997

[19] Yunfan Gao, Yun Xiong, Meng Wang, and Haofen Wang. 2024. Modular rag: Transforming rag systems into lego-like reconfigurable frameworks. *arXiv preprint arXiv:2407.21059* (2024).

[20] Jingzhi Gong, Tao Chen, and Rami Bahsoon. 2025. Dividable Configuration Performance Learning. *IEEE Trans. Software Eng.* 51, 1 (2025), 106–134. doi:10.1109/TSE.2024.3491945

[21] Kelvin Guu, Kenton Lee, Zora Tung, Panupong Pasupat, and Mingwei Chang. 2020. Retrieval augmented language model pre-training. In *International conference on machine learning*. PMLR, 3929–3938.

[22] Xue Han and Tingting Yu. 2016. An empirical study on performance bugs for highly configurable software systems. In *Proceedings of the 10th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement*. 1–10.

[23] Xue Han, Tingting Yu, and David Lo. 2018. Perflearner: Learning from bug reports to understand and generate performance test frames. In *Proceedings of the 33rd ACM/IEEE international conference on automated software engineering*. 17–28.

[24] Haochen He. 2019. Tuning backfired? Not (always) your fault: Understanding and detecting configuration-related performance bugs. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 1229–1231.

[25] Haochen He, Zhouyang Jia, Shanshan Li, Erci Xu, Tingting Yu, Yue Yu, Ji Wang, and Xiangke Liao. 2020. Cp-detector: Using configuration-related performance properties to expose performance bugs. In *Proceedings of the 35th IEEE/ACM International Conference on Automated Software Engineering*. 623–634.

[26] Hangfeng He, Hongming Zhang, and Dan Roth. 2022. Rethinking with retrieval: Faithful large language model inference. *arXiv preprint arXiv:2301.00303* (2022).

[27] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *ACM computing surveys* 55, 12 (2023), 1–38.

[28] Guoliang Jin, Linhai Song, Xiaoming Shi, Joel Scherpelz, and Shan Lu. 2012. Understanding and detecting real-world performance bugs. *SIGPLAN Not.* 47, 6 (June 2012), 77–88. doi:10.1145/2345156.2254075

[29] J Richard Landis and Gary G Koch. 1977. The measurement of observer agreement for categorical data. *biometrics* (1977), 159–174.

[30] Kin Kwan Leung, Mouloud Belbahri, Yi Sui, Alex Labach, Xueying Zhang, Stephen Anthony Rose, and Jesse C. Cresswell. 2026. Classifying and Addressing the Diversity of Errors in Retrieval-Augmented Generation Systems. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, Vera Demberg, Kentaro Inui, and Lluís Marquez (Eds.). Association for Computational Linguistics, Rabat, Morocco, 3185–3207. doi:10.18653/v1/2026.eacl-long.147

[31] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, Hugo Larochelle, Marc'Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin (Eds.). https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html

[32] Siran Li, Linus Stenzel, Carsten Eickhoff, and Seyed Ali Bahrainian. 2025. Enhancing retrieval-augmented generation: a study of best practices. In *Proceedings of the 31st International Conference on Computational Linguistics*. 6705–6717.

[33] Yuhao Liu, Yingnan Zhou, Hanfeng Zhang, Zhiwei Chang, Sihan Xu, Yan Jia, Wei Wang, and Zheli Liu. 2024. Rethinking software misconfigurations in the real world: an empirical study and literature analysis. *arXiv preprint arXiv:2412.11121* (2024).

[34] Guoming Long and Tao Chen. 2022. On Reporting Performance and Accuracy Bugs for Deep Learning Frameworks: An Exploratory Study from GitHub. In *Proceedings of the 26th International Conference on Evaluation and Assessment in Software Engineering* (Gothenburg, Sweden) *(EASE '22)*. Association for Computing Machinery, New York, NY, USA, 90–99. doi:10.1145/3530019.3530029

[35] Junliang Luo, Tianyu Li, Di Wu, Michael Jenkin, Steve Liu, and Gregory Dudek. 2024. Hallucination detection and hallucination mitigation: An investigation. *arXiv preprint arXiv:2401.08358* (2024).

[36] Youpeng Ma, Tao Chen, and Ke Li. 2025. *Faster Configuration Performance Bug Testing with Neural Dual-Level Prioritization*. IEEE Press, 988–1000. https://doi.org/10.1109/ICSE55347.2025.00201

[37] Humza Naveed, Asad Ullah Khan, Shi Qiu, Muhammad Saqib, Saeed Anwar, Muhammad Usman, Naveed Akhtar, Nick Barnes, and Ajmal Mian. 2025. A comprehensive overview of large language models. *ACM Transactions on Intelligent Systems and Technology* 16, 5 (2025), 1–72.

[38] Joanne Neale. 2016. Iterative categorization (IC): a systematic technique for analysing qualitative data. *Addiction* 111, 6 (2016), 1096–1106. arXiv:https://onlinelibrary.wiley.com/doi/pdf/10.1111/add.13314 doi:10.1111/add.13314

[39] Agada Joseph Oche, Ademola Glory Folashade, Tirthankar Ghosal, and Arpan Biswas. 2025. A systematic review of key retrieval-augmented generation (rag) systems: Progress, gaps, and future directions. *arXiv preprint arXiv:2507.18910* (2025).

[40] Vipula Rawte, Amit Sheth, and Amitava Das. 2023. A survey of hallucination in large foundation models. *arXiv preprint arXiv:2309.05922* (2023).

[41] Siddhant Ray, Rui Pan, Zhuohan Gu, Kuntai Du, Shaoting Feng, Ganesh Ananthanarayanan, Ravi Netravali, and Junchen Jiang. 2025. Metis: fast quality-aware rag systems with configuration adaptation. In *Proceedings of the ACM SIGOPS 31st symposium on operating systems principles*. 606–622.

[42] Tolga Şakar and Hakan Emekci. 2025. Maximizing RAG efficiency: A comparative analysis of RAG methods. *Natural Language Processing* 31, 1 (2025), 1–25.

[43] Yuchen Shao, Yuheng Huang, Jiawei Shen, Lei Ma, Ting Su, and Chengcheng Wan. 2025. Are LLMs Correctly Integrated into Software Systems?. In *47th IEEE/ACM International Conference on Software Engineering, ICSE 2025, Ottawa, ON, Canada, April 26 - May 6, 2025*. IEEE, 1178–1190. doi:10.1109/ICSE55347.2025.00204

[44] Xinyue Shen, Zeyuan Chen, Michael Backes, and Yang Zhang. 2023. In chatgpt we trust? measuring and characterizing the reliability of chatgpt. *arXiv preprint arXiv:2304.08979* (2023).

[45] Norbert Siegmund, Alexander Grebhahn, Sven Apel, and Christian Kästner. 2015. Performance-influence models for highly configurable systems. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*. 284–294.

[46] Simin Wang, Liguo Huang, Amiao Gao, Jidong Ge, Tengfei Zhang, Haitao Feng, Ishna Satyarth, Ming Li, He Zhang, and Vincent Ng. 2023. Machine/Deep Learning for Software Engineering: A Systematic Literature Review. *IEEE Transactions on Software Engineering* 49, 3 (2023), 1188–1231. doi:10.1109/TSE.2022.3173346

[47] Siyu Wu, Yulong Ye, Zezhen Xiang, Pengzhou Chen, Gangda Xiong, and Tao Chen. 2026. LLMSYS-HPOBench: Hyperparameter Optimization Benchmark Suite for Real-World LLM Systems. *CoRR* abs/2605.08305 (2026). arXiv:2605.08305 doi:10.48550/ARXIV.2605.08305

[48] Zezhen Xiang, Jingzhi Gong, and Tao Chen. 2026. Dually Hierarchical Drift Adaptation for Online Configuration Performance Learning. In *48th IEEE/ACM International Conference on Software Engineering (ICSE)*. ACM.

[49] Gangda Xiong and Tao Chen. 2025. CoTune: Co-evolutionary Configuration Tuning. In *40th IEEE/ACM International Conference on Automated Software Engineering, ASE 2025, Seoul, Korea, Republic of, November 16-20, 2025*. IEEE, 1490–1502. doi:10.1109/ASE63991.2025.00126

[50] Tianyin Xu, Long Jin, Xuepeng Fan, Yuanyuan Zhou, Shankar Pasupathy, and Rukma Talwadker. 2015. Hey, you have given me too many knobs!: Understanding and dealing with over-designed configuration in system software. In *Proceedings of the 2015 10th Joint Meeting on Foundations of Software Engineering*. 307–319.

[51] Tianyin Xu, Jiaqi Zhang, Peng Huang, Jing Zheng, Tianwei Sheng, Ding Yuan, Yuanyuan Zhou, and Shankar Pasupathy. 2013. Do not blame users for misconfigurations. In *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles* (Farminton, Pennsylvania) *(SOSP '13)*. Association for Computing Machinery, New York, NY, USA, 244–259. doi:10.1145/2517349.2522727

[52] Ziluo Xue, Yanjie Zhao, Shenao Wang, Kai Chen, and Haoyu Wang. 2025. A Characterization Study of Bugs in LLM Agent Workflow Orchestration Frameworks. In *2025 40th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 3369–3380.

[53] Yulong Ye, Tao Chen, and Miqing Li. 2025. Distilled Lifelong Self-Adaptation for Configurable Systems. In *47th IEEE/ACM International Conference on Software Engineering, ICSE 2025, Ottawa, ON, Canada, April 26 - May 6, 2025*. IEEE, 1333–1345. doi:10.1109/ICSE55347.2025.00094

[54] Yulong Ye, Miqing Li, and Tao Chen. 2026. Less Is More: Tuning Configurable Systems with Imperfect Fidelity *(ASE '26)*. Association for Computing Machinery, New York, NY, USA, 575–587. doi:10.1145/3832783.3834376

[55] Yulong Ye, Hongyuan Liang, Chao Jiang, Miqing Li, and Tao Chen. 2026. Revealing Domain-Spatiality Patterns for Configuration Tuning: Domain Knowledge Meets Fitness Landscapes. *ACM Trans. Softw. Eng. Methodol.* (March 2026). doi:10.1145/3803859 Just Accepted.

[56] Sai Zhang and Michael D. Ernst. 2013. Automated diagnosis of software configuration errors. In *Proceedings of the 2013 International Conference on Software Engineering* (San Francisco, CA, USA) *(ICSE '13)*. IEEE Press, 312–321.

[57] Tianyi Zhang, Cuiyun Gao, Lei Ma, Michael Lyu, and Miryung Kim. 2019. An Empirical Study of Common Challenges in Developing Deep Learning Applications. In *2019 IEEE 30th International Symposium on Software Reliability Engineering (ISSRE)*. 104–115. doi:10.1109/ISSRE.2019.00020

[58] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, et al. 2023. A survey of large language models. *arXiv preprint arXiv:2303.18223* 1, 2 (2023), 1–124.