Evaluating Agentic Capabilities Before Post-Training: An End-to-End Approach

Jiarui Qin[1†], Junjie Huang[2†], Yunjia Xi[1], Di Yin[1], Yong Yu[2],
Weinan Zhang[2], Xing Sun[1], Weiwen Liu[2‡]

[1]Tencent Youtu Lab
[2]Shanghai Jiao Tong University

arXiv:2610.19425v1 [cs.AI] 14 Oct 2026

Agentic capabilities are increasingly important for modern LLMs, yet remain difficult to evaluate during pretraining. Existing approaches typically rely on static trajectory-derived tasks or likelihood-based proxies, which are lightweight but provide limited visibility into how base models behave in interactive environments. We introduce **ASAP (Agent Scaffold Adaptation for Pretraining Evaluation)**, a lightweight framework for eliciting and evaluating agentic behavior from pretrained base models on end-to-end tasks. Rather than converting base models into chat-style agents, ASAP preserves the completion interface and adapts existing agent scaffolds through structured trajectory continuation, short reference prefixes, and minimal recovery mechanisms. Across software-engineering and deep-research benchmarks, ASAP yields discriminative end-to-end performance across base models and pretraining checkpoints. Beyond aggregate scores, ASAP exposes model-generated execution trajectories, enabling analysis of exploration, editing, testing, recovery, and task completion across training stages. These trajectories provide a behavioral view of how agent-relevant capabilities change during pretraining and how execution patterns differ after post-training. Our results suggest that end-to-end agent evaluation can complement conventional pretraining benchmarks by providing a practical diagnostic of interactive model behavior.

1 Introduction

Agentic evaluation is increasingly centered on long-horizon, interactive tasks. Benchmarks such as SWE-bench [Jimenez et al., 2024], DeepSWE [Huang et al., 2026], and Terminal-Bench [Merrill et al., 2026] evaluate sustained software-engineering and terminal interaction, while challenging knowledge-oriented benchmarks such as Humanity's Last Exam (HLE) [Phan et al., 2025] continue to raise the capability frontier. Yet these evaluations are used almost exclusively *after* post-training.

During pretraining, however, there remains no established way to evaluate how base models behave on end-to-end agentic tasks. Such an evaluation could help track changes across pretraining checkpoints, characterize differences between training stages, and expose behavioral deficiencies that are difficult to observe through conventional static benchmarks alone.

The main challenge is that existing agent benchmarks assume instruction-following models: their scaffolds rely on system prompts, chat templates, prescribed tool formats, and stable multi-turn interaction. Pretrained base models are not explicitly trained for these protocols, and directly inserting them into standard agent scaffolds often produces malformed or unstable trajectories. One alternative is to fine-tune each checkpoint before evaluation, but repeated SFT introduces additional cost and changes the model being evaluated. As a result, existing pretraining-side evaluations largely rely on static or likelihood-based proxies, such as

† Equal contribution. jaredqin@tencent.com, huangjunjie2019@sjtu.edu.cn
‡ Corresponding author.

trajectory-derived QA tasks [Qin et al., 2025] or code and trajectory likelihood [Cursor Research et al., 2026, DeepSeek-AI, 2026]. These signals are useful for model comparison, but provide limited visibility into how the evaluated model itself interacts with an environment.

This distinction matters because scalar proxy scores cannot show whether a model explores the environment, makes useful edits, tests its hypotheses, revises after failed attempts, or simply fails to make progress. End-to-end interaction therefore provides a complementary view: in addition to final task success, it exposes the execution process through which that outcome is reached. This leads to the central question of our work:

Can we evaluate pretrained base models on end-to-end agentic tasks without first turning them into instruction-following models?

We introduce **ASAP (Agent Scaffold Adaptation for Pretraining Evaluation)**, a lightweight framework for adapting existing agent scaffolds to pretrained base models. Our key principle is that **base models should be evaluated as continuation models rather than incomplete chat models**. ASAP preserves the completion interface and modifies only the model-facing interaction protocol through three simple components: a Markdown+XML trajectory format, a short reference trajectory prefix, and minimal recovery from structural failures such as premature submission or repeated actions.

Across software-engineering and deep-research settings, ASAP yields discriminative end-to-end performance across base models and remains sensitive to differences between checkpoints within the same pretraining run. In our pretraining-stage study, ASAP continues to separate checkpoints even when conventional pretraining benchmarks change only marginally. More importantly, the resulting execution trajectories make it possible to characterize these differences in behavioral terms, including changes in exploration, editing, testing, recovery, and task completion. Comparing base and post-trained models further illustrates how execution patterns differ after post-training.

These results suggest that end-to-end agent evaluation can complement conventional pretraining benchmarks by providing a practical behavioral diagnostic of model behaviors.

Our contributions are:

- **End-to-end evaluation of pretrained base models.** We introduce ASAP, a lightweight adaptation of existing agent scaffolds that enables pretrained base models to execute agentic tasks through their native completion interface, without weight updates.
- **A sensitive probe of agent-relevant behavior during pretraining.** ASAP distinguishes both different base models and checkpoints within the same training run, including stages where conventional pretraining benchmarks show only small changes.
- **Trajectory-level diagnosis of behavioral change.** ASAP exposes model-generated execution trajectories, enabling fine-grained analysis of how exploration, editing, testing, recovery, and task completion vary across pretraining checkpoints and after post-training.

2 Related Work

Existing efforts to evaluate agentic capability before post-training have largely relied on proxy tasks rather than end-to-end agent execution. APTBench [Qin et al., 2025] converts successful agent trajectories into static multiple-choice and text-completion problems suitable for base models, while PACE [Song et al., 2026] predicts agent performance from a set of static benchmark items through cross-model calibration. Similarly, model-development pipelines often use coding knowledge, state tracking, or trajectory likelihood as pretraining-side signals. For example, Composer 2 [Cursor Research et al., 2026] uses codebase negative

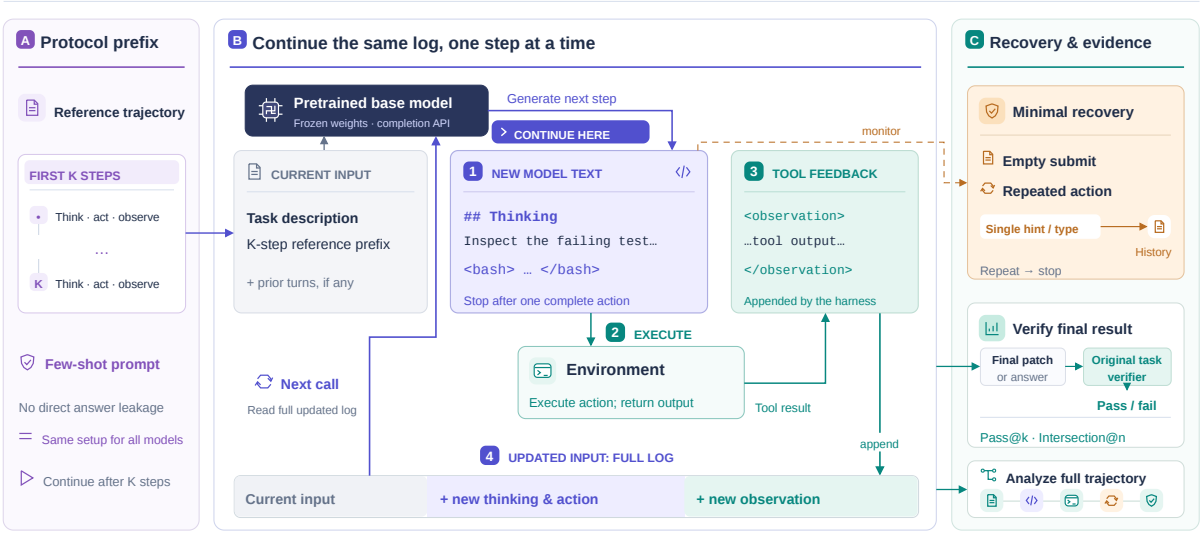

Figure 1. ASAP enables end-to-end evaluation of pretrained base models. (A) A short prefix from a reference trajectory. **(B)** The frozen base model continues the trajectory through its native completion interface. **(C)** Minimal recovery handles premature empty submissions and repeated actions.

log-likelihood for base-model selection, and trajectory BPB/NLL has also been explored as a proxy for downstream agentic capability [DeepSeek-AI, 2026].

These approaches provide lightweight signals for model selection, but they evaluate fixed examples or reference trajectories rather than the behavior generated by the evaluated model itself. ASAP instead adapts existing agent scaffolds so that pretrained base models can execute the original tasks end-to-end through their native completion interface. This enables not only direct evaluation before post-training, but also trajectory-level analysis of how agentic behavior evolves across pretraining checkpoints and is subsequently affected by post-training.

3 Methodology

We evaluate a pretrained base model directly through a simple text completions task, treating agent execution as a continuation problem rather than converting the model into a chat-style interface. In particular, we do not apply a chat template or system prompt. Instead, we adapt the underlying agent scaffold with a small set of simple design principles that make base-model generation sufficiently stable for end-to-end agent evaluation. ASAP is built around three components: (1) representing each interaction step in a structured Markdown+XML trajectory format, (2) bootstrapping the model with a short reference prefix that serves as in-context guidance for how to continue the trajectory, and (3) lightweight recovery from structural failures such as repeated actions and empty submit to prevent degenerate loops. Figure 1 provides an overview of the complete ASAP evaluation pipeline.

3.1 Markdown and XML Shell-Log Format

We formulate agent interaction as a *structured text continuation* task rather than a conversational one. Each step uses Markdown for free-form thinking and XML tags (e.g., `<bash>` and `<observation>`) to delimit

actions and environment feedback. This representation resembles structured text and shell-log patterns commonly seen during pretraining, making it natural for base models to continue. Among alternative formats including pure JSON, XML, and Markdown, the hybrid Markdown+XML format yields the most stable trajectories in our experiments (Appendix A.4.1). An illustrative format example is provided in Appendix A.2.

Thought-action generation. At each agent step, a single completion call generates both the model's thinking and its subsequent action. Generation starts from `## Thinking` and stops after a complete action (e.g., `</bash>` or a terminal action). The harness executes the action, appends the resulting output as an `<observation>` block, and continues the trajectory. Execution terminates upon `<submit/>`, the maximum number of steps, or detection of a degenerate repeated-action loop.

3.2 Reference Trajectory Prefix Bootstrapping

Unlike instruction-tuned models, pretrained base models are not explicitly trained to initiate a prescribed agent workflow from a standalone task instruction. We therefore prepend a short segment from a reference agent trajectory, turning the evaluation into continuation of an already established interaction process. The prefix plays a role similar to few-shot context in conventional base-model evaluation: it helps the model enter the intended interaction pattern and provides lightweight task navigation, while leaving the substantive task solving to the evaluated model.

For each instance, we take the initial steps from a reference trajectory produced by an instruction-tuned model and use them as the bootstrap prefix. These steps may provide limited task-relevant guidance, such as exposing the repository structure or inspecting potentially relevant files, but we verify that they do not directly reveal the solution or contain the target patch. The remaining trajectory is then generated autonomously by the evaluated base model. For each instance, the same prefix is used across all evaluated models, providing a controlled basis for comparison across models and pretraining checkpoints.

Additional details on prefix construction, length ablations, and robustness to reference-source resampling are provided in Appendix A.4.2.

3.3 Lightweight Recovery from Structural Failures

Even with trajectory-format guidance and reference-prefix bootstrapping, base models can occasionally fall into structural failure modes from which they do not recover autonomously. We observe two recurring patterns during harness development: premature submission before making any effective edit, and repeated execution of the same action. We apply a lightweight intervention to each case, providing the model with a single opportunity to recover before terminating the trajectory.

Premature submission recovery. Some base models exhibit a *submit bias*, emitting `<submit/>` before modifying any source file. On the first such attempt, we reject the submission and append a short hint indicating that at least one source file should be edited before submission. The model is then allowed to continue autonomously. If it again attempts to submit without producing an effective edit, we accept the submission as an empty patch and terminate the trajectory.

Repeated-action recovery. Another common failure mode is repeated actions, in which the model issues the same action without making progress. When such repetition is first detected, we suppress the redundant action and append a short hint asking the model to try a different approach. If the model subsequently falls into the same repeated-action pattern again, we terminate the trajectory.

These interventions are deliberately minimal: each failure mode triggers at most one textual hint, and the harness never inserts task-specific actions or edits on behalf of the model. If the base model cannot recover after this hint, the trajectory is simply terminated. The intervention therefore serves only to prevent trivial structural degeneration from dominating the evaluation, while leaving task solving to the base model itself.

3.4 Implementation & Evaluation

Patching Agent Scaffolds. We implement our method as lightweight patches to existing agent scaffolds, including SWE-agent [Yang et al., 2024], OpenHands [Wang et al., 2024], Terminus, and BrowseComp. These patches adapt the model-facing interaction loop to support base-model continuation while leaving the underlying task environment and scaffold logic largely unchanged, allowing the same evaluation approach to be applied across different agentic task settings.

Metrics. Our primary metric is Pass@k, computed using the unbiased estimator [Chen et al., 2021]. Pass@k captures whether a task is solvable by the model under repeated sampling. To additionally measure execution stability, we report **Intersection@n**, defined as the number of instances successfully resolved in all n runs. While Pass@k reflects the model's attainable capability, Intersection@n captures how consistently that capability is realized.

4 Main Evaluation Results

We first compare six open-source base models on SWE-bench Verified [Jimenez et al., 2024], BrowseComp-Plus [Chen et al., 2025] and Terminal Bench 2.1 [Merrill et al., 2026]. We then conduct ablation study of ASAP and extensive analysis by using ASAP as an effective pre-training probe.

4.1 Experimental Setup

We evaluate six open-source base models spanning the Qwen [Qwen Team, 2025, 2026], GLM [Zeng et al., 2025], Gemma [Gemma Team, 2026], and Mimo [MiMo Team, 2026] families. The selection covers same-family size scaling (Qwen3.5-9B vs. Qwen3.5-35B), cross-generation comparison (Qwen3-30B vs. Qwen3.5-35B), and cross-family comparisons. Every model is evaluated on SWEBench Verified, Terminal Bench 2.1 and BrowseComp-Plus using identical ASAP configuration.

4.2 Main Results

Table 1 summarizes the performance of six open-weight base models across different agentic benchmarks under ASAP . Each model is evaluated with eight independent rollouts.

Table 1. Main results (%) across agentic benchmarks. We report Pass@1, Pass@8, and the number of instances resolved in all eight runs (#Inter). Full Pass@k results are provided in Appendix A.3. **Bold**: best result within each benchmark–scaffold setting.

	SWE-bench Verified		*SWE-agent*	**SWE-bench Verified**		*OpenHands*
Model	**Pass@1**	**Pass@8**	**#Inter**	**Pass@1**	**Pass@8**	**#Inter**
Qwen3-30B-A3B-Base	7.38	24.20	0	4.00	14.40	0
Qwen3.5-9B-Base	21.85	46.60	16	32.95	55.40	51
Qwen3.5-35B-A3B-Base	34.23	56.20	67	**39.23**	**61.80**	69
Gemma4-26B-A4B-Base	7.12	23.00	0	4.67	18.20	0
GLM-4.5-Base	15.97	41.40	4	19.73	44.80	11
MiMo-V2.5-Base	**37.70**	**59.40**	**82**	37.52	60.80	**70**
	Terminal-Bench 2.1		*Terminus-2*	**BrowseComp-Plus**		*Native*
Model	**Pass@1**	**Pass@8**	**#Inter**	**Pass@1**	**Pass@8**	**#Inter**
Qwen3-30B-A3B-Base	13.76	20.22	5	17.36	32.89	29
Qwen3.5-9B-Base	19.80	30.34	10	20.53	34.94	54
Qwen3.5-35B-A3B-Base	**26.69**	**44.94**	**11**	26.31	43.01	97
Gemma4-26B-A4B-Base	12.50	20.22	4	14.64	30.72	20
GLM-4.5-Base	16.01	33.71	6	23.34	37.47	80
MiMo-V2.5-Base	25.84	41.57	10	**29.44**	**46.51**	**125**

Key observations. (1) ASAP provides meaningful discrimination among base models across diverse agentic settings. All evaluated models achieve non-trivial end-to-end success, while their performance differs substantially across benchmarks. On SWE-bench Verified, Pass@1 ranges from 7.12% to 37.70% with SWE-agent and from 4.00% to 39.23% with OpenHands. The separation persists beyond software-repository repair: on Terminal-Bench 2.1, Pass@1 ranges from 12.50% to 26.69%, while on BrowseComp-Plus it ranges from 14.64% to 29.44%. Qwen3.5-35B-A3B-Base and MiMo-V2.5-Base consistently achieve strong performance across these heterogeneous settings, suggesting that ASAP captures meaningful differences in base-model agentic capability across both benchmarks and scaffolds.

(2) Broad capability tiers are preserved across settings, while exact rankings remain scaffold- and domain-dependent. On SWE-bench Verified, the Pass@1 rankings under SWE-agent and OpenHands are strongly correlated (Spearman $\rho \approx 0.89$), and the stronger and weaker models remain broadly separated across both scaffolds. At the same time, exact rankings and absolute performance can change across scaffolds and benchmarks. For example, Qwen3.5-9B-Base improves from 21.85% under SWE-agent to 32.95% under OpenHands, while Qwen3-30B-A3B-Base decreases; similarly, the relative ordering of Qwen3.5-35B-A3B-Base and MiMo-V2.5-Base varies across Terminal-Bench 2.1, SWE-bench, and BrowseComp-Plus. These results suggest that agentic capability contains both transferable components and setting-specific behaviors: strong models tend to remain strong overall, but different scaffolds and domains elicit different parts of that capability. This motivates training for more scaffold-invariant agentic behaviors and broader domain coverage, rather than optimizing for a particular interaction protocol or task distribution.

(3) Repeated sampling separates occasional solvability from stable capability. The strongest models combine high Pass@1 with large all-run intersections. Under SWE-agent, MiMo-V2.5-Base reaches 37.70% Pass@1 with 82 instances solved in all eight runs, while Qwen3.5-35B-A3B-Base reaches 34.23% with an intersection of 67. In contrast, weaker models can benefit substantially from repeated sampling without exhibiting a comparable stable core. Qwen3-30B-A3B-Base, for example, rises from 7.38% Pass@1 to 24.20% Pass@8, yet solves no instance in all eight runs. These results show that multi-run evaluation reveals an important distinction between capabilities that can be elicited occasionally and those that can be executed reliably.

Table 2. Training stages of the internal 30B-A3B MoE base model *M*. Trajectory-like data is first introduced at stage3.

	Stage1	**Stage2**	**Stage3**	**Stage4**	**Stage5**
Training phase	General	LR decay		Long-context training	
Context length	8K	8K	32K	128K	256K
Trajectory data	×	×	✓	✓	✓

The detailed ablation studies on bootstrap prefix steps, input format design are elaborated in Appendix A.4.

5 Extensive Analysis

In this section, we use ASAP to study how agentic capabilities evolve across pretraining and how they are subsequently affected by post-training. We first analyze capability changes across pretraining stages and decompose them into specific behavioral transitions in Section 5.1, and then examine the changes from base to post-trained models in Section 5.2. We further find that the model ordering exposed by ASAP is preserved after applying the same lightweight SFT recipe, and the potential evaluation hacking through external network access is scarce for base models. These details can be found in Appendix A.5 and A.6.

5.1 Monitoring Pretraining Progress and Behavioral Transitions

We evaluate checkpoints from different pretraining stages of our self-trained 30B-A3B MoE base model *M* with ASAP. By using ASAP we could examine how agentic capability evolves over the course of pretraining. Specifically, we study whether ASAP can sensitively distinguish checkpoints at different stages of pretraining, and, beyond aggregate performance gains, identify which agentic behaviors improve as pretraining progresses. The latter analysis is enabled by the execution trajectories collected through our end-to-end evaluation: rather than observing only changes in final scores, we can decompose these gains into changes in specific agent behaviors. This demonstrates that ASAP serves not only as a discriminative benchmark for pretraining checkpoints, but also as a diagnostic tool for understanding how agentic capabilities emerge and evolve during pretraining.

Experimental setup. We evaluate nine checkpoints spanning five stages with ASAP on SWE-bench Verified. Mid-stage checkpoints are included to measure within-stage progress (stagex-mid means $\frac{2}{3}$ training tokens of stage x). Stage1 is general pretraining; stage2 applies learning-rate decay over a higher-quality mixture with QA and instruction-following data but no agentic trajectories; stage3 first introduces trajectory-like data together with 32K long-context training; and stages 4 and 5 continue trajectory training while extending the training context to 128K and 256K, respectively. For comparison, we use the aggregate of our conventional pretraining evaluations , which cover knowledge & factual QA, mathematics, reasoning and code.

5.1.1 Cross-stage performance on ASAP and general benchmarks

Figure 2(a) shows how performance on ASAP evolves across pretraining stages. The largest transition occurs from stage2 to stage3, where trajectory-like data is introduced together with native 32K long-context training.

MMLU (-pro,-redux), SimpleQA, INCLUDE, GPQA, GSM8K, MGSM, BBH, MATH500, EvalPlus, Livecodebench, CruxEval, RepoBench, etc

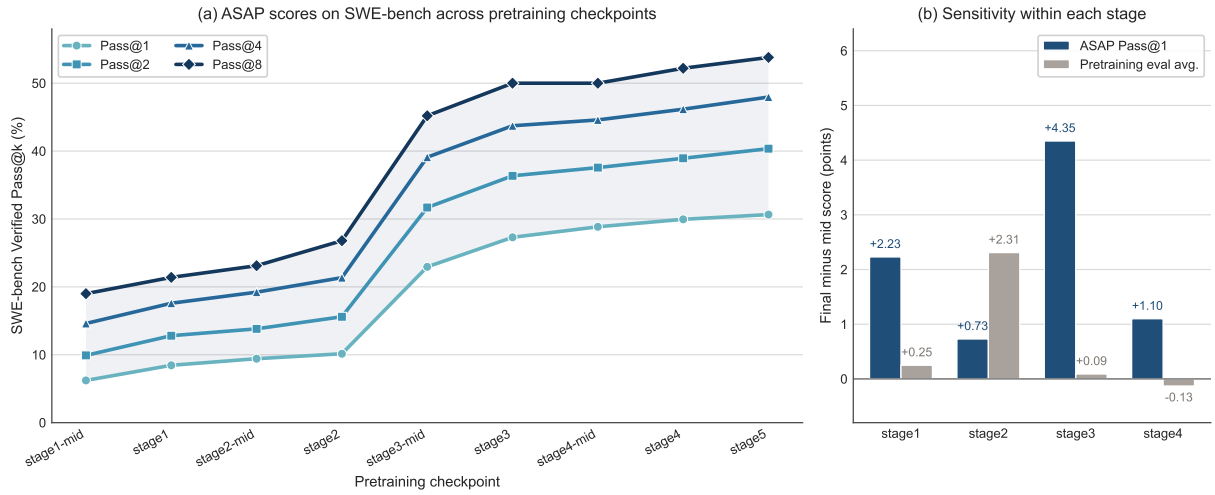

Figure 2. Pretraining progress of internal model *M*. **(a)** ASAP Pass@k with SWE-agent on SWE-bench Verified across nine checkpoints. **(b)** ASAP scores sensitivity within stage1 to stage4, comparing with the conventional pretraining evaluations.

While the stage2 model is evaluated beyond its native 8K context using YaRN extrapolation, the stage3 transition still jointly changes both the training data and long-context capability.

Across this transition, Pass@1 increases from 10.15% at the end of stage2 to 27.30% at the end of stage3, with stage3-mid already reaching 22.95%. Similar improvements are observed across different values of k. Later stages bring more moderate improvements, with Pass@1 reaching 29.95% at stage4 and 30.65% at stage5. Overall, ASAP sensitively captures the pronounced agent-relevant capability transition around stage3 and the more gradual improvements that follow.

We next examine the within-stage comparisons in Figure 2(b), which provide a more controlled view of how sensitively different evaluations distinguish checkpoints during pretraining. ASAP consistently separates the midpoint and final checkpoint within each stage. This distinction is particularly pronounced in stages 1, 3, and 4: ASAP Pass@1 improves by 2.23, 4.35, and 1.10 percentage points, respectively, whereas the conventional benchmarks scores changes by only +0.25, +0.09, and -0.13 points. In other words, once training enters the latter half of these stages, the aggregate conventional benchmark score changes only marginally, while ASAP continues to show measurable improvements.

Stage2 exhibits a different pattern, with the conventional benchmark improving more than ASAP (+2.31 points versus +0.73pp in Pass@1). Stage2 uses a higher-quality annealing mixture containing QA and instruction-following data together with learning-rate decay, and is therefore specifically optimized to improve conventional knowledge, reasoning, and coding performance without introducing agent trajectory data.

Taken together, these results show that ASAP provides strong discrimination both across pretraining stages and among checkpoints within the same stage, and is substantially more sensitive than conventional base-model benchmarks to continued improvements in agentic capability, even in the general pretraining stage (stage1).

5.1.2 Behavioral decomposition

To further investigate what changes at each stage, we compare checkpoints of each stage with the adjacent stages. Each pair retains trajectories with the same SWE-bench instance and run index, 4,000 (8 runs) paired trajectories per checkpoint across all 500 tasks; prefix source and prefix length K are also matched. We group diagnostics into *interface* behavior (turn parsability, valid tool use), *agentic execution* (relevant-file access, autonomous source edits, active testing), and *outcomes* (applicable and resolving patches). The results are reported in Figure 3.

Stage1→stage2: instruction and interface acquisition. Parsable turns increase from 88.81% to 99.41%, and valid tool calls from 79.78% to 97.06%. Yet only 36.33% of stage2 trajectories produce an applicable patch and 10.15% resolve the task. Together with the introduction of QA/instruction-following data and the absence of agentic trajectories, this separation indicates that stage2 primarily improves interface reliability rather than agentic abilities. Relevant-file access also improves, but autonomous edits and active testing improves in a relatively slow manner, resulting in modest resolved rate improvements.

Stage2→stage3: agentic task execution. After trajectory data and 32K training are introduced, parsability and valid tool calls remain high, but autonomous edits rise to 42.86%, active testing to 65.64%, applicable patches to 63.88%, and resolution to 27.30%. We also find that repetitive command loops and abnormal termination also become substantially less common in stage3. The score increase is accompanied primarily by improvements in execution behaviors rather than interface compliance.

Stages 4 and 5 exhibit modest changes across the individual capability dimensions, which is consistent with the smaller improvement in resolved rate observed at these stages. More broadly, these results highlight an important advantage of ASAP : it allows us to decompose aggregate performance changes into more fine-grained behavioral changes. Such trajectory-level analysis provides a more informative view of how capabilities evolve during pretraining, and can in turn help diagnose deficiencies and guide the design of pretraining data. This type of behavioral attribution is fundamentally unavailable to conventional accuracy-based pretraining benchmarks, which observe only the correctness of the final prediction.

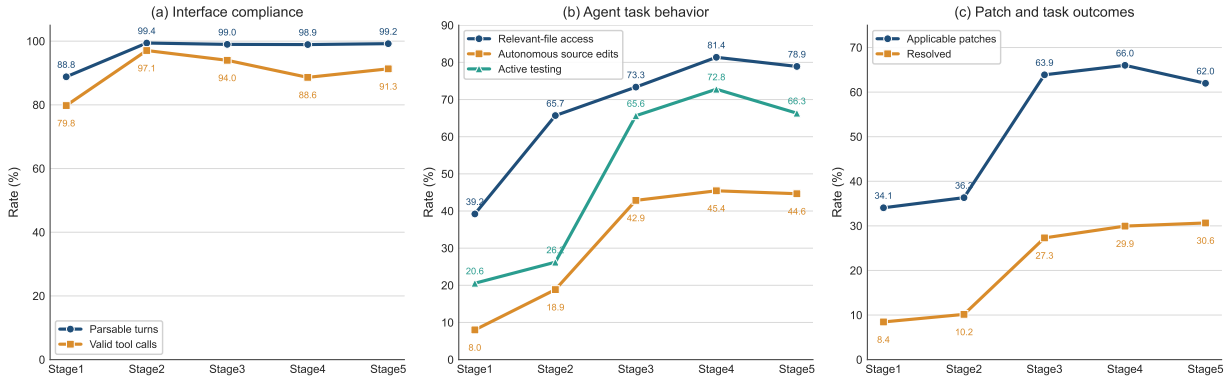

Figure 3. Behavioral change across different pretraining stages. Stage2 brings parsability and valid tool calls close to saturation, whereas the main transition in autonomous editing, active testing, applicable patches, and task resolution occurs at stage3.

5.2 Behavioral Differences Before and After Post-Training

We compare base and post-trained models from two families, Qwen3.5-35B-A3B and MiMo-V2.5, on SWE-bench Verified. Base models are evaluated with ASAP , while post-trained models use the original SWE-agent scaffold. Because the protocols differ by design, this comparison is descriptive rather than a scaffold-controlled estimate of the effect of post-training. Table 3 summarizes the behavioral differences.

Table 3. Behavioral differences between base and post-trained models under their respective evaluation protocols. Δ is in percentage points.

	Qwen3.5-35B-A3B			MiMo-V2.5		
Metric	**Base**	**Post**	**Δ**	**Base**	**Post**	**Δ**
Pass@1 / success rate	34.23	49.40	+15.17	37.70	74.00	+36.30
Valid source edit	69.99	79.41	+9.42	62.12	96.34	+34.22
Active testing	85.90	96.80	+10.90	76.80	99.77	+22.97
Valid re-edit after a failed test	31.80	63.78	+31.98	28.28	57.27	+28.99
Explicit submission	20.57	52.17	+31.60	26.81	89.02	+62.21
Empty patch	22.63	13.04	−9.59	34.67	1.60	−33.07
Exact match to gold modified-file set	44.92	61.05	+16.13	57.42	72.09	+14.67

More reliable execution after post-training. Across both families, post-trained models exhibit a more reliable edit-test-recover-submit pattern. They are more likely to revise their implementation after failed tests, produce fewer empty patches, and explicitly submit a final solution. Their modified-file sets also align more closely with the gold patch scope. Taken together, these differences suggest that several components of the agent workflow are already present in base models, but are executed more consistently and with better completion reliability in the post-trained models.

Different bottlenecks in base models. The two base models nevertheless exhibit different failure patterns. Qwen-Base already performs source editing and active testing in a large fraction of trajectories, suggesting that the basic execution workflow can often be elicited before post-training; its main weaknesses lie in recovery, refinement, and reliable completion. MiMo-Base, in contrast, more often remains in the exploration stage and fails to transition into effective editing. These patterns illustrate how ASAP can reveal not only differences in overall task success, but also where the execution pipeline breaks down for different pretrained models.

6 Conclusion

We introduced ASAP, a lightweight framework for directly evaluating pretrained base models on end-to-end agentic tasks without weight updates or conversion to a chat-style interface. By adapting existing agent scaffolds through structured continuation, short trajectory prefixes, and minimal recovery mechanisms, ASAP enables stable and discriminative evaluation across models, checkpoints, and agent environments. More importantly, end-to-end execution exposes the models' own interaction trajectories, allowing capability gains during pretraining to be decomposed into concrete behavioral changes in exploration, editing, testing, and recovery, and enabling direct comparison of these behaviors between pre-training and post-training. Our results suggest that agent evaluation can serve not only as a downstream benchmark, but also as a practical diagnostic during pretraining, providing richer signals for understanding capability development and guiding future data and training design.

References

[1] Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? *arXiv preprint arXiv:2310.06770*, 2024.

[2] Wenqi Huang, Charley Lee, Leonard Tng, and Serena Ge. Deepswe: Measuring frontier coding agents on original, long-horizon engineering tasks, 2026. URL https://arxiv.org/abs/2607.07946.

[3] Mike A. Merrill, Alexander G. Shaw, Nicholas Carlini, Boxuan Li, Harsh Raj, Ivan Bercovich, Lin Shi, Jeong Yeon Shin, Thomas Walshe, E. Kelly Buchanan, et al. Terminal-bench: Benchmarking agents on hard, realistic tasks in command line interfaces. In *The Fourteenth International Conference on Learning Representations*, 2026. URL https://openreview.net/forum?id=a7Qa4CcHak.

[4] Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, et al. Humanity's last exam, 2025. URL https://arxiv.org/abs/2501.14249.

[5] Jiarui Qin, Yunjia Xi, Junjie Huang, Renting Rui, Di Yin, Weiwen Liu, Yong Yu, Weinan Zhang, and Xing Sun. Aptbench: Benchmarking agentic potential of base llms during pre-training, 2025. URL https://arxiv.org/abs/2510.24397.

[6] Cursor Research, Aaron Chan, Ahmed Shalaby, Alexander Wettig, Aman Sanger, Andrew Zhai, Anurag Ajay, Ashvin Nair, Charlie Snell, Chen Lu, et al. Composer 2 technical report, 2026. URL https://arxiv.org/abs/2603.24477.

[7] DeepSeek-AI. Deepseek-v4.1-flash: Pushing the limits of kv cache compression, 2026.

[8] Yueqi Song, Lintang Sutawika, Jiarui Liu, Lindia Tjuatja, Jiayi Geng, Yunze Xiao, Daniel Lee, Aditya Bharat Soni, Vincent Lo, Xiang Yue, and Graham Neubig. Pace: A proxy for agentic capability evaluation, 2026.

[9] John Yang, Carlos E Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering. *arXiv preprint arXiv:2405.15793*, 2024.

[10] Xingyao Wang, Boxuan Gao, Yufan Tian, Hao Chen, Yueqi Zhao, et al. OpenHands: An open platform for AI software developers as generalist agents. *arXiv preprint arXiv:2407.16741*, 2024.

[11] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

[12] Zijian Chen, Xueguang Ma, Shengyao Zhuang, Ping Nie, Kai Zou, Andrew Liu, Joshua Green, Kshama Patel, Ruoxi Meng, Mingyi Su, Sahel Sharifymoghaddam, Yanxi Li, Haoran Hong, Xinyu Shi, Xuye Liu, Nandan Thakur, Crystina Zhang, Luyu Gao, Wenhu Chen, and Jimmy Lin. Browsecomp-plus: A more fair and transparent evaluation benchmark of deep-research agent. *arXiv preprint arXiv:2508.06600*, 2025.

[13] Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

[14] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL https://qwen.ai/blog?id=qwen3.5.

[15] Aohan Zeng, Xin Lv, Qinkai Zheng, Zhenyu Hou, Bin Chen, Chengxing Xie, Cunxiang Wang, Da Yin, Hao Zeng, Jiajie Zhang, et al. Glm-4.5: Agentic, reasoning, and coding (arc) foundation models. *arXiv preprint arXiv:2508.06471*, 2025.

[16] Gemma Team. Gemma 4 technical report, 2026. URL https://arxiv.org/abs/2607.02770.

[17] MiMo Team. Mimo-v2.5. https://huggingface.co/collections/XiaomiMiMo/mimo-v25, 2026.

[18] Kimi Team, Tongtong Bai, Yifan Bai, Yiping Bao, S. H. Cai, Yuan Cao, Y. Charles, H. S. Che, Cheng Chen, Guanduo Chen, Huarong Chen, Jia Chen, Jiahao Chen, Jianlong Chen, Jun Chen, Kefan Chen, Liang Chen, Ruijue Chen, Xinhao Chen, Yanru Chen, Yanxu Chen, Yicun Chen, Yimin Chen, Yingjiang Chen, Yuankun Chen, Yujie Chen, Yutian Chen, Zhirong Chen, Ziwei Chen, Dazhi Cheng, Minghan Chu, Jialei Cui, Jiaqi Deng, Muxi Diao, Hao Ding, Mengfan Dong, Mengnan Dong, Yuxin Dong, Yuhao Dong, Angang Du, Chenzhuang Du, Dikang Du, Lingxiao Du, Yulun Du, Yu Fan, Shengjun Fang, Qiulin Feng, Yichen Feng, Garimugai Fu, Kelin Fu, Hongcheng Gao, Tong Gao, Yuyao Ge, Shangyi Geng, Chengyang Gong, Xiaochen Gong, Zhuoma Gongque, Qizheng Gu, Xinran Gu, Yicheng Gu, Longyu Guan, Yuanying Guo, Xiaoru Hao, Weiran He, Wenyang He, Yunjia He, Chao Hong, Hao Hu, Jiaxi Hu, Yangyang Hu, Zhenxing Hu, Ke Huang, Ruiyuan Huang, Weixiao Huang, Zhiqi Huang, Tao Jiang, Zhejun Jiang, Xinyi Jin, Yu Jing, Guokun Lai, Aidi Li, C. Li, Cheng Li, Fang Li, Guanghe Li, Guanyu Li, Haitao Li, Haoyang Li, Jia Li, Jingwei Li, Junxiong Li, Lincan Li, Mo Li, Weihong Li, Wentao Li, Xinhang Li, Xinhao Li, Yang Li, Yanhao Li, Yiwei Li, Yuxiao Li, Zhaowei Li, Zheming Li, Weilong Liao, Jiawei Lin, Xiaohan Lin, Zhishan Lin, Zichao Lin, Cheng Liu, Chenyu Liu, Hongzhang Liu, Liang Liu, Shaowei Liu, Shudong Liu, Shuran Liu, Tianwei Liu, Tianyu Liu, Weizhou Liu, Xiangyan Liu, Yangyang Liu, Yanming Liu, Yibo Liu, Yuanxin Liu, Yue Liu, Zhengying Liu, Zhongnuo Liu, Enzhe Lu, Haoyu Lu, Zhiyuan Lu, Junyu Luo, Tongxu Luo, Yashuo Luo, Long Ma, Yingwei Ma, Shaoguang Mao, Yuan Mei, Xin Men, Fanqing Meng, Zhiyong Meng, Yibo Miao, Minqing Ni, Kun Ouyang, Siyuan Pan, Bo Pang, Yuchao Qian, Ruoyu Qin, Zeyu Qin, Jiezhong Qiu, Bowen Qu, Zeyu Shang, Youbo Shao, Tianxiao Shen, Zhennan Shen, Juanfeng Shi, Lidong Shi, Shengyuan Shi, Feifan Song, Pengwei Song, Tianhui Song, Xiaoxi Song, Hongjin Su, Jianlin Su, Zhaochen Su, Lin Sui, Jinsong Sun, Junyao Sun, Tongyu Sun, Flood Sung, Yunpeng Tai, Chuning Tang, Heyi Tang, Xiaojuan Tang, Zhengyang Tang, Jiawen Tao, Shiyuan Teng, Chaoran Tian, Pengfei Tian, Ao Wang, Bowen Wang, Chensi Wang, Chuang Wang, Congcong Wang, Dingkun Wang, Dinglu Wang, Dongliang Wang, Feng Wang, Hailong Wang, Haiming Wang, Hengzhi Wang, Huaqing Wang, Hui Wang, Jiahao Wang, Jinhong Wang, Jiuzheng Wang, Kaixin Wang, Linian Wang, Qibin Wang, Shengjie Wang, Shuyi Wang, Si Wang, Wei Wang, Xiaochen Wang, Xinyuan Wang, Yao Wang, Yejie Wang, Yipu Wang, Yiqin Wang, Yucheng Wang, Yuzhi Wang, Zhaoji Wang, Zhaowei Wang, Zhengtao Wang, Zhexu Wang, Zihan Wang, Zizhe Wang, Chu Wei, Ming Wei, Chuan Wen, Zichen Wen, Chengjie Wu, Haoning Wu, Junyan Wu, Rucong Wu, Wenhao Wu, Yuefeng Wu, Yuhao Wu, Yuxin Wu, Zijian Wu, Chenjun Xiao, Jin Xie, Xiaotong Xie, Yuchong Xie, Yifei Xin, Bowei Xing, Boyu Xu, Jianfan Xu, Jing Xu, Jinjing Xu, L. H. Xu, Lin Xu, Suting Xu, Weixin Xu, Xinbo Xu, Xinran Xu, Yangchuan Xu, Yichang Xu, Yuemeng Xu, Zelai Xu, Ziyao Xu, Junjie Yan, Yuzi Yan, Guangyao Yang, Hao Yang, Junwei Yang, Kai Yang, Ningyuan Yang, Ruihan Yang, Xiaofei Yang, Xinlong Yang, Ying Yang, Yi Yang, Yi Yang, Zhen Yang, Zhilin Yang, Zonghan Yang, Haotian Yao, Dan Ye, Wenjie Ye, Zhuorui Ye, Bohong Yin, Chengzhen Yu, Longhui Yu, Tao Yu, Tianxiang Yu, Enming Yuan, Mengjie Yuan, Xiaokun Yuan, Yang Yue, Weihao Zeng, Dunyuan Zha, Haobing Zhan, Dehao Zhang, Hao Zhang, Jin Zhang, Puqi Zhang, Qiao Zhang, Rui Zhang, Xiaobin Zhang, Y. Zhang, Yadong Zhang, Yangkun Zhang, Yichi Zhang, Yizhi Zhang, Yongting Zhang, Yu Zhang, Yushun Zhang, Yutao Zhang, Yutong Zhang, Zheng Zhang, Chenguang Zhao, Feifan Zhao, Jinxiang Zhao, Shuai Zhao, Xiangyu Zhao, Yikai Zhao, Zijia Zhao, Huabin Zheng, Ruihan Zheng, Shaojie Zheng, Tengyang Zheng, Junfeng Zhong, Longguang Zhong, Weiming Zhong, M. Zhou, Runjie Zhou, Xinyu Zhou, Zaida Zhou, Jinguo Zhu, Liya Zhu, Xinhao Zhu, Yuxuan Zhu, Zhen Zhu, Jingze Zhuang, Weiyu Zhuang, Ying Zou, and Xinxing Zu. Kimi k2.5: Visual agentic intelligence, 2026. URL https://arxiv.org/abs/2602.02276.

[19] GLM-5-Team, Aohan Zeng, Xin Lv, Zhenyu Hou, Zhengxiao Du, Qinkai Zheng, Bin Chen, Da Yin, Chendi Ge, Chenghua Huang, Chengxing Xie, Chenzheng Zhu, Congfeng Yin, Cunxiang Wang, Gengzheng Pan, Hao Zeng, Haoke Zhang, Haoran Wang, Huilong Chen, Jiajie Zhang, Jian Jiao, Jiaqi Guo, Jingsen Wang, Jingzhao Du, Jinzhu Wu, Kedong Wang, Lei Li, Lin Fan, Lucen Zhong, Mingdao Liu, Mingming Zhao, Pengfan Du, Qian Dong, Rui Lu, Shuang-Li, Shulin Cao, Song Liu, Ting Jiang, Xiaodong Chen, Xiaohan Zhang, Xuancheng Huang, Xuezhen Dong, Yabo Xu, Yao Wei, Yifan An, Yilin Niu, Yitong Zhu, Yuanhao Wen, Yukuo Cen, Yushi Bai, Zhongpei Qiao, Zihan Wang, Zikang Wang, Zilin Zhu, Ziqiang Liu, Zixuan Li, Bojie Wang, Bosi Wen, Can Huang, Changpeng Cai, Chao Yu, Chen Li, Chengwei Hu, Chenhui Zhang, Dan Zhang, Daoyan Lin, Dayong Yang, Di Wang, Ding Ai, Erle Zhu, Fangzhou Yi, Feiyu Chen, Guohong Wen, Hailong Sun, Haisha Zhao, Haiyi Hu, Hanchen Zhang,

Hanrui Liu, Hanyu Zhang, Hao Peng, Hao Tai, Haobo Zhang, He Liu, Hongwei Wang, Hongxi Yan, Hongyu Ge, Huan Liu, Huanpeng Chu, Jia'ni Zhao, Jiachen Wang, Jiajing Zhao, Jiamin Ren, Jiapeng Wang, Jiaxin Zhang, Jiayi Gui, Jiayue Zhao, Jijie Li, Jing An, Jing Li, Jingwei Yuan, Jinhua Du, Jinxin Liu, Junkai Zhi, Junwen Duan, Kaiyue Zhou, Kangjian Wei, Ke Wang, Keyun Luo, Laiqiang Zhang, Leigang Sha, Liang Xu, Lindong Wu, Lintao Ding, Lu Chen, Minghao Li, Nianyi Lin, Pan Ta, Qiang Zou, Rongjun Song, Ruiqi Yang, Shangqing Tu, Shangtong Yang, Shaoxiang Wu, Shengyan Zhang, Shijie Li, Shuang Li, Shuyi Fan, Wei Qin, Wei Tian, Weining Zhang, Wenbo Yu, Wenjie Liang, Xiang Kuang, Xiangmeng Cheng, Xiangyang Li, Xiaoquan Yan, Xiaowei Hu, Xiaoying Ling, Xing Fan, Xingye Xia, Xinyuan Zhang, Xinze Zhang, Xirui Pan, Xu Zou, Xunkai Zhang, Yadi Liu, Yandong Wu, Yanfu Li, Yidong Wang, Yifan Zhu, Yijun Tan, Yilin Zhou, Yiming Pan, Ying Zhang, Yinpei Su, Yipeng Geng, Yong Yan, Yonglin Tan, Yuean Bi, Yuhan Shen, Yuhao Yang, Yujiang Li, Yunan Liu, Yunqing Wang, Yuntao Li, Yurong Wu, Yutao Zhang, Yuxi Duan, Yuxuan Zhang, Zezhen Liu, Zhengtao Jiang, Zhenhe Yan, Zheyu Zhang, Zhixiang Wei, Zhuo Chen, Zhuoer Feng, Zijun Yao, Ziwei Chai, Ziyuan Wang, Zuzhou Zhang, Bin Xu, Minlie Huang, Hongning Wang, Juanzi Li, Yuxiao Dong, and Jie Tang. Glm-5: from vibe coding to agentic engineering, 2026. URL https://arxiv.org/abs/2602.15763.

[20] Aili Chen, Aonian Li, Baichuan Zhou, Bangwei Gong, Binyang Jiang, Boji Dan, Changhao Zhang, Changqing Yu, Chao Wang, Cheng Ma, et al. The minimax-m2 series: Mini activations unleashing max real-world intelligence. *arXiv preprint arXiv:2605.26494*, 2026.

A Appendix

A.1 Configuration Parameters

Table 4 lists the key configuration parameters used across our evaluations.

Table 4. Key ASAP configuration parameters.

Parameter	Value
Temperature	0.7
top_p	0.95
Completions calls per step	1
Generated span per step	Thinking + action
max_tokens (per step)	2,048
max_steps	30
Prefix K	3
max_observation_length	100,000 chars
max_observation_in_prefix	3,000 chars
max_prefix_tokens	64,000

A.2 Illustration of the selected format

The following snippet is an example of our Markdown+XML format for the base model evaluation.

```
## Thinking
I need to find the relevant file and understand the bug...

<bash>
find . -name "*.py" | grep models
</bash>

<observation>
./src/models.py
./tests/test_models.py
</observation>

## Thinking
Let me examine the failing test to understand the expected behavior...

<bash>
cat tests/test_models.py
</bash>
```

A.3 Full evaluation results

In this section, we show all the Pass@k results in Table 5.

Table 5. Main results (%) across agentic benchmarks. #Inter is the number of instances resolved in all 8 runs. For SWE-bench Verified, we evaluate each model with both SWE-agent and OpenHands; Terminal-Bench 2.1 uses the Terminus-2 scaffold; BrowseComp-Plus uses its native scaffold. **Bold**: best in panel.

Model	Pass@1	Pass@2	Pass@4	Pass@8	#Inter
SWE-bench Verified *Scaffold: SWE-agent*					
Qwen3-30B-A3B-Base	7.38	12.32	18.37	24.20	0
Qwen3.5-9B-Base	21.85	30.61	38.81	46.60	16
Qwen3.5-35B-A3B-Base	34.23	42.82	50.15	56.20	67
Gemma4-26B-A4B-Base	7.12	11.84	17.60	23.00	0
GLM-4.5-Base	15.97	24.46	33.47	41.40	4
Mimo-V2.5-Base	**37.70**	**46.51**	**53.80**	**59.40**	**82**
SWE-bench Verified *Scaffold: OpenHands*					
Qwen3-30B-A3B-Base	4.00	6.64	10.11	14.40	0
Qwen3.5-9B-Base	32.95	42.04	49.32	55.40	51
Qwen3.5-35B-A3B-Base	**39.23**	**48.52**	**55.79**	**61.80**	69
Gemma4-26B-A4B-Base	4.67	8.05	12.70	18.20	0
GLM-4.5-Base	19.73	28.58	37.07	44.80	11
Mimo-V2.5-Base	37.52	46.99	54.53	60.80	**70**
Terminal-Bench 2.1 *Scaffold: Terminus-2*					
Qwen3-30B-A3B-Base	13.76	16.65	18.68	20.22	5
Qwen3.5-9B-Base	19.80	23.52	27.35	30.34	10
Qwen3.5-35B-A3B-Base	**26.69**	**33.19**	**39.78**	**44.94**	**11**
Gemma4-26B-A4B-Base	12.50	15.33	17.40	20.22	4
GLM-4.5-Base	16.01	20.22	25.75	33.71	6
Mimo-V2.5-Base	25.84	31.58	36.34	41.57	10
BrowseComp-Plus *Scaffold: Native*					
Qwen3-30B-A3B-Base	17.36	22.89	27.79	32.89	29
Qwen3.5-9B-Base	20.53	26.26	31.00	34.94	54
Qwen3.5-35B-A3B-Base	26.31	32.22	37.60	43.01	97
Gemma4-26B-A4B-Base	14.64	20.19	25.50	30.72	20
GLM-4.5-Base	23.34	28.68	33.11	37.47	80
Mimo-V2.5-Base	**29.44**	**35.32**	**40.69**	**46.51**	**125**

A.4 Ablation Studies

This section provides the details and ablation studies of the major design of ASAP : the format choice and the bootstrap prefix.

A.4.1 Input Format

The input representation can substantially affect the ability of pretrained base models to interact with an agent scaffold. Unlike instruction-tuned models, base models are not explicitly trained to follow a particular tool-use protocol, and therefore the choice of serialization format may introduce a non-trivial source of variance. We study four commonly used representations: pure JSON, pure Markdown, pure XML, and our default Markdown+XML format. The latter uses Markdown for free-form reasoning and natural-language context, while reserving XML tags for structured actions and observations.

We conduct the ablation on a randomly sampled 100-instance subset of SWE-bench Verified using the SWE-agent scaffold. For each format, we sample three independent trajectories per instance and compute

Pass@1 and Pass@3 using the unbiased Pass@k estimator.

Table 6. Input-format ablation on a randomly sampled 100-instance subset of SWE-bench Verified using the SWE-agent scaffold. We sample three trajectories per instance and report Pass@1 and Pass@3 using the unbiased Pass@k estimator.

Model	Format	Pass@1	Pass@3
Qwen3.5-35B-A3B	JSON	4.33	13.00
	Markdown	21.67	35.00
	XML	27.00	38.00
	Markdown+XML	**33.00**	**50.00**
MiMo-V2.5-Base	JSON	14.67	30.00
	Markdown	30.67	47.00
	XML	38.00	52.00
	Markdown+XML	**42.00**	**59.00**

As shown in Table 6, the representation format has a substantial effect on agent performance. Pure JSON performs particularly poorly for both models, suggesting that requiring the entire interaction to conform to a rigid structured representation is difficult for pretrained base models. Markdown performs considerably better, but remains consistently below formats that explicitly delimit structured interaction fields with XML tags.

Pure XML further improves performance, indicating that explicit structural boundaries are useful for separating actions, observations, and other components of the agent trajectory. However, combining Markdown with XML performs best for both models. Compared with pure XML, Markdown+XML improves Pass@1 from 27.00% to 33.00% for Qwen3.5-35B-A3B and from 38.00% to 42.00% for MiMo-V2.5-Base. The same trend is also reflected in Pass@3.

These results suggest that Markdown and XML provide complementary benefits. Markdown offers a natural continuation space for free-form reasoning and textual context, while XML provides explicit local structure for tool actions and observations. We therefore adopt Markdown+XML as the default representation throughout our experiments.

A.4.2 Prefix Length and Source Robustness

Prefix construction. For each instance, we construct the reference prefix from trajectories generated by instruction-tuned models. Our trajectory pool contains trajectories from Kimi K2.5 [18], GLM-5.2 [19], and MiniMax M2.7 [20]. We randomly select one trajectory and use its initial interaction steps as the prefix. The prefix is intended to provide *protocol-level* few-shot guidance and demonstrates the expected turn structure, tool-use convention, and Markdown+XML interaction format. For the default $K = 3$ setting, the prefix remains in the initial interaction phase and does not contain a golden patch modification or reveal the ground-truth directly. The evaluated base model therefore performs the subsequent task-solving trajectory autonomously from step $K + 1$ onward.

Prefix-length ablation. We evaluate $K \in \{0, 1, 3\}$ on 100 randomly sampled instances from SWE-bench Verified, using three independent runs to report Pass@1 and Pass@3. As shown in Table 7, increasing the number of prefix steps consistently improves performance for both models. This is expected for pretrained base models, as additional few-shot context provides a clearer demonstration of the interaction protocol and makes the intended agent workflow easier to continue.

The same trend is observed for both Qwen3.5-35B-A3B and MiMo-V2.5-Base across different values of K. Moreover, their relative ordering remains unchanged: MiMo-V2.5-Base consistently outperforms Qwen3.5-35B-A3B under all prefix lengths and for both Pass@1 and Pass@3. This suggests that the prefix mainly improves capability elicitation while preserving the relative comparison between base models.

We use $K = 3$ as the default setting. In preliminary experiments with longer prefixes ($K > 3$), we find that some prefixes begin to expose solution-bearing information, including parts of the gold patch. We therefore limit the prefix to three interaction steps to provide sufficient few-shot guidance while avoiding direct solution leakage.

Table 7. Prefix-length ablation on 100 randomly sampled SWE-bench Verified instances. Results (%) are computed from three independent runs. Increasing the prefix length consistently improves performance while preserving the relative model ordering.

Model	Prefix length K	Pass@1	Pass@3
Qwen3.5-35B-A3B	0	24.00	41.00
	1	27.67	43.00
	3 (selected)	33.00	50.00
MiMo-V2.5-Base	0	28.00	46.00
	1	32.00	53.00
	3 (selected)	42.00	59.00

Robustness to prefix source. We next examine whether the evaluation conclusions depend on the particular set of reference prefixes used. We compare two independently constructed prefix sets while fixing the prefix length to $K = 3$. Prefix A is constructed using our default strategy, whereas Prefix B is newly generated by GLM-5.3-Flash for the same set of evaluation instances.

We evaluate Qwen3.5-35B-A3B and MiMo-V2.5-Base on the same 100-instance subset of SWE-bench Verified used in the prefix-length ablation. For each model–prefix combination, we sample three independent trajectories per instance under the same SWE-agent scaffold, decoding configuration, and interaction budget. Table 8 reports Pass@1, Pass@3, and the score difference between Prefix B and Prefix A.

Changing the reference prefixes produces only modest changes in absolute performance. Across the two models and two metrics, the maximum absolute difference between Prefix A and Prefix B is 3 percentage points. Moreover, the relative model comparison remains unchanged: MiMo-V2.5-Base outperforms Qwen3.5-35B-A3B under both prefix sets.

These results indicate that the model separation observed by ASAP is not specific to the particular reference prefixes used in our default configuration. Instead, independently generated prefixes that establish the same interaction protocol yield similar absolute performance and preserve the same relative model ordering.

A.5 Consistency with Lightweight SFT

To examine whether the model ordering exposed by direct base-model evaluation remains informative after instruction following is introduced, we apply the same lightweight 70K-example SFT recipe to four base models. The goal of this experiment is not to maximize downstream performance, but to provide a controlled comparison under a shared post-training recipe. We evaluate each post-SFT model with SWE-agent for three runs on all 500 SWE-bench Verified instances.

The model ordering is identical across base Pass@1, base Pass@8, post-SFT Pass@1, and post-SFT Pass@3, yielding Spearman $\rho = 1.0$ for all four pairwise comparisons. The corresponding Pearson correlations are

Table 8. Robustness to independently constructed reference prefixes on the same 100-instance subset of SWE-bench Verified. Prefix A denotes our default prefix strategy, while Prefix B is independently generated by GLM-5.3-Flash. Δ denotes Prefix B minus Prefix A.

Model	Metric	Prefix A	Prefix B	Δ
Qwen3.5-35B-A3B	Pass@1	33.00	34.00	+1.00
Qwen3.5-35B-A3B	Pass@3	50.00	48.00	−2.00
MiMo-V2.5-Base	Pass@1	42.00	44.33	+2.33
MiMo-V2.5-Base	Pass@3	59.00	56.00	−3.00

Table 9. Direct base-model evaluation with ASAP and end-to-end SWE-agent evaluation after the same 70K-example lightweight SFT. Base results use eight runs and SFT results use three runs. All values are percentages. **Bold**: best in column.

	Base / ASAP		70K SFT	
Model	Pass@1	Pass@8	Pass@1	Pass@3
Qwen3-30B-A3B	7.38	24.20	19.60	31.80
Qwen3.5-9B	21.85	46.60	24.60	37.60
Qwen3.5-35B-A3B	**34.23**	**56.20**	**37.47**	**53.20**
Gemma4-26B-A4B	7.12	23.00	18.70	31.60

also positive. Given the small number of models and the potentially large effect of post-training, we treat this result only as an aggregate consistency check rather than evidence that ASAP fully predicts post-trained performance. The primary use of ASAP remains direct evaluation and behavioral analysis of pretrained base models, as studied in Section 5.1 and Section 5.2.

A.6 External Network Access Audit

In this section, we investigate the evaluation hacking by external network access of base models with ASAP . If an execution environment permits outbound network requests, a model could in principle search for the original issue, pull request, or a known patch instead of deriving the fix from the local repository. We therefore audit this observable shortcut by scanning the executable action fields of 4,000 trajectories per model (8 runs × 500 instances). We count trajectories containing any attempted external network access and separately flag accesses that appear to target a gold solution or external documentation. Table 10 reports the results for the six models with observed network activity.

Table 10. Trajectory-level audit of attempted external network access. "Gold" denotes accesses that appear to seek a reference solution or patch, while "Docs" denotes accesses to external documentation.

Model	Trajectories	Network	Network Rate	Gold	Gold Rate	Docs
Qwen3-30B-A3B-Base	4,000	86	2.15%	9	0.23%	23
Gemma4-26B-A4B-Base	4,000	25	0.63%	5	0.13%	8
Mimo-V2.5-Base	4,000	9	0.23%	1	0.03%	0
Qwen3.5-35B-A3B-Base	4,000	9	0.23%	1	0.03%	2
Qwen3.5-9B-Base	4,000	8	0.20%	1	0.03%	2
GLM-4.5-Base	4,000	7	0.18%	0	0.00%	6
Total	24,000	144	0.60%	17	0.07%	41

Across the 24,000 audited trajectories, only 0.6% attempt any network access and 0.07% are flagged as

Table 11. Performance and failure patterns on SWEBench Verified using ASAP . "Non-empty success rate" denotes the fraction of submitted non-empty patches that pass the official grader.

Metric	Qwen3.5-35B-A3B-Base	MiMo-V2.5-Base
Pass@1	99/300 (33.0%)	**126/300 (42.0%)**
Pass@3	50/100 (50.0%)	**59/100 (59.0%)**
Empty patch	**73/300 (24.3%)**	101/300 (33.7%)
Non-empty unsuccessful/error	128/300 (42.7%)	**73/300 (24.3%)**
Non-empty success rate	43.6%	**63.3%**
Repetition-based termination	**9/300 (3.0%)**	24/300 (8.0%)
Non-empty failures without product-code edits	19	**5**
Non-empty failures with regression	44	**20**
Median failed-patch size (changed lines)	16	**11**

potentially seeking a gold solution. Even for Qwen3-30B-A3B-Base, which has the highest rates, the corresponding values are 2.15% and 0.23%. Under the deliberately pessimistic assumption that every gold-flagged trajectory succeeds solely because of external lookup, its direct contribution to Pass@1 is therefore at most 0.23 percentage points for any audited model. Documentation access is also rare (41 trajectories in total). Overall, network activity is rare among base models, indicating that external answer lookup is unlikely to pose a meaningful risk to the validity of our evaluation.

A.7 Failure Pattern Analysis

We compare Qwen3.5-35B-A3B-Base and MiMo-V2.5-Base on the same 100 SWE-bench Verified instances, with three independent runs per instance. Both models use the same SWE-agent scaffold, decoding configuration, and interaction budget. Table 11 summarizes their main performance and failure characteristics.

Distinct execution profiles. The two models fail in notably different ways. Qwen produces non-empty patches more often, but these patches are substantially less precise: only 43.6% of its non-empty submissions are correct, compared with 63.3% for MiMo. Its failed patches are also larger and more frequently introduce regressions or fail to modify the relevant product code. Manual inspection shows a recurring tendency to edit at the wrong abstraction level, implement only part of the required behavior, or make unnecessarily broad changes.

MiMo exhibits the opposite pattern. Once it commits to a source-code modification, its patches are typically smaller and more likely to satisfy both the target and regression tests. However, it is also more likely to terminate without producing a patch, often after prolonged or repetitive exploration. The main contrast is therefore not simply overall capability, but execution style: Qwen acts more readily but with lower precision, whereas MiMo is more selective but sometimes fails to transition from exploration to implementation.

Implications for targeted data synthesis. These failure profiles suggest that pretraining data could be tailored to the dominant weakness of each model. For Qwen, the main target is *repair precision*: training examples should favor minimal, correct patches over alternatives that modify the wrong layer, introduce regressions, or make unnecessarily broad changes. For MiMo, the main target is *progress under uncertainty*: trajectories should encourage the model to move from bounded exploration to a concrete hypothesis, focused edit, validation, and timely submission, while discouraging repetitive searches with little new information.

More generally, the trajectory-level diagnostics exposed by ASAP make it possible to construct model-

specific training signals rather than relying on a single generic trajectory mixture. For example, candidate continuations from the same intermediate repository state can be evaluated with the task grader and used to construct preference pairs that distinguish precise, successful progress from empty, regressive, or unnecessarily broad continuations.