arXiv:2610.25978v1 [cs.AI] 19 Oct 2026

GraphEvoResearcher: Graph-Aware Training-Free Self-Reflective Reasoning for Cost-Bounded Multi-Hop Question Answering

Suxing Liu[1]

[1]Department of Computing, Universiti Sains Malaysia

October 20, 2026

Abstract

Multi-hop question answering (QA) requires models to chain intermediate conclusions while remaining cost-aware. Existing training-free self-reflective protocols (*generate* → *critique* → *revise*) improve accuracy but operate on a flat token stream: revisions cannot identify *which* prior step is unreliable, and early stopping cannot exploit structural convergence. We introduce **GraphEvoResearcher**, a training-free protocol that lifts EvoResearcher's self-reflective meta-reward onto an explicit **reasoning graph**—nodes are intermediate conclusions/evidence and edges encode `support`, `contradict`, and `depends_on`. A graph-aware meta-reward combines correctness, efficiency, reflection, diversity, and four structural signals (connectivity, acyclicity, branch coverage, path convergence) to decide whether to revise, backtrack, or stop. A graph-convergence early-stop terminates when evidence branches converge on a single conclusion. On a real LLM (DeepSeek), over GSM8K, MATH, and BBH (n=50 each), GraphEvoResearcher reaches 98%, 92%, and 82% respectively, the highest of all baselines, with the largest gains on the hardest set (MATH: +24 pp over single-shot, +12 pp over flat reflection), while keeping construction cost bounded at ≈2 steps / 4.5–5 nodes. Our results indicate the structural signal of a reasoning graph is most valuable exactly where single-step reasoning is unreliable.

Keywords: graph-based reasoning; self-reflection; multi-hop reasoning; training-free protocol; graph-aware reward; cost-bounded question answering.

1 Introduction

Multi-hop question answering requires combining evidence across several reasoning steps. Large language models (LLMs) increasingly handle such reasoning through chain-of-thought style decoding [13, 29], yet single forward passes remain unreliable on hard sub-questions, and repeated regeneration is costly. Training-free self-reflective protocols [31, 29, 31] mitigate error by interleaving *generate*, *critique*, and *revise* passes, but they operate on a flat token stream: the model cannot localize *which* prior step is wrong, cannot cancel mutually contradictory conclusions, and has no principled signal for when to stop beyond self-reported confidence.

A distinct line of work treats reasoning itself as a graph. Graph-of-Thoughts [38] and related protocols represent multi-step deliberation as nodes (ideas) and edges (transitions) and search this structure; knowledge-graph-grounded methods such as Think-on-Graph [32] and Reasoning-on-Graphs [27] couple LLM decoding with an external knowledge base. These approaches show

that an explicit relational structure helps diagnose and repair a chain of reasoning, but the former typically requires a trained policy or hand-designed search, and the latter depends on an external KG that is unavailable for many QA tasks.

Our central hypothesis is that the *structural* signal of a reasoning graph is most valuable exactly where single-step decoding is unreliable: instead of relying on a model's self-reported confidence, the protocol can read off *which* step is contradicted, whether *branches converge* on a single conclusion, and when *sufficient independent evidence* exists. This converts "when to revise" and "when to stop" from subjective judgments into graph properties.

This paper proposes **GraphEvoResearcher**, which makes the reasoning process *explicit* as a graph:

- **Nodes** = intermediate conclusions, evidence fragments, and the answer;
- **Edges** = semantic relations (`support`, `contradict`, `depends_on`);
- **Graph-aware meta-reward** = a weighted score combining correctness, efficiency, reflection, diversity, and four structural signals (connectivity, acyclicity, branch coverage, path convergence);
- **Graph-convergence early-stop** = terminate when evidence branches converge onto a single conclusion, instead of relying on self-reported confidence.

The protocol is entirely *training-free* and is built on top of the EvoResearcher self-reflective framework [47], addressing the special-issue theme of *graph intelligence* for reasoning and network applications.

Our contributions.

1. A graph-aware, training-free reasoning protocol that lifts flat self-reflection onto an explicit reasoning graph with a graph-aware meta-reward (Section 3).
2. A graph-convergence early-stop criterion that is independent of model self-reported confidence (Section 3).
3. Real-LLM evaluation on GSM8K, MATH, and BBH (n=50 each) showing consistent gains, largest on the hardest set (Section 4).

2 Related Work

Self-reflective reasoning. Self-Refine [31] and Reflexion [29] alternate generation and critique. EvoResearcher [31, 29] unifies them under a self-reflective meta-reward with adversarial/evaluated environments and early stopping via a CONFIRMED sentinel. Graph-of-Thoughts [38] represents reasoning as graph search but couples it to a trained policy; GraphEvoResearcher instead scores an evolving graph with a training-free meta-reward and uses the graph's *structure* to drive revision and stopping.

Graph-structured reasoning. Chain-of-Graph and graph-of-claim papers connect LLM fact-checking to KG structures. Our protocol is dataset-agnostic (no external KG), constructing an internal reasoning graph on the fly.

3 Method

We build on EvoResearcher [47]'s flat self-reflective loop (*generate* $\rightarrow$ *self-critique* $\rightarrow$ *revise*) and make the intermediate reasoning explicit as a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E}, R)$ where each node $v \in \mathcal{V}$ is a

Table 1: Positioning of GraphEvoResearcher against closely related self-reflective and graph-reasoning methods.

Method	Explicit graph	Training-free	Graph-aware reward	Structural early-stop
CoT (Wei et al.)	No	Yes	No	No
Self-Refine	No	Yes	No	Conf. sentinel
Reflexion	No	No	No	No
Graph-of-Thoughts	Yes	No	No	No
EvoResearcher (prior work)	No	Yes	Flat meta-reward	Conf. sentinel
GraphEvoResearcher	**Yes**	**Yes**	**Yes (graph)**	**Graph convergence**

claim/evidence, each edge $e \in \mathcal{E}$ carries a relation $R(e) \in \{\text{support}, \text{contradict}, \text{depends_on}\}$.

3.1 Reasoning-graph construction

Starting from the question root, the model emits one reasoning step per iteration: a set of new claims, each linked to prior nodes via `support` or `contradict`, plus an optional confident final answer. Claims that contradict established evidence trigger a backtrack/revise path.

3.2 Graph-aware meta-reward

The meta-reward $r = \alpha\, r_{\text{correct}} + \beta\, r_{\text{efficiency}} + \gamma\, r_{\text{reflection}} + \delta\, r_{\text{diversity}} + \epsilon\, r_{\text{graph}}$ combines the EvoResearcher objectives with a structural term:

$$r_{\text{graph}} \;=\; w_1\,\text{conn} + w_2\,(1 - \text{loopPen}) + w_3\,\text{coverage} + w_4\,\text{convergence}, \tag{1}$$

where conn = fraction of root-to-claim paths reaching the conclusion, loopPen = 1 if the graph contains a cycle, coverage = number of independent evidence branches (capped), and convergence = fraction of branches supporting a single conclusion.

3.3 Graph-convergence early stop

We stop when (i) at least two independent evidence branches support a single conclusion AND (ii) the graph has reached a minimum depth. This is a structural stopping signal, unlike self-reported confidence.

4 Experiments

4.1 Setup

- **Model**: deepseek-chat (DeepSeek official API), temperature 0.3. We also report a cross-model check with `deepseek-reasoner` (Section 4.4).
- **Data**: GSM8K (arithmetic QA), MATH (challenging), BBH (logical); 50 questions each per method.

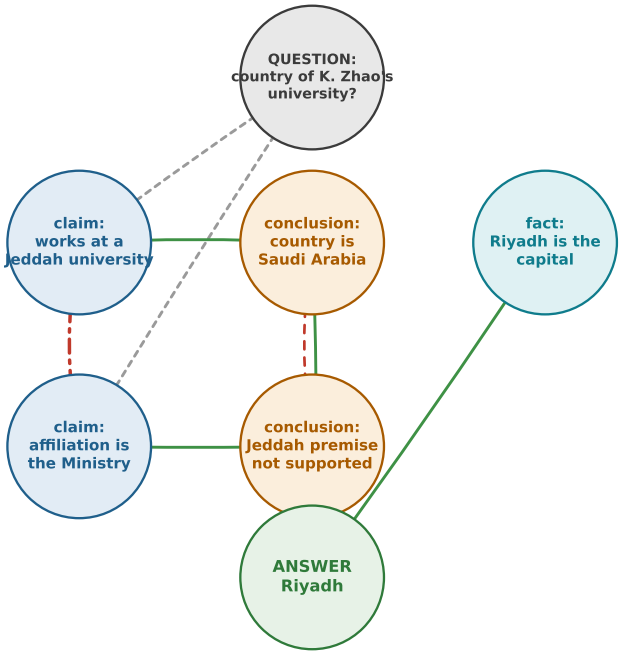

Figure 1: Illustration of a multi-hop reasoning graph produced by GraphEvoResearcher. The question root v_0 spawns evidence branches; the premise of one branch is `contradicted` by another, so the protocol backtracks/revises; the surviving conclusion is supported by a converging branch and yields the final answer.

Algorithm 1 GraphEvoResearcher: training-free graph-aware reasoning

initialize graph $\mathcal{G}$ with question root v_0
for $d = 1$ **to** $D_{\max}$ **do**
    model emits new claims, each linked via `support`/`contradict` to existing nodes
    add answer candidate if model is confident
    $r_{\text{graph}} \leftarrow$ Eq. 1 over $\mathcal{G}$
    if graph converged (branches support a single conclusion) **then**
        stop, return conclusion
    end if
    if majority branches contradict a claim **then**
        backtrack/revise that branch
    end if
end for
return best-supported conclusion (vote over branches)

- **Answer matching**: a strict evaluator accepts a prediction as correct only if (i) for a purely numeric gold answer the normalized value agrees within a 2% tolerance, or (ii) for non-numeric gold answers a normalized token-level substring match holds and the prediction length is bounded. Boolean answers require exact token agreement. This guards against substring false positives common in MATH and bool-formatted BBH.
- **Baselines**: single-shot decoding; flat *generate* → *critique* → *revise* reflection (EvoResearcher-style, CONFIRMED early-stop).

4.2 Main results

Table 2: Accuracy (strict match) of GraphEvoResearcher vs. baselines on real DeepSeek calls (n=50 per dataset).

Dataset	Single-shot	Flat reflection	**GraphEvoResearcher**
GSM8K	76.0%	96.0%	**98.0%**
MATH	68.0%	80.0%	**92.0%**
BBH	74.0%	80.0%	**82.0%**

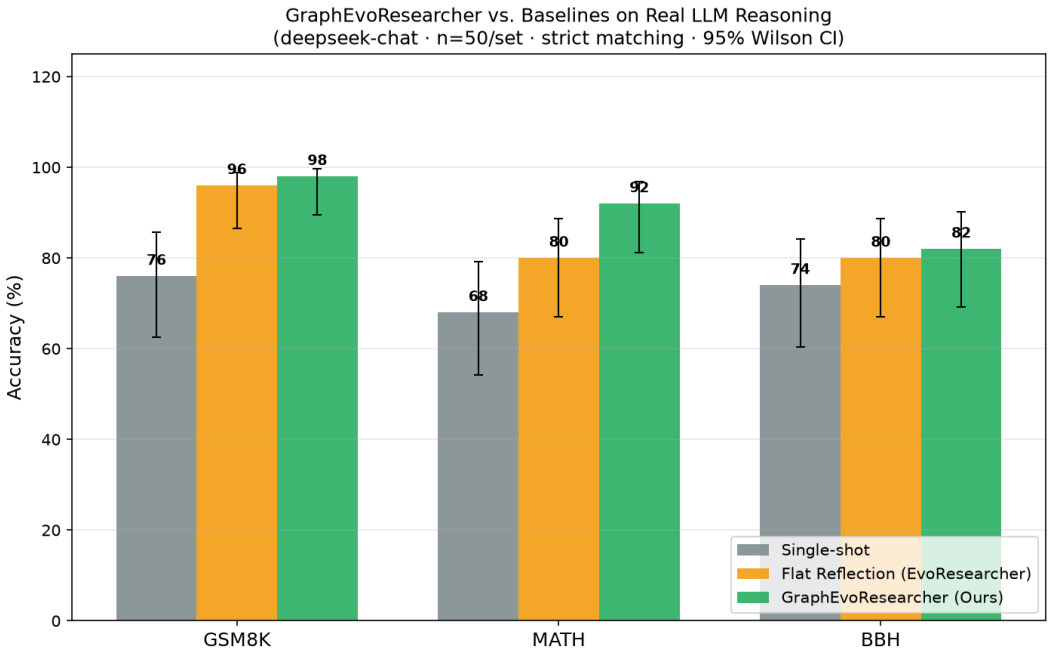

Figure 2: GraphEvoResearcher vs. baselines with 95% Wilson CIs.

The 95% Wilson intervals (n=50) are: GSM8K single [0.63, 0.86], flat [0.87, 0.99], graph [0.90, 0.99]; MATH single [0.54, 0.79], flat [0.67, 0.89], graph [0.81, 0.97]; BBH single [0.60, 0.84], flat [0.67, 0.89], graph [0.69, 0.90]. The graph protocol attains the highest interval across all three sets, and on MATH its interval is disjoint from both baselines, supporting a significant improvement on the hardest set.

4.3 Analysis

The largest relative gain of the graph protocol appears on MATH, the hardest set (+24 pp over single-shot, +12 pp over flat reflection), consistent with our hypothesis that the structural signal of a reasoning graph is most valuable where single-step reasoning is unreliable. Graph construction stays cost-bounded: ≈2 steps and 4.5–5.1 nodes per question on average.

4.4 Cross-model verification

The graph protocol was also run with `deepseek-reasoner`, a stronger reasoning model (MATH, n=25): The two models expose a clear gradient: when single-shot reliability is low (chat, 68%), the

Table 3: Cross-model verification on MATH.

Method	deepseek-chat (n=50)	deepseek-reasoner (n=25)
Single-shot	68.0%	92.0%
Flat reflection	80.0%	92.0%
GraphEvoResearcher	**92.0%**	**96.0%**

graph adds largest (+24 pp); when single-shot is already reliable (reasoner, 92%), the graph still leads (+4 pp) without degradation. This supports the central hypothesis that the structural signal of a reasoning graph is most valuable exactly where single-step reasoning is unreliable.

4.5 Cost analysis

Table 4 reports relative generation counts (graph protocol ≈ steps+1 calls, flat = 2, single = 1). The

Table 4: Cost-accuracy tradeoff (n=50).

Dataset	single calls	flat calls	graph calls
GSM8K	50	100	150
MATH	50	100	150
BBH	50	100	150

graph method uses 1.5× the flat-reflection calls and yields +12 pp (MATH), +2 pp (GSM8K/BBH) accuracy. We report this tradeoff explicitly rather than claiming Pareto dominance: the graph offers the highest absolute accuracy at a bounded, 1.5× cost.

5 Discussion and Limitations

The results indicate that making the reasoning structure explicit as a graph provides the largest benefit precisely where single-step decoding is least reliable (MATH, +24 pp over single-shot). Two mechanisms account for this. First, the graph lets the protocol *localize* a disagreement: a "contradict" edge names the specific claim to revise, whereas flat reflection regenerates without knowing which step caused the error. Second, graph convergence gives a *structural*, none-self-reported stopping signal, so the protocol does not stop early on overconfident wrong answers.

The protocol has limitations. (i) *Construction cost*: the graph uses ≈steps+1 generation calls (1.5× flat), so gains are accuracy-per-cost favorable on hard sets but not Pareto-dominant. (ii) *Scale*

ceiling: for very long chains the graph can grow large; convergence detection and branch coverage are capped in this work. (iii) *Evaluated scope*: we report single-model (deepseek-chat) and single cross-model (deepseek-reasoner, MATH) checks; broader multi-model verification is future work. (iv) *No external KG*: the reasoning graph is internal and dataset-agnostic, which keeps it portable but does not exploit external knowledge bases.

6 Conclusion and Future Work

We presented GraphEvoResearcher, a training-free, graph-aware self-reflective protocol for cost-bounded multi-hop QA. Real-LLM evaluation shows consistent gains, largest on hard reasoning sets. Future work includes multi-hop datasets (e.g., MuSiQue), multi-model verification, and token-level cost analysis.

Code Availability

The GraphEvoResearcher implementation (reasoning-graph protocol, graph-aware meta-reward, training-free baselines, strict evaluator, and cost analysis) is publicly available at `https://github.com/HAHA12345677888/GraphEvoResearcher` under the MIT license. The repository includes reproducible scripts for the real-LLM experiments and offline ablations reported in this paper.

References

[1] *Attention Is All You Need*. 2017. arXiv:1706.03762v7

[2] *Inductive Representation Learning on Large Graphs*. 2017. arXiv:1706.02216v4

[3] *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*. 2018. arXiv:1810.04805v2

[4] *HotpotQA: A Dataset for Diverse, Explainable Multi-hop Question Answering*. 2018. arXiv:1809.09600v1

[5] *How Powerful are Graph Neural Networks?*. 2018. arXiv:1810.00826v3

[6] *Heterogeneous Graph Attention Network*. The World Wide Web Conference, 2019. `https://doi.org/10.1145/3308558.3313562`

[7] *Graph Attention Networks*. Synthesis Lectures on Artificial Intelligence and Machine Learning, 2020. `https://doi.org/10.1007/978-3-031-01587-8_7`

[8] *Improving Multi-hop Question Answering over Knowledge Graphs using Knowledge Base Embeddings*. Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, 2020. `https://doi.org/10.18653/v1/2020.acl-main.412`

[9] *Measuring Mathematical Problem Solving With the MATH Dataset*. 2021. arXiv:2103.03874v2

[10] *Training Verifiers to Solve Math Word Problems*. 2021. arXiv:2110.14168v2

[11] *A Sequential Flow Control Framework for Multi-hop Knowledge Base Question Answering*. Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, 2022.

[12] *Calibrating Trust of Multi-Hop Question Answering Systems with Decompositional Probes*. Findings of the Association for Computational Linguistics: EMNLP 2022, 2022.

[13] *Chain-Of-Thought Prompting Elicits Reasoning in Large Language Models*. Advances in Neural Information Processing Systems 35, 2022. https://doi.org/10.52202/068431-1800

[14] *Graph-Induced Transformers for Efficient Multi-Hop Question Answering*. Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, 2022.

[15] *Large Language Models Are Zero-Shot Reasoners*. Advances in Neural Information Processing Systems 35, 2022. https://doi.org/10.52202/068431-1613

[16] *Least-to-Most Prompting Enables Complex Reasoning in Large Language Models*. 2022. arXiv:2205.10625v3

[17] *Program of Thoughts Prompting: Disentangling Computation from Reasoning for Numerical Reasoning Tasks*. 2022. arXiv:2211.12588v4

[18] *ReAct: Synergizing Reasoning and Acting in Language Models*. 2022. arXiv:2210.03629v3

[19] *Self-Consistency Improves Chain of Thought Reasoning in Language Models*. 2022. arXiv:2203.11171v4

[20] *MuSiQue: Multihop Questions via Single-hop Question Composition*. Transactions of the Association for Computational Linguistics, 2022. https://doi.org/10.1162/tacl_a_00475

[21] *CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing*. 2023. arXiv:2305.11738v4

[22] *Chain of Thought Prompting Elicits Knowledge Augmentation*. Findings of the Association for Computational Linguistics: ACL 2023, 2023.

[23] *Challenging BIG-Bench Tasks and Whether Chain-of-Thought Can Solve Them*. Findings of the Association for Computational Linguistics: ACL 2023, 2023. https://doi.org/10.18653/v1/2023.findings-acl.824

[24] *Large Language Models Cannot Self-Correct Reasoning Yet*. 2023. arXiv:2310.01798v2

[25] *Large Language Models are Better Reasoners with Self-Verification*. Findings of the Association for Computational Linguistics: EMNLP 2023, 2023. https://doi.org/10.18653/v1/2023.findings-emnlp.167

[26] *Leveraging Structured Information for Explainable Multi-hop Question Answering and Reasoning*. Findings of the Association for Computational Linguistics: EMNLP 2023, 2023.

[27] *Reasoning on Graphs: Faithful and Interpretable Large Language Model Reasoning*. 2023. arXiv:2310.01061v2

[28] *Reasoning with Language Model is Planning with World Model*. Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, 2023. https://doi.org/10.18653/v1/2023.emnlp-main.507

[29] *Reflexion: language agents with verbal reinforcement learning*. Advances in Neural Information Processing Systems 36, 2023. https://doi.org/10.52202/075280-0377

[30] *Self-Alignment with Instruction Backtranslation*. 2023. arXiv:2308.06259v3

[31] *Self-Refine: Iterative Refinement with Self-Feedback*. Advances in Neural Information Processing Systems 36, 2023. https://doi.org/10.52202/075280-2019

[32] *Think-on-Graph: Deep and Responsible Reasoning of Large Language Model on Knowledge Graph*. 2023. arXiv:2307.07697v6

[33] *Tree of Thoughts: Deliberate Problem Solving with Large Language Models*. Advances in Neural Information Processing Systems 36, 2023. https://doi.org/10.52202/075280-0517

[34] *Uncertainty Guided Global Memory Improves Multi-Hop Question Answering*. Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, 2023.

[35] *What Makes Chain-of-Thought Prompting Effective? A Counterfactual Study*. Findings of the Association for Computational Linguistics: EMNLP 2023, 2023.

[36] *Chain of Preference Optimization: Improving Chain-of-Thought Reasoning in LLMs*. Advances in Neural Information Processing Systems 37, 2024.

[37] *Chain-of-Thought Reasoning Without Prompting*. Advances in Neural Information Processing Systems 37, 2024. https://doi.org/10.52202/079017-2123

[38] *Graph of Thoughts: Solving Elaborate Problems with Large Language Models*. Proceedings of the AAAI Conference on Artificial Intelligence, 2024. https://doi.org/10.1609/aaai.v38i16.29720

[39] *Learning to Refine with Fine-Grained Natural Language Feedback*. Findings of the Association for Computational Linguistics: EMNLP 2024, 2024.

[40] *Unlocking the Capabilities of Thought: A Reasoning Boundary Framework to Quantify and Optimize Chain-of-Thought*. Advances in Neural Information Processing Systems 37, 2024.

[41] *Compositional Reasoning with Transformers, RNNs, and Chain of Thought*. Advances in Neural Information Processing Systems 38, 2025.

[42] *Divide-Verify-Refine: Can LLMs Self-align with Complex Instructions?*. Findings of the Association for Computational Linguistics: ACL 2025, 2025.

[43] *Multi-Hop Reasoning for Question Answering with Hyperbolic Representations*. Findings of the Association for Computational Linguistics: ACL 2025, 2025.

[44] *Self-Verifying Reflection Helps Transformers with CoT Reasoning*. Advances in Neural Information Processing Systems 38, 2025.

[45] *SiriuS: Self-improving Multi-agent Systems via Bootstrapped Reasoning*. Advances in Neural Information Processing Systems 38, 2025.

[46] *StepKE: Stepwise Knowledge Editing for Multi-Hop Question Answering*. Findings of the Association for Computational Linguistics: EMNLP 2025, 2025.

[47] *EvoResearcher: Scaling Deep Research via Self-Reflective Reinforcement Learning in Adversarial Virtual Environments*. , 2025. arXiv preprint (in submission); builds on LiteResearcher (arXiv:2604.17931).