arXiv:2610.29669v1 [cs.DB] 21 Oct 2026

Conflicted Records: Typed Adversarial Contracts for Auditable LLM Output Quality Enforcement

Kardelen Subasi and Jarosław A. Chudziak

Warsaw University of Technology, Warsaw, Poland
{kardelen.subasi.stud, jaroslaw.chudziak}@pw.edu.pl

Abstract. LLM-generated records increasingly populate production information systems where they are treated as data assets subject to the same quality expectations as records from operational databases. A single such record can simultaneously satisfy and violate a quality rule: a grounded fact adjacent to a hallucination, a useful answer adjacent to a stereotype. Classical data cleaning cannot resolve this because the two judgments are incompatible over the same span; a single LLM judge cannot either, since it shares the training biases that produced the conflict. Existing interventions such as prompt engineering, chain-of-thought, and RLHF reduce surface error rates but provide no audit trail and cannot distinguish genuine remediation from blanket suppression. The question is whether a deterministic check over structured inter-agent records can catch the defects that referential checks alone miss.

We present BiasGuard-D: two adversarial roles drawn from distinct model families exchange typed claim documents, a three-model quorum adjudicates by majority vote, and a combined referential-plus-semantic invariant validates the verdict. The semantic component fires on verdicts referential checks miss; removing the iterative loop drops invariant satisfaction sharply. The deliberation layer transfers from stereotype to hallucination remediation on HaluEval QA with only role-framing changes, raising right-answer recovery by 49 percentage points; the triage gate required a cross-family second auditor to maintain recall on objective hallucination. Every claim, vote, and instruction is stored as a typed record giving operators a provenance trail auditable under emerging AI regulations.

Keywords: LLM Output Quality Enforcement · Typed Adversarial Contracts · Multi-agent Systems · Bias Remediation · Hallucination Detection · Large language models · Artificial intelligence.

1 Introduction

LLM outputs are no longer transient. They are persisted into search indexes, written into knowledge bases, embedded in recommendation feeds, and surfaced through customer-facing information systems where they act as data [17,8]. As deployments scale into regulated domains, operators need more than a quality score; they need to know *why* an output passes or fails a constraint, *which span*

triggered the verdict, and *whether the remediation instructions are internally consistent*.

What makes this hard is that an LLM-generated record can satisfy and violate a quality rule in the same sequence of text. A hiring recommendation may accurately ground its reasoning in a candidate's qualifications while simultaneously making an age generalisation the input does not support. A knowledge-base entry may pair a faithfully lifted fact with one the model invented. In neither case are the two things in separate fields or even separate sentences; they share the same characters. Classical data cleaning, which reads a record and refines it, cannot resolve this, because the two issues require opposite judgments about the same text. A single LLM judge cannot either, since it shares the same training biases that produced the conflict [15,25,26].

The question this paper asks is: *can a deterministic check over structured inter-agent records catch the quality defects that referential checks alone miss, and do those records carry enough information to audit pipeline quality without calling an additional model?* We evaluate this across three dimensions: whether a semantic check catches a recurring class of contradictions that referential checks miss; whether stored records carry enough structure to tell genuine resolution from suppression without calling a further model; and whether the deliberation layer transfers across different quality domains with only role-framing changes.

We address this by treating enforcement as a data-quality problem rather than a single-judge problem: two adversarial roles with opposing biases produce structured claims about specific text spans; a multi-family quorum adjudicates; and a deterministic invariant validates the result before any downstream action. Records with a consistent verdict pass with a full provenance trail; contradictory verdicts are quarantined rather than auto-repaired, because auto-repair would bring back the model judgment the architecture is designed to avoid. The experimental study confirms all three evaluation questions and surfaces a finding we did not anticipate: the iterative deliberation loop matters far more than any individual model choice, and the audit trail is what makes this visible.

This paper makes three contributions. We present a five-stage typed adversarial contract pipeline in which two adversarial roles from distinct model families exchange typed claim documents over bounded peer-to-peer rounds, a three-model quorum adjudicates, and a combined deterministic invariant ($\mathcal{I}_{\text{ref}} \wedge \mathcal{I}_{\text{sem}}$) validates each verdict, making the full lineage from any instruction back to its originating span queryable as data. An ablation shows that the iterative deliberation loop accounts for most of the invariant-satisfaction gain: removing the loop drops satisfaction by nearly 30 percentage points on the BBQ Disability benchmark, while schema-level coordination metrics stored as typed records make the degradation auditable without calling an additional model. Finally, we provide cross-domain portability evidence showing that the deliberation layer, data model, and invariant transfer from subjective stereotype remediation to objective hallucination remediation by changing only role-framing prompts; the triage gate required a cross-family second auditor to maintain recall on objective hallucination, isolating the gate as the one domain-sensitive component.

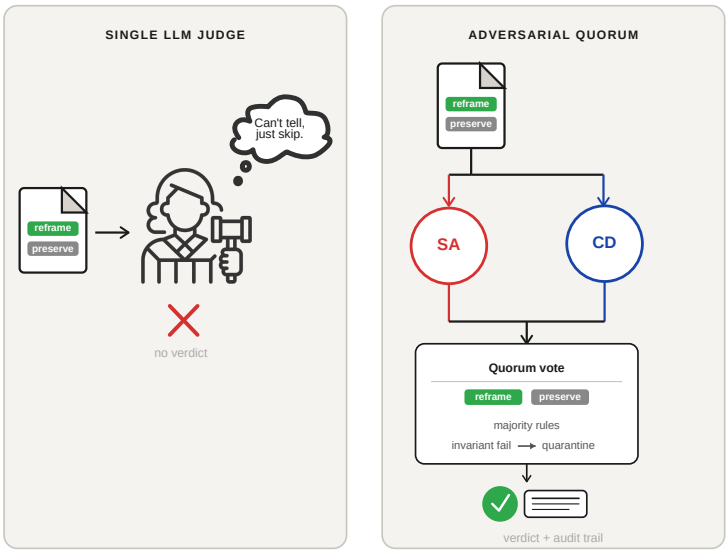

Fig. 1. A conflicted record (input): co-located spans require opposite actions (one contains demographic reasoning that must be reframed or dropped; another is contextually grounded and must be preserved) that a single LLM judge cannot resolve consistently and returns no verdict. BiasGuard-D returns a typed per-span instruction record (output) with a full provenance trail.

Section 2 reviews related work and shows where existing approaches fall short. Section 3 defines the problem and describes the data model, pipeline, and invariant. Section 4 evaluates the system on stereotype and hallucination benchmarks; Sections 5 and 6 discuss the findings, limitations, and what the provenance trail offers in practice. Source code, configuration, and run artefacts are at https://github.com/ksubasi/biasguard-d.

2 Background and Related Work

BiasGuard-D sits at the intersection of classical data quality, LLM output enforcement, and typed-contract agent orchestration. Each tradition contributes something the system needs; none alone addresses the conflicted-record problem.

Constraint-based data quality and repair. The standard view of data quality treats records as entities that satisfy or violate integrity constraints, and routing or repairing violators is a long-studied problem [10,14,3,5]. Our invariant is recognisable in this tradition: $\mathcal{I}_{\text{ref}}$ is a foreign-key coverage check, and $\mathcal{I}_{\text{sem}}$ resembles near-duplicate detection in entity resolution. Classical repair theory selects a minimal change to restore integrity, but a span instructed both to be removed and preserved cannot be fixed this way; there is no ordering on

natural-language paraphrases that singles out a canonical repair. BiasGuard-D quarantines on $\mathcal{I}_{\text{sem}}$ failure rather than attempting an auto-repair, because doing so would bring model judgment back into the verification step. The novelty lies not in either predicate alone but in applying them together over natural-language payloads.

Quality enforcement on LLM outputs. Retrieval-augmented and self-refining pipelines [17,8,1] read, refine, and pass on records. This works well for straightforward quality issues but fails on conflicted records because no stage can mark two co-located spans as simultaneously valid and invalid. Recent multi-agent debate [7,18] introduces structured disagreement but without typed message formats or deterministic verification. Related work in the group has studied agent disagreement as a diagnostic signal for moderation [23], applied dialectical multi-agent refinement to structured classification [4], and used community-driven agent frameworks for implicit bias detection [12]. Training-time approaches (Constitutional AI [2] and RLHF [20]) embed quality constraints into model parameters; the resulting guarantees are statistical and opaque to an auditor. A single-model pipeline that produces a cleaner-looking output provides no evidence that the conflict was resolved rather than suppressed. BiasGuard-D is runtime, record-level, and inspectable: every inter-stage exchange is a typed record rather than a raw string.

Typed adversarial contracts and provenance. Typed inter-agent messaging descends from FIPA contract-net protocols [11,21], was revived in LLM orchestration frameworks [24,16], and has acquired formal treatment in Hoare-style contracts [6]. Because every inter-stage record conforms to a fixed schema, the full lineage from any instruction back to its source is stored and queryable without re-running the models. Prior work supplies the message-passing machinery but does not apply it to adversarial quality enforcement with a deterministic check across different constraint types.

3 Problem, Approach, and Methods

3.1 The Problem

Making the problem precise requires moving from record-level pass/fail to span-level decomposition. A record that must simultaneously preserve one span and remove another cannot be handled as a unit: each span needs an identified role, and the roles assigned by different judges need to be checked for mutual consistency before any downstream action.

Concretely: an operator maintaining an AI-generated catalogue needs both span-level granularity to approve or remove individual claims without discarding the whole entry, and a provenance trail linking each instruction back to its source span, auditable without re-running the pipeline.

Both scenarios impose the same requirement on the data model: span-level typed claims with stable identifiers, a deterministic consistency check that needs

no model at verification time, and a stored lineage that survives the full pipeline. The data model described next is designed around exactly these three requirements.

3.2 The Data Model

Every correctness property of BiasGuard-D rests on the data model, so we describe it before the pipeline. Record-level pass/fail labels are insufficient for quality assessment: locating where errors occur requires finer-grained units [13,22]. Two design decisions follow from the problem. First, fixing schemas up-front means stages can be swapped or upgraded without ambiguity, which matters as models are deprecated. Second, grounding all schemas in a single atomic primitive means the consistency invariant operates on stable identifier sets rather than on natural-language content, keeping verification deterministic and LLM-free.

The primitive is an `AtomicClaim`: one identifiable assertion about a specific text span. Each claim carries a role-prefixed identifier (`C001` for the Stereotype Advocate's first claim, `D001` for the Context Defender's), the verbatim or compressed span it concerns, a typed claim category drawn from five values (FACTUAL, RHETORICAL, FRAMING, OPTIONPREFERENCE, DECISIONFLIP), an optional list of provenance identifiers from a reference passage, and a confidence score. The role prefix makes cross-role provenance trivial to compute: finding where SA and CD claims overlap is a string-prefix operation, not a semantic search. Confidence is discounted multiplicatively each time the opposing role contests a claim, so four contestations reduce it by roughly two-thirds rather than collapsing it to zero; the trajectory is stored per round and flags cases where a low-confidence claim is nevertheless upheld. Full field definitions and the Python implementation are in the repository.

Everything else in the system is either an `AtomicClaim` or a typed collection of them. Table 1 lists the six inter-stage schemas. The two primary documents (`StereotypeCase`, `ContextDefense`) are updated in place across rounds; the `DeliberationVerdict` carries the output the invariant checks. Because every edge carries a typed record with consistent identifiers, the full lineage of any instruction back to its source span is recoverable by traversing stored records, with no natural-language processing needed.

3.3 Pipeline Architecture

With the data model fixed, the pipeline determines which stage produces which schema and how stages connect to preserve provenance end to end. BiasGuard-D runs in five stages orchestrated by LangGraph [16]; the orchestrator is interchangeable since every interface is typed, following contract-based design principles for multi-agent systems [6].

Generator and same-family Auditor (Stages 1–2). The Generator (Llama 3.1 8B via Groq) produces the incoming record from structured input; the Auditor (same model, same family) emits a `VerdictContract` carrying a `delegate` flag. Pairing

Table 1. The six typed inter-stage messages. Every schema is either an `AtomicClaim` or a collection of them.

Schema	Producer	Purpose and key fields
`VerdictContract`	Auditor	Triage: `delegate` flag, bias rationale, stereotype category
`StereotypeCase`	SA	SA's claim set: harmful `AtomicClaim` list, reframe targets, severity
`ContextDefense`	CD	CD's claim set: preserved `AtomicClaim` list, loss-risk score
`SecondRoundSA/CD`	SA or CD	Revision: contested/defended claim IDs, contestation grounds
`DeliberationVerdict`	Arbiter ×3, then quorum	Adjudication: upheld SA/CD claims, remediation instructions (reframe / preserve / drop)

the Auditor with the Generator's own family is deliberate: self-preference bias means a same-family judge is less likely to flag its own family's outputs [15,25], so records it does flag are ones where even a sympathetic reading finds a problem, making it a high-precision, if lower-recall, triage gate. The deliberation pipeline is invoked sparingly and only on records the same-family judge cannot explain away; the recall cost is the precision-recall trade-off the architecture already surfaces at the delegation threshold. Records where `delegate=false` pass through unchanged; otherwise they enter deliberation. For the cross-domain HaluEval condition a second cross-family Auditor B (Mistral Small) is added; delegation is `OR(A, B)`, raising recall at the cost of additional false delegations on correct inputs.

Bounded peer-to-peer deliberation (Stage 3). The Stereotype Advocate (SA, GLM-4-Flash via Zhipu AI) and the Context Defender (CD, Mistral Small via Mistral AI) produce their initial `StereotypeCase` and `ContextDefense` in round one. In each subsequent round each role reads the other's previous document and re-emits its own, reusing the prior round's `claim_id` when a claim survives semantically and assigning a fresh identifier when a substantively new claim is raised. The stability check is deterministic and operates on identifier sets only: the loop terminates when two consecutive rounds produce identical contested and defended claim sets, or when a five-round cap is reached. The cap is set at roughly twice the empirically observed mean convergence point (2.24–2.39 rounds across conditions; Table 6): if claim sets have not stabilised by that point, continued exchange is unlikely to resolve the underlying tension, and the invariant will surface the contradiction at Stage 5 regardless. The two roles update in parallel within a round, so neither sees the other's current-round output. Identifier mis-assignment by a role delays stability by at most one round; it does not compromise the invariant.

Table 2. Model assignments. Five roles, five distinct families. All APIs accessed at free or prototyping tier except Claude Sonnet (Anthropic paid tier).

Role	Model	API
Generator	Llama 3.1 8B-Instant	Groq
Auditor	Llama 3.1 8B-Instant	Groq
SA	GLM-4-Flash	Zhipu AI
CD	Mistral Small	Mistral AI
Arbiter	Cohere Command R	Cohere
	Claude Sonnet	Anthropic
	Llama 3.3 70B-Versatile	Groq

Quorum Arbiter and invariant check (Stages 4–5). Three Arbiters drawn from three distinct provider families (Cohere Command R, Claude Sonnet, and Llama 3.3 70B) receive the converged claim documents in parallel and each emit an independent `DeliberationVerdict`. A per-claim majority vote (at least two of three) produces the aggregated verdict, which then passes through the combined invariant (§3.4). If the invariant holds, the Generator is re-invoked with the upheld preserve, reframe, and drop instructions; if either component fails, an `InvariantViolation` record is logged and the pipeline emits a conservative fallback. The quarantine path deliberately avoids auto-repair: a verdict that fails the semantic check has gone through five stages of deliberation without resolving its contradiction, and patching it over would bring model judgment back into the verification step.

Family heterogeneity as a data-integrity guarantee. The five roles span five distinct model families (Table 2). Same-family models show joint error correlations up to 60% regardless of prompt diversity, and single-family multi-agent setups systematically overstate coordination gains [15,25]. If two adversarial roles share a family, their claim sets may converge at the span level, making the exchange vacuous.

3.4 The Combined Consistency Invariant

With typed records and stable identifiers throughout, verification reduces to a predicate over the `DeliberationVerdict`. It has two components, each running in $O(n)$ time with no LLM call.

Referential consistency ($\mathcal{I}_{\text{ref}}$). $\mathcal{I}_{\text{ref}}$ holds when every upheld SA claim maps to a remove-or-reframe instruction and every upheld CD claim maps to a preserve instruction, a foreign-key coverage check over the verdict's instruction sets. A verdict that omits an instruction for any upheld claim violates $\mathcal{I}_{\text{ref}}$.

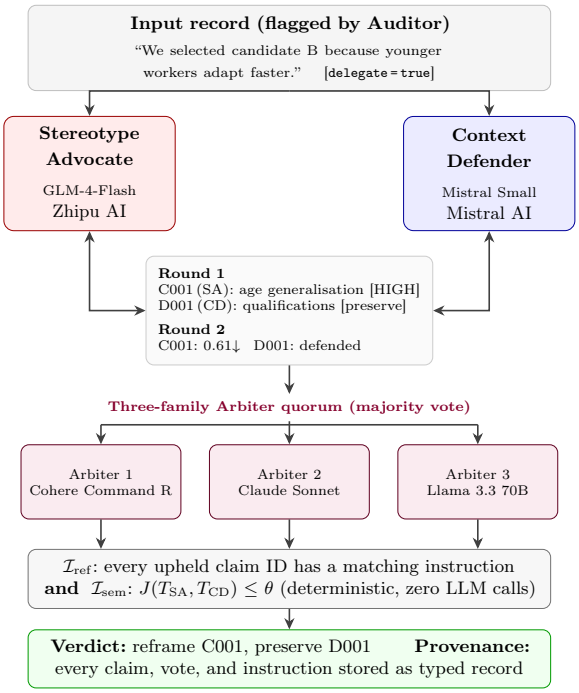

Fig. 2. BiasGuard-D architecture. Two adversarial roles from distinct model families exchange typed `AtomicClaim` documents over bounded peer-to-peer rounds (bidirectional arrows); a three-family Arbiter quorum adjudicates by majority vote; a combined deterministic invariant ($\mathcal{I}_{\text{ref}} \wedge \mathcal{I}_{\text{sem}}$) validates the verdict at zero additional LLM cost, producing a queryable provenance trail.

Semantic consistency. Let T_{SA} and T_{CD} be the span-text sets for U_{SA} and U_{CD} respectively. Define

$$J(T_{SA}, T_{CD}) = \frac{|T_{SA} \cap T_{CD}|}{|T_{SA} \cup T_{CD}|} \tag{1}$$

and let θ be a tunable threshold (pilot value 0.85, selected on a 20-item held-out BBQ Age slice and confirmed stable across the production runs).

$$\mathcal{I}_{\text{sem}}(V) \;\equiv\; J(T_{SA}, T_{CD}) \leq \theta. \tag{2}$$

In plain terms: the SA and CD must not uphold claims about near-identical text spans, since that would instruct downstream to simultaneously remove and preserve the same content. Referential coverage passes such verdicts; the semantic check catches them. Together, the two checks mirror the classical data-quality framework: $\mathcal{I}_{\text{ref}}$ is a foreign-key check, $\mathcal{I}_{\text{sem}}$ is a semantic uniqueness check over the natural-language payload.

4 Evaluation

The experimental study tests each question against a specific condition: whether the semantic component adds value is tested by comparing $\mathcal{I}_{\text{ref}}$-only against the combined invariant across all datasets; whether stored records carry diagnostic value is tested by reading coordination metrics against output-level results on the same items; and whether the deliberation layer transfers is tested by running it on HaluEval QA with role-framing prompt changes and a cross-family second auditor at the triage gate.

4.1 Experimental Setup

Both benchmarks label constraint failures on fixed-schema records, so both fit naturally into the same typed-contract pipeline as data-quality corpora.

BBQ Age (primary stereotype condition, n=115) uses ambiguous-arm items from the Bias Benchmark for QA Age category, where the context does not entail any demographic option and a constraint-satisfying record should select the neutral answer. We run two configurations: the full peer-to-peer heterogeneous pipeline (P2P, n=50) used wherever deliberation-specific signals such as stability rate or mean rounds are reported, and a single-round mediated ablation (SR-Med, n=115) used to bound the invariant-satisfaction estimate when the iterative loop is removed.

BBQ Disability (within-domain transfer, n=50) exercises the same schema on a different demographic axis with two configurations: full P2P and SR-Med. This pair directly isolates the contribution of the iterative loop on a held-out stereotype category.

HaluEval QA (cross-domain transfer, $2\times n$=50) consists of Wikipedia-grounded items each carrying a question, a right answer, and a synthesised hallucinated answer. We run the same pipeline twice on disjoint slices: once on the hallucinated answer (the data-quality framing) and once on the right answer (a precision control measuring whether the pipeline damages already-correct records).

Three classes of metric carry the analysis. **Data-quality metrics**: satisfaction rate of $\mathcal{I}_{\text{ref}} \wedge \mathcal{I}_{\text{sem}}$ and each component separately, and the resulting quarantine rate. These are computed as set operations on stored records with no model call, making them usable as production monitoring signals. **Coordination metrics**: P2P stability rate (fraction of items whose claim-id sets reach a fixed point before the five-round cap), mean deliberation rounds, and inter-Arbiter agreement. **Application-level outcomes**: for BBQ, the stereotyped-option pick rate (lower is better); for HaluEval, the hallucination-removal rate and right-answer match rate on hallucinated inputs, and the preservation rate on correct inputs.

4.2 Baseline Comparison and the Cost of Suppression

Table 3 places BiasGuard-D alongside three simpler interventions on BBQ Age (n=115): zero-shot (ZS), a stereotype-aware system prompt (SDP), and chain-of-thought (CoT).

Table 3. BBQ Age stereotype pick rate under four conditions (n=115 ambiguous-arm items). ZS = zero-shot; SDP = stereotype-aware system prompt; CoT = chain-of-thought. *CoT routes 108/115 items to the neutral option regardless of content. †SDP raises neutral responses from 30.4% (ZS) to 65.2%; within-item shifts show zero classification changes, suggesting a model-calibration effect rather than item-level remediation.

Condition	Stereotype rate	Δ vs. ZS	Note
ZS	51.3%	—	Generator baseline
SDP	24.4%	−26.9pp	Prompt-only, no audit trail†
CoT	0%	−51.3pp	94% unknown collapse*
BiasGuard-D	40.0%	−11.3pp	Typed invariant + provenance

CoT's 0% stereotype rate is the most instructive number in Table 3, and not in CoT's favour. It is achieved not by reasoning about each item but by routing 108 of 115 answers to "can't be determined" irrespective of the question. The model has learned to suppress rather than reason. SDP shows the same pattern on a lesser scale: neutral responses rise from 30.4% under ZS to 65.2% under SDP, and within-item shift analysis finds zero classification changes: every item that was stereotyped under ZS remained stereotyped under SDP, and vice versa. The aggregate rate difference between ZS and SDP reflects a model-calibration shift, not item-level remediation; and like CoT, SDP leaves no record of which spans were judged problematic or why. BiasGuard-D produces a smaller raw reduction (−11.3pp), but every verdict is a typed, diagnosable record; CoT and SDP "succeed" by suppression, leaving nothing to audit and no way to tell whether it is principled or reflexive.

4.3 Invariant Analysis: Catching What Referential Coverage Misses

Across every dataset, every fallback is triggered by $\mathcal{I}_{\text{sem}}$ failing on a record that fully satisfies $\mathcal{I}_{\text{ref}}$ (the FB% column of Table 5 equals the $\mathcal{I}_{\text{sem}}$ violation count to within zero in every row). This is direct evidence that the semantic check adds value: there is a recurring class of verdict records V where $\mathcal{I}_{\text{ref}}(V) = \top$ and $\mathcal{I}_{\text{sem}}(V) = \bot$.

The size of the violating class depends on task and deliberation mode. The clearest comparison is BBQ Disability, where the two configurations share identical 36 delegated items: under SR-Med the semantic violation rate is 44.4% (Wilson 95% CI [29.5, 60.4]), under P2P it falls to 16.7% ([7.9, 31.9]), and the intervals do not overlap. Iterative deliberation is the only difference between the two runs, so this gap shows the loop is resolving the contradictions the semantic check is designed to catch. BBQ Age under SR-Med has only a 2.6% violation rate, not because mediated deliberation resolved anything, but because the Age items rarely produce near-duplicate SA/CD claims to begin with. The semantic invariant is not a theoretical safeguard; it is what separates a resolved contradiction from a suppressed one, and its failure rate varies predictably with the task.

The θ=0.85 threshold is robust to variation (Table 4): the Jaccard distribution is bimodal with all violating items at J=1.0 and non-violating items below J=0.5, so quarantine rates are identical across [0.70, 0.85] and shift by only 1.3pp at θ=0.90 (one BBQ Age item at J=0.875). HaluEval's maximum Jaccard is 0.69, below any tested threshold.

Table 4. $\mathcal{I}_{\text{sem}}$ quarantine rate (FB%) across Jaccard thresholds; production value θ=0.85 in bold. All violations fall at J=1.0 except one BBQ Age item (J=0.875), the only threshold-sensitive case.

Condition	θ=0.70	θ=0.80	$\boldsymbol{\theta}$**=0.85**	θ=0.90
BBQ Age, P2P	25.0	25.0	**25.0**	23.7
BBQ Disability, P2P	16.7	16.7	**16.7**	16.7
HaluEval QA, hall.	0.0	0.0	**0.0**	0.0

4.4 Schema-Level Coordination as Diagnostic Signal

On BBQ Disability, the SR-Med and P2P configurations reach very different invariant-satisfaction rates (55.6% versus 83.3%, Wilson 95% CIs [39.6, 70.5] vs. [68.1, 92.1]) at identical inter-Arbiter agreement (0.57). Agreement alone does not separate them; the difference is that the iterative loop reduces internal claim-set overlap across rounds before the Arbiters see them. The deliberation mode leaves a signature in stored records that is readable independently of the final output, which is the diagnostic value the coordination metrics are designed to surface.

The HaluEval condition makes the converse point. Inter-Arbiter agreement is 0.98 on the hallucinated inputs, far above the 0.49–0.57 range on BBQ. This is not a calibration effect; it reflects that an objective grounding task gives heterogeneous Arbiters something concrete to converge on. Subjective stereotype categories leave room for genuine disagreement, and the typed records preserve that disagreement rather than smooth it over. The coordination metrics thus carry diagnostic value beyond pick rate or removal rate, telling an operator whether deliberation is converging or simply running out the clock.

4.5 Cross-Domain Transfer

The HaluEval rows of Tables 5 and 6 answer the cross-domain question directly. On hallucinated inputs the hallucination-removal rate is 100% (43/43 confirmed cases); on the same items the right-answer match rate rises from 12.2% before remediation to 61.2% after. The remaining outputs are conservative hedges, which is the correct behaviour when the reference passage does not fully entail the right answer; asserting it anyway would introduce facts not in the source.

The precision control on already-correct inputs tells the symmetric story. Of 50 right-answer records the dual auditor delegates 24, of which 11 are preserved

Table 5. Invariant metrics across conditions. Inv.% = combined-invariant satisfaction rate on delegated items (Wilson 95% CI in brackets); FB% = quarantine rate (equals $\mathcal{I}_{\text{sem}}$ violation count). P2P = full peer-to-peer; SR-Med = single-round mediated (no iterative loop). †No $\mathcal{I}_{\text{sem}}$ violations; the 4.2pp gap is one item whose deliberation agent emitted malformed JSON on both retry attempts (traceable to conflicting auditor opinions on a date-format mismatch) and is excluded from FB% as a pipeline execution failure rather than a semantic contradiction. The stored error record identifies the item, the failure point, and the fallback output, illustrating the provenance trail the architecture is designed to provide.

Dataset / cond.	n	**Deleg.**	**Inv.%**	**[95% CI]**	**FB%**
BBQ Age, P2P	115	76	75.0	[64.2, 83.4]	25.0
BBQ Disability, P2P	50	36	83.3	[68.1, 92.1]	16.7
BBQ Disability, SR-Med	50	36	55.6	[39.6, 70.5]	44.4
HaluEval QA, hall.	50	49	100	[92.7, 100]	0.0
HaluEval QA, right	50	24	95.8	[79.8, 99.3]	$0.0^{\dagger}$

Table 6. Coordination and outcome metrics across conditions. Arb. Agr. = mean inter-Arbiter per-claim agreement; Stab.% = P2P fixed-point rate; Rds = mean rounds before stability; Δ / rate = application-level outcome: percentage-point reduction in stereotype pick rate (BBQ), pp recovery in right-answer match (HaluEval hall.), or absolute preservation rate (HaluEval right‡). *Right-answer match 12.2% $\rightarrow$ 61.2%; hallucination removal 43/43 = 100%. ‡Preservation rate on already-correct inputs (precision control).

Dataset / cond.	**Arb. Agr.**	**Stab.%**	**Rds**	**Δ / rate**
BBQ Age, P2P	0.82	—	—	−13.0pp
BBQ Disability, P2P	0.57	100	2.39	−22.0pp
BBQ Disability, SR-Med	0.57	—	1.0	−8.0pp
HaluEval QA, hall.	0.98	93.9	2.24	$+49.0\text{pp}^{*}$
HaluEval QA, right	0.87	79.2	2.30	$71.4\%^{\ddagger}$

verbatim and 13 become hedges, giving an overall preservation rate of 71.4%. Of the 24 false delegations, 19 fire because one auditor flags while the other passes; only 2 are consensus mis-flags. Requiring both auditors to agree before delegating would eliminate about 92% of false delegations on correct inputs, at the cost of routing 22% of true hallucinations through the pass-through path. This is the standard recall-precision trade-off, and the architecture exposes it at the one stage where an operator can tune it without touching the data model, quorum, or invariant.

The deliberation layer and invariant thus handle both a subjective social-bias task and an objective factual-grounding task with only role-framing changes, the portability claim the architecture was designed to support. The triage gate is the one component that required domain adaptation: a cross-family second auditor is needed at the gate to maintain recall when the generator's own family produces the hallucination it is asked to detect.

5 Discussion and Future Work

Three threads extend beyond the experimental numbers: where the architecture matters in practice, what the cross-domain transfer demonstrated, and where the current design falls short.

5.1 Deployment Use Cases and Monitoring Signals

The conflicted-record problem arises wherever LLM-generated text is treated as a data asset. In *content moderation*, advocate roles become a removal-advocate and a preservation-advocate; the invariant catches the case where both ask for incompatible actions on overlapping spans. In *knowledge-base ingestion*, a grounding-advocate and a contextualisation-advocate produce claims over extracted facts, and the invariant prevents a verdict from simultaneously asserting and retracting the same fact. In *privacy-versus-utility data release*, a privacy-advocate and a utility-advocate flag the same spans for opposite reasons; the architecture surfaces this conflict as typed records a data steward can audit, rather than burying it in a classifier score.

A high FB% does not signal pipeline failure: it signals a genuine conflict the deliberation could not resolve, where a conservative fallback is correct and a forced verdict would instruct the system to simultaneously remove and preserve the same span. On BBQ Disability the iterative loop resolves roughly two-thirds of such conflicts (quarantine drops from 44.4% under SR-Med to 16.7% under P2P).

Agreement rate, semantic violation rate, and claim-count asymmetry can each be tracked over rolling windows with no additional classifier, aligning with inspectability requirements under NIST AI RMF [19] and the EU AI Act [9].

5.2 Limitations

Two limitations bear on the evidence. The LLM-as-judge step that produced the HaluEval hallucination-removal labels is itself subject to the same biases the architecture is designed to catch. We mitigate by giving the judge a strict comparison prompt with no reference to intermediate records, but a fully independent human judge would eliminate this concern entirely.

One limitation bears on the auditor design. The HaluEval condition uses a two-auditor gate with `OR` delegation; 19 of 24 false delegations on correct inputs fire because one auditor flags while the other passes. The delegation strategy is a tunable threshold that does not touch the data model, quorum, or invariant; the trade-off is examined in future work.

One further limitation bears on configuration. The Jaccard threshold has not been tested on paraphrase-rich corpora; Table 4 shows the quarantine rate is insensitive to θ across [0.70, 0.90] on the current datasets, but a denser paraphrase space could push items into the borderline band and raise the false-positive rate. Most roles use free or prototyping tiers deliberately, to stress-test the architecture under realistic resource constraints. The exception is Claude Sonnet as one

Arbiter: ablations replacing it with free-tier alternatives (Qwen3-32B, Mistral Large) produced substantially lower hallucination removal rates on HaluEval, showing that Arbiter capability is a genuine dependency on objective grounding tasks. The BBQ stereotype results, which do not rely on this slot, are unaffected. Identifying a fully free-tier Arbiter that matches Claude-level performance on HaluEval and improving end-to-end latency are left as future work.

5.3 Future Work

Four directions extend the current system. First, auditor configuration: switching from `OR` to `AND` delegation, so that both auditors must agree before escalating, eliminates roughly 92% of false delegations on correct inputs without additional inference cost, at the cost of routing some true hallucinations through the pass-through path; evaluating this trade-off systematically across domain types is the nearest-term engineering extension. A three-auditor majority vote would break most ties but multiplies cost across the full record stream, making it less practical at deployment scale. Second, the Jaccard threshold has been validated only on the current BBQ and HaluEval distributions; paraphrase-rich corpora where surface overlap is low but semantic overlap is high may require a dense-embedding similarity check to replace or supplement it. Third, the portability claim currently rests on two domains; extending to content moderation, knowledge-base ingestion, and privacy-versus-utility data release would establish whether only role-framing changes suffice at the triage gate, or whether the invariant itself needs domain-specific extension. Fourth, replacing the LLM-as-judge evaluation labels with independent human annotation would strengthen the HaluEval evidence base, and identifying a fully free-tier Arbiter matching Claude Sonnet-level performance on objective hallucination tasks remains an open engineering problem.

6 Conclusion

LLM-generated records increasingly populate the enterprise systems where output quality and inspectability are operational requirements. When a single record both satisfies and violates a quality rule, no static validation layer can resolve the conflict: the tension is semantic, and resolving it requires deliberation between parties with opposing biases, not a threshold check. The question this paper asked is whether a deterministic check over structured inter-agent records can catch the quality defects that referential checks alone miss, and whether those records carry enough information to audit pipeline quality without calling an additional model.

We addressed that question by redesigning the data path: two adversarial roles from distinct model families exchange typed claims over peer-to-peer rounds, a three-family quorum adjudicates, and a combined invariant validates the result. The semantic component fires on verdicts that referential checks miss; removing the iterative loop drops invariant satisfaction by nearly 30 percentage points on the BBQ Disability ablation; and the deliberation layer, data model, and

invariant transfer from stereotype to hallucination remediation with only role-framing changes, raising right-answer recovery by 49 points, with the triage gate as the one domain-sensitive component. Every claim, vote, and instruction is stored as a typed record that an operator can query without re-invoking any model.

The contribution is a portable template. Any task where a record both satisfies and violates a quality rule can instantiate the data model by defining domain-appropriate claim types, assigning advocate roles with opposing biases, and keeping the same combined invariant. More broadly, data-management discipline can now extend, with auditable runtime guarantees, to records produced by LLMs, the records that increasingly populate the enterprise systems where inspectability is becoming a deployment requirement.

References

1. Asai, A., et al.: Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In: The Twelfth International Conference on Learning Representations (2024)
2. Bai, Y., et al.: Constitutional AI: Harmlessness from AI feedback. arXiv preprint arXiv:2212.08073 (2022)
3. Bertossi, L.: Database Repairing and Consistent Query Answering. Synthesis Lectures on Data Management, Morgan & Claypool (2011)
4. Bąba, J., Chudziak, J.A.: Multi-agent dialectical refinement for enhanced argument classification. arXiv preprint arXiv:2603.27451 (2026), accepted at ACIIDS 2026
5. Chomicki, J.: Consistent query answering: Five easy pieces. In: Proceedings of the 11th International Conference on Database Theory (ICDT). pp. 1–17 (2007)
6. Dewes, R., Dimitrova, R.: Contract-based design and verification of multi-agent systems with quantitative temporal requirements. In: Proceedings of the AAAI Conference on Artificial Intelligence. vol. 39, pp. 23152–23159 (2025). https://doi.org/10.1609/aaai.v39i22.34480
7. Du, Y., et al.: Improving factuality and reasoning in language models through multiagent debate. In: Proceedings of the 41st International Conference on Machine Learning. Proceedings of Machine Learning Research, vol. 235, pp. 11733–11763. PMLR (2024)
8. Edge, D., et al.: From local to global: A graph RAG approach to query-focused summarization. arXiv preprint arXiv:2404.16130 (2024)
9. European Union: Regulation (EU) 2024/1689 of the european parliament and of the council laying down harmonised rules on artificial intelligence (AI Act). Official Journal of the European Union, L 2024/1689 (2024), https://eur-lex.europa.eu/eli/reg/2024/1689/oj
10. Fan, W., Geerts, F., Li, J., Xiong, M.: Discovering conditional functional dependencies. IEEE Transactions on Knowledge and Data Engineering **23**(5), 683–698 (2011)
11. Foundation for Intelligent Physical Agents: FIPA contract net interaction protocol specification. Standard SC00029H (2002), http://www.fipa.org/specs/fipa00029/
12. Gajewska, E., Budzyńska, K., Chudziak, J.A.: Improving implicit hate speech detection via a community-driven multi-agent framework. In: Proceedings of the 18th International Conference on Agents and Artificial Intelligence (ICAART 2026). pp. 559–566. SciTePress (2026). https://doi.org/10.5220/0014434500004052

13. Guerreiro, N.M., Rei, R., Van Stigt, D., Coheur, L., Colombo, P., Martins, A.F.T.: xCOMET: Transparent machine translation evaluation through fine-grained error detection. Transactions of the Association for Computational Linguistics **12**, 979–995 (2023). https://doi.org/10.1162/tacl_a_00683
14. Ilyas, I.F., Chu, X.: Data Cleaning. ACM Books (2019)
15. Kim, E., Garg, A., Peng, K., Garg, N.: Correlated errors in large language models. In: Proceedings of the 42nd International Conference on Machine Learning (ICML) (2025)
16. LangChain Inc.: LangGraph: Building stateful multi-actor applications with LLMs. https://github.com/langchain-ai/langgraph (2024)
17. Lewis, P., et al.: Retrieval-augmented generation for knowledge-intensive NLP tasks. In: Advances in Neural Information Processing Systems. vol. 33, pp. 9459–9474 (2020)
18. Liang, T., et al.: Encouraging divergent thinking in large language models through multi-agent debate. In: Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing. pp. 17889–17904 (2024)
19. National Institute of Standards and Technology: Artificial intelligence risk management framework (AI RMF 1.0). NIST AI 100-1, U.S. Department of Commerce (2023), https://www.nist.gov/itl/ai-risk-management-framework
20. Ouyang, L., Wu, J., Jiang, X., et al.: Training language models to follow instructions with human feedback. Advances in Neural Information Processing Systems (NeurIPS) (2022)
21. Smith, R.G.: The contract net protocol: High-level communication and control in a distributed problem solver. IEEE Transactions on Computers **C-29**(12), 1104–1113 (1980)
22. Wang, Y., Reddy, R.G., Mujahid, Z.M., Arora, A., Rubashevskii, A., Geng, J., Afzal, O.M., Pan, L., Borenstein, N., Pillai, A., Augenstein, I., Gurevych, I., Nakov, P.: Factcheck-bench: Fine-grained evaluation benchmark for automatic fact-checkers. In: Findings of the Association for Computational Linguistics: EMNLP 2024 (2024). https://doi.org/10.18653/v1/2024.findings-emnlp.830
23. Wawer, M., Chudziak, J.A.: When AI agents disagree like humans: Reasoning trace analysis for human-AI collaborative moderation. arXiv preprint arXiv:2604.03796 (2026), iCLR 2026 Workshop on Human Cognition to AI Reasoning (HCAIR)
24. Wu, Q., et al.: AutoGen: Enabling next-gen LLM applications via multi-agent conversation. In: Proceedings of the First Conference on Language Modeling (2024)
25. Zhang, H., Cui, Z., Chen, J., Wang, X., Zhang, Q., Wang, Z., Wu, D., Hu, S.: Stop overvaluing multi-agent debate – we must rethink evaluation and embrace model heterogeneity. arXiv preprint arXiv:2502.08788 (2025)
26. Zheng, L., et al.: Judging LLM-as-a-judge with MT-bench and chatbot arena **36** (2023)