arXiv:2610.29435v1 [cs.AI] 21 Oct 2026

# Events, Not Pixels: Where World Models Get the Return Wrong

Haoyue Liu[1,2] Xiaoying Tang[1,2,3,*]

[1]School of Science and Engineering, The Chinese University of Hong Kong, Shenzhen

[2]Shenzhen Future Network of Intelligence Institute (FNii-Shenzhen)

[3]Guangdong Provincial Key Laboratory of Future Networks of Intelligence, CUHK(SZ)

[*]Corresponding author.

## Abstract

*World models are increasingly used for planning, where a planner ranks candidate actions by returns imagined by the model. However, these models are judged mainly by how well they render frames, even though a visually faithful rollout can still assign the wrong return. To identify where this mismatch arises, we introduce a transition-level diagnostic that replaces imagined transitions with real ones and measures the reduction in return error against a budget-matched random control. Our analysis yields two findings. (1) Event transitions, the few steps where a reward is earned or an episode ends, are rendered as well as other steps (median pixel-error ratio 0.85), yet restoring them removes a median of 70% of the return error. In contrast, restoring transitions with the largest MSE, SSIM or LPIPS error removes less than a random choice. Across the 44 Atari-100k model–collection pairs that pass our reliability checks, including three released families, event patches outperform random patches in 39 (sign test $p = 1.4 \times 10^{-7}$), with a median gain of +44 percentage points. (2) The remaining error is structured: a complementary patch separates errors at true events from those elsewhere, while an eight-step timing tolerance recovers few events missed by our deterministic models. This structure also guides repair: a rule recorded before four repair outcomes were known predicts when a diffusion model repairs these failures. Event patches improve plan ranking, and the same concentration appears on control tasks with rare reward events.*

## 1. Introduction

World models let an agent imagine the consequences of an action before taking it [18]. A planner rolls candidate action sequences forward inside the model and acts on the ordering of their imagined returns [2, 11, 20, 22, 23, 36, 48]. Released pixel world models support this pipeline in Atari, continuous control, and manipulation [12, 52, 60], and their progress is mainly measured by how faithfully they render the future [53, 59, 65].

Rendering quality, however, is only loosely tied to what a planner needs. A return often hinges on a few transitions that carry a reward or end an episode, so a model can render a rollout well yet misprice it by missing a single paddle contact. Model-level studies expose this mismatch: rankings by visual quality can disagree with rankings by task success [33, 64], and whole-model oracle substitutions find that average prediction accuracy does not consistently explain planning performance [7, 41, 49]. Yet neither these analyses nor compounding-error bounds [26] reveal where that error occurs: is it spread across all sixty steps of a rollout, or concentrated in only a few? We therefore ask:

*Where inside an imagined rollout does the error that changes a planner's decision occur?*

To answer this question, we introduce a transition-level substitution diagnostic (Fig. 1). Each imagined rollout is paired with a real rollout that starts from the same context and follows the same recorded actions, so the two differ only in who drew the frames. A frozen, memoryless detector, trained once per game, reads reward and termination from each pair of consecutive frames; replacing one imagined pair with the real one therefore changes the return only at that step. We patch the pairs at true events, the same number of random pairs, all non-event pairs, or every pair. These four patches isolate the targeting effect, its budget-matched control, the error away from events, and an all-real reference.

Our analysis yields two findings. **First, return error concentrates at events that do not look wrong.** In Fig. 1D, restoring three transitions around a Seaquest reward removes almost all return error; three random ones remove little. Across 44 reliable Atari model–collection pairs, event patches beat random patches in 39 by a median of +44 points, and the advantage appears within a few patches (Fig. 2). Event frames are rendered as well as others (median pixel-error ratio 0.85), yet restoring them removes a

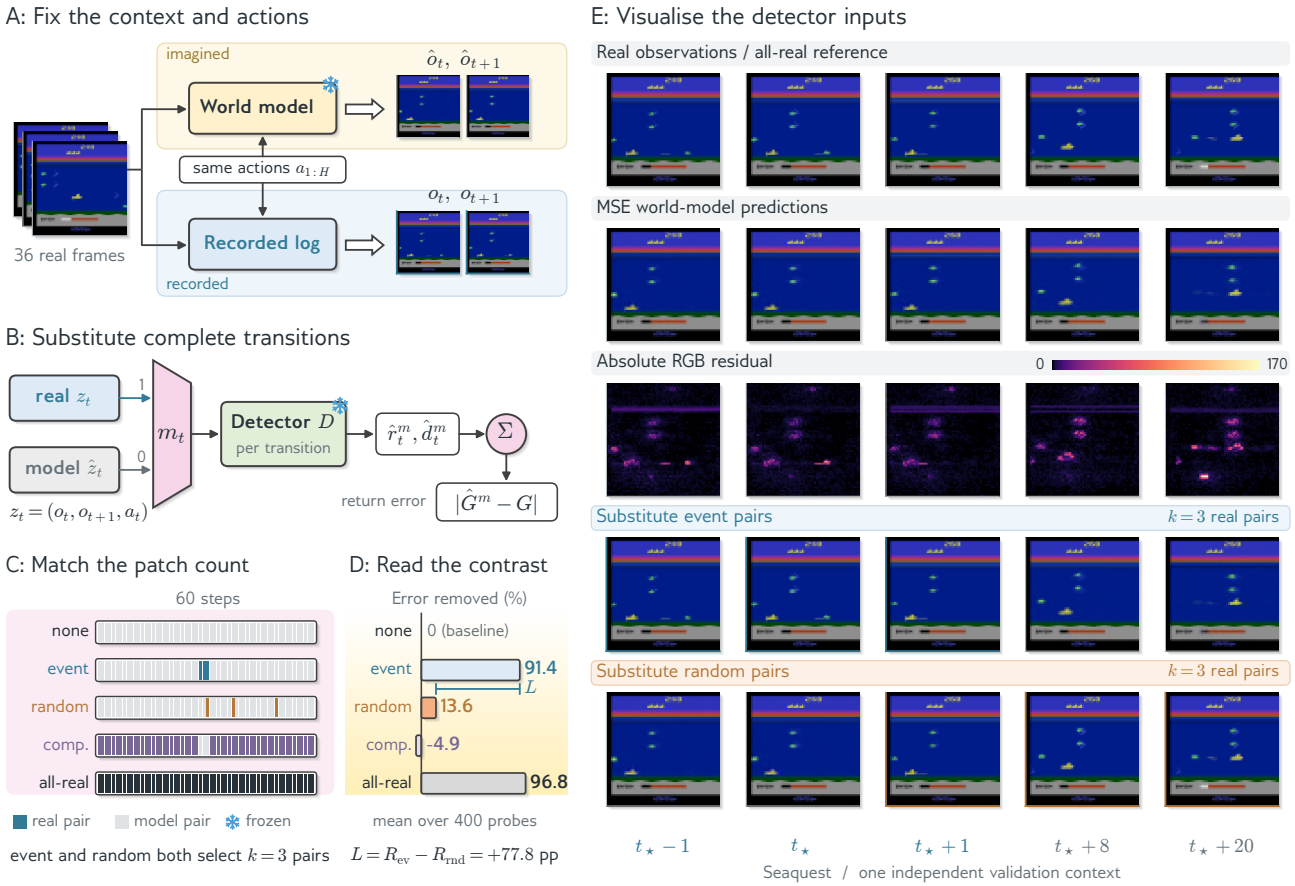


Figure 1. **Return error concentrates at event transitions that pixel error does not single out. A:** A saved Seaquest MSE model predicts a 60-step future from 36 real context frames and recorded actions. **B:** Each selected detector input substitutes both frames of a transition, keeping its action fixed. **C:** Event and random masks each select three of sixty steps here; complement and all-real inputs complete the controls. **D:** Saved return-error reductions over 400 probes. **E:** Five selected times show real frames, predictions, absolute RGB residuals on a shared square-root colour scale, and the frames supplied by the two substitution arms. Coloured borders mark real-pair substitutions.

median of 70% of the return error; the same pattern holds on eighteen games added after our predictions were recorded. **Second, the remaining error is structured rather than merely mistimed.** Complementary patches separate errors at true events from errors elsewhere, and timing-tolerant recall shows that temporal offset explains little of the events our deterministic models miss. This structure also predicts where a diffusion model helps; event patches improve plan ranking and carry over to control tasks with rare reward events.

We summarise our contributions as follows:

- We introduce a per-transition substitution diagnostic with a budget-matched random control. It treats the world model as a black box, and we apply it to three released model families on all 26 Atari-100k games.
- We show, for the first time on released pixel world models, that return error concentrates at reward and termination transitions, also under each model's own reward head and without the score display, and that a complementary patch separates errors at true events from other errors.
- We show that a temporal offset explains little of the events missed by deterministic models, predict all four repair outcomes with a rule recorded in advance, and show that event patches improve plan ranking in every counted pair and carry over to control tasks with rare reward events.

## 2. Related Work

Pixel world models support prediction and control in games [2, 10, 12, 21, 29, 36, 37, 40, 46, 66] and interactive video simulators [6, 54, 60, 62]. Yet visual fidelity can diverge from physical correctness [38] and task success [33, 58, 64]; the same disconnect appears for training likelihood [13, 16, 31], synthetic training data [1], and token pruning [57]. Concurrent work weights error by decision-relevant state dimensions [56], while DIAMOND shows that visual details matter [2]. We instead locate where decision-relevant error occurs inside a rollout.

Oracle substitution localises such failures. Cloning De-

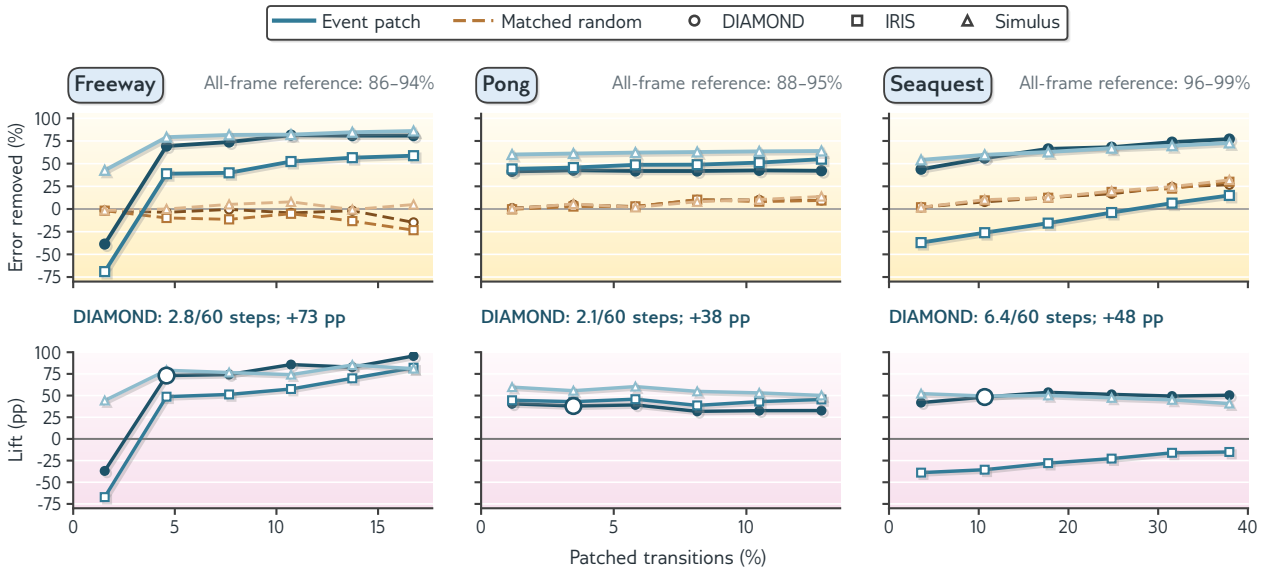


Figure 2. **Targeting advantage across patch budgets.** 9 sweeps, three games, 3 released families. Solid/dashed curves denote event/random patches; markers identify families. Lower panels show their difference, $L$, at matched coverage. Dotted guides mark the second DIAMOND budget; annotations give its mean patched-step count and lift. All-frame reference ranges are listed above each panel. Sweeps use one sampling seed; Table 1 reports three-seed means.

terministic Worlds [61] substitutes true states for learned representations, while whole-model oracle interventions relate accuracy to planning [7, 41, 49]; a concurrent diagnostic uses a kinematic null [47]. Closest to ours, counterfactual debugging [32] assigns Shapley credit over time steps to a policy's sim-to-real gap under injected faults, and agent explanations identify steps that decide a return [3, 8, 17]. We instead operate directly on released world-model rollouts, replacing individual transition pairs under a budget-matched random control. This brings perturbation-based attribution [14, 25, 44, 63], used and tested on Atari agents [4, 15], from pixels to time steps.

Other work improves prediction with multi-step training [50, 55], goal-aware objectives [39], keyframe targets [27, 43], region- and event-aware objectives [34, 42], causal rollout structure [67], and architecture-level probing [28]. Conditional-mean effects are also studied in video prediction [35] and perception and driving models [45]. We test which explanation accounts for missed events using event survival, a diffusion model, and the same network trained deterministically.

## 3. Diagnostic Setup

The diagnostic needs only rendered frames, so it applies to released checkpoints without access to their internals.

**Return from a Predicted Trajectory.** A *probe* is a real context with fixed actions $a_{1:H}$ (Fig. 1A); from it, a model produces imagined observations $\hat{o}_{1:H}$. A frozen detector, shared by all world models of a game, reads each transition pair and its action as a reward sign and a termination flag, $(\hat{r}_t, \hat{d}_t) = D(\hat{o}_t, \hat{o}_{t+1}, a_t)$ for $t < H$; the last frame has no successor, so $\hat{r}_H = \hat{d}_H = 0$. The emulator, replaying the same actions, supplies $r_{1:H}$ and $d_{1:H}$, and both returns stop at the first termination. Rewards enter through their sign, as in Atari-100k training, so all games share one scale:

$$\hat{G} = \sum_{t=1}^{H} \gamma^{t-1} \operatorname{sgn}(\hat{r}_t) \prod_{j<t} (1 - \hat{d}_j), \quad E = \mathbb{E}|\hat{G} - G|, \tag{1}$$

where $G$ is computed in the same way from $r$ and $d$, and $\mathbb{E}$ averages over probes. Atari uses $H = 60$ and $\gamma = 0.99$.

**Substitutions and Their Controls.** For channel $c$ (observation $o$, reward $r$ or termination $d$) and a binary mask $m$, let $\hat{G}^{c,m}$ be the return after substituting the selected real information (Fig. 1). Observation patches replace *both* frames of the selected transition before detection, so no pair mixes real and generated frames; unselected pairs keep their predictions. Reward and termination interventions replace $\hat{r}$ or $\hat{d}$. We report the fractional reduction in mean absolute return error:

$$R(c, m) = 1 - \frac{\mathbb{E}[|\hat{G}^{c,m} - G|]}{\mathbb{E}[|\hat{G} - G|]}. \tag{2}$$

The all-frame value $R(o, \mathbf{1})$ measures performance when the detector sees only real transition pairs.

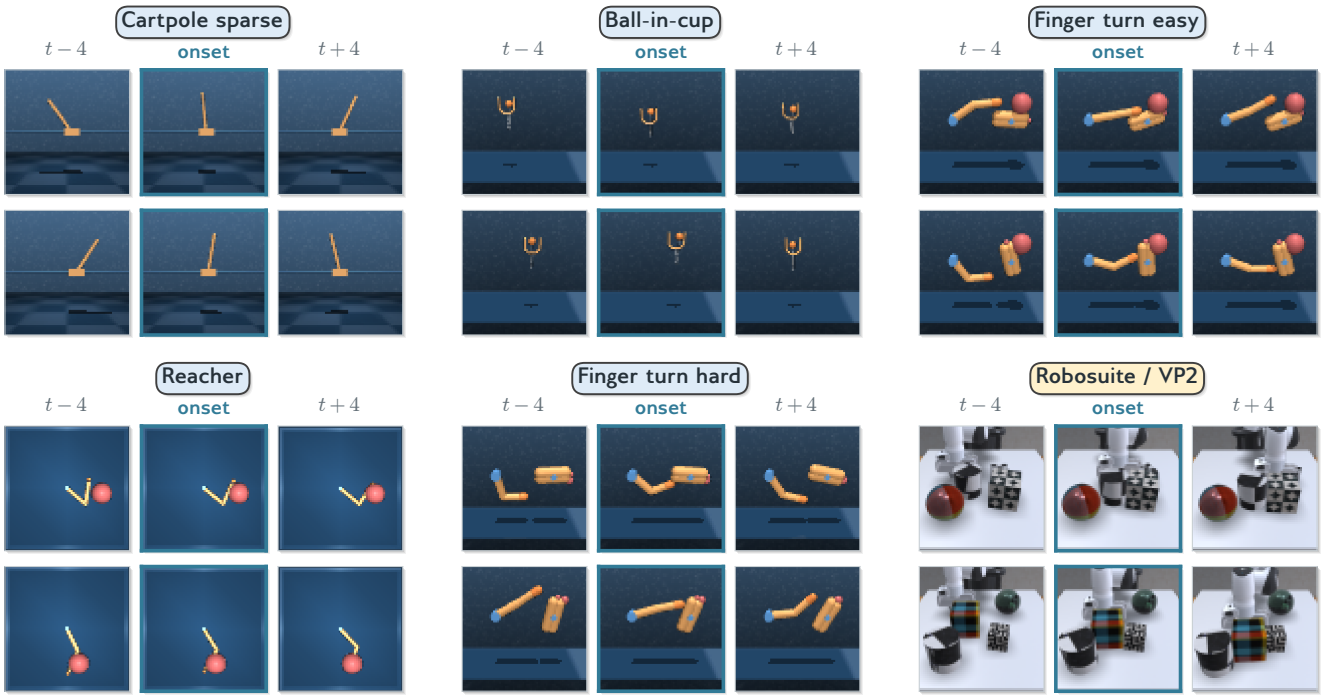


Figure 3. **Task-event transitions beyond Atari.** Six tasks, two real trajectory segments per task, and three frames per segment. Blue borders mark reward onset, the first rewarded step after at least four reward-free steps. Cases are selected from real events, separated by more than 96 steps within each task. Aggregate targeting and detector results are in the supplement.

The event mask starts from nonzero real reward or termination and is dilated by one step, i.e., the transitions just before and after are also selected, to cover delays in the visual markers. We report its realised coverage $q = \mathbb{E}[H^{-1} \sum_t m_t^{\mathrm{ev}}]$, which includes dilation. The random mask selects the same number of steps in each rollout. The targeting advantage is

$$L = R(o, m^{\mathrm{ev}}) - R(o, m^{\mathrm{rnd}}). \tag{3}$$

The complementary mask $\bar{m}^{\mathrm{ev}}$ measures errors away from the selected events, and its contrast with the event patch classifies the failure pattern of each model.

## 4. Localising the Return Error

We localise the return error with one frozen, memoryless detector per game, shared by every world model of that game. A random patch also hits true events, and more of them as coverage grows, so the lift of Eq. (3) credits the event patch only with what it adds beyond chance.

Figure 2 widens the event neighbourhood step by step and plots the lift against the resulting coverage. In 8 of 9 sweeps over three games and three released families, the event curve stays well above the random curve. For DIAMOND on Freeway, patching 2.8 of sixty steps removes 69% of the error, against $-4\%$ for random patches and 86% when every frame is real; on Pong, 2.1 steps remove 43% against 5%. The sweep also exposes a timing effect: without dilation the Freeway event patch is negative, consistent with the reward digit appearing after the transition that earns it, so we treat the neighbourhood as part of the measurement.

The sweep also shows why the random control is needed. For DIAMOND and Simulus on Seaquest, both patches remove more error as coverage grows while the lift stays flat, so the event-patch reduction alone would overstate the targeting effect. Table 1 reports every game at a one-step neighbourhood, and Fig. 3 shows events beyond Atari.

## 5. Missing Events and a Model-Class Repair

Localisation shows where the return error lies, not why. We ask why imagined events are missed and whether another model class draws them.

### 5.1. Absent or Mistimed Events

A detector can miss an event in imagined frames for two reasons: the model draws it at a neighbouring step, or it does not draw it in a form the detector reads. The two predict different curves. If events are only mistimed, recall should rise with the alignment tolerance faster than for a displaced-prediction null; if they are not drawn readably, recall should stay low at every tolerance.

Table 1. **Event targeting separates concentrated errors from diffuse failures.** Percent of return error removed by patching event (Ev), complementary (Non) or budget-matched random (Rnd) transitions (mean of three seeds); shading grows with the lift $L = \mathrm{Ev} - \mathrm{Rnd}$. Rows are sorted by event coverage $q$; grey pairs are outside the counted set, and games whose detector fails the gate are omitted (supplement).

| **Game** | **DIAMOND [2]** | | | | | **IRIS [36]** | | | | **Simulus [12]** | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| | $q$ (%) | Ev | Non | Rnd | $L$ | Ev | Non | Rnd | $L$ | Ev | Non | Rnd | $L$ |
| Pong | 3.5 | 41 | 46 | 5 | +36 | 48 | 38 | 4 | +45 | 63 | 23 | 4 | +59 |
| Kangaroo | 4.1 | 79 | 4 | 5 | +74 | 98 | 0 | 5 | +92 | 38 | −1 | 2 | +36 |
| Freeway | 4.6 | 68 | −6 | −4 | +72 | 35 | −38 | −2 | +38 | 72 | −16 | 1 | +71 |
| Seaquest | 10.7 | 57 | −19 | 9 | +48 | −26 | 43 | 7 | −33 | 54 | −28 | 8 | +45 |
| Frostbite | 12.0 | 57 | 17 | −181 | +238 | 71 | −18 | 5 | +65 | 83 | −4 | 16 | +66 |
| RoadRunner | 12.3 | 82 | −2 | 20 | +62 | 91 | 1 | 21 | +70 | 81 | −2 | 19 | +62 |
| KungFuMaster | 13.9 | 93 | −3 | 17 | +75 | 83 | −7 | 16 | +67 | 92 | −2 | 17 | +75 |
| Asterix | 16.1 | 71 | 2 | 12 | +59 | 35 | 1 | 12 | +23 | 60 | −11 | 16 | +44 |
| Hero | 16.4 | 70 | 25 | 53 | +18 | 30 | 36 | 24 | +7 | 9 | 19 | 14 | −5 |
| Boxing | 19.6 | 12 | 72 | 19 | −7 | 79 | −3 | 17 | +62 | 70 | −5 | 17 | +53 |
| Gopher | 24.8 | 77 | −4 | 33 | +44 | 85 | −7 | 34 | +51 | 92 | −4 | 38 | +55 |
| Qbert | 25.3 | 46 | 38 | 27 | +19 | 92 | −2 | 47 | +45 | 77 | 9 | 42 | +35 |
| CrazyClimber | 38.2 | 54 | 33 | 33 | +20 | 58 | 26 | 34 | +23 | 88 | 3 | 49 | +40 |
| MsPacman | 42.1 | 43 | 8 | 37 | +7 | 30 | 69 | 33 | −3 | 33 | 65 | 36 | −3 |
| Krull | 58.5 | 71 | 15 | 60 | +11 | 90 | −1 | 76 | +14 | 93 | −3 | 78 | +16 |
| **Median, counted** | | | | 13/14 positive | **+40** | | | 11/13 positive | **+38** | | | 12/14 positive | **+49** |
| *ours, matched data* | | | *DIAMOND arch.* | | | | *RSSM* | | | | *Transformer* | | |
| Breakout (random) | 8.3 | 13 | 35 | 0 | +13 | 44 | 29 | 6 | +38 | 92 | 1 | 6 | +87 |

## 5.2. A Rule Recorded Before the Outcomes

If events are not drawn readably, a model that draws them should reduce the return error, most where the baseline misses most events. As a repair, we replace our deterministic MSE predictor with a diffusion model trained on the same data [2, 24, 30], which samples each next observation instead of regressing it ("sampling" in figures and tables). A sampled future can differ from the realised one, so we evaluate event recall, return error, and plan ranking.

Before observing the remaining outcomes, we prespecified this rule: a baseline event recall at $\tau = 0$ below 0.50 predicts a return-error reduction larger than the declared floor of that game, and a recall at or above 0.50 predicts no such reduction. The threshold is a heuristic for the headroom left by missing events; floors and failure conditions were fixed at the same stage. Freeway had already been scored, so the prospective score covers Pong and the three unrun sampling arms. The MSE baseline gets more training, 30k updates against 8k, because at 8k updates it detects no return on Freeway and Pong; data, sequence length, batch size, learning rate, and seeds match (supplement).

# 6. Experiments

In this section, we carry out experiments to address the following questions: **Q1:** Is the return error concentrated at event transitions (Sec. 6.2)? **Q2:** Are missed events absent or only mistimed (Sec. 6.3)? **Q3:** Does a diffusion model repair them as recorded (Sec. 6.4)? **Q4:** Do the same transitions drive plan selection (Sec. 6.5)?

Table 2. **Pixel error does not find the event transitions.** Every row patches as many transitions as the event patch (seed 0). A random patch removes 17% and hits events at rate $q$ (median 16%).

| **Selection** | **Medians over the 44 counted pairs** | | | |
|---|---|---|---|---|
| Patch chosen by | Removed | $L$ | $L > 0$ | On events |
| Event transitions | 70% | +43.7 | 39/44 | 100% |
| Largest MSE | 15% | −1.1 | 17/44 | 19% |
| Largest 1−SSIM | 11% | −1.2 | 15/44 | 15% |
| Largest LPIPS | 15% | −0.8 | 17/44 | 16% |

## 6.1. Experimental Setup

We study 26 Atari games [5] with $64 \times 64$ observations and frameskip four, using Atari-100k [29] model checkpoints [2, 12, 36]. Across three released families and three locally trained random-policy Breakout models, 44 pairs clear reliability checks. Transfer evaluation adds 9 rows across 6 DeepMind Control tasks [51] and released iVideoGPT on robosuite [52, 60, 68].

Atari attribution uses 400 probes, each paired with a real trajectory under the same actions. Each frozen detector must reach a reward-event F1 of 0.90 on an independent collection with environment seeds disjoint from its training data. Terminations are read but not gated; replacing them removes a median 0.7% of the return error. The real reward at the last step stays in the all-real residual (median 1.5% of the error; supplement). Figure 6 adds matched-action RSSM [20] and Transformer predictions on MiniGrid [9] and Crafter [19].

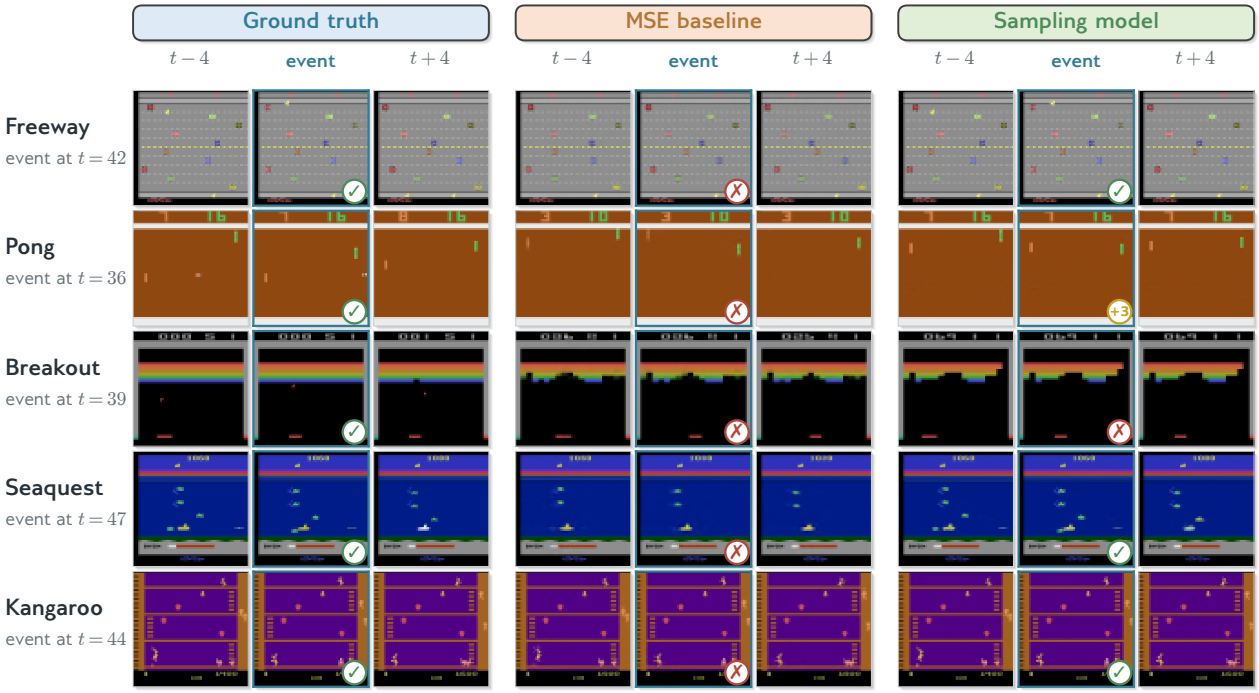


Figure 4. **Testing the proposed repair across five games.** Five games, each shown as ground truth, deterministic MSE prediction, and sampling prediction under the same actions. Blue borders mark the real-event time; marks show whether the repair study's detector reads the reward there (✓), a few steps off (+3) or not (×). Each row uses a validation window selected from real-event timing alone.

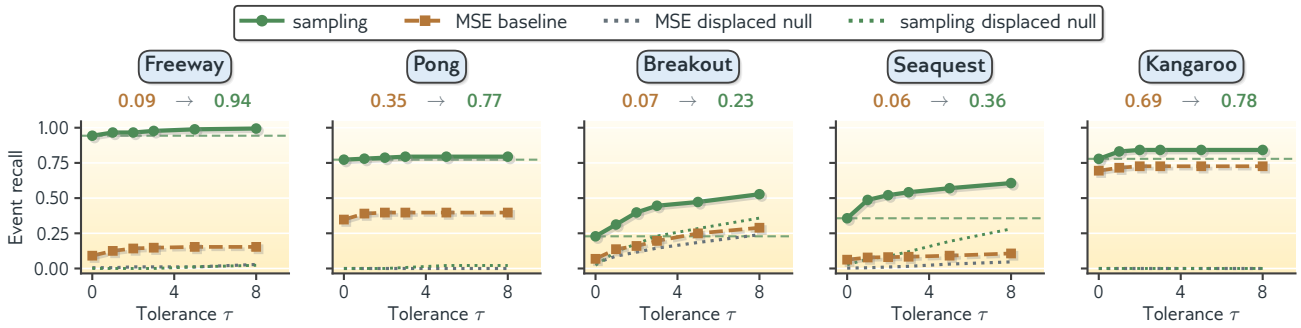


Figure 5. **The event deficit of the deterministic model persists under temporal tolerance.** Reward-event recall versus alignment tolerance across five games. Header values compare MSE → sampling at $\tau = 0$; green guides mark the sampling value. Each displaced null matches its model's predicted event rate. Kangaroo is the pre-specified no-effect case for return-error improvement.

## 6.2. A1: Return Error Concentrates at Events

**Main Results.** As shown in Table 1, the event patch has positive lift in 39 of 44 counted pairs (sign test $p = 1.4 \times 10^{-7}$), with a median of +44 points (probe-bootstrap 95% intervals: 38 above zero, 2 spanning it, 4 below). Each released family and every counted game except MsPacman has a positive median lift (per-game sign test, $p = 5.2 \times 10^{-4}$); three sampling seeds change lift by a median of 2.0 points, and other discounts and horizons keep 39–41 pairs positive (supplement). As a patch changes only the selected transitions, event transitions cause most of the return error.

**Robustness to the Reward Reader.** To rule out our detector, we use each model's reward head: all 25 pairs whose head reads real futures reliably stay positive (median +46; rank correlation 0.85 with our detector), as do 6 pairs from games it cannot read. Hiding the score display keeps all 6 Pong and Freeway pairs positive (median +64; supplement).

**Comparison with Pixel Error.** One may expect the error to sit in poorly rendered frames. Yet event frames are rendered as well as other frames (median ratio 0.85), and as shown in Table 2, the transitions with the largest MSE, SSIM or LPIPS error are events only about as often as random picks, so patching them removes less than random

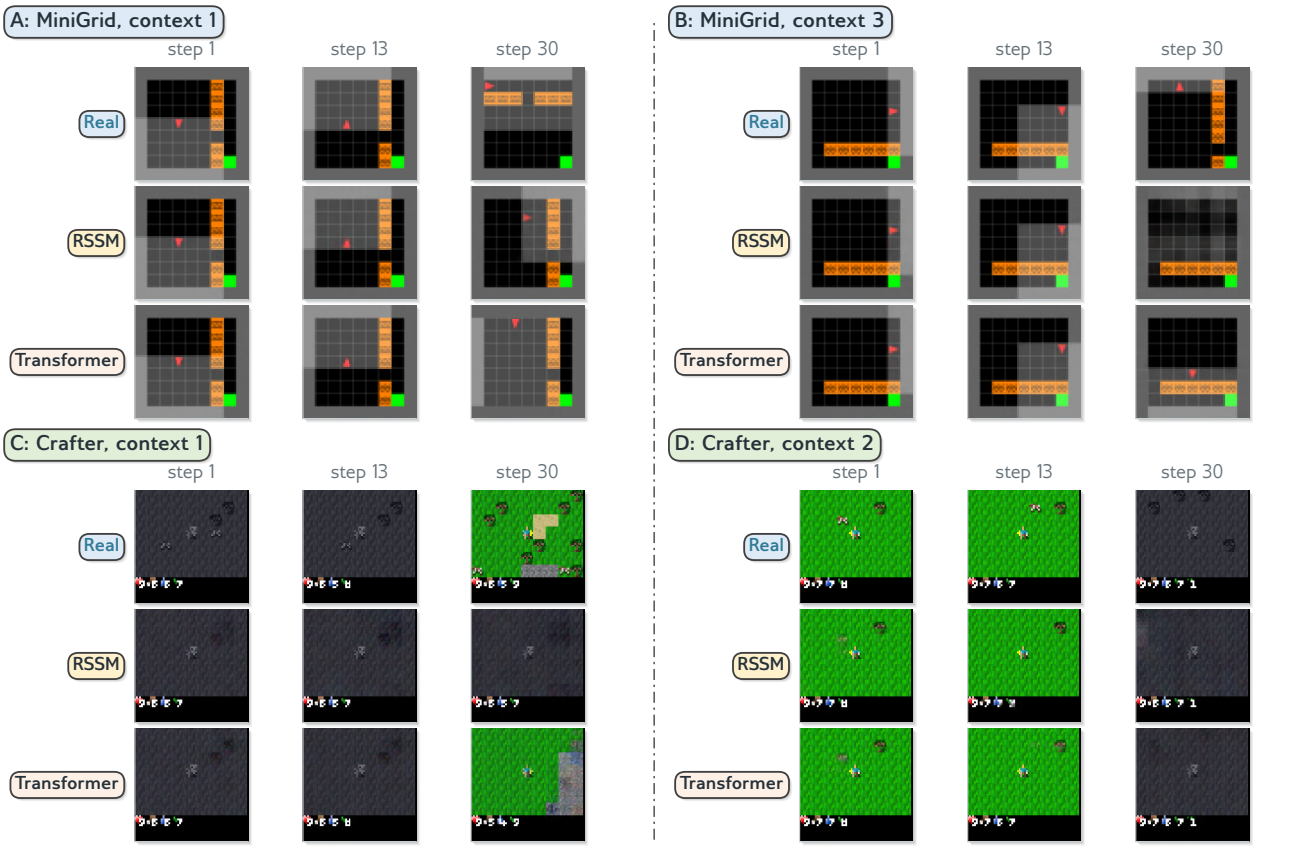


Figure 6. **Scene transitions in additional environments. Figures A–B** (MiniGrid) and **C–D** (Crafter) show real frames and saved RSSM/Transformer rollouts under the same recorded actions at steps 1, 13 and 30. Map orientation and night/day changes make trajectory divergence visible. Each rollout starts from 36 real frames; futures may cross episode boundaries. The supplement shows additional, nonoverlapping contexts.

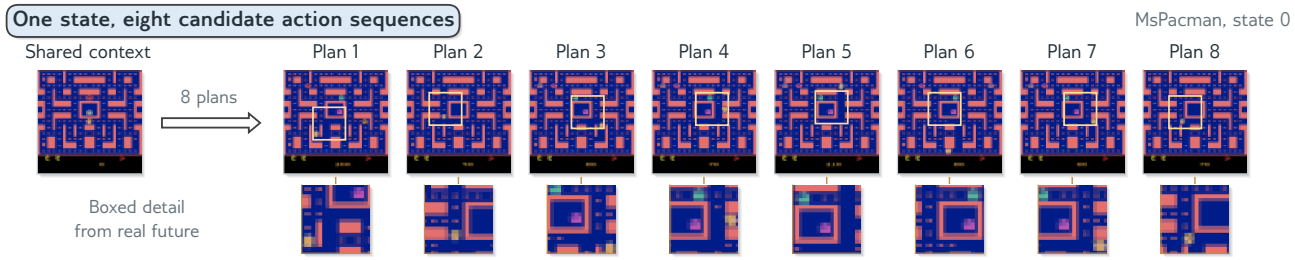


| | Plan 1 | Plan 2 | Plan 3 | Plan 4 | Plan 5 | Plan 6 | Plan 7 | Plan 8 |
|---|---|---|---|---|---|---|---|---|
| True return | 7.33 | 3.62 | 6.09 | 5.05 | **8.15** | 6.37 | 6.34 | 5.98 |
| DIAMOND (tie) | 0.00 | 0.00 | 0.00 | 0.00 | 0.00 | 0.00 | 0.00 | 0.00 |
| IRIS | 3.98 | 0.00 | 1.20 | 2.95 | **5.15** | 0.00 | 3.88 | 0.00 |
| Simulus | 3.71 | 2.77 | 11.98 | 3.17 | 10.54 | 2.35 | **12.05** | 0.00 |

Higher return is better; outlines mark the highest value in each row.

Figure 7. **One MsPacman state, eight possible plans.** The first state with nonconstant true returns shows all eight terminal-or-last real outcomes. Insets enlarge the boxed $20 \times 20$ regions with greatest pixel change from the context, excluding the HUD. Scores are unpatched; outlines mark row maxima. DIAMOND ties all plans; IRIS selects the true-best plan 5, Simulus plan 7 (true return 6.34). Complete planner statistics and different probe states are in the supplement.

Table 3. **Selected repair comparisons.** $\Delta E = E_{\text{MSE}} - E_{\text{sample}}$. Errors compare competent-seed MSE means (30k updates) with designated sampling runs (8k). Freeway is the development result; shaded rows are prospective tests. **Red** gains exceed declared seed floors; ✓ marks matched predictions. The complete original study, including the fourth prospective outcome, is in the supplement.

| **Evaluation setting** | | **Event recall ↑** | | **Return error** | | | **Decision margin** | | |
|---|---|---|---|---|---|---|---|---|---|
| Game | Split | MSE | Sample | MSE ↓ | Sample ↓ | $\Delta E$ ↑ | Floor | ×floor | Pred. / obs. |
| Freeway | development | 0.091 | 0.943 | 0.6048 | 0.1545 | **+0.4503** | 0.050 | 9.0 | helps ✓ |
| Pong | prospective | 0.348 | 0.773 | 0.2981 | 0.1421 | **+0.1560** | 0.079 | 2.0 | helps ✓ |
| Breakout | prospective | 0.069 | 0.229 | 1.4138 | 0.8822 | **+0.5316** | 0.146 | 3.6 | helps ✓ |
| Seaquest | prospective | 0.062 | 0.357 | 1.3490 | 0.5688 | **+0.7802** | 0.134 | 5.8 | helps ✓ |

patches. Pixel metrics thus cannot locate the return error.

**Effect of Coverage.** Random patches hit more events as $q$ grows, so a lift that only reflected sparse events would vanish. Yet the median lift is +40 above $q = 12\%$ (+59 below $6\%$), and all Krull pairs stay positive at $q = 58.5\%$.

**Failure Patterns.** With $R_{\text{ev}} = R(o, m^{\text{ev}})$ and the complementary patch's $R_{\text{non}} = R(o, \bar{m}^{\text{ev}})$, a pair is event-dominated if $R_{\text{ev}} - R_{\text{non}} > 0.20$, non-event-dominated if it is below $-0.20$, and unclassified otherwise; the counts are 33, 5 and 6. Simulus on Freeway is event-dominated: event patches remove $72\%$ and the complement $-16\%$. DIAMOND on Boxing is the opposite: it imagines a return of 6.67 against a true 2.31, and event patches remove $12\%$ and the complement $72\%$, so under our detector its excess return lies away from true scoring events. Under its own reward model, it is event-dominated (supplement). Most pairs thus fail by missing true events.

### 6.3. A2: Missed Events Persist Under Tolerance

We next ask whether missed events are merely mistimed (Sec. 5.1). To investigate this, we compare reward-event recall at alignment tolerance $\tau$ with a circularly displaced copy of the predictions, which keeps their event rate. If events were only mistimed, widening $\tau$ would recover most of them. As shown in Figs. 4 and 5, it recovers few: on Freeway, the MSE baseline recovers 0.091 of the events at $\tau = 0$ and 0.153 at $\tau = 8$, against 0.028 for the null, and across the five games at least 76% of its recall deficit remains at $\tau = 8$. Most missed events are thus absent rather than mistimed. DIAMOND on Freeway recovers more (0.767 to 0.932), and the diffusion model already reaches 0.943 at $\tau = 0$.

### 6.4. A3: A Diffusion Model Repairs as Predicted

If missing events cause the return error, a model that draws them should reduce it. As shown in Table 3, all four comparisons clear their floors, and the pre-specified rule scores 4/4, including 3/3 on the games whose detector passes the attribution gate. Three sampling seeds on Freeway and Pong give error ranges of 0.0356 and 0.0278, both below $\Delta E$.

**Ablation on Sampling.** A natural question is whether the gain comes from sampling. To isolate it, we train the same diffusion network with a deterministic squared-error objective (same data, 8k updates and seeds). It still has lower error than our MSE baseline on all five games and keeps Freeway's events (recall 0.95 against 0.94); its seed-mean error is lower than with sampling on Freeway (0.096 against 0.171) and within the larger seed span on the other four, while sampling recalls more events on Pong (0.80 against 0.69; supplement). The repair thus comes from the diffusion network, not from sampling or the training loss.

### 6.5. A4: Event Transitions Drive Plan Selection

Finally, if events decide the imagined return, restoring them should also improve plan choice beyond a matched random patch. To test this, we rank eight plans from each of 150 emulator states, as a sampling-based planner does [11, 20], and score pairwise accuracy and normalised top-1 regret; 7 of the twelve pairs, including the diffusion model, pass the reliability rules (supplement), and Fig. 7 shows one. The event patch improves pairwise accuracy in all 7 pairs, whereas a random patch changes it by a median of +0.019, and median top-1 regret falls from 0.565 to 0.141. Under pre-specified rules, 15 further games add 8 pairs from Freeway, CrazyClimber and Gopher, whose detectors also pass on plan states, and the event patch improves all 8 of them (median +0.156, random +0.012; supplement). On Pong, the diffusion model raises pairwise accuracy from 0.5000 to 0.6110 (random-weighted control 0.5007). The transitions carrying the return error thus also decide plan selection.

## 7. Cross-Task Transfer and Discussion

Beyond Atari, we apply the diagnostic to six DeepMind Control tasks and robosuite, and compare matched-action rollouts on MiniGrid and Crafter (Fig. 6). On the control tasks of Fig. 3, the three comparisons whose detector passes yield +54/+54 points of lift on sparse cartpole across two architectures and +8 on ball-in-cup (supplement).

Concurrent work finds ball disappearance and invalid ball–paddle interactions in five frozen Pong world models [34]; we locate the return error across 26 games.

## 8. Conclusion

In this paper, we introduce a per-transition substitution diagnostic that locates where the imagined return of a world model goes wrong. In 39 of 44 reliably measured model–collection pairs on Atari-100k, the return error concentrates at reward and termination transitions that pixel metrics do not single out: world models get the return wrong at events, in frames that look no worse than the rest. A rule recorded in advance predicts where a diffusion model class helps, and event patches improve plan ranking in every counted pair; the same concentration appears on control tasks with rare reward events. These results suggest evaluating world models for planning by where their return error occurs, in addition to how well they render frames.

## References


[1] Krzysztof Adamkiewicz, Brian B Moser, Stanislav Frolov, Tobias Christian Nauen, Federico Raue, and Andreas Dengel. When pretty isn't useful: Investigating why modern text-to-image models fail as reliable training data generators. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 36660–36669, 2026. 2

[2] Eloi Alonso, Adam Jelley, Vincent Micheli, Anssi Kanervisto, Amos Storkey, Tim Pearce, and François Fleuret. Diffusion for world modeling: Visual details matter in atari. *Advances in Neural Information Processing Systems*, 37: 58757–58791, 2024. 1, 2, 5

[3] Jose A Arjona-Medina, Michael Gillhofer, Michael Widrich, Thomas Unterthiner, Johannes Brandstetter, and Sepp Hochreiter. Rudder: Return decomposition for delayed rewards. *Advances in Neural Information Processing Systems*, 32, 2019. 3

[4] Akanksha Atrey, Kaleigh Clary, and David Jensen. Exploratory not explanatory: Counterfactual analysis of saliency maps for deep reinforcement learning. *arXiv preprint arXiv:1912.05743*, 2019. 3

[5] Marc G Bellemare, Yavar Naddaf, Joel Veness, and Michael Bowling. The arcade learning environment: An evaluation platform for general agents. *Journal of artificial intelligence research*, 47:253–279, 2013. 5

[6] Jake Bruce, Michael D Dennis, Ashley Edwards, Jack Parker-Holder, Yuge Shi, Edward Hughes, Matthew Lai, Aditi Mavalankar, Richie Steigerwald, Chris Apps, et al. Genie: Generative interactive environments. In *Forty-first international conference on machine learning*, 2024. 2

[7] Wei-Di Chang, Mikael Henaff, Brandon Amos, Gregory Dudek, and Scott Fujimoto. The surprising difficulty of search in model-based reinforcement learning. *arXiv preprint arXiv:2601.21306*, 2026. 1, 3

[8] Zelei Cheng, Xian Wu, Jiahao Yu, Wenhai Sun, Wenbo Guo, and Xinyu Xing. Statemask: Explaining deep reinforcement learning through state mask. *Advances in Neural Information Processing Systems*, 36:62457–62487, 2023. 3

[9] Maxime Chevalier-Boisvert, Bolun Dai, Mark Towers, Rodrigo Perez-Vicente, Lucas Willems, Salem Lahlou, Suman Pal, Pablo Samuel Castro, and Jordan Terry. Minigrid & miniworld: Modular & customizable reinforcement learning environments for goal-oriented tasks. *Advances in neural information processing systems*, 36:73383–73394, 2023. 5

[10] Silvia Chiappa, Sébastien Racaniere, Daan Wierstra, and Shakir Mohamed. Recurrent environment simulators. *arXiv preprint arXiv:1704.02254*, 2017. 2

[11] Kurtland Chua, Roberto Calandra, Rowan McAllister, and Sergey Levine. Deep reinforcement learning in a handful of trials using probabilistic dynamics models. *Advances in neural information processing systems*, 31, 2018. 1, 8

[12] Lior Cohen, Kaixin Wang, Bingyi Kang, Uri Gadot, and Shie Mannor. Simulus: Combining improvements in sample-efficient world model agents. *arXiv preprint arXiv:2502.11537*, 2025. 1, 2, 5

[13] Amir-massoud Farahmand, Andre Barreto, and Daniel Nikovski. Value-aware loss function for model-based reinforcement learning. In *Artificial Intelligence and Statistics*, pages 1486–1494. PMLR, 2017. 2

[14] Ruth C Fong and Andrea Vedaldi. Interpretable explanations of black boxes by meaningful perturbation. In *2017 IEEE international conference on computer vision (ICCV)*, pages 3449–3457. IEEE, 2017. 3

[15] Samuel Greydanus, Anurag Koul, Jonathan Dodge, and Alan Fern. Visualizing and understanding atari agents. In *International conference on machine learning*, pages 1792–1801. PMLR, 2018. 3

[16] Christopher Grimm, André Barreto, Satinder Singh, and David Silver. The value equivalence principle for model-based reinforcement learning. *Advances in neural information processing systems*, 33:5541–5552, 2020. 2

[17] Wenbo Guo, Xian Wu, Usmann Khan, and Xinyu Xing. Edge: Explaining deep reinforcement learning policies. *Advances in Neural Information Processing Systems*, 34: 12222–12236, 2021. 3

[18] David Ha and Jürgen Schmidhuber. Recurrent world models facilitate policy evolution. *Advances in neural information processing systems*, 31, 2018. 1

[19] Danijar Hafner. Benchmarking the spectrum of agent capabilities. *arXiv preprint arXiv:2109.06780*, 2021. 5

[20] Danijar Hafner, Timothy Lillicrap, Ian Fischer, Ruben Villegas, David Ha, Honglak Lee, and James Davidson. Learning latent dynamics for planning from pixels. In *International conference on machine learning*, pages 2555–2565. PMLR, 2019. 1, 5, 8

[21] Danijar Hafner, Timothy Lillicrap, Mohammad Norouzi, and Jimmy Ba. Mastering atari with discrete world models. *arXiv preprint arXiv:2010.02193*, 2020. 2

[22] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640(8059):647–653, 2025. 1

[23] Nick Hansen, Hao Su, and Xiaolong Wang. Td-mpc2: Scalable, robust world models for continuous control. In *International Conference on Learning Representations*, pages 47376–47405, 2024. 1

[24] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020. 5

[25] Sara Hooker, Dumitru Erhan, Pieter-Jan Kindermans, and Been Kim. A benchmark for interpretability methods in deep neural networks. *Advances in neural information processing systems*, 32, 2019. 3

[26] Michael Janner, Justin Fu, Marvin Zhang, and Sergey Levine. When to trust your model: Model-based policy optimization. *Advances in neural information processing systems*, 32, 2019. 1

[27] Dinesh Jayaraman, Frederik Ebert, Alexei A Efros, and Sergey Levine. Time-agnostic prediction: Predicting predictable video frames. *arXiv preprint arXiv:1808.07784*, 2018. 3

[28] Sonia Joseph, Quentin Garrido, Randall Balestriero, Matthew Kowal, Thomas Fel, Shahab Bakhtiari, Blake Richards, and Mike Rabbat. Interpreting physics in video world models. *arXiv preprint arXiv:2602.07050*, 2026. 3

[29] Lukasz Kaiser, Mohammad Babaeizadeh, Piotr Milos, Blazej Osinski, Roy H Campbell, Konrad Czechowski, Dumitru Erhan, Chelsea Finn, Piotr Kozakowski, Sergey Levine, et al. Model-based reinforcement learning for atari. *arXiv preprint arXiv:1903.00374*, 2019. 2, 5

[30] Tero Karras, Miika Aittala, Timo Aila, and Samuli Laine. Elucidating the design space of diffusion-based generative models. *Advances in neural information processing systems*, 35:26565–26577, 2022. 5

[31] Nathan Lambert, Brandon Amos, Omry Yadan, and Roberto Calandra. Objective mismatch in model-based reinforcement learning. *arXiv preprint arXiv:2002.04523*, 2020. 2

[32] Mingxuan Li, Kai-Zhan Lee, Michael Dennis, and Elias Bareinboim. Counterfactual debugging the world model transfer gap. 3

[33] Ao Liang, Lingdong Kong, Tianyi Yan, Hongsi Liu, Yu Yang, Ziqi Huang, Wei Yin, Jialong Zuo, Yixuan Hu, Dekai Zhu, et al. Worldlens: Full-spectrum evaluations of driving world models in real world. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 36385–36399, 2026. 1, 2

[34] Yukuan Lu, Zaishuo Xia, Weyl Lu, and Yubei Chen. Improving weak world models behind strong agents in atari pong, 2026. 3, 8

[35] Michael Mathieu, Camille Couprie, and Yann LeCun. Deep multi-scale video prediction beyond mean square error. *arXiv preprint arXiv:1511.05440*, 2015. 3

[36] Vincent Micheli, Eloi Alonso, and François Fleuret. Transformers are sample-efficient world learners. In *The Eleventh International Conference on Learning Representations (ICLR)*, 2023. 1, 2, 5

[37] Vincent Micheli, Eloi Alonso, and François Fleuret. Efficient world models with context-aware tokenization. *arXiv preprint arXiv:2406.19320*, 2024. 2

[38] Saman Motamed, Laura Culp, Kevin Swersky, Priyank Jaini, and Robert Geirhos. Do generative video models understand physical principles? In *2026 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 948–958. IEEE, 2026. 2

[39] Suraj Nair, Silvio Savarese, and Chelsea Finn. Goal-aware prediction: Learning to model what matters. In *International conference on machine learning*, pages 7207–7219. PMLR, 2020. 3

[40] Junhyuk Oh, Xiaoxiao Guo, Honglak Lee, Richard L Lewis, and Satinder Singh. Action-conditional video prediction using deep networks in atari games. *Advances in neural information processing systems*, 28, 2015. 2

[41] Daniel Palenicek, Michael Lutter, Joao Carvalho, and Jan Peters. Diminishing return of value expansion methods in model-based reinforcement learning. *arXiv preprint arXiv:2303.03955*, 2023. 1, 3

[42] Zhao-Han Peng, Shaohui Li, Zhi Li, Shulan Ruan, Yu Liu, and You He. From observations to events: Event-aware world model for reinforcement learning. *arXiv preprint arXiv:2601.19336*, 2026. 3

[43] Karl Pertsch, Oleh Rybkin, Jingyun Yang, Shenghao Zhou, Konstantinos Derpanis, Kostas Daniilidis, Joseph Lim, and Andrew Jaegle. Keyframing the future: Keyframe discovery for visual prediction and planning. In *Learning for Dynamics and Control*, pages 969–979. PMLR, 2020. 3

[44] Vitali Petsiuk, Abir Das, and Kate Saenko. Rise: Randomized input sampling for explanation of black-box models. *arXiv preprint arXiv:1806.07421*, 2018. 3

[45] Francois Porcher, Nicolas Carion, Karteek Alahari, and Shizhe Chen. Flow matching in feature space for stochastic world modeling. *arXiv preprint arXiv:2606.29059*, 2026. 3

[46] Jan Robine, Marc Höftmann, Tobias Uelwer, and Stefan Harmeling. Transformer-based world models are happy with 100k interactions. *arXiv preprint arXiv:2303.07109*, 2023. 2

[47] Finn Rasmus Schäfer, Korbinian Moller, Yuan Gao, Christian Oefinger, Sebastian Schmidt, and Johannes Betz. Imagined rollouts are kinematic, not dynamic: A diagnosis of long-horizon world-model failure. *arXiv preprint arXiv:2607.05966*, 2026. 3

[48] Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, et al. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588(7839):604–609, 2020. 1

[49] Anya Sims, Cong Lu, Jakob N Foerster, and Yee W Teh. The edge-of-reach problem in offline model-based reinforcement learning. *Advances in Neural Information Processing Systems*, 37:63029–63056, 2024. 1, 3

[50] Erik Talvitie. Model regularization for stable sample rollouts. In *UAI*, pages 780–789, 2014. 3

[51] Yuval Tassa, Yotam Doron, Alistair Muldal, Tom Erez, Yazhe Li, Diego de Las Casas, David Budden, Abbas Abdolmaleki, Josh Merel, Andrew Lefrancq, et al. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018. 5

[52] Stephen Tian, Chelsea Finn, and Jiajun Wu. A control-centric benchmark for video prediction. *arXiv preprint arXiv:2304.13723*, 2023. 1, 5

[53] Thomas Unterthiner, Sjoerd Van Steenkiste, Karol Kurach, Raphael Marinier, Marcin Michalski, and Sylvain Gelly. Towards accurate generative models of video: A new metric & challenges. *arXiv preprint arXiv:1812.01717*, 2018. 1

[54] Dani Valevski, Yaniv Leviathan, Moab Arar, and Shlomi Fruchter. Diffusion models are real-time game engines. In *International Conference on Learning Representations*, pages 73754–73776, 2025. 2

[55] Arun Venkatraman, Martial Hebert, and J Bagnell. Improving multi-step prediction of learned time series models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2015. 3

[56] Linhao Wang, Yiyan Fan, and Dongjin Huang. Not all errors matter: Decision-relevant prediction error predicts planning quality. *arXiv preprint arXiv:2609.32322*, 2026. 2

[57] Yahong Wang, Juncheng Wu, Zhangkai Ni, Longzhen Yang, Yihang Liu, Chengmei Yang, Ying Wen, Lianghua He, Xianfeng Tang, Hui Liu, et al. When token pruning is worse than random: Understanding visual token information in vllms. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 31910–31919, 2026. 2

[58] Zhaohui Wang. Learning from world feedback: Why model uncertainty fails as a risk signal in model-based rl. *arXiv preprint arXiv:2607.16591*, 2026. 2

[59] Zhou Wang, Alan C Bovik, Hamid R Sheikh, and Eero P Simoncelli. Image quality assessment: from error visibility to structural similarity. *IEEE transactions on image processing*, 13(4):600–612, 2004. 1

[60] Jialong Wu, Shaofeng Yin, Ningya Feng, Xu He, Dong Li, Jianye Hao, and Mingsheng Long. ivideogpt: Interactive videogpts are scalable world models. *Advances in Neural Information Processing Systems*, 37:68082–68119, 2024. 1, 2, 5

[61] Zaishuo Xia, Yukuan Lu, Xinyi Li, Yifan Xu, and Yubei Chen. Cloning deterministic worlds: the critical role of latent geometry in long-horizon world models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 32602–32612, 2026. 3

[62] Mengjiao Yang, Yilun Du, Kamyar Ghasemipour, Jonathan Tompson, Dale Schuurmans, and Pieter Abbeel. Learning interactive real-world simulators. *arXiv preprint arXiv:2310.06114*, 1(2):6, 2023. 2

[63] Matthew D Zeiler and Rob Fergus. Visualizing and understanding convolutional networks. In *European conference on computer vision*, pages 818–833. Springer, 2014. 3

[64] Jiahan Zhang, Muqing Jiang, Nanru Dai, Taiming Lu, Arda Uzunoglu, Shunchi Zhang, Yana Wei, Jiahao Wang, Vishal Patel, Paul Liang, et al. World-in-world: World models in a closed-loop world. In *International Conference on Learning Representations*, pages 55660–55699, 2026. 1, 2

[65] Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *2018 IEEE/CVF conference on computer vision and pattern recognition*, pages 586–595. IEEE, 2018. 1

[66] Weipu Zhang, Gang Wang, Jian Sun, Yetian Yuan, and Gao Huang. Storm: Efficient stochastic transformer based world models for reinforcement learning. *Advances in Neural Information Processing Systems*, 36:27147–27166, 2023. 2

[67] Chenran Zhao, Dianxi Shi, Yaowen Zhang, Chunping Qiu, and Shaowu Yang. Transferable delay-aware reinforcement learning via implicit causal graph modeling. *arXiv preprint arXiv:2605.12312*, 2026. 3

[68] Yuke Zhu, Josiah Wong, Ajay Mandlekar, Roberto Martín-Martín, Abhishek Joshi, Kevin Lin, Abhiram Maddukuri, Soroush Nasiriany, and Yifeng Zhu. robosuite: A modular simulation framework and benchmark for robot learning. *arXiv preprint arXiv:2009.12293*, 2020. 5