arXiv:2610.30004v1 [cs.NI] 21 Oct 2026

Every Alarm Has a Cause: An Open Generator of Telecom Alarm Storms with Exact Root Causes, Bound to NORIA-O, RFC 8632 and TMF642

Madhulatha Mandarapu* Sandeep Kunkunuru†

VaidhyaMegha Private Limited, India
https://samyama.ai/

October 2026

Abstract

Alarm root-cause analysis is studied on data that cannot be shared. Operator logs are confidential, so what the field publishes is a causal graph between alarm *types* rather than per-incident causes, a fixed scenario set, gated, or never released. Two recent benchmarks carry per-incident root causes at scale; neither states a download location or a licence.

We give a generator rather than a dataset. It runs on real backbone topologies from the Internet Topology Zoo and records the exact root cause *and the propagation path* of every alarm — a by-product of generation, not a labelling pass. Every emitted field binds to a term in NORIA-O, RFC 8632, RFC 8345 or TMF642, and an unknown term fails the run. We ship fixed configurations, splits, a scorer and nine baselines.

A one-line rule — rank the device that raised the earliest critical alarm — takes top-1 at 0.650 against 0.575 for our own temporal method, and wins 13 of 15 runs over three backbones and five seeds. An LLM given the same view reaches 0.600, the same model allowed to query the graph 0.625, and the rule wins on real operator alarms too, at 0.773: a property of alarm streams, not of our generator. One fault class separates the table: when a site loses power its PoP reports nothing, so every method that ranks alarmed devices scores 0.000 or near it, while the two that walk the topology to a silent candidate take 1.000 and 0.833.

1 Introduction

A network operations centre sees a storm of alarms and must name the one resource that caused them. Research on this problem is evaluated on data nobody else can obtain, because an operator's alarm log carries its customers, its topology and its failures. The consequence is not that the work is wrong but that it is incomparable: two methods reporting 0.7 on two private logs have not been compared.

The resources the field has published do not close that gap, and Section 2 goes through them one at a time. Briefly: the most-used open alarm data gives a causal graph between alarm *types*, with no incident identifier and no topology [10, 7]; the benchmarks that do carry a cause per incident are either announced without a download [5, 6] or gated with terms that contradict their own licence

*madhulatha@samyama.ai
†sandeep@samyama.ai

tag [8]; and the open RCA benchmarks that are genuinely downloadable are microservice systems, not telecom alarms [3, 2]. One paper advertises its data at a repository that contains a single 378-byte README [15], and is still cited as a data source five years later.

A generator, not a dataset. We take the other route. A generator sidesteps confidentiality in a way no release process can, and it gives something a released log cannot: the ground truth is not an annotation added afterwards by an expert who was not there, it is what the generator did. That distinction matters most for the *propagation path*. An annotator can reasonably name the resource that caused an incident; reconstructing which alarm caused which other alarm, through what, is a different task, and the largest prior benchmark states plainly that it does not have one [6].

The claim. We do not claim the first per-incident root causes for telecom alarms. Two 2025–2026 benchmarks claim that, and at a realism we cannot match — real production networks, expert annotation, hundreds of thousands of alarms [6, 5]. Our claim is narrower:

> No *downloadable, openly licensed* resource gives per-incident root causes *and propagation paths* for telecom alarms at scale, and no open generator exists.

Both halves were checked against the sources rather than assumed, on 2026-10-21, and the check is published with the paper.

Contributions.

1. A generator of telecom alarm storms that records the root cause and the full propagation path of every alarm, parameterised by topology and seeded (Section 3).
2. A binding of every emitted type and field to NORIA-O, RFC 8632, RFC 8345 or TMF642, enforced at generation time (Section 4). NORIA-O itself aligns to ITU-T X.733/X.721 and TMF638 and references neither RFC 8632 nor TMF642; the binding to those two is ours, not inherited.
3. A benchmark — fixed configurations, splits, a scoring script, and baselines we did not design (Sections 5–7).
4. Nine scored baselines — four one-line rules, a fitted topological Hawkes process, personalised PageRank, a language model given the same view, the same model allowed to query the graph, and a control that isolates what the winning rule actually uses: severity, not time (Section 8).
5. A realism study against public real alarm data, including the test that decides whether the resource is useful: the ranking of methods on generated data is the ranking on real data (Kendall $\tau = 1.0$ over four rules on 4,539 real bursts), while three realism metrics show the generated stream matches real alarms on type vocabulary and chaining and diverges on time, arriving in burstier volleys.

2 Related work

Every row of Table 1 was read from the source on 2026-10-21 rather than from a secondary description, and the record of what was checked ships with the paper. Two things we expected to be true were not, and both are stated here rather than quietly corrected.

Resource	Obtainable	Labels	Topology
PCIC 2021 [7]	yes, MIT	alarm-*type* graph	none
Influence-based [15]	no	per-transaction (unreleased)	none
TN-RCA530 [5]	announced only	per-incident	per-scenario KG
TeleRCA [6]	announced only	per-incident, no path	none published
TeleLogs [8]	gated	8 fixed categories	none
RCAEval [3]	yes	per-case service	microservice
LEMMA-RCA [2]	yes, ND	per-case pods	microservice
OpenRCA [4]	yes	per-case	cloud telemetry
This work	**yes, open**	**per-incident + path**	**RFC 8345**

Table 1: Public resources for alarm root-cause analysis. "Announced only" means the paper states the resource is released and contains no repository, download or licence.

The closest prior work is TeleRCA, and it is better than ours on realism. TeleRCA [6] gives 10,922 production incidents, 485,681 alarm events and 5,621 network resources, each incident carrying an expert annotation naming the root-cause resource. That is a real network and we do not have one. We differ on three things its own paper establishes: it states twice that it has *no ground-truth propagation graph*, and reports no graph-reconstruction metric in consequence; it publishes no topology or inventory; and although it says four times that it is released, the paper contains no repository, no download and no licence. TN-RCA530 [5] is closer to us on topology — each of its 530 scenarios is a knowledge graph — and is likewise announced without a download or a dataset licence.

The most-used open alarm data labels types, not incidents. The Huawei PCIC 2021 data [10] is live under MIT, though not at the URL its paper prints [7]: 34,838 alarms, 18 alarm types, 55 devices over 243 days. Its ground truth is an 18×18 adjacency between alarm *types*. There is no incident identifier, no storm grouping, no root-cause column and no topology. It is a good realism reference and not a source of per-incident labels. The nearest open artefact to such labels is a 49-row prior shipped with a later competition sample, whose "root cause" is itself an alarm type.

Availability is the field's recurring failure, and it is the argument for a generator. Zhang et al. [15] states that its data and code are available at a named repository; that repository holds one 378-byte README saying the release is going through internal procedures. The paper is still cited as a data source. TeleRCA and TN-RCA530 repeat the pattern. A generator is not a better dataset — it is a way of not depending on someone else's release process.

Open RCA benchmarks exist, for other systems. RCAEval [3] and LEMMA-RCA [2] are downloadable and well constructed, and are microservice and water-treatment systems rather than telecom alarms; LEMMA-RCA carries CC BY-ND on its own site, which would block derived works. OpenRCA [4] has a domain named *Telecom*, which invites confusion: it is cloud observability telemetry — KPIs, traces and logs of enterprise systems — not NOC alarms, and the repository states no licence for the data.

No generator on the Internet Topology Zoo. We found none that emits alarms with labelled root causes. The closest prior work describes an optical alarm-propagation simulator with an ITU-T

G.798 rule engine producing alarm correlation trees; no code repository was found.

3 The generator

Topologies are real backbones taken from the Internet Topology Zoo [9], which is CC BY 4.0. A backbone's GraphML gives the PoPs and the links between them, and those are the only real objects in the output; everything below carries `synthetic: true`.

Inventory and services. Each PoP gets two core routers, two or three aggregation routers dual-homed to both, and six to twelve access or radio sites homed to aggregation, then cards, then ports, then cabling. Backbone links land on core-router ports, so the real topology constrains the synthetic one rather than sitting beside it. Services follow the inventory: mobile data per radio site, broadband per access site, and enterprise VPNs across PoPs with a primary and a backup path. Customers carry a market segment and nothing else — there is no personal data to anonymise.

Faults and propagation. There are 7 fault types: card, device, port and software failure, fibre cut, site power loss, and a configuration change, which is always preceded by a `Change` whose outcome is *failed*. What follows from a fault is computed from the topology, not scripted: a dead card takes its ports down, the far end of each down port loses signal, a breadth-first search from the NOC's own PoP over surviving links decides which devices still have a path home, and a service whose endpoint is unreachable is out. Redundancy is what stops a failure spreading, so dual homing and two core routers per PoP are what make most faults local and a few of them storms.

Alarms. Each of those states raises an alarm on the resource that would report it, with a per-mechanism severity and a per-mechanism emission probability: the root alarm when the root can still speak, loss of signal at the far ends, an unreachable-device alarm after the next poll, a service alarm per lost service, neighbour-cell link failures on live radio sites near a dead one, and congestion where traffic reroutes. The dataset uses 4 X.733 event types and 21 NORIA-O probable causes over 4 RFC 8632 severities.

What makes it hard is in the config, not in a script. The root resource often raises no alarm at all: in 50 of 70 faults it does, and in the rest the method has to infer the root from what its neighbours said. A site power loss never raises one — the PoP it kills cannot report itself, which is also why the fault model's fifth event type, `environmentalAlarm`, appears nowhere in the data. On top of that, alarms go missing, flap through raise/clear cycles, duplicate, and carry per-device clock skew drawn once per device, so the order alarms arrive in is not the order the faults happened in. Background alarms arrive throughout, unrelated to any fault. Two faults can overlap inside one PoP. Every one of those is a knob with a number in the config, which is what lets the ablation in Section 8 say which of them the methods are actually reading.

Section 8.3 removes each of these in turn and measures what it was worth. Alarms that never arrive are worth the most by a wide margin; mis-grouping is worth the least, and was worth nothing at all before the arrival times were fixed.

The distribution, not the mean. Alarms per fault have a median of 3 and a 90th percentile of 41.1, with a maximum of 103. Most faults are quiet and a few are storms, so a method evaluated at the mean is evaluated at neither end.

What is recorded. For every injected fault the generator writes the fault, its onset and repair, the resources it hit, the alarms it caused and the window it spans. This is the ground truth, and it is a by-product of generation rather than a labelling pass.

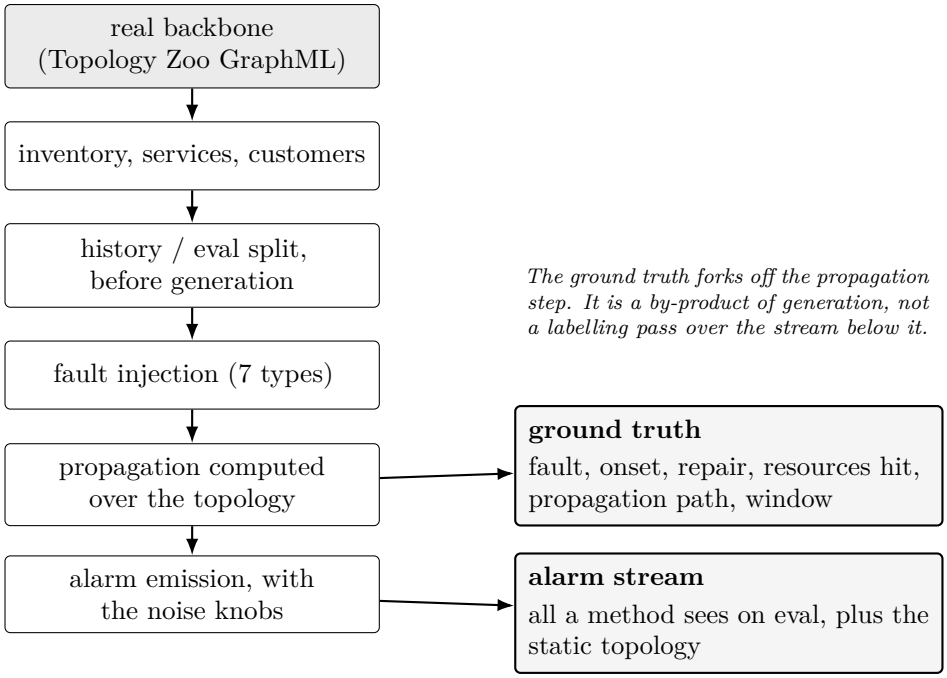

Figure 1: How a run is produced. Only the top box is real; everything below it carries `synthetic: true`. The two outputs are written by the same pass, which is why every alarm's cause and path are exact rather than inferred.

Splits. Scenarios are split into *history* and *eval* before generation, not after, and each scenario owns its own window, so the two never interleave in time. A method may use everything about history — its alarms, its incidents, its resolved tickets, its changes, and the ground truth — to fit whatever it likes. On eval it sees the static topology and the incident's alarms, and nothing else. The GÉANT small configuration has 30 history faults and 40 eval faults in three scenario shapes: a single fault, a storm, and two faults overlapping inside one PoP.

The split is only a split if a method cannot see across it, which is a property of the method and not of the data. The one baseline here that fits anything — the topological Hawkes process in Section 7 — is given a hard cutoff at the onset of the first eval incident, and a test asserts that every alarm it was trained on falls below that cutoff and that the cutoff excludes a non-zero number of them. A cutoff that excludes nothing proves nothing.

4 Standards binding

Every emitted type and field resolves to a term in one of four published vocabularies: NORIA-O [11], the alarm-management YANG model of RFC 8632 [14], the network-topology model of RFC 8345 [1], and the TM Forum alarm API TMF642 [13]. Generation fails on an unknown term rather than emitting it. The binding is a build step rather than a convention. Each vocabulary is converted to a graph of its own terms — an OWL/SKOS pass for NORIA-O, a YANG pass for the two RFCs,

an OpenAPI pass for TMF642 — and a catalog entry records where each build came from. The generator then declares what it needs: for each node label, the ontology classes it claims to be, and for each enumerated field, the vocabulary and the owning term the values must come from. At generation time every label and every enumeration value is resolved against those builds, 11 labels and 4 vocabularies in this configuration, and an unresolved term raises rather than being emitted. The release carries `binding.json`: the terms used, and the build files they came from with their sha256, so a reader can check the vocabulary a value claims without trusting the paper.

The grounding also travels in the graph itself, as one `SchemaLabel` node per label with a `GROUNDED_IN` edge to each ontology concept, rather than as a table in a README. A query can therefore ask what an `Alarm` is in NORIA-O's terms.

What is ours and what is inherited. NORIA-O's own alarm vocabulary aligns to ITU-T X.733/X.721 and to TMF638 Service Inventory; neither it nor its IETF draft [12] references RFC 8632 or TMF642. Binding to those two is this paper's contribution rather than an alignment we inherited, and we state it so that a reader who opens the ontology does not have to discover it.

Licence. NORIA-O is BSD-4-Clause, which carries an advertising clause and a no-endorsement clause. What we ship from it is term names — 6 class IRIs and 21 of the 173 concepts in its probable-cause scheme — and not the ontology files, so the clauses requiring a redistribution to retain the notice do not reach the dataset; the advertising clause does, because this paper mentions the use, and satisfying it is an acknowledgement sentence, which the release carries. RFC 8632 and RFC 8345 are Simplified BSD; TMF642 is Apache-2.0; Topology Zoo is CC BY 4.0.

Acknowledgement required by NORIA-O's licence. This product includes software developed by Orange. That sentence is the acknowledgement NORIA-O's BSD-4-Clause licence requires of material describing use of it. It is a statement about the ontology and not about any commercial relationship, and Orange does not endorse this work.

5 The benchmark

The task. One incident at a time, rank the resources that could hold its root cause, best first. An incident is the NOC's grouping of alarms — imperfect on purpose, since some alarms are attached to the wrong one — so the method is given the storm an operator would actually see rather than a clean partition.

Scoring. A ranking is correct at k if one of its first k entries is the device of any resource the fault was at. Two details matter and both are properties of the faults rather than choices: a fibre cut has two ends and either is correct, and a site power loss names the PoP, so any device hosted in it counts. We report $k = 1$ and $k = 3$: top-1 is what an automated remediation step needs, and top-3 is what an operator who will check three devices needs. They are different questions and a method can be good at one and not the other.

An eval incident that raises no alarm anywhere stays in the denominator and is charged a miss to every method. There are two such incidents in this configuration, and dropping them — which both of our scorers did until 2026-10-21 — reports a rate over 38 trials under a heading that says 40 (Section 9).

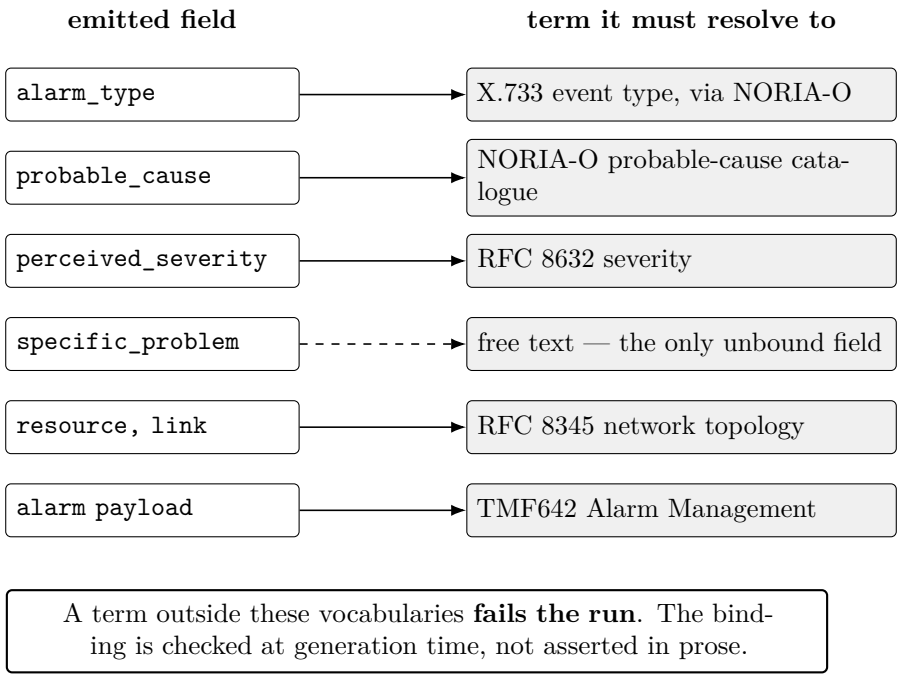

Figure 2: What each emitted field binds to. NORIA-O references neither RFC 8632 nor TMF642; that binding is ours. `specific_problem` is deliberately free text, and Section 5 shows why it had to be withheld from the methods.

What a method sees. The static topology and inventory, and the alarms in the incident window. Not the fault, not the propagation path, not the repair.

What a method must not see, and how we check. Ground truth travels with the data — `Fault` nodes, their `AT` edges to the true root resources, and a `CAUSED` edge to each alarm they produced. That is deliberate and labelled, and a method may not read it. The risk a labelled boundary does not cover is ground truth reaching an *ordinary* property, where nothing marks it, and this dataset shipped exactly that: every alarm carried its role in the fault — root, neighbour, background — in the X.733 `specificProblem` field. Reading that one field scored 0.700 top-1 on the dataset as it then stood — level with the best baseline that dataset had — while ranking nothing.

So the boundary is now measured rather than asserted. We score every (property, value) pair that appears on an alarm in the released files as a one-line method — rank the devices whose alarm carries that value, earliest first — in the benchmark's own units: 112 pairs over the 40 eval incidents. Table 2 gives the fields that matter. A property is *supposed* to correlate with the fault; what must not exist is a value that names the root almost every time it appears *and* appears in most incidents, which is the pair of conditions that made the old field free. The role now travels only on `CAUSED.via` and in `ground_truth.jsonl`, and `specificProblem` is a function of `probableCause` alone, so it cannot carry side information; a test asserts the two score identically.

Alarm property	best value	Top-1	Prec.	Inc.
`specific_problem` *(before)*	`root`	0.700	1.000	28
`perceived_severity`	`critical`	0.625	0.893	28
`specific_problem`	`Transmit module failure`	0.225	1.000	9
`probable_cause`	`transmitterFailure`	0.225	1.000	9
`alarm_type`	`communicationsAlarm`	0.225	0.250	36

Table 2: Each alarm field's best single value read as a one-line RCA rule, GÉANT small, $n = 40$. *Prec.* is the share of the incidents where such an alarm exists in which its earliest one sits on a true root; *Inc.* is how many of the 40 those are. The top row is the leak, measured with the same tool on the dataset as it was shipped before 2026-10-21, and it is the only row not measured on the released data. The row below it is not a leak and is the highest on the current dataset, at 0.625 — level with the best method in Table 4: `perceived_severity` is an RFC 8632 field a real console shows, and what its score says is that our fault model marks the root critical too reliably (Section 9). The two rows below that reach 1.000 precision on 9 incidents each, which the gate deliberately allows: a value that names the root whenever it appears but appears in a fifth of incidents is a narrow cause doing its job, not a free answer.

6 Realism study

A generator is only worth using if its difficulty is the real task's difficulty. The reference is the Huawei PCIC 2021 competition data, released under the MIT licence [7, 10]: subset `18V_55N_Wireless` holds 34,838 timestamped alarms over 18 alarm types and 55 devices across 240 days, with a ground-truth causal graph over alarm *types* and no per-incident root-cause labels. It can therefore say whether our stream is plausible; it cannot score a method.

Four quantities are computable on both sides from (type, device, time) alone, and the same definition is applied to both. Bursts split on a 60 s quiet gap. Table 3 reports them.

One divergence left, and the one we fixed changed the results. The first version of this study found the generated stream firing in near-simultaneous volleys: 0.205 of consecutive alarms under a second apart against 0.019 of real ones. The cause was one line of the fault model. Every neighbour-cell alarm in a radio storm was drawn from a single lognormal around the fault's onset, so a dead site produced a cloud rather than a sequence: 283 of the 332 sub-second gaps in the stream were neighbour alarms. A live site does not notice a dead neighbour the moment it dies — it detects the degradation on its own measurement period — so each site now draws an offset into a 15-minute period. The sub-second share fell to 0.041, the median inter-arrival rose from 40.6 s to 60.4 s against 132.0 s real, and the singleton-burst share moved from 0.824 to 0.756 against 0.712 real. Section 8 is measured after that fix, and Section 9 records what it moved.

What remains is burst size: our largest burst is 82 alarms against 30. We tested the obvious explanation — that a bigger network gives one incident more elements to affect — and it is wrong. Across the sweep of Section 8.2, Abilene at 144 devices has a larger mean largest-burst than Cogentco at 2,625, so burst size here is set by the fault model and the alarm rates and not by the size of the graph. We have no mechanism for the remaining factor of 2.7 and do not claim one.

On the other two axes the generated stream matches. The type vocabulary is 21 named causes against 18, slightly flatter than the real one rather than more concentrated (normalised entropy 0.784 against 0.839; commonest type 0.167 against 0.212). The share of ordered type pairs exciting

	Generated	Real (PCIC)
alarms	1,572	34,838
alarm types	21	18
devices	717	55
inter-arrival median (s)	60.4	132.0
share of gaps under 1 s	0.041	0.019
largest burst (alarms)	**82**	**30**
singleton bursts	0.756	0.712
share of the commonest type	0.167	0.212
type entropy, normalised	0.784	0.839
excitation lift, maximum	524.0	69.48
pairs at lift ≥ 2, share	0.150	0.145

Table 3: Generated GÉANT small against real operator alarms. The type field is `probable_cause`, X.733's catalogue of named causes, which is the analogue of PCIC's `alarm_id`; see the caveat below on the field this table used first. Excitation is a lift-over-base-rate *proxy*: for each ordered type pair, how often B follows A within 300 s divided by B's base rate. No Hawkes kernel is fitted and no parameter is estimated on either side; a fit remains unmeasured. Bold marks the one axis on which the two streams differ.

each other above 2× is 0.150 against 0.145, a ratio of 1.03. The maximum lift does not match and is not comparable: ours is 524.0× against 69.48×, because spreading the storms left some ordered pairs rare enough that a handful of occurrences divides by a very small base rate. A maximum over 441 pairs is a tail statistic on a 1,572-alarm stream; the share above a threshold is the one to read, and it is the one we read.

The field this table measures is a choice, and the first version chose wrong. The analogue of PCIC's `alarm_id` — a named vendor alarm — is X.733's `probable_cause`, of which this dataset uses 21 values. The first version of this study used `alarm_type`, X.733's five-value *event type* enum, and so compared a severity-class taxonomy against a vendor alarm name. On that field the same alarms at the same times read 4 types against 18, a commonest-type share of 0.503 against 0.212, and a strongest excitation of 1.95× against 69.5× — from which we concluded that our vocabulary was too small and that our alarm types did not chain at all. Both conclusions were artefacts of the field. With four types the commonest has base rate 0.51, so a lift of 2× is near the arithmetic ceiling and the measure had almost no room to move. Same data, different label column, opposite finding. The results file carries both fields so the choice is auditable rather than buried.

The test that matters: the ranking mostly transfers. Whether the *order* of methods on generated data holds on real data decides whether the generator is useful for method development. The four topology-free rules, scored on 4,539 real alarm bursts and on our 40 eval incidents, agree on the top two and disagree on the bottom two: earliest then most-frequent on both sides, then real data puts random above most-recent and ours does not, for a Kendall τ of 0.667. The two it swaps are tied on our side at 0.225 each, so the disagreement is between two rows our own data cannot order. An earlier version of this study read $\tau = 1.0$; that was on the generated stream before the arrival-time fix of Section 6, and we report the weaker number measured after it.

What does transfer is the finding that matters. The earliest alarm names a root cause in 0.773 of real bursts against 0.625 of ours, so the rule that beats our methods is not an artefact of the generator; it is how operator alarm streams behave.

Four caveats, all load-bearing. The per-burst label on the real side is *derived by us* from the dataset's type-level causal graph — a type is a root if no other type present in the burst is its ancestor — because PCIC ships no per-incident labels; that definition encodes the same "a cause precedes its effects" assumption our generator does, and can flatter a rule that shares it. The real task is type-level and ours is device-level, so only the ordering is comparable, never the rates. Our own two graph methods need a propagation graph built from device topology, which `18V_55N` does not carry, so they are absent from the comparison: the transfer claim is about the baselines' order, not our margin over them. And four rules admit 24 orderings, so an exact match carries $p = 0.042$ under a uniform null — one ranking of four things, not a strong test. Of the 21,427 real bursts, 16,071 contain a single alarm type and 817 have no causal edge among the types present; both classes make every rule identical, and the 4,539 that remain are the comparison. That skip is a property of the real stream, not a filter chosen to help a rule.

The second PCIC subset is labelled and still unusable. `espk` carries 141,853 alarms, a 35-device topology, and 49 incidents labelled (alarmed device set → root-cause device) — device-level and labelled, which is exactly our task. Its snapshots carry no timestamps, so no time-ordering method can run on them, and its 49 incidents share only 4 distinct root causes over a median snapshot of 28 of 35 devices, which a constant predictor exploits. We report it rather than use it.

7 Baselines

Nine rows are scored. Four are one-line rules, two are the graph methods we designed, and three are baselines we did not design — the last group matters because a resource paper whose own method is the best row has not shown the resource is useful to anyone else.

Rules. Earliest critical alarm; most-alarmed device; most recent change within two hours of the first alarm; and a uniform draw among the alarmed devices, which is the floor.

Ours. *temporal* ranks candidates by how many symptoms they explain under a backwards walk over edges derived from topology and alarm times, breaking ties by onset. *neighbourhood* is the same candidate set with time removed, ranking by how many alarmed devices a candidate neighbours; it is in the table to show what the time ordering contributes.

Both run on our own graph engine, Samyama, through its `algo.symptomExplanation` procedure: the walk and the explanation count are the engine's, and what we wrote is the construction of the edges it walks. The runs reported here used the embedded Python SDK **1.9.0**. We name the version because these two rows cannot be reproduced without it, and because the engine is ours — a resource paper that leaves its own dependency unversioned is asking the reader to take the rows on trust. A later engine release would be a different measurement and we have not made it.

External. *thp_causal* is gCastle's `TTPM` [7], the reference implementation of the topological Hawkes process of Cai et al. [10], under Apache-2.0. It is fitted once, on the 697 alarms raised before the first eval incident begins, and sees no eval alarm, no incident grouping and no ground truth; it is given the same device adjacency the other methods get. It learns 41 edges over the 21 `probable_cause`

values. That graph is type-level, so a reduction is needed to answer a device question, and we use the published one: among the cause types present in an incident, keep those with no parent also present, and rank the devices that raised one of them, earliest first.

thp_nograph is that same reduction with the learned graph removed, so every cause type present counts as root — which makes it "the device that alarmed earliest". It is the control. Without it the *thp_causal* row cannot be read at all, because a change could equally be the causal model or the reduction wrapped around it.

pagerank_alarmed is personalised PageRank over the same device adjacency, teleporting uniformly to the alarmed devices — the MonitorRank and CloudRanger shape. It uses no time and no causal model, ranking purely by structural centrality within the alarmed set. It cannot name a PoP, so a site-wide cause is reachable for it only through the scoring rule that a device at the root PoP counts as a hit.

An LLM, given the same view. *llm_ranker* sends one prompt per incident to a general-purpose model (claude-sonnet-5-5) containing exactly what Section 5 allows: the incident's alarms with their times relative to the first, the role of each alarmed device, the adjacency around them, and the candidate list. It asks for three candidate ids, best first. It is the only row whose input is free text, so what reaches it cannot be audited by reading a query; a test asserts over the generated prompt that no fault id, scenario id, mechanism name or absolute timestamp appears in it.

Two of its properties are unlike every other row and both are recorded with the number. It is not deterministic and this is one run, so 0.600 is a sample from a distribution rather than a repeatable figure. And a reply can be unusable: an answer that names no candidate is charged a miss, never retried, because a method that sometimes returns nothing scores lower and a retry loop would measure a different method. The run cost $2.30, which makes it and the agent below the only baselines here with a price per query.

The same model, allowed to ask. *llm_agent* is the agent row. It gets ten read-only verbs over the graph — an incident's alarms, a node's properties, a device's cable links, a path between two devices, the devices at a site, the changes targeting a device, the services depending on it, how many alarms it raised before this incident — and 6 rounds of queries before it must answer. It has no shell and no file access: it emits query lines, we execute them, and the verb set is closed, so a request we do not implement is answered with the verb list.

The graph it queries is sanitised before it sees anything: `Fault` nodes, `AT` and `CAUSED` edges are removed and the `scenario` field is stripped from every alarm and incident. That last one matters more than it looks — an incident's id is `inc:flt:eval-overlap-0050:0`, which names its own fault, and the scenario names the split and the index. A ranker is safe because its prompt is a fixed string a test can read; an agent chooses its own queries, so the only durable guarantee is that the answer is not in the graph. We assert that in a test rather than trusting the verb list.

8 Results

Forty incidents was not enough, and we can say exactly how. Table 4 reports every method at two sample sizes. The $n = 40$ column is the configuration this study first used; the $n = 200$ column is five times as many incidents from the same generator at the same seed with the same rates, so the two differ in how much was measured and in nothing else. Figure 3 draws both with the interval each carries.

The smaller sample was not merely imprecise, it was wrong about the order. Three rows move by 0.1 or more — our temporal method by 0.125, our neighbourhood method by 0.105, the most-alarmed rule by 0.125 — and the ordering changes with them: at $n = 40$ our temporal method sits below *thp_causal*, and at $n = 200$ it is above it. Forty incidents can separate the top of this table from the bottom and cannot order the middle. We report the larger column as the result and keep the smaller one visible, because a reader deserves to see what a 5× smaller sample of the same data said.

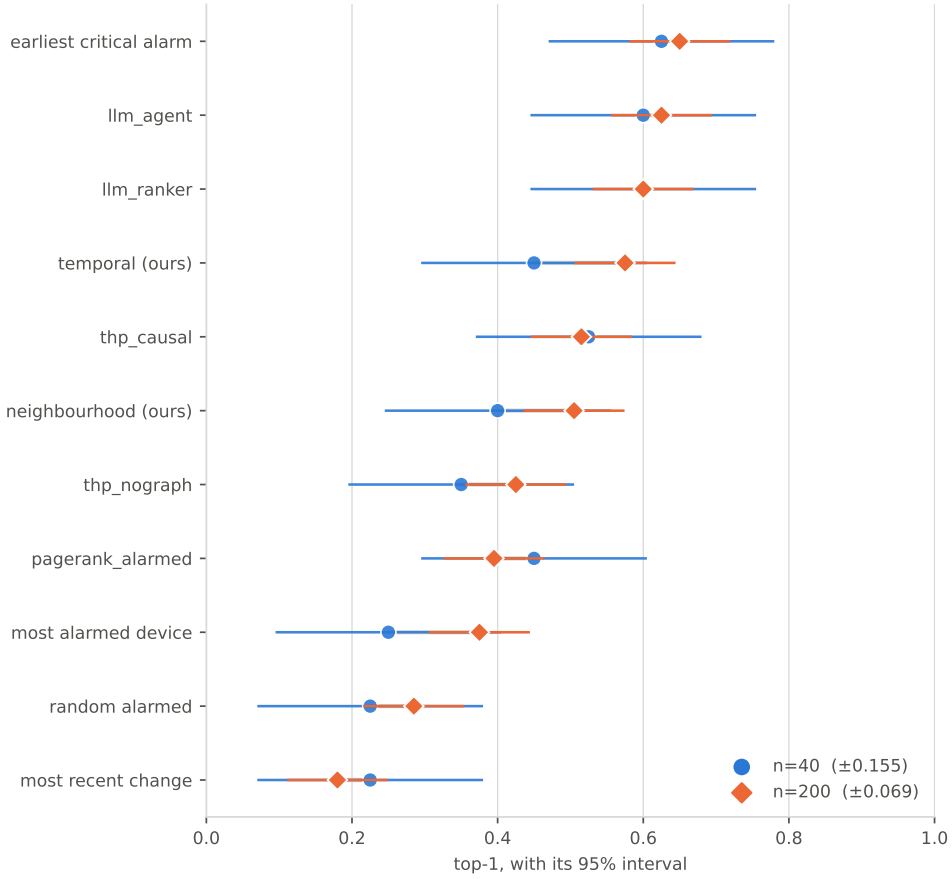

Figure 3: Every method at both sample sizes, with its 95% interval. The $n = 40$ intervals (circles) are more than twice the width of the $n = 200$ ones (diamonds), and most of the table sits inside a single one of them.

The simplest rule usually wins at top-1. Ranking the device that raised the earliest critical alarm takes 0.625, against 0.450 for our temporal method and 0.525 for *thp_causal*, the best non-LLM baseline. Across the 15 runs of Section 8.2 it takes top-1 in 13 of them, so "usually" is the strongest form of this claim we can make: in the other 2 our temporal method does. Its margin is narrower than it was: on the pre-fix stream it took 0.700 against 0.475, and spreading the storms cost the rule 0.075 while costing our temporal method 0.025. A rule that reads the front of the queue loses more when the queue stops arriving all at once. Section 6 tested whether this is an artefact of the generator, and it is not: the same rule takes 0.773 on real operator alarms, higher than on ours. Part of the margin is still ours to answer for — our storms are 10× more likely to

	$n = 40$	$n = 200$	
Method	Top-1	Top-1	Top-3
rules we did not design			
earliest critical alarm	0.625	**0.650**	0.650
most alarmed device	0.250	0.375	0.535
random among alarmed devices	0.225	0.285	0.545
most recent change	0.225	0.180	0.200
external, not designed by us			
llm_ranker (claude-sonnet-5-5)†	0.600	0.600	0.670
llm_agent (same model, 10 graph verbs)†	0.600	0.625	0.670
thp_causal (gCastle TTPM)	0.525	0.515	0.655
thp_nograph *(control)*	0.350	0.425	0.625
pagerank_alarmed	0.450	0.395	0.600
ours			
neighbourhood, no time	0.400	0.505	0.670
temporal	0.450	0.575	**0.790**

Table 4: GÉANT, the same generator and seed at two sample sizes: 40 eval incidents and 200. The configurations differ in scenario counts and in nothing else. Two of them raise no alarm on any device, so no method has anything to rank; each is charged a miss to every row rather than dropped (Section 9). A rate near 0.5 carries ±0.155 at 95% on 40 trials and ±0.069 on 200; Figure 3 draws both. No paired test has been run. † one run of a non-deterministic method, so that row is a sample and not a repeatable number; the others are. Every row is measured after the arrival-time fix of Section 6; the numbers in the first version of this paper were measured before it and all of them moved.

arrive inside one second (Table 3), which shortens the queue the rule has to pick from — but the finding survives the generator. Tuning the method against this number would be fitting to our own data and is not done.

The word doing the work is "critical", not "earliest". *thp_nograph* is the earliest-alarm rule without the severity filter, and it takes 0.350 against the earliest *critical* rule's 0.625. The two differ in nothing else: same candidate set, same scoring function, same 40 incidents. The gap widened when the storms were spread: on the pre-fix stream the control took 0.475 against the rule's 0.700, so dropping the severity filter costs more, not less, once alarms stop arriving in one volley. So the one-line rule's strength is not that it reads the clock; it is that the generator marks the root alarm critical often enough for a severity filter to find it. That is a property of our fault model and the sharpest single thing to attack in the next version of the generator.

The learned causal graph helps, slightly, and only at top-1. *thp_causal* takes 0.525 against its own control's 0.350: restricting the earliest-first ranking to causally-root cause types promotes the right device in seven more incidents out of 40, and it is now the best method in the table that is neither a one-line rule nor an LLM. On the pre-fix stream the same comparison was 0.525 against 0.475, so the whole of this gain appeared when the arrival times were fixed: the control degrades without the volley and the causal filter does not. At top-3 the two are level at 0.625. The top-1

difference is seven incidents wide against a ±0.155 interval, which is the first result here where the learned structure does something a single draw can see — and it is still one draw, from a kernel fitted on 688 alarms.

A model that knows what an alarm means still loses to the severity field. *llm_ranker* takes 0.600 top-1 — above every graph method in the table, above both of ours, and below the 0.625 of a rule that reads one field. Whatever telecom knowledge the model brings is worth something over structure-only ranking and is not worth as much as knowing that the generator marks root alarms critical. Its top-3 is the more interesting number: 0.625, barely above its own top-1, where our temporal method rises from 0.450 to 0.725 given the same two extra slots. Of 40 incidents it is right first in 24, second in 1, and third in none. It returns one confident answer and a tail it has not really ranked, which is exactly the shape that does not help an operator who will check three devices.

Its per-fault-type scores say where the knowledge pays, and one class where it buys nothing at all: 0.000 on every site power loss, at both sample sizes. That class is the subject of the next paragraph, and it is the one place in this paper where the methods we designed are not merely competitive but alone.

Giving the model tools changed what it looked at and not what it answered. *llm_agent* is the same model with ten read-only graph verbs and six rounds of queries before it must answer: it can list an incident's alarms, walk cable links, find a path between two devices, pull the change records near a device, and ask how many alarms a device raised before this incident. It used them — 263 queries, 6.58 per incident, never exhausting its turns — and scored 0.600 top-1 and 0.625 top-3, identical to the ranker that saw one fixed prompt, at 2.7× the cost ($6.14 against $2.30). Its per-fault-type scores are identical too, in all six classes.

The two agree on the first pick in 38 of 40 incidents; they differ only in the second and third slots, which this benchmark's scoring barely rewards. What the agent chose to look at is the more useful result: 101 of its queries asked for change records and 29 for a device's alarm history, against 66 for the topology and 35 for the incident's own alarms. It spent most of its budget on the hypothesis that a recent change caused the fault — which is the worst-scoring baseline in Table 4 at 0.225 — and on 7 incidents it answered without issuing a single query. We read this as evidence that the task is not tool-limited: the information needed to beat 0.625 is in the alarm list, and no amount of graph querying supplies what is missing.

When the root is silent, only a method that reads the topology works. A site power loss kills the PoP, and a dead PoP raises no alarm. Every alarm in the storm that follows belongs to a device that is working, so a method that *ranks alarmed devices* cannot name the root — not because it ranks them badly, but because the answer is not in the list it ranks. The 12 power-loss incidents at $n = 200$ say exactly that:

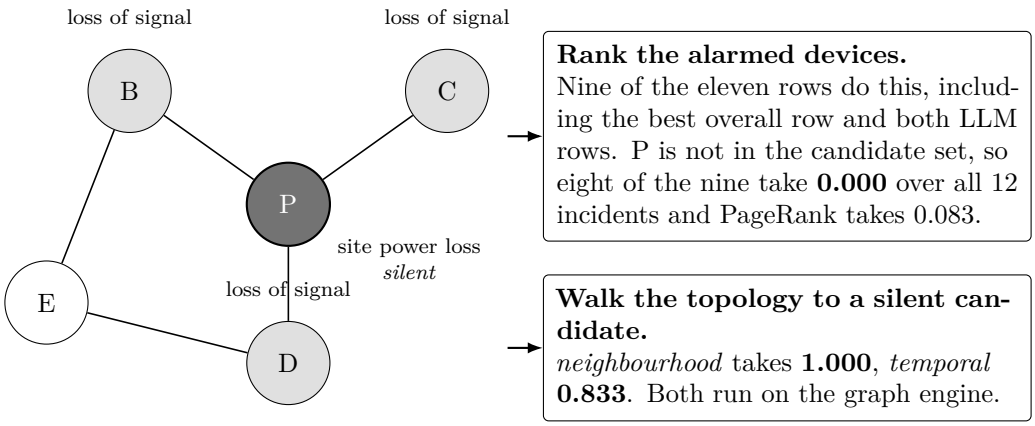

Figure 4: The fault class that separates the table. A dead PoP cannot report itself, so the answer is the one device that emitted nothing. Reaching it is a traversal over the topology, and no amount of reading the alarm stream substitutes for it.

On the 12 site power losses	Top-1
neighbourhood (ours)	**1.000**
temporal (ours)	0.833
pagerank_alarmed	0.083
earliest critical alarm	0.000
most alarmed device	0.000
thp_causal	0.000
llm_ranker	0.000
llm_agent	0.000

Our two methods nominate candidates that never alarmed, by walking the topology from the devices that did: a PoP whose hosted devices all went quiet at once is itself a candidate. Everything else in the table — four rules, a fitted topological Hawkes process, personalised PageRank, a language model with the alarms in front of it, and the same model allowed ten graph queries — scores zero or near it. This is not a tuning difference; it is the difference between selecting from a list and inferring something absent from it.

It is also the result that 40 incidents could not show. At $n = 40$ this class held 3 incidents, the paper reported that no method did well on it, and that statement was wrong: it was three samples of a class our methods solve and nothing else touches. Figure 5 shows the whole grid, where the power-loss column is the only one with a clean split.

Where else the methods win. Our temporal method takes top-3 at 0.790, the best in the table — the next row is 0.670 — and the earliest-critical rule cannot improve on its own top-1 at all — 0.650 at both $k = 1$ and $k = 3$, because it returns one candidate. That is a comparison between a ranker and a classifier, and the gain should be read as such: a top-3 list is worth something to an operator who will check three devices, and nothing to an automated remediation step that acts on one. We claim the former only.

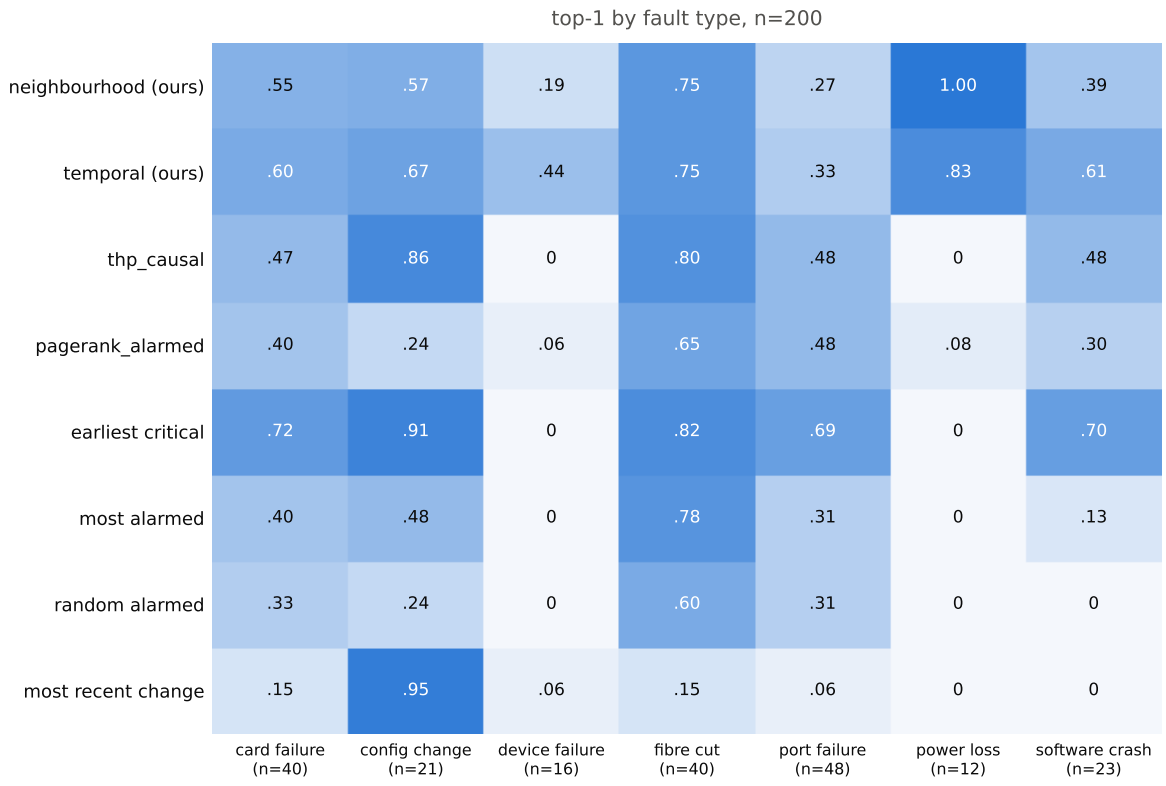

Figure 5: Top-1 per fault type at $n = 200$. Reading down the *power loss* column: the two methods that walk the topology take 1.000 and 0.833, and every rule, both Hawkes variants and both LLM rows take 0.000. No other column separates the table that way.

	GÉANT (40)		Abilene (11)		Cogentco (197)	
Method	mean	range	mean	range	mean	range
earliest critical alarm	0.585 (±0.038)	0.550–0.625	0.620 (±0.048)	0.575–0.700	0.595 (±0.054)	0.525–0.650
temporal (ours)	0.485 (±0.055)	0.450–0.575	0.565 (±0.076)	0.475–0.650	0.535 (±0.029)	0.500–0.575
neighbourhood (ours)	0.420 (±0.062)	0.375–0.525	0.480 (±0.084)	0.350–0.550	0.485 (±0.065)	0.400–0.575
thp_causal	0.395 (±0.116)	0.225–0.525	0.420 (±0.076)	0.350–0.525	0.380 (±0.021)	0.350–0.400

Table 5: Top-1 over five seeds per topology, mean (± sd) and observed range; $n = 40$ in every run. The number in each heading is the backbone's PoP count. *llm_ranker* is absent: it is a paid, non-deterministic method and was run once.

8.1 What the graph engine contributes, and what it does not

The *neighbourhood* and *temporal* rows are the only two that run on a graph database, and Figure 4 is the only place in the grid where that is decisive. The reason is structural rather than incidental: every other row consumes the alarm stream and returns something from it, so its answer is bounded by the set of devices that spoke. A traversal is not. `algo.symptomExplanation` walks outwards from the alarmed set over edges we construct from topology and alarm times, and a node it reaches is a candidate whether or not it emitted anything — which is why a PoP whose hosted devices all went quiet at once can be nominated at all. That is one capability, stated once: the ability to name a vertex that is absent from the evidence.

We claim nothing about speed or scale here, and the reader should not infer it. The graph these two rows run on is 11,659 nodes and 14,793 edges, which fits anywhere — and the largest backbone in Section 8.2 is bigger only in proportion to its device count, which is 2,625 against 717; no query latency, memory figure or throughput comparison appears in this paper, and none should be read into the two rows. The engine is reported as a dependency with a version, not as a result. Equally, the capability is not unique to our engine — any system that can traverse the topology can express the same walk, and we would expect a reimplementation to land on the same two numbers. What Figure 4 measures is that the traversal is worth doing, not that ours is the only way to do it.

8.2 Across seeds and topologies

We re-ran the whole benchmark 15 times: three Internet Topology Zoo backbones — GÉANT (40 PoPs), Abilene (11) and Cogentco (197) — at five seeds each, one config throughout, scored by the same script. A seed redraws the network's failure history *and* the 40 eval incidents, so each run is a different benchmark instance of the same generator.

The rule wins most runs, not every run. The earliest-critical rule takes top-1 in 13 of the 15 runs. In the other 2 our temporal method does, and both are the same seed on two different backbones. No method wins everywhere, and the full ordering of Table 4 reproduces in 1 of the 15 runs — the shipped seed itself. Every other run permutes the middle of the table somewhere. The headline claim survives in the form "usually" and does not survive in the form "always".

The seed spread was itself an artefact of the arrival times. At $n = 40$ one run's rate carries a standard error of 0.079, and ±0.155 at 95% is 1.96 of those. The seed-to-seed sd of the best row is now 0.038 on GÉANT, 0.048 on Abilene and 0.054 on Cogentco — that is 0.48, 0.61 and 0.68 of one run's standard error, so redrawing the whole network and all 40 incidents moves the rate *less* than resampling 40 incidents would. The ±0.155 interval is the conservative one to quote.

Source removed	earliest crit.	temporal	neighbourhood	thp_causal
alarms that never arrive	+0.150	+0.150	+0.200	−0.025
background alarms‡	+0.000	+0.025	+0.075	+0.075
per-device clock skew	+0.000	+0.100	+0.075	−0.075
neighbour-cell storms	+0.000	+0.100	+0.175	−0.100
mis-grouped incidents	+0.050	+0.025	+0.050	+0.075

Table 6: Signed change in top-1 when one noise source is removed; positive means the method got better without it. GÉANT small, seed as in Table 4, $n = 40$ throughout. ‡ removing background alarms also changes how many incidents have any alarm to rank, from 0 to 3: for three incidents the only alarms the NOC had attached were background ones. Three forced misses then sit inside every method's delta on that row, so it is not a difference in difficulty and we do not read it as one.

That is the opposite of what the same sweep said before the arrival-time fix of Section 6, and the reversal is the finding. On the pre-fix stream the GÉANT sd was 0.127, or 1.61 standard errors, and the rule's range across five seeds was 0.300 wide. Spreading the storms cut that spread to a quarter. The instability was not seed sensitivity in any interesting sense: it was that a rule reading the front of a simultaneous volley is at the mercy of which alarm happened to land first, and a few hundred milliseconds of draw decided whole incidents. Fixing the mechanism removed most of the variance along with some of the accuracy.

What does hold across all three backbones. The earliest-critical rule has the highest mean on all three; our temporal method is second on all three; and *thp_causal*, the only method given the topology to learn from, is last of these four on all three, with the widest spread of any row on GÉANT (±0.116) and the narrowest on Cogentco (±0.021). Network size moves the absolute rates very little: Cogentco has 197 PoPs against Abilene's 11, and the best mean differs by 0.025 between them. The generator's difficulty travels with the fault model and the noise process, not with the graph — which is also why the remaining burst-size divergence in Section 6 cannot be explained by network size.

8.3 Which noise source makes it hard

The generator asserts five sources of difficulty, each a config knob. An ablation turns exactly one off, on the same topology and the same seed, and re-scores: what moves is attributable to that knob, and what does not move is difficulty we claimed and did not add. Six runs, the shipped config plus five.

Missing alarms are the difficulty; the severity rule is immune to the rest. Delivering every alarm that could be raised is the largest move in the table: +0.150 for the earliest-critical rule, +0.150 for our temporal method and +0.200 for our neighbourhood method. Alarms that never arrive are what the benchmark is mostly made of. Three of the other four rows move the earliest-critical rule by exactly 0.000: neither clock skew, nor the neighbour storms, nor the background noise changes which device raised the earliest critical alarm. That is the clearest statement of what the one-line rule is doing, and of why the noise we add does not trouble it.

Clock skew and the neighbour storms cost the two methods that read order and count neighbours: +0.100 each to our temporal method, +0.075 and +0.175 to our neighbourhood method. Both *hurt thp_causal* when removed (−0.075 and −0.100), which we do not have a mechanism for and do not read as a result at two or four incidents wide. Mis-grouping, the source an earlier version of this

study found worthless, now moves four methods by 0.025 to 0.075: it was worth nothing while the storms arrived in a volley and is worth something now that they do not.

One row is not a difference, and a check that watched the wrong thing said it was. Removing background alarms takes the count of incidents with no alarm to rank from 0 to 3, because for three incidents every alarm the NOC had attached was a background one. Three forced misses sit inside that row's deltas. The first version of this check watched the denominator n, which cannot move — an unrankable incident is charged a miss, not dropped — and so reported all five ablations as clean subtractions, including the two largest moves in the table at the time.

9 Limitations

An incident with no alarm is charged as a miss, not dropped. A fault whose every symptom went unreported raises no alarm on any device, and no method has anything to rank. Both of our scoring scripts dropped such an incident until 2026-10-21, which scored 38 trials while the split size printed in the caption was 40. An invisible fault is still a fault the operator did not localise, and a benchmark that deletes what nothing can answer reports a denominator it does not have, so each is now a miss for every method and every baseline. The dataset in Table 4 happens to have 0 of them, against 2 before the arrival-time regeneration, so the rule costs nothing here — it is stated because it is the rule, not because it binds. The ablation in Table 6 shows why it matters: switching off background alarms takes that count to 3.

Nearly half the alarms are background. 0.471 of alarms are unrelated to any injected fault. A method scoring above the random-among-alarmed floor of 0.225 is therefore not reading a clean signal — but neither is the task as hard as a real console, where background rates vary by time of day and our noise process is stationary.

One configuration, and one seed per row of the main table. Every number in Table 4 is GÉANT small at one seed. Section 8.2 measures what that costs and the answer is less than we expected: the best row's seed-to-seed sd is 0.038 on that topology, under half of one run's standard error, so the ±0.155 sampling interval already covers it. What does not reproduce is the *ordering*: the full order of Table 4 appears in 1 of the 15 runs, because the methods sit closer together than the spread of any one of them. Read the table as one instance of the rates and as a weak claim about their order, with Table 5 as the error bar on both. Three backbones and five seeds are still one config: the fault rates, the noise rates and the split policy are fixed throughout, and nothing here says what happens when those change.

We built the data and two of the methods. The resource is the claim, and the methods we designed are not the best rows in Table 4: a one-line rule beats both at top-1. Five baselines we did not design are now scored — three graph or statistical methods, an LLM ranker and an LLM agent that queries the graph itself — and none of them beats that rule either, which is some evidence the task is as it looks rather than tuned to us. The strongest remaining objection is not a missing baseline but the fault model: the rule works because our generator marks root alarms critical, and no baseline can tell us whether a real network does so as often.

Our storms were too bursty, and fixing it moved every number in this paper. The version of this study before 2026-10-21 reported 0.205 of consecutive alarms arriving within a second against 0.019 of real ones, and said that fixing the fault model's emission timing would invalidate every number in its results table. It did, and this is that paper. The mechanism was one line — every neighbour-cell alarm in a radio storm drawn from a single lognormal around the fault's onset — and the fix moved the sub-second share to 0.041. What it cost is recorded throughout: the severity rule fell from 0.700 to 0.625, our temporal method from 0.475 to 0.450 at top-1 and from 0.775 to 0.725 at top-3, the LLM ranker from 0.625 to 0.600, the transfer test's Kendall τ from 1.0 to 0.667, and the seed-to-seed spread to roughly a quarter of what it was. Two findings survived and one appeared: the one-line rule still wins, mis-grouping stopped looking worthless, and *thp_causal*'s advantage over its own control grew from two incidents to seven.

What is still wrong is burst size, at 82 alarms against 30, and we have no mechanism for it. The explanation we expected — that a larger network gives one incident more elements to affect — is measurably false across the three backbones of Section 8.2.

We shipped the answer in an alarm field. Until 2026-10-21 every alarm carried its role in the fault as `specificProblem`, so `specific_problem = 'root'` named a true root resource in all 28 of the 40 eval incidents where such an alarm existed — 1.000 precision, 0.700 top-1, from 59 of 1,617 alarms and no reasoning. No number in this paper came from it: nothing in the benchmark or the baselines read the field, and regenerating the dataset without it changed no scored rate, because the fix draws no extra randomness and the two datasets differ in that one field. The cost was to a reader, who could have taken the benchmark's best score with a single `WHERE` clause. We report it rather than quietly reissuing the data, and the check that found it is now part of the release (Section 5). What it does not fix is the row below it in Table 2: a severity filter still scores 0.625 — the same as the best method in Table 4 — and severity is a field the data is supposed to have.

One realism finding here was previously wrong in our favour's opposite direction. Two of the three divergences this section reported before 2026-10-21 were artefacts of measuring over the wrong alarm field, and the error made the generator look worse than it is. The correction is in Section 6; we record it here because a resource paper's realism section is the part a reader cannot re-derive without the data, and the field it is computed over is part of the measurement.

The excitation comparison is a proxy, not a fit. What Table 3 reports is a lift over base rate within a 300 s window. No Hawkes kernel is estimated, on either side, and the shape of the real excitation decay is therefore unknown to us. A fitted model could change the sign of nothing in that table and would still be the right measurement.

10 Release

What ships, and where. The generator is public, Apache-2.0, at https://github.com/samyama-ai/telecom-synthetic-data-generator. It is a squashed export carrying no development history, regenerated per release rather than maintained by hand, so what is published is the state of a release and nothing about how it was reached. The benchmark and the scoring scripts follow in a second repository; they are cited here as forthcoming rather than listed as shipped, because a URL that does not resolve is worse than none. What ships today is the generator, plus one generated dataset per configuration. The dataset is not the artifact; the generator is. A directory is a function of (pack version, config, seed) and the same seed reproduces its files byte for

byte, which is a tested property and not an aspiration — so a reader who disagrees with a number here can regenerate the input to it rather than download ours. Every emitted directory carries its own `manifest.json` (seed, config, every input and output with sha256) and `binding.json` (the vocabulary terms used and the build files they came from).

Croissant. Each directory carries MLCommons Croissant 1.0 metadata describing its files, their checksums and the fields inside them, generated from that manifest so the checksums in the metadata are the checksums of the files beside it rather than numbers typed a second time. It names the ground-truth fields explicitly, including where each alarm's mechanism lives, because a downloader who does not know which fields are the answers will train on them.

Licence, settled. The data is CC BY 4.0, which carries the Internet Topology Zoo's own terms for the backbone it is built on; the generator and the scoring code are Apache-2.0. RFC 8632 and RFC 8345 are Simplified BSD and TMF642 is Apache-2.0. NORIA-O's advertising clause, the one part that was open, is satisfied by the acknowledgement in Section 4: what we redistribute from NORIA-O is term names, not the ontology, so the clauses binding a redistribution to retain its notice apply to the generator's own vocabulary fixture — where that notice now sits — and not to the dataset. The Croissant metadata asserts those two licences and carries the acknowledgement in a generated field, so a release cannot drop it by editing prose. Until the review closed the same metadata asserted no licence at all, which was the right default: the field a downstream tool reads with no human present is the worst place to put a guess.

A DOI, and what this version does not yet carry. The dataset is deposited as a prepared, unpublished record: the archive, its checksum and the full deposit metadata exist, and publishing it mints the DOI. That step is deliberately not taken in time for this version, because a DOI is permanent and the record it names should be the one a venue will review. Until then the resource is cited by this paper and by the public generator repository above, which is what a reader needs to reproduce it; a revision will carry the DOI. We state this rather than cite an identifier that does not yet resolve.

References

[1] A YANG data model for network topologies. https://www.rfc-editor.org/rfc/rfc8345, 2018.

[2] LEMMA-RCA: A large multi-modal multi-domain dataset for root cause analysis. arXiv preprint arXiv:2406.05375, 2024. Microservice and water-treatment domains. CC BY-ND 4.0 on the site and in the paper; the Hugging Face header tags it cc-by-nc-4.0.

[3] RCAEval: A benchmark for root cause analysis of microservice systems. arXiv preprint arXiv:2412.17015, 2024. 735 failure cases over three microservice systems. Zenodo copy CC BY 4.0, Hugging Face copy tagged MIT.

[4] OpenRCA: Can large language models locate the root cause of software failures? ICLR 2025, 2025. 335 cases over Bank, Market and Telecom domains. Cloud observability telemetry, not NOC alarms. MIT on the code; no licence stated on the data.

[5] TN-AutoRCA: Benchmark construction and agentic framework for telecom network root cause analysis. arXiv preprint arXiv:2507.18190, 2025. Announces TN-RCA530, 530 scenarios as

knowledge graphs with per-incident labels. No repository, download, availability statement or dataset licence appears in the paper.

[6] EvoCause: Evolutionary causal discovery for root cause analysis in telecom networks. arXiv preprint arXiv:2607.27290, 2026. Introduces TeleRCA: 10,922 production incidents, 485,681 alarms, 5,621 resources, 194 alarm types, expert `causeNode` annotation. States twice that TeleRCA has no ground-truth propagation graph. No URL in the paper.

[7] Huawei Noah's Ark Lab. gcastle-hub/dataset: alarm datasets from the PCIC 2021 causal discovery competition. https://github.com/gcastle-hub/dataset, 2026. MIT. The URL printed in Qiao et al. [10] no longer serves data; this mirror does. Subset `18V_55N_Wireless`: 34,838 alarms, 18 types, 55 devices, 240 days of alarms, plus a type-level ground-truth causal graph and no per-incident labels. Subset `espk`: 141,853 alarms, a 35-device topology and 49 incidents labelled (alarmed device set -> root-cause device), without timestamps. Accessed 2026-10-21.

[8] Huawei Paris Research Center, NetOp Team. TeleLogs. https://huggingface.co/datasets/netop/TeleLogs, 2025. Gated. Licence tag MIT while the card asks that the data not be shared or redistributed. 5G KPI/drive-test records, not alarms.

[9] Simon Knight, Hung X. Nguyen, Nickolas Falkner, Rhys Bowden, and Matthew Roughan. The Internet Topology Zoo. *IEEE Journal on Selected Areas in Communications*, 29(9):1765–1775, 2011. CC BY 4.0.

[10] Jie Qiao et al. Structural hawkes processes for learning causal structure from discrete-time event sequences. *arXiv preprint arXiv:2305.05986*, 2023. IJCAI 2023. Uses the Huawei PCIC 2021 alarm data; ground truth is an 18x18 adjacency between alarm *types*.

[11] Lionel Tailhardat, Raphaël Troncy, and Yoan Chabot. NORIA-O: An ontology for anomaly detection and incident management in ICT systems. In *ESWC*, 2024. BSD-4-Clause with the advertising clause. Alarm vocabulary aligns to ITU-T X.733/X.721 and TMF638; it references neither RFC 8632 nor TMF642.

[12] Lionel Tailhardat, Raphaël Troncy, Yoan Chabot, Fano Ramparany, Pauline Folz, and Anne Kavanagh. Knowledge graphs for enhanced cross-operator incident management and network design. Internet-Draft draft-tailhardat-nmop-incident-management-noria-05, 2026. Updated 2026-08-07, expires 2027-02-08. Cites RFC 8345; no occurrence of RFC 8632 or TMF642.

[13] TM Forum. TMF642 alarm management API, 2025. v5.0, directory release 2025-01-23. Apache-2.0.

[14] Stefan Vallin and Martin Bjorklund. A YANG data model for alarm management. https://www.rfc-editor.org/rfc/rfc8632, 2019.

[15] Keli Zhang, Marcus Kalander, et al. An influence-based approach for root cause alarm discovery in telecom networks. *arXiv preprint arXiv:2105.03092*, 2021. States data and code are available at github.com/shaido987/alarm-rca; that repository contains a single 378-byte README and no licence.