arXiv:2610.30611v1 [cs.AI] 22 Oct 2026

Better Reasoning, Not Just Bigger Models: A Structure-Grounded Framework for Efficient and Reliable Inference

Amirhossein Aghaei*
Huawei Research Canada
Markham, Canada
amirhossein.aghaei@h-partners.com

Liangyan Li*
Huawei Research Canada
Waterloo, Canada
lillian.li2@h-partners.com

Shiyi Yang
Huawei Research Canada
Waterloo, Canada
shiyi.yang1@huawei.com

Shaohua Wang
Huawei Technologies Co., Ltd.
Beijing, China
wangshaohua25@huawei.com

Yan Wang
Central University of Finance and Economics
Beijing, China
dayanking@gamil.com

Abstract

Scaling large language models has substantially improved linguistic fluency, yet reliable autonomous agents require more than fluent generation. In multi-turn inference, current agentic systems often produce semantically plausible but structurally ungrounded reasoning, especially under incomplete observations. We propose **Structure-Grounded Uncertainty Reduction (SGUR)**, a training-free framework that complements token-level generation with active uncertainty reduction over structured hypothesis spaces. SGUR represents incomplete observations as a query graph and aligns them with candidate hypothesis graphs through a prior-weighted partial graph matching mechanism, where database-derived empirical priors assign greater importance to decision-relevant relations rather than treating all graph edges uniformly. The resulting inference state is iteratively refined by an active information-gain strategy that selects the next most informative node to query, systematically pruning the candidate space. We instantiate SGUR on a challenging PrimeKG-derived sequential medical diagnosis benchmark, demonstrating that structural grounding can partially substitute for model scale: starting from only two observed symptoms, SGUR achieves **92.5% Hit@3**, outperforming the strongest Graph-RAG baseline by **9.4 percentage points** while using approximately $5\times$ **fewer tokens**. Notably, SGUR enables smaller backbone models to match or exceed the performance of models up to $10\times$ **larger**, suggesting that structural grounding provides a more robust and efficient paradigm for scaling agent performance than relying exclusively on increasing model size.

*Equal contribution.

Preprint.

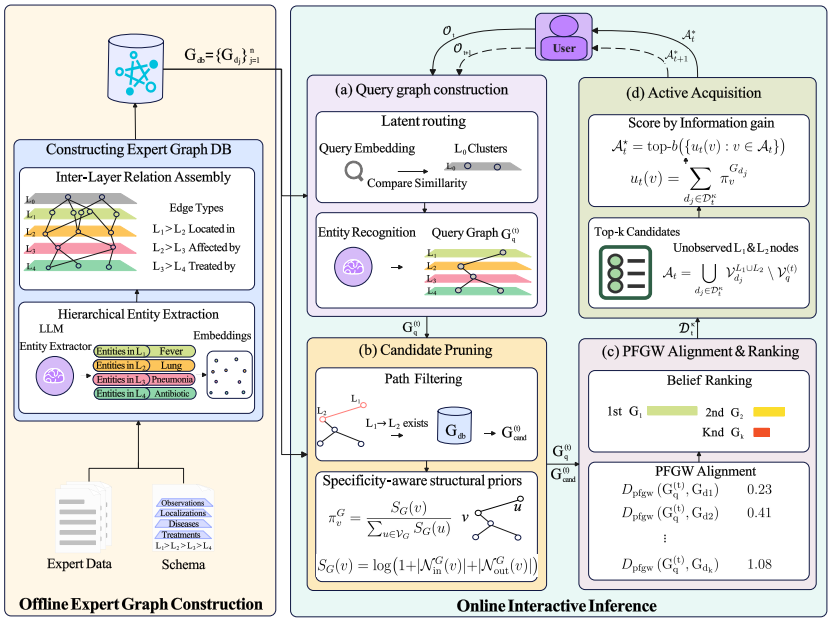

Figure 1: Overview of the SGUR framework. The offline phase constructs a hierarchical expert graph database G_{db} from expert data under a shared schema. The online phase performs closed-loop multi-turn inference through four stages: (a) query graph construction, which grounds observations $\mathcal{O}_t$ into $G_q^{(t)}$ with latent routing; (b) candidate pruning, which applies path filtering and specificity-aware structural priors to obtain $G_{\mathrm{cand}}^{(t)}$; (c) prior-weighted partial FGW alignment, which ranks candidate graphs by structural consistency with the query; and (d) active acquisition, which selects informative unobserved nodes to obtain $\mathcal{O}_{t+1}$. Across rounds, SGUR progressively contracts the candidate space under partial observation.

1 Introduction

The design of autonomous agents capable of reasoning under *partial, noisy, and sequentially revealed information* remains a central challenge for building reliable and efficient agentic systems. Although scaling large language models (LLMs) has substantially improved linguistic fluency[OpenAI, 2024b, 2025b,c], reliable multi-turn inference requires more than fluent generation. In complex decision-making environments, an agent must maintain a coherent hypothesis space, strategically acquire missing evidence, and update its belief state as new information arrives. However, many current agentic systems still rely primarily on token-level generation and linguistic similarity, which can lead to weakly grounded multi-turn inference when semantic evidence is not jointly constrained by relational structure[Fastowski and Kasneci, 2025, Zheng et al., 2026]. Under iterative interaction, such systems may exhibit *heuristic substitution*[Kahneman, 2011], relying on associative pattern matching rather than refining hypotheses against jointly grounded semantic evidence and relational structure.

Existing paradigms only partially address this challenge. Chain-of-Thought and Tree-of-Thought improve reasoning decomposition, but continue to operate largely within the model's internal probabilistic space[Wei et al., 2022, Yao et al., 2023a]. Recent memory-augmented systems also compress conversational history, yet topology-insensitive compression can discard decision-relevant cues[You et al., 2026]. GraphRAG frameworks introduce additional structure by retrieving graph-organized evidence[Edge et al., 2024a, Guo et al., 2024, Peng et al., 2024, Zhang et al., 2025a]. Recent hierarchical GraphRAG methods further strengthen this line by organizing knowledge into attributed communities and hierarchical indexes, improving retrieval relevance and reducing online token cost[Wang et al., 2026]. However, their main contribution remains on the retrieval side. The graph is primarily used to access supporting evidence, which is then flattened into textual context for downstream generation. Even with hierarchical indices such as ArchRAG[Wang et al., 2026], the reasoning process can remain

structurally underconstrained. Retrieved communities are typically converted into textual context for generation, while graph topology and relational distances do not directly define the inference objective. This limits the agent's ability to use the knowledge graph as an active hypothesis space rather than only as a retrieval substrate.

This limitation is especially consequential in high-stakes sequential settings, where deployment-relevant agent performance depends on uncovering missing evidence through iterative, goal-directed interaction[Bobrova, 2026]. Sequential, multi-hop clinical diagnosis provides a representative example, where each intermediate inference can shape subsequent evidence acquisition and hypothesis revision, a process recently formalized as the *Sequential Diagnosis* paradigm[Nori et al., 2025, Trivedi et al., 2023]. As a result, early reasoning errors may propagate across turns and compromise the final decision[Zhao et al., 2025]. Recent empirical analyses further support this concern. LLMs often become path-dependent in extended interactions, committing early to brittle assumptions and failing to recover once the dialogue drifts[Laban et al., 2025]. Moreover, from an information-theoretic perspective, multi-turn interaction can exhibit diminishing information gain and increasing redundancy across turns[Gorle et al., 2025]. Although recent methods introduce iterative exploration, their step-wise search often remains heuristic, making them vulnerable to shallow retrieval, error accumulation, and task-shift failures[Luo et al., 2023, Sun et al., 2023, Lv et al., 2026]. These challenges motivate a shift from likelihood-driven generation and retrieval-augmented context construction toward semantic-structural uncertainty reduction over organized hypothesis spaces[Scanlin et al., 2026].

In this paper, we introduce **S**tructure-**G**rounded **U**ncertainty **R**eduction (**SGUR**), a training-free framework for efficient and reliable multi-turn inference. SGUR reformulates **multi-step reasoning as uncertainty reduction over structured hypothesis spaces**, rather than unconstrained free-form generation. At its core, SGUR couples geometric structural alignment with probabilistic belief updating. It maps the alignment cost between fragmented observations and candidate knowledge graphs into a likelihood score, enabling a principled update of the agent's belief state. By jointly grounding each inference step in semantic evidence, relational structure, and domain-specific priors, SGUR constrains the evolving reasoning trajectory to remain both semantically aligned and structurally coherent under incomplete observations.

Operationally, SGUR uses two coupled mechanisms. (i) a **prior-weighted partial alignment** that enables robust semantic-structural matching under sparse, incomplete evidence without requiring full graph correspondence; and (ii) an **Expected Information Gain (EIG)-driven acquisition strategy** that selects the next most informative query node. This cycle ensures that both hypothesis refinement and evidence acquisition are governed by shared semantic-structural constraints, leading to a rapid and efficient collapse of the candidate hypothesis space.

We evaluate SGUR in sequential clinical diagnosis over a large-scale hierarchical biomedical knowledge graph, a demanding setting in which evidence is incomplete, observations are revealed iteratively, and reliability is essential. Unlike small single-turn question-answering benchmarks, this setting requires the agent to reason over a structured hypothesis space, acquire missing evidence across turns, and maintain consistency over long-horizon inference. Empirically, SGUR delivers significant system-level gains in diagnostic accuracy, inference efficiency, and token cost across strong baselines. Crucially, when paired with smaller backbone models, SGUR achieves performance comparable to or better than baselines powered by substantially larger state-of-the-art models. These results show that reliable agent performance depends not only on model scale, but also on **the quality of the reasoning framework that organizes inference**.

In summary, our contributions are threefold. First, we introduce a training-free **inference paradigm** that reformulates multi-turn reasoning as Bayesian uncertainty contraction over structured hypothesis spaces, providing a principled alternative to unconstrained free-form generation. Second, we develop a unified **semantic-structural reasoning mechanism** that couples prior-weighted partial graph alignment with information-theoretic evidence acquisition, enabling robust belief refinement and goal-directed probing under sparse and sequentially revealed observations. Third, we provide a large-scale evaluation on a hierarchy-structured sequential diagnosis benchmark, demonstrating an empirical **structure-scale substitution effect**: SGUR improves Hit@3 by 9.4 percentage points over the strongest GraphRAG baseline while reducing token overhead by approximately $5\times$, and enables smaller backbone models to match or exceed substantially larger LLMs.

2 Related Work

Reasoning and Retrieval under Weak Structural Control. Recent advances in LLM reasoning have improved performance through structured prompting, retrieval augmentation, and graph-organized evidence. Structured prompting methods, including CoT, ToT, GoT, least-to-most prompting, and self-consistency, improve decomposition and deliberative search[Wei et al., 2022, Yao et al., 2023a, Besta et al., 2024, Zhou et al., 2022, Wang et al., 2022]. In parallel, RAG-style systems improve factual grounding by incorporating external evidence into generation[Lewis et al., 2020, Gao et al., 2023, Borgeaud et al., 2022]. Hybrid approaches such as RAT and BayesRAG variants further interleave retrieval with reasoning or verification[Wang et al., 2024a, Li et al., 2026, Yang et al., 2026, Wan and Yu, 2025, Jiang et al., 2023]. GraphRAG methods extend retrieval-augmented reasoning by organizing evidence into explicit graph structures for local and global querying[Edge et al., 2024a, Guo et al., 2024, Hu et al., 2025, Zhu et al., 2025], including medical variants such as MedGraphRAG and KARE[Wu et al., 2025, Jiang et al., 2024]. Recent hierarchical GraphRAG variants further improve retrieval relevance and efficiency through community-level organization and hierarchical indexing[Wang et al., 2026].

However, these methods still organize reasoning primarily in prompt, text, or retrieval space. Even when graphs are introduced, they are typically used as retrieval scaffolds whose outputs are flattened into textual context for downstream generation[Edge et al., 2024a, Wu et al., 2025, Li et al., 2024, Jiménez Gutiérrez et al., 2024, Wang et al., 2026]. As a result, they may improve evidence access without explicitly constraining the evolving reasoning trajectory itself[Guo et al., 2024, Li et al., 2024, Chen et al., 2024, Liang et al., 2025, Han et al., 2025, Xiang et al., 2025]. The central gap is therefore not only how to retrieve better evidence, but how to make graph structure part of the inference state itself.

Geometric Alignment and Active Uncertainty Reduction. Optimal Transport (OT) provides a principled foundation for graph alignment under soft correspondence. Gromov–Wasserstein (GW) and Fused Gromov–Wasserstein (FGW) compare structured objects by jointly accounting for node semantics and relational structure[Mémoli, 2011, Vayer et al., 2019, 2020, 2018], while scalable approximations improve practicality for larger graphs[Scetbon et al., 2022, Nguyen and Tsuda, 2023]. Partial transport methods relax the full-correspondence requirement and instead focus alignment on overlapping substructures, making them especially relevant under incomplete observations[Chapel et al., 2020, Bai et al., 2025].

A complementary line of work treats reasoning as an interactive process of sequential uncertainty reduction rather than one-shot generation[Yao et al., 2023b, Trivedi et al., 2023]. In medical settings, this appears in active question selection and information-seeking diagnosis[Hu et al., 2024, Wong et al., 2026, Wang et al., 2024b, Nori et al., 2025], while related systems explore adaptive inference control, solver arbitration, and memory management[Bergamaschi Ganapini et al., 2025, Zhang et al., 2025b, You et al., 2026, Chung et al., 2025, Hu et al., 2026]. Yet these approaches generally manage alignment, questioning, and belief refinement as separate procedures in text, memory, or retrieval space. SGUR differs by coupling semantic-structural partial alignment with EIG-driven evidence acquisition within a shared belief state, so that both hypothesis refinement and question selection are governed by the same structure-grounded inference objective.

3 Methodology: The SGUR Framework

SGUR follows the two-phase framework in Figure 1. Section 3.1 describes the offline phase on the left, where expert data are organized into a fixed hierarchical graph database G_{db} under a shared diagnostic schema. Section 3.2 describes the online phase on the right, where each interaction round follows four figure-aligned stages: (a) latent routing first aligns the current observations $\mathcal{O}_t$ with database-level observation anchors and uses the matched anchors to construct the query graph $G_q^{(t)}$; (b) candidate pruning retrieves the associated candidate hypothesis graphs $G_{\mathrm{cand}}^{(t)}$ using path filtering and specificity-aware structural priors; (c) prior-weighted partial Fused Gromov–Wasserstein (PFGW) alignment compares $G_q^{(t)}$ with each candidate graph and induces a belief ranking; and (d) active acquisition selects informative unobserved nodes to obtain $\mathcal{O}_{t+1}$. Together, the fixed expert database and the online loop define SGUR as a structure-grounded inference procedure that progressively contracts the candidate hypothesis space under partial observation.

3.1 Offline Expert Graph Construction

The offline phase constructs the fixed expert graph database used by the online inference loop. As shown on the left side of Figure 1, SGUR extracts hierarchical entities from expert data, maps them into a shared diagnostic schema, and assembles directed inter-layer relations between adjacent semantic layers. In our PrimeKG-derived clinical instantiation, the schema contains observations (L_1), localizations (L_2), candidate diagnoses (L_3), and treatments (L_4). SGUR also builds an auxiliary L_0 routing layer by clustering observation embeddings; L_0 serves as a routing index for fast query-graph construction, not as a diagnostic semantic layer.

Formally, the expert graph database is

$$G_{\mathrm{db}} = \{G_{d_j}\}_{j=1}^{n}, \tag{1}$$

where each diagnosis-centered graph is

$$G_{d_j} = (\mathcal{V}_{d_j}, \mathcal{E}_{d_j}, \omega_{d_j}, \mathbf{X}_{d_j}, \ell). \tag{2}$$

Here, $\mathcal{V}_{d_j}$ is the node set, $\mathcal{E}_{d_j} \subseteq \mathcal{V}_{d_j} \times \mathcal{V}_{d_j}$ is the directed edge set, $\mathbf{X}_{d_j}$ stores node features, and $\ell : \mathcal{V}_{d_j} \to \{L_1, L_2, L_3, L_4\}$ maps nodes to semantic layers. In our benchmark, each d_j corresponds to a disease, but the construction only requires a structured diagnostic schema.

The edge weights ω_{d_j} instantiate the specificity-aware structural priors. For each outgoing relation $(u, v) \in \mathcal{E}_{d_j}$, SGUR uses

$$\omega_{uv} = \frac{1}{|\mathcal{N}_{\mathrm{out}}(u)|}, \tag{3}$$

where $\mathcal{N}_{\mathrm{out}}(u)$ denotes the outgoing neighbors of u. This down-weights generic nodes with many downstream continuations and emphasizes more discriminative relations during candidate pruning and prior-weighted partial FGW alignment. During online inference, G_{db} is kept fixed: SGUR grounds current observations, routes them to a candidate subset, aligns the query graph with candidate graphs, and updates the belief ranking without reconstructing or updating the expert database.

3.2 Online Interactive Inference

The online phase performs closed-loop inference over the fixed database G_{db}. At round t, the accumulated evidence $\mathcal{O}_t$ is represented as a partial query graph

$$G_q^{(t)} = \Gamma(\mathcal{O}_t) = (\mathcal{V}_q^{(t)}, \mathcal{E}_q^{(t)}, \omega_q^{(t)}, \mathbf{X}_q^{(t)}, \ell_q^{(t)}), \tag{4}$$

where $\Gamma(\cdot)$ maps extracted entities into the shared schema. The online loop refines a belief ranking over candidate graphs while acquiring the next evidence that is most useful under the current ranking.

(a) Query graph construction. SGUR first grounds the current observations into $G_q^{(t)}$. To avoid aligning every query against the full database, the auxiliary L_0 routing layer localizes the query to a coarse observation neighborhood before fine-grained matching. After passing through L_0, entities are extracted from the query (see Appendix A.3 for the entity recognition prompt) and used to construct the query graph $G_q^{(t)}$. This corresponds to the query graph construction block in Figure 1.

(b) Candidate pruning and structural priors. After latent routing, SGUR applies path filtering to retain only candidate graphs that are structurally compatible with the current query:

$$G_{\mathrm{cand}}^{(t)} = \left\{ G_{d_j} \in G_{\mathrm{db}} : R_0(G_q^{(t)}, G_{d_j}) = 1,\ F_{\mathrm{path}}(G_q^{(t)}, G_{d_j}) = 1 \right\}, \tag{5}$$

where R_0 denotes the L_0-based routing condition and F_{path} denotes inter-layer path compatibility. SGUR reuses the specificity-aware relation weights in Eq. (3). For node-level saliency, for any graph G we define

$$S_G(v) = \log\big(1 + |\mathcal{N}_{\mathrm{in}}^G(v)| + |\mathcal{N}_{\mathrm{out}}^G(v)|\big), \qquad \pi_v^G = \frac{S_G(v)}{\sum_{u \in \mathcal{V}_G} S_G(u)}. \tag{6}$$

The normalized saliency π_v^G provides node masses for alignment and utility scores for active acquisition.

(c) Prior-weighted partial FGW alignment and belief ranking. For each $G_{d_j} \in G_{\text{cand}}^{(t)}$, SGUR compares $G_q^{(t)}$ with G_{d_j} using prior-weighted PFGW alignment. Let $\mathbf{M}^{(j,t)}$ be the node feature discrepancy matrix, instantiated as cosine distances between query-node and candidate-node embeddings. Let $C_q^{(t)}$ and C_{d_j} denote weighted shortest-path distance matrices encoding internal graph structure. We set the candidate and query node masses as $\mathbf{p}^{(j)} = (\pi_a^{G_{d_j}})_{a \in \mathcal{V}_{d_j}}$ and $\mathbf{q}^{(t)} = (\pi_b^{G_q^{(t)}})_{b \in \mathcal{V}_q^{(t)}}$. SGUR computes

$$D_{\text{PFGW}}(G_q^{(t)}, G_{d_j}) = \min_{\gamma \in \Pi_m(\mathbf{p}^{(j)}, \mathbf{q}^{(t)})} \Big\{ (1-\alpha)\langle \gamma, \mathbf{M}^{(j,t)} \rangle + \alpha \sum_{a,a',b,b'} \Big| C_{d_j}(a,a') - C_q^{(t)}(b,b') \Big|^2 \gamma_{ab} \gamma_{a'b'} \Big\}. \tag{7}$$

where $\Pi_m(\mathbf{p}^{(j)}, \mathbf{q}^{(t)})$ is the set of partial couplings with transported mass m. The partial coupling allows unmatched query evidence and unmatched expert structure to remain uncoupled, which is essential when the query covers only a sparse substructure of the candidate graph.

The resulting discrepancies are converted into a normalized belief ranking:

$$P^{(t)}(d_j) = \frac{\exp(-D_{\text{PFGW}}(G_q^{(t)}, G_{d_j})/\beta)}{\sum_{G_{d_k} \in G_{\text{cand}}^{(t)}} \exp(-D_{\text{PFGW}}(G_q^{(t)}, G_{d_k})/\beta)}, \tag{8}$$

where $\beta > 0$ controls ranking sharpness. The induced rank is

$$r_t(d_j) = \text{rank}\big(P^{(t)}(d_j)\big), \qquad G_{d_j} \in G_{\text{cand}}^{(t)}. \tag{9}$$

We measure remaining uncertainty by the entropy of this candidate belief distribution:

$$H^{(t)} = -\sum_{G_{d_j} \in G_{\text{cand}}^{(t)}} P^{(t)}(d_j) \log P^{(t)}(d_j). \tag{10}$$

Thus, candidate-space contraction is realized as progressive concentration of $P^{(t)}$ across interaction rounds.

(d) Active acquisition. SGUR selects the next evidence from unobserved L_1/L_2 nodes in the top-κ candidate graphs. The ideal objective is to choose the query that most reduces uncertainty in the belief distribution:

$$\text{EIG}(s) = \mathbb{E}_{r \sim p(r|s,P^{(t)})}\big[H^{(t)} - H(P^{(t+1)} \mid r, s)\big]. \tag{11}$$

Exact EIG requires a response model over all possible user answers. SGUR therefore uses the structure-guided proxy shown in Figure 1. Let $\mathcal{D}_t^\kappa$ be the top-κ candidates under $P^{(t)}$, and define the acquisition pool

$$\mathcal{A}_t = \bigcup_{d_j \in \mathcal{D}_t^\kappa} \mathcal{V}_{d_j}^{L_1 \cup L_2} \setminus \mathcal{V}_q^{(t)}. \tag{12}$$

For each $v \in \mathcal{A}_t$, SGUR scores

$$u_t(v) = \sum_{d_j \in \mathcal{D}_t^\kappa} \pi_v^{G_{d_j}}, \tag{13}$$

and selects the next query batch as

$$\mathcal{A}_t^\star = \text{top-}b\big(\{u_t(v) : v \in \mathcal{A}_t\}\big). \tag{14}$$

The selected query is presented to the user, the returned observation is incorporated into $\mathcal{O}_{t+1}$, and the pruning–alignment–ranking–acquisition loop repeats.

Termination and output. The interaction terminates when the turn budget T is exhausted or when no informative unobserved node remains. SGUR outputs the top-k candidate diagnoses under the terminal belief ranking:

$$\widehat{\mathcal{C}}(\mathcal{O}_\tau) = \text{top-}k\big(\{P^{(\tau)}(d_j) : G_{d_j} \in G_{\text{cand}}^{(\tau)}\}\big). \tag{15}$$

For evaluation, we record the first round in which the ground-truth diagnosis enters the top-k set, yielding the MTTC metric used in Section 4.

Table 1: Main accuracy, convergence, and efficiency results. Final retrieval metrics are measured after round 3 in the complex two-symptom regime; convergence reports complex-regime Hit@3 across rounds; efficiency reports three-round average total tokens and latency. All Hit@K values are percentages. SGUR achieves the best final ranking quality, lowest MTTC, and lowest inference cost.

Method	Complex regime after R3					Complex Hit@3			Efficiency Avg.	
	MTTC ↓	MRR ↑	Hit@1 ↑	Hit@2 ↑	Hit@3 ↑	R1	R2	R3	Tokens ↓	Latency ↓
GraphRAG	1.74	0.71	61.20	77.60	83.10	59.60	76.90	83.10	10454.00	34.90
NodeRAG	1.90	0.71	60.20	75.20	80.10	54.20	75.00	80.10	8461.10	15.20
HippoRAG2	1.83	0.68	59.50	74.50	81.70	60.40	77.40	81.70	2352.80	11.40
ArchRAG	2.01	0.65	56.70	70.40	76.10	56.70	70.40	76.10	2683.40	10.40
SGUR (Ours)	**1.62**	**0.77**	**64.90**	**83.40**	**92.50**	**61.40**	**84.70**	**92.50**	**1845.40**	**8.90**

4 Experiments

We evaluate SGUR on a multi-round diagnostic retrieval task that stresses inference under incomplete and progressively revealed evidence. The experiments are organized into two parts. Section 4.1 compares SGUR with graph-based RAG baselines in terms of accuracy, convergence, efficiency, and robustness to model scale. Section 4.2 isolates the contribution of SGUR's core mechanisms.

Experimental setup. We construct a controlled multi-round diagnostic retrieval benchmark from **PrimeKG** [Chandak et al., 2023]. Each case is anchored by a ground-truth diagnosis $d^\star$, and evidence is progressively revealed over three interaction rounds. PrimeKG provides typed symptom–anatomy–disease–treatment relations, enabling hierarchical graph construction and database-derived edge-prior initialization. By contrast, standard agent, tool-use, web-interaction, and GraphRAG benchmarks mainly evaluate final-answer accuracy, tool execution, short-horizon task completion, or retrieval-generation quality [Liu et al., 2024, Mialon et al., 2023, Qin et al., 2023, Zhou et al., 2023, Xiao et al., 2025, Xiang et al., 2025], but typically do not expose the typed hierarchy and relational statistics required to evaluate structure-grounded uncertainty reduction.

For each disease, we randomly permute its associated symptoms once and reveal the same sequence to all methods to remove ordering randomness. We evaluate two regimes: **Simple Diagnosis**, with $b = 4$ newly revealed symptoms per round, and **Complex Diagnosis**, with $b = 2$ newly revealed symptoms per round. The simple regime contains 854 test cases and the complex regime contains 615 test cases, yielding 1,469 reasoning cases in total. We emphasize the complex regime in the main results because it best captures sparse and ambiguous evidence under partial observation. Additional construction details are provided in Appendix A.1.

4.1 Comparison with Graph-RAG Baselines

Baselines and Multi-round Adaptation. To verify the superiority of structure-grounded reasoning, we compare **SGUR** against four representative graph-RAG paradigms: the PageRank-based **HippoRAG2** [Jiménez Gutiérrez et al., 2025], node-centric **NodeRAG** [Xu et al., 2025], community-level **GraphRAG** [Edge et al., 2024b], and the hierarchical **ArchRAG** [Wang et al., 2026]. Since these baselines are primarily designed for single-shot inference, we adapt them to the three-round setting using a **textual context re-injection** scheme. In this setup, each baseline receives a compact textual summary of prior predictions and supporting evidence as conversational context in subsequent rounds. This ensures all methods access the same revealed evidence while preserving their native retrieval mechanisms. In contrast, SGUR facilitates a transition from context stacking to structured state evolution. By maintaining an incrementally updated query graph $G_q^{(t)}$, SGUR accumulates evidence directly as graph state, circumventing the semantic drift and computational redundancy inherent in repeatedly re-injecting dialogue history as text.

Metrics, Results, and Implementation Details. We assess retrieval effectiveness through **Hit@**K ($K \in \{1, 2, 3\}$), **MRR@R3**, and **Mean Turns to Confidence (MTTC)**, defined as the first round in which the ground-truth diagnosis enters the top-3 candidate set, with a penalty of $T + 1 = 4$ for unretrieved cases. System efficiency is quantified via total token consumption and end-to-end latency. All models utilize `text-embedding-3-large` OpenAI [2024a] for indexing and retrieval. Unless

Table 2: Model-scaling stress test in the complex regime. Each method is evaluated with large and small backbones; retention reports the small/large ratio. SGUR preserves over 99% Hit@3 after backbone scaling and remains above all large-backbone graph-RAG baselines.

Method	Backbone	Hit@3 ↑	MRR ↑	MTTC ↓
GraphRAG	GPT-5.2 (L)	83.10	0.71	1.74
	GPT-5.4-nano (S)	74.70	0.62	2.02
	Retention (%)	90.30	87.10	86.10
NodeRAG	GPT-5.2 (L)	80.10	0.70	1.90
	GPT-5.4-nano (S)	73.60	0.60	2.12
	Retention (%)	91.90	84.50	89.60
HippoRAG2	GPT-5.2 (L)	81.70	0.68	1.83
	GPT-5.4-nano (S)	79.80	0.60	2.11
	Retention (%)	97.70	87.30	86.70
ArchRAG	GPT-5.2 (L)	76.10	0.65	2.01
	GPT-5.4-nano (S)	68.70	0.57	2.30
	Retention (%)	90.20	87.40	87.30
SGUR	GPT-5.2 (L)	**92.50**	**0.77**	**1.62**
	GPT-5.4-nano (S)	**92.40**	**0.73**	**1.63**
	Retention (%)	**99.80**	**94.60**	**99.40**

otherwise specified, online inference is evaluated with `GPT-5.2` OpenAI [2025a], and robustness to backbone scale is tested with `GPT-5.4-nano` OpenAI [2026].

As detailed in Table 1, SGUR achieves the strongest accuracy–efficiency profile in the complex diagnostic regime. It reaches **92.50% Hit@3**, outperforming the strongest graph-RAG baseline by **9.40 percentage points**, while also obtaining the lowest MTTC. Its Hit@3 improves from 61.40% to 92.50% across rounds, indicating that the structured query state progressively narrows the candidate hypothesis space as new evidence is acquired. Furthermore, SGUR reduces token consumption by **5.70**× and latency by **3.90**× compared with GraphRAG. Together, these results show that SGUR improves diagnostic reliability and lowers the online cost of multi-turn inference by organizing reasoning through structure-grounded state updates, rather than relying solely on larger backbone models.

Robustness to backbone scale. We further test whether SGUR's advantage persists when the backbone model is scaled down. Table 2 shows that SGUR with the smaller **GPT-5.4-nano** backbone remains above all graph-RAG baselines using the larger **GPT-5.2** backbone. Moreover, SGUR retains over 99% of its large-backbone Hit@3 after scaling down, while the baselines exhibit larger drops in ranking quality and convergence. This result supports the central claim that reliable multi-turn inference depends not only on backbone scale, but also on how the reasoning process is structurally organized. Full scaling results, including Hit@1 and Hit@2, are provided in Table 7 in Appendix A.6.

Accuracy–efficiency trade-off. Figure 2 visualizes the system-level trade-off between final Hit@3 and average token cost. SGUR occupies the upper-left Pareto-optimal region, achieving the highest retrieval accuracy with the lowest inference budget among all methods. This demonstrates that SGUR shifts the accuracy–efficiency frontier, rather than improving accuracy by using larger prompts, longer interaction histories, or higher online computation.

4.2 Ablation Study

Component contribution under sparse evidence. Table 3 jointly evaluates component necessity and evidence sufficiency. Removing the L_0 latent routing layer causes the largest degradation, reducing final Hit@3 from 92.50% to 68.90% and increasing MTTC from 1.62 to 2.31. This shows that latent routing is not merely a computational shortcut: by aligning raw observations to database-level anchors before query-graph construction, it provides the grounding needed to build focused query graphs and retrieve relevant candidate hypotheses. Replacing partial transport with full FGW also causes a substantial drop, reducing Hit@3 to 73.30% and increasing MTTC to 2.45. This confirms that

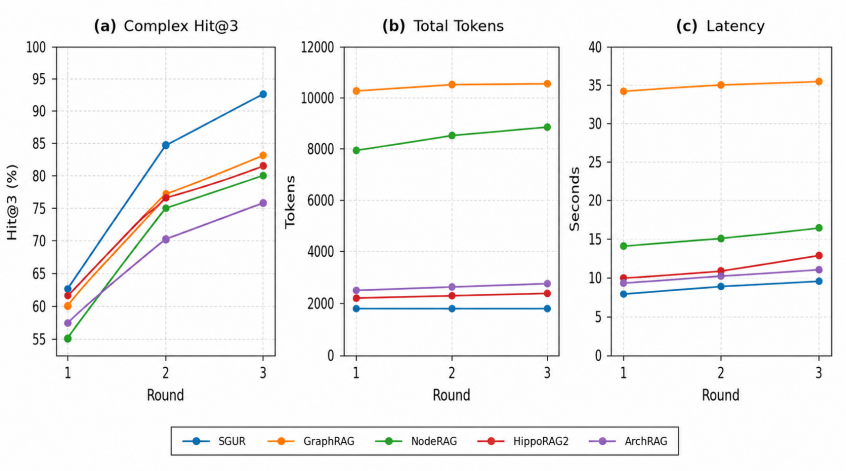

Figure 2: Accuracy–efficiency Pareto frontier in the complex regime. The x-axis reports average total tokens over three interaction rounds, and the y-axis reports final Hit@3 after round 3. Higher is better on the y-axis and lower is better on the x-axis. SGUR achieves the highest retrieval accuracy with the lowest inference cost.

Table 3: Combined ablation study in the complex regime. The left block isolates each SGUR component at the final round, while the right block reports Hit@3 as observed evidence increases. All Hit@3 and ΔHit@3 values are percentages.

Variant	Final complex-regime performance				Evidence sufficiency: Hit@3		
	Hit@3 ↑	ΔHit@3	MRR ↑	MTTC ↓	2 Obs.	4 Obs.	6 Obs.
Full SGUR	**92.50**	–	**0.77**	**1.62**	**61.40**	**85.70**	**92.50**
Semantic-only Alignment	91.90	−0.60	0.75	1.64	59.00	85.30	91.90
Structure-only Alignment	91.60	−0.90	0.76	1.67	56.20	84.40	91.60
Uniform Importance	83.40	−9.10	0.64	2.01	38.20	71.80	83.40
Full FGW	73.30	−19.20	0.53	2.45	15.30	59.50	73.40
w/o Latent Routing	68.90	−23.60	0.57	2.31	35.70	59.80	68.90

incomplete observations should be matched only to supported substructures rather than forced into full graph correspondence. Uniform importance reduces Hit@3 to 83.40%, indicating that database-derived specificity priors help distinguish generic relations from diagnostically informative ones. Finally, semantic-only and structure-only variants remain competitive but consistently underperform the full model, showing that SGUR benefits from coupling node-level evidence with relational consistency.

Effect of evidence sufficiency. The right block of Table 3 shows that SGUR's advantage is most pronounced when evidence is sparse. With only two observed symptoms, Full SGUR reaches 61.40% Hit@3, whereas removing latent routing drops performance to 35.70%, and replacing partial transport with full FGW drops it further to 15.30%. As additional observations are revealed, all variants improve, but the full model remains consistently strongest. This pattern supports the central mechanism of SGUR: structure-grounded routing, partial alignment, and specificity-aware weighting are most valuable when the agent must reduce uncertainty from incomplete evidence rather than simply rank hypotheses from abundant observations.

5 Discussion and Conclusion

Our results support a central thesis: reliable multi-turn inference depends not only on the scale of the backbone model, but also on how the reasoning process is organized. SGUR improves sequential inference by turning fragmented observations into a structured query state, aligning this state with

candidate hypothesis graphs through prior-weighted partial matching, and actively selecting the next evidence node to reduce uncertainty. Across accuracy, convergence, latency, token cost, and model-scaling experiments, this structure-grounded design consistently improves performance under sparse and progressively revealed evidence. These findings suggest that reasoning frameworks can partially substitute for model scale by providing the structural discipline that free-form generation often lacks.

SGUR should nevertheless be interpreted as a research framework, not as a deployed clinical decision system. We instantiate the framework in sequential clinical diagnosis because it provides a demanding testbed with incomplete evidence, multi-round interaction, and structured hypothesis spaces. The system is not intended for real-world diagnosis, treatment recommendation, or medical decision-making without expert validation, clinician oversight, and appropriate regulatory review.

The framework also relies on several assumptions. It assumes that domain knowledge can be organized into a meaningful hierarchy, that observations can be mapped to standardized evidence nodes, and that the expert graph is sufficiently reliable during inference. Although PrimeKG provides a controlled large-scale benchmark, real-world settings may contain noisy observations, missing relations, evolving knowledge, and heterogeneous evidence sources. Future work should evaluate SGUR on broader long-horizon agent benchmarks and extend it to domains such as technical support, device maintenance, and software debugging, where agents must iteratively gather evidence, refine hypotheses, and make decisions under structural constraints.

References

Yikun Bai, Shuang Wang, Huy Tran, Hengrong Du, Juexin Wang, and Soheil Kolouri. Fused partial gromov–wasserstein for structured objects. *arXiv preprint arXiv:2502.09934*, 2025. URL https://arxiv.org/abs/2502.09934.

M Bergamaschi Ganapini, M Campbell, F Fabiano, L Horesh, J Lenchner, A Loreggia, N Mattei, F Rossi, B Srivastava, and K B Venable. Fast, slow, and metacognitive thinking in AI. *npj Artificial Intelligence*, 1(1):27, 2025. doi: 10.1038/s44387-025-00027-5. URL https://www.nature.com/articles/s44387-025-00027-5.

Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Niewiadomski, Piotr Nyczyk, and Torsten Hoefler. Graph of thoughts: Solving elaborate problems with large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 17682–17690, 2024. URL https://ojs.aaai.org/index.php/AAAI/article/view/29720.

Natalia Bobrova. Clinical agentic systems: Methodologies for development, orchestration, and governance of autonomous AI agents in healthcare workflows—alignment with 2026 FDA/EMA guiding principles. SSRN, 2026. URL https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6350259. Preprint, SSRN abstract 6350259.

Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George Bm Van Den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, Diego De Las Casas, Aurelia Guy, Jacob Menick, Roman Ring, Tom Hennigan, Saffron Huang, Loren Maggiore, Chris Jones, Albin Cassirer, Andy Brock, Michela Paganini, Geoffrey Irving, Oriol Vinyals, Simon Osindero, Karen Simonyan, Jack Rae, Erich Elsen, and Laurent Sifre. Improving language models by retrieving from trillions of tokens. In *International Conference on Machine Learning*, pages 2206–2240. PMLR, 2022. URL https://proceedings.mlr.press/v162/borgeaud22a.html.

Payal Chandak, Kexin Huang, and Marinka Zitnik. Building a knowledge graph to enable precision medicine. *Scientific Data*, 10(1):67, 2023. doi: 10.1038/s41597-023-01960-3. URL https://doi.org/10.1038/s41597-023-01960-3.

Laetitia Chapel, Mokhtar Z Alaya, and Gilles Gasso. Partial optimal transport with applications on positive-unlabeled learning. *Advances in Neural Information Processing Systems*, 33:2903–2913, 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/hash/1e6e25d952a0d639b676ee20d0519ee2-Abstract.html.

Rubing Chen, Xulu Zhang, Jiaxin Wu, Wenqi Fan, Xiao-Yong Wei, and Qing Li. Knowledge pyramid construction for multi-level retrieval-augmented generation. *arXiv preprint arXiv:2407.21276*, 2024. URL https://arxiv.org/abs/2407.21276. v1 (2024) was titled "Multi-Level Querying using A Knowledge Pyramid"; retitled in v3 (2025).

Stephen Chung, Wenyu Du, and Jie Fu. Thinker: Learning to think fast and slow. *arXiv preprint arXiv:2505.21097*, 2025. URL https://arxiv.org/abs/2505.21097.

Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitansky, Robert Osazuwa Ness, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024a. URL https://arxiv.org/abs/2404.16130.

Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, Dasha Metropolitansky, Robert Osazuwa Ness, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024b. URL https://arxiv.org/abs/2404.16130.

Alina Fastowski and Gjergji Kasneci. Understanding knowledge drift in LLMs through misinformation. In *International Workshop on Discovering Drift Phenomena in Evolving Landscapes*, pages 74–85. Springer, 2025. doi: 10.1007/978-3-031-82346-6_5. URL https://link.springer.com/chapter/10.1007/978-3-031-82346-6_5.

Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2023. URL https://arxiv.org/abs/2312.10997.

Abhiram Rao Gorle, Amit Kumar Singh Yadav, and Tsachy Weissman. Quantifying information gain and redundancy in multi-turn LLM conversations. In *NeurIPS 2025 Workshop on Multi-Turn Interactions in Large Language Models*, 2025. URL https://openreview.net/forum?id=5gpABTkcUJ.

Zirui Guo, Lianghao Xia, Yanhua Yu, Tu Ao, and Chao Huang. LightRAG: Simple and fast retrieval-augmented generation. *arXiv preprint arXiv:2410.05779*, 2024. URL https://arxiv.org/abs/2410.05779v1.

Haoyu Han, Li Ma, Harry Shomer, Yu Wang, Yongjia Lei, Kai Guo, Zhigang Hua, Bo Long, Hui Liu, Charu C. Aggarwal, and Jiliang Tang. RAG vs. GraphRAG: A systematic evaluation and key insights. *arXiv preprint arXiv:2502.11371*, 2025. URL https://arxiv.org/abs/2502.11371v2.

Yuntong Hu, Zhihan Lei, Zheng Zhang, Bo Pan, Chen Ling, and Liang Zhao. GRAG: Graph retrieval-augmented generation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 4145–4157, 2025. URL https://aclanthology.org/2025.findings-naacl.232/.

Zhanghao Hu, Qinglin Zhu, Runcong Zhao, Di Liang, Hanqi Yan, Yulan He, and Lin Gui. Beyond RAG for agent memory: Retrieval by decoupling and aggregation. *arXiv preprint arXiv:2602.02007*, 2026. URL https://arxiv.org/abs/2602.02007.

Zhiyuan Hu, Chumin Liu, Xidong Feng, Yilun Zhao, See-Kiong Ng, Anh Tuan Luu, Junxian He, Pang Wei Koh, and Bryan Hooi. Uncertainty of thoughts: Uncertainty-aware planning enhances information seeking in large language models. In *Advances in Neural Information Processing Systems*, volume 37, pages 24181–24215, 2024. URL https://arxiv.org/abs/2402.03271.

Pengcheng Jiang, Cao Xiao, Minhao Jiang, Parminder Bhatia, Taha Kass-Hout, Jimeng Sun, and Jiawei Han. Reasoning-enhanced healthcare predictions with knowledge graph community retrieval. *arXiv preprint arXiv:2410.04585*, 2024. URL https://arxiv.org/abs/2410.04585.

Zhengbao Jiang, Frank F Xu, Luyu Gao, Zhiqing Sun, Qian Liu, Jane Dwivedi-Yu, Yiming Yang, Jamie Callan, and Graham Neubig. Active retrieval augmented generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7969–7992, 2023. URL https://aclanthology.org/2023.emnlp-main.495/.

Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. *Advances in Neural Information Processing Systems*, 37:59532–59569, 2024. URL https://arxiv.org/abs/2405.14831.

Bernal Jiménez Gutiérrez, Yiheng Shu, Weijian Qi, Sizhe Zhou, and Yu Su. From RAG to memory: Non-parametric continual learning for large language models. *arXiv preprint arXiv:2502.14802*, 2025. doi: 10.48550/arXiv.2502.14802. URL https://arxiv.org/abs/2502.14802.

Daniel Kahneman. *Thinking, Fast and Slow*. Macmillan, 2011. URL https://us.macmillan.com/books/9780374275631/thinkingfastandslow/.

Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. LLMs get lost in multi-turn conversation. *arXiv preprint arXiv:2505.06120*, 2025. URL https://arxiv.org/abs/2505.06120.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020. URL https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html.

Xuan Li, Yining Wang, Haocai Luo, Shengping Liu, Jerry Liang, Ying Fu, Weihuang, Jun Yu, and Junnan Zhu. BayesRAG: Probabilistic mutual evidence corroboration for multimodal retrieval-augmented generation. *arXiv preprint arXiv:2601.07329*, 2026. URL https://arxiv.org/abs/2601.07329.

Zhuoqun Li, Xuanang Chen, Haiyang Yu, Hongyu Lin, Yaojie Lu, Qiaoyu Tang, Fei Huang, Xianpei Han, Le Sun, and Yongbin Li. StructRAG: Boosting knowledge intensive reasoning of LLMs via inference-time hybrid information structurization. *arXiv preprint arXiv:2410.08815*, 2024. URL https://arxiv.org/abs/2410.08815.

Lei Liang, Zhongpu Bo, Zhengke Gui, Zhongshu Zhu, Ling Zhong, Peilong Zhao, Mengshu Sun, Zhiqiang Zhang, Jun Zhou, Wenguang Chen, et al. KAG: Boosting LLMs in professional domains via knowledge augmented generation. In *Companion Proceedings of the ACM on Web Conference 2025*, pages 334–343, 2025.

Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. Agentbench: Evaluating llms as agents. In *International Conference on Learning Representations*, 2024. URL https://arxiv.org/abs/2308.03688.

Linhao Luo, Yuan-Fang Li, Gholamreza Haffari, and Shirui Pan. Reasoning on graphs: Faithful and interpretable large language model reasoning. *arXiv preprint arXiv:2310.01061*, 2023. URL https://arxiv.org/abs/2310.01061.

Can Lv, Heng Chang, Shengyu Tao, Mingju Chen, Zhaoxin Fan, Ziwei Zhang, Yuchen Guo, and Shiji Zhou. All-mem: Agentic lifelong memory via dynamic topology evolution. *arXiv preprint arXiv:2603.19595*, 2026. URL https://arxiv.org/abs/2603.19595.

Facundo Mémoli. Gromov–wasserstein distances and the metric approach to object matching. *Foundations of Computational Mathematics*, 11:417–487, 2011. URL https://www.researchwithrutgers.org/en/publications/gromov-wasserstein-distances-and-the-metric-approach-to-object-ma/.

Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. GAIA: A benchmark for general AI assistants. *arXiv preprint arXiv:2311.12983*, 2023. URL https://arxiv.org/abs/2311.12983.

Dai Hai Nguyen and Koji Tsuda. On a linear fused gromov–wasserstein distance for graph structured data. *Pattern Recognition*, 138:109351, 2023. URL https://www.sciencedirect.com/science/article/abs/pii/S0031320323000523.

Harsha Nori, Mayank Daswani, Christopher Kelly, Scott Lundberg, Marco Tulio Ribeiro, Marc Wilson, Xiaoxuan Liu, Viknesh Sounderajah, Jonathan Carlson, Matthew P. Lungren, Bay Gross, Peter Hames, Mustafa Suleyman, Dominic King, and Eric Horvitz. Sequential diagnosis with language models. *arXiv preprint arXiv:2506.22405*, 2025. URL https://arxiv.org/abs/2506.22405.

OpenAI. New embedding models and API updates, 2024a. URL https://openai.com/index/new-embedding-models-and-api-updates/. Introduces text-embedding-3-large.

OpenAI. New embedding models and API updates, 2024b. URL https://openai.com/index/new-embedding-models-and-api-updates/. Introduces text-embedding-3-large.

OpenAI. Introducing GPT-5.2, 2025a. URL https://openai.com/index/introducing-gpt-5-2/.

OpenAI. Introducing GPT-5.2, 2025b. URL https://openai.com/index/introducing-gpt-5-2/.

OpenAI. Update to GPT-5 system card: GPT-5.2, 2025c. URL https://openai.com/index/gpt-5-system-card-update-gpt-5-2/.

OpenAI. Introducing GPT-5.4 mini and nano, March 2026. URL https://openai.com/index/introducing-gpt-5-4-mini-and-nano/.

Boci Peng, Yun Zhu, Yongchao Liu, Xiaohe Bo, Haizhou Shi, Chuntao Hong, Yan Zhang, and Siliang Tang. Graph retrieval-augmented generation: A survey. *arXiv preprint arXiv:2408.08921*, 2024. URL https://arxiv.org/abs/2408.08921.

Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. *arXiv preprint arXiv:2307.16789*, 2023. URL https://arxiv.org/abs/2307.16789.

Joe Scanlin, Armando Cuesta, and Martín Varsavsky. Representation before retrieval: Structured patient artifacts reduce hallucination in clinical AI systems. *medRxiv*, 2026. doi: 10.64898/2026.02.13.26346256. URL https://www.medrxiv.org/content/10.64898/2026.02.13.26346256v1. Preprint, not peer reviewed.

Meyer Scetbon, Gabriel Peyré, and Marco Cuturi. Linear-time Gromov–Wasserstein distances using low rank couplings and costs. In *International Conference on Machine Learning*, pages 19347–19365. PMLR, 2022. URL https://proceedings.mlr.press/v162/scetbon22b.html.

Jiashuo Sun, Chengjin Xu, Lumingyuan Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Lionel M Ni, Heung-Yeung Shum, and Jian Guo. Think-on-graph: Deep and responsible reasoning of large language model on knowledge graph. *arXiv preprint arXiv:2307.07697*, 2023. URL https://arxiv.org/abs/2307.07697.

Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10014–10037, 2023. URL https://aclanthology.org/2023.acl-long.557/.

Titouan Vayer, Laetitia Chapel, Rémi Flamary, Romain Tavenard, and Nicolas Courty. Optimal transport for structured data with application on graphs. *arXiv preprint arXiv:1805.09114*, 2018. URL https://arxiv.org/abs/1805.09114.

Titouan Vayer, Laetitia Chapel, Rémi Flamary, Romain Tavenard, and Nicolas Courty. Optimal transport for structured data with application on graphs. In *International Conference on Machine Learning*, pages 6275–6284. PMLR, 2019. URL https://proceedings.mlr.press/v97/titouan19a.html.

Titouan Vayer, Laetitia Chapel, Rémi Flamary, Romain Tavenard, and Nicolas Courty. Fused Gromov–Wasserstein distance for structured objects. *Algorithms*, 13(9):212, 2020. doi: 10.3390/a13090212. URL https://www.mdpi.com/1999-4893/13/9/212.

Xueyao Wan and Hang Yu. MMGraphRAG: Bridging vision and language with interpretable multimodal knowledge graphs. *arXiv preprint arXiv:2507.20804*, 2025. URL https://arxiv.org/abs/2507.20804.

Shu Wang, Yixiang Fang, Yingli Zhou, Xilin Liu, and Yuchi Ma. ArchRAG: Attributed community-based hierarchical retrieval-augmented generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 15868–15876, 2026. doi: 10.1609/aaai.v40i19.38619. URL https://ojs.aaai.org/index.php/AAAI/article/view/38619.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022. URL https://arxiv.org/abs/2203.11171.

Zihao Wang, Anji Liu, Haowei Lin, Jiaqi Li, Xiaojian Ma, and Yitao Liang. RAT: Retrieval augmented thoughts elicit context-aware reasoning in long-horizon generation. *arXiv preprint arXiv:2403.05313*, 2024a. URL https://arxiv.org/abs/2403.05313.

Ziyu Wang, Hao Li, Di Huang, and Amir M. Rahmani. HealthQ: Unveiling questioning capabilities of LLM chains in healthcare conversations. *arXiv preprint arXiv:2409.19487*, 2024b. URL https://arxiv.org/abs/2409.19487v1.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H Chi, Quoc V Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837, 2022. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html.

Hui Min Wong, Philip Heesen, Pascal Janetzky, Martin Bendszus, and Stefan Feuerriegel. MedClarify: An information-seeking AI agent for medical diagnosis with case-specific follow-up questions. *arXiv preprint arXiv:2602.17308*, 2026. URL https://arxiv.org/abs/2602.17308.

Junde Wu, Jiayuan Zhu, Yunli Qi, Jingkun Chen, Min Xu, Filippo Menolascina, Yueming Jin, and Vicente Grau. Medical graph RAG: Evidence-based medical large language model via graph retrieval-augmented generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 28443–28467, 2025. URL https://aclanthology.org/2025.acl-long.1381/.

Zhishang Xiang, Chuanjie Wu, Qinggang Zhang, Shengyuan Chen, Zijin Hong, Xiao Huang, and Jinsong Su. When to use graphs in RAG: A comprehensive analysis for graph retrieval-augmented generation. *arXiv preprint arXiv:2506.05690*, 2025. URL https://arxiv.org/abs/2506.05690.

Yilin Xiao, Junnan Dong, Chuang Zhou, Su Dong, Qian-wen Zhang, Di Yin, Xing Sun, and Xiao Huang. GraphRAG-Bench: Challenging domain-specific reasoning for evaluating graph retrieval-augmented generation. *arXiv preprint arXiv:2506.02404*, 2025. URL https://arxiv.org/abs/2506.02404.

Tianyang Xu, Haojie Zheng, Chengze Li, Haoxiang Chen, Yixin Liu, Ruoxi Chen, and Lichao Sun. NodeRAG: Structuring graph-based RAG with heterogeneous nodes. *arXiv preprint arXiv:2504.11544*, 2025. doi: 10.48550/arXiv.2504.11544. URL https://arxiv.org/abs/2504.11544.

Jiashu Yang, Chi Zhang, Abudukelimu Wuerkaixi, Xuxin Cheng, Cao Liu, Ke Zeng, Xu Jia, and Xunliang Cai. AutoThinkRAG: Complexity-aware control of retrieval-augmented reasoning for image-text interaction. *arXiv preprint arXiv:2603.05551*, 2026. URL https://arxiv.org/abs/2603.05551.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Tom Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems*, volume 36, pages 11809–11822, 2023a. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/271db9922b8d1f4dd7aaef84ed5ac703-Abstract.html.

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations*, 2023b. URL https://arxiv.org/abs/2210.03629.

Zhixing You, Jiachen Yuan, and Jason Cai. D-mem: A dual-process memory system for LLM agents. *arXiv preprint arXiv:2603.18631*, 2026. URL https://arxiv.org/abs/2603.18631.

Qinggang Zhang, Shengyuan Chen, Yuanchen Bei, Zheng Yuan, Huachi Zhou, Zijin Hong, Hao Chen, Yilin Xiao, Chuang Zhou, Junnan Dong, Yi Chang, and Xiao Huang. A survey of graph retrieval-augmented generation for customized large language models. *arXiv preprint arXiv:2501.13958*, 2025a. URL https://arxiv.org/abs/2501.13958.

Shao Zhang, Xihuai Wang, Wenhao Zhang, Chaoran Li, Junru Song, Tingyu Li, Lin Qiu, Xuezhi Cao, Xunliang Cai, Wen Yao, Weinan Zhang, Xinbing Wang, and Ying Wen. Leveraging dual process theory in language agent framework for real-time simultaneous human-AI collaboration. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4081–4108, 2025b. URL https://aclanthology.org/2025.acl-long.206/.

Xinjie Zhao, Fan Gao, Xingyu Song, Yingjian Chen, Rui Yang, Yanran Fu, Yuyang Wang, Yusuke Iwasawa, Yutaka Matsuo, and Irene Li. ReAgent: Reversible multi-agent reasoning for knowledge-enhanced multi-hop QA. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 4067–4089, 2025. URL https://aclanthology.org/2025.emnlp-main.202.pdf.

Lifan Zheng, Jiawei Chen, Qinghong Yin, Jingyuan Zhang, Xinyi Zeng, and Yu Tian. Rethinking the reliability of multi-agent systems: A perspective from byzantine fault tolerance. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 35012–35020, 2026. URL https://ojs.aaai.org/index.php/AAAI/article/view/40806. Extended version: arXiv:2511.10400.

Denny Zhou, Nathanael Schärli, Le Hou, Jason Wei, Nathan Scales, Xuezhi Wang, Dale Schuurmans, Claire Cui, Olivier Bousquet, Quoc Le, and Ed Chi. Least-to-most prompting enables complex reasoning in large language models. *arXiv preprint arXiv:2205.10625*, 2022. URL https://arxiv.org/abs/2205.10625.

Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. WebArena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023. URL https://arxiv.org/abs/2307.13854.

Xiangrong Zhu, Yuexiang Xie, Yi Liu, Yaliang Li, and Wei Hu. Knowledge graph-guided retrieval augmented generation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 8912–8924, 2025. URL https://aclanthology.org/2025.naacl-long.449/.

A Additional Experimental Details and Results

This section provides the additional implementation details needed to reproduce our main experimental results, including benchmark construction, test case sampling, hyperparameters, prompts, the compute environment, code and data release, and comprehensive results.

A.1 Source Data, Schema Construction, and Test Case Sampling

Source data. We use **PrimeKG** Chandak et al. [2023] as the source of clinical knowledge. From the full graph we retain four node types corresponding to our hierarchical schema: disease (L_3), phenotype (L_1), anatomy (L_2), and drug (L_4). All other node types and their incident edges are removed.

Phenotype–anatomy linkage. PrimeKG does not directly link phenotypes to the body parts where they appear. The closest available connection goes through genes: a phenotype is linked to genes, and those genes are linked to the organs where they are expressed. However, this only tells us where the relevant genes happen to be active at the molecular level, not where the phenotype is actually felt or observed in the clinic. For example, a phenotype like headache may be associated with genes expressed throughout the body, which does not mean the phenotype manifests in all of those locations. For this reason, we do not use this indirect path. Instead, we build the phenotype–anatomy edges directly using a prompted LLM as a medical knowledge source. For each phenotype $s \in L_1$, we ask the LLM to list the body parts where s typically appears in clinical practice, restricting its answers to the anatomy vocabulary already present in PrimeKG. These (phenotype, anatomy) pairs are then added to G_{db} as $L_1 \to L_2$ edges, completing the four-layer schema.

Test case sampling. We sample disjoint test sets for the simple (b=4) and complex (b=2) regimes. For each regime, we draw diseases uniformly at random from the disease pool. For each sampled disease, symptoms are revealed over three rounds according to the disease's associated symptom set. If a disease has fewer than $3b$ associated symptoms, symptoms are re-sampled from its symptom set in subsequent rounds to complete the required observation budget, while preserving diversity within each round whenever possible.

A.2 Hyperparameters

Table 4 summarizes the main SGUR hyperparameters used across all experiments. Below we provide the additional implementation-level settings required for full reproduction.

Table 4: Hyperparameter settings used in all SGUR experiments unless stated otherwise.

Symbol	Description	Value
T	Turn budget	3
b	Per-turn question batch size (complex regime)	2
b	Per-turn question batch size (simple regime)	4
k	Top-k return list size	3
κ	Top-κ candidate pool for evidence acquisition	5
τ_{sim}	Entity-matching cosine-similarity threshold	0.85
α	PFGW feature–structure trade-off (Eq. 7)	0.5
c_{mis}	PFGW feature-mismatch penalty	1
β	Belief-distribution temperature (Eq. 8)	1.0
N_{L_0}	Number of L_0 routing clusters	256

PFGW solver. The PFGW in Eq. (7) is solved with the Frank–Wolfe algorithm from the POT library, using the squared-loss coupling $L(x, y) = |x - y|^2$ and a maximum of 2000 outer iterations. The structural cost matrices C_{d_j} and $C_q^{(t)}$ are computed as directed shortest-path distances under the specificity-aware edge weights of Eq. (3), then jointly rescaled to $[0, 1]$. The partial transport mass m is set to the inherited overlap mass between query and candidate marginals, allowing unmatched evidence and unmatched expert structure to remain uncoupled. The trade-off α and feature-mismatch penalty c_{mis} follow Table 4.

L_0 **pre-filter.** The auxiliary routing layer L_0 is constructed offline over the L_1 observation embeddings produced by **text-embedding-3-large**, in two stages. We first reduce dimensionality with UMAP (`n_neighbors`$= 30$, controlling the local-vs-global trade-off; `n_components`$= 10$ target dimensions; cosine metric), then cluster the projected vectors with HDBSCAN (`min_cluster_size`$=$ 15, the minimum cluster size; `min_samples`$= 5$, controlling cluster density). At query time, the top 10 clusters ranked by cosine similarity between the query embedding and each cluster's centroid are forwarded to the schema-constrained extractor as the candidate observation pool, which then drives the path-compatibility filter that yields $G^{(t)}_{\text{cand}}$.

Entity matching. Extracted entities are aligned to graph nodes by cosine similarity at threshold τ_{sim}, retaining the top 2 observation (L_1) matches per extracted symptom and the top 1 localization (L_2) match per observation.

Baseline hyperparameters. We use the official open-source implementations and default hyperparameters of each baseline, with two exceptions made for fair comparison: (i) the same embedding model (**text-embedding-3-large**) is used across all methods, and (ii) the same indexing LLM (**GPT-5.2**) is used during corpus construction.

A.3 Prompts

We document the prompts used for each LLM-invoked component below. All prompts are released in full with our code.

SGUR: schema-constrained extractor. A single prompt asks the LLM to read the patient's message and pull out two things: the observations being described (L_1) and the localizations where they are felt (L_2). To keep the model from guessing, we also pass in a shortlist of likely observations produced by the L_0 pre-filter and instruct the LLM to choose only from this list. The output must be returned as JSON with three fields: observations and localizations.

Entity recognition prompt

System message:

You are a medical Named Entity Recognizer (NER).

Task. Extract symptoms mentioned or strongly implied by the user query, and list exactly 3 organs most specifically associated with each symptom.

Matching.

- The user may describe symptoms directly, indirectly, or in lay language.
- You MUST semantically match to the closest symptom names from the Symptom List.
- Use meaning-based matching, not exact keywords.

Hard rules.

1. Choose ONLY from the Symptom List (never invent symptoms).
2. Return minimum 1 and AT MOST 3 symptoms total.
3. For each symptom, list 1 of the most directly affected organs that are implicated by that symptom.

Symptom List (flat list): `{routed_observation_shortlist}`

Organ List: `{organ_vocabulary}`

Output format (JSON only):

```
{"symptom name from list": ["organ1", "organ2", "organ3"], ...}
```

User message: `{query}`

Baselines: multi-round adaptation prompt. As described in Section 4, single-shot baselines are adapted to the multi-round protocol by prompting them at each round to summarize their top-ranked diseases together with the supporting observations and localizations. The exact prompt template used for this adaptation is shown in the box below.

Multi-Round Adaptation Prompt for Single-Shot Baselines

ROLE. You are a clinical reasoning assistant performing differential disease identification based *only* on retrieved medical evidence.

TASK. The user's query describes symptoms experienced by a patient. Using only the retrieved information:

1. Identify which diseases are supported by the patient's symptoms.
2. For each candidate disease, extract the supporting *observations* (matched symptoms) and *localizations* (affected organs or anatomical sites) from the retrieved context.
3. Select the ten most likely diseases.

RULES.

- Return exactly the top 10 diseases.
- Base reasoning solely on the retrieved context.
- Do not introduce diseases, observations, or localizations absent from the retrieved context.
- Do not fabricate symptoms, organs, or clinical facts.
- A disease is more likely if more of its phenotypes match the patient's symptoms and its associated organs are consistent with the described localizations.
- If observation or localization support is weak or ambiguous, still return the best 10 candidates with whatever evidence is available (use an empty list if none).
- Do not explain general medical knowledge or include disclaimers.

OUTPUT. Return the following raw JSON only:

```
{
  "top_10_diseases": [
    {
      "disease_name": "<name>",
      "supporting_observations": ["<symptom>", "<symptom>"],
      "supporting_localizations": ["<organ>", "<organ>"]
    },
    ...
  ]
}
```

A.4 Compute Environment

All experiments are run on a single machine with an 11th Gen Intel Core i7-11370H processor (4 cores, 8 threads) and 16 GB of RAM. LLM inference is performed through the TensorBlock API gateway, which provides unified access to the evaluated backbones (GPT-5.2 and GPT-5.4-nano from OpenAI). Graph operations and PFGW solving are implemented in Python using `networkx` and `POT`. The end-to-end latency reported in Table 6 includes API round-trip time and is measured under the same network conditions for all methods to avoid bias.

A.5 Code and Data Release

To support reproducibility, we will release the source code, experimental scripts, processed data splits, and evaluation protocols upon acceptance of the paper. The released materials will include the implementation of SGUR, baseline configuration files, and instructions for reproducing the main results reported in the paper.

A.6 Full Accuracy, Convergence, and Efficiency Results

Table 5 expands the main results by reporting the full retrieval metrics and round-wise Hit@3 trajectories in both simple and complex regimes. This table complements the compact summary in the main text and shows that SGUR's advantage is not limited to the final round: in the complex two-symptom regime, SGUR obtains the best Hit@3 at every interaction round and reaches the correct diagnosis in fewer turns, as reflected by the lowest MTTC. In the simple regime, SGUR is slightly below HippoRAG2 at the first round but improves more rapidly as additional evidence is acquired, achieving the best R2 and R3 Hit@3. These trends support the claim that structure-

Table 5: Full retrieval performance and multi-round convergence. Final metrics are measured after round 3 in the complex two-symptom regime, while the simple and complex columns report Hit@3 across interaction rounds. All Hit@K values are percentages. SGUR achieves the best final ranking quality and the lowest MTTC, and its advantage becomes more pronounced as evidence accumulates.

Method	Final complex-regime performance (R3)					Simple Hit@3			Complex Hit@3		
	MTTC ↓	MRR@R3 ↑	Hit@1 ↑	Hit@2 ↑	Hit@3 ↑	R1	R2	R3	R1	R2	R3
GraphRAG	1.74	0.71	61.20	77.60	83.10	64.20	81.30	84.50	59.60	76.90	83.10
NodeRAG	1.90	0.71	60.20	75.20	80.10	63.00	79.00	86.00	54.20	75.00	80.10
HippoRAG2	1.83	0.68	59.50	74.50	81.70	**68.30**	82.90	87.70	60.40	77.40	81.70
ArchRAG	2.01	0.65	56.70	70.40	76.10	58.20	74.10	80.90	56.70	70.40	76.10
SGUR (Ours)	**1.62**	**0.77**	**64.90**	**83.40**	**92.50**	67.10	**86.70**	**93.40**	**61.40**	**84.70**	**92.50**

Table 6: Full per-round inference efficiency breakdown. We report input tokens, output tokens, total tokens, and end-to-end latency across three interaction rounds. SGUR keeps total token usage nearly flat from R1 to R3 and achieves the lowest latency in every round, confirming that structured query-state updates avoid the growing cost of re-injecting accumulated dialogue history.

Method	Input Tokens ↓			Output Tokens ↓			Total Tokens ↓			Latency (s) ↓		
	R1	R2	R3	R1	R2	R3	R1	R2	R3	R1	R2	R3
GraphRAG	9886.00	10076.00	10080.00	376.00	452.00	491.00	10263.00	10529.00	10570.00	34.10	35.00	35.50
NodeRAG	7609.00	8118.40	8394.70	337.40	427.90	496.00	7946.40	8546.30	8890.70	14.10	15.10	16.40
HippoRAG2	2014.20	2058.90	2116.10	249.40	293.50	326.20	2263.70	2352.40	2442.40	10.10	11.00	13.00
ArchRAG	2223.40	2283.10	2384.40	326.10	402.30	431.00	2549.50	2685.40	2815.40	9.50	10.40	11.20
SGUR (Ours)	**1755.50**	**1759.70**	**1765.50**	**78.60**	**79.40**	**85.40**	**1838.00**	**1843.20**	**1854.90**	**8.10**	**9.10**	**9.70**

grounded alignment and active acquisition convert newly revealed evidence into faster candidate-space contraction.

Table 6 provides the full per-round efficiency breakdown used to compute the average token and latency results reported in the main text. SGUR keeps input tokens nearly constant from R1 to R3 and produces substantially fewer output tokens than all baselines, indicating that the online loop relies on structured query-state updates rather than repeated free-form generation. In contrast, graph-RAG baselines accumulate larger input contexts as previous dialogue is re-injected. This explains why SGUR is both more accurate and more efficient in the main results.

Table 7: Full model-scaling stress test in the complex regime. Each method is evaluated under a Large (L) and Small (S) backbone, with retention reported in the third row of each method. Hit@k and retention values are percentages. Retention is computed as Small/Large for Hit@k and MRR; for MTTC, where lower is better, retention is computed as $\min(\mathrm{L},\mathrm{S})/\max(\mathrm{L},\mathrm{S})$. All metrics are computed after round 3.

Method	Backbone	Hit@1	Hit@2	Hit@3	MRR	MTTC
GraphRAG	GPT-5.2 (L)	61.20	77.60	83.10	0.71	1.74
	GPT-5.4-nano (S)	52.10	64.40	74.70	0.62	2.02
	Retention (%)	85.10	83.00	90.30	87.10	86.10
NodeRAG	GPT-5.2 (L)	60.20	75.20	80.10	0.70	1.90
	GPT-5.4-nano (S)	48.50	65.40	73.60	0.60	2.12
	Retention (%)	80.60	87.00	91.90	84.50	89.60
HippoRAG2	GPT-5.2 (L)	59.50	74.50	81.70	0.68	1.83
	GPT-5.4-nano (S)	54.30	71.60	79.80	0.60	2.11
	Retention (%)	91.40	96.10	97.70	87.30	86.70
ArchRAG	GPT-5.2 (L)	56.70	70.40	76.10	0.65	2.01
	GPT-5.4-nano (S)	49.80	63.80	68.70	0.57	2.30
	Retention (%)	87.80	90.60	90.20	87.40	87.30
SGUR	GPT-5.2 (L)	**64.90**	**83.40**	**92.50**	**0.77**	**1.62**
	GPT-5.4-nano (S)	**60.40**	**80.50**	**92.40**	**0.73**	**1.63**
	Retention (%)	**93.10**	**96.50**	**99.80**	**94.60**	**99.40**