Beyond Accuracy: Disentangling Decomposition from Typed-State Quality Assurance in Multi-Agent LLM Pipelines

Maria Gabriela Alves Zuppardo
Centro de Informática, Universidade Federal de Pernambuco (UFPE)
Recife, PE, Brazil
mgaz@cin.ufpe.br

Vinicius Cardoso Garcia
Centro de Informática, Universidade Federal de Pernambuco (UFPE)
Recife, PE, Brazil
vcg@cin.ufpe.br

arXiv:2610.31280v1 [cs.SE] 22 Oct 2026

ABSTRACT

Multi-agent LLM pipelines are often credited with gains in software engineering tasks, but it is rarely clear whether a gain stems from the typed contract between agents or from the specialized prompts that decomposition enables. We study this attribution problem in Requirements Engineering, comparing three architectures for classifying requirements as functional or non-functional and then prioritizing them with MoSCoW: a single-call baseline; a two-call variant that reuses the pipeline's specialized prompts but passes an unvalidated intermediate; and a pipeline with a Pydantic-validated typed state. The 18 conditions cover PROMISE NFR+ (625 requirements), three locally hosted small language models, and two languages. The F1 gain over the baseline is associated with the two-call, specialized-prompt variant (up to +0.075), whereas the typed state adds a small difference (−0.003 to +0.013) that is significant in one of six comparisons. Beyond accuracy, the typed pipeline contains out-of-range model outputs at the agent boundary, while the unvalidated variant fails later and drops the requirement, and it shows a small, reproducible increase in inter-model MoSCoW consistency whose cause our design cannot isolate. The study is associative: the contrasts bundle prompts, call count, and staged context, and the results come from three 7–8B models on one dataset. The implementation and replication package are archived on Zenodo (https://doi.org/10.5281/zenodo.23380444).

KEYWORDS

Software Quality Assurance, LLM Agents, Multi-Agent Systems, Requirements Engineering, Schema Validation, Empirical Study, MoSCoW Prioritisation, PROMISE NFR+

1 Introduction

Large Language Models (LLMs) are increasingly deployed as components of larger software pipelines, where their outputs feed downstream decisions with limited human review [10]. This raises a quality assurance question that has received little controlled empirical attention in the Brazilian Software Quality community: when a pipeline is decomposed into specialised agents connected by a typed contract, what does that contract buy beyond the accuracy of the underlying model calls? Recent SBQS work shares the evaluate-beyond-accuracy concern in adjacent settings, such as judging LLM-generated unit tests against mutation-testing criteria [1] and assessing generative AI as a substitute for usability inspectors [2]. Practitioners likewise cite reliability and weak contextual grounding as primary barriers to adopting LLMs in Requirements Engineering [14].

None of the studies we surveyed isolates the contribution of an architectural mechanism, such as a typed inter-agent contract, from that of prompt design. When a pipeline outperforms a single-agent baseline, the gain is often read as evidence that the contract improves output quality; whether prompting alone would produce the same gain is rarely checked, and systematic reviews confirm that this conflation remains unresolved [6, 18] (§3). F1 is not this paper's object of interest but the instrument we use to rule out one explanation before asking a more specific question: we show that the typed contract barely moves F1 ($\Delta_s \approx 0$), which points to decomposition and its task-specific prompts as the accuracy driver and leaves open what differences accompany the typed contract, a question F1 cannot answer (§6.5, §6.6).

We therefore ask not only whether a multi-agent pipeline refines classification accuracy, but *which* architectural element drives any gain and what differences accompany the typed contract. We operationalise *quality assurance* narrowly, as early structural validation and cross-model behavioural consistency, loosely related to ISO/IEC 25010's *Reliability* (we do not measure *Maintainability*) rather than the *Functional Suitability* that F1 already captures. We compare three architectures on FR/NFR classification followed by MoSCoW prioritisation, using PROMISE NFR+ [15, 16] (625 requirements): a **baseline** that issues a single LLM call per requirement; a **pipeline** of two specialised agents orchestrated with LangGraph [9] over a Pydantic-validated typed state; and a **two-call baseline** that reuses the pipeline's task-specific prompts but passes an unvalidated intermediate object. Because the second call also receives the first call's decided type and category, the baseline-to-two-call difference is a *decomposition effect* that bundles prompt wording, call count, and staged context (§4.5, §6.5). All three run on the same three locally served SLMs and decoding settings, yielding an 18-condition grid (3 models × 3 architectures × 2 languages).

Here, "disentangling" refers to comparing architectural contrasts; it does not imply causal isolation of the typed contract's effect.

Contributions. (1) A controlled empirical study (18 conditions, $n = 625$; paired McNemar and Wilcoxon tests, Cohen's h and paired g, Bonferroni; Fleiss' κ with a paired bootstrap; and a sensitivity analysis retaining parse-failure fallbacks) that separates the accuracy effect of decomposition from the quality-assurance properties that accompany a typed contract, corroborated by a direct baseline-vs-two-call comparison (§6.5). (2) Evidence that a near-zero F1 difference does not make the typed contract dispensable, restricted to two properties: containment of out-of-range fields at the agent boundary (classification faults one call earlier than the unvalidated ablation), verified by fault injection (§6.6), and a small,

statistically confirmed increase in inter-model MoSCoW consistency (§6.3). Confidence-aware routing, the other mechanism the design attributes to the typed state, fired for only 3 of 3,750 pipeline predictions (all parse fallbacks), so it is validated only through the synthetic fault-injection test, not on the PROMISE NFR+ grid. (3) An open-source pipeline built on a reusable `BaseAgent[Item, Output]` abstraction. (4) A replication package (prompts, runner, manifests, raw predictions for all 36 executions, analysis scripts, tests) archived on Zenodo (DOI `10.5281/zenodo.23380444`).

2 Background

LLM-based agents. LLMs have evolved from passive text generators into agents that invoke tools, maintain state, and work with other agents [10]. Multi-agent systems compose such agents into pipelines in which each handles a sub-task and communicates through typed messages or shared state, reflecting classical design principles (task specialisation, typed contracts, composability). Their empirical benefit remains open, since shared context may offset the gains of decomposition, and it splits into two facets that this paper treats separately: *functional accuracy*, and *quality attributes* (e.g., ISO/IEC 25010 *Reliability*) that an accuracy metric does not capture.

RE and NFR classification. PROMISE NFR+ [15, 16] provides 625 requirements across one functional class and ten NFR categories, the canonical dataset for controlled classification studies in the Brazilian community [11]. MoSCoW prioritisation [3] (Must, Should, Could, Won't Have) lacks a public ground truth, which motivates treating MoSCoW consistency as a quality proxy rather than an accuracy metric (§6.3).

Small language models. SLMs ($\leq$ 8B parameters) are competitive on SE classification tasks [13, 19], and local deployment enables controlled, versioned experiments. Structured-output compliance, however, varies across SLMs [7], a failure mode that matters inside pipelines, where a malformed output at one agent's boundary can spread into the next agent's context (§6.6).

3 Related Work

Recent work has investigated the use of Large Language Models (LLMs) in Requirements Engineering (RE) for tasks such as the classification of functional and non-functional requirements (FR/NFR), prioritization, traceability, and requirements generation. However, most studies evaluate the models as standalone predictors and emphasize classification accuracy, paying less attention to the architectural mechanisms that coordinate multiple interactions between LLMs. FR/NFR classification has been extensively investigated using the PROMISE NFR+ benchmark [15, 16]. Recent studies compare proprietary and open-source models under various prompting strategies [11], combine fine-tuned Small Language Models (SLMs) with explainability techniques [13], and analyze performance differences between compact and proprietary models [19]. In general, these works aim to improve predictive performance while keeping the underlying processing architecture unchanged. From a broader perspective, systematic reviews by Zadenoori et al. [18] and Hemmat et al. [6] identify growing interest in the use of LLMs in RE but also point out a lack of rigorous, reproducible evaluations of architectural design decisions. Concurrently, LLM-based multi-agent systems are gaining traction in software engineering tasks, driven by the decomposition of work among specialized agents and explicit orchestration mechanisms [10]. Although existing studies often report end-to-end performance gains, few isolate the effects of decomposition and the architectural decisions governing agent interaction. This work addresses this attribution problem through a controlled experimental design. We compare three configurations: a single-call baseline, a multi-agent pipeline with typed contracts between agents, and a two-call baseline without typed state. By keeping the underlying models and decoding configurations constant, the design makes it possible to distinguish the effects of task-specific decomposition and prompting from those associated with the typed contract between agents. The gap we claim is relative to the two systematic reviews above and to SBQS/SBES proceedings of 2024–2025; our survey was not systematic.

4 Implementation and Experimental Platform

All three architectures address the same compound RE task, FR/NFR classification followed by MoSCoW prioritisation, and share the dataset adapter, model factory, decoding parameters, retry policy, concurrency level, and metric computation. As Table 1 shows, they differ in more than the intermediate representation; we state each difference so that effects are not attributed to a single factor (§4.5).

4.1 Architectural Overview

The **baseline** classifies and prioritises in a single LLM call (`BaselineAgent`, joint prompt, no intermediate state). The **two-call baseline** (`TwoCallBaselineAgent`) uses the same task-specific prompt builders as the pipeline but passes the intermediate result in an unvalidated Python dataclass, so Pydantic validation happens only when the final output is assembled. The **pipeline** divides the work between two agents that exchange a validated Pydantic object, orchestrated by LangGraph. In both two-call arms the prioritisation prompt receives the type and category decided by the first call.

Comparing the baseline with the two-call baseline bundles decomposition, call count, task-specific prompts, and staged context. Comparing the two-call baseline with the pipeline bundles the typed contract, the validation boundary, the handling of confidence, and LangGraph orchestration. Neither comparison isolates a single component.

4.2 Design and Pipeline

The code follows a layered structure with a one-way dependency rule (presentation, application, domain, infrastructure; the domain layer imports from none of the others). Four patterns carry the experimental control: an *Abstract Factory* (`build_llm`) selects the provider by model-name prefix; an *Adapter* (`PromiseAdapter`) isolates the dataset schema from the prompts; a *Template Method* (`BaseAgent[Item, Output]`) fixes the agent contract; and a *Strategy* (`OrchestrationStrategy`) lets one runner execute every architecture from the same configuration. Prompts are versioned files whose version is recorded in each manifest.

The pipeline consists of `ClassificationAgent` and `PrioritizationAgent`. Agents process requirements in a pool of three workers, retry each call up to three times with

Table 1: Differences between the three architectures. Shared by all: model, temperature, workers, retry policy, dataset adapter, response parsing, and metrics. EN conditions use the original English text and an English prompt; PT conditions use the translated text and a Portuguese prompt.

Property	Baseline	Two-call	Pipeline
LLM calls per requirement	1	2	2
Prompts	joint	task-specific	task-specific (same as two-call)
Intermediate result	none	plain dataclass	validated Pydantic object
First Pydantic validation	after the single call	after the second call	after the first call (agent boundary)
Confidence given to the prioritiser	—	fixed at 1.0	passed; below 0.70 the prompt asks for conservative labels
Orchestration	direct agent call	direct agent call	LangGraph `StateGraph`

exponential back-off, and preserve input order; the global rank is recomputed by sorting on `priority_score`. The prioritisation prompt receives the classifier's confidence and, below 0.70, instructs the model to be conservative; the threshold is a design choice, not a calibrated parameter. The workflow is a LangGraph `StateGraph` with the fixed topology `classifier` → `prioritizer` → `END`; only the pipeline uses LangGraph.

The agents share a `PipelineState` (a Pydantic `BaseModel`) whose requirement entity grows by inheritance (`Requirement` → `ClassifiedRequirement` → `PrioritizedRequirement`). Agent outputs have bounded fields: `confidence` and `priority_score` lie in $[0, 1]$, `priority_rank` is at least 1, and the requirement type and MoSCoW label are enumerations. A value outside these bounds raises a `ValidationError` when the typed object is built, inside the agent and immediately after the call; validation is thus a property of these constructors, not of LangGraph's state assignment. The agent's parser catches the error and substitutes a flagged fallback (§4.3).

4.3 Robustness Layer

Small models frequently return malformed or off-schema JSON. Every agent extracts the outermost JSON object with a regular expression (duplicated in each agent rather than centralised). The classification-producing agents additionally normalise the category (mapping Portuguese or English prose to the PROMISE code) and repair a recurrent off-schema pattern in which the model returns an NFR category code (e.g., `PE`) in the requirement-type field, reading it as NF. The repair was needed only by `mistral:7b`, in 16 of 625 PT and 7 of 625 EN outputs, with identical counts in the pipeline classifier and the two-call baseline. If parsing or validation still fails, the agent returns a deterministic fallback instead of raising: `FUNCTIONAL` with confidence zero for classification and `COULD_HAVE` with score 0.5 for prioritisation, so a batch is never aborted. Fallbacks are identifiable (confidence zero and an error string) and are the outputs excluded or retained in the sensitivity analysis (§5.6).

4.4 Execution, Traceability, and Verification

One `ExperimentRunner` executes every architecture from a frozen configuration and writes a manifest (strategy, model, language, seed, temperature, prompt version, dataset path and MD5 hash, LangGraph version, elapsed time, code commit) and a results file with the metrics and every prediction. Per-call JSONL traces were retained only for the two-call supplement runs. The main grid has **12 conditions**: six for the baseline and six for the pipeline (three models, two languages); the two-call ablation adds six, for **18 in total**.

The unit tests cover agent parsing, category normalisation and fallbacks, the domain models, the runner and manifest, the statistical routines, the language rule (a test fails if any agent sends Portuguese text in an English condition), and the confidence rule (the pipeline's prioritiser receives the classifier's confidence; the two-call baseline fixes it at 1.0). A fault-injection campaign injects 13 faults (six value-range violations, e.g. `confidence` 1.4 or `priority_score` −0.1, and seven parse-level faults); the tests assert that the typed pipeline contains each value-range violation at the agent boundary, whereas the unvalidated intermediate fails only when the output model is assembled (§6.6).

4.5 Known Limitations of the Implementation

Five aspects bound what the experiments can attribute. First, confidence-aware routing is wired but almost never fires: only three of 3,750 pipeline predictions (two for `qwen2.5:7b` PT, one for `mistral:7b` EN, all parse fallbacks excluded from the metrics) have confidence below 0.70, so its isolated effect cannot be evaluated; the two-call baseline fixes the confidence at 1.0. Second, response parsing is regex-based and duplicated across agents, so a response with several brace-delimited blocks fails extraction and becomes a fallback. Third, `seed` governs dataset sampling only and is not forwarded to the model server, so determinism rests on temperature 0. Fourth, agents write their results into the shared state object in place. Fifth, we evaluated the design only in RE, although we expect the principles to apply to other multi-stage SE tasks.

5 Methodology

5.1 Research Questions

Two research questions guide this study, complemented by a language-effect analysis (§6.2) that is reported alongside the RQs but is not itself a formal RQ, since it varies text language and prompt language jointly rather than isolating either as a controlled factor.

RQ1. To what extent does the pipeline outperform the single-agent baseline in FR/NFR binary classification? We report Macro F1 and test the paired per-requirement outcomes across the six model–language combinations with the Wilcoxon signed-rank test (confidence-zero fallbacks excluded) and the exact McNemar test

(fallbacks retained), with Bonferroni correction ($\alpha' = 0.05/6 = 0.0083$).

RQ2 (central). To what extent does decomposition with its task-specific prompts, rather than the typed inter-agent state contract, account for the F1 difference between the pipeline and the single-agent baseline, and what differences accompany the typed state contract that F1 does not capture? We examine this through two architectural contrasts, Δ_p = F1(two-call) − F1(baseline) and Δ_s = F1(pipeline) − F1(two-call). The first bundles task-specific prompts, an additional call, and staged context. The second bundles the typed Pydantic state, its validation boundary, confidence-conditioned prompting (which fired for only 3 of 3,750 pipeline predictions, §4.5), and LangGraph orchestration. Neither contrast isolates a single mechanism, and because the unified baseline prompt and the task-specific prompts cannot be made strictly identical (§7), Δ_s is read as an upper bound on the F1 difference associated with the typed state rather than as a point estimate. Validation and error-propagation behaviour are examined separately (§6.5, §6.6).

Inter-model MoSCoW consistency is reported as a software quality observation (Fleiss' κ with a paired bootstrap of $\Delta\kappa$; §6.3) rather than as accuracy, because PROMISE NFR+ provides no priority ground truth.

Complementary analysis: language effect. We compare an English condition (original text and English prompt) with a Portuguese condition (translated text and Portuguese prompt) across all three architectures. Because text and prompt language change together, the results capture their joint effect rather than either in isolation; the residual confound is acknowledged in §7 (§6.2).

5.2 Dataset

We use PROMISE NFR+ [15, 16], the canonical benchmark for requirements classification, comprising 625 requirements: 255 functional (F) and 370 non-functional (NF) across 10 categories (A, FT, LF, MN, O, PE, PO, SC, SE, US).

The original dataset is in English. To support the language-effect analysis, we produced a Portuguese variant by machine-translating each requirement with `deep_translator` (Google Translate, 2024). Both variants share identical ground-truth labels. We checked translation quality via back-translation (PT→EN, BERTScore F1 with `roberta-large`): the mean score was 0.9646 across all 625 pairs, and 624 of 625 scored ≥ 0.90. The single below-threshold item was manually inspected and confirmed semantically equivalent.

5.3 Architectural Variants

The architectural variant is the primary independent variable for RQ1: baseline, pipeline, and two-call baseline (§4) all accept the same 625 requirements and produce outputs of identical structure, so that architecture is the main manipulated factor (the remaining implementation differences are listed in Table 1 and §4.5). Their mechanics are described in §4.1–4.2; in brief, the baseline issues one unified LLM call, the pipeline decomposes the task via a typed `PipelineState`, and the two-call baseline reuses the pipeline's prompts without the typed contract.

5.4 Models and Configuration

We evaluate three SLMs served locally via Ollama 0.24.0 (Table 2). Local deployment removes network and API dependencies, although decoding at temperature 0 is not bit-deterministic (§7). Prior work reports that SLMs achieve approximately 2 pp lower F1 than large proprietary models on PROMISE-family datasets [13, 19].

Table 2: Models evaluated. All use Q4_K_M quantization.

Model	Params	Quant.
qwen2.5:7b	7.6B	Q4_K_M
llama3.1:8b	8.0B	Q4_K_M
mistral:7b	7.2B	Q4_K_M

Fixed hyperparameters: temperature = 0.0 (greedy decoding); seed = 42 fixes the dataset sample only and is not forwarded to the model server. All experiments were run on a single machine (Intel Core i5-13450HX, 16 GB RAM, NVIDIA RTX 3050 6 GB). Each agent processes requirements in parallel (`max_workers` = 3) with exponential-backoff retry up to three attempts (30–240,s, `tenacity`).

5.5 Experimental Design

The grid comprises 18 conditions: 3 models × 3 architectures × 2 languages (Table 3). All conditions use the complete dataset ($n = 625$, temperature = 0.0); seed 42 only fixes the dataset-sampling routine, which here selects every requirement. Conditions run sequentially on a single machine to avoid GPU contention, and the runner supports resumption via `–resume`. The replication package (prompts, manifests, raw predictions, and analysis scripts) is archived on Zenodo (https://doi.org/10.5281/zenodo.23380444); the dataset is not redistributed, and its file checksum is recorded in every manifest and in the package notes.

Table 3: Experimental grid (18 conditions).

Factor	Levels
Architecture	baseline, pipeline, two-call baseline
Model	qwen2.5:7b, llama3.1:8b, mistral:7b
Language	EN (original), PT (machine-translated)

$n = 625$, temperature = 0.0, seed = 42

5.6 Evaluation Metrics and Statistical Analysis

Classification (RQ1, language effect). The primary outcome is binary FR/NFR classification against PROMISE NFR+ labels, reported as Macro F1 (primary; it weights the F and NF classes equally despite the 255 vs. 370 imbalance), Accuracy, and the Matthews Correlation Coefficient (MCC). NFR subcategory F1 is diagnostic and computed over the ground-truth NF requirements also classified correctly as NF; subcategory precision is computed against every emission of that label in the condition. Categories with $n < 30$ (PO $n = 1$, FT $n = 10$, MN $n = 17$, A $n = 21$, SC $n = 21$) are excluded from the primary conclusions (§6.4).

MoSCoW consistency (quality observation). As no priority ground truth exists, we measure inter-model agreement with Fleiss' κ [5], treating the three models as raters, within each of the 6 (architecture × language) strata ($n_{\text{aligned}} = 623$ in all; the 2 dropped texts are the same alignment exclusion used for RQ1), interpreted with the Landis–Koch scale [8]. To test whether the pipeline's increment over the two-call baseline is distinguishable from zero, we run a paired percentile bootstrap ($B = 10{,}000$, same resampled indices for both architectures) on $\Delta\kappa = \kappa_{\text{pipeline}} - \kappa_{\text{two-call}}$, with Bonferroni correction over $N = 2$ language strata ($\alpha' = 0.025$). Because the bootstrap resamples requirements but not executions, we also re-executed the twelve pipeline and two-call conditions once and compare $\Delta\kappa$ across the two executions and against the κ variation between repeated runs (§6.3).

RQ1 paired comparison. For each of the 6 (model × language) pairs, the exact McNemar test [12] on the discordant pairs compares per-requirement outcomes of pipeline and baseline over $n = 623$ aligned requirements (joined by requirement text, since `Requirement.id` is regenerated on every load; two texts occur twice verbatim in PROMISE NFR+ and collapse to one key each, so $n = 623$; under exclusion a pair also needs a non-fallback prediction on both sides, giving effective n of 621, 623, 623, 590, 623, and 622 for the six conditions). Because the outcomes are binary, the Wilcoxon signed-rank test [17] on the same paired differences reduces to the same discordant-pair counts; it is the primary test on the fallback-excluded data, and we report both. Both test per-requirement correctness, not Macro-F1, so we also test the Macro-F1 differences directly with a paired bootstrap resampling requirements within each true class ($B = 10{,}000$; Bonferroni-adjusted intervals, $\alpha' = 0.05/6$, over the 6 comparisons per contrast). The effect size is Cohen's h [4] (thresholds 0.20/0.50/0.80). Because h uses only marginal accuracies, we also report the paired $g = b/(b+c) - 0.5$ on the $b + c$ discordant pairs with an exact (Clopper–Pearson) 95% interval and thresholds 0.05/0.15/0.25; a magnitude is estimable only if its interval lies within one band. Bonferroni correction over 6 comparisons gives $\alpha' = 0.05/6 = 0.0083$.

Language effect. For each of the 9 (architecture × model) pairs, PT and EN outcomes are aligned by requirement text and compared with the exact McNemar test ($n = 623$; $\alpha' = 0.05/9 = 0.0056$). An earlier unpaired analysis (Mann–Whitney U) is inappropriate, since both languages classify the same requirements, and is reported only for comparison. Uncorrected p-values are reported throughout.

Parse-failure sensitivity. The primary figures exclude the 38 confidence-zero fallbacks (§6.6). As a sensitivity analysis we repeat the RQ1, Δ_p and Δ_s, and language-effect analyses with those fallbacks retained as incorrect predictions and report both readings; the excluded readings are in `statistical_results.json`, the retained ones in `statistical_results_retained.json`.

Ablation contrasts. Δ_s (pipeline vs. two-call) uses the same protocol as RQ1 (Wilcoxon, Cohen's h, $n = 623$, $\alpha' = 0.0083$) to test whether it is distinguishable from zero. As a direct, non-subtractive check we also compare the baseline against the two-call baseline, where the typed contract is absent on both sides, so a significant difference reflects decomposition and its accompanying prompts rather than the typed state (§6.5).

6 Results

We report RQ1 (§6.1), the language-effect analysis (§6.2), MoSCoW consistency (§6.3), NFR subcategory diagnostics (§6.4), RQ2 (§6.5), and validation timing (§6.6).

In short, most of the pipeline's F1 gain is associated with the two-call variant with specialised prompts rather than with the typed state, which is accompanied by earlier containment of out-of-range values and a small, unisolated consistency difference.

Uniform exclusion policy. All accuracy, F1, and MCC figures are computed after excluding confidence-zero parse-failure fallbacks (38 outputs across 4 of the 18 conditions), which are infrastructure artefacts of the shared parsing layer rather than model predictions, applied identically to all conditions (§6.6 details which are affected). Because exclusion is an analyst's choice, we also report the reading that *retains* them as incorrect predictions: McNemar tests and Cohen's g are computed on all 623 aligned pairs, and the Wilcoxon sensitivity is stated wherever it changes a conclusion. Exclusion is the primary analysis because it isolates model behaviour; the retained reading is the end-to-end sensitivity analysis, in which a parsing failure is an unsuccessful system outcome, and it is reported in full beside it. The conclusions of 5 of the 6 RQ1 conditions and of 8 of the 9 language-effect cells are unaffected; the exceptions are discussed in §6.1–§6.5.

6.1 RQ1: Pipeline vs. Baseline Classification

The pipeline achieves statistically significant improvements in 3 of 6 conditions after Bonferroni correction under the primary (exclusion) analysis, and in 4 of 6 under the paired McNemar test that retains the fallbacks (Table 4). These tests compare per-requirement correctness; a paired, class-stratified bootstrap of the macro-F1 difference (Table 5) gives the same verdict under exclusion: the Bonferroni-adjusted interval excludes zero for `qwen2.5:7b` PT, `llama3.1:8b` PT, and `mistral:7b` EN, and for Δ_p in the first two. For Δ_s it excludes zero only for `mistral:7b` EN. The fourth is `llama3.1:8b` EN, non-significant under exclusion ($p = 0.680$) but significant under McNemar ($p = 0.0073$) and under Wilcoxon with fallbacks retained ($p = 0.0055$): the 33 fallbacks removed from the baseline in that cell hid the gap. Cohen's h places every effect in the negligible range ($|h| < 0.20$), but h ignores the pairing. The paired g is large for `qwen2.5:7b` PT and `llama3.1:8b` PT and medium for `llama3.1:8b` EN and `mistral:7b` EN, and its 95% interval stays within one magnitude band only for the first two and for `mistral:7b` EN; for the other three, the 50–97 discordant pairs do not support a confident magnitude. We therefore do not claim that all effects are negligible: they are modest by h and non-negligible by g where g is estimable. The largest gain is `qwen2.5:7b` PT ($\Delta = +7.9$ pp, $g = +0.253$, F1 0.812 vs. 0.733).

`mistral:7b` PT is the only case where the pipeline leans toward underperformance ($\Delta = -2.6$ pp, $p = 0.071$), consistent with a capacity-floor hypothesis. The zero-confidence count does not support it (`mistral:7b` has at most 2/625 such fallbacks in any condition), so we use a finer compliance proxy: the share of predicted-NF outputs without a valid `nfr_category` is 7.3% (pipeline) and 8.1% (two-call) for `mistral:7b` PT, against 1.3–1.4% in English and 0.3–2.6% in the other cells with a non-zero rate, i.e. about 5–6× its

Table 4: RQ1, pipeline vs. baseline (primary metric: macro F1). Acc, F1 and MCC exclude confidence-zero parse-failure fallbacks (§6.6). p_W: Wilcoxon signed-rank on the same data. p_M: exact McNemar on 623 pairs retaining the fallbacks (Bonferroni $\alpha' = 0.0083$). ✓: significant in per-requirement correctness in the primary analysis (fallbacks excluded, $p_W < \alpha'$); it does not mark the macro-F1 bootstrap of Table 5. ‡: significant only in the correctness test that retains fallbacks ($p_M < \alpha'$). g: paired Cohen's g; †: interval spans several magnitude bands, not reliably estimable.

		Baseline			Pipeline								
Model	**L**	Acc	F1	MCC	Acc	F1	MCC	ΔF1	h	p_W	p_M	g	Sig.
qwen2.5:7b	PT	0.733	0.733	0.527	0.814	0.812	0.638	+0.079	+0.193	3.3×10^{-7}	6.4×10^{-7}	+0.253	✓
qwen2.5:7b	EN	0.784	0.784	0.613	0.808	0.807	0.639	+0.023	+0.060	0.059	0.077	+0.119†	–
llama3.1:8b	PT	0.640	0.630	0.446	0.696	0.693	0.511	+0.063	+0.119	7.4×10^{-6}	7.7×10^{-6}	+0.287	✓
llama3.1:8b	EN	0.728	0.724	0.557	0.730	0.728	0.556	+0.004	−0.011	0.680	0.0073	+0.154†	✓‡
mistral:7b	PT	0.602	0.586	0.380	0.579	0.560	0.345	−0.026	−0.046	0.071	0.092	−0.117†	–
mistral:7b	EN	0.606	0.591	0.394	0.638	0.628	0.440	+0.037	+0.066	0.005	0.0066	+0.200	✓

Table 5: Paired, class-stratified bootstrap of macro-F1 differences ($B = 10{,}000$). Shown: point estimate and Bonferroni-adjusted interval ($\alpha' = 0.05/6$, i.e. a 99.17% percentile interval); bold when that interval excludes zero. n: requirements valid (not a fallback) for both sides of that contrast. Point estimates can differ slightly from Tables 4 and 9, which compute each F1 on its own valid set. Unadjusted 95% intervals are in `macro_f1_bootstrap.json`.

		Pipeline − Baseline		Δ_p (2-call − Base)		Δ_s (Pipe − 2-call)	
Model	**L**	Δ [adj. interval]	n	Δ [adj. interval]	n	Δ [adj. interval]	n
qwen2.5:7b	PT	**+0.079 [+0.042, +0.119]**	621	**+0.076 [+0.034, +0.115]**	623	+0.005 [−0.011, +0.020]	621
qwen2.5:7b	EN	+0.023 [−0.009, +0.057]	623	+0.020 [−0.013, +0.054]	623	+0.003 [−0.005, +0.013]	623
llama3.1:8b	PT	**+0.063 [+0.028, +0.100]**	623	**+0.053 [+0.018, +0.087]**	623	+0.011 [−0.003, +0.026]	623
llama3.1:8b	EN	−0.005 [−0.041, +0.028]	590	−0.003 [−0.036, +0.029]	590	−0.002 [−0.013, +0.010]	623
mistral:7b	PT	−0.027 [−0.066, +0.011]	623	−0.024 [−0.062, +0.015]	623	−0.003 [−0.014, +0.006]	623
mistral:7b	EN	**+0.037 [+0.003, +0.073]**	622	+0.024 [−0.008, +0.058]	621	**+0.013 [+0.002, +0.027]**	621

English rate (§6.6). We present the capacity floor as a hypothesis, not a tested mechanism.

Answer to RQ1: On requirements valid for both sides, the paired bootstrap supports a Macro-F1 difference in favour of the pipeline in 3 of 6 conditions (Bonferroni-adjusted). In per-requirement correctness, the Wilcoxon test agrees in the same 3, and with the fallbacks retained the McNemar test finds a difference in 4 of 6. Effect sizes are modest by Cohen's h but medium to large by the paired g where estimable. The difference depends on the model's ability to produce structured output; models below a compliance threshold show no gain from the specialised two-call pipeline.

6.2 Complementary Analysis: PT vs. EN Conditions

Because PT and EN classify the *same* requirements, the paired analysis is the appropriate one: aligning them by requirement text and applying the exact McNemar test (column p_M of Table 6; $n = 623$, fallbacks retained, $\alpha' = 0.0056$) shows a significant PT<EN gap in five of nine cells: the baseline with `qwen2.5:7b` and `llama3.1:8b`, the pipeline with `mistral:7b`, and the two-call baseline with `llama3.1:8b` and `mistral:7b`. English is the better language in every significant cell, and only `qwen2.5:7b` is language-insensitive in both decomposed architectures ($p = 0.91$ and 0.82). An earlier unpaired Mann–Whitney analysis, kept in Table 6 for comparison only, reached significance after correction only for the baseline with `llama3.1:8b` ($h = -0.191$, $p < 0.001$). That baseline `llama3.1:8b` result also depends on the exclusion policy: with the 33 EN fallbacks retained, the Mann-Whitney test gives $p = 0.063$ and no longer survives correction, while the paired test is significant ($p = 0.0009$).

Table 6: Language-effect results. p_{MW}: Mann-Whitney U (unpaired, original analysis; fallbacks excluded). p_M: exact McNemar on 623 PT/EN-aligned pairs (fallbacks retained); bold = significant after Bonferroni over 9 comparisons ($\alpha' = 0.0056$). Accuracies follow the exclusion policy of §6; four of the 18 cells are affected (§6.6).

Arch.	**Model**	PT	EN	h	p_{MW}	p_M
Baseline	qwen2.5:7b	0.734	0.785	−0.120	0.034	3.0×10^{-5}
Baseline	llama3.1:8b	0.640	0.729	−0.191	< 0.001	**0.0009**
Baseline	mistral:7b	0.602	0.607	−0.010	0.862	0.801
Pipeline	qwen2.5:7b	0.815	0.809	+0.015	0.793	0.913
Pipeline	llama3.1:8b	0.697	0.730	−0.075	0.188	0.011
Pipeline	mistral:7b	0.579	0.638	−0.121	0.034	9.1×10^{-6}
Two-call	qwen2.5:7b	0.811	0.806	+0.012	0.829	0.824
Two-call	llama3.1:8b	0.687	0.732	−0.099	0.081	**0.0004**
Two-call	mistral:7b	0.583	0.626	−0.090	0.115	**0.0004**

Language-effect finding: Under the paired analysis, Portuguese is associated with significantly lower accuracy than English in 5 of 9 architecture–model cells, including decomposed architectures for `mistral:7b` and `llama3.1:8b`. We therefore do not claim

that decomposition attenuates the language effect, nor present the results as evidence of multilingual portability. The data support only that `qwen2.5:7b` is language-insensitive under both decomposed architectures and that the baseline `llama3.1:8b` gap is not specific to that architecture. Because requirement and prompt language change together, and the Portuguese text is a machine translation, the gap cannot be attributed to either (§7).

6.3 MoSCoW Consistency: A Software Quality Observation

Because PROMISE NFR+ provides no MoSCoW ground truth, we assess prioritisation as **inter-model behavioural consistency**, operationalised as Fleiss' κ over the three models within each (architecture × language) stratum (Table 7). It is a potential proxy for the *Reliability* characteristic of §2, not a measure of prioritisation quality: models may agree and still be wrong, and we have neither priority ground truth nor practitioner evaluation.

Table 7: Inter-model MoSCoW consistency, Fleiss' κ (3 raters, 6 strata), full four-label priority set. n: requirements aligned across all three raters (same filter as RQ1).

Architecture	L	n	κ	Interp.
Baseline	PT	623	−0.1406	poor
Baseline	EN	623	−0.1220	poor
Two-call	PT	623	+0.1360	slight
Two-call	EN	623	+0.1695	slight
Pipeline	PT	623	+0.1768	slight
Pipeline	EN	623	+0.2388	fair

The negative κ in both baseline strata reflects systematic strategy divergence: in the PT baseline, `llama3.1:8b` and `mistral:7b` collapse to a Must rate of 99% and 97% while `qwen2.5:7b` spreads evenly (46%), a configuration Fleiss' κ scores as random disagreement but which shows *which* models fail to discriminate. The Must rate of those two models falls to 0.90/0.86 (two-call) and 0.87/0.85 (pipeline), while `qwen2.5:7b` stays at 0.41–0.42; the whole progression is driven by the two weaker models reducing their Must-collapse at each step.

The six strata show a monotone progression: baseline ($\kappa \leq -0.12$) → two-call ($\approx +0.15$) → pipeline ($\leq +0.24$). The two-call baseline, with the pipeline's prompts but no typed state, already recovers most of the improvement; the pipeline adds $\Delta\kappa = +0.04$ to +0.07. Unlike Δ_s for F1, this increment is distinguishable from resampling noise: a paired percentile bootstrap ($B = 10{,}000$, Bonferroni over $N = 2$, $\alpha' = 0.025$; §5.6) gives PT $\Delta\kappa = +0.0408$, 97.5% CI $[+0.0133, +0.0703]$, $p = 0.0004$, and EN +0.0694, CI $[+0.0310, +0.1083]$, $p < 0.0001$. This holds only within this three-model panel (§7). The primary driver of the κ gain is thus decomposition with its task-specific prompts, matching the F1 ablation (§6.5).

The residual pipeline-vs-two-call difference is reproducible: a second execution of the twelve pipeline and two-call conditions (same code, seed, and configuration; the model server was updated in between, §7) reproduced the sign of $\Delta\kappa$ (PT +0.022, EN +0.079) and of all four pairings of the two executions (PT +0.022 to +0.041, EN +0.052 to +0.097), each exceeding the κ variation between repeated runs (at most 0.011 in PT and 0.027 in EN). With two executions the noise estimate is coarse, and the PT increment roughly halved (item-bootstrap 95% interval $[+0.003, +0.041]$), so we regard it as a reproducible association, not a precisely estimated effect. We do not attribute it to the typed contract: the validation boundary rejected no real output, confidence routing affected at most three of 3,750 prioritisations (§4.5), and the two architectures differ on 0.6–2.1% of types and up to 13% of priorities, which points to execution-order and run-to-run effects that a single-run design cannot separate from the contract. We did not test whether the increment translates into an operational benefit.

6.4 NFR Subcategory F1 (Diagnostic)

Subcategory recall is computed over the 272 (of 370) NF requirements also classified correctly as NF, and precision against all 625 predictions, so a subcategory label emitted on a requirement misclassified as FR counts against that label's precision (§5.6). In the best condition (`qwen2.5:7b` pipeline EN), among categories with $n \geq 30$, SE achieves the highest F1 (0.916; $P = 0.927$, $R = 0.905$, $n = 42$), followed by PE (0.701) and LF (0.566), while O (0.333) and US (0.429; $P = 1.000$, $R = 0.273$) are weakest. The US imbalance indicates a conservative model that labels Usability only when confident, consistently across architectures, so it reflects the model's representation of the class rather than an effect of decomposition. These patterns do not affect the RQ1 and ablation conclusions (binary macro F1) but mark a boundary for our quality claims: validation timing (§6.6) targets *structural* conformance, not *semantic* correctness such as category ambiguity. Per-category values (all conditions) and the Must-label rates per model are in `subcategory_errorprop.json` and `kappa_table.json`.

6.5 RQ2: Decomposition vs. Typed-State Quality Assurance

Table 8: What each contrast changes and permits us to conclude.

Contrast	What changes	What it supports
Δ_p: baseline → two-call	calls, task prompts, staged context	F1 associated with the variant (not which element)
Δ_s: two-call → pipeline	validation boundary, confidence, LangGraph	bound on the F1 difference; QA properties (not their cause)

Table 9 reports the three-way comparison. Δ_p is a difference associated with the two-call specialised-prompt variant: its second call receives the already-decided `requirement_type` and `nfr_category`, so it conflates prompt wording, an additional call, and staged context (a call-count control is future work). Read with that caveat, it is substantial in 4 of 6 conditions (up to +0.075 F1), so most of the pipeline's accuracy advantage over the single-call baseline is associated with decomposition. Under exclusion it is

negligible for `llama3.1:8b` EN (+0.006), but retaining the baseline's 33 fallbacks raises it to +0.045; it is negative for `mistral:7b` PT, and largest for `qwen2.5:7b` PT, the strongest instruction follower. Δ_s is the residual F1 difference between the typed pipeline and the two-call baseline; it bundles the earlier validation point, execution-order effects, and confidence routing (triggered on only three prioritisations, §4.5). It is small in all conditions ($|\Delta_s| \leq 0.013$), inconsistent in direction, and essentially unchanged when fallbacks are retained (largest shift 0.003).

On per-requirement correctness, a paired Wilcoxon test (Bonferroni $\alpha' = 0.0083$) finds five of six Δ_s comparisons not significant; the sixth (`mistral:7b` EN, $p = 0.0082$) narrowly reaches significance despite a small effect ($h = +0.023$). The Macro-F1 bootstrap of Table 5 agrees: only that condition has an adjusted interval excluding zero. The pipeline and the two-call baseline disagree on only 4–15 requirements per condition, so the interval of g spans several magnitude bands in five of six conditions; only `mistral:7b` EN (7 vs. 0, $g = +0.50$, $[+0.09, +0.50]$) is reliably non-negligible. We therefore report these effects as small in F1 and not distinguishable, at this n, from effects of the size the interval allows.

A small Δ_s does not show that the typed contract has no quality value; it shows that F1 registers no difference, while other instruments do (§6.3 for κ and §6.6 for validation timing). It does not establish that the typed contract causes those differences.

As a direct, non-subtractive check, we compare the single-call baseline with the two-call baseline head-to-head, so that the typed contract is absent on both sides (same protocol as RQ1, $\alpha' = 0.0083$). Decomposition alone is significant in 2 of 6 conditions with fallbacks excluded (`qwen2.5:7b` PT, $h = +0.184$, $g = +0.250$; `llama3.1:8b` PT, $h = +0.099$, $g = +0.254$; both negligible by h and large by g), and in 3 of 6 when the 33 baseline fallbacks are retained (`llama3.1:8b` EN, $h = +0.092$, McNemar $p = 0.0038$), the same flip that RQ1 shows. This reproduces the RQ1 pattern without the subtraction logic of Δ_p and Δ_s, strengthening the case that where the pipeline beats the baseline, the gain is associated with decomposition and its prompts rather than with the typed contract.

Answer to RQ2: Most of the F1 difference between the pipeline and the single-agent baseline is associated with decomposition and its task-specific prompts ($\Delta_p \in [+0.020, +0.075]$ in 4 of 6 conditions), corroborated by the direct baseline-vs-two-call comparison; we cannot separate prompt wording from the extra call and staged context. The pipeline-minus-two-call difference is small and directionally inconsistent ($\Delta_s \in [-0.003, +0.013]$), and the discordant pairs are too few to bound its magnitude. Two narrow pieces of evidence concern the typed contract's quality-assurance properties: earlier structural validation (§6.6, 13 injected faults) and a statistically confirmed, reproducible difference in inter-model MoSCoW consistency (§6.3) whose cause the design does not isolate.

Supplementary check with a larger model. To test whether the pipeline's gain depends on the small models' compliance, we ran the three architectures on `gpt-4.1-mini` (Azure AI Foundry), outside the 18-condition grid, with the same prompts, seed, and temperature; no classification output fell back. The pipeline outperforms the baseline in accuracy (PT +0.019, McNemar $p = 0.043$; EN +0.034, $p < 10^{-4}$, $n = 623$; only EN survives Bonferroni over two languages, $\alpha' = 0.025$). The gain is associated with decomposition: Δ_p is +0.013 (PT, $p = 0.17$) and +0.037 (EN, $p < 10^{-4}$), whereas Δ_s is +0.006 (PT, $p = 0.42$) and −0.003 (EN, $p = 0.63$), reproducing the small-model pattern. Caveats: one model, one execution per condition; the PT baseline and pipeline and both two-call runs were executed on an adjacent development line whose agents differ in category normalisation, JSON extraction, and confidence forwarding (none affects the binary accuracy reported, as no classification fallbacks occurred), and only the EN baseline and pipeline were re-executed with the artifact's agents; and 32 pipeline and 37 two-call English prioritisations (5–6%) fell back to Could Have because of a trailing comma in the model's JSON, which does not enter the accuracy figures but would affect MoSCoW analyses.

6.6 Validation Timing and Schema Compliance

The clearest evidence concerning the typed state's quality-assurance properties is a difference in *when* structural validation occurs, not a metric. In all three architectures, a parse or enum-conversion failure is caught inside the responsible agent's `_parse_response` and replaced with a conservative default (`FUNCTIONAL`, `confidence=0.0`); no malformed value reaches a downstream agent's context. This robustness is not attributable to the typed contract: all four agents use the same regex-based JSON extraction and fallback, duplicated across the codebase.

The typed contract's verifiable contribution is validation *timing*. In the pipeline, `ClassificationOutput`'s Pydantic constraints (e.g., `confidence` in $[0, 1]$) are enforced immediately after the `ClassificationAgent` call, before the result enters `PipelineState`. In the two-call ablation the intermediate is a plain `@dataclass`, and Pydantic validation occurs only when the final output is assembled after the second call. The typed contract thus validates one LLM call earlier, a narrow V&V property: a violation could be acted on before the second, more expensive call.

On the PROMISE NFR+ grid this difference never surfaced as a failure: no confidence or type value fell outside its bounds in any of the 18 runs. Parse-failure fallbacks are rare: 38 outputs across 4 runs (33 baseline `llama3.1:8b` EN, 5.3%; 2 pipeline `qwen2.5:7b` PT; 1 pipeline `mistral:7b` EN; 2 two-call `mistral:7b` EN). We exclude them from Tables 4, 6, and 9 and report the retained reading as a sensitivity analysis (§6). The zero-confidence count is a poor proxy for structured-output compliance; a finer one is the share of predicted-NF outputs with no valid `nfr_category`: for `mistral:7b` it is 7.3% (pipeline) and 8.1% (two-call) in Portuguese against 1.3–1.4% in English, and no other condition exceeds 2.6%, which is what the capacity-floor reading of §6.1 relies on. The three pipeline fallbacks are the only outputs with confidence below 0.70, so only 3 of 3,750 prioritisations received the conservative-label instruction; routing was effectively inert, and no reported result depends on it.

To verify the boundary independently of the grid, we ran a systematic fault-injection campaign through the real agents of each architecture with a scripted LLM (`run_fault_injection_campaign.py`). It injects 13 faults into the model's JSON, in either call: six value-range violations (`confidence` 1.4, −0.5, NaN; `priority_score` 1.5, −0.1; `priority_rank` 0) and seven parse-level faults (invalid or missing enumeration, wrong type, truncated JSON). The typed contract adds nothing for parse-level faults, which the two-call baseline and the pipeline handle identically (a flagged fallback after two calls).

Table 9: Decomposition ablation (macro F1). Δ_p = F1(two-call) − F1(baseline), a difference associated with the two-call variant (prompt wording, extra call, and staged context are confounded); Δ_s = F1(pipeline) − F1(two-call). *Excluded* drops the confidence-zero fallbacks (§6.6); *Retained* keeps them as errors ($n = 625$; Δ computed from unrounded F1). b/c: requirements correct only for the pipeline / only for the two-call baseline. g: paired Cohen's g for Δ_s; †: not reliably estimable (4–15 discordant pairs).

					Excluded		Retained			
Model	**L**	Base	2Call	Pipe	Δ_p	Δ_s	Δ_p	Δ_s	b/c	g
qwen2.5:7b	PT	0.733	0.808	0.812	+0.075	+0.004	+0.075	+0.001	8/7	+0.033†
qwen2.5:7b	EN	0.784	0.804	0.807	+0.020	+0.003	+0.020	+0.003	3/1	+0.250†
llama3.1:8b	PT	0.630	0.682	0.693	+0.053	+0.011	+0.053	+0.011	8/2	+0.300†
llama3.1:8b	EN	0.724	0.730	0.728	+0.006	−0.002	+0.045	−0.002	3/4	−0.071†
mistral:7b	PT	0.586	0.563	0.560	−0.023	−0.003	−0.023	−0.003	1/3	−0.250†
mistral:7b	EN	0.591	0.614	0.628	+0.023	+0.013	+0.024	+0.013	7/0	+0.500

The difference lies in the value-range violations (Table 10). The typed pipeline rejects each where the typed object is built, detecting classification-stage faults after the first call and prioritisation-stage faults after the second, and keeps the requirement as a flagged fallback after two calls. The unvalidated two-call baseline passes the same values through its dataclass and fails when the output model is assembled, after the second call; because that failure propagates through the retry wrapper, both calls are repeated up to three attempts, consuming six calls before the requirement is dropped without a failure record. The single-call baseline contains all 13 faults at its one call, and no architecture let an out-of-range value reach the output. The campaign is synthetic: it characterises where each architecture detects a fault and what that costs, not how often such faults occur; that earlier validation prevents an operational failure remains an untested hypothesis. The six-call cost and the loss in the unvalidated arm partly reflect the shared retry policy, which retries a deterministic validation error, so the contrast is between a contained and an uncontained failure rather than between detection and non-detection.

Table 10: Fault-injection campaign: outcome and LLM calls consumed per fault (13 faults: 6 value-range, 7 parse-level).

Architecture	Value-range (6)	Parse-level (7)
Baseline	flagged, 1 call	flagged, 1 call
Two-call	dropped, 6 calls	flagged, 2 calls
Pipeline	flagged, 2 calls	flagged, 2 calls

Two conclusions follow. Structured-JSON compliance is achieved independently in each architecture, not through one centralised component. And validation *timing*, not failure containment, is what distinguishes the typed pipeline from the unvalidated ablation; the isolated effect of confidence routing could not be evaluated, since it fired on only three fallbacks.

7 Discussion and Threats to Validity

The experiments indicate that the typed contract is accompanied by little difference in classification performance ($\Delta_s \in [-0.003, +0.013]$). The evidence for what it is accompanied by is limited to two properties: earlier structural validation, verified by a 13-fault injection campaign, and a statistically confirmed difference in inter-model MoSCoW consistency whose cause we did not isolate. We do not claim, because we did not test, benefits such as predictable behaviour under model substitution or auditable intermediate decisions, and even these would depend on the model satisfying the structured-output contract. For `mistral:7b` PT (§6.1, §6.6), parse failures fall back to `FUNCTIONAL` rather than NF, silently biasing misclassifications toward the functional class; since the model has no more zero-confidence outputs than the others but 5–6× more NF predictions without a valid category in Portuguese, these failures are subtler than malformed responses. Structured-output compliance is therefore a prerequisite for benefiting from decomposition.

Practically, teams interested mainly in F1 should invest in better prompts rather than a typed contract, whereas applications that value earlier structural validation or cross-model consistency may prefer the pipeline despite its limited accuracy gain over the two-call baseline. Our paired analysis does not support deploying the same architecture bilingually without adaptation (§6.2). Although the pipeline makes two sequential calls, local execution times were comparable to the baseline (`qwen2.5:7b` PT: 3529 s vs. 3415 s for $n = 625$); hosted APIs may approach a 2× wall-clock cost, which we did not test.

Internal threats. PROMISE NFR+ has no MoSCoW ground truth, so prioritisation quality is operationalised as inter-model stability, reflecting reliability rather than correctness; some requirements also admit several plausible NFR categories while the dataset gives one label. The principal threat is the prompt–architecture confound: the baseline combines both subtasks in one prompt, so Δ_p is a difference associated with the two-call variant, bundling prompt wording, an extra call, and staged context, and a call-count-controlled arm would be needed to separate them. Δ_s likewise bundles the validation boundary, confidence handling (which fired for 3 of 3,750 predictions), and LangGraph orchestration, so it is a pipeline-versus-two-call contrast rather than an isolated typed-contract effect, and its range is an upper bound on the F1 difference associated with the typed contract. Cohen's h places all RQ1 effects below 0.20 but ignores the pairing; the paired g is medium or large where estimable, so we do not call the effects negligible. For Δ_s the 4–15 discordant pairs are too few for any reliable paired effect size. Excluding 38 fallbacks is an analyst's decision, hence both readings; it changes one RQ1 verdict (`llama3.1:8b` EN) and one

Δ_p estimate. Decoding at temperature 0 is not bit-deterministic: primary results come from one execution per condition, and the one repeated execution (pipeline and two-call) ran on a newer Ollama release (0.35.1 vs. 0.24.0), so its variation mixes run-to-run and server-version effects; differences below about one point of F1 (and 0.03 in κ) should not be interpreted. The fault-injection evidence rests on 13 synthetic faults (six of them value-range violations, which extend the earlier five-value probe): it shows where validation occurs and what it costs, not how often such faults arise. Per-call traces were retained only for the two-call supplement runs, which limits audit of individual failures.

External threats. All experiments use PROMISE NFR+ and three decoder-only SLMs (7–8B parameters) served through versioned Ollama tags, whose behaviour may change; compliance rates and the contribution of typed contracts may differ for other datasets, domains, or models, and the κ gain reflects only this three-model panel. The single larger model (§6.5) is one model with one execution per condition. The language comparison varies requirement and prompt language jointly, and the Portuguese text is a machine translation (back-translation BERTScore F1 0.9646; 624 of 625 pairs ≥ 0.90), so translation artefacts may contribute to the PT–EN gap. Only the pipeline used LangGraph, although the typed-state pattern is not specific to it. Bonferroni correction was applied throughout, and categories with fewer than 30 instances were excluded from the primary analyses. Our conclusions are therefore intentionally narrow: for the evaluated dataset, task, and model panel, the measurable differences that accompany the typed contract are earlier structural validation and a small difference in inter-model consistency, rather than improvements in classification accuracy.

8 Conclusion and Future Work

In the evaluated setting, the accuracy difference is associated with decomposition into specialised prompts, not with the typed state (§6.1–§6.5). The typed state is accompanied by two differences beyond accuracy: containment of out-of-range values at the agent boundary, with classification faults detected one call earlier in a 13-fault campaign (§6.6), and a small, reproducible increase in inter-model MoSCoW consistency (paired bootstrap, $p < 0.001$ in both languages, §6.3). The design does not isolate their causes, and the limits are detailed in §7.

Future Work. (i) Calibrating the confidence threshold, to isolate routing from validation timing. (ii) A call-count-controlled arm to separate prompt wording from the extra call and staged context. (iii) Larger models and a larger rater panel, to test whether Δ_s grows with capacity and the consistency gain generalises.

ARTIFACT AVAILABILITY

The implementation, prompts, run manifests, raw predictions, analysis scripts, and unit tests are archived on Zenodo (https://doi.org/10.5281/zenodo.23380444). The PROMISE NFR+ file is not redistributed; its MD5 hash is recorded in every manifest.

ACKNOWLEDGMENTS

For partially supporting this work, we would like to thank INES.IA (National Institute of Science and Technology for Software Engineering Based on and for Artificial Intelligence), https://www.ines.org.br, CNPq grant 408817/2024-0.

REFERENCES

[1] Pedro Cabral, João Arruda, Cleidson Souza, and Victor Pinto. 2025. Evaluating LLM-Generated Unit Tests with Mutation Testing: ChatGPT vs DeepSeek. In *Anais do XXIV Simpósio Brasileiro de Qualidade de Software* (São José dos Campos/SP). SBC, Porto Alegre, RS, Brasil, 120–130. doi:10.5753/sbqs.2025.13792

[2] Luis F. G. Campos, Leonardo C. Marques, and Walter T. Nakamura. 2025. Will AI also replace inspectors? Investigating the potential of generative AIs in usability inspection. In *Anais do XXIV Simpósio Brasileiro de Qualidade de Software (SBQS 2025)*. SBC. https://sol.sbc.org.br/index.php/sbqs/article/view/38990

[3] Dai Clegg and Richard Barker. 1994. *Case Method Fast-Track: A RAD Approach*. Addison-Wesley.

[4] Jacob Cohen. 1988. *Statistical Power Analysis for the Behavioral Sciences* (2nd ed.). Lawrence Erlbaum Associates.

[5] Joseph L. Fleiss. 1971. Measuring Nominal Scale Agreement Among Many Raters. *Psychological Bulletin* 76, 5 (1971), 378–382.

[6] Arshia Hemmat, Mohammadreza Sharbaf, Shekoufeh Kolahdouz-Rahimi, Kevin Lano, and Sobhan Y. Tehrani. 2025. Research Directions for Using LLM in Software Requirement Engineering: A Systematic Review. *Frontiers in Computer Science* (2025). doi:10.3389/fcomp.2025.1519437

[7] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Transactions on Information Systems* 43, 2 (2025), 1–55. Preprint: arXiv:2311.05232 (2023). doi:10.1145/3703155

[8] J. Richard Landis and Gary G. Koch. 1977. The Measurement of Observer Agreement for Categorical Data. *Biometrics* 33, 1 (1977), 159–174.

[9] LangChain. 2024. LangGraph: Building Stateful, Multi-Actor Applications with LLMs. https://www.langchain.com/langgraph. Software documentation, accessed 2026.

[10] Junwei Liu, Kaixin Wang, Yixuan Chen, Xin Peng, Zhenpeng Chen, Lingming Zhang, and Yiling Lou. 2026. Large Language Model-Based Agents for Software Engineering: A Survey. (2026). doi:10.1145/3796507

[11] Murilo Martin, Daniel Coutinho, Anderson Uchôa, and Juliana Alves Pereira. 2025. Evaluating the Potential of Large Language Models in Security-Related Software Requirements Classification. In *Anais do XXXIX Simpósio Brasileiro de Engenharia de Software (SBES 2025)*. SBC. https://sol.sbc.org.br/index.php/sbes/article/view/37009

[12] Quinn McNemar. 1947. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika* 12, 2 (1947), 153–157.

[13] Gokul Rejithkumar and Preethu Rose Anish. 2025. NICE: Non-Functional Requirements Identification, Classification, and Explanation Using Small Language Models. In *Proceedings of the IEEE/ACM International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP 2025)*. IEEE/ACM. doi:10.1109/ICSE-SEIP66354.2025.00031

[14] Reine Santos, Igor Steinmacher, Tayana Conte, Ana Carolina Oran, and Bruno Gadelha. 2025. Adoption of LLMs in Requirements Engineering: What Practitioners Are Worried About?. In *Anais do XXIV Simpósio Brasileiro de Qualidade de Software (SBQS 2025)*. SBC. https://sol.sbc.org.br/index.php/sbqs/article/view/39013

[15] Bruno Silva, Rodrigo Nascimento, Luis Rivero, Geraldo Braz Júnior, Rodrigo Pereira dos Santos, Luiz Eduardo Galvão Martins, and Davi Viana. 2024. Promise+: A Dataset of Labeled Software Requirements. doi:10.5281/zenodo.12805484

[16] Bruno Silva, Rodrigo Nascimento, Luis Rivero, Geraldo Braz Júnior, Rodrigo Pereira dos Santos, Luiz Eduardo Galvão Martins, and Davi Viana. 2024. Promise+: Expandindo a Base de Dados de Requisitos de Software Promise_exp. In *Anais do XXXVIII Simpósio Brasileiro de Engenharia de Software (SBES 2024)*. SBC. https://sol.sbc.org.br/index.php/sbes/article/view/30370

[17] Frank Wilcoxon. 1945. Individual Comparisons by Ranking Methods. *Biometrics Bulletin* 1, 6 (1945), 80–83.

[18] Mohammad Amin Zadenoori, Jacek Dąbrowski, Waad Alhoshan, Liping Zhao, and Alessio Ferrari. 2025. Large Language Models (LLMs) for Requirements Engineering (RE): A Systematic Literature Review. arXiv:2509.11446. https://arxiv.org/abs/2509.11446

[19] Mohammad Amin Zadenoori, Vincenzo De Martino, Jacek Dąbrowski, Xavier Franch, and Alessio Ferrari. 2025. Does Model Size Matter? A Comparison of Small and Large Language Models for Requirements Classification. arXiv:2510.21443. https://arxiv.org/abs/2510.21443