arXiv:2610.35598v1 [cs.AI] 26 Oct 2026

Evolve on Graph: Self-Evolving Knowledge Graph Reasoning without Question Annotations

Zongyuan Sun[1,*], Shuai Zhao[1,†], Ke Shu[2,*], Ziyue Zhu[1], Shangyang Wu[1], Yuanlin Lin[1], Tao Wang[1], Shiyue Li[1], Yingjun Du[3], Anh Tuan Luu[4], Haoran Luo[4,†]

[1]Beijing University of Posts and Telecommunications, Beijing, China [2]University of Helsinki, Helsinki, Finland
[3]University of Amsterdam, Amsterdam, Netherlands [4]Nanyang Technological University, Singapore
[*]Equal contribution. [†]Corresponding authors.

Homepage GitHub Model

Abstract

Knowledge graphs (KGs) provide reliable external knowledge for multi-hop factual reasoning with large language models (LLMs), yet most KGs in real-world domains such as finance and industry lack question–answer (QA) annotations. Without such annotations, existing methods struggle to improve LLM graph reasoning: training-free agents cannot improve their exploration through interaction; reinforcement learning (RL) agents depend on QA data annotated by humans or synthesized by closed-source models; and training data is typically built offline, so its topology mix is fixed and cannot follow the agent's ability. We propose **Evolve on Graph (EoG)**, a framework that lets an LLM self-evolve its graph reasoning without any QA annotations. Through Graph Self-Play, a single LLM alternates between a Proposer that poses questions and a Solver that answers them. To make questions verifiable, Symbolically Verifiable Topology Synthesis (SVTS) samples reasoning topologies on the KG and executes the corresponding first-order logic queries, yielding topology skeletons with correct and complete answer sets; guided by a skeleton, the Proposer explores the graph to pose a multi-hop question. To let task topologies follow the Solver, a Boundary-Aware Graph Curriculum rewards the Proposer by the Solver's score variance, keeping questions at the Solver's capability boundary so that complex topologies enter training as the Solver improves. On six benchmarks over Freebase, Wikidata, and an industrial KG, EoG consistently outperforms prompting-, retrieval-, and agent-based methods, showing that LLM graph reasoning can improve through self-play without QA annotations. Our code is available at https://github.com/Szy5/EoG.

CCS Concepts

• **Information systems** → **Question answering**; • **Computing methodologies** → **Semantic networks**.

Keywords

Knowledge Graph, Large Language Models, Agentic Reasoning, Self-Play, Graph Reasoning

1 Introduction

Large language models (LLMs) [46] have shown remarkable abilities across a wide range of natural language processing tasks [51, 53], yet in complex reasoning that requires deep, multi-hop, and structured factual knowledge, they remain limited by insufficient coverage of long-tail knowledge and by factual hallucinations [12].

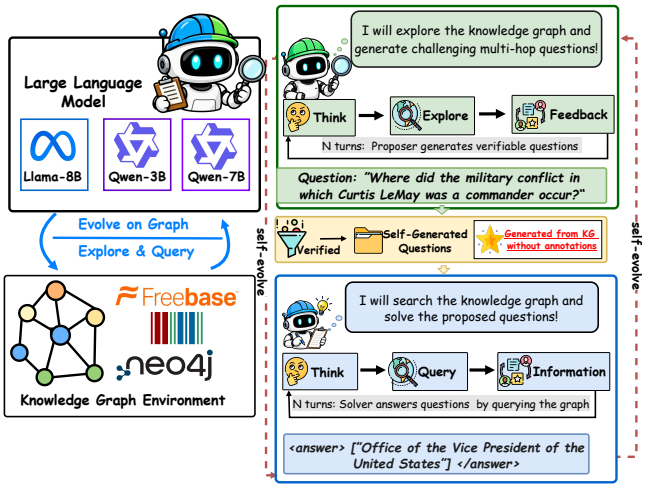

Figure 1: With only a KG and no question annotations, EoG lets an LLM self-evolve by proposing questions verified on the graph and then learning to solve them.

Knowledge graphs (KGs) [31] store large-scale structured facts and provide reliable external knowledge for LLM reasoning, but this requires the LLM to locate the facts needed at each reasoning step in a large and heterogeneous graph. Research on this problem has moved from one-shot graph retrieval [7, 16, 18, 28, 29] toward multi-turn graph exploration (Figure 2). Training-free agents [4, 38, 43, 54] prompt a frozen LLM to select relations and entities step by step, while RL-based agents [10, 21, 36, 56] treat the KG as an interactive environment and train the exploration policy on question–answer data with reinforcement learning (RL). However, most KGs in real-world settings such as industry and healthcare [20, 71] come without any question–answer annotations, and their data often cannot be handed to closed-source models [44]. This raises a question: *can an LLM improve its graph reasoning on a KG using only the KG itself, without any question–answer annotations?*

Answering this question, existing methods still face three challenges. (1) **The exploration ability of training-free agents cannot improve through interaction.** Training-free agents need no annotations and can be applied to a new KG directly, but their exploration strategy is fixed by the prompt and the frozen backbone, so they cannot accumulate experience from interacting with the graph, and their performance drops sharply on small open-source models [6, 47]. (2) **Training data rely on human annotation or closed-source models.** RL-based agents need large amounts

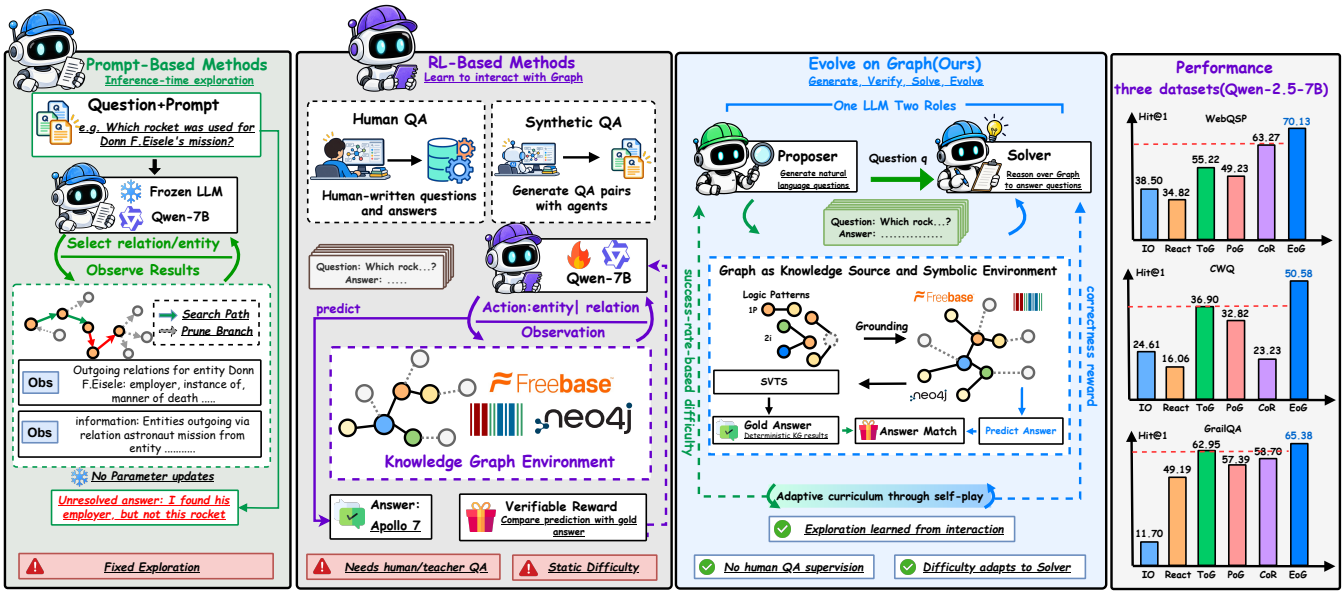

Figure 2: Comparison of KG reasoning paradigms. Prompt-based agents explore with a frozen LLM and RL-based agents learn from fixed QA written by humans or closed-source models, whereas EoG proposes and solves KG-verified tasks of adaptive difficulty. With Qwen-2.5-7B-Instruct, EoG outperforms all shown baselines on three Freebase benchmarks.

of question–answer data: human annotation [9, 17, 42, 62, 69] requires experts familiar with the KG schema and has to be redone for every new KG, while automatic synthesis [3, 56, 63] relies on closed-source models to synthesize question data. On a KG without annotations that cannot be exposed to closed-source models, neither kind of data is easy to obtain. (3) **The topology of training tasks cannot follow the agent's ability.** The difficulty of a graph reasoning question lies mainly in its reasoning topology. Existing training data fix the mix of topologies offline before training [1, 5] and do not adjust it as the agent learns. As the agent grows stronger, simple topologies are already mastered while complex ones remain unsolvable, and neither provides effective learning signals.

To address these challenges, we propose **Evolve on Graph (EoG)**, a framework in which a single LLM self-evolves on a KG using only the KG itself (Figure 1). First, we propose **Graph Self-Play**, in which the same LLM alternates between a *Proposer* that poses questions and a *Solver* that answers them. The Proposer explores the graph from real anchor entities, recovers the facts, and writes a question, while the Solver explores the graph to answer it. The two roles share parameters. Second, we propose **Symbolically Verifiable Topology Synthesis (SVTS)**, which samples reasoning topologies on the graph and executes the corresponding first-order logic queries, giving each task a topology skeleton and a sound and complete answer set. The Proposer explores the graph according to the topology skeleton and writes a multi-hop question. In this way, training tasks and their answers come only from the KG itself, without human annotation or closed-source models. Third, we propose the **Boundary-Aware Graph Curriculum**, which reinforces the Proposer according to the Solver's score variance, so that it keeps posing questions on the Solver's *capability boundary*, i.e., questions the Solver has not mastered but can still solve. As the Solver grows stronger, more questions on complex topologies fall on its boundary and enter training, forming a graph curriculum that evolves with the Solver.

On six benchmarks over Freebase, Wikidata, and an industrial KG, EoG outperforms prompting-, agent-, and retrieval-based methods without any annotations. With Qwen-2.5-7B-Instruct, EoG exceeds the strongest baseline by +7.46 (Hit@1) and +12.05 (F1) on average, and by +17.80 (Hit@1) on the industrial KG. EoG thus offers a new self-evolving path for training LLM graph reasoning on real-world KGs without any question–answer annotations.

2 Related Work

Knowledge-graph-augmented reasoning. Retrieval-based methods [7, 8, 16, 18, 25, 28–30, 34, 47, 64, 65] retrieve paths, triples, or subgraphs in one shot, training-free agents [4, 13, 15, 19, 24, 32, 38, 39, 43, 48, 52, 54, 57, 60, 68] prompt a frozen LLM to explore the graph over multiple turns, and trained agents [10, 14, 21, 22, 27, 36, 37, 40, 41, 49, 50, 56, 58, 63] learn exploration from question–answer data with RL. These methods either keep exploration fixed or learn it from annotations on one KG, whereas EoG trains graph reasoning on the target KG alone, without question–answer annotations.

Training data for graph reasoning. Benchmarks such as WebQSP [62], CWQ [42], and GrailQA [9] are annotated by experts familiar with the KG schema at a high cost, and GraphWalker [56] scales up data with random walks and closed-source LLMs; both fix their topology mix before training. Self-play lets LLMs pose and solve their own questions in code and math reasoning [11, 70] and search agents [26, 67], but their knowledge comes from model parameters or unstructured text. Without annotations or closed-source models, EoG generates its own tasks on the KG through self-play, with exact answers from query execution and topologies that advance with the Solver (Table 1), improving the graph reasoning of LLMs.

Table 1: Comparison with representative methods. Training-free methods keep θ_0, RL agents learn from fixed tasks written by humans or closed-source models, and EoG learns from KG-verified tasks that adapt to the Solver.

Method	**Exploration** (§3.2)	**Training Tasks and Answers** (§3.3)	**Task Difficulty** (§3.4)
ReAct [61]	$a_t \sim \pi_{\theta_0}(\cdot \mid s_t)$ Frozen LLM calls the graph tools.	$\theta = \theta_0$ No training.	No training tasks.
ToG / PoG [4, 38]	$\mathcal{B}_{t+1} = \text{top-}k\ \pi_{\theta_0}(\cdot \mid q, \mathcal{B}_t)$ Prompted beam search or sub-goal planning.	$\theta = \theta_0$ No training.	No training tasks.
Graph Reranker	$o = \text{top-}k\ s\big(q, \mathcal{T}_L(\mathcal{E}_Q)\big)$ Reranks triples near the anchors once.	$\theta = \theta_0$ No training.	No training tasks.
KG-R1 [21]	$a_t \sim \pi_\theta(\cdot \mid s_t)$ Multi-turn exploration trained with RL.	$(q, \mathcal{A}) \sim \mathcal{D}_{\text{human}}$ Human-annotated QA on one KG.	$d \sim \mathcal{D}$, fixed Given before training.
GraphWalker [56]	$\pi_\theta \leftarrow \text{SFT}(\tau_{\text{teacher}}) \rightarrow \text{RL}$ Imitates closed-source trajectories, then RL.	$(q, \mathcal{A}) \sim \mathcal{D}_{\text{teacher}} \cup \mathcal{D}_{\text{CWQ}}$ Closed-source walk questions plus CWQ.	$d \sim \mathcal{D}$, fixed Synthesized offline.
EoG (Ours)	$\pi_\theta^{\text{P}},\ \pi_\theta^{\text{S}}$ with shared θ Both roles explore the KG and learn.	$q \sim \pi_\theta^{\text{P}}(\cdot \mid G_Q, \mathcal{E}_Q, \mathcal{G}),\ \mathcal{A}_Q = [\![\phi]\!]_{\mathcal{G}}$ No annotation; answers executed on the KG.	$R_{\text{P}}\big(\sigma^2_{\theta_t}(q)\big)$, Gaussian Peaks at the Solver's boundary.

3 Methods

As shown in Table 1, existing methods either keep graph exploration fixed, retrieve only once, or train on question–answer data written by humans or closed-source models with fixed difficulty. EoG (Figure 3) consists of Graph Self-Play (Section 3.2), SVTS (Section 3.3), and the Boundary-Aware Graph Curriculum (Section 3.4), which address challenges (1)–(3) of Section 1, respectively.

3.1 Problem Formulation

Definition 1: Logic Query. A knowledge graph $\mathcal{G} = (\mathcal{E}, \mathcal{R}, \mathcal{T})$ has an entity set $\mathcal{E}$, a relation set $\mathcal{R}$, and a triple set $\mathcal{T} \subseteq \mathcal{E} \times \mathcal{R} \times \mathcal{E}$. A multi-hop question corresponds to a first-order logic query

$$\phi(x) = \exists y_1, \ldots, y_k : \ e_1 \wedge e_2 \wedge \cdots \wedge e_n, \tag{1}$$

where x is the answer variable, $y_1, \ldots, y_k$ are intermediate variables, and each atom $e_i = r(u, u')$ links an anchor entity or a variable to a variable and is true if and only if $(u, r, u') \in \mathcal{T}$. The answer set of ϕ is its denotation $[\![\phi]\!]_{\mathcal{G}} = \{a \in \mathcal{E} \mid \phi(a) = \text{True}\}$.

Definition 2: Graph Environment. Following GraphWalker [56], an agent interacts with $\mathcal{G}$ in a think–query–observe loop through two tools: `get_relations` returns the relations of an entity, and `get_triples` returns its triples under given relations. At step t, given the state s_t of all previous turns, the agent takes an action a_t, a tool call or a final answer, and observes triples $o_t \subseteq \mathcal{T}$. A trajectory $\tau = (a_1, o_1, \ldots, a_T)$ collects the evidence

$$\mathcal{O}(\tau) = \bigcup_{t=1}^{T} o_t \subseteq \mathcal{T}, \tag{2}$$

where the last action a_T gives the final answer $\mathcal{A}(\tau) \subseteq \mathcal{E}$.

Problem Statement. Given $\mathcal{G}$ and an initial policy π_{θ_0}, we learn a policy π_θ that maximizes the expected F1 of answering multi-hop questions on $\mathcal{G}$, without any question–answer pair or closed-source model. Training uses only symbolic tasks $d = (G_Q, \mathcal{E}_Q, \mathcal{A}_Q)$ built from $\mathcal{G}$, where the query graph G_Q grounds a query ϕ on $\mathcal{G}$, $\mathcal{E}_Q$ is its set of anchor entities, and $\mathcal{A}_Q = [\![\phi]\!]_{\mathcal{G}}$ is its gold answer set.

3.2 Graph Self-Play

To let the exploration ability improve through interaction, we propose Graph Self-Play, in which a single policy π_θ plays two roles under different role prompts with shared parameters: the Proposer π_θ^{P} turns a task into a natural-language question, and the Solver π_θ^{S} answers it. Both roles reason inside `<think>`, call the graph tools of Definition 2 inside `<kg-query>`, and read the returned triples inside `<information>` blocks, so both act by exploring the same graph. The Solver aims to answer correctly, while the Proposer aims to pose valid questions that the Solver can learn from:

$$\begin{aligned} &\max_{\pi_\theta^{\text{S}}} \ \mathbb{E}_{d,\, q\sim\pi_\theta^{\text{P}},\, \rho\sim\pi_\theta^{\text{S}}}\big[R_{\text{S}}(\rho)\big], \\ &\max_{\pi_\theta^{\text{P}}} \ \mathbb{E}_{d,\, q\sim\pi_\theta^{\text{P}}}\big[R_{\text{P}}(q)\big] \quad \text{s.t.} \quad v(q) = 1, \end{aligned} \tag{3}$$

where d is a task synthesized by SVTS, R_{S} is the Solver reward defined below, $v(q)$ is the validity check of Section 3.3, and R_{P} of Section 3.4 depends on the Solver's performance on q. Since the two roles share θ, the policy gains exploration experience both when it poses a question and when it searches for the answers. It includes question generation, question solving, and Solver optimization.

Question generation. For each task $d = (G_Q, \mathcal{E}_Q, \mathcal{A}_Q)$, the Proposer reads the reasoning skeleton of G_Q, in which only the anchors appear as real entity names, while the intermediate and answer entities appear as placeholders. For example, the skeleton Goodison Park $\xleftarrow{\texttt{arena_stadium}}$ `Mid1@` $\xrightarrow{\texttt{owner_s}}$ `Ans@` should become "Who is the owner of the team that plays at Goodison Park?". Since the tools cannot be called on placeholders, the Proposer explores the graph over multiple turns from the real anchors $\mathcal{E}_Q$ to uncover the facts, and produces a trajectory τ^{P} that ends with a question

$$q \sim \pi_\theta^{\text{P}}(\cdot \mid G_Q, \mathcal{E}_Q, \mathcal{G}), \tag{4}$$

where q is extracted from the final `<result>` block of τ^{P}. The prompt asks q to name the anchors, keep the full reasoning chain of the skeleton, and hide intermediate and answer entities, so that answering q needs the same multi-hop reasoning as the task.

Question solving. For a question q, the Solver starts from the anchors $\mathcal{E}_Q$ given in its prompt and samples a group of n trajectories through multi-turn interaction with the graph,

$$\rho^j \sim \pi_\theta^{\text{S}}(\cdot \mid q, \mathcal{E}_Q, \mathcal{G}), \qquad j = 1, \ldots, n, \tag{5}$$

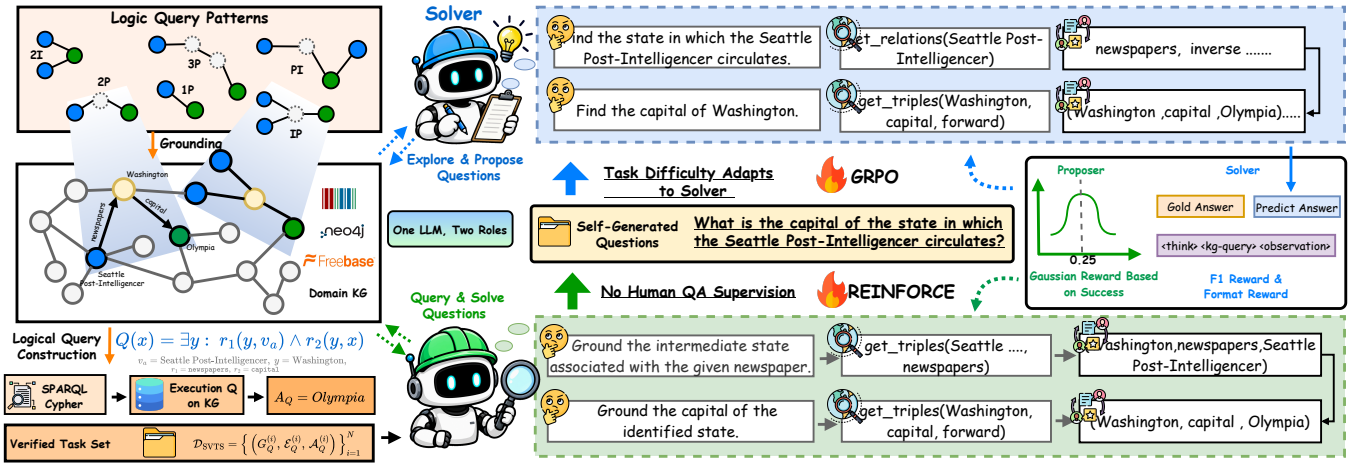

Figure 3: Overview of EoG. One LLM self-plays as the Proposer and the Solver, synthesizes verifiable tasks with SVTS, and forms a boundary-aware curriculum by rewarding the Proposer for questions on the Solver's capability boundary.

where the Solver decides which relations and entities to explore and when to stop, and each trajectory ends with a JSON list of entity names in an <answer> block, which gives the prediction $\hat{y}^j = \mathcal{A}(\rho^j)$. Each trajectory is scored by answer and format:

$$R_S(\rho^j) = \lambda_{\text{ans}}\, \text{F1}\big(\hat{y}^j, \mathcal{A}_Q\big) + \lambda_{\text{fmt}}\, R_{\text{fmt}}(\rho^j), \tag{6}$$

where the first term is the F1 between the predicted and gold answer sets, $R_{\text{fmt}} \in [0, 1]$ measures how well the trajectory follows the think–query–observe structure, and the weights are $\lambda_{\text{ans}} = 0.8$ and $\lambda_{\text{fmt}} = 0.2$. We use F1 rather than exact match because a task may have several gold answers, so a prediction that recovers only part of $\mathcal{A}_Q$ still receives partial credit. Since R_S scores only the final answer, any exploration path that reaches $\mathcal{A}_Q$ is rewarded, so the Solver learns its own exploration strategy from interaction.

Solver optimization. The Solver is updated with GRPO. For a group $\{\rho^j\}_{j=1}^n$ on the same q, the relative advantage is $\hat{A}^j = \big(R_S(\rho^j) - \text{mean}\big)/\text{std}$, with mean, std over $\{R_S(\rho^k)\}_{k=1}^n$, and the objective is

$$\mathcal{J}_S(\theta) = \mathbb{E}\left[\frac{1}{n}\sum_{j=1}^{n}\frac{1}{|\rho^j|}\sum_{t=1}^{|\rho^j|}\min\Big(w_t^j \hat{A}^j,\quad \text{clip}(w_t^j, 1-\epsilon, 1+\epsilon)\hat{A}^j\Big)\right] - \beta\, \mathbb{D}_{\text{KL}}\big[\pi_\theta^S \,\|\, \pi_{\text{ref}}\big], \tag{7}$$

where $w_t^j = \pi_\theta^S(a_t^j \mid s_t^j)/\pi_{\theta_{\text{old}}}^S(a_t^j \mid s_t^j)$ is the token-level importance ratio at step t, ϵ is the clipping threshold, and β is the KL coefficient.

3.3 Symbolically Verifiable Topology Synthesis

To obtain reliable training data without human annotation or closed-source models, we propose SVTS, which views the KG as a first-order logic structure and takes query execution on the graph as the source of gold answers. It includes three steps: defining query patterns, grounding and execution, and question verification.

Query patterns. Following complex logical queries over KGs [33], we define six query patterns as reasoning topologies:

$$\mathcal{P} = \{1p,\ 2p,\ 3p,\ 2i,\ ip,\ pi\}, \tag{8}$$

where p and i denote projection (one hop along a relation) and intersection (multiple constraints on one variable), respectively; $1p$–$3p$ are 1- to 3-hop chain queries, $2i$ is a two-way intersection, and ip, pi are nested compositions of projection and intersection. For example, the $2p$ and $2i$ query patterns are written as

$$\begin{aligned}\phi_{2p}(x) &= \exists y:\ r_1(v_a, y) \wedge r_2(y, x),\\ \phi_{2i}(x) &= r_1(v_{a_1}, x) \wedge r_2(v_{a_2}, x),\end{aligned} \tag{9}$$

where ϕ_{2p} follows a two-hop chain from one anchor, whereas ϕ_{2i} requires the answer to satisfy constraints from two anchors; the logical forms of all six patterns are listed in Appendix B. On Freebase, the synthesized questions cover all six patterns, with shares ranging from 9.5% (pi) to 22.4% ($1p$), so the Solver must learn path expansion, constraint intersection, and intermediate-variable tracking rather than fit a single pattern (Appendix C.4).

Grounding and execution. Given a pattern $pat \in \mathcal{P}$, grounding instantiates pat on real facts and extracts its grounded query graphs $\mathcal{I}_{pat} = f_{pat}(\mathcal{G})$, where f_{pat} is the corresponding extraction procedure: it samples a seed entity from a seed pool and walks on $\mathcal{G}$ to bind every variable and relation of pat to a real triple, and it repeats until a target number of tasks per pattern is collected or an attempt budget runs out. The walk only proposes the structure of a task; its answers are decided by execution below. For example, a $2p$ grounded query graph sampled from Freebase is Shaping $\xrightarrow{\text{recipes}}$ Country Sourdough Bread $\xrightarrow{\text{techniques}}$ x, whose execution returns two answers, Kneading and Slashing dough. We then compile each $G_Q \in \mathcal{I}_{pat}$ deterministically into an executable KG query Q for the logic query ϕ of Equation 1, and evaluate it on $\mathcal{G}$ to obtain the complete gold answer set, i.e., the denotation of ϕ:

$$\mathcal{A}_Q = \text{Execute}(Q, \mathcal{G}) = [\![\phi]\!]_{\mathcal{G}}, \tag{10}$$

where queries whose answers are empty, abnormally large, or unnameable are discarded rather than truncated, so that every kept task has an answer set that can be checked objectively. Unlike random-walk synthesis, which takes a single entity of the walk as

the answer, execution returns every entity that satisfies the query. This entity is always one of these answers, but when $|\mathcal{A}_Q| > 1$, scoring a complete prediction against it alone would give an F1 of only $2/(1+|\mathcal{A}_Q|)$. SVTS runs once per KG before training and only needs query access to the KG, so applying EoG to a new KG requires only its query endpoint. It finally outputs the task set

$$\mathcal{D}_{\text{SVTS}} = \left\{(G_Q^{(i)}, \mathcal{E}_Q^{(i)}, \mathcal{A}_Q^{(i)})\right\}_{i=1}^{N}, \tag{11}$$

where N is the task count (Algorithm 1 in Appendix B).

Question verification. Execution guarantees that $\mathcal{A}_Q$ answers the query Q, but not that the question q written by the Proposer expresses the same query. To keep q faithful to its task, we collect the evidence $O(\tau^{\text{P}})$ retrieved by the Proposer and let the Solver answer q using *only* that evidence as a verifier:

$$\sigma \sim \pi_\theta^{\text{S}}\big(\cdot \mid q, \mathcal{E}_Q, O(\tau^{\text{P}})\big). \tag{12}$$

A question is valid only if it passes a rule check $v_{\text{rule}}(q)$, which requires that q can be extracted from the `<result>` block, has a reasonable length, and names no intermediate or answer entity, and if the verification answer matches the gold answer:

$$v(q) = v_{\text{rule}}(q) \cdot \mathbb{I}\big[\, \mathcal{A}(\sigma) = \mathcal{A}_Q \,\big] \in \{0, 1\}, \tag{13}$$

where $\mathbb{I}[\cdot]$ is the indicator function. Only questions with $v(q) = 1$ enter reward computation and policy update. The rule check removes questions that leak intermediate or answer entities. The evidence check removes questions that are ambiguous, under-specified, or inconsistent with their task: since the verifier sees only the Proposer's triples, it recovers exactly $\mathcal{A}_Q$ only if these triples, read through q, determine the complete answer set. The check is conservative: a faithful question may be rejected when the single verification trajectory σ fails, lowering the yield of valid questions, but a question whose evidence does not determine $\mathcal{A}_Q$ is never accepted.

3.4 Boundary-Aware Graph Curriculum

To let task topologies follow the Solver's ability, we propose the Boundary-Aware Graph Curriculum, which rewards the Proposer by the Solver's score variance, so that it keeps posing questions on the Solver's capability boundary. It includes three parts: the capability boundary, the Proposer reward, and Proposer optimization.

Capability boundary. For a valid question q with Solver trajectories $\{\rho^j\}_{j=1}^n$, let $s_j = \text{F1}(\hat{y}^j, \mathcal{A}_Q) \in [0, 1]$ be the answer score of ρ^j in Equation 6 and $\bar{s}(q)$ their mean. The Solver's score variance on q, which measures how differently its trajectories score, is

$$\sigma^2(q) = \frac{1}{n}\sum_{j=1}^{n}\big(s_j - \bar{s}(q)\big)^2. \tag{14}$$

Since the advantages in Equation 7 are normalized within a group, a question on which all n trajectories receive the same score ($\sigma^2(q) = 0$), such as one the Solver always or never solves, gives almost no learning signal [66]. The Solver learns mainly from questions with $\sigma^2(q) > 0$, which it has not mastered but can still partly solve, such as one it answers in full on some trajectories and only in part on others; we call them its *capability boundary*.

Proposer reward. We therefore reward the Proposer with a Gaussian boundary reward on the Solver's score variance:

$$R_{\text{P}}(q) = \exp\left(-\frac{\left(\sigma^2(q) - \frac{1}{4}\right)^2}{2h}\right), \tag{15}$$

where $h > 0$ controls the peak width. Since scores lie in $[0, 1]$, $\sigma^2(q) \le \frac{1}{4}$, with equality only when half of the trajectories score 1 and the rest 0. The reward is thus minimal, $\exp(-\frac{1}{32h})$, when all scores are equal, higher on the capability boundary, and maximal at such an even split. As the scores are F1 against the complete answer sets of SVTS, a question the Solver answers only in part stays on the boundary instead of counting as solved. The reward also stays low when every trajectory fails, so the Proposer is not reinforced by questions the Solver cannot learn from.

Proposition 1. *With equal format scores, R_{P} is maximal only where the learning signal peaks and minimal where it vanishes.*

Proof. We provide experimental results on the Proposer rewards in Section 4.6 and a proof in Appendix A. □

Proposer optimization. Since the Proposer writes a single question per task, it has no group of samples for a relative advantage; we therefore update it with REINFORCE, taking the Proposer reward as the return of its whole question trajectory:

$$\mathcal{J}_{\text{P}}(\theta) = \mathbb{E}\left[R_{\text{P}}(q)\sum_{t=1}^{|\tau^{\text{P}}|}\log \pi_\theta^{\text{P}}(a_t^{\text{P}} \mid s_t^{\text{P}})\right], \tag{16}$$

where a_t^{P} and s_t^{P} are the action and state of the Proposer at step t. The two roles update alternately: as the Solver masters a kind of question, its scores approach 1, so σ^2 and R_{P} fall, driving the Proposer toward harder questions; the boundary thus moves outward and the two roles co-evolve into an adaptive curriculum.

Training loop. At each training step, the shared policy first acts as the Proposer on a batch of SVTS tasks and writes one question per task. Each question is checked by Equation 13; questions that fail the check are skipped by both updates, and for every valid question the Solver samples n trajectories on the graph, which yield both the Solver rewards R_{S} and, through $\sigma^2(q)$, the Proposer reward R_{P}. The same parameters θ are then updated with GRPO on the Solver trajectories and with REINFORCE on the Proposer trajectories. Because both roles are updated at every step, a change in the Solver's ability is reflected in R_{P} at the next step, so the curriculum follows the Solver without a manually designed difficulty schedule. At test time, only the Solver is used, and it answers human-written questions with the same tools. The whole loop thus uses only the graph: tasks come from SVTS, answers from query execution, and difficulty from the Solver's own score variance on each question.

4 Experiments

We answer five research questions (RQs). **RQ1:** Without any annotations, can EoG improve LLM graph reasoning on different KGs? **RQ2:** Is each component of EoG necessary? **RQ3:** Can Graph Self-Play improve exploration through interaction? **RQ4:** Can SVTS replace human annotation as the training data source? **RQ5:** Can the curriculum advance from simple to complex topologies?

Table 2: Main results on five public KGQA benchmarks (Hit@1 and F1, %; IndustryCarBench in Figure 4). Full-supervised methods are for reference, with icons marking training sets; bold marks the best annotation-free result; "–" means not evaluated.

Method	LLM	WebQSP W		CWQ C		GrailQA G		GTSQA T		WikiWebQuestions		Avg.	
		Hit@1	F1	Hit@1	F1	Hit@1	F1	Hit@1	F1	Hit@1	F1	Hit@1	F1
Full-Supervised Methods													
KG-R1† W	Qwen-2.5-7B	78.42	69.02	50.92	46.88	61.64	59.29	–	–	–	–	63.66	58.40
KG-R1† C	Qwen-2.5-7B	80.71	68.57	58.03	53.93	67.81	65.52	–	–	–	–	68.85	62.67
KG-R1† G	Qwen-2.5-7B	71.06	61.64	46.02	42.43	69.13	64.88	–	–	–	–	62.07	56.32
KG-R1† T	Qwen-2.5-7B	–	–	–	–	–	–	48.03	49.23	75.39	66.28	61.71	57.76
Annotation-Free Methods													
IO	Qwen-2.5-3B	32.21	20.60	12.42	13.72	11.00	8.70	7.83	7.45	37.43	23.00	20.18	14.69
	Qwen-2.5-7B	38.50	24.20	24.61	20.25	11.70	8.96	10.11	9.16	45.87	28.45	26.16	18.20
	LLaMA-3.1-8B	48.14	31.79	29.85	24.30	16.30	11.13	12.95	11.56	51.99	34.09	31.85	22.57
CoT [53]	Qwen-2.5-3B	34.96	22.20	17.76	15.00	11.00	8.70	5.73	4.92	39.83	23.41	21.86	14.85
	Qwen-2.5-7B	40.94	26.98	28.38	23.70	15.40	12.25	10.23	9.15	47.30	30.27	28.45	20.47
	LLaMA-3.1-8B	43.99	30.01	23.93	20.18	15.20	11.16	9.06	8.27	51.86	33.87	28.81	20.70
SC [51]	Qwen-2.5-3B	50.64	23.50	29.31	14.99	22.10	9.81	13.13	6.40	40.50	25.47	31.14	16.03
	Qwen-2.5-7B	50.58	27.94	38.97	26.08	23.00	12.67	17.76	9.90	60.46	31.19	38.15	21.56
	LLaMA-3.1-8B	61.81	33.77	43.39	25.78	28.00	13.80	20.84	11.85	66.71	35.70	44.15	24.18
ReAct [61]	Qwen-2.5-3B	22.63	18.25	12.06	10.63	21.26	18.61	15.47	14.72	36.46	29.56	21.58	18.35
	Qwen-2.5-7B	34.82	28.42	16.06	13.45	49.19	50.07	29.41	26.52	49.14	39.63	35.72	31.62
	LLaMA-3.1-8B	48.42	40.53	17.73	15.00	27.33	28.21	18.87	17.99	33.83	28.73	29.24	26.09
ToG [38]	Qwen-2.5-7B	55.22	43.25	36.90	35.70	62.95	56.07	38.04	34.21	64.52	49.07	51.53	43.66
PoG [4]	Qwen-2.5-7B	49.23	36.83	32.82	28.57	57.39	50.79	35.02	32.40	53.26	40.14	45.54	37.75
CoR [23]	Qwen-2.5-7B	63.27	57.62	23.23	21.32	58.70	54.01	26.63	24.48	66.39	54.00	47.64	42.29
Text-Based Retrieval	Qwen-2.5-7B	59.86	42.78	31.61	28.78	53.64	46.60	37.73	36.44	67.22	50.00	50.01	40.92
Graph Reranker	Qwen-2.5-7B	68.27	48.89	37.69	33.28	56.28	48.99	37.85	36.65	70.37	53.48	54.09	44.26
EoG (Qwen-2.5)	Qwen-2.5-3B	63.57	55.22	38.14	36.84	56.17	55.05	30.95	28.76	41.26	34.30	46.02	42.03
	Qwen-2.5-7B	70.13	**62.10**	**50.58**	**49.27**	**65.38**	**63.01**	**48.39**	**46.36**	**73.28**	60.82	**61.55**	**56.31**
EoG (LLaMA-3.1)	LLaMA-3.1-8B	**71.82**	61.51	50.36	48.25	63.87	62.59	36.37	34.67	70.53	**62.55**	58.59	53.91

4.1 Experimental Setup

Datasets. We evaluate on WebQSP [62], CWQ [42], and GrailQA [9] over Freebase, GTSQA [2] and WikiWebQuestions [55] over Wikidata, and IndustryCarBench, built on a real-world automotive-aesthetics KG of car models, design elements, colors, materials, and brands, with 500 expert-written test questions and no training set. Details of all six datasets are in Appendix C.1.

Baselines. We compare EoG with prompting-based (IO, CoT [53], SC [51]), agent-based (ReAct [61], ToG [38], PoG [4], CoR [23]), and retrieval-based methods (Text-Based Retrieval, Graph Reranker), all using the same backbone and KG as EoG. GraphWalker [56] synthesizes data with closed-source models, which our setting excludes; the random-walk variant in Table 3 approximates it. KG-R1† (Section 4.5) is a fully supervised reference trained on the training set of each benchmark, whose annotations EoG never uses.

Implementation Details. Following prior work [21, 38], we report Hit@1 and F1. We build EoG on Qwen-2.5-3B/7B-Instruct [59] and LLaMA-3.1-8B-Instruct [45], and train it with Verl [35] for 140 steps on 8 GH200 GPUs. The Solver uses GRPO and the Proposer REINFORCE, with 5 rollouts per task, a learning rate of 1×10^{-6}, and at most 15 assistant turns. Unless otherwise stated, analyses use Qwen-2.5-7B-Instruct on the three Freebase benchmarks. Appendix C.3 lists the full set of hyperparameters.

4.2 Main Results (RQ1)

As shown in Table 2 and Figure 4, without any question–answer annotations, EoG enables open-source LLMs to outperform prompting-, agent-, and retrieval-based methods on all three KGs. With Qwen-2.5-7B-Instruct, it exceeds the strongest baseline by **7.46** points in average Hit@1, and by **11.61** and **10.35** points on the multi-hop CWQ and GTSQA, showing that exploration learned through interaction handles complex questions better than retrieval or prompting. On Wikidata, EoG trains directly on the target KG and leads all annotation-free methods in Hit@1 on GTSQA and WikiWebQuestions, so the same pipeline carries over unchanged to a KG with a different schema. On IndustryCarBench, where annotation-based RL agents cannot be trained, EoG trains on the KG alone and exceeds the strongest baseline, ToG, by **17.80** points in Hit@1, so it applies directly to new KGs that lack annotations. On Qwen-2.5-3B and LLaMA-3.1-8B, against prompting baselines and ReAct, EoG leads in average Hit@1 and F1 (e.g., by 14.44 Hit@1 over SC on LLaMA-3.1-8B), and EoG on Qwen-2.5-3B (46.02) already surpasses PoG on Qwen-2.5-7B (45.54). On schema unseen in the annotations, EoG even outperforms the fully supervised KG-R1† trained on WebQSP or GrailQA by 2.65 and 5.84 points in F1 (Figure 6), showing that the exploration it learns through self-play over the whole KG generalizes well beyond the schema covered by annotations.

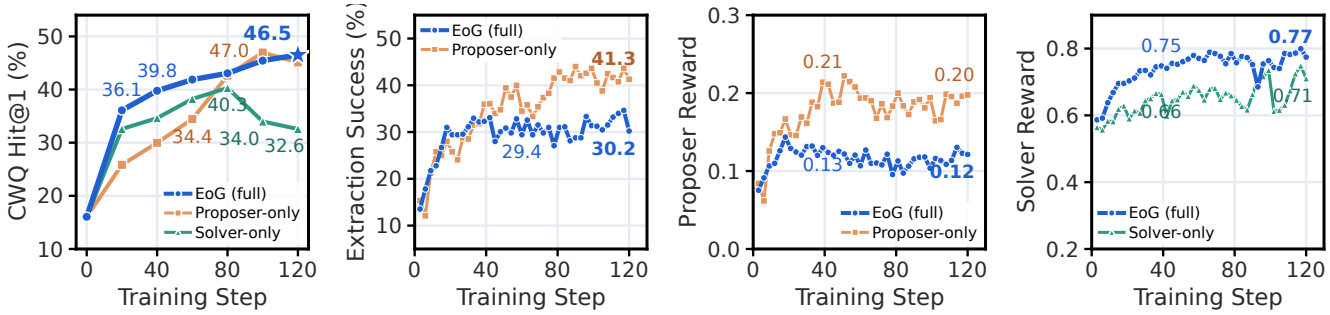

(a) CWQ Hit@1 (b) Extraction rate (c) Proposer reward (d) Solver reward

Figure 5: Training dynamics of role ablations. Proposer-only and Solver-only denote w/o Solver update and w/o Proposer update. (a) Validation Hit@1 on CWQ. (b) Question-extraction rate of the Proposer. (c–d) Proposer and Solver rewards.

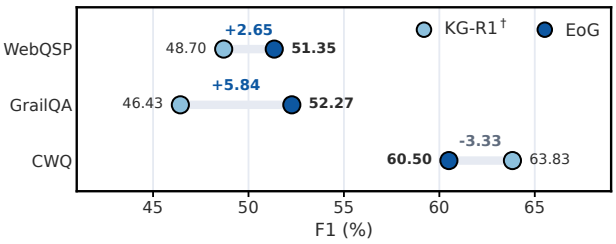

Figure 6: EoG vs. KG-R1† on unseen schema. F1 (%) of EoG and KG-R1† trained on each dataset with Qwen-2.5-7B-Instruct.

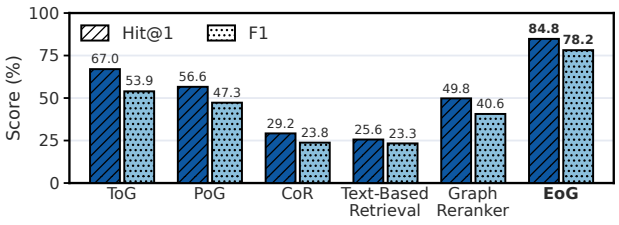

Figure 4: Results on IndustryCarBench. Hit@1 and F1 (%) of EoG and baselines with Qwen-2.5-7B-Instruct.

4.3 Ablation Study (RQ2)

As shown in Table 3, we conduct an ablation study on Qwen-2.5-7B-Instruct-based EoG with five ablation settings: removing the Solver update; removing the Proposer update; replacing the query execution of SVTS with random walks [56], which walk backward from an answer entity and take the start entity as the only answer of the question; and replacing the Gaussian boundary reward with the adversarial reward $1 - \bar{c}$ or the triangular boundary reward $1 - 2|\bar{c} - \frac{1}{2}|$, where $\bar{c}$ is the fraction of Solver trajectories that hit a gold answer. All five settings lower the Hit@1 on the multi-hop CWQ by 4.11–14.65 points, showing that Graph Self-Play, SVTS, and the curriculum each matter most on multi-hop questions.

4.4 Analysis of Graph Self-Play (RQ3)

Comparison with Training-Free Agents. To examine whether exploration comes from interaction rather than the prompt, we compare EoG with its untrained starting point, ReAct. EoG improves the average Hit@1 of ReAct by more than 24 points on all three backbones and also surpasses ToG (Table 2), because the exploration strategy of a training-free agent cannot improve from interaction.

Impact of Removing the Solver Update. We observe an interesting phenomenon: without Solver updates, the CWQ Hit@1 still rises from 16.06 to 46.39. We believe that proposing a question is itself a multi-hop graph exploration, which transfers to solving via the shared parameters. However, the variant still trails EoG by 16.39 and 6.88 on WebQSP and GrailQA, so the exploration gained from proposing only partly transfers; only when both roles are updated alternately do proposing and solving reinforce each other.

Table 3: Ablation study of EoG with Qwen-2.5-7B-Instruct on Freebase. Bold marks the best result in each column.

Variant	WebQSP		CWQ		GrailQA	
	Hit@1	F1	Hit@1	F1	Hit@1	F1
ReAct (untrained)	34.82	28.42	16.06	13.45	49.19	50.07
w/o Solver update	53.74	47.08	46.39	42.15	58.50	55.69
w/o Proposer update	63.14	56.91	36.22	35.83	55.67	55.18
w/o query exec. (Random Walk)	64.87	57.67	35.93	35.64	60.12	60.41
Adversarial reward	69.51	60.40	42.59	41.85	63.26	61.01
Triangular reward	**71.12**	**63.73**	46.47	44.03	63.46	62.10
EoG (full, Gaussian reward)	70.13	62.10	**50.58**	**49.27**	**65.38**	**63.01**

4.5 Analysis of SVTS (RQ4)

To examine whether SVTS can replace human annotation, we design KG-R1†, a variant of KG-R1 [21] that explores the full Freebase instead of a 2-hop subgraph pre-extracted for each question and differs from EoG only in being trained on annotated questions.

Comparison with Human Annotation. As shown in Table 2, KG-R1† is stronger in distribution, but out of distribution, the annotation-free EoG achieves an average F1 3.05 and 3.65 higher than KG-R1† trained on WebQSP and GrailQA, and trails KG-R1† trained on 27,623 annotated CWQ questions by only 4.49. To examine generalization beyond the annotated schema, we split the Freebase test sets into seen and unseen schema by whether all relations in their SPARQL query appear in the training set of KG-R1†, and compare them on unseen schema (Figure 6). EoG exceeds

(a) Synthesis success

(b) Simple vs. complex

(c) Topology heatmap

Figure 7: Evolution of task difficulty. Synthesis success rate of the Proposer over training steps (a) across three backbones, (b) on simple and complex topologies, and (c) per topology; complex topologies keep improving after simple ones plateau.

KG-R1† trained on WebQSP and GrailQA by 2.65 and 5.84 in F1, although both lead on seen schema; KG-R1† trained on CWQ still leads, but by 3.33 instead of 5.05. We believe that the gains of KG-R1† from annotation are confined to the schema covered by its training set, whereas EoG trains on SVTS tasks sampled across the whole KG, so its graph reasoning transfers better to unseen schema.

Comparison with Random-Walk Synthesis. SVTS outperforms random walks on all three datasets, by 14.65 in Hit@1 on CWQ (Table 3), for two reasons. First, random walks favor chain paths and rarely yield the intersection or nested questions that multi-constraint reasoning needs. Second, a random walk takes only its start entity as the answer, although the same relations often reach several entities; correct Solver answers are then judged wrong, so the Proposer can hack the boundary reward with questions whose answers are not unique. SVTS instead executes the query to obtain the complete answer set, which removes such incorrect rewards. Moreover, the synthesized questions rarely resemble the test questions: only 2.28% of them have a nearest Freebase test question with cosine similarity above 0.90 (Appendix C.4), so the gains of EoG do not come from test-set leakage into training.

4.6 Analysis of the Boundary-Aware Graph Curriculum (RQ5)

Impact of Removing the Proposer Update. Without Proposer updates, task topologies no longer follow the Solver: the Solver reward of this variant keeps rising, yet its CWQ Hit@1 drops after step 80 (Figure 5(a, d)), i.e., the Solver overfits the static task distribution of a Proposer without updates. This indicates that a Proposer without updates cannot shift to topologies that the Solver has not mastered, so the Solver stalls on early, simple questions and trails EoG by 14.36 points in Hit@1 on the multi-hop CWQ.

Impact of the Boundary Reward. We compare three Proposer rewards: the adversarial reward, the triangular boundary reward, and the Gaussian boundary reward. During training, the Solver occasionally produces overlong responses, which are truncated and temporarily lower its reward. With the adversarial reward, the Proposer cannot tell a hard question from such a truncated response and is rewarded most when the Solver performs worst: its reward drop lasts the longest among the three rewards, and its CWQ Hit@1 is 7.99 below that of the Gaussian boundary reward (Table 3). This shows that the Gaussian boundary reward is not reinforced by the Solver's truncated failures, so the curriculum returns to the Solver's capability boundary sooner after its outputs break down.

Evolution of Task Difficulty. As shown in Figure 7, the synthesis success rate of the Proposer rises on all three backbones (a), complex topologies rise from 17.9 to 31.3 while simple ones level off near 34, narrowing the gap to 2.5 points (b), and 3p keeps rising from 16 to 41 after simple patterns plateau (c), while the Solver reward does not decline (Figure 5(d)). This shows that the Proposer shifts toward complex topologies as the Solver improves, forming an adaptive graph curriculum that advances dynamically from easy to hard questions over time, without a manual schedule.

4.7 Case Study

Appendix D shows one complete self-play round on a 2i task whose executed answer is Jersey. The Proposer finds that Jèrriais is spoken in Jersey, checks that French is also spoken there, and asks "In which country are both Jèrriais and French spoken?", which names both anchors but hides the answer. The Solver retrieves the countries of each language, among them Dominica, Mali, and Chad for French, and keeps only Jersey, the one country shared with Jèrriais. The Solver thus answers by intersecting two entity sets on the graph rather than recalling one fact, the reasoning that 2i tasks train.

5 Conclusion

We present EoG, in which a single LLM self-evolves its graph reasoning on a KG using only the KG itself, aimed at KGs that have no question–answer annotations and cannot be handed to closed-source models. Graph Self-Play lets one policy explore the graph as both the Proposer and the Solver, SVTS derives verifiable tasks with complete answer sets by executing logic queries on the KG, and the Boundary-Aware Graph Curriculum keeps the questions on the Solver's capability boundary as it shifts from simple to complex topologies. On six benchmarks over three KGs, EoG outperforms prompting-, agent-, and retrieval-based methods, and on schema unseen in the annotations it exceeds KG-R1† trained on WebQSP or GrailQA. These results show that a KG can supply its own training tasks, answers, and curriculum, offering a direction for self-evolving graph reasoning on real-world KGs that lack annotations, requiring nothing beyond the query endpoint of each new target KG.

References

[1] Yuyang Bai, Zhuofeng Li, Ping Nie, Jianwen Xie, and Yu Zhang. 2026. GraphDancer: Training LLMs to Explore and Reason over Graphs via Two-Stage Curriculum Post-Training. *arXiv preprint arXiv:2602.02518* (2026).

[2] Alberto Cattaneo, Carlo Luschi, and Daniel Justus. 2025. Ground-truth subgraphs for better training and evaluation of knowledge graph augmented llms. *arXiv preprint arXiv:2511.04473* (2025).

[3] Ge Chang, Jinbo Su, Jiacheng Liu, Pengfei Yang, Yuhao Shang, Huiwen Zheng, Hongli Ma, Yan Liang, Yuanchun Li, and Yunxin Liu. 2026. Enhancing Agentic Textual Graph Retrieval with Synthetic Stepwise Supervision. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 25510–25525.

[4] Liyi Chen, Panrong Tong, Zhongming Jin, Ying Sun, Jieping Ye, and Hui Xiong. 2024. Plan-on-Graph: Self-Correcting Adaptive Planning of Large Language Model on Knowledge Graphs. In *Advances in Neural Information Processing Systems 37: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*.

[5] Zhuo Chen, Fei Wang, Zixuan Li, Zhao Zhang, Weiwei Ding, Chuanguang Yang, Yongjun Xu, Xiaolong Jin, and Jiafeng Guo. 2025. Knowcoder-a1: Incentivizing agentic reasoning capability with outcome supervision for KBQA. *arXiv preprint arXiv:2510.25101* (2025).

[6] Yi-Jie Cheng, Oscar Chew, and Yun-Nung Chen. 2025. The role of exploration modules in small language models for knowledge graph question answering. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 4: Student Research Workshop)*. 919–928.

[7] Rong Fu, Yemin Wang, Tianxiang Xu, Yongtai Liu, Weizhi Tang, Wangyu Wu, Xiaowen Ma, and Simon Fong. 2026. S-Path-RAG: semantic-aware shortest-path retrieval augmented generation for multi-hop knowledge graph question answering. In *Proceedings of the ACM Web Conference 2026*. 4057–4068.

[8] Guangze Gao, Zixuan Li, Chunfeng Yuan, Jiawei Li, Wu Jianzhuo, Yuehao Zhang, Xiaolong Jin, Bing Li, and Weiming Hu. 2025. D-RAG: Differentiable Retrieval-Augmented Generation for Knowledge Graph Question Answering. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing, EMNLP 2025, Suzhou, China, November 4-9, 2025*. Association for Computational Linguistics, 35398–35417.

[9] Yu Gu, Sue Kase, Michelle Vanni, Brian Sadler, Percy Liang, Xifeng Yan, and Yu Su. 2021. Beyond iid: three levels of generalization for question answering on knowledge bases. In *Proceedings of the web conference 2021*. 3477–3488.

[10] Haoyu Han, Kai Guo, Harry Shomer, Yu Wang, Yucheng Chu, Hang Li, Li Ma, and Jiliang Tang. 2026. Reasoning by exploration: A unified approach to retrieval and generation over graphs. In *Proceedings of the ACM Web Conference 2026*. 2408–2418.

[11] Chengsong Huang, Wenhao Yu, Xiaoyang Wang, Hongming Zhang, Zongxia Li, Ruosen Li, Jiaxin Huang, Haitao Mi, and Dong Yu. 2026. R-zero: Self-evolving reasoning llm from zero data. In *International Conference on Learning Representations*, Vol. 2026. 130770–130790.

[12] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2025. A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Trans. Inf. Syst.* 43, 2 (2025), 42:1–42:55.

[13] Jinhao Jiang, Kun Zhou, Zican Dong, Keming Ye, Xin Zhao, and Ji-Rong Wen. 2023. Structgpt: A general framework for large language model to reason over structured data. In *Proceedings of the 2023 conference on empirical methods in natural language processing*. 9237–9251.

[14] Jinhao Jiang, Kun Zhou, Xin Zhao, Yang Song, Chen Zhu, Hengshu Zhu, and Ji-Rong Wen. 2025. KG-Agent: An Efficient Autonomous Agent Framework for Complex Reasoning over Knowledge Graph. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*. Association for Computational Linguistics, 9505–9523.

[15] Bowen Jin, Chulin Xie, Jiawei Zhang, Kashob Kumar Roy, Yu Zhang, Zheng Li, Ruirui Li, Xianfeng Tang, Suhang Wang, Yu Meng, and Jiawei Han. 2024. Graph Chain-of-Thought: Augmenting Large Language Models by Reasoning on Graphs. In *Findings of the Association for Computational Linguistics, ACL 2024, Bangkok, Thailand and virtual meeting, August 11-16, 2024 (Findings of ACL, Vol. ACL 2024)*. Association for Computational Linguistics, 163–184.

[16] Dongcheon Lee, Ji-Yeon Park, Hye-Yoon Baek, Jimyeung Seo, Seyeong Kim, and Byungkook Oh. 2026. Diversifying Differentiable Graph Retrieval with Topic-Adaptive Multi-Intent Learning. In *Proceedings of the ACM Web Conference 2026*. 2273–2284.

[17] Chaofan Li, Lixing Chen, Junhua Tang, Yang Bai, Yutong Zhang, Zhi Zheng, Pan Zhou, and Zhe Qu. 2026. How Human Experts Educate Specialized LLMs: Filling Knowledge Gaps in KG-Augmented Generation through Hallucination Detection. In *Proceedings of the ACM Web Conference 2026*. 4172–4183.

[18] Mufei Li, Siqi Miao, and Pan Li. 2025. Simple is Effective: The Roles of Graphs and Large Language Models in Knowledge-Graph-Based Retrieval-Augmented Generation. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*. OpenReview.net.

[19] Songze Li, Zhiqiang Liu, Zhengke Gui, Huajun Chen, and Wen Zhang. 2025. Enrich-on-Graph: Query-Graph Alignment for Complex Reasoning with LLM Enriching. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing, EMNLP 2025, Suzhou, China, November 4-9, 2025*. Association for Computational Linguistics, 7672–7692.

[20] Lei Liang, Zhongpu Bo, Zhengke Gui, Zhongshu Zhu, Ling Zhong, Peilong Zhao, Mengshu Sun, Zhiqiang Zhang, Jun Zhou, Wenguang Chen, et al. 2025. Kag: Boosting llms in professional domains via knowledge augmented generation. In *Companion Proceedings of the ACM on Web Conference 2025*. 334–343.

[21] Junhong Lin, Shicheng Liu, Jinyeop Song, Song Wang, Julian Shun, and Yada Zhu. 2025. Efficient and transferable agentic knowledge graph rag via reinforcement learning. *arXiv preprint arXiv:2509.26383* (2025).

[22] Ben Liu, Jihai Zhang, Fangquan Lin, Cheng Yang, Min Peng, and Wotao Yin. 2025. Symagent: A neural-symbolic self-learning agent framework for complex reasoning over knowledge graphs. In *Proceedings of the ACM on Web Conference 2025*. 98–108.

[23] Chenhui Liu, Jianpeng Zhou, and Jiahai Wang. 2026. Chain-of-Relations: Faithful and Efficient LLM Reasoning over Knowledge Graphs via Relation-Centric Exploration. In *Findings of the Association for Computational Linguistics, ACL 2026, San Diego, California, United States, July 2-7, 2026*. Association for Computational Linguistics, 43108–43119.

[24] Yuanxiang Liu, Songze Li, Xiaoke Guo, Zhaoyan Gong, Qifei Zhang, Huajun Chen, and Wen Zhang. 2026. CoG: Controllable Graph Reasoning via Relational Blueprints and Failure-Aware Refinement over Knowledge Graphs. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026*. Association for Computational Linguistics, 15547–15568.

[25] Yu Liu, Xixun Lin, Yanmin Shang, Yangxi Li, Shi Wang, and Yanan Cao. 2026. PathMind: A Retrieve-Prioritize-Reason Framework for Knowledge Graph Reasoning with Large Language Models. In *Fortieth AAAI Conference on Artificial Intelligence, Thirty-Eighth Conference on Innovative Applications of Artificial Intelligence, Sixteenth Symposium on Educational Advances in Artificial Intelligence, AAAI 2026, Singapore, January 20-27, 2026*. AAAI Press, 15386–15393.

[26] Hongliang Lu, Yuhang Wen, Pengyu Cheng, Ruijin Ding, Jiaqi Guo, Haotian Xu, Chutian Wang, Haonan Chen, et al. 2026. Search self-play: Pushing the frontier of agent capability without supervision. In *International Conference on Learning Representations*, Vol. 2026. 65132–65165.

[27] Haoran Luo, Haihong E, Yikai Guo, Qika Lin, Xiaobao Wu, Xinyu Mu, Wenhao Liu, Meina Song, Yifan Zhu, and Anh Tuan Luu. 2025. KBQA-o1: Agentic Knowledge Base Question Answering with Monte Carlo Tree Search. In *Forty-second International Conference on Machine Learning, ICML 2025, Vancouver, BC, Canada, July 13-19, 2025 (Proceedings of Machine Learning Research, Vol. 267)*. PMLR / OpenReview.net.

[28] Linhao Luo, Yuan-Fang Li, Gholamreza Haffari, and Shirui Pan. 2024. Reasoning on Graphs: Faithful and Interpretable Large Language Model Reasoning. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.

[29] Linhao Luo, Zicheng Zhao, Gholamreza Haffari, Yuan-Fang Li, Chen Gong, and Shirui Pan. 2024. Graph-constrained reasoning: Faithful reasoning on knowledge graphs with large language models. *arXiv preprint arXiv:2410.13080* (2024).

[30] Costas Mavromatis and George Karypis. 2025. GNN-RAG: Graph Neural Retrieval for Efficient Large Language Model Reasoning on Knowledge Graphs. In *Findings of the Association for Computational Linguistics, ACL 2025, Vienna, Austria, July 27 - August 1, 2025 (Findings of ACL, Vol. ACL 2025)*. Association for Computational Linguistics, 16682–16699.

[31] Shirui Pan, Linhao Luo, Yufei Wang, Chen Chen, Jiapu Wang, and Xindong Wu. 2024. Unifying large language models and knowledge graphs: A roadmap. *IEEE Transactions on Knowledge and Data Engineering* 36, 7 (2024), 3580–3599.

[32] Joaquín Polonuer, Lucas Vittor, Iñaki Arango, Ayush Noori, David A. Clifton, Luciano Del Corro, and Marinka Zitnik. 2026. Autonomous Knowledge Graph Exploration with Adaptive Breadth-Depth Retrieval. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026*. Association for Computational Linguistics, 15694–15709.

[33] Hongyu Ren and Jure Leskovec. 2020. Beta embeddings for multi-hop logical reasoning in knowledge graphs. *Advances in Neural Information Processing Systems* 33 (2020), 19716–19726.

[34] Jing Ren, Bowen Li, Ziqi Xu, Xikun Zhang, Haytham Fayek, and Xiaodong Li. 2026. When to Trust: A Causality-Aware Calibration Framework for Accurate Knowledge Graph Retrieval-Augmented Generation. In *Proceedings of the ACM Web Conference 2026*. 3859–3869.

[35] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. 2025. Hybridflow: A flexible and efficient rlhf framework. In *Proceedings of the Twentieth European Conference on*

Computer Systems. 1279–1297.

[36] Yanlin Song, Ben Liu, Víctor Gutiérrez-Basulto, Zhiwei Hu, Qianqian Xie, Min Peng, Sophia Ananiadou, and Jeff Z Pan. 2026. Plan then retrieve: reinforcement learning-guided complex reasoning over knowledge graphs. In *Proceedings of the ACM Web Conference 2026*. 3666–3676.

[37] Yurun Song, Xiangqing Shen, Jianfei Yu, and Rui Xia. 2026. Interactive Semantic Parsing with Reinforcement Learning for Knowledge Graph Reasoning. In *Findings of the Association for Computational Linguistics: ACL 2026*. 20464–20484.

[38] Jiashuo Sun, Chengjin Xu, Lumingyuan Tang, Saizhuo Wang, Chen Lin, Yeyun Gong, Lionel M. Ni, Heung-Yeung Shum, and Jian Guo. 2024. Think-on-Graph: Deep and Responsible Reasoning of Large Language Model on Knowledge Graph. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.

[39] Jia Ao Sun, Hao Yu, Fabrizio Gotti, Fengran Mo, Yihong Wu, Yuchen Hui, Zhan Su, Lingfeng Xiao, and Jian-Yun Nie. 2026. Search-on-graph: Iterative informed navigation for large language model reasoning on knowledge graphs. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*. 4475–4485.

[40] Jia Ao Sun, Hao Yu, Fengran Mo, Zhan Su, Yuchen Hui, Bang Liu, and Jian-Yun Nie. 2026. Search-on-Graph-R1: Training Large Language Models to Search Knowledge Graphs with Reinforcement Learning. *arXiv preprint arXiv:2607.18481* (2026).

[41] Xin Sun, Zhongqi Chen, Xing Zheng, Qiang Liu, Shu Wu, Bowen Song, Zilei Wang, Weiqiang Wang, and Liang Wang. 2025. Kbqa-r1: Reinforcing large language models for knowledge base question answering. *arXiv preprint arXiv:2512.10999* (2025).

[42] Alon Talmor and Jonathan Berant. 2018. The Web as a Knowledge-Base for Answering Complex Questions. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT 2018, New Orleans, Louisiana, USA, June 1-6, 2018, Volume 1 (Long Papers)*. Association for Computational Linguistics, 641–651.

[43] Xingyu Tan, Xiaoyang Wang, Qing Liu, Xiwei Xu, Xin Yuan, and Wenjie Zhang. 2025. Paths-over-graph: Knowledge graph empowered large language model reasoning. In *Proceedings of the ACM on Web Conference 2025*. 3505–3522.

[44] Xingyu Tan, Xiaoyang Wang, Qing Liu, Xiwei Xu, Xin Yuan, Liming Zhu, and Wenjie Zhang. 2026. Privgemo: Privacy-preserving dual-tower graph retrieval for empowering llm reasoning with memory augmentation. *arXiv preprint arXiv:2601.08739* (2026).

[45] Llama Team. 2024. The Llama 3 Herd of Models. *CoRR* abs/2407.21783 (2024).

[46] Qwen Team. 2025. Qwen3 Technical Report. *CoRR* abs/2505.09388 (2025).

[47] Kaehyun Um, KyuHwan Yeom, Haerim Yang, Minyoung Choi, Hyeongjun Yang, and Kyong-Ho Lee. 2026. RPO-RAG: Aligning Small LLMs with Relation-aware Preference Optimization for Knowledge Graph Question Answering. In *Proceedings of the ACM Web Conference 2026*. 2513–2523.

[48] Shuai Wang and Yinan Yu. 2025. iQUEST: An Iterative Question-Guided Framework for Knowledge Base Question Answering. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2025, Vienna, Austria, July 27 - August 1, 2025*. Association for Computational Linguistics, 15616–15628.

[49] Shuai Wang and Yinan Yu. 2026. KG-Hopper: Empowering Compact Open LLMs with Knowledge Graph Reasoning via Reinforcement Learning. *CoRR* abs/2603.21440 (2026).

[50] Shuai Wang and Yinan Yu. 2026. KG-Reasoner: A Reinforced Model for End-to-End Multi-Hop Knowledge Graph Reasoning. *CoRR* abs/2604.12487 (2026).

[51] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net.

[52] Yingxu Wang, Shiqi Fan, Mengzhu Wang, Siyang Gao, Chao Wang, and Nan Yin. 2026. DAMR: Efficient and Adaptive Context-Aware Knowledge Graph Question Answering with LLM-Guided MCTS. In *The Fourteenth International Conference on Learning Representations, ICLR 2026, Rio de Janeiro, Brazil, April 23-27, 2026*. proceedings.iclr.cc.

[53] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*.

[54] Jiasheng Xu, Mingda Li, Yongqiang Tang, Peijie Wang, and Wensheng Zhang. 2026. Towards open-world retrieval-augmented generation on knowledge graph: A multi-agent collaboration framework. In *Proceedings of the ACM Web Conference 2026*. 2171–2182.

[55] Silei Xu, Shicheng Liu, Theo Culhane, Elizaveta Pertseva, Meng-Hsi Wu, Sina Semnani, and Monica Lam. 2023. Fine-tuned LLMs know more, hallucinate less with few-shot sequence-to-sequence semantic parsing over Wikidata. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. 5778–5791.

[56] Shuwen Xu, Yao Xu, Jiaxiang Liu, Chenhao Yuan, Wenshuo Peng, Jun Zhao, and Kang Liu. 2026. GraphWalker: Agentic Knowledge Graph Question Answering via Synthetic Trajectory Curriculum. *CoRR* abs/2603.28533 (2026).

[57] Yao Xu, Shizhu He, Jiabei Chen, et al. 2024. Generate-on-Graph: Treat LLM as both Agent and KG in Incomplete Knowledge Graph Question Answering. *CoRR* abs/2404.14741 (2024).

[58] Shiqi Yan, Yubo Chen, Ruiqi Zhou, Zhengxi Yao, Shuai Chen, Tianyi Zhang, Shijie Zhang, Wei-Qiang Zhang, Yongfeng Huang, Haixin Duan, et al. 2026. Explore-on-graph: Incentivizing autonomous exploration of large language models on knowledge graphs with path-refined reward modeling. In *International Conference on Learning Representations*, Vol. 2026. 56761–56789.

[59] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, et al. 2024. Qwen2.5 Technical Report. *CoRR* abs/2412.15115 (2024).

[60] Ruiyi Yang, Hao Xue, Imran Razzak, and Flora D Salim. 2026. Beyond single pass, looping through time: Kg-irag with iterative knowledge retrieval. In *Proceedings of the ACM Web Conference 2026*. 4460–4471.

[61] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2022. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629* (2022).

[62] Wen-tau Yih, Matthew Richardson, Christopher Meek, Ming-Wei Chang, and Jina Suh. 2016. The Value of Semantic Parse Labeling for Knowledge Base Question Answering. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics, ACL 2016, August 7-12, 2016, Berlin, Germany, Volume 2: Short Papers*. The Association for Computer Linguistics.

[63] Yuchen Ying, Weiqi Jiang, Tongya Zheng, Yu Wang, Shunyu Liu, Kaixuan Chen, and Mingli Song. 2026. GraphScout: Empowering Large Language Models with Intrinsic Exploration Ability for Agentic Graph Reasoning. In *Proceedings of the 32nd ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*. 6258–6269.

[64] Junchi Yu, Yujie Liu, Jindong Gu, Philip H. S. Torr, and Dongzhan Zhou. 2025. Can Knowledge-Graph-based Retrieval Augmented Generation Really Retrieve What You Need?. In *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2025, NeurIPS 2025, San Diego, CA, USA, December 2-7, 2025 / Mexico City, Mexico, November 30 - December 5, 2025*.

[65] Peng Yu, En Xu, Bin Chen, Haibiao Chen, and Yinfei Xu. 2026. STEM: Structure-Tracing Evidence Mining for Knowledge Graphs-Driven Retrieval-Augmented Generation. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026*. Association for Computational Linguistics, 7252–7285.

[66] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. 2026. Dapo: An open-source llm reinforcement learning system at scale. *Advances in Neural Information Processing Systems* 38 (2026), 113222–113244.

[67] Zhenrui Yue, Kartikeya Upasani, Xianjun Yang, Suyu Ge, Shaoliang Nie, Yuning Mao, Zhe Liu, and Dong Wang. 2026. Dr. Zero: Self-Evolving Search Agents without Training Data. *CoRR* abs/2601.07055 (2026).

[68] Chao Zhang, Minghan Li, Tianrui Lv, and Guodong Zhou. 2026. RFKG-CoT: Relation-Driven Adaptive Hop-count Selection and Few-Shot Path Guidance for Knowledge-Aware QA. In *Fortieth AAAI Conference on Artificial Intelligence, Thirty-Eighth Conference on Innovative Applications of Artificial Intelligence, Sixteenth Symposium on Educational Advances in Artificial Intelligence, AAAI 2026, Singapore, January 20-27, 2026*. AAAI Press, 34647–34655.

[69] Liangliang Zhang, Zhuorui Jiang, Hongliang Chi, Haoyang Chen, Mohammed Elkoumy, Fali Wang, Qiong Wu, Zhengyi Zhou, Shirui Pan, Suhang Wang, et al. 2026. Diagnosing and addressing pitfalls in KG-RAG datasets: Toward more reliable benchmarking. *Advances in Neural Information Processing Systems* 38 (2026).

[70] Andrew Zhao, Yiran Wu, Tong Wu, Quentin Xu, Yang Yue, Matthieu Lin, Shenzhi Wang, Qingyun Wu, Zilong Zheng, and Gao Huang. 2026. Absolute zero: Reinforced self-play reasoning with zero data. *Advances in Neural Information Processing Systems* 38 (2026), 105816–105879.

[71] Xuejiao Zhao, Siyan Liu, Su-Yin Yang, and Chunyan Miao. 2025. Medrag: Enhancing retrieval-augmented generation with knowledge graph-elicited reasoning for healthcare copilot. In *Proceedings of the ACM on Web Conference 2025*. 4442–4457.

A Proof of Proposition 1

Proposition 1. *With equal format scores, R_{P} is maximal only where the learning signal peaks and minimal where it vanishes.*

Proof. Let q be a valid question with Solver trajectories $\{\rho^j\}_{j=1}^n$, answer scores $s_j \in [0,1]$ and mean $\bar{s}$ as in Equation 14, and format scores that all equal f. Then $R_{\mathrm{S}}(\rho^j) = \lambda_{\mathrm{fmt}} f + \lambda_{\mathrm{ans}} s_j$ is a positive affine map of s_j, and since the group-normalized advantage is invariant to such maps, $\hat{A}^j = (s_j - \bar{s})/\sigma$ with $\sigma = \sqrt{\sigma^2(q)}$ when $\sigma > 0$, and $\hat{A}^j = 0$ when $\sigma = 0$.

Learning signal. We measure the Solver's learning signal on q by the mean absolute advantage of its group,

$$S(q) = \frac{1}{n}\sum_{j=1}^{n} |\hat{A}^j|. \tag{17}$$

If $\sigma^2(q) = 0$, all scores are equal, every $\hat{A}^j = 0$, and $S(q) = 0$. Otherwise the normalized advantages satisfy $\frac{1}{n}\sum_j (\hat{A}^j)^2 = 1$, so by the Cauchy–Schwarz inequality $S(q) \le 1$, with equality if and only if all $|\hat{A}^j|$ are equal, i.e., half of the scores lie at $\bar{s} + \sigma$ and the other half at $\bar{s} - \sigma$.

Maximum and minimum. Since $0 \le s_j \le 1$, we have $s_j^2 \le s_j$ and hence $\sigma^2(q) = \frac{1}{n}\sum_j s_j^2 - \bar{s}^2 \le \bar{s}(1-\bar{s}) \le \frac{1}{4}$, where the first inequality is tight only if every $s_j \in \{0,1\}$ and the second only if $\bar{s} = \frac{1}{2}$. Hence $R_{\mathrm{P}}(q)$ reaches its maximum 1 only when half of the trajectories score 1 and the other half score 0; there $\hat{A}^j = \pm 1$ and $S(q) = 1$, its largest possible value. Conversely, $R_{\mathrm{P}}(q)$ takes its minimum $\exp(-\frac{1}{32h})$ exactly when $\sigma^2(q) = 0$, which is exactly when $S(q) = 0$. With $n = 5$, an even split is impossible and $\sigma^2(q) \le 0.24$, attained only by a 2:3 split between scores 0 and 1. □

Remark. The adversarial and triangular rewards of Section 4.3 are functions of the success rate $\bar{c}$. The adversarial reward peaks at $\bar{c} = 0$, where every trajectory misses, all scores are 0, and $S(q) = 0$. The triangular reward peaks at $\bar{c} = \frac{1}{2}$ but, like the adversarial reward, does not reflect how much of the answer set each trajectory recovers.

B Algorithm Details

SVTS. As shown in Table 6 and Algorithm 1, SVTS repeats grounding, compilation, and execution for each of the six query patterns until it collects N tasks or exhausts T attempts. Grounding walks from a seed entity to bind the variables and relations of the pattern to real facts; the resulting query graph is compiled into an executable query and evaluated on the graph, and queries with empty, at least κ, or unnameable answers are discarded before training.

Algorithm 1: Symbolically Verifiable Topology Synthesis

Input: KG $\mathcal{G}$; pattern set $\mathcal{P}$; seed pool $\mathcal{E}_{\mathrm{seed}}$; per-pattern target N; attempt budget T; answer cap κ
Output: task set $\mathcal{D}_{\mathrm{SVTS}}$

```
D ← ∅
foreach pat ∈ P do
    n ← 0;  t ← 0
    while n < N and t < T do
        t ← t + 1;  sample seed s ∈ E_seed
        π ← random walk realizing pat from s in G
        if π = ⊥ then continue
        G_Q ← assemble(pat, π);  Q ← Compile(G_Q)
        A_Q ← Execute(Q, G)
        if A_Q = ∅ or |A_Q| ≥ κ or ¬Named(A_Q) then
            continue
        D ← D ∪ {(G_Q, E_Q, A_Q)};  n ← n + 1
    end
end
return D_SVTS ← D
```

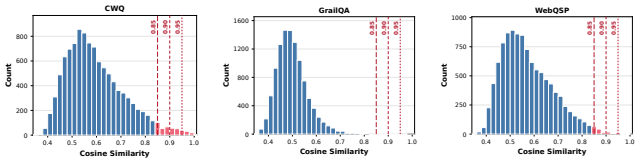

(a) CWQ (b) GrailQA (c) WebQSP

Figure 8: Semantic contamination analysis. Max cosine similarity of synthesized questions to each test set.

C Experimental Details

C.1 Datasets

On Freebase, **WebQSP** [62] contains single-hop and two-hop questions, **CWQ** [42] extends WebQuestions with multi-entity constraints, conjunctions, and superlatives, and **GrailQA** [9] tests i.i.d., compositional, and zero-shot generalization. On Wikidata, **GTSQA** [2] provides 1,622 test questions with gold answer subgraphs, and **WikiWebQuestions** [55] migrates WebQuestions to Wikidata with re-annotated SPARQL queries. **IndustryCarBench** is a domain benchmark that we build on a real-world automotive-aesthetics KG covering car models, design elements, attributes, colors, materials, and brands; it contains 500 test questions written by domain experts, with gold answers executable on the graph and has no annotated training set. EoG never uses the annotated training questions of any benchmark; only KG-R1† does.

C.2 Baselines and Metrics

All baselines use the same backbone as the compared EoG model. IO answers directly, CoT [53] reasons step by step, and SC [51] takes a majority vote over sampled CoT chains. ReAct [61] interleaves thoughts, graph-tool calls, and observations with the tools and Solver prompt of EoG. ToG [38] runs LLM-guided beam search over entities and relations, PoG [4] decomposes a question into subgoals with adaptive exploration and self-correction, and CoR [23] builds a relation chain and filters the reachable candidates with the question constraints. Text-Based Retrieval selects the top-k verbalized triples by embedding similarity, and Graph Reranker reranks the triples within the L-hop neighborhood of the anchors. Following [4, 38], Hit@1 checks whether the normalized prediction exactly matches any gold answer, and F1 measures the overlap between the predicted and gold answer lists of each question.

Table 4: Proposer on a $2i$ task. It explores both anchors and asks for the country they share.

Role: Proposer **Task:** Jèrriais ← location.country.languages_spoken– Ans@ ; French –language. human_language.countries_spoken_in → Ans@ **Executed answer:** Jersey
<think> The Path suggests that Jèrriais is a language spoken in a country, and French is also spoken in that same country. The question should ask for the country where both Jèrriais and French are spoken. Let's explore the knowledge graph to identify the country. </think>
<kg-query> get_relations("Jèrriais") </kg-query>
<information> location.country.languages_spoken, language.human_language.main_country, language.human_language.language_family, language.language_family.languages, people.ethnicity. languages_spoken, language.human_language.countries_spoken_in </information>
<think> The relation 'location.country.languages_spoken' confirms that Jèrriais is spoken in a country. Let's use 'get_triples' to find out which country it is. </think>
<kg-query> get_triples("Jèrriais", ["location.country.languages_spoken"]) </kg-query>
<information> [Jersey, location.country.languages_spoken, Jèrriais] </information>
<think> We have determined that Jèrriais is spoken in Jersey. Now, we need to find out if French is also spoken in Jersey. Let's check the languages spoken in Jersey using the relation 'language. human_language.countries_spoken_in'. </think>
<kg-query> get_triples("Jersey", ["language.human_language.countries_spoken_in"]) </kg-query>
<information> [French, language.human_language.countries_spoken_in, Jersey], [English Language, language.human_language.countries_spoken_in, Jersey], [Jèrriais, language.human_ language.countries_spoken_in, Jersey] </information>
<think> The information confirms that French is indeed spoken in Jersey. Thus, the answer to the question is Jersey. </think>
<result> {"question": "In which country are both Jèrriais and French spoken?"} </result>

Table 5: Solver on the question of Table 4. It intersects the countries of the two languages.

Role: Solver **Question:** In which country are both Jèrriais and French spoken?
<think> 1. The initial entities are Jèrriais and French. I need to find the country where both Jèrriais and French are spoken. I will start by querying the knowledge graph to find related information about Jèrriais and its spoken locations. </think>
<kg-query> get_relations("Jèrriais") </kg-query>
<information> location.country.languages_spoken, language.human_language.main_country, language.human_language.language_family, language.language_family.languages, people.ethnicity. languages_spoken, language.human_language.countries_spoken_in </information>
<think> I now have some potential relations to explore. I will query the knowledge graph with the first relation "location.country.languages_spoken" to see if it can provide the country where both languages are spoken. </think>
<kg-query> get_triples("Jèrriais", ["location.country.languages_spoken"]) </kg-query>
<information> [Jersey, location.country.languages_spoken, Jèrriais] </information>
<think> The knowledge graph has shown that Jèrriais is spoken in Jersey. Now, to find out if French is also spoken there, I will query the knowledge graph with the relation "language.human_language. countries_spoken_in" to verify if French is spoken in Jersey. </think>
<kg-query> get_triples("French", ["language.human_language.countries_spoken_in"]) </kg-query>
<information> [French, language.human_language.countries_spoken_in, Dominica], [French, language.human_language.countries_spoken_in, Mali], [French, language.human_language. countries_spoken_in, Chad], [French, language.human_language.countries_spoken_in, Jersey], [French, language.human_language.countries_spoken_in, Ivory Coast] </information>
<think> I have confirmed that French is spoken in Jersey alongside Jèrriais. Therefore, the country where both languages are spoken is Jersey. </think>
<answer> ["Jersey"] </answer>

Table 6: Basic logic query patterns. v_a: anchor entity; x: answer variable; y_i: intermediate variables; r_i: relations.

Pattern	Structure	Logical query
$1p$		$\phi(x) = r_1(v_a, x)$
$2p$	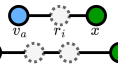	$\phi(x) = \exists y : \ r_1(v_a, y) \wedge r_2(y, x)$
$3p$	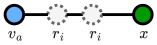	$\phi(x) = \exists y_1, y_2 : \ r_1(v_a, y_1) \wedge r_2(y_1, y_2) \wedge r_3(y_2, x)$
$2i$		$\phi(x) = r_1(v_{a_1}, x) \wedge r_2(v_{a_2}, x)$
ip	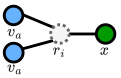	$\phi(x) = \exists y : \ r_1(v_{a_1}, y) \wedge r_2(v_{a_2}, y) \wedge r_3(y, x)$
pi	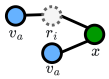	$\phi(x) = \exists y : \ r_1(v_{a_1}, y) \wedge r_2(y, x) \wedge r_3(v_{a_2}, x)$

Table 7: Settings of EoG. Hyperparameters of self-play training with Qwen-2.5-7B-Instruct on Freebase.

Parameter	Value
Training data	SVTS tasks
Training steps	140
Train / PPO mini-batch size	256 / 128
Actor learning rate	1×10^{-6}
KL loss coefficient	0.01
Rollout samples (n)	5
Max response length	2560
Max assistant turns	15
Rollout temperature	0.8
Solver / Proposer estimator	GRPO / REINFORCE
Solver reward	$0.8\,\mathrm{F1} + 0.2\,R_{\mathrm{fmt}}$

C.3 Implementation Details

We train EoG with Verl on 8 GH200 GPUs across two nodes, using FSDP with CPU offloading of parameters and optimizer states and SGLang for multi-turn rollouts; Table 7 lists the hyperparameters. KG-R1† converts the official training splits of WebQSP, CWQ, and GrailQA into the Solver format, which yields 3,058, 27,623, and 36,954 questions, and trains one model per dataset. It uses the same backbone, Solver prompt, graph tools, and Freebase endpoint as EoG, so only the source of the training questions differs. Its reward is the F1 of EoG without the format term.

C.4 Synthetic Data Analysis

Topology coverage. Synthesized questions on Freebase cover all six patterns: the chain patterns $1p$ (22.4%), $2p$ (20.6%), and $3p$ (16.5%), and the intersection patterns $2i$ (17.7%), pi (9.5%), and ip (13.3%). The Solver must learn path expansion, constraint intersection, and intermediate-variable tracking rather than fit a single pattern.

Domain coverage. The questions cover 20 domains, including Places & Geography (14.2%), Film & Television (13.2%), Government, Law & Military (12.6%), People & Society (8.7%), and Sports (8.4%); the top three together account for only 40.0%. Training signals are therefore not confined to a few dense subgraphs of Freebase.

Test-set leakage. To rule out test-set leakage, we encode each question with BGE-M3 and take the maximum cosine similarity between each synthetic question and the questions of each Freebase test set (Figure 8). Only 4.57%, 2.28%, and 0.71% of the synthetic questions exceed 0.85, 0.90, and 0.95, respectively, and the mean similarity is 0.634, so there is no meaningful leakage between the synthetic training data and the Freebase evaluation sets.

D Case Study

Tables 4 and 5 give the full trajectories of the case in Section 4.7. The Proposer starts from Jèrriais, finds that it is spoken in Jersey, and queries the languages of Jersey to confirm that French is spoken there before asking. The Solver instead retrieves every country where French is spoken, among them Mali, Chad, and Jersey, and keeps the one shared with Jèrriais, the executed answer.