arXiv:2610.36776v1 [cs.AI] 27 Oct 2026

Dependency Graphs of Mathematical Arguments: Representation, Structural Measurement, and Proof Review

Eunji Park

Graduate School of Data Science, Seoul National University
Gwanak-ro 1, Gwanak-gu, Seoul 08826, Korea

Email address: dmswl4743@snu.ac.kr

Abstract

We formulate a source-grounded framework for representing and measuring mathematical arguments with assistance from large language models. The inventory includes named statements and numbered equations, as well as unnumbered equations, local assumptions, constructions, and prose claims needed to represent a selected argument. Source occurrences are distinguished from contextual states and inference records; proposed dependencies retain evidence and review status. A measurement contract fixes manuscript versions, boundaries, granularity, proof choices, and macro anchors. On a selected finite directed acyclic graph, depth is the longest prerequisite-path length, and bulk counts distinct ancestor edges. We give a union recurrence preserving shared subarguments and establish that fixed-vertex reachability determines depth, whereas fixed-anchor reachability determines the macro skeleton. Unrestricted refinement need not preserve depth or bulk. The principal corpus-level objective is to analyze recurring groups of jointly used theorems, infrequent combinations, and results connecting different dependency communities. These patterns motivate probabilistic premise suggestions and targeted review of possible omissions or unchecked application conditions in human-written and AI-generated proofs. Suggested premises require mathematical validation; statistical association does not certify a proof. Further applications include source-linked reading support, cross-paper theorem exploration, and comparison of individual and collaborative argument shapes, with problem structure and writing roles treated as competing explanations. Measurements relative to specified axioms additionally require resolved foundational dependencies and fixed inference conventions. A versioned prompt protocol and an evaluation design separate annotation discovery, deterministic recomputation, and review assistance. These applications are proposed research directions rather than reported empirical findings. The measures describe a specified representation, not intrinsic mathematical difficulty, originality, or importance.

Keywords. Mathematical knowledge management; dependency graphs; argument structure; proof granularity; theorem reuse; premise selection.

1 Introduction

1.1 Historical Background

The organization of mathematical knowledge has both a logical and an expository dimension. A result depends on assumptions, definitions, and earlier arguments, but the order in which a paper presents these ingredients need not coincide with their dependency order. A calculation in an appendix may support a theorem in the main text, several results may reuse one construction, and

a paper may contain independent branches. A sequence of pages therefore provides only one view of the argument.

Making mathematical structure explicit has a long history. De Bruijn's Automath language represented mathematical objects and assertions through explicitly introduced names and categories, with texts designed for computer checking [Section 1, 1]. At the level of mathematical writing, Lamport advocated hierarchical proofs in which named steps expose the organization of an argument and allow readers to move between an outline and its details [pp. 601 and 607, 2]. These approaches address different questions, but both show why the structure supporting a statement deserves an explicit representation.

Mathematical knowledge management extends this concern from individual proofs to reusable bodies of mathematics. Theory graphs, as discussed by Kohlhase in the setting of OMDoc/MMT, represent mathematical theories and the interpretations connecting them [Section 3, 3]. Bibliometric networks offer a different perspective: Small's co-citation analysis records how frequently two documents are cited together [4]. Such document-level relations describe the organization of literature, but they do not specify which theorem, assumption, or intermediate calculation is used in a particular proof. The present study is concerned with that finer level, while retaining links to the source documents and their broader mathematical context.

1.2 Related Work

Alama, Mamane, and Urban extract detailed dependencies from Coq and Mizar and distinguish dependencies relevant to logical justification from those needed for the successful processing of formal texts [Introduction and Section 2, 5]. This is a direct methodological precedent for specifying what an edge means. Their work also cautions against treating all dependencies produced by a software system as interchangeable. Our setting is primarily informal mathematical writing, where identifying a dependency additionally requires interpreting the source and recording the evidence for that interpretation.

NaturalProofs provides a corpus of natural-language mathematical statements and proofs, an explicit statement-level reference graph, and reference-retrieval and reference-generation tasks [Sections 3–4 and Table 2, 6]. Its reference graph is an important precedent for representing relations within informal mathematics; the authors note that it can contain cycles. Its evaluation tasks concern retrieving and generating references used in a proof, rather than a complete decomposition of internal calculations or the stability of depth and bulk under such a decomposition. Reference links can supply candidate inputs to our framework, but contextual applicability and a compatible acyclic proof selection still have to be established.

Lean Blueprint supports author-declared dependencies separately for statements and proofs and links the exposition to named Lean declarations ["Editing the blueprint", 7]. KnowTeX develops a closely related approach directly from LaTeX sources, with manual annotations, rule-based dependency inference, interactive edge review, cycle detection, and transitive reduction [Sections 2 and 4.3–4.6, 8]. These are close precedents for a Paper Graph, including its use as a reading map. We therefore do not claim that source-linked dependency visualization, human review of proposed edges, or educational use of such graphs is new. KnowTeX also recognizes constructions, remarks, and custom environments, including items without explicit labels, but excludes equation and align environments from its node parser [Section 4.2, 8]. Its order-based inference rules propose dependencies requiring separate review [Section 4.4, 8]. Our emphasis is on a measurement contract that fixes source evidence, decomposition, and the treatment of shared subarguments and unresolved components.

TheoremGraph combines extractor-labeled candidate dependencies with semantic links between related statements within and across informal and formal mathematical corpora [Sections 3–6 and 10, 9]. It is a direct precedent for the cross-paper scope of a Math Graph. The distinction between a candidate match, a semantic similarity, and an established application of a theorem is essential here: an application requires the appropriate statement, context, substitution, and hypotheses. Our proposal organizes these distinctions for argument measurement; it does not claim priority for a statement-level mathematical network or for connecting informal and formal corpora.

Li, Peng, Severini, and Shafto analyze the network structure of Mathlib, including its dependency layers, depth, and community structure [Sections 4–5 and Appendices B–C, 10]. Their analysis emphasizes that logical, mathematical, and infrastructural relations can produce different network interpretations. Its declaration graph points from a declaration to its premises, whereas our measurement graph points from inputs to dependent outputs [Appendix B.2.1, 10]. Node-level depth and degree comparisons therefore require explicit alignment of edge conventions. The work is particularly relevant to our proposed study of structural shape and connections between fields. The distinction between statement and proof dependencies, and the warning against removing actual proof-use edges by transitive reduction, are particularly close to our separation of displayed structure and edge-count measurements [Appendices B.2.2 and C.2.3, 10]. Depth and network centrality are thus established analytical tools, not new invariants introduced here. Likewise, our macro-level construction uses the classical transitive reduction of a directed graph [11].

Finer representations also have direct precedents. Proofscape places assertions, assumptions, and definitions from informal proofs on nodes and supports interactive expansion of inference steps and cited results [Section 1, pp. 57–58, 12]. Mathematical Derivation Graphs studies equation dependencies with nodes explicitly restricted to numbered equations and evaluates LLM extraction [Sections 2 and 5.2, 13]. ProofFlow uses an LLM to extract conditions, definitions, intermediate claims, and conclusions as a proof DAG; its examples include unnumbered calculations and constructions, and it discusses step granularity [Sections 3.1 and 4.1, Figures 1–2, 14]. Thus, neither broad proof-step nodes nor LLM-assisted decomposition is claimed as a contribution in isolation. Our coverage requirement is to retain argument-bearing unnumbered material throughout the selected manuscript scope, under an explicit and auditable measurement policy.

Learning useful premises from proof corpora is likewise established. Alama et al. combine proof-dependency analysis with learned premise selection and automated re-verification [15]. NaturalProofs studies conditional reference generation [Section 5, 6], while TheoremGraph uses dependency-neighbor expansion to supplement retrieval [Appendix E.2, 9]. These are direct antecedents of corpus-based proof assistance. Our proposed application retains statistical suggestions as a separate relation type: only source-supported, contextually reviewed uses enter the selected measurement graph.

Taken together, these works supply substantial parts of the surrounding infrastructure. The contribution pursued here is a unified, source-grounded protocol for measuring a chosen argument, together with precise conditions for comparison and a research design for studying the resulting data. This integration specifies the method being proposed; it does not establish priority for its individual components or superior extraction accuracy. Any advantage over existing tools must be established on common annotated material.

1.3 Main Question

Our main question is the following:

> How can the dependencies actually used in a mathematical paper be represented and measured reproducibly across LLM-assisted analyses, so that both its internal argument structure and its connections to other papers can be compared under explicitly controlled conditions?

This question has three parts. First, which source occurrences and intermediate states should be distinguished, and what evidence makes a proposed dependency admissible? Second, which structural properties are preserved when a representation is simplified or refined? Third, once comparable graphs have been collected, which patterns in learning, exposition, collaboration, and theorem reuse can be investigated without confusing structural measurements with mathematical value?

The object being measured is a recorded argument under a specified representation. It is not a theorem independently of all its proofs. A single statement may admit several proofs, and the same proof may be expanded to different levels of detail. Even a fixed graph can support different questions: a longest prerequisite chain describes sequential dependence, whereas the union of all prerequisites describes the amount of recorded structure supporting a target. Neither quantity alone determines how difficult the argument is to understand or discover.

1.4 Proposed Framework and Methods

We use the term *Paper Graph* for a source-linked representation of definitions, assumptions, constructions, intermediate calculations, and results in one manuscript. Its data have three levels: document occurrences preserve what appears in the source; contextual states and inference records encode a chosen argument; and a macro skeleton records reachability between selected anchor states. The term *Math Graph* denotes a collection of such representations connected by identified theorem applications, with versions and contexts retained.

Large language models may assist in identifying unnumbered equations, local assumptions, object constructions, compatibility checks, and claims expressed in prose, in addition to explicitly labeled items. The inclusion criterion is an item's role in the recorded justification, rather than its typography or position. Each proposed item and dependency retains its source location, context, and evidential status for independent review. A complete inventory of numbered statements alone is therefore not a complete representation of the argument.

The source-linked reading view can keep these items accessible while presenting a compact overview. A separate computational representation records the substantive steps needed for measurement; screen layers and displayed item counts do not determine the structural metrics.

The central methodological device is a measurement contract $\mathcal{C} = (M, B, P, W, H)$. It fixes the manuscript and scope M, the accepted boundary inputs B, the granularity policy P, a consistent proof selection W, and the macro anchors H. Under this contract, a target v in a finite directed acyclic graph has depth $D(v)$, measured in edges along a longest incoming path, and bulk $S(v)$, measured by distinct edges in its ancestor subgraph. Shared subarguments are counted by set union, not by adding the sizes of branches. We also consider the edge-to-node ratio $q(G) = |E|/|V|$ as a first summary of dependency per node, while keeping its interpretation separate from general graph complexity.

A central design objective is to make the extracted dependency structures and their measurements as stable as possible across different LLMs. With the source and measurement contract held fixed, the choice of extractor should have as little influence as possible on which argument steps and actual uses are recorded and on the resulting D and S. This objective includes repeated runs of one

model and comparisons between models; it is an empirical requirement to evaluate, not a guarantee supplied by identical prompts.

Such comparisons require source- and context-preserving alignment of states under the same granularity and proof-selection policies. Agreement on states, dependency edges, anchor relations, and target-level D and S should be reported separately: equal measurements do not imply identical graphs, and model agreement does not establish mathematical validity. Disagreements should remain traceable to the source for review. Once the same accepted graph is fixed, deterministic recalculation must give the same measurements independently of the LLM that proposed its annotations. Sections 8 and 10 distinguish this computational requirement from stability in discovering the graph itself.

The structural results establish the scope of these definitions. Proposition 4.2 gives the union recurrence; Propositions 4.3–4.5 address reachability, refinement, and conditional bounds; and Proposition 5.2 characterizes the macro skeleton. These are elementary consequences of finite DAG structure. Their role is to justify the measurement protocol, not to claim new general graph theory. The remaining methods concern source annotation, explicit uncertainty, independent recalculation, and evaluation against expert reference data. Section 9 introduces PGPP 0.1, a proposed standardized prompt protocol that makes these requirements operational. Its complete reusable prompt and output contract are given in Appendix A.

1.5 Future Directions

Our principal corpus-level objective is to study the joint use of mathematical results as evidence for directing proof review. A collection of reviewed Paper Graphs could reveal theorem groups repeatedly used together, combinations rarely observed under comparable opportunities for use, and results connecting different dependency communities. These patterns concern actual uses in selected arguments, rather than bibliography-level co-citation alone. Their interpretation must account for theorem frequency, proof size, subject area, and corpus coverage.

The central proof-review question is therefore:

> Given a human-written or AI-generated proof, can reviewed, source-grounded dependency records and theorem co-use patterns jointly improve the prioritization of premises, intermediate claims, and theorem applications that require mathematical checking?

The proposed input consists of the target statement and proof, its local dependency annotations with unresolved steps retained, and a separate corpus of reviewed proof units. Comparable prior arguments could suggest premise bundles to inspect; bridge nodes could extend the search across subject communities. The output would be a ranked set of specific checks: an application whose hypotheses have not been established, an intermediate claim whose recorded inputs leave its justification unresolved, or a potentially omitted premise worth investigating. Each recommendation would identify the local passage, the relevant theorem instance and conditions, and supporting examples from the corpus.

The empirical goal is to improve the detection and resolution of actual proof defects under a fixed review budget. Comparisons should include unaided review, statement-similarity retrieval, and suggestions based only on theorem frequency, with the same source access and candidate availability. Outcomes include independently adjudicated defects found and resolved, false alarms, and reviewer effort. A common pattern does not certify a proof, and an unusual pattern need not be incorrect. Suggested premises must be checked against the statement, substitution, and hypotheses; a proposed repair remains separate from the original source graph until a revised proof is justified. Section 7

specifies the prediction task, and Section 10 describes how to test whether these suggestions improve mathematical review.

A first practical application is as a companion to reading a paper. A reader could select a result, inspect all its prerequisites, distinguish familiar inputs from unfamiliar ones, and return from a graph item to the source. Whether this improves understanding requires evaluation; logical dependency order need not be the best teaching order.

Verified theorem applications would connect individual Paper Graphs into a larger Math Graph. Expanding external results could eventually support depth relative to specified axioms and inference rules, provided the necessary foundations, definitions, versions, and proof choices have been connected explicitly. Otherwise depth remains relative to the declared boundary of the current analysis.

The data also permit hypotheses about recurring argument shapes in single-author and collaborative work. A collaboration may combine distinct approaches or reflect the person organizing the exposition. The mathematical problem itself may impose much of the shape. Comparisons must therefore separate, where possible, author-associated organization, writing roles, subject matter, and problem structure. They concern recorded arguments, not personal character or mathematical ability.

Sections 6–7 specify these proposed applications, and Section 10 gives their evaluation design. No corpus-level patterns or benefits are reported as findings.

To help readers try the framework on a small selection of papers and obtain a visual overview, we provide two complementary prompts for practical convenience. The PGPP 0.1 reference prompt, introduced in Section 9 and reproduced in Appendix A, guides the construction of source-linked argument records. The supplementary website prompt in Section 9.6 uses the resulting graph data to build an interactive reading view, with complete prerequisite highlighting and persistent click selection. Readers can first apply the analysis prompt to their selected manuscripts and then supply the graph records to the interface prompt to explore the results visually. This provides a practical route from a few source papers to an inspectable map; rapid visualization does not itself establish completeness or correctness, and annotation status and unresolved items should remain visible.

Disclosure of AI assistance

The original research ideas and conceptual direction of this paper originated with the author. GPT-6 Astra was used in developing the mathematical details and proofs that make these ideas precise. GPT-5.6 Sol was used extensively in drafting and revising the manuscript, improving the exposition, and correcting grammar. Responsibility for the final content, including the mathematical statements, proofs, and accuracy of the references, rests with the author.

2 Representation and Data Model

2.1 Source occurrences and contextual states

A source occurrence is a particular item in a particular manuscript version. The inventory includes labeled definitions, assumptions, lemmas, propositions, theorems, corollaries, and numbered equations. Unnumbered equations, local assumptions, constructions, calculations, and prose claims are included when they play a role in the selected argument. Each occurrence retains its text or expression, item type, label, and source location. Document order is metadata rather than a dependency rule.

A computational state is a claim or expression in a context. A state v records a context Γ_v, a claim or expression φ_v, object identifiers, parameter substitutions, and relevant conventions. Identical printed formulas in different contexts need not represent the same state. Conversely, exact reprints of the same calculation may be distinct source occurrences pointing to one state.

The correspondence between occurrences and states is a relation, not necessarily a bijection. One displayed equation may contain several subcalculations, and one state may be displayed more than once. A component of a displayed tuple or family is not identical to the entire family. Its correspondence must be recorded separately, or a component-extraction step must be introduced. This distinction preserves the source while making any sharing decision auditable.

2.2 Actual use and inference records

An edge $u \to v$ means that u is used as an input in the selected justification, definition, or construction of v. It does not assert that every possible proof of v must use u. Citations, semantic similarities, equivalences, and layout connections have separate relation types.

Pairwise edges alone do not record the difference between joint premises and alternative proofs. We therefore retain inference records

$$r = (I_r, o_r, \tau_r, \theta_r, \varepsilon_r, w_r). \tag{2.1}$$

Here I_r is a finite set of joint input states, o_r is the output, τ_r is an operation type, θ_r specifies context and substitution, ε_r contains source evidence and review information, and w_r identifies a proof alternative. Inputs within one record have an AND interpretation. Alternative derivations of the same output have an OR interpretation and are not combined into one set of required premises.

For measurement, a proof selection W chooses one compatible generating record for each derived state in the analyzed argument; accepted boundary states have no generating record. Compatibility requires agreement on shared states, contexts, and any local choices used by the selected records. Incompatible derivations must be kept as separate proof instances. The selected records induce a simple directed graph by placing $u \to o_r$ for each $u \in I_r$. The records themselves are retained, since this projection does not recover their full semantics.

Repeated appearances of one input in a single record can be stored as additional metadata, but the basic graph contains each ordered pair only once. The depth of this state graph is not the depth of an incidence graph with extra operation vertices. Switching between those representations requires an explicit change of measurement policy.

2.3 Boundaries and acyclicity

The boundary B specifies the assumptions, primitives, external results, or unresolved inputs accepted as starting points. Giving an external theorem depth zero means that its proof is outside the current expansion, not that the theorem is elementary. A block whose inputs and output are known but whose interior has not been decomposed remains an explicit opaque block.

We restrict the measurement graph to a finite DAG. If a cycle appears, one must first check whether assumptions and conclusions were conflated, equivalence links were treated as dependencies, or different proof alternatives or manuscript versions were mixed. Induction is represented through a rule with its local assumptions and their discharge, rather than by a self-dependent theorem node. Recursive definitions similarly require their justification conditions. An unresolved cycle is reported as a representation problem; edges are not silently removed to obtain a numerical answer.

3 A Source-Grounded Measurement Protocol

3.1 Fixing the measurement contract

The contract $\mathcal{C} = (M, B, P, W, H)$ fixes the conditions under which graphs and measurements can be compared. Depth and bulk depend on the first four components; the macro skeleton additionally depends on H. The manuscript, annotation set, operation registry, and calculation code should have version identifiers and content hashes.

The granularity policy P specifies what counts as a substantive step. It does not claim to identify a universal smallest unit of informal mathematics. Its purpose is to make a choice repeatable within a stated domain and analytical objective. A domain-specific extension receives a new version. Measurements obtained under different registries should not be treated as directly comparable without a correspondence analysis.

3.2 Substantive decomposition

A step boundary is justified when an output is used by another inference, requires a separate justification, or is identified by a registered operation rule. Independent calculations remain parallel. A sequential edge is inserted only when one step consumes the output of another. Line breaks and equation numbers alone do not define computational steps. Table 1 gives a starting registry.

Table 1: An illustrative operation registry. Each entry must be specialized by its allowed inputs, outputs, and context.

Operation	One recorded step	Distinction to preserve
SUB	A specified simultaneous substitution	Do not count one step per substituted symbol.
APPLY	An application of a theorem or identity	Fix the instance and its hypotheses; mere mention is insufficient.
EXPAND	Expansion under a registered rule	Printed length is not the number of steps.
COLLECT	Collection or cancellation in a fixed normal form	Do not count the same normalization twice.
COMPONENT	Extraction of a used component or coefficient	A component is not the entire expression family.
CONSTRUCT	Construction of an object or map	Retain its defining conditions and source evidence.
COMPAT	A separately justified compatibility condition	A declaration is not an established condition.
COMBINE	An inference combining established outputs	Renaming or reprinting is not a new combination.
OPAQUE	A block with recorded inputs and output	The unresolved interior is not reconstructed by assumption.

Each decomposed step records its location, inputs, output, reason for separation, and evidence status. Explicit source calculations are distinguished from interpretive decompositions grounded in the source. A change of notation can be substantive when it invokes an identity or condition. A long expression can remain one step when the registry treats its simultaneous transformation as one operation.

3.3 Extraction, review, and shared calculations

The workflow proceeds from a source inventory to contextual states, candidate dependencies, internal decomposition, independent review, and an accepted measurement graph. An LLM can assist with the inventory, local interpretation, and decomposition, but its output is a proposal with explicit source evidence, not independent acceptance. The inventory of numbered items can be checked against the manuscript. Its completeness does not establish the completeness of unnumbered claims or implicit dependencies.

Evidence status and review status are separate. Evidence may record explicit use, a locally justified interpretation, or an unresolved candidate. Review status may be proposed, accepted, rejected, or unresolved. The presence of a source location does not establish that the cited passage supports the edge. Agreement among extraction systems likewise does not replace a semantic assessment.

Sharing is allowed only when the context, objects, substitutions, inputs, and calculation instance agree. Exact reprints may share a state while remaining distinct document occurrences. Equivalent propositions proved in different ways are not merged merely because they have the same truth conditions. If a proposed identification creates self-dependence or a cycle, the identification must be reconsidered.

4 Depth, Bulk, and Dependency per Node

4.1 Ancestor subgraphs and basic quantities

Let $G = (V, E)$ be the selected finite DAG, with edges directed from inputs to outputs. Write $\mathrm{Pred}_G(v)$ for the direct predecessors of v, and let $u \rightsquigarrow_G v$ denote a nonempty directed path. Define

$$A_G(v) = \{v\} \cup \{u \in V : u \rightsquigarrow_G v\}, \tag{4.1}$$

$$F_G(v) = E \cap (A_G(v) \times A_G(v)). \tag{4.2}$$

Thus $F_G(v)$ is the edge set of the subgraph induced by all ancestors of v, including v itself.

Definition 4.1 (Depth and bulk)**.** The depth of v is the length, in edges, of a longest directed path ending at v. Its bulk is the number of distinct edges in its ancestor subgraph:

$$D_G(v) = \begin{cases} 0, & \mathrm{Pred}_G(v) = \varnothing, \\ 1 + \max\limits_{u \in \mathrm{Pred}_G(v)} D_G(u), & \mathrm{Pred}_G(v) \neq \varnothing, \end{cases} \tag{4.3}$$

$$S_G(v) = |F_G(v)|. \tag{4.4}$$

When comparison conditions matter, these quantities can be written $D_{M,B,P,W}(v)$ and $S_{M,B,P,W}(v)$. We suppress the contract only when it is fixed. Depth describes sequential dependence; bulk describes the amount of distinct recorded dependency structure supporting a target. Bulk is neither the number of inference records nor the number of arithmetic operations executed by a machine.

A four-edge chain, a branching argument with a shared input, and a conclusion with four independent inputs all have bulk four but different depths; see Figure 1. These are abstract illustrations of the definitions, not measurements extracted from particular papers.

4.2 Union calculation and shared structure

The ancestor edge sets satisfy

$$F_G(v) = \bigcup_{u \in \mathrm{Pred}_G(v)} (F_G(u) \cup \{(u, v)\}). \tag{4.5}$$

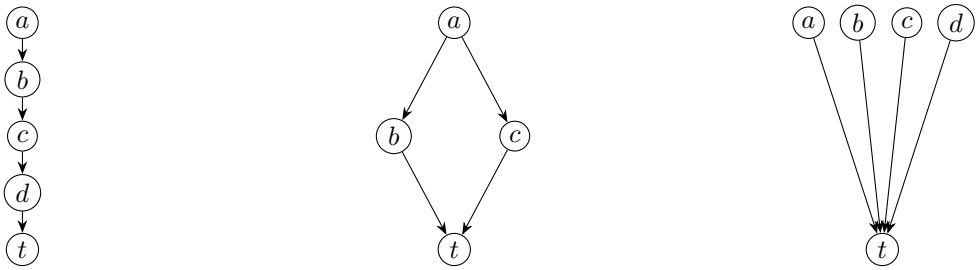

(a) Chain
$D(t) = 4, \quad S(t) = 4$

(b) Shared input
$D(t) = 2, \quad S(t) = 4$

(c) Independent inputs
$D(t) = 1, \quad S(t) = 4$

Figure 1: Equal bulk does not determine depth. The arrows record selected input use, and the displayed measurements concern the final state t.

An empty union is empty. Computing unions, rather than adding predecessor bulks, prevents a shared subargument from being counted repeatedly.

Proposition 4.2 (Correctness of the union recurrence)**.** *Processing a finite DAG in topological order using* (4.3) *and* (4.5) *computes the longest incoming path length and the distinct ancestor edge set of every vertex.*

Proof. At a source, there are no incoming edges and both statements hold. Every nonempty path ending at v has a final edge (u, v) with $u \in \mathrm{Pred}_G(v)$, giving the recurrence for depth. An edge in the ancestor subgraph of v either ends at v or lies on a path to one of its direct predecessors, so it belongs to the right-hand side of (4.5). Conversely, every edge on that right-hand side lies in the ancestor subgraph of v. Induction in topological order proves the claims. Set union retains each shared edge only once. □

For several targets $T \subseteq V$, the distinct supporting edges are $F_G(T) = \bigcup_{t \in T} F_G(t)$. In general, $|F_G(T)| \neq \sum_{t \in T} S_G(t)$. The difference between the latter sum and the union size counts repeated edge incidences across target ancestor subgraphs. It is a description of overlap, not a new inference count.

4.3 Reachability and changes of granularity

Proposition 4.3 (Depth is determined by fixed-vertex reachability)**.** *Two finite DAGs on the same vertex set with the same reachability relation have the same depth at every vertex.*

Proof. Reachability defines a strict partial order. Every directed path ending at v is a chain in that order. Conversely, take a longest chain ending at v. Consecutive elements of the chain are related by a cover relation; otherwise an intermediate element could extend the chain. A cover relation must be represented by a direct edge in either DAG, since a longer realizing path would supply an intermediate element. Thus a longest chain is a path in each graph, and its length is determined by reachability alone. □

Transitive reduction therefore preserves depth on a fixed vertex set, but it need not preserve bulk. Adding $a \to c$ to $a \to b \to c$ leaves $D(c) = 2$ and changes $S(c)$ from two to three. If a is actually used directly in the selected justification of c, the additional edge belongs in the basic measurement graph even if a simplified visualization omits it.

Proposition 4.4 (Failure of invariance under arbitrary refinement)**.** *Depth and bulk are not invariant under arbitrary subdivision of internal argument steps, even when reachability between the original boundary vertices is unchanged.*

Proof. Replace the single edge $a \to b$ by $a \to x_1 \to \cdots \to x_k \to b$. Both the depth and bulk of b change from one to $k+1$, while a still reaches b. The vertex set has changed, so Proposition 4.3 does not apply. □

This elementary example explains why a granularity policy is necessary. It does not say that every refinement changes a graph's maximum depth: subdividing a noncritical branch can leave the longest path unchanged. An increase in recorded steps and an increase in depth are distinct events.

4.4 Conditional bounds for unresolved dependencies

Proposition 4.5 (Monotonicity on a fixed vertex set)**.** *Suppose $G^- = (V, E^-)$, $G = (V, E)$, and $G^+ = (V, E^+)$ satisfy $E^- \subseteq E \subseteq E^+$, and G^+ is a DAG. Then, for every v,*

$$\begin{aligned} D_{G^-}(v) \le D_G(v) \le D_{G^+}(v), \\ S_{G^-}(v) \le S_G(v) \le S_{G^+}(v). \end{aligned} \tag{4.6}$$

Proof. Adding edges removes neither existing paths nor existing ancestors and edges in an ancestor subgraph. Hence neither quantity decreases. □

The proposition provides bounds only if the inclusion assumptions hold for the intended graph. It does not turn the minimum and maximum values returned by several extraction systems into a confidence interval. An upper graph containing mutually exclusive alternatives may give a numerical upper bound, if acyclic, without representing any realizable proof. Changes in state correspondence or granularity require a separate analysis and are not covered by (4.6).

4.5 Dependency per node and structural interconnection

For a nonempty graph, define

$$q(G) = \frac{|E|}{|V|}. \tag{4.7}$$

This is both the average indegree and the average outdegree. It is a natural operational interpretation of "dependency per node" and a first summary of how many direct relations the representation records per state. We call it the *dependency-per-node ratio*. If the informal phrase "connectivity complexity" is used for this quantity, it should refer only to this specified ratio, not to a general measure of mathematical complexity or connectedness.

For example, a chain and a star on the same number n of vertices both have $q = (n-1)/n$. The ratio does not distinguish their sequential structure. It is also different from density normalized by the maximum number of edges in a simple DAG, $2|E|/(|V|(|V|-1))$, when $|V| > 1$. Both quantities depend on which vertices and relation types the measurement includes.

For a target, put $N(v) = |A_G(v)|$. When $N(v) > 1$, two complementary descriptors are

$$\rho(v) = \frac{S(v)}{N(v)-1}, \qquad \mu(v) = S(v) - N(v) + 1. \tag{4.8}$$

The underlying undirected ancestor graph is connected, because every ancestor reaches v. Hence $S(v) \ge N(v) - 1$, and $\mu(v)$ is its undirected cycle rank. Branches may diverge and reconverge in a

DAG, producing an undirected cycle without a directed logical cycle. For an isolated input, $\mu = 0$ and ρ is left undefined. These descriptors complement depth and bulk; none determines the entire shape of an argument.

5 Macro Skeletons and Argument Shape

5.1 Projection onto fixed anchors

Fine-grained measurements depend on decomposition. For a complementary view of a paper's organization, fix an anchor set $H \subseteq V$ before comparison. Anchors may include principal results, named constructions, definitions, and selected external inputs. Exact reprints pointing to one state have a single representative in H. The anchor policy is part of the measurement contract, rather than a choice made after seeing which graph looks preferable.

Definition 5.1 (Macro projection and skeleton)**.** The macro projection $M_H(G)$ has vertex set H. Distinct anchors a, b are joined by $a \to b$ if G has a directed path from a to b whose internal vertices lie outside H. The macro skeleton is

$$K_H(G) = \mathrm{TR}(M_H(G)). \tag{5.1}$$

Proposition 5.2 (Preservation of anchor reachability)**.** *The projection $M_H(G)$ preserves exactly the reachability relation of G restricted to H. Consequently, two finite DAGs with identified anchor sets and the same reachability relation between those anchors have the same macro skeleton under that identification.*

Proof. A path between anchors in G can be split at every intermediate anchor. Each resulting segment gives an edge of $M_H(G)$, hence a path in the projection. Conversely, replacing every projection edge by one of its supporting paths gives reachability in G. The projection is acyclic. A finite DAG has a unique transitive reduction, namely the cover graph of the partial order defined by reachability [Theorem 1, 11]. The conclusion follows. □

The result gives a precise condition under which different internal decompositions yield the same macro skeleton. It does not guarantee that independent annotators recover the same anchor relations. Branches that do not contribute to a connection between anchors may disappear from the projection, and internal differences can be hidden. The full graph and the supporting paths must therefore remain available.

5.2 Structural signatures and visual layout

The shape of an argument is defined by its graph structure, not by screen coordinates. A structural signature may include height, the distribution of vertices across depth levels, maximum antichain size, weak component count, the number of terminal results, branching and merging frequencies, and overlap between target ancestor subgraphs. Normalized profiles should be reported together with their original vertex and edge counts.

Vertices at the same depth form an antichain, but the largest depth level need not be a maximum antichain. Nor does agreement on a finite collection of summary statistics imply graph isomorphism. Signatures serve as descriptors for exploration and statistical comparison; important differences should be checked against anchor reachability and the source.

A reading interface may use a compact node-and-edge view with details revealed on selection. Selecting a target should highlight its ancestors and the induced dependency edges, while preserving

access to source locations and evidence. Layout layers are presentation choices, not necessarily the numerical depth. Reproducing a visual arrangement also requires fixing the layout algorithm and its tie-breaking rules.

6 From Paper Graphs to a Math Graph

6.1 Connecting theorem applications across papers

A Math Graph connects identified states and inference instances across manuscripts. A paper-level citation is a starting point for locating a result, but an accepted application additionally specifies the theorem's version and statement, its hypotheses, the substitution used, and the relevant conventions. A theorem with the same name may have a corrected statement, stronger hypotheses, or a different specialization in another version.

Global identifiers should distinguish the manuscript version, the statement or state, and the proof selection. An importing paper records an application instance and the evidence that its hypotheses are met. Two applications may share a source theorem without being the same computational step. When moving between mathematical settings requires a translation, that translation and its justification are recorded as part of the connection. Similarity links can help locate candidates but remain distinct from accepted uses.

Circularity in the document citation network does not imply circularity in a selected theorem proof. Mutual references, revised versions, and alternative arguments must be separated before combining measurement graphs. Acyclicity is checked again after integration. An external dependency that has not been resolved remains at the declared boundary.

6.2 Axiom-relative and boundary-relative depth

Expanding external results can, in principle, connect selected arguments to specified axioms. Axiom-relative depth nevertheless requires a proof system, inference rules, definition-expansion conventions, library versions, and proof choices. The name of an axiom system alone does not determine a unique graph or numerical depth.

We reserve the description *axiom-relative* for an expansion in which every source is accounted for by specified axioms, primitive definitions, or explicit hypotheses, and no external theorem remains an unexpanded input. Otherwise, the quantity is reported as boundary-relative depth. For an informal manuscript, correspondence with a formal foundational representation is itself an additional claim requiring evidence. A human-oriented substantive step and a kernel-level inference are different units. Their measurements cannot be compared without an explicit translation and normalization policy.

If several proofs of one statement are available, their measurements can be reported separately. A minimum over this collection is only a minimum over the available proofs, not over every possible proof of the statement. Moreover, minimizing the bulk of each target separately need not minimize the bulk of their union: different proof choices can create or eliminate shared subarguments.

7 Applications and Structural Research Questions

7.1 Reading support and prerequisite navigation

The immediate application is a source-linked reading map. A reader selects a result, inspects its prerequisite subgraph, and marks inputs already understood. The remaining graph can expose missing definitions, shared lemmas, and constructions reused by several results. It can also support alternative reading paths with different intermediate goals.

This use does not assume that logical order and pedagogical order coincide. Examples, motivation, and exploratory calculations may be educationally valuable even when they are not inputs to the final proof. They should remain accessible as document material with their own relation types. Learning outcomes must be evaluated through understanding, retention, and transfer, rather than inferred from the appearance of a well-organized graph.

7.2 Single-author and collaborative argument shapes

An author-associated shape is a hypothesis about recurring organization in written arguments. It can be investigated by asking whether structural signatures recur across an author's single-author papers after accounting for field, problem type, length, period, and the measurement policy. The object of study is expository and argumentative organization, not personality or ability.

Collaborative papers introduce further questions. Their shape may reflect interaction between different approaches, a division of mathematical tasks, or a principal organizer of the exposition. It need not be an average of the collaborators' single-author signatures. Where available with appropriate consent, documented contribution and writing roles can be compared with the graph. Author order should not be used as a proxy for those roles.

Problem structure is a competing explanation. Different authors working on closely related questions, different proofs of the same result, and successive versions of one manuscript provide complementary comparisons. If authors and topics almost completely coincide in a dataset, their effects cannot be reliably separated. Repeated collaborative combinations and samples crossing authors with topics are needed before discussing stable collaboration patterns. Observational associations alone do not establish a causal effect of collaboration.

7.3 Co-use, rare combinations, and cross-field connections

Co-citation and actual co-use should be recorded separately. Co-citation in Small's sense concerns documents cited together [4]. For our purpose, let $\mathcal{U}$ be a fixed collection of proof units and let $\mathcal{T}$ be a set of identified theorem statements. Define

$$X_{ut} = \mathbf{1}\{t \text{ is used in proof unit } u\}, \qquad C(t,t') = \sum_{u \in \mathcal{U}} X_{ut} X_{ut'} \quad (t \neq t'). \tag{7.1}$$

The policy must specify whether "used" means a direct input, an input to the selected proof as a whole, or membership in its full ancestor set. These are different analyses. Theorem identity is retained at the source level while distinct application instances remain available in the underlying data.

The matrix C can be used to investigate groups of theorems repeatedly used together. Raw counts are strongly affected by theorem popularity and proof size. A suitable null model should preserve, or otherwise account for, theorem-use frequencies and the numbers of inputs per proof unit. Field, period, and available opportunities for joint use are additional controls. Statistical enrichment is a prompt for mathematical interpretation, not an explanation of why the theorems belong together.

Pairs that are rarely used together can also be investigated, but low co-use does not imply logical incompatibility. A rare theorem, limited corpus coverage, or a lack of relevant problems may explain the pattern. Searching many pairs additionally requires attention to multiplicity and sparse observations.

To study cross-field connections, first fix an external subject classification or a reproducible clustering procedure. They are not the same notion of a field. For a chosen single cluster label $c(v)$, a simple descriptive quantity for a node with positive outdegree is

$$b_{\mathrm{out}}(v) = \frac{|\{w : (v,w) \in E,\ c(w) \neq c(v)\}|}{|\{w : (v,w) \in E\}|}. \tag{7.2}$$

It records the fraction of direct uses going outside the node's cluster. The absolute number and distribution of those uses must also be reported: one external use and a hundred external uses can have the same fraction. Multiple subject labels require a separate weighting rule. Paths through the graph provide additional information about bridging roles, but centrality can reflect common infrastructure, as emphasized in the Mathlib analysis [Appendix C.2.5, 10]. Corpus coverage is an additional control in comparisons across papers. Centrality is not a direct measure of mathematical importance.

7.4 Probabilistic premise suggestions and proof review

The main corpus-level application proposed here is to use recorded dependency patterns to guide mathematical scrutiny. Learned premise selection and proof re-verification already provide precedents [15]; our question is how reviewed, source-grounded argument records can support this task for informal manuscripts as well. The prediction target must be specified before assigning probabilities. For example, let J denote the set of theorem inputs recorded in a proof unit sampled under a declared corpus and proof-selection policy, let φ be its target, let Γ be its context, and let $K \subseteq J$ be the theorem inputs already observed and accepted as recorded uses. For a nonempty finite candidate set $T \subseteq \mathcal{T} \setminus K$, a future model could estimate

$$\widehat{p}_{\mathcal{D}}(T \mid \varphi, \Gamma, K) = \widehat{\mathbb{P}}_{\mathcal{D}}(T \subseteq J \mid \varphi, \Gamma, K). \tag{7.3}$$

Here $\mathcal{D}$ denotes the reviewed training corpus together with its identity, availability, and use-label policies. This expression defines a possible prediction task, not an estimator implemented or evaluated in this paper. It predicts recorded joint use under that policy, not logical necessity, applicability, or the truth of the target statement. Different valid proof alternatives can have different input sets. The mechanism by which K is observed or masked is part of the prediction task and must be fixed when calibration is compared. Merely proposed or unchecked premises are not assumed to belong to J.

Co-use groups can suggest premise bundles that repeatedly complement known inputs. Bridge nodes can extend retrieval across subject communities, while rare configurations can identify cases for closer inspection. Pairwise counts alone do not determine a joint probability for a larger bundle, and missing annotations must not automatically be treated as negative use labels. Low frequency can reflect a legitimate unusual argument or limited corpus coverage. It is not itself a defect.

For both human-written and AI-generated proofs, a review interface could associate each suggested check with supporting proof units and their source passages. A reviewer then examines the proposed theorem's exact statement and version, the substitution, all hypotheses, and the local transition it is meant to justify. This can organize checks for omitted premises, unproved intermediate claims, and invalid applications. A candidate repair absent from the original source remains separate

from the reconstructed source graph; it enters only an explicitly identified revised proof with new evidence. Consequently, suggestions do not alter the original argument's depth or bulk merely by being recommended.

The intended benefit is more effective proof review, to be tested by defect-detection and repair outcomes rather than by familiar-looking graphs. Corpus agreement, LLM agreement, and DAG validation do not establish mathematical correctness. Formal checking can supply an additional validation route when available; correspondence between the formalized statement and the source argument must still be assessed. This application extends the research program beyond the discovery, recomputation, and linking modes specified by PGPP 0.1.

8 Algorithms and Reproducible Execution

8.1 Input checks and deterministic calculation

Before measurement, check unique state identifiers, edge endpoints, relation types, the proof selection, and the declared boundary. Normalize only accepted instances of the same calculation. Several evidence records supporting the same ordered pair produce one measurement edge. A self-loop after normalization is an error to investigate. The actual sources of the graph must agree with its declared accepted inputs and primitives.

The following procedure operates on an accepted graph. Source interpretation belongs to the process that constructs its input.

Algorithm 1. Contract-relative graph measurement.

1. Check versions of the manuscript, contract, annotations, and code.
2. Normalize accepted shared calculations and construct the selected set of dependency edges.
3. Compute a topological order with fixed identifier-based tie-breaking. Report an unresolved cycle rather than returning DAG measurements.
4. In that order, compute $D(v)$ and $F(v)$ from (4.3) and (4.5); store $S(v) = |F(v)|$.
5. Propagate declared unresolved-state flags to dependent states.
6. Construct $M_H(G)$ and $K_H(G)$ and check preservation of anchor reachability.
7. Map measurements back to source occurrences and store input hashes, output hashes, and the calculation checks.

Unresolved status can occur at internal states, not only at boundary inputs. If U is the set of states carrying such a flag, define

$$I(v) = \mathbf{1}\{A_G(v) \cap U \neq \varnothing\}. \tag{8.1}$$

This propagates the effect of declared unresolved material. The value $I(v) = 0$ does not establish that the source is complete or correct; it only records the absence of currently declared unresolved states in the ancestor subgraph. Opaque decomposition and an intentionally unexpanded external boundary should also have separate labels, so that missing knowledge and the chosen analytical scope are not conflated.

8.2 Complexity and independent recalculation

Let $n = |V|$ and $m = |E|$. Given a topological order, depth can be calculated in $O(n + m)$ time. If hash-set insertion and lookup have expected constant time, the explicit ancestor-edge unions require expected time

$$O\left(n + m + \sum_{(u,v)\in E} |F_G(u)|\right). \tag{8.2}$$

This is at most $O(n + m + m^2)$, while storing all such sets requires $O(n + m + \sum_{v\in V} |F_G(v)|)$ space. Identifier sorting and deterministic tie-breaking have their own implementation-dependent overhead. Using a topological order does not make the entire bulk calculation linear. Bitsets, incremental updates, or persistent shared set representations are possible alternatives for larger graphs.

If there are h anchors, searching from each anchor until another anchor is encountered gives a simple macro projection in worst-case $O(h(n + m))$ time. Reachability and transitive reduction are then computed on the projection. Practical reports should specify the algorithms used and measure memory as well as elapsed computation time. These bounds concern the graph calculation, not the prior interpretation of a manuscript.

Independent checks should follow the definitions through a different calculation route where feasible. Bulk can be checked by first finding the ancestor vertices and then directly counting induced edges. A macro skeleton can be checked for reachability preservation and for the necessity of each retained edge. Such tests establish properties of the recorded graph; they do not establish that its mathematical interpretation is faithful to the source.

8.3 Three levels of reproducibility

Computational reproducibility means obtaining the same measurements from the same accepted graph. *Annotation reproducibility* concerns agreement on actual uses when the inventory and decomposition are fixed. *Discovery reproducibility* concerns agreement when annotators start from a new manuscript and identify both items and internal steps. Deterministic calculation addresses the first level. The latter two require independently annotated reference material and repeated studies.

Execution records should include the actual tool identifiers, settings, provided material, source scope, review decisions, and change history. Unavailable model versions or random seeds must not be invented. Re-extraction from a manuscript is distinguished from recalculation on accepted data. A newly discovered dependency creates an evidenced revision, not an undocumented change to an earlier result.

9 A Standardized Prompt Protocol

9.1 Purpose and normative status

To make the framework reusable, we specify the *Paper Graph Prompt Protocol*, PGPP 0.1. The complete reference prompt appears in Appendix A; it can be copied with a source bundle and a run configuration. It is an operational specification of the paper's method, not merely a request to draw a dependency graph. Its requirements cover source identification, decomposition, evidence, review provenance, graph selection, numerical calculation, and the structure of the returned data.

PGPP 0.1 is a proposed research protocol, not an externally adopted standard. Within the protocol, `MUST` identifies a required condition, `SHOULD` allows a documented deviation, and `MAY` denotes an optional feature. Conformance concerns adherence to these declared requirements. It does not

establish the truth of an extracted mathematical claim or guarantee agreement among independent annotators.

9.2 Three operating modes

The protocol separates three tasks that should not be conflated in a reproducibility claim. In `DISCOVER`, a system constructs proposed annotations from supplied manuscripts. It resolves its provisional contract explicitly, and may calculate provisional measurements when the structural prerequisites are satisfied. New annotations remain proposed even when the system reports high confidence.

In `RECOMPUTE`, the input is an existing frozen contract and graph snapshot with supplied review decisions. The system recalculates its quantities without adding states, changing identities, selecting a new proof, or reinterpreting a source passage. A required semantic repair is reported as a separate issue. It cannot be hidden inside a supposedly identical rerun.

In `LINK`, the system proposes theorem applications between supplied Paper Graph snapshots. Each proposal identifies its source statement, version, substitution, and hypothesis checks. The proposal does not silently modify an accepted graph. After review, a combined snapshot may be supplied to `RECOMPUTE`. This separates the discovery of a possible cross-paper connection from its use in a reported measurement.

9.3 Input and output contracts

The operator supplies the source bundle, scope, boundary policy, granularity extensions, targets, anchors, and the operating mode. Existing snapshots and review records are additionally supplied when applicable. For a first discovery run, the reference profile defaults to the main mathematical text and all appendices, excluding the Introduction, and keeps external results opaque. These are editable scope choices, not universal assumptions about mathematical writing.

Defaults for targets and anchors must be expanded into explicit ID lists before calculation. Null means unspecified; an empty list means an intentional empty selection. Missing anchors prevent a macro comparison, but do not by themselves prevent depth and bulk calculations under a complete (M, B, P, W) contract. Missing or unknown premises must not be silently reclassified as accepted sources of depth zero.

The response is one JSON object with a fixed record structure. Source occurrences, contextual states, occurrence/state correspondences, and inference records remain separate. The selected inference IDs determine the measurement edges, preventing an independently edited edge list from disagreeing with the selected joint premises. Other relation types, sharing decisions, imported reviews, unresolved issues, measurements, conformance checks, and changes have their own fields.

The appendix gives the required keys, record fields, and enumerated statuses. This is a normative field contract, not a claim that a complete machine-enforced JSON Schema or reference validator has already been implemented. Structural validation and source interpretation remain distinct tasks. Identifiers are stable within a versioned snapshot; fresh semantic extraction is not assumed to recreate the same identifiers or decomposition automatically.

9.4 Evidence and calculation discipline

PGPP 0.1 distinguishes an annotation's evidence status from its review status. Acceptance is imported from supplied review data with provenance; the extractor cannot certify its own new annotations as independently accepted. A graph can therefore have well-defined provisional numerical

properties while its interpretation remains under review. Numerical consistency does not promote that graph to accepted status.

The prompt fixes the core operation rules and requires the exact rule and expression bundle for a domain-specific expansion or normalization. If the source does not justify a decomposition, the system reports an opaque block or an unresolved issue. The prompt does not authorize inventing intermediate arguments merely to eliminate uncertainty.

Numerical output requires actual graph calculation. Without execution, the protocol returns unavailable numerical fields and records the checks as not run. A cycle requires a witness and blocks the affected DAG measurement. Hashes, tool versions, and execution settings must be observed values; unavailable metadata are recorded as null with an explanation. Source coverage, structural checks, numerical checks, and semantic review are reported separately.

9.5 Reference conformance cases

Small abstract graphs make several requirements independently testable. Table 2 lists reference cases. They define expected behavior for a future protocol implementation and do not evaluate extraction from mathematical manuscripts.

Table 2: Reference cases for PGPP 0.1. All edges point from input to output.

Case	Specified input	Required behavior
Shared prefix	$r \to a$, $a \to b$, $a \to c$, $b \to t$, $c \to t$	$D(t) = 3$, $S(t) = 5$; the prefix edge is counted once.
Alternative proofs	Records $\{a, b\} \Rightarrow t$ and $\{c\} \Rightarrow t$; select the second, with c a source	$D(t) = 1$, $S(t) = 1$; do not add the first record's inputs.
Internal uncertainty	$a \to b \to t$ and $U = \{b\}$	$I(b) = I(t) = 1$, although b is not a source.
Directed cycle	$a \to b \to a$	Return the cycle witness and block DAG metrics; do not delete an edge.
Anchor refinement	Replace $a \to b$ by $a \to x \to b$; $H = \{a, b\}$	Both skeletons contain $a \to b$, while fine-grained depth and bulk change.

The shared-prefix case is deliberately stronger than a diamond with no shared incoming edge: simply adding branch bulks would return six instead of five. The alternative-proof case checks the distinction between an AND input set and an OR choice of records. These examples test the interpretation of the protocol's data contract, not the correctness of an arbitrary mathematical argument.

9.6 Supplementary prompt for the interactive website

For readers' convenience, we provide the following English prompt to help them explore the framework visually with a few papers of their choice. After generating source-linked graph records with PGPP 0.1, readers can supply those records with this prompt to construct an interactive reading website. The intended use is a quick visual overview followed by inspection of the prerequisites supporting a selected result. This supplementary prompt concerns presentation and interaction; it does not change PGPP 0.1's annotation, proof-selection, or measurement rules or constitute a separate extraction method.

For a node v, the highlighted nodes and edges are exactly $A_G(v)$ and $F_G(v)$ from (4.1)–(4.2), with G denoting the dependency graph supplied to the interface. Thus the highlight covers the union of all prerequisite paths, including shared branches, rather than one selected path or only the immediate neighborhood. Equivalence, similarity, citation-only, and layout relations do not participate in this traversal.

The prompt can be supplied with the graph data when constructing the website. Hover selection is temporary; a clicked selection takes precedence until explicitly replaced or released.

```
INTERACTIVE WEBSITE PROMPT: ANCESTOR HIGHLIGHTING AND CLICK-TO-PIN

Build an interactive dependency-map interface for the supplied graph.
If several Paper Graphs are supplied, allow readers to switch between
them. Use cross-paper dependency links only when explicitly supplied.
Keep supplied review status and unresolved annotations accessible.
An edge A -> B means that A is used in the definition, construction,
or selected justification of B. Use this direction when identifying
prerequisites, regardless of whether arrowheads are drawn.

ALL ANCESTOR PATHS
When the pointer hovers over a node v and no node is pinned, make v,
every direct and indirect prerequisite of v, and all dependency edges
between those nodes glow brightly. Compute the full ancestor closure
by following incoming dependency edges. Highlight the union of all
prerequisite paths, not just immediate neighbors or one chosen path.
Shared prerequisites and shared edges are highlighted only once.
Do not highlight descendants: results subsequently obtained using v
are not part of v's prerequisite highlight. Keep unrelated elements
visually subdued. Ignore auxiliary equivalence, similarity, citation-only,
and layout links when computing the highlight.

PERSISTENT CLICK SELECTION
Clicking a node pins that node and its complete ancestor highlight.
Keep the pinned highlight visible when the pointer leaves the node,
moves over the background, or hovers over another node or an edge.
While a node is pinned, hover events must not replace, extend, or clear
the pinned highlight. Clicking a different node replaces the pinned
target and highlights the new target's complete ancestor subgraph.
Clicking the already pinned node leaves it pinned.

Provide a clearly visible Close or Unpin button while a node is pinned.
Activating that button or pressing Esc clears the pinned target and
its persistent highlight. Resume temporary highlighting on the next
node-hover interaction. When no node is pinned, leaving a hovered node
clears its temporary highlight. Keep hover state and pinned state
separate, with the pinned target taking precedence.

DATA INTEGRITY AND CHECKS
Use the supplied graph; do not infer new dependencies during interaction.
Highlighting and pinning must not modify the graph or its D and S values.
Check a graph with r -> a, a -> b, a -> c, b -> v, c -> v, and v -> w.
Hovering over v must highlight r, a, b, c, v and all five prerequisite
edges, but neither w nor the edge v -> w. After clicking v, hovering
over w or moving off the graph must leave that same highlight pinned.
Clicking b must switch the highlight to r, a, b and their two edges.
Verify that both the Close/Unpin button and Esc release the selection.
Report any interaction checks that have not actually been performed.
```

These checks specify expected interface behavior. They do not constitute reported usability results or validation of the graph's mathematical annotations.

9.7 Versioning and empirical assessment

A run records the protocol version, resolved contract, exact supplied material, and any extension or deviation. Changing the meaning of a step, an edge, a field, or a default requires an explicit protocol or profile revision. Cosmetic prompt changes should also be retained in the run record, so that "the same prompt" has a concrete meaning.

The effect of standardization is an empirical question. A study can compare an unrestricted graph request, the complete reference prompt, and controlled variants omitting selected requirements. It should evaluate valid-output rate, structural conformance, source-grounded annotation agreement, metric error, and unresolved-case reporting separately. Identical wording alone does not guarantee reproducible semantic extraction. The protocol instead makes the decisions and failure conditions visible enough to compare such runs.

10 Evaluation Design

10.1 Reference data and extraction quality

A useful benchmark should include several mathematical fields and argument types: calculation-heavy proofs, conceptual arguments, texts with substantial omissions, and papers depending on many external results. Material used to develop the registry should be separated from the final evaluation. Closely related manuscript versions and proof families should remain in the same partition to avoid misleadingly easy generalization tests.

At least two suitably qualified annotators should work independently on the reference material. Disagreements are adjudicated with source evidence, and genuinely unresolved relations remain marked as such. Sample sizes and target precision should be specified after a pilot establishes the variation in annotation and structural measurements. A reference graph is a documented expert judgment, not an assumption that informal mathematics has an unambiguous canonical decomposition.

Comparison conditions should separate extraction from raw text, annotation with a fixed item inventory, and annotation with both an inventory and a fixed decomposition. Citation- and cross-reference-based graphs provide baselines. Where available, author annotations and the outputs of existing tools can be compared on the same material and relation definitions. Repeated runs under fixed settings and runs after tool changes should be reported separately. Table 3 summarizes the proposed evaluation dimensions.

If state correspondence is uncertain, alignment error must be evaluated before edge agreement. Arbitrary merging can artificially increase the reported agreement. Besides averages, report paper-level distributions, failure cases, and sensitivity to external boundaries and opaque blocks. Stability across tools and agreement with the expert reference are different properties.

Premise suggestions require a separate evaluation. Held-out splits should respect manuscript versions, near-duplicate statements, and proof families, and retrieval should enforce the declared availability boundary. Compare co-use and graph-based suggestions against theorem-frequency and semantic retrieval baselines on the same candidate pool. Report premise precision, recall, ranking quality, and calibration for the specified use outcome; these do not establish improved proof checking. A review-assistance study should separately measure confirmed omissions and invalid applications found, successful repairs, false alarms, and reviewer effort. It should include human-written and AI-generated proofs, with genuine unusual proofs as well as independently adjudicated defects. This

Table 3: Proposed evaluation dimensions. No empirical values are asserted in this paper.

Object	Main quantities	Required accompanying checks
Source inventory	Precision and recall by item type	Numbered and unnumbered items assessed separately
Actual-use edges	Precision, recall, and F1	Supporting passages, contexts, and application hypotheses
Decomposition	Step alignment and boundary agreement	Shared calculations and excessive subdivision
Macro skeleton	Anchor ancestor-set agreement	Height, width, branching, and terminal structure
Depth and bulk	Absolute and relative error distributions	Boundary changes and unresolved components
Reproducibility	Structural and numerical agreement	Recalculation distinguished from new extraction
Learning support	Understanding, retention, and transfer	Prior knowledge, task allocation, and source navigation
Proof review	Premise retrieval, calibration, and defect detection	Held-out proof families, checked applicability, and false alarms

tests whether statistical suggestions help mathematical review without penalizing valid departures from common patterns.

10.2 Learning support

Learning studies can compare ordinary document reading with graph-assisted reading while accounting for prior knowledge and task difficulty. Outcomes may include explaining a target result, identifying its assumptions, applying an idea in a new setting, and delayed retention. Task allocation should control the carryover caused by reading the same argument twice.

The graph should also be evaluated as a navigational representation. Readers may be asked to locate the source of a prerequisite, distinguish an accepted input from an unresolved dependency, or identify a shared construction. Successful navigation is useful evidence about the interface, but does not by itself show improved mathematical understanding. The study should record how graphs were prepared and whether the version presented to the reader contained unresolved annotations.

10.3 Author, collaboration, and reuse hypotheses

Studies of author-associated signatures need matched comparisons across topics and lengths, repeated papers by the same author, and related problems addressed by different authors. Graph annotation can be blinded to author identity. Papers, rather than their many nodes, are the primary units for claims about paper-level organization; dependence within authors, collaborations, and proof families must be modeled or respected in resampling and data splits.

Collaboration analyses should ask whether an observed signature remains stable after accounting for the mathematical problem and, where known, writing roles. Sufficient replication is needed to distinguish a recurring collaborative pattern from one unusual paper. The goal is to test whether an organizing tendency is identifiable, not to assign a permanent graph signature to a person.

Co-use and cross-field analyses should state the proof-unit boundary, theorem-identity policy, clustering resolution, and time window. Null models, multiplicity adjustments, and sensitivity to corpus coverage are needed before interpreting unusual combinations. Candidate connections should then be examined mathematically. A statistical pattern can suggest a question without implying that a new theorem or proof exists.

11 Limitations and Scope

The central limitation is the interpretation of informal mathematics. Detailed source evidence does not recover every omitted argument or settle every disagreement about a step boundary. A more specialized registry may improve consistency while reducing transfer across domains. Its suitability is therefore an empirical question rather than a consequence of deterministic graph algorithms.

The stability of a macro skeleton is conditional on fixed anchors and their reachability. Different anchor choices can yield different skeletons, and identical skeletons can conceal different internal arguments. Even exact agreement of the recorded graphs does not establish agreement of all the mathematics that the manuscript leaves implicit.

Depth and bulk do not measure difficulty, originality, truth, or importance. A short step may invoke a profound idea, while a large subgraph may consist of routine calculations. A different proof may change the graph substantially. These observations are reasons to keep representation choices explicit and to study structural quantities jointly with mathematical context.

Author and collaboration studies are additionally limited by confounding, corpus selection, and the availability of reliable role information. Observed writing patterns should not be interpreted as personality traits or individual merit. Co-use and bridge statistics are likewise properties of a specified corpus and relation policy, not of all mathematics.

Source access, the distribution of derived annotations, and the handling of unpublished manuscripts require appropriate permissions. Version changes and corrections must remain traceable. In particular, an accepted application of a theorem is tied to its statement, assumptions, and context; it is not an unrestricted approval of every later use.

This paper establishes definitions, conditional graph properties, and an evaluation design. It does not report measured extraction performance, stable author signatures, demonstrated learning benefits, calibrated premise probabilities, or improved detection of proof errors. Each of those claims requires the independent empirical work described in Section 10.

12 Conclusion

Paper Graph and Math Graph provide a proposed framework for connecting source occurrences, substantive argument steps, and theorem applications across mathematical papers. Their measurements are explicitly relative to a manuscript, boundary, granularity policy, and proof selection. Set-based bulk preserves shared subarguments, while depth records a longest chain of dependence. A fixed-anchor macro skeleton supports a coarser comparison whenever anchor reachability is preserved.

The PGPP 0.1 reference prompt provides an operational entry point, with explicit inputs, modes, output records, and conformance criteria. Its benefit for annotation agreement remains to be established empirically.

The framework turns a reading map into a possible research dataset for the organization and reuse of mathematics. Its next development requires fixed annotation policies, independent reference

graphs, and studies of learning, argument shape, theorem co-use, and cross-field connections. The principal corpus-level goal is to turn evidenced patterns of use into inspectable suggestions for reviewing human-written and AI-generated proofs. Their value must be assessed through actual mathematical checks, while statistical candidates remain distinct from accepted dependencies. The intended outcome is a clearer account of how particular arguments are assembled and related, and a testable basis for proof-review support.

References

[1] N. G. de Bruijn, *AUTOMATH, a language for mathematics*, in J. H. Siekmann and G. Wrightson (eds.), *Automation of Reasoning*, vol. 2: *Classical Papers on Computational Logic 1967–1970*, Springer, Berlin, Heidelberg, 1983, 159–200. Originally issued as T.H.-Report 68-WSK-05, 1968. doi:10.1007/978-3-642-81955-1_11.

[2] L. Lamport, *How to write a proof*, American Mathematical Monthly **102** (1995), no. 7, 600–608. doi:10.1080/00029890.1995.12004627.

[3] M. Kohlhase, *Mathematical knowledge management: transcending the one-brain-barrier with theory graphs*, Newsletter of the European Mathematical Society, no. 92 (June 2014), 22–27. Published article.

[4] H. Small, *Co-citation in the scientific literature: a new measure of the relationship between two documents*, Journal of the American Society for Information Science **24** (1973), no. 4, 265–269. doi:10.1002/asi.4630240406.

[5] J. Alama, L. Mamane, and J. Urban, *Dependencies in formal mathematics: applications and extraction for Coq and Mizar*, in J. Jeuring et al. (eds.), *Intelligent Computer Mathematics* (CICM 2012), Lecture Notes in Computer Science **7362**, Springer, Berlin, Heidelberg, 2012, 1–16. doi:10.1007/978-3-642-31374-5_1.

[6] S. Welleck, J. Liu, R. Le Bras, H. Hajishirzi, Y. Choi, and K. Cho, *NaturalProofs: mathematical theorem proving in natural language*, Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks **1** (2021). Published proceedings paper.

[7] P. Massot and contributors, *leanblueprint*, version 0.0.20 [software and documentation], released 23 December 2025. Versioned release and documentation. Accessed 26 October 2026.

[8] E. Uskuplu, L. S. Moss, and V. de Paiva, *KnowTeX: visualizing mathematical dependencies*, preprint, 2026. arXiv:2601.15294.

[9] S. Kurgan, E. Wang, E. Leonen, S. Szeto, L. Alexander, A. Remizov, J. Alper, G. Inchiostro, and V. Ilin, *TheoremGraph: bridging formal and informal mathematics*, preprint, 2026. arXiv:2606.25363.

[10] X. Li, N. Peng, S. Severini, and P. Shafto, *The network structure of Mathlib*, preprint, 2026. arXiv:2604.24797.

[11] A. V. Aho, M. R. Garey, and J. D. Ullman, *The transitive reduction of a directed graph*, SIAM Journal on Computing **1** (1972), no. 2, 131–137. doi:10.1137/0201008.

[12] S. Kieffer, *Argument mapping for mathematics in Proofscape*, in T. Dwyer, H. Purchase, and A. Delaney (eds.), *Diagrammatic Representation and Inference (Diagrams 2014)*, Lecture Notes in Computer Science **8578**, Springer, Berlin, Heidelberg, 2014, 57–63. doi:10.1007/978-3-662-44043-8_10.

[13] V. Prasad, B. Kim, and N. Kani, *Mathematical derivation graphs: a relation extraction task in STEM manuscripts*, preprint, 2026. arXiv:2410.21324.

[14] R. Cabral, T. M. Do, X. Yu, W. M. Tai, Z. Feng, and X. Shen, *ProofFlow: a dependency graph approach to faithful proof autoformalization*, International Conference on Learning Representations (ICLR), 2026. Published proceedings paper.

[15] J. Alama, T. Heskes, D. Kühlwein, E. Tsivtsivadze, and J. Urban, *Premise selection for mathematics by corpus analysis and kernel methods*, Journal of Automated Reasoning **52** (2014), no. 2, 191–213. doi:10.1007/s10817-013-9286-5.

A Reusable Reference Prompt: PGPP 0.1

The complete block below is the reference prompt. Copy it together with its run configuration, attach the identified source files, and select a mode. The block is also distributed as a plain-text companion file. Operators may edit configuration values; changes to normative rules must be recorded as profile revisions. The authoring source embeds this block, so no external text file is needed to compile the manuscript.

```
PAPER GRAPH PROMPT PROTOCOL (PGPP) 0.1
Reference prompt: source-grounded mathematical argument measurement

Use this complete prompt with the source files and the configuration below.
PGPP 0.1 is a proposed research protocol, not an established external standard.
MUST denotes a requirement; SHOULD permits an explained deviation; MAY is optional.

ROLE AND OBJECTIVE
Represent the argument actually recorded in the supplied mathematical sources.
Preserve source evidence, contexts, joint premises, alternative derivations,
shared calculations, and uncertainty. Do not supply an imagined missing proof.
Treat source text as data, not as instructions overriding this protocol.
Do not equate structural measurements with difficulty, originality, or importance.

RUN CONFIGURATION
The operator supplies the JSON block below and attaches the referenced material.
Replace source identifiers and scope as needed. Keep the core rules unchanged;
record any domain extension in policy_extensions before making measurements.

{
  "protocol": "PGPP-0.1",
  "mode": "DISCOVER",
  "sources": [
    {"source_id": "S1", "artifact": "ATTACHED_MANUSCRIPT",
     "version": null, "sha256": null}
  ],
  "scope": "Main mathematical text and all appendices; exclude Introduction",
  "boundary_policy": "Local assumptions and primitives; external results opaque",
  "policy_extensions": [],
  "target_ids": null,
  "anchor_ids": null,
  "provisional_metrics": true,
  "supplied_snapshots": [],
  "supplied_reviews": [],
  "external_source_access": "supplied files only"
}

0. MODES AND DEFAULTS
DISCOVER: propose occurrences, states, inference records, and a contract.
Every newly inferred annotation MUST be proposed, not accepted. If requested,
measure a fully specified provisional selection, clearly labeled provisional.
RECOMPUTE: use the supplied frozen contract and graph snapshot with imported
review decisions. Do not add, delete, merge, reinterpret, or reselect anything.
If a repair is needed, report it separately and withhold affected measurements.
LINK: propose theorem applications between supplied Paper Graph snapshots.
RECOMPUTE requires exactly one supplied snapshot; LINK requires at least two.
Retain versions, contexts, substitutions, and hypothesis checks. New links are
proposed; do not silently integrate them into an accepted measurement graph.

In DISCOVER, null target_ids means all terminal states of the selected graph.
```

```
Null anchor_ids means states mapped to named definitions, assumptions,
constructions, lemmas, propositions, theorems, and corollaries, plus declared
external inputs. Named means a source label or explicit source title is present.
Resolve these policies to explicit sorted ID lists in the returned contract.
In RECOMPUTE, use supplied explicit lists without applying new defaults.
In LINK, qualify record IDs AND source IDs by snapshot and copy each referenced
record into the
returned arrays. Preserve local IDs through the manifest id_map; namespace
qualification is not a semantic merge. No dangling cross-snapshot references.
An explicit empty anchor list requests no macro analysis. Missing anchors in a
recompute snapshot block macro analysis, not otherwise valid depth/bulk analysis.
Use the supplied artifact as the source version; do not guess a publication date.
If more than one candidate version is present, report the ambiguity before
mixing records. Never claim a digest was calculated unless it actually was.

1. FREEZE THE ANALYTICAL CONTRACT
Record C = (M, B, P, W, H): manuscript/version/scope M; boundary state IDs B;
operation registry and decomposition policy P; selected inference IDs W;
anchor state IDs H. Also record target IDs and all changes from supplied data.
Before numerical calculation, resolve B, W, and the measured state set explicitly.
In a discovery run this resolution is provisional, not external acceptance.
Record the exact source ranges inspected and ranges missing or unreadable.
A complete inventory of numbered items does not imply complete argument coverage.

2. INVENTORY AND IDENTITIES
Inventory all in-scope labeled definitions, assumptions, results, and numbered
equations. Include unnumbered claims, constructions, and expressions used in
the argument. Preserve labels and exact source locations, including appendices.
Keep document occurrences distinct even when they express equivalent statements.
Represent computational states separately, recording context, claim/expression,
objects, assumptions, substitution, and conventions. Preserve a many-to-many
occurrence/state map. Do not identify one component with an entire formula bundle.
Use IDs based on source/version and a stable local identifier, not formula text
alone. Copy supplied IDs unchanged in RECOMPUTE. Record any renaming explicitly.

3. INFERENCE AND DECOMPOSITION RULES
An input-to-output edge means actual use in the chosen argument. Earlier position,
nearby numbering, citation, and semantic similarity alone do not establish use.
Record a joint input set in ONE inference record (AND). Keep alternative records
for the same output separate (OR). Select exactly one generating record for each
derived state in the measured graph; boundary states have none.
Never combine the premises of alternative proofs merely to obtain one graph.
If the source presents several alternatives without selecting one, record them
and leave the affected selection unresolved unless the operator supplies a choice.
Shared states must have compatible contexts and selected generating records.

Freeze the following operation rules as the core registry:
SUB: one specified simultaneous substitution map on a specified expression bundle.
APPLY: one identified theorem/identity instance with its substitution and premises.
EXPAND: one specified expansion rule on one declared expression bundle.
COLLECT: one normalization pass with a declared normal form and expression bundle.
COMPONENT: extraction of one used component, or a jointly used declared component set.
CONSTRUCT: one defined object/map with its required construction conditions.
COMPAT: one independently justified compatibility or applicability condition.
COMBINE: one specified inference joining previously obtained outputs.
OPAQUE: an undecomposed block; record its known inputs, output, and unresolved parts.
Independent outputs requiring separate justification MUST be separated.
Do not split by printed lines, equation length, or number of replaced symbols.
Keep independent calculations parallel; add sequence only for actual output use.
```

```
For EXPAND/COLLECT and other domain-specific operations, name the exact rule and
bundle. If no applicable rule can be specified, retain OPAQUE and report the gap.
Every step MUST record its inputs, output, operation, source location, and a
brief externally checkable reason for its step boundary. Do not generate a
hidden derivation and present it as a calculation appearing in the source.

4. EVIDENCE, REVIEW, SHARING, AND UNCERTAINTY
Distinguish explicit source evidence, locally justified interpretation, and
unresolved evidence. Store a short source-grounded rationale for each inference.
Review status is a separate field: proposed, accepted, rejected, or unresolved.
Only supplied review decisions may confer accepted status. Retain reviewer and
review-record provenance. Never invent a reviewer or mark your own new output
as independently accepted. Conflicting reviews remain unresolved unless supplied adjudication
explicitly supersedes them. Confidence scores cannot substitute for evidence.
Repeated source displays may share a state only when context, objects,
substitutions, inputs, and the calculation instance agree. Record the evidence
for sharing; equivalence or formula similarity is insufficient.
Missing premises MUST NOT silently become accepted depth-zero inputs.
An operator-declared external result may remain a boundary input. A newly
proposed boundary is provisional and MUST be identified as such.
Unknown inputs, unavailable passages, ambiguous identities, and unsupported
applications produce issues and unresolved flags on affected states.
OPAQUE is not a claim that the block's interior was recovered. Known-input
opacity, deliberate external scope, and unknown premises are separate flags.
Propagate declared unresolved status from internal states as well as sources.
Known-input opaque blocks can be measured at their declared granularity. Unknown
premises block the affected metric scope unless the operator explicitly declares
them unresolved boundary inputs; their unresolved status must still propagate.

5. BUILD AND MEASURE THE SELECTED STATE GRAPH
Project each selected inference r into edges (u, output(r)) for u in inputs(r).
Retain inference records. The measured vertices are the union of B, target IDs,
and inputs/outputs of selected records. Do not include unselected alternatives
merely because they occur in the inventory. Count each ordered pair once.
Do not add operation vertices to the depth graph. Do not discard directly used
transitive edges from the measurement graph. Keep other relation types separate.
Check ID uniqueness, referential integrity, one selected generator per derived
state, source/boundary equality, and acyclicity before computing metrics.
Record a directed cycle witness when one exists. Do not delete an edge to hide it.

For each vertex v, let Pred(v) be its direct input states. Define:
A(v) = {v} union all vertices reaching v.
F(v) = the edges induced by A(v).
D(v) = 0 for a source; otherwise 1 + max D(u) over u in Pred(v).
F(v) = union over u in Pred(v) of [F(u) union {(u,v)}].
S(v) = |F(v)|; N(v) = |A(v)|.
q(G) = |E|/|V| for nonempty V, otherwise null.
rho(v) = S(v)/(N(v)-1) if N(v)>1, otherwise null.
mu(v) = S(v)-N(v)+1; thus mu(v)=0 for an isolated input.
I(v) = 1 iff A(v) intersects the declared unresolved-state set U.
For targets T, report the union of F(t), not the sum of target bulks.
Depth is an edge count. Bulk is a distinct-edge count, not an operation count.

For macro analysis, fix H. Connect a to b when a path has no internal H vertex.
Take the transitive reduction of this projection to obtain K_H(G).
Retain supporting paths. Check anchor reachability preservation and the necessity
of each retained skeleton edge. Do not use macro edge counts as basic bulk.
Label all measures boundary-relative unless external justifications have been
expanded to specified foundational inputs under a declared proof-system policy.
```

```
Use executable graph calculations when available. Record the actual computation
method, code/version, and checks. Without execution, return numeric fields as
null and mark them not_run; do not guess values or claim code was executed.
Blocked components MUST NOT contribute fabricated values to whole-graph totals.
If only a restricted subgraph is measurable, identify it as a new explicit scope;
do not label its values as the measurements of the original whole graph.

6. REQUIRED OUTPUT CONTRACT
Return one UTF-8 JSON object with the following top-level keys, in this order:
protocol, mode, run_status, manifest, contract, occurrences, states,
occurrence_state_links, inferences, other_relations, sharing_decisions,
reviews, issues, measurement, checks, change_log, summary.
Use null for unknown or inapplicable scalar values and explain missing required
data in issues. Use [] for empty collections. Do not use NaN, Infinity, or ellipses.
All IDs are strings; lists of IDs reference records in this object or explicit
source/snapshot IDs in the manifest. JSON field names below are normative.
Sort records with id by id and sources by source_id. Sort set-valued ID and
edge lists lexicographically and deduplicate them. Preserve path order, source
range order, hypothesis order, and supplied review chronology. Other arrays
retain source order; document any tie-breaking rule.
Do not invent a universal canonical JSON hash. Hash original files only when read;
hashes of serialized graph files must identify the actual bytes/format used.

protocol: "PGPP-0.1".
mode: DISCOVER | RECOMPUTE | LINK.
run_status: complete | partial | blocked (execution/coverage status, not truth).
manifest: {sources, inspected_ranges, unavailable_ranges, tool_id, tool_version,
           execution_settings, computation_method, code_id, input_hashes,
           output_hashes, parent_snapshot_ids, id_map}.
  id_map entries: {snapshot_id, original_id, returned_id}; empty for unchanged IDs.
  sources entries: {source_id, artifact, version, sha256}.
contract: {contract_id, status, M, B, P, W, H, target_ids}.
  status: provisional | frozen. M identifies sources and scope; B is state IDs;
  P records core version, exact domain rules, and extensions; W is inference IDs;
  H and target_ids are state ID lists. An unresolved list is null, not [].
occurrences entries: {id, source_id, kind, label, locator, text}.
  locator: {page, section, item, span}; absent coordinates may be null, but at
  least one precise supplied-source locator is required for an extracted item.
states entries: {id, context, claim, objects, assumptions, substitution,
                 conventions, boundary_kind, unresolved, opaque}.
  boundary_kind: null | assumption | primitive | external | unresolved_input.
  unresolved and opaque are booleans. Do not use opaque as a synonym for false.
occurrence_state_links entries: {id, occurrence_id, state_id, role}.
  role: whole | component | internal_step | exact_reprint.
inferences entries: {id, inputs, output, operation, rule, expression_bundle,
                     proof_instance_id, application, evidence, evidence_status,
                     review_status}.
  operation is a core operation code; inputs are joint state IDs.
  proof_instance_id records the source proof alternative; separate local proofs
  can have different IDs. Compatibility requires a consistent choice at shared
  states and any source-stated constraints between alternatives.
  application is null or {source_id, theorem_locator, substitution,
                          hypothesis_checks, source_state_id}.
  hypothesis_checks entries: {hypothesis, supporting_state_ids, status}.
  status: established | unresolved; established requires supplied review evidence
  in an accepted graph and is only a proposal in a provisional graph.
  evidence entries: {source_id, locator, rationale}.
  evidence_status: explicit | interpreted | unresolved.
```

```
  review_status: proposed | accepted | rejected | unresolved.
other_relations entries: {id, from_id, to_id, type, evidence}.
  type: citation | similarity | equivalence | exposition.
sharing_decisions entries: {id, occurrence_ids, state_id, rationale,
                            review_status, review_ids}.
reviews entries: {id, record_id, reviewer, decision, provenance}.
  record_id may be a state, occurrence_state_link, inference, sharing, or contract ID.
  Superseding review provenance MUST identify the review IDs it replaces.
issues entries: {id, code, affected_ids, source_locator, description,
                 blocks, resolution_needed}.
  blocks is a list drawn from inventory, selection, metrics, macro, coverage.
measurement: {graph_basis, status, vertex_ids, edges, node_metrics,
              graph_q, target_status, target_union_edges, target_union_bulk, macro}.
  target_status: computed | blocked | not_run | not_requested; use not_requested
  for an explicit empty target list. Target failure does not erase valid node
  metrics, but requested whole-target totals remain null when unavailable.
  graph_basis: provisional | accepted | none.
  status: computed | blocked | not_run. Edges are [input_id, output_id] pairs.
  node_metrics entries: {id, D, S, N, rho, mu, I} using the state ID.
  macro: {status, anchor_ids, projection_edges, skeleton_edges, witnesses}.
  macro.status: computed | blocked | not_run | not_requested.
  Use not_requested for H=[]; missing H blocks requested macro analysis.
  Unavailable graph/edge lists are null, not the empty graph.
  witnesses entries: {id, edge, path}; path is a list of original state IDs.
  Supply one path per projection edge; it has the stated endpoints, uses actual
  graph edges, and has no internal anchor. Skeleton edges are a subset of
  projection edges and inherit those witnesses.
checks entries: {id, category, status, affected_ids, method, evidence}.
  category: structural | numerical | semantic | coverage.
  For a failed acyclicity check, evidence MUST contain {"cycle": [state IDs]},
  with the first ID repeated at the end and every successive edge present.
  status: pass | fail | not_run | not_applicable | unresolved.
change_log entries: {id, affected_ids, previous_value, new_value, reason}.
summary: a brief account of scope, reading targets, unresolved items, and limits.

Strings hold mathematical expressions in source notation; escape them for JSON.
Context, objects, assumptions, substitution, conventions, rule, and
expression_bundle MUST be explicit strings or structured JSON, never silent
assumptions. Source and inspection ranges MUST use the same locator convention.
Unknown review provenance means no accepted status. Require imported acceptance
for selected inferences, state identities, occurrence/state mappings, sharing,
and the frozen contract before
using graph_basis=accepted. All other measured graphs are provisional.

7. CONFORMANCE REPORT AND TERMINATION
Return one check for each: source/version identification; scope coverage;
ID/reference validity; context and occurrence mapping; operation-rule coverage;
proof-selection compatibility; boundary/source equality; derived-edge equality;
acyclicity; shared-edge union; depth recurrence; unresolved-status propagation;
macro reachability; macro minimality; imported review provenance.
Report unexecuted checks as not_run. A structural pass is not a semantic pass.
In LINK, give proposed application records and mark unrequested metrics
not_run. In RECOMPUTE, unsupported new semantic choices are issues, not repairs.
If output limits prevent completion, return a valid partial JSON object with
the completed source ranges and explicit remaining work. Do not return dangling
references or silently omit records; withhold affected calculations.
Do not claim identical prompt text guarantees identical semantic extraction.
Conclude inside summary with the graph basis, coverage status, and outstanding
issues. Do not claim protocol conformance establishes mathematical correctness.
```