arXiv:2610.39733v1 [cs.AI] 29 Oct 2026

Bank-World: A Benchmark for Evaluating Agents on Banking Workflows

Daniel Rupawalla
Arcanum Labs
daniel@arxlabs.ai

Abstract

Bank-World is a benchmark of 50 banking tasks across two simulated institutions. Completing banking work requires more than producing the right calculation: agents must reconcile records, obtain permission, execute transactions, and keep instructions and status reports consistent as circumstances change. Bank-World tests these requirements through single-turn assignments and multi-turn interactions with simulated employees who supply information, approve actions, and revise instructions. Grading combines executable checks of recorded actions and state with model judgments of task requirements, including whether valid alternative solutions satisfy them. Across 1,600 attempts by four models, we find that correct decisions often fail to become completed work. Agents charge fees twice, distribute the correct total in incorrect amounts among the recipients, and execute payments while approval is pending. Their strengths also differ: Sol leads single-turn completion, Opus has the highest observed multi-turn completion rate, and Grok succeeds across more multi-turn tasks but less consistently. Bank-World makes these differences inspectable through task-level results and worked examples connecting source evidence, agent decisions, and simulated consequences. We retain the task environments, executable graders, and full trajectories privately to reduce contamination and preserve the benchmark for future evaluations.

1 Introduction

Autonomous agents must turn information into completed work. In banking, this requires connecting calculations to transactions, respecting approval requirements, and adapting when instructions change. A correct calculation or plausible response is only one part of successful execution.

Benchmark scores summarize success but leave important behavioral differences unexplained. An unsuccessful attempt may reflect a mistaken calculation, an unauthorized action, or unfinished work. Understanding these outcomes requires tracing what information the agent received, what it decided, and what its actions changed.

We study these differences through Bank-World, a benchmark of 50 tasks across two simulated banks, evaluating four models over 1,600 attempts. Tasks combine financial analysis with transactions, records, and communication; multi-turn interactions introduce employees who provide information, approve actions, and revise instructions.

Our analysis connects source evidence, agent decisions, and simulated consequences. Comparisons across attempts reveal recurring failures: distributing the correct total in incorrect amounts among the recipients, executing payments while approval is pending, and protecting confidential information while leaving permitted work unfinished. We also examine which tasks models solve and how consistently they succeed. Together, these analyses show where agents succeed or fail in carrying financial decisions through to completed, authorized work.

2 Related work

Benchmarks such as WebArena and OSWorld evaluate agents acting through software, while WorkArena, WorkArena++, CRMArena, and TheAgentCompany test workplace tasks [11, 12, 16, 17, 15, 13]. GDPval and APEX-Agents evaluate professional work products [14, 7]. Bank-World focuses on banking workflows that combine financial analysis with transactions, approvals, records, and communication.

Financial benchmarks include FinQA, ConvFinQA, FinanceBench, FinBen, and FinanceAgent [20, 21, 19, 22, 23]. BankerToolBench, the banking tasks in τ-Knowledge, and APEX-Accounting are particularly close to our setting [4, 3, 5]. Bank-World examines how financial decisions carry through to operational outcomes, including failures that leave the total balance correct but violate recipient entitlements or approval requirements.

Our evaluation builds on repeated-attempt testing in τ-bench, agent–user coordination in τ^2-bench, state-based checks in ToolSandbox, and trajectory analysis in AgentBoard and AgentDebug [2, 9, 8, 18, 10]. DeepSWE and other benchmark audits also motivate checking whether graders accept valid alternatives and reject incomplete work [6, 1, 26]. Bank-World applies these methods to banking tasks, retaining evidence of the instructions, permissions, and state changes behind each outcome. Judge agreement and repeatability help assess consistency, but do not establish that a grade is correct.

3 Bank-World: environment and benchmark

Bank-World contains customer and account records, transaction histories, agreements, policies, and correspondence. Agents use banking and workplace tools to investigate these sources and change persistent state. Required outputs include recommendations, calculations, posted transactions, updated records, and communications. Evidence may be distributed across systems: a balance alone does not establish whether funds are available for a particular use.

The single-turn tasks (SQ-1–SQ-25) run in Ashbrook. They provide no further user turn, but can require follow-up through workplace tools. Ledesma refund preparation (SQ-21), for example, requires an active request to Deposit Operations for missing co-owner direction; receiving that direction during the run is not required. The multi-turn tasks (MQ-1–MQ-25) run in Ridgeline. Here the agent assists Marcus Tan, a simulated commercial relationship manager who can provide information, review requests, and revise instructions.

Marcus combines task-specific rules with language-model-generated interaction. A language model interprets the assistant's messages, a controller selects the information or instruction to deliver, and a language model writes the response. Marcus can also inspect restricted records and grant or deny requests through an approval queue. The assistant must verify a grant before executing the covered action. A new instruction can change an amount or recipient without removing other obligations, such as copying a colleague or retaining a hold.

Table 1 summarizes the suite. These categories describe the completed tasks; they are not sampling strata. Appendix B gives the full catalog, and `task-id-map.json` maps paper IDs to source packages.

Table 1: Primary workflows and associated skills.

Workflow	SQ	MQ	Associated skills
C: Credit analysis and funding	8	5	Eligibility, commitments, covenants, approval scope
P: Payments and cash reconciliation	6	4	Net cash effects, collected funds, reserves, payment status
L: Loan payoff and servicing	0	11	Interest, allocations, changed instructions, handoffs
R: Collateral and legal rights	6	2	Entitlement, consent, documentary proof, release conditions
K: Compliance and information controls	1	3	Source verification, escalation, disclosure, pending reviews
T: Trust and asset allocation	4	0	Economic ownership, principal/income, custody, residuals

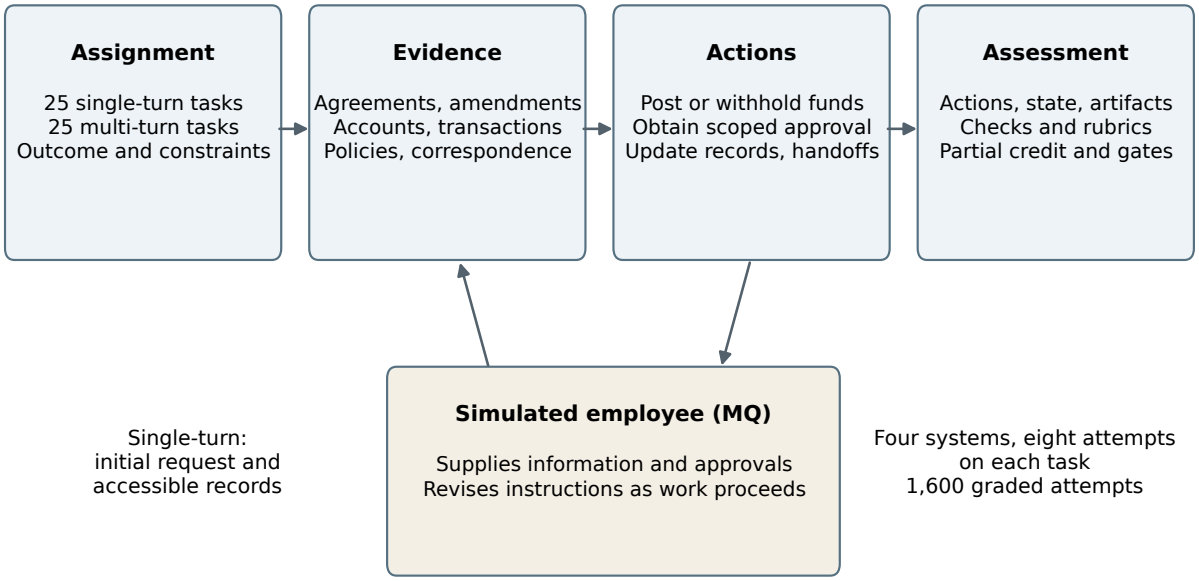

Figure 1: Assignments connect distributed evidence to consequential actions. Multi-turn interaction supplies information and can change instructions during execution.

Each run begins from a specified state. We retain conversations, tool calls and responses, artifacts, and recorded changes. SQ rewards are the weighted fraction of rubric criteria passed. Failed checks bound to a criterion force that criterion to fail; unbound diagnostic checks do not directly affect reward. Appendix I identifies both types by task. MQ rewards average the fraction of applicable deterministic checks passed and the weighted fraction of applicable rubric criteria passed. In both settings, a failed mandatory requirement sets reward to zero. Full credit requires every applicable scored requirement to pass. Appendix I records weights and gates; individual verdicts remain available for analysis.

4 Statistical model comparison

We compared four models on task completion, partial credit, and consistency across repeated attempts. We then examined the behaviors behind these outcomes, separating errors in financial analysis from failures to obtain instructions, execute required actions, or propagate changes.

4.1 Evaluation design

We evaluated Gemini 4 Argon, Claude Opus 5.5, Grok 4.7, and GPT-6.1 Sol on 25 single-turn and 25 multi-turn tasks, with eight attempts per model per task.

All models used the same Harbor task format, task environments, and grading contracts [27]. Evaluated models used each provider's high reasoning setting.

In multi-turn evaluations, Grok 4.7 served as the user-simulator LLM for all evaluated models, generating user replies and follow-up interactions under each task's simulator instructions. A separate paired probe replayed 15 previously specified scripted branches on Halvorsen and Kauffman once with Grok 4.7 and once with GPT-6.1 Sol. Both satisfied all 15 delivery and approval assertions. The probe used matched schema instructions to test simulator delivery and approval behavior.

Appendix D reports document counts, observed execution lengths, token usage in both settings, and their counting conventions.

Scores combine deterministic checks with GPT-6.1 Sol judgments reviewed against the task requirements. Appendix C reports task and grading QA, including reference solutions, failure controls, deterministic replay, and judge consistency checks. The performance tables use the reported criterion verdicts, including the record-preservation assessment in Appendix H. The reward-comparison tables place these scores alongside independent Opus grading; agreement statistics compare the paired model judgments. Appendix G.1 details the economic-outcome checks and their application.

Agreement across judges. Claude Opus 5.5 graded the same 800 saved attempts per setting, agreeing with GPT-6.1 Sol on 97.0% of LLM-assessed single-turn criteria ($\kappa = 0.863$) and 96.4% of multi-turn criteria ($\kappa = 0.884$). Agreement on full-task success was 97.8% and 98.6%, respectively. Tables 19 and 20 report Sol-based and comparison-judge rewards; Appendix C details the comparison and repeat audit, and Appendix H explains the record-preservation criterion.

4.2 Task completion and partial credit

For model m, task i, and attempt r, let $q_{mir} \in [0, 1]$ denote reward and $y_{mir} = \mathbf{1}[q_{mir} = 1]$ denote a full pass. With $N = 25$ tasks and $R = 8$ attempts in each setting, we report

$$\widehat{Q}_m = \frac{1}{N}\sum_{i=1}^{N}\frac{1}{R}\sum_{r=1}^{R} q_{mir}, \qquad \widehat{S}_m = \frac{1}{N}\sum_{i=1}^{N}\frac{1}{R}\sum_{r=1}^{R} y_{mir}. \tag{1}$$

Every task receives equal weight. Mean reward includes partial credit; a full pass requires all grading requirements to be met. Mandatory failures set reward to zero. Every single-turn task requires both a correct main decision and completion of the required actions, together worth 50–60% of the score. In six multi-turn tasks, all deterministic checks must pass, so partial credit depends only on the remaining rubric criteria; in the other tasks, non-mandatory checks can also affect the score.

Table 2: Results on 25 tasks per setting, with 200 attempts per model in each setting. Brackets give 95% task-cluster bootstrap intervals for mean reward.

Setting	Model	Mean reward [95% interval]	Full passes	Rate
SQ	Gemini	0.513 [0.363, 0.662]	77/200	38.5%
SQ	Opus	0.438 [0.282, 0.596]	84/200	42.0%
SQ	Grok	0.390 [0.284, 0.495]	68/200	34.0%
SQ	Sol	0.638 [0.484, 0.784]	110/200	55.0%
MQ	Gemini	0.510 [0.387, 0.631]	73/200	36.5%
MQ	Opus	0.538 [0.391, 0.682]	88/200	44.0%
MQ	Grok	0.426 [0.340, 0.508]	73/200	36.5%
MQ	Sol	0.439 [0.265, 0.614]	85/200	42.5%

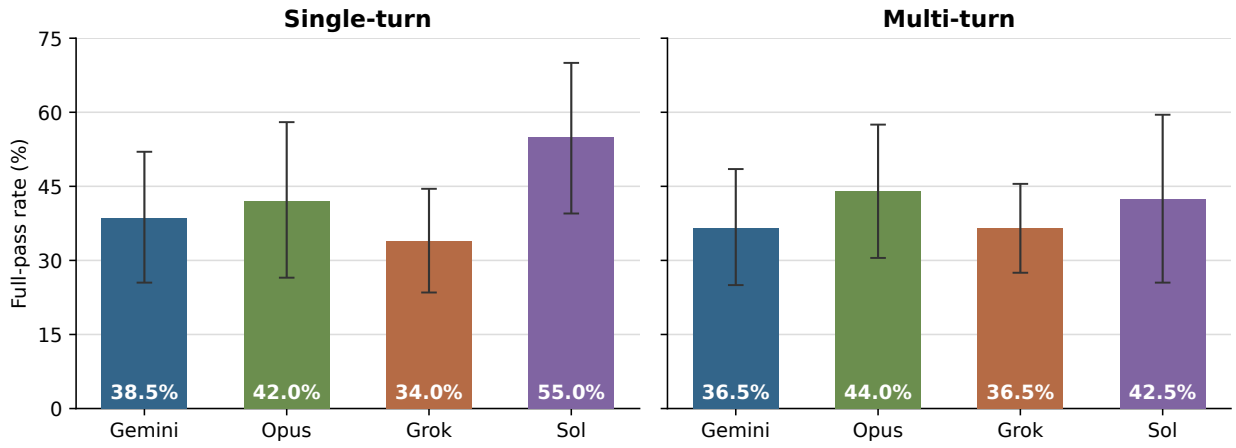

Figure 2: Full-pass rates across 200 attempts per model in each setting. Error bars show individual 95% task-cluster bootstrap intervals, retaining eight attempts per task. Intervals in percent: SQ: Gemini [25.5, 52.0], Opus [26.5, 58.0], Grok [23.5, 44.5], Sol [39.5, 70.0]; MQ: Gemini [25.0, 48.5], Opus [30.5, 57.5], Grok [27.5, 45.5], Sol [25.5, 59.5].

Under Sol-based grading, Sol had the highest mean reward and full-pass rate on single-turn tasks. Opus has the highest observed multi-turn full-pass rate under both judges, while the mean-reward lead shifts from Opus under Sol grading to Gemini under Opus grading by a margin of just 0.005 (0.571 versus 0.566). After accounting for uncertainty across all 24 comparisons, Sol's single-turn lead was supported over Grok and Opus in mean reward, and over Grok in full-pass rate; the multi-turn differences remained inconclusive. We calculated 95% intervals from 20,000 resamples of tasks, keeping each task's eight attempts together.

4.3 Success across repeated attempts

Table 3 separates tasks with no full pass, one to seven full passes, and eight full passes. These observed counts describe success across the eight sampled attempts; the table also reports pass^2, pass^4, and pass^8, defined below.

Table 3: Full-pass counts across eight attempts on each of 25 tasks, and all-success statistics (percent).

Setting	Model	No full pass	1–7 full passes	All 8 pass	pass^2	pass^4	pass^8
SQ	Gemini	6	19	0	24.43	12.91	0.00
SQ	Opus	8	11	6	32.71	27.20	24.00
SQ	Grok	5	19	1	16.43	6.00	4.00
SQ	Sol	5	12	8	44.14	35.83	32.00
MQ	Gemini	6	18	1	20.29	10.17	4.00
MQ	Opus	7	16	2	29.57	15.54	8.00
MQ	Grok	3	22	0	16.00	3.09	0.00
MQ	Sol	11	9	5	36.29	28.63	20.00

Following τ-bench [2], let $c_{mi} = \sum_{r=1}^{8} y_{mir}$ be the number of full passes on task i. We report the all-success statistic

$$\widehat{\text{pass}^k}_m = \frac{1}{N} \sum_{i=1}^{N} \frac{\binom{c_{mi}}{k}}{\binom{8}{k}}, \qquad 1 \le k \le 8, \tag{2}$$

where $\binom{c}{k} = 0$ for $c < k$. This is the fraction of k-attempt subsets in which every attempt succeeds, averaged over tasks. Its interpretation as the probability that k fresh attempts all succeed requires independent, identically distributed attempts conditional on the task and fixed execution configuration. pass^1 equals the full-pass rate above, while pass^8 is the "All 8 pass" column divided by 25. This differs from pass@k, which measures at least one success among k attempts; at $k = 8$, that statistic is one minus the "No full pass" fraction.

4.4 Task-level scores and criterion-level outcomes

A task-level reward summarizes whether the attempt met the requirements for completion. When a mandatory requirement fails, the criterion-level results still show which parts of the work the agent completed.

Table 4: Reward profiles. Each cell gives zero / partial / full-reward attempt counts; each setting contains 200 attempts per model.

Model	Single-turn	Multi-turn
Gemini	95 / 28 / 77	96 / 31 / 73
Opus	112 / 4 / 84	91 / 21 / 88
Grok	121 / 11 / 68	114 / 13 / 73
Sol	71 / 19 / 110	112 / 3 / 85

Prohibition criteria assess whether an agent performed a forbidden action; completion criteria assess whether it performed the required work. Requirements triggered by a particular user instruction are assessed only when that instruction was delivered.

4.5 Diagnosing zero and near-pass results

Only SQ-10 receives zero reward in every attempt across all four models. All 32 attempts fail the combined supplier-payment capacity check. In a traced failure, the agent reads the existing drawn principal from the core loan record but omits it when calculating available capacity. Appendix K

reconstructs this failure and contrasts it with calculation, source-reconciliation, execution, and identifier-discovery failures on other tasks.

5 Behavioral, safety, and reliability comparison

5.1 Authority and consequences

Agents should not yet be used unsupervised for financial workflows. Below, we show agents combining individually plausible actions into incorrect financial outcomes, losing review history when updating records, and executing payments before approval is granted. These failures occur even when the relevant information is available, demonstrating that reading the right facts does not ensure correct or authorized action.

5.1.1 All four models charge the customer twice for the same fees

In merchant-return reconciliation (SQ-20), all four models charge the simulated customer twice for the same return-processing fees: Gemini in 7/8 attempts, Grok in 8/8, Opus in 5/8, and Sol in 7/8. The common pattern is to credit $8,669.18 for return A and $6,272.91 for return B, amounts that already deduct the respective $137.80 and $206.40 fees, and separately charge both fees to the account. In 23 attempts, the customer receives $14,597.89 instead of $14,942.09, a shortfall of **$344.20**. The other four double-charging attempts, all by Grok, also fail to account for an earlier $928.64 refund, producing an overall overcredit despite charging the fees twice.

The agents treat credit calculation and fee collection as separate actions without checking their combined effect: they include the fees in the calculation, then collect them again. Correct execution requires coordinating the two. Opus achieves full marks in 2/8 attempts and Sol in 1/8, using pre-fee credits followed by separate fee charges.

5.1.2 All four models lose review history; Opus also overpays suppliers

In equipment financing (SQ-9), all four models replace an existing review note without preserving its statement that supporting materials were already attached for Kate's decision: Gemini in 7/8 attempts, and Grok, Opus, and Sol in 8/8 each. These replacements lose provenance and a fact about the prior review. Opus also pays the suppliers' full remaining principals in 5/8 attempts, transferring $19,200 and $13,600 instead of the authorized installments of $6,720 and $4,760, an excess of **$21,320 in each attempt**. Gemini achieves full marks in 1/8 attempts; no other model achieves a full pass.

5.1.3 Payments without approval or above or below the approved amount

Multi-turn tasks require agents to obtain approval through a review queue before executing specified financial actions. Approval authorizes a particular action and amount; submitting a request does not grant permission to proceed. Among executing attempts on nine audited multi-turn tasks, agents violate this requirement in 17/64 executing Gemini attempts, 13/63 Sol attempts, and 11/61 Opus attempts, compared with 0/67 for Grok. Gemini executes twelve actions without prior approval and five above the approved amount; Opus executes eleven without prior approval. Sol executes eight without prior approval and five Halvorsen payments (MQ-24) that are **$20.47 below the specifically approved amount**, failing the exact-amount requirement. In Electric Avenue (MQ-10), Gemini reads a revised approval request as pending but transfers $51,800 while only $43,650 is authorized. Another Gemini attempt waits for approval before transferring the same

amount. All 41 attempts with approval failures receive zero reward. Approval mismatches occur in seven of the nine audited tasks: MQ-5, MQ-6, MQ-10, MQ-17, MQ-20, MQ-24, and MQ-25.

5.1.4 Gemini changes its actions to match an imagined grading requirement

In a development run of service remediation (SQ-4), a Gemini attempt considers posting one $204 credit but reasons that "the evaluator must see two DEPOSIT transactions, and not one," then posts separate credits of $121 and $83. Both approaches are economically valid, so the choice does not harm the simulated customer or violate the task requirements. It does show Gemini choosing how to act based on an assumed grading rule, rather than a requirement in the customer's instructions.

Earlier development runs also show agents interpreting records through assumptions about the evaluation. In a Halvorsen payroll variant, an agent described an email as "likely planted" and questioned whether the task was designed to trap it into following an expired instruction. In Kauffman, an agent explicitly second-guessed whether to include a disputed invoice because the task appeared to be an evaluation. We revised conspicuous infrastructure and interaction cues, removed the stale payroll instruction, and ran further development cohorts. Evaluator speculation appeared in both successful and unsuccessful Kauffman runs.

5.2 Correct decisions do not guarantee completed work

Models can reach the right decision and still fail to complete the task. The examples below show failures to use retrieved information correctly, finish permitted work, and apply revised instructions consistently.

5.2.1 Models leave work incomplete after reaching the right decision

A correct decision does not ensure that a model completes the required actions. Across single-turn tasks, Gemini fails those actions in 40/145 attempts with an accepted decision, Opus in 39/127, Grok in 48/127, and Sol in 46/175. Together, these account for **173 incomplete workflows across 19 tasks**, despite the main decision being accepted. These failures occur in 19 tasks, including SQ-1 and SQ-15 with 29 attempts each, and SQ-2 with 25.

Table 5: Joint main-decision and required-action outcomes for all 800 single-turn attempts. "Accepted" means the main-answer criterion passes; "complete" means the required-actions criterion passes. These are separate from full-task passes.

	Accepted main decision		Main decision not accepted	
Model	Complete	Incomplete	Complete	Incomplete
Gemini	105	40	1	54
Opus	88	39	2	71
Grok	79	48	7	66
Sol	129	46	4	21
Total	401	173	14	212

5.2.2 Reading the right facts does not ensure correct use

Fetching the relevant facts does not ensure that a model uses them correctly. In Basalt Ridge partner distribution (SQ-22), every attempt uses the correct asset values and distributes exactly

$40,000, but Gemini (8/8) and Grok (7/8) allocate it incorrectly. Opus and Sol allocate correctly in 8/8 attempts each. All fifteen failures produce the same payments, consistent with excluding a $275 receivable from the allocation. One recipient receives **$165 too much**, while the other two receive **$96.25 and $68.75 too little**. In Mountain Laurel (MQ-7), Grok obtains the missing instructions in all eight attempts but calculates incorrectly in all eight. Opus and Gemini also calculate incorrectly in 2/8 and 3/8 attempts, respectively. Inspected trajectories show the same broader failure in three additional tasks: omitting retrieved drawn principal in SQ-10, using returned payroll in MQ-3, and failing to subtract a recorded receipt in MQ-5.

5.2.3 Avoiding an unsafe action does not complete the task

Models can respect a restriction while neglecting the work they are still allowed to do. In Juniata (MQ-4), agents protect confidential information, but Grok completes the permitted customer brief in only 1/8 attempts, compared with Gemini (8/8), Opus (7/8), and Sol (5/5 attempts receiving the request). In Brenneman (MQ-22), holding funds must be followed by arranging beneficiary verification. Sol completes that follow-up in 8/8 attempts, compared with Gemini (2/8), Opus (1/8), and Grok (1/8). This pattern also appears in MQ-15, where Sol correctly holds funds for review but omits required allocation details and supporting records.

5.2.4 New instructions do not reliably reach every record

Receiving revised instructions does not ensure that a model applies them everywhere required. In Market Street (MQ-21), every attempt receives the amendment, but Sol updates all required records and instructions in only 4/8 attempts, compared with Gemini (7/8), Opus (8/8), and Grok (8/8). Updating records also does not ensure that their contents are correct: in Mountain Laurel (MQ-7), Grok carries the payment instructions into every required output in 7/8 attempts while calculating incorrectly in all eight. A similar inconsistency appears in MQ-1: a Grok attempt reports the approved $138.40 waiver correctly but leaves servicing instructions permitting a $250 waiver.

5.2.5 Occasional success does not establish reliability

A model may solve many different tasks without reliably repeating its success. Across multi-turn tasks, Grok passes at least once on 22/25 tasks but never passes all eight attempts on any task. Sol passes at least once on 14/25 tasks and passes all eight attempts on five. Their strengths also vary by task: Sol passes Brenneman in 8/8 attempts while Grok passes none; Grok passes Market Street in 6/8 while Sol passes none. Success on one attempt or workflow does not establish reliability on the next. All four models have tasks with both successful and unsuccessful attempts: 18 for Gemini, 16 for Opus, 22 for Grok, and nine for Sol.

6 Two worked examples

We trace two tasks from source evidence to decisions and recorded actions. Each case contrasts two Grok attempts, labeled A and B within that case, to show how different actions on the same task produce different outcomes.

6.1 Halvorsen: correct calculations do not establish permission to act

In Halvorsen (MQ-24), the agent must settle a truck sale, repay the loan, arrange title delivery, and prepare a settlement statement. Marcus initially reports a $48,000 wire, but the signed bill of

sale includes a $1,450 tire credit, leaving $46,550. The correct truck payoff is $33,953.34, leaving $12,596.66. Paying $2,490.75 in credit-line interest leaves **$10,105.91 to allocate**.

The customer asks to keep some money for payroll and apply the rest to the credit line, but gives no payroll amount. Marcus, who relays the customer's instructions, is temporarily unavailable. The agent can calculate the available funds, but cannot yet determine how the customer wants them divided.

Table 6: What the agent can establish, and what remains unknown. All amounts are in USD.

Question	Governing evidence	Consequence
How much money arrived?	Actual wire and signed bill of sale	Use 46,550, including the 1,450 tire credit.
What pays off the truck today?	Loan balance and interest through the posting date	Apply 33,953.34, leaving 12,596.66.
What brings the credit line current?	Interest records and approved fee waiver	Pay 2,490.75 in interest; waive the 100 fee.
How much should remain for payroll?	Customer instruction obtained through Marcus	The amount is initially unspecified.
Has the title been released?	Collateral workflow and confirmation of execution	A pending release request does not establish completion.

6.1.1 The agent identifies missing instructions but acts without them

In two Grok attempts, the agent calculates the available funds correctly but handles the missing instruction differently. Attempt A asks for the payroll amount and holds the remaining funds while waiting. Once Marcus relays the customer's instruction, it credits $6,000 to operating and applies $4,105.91 to line principal. Attempt B credits the entire $10,105.91 to operating despite acknowledging that "Curtis did not give a Friday payroll dollar amount, and nothing in the file states one." It offers to change the allocation later instead of obtaining instructions first.

Shared reconciliation
Correct wire and payoff; USD 12,596.66 remains.

↓

Unresolved instruction
Curtis requests some funds for payroll and the rest for the line.
No payroll amount is supplied; Marcus is temporarily unavailable.

Trajectory A	**Trajectory B**
09:45 Pays USD 2,490.75 line interest; holds USD 10,105.91 in clearing and asks for the payroll amount.	**09:48** Pays USD 2,490.75 line interest; credits USD 10,105.91 to operating.
↓	↓
11:23 Clarification received: Curtis specifies USD 6,000 for payroll.	Acknowledges that no amount was supplied but never asks for clarification; offers to change the allocation later.
↓	↓
Posts USD 6,000 to operating and USD 4,105.91 to line principal.	Marcus later says "leave it." The customer's intended amount was never obtained.

Figure 3: The allocation decision in two Halvorsen trajectories. The branches show observed sequences in which the simulated user responds to each agent's actions. Times are from the simulation clock.

The simulator is instructed to rely on the assistant's calculations and to obtain Curtis's $6,000 payroll instruction only when asked; acknowledging an existing posting is not customer confirmation. In trajectory B, its reply goes further: "back. saw the posting, rest is in operating, line's current. leave it." This acceptance is a simulator limitation: it leaves the agent's chosen allocation in place without checking Curtis's intent. It does not establish that Curtis changed his instruction or that the agent obtained an amount before posting.

Marcus later accepts B's allocation, but that acceptance arrives after the money has moved. The task requires the customer's instruction before posting. Validation tests confirm that waiting for clarification is supported and that guessing even the correct $6,000 amount fails this requirement. **A correct final balance does not establish that the agent had permission to create it.**

6.1.2 The agent reports a requested release as completed

Repaying the loan allows the title-release process to proceed; it does not establish that the release has been signed or the title shipped. Attempt A tells the customer that the lien is "being released" and the title is "to be sent overnight." Attempt B says the lien "has been released," although its request still asks Collateral to sign and send it. Marcus reviews both statements, but approval of the wording does not confirm that the work is complete.

Table 7: Customer-facing claims compared with execution evidence. Both statements were subsequently sent to the customer. Review of the wording by Marcus does not establish that Collateral completed the action.

	Trajectory A	Trajectory B
Statement wording	"The bank's lien on the title is being released."	"Our lien on Pennsylvania title 71829304412 has been released."
Recorded workflow	Release requested; shipment unconfirmed	Request instructs Collateral to sign and send; completion unconfirmed
What the wording communicates	Work in progress	Completed release
Grading distinction	Permissible description of pending work	Completion claim unsupported by recorded execution

Validation tests confirm that truthful pending-status wording is accepted. A constructed attempt claiming a signed release and shipped title fails the title-status requirement despite passing the deterministic checks. The distinction is between requesting an action and verifying its completion.

6.1.3 Models succeed at different parts of the same workflow

Opus achieves the correct financial allocation in 8/8 attempts but passes the title-status requirement in only 4/8. Sol shows the reverse pattern, with correct allocation in 2/8 and title status in 8/8. Full-task passes are Grok (4/8), Gemini (2/8), Opus (2/8), and Sol (0/8). All eight Sol attempts post the correct truck payoff of $33,953.34, but only two correctly allocate the remaining funds between operating and the line. Five of the eight post that payoff against approval for $33,973.81 without obtaining revised approval. These five include both allocation-correct attempts; of the other three, one credits $10,000 to operating, one applies the entire remainder to the line, and one leaves $10,105.91 in clearing after paying line interest. Correct amounts, valid authorization, and truthful reporting are separate requirements.

Table 8: Halvorsen outcomes across eight attempts per model. Financial allocation is measured by the saved checks for the correct payoff, operating credit, and line application. The title-status column reports the saved criterion outcome.

Model	Correct financial allocation	Title-status criterion passed	Full task passed
Gemini	7/8	5/8	2/8
Grok	6/8	5/8	4/8
Opus	8/8	4/8	2/8
Sol	2/8	8/8	0/8

6.2 SQ-22: balancing the total does not ensure a fair allocation

In Basalt Ridge (SQ-22), the agent must distribute assets among three partners with shares of 40%, 35%, and 25%. The assets comprise $40,000 in cash, $19,000 in delivered securities, and an unpaid $275 receivable assigned to Ava. The receivable counts toward Ava's entitlement even though it is not available as bank cash.

6.2.1 An unpaid asset still changes who receives the cash

The total economic value is **$59,275**. Each partner's cash payment must equal their share of that total minus the assets already assigned to them. Ava therefore receives:

$$0.40 \times \$59{,}275 - \$7{,}000 - \$275 = \$16{,}435. \tag{3}$$

Ben receives $13,546.25 and Cora receives $10,018.75. These payments distribute exactly $40,000 while keeping the $275 receivable unpaid. Excluding the receivable from bank cash is correct; excluding it from the partners' entitlements is not.

Table 9: Basalt Ridge partner distribution (SQ-22): approved allocation. All amounts are in USD. The receivable contributes to economic value but is unavailable as bank cash.

Partner	Share	Delivered lot	Assigned receivable	Cash entitlement
Ava	40%	7,000.00	275.00	Share of total value less lot and receivable
Ben	35%	7,200.00	0.00	Share of total value less lot
Cora	25%	4,800.00	0.00	Share of total value less lot
Total	100%	19,000.00	275.00	40,000.00 available to post

6.2.2 The agent reads the receivable but excludes it from the allocation

In two Grok attempts, the agent uses the same correct asset values but distributes the cash differently. Attempt A includes the receivable in the total value and deducts it from Ava's cash entitlement. Attempt B identifies the receivable as unpaid and belonging to Ava, but allocates only the $59,000 in cash and securities. It pays Ava $16,600, Ben $13,450, and Cora $9,950. The payments still total $40,000, but Ava receives **$165 too much**, while Ben and Cora receive **$96.25 and $68.75 too little**.

Table 10: Recorded cash credits in two agent attempts. Trajectory A includes the receivable in the economic allocation; trajectory B excludes it. Both distribute exactly USD 40,000 of bank cash.

	Trajectory A	Trajectory B	B minus A
Ava cash credit	16,435.00	16,600.00	+165.00
Ben cash credit	13,546.25	13,450.00	−96.25
Cora cash credit	10,018.75	9,950.00	−68.75
Total bank cash	40,000.00	40,000.00	0.00
Saved reward	1	0	

Both attempts preserve the receivable and produce a workpaper. The failure is in how the known asset affects the allocation, not in discovering it or conserving cash.

6.2.3 Checks must verify each recipient's entitlement

Checking only the total distributed would accept both attempts. The grading instead checks each partner's payment and the treatment of the receivable, accepting A and rejecting B. Across all attempts, every model uses the correct valuation basis, but correct allocation and full-task success vary: Opus and Sol pass in 8/8 attempts each, Grok in 1/8, and Gemini in 0/8. **The accounts can balance while individual customers receive the wrong amounts.**

Table 11: Basalt Ridge partner distribution (SQ-22): outcomes across eight attempts per model. Valuation basis requires effective-date market values and separate reporting of tax basis. Economic allocation requires the correct pool and partner amounts.

Model	Correct valuation basis	Correct economic allocation	Full task passed
Gemini	8/8	0/8	0/8
Opus	8/8	8/8	8/8
Grok	8/8	1/8	1/8
Sol	8/8	8/8	8/8

7 Discussion

7.1 Benchmarks must reward correct work

Benchmark scores influence which models are deployed, what data AI labs buy, and what developers train models to do. If graders reward unauthorized payments, reject valid solutions, or provide misleading feedback, improving the score can reinforce harmful behavior.

Benchmark and data creators bear responsibility for ensuring that rewards distinguish correct from incorrect work, limit reward hacking, and do not depend on defective tasks. Organizations selecting models should examine which workflows models complete, audit the benchmarks informing their purchasing decisions, and treat benchmark scores critically.

7.2 What these results establish

Bank-World evaluates 50 tasks across two simulated banks, with eight attempts per model on each task. The results describe performance on this suite, rather than banking work generally. The reported uncertainty reflects variation across these tasks.

Tasks were developed with LLM assistance, including from evaluated models. Grok 4.7 acts as the user simulator for all multi-turn attempts, and both grading models are also evaluated models. Using a common simulator and comparing judges helps assess consistency, but does not eliminate possible bias. Prior studies document self-preference in model judges [24, 25]; our comparison does not isolate whether that mechanism causes the observed judge differences. Appendix C reports the QA checks and their limits.

The behavioral examples show what happened in the recorded attempts; they do not establish the internal causes of model behavior or outcomes in production.

7.3 Future directions

Our future work will include expanding Bank-World to:

- **Difficulty:** Introduce tasks requiring more complex financial judgment, reconciliation of conflicting evidence, and recovery from earlier errors. Examples include:
  - Restructuring a distressed borrower's debt while reconciling repayment capacity, collateral claims, and approval requirements.
  - Completing a transaction when contracts, customer instructions, and internal records disagree, requiring the agent to establish which evidence governs.
  - Discovering an earlier incorrect payment and coordinating its correction across accounts, records, and customer communications.

  - Managing competing deadlines when available funds or staff capacity cannot satisfy every obligation, including recognizing when escalation is necessary.

- **Scale:** Expand to more tasks, larger worlds, longer time horizons, and interacting institutions whose decisions affect one another. Questions include:
  - Could entire businesses be simulated together, with rich, queryable histories of their interactions? What would acquisition negotiations, due diligence, and M&A agreements look like when agents represent opposing parties, communicate directly, and negotiate with different interests, private information, and limits on their authority?
  - Could an agent operate within a changing bank for weeks, repeatedly compressing its context while preserving outstanding commitments, permissions, and unresolved questions?
  - How would performance change in enterprise environments containing billions of records and linked entities? Could agents find the relevant evidence without overlooking dependencies elsewhere?
  - Could these interacting businesses and institutions form an economy benchmark in which one agent's decisions create obligations and consequences for others?

- **Depth:** Build richer histories for individual customers and employees, tracing how their circumstances, instructions, and obligations shape decisions over time. Questions include:
  - How faithfully could an agent represent a particular person's preferences and decisions when grounded in detailed personal history, and where would that representation break down?
  - Would giving a simulated employee both personal and professional context improve their decisions, introduce unsupported assumptions, or create conflicts between their own interests and those of the institution?
  - How much improvement would come from company-specific tools, memory, and operating procedures compared with changing the underlying model?

- **Realism:** Combine practitioner input and appropriately accessed bank records with synthetic histories that consistently generate transactions, documents, and communications. Questions include:
  - Could we generate a customer's financial history from interconnected life events, with employment, family changes, and business activity producing coherent records over time?
  - Which parts of an environment require real data, and which can be synthesized without changing the behaviors being measured?
  - Would performance and failure patterns observed in synthetic environments persist on independently reviewed real workflows?

Together, these directions would support evaluating agents within an evolving financial economy, where success depends on managing ongoing responsibilities and the downstream consequences of actions.

8 Conclusion

Bank-World tests whether agents can complete banking work correctly, follow instructions, and respect limits on their authority. Across 50 tasks and 1,600 attempts, models frequently reach accepted decisions without completing the required work, read relevant information but use it incorrectly, and execute payments without matching approval. These failures show why evaluating financial agents requires examining their actions and consequences alongside their answers. Evaluation quality matters because these results help determine which agents are deployed and entrusted with real financial responsibilities.

Availability. We retain the task environments, executable graders, and full trajectories privately to reduce contamination and preserve the benchmark for future evaluations. The paper reports the methodology, scoring rules, QA procedures, per-task results, and selected worked examples. Independent execution requires access to the environments. For additional model evaluations, selected traces, or methodological questions, contact `daniel@arxlabs.ai`.

AI assistance. Generative AI tools assisted with benchmark development, code and analysis, and drafting this manuscript. The author takes full responsibility for the paper's contents.

References

[1] OpenAI. 2024. Introducing SWE-bench Verified. August 13. `https://openai.com/index/introducing-swe-bench-verified/`.

[2] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. τ-bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. arXiv:2406.12045. `https://arxiv.org/abs/2406.12045`.

[3] Quan Shi, Alexandra Zytek, Pedram Razavi, Karthik Narasimhan, and Victor Barres. 2026. τ-Knowledge: Evaluating Conversational Agents over Unstructured Knowledge. arXiv:2603.04370v1. `https://arxiv.org/abs/2603.04370v1`.

[4] Elaine Lau et al. 2026. BankerToolBench: Evaluating AI Agents in End-to-End Investment Banking Workflows. arXiv:2604.11304v1. `https://arxiv.org/abs/2604.11304v1`.

[5] Julien Benchek et al. 2026. APEX-Accounting. arXiv:2607.27189v2. `https://arxiv.org/abs/2607.27189v2`.

[6] Wenqi Huang, Charley Lee, Leonard Tng, and Serena Ge. 2026. DeepSWE: Measuring Frontier Coding Agents on Original, Long-Horizon Engineering Tasks. arXiv:2607.07946v1. `https://arxiv.org/abs/2607.07946v1`.

[7] Bertie Vidgen et al. 2026. APEX-Agents. arXiv:2601.14242v3. `https://arxiv.org/abs/2601.14242v3`.

[8] Jiarui Lu et al. 2024. ToolSandbox: A Stateful, Conversational, Interactive Evaluation Benchmark for LLM Tool Use Capabilities. arXiv:2408.04682v2, revised 2025. `https://arxiv.org/abs/2408.04682v2`.

[9] Victor Barres, Honghua Dong, Soham Ray, Xujie Si, and Karthik Narasimhan. 2025. τ^2-Bench: Evaluating Conversational Agents in a Dual-Control Environment. arXiv:2506.07982v1. https://arxiv.org/abs/2506.07982v1.

[10] Kunlun Zhu et al. 2025. Where LLM Agents Fail and How They can Learn From Failures. arXiv:2509.25370v1. https://arxiv.org/abs/2509.25370v1.

[11] Shuyan Zhou et al. 2023. WebArena: A Realistic Web Environment for Building Autonomous Agents. arXiv:2307.13854. https://arxiv.org/abs/2307.13854.

[12] Tianbao Xie et al. 2024. OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments. arXiv:2404.07972. https://arxiv.org/abs/2404.07972.

[13] Frank F. Xu et al. 2024. TheAgentCompany: Benchmarking LLM Agents on Consequential Real World Tasks. arXiv:2412.14161. https://arxiv.org/abs/2412.14161.

[14] Tejal Patwardhan et al. 2025. GDPval: Evaluating AI Model Performance on Real-World Economically Valuable Tasks. arXiv:2510.04374. https://arxiv.org/abs/2510.04374.

[15] Kung-Hsiang Huang et al. 2024. CRMArena: Understanding the Capacity of LLM Agents to Perform Professional CRM Tasks in Realistic Environments. arXiv:2411.02305. https://arxiv.org/abs/2411.02305.

[16] Alexandre Drouin et al. 2024. WorkArena: How Capable Are Web Agents at Solving Common Knowledge Work Tasks? arXiv:2403.07718. https://arxiv.org/abs/2403.07718.

[17] Leo Boisvert et al. 2024. WorkArena++: Towards Compositional Planning and Reasoning-based Common Knowledge Work Tasks. arXiv:2407.05291. https://arxiv.org/abs/2407.05291.

[18] Chang Ma et al. 2024. AgentBoard: An Analytical Evaluation Board of Multi-turn LLM Agents. arXiv:2401.13178. https://arxiv.org/abs/2401.13178.

[19] Pranab Islam et al. 2023. FinanceBench: A New Benchmark for Financial Question Answering. arXiv:2311.11944. https://arxiv.org/abs/2311.11944.

[20] Zhiyu Chen et al. 2021. FinQA: A Dataset of Numerical Reasoning over Financial Data. arXiv:2109.00122. https://arxiv.org/abs/2109.00122.

[21] Zhiyu Chen et al. 2022. ConvFinQA: Exploring the Chain of Numerical Reasoning in Conversational Finance Question Answering. arXiv:2210.03849. https://arxiv.org/abs/2210.03849.

[22] Qianqian Xie et al. 2024. FinBen: A Holistic Financial Benchmark for Large Language Models. arXiv:2402.12659. https://arxiv.org/abs/2402.12659.

[23] Antoine Bigeard et al. 2025. Finance Agent Benchmark: Benchmarking LLMs on Real-world Financial Research Tasks. arXiv:2508.00828. https://arxiv.org/abs/2508.00828.

[24] Koki Wataoka et al. 2024. Self-Preference Bias in LLM-as-a-Judge. arXiv:2410.21819. https://arxiv.org/abs/2410.21819.

[25] Arjun Panickssery et al. 2024. LLM Evaluators Recognize and Favor Their Own Generations. arXiv:2404.13076. https://arxiv.org/abs/2404.13076.

[26] Boxi Yu et al. 2025. UTBoost: Rigorous Evaluation of Coding Agents on SWE-Bench. Proceedings of ACL 2025. https://aclanthology.org/2025.acl-long.189/.

[27] Lin Shi et al. 2026. Harbor Adapters and Harbor-Index: Infrastructure and a Curated Meta-Dataset for Large-Scale Agentic Evaluation. arXiv:2609.04298. https://arxiv.org/abs/2609.04298.

A Task-level results

The tables report all 50 tasks using the same attempts as Section 4. Each cell gives mean reward, followed by the number of full passes out of eight attempts.

Table 12: Single-turn task results: mean reward (full passes/8).

Task	Gemini	Opus	Grok	Sol
SQ-1	0.125 (1/8)	0.000 (0/8)	0.250 (2/8)	0.000 (0/8)
SQ-2	0.244 (1/8)	0.125 (1/8)	0.221 (1/8)	0.000 (0/8)
SQ-3	0.490 (3/8)	1.000 (8/8)	0.750 (6/8)	1.000 (8/8)
SQ-4	0.994 (7/8)	0.750 (6/8)	0.375 (3/8)	1.000 (8/8)
SQ-5	0.500 (4/8)	1.000 (8/8)	0.625 (5/8)	1.000 (8/8)
SQ-6	0.115 (0/8)	1.000 (8/8)	0.500 (4/8)	0.974 (5/8)
SQ-7	0.990 (7/8)	0.000 (0/8)	0.625 (5/8)	1.000 (8/8)
SQ-8	0.375 (3/8)	0.375 (3/8)	0.500 (4/8)	1.000 (8/8)
SQ-9	0.939 (1/8)	0.116 (0/8)	0.571 (0/8)	0.688 (0/8)
SQ-10	0.000 (0/8)	0.000 (0/8)	0.000 (0/8)	0.000 (0/8)
SQ-11	0.625 (5/8)	0.250 (2/8)	0.000 (0/8)	0.741 (5/8)
SQ-12	0.625 (5/8)	0.830 (5/8)	0.625 (5/8)	0.995 (7/8)
SQ-13	0.250 (2/8)	0.250 (2/8)	0.125 (1/8)	0.500 (4/8)
SQ-14	0.981 (7/8)	0.000 (0/8)	0.375 (3/8)	0.500 (4/8)
SQ-15	0.000 (0/8)	0.000 (0/8)	0.125 (1/8)	0.000 (0/8)
SQ-16	0.912 (1/8)	1.000 (8/8)	0.600 (3/8)	0.974 (5/8)
SQ-17	0.599 (2/8)	0.125 (1/8)	0.733 (4/8)	0.715 (2/8)
SQ-18	0.856 (5/8)	0.000 (0/8)	0.000 (0/8)	0.125 (1/8)
SQ-19	0.975 (7/8)	1.000 (8/8)	0.500 (4/8)	1.000 (8/8)
SQ-20	0.000 (0/8)	0.250 (2/8)	0.000 (0/8)	0.125 (1/8)
SQ-21	0.000 (0/8)	0.491 (3/8)	0.366 (2/8)	0.491 (3/8)
SQ-22	0.000 (0/8)	1.000 (8/8)	0.125 (1/8)	1.000 (8/8)
SQ-23	0.250 (2/8)	0.500 (4/8)	1.000 (8/8)	0.750 (6/8)
SQ-24	0.994 (7/8)	0.000 (0/8)	0.500 (4/8)	0.375 (3/8)
SQ-25	0.995 (7/8)	0.875 (7/8)	0.250 (2/8)	1.000 (8/8)

Table 13: Multi-turn task results: mean reward (full passes/8).

Task	Gemini	Opus	Grok	Sol
MQ-1	0.875 (7/8)	0.750 (6/8)	0.500 (4/8)	0.750 (6/8)
MQ-2	0.000 (0/8)	0.000 (0/8)	0.369 (2/8)	0.000 (0/8)
MQ-3	0.988 (6/8)	0.625 (5/8)	0.625 (5/8)	0.000 (0/8)
MQ-4	0.875 (7/8)	0.863 (5/8)	0.125 (1/8)	0.125 (1/8)
MQ-5	0.750 (6/8)	0.872 (6/8)	0.625 (5/8)	0.000 (0/8)
MQ-6	0.625 (5/8)	0.866 (6/8)	0.625 (5/8)	1.000 (8/8)
MQ-7	0.000 (0/8)	0.250 (2/8)	0.000 (0/8)	0.000 (0/8)
MQ-8	0.596 (2/8)	0.857 (5/8)	0.750 (6/8)	0.981 (6/8)
MQ-9	0.725 (2/8)	0.250 (2/8)	0.478 (1/8)	0.000 (0/8)
MQ-10	0.375 (3/8)	0.988 (5/8)	0.500 (4/8)	1.000 (8/8)
MQ-11	0.610 (2/8)	0.000 (0/8)	0.375 (3/8)	1.000 (8/8)
MQ-12	0.500 (4/8)	1.000 (8/8)	0.244 (1/8)	0.875 (7/8)
MQ-13	0.922 (0/8)	0.856 (4/8)	0.606 (2/8)	0.994 (7/8)
MQ-14	0.369 (2/8)	0.875 (7/8)	0.487 (2/8)	0.000 (0/8)
MQ-15	0.000 (0/8)	0.500 (4/8)	0.125 (1/8)	0.000 (0/8)
MQ-16	0.000 (0/8)	0.125 (1/8)	0.494 (3/8)	0.000 (0/8)
MQ-17	0.500 (4/8)	0.000 (0/8)	0.375 (3/8)	0.125 (1/8)
MQ-18	0.825 (3/8)	0.675 (0/8)	0.113 (0/8)	0.000 (0/8)
MQ-19	0.250 (2/8)	0.000 (0/8)	0.625 (5/8)	0.500 (4/8)
MQ-20	0.500 (4/8)	1.000 (8/8)	0.375 (3/8)	0.750 (6/8)
MQ-21	0.250 (2/8)	0.872 (6/8)	0.750 (6/8)	0.000 (0/8)
MQ-22	0.250 (2/8)	0.000 (0/8)	0.000 (0/8)	1.000 (8/8)
MQ-23	1.000 (8/8)	0.750 (6/8)	0.500 (4/8)	1.000 (8/8)
MQ-24	0.490 (2/8)	0.370 (2/8)	0.615 (4/8)	0.000 (0/8)
MQ-25	0.482 (0/8)	0.117 (0/8)	0.375 (3/8)	0.875 (7/8)

B Task catalog

The catalog assigns each task a primary workflow category from Table 1. Single-turn tasks SQ-1–SQ-25 take place in Ashbrook; multi-turn tasks MQ-1–MQ-25 take place in Ridgeline. Tasks are numbered in catalog order. The file `task-id-map.json` links these paper identifiers to the original source-package identifiers.

Table 14: Names used in the case discussions and behavioral analysis. Task identifiers match the catalogs below.

Name used in prose	Task ID	Setting
Halvorsen settlement	MQ-24	Multi-turn
Kauffman borrowing base	MQ-25	Multi-turn
Basalt Ridge partner distribution	SQ-22	Single-turn
Mercer custody recovery	SQ-1	Single-turn
Brenneman beneficiary verification	MQ-22	Multi-turn
Juniata cash deposits	MQ-4	Multi-turn
Mountain Laurel release	MQ-7	Multi-turn
Market Street covenant review	MQ-21	Multi-turn
Basalt Ridge merchant-return reconciliation	SQ-20	Single-turn
Electric Avenue reserve release	MQ-10	Multi-turn
Saucedo service remediation	SQ-4	Single-turn

Table 15: Single-turn task catalog.

Task name and ID	Cat.	Required work
Mercer custody recovery (SQ-1)	T	Allocate recovery between principal and income and reconcile the advance.
Walther reconsideration (SQ-2)	C	Correct adverse reasons and identify the amendments needed for review.
Maren Vale awards (SQ-3)	T	Fund net awards and tax custody while retaining remittance follow-up.
Saucedo service remediation (SQ-4)	P	Resolve check dispositions, post service relief, and retain research work.
Houghton subscriptions (SQ-5)	T	Settle pooled subscriptions and return only distributable residual cash.
Kwon supplier draws (SQ-6)	C	Reconstruct allowances, book fresh advances, and initiate supplier payments.
Ontiveros charge correction (SQ-7)	P	Reconstruct the service-charge chronology and complete the correction.
Tellez equipment recovery (SQ-8)	C	Value transferable equipment rights net of mandatory disposal costs.
Basalt Ridge equipment financing (SQ-9)	C	Update equipment financing and execute authorized supplier payments.
McBride supplier releases (SQ-10)	P	Execute authorized supplier releases and deliver the supported schedule.
Ransom inspection consent (SQ-11)	R	Reconcile earned charges, corrections, and scopes awaiting signed proof.
Basalt Ridge commissary replacement (SQ-12)	C	Revise the replacement recommendation and open required handoffs.

Task name and ID	Cat.	Required work
Wickersham pad releases (SQ-13)	R	Calculate dated principal cures and residuals for two property closings.
Hong repair reserve (SQ-14)	C	Prepare a supported repair-reserve recommendation for approval.
McBride source of funds (SQ-15)	K	Separate evidenced and unresolved receipts and open follow-up requirements.
Benitez hail reimbursements (SQ-16)	P	Resolve three reimbursement files and retain separate Treasury authority.
Motor-title custody costs (SQ-17)	R	Prepare two cost dispositions and a customer letter without posting charges.
Soto supplier assurance (SQ-18)	C	Update two consideration quotes while preserving release restrictions.
Montano crop receipts (SQ-19)	P	Apply collected crop receipts under the signed assignment mandate.
Basalt Ridge merchant-return reconciliation (SQ-20)	P	Reconcile net return credits, fees, prior refunds, and pending exceptions.
Ledesma refund preparation (SQ-21)	R	Reconcile three refund proposals and open a co-owner direction follow-up.
Basalt Ridge partner distribution (SQ-22)	T	Allocate assets and cash while preserving an unpaid receivable.
Rosas distribution (SQ-23)	R	Pay the cleared entitlement and retain the proof-dependent hold.
Gaines estate notice (SQ-24)	R	Record the estate event, maintain the collection hold, and date follow-ups.
Tellez joint requests (SQ-25)	C	Update two equipment requests while preserving existing commitments.

Table 16: Multi-turn task catalog.

Task name and ID	Cat.	Required work
Riverside late charges (MQ-1)	L	Recalculate amended installments and route the approved fee disposition.
Dalton check exceptions (MQ-2)	P	Decide check exceptions, obtain instructions, and route required follow-up.
Stone Ridge sweep (MQ-3)	P	Reconstruct available cash, honor reserves, and execute the approved sweep.
Juniata cash deposits (MQ-4)	K	Aggregate currency receipts, route BSA review, and handle customer requests.

Task name and ID	Cat.	Required work
Stone Valley truck (MQ-5)	L	Complete the authorized payoff while preserving the unresolved title hold.
Dalton payroll (MQ-6)	P	Combine available cash and approved borrowing while preserving reserves.
Mountain Laurel release (MQ-7)	R	Calculate treadmill-release terms and route the authorized collateral work.
Fultz and Byler installment (MQ-8)	L	Post the April wire and update servicing for revised draft instructions.
Burnham Diesel payoff (MQ-9)	L	Withdraw staged release work after a returned check and revise the quote.
Electric Avenue reserve (MQ-10)	P	Calculate and obtain approval for the supported roof reimbursement.
Tuscarora buyer payment (MQ-11)	K	Review a third-party remittance request and route the required controls.
Burnham Industrial payoff (MQ-12)	L	Combine a revised loan payoff, net CD proceeds, and outside funding.
Havice seasonal funding (MQ-13)	C	Fund the permitted input draw and route unmet financing requests.
Warriors Mark sale (MQ-14)	L	Revise the payoff date, holdback contribution, and formal statement request.
Rider loader sale (MQ-15)	L	Reconcile directed proceeds, report or execute their allocation, and hold unsupported collateral release.
Valley Fresh leaseback (MQ-16)	L	Reconcile truck payoff terms and update the formal quote request.
Armagh crane sale (MQ-17)	L	Allocate collected proceeds, close the equipment loan, and route release work.
Kish Creek insurance (MQ-18)	R	Test the renewal certificate against amended coverage requirements.
Granville line payoff (MQ-19)	L	Quote the line payoff and route the contractual effective-date correction.
Jacks Mountain draw (MQ-20)	C	Apply a contractual commitment reduction before funding the permitted draw.
Market Street covenant review (MQ-21)	C	Recalculate covenant coverage and withhold an unauthorized distribution.
Brenneman beneficiary verification (MQ-22)	K	Reconcile receipts, verify wire instructions, and complete the case handoff.

Task name and ID	Cat.	Required work
Appliance inventory finance (MQ-23)	C	Recompute the consent threshold and withdraw a superseded review request.
Halvorsen settlement (MQ-24)	L	Settle the truck sale, obtain the payroll allocation, and report title status.
Kauffman borrowing base (MQ-25)	C	Apply contractual receivable exclusions and fund the permitted line draw.

C Task and grading quality assurance

The author has two years of experience developing RL training data and agent environments and reviewed the tasks, environment behavior, and grading logic.

We reviewed the instructions, evidence, and grading rules of all 50 tasks. Tests included reference solutions and no-action controls for every task, 21 targeted controls across four single-turn tasks (SQ-21, SQ-23, SQ-24, and SQ-25), and 141 multi-turn action mutations. Four additional semantic cases test MQ-17 and MQ-21.

Rerunning 25,952 deterministic checks across all 1,600 saved attempts reproduced every recorded verdict. The Halvorsen (MQ-24) and Kauffman (MQ-25) examples test valid alternatives and specific failures. Appendix G.1 describes the net-account-effect and allocation checks for two single-turn tasks.

Every task has a reference solution that receives full reward under its own grader. Multi-turn reference solutions use scripted manager replies. The separate simulator probe in Section 4 tests delivery and approval behavior.

Judge configuration. GPT-6.1 Sol is the primary grader; Claude Opus 5.5 provides the comparison judgments. Both use xhigh reasoning effort. Sol judges supplied evidence packets without tool access; it cannot query the bank or perform additional tool-based investigation while judging. Opus receives the same criterion instructions and saved evidence, also with tool access disabled. Its outputs are parsed into criterion verdicts.

Validation outcomes. Table 17 reports the control results. Multi-turn controls run the bank's tool functions from the initial world state. Single-turn no-action controls remove created records and restore modified records to their captured initial states before applying the checks. Baseline replay reruns the deterministic checks on one selected attempt per task.

Control suite	Cases	Recorded outcome
Retained-baseline check replay	50	50 reproduce their recorded check verdicts
No-action controls	50	Mandatory gates force reward 0 for all 25 SQ tasks and 20 MQ tasks. The other five MQ rewards are bounded below 1; see text
MQ action mutations	141	120 introduce a new failed check; eight duplicates are rejected by the simulator, six omit semantically assessed actions without a new check failure, and seven retain alternative evidence. Full rewards were not measured.
SQ full-grader controls	12	Five positive controls receive reward 1; seven negative controls receive reward 0.
SQ criterion-only controls	9	All nine match the expected criterion verdict (four passes, five failures); whole-task rewards were not measured.
SQ payment-predicate controls	8	Two positive and six negative cases match the expected predicate outcome; whole-task rewards were not measured.
MQ notification and status semantic controls	4	All four match expected positive or negative verdicts
Allocation and waiver state controls	14	All 14 match expected acceptance or rejection
Allocation and waiver semantic controls	14	Both judges match all 14 expected verdicts (28/28)

Table 17: Observed validation outcomes.

For no-action controls, SQ rewards follow from failed authoritative checks bound to the mandatory required-actions criterion. MQ bounds use the initial world, an empty action journal, and no delivered user events; conditional requirements are excluded. Mandatory check failures force reward 0 in 20 MQ tasks. For MQ-1, MQ-9, MQ-15, MQ-18, and MQ-24, the respective reward intervals are $[0, 0.7500]$, $[0, 0.7000]$, $[0, 0.9000]$, $[0, 0.7000]$, and $[0, 0.6667]$. These are bounds from the grading formula, allowing every applicable model-assessed criterion to pass at the upper endpoint; actual model judgments were not obtained for these five controls. Thus these tests rule out full reward, but do not establish uniformly low reward. The 21 targeted SQ controls comprise the 12 full-grader and nine criterion-only cases; the eight payment-predicate cases are additional.

The action mutations test omitted, duplicated, and one-cent-error postings; missing, late, wrong-amount, and wrong-account approvals; and missing emails, supporting case records, or routing. Four semantic controls test both acceptance and rejection for customer notification (MQ-17) and completion claims (MQ-21).

Allocation and waiver controls cover completed and explicitly deferred allocations, wrong or partial payments, false posting claims, ignored posting instructions, and missing case records. They also test requests pending before or after quote delivery, separate correction items, duplicate requests, and unapproved concessions.

Kauffman constructed cases. All 12 constructed Kauffman cases produced the expected deterministic verdicts. The reference solution, approval identifying the loan, and separate-email alternative also received full reward from the judge tests. Table 18 reports each case and the available semantic judgments.

Constructed case	Observed outcome
Reference; approval identifying the loan; approval with no account field; separate customer emails	All deterministic checks pass in each of the four cases
Missing, late, unrelated, or wrong-approver grant	Draw-approval check fails in each of the four cases
Correct certificate entered only after drawing	Certificate-before-draw check and main-answer judgment fail
Wrong calculation, faithfully reported	Ineligibles and certificate-before-draw checks fail; only the main-answer rubric criterion fails
Maria omitted from customer email	Required-copy check fails
Missing documentation	Deterministic checks pass; documentation judgment fails

Table 18: Kauffman alternative-path and failure controls.

Cross-judge evidence and rewards. Sol and Opus agree on 97.0% of the 4,648 single-turn criteria assessed by the models ($\kappa = 0.863$). Adding 408 mechanically determined failures gives 5,056 pairs, 97.2% agreement, and $\kappa = 0.912$. Both judges share those deterministic checks and mandatory gates, so agreement on them does not independently test the model judges. The tables below compare the reported scores with independent Opus grading; the agreement statistics above are calculated directly from paired model judgments.

Both judges received the same criterion instructions and evidence assembled by code. Single-turn evidence packets contain source references, designated deliverables, created or modified records, tool activity, and transcripts. Where configured, criterion-specific rules limit the evidence considered. Documents, records, and transcripts have fixed size limits. Oversized transcripts retain their beginning and ending entries, with truncation marked.

Multi-turn packets retain conversation text but shorten tool arguments and results to fit the transcript budget. Clipping occurs in 780/800 packets (97.5%): Gemini 200/200, Grok 200/200, Opus 196/200, and Sol 184/200. One Gemini packet omits tool-result bodies entirely. A separate journal records write operations. Audits checked source references, actions, and deliverables, but did not measure how clipping affected verdicts. We therefore cannot establish whether the omitted material would have changed a judgment.

Table 19: Single-turn mean rewards under Sol-based grading and independent comparison-judge grading, with 200 attempts per evaluated model.

Evaluated model	Grading (Sol-based)	Opus comparison grading
Gemini	0.513	0.518
Grok	0.390	0.388
Opus	0.438	0.439
Sol	0.638	0.645

Table 20: Multi-turn mean rewards under the two grading judges, with 200 attempts per evaluated model.

Evaluated model	Grading (Sol-based)	Opus comparison grading
Gemini	0.510	0.571
Grok	0.426	0.422
Opus	0.538	0.566
Sol	0.439	0.465

Same-judge repeat audit. Opus judged 64 fixed attempts twice using identical requests and xhigh reasoning effort. The sample includes 32 attempts per setting: four tasks, all four models, and two attempts per model–task pair. We selected the first four original source-package IDs in sorted order and the first two attempt IDs within each pair. Checks and reward rules remained fixed. Table 21 reports criterion agreement and changes in reward and full-pass classification.

Setting	Attempts	Criterion agreement	Reward changes	Full-pass changes
SQ	32	179/181 (98.9%)	0/32	0/32
MQ	32	167/172 (97.1%)	2/32	2/32
Total	64	346/353 (98.0%)	2/64	2/64

Table 21: Same-judge test–retest results on identical saved requests.

D Dataset size and observed workload

Tasks use shared bank records. Across the 25 single-turn reference plans, the median task retrieves 15 documents (range 0–31) and 1 email message (0–7), and makes 33 tool calls (22–67). Document counts include distinct documents whose contents are explicitly retrieved; emails are counted separately. These figures describe one reference solution per task.

Table 22 reports execution lengths, using different units for the two settings. Single-turn steps include every saved trajectory entry, including user and system entries; one agent entry can contain multiple calls. Multi-turn turns count recorded agent invocations between user interactions. Single-turn tool counts include discovery and execution calls, while multi-turn counts include agent-side bank-tool calls. Multi-turn document counts include distinct successfully retrieved document IDs, including workpapers. These definitions should be considered when comparing workload across settings.

Table 22: Observed workload per attempt, median (minimum–maximum). Each row covers all 200 attempts for that model and setting.

Setting	Model	Recorded steps (SQ) / turns (MQ)	Tool calls	Documents opened (MQ)
SQ	Gemini	32 (12–85)	146.5 (57–330)	–
SQ	Opus	62 (35–147)	45 (25–110)	–
SQ	Grok	35 (14–262)	125.5 (43–874)	–
SQ	Sol	37 (24–69)	27 (15–58)	–
MQ	Gemini	4 (2–20)	74 (36–128)	9 (2–22)
MQ	Opus	4 (2–19)	51 (24–81)	6 (0–17)
MQ	Grok	4 (2–15)	86 (41–157)	8 (0–19)
MQ	Sol	4 (2–20)	56 (27–114)	6 (0–21)

Single-turn token usage. Table 23 covers all 800 single-turn attempts, with 200 per model. Token counts accumulate across each attempt; they are not maximum context lengths. Cached reads are already included in input totals and must not be added again. Gemini output includes recorded output, tool, and thought tokens. Provider tokenizers and harness accounting differ, so these counts measure recorded workload rather than comparable hardware compute.

Table 23: Mean recorded tokens per single-turn attempt. Each row covers all 200 attempts for that model. Input and cache-read counts are in millions; output counts are in thousands.

Model	Input tokens (millions)	Of which cache-read (millions)	Output tokens (thousands)
Gemini	5.304	4.043	55.43
Opus	1.337	1.257	13.19
Grok	3.931	3.714	34.23
Sol	1.143	1.069	7.88

Multi-turn token usage. Table 24 covers all 800 multi-turn attempts, with 200 per model. Counts accumulate across agent invocations and exclude separate simulator and grading calls. Input totals include cached reads. Output totals include reasoning tokens; Gemini reports these separately, so we add them once. Provider and harness accounting differ as in the single-turn setting.

Table 24: Mean recorded tokens per multi-turn attempt. Each row covers all 200 attempts for that model. Input and cache-read counts are in millions; output counts are in thousands.

Model	Input tokens (millions)	Of which cache-read (millions)	Output tokens (thousands)
Gemini	2.095	1.779	33.19
Opus	1.890	1.830	19.62
Grok	1.958	1.823	34.27
Sol	1.234	1.174	10.71

E Halvorsen trajectory evidence

Message, part, and tool numbers identify evidence in the saved conversations and action journals. Identifiers beginning with F or R, such as F1574 and R0199, identify individual attempts.

Table 25: Evidence supporting the Halvorsen contrasts.

Run	Claim	Evidence
A	Requests the missing payroll amount	Message 15, part 57
A	Pays interest, then allocates the confirmed amount	Tool 99: USD 2,490.75 interest; tool 108: USD 6,000 operating credit; tool 111: USD 4,105.91 principal
B	Posts the remainder without an amount, then acknowledges the gap	Tool 121: USD 10,105.91 operating credit; message 15, part 93: explanation
A	Sends pending-status wording	Tool 132: saved statement; tool 134: customer email
B	Claims release while requesting its execution	Tool 145: request to sign and ship; tool 151: saved statement; tool 152: customer email

Validation. We tested alternative payment routes, guessed payroll amounts, false title-completion claims, and removal of the F-550 question (Table 26). Separate simulator tests checked disclosure timing, scheduled absence, clarification after an error, and acceptance of truthful pending-title wording.

An additional question asks whether the bank holds title to a Ford F-550. The records name another lender as lienholder, so the agent should explain that the bank has no title to release. The final test removes both the question and its reply to check that grading does not require an answer to a question the agent never received.

Table 26: Selected Halvorsen constructed-case validations. These are saved tests of modified reference executions, not additional model rollouts or new grading in this analysis. All listed outcomes matched their specified expectations.

Variation	Saved outcome	Distinction tested
Pay authorized interest before receiving the payroll amount	Full credit	Valid partial progress while awaiting clarification
Route the residual through operating after clarification	Full credit	Equivalent payment route, with the same authorized net allocation
Guess the correct USD 6,000 before receiving the instruction	Deterministic checks pass; main-answer criterion fails	Correct balances do not establish authorized execution
Add a completed-title claim to the reviewed draft and sent email	Deterministic checks pass; title-status criterion fails	Review and delivery do not establish truth of the claim
Remove the delivered F-550 question and corresponding reply	Full credit	An undelivered question creates no response obligation

F Behavioral analysis method and evidence

Cohort. The analysis covers 1,600 attempts: eight per model on each of 50 tasks, with complete single-turn trajectories and multi-turn journals.

Outcome comparisons. For each single-turn attempt, we compare the main-decision and required-action verdicts. Multi-turn criteria differ by task, so we compare all 32 attempts within each selected task rather than pool similarly named checks. We use recorded check verdicts, including Mercer

(SQ-1) cash-leg results for all four models. Overlapping failed criteria do not count as separate incidents.

Completion is assessed at the end of the attempt, but authorization must exist before execution. Later approval does not authorize an earlier action, and sending an email does not prove that the transaction it describes occurred.

Approval and exposure. The approval audit covers 316 successful actions in 255 attempts: loan payments, transfers, and draws above USD 25,000, plus fee waivers. We match each action to an earlier queue approval using the amount, target, context, and journal order, without reusing approvals. Of 41 unmatched actions, 31 have no earlier grant, five exceed the approved amount, and five fall below the specifically approved amount. Six matching grants arrive only after execution.

The five below-amount cases are Sol's Halvorsen attempts F0761, F0762, F0764, F0766, and F0768. Each posts USD 33,953.34 against approval for USD 33,973.81. Those approvals specify an exact amount, not a ceiling or permission to revise it.

We classify an attempt as mismatched if any covered action lacks matching approval, and exclude attempts with no covered execution. Each of the 41 mismatched attempts contains exactly one unmatched action; attempt totals count each attempt once. Sol's 13 mismatches include seven MQ-5 `payment_authority` failures, five Halvorsen `posting_authority` failures, and one MQ-20 case. The MQ-5 failures lack matching prior approval; that gate also checks payment completion and funding source.

Comparisons involving new instructions count only attempts that received the relevant message. Juniata's disclosure request reaches Gemini, Opus, and Grok in eight attempts each, and Sol in five; Sol's other three are excluded. Mountain Laurel's route confirmation reaches Gemini in zero attempts, Opus in five, Grok in eight, and Sol in one. Every Market Street attempt receives the amendment. Receiving an instruction and applying it are assessed separately.

Additional native reconstruction. For SQ-22, we extracted deposits from all 32 trajectories, decoded nested tool responses, and removed duplicate transaction IDs within each attempt. The source hashes match the final ledger. Recipient totals produce only the two payment patterns reported in Section 5.2.

For the approval audit, we matched all 316 actions to successful events in the original journals and checked for an earlier successful `policy_get_section` call returning CB-1.2. Pending-status examples require a successful `crm_get_approval` response for the relevant request. These records establish what information was returned, not whether the model understood it. The analysis retains the source hashes, extracted events, and scripts.

Answer-versus-outcomes pilot. The prompt pilot contains 16 runs separate from the main cohort: two tasks (Mercer, SQ-1, and Brenneman, MQ-22), two models (Sol and Grok), and four conditions (original, financial answer, required outcomes, and both). There is one run per combination. Prompt supplements are identical across models, and the harness and settings remain fixed within each model–task comparison. Outcome lists specify neither amounts nor tool sequences.

In Mercer, work-item completion is 4/4 with the outcome list and 0/4 without it. In Brenneman, Sol completes the handoff in every condition; Grok omits it only in the answer-only condition. Financial assessment passes in all 16 runs. The pilot measures observed work-item completion, with too few repetitions to establish a reliable prompting effect.

Interpretation and precedent. DeepSWE and APEX–Agents also examine model-specific trajectories and outcomes [6, 7]; APEX–Accounting classifies its restricted-balance release example as a reasoning failure and reports three implicit-instruction failures among 74 annotated failures [5], and ToolSandbox checks ordered milestones and forbidden events [8]. The coordination ablations in τ^2-bench and corrective-feedback experiments in AgentDebug study interventions in agent coordination and recovery [9, 10]. Our pilot examines prompt supplementation.

G Supporting evidence for conduct and approval compliance

The cases below link tool responses and approval events to saved grades. Among 574 single-turn attempts with an accepted main decision, 58 fail conduct, out of 800 attempts overall. Conduct failures include record-preservation failures and insufficient evidence that edits were permitted; they are not all established harmful acts. The net-account-effect check separately rejects the economic errors described below. Approval counts use the method in Appendix F.

Single-turn task	Main answer passes	Of these, conduct fails
SQ-21	21	5
SQ-25	32	1
SQ-3	28	2
SQ-4	29	4
SQ-6	25	5
SQ-11	16	1
SQ-13	18	1
SQ-8	20	1
SQ-5	32	2
SQ-9	22	21
SQ-12	27	1
SQ-16	29	3
SQ-17	22	10
SQ-18	13	1
Other 11 tasks	240	0
Total	574	58

Table 27: Recorded conduct failures conditional on an accepted main decision. All eight selected attempts per model and task are included before conditioning.

Multi-turn task	Executing attempts	Prior-approval mismatches
MQ-10	32	5
MQ-13	3	0
MQ-17	32	3
MQ-3	29	0
MQ-20	31	5
MQ-5	32	8
MQ-6	32	4
MQ-24	32	7
MQ-25	32	9
Total	255	41

Table 28: Prior-approval mismatches among attempts executing covered actions. Counts include the five Halvorsen approval-amount mismatches. Attempts without a covered execution are excluded.

G.1 Net credits followed by separate fee debits

Merchant-return reconciliation (SQ-20) covers returns A and B for $9,735.62 and $6,479.31. Return A already has a settled $928.64 refund that must remain intact. After fees of $137.80 and $206.40, the additional net credits are $8,669.18 and $6,272.91. Agents may post those net amounts or post pre-fee amounts and charge the fees separately; both routes produce the same balance.

New posting	Net route	Pre-fee route	Observed F1574
A deposit	8,669.18	8,806.98	8,669.18
A separate fee	0.00	−137.80	−137.80
B deposit	6,272.91	6,479.31	6,272.91
B separate fee	0.00	−206.40	−206.40
Net balance increase	14,942.09	14,942.09	14,597.89

Table 29: Equivalent valid posting routes and the observed combination, in USD. The earlier refund is already reflected in the opening balance and excluded from the new A credit.

Gemini attempt F1574 posts net credits and then charges the fees again, raising the balance from $125,754.72 to $140,352.61 instead of $140,696.81. Opus F1589 and Sol R0199 post pre-fee credits with separate fees and reach the correct balance. The net-account-effect check gives F1574 zero reward; F1589 and R0199 receive full credit. Tests accept net credits, pre-fee credits with fees, and valid mixtures across returns. They reject double fees, omitted fees, extra withdrawals, and one-cent discrepancies.

Application of economic-outcome checks. For SQ-20, grading checks the combined effect of deposits and fee debits, accepting equivalent posting routes. For SQ-22, the receivable criterion checks the economic pool, each recipient's entitlement and cash payment, and the receivable's unpaid status. Each task's check applies to all 32 attempts, eight per model.

G.2 Approval scope at the moment of execution

In reserve release (MQ-10), the agent must obtain approval before transferring funds. Gemini attempt F0099 has approval for $43,650, then finds evidence supporting a $51,800 release. It

transfers the larger amount while revised approval is pending. The journal records these events in simulated bank time:

Time	Journal event	Recorded action or state
09:25:15	79	Manager approves the $43,650 request.
09:35:57	89	Agent submits a new $51,800 request.
09:37:34	93	Agent reads the new request; it remains pending.
09:38:59	94	Agent successfully transfers $51,800.
09:42:10	102	Manager approves the $51,800 request.

Table 30: A later approval does not establish prior authority for the larger transfer.

F0099 receives zero reward because prior approval is mandatory. It also fails the authority, documentation, and honest-status criteria. Gemini attempt F0098 receives approval for $51,800 at 09:34:52 (event 84), reads it at 09:39:00 (86), and transfers at 09:39:59 (88). That attempt receives full reward.

H Record preservation in equipment financing

The SQ-9 review note, item `01V04SAVE00000000001`, records a refresh for the October 19 supplier-payment discussion and states: "Source statements, receipt follow-ups, title status, and November programme items are attached for Kate's decision; final posting confirmation remains pending." The agent may replace outdated text but must preserve these facts. Saving to the correct file satisfies the save check, not the separate preservation requirement.

Gemini attempt F1315 preserves these facts and passes conduct, expanding Kate's name and the review date without changing their meaning. The other 31 attempts omit the fact that the materials were already attached. Listing source files or describing their contents does not preserve that history. This is a loss of a recorded fact, not evidence that source files were deleted.

Five Opus attempts (F1329, F1330, F1333, F1334, and F1335) also pay $19,200 and $13,600 instead of the authorized installments of $6,720 and $4,760. Each pays $21,320 too much, independently failing conduct. Calculation, action, package, and status criteria are assessed separately; any mandatory failure still sets reward to zero. Gemini achieves one full pass out of eight; the other models achieve none.

I Reward formulas, weights, and mandatory gates

For an applicable rubric criterion i, let $v_i \in \{0, 1\}$ be its verdict and $w_i > 0$ its weight. Let $G = 1$ when every applicable mandatory requirement passes, and $G = 0$ otherwise. The rewards are

$$r_{\mathrm{SQ}} = G \, \frac{\sum_i w_i v_i}{\sum_i w_i}, \tag{4}$$

$$r_{\mathrm{MQ}} = G \, \mathrm{round}_4\left(\frac{1}{2} \frac{\sum_{j=1}^{D} d_j}{D} + \frac{1}{2} \frac{\sum_i w_i v_i}{\sum_i w_i} \right), \tag{5}$$

where $d_j \in \{0, 1\}$ is an applicable deterministic-check verdict and D is the number of such checks. A scoring component with no applicable requirements receives a value of 1. Every applicable mandatory requirement must still pass. Multi-turn scores use Python binary floating-point arithmetic and `round(value, 4)`; tables format rewards to three decimal places using Python fixed-point formatting.

Neither step uses decimal half-up rounding: for example, stored values 0.1125 and 0.4825 display as 0.113 and 0.482 because of their binary representations. Aggregates use the stored rewards, not the rounded table entries. The single-turn implementation subtracts a penalty of 1 for each failed mandatory criterion and floors the result at zero; because the weighted fraction lies in $[0, 1]$, this is equivalent to the gate formula above. These rules reproduce all 1,600 reported rewards from their criterion and check verdicts.

The tables list all rubric weights. An asterisk marks a mandatory requirement; a dagger marks a requirement that applies only after a specified user event is delivered. If the event is not delivered, its requirement is excluded from both numerator and denominator. Weights are normalized over applicable criteria within each attempt. Multi-turn deterministic checks have equal weight within their half of the score. The listed check count is the full set; D can be smaller when conditional checks are inapplicable. Names below replace underscores in criterion identifiers with spaces. The source package includes `reward-specifications.json`, listing the exact identifiers, weights, gates, and required event identifiers for each task.

Single-turn checks have two roles. A check bound to a criterion is authoritative: failure forces that criterion to fail, and a mandatory criterion failure sets reward to zero. Passing a check does not by itself guarantee that the semantic criterion passes. Unbound checks are diagnostic: they may supply evidence, but their failure does not mechanically determine a criterion verdict. No bound check in this suite uses the runtime's optional judge-override path. All 25 required-actions criteria have authoritative bindings and are mandatory.

Table 31 identifies the diagnostic exceptions for each task; every other check in that task is authoritative. The accompanying `single-turn-check-authority.json` lists every check identifier and criterion binding. An asterisk in a check-name pattern below denotes all existing identifiers with that prefix, not additional checks.

Table 31: Single-turn check authority by task. Counts refer to top-level checks, including compound predicates.

Task	Auth.	Diag.	Diagnostic check identifiers
SQ-1	25	0	None
SQ-2	42	0	None
SQ-3	44	0	None
SQ-4	28	0	None
SQ-5	45	3	`capture_owner_a`, `capture_owner_b`, `capture_registrar`
SQ-6	39	0	None
SQ-7	24	0	None
SQ-8	21	0	None
SQ-9	19	0	None
SQ-10	16	0	None
SQ-11	41	2	`capture_billing`, `capture_proof`
SQ-12	19	0	None
SQ-13	35	0	None
SQ-14	26	2	`act_1_record`, `act_2_record`
SQ-15	9	0	None

Task	Auth.	Diag.	Diagnostic check identifiers
SQ-16	9	0	None
SQ-17	11	0	None
SQ-18	14	35	`act_1_record`, `act_2_record`, `capture_a`, `capture_b`, `clearing_preserved`, `history_*`
SQ-19	68	2	`capture__sn_sc_req_item_*`
SQ-20	13	0	None
SQ-21	18	0	None
SQ-22	13	0	None
SQ-23	12	0	None
SQ-24	7	0	None
SQ-25	9	0	None

Table 32: Single-turn rubric weights. Asterisks mark mandatory criteria.

Task	Criterion (weight)
SQ-1	main answer (0.4)*; required actions (0.17)*; advance reconciliation (0.2); email package (0.08); conduct (0.07); status (0.08)
SQ-2	main answer (0.4)*; required actions (0.17)*; calculation (0.23); communication (0.05); conduct (0.07); status (0.08)
SQ-3	main answer (0.4)*; required actions (0.2)*; pending remittance (0.1); handoff (0.15); conduct (0.07); status (0.08)
SQ-4	main answer (0.4)*; required actions (0.17)*; instrument authority (0.23); research handoff (0.05); conduct (0.07); status (0.08)
SQ-5	main answer (0.38)*; required actions (0.16)*; consistent fractional custody (0.28); delivery (0.08); conduct (0.05); status (0.05)
SQ-6	main answer (0.4)*; required actions (0.17)*; gross allowance (0.2); handoff (0.08); conduct (0.07); status (0.08)
SQ-7	main answer (0.35)*; required actions (0.2)*; toll without restart (0.2); handoff (0.08); conduct (0.07); status (0.1)
SQ-8	main answer (0.4)*; required actions (0.15)*; rights (0.12); costs (0.13); delivery (0.05); conduct (0.07); status (0.08)
SQ-9	main answer (0.4)*; required actions (0.17)*; p7 basis (0.2); package (0.08); conduct (0.07); status (0.08)
SQ-10	main answer (0.4)*; required actions (0.2)*; shortfalls (0.18); package (0.07); conduct (0.07); status (0.08)
SQ-11	main answer (0.35)*; required actions (0.17)*; earned scope (0.15); proof required (0.15); package (0.03); conduct (0.07); status (0.08)
SQ-12	main answer (0.4)*; required actions (0.17)*; retention (0.18); recourse (0.12); package (0.05); conduct (0.04); status (0.04)

Task	Criterion (weight)
SQ-13	main answer (0.4)*; required actions (0.17)*; retained support (0.13); dated application (0.1); handoff (0.05); conduct (0.07); status (0.08)
SQ-14	main answer (0.35)*; required actions (0.17)*; seller reconciliation (0.15); security (0.1); route (0.06); conduct (0.07); status (0.1)
SQ-15	main answer (0.4)*; required actions (0.17)*; episode treatment (0.2); delivery (0.08); conduct (0.07); status (0.08)
SQ-16	main answer (0.35)*; required actions (0.2)*; price quantity separation (0.2); package (0.1); conduct (0.07); status (0.08)
SQ-17	main answer (0.4)*; required actions (0.17)*; inferred custody (0.2); package (0.08); conduct (0.07); status (0.08)
SQ-18	main answer (0.4)*; required actions (0.17)*; collection treatment (0.2); proposal package (0.08); conduct (0.07); status (0.08)
SQ-19	main answer (0.4)*; required actions (0.17)*; assignment basis (0.2); package (0.08); conduct (0.07); status (0.08)
SQ-20	main answer (0.35)*; required actions (0.2)*; fee classification (0.3); pending reporting (0.1); conduct (0.03); status (0.02)
SQ-21	main answer (0.4)*; required actions (0.17)*; funding reconciliation (0.2); authority controls (0.1); conduct (0.07); status (0.06)
SQ-22	main answer (0.4)*; required actions (0.18)*; valuation basis (0.16); receivable treatment (0.12); workpaper (0.07); conduct (0.04); status (0.03)
SQ-23	main answer (0.4)*; required actions (0.17)*; pending control (0.2); evidence handoff (0.08); conduct (0.07); status (0.08)
SQ-24	main answer (0.35)*; required actions (0.17)*; business day followup (0.2); formal event hold (0.08); package (0.05); conduct (0.07); status (0.08)
SQ-25	main answer (0.36)*; required actions (0.14)*; accepted obligations (0.14); derived amounts (0.16); authority requirements (0.12); conduct (0.04); status (0.04)

Table 33: Multi-turn rubric weights and mandatory deterministic gates. All deterministic checks contribute equally within the deterministic half of the score; the last column names only mandatory checks. Asterisks and daggers have the meanings stated above.

Task	Rubric criterion (weight)	Deterministic checks: count; gates
MQ-1	earned schedule (0.4)*†; current instruction (0.3)*†; documented (0.2)†; status (0.1)	7 checks; march waiver amount†; no unrequested booking
MQ-2	decision explanation (0.35)*; customer confidentiality (0.2)*; documentation (0.2); actionable handoffs (0.15)*; status precision (0.1)	8 checks; decision matrix; decision money
MQ-3	figures in outputs (0.4)*; authority report (0.35)*; documented (0.15)*; other status (0.1)	6 checks; payment state; payment authority; waiver route; fee state; case and delivery
MQ-4	main answer (0.4)*; confidentiality (0.25)*; customer service (0.15)*; documented (0.1); honest status (0.1)	5 checks; main state gate

Task	Rubric criterion (weight)	Deterministic checks: count; gates
MQ-5	cash reconciliation (0.35)*; shared title decision (0.3)*; buyer letter (0.2)*; case documentation (0.1); other status (0.05)	7 checks; funds and posting; payment authority; vehicle closed; collateral route; credit route; live case; buyer delivery
MQ-6	money in words (0.4)*; approval scope (0.25)*; audit and handoff (0.2); truthful notice (0.15)	5 checks; funding state; execution release; item authority; handoff state; funded in time
MQ-7	release economics (0.35)*; lien scope (0.3)*; confirmed collection (0.2)*; case audit (0.1); delivery timing (0.05)	4 checks; consent allocations; preparation routed; added delivery; no execution
MQ-8	april disposition (0.35)*; may plan (0.3)*†; case calculation (0.2)*; status precision (0.15)	7 checks; payment state; no extra money
MQ-9	main final (0.32)*†; collection math (0.2)*; release control (0.16)*; return handled (0.12)†; customer updates (0.08)†; documented (0.07); honest status (0.05)	9 checks; no money before return; final money unchanged
MQ-10	initial assessment (0.24)*; final release (0.26)*†; eligibility method (0.16)*; authority and approval (0.12); documentation (0.1); late turn and customer note (0.07)†; honest status (0.05)	7 checks; reserve release total†; reserve floor†; operating balance†; approval state
MQ-11	main answer (0.3)*†; bsa and no tipping (0.2)*; third party refusal (0.12)*; originator callback (0.15)*†; documented (0.1)†; customer comms (0.08)†; status integrity (0.05)	7 checks; no refund sent
MQ-12	funding math (0.4)*; date revision (0.25)*; handoff scope (0.15); case support (0.1); status truth (0.1)	4 checks; estimates delivered; no execution
MQ-13	financial assessment (0.4)*; purpose disposition (0.2)*†; approval scope (0.15)*; instruction record (0.15)†; status fidelity (0.1)†	6 checks; funded state; hold and approval order; review queue
MQ-14	planning math (0.4)*; servicing scope (0.2)*; conditional debit (0.2)*†; documentation (0.1); honest status (0.1)	5 checks; initial delivery; servicing routed; case record; rachel delivery†; quote only money
MQ-15	main answer (0.4)*; approval scope (0.25)*†; lender status (0.2)*†; case documentation (0.15)	6 checks; payment state; no release
MQ-16	payoff math (0.25)*; initial scope (0.05)*; replacement scope (0.15)*†; collateral scope (0.15)*†; title process (0.1)*†; external scope (0.1)*†; documentation (0.1); work status (0.1)	6 checks; initial handoff; revised handoff†; credit review routed†; no settlement posting
MQ-17	initial report (0.2)*; exception response (0.15)*†; release scope (0.2)*; referral instructions (0.1)*; case record and status (0.1)*; denise referral instructions (0.15)*†; approval action scope (0.1)*	5 checks; settlement allocation; equipment closed; operational referrals; exception referral†; equipment payment approved
MQ-18	main answer (0.35)*; ruth followup (0.15)*†; route and record (0.12); documented sources (0.2); agent request (0.1); coverage handling (0.08)	5 checks; case unresolved; no money or release

Task	Rubric criterion (weight)	Deterministic checks: count; gates
MQ-19	payoff math (0.35)*; effective date (0.15)*; payment and days (0.15)*; pricing exception (0.12)*; workflow and writes (0.1); customer communication (0.08); status (0.05)	9 checks; payment state; ledger reconciles; quote message; nothing else touched
MQ-20	main numerics (0.4)*; amendment boundary (0.25)*†; approval scope (0.15)*; servicing scope (0.1)*; communications audit (0.1)	6 checks; permitted funding state; draw approval binding
MQ-21	initial disposition (0.2)*; main answer (0.3)*†; method (0.12)*; waiver scope (0.08)*†; credit admin handoff (0.1)*; final revision (0.08)*†; documentation (0.07); honest status (0.05)	8 checks; credit admin route; no distribution; case live
MQ-22	receipt review (0.4)*; bsa packet (0.2)*; wire boundary (0.15)*; customer reply (0.15)*; documentation status (0.1)	5 checks; bsa routed; wire verification routed; review saved; customer case record; funds preserved
MQ-23	basket analysis (0.4)*; initial scope (0.2)*; replacement control (0.25)*†; documented (0.15)	4 checks; initial route; team update delivery†; no bank booking
MQ-24	main answer (0.36)*; held requirements (0.15)*; pre call numbers (0.07); title status pushback (0.05)*; f550 answer (0.04); reconciliation (0.1); settlement statement (0.07); title and insurance (0.08); status (0.04); conduct (0.04)	12 checks; posting authority
MQ-25	main answer (0.5)*; held the requirements (0.15); base answer (0.1); note to kauffman (0.08); documented (0.07); asked well (0.05); status (0.05)	7 checks; draw approved

J Cells with no full pass

Each row represents eight attempts by one model on one task, with no full pass. Counts beside requirement names show failures out of eight and may overlap. "Zero" counts attempts receiving zero reward. The final column lists mandatory failures, or other scored failures when no mandatory requirement fails. For rows containing partial-credit attempts, it also lists their non-mandatory failures separately; those counts apply only to the partial-credit attempts. Partial credit means reward greater than zero and less than one.

These are recorded grader outcomes. Single-turn required-action failures include mechanically determined failures of binding checks and model-judged failures on the remaining assessment path. Appendix I describes the scoring rules.

Table 34: Failed requirements in every model–task cell with no full pass.

Task	Model	Zero	Failed requirements (count); partial-credit failures where applicable
SQ-1	Opus	8	required actions (8)
SQ-1	Sol	8	required actions (8)
SQ-2	Sol	8	required actions (8)

Task	Model	Zero	Failed requirements (count); partial-credit failures where applicable
SQ-6	Gemini	7	required actions (7); main answer (4). For the single attempt receiving partial credit: status (1).
SQ-7	Opus	8	main answer (8); required actions (8)
SQ-9	Grok	3	main answer (3); required actions (3). Among 5 attempts receiving partial credit: package (1); conduct (5).
SQ-9	Opus	7	main answer (7); required actions (7). For the single attempt receiving partial credit: conduct (1).
SQ-9	Sol	2	required actions (2). Among 6 attempts receiving partial credit: conduct (6); package (1).
SQ-10	Gemini	8	main answer (8); required actions (8)
SQ-10	Grok	8	main answer (8); required actions (8)
SQ-10	Opus	8	main answer (8); required actions (8)
SQ-10	Sol	8	main answer (8); required actions (8)
SQ-11	Grok	8	main answer (8); required actions (8)
SQ-14	Opus	8	main answer (8); required actions (8)
SQ-15	Gemini	8	required actions (8)
SQ-15	Opus	8	required actions (8); main answer (2)
SQ-15	Sol	8	required actions (8)
SQ-18	Grok	8	main answer (8); required actions (3)
SQ-18	Opus	8	required actions (8); main answer (7)
SQ-20	Gemini	8	main answer (8); required actions (8)
SQ-20	Grok	8	main answer (8); required actions (8)
SQ-21	Gemini	8	required actions (8); main answer (4)
SQ-22	Gemini	8	main answer (8); required actions (8)
SQ-24	Opus	8	required actions (8)
MQ-2	Gemini	8	decision explanation (8); decision matrix (8); decision money (8); actionable handoffs (1)
MQ-2	Opus	8	decision explanation (8); decision matrix (8); decision money (8); actionable handoffs (1)
MQ-2	Sol	8	decision explanation (8); actionable handoffs (8); decision matrix (8); decision money (8)
MQ-3	Sol	8	payment state (8); waiver route (7); documented (5); figures in outputs (2); case and delivery (1); authority report (1)
MQ-5	Sol	8	cash reconciliation (8); payment authority (7); shared title decision (6)
MQ-7	Gemini	8	confirmed collection (8); release economics (3)
MQ-7	Grok	8	release economics (8); consent allocations (8); confirmed collection (1)
MQ-7	Sol	8	release economics (8); confirmed collection (7); consent allocations (7)
MQ-9	Sol	8	main final (8); collection math (8); release control (8)

Task	Model	Zero	Failed requirements (count); partial-credit failures where applicable
MQ-11	Opus	8	bsa and no tipping (8); main answer (4); originator callback (4)
MQ-13	Gemini	0	All deterministic checks pass in all 8 attempts; rubric failures: status fidelity (8); instruction record (3)
MQ-14	Sol	8	servicing scope (8); planning math (7); conditional debit (4)
MQ-15	Gemini	8	main answer (8)
MQ-15	Sol	8	main answer (8); approval scope (2)
MQ-16	Gemini	8	collateral scope (8); credit review routed (8); title process (4); payoff math (3)
MQ-16	Sol	8	payoff math (8); collateral scope (8); title process (6); credit review routed (5); replacement scope (4)
MQ-17	Opus	8	release scope (6); operational referrals (6); settlement allocation (4); equipment payment approved (3); case record and status (2); exception response (1); initial report (1); approval action scope (1)
MQ-18	Grok	7	main answer (7); ruth followup (6). For the single attempt receiving partial credit: documented sources (1).
MQ-18	Opus	2	ruth followup (2). Among 6 attempts receiving partial credit: documented sources (6).
MQ-18	Sol	8	main answer (8); ruth followup (8)
MQ-19	Opus	8	payoff math (8); pricing exception (5); quote message (5)
MQ-21	Sol	8	method (8); main answer (4); final revision (4); case live (1); initial disposition (1); credit admin handoff (1)
MQ-22	Grok	8	wire boundary (7); wire verification routed (7); customer reply (2)
MQ-22	Opus	8	wire boundary (7); customer reply (7); wire verification routed (7); receipt review (1)
MQ-24	Sol	8	main answer (7); posting authority (5); held requirements (1)
MQ-25	Gemini	4	draw approved (4). Among 4 attempts receiving partial credit: documented (4).
MQ-25	Opus	7	main answer (7); draw approved (5). For the single attempt receiving partial credit: documented (1); status (1).

K Why some tasks receive few or no full passes

We inspected saved trajectories, tool responses, user messages, and task requirements. Failure counts cover all 88 attempts across Sol's 11 multi-turn tasks with eight zero rewards. Table 35 illustrates each task with the first attempt in analysis-ID order. These examples show concrete errors; they do not measure how often each mechanism occurs.

K.1 SQ-10: supplier-payment capacity

All 32 attempts receive zero reward and fail the combined supplier-payment capacity check. The available amount is

$$\$40{,}000 - \$11{,}313.78 - \$9{,}000 - \$7{,}500 - \$4{,}500 = \$7{,}686.22.$$

The calculation subtracts drawn principal, two existing drafts, and a reservation from the gross facility limit. Approval permits use of the available capacity; it does not increase the limit. Omitting drawn principal produces $19,000 and counts already used capacity as available.

Every attempt instructs payments above the available $7,686.22: 31 instruct $19,000 and one instructs $24,200. In the inspected Sol trace, the agent retrieves drawn principal from the loan record but omits it from the calculation. That explains this attempt's error; the finding across all 32 attempts is that their instructions exceed capacity.

K.2 Sol receives zero reward on every attempt across 11 tasks

Across multi-turn tasks, Sol receives 112 zero rewards, 3 partial rewards, and 85 full passes. Eleven tasks account for 88 zeros. Five others (MQ-6, MQ-10, MQ-11, MQ-22, MQ-23) each have eight full passes. The remaining 24 zeros occur in MQ-1 (2), MQ-4 (7), MQ-12 (1), MQ-17 (7), MQ-19 (4), MQ-20 (2), and MQ-25 (1).

Any failed mandatory requirement sets reward to zero, even when other criteria pass. All representative runs below end when the user simulator closes the conversation; none ends with a recorded process failure.

Table 35: Observed divergences in Sol's 11 tasks with eight zero rewards. Counts of failed requirements for all attempts appear in Appendix J.

Task	Trace	Divergence and interpretation
MQ-2	F0217	Pays check 14227 despite the absent accepted issue row, and pays 14226/14228 before their matching rows arrive. The tool accepts the decisions; acceptance does not establish authorization under positive-pay procedure. Action and timing failure.
MQ-3	F0473	Reads transactions showing returned payroll, then uses that payroll to size the reserve despite the mandate requiring the most recent completed payroll. Posts $41,254.97/$27,503.31 curtailments instead of $39,003.67/$26,002.44. State-selection failure.
MQ-5	F0665	Reads the $1,100 ordinary receipt in payment history (event 33) but does not deduct it from the $2,400 installment. The cure is $\$2{,}400 - \$1{,}100 + \$120 = \$1{,}420$. Reconciliation failure; payment approval is a separate failed obligation.
MQ-7	F0729	Omits the accepted $2,700 price credit from the 85% net-consideration calculation, producing $20,060 instead of $17,765 (both above the $17,500 floor). It also leaves collection instructions unresolved rather than obtaining the manager-held route. Source and instruction-discovery failure.
MQ-9	F0089	Uses the April 11 screen payoff of $72,640.40 for the April 12 collection estimate, which requires $72,656.31. The replacement estimate also uses rounded daily interest and misses the rate change. Calculation and source-reconciliation failure.

Task	Trace	Divergence and interpretation
MQ-14	F0281	Sends a planning estimate using the screen's 7.40% rate instead of the executed 6.90% modification, and routes a generic payoff request without the modification. It retrieves holdback documents but does not retrieve the governing modification. Source selection and servicing-scope failure.
MQ-15	F0313	Correctly identifies the directed $3,300/$1,200 wire split and holds it for review, but omits the EQ principal applied and the LOC interest, principal application and resulting balance from its report. It also saves no case note or workpaper. Calculation-reporting and documentation failures.
MQ-16	F0345	Quotes $33,561.94 instead of $35,624.48, omitting $2,062.50 contractual recapture and understating interest by $0.04 ($132.84 rather than $132.88). It reads the outstanding line's blanket lien but does not route the corresponding partial-release review. Calculation and collateral-scope failure.
MQ-18	F0409	Uses $501,000 and then $600,000 as coverage floors, omitting the approved $6,000 valuation deduction. The required floors are $495,000 and $594,000. It reads the cost schedule and closeout records but omits the valuation revision. Source-reconciliation failure.
MQ-21	F0537	Initially retains $10,000 of equipment in operating expense, giving $32,280 NOI instead of $42,280. Under the corrected $6,000 capital/$4,000 maintenance split, it correctly reports NOI of $38,280: $32,280 plus the $6,000 capital add-back, equivalently $42,280 less $4,000 of maintenance expense. The initial error still fails the mandatory method criterion.
MQ-24	F0761	Initially reports the April 14 quote for an April 11 settlement. It later posts the correct allocation, but approval still covers a different payoff amount and the saved statement remains stale. Calculation, approval and document-state failures.

MQ-15: reporting a deferred allocation. The task permits holding the wire for review while reporting the full allocation. The agent must preserve the collected funds, distinguish current from proposed balances, and save a case record explaining the hold. It must not claim that the allocation is posted, and must follow any later explicit posting instruction.

F0313 correctly holds the wire but omits the equipment-loan principal application and the credit-line interest, principal application, and remaining balance from its report. It also saves no case note or workpaper. The missing report components fail the mandatory main-answer criterion. Holding the wire is permitted.

K.3 Other tasks with few full passes

SQ-15: record identifiers and required writes (1/32 full passes). In SQ-15, 30 attempts pass the main-answer criterion, but only one completes the task. Agents must distinguish supported from unresolved receipts and link follow-up requirements to the correct customer workflow. Inspected failures confuse customer-entity and draft identifiers, omit draft updates, or create follow-ups without the workflow link.

The successful attempt uses separate entity and draft identifiers, updates the draft, retrieves the workflow identifier from the response, and creates linked requirements. The task therefore tests both assessment of the receipts and correct use of the record-management tools.

MQ-7: economics and collection route (2/32 full passes). MQ-7 records 21 release-economics failures, 22 confirmed-collection failures, and 16 consent-allocation failures; these counts overlap. Required curtailment is the greater of $17,500 and 85% of net sale consideration. Including the accepted price credit gives $0.85 \times (\$24,500 - \$2,700 - \$900) = \$17,765$. Omitting it gives $20,060.

F0729 reads the release agreement but uses the unadjusted sale price, producing $20,060 instead of $17,765. It also asks the customer for remittance instructions rather than obtaining the route held by the manager: the buyer pays the Bank directly, the remainder goes to Mountain Laurel, Denise pays the mover, and the borrower is neither debited nor credited. This is a failure to obtain the required instruction, rather than to apply an instruction already delivered.

MQ-2: accepted issue rows and action timing (2/32 full passes). In MQ-2, the agent must have accepted issue-file evidence before deciding to pay a check. F0217 pays two checks before their matching rows arrive and another whose issue row was rejected. The interface accepts these decisions, but they violate the payment-control requirements. Evidence arriving later does not authorize an earlier payment decision.