Hybrid AI Patterns

Thirteen Ways to Couple Language Models with Rules, and How to Choose Among Them

Siegfried Handschuh
Institute of Computer Science
University of St. Gallen (HSG)
siegfried.handschuh@unisg.ch

Götz-Henrik Wiegand
Institute of Computer Science
University of St. Gallen (HSG)
goetz-henrik.wiegand@unisg.ch

Lorena Raichle
Institute of Computer Science
University of St. Gallen (HSG)
lorena.raichle@unisg.ch

October 2026

arXiv:2610.40036v1 [cs.AI] 29 Oct 2026

Abstract

Language models handle text well, and rule engines give conclusions with proofs. Systems that combine the two are usually described by example, and the same coupling turns up under different names in agent frameworks, guardrail toolkits and neuro-symbolic research. This document defines thirteen patterns for coupling a language model with an executable rule language, in the catalogue form introduced for object-oriented design: intent, motivation, applicability, participants, collaborations, consequences, implementation, known uses and related patterns. Four patterns concern what crosses the boundary between model and rules, five who controls whom, three how the model takes part in the reasoning, and one how knowledge is kept over time. Five invariants hold for every combination, and five failure patterns name the couplings that break them. We relate the catalogue to published pattern catalogues and taxonomies of hybrid and neuro-symbolic systems and show where they overlap with it. A decision guide leads from a problem to a pattern. It rests on our measurements, including the cases in which a model alone did as well as the hybrid or better, and it reports what each design cost.

Contents

1 Introduction **3**

2 Foundations **4**
2.1 What FrameX adds for coupling 4
2.2 Five invariants 6
2.3 Four elementary moves 7
2.4 The catalogue form 7

3 Patterns at the boundary: what crosses it **10**
3.1 Pattern 1: Model as Sensor 10
3.2 Pattern 2: Model as Author 14
3.3 Pattern 3: Model as Translator 18
3.4 Pattern 12: Model Selects, Rules Write 23

4 Patterns of control: who decides **27**
4.1 Pattern 4: Model Proposes, Rules Dispose 27
4.2 Pattern 5: Rules Steer the Model 32
4.3 Pattern 6: Critique and Repair 36
4.4 Pattern 7: Division by Competence 38
4.5 Pattern 13: Fail-Safe Union 42

5 Patterns in the reasoning itself **47**
5.1 Pattern 8: Model as Hypothesis Generator 47
5.2 Pattern 9: Model as Critic of the Rules 51
5.3 Pattern 10: Rules Teach the Model 55

6 Patterns over time **59**
6.1 Pattern 11: Rule-Governed Memory 59

7 Compositions **64**

8 Failure patterns **65**

9 Related catalogues **66**
9.1 Mapping 66
9.2 The patterns in boxology notation 66
9.3 What the catalogue adds 67
9.4 What the catalogue lacks 68
9.5 What the catalogue learns from them 69

10 Choosing a pattern **70**
10.1 First question: is a hybrid needed? 70
10.2 From problem to pattern 70
10.3 What the guide rests on 71

11 Limits **74**

A Boxology legend **75**

B Pattern libraries **76**

C Rule graph legend **84**

1 Introduction

More and more organisations let language models act on documents, conversations and transactions that are governed by explicit rules: a regulation, a protocol, an authority limit. Two designs are possible. The rules can be written into the prompt and the model asked to apply them, or the rules can stay executable and the model be given the parts of the task that rules cannot do: reading text, phrasing a question, proposing a candidate. Between these two designs lie many couplings, and practitioners tend to reinvent them project by project.

A pattern names a recurring problem in a context and the core of a solution that can be reused without being copied (Alexander et al. 1977). The catalogue form of Gamma et al. (1994) made patterns a working vocabulary for object-oriented design, and Buschmann et al. (1996) carried it to the architecture of whole systems: each pattern has a name, an intent, a motivating example, conditions of applicability, a structure of participants and collaborations, consequences, notes on implementation, known uses and related patterns. We use the same form for couplings of language models with rules. A shared name lets designers say which coupling they mean. The consequences make the trade-offs explicit, including the cases where a model alone is the better choice, and the known uses tie each pattern to running systems and to measurements.

The patterns were found in the applications of FrameX, a rule and knowledge engine based on F-Logic (Kifer et al. 1995), and checked against published catalogues (Section 9). They do not depend on FrameX: any rule engine that can request a model call, store the answer as data and derive conclusions over it can host them. The examples use FrameX's notation because we made the measurements in it. An early version of the FrameX engine and framework is publicly available at `https://gitlab.com/ds-unisg/framex-core`; a full release is forthcoming.

Contributions. The following contributions are made in this document:
1: Thirteen patterns in catalogue form, composed from four elementary moves, with five invariants that hold for all of them (Sections 2 to 6), and five failure patterns that break the invariants (Section 8).
2: A mapping onto published pattern catalogues and taxonomies that shows which patterns are known, which are new, and which published patterns this catalogue lacks (Section 9).
3: A decision guide from problem to pattern, grounded in pre-registered measurements (Section 10).

2 Foundations

2.1 What FrameX adds for coupling

The listings use FrameX, a rule language in the F-Logic tradition. We assume that the reader knows its basic syntax: frames `?X:Class[slot -> ?V]`, rules `head <- body`, `AND`, `NOT`, queries `?- ...`, and the open-world reading of *unknown*. Every program declares its world. Under `world open.` a missing fact is *unknown*, and `NOT` gives *true* only over a slot, method or class declared closed (`closed slot allowsSlot.`, `closed method refuted/1.`, where `/1` is the number of parameters, `closed class Admitted.`); over anything else the negation stays *unknown*. The programs also compute values (`?M := ?N + 1`), compare them (<, >, !=), use methods with parameters (`slot(?I) -> ?V`) and strings as objects (`"mbi":Programme`), and write an object whose name is a term in backquotes (`` `read(d1)` ``). This section introduces only what coupling with a model adds. All thirteen patterns rest on it.

External requests. A rule asks for something outside the engine with the keyword `external`:

```
external read(?D) using _llm[prompt -> "Report what the document states ...";
    document -> ?T]
  <- ?D:Document[pending -> true; text -> ?T].
```

The engine evaluates this rule like any other. Its conclusion is a request: an object `` `read(d1)` `` with the fields in brackets as attributes, the tool (`tool#name -> "_llm"`), a mark `tool#requested -> true`, and link facts for the request's arguments (`tool#arg1 -> d1`, and `tool#arg2` for a second argument). The engine never calls a network. A *host*, a program beside the engine, reads the open requests, runs the tool and adds its answer as a *receipt*; the engine then recomputes, and rules that read the receipt can fire, including rules that issue the next request. The keyword is kept although the tool name would identify a request, because such a rule follows other laws than a derivation: a request is carried out at most once in a run and is never withdrawn, whereas a derived fact disappears when its premises do. Every point at which a program reaches outside the engine is marked by it.

Receipts. An answer returns as facts about the request object:

```
// Written by the host after the call.
`read(d1)`[tool#status -> "succeeded"; tool#completed -> true;
   row(1, "slot") -> "customer"; row(1, "value") -> "Acme AG";
   row(1, "quote") -> "Acme AG (the Customer)"; row(1, "quoteFound") -> true].

// Read by the rules, joined to the document through the request's argument.
?D[slot(?I) -> ?S]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?D; row(?I, "slot") -> ?S].
```

A scalar answer field `key` becomes `result("key")`; the reserved field `rows`, a list of at most 64 flat objects, becomes `row(N, "key")`. Every receipt ends with `tool#completed -> true` and `tool#status`, which is `"succeeded"` or `"failed"`. A failed call (the model did not answer in time, the reply was not valid JSON, a field the host needs was missing) adds no `result` and no `row`, only an attribute `error` with the reason, `` `read(d1)`[error -> "model response is not valid JSON; no facts accepted"] ``. A rule that reads `tool#status -> "failed"` keeps such a case apart from a rejection (invariant 5). Answers must be flat: values are stored as literals and never evaluated, nested objects and lists are refused, and a `null` states nothing. Rules usually key rows by method parameters (`?D[slot(?I) -> ...]`); since FrameX allows function terms in rule heads, a rule can also give each row an object of its own, `st(?U, ?I):Statement[...]` (Section 6.1). Table 2 lists what each tool reads from a request and writes into its receipt.

Identity and order. A request is identified by its name and arguments. `read(d1)` runs once; a second attempt needs a new identity, as in `propose(?T, ?N)` with a step number (Section 4.1). A request

without arguments, `external propose using ...`, is the object `propose` itself, and its receipt is read as `propose[...]` (Section 5.1). Instances of one request name form a queue; requests of different names that are open at the same time need `tool#order -> N`, and the lowest goes first. Two values for one request field are an error. A run ends when a rule derives `workflow[status -> ...]` or when no request is open.

Tools. The tool after `using` names what the host does (Table 1). Tools beyond those shipped with FrameX come from adapter files that the operator names when starting the host; a program can never load code.

Table 1: Tools used in this document.

Tool	Provided by	What the host does
`_llm`	`framex.tools`	one model call; see below
`_query`, `_evidence_needs`	`framex.tools`	a query on the rules themselves; the evidence that would settle a goal
`_turn`, `_judge`, `_repairTurn`	`framex.dialogue`	the dialogue model phrases a question; a judge decides a criterion; a repeated turn after a broken reply
`_ask`	`framex.dialogue`	one question to a person, as written, without a model
`_checkCandidate`	`framex.candidates`, by adapter	checks a drafted rule against cases, before and after
`_forgetAnswers`	`framex.tools`	removes stored model answers that contain a forgotten text

Table 2: What each tool reads from a request and writes into its receipt. Every receipt also has `tool#status` and `tool#completed`; a failed call has `error` instead of results.

Tool	Request fields	Receipt
`_llm`	`prompt` (or `with P`) and any named fields; `groundIn`, `mustSay`, `contextKey`, `tool#maxContext` for the host	the reply's fields as `result("k")` and `row(N, "k")`; with `groundIn` `quoteFound` and `valueQuoted` per row; with `mustSay` `saysDecision`; with `contextKey` `tool#sent("id")`
`_query`	`query`	`result("status")` (`true`, `false`, `unknown` or `bindings`) and the whole answer as JSON in `result("text")`
`_checkCandidate`	`candidate`	`status` (`checked`, `rejected`, `error`), `feedback`, and after evaluation `regressions`, `failedAfter`, `newlyFailing`, `newlyPassing`, `newViolations`, `changesConclusions`, all as `result("k")`
`_turn`	`question`, `attempt`	`answer`, `utterance`; after a broken reply `formatError` and `received`
`_judge`	`criterion`, `question`, `answer`	`verdict` (`yes`, `no`, `unclear`), `reason`, `label`; after a broken reply `formatError` and `received`
`_ask`	`question`	`question`, `answer`
`_forgetAnswers`	`text`	`removed`, the number of stored answers deleted

The model call. `_llm` sends `prompt` as the instruction and the other fields as one JSON message, so a rule hands over separate, traceable slots. Fields that start with `tool#` and the fields read by the host (`groundIn`, `mustSay`, `contextKey`) are not sent. Long prompts live in prompt objects (`readPrompt:Prompt[text -> "..."]`) that a request names with `using _llm with readPrompt[...]`. Every answer is stored with its date, model and messages and replayed on the next identical request until a refresh is asked for; a run can therefore be repeated without the model. Which

model answers is set in the host's environment (a local server or a hosted endpoint), and keys never enter programs or receipts.

Measurements on text. Two request fields are read by the host and never shown to the model. With `groundIn -> ?Text`, every answer row gets `quoteFound` (its quote occurs verbatim in that text) and `valueQuoted` (its value occurs in its quote). With `mustSay -> ?Word`, the receipt gets `saysDecision` (the reply contains the word). Both are measurements; whether a row or a reply counts is decided by rules.

Rule-selected context. A request with `contextKey -> ?Step` receives in its field `context` exactly the statements that rules select with `?S[selectedFor -> step]`, ordered by identifier and at most `tool#maxContext` of them; more makes the call fail. Each statement goes to the model as one JSON object with its identifier (`id`, the object's printed name, such as `st(u2, 1)`) and its slots; the default limit is 32. The receipt records only the identifiers of what was sent, as `tool#sent("st(u2, 1)") -> st(u2, 1)`. A reply cites statements by identifier in its rows, so rules can check that every cited statement was sent and is current (Section 6.1).

One trap. Because the fields of a request become attributes of its object, a condition without a class can match the request itself: `?Q[decision -> ?D]` also matches `` `phrase(q1)` `` when that request has a field `decision`. FrameX stops such a run and names the cause. A field named like a slot that rules derive for the receipt gives it two values; `framex validate` warns about it. Give rule subjects a class (`?Q:Request[...]`) and name request fields apart from derived slots.

The other direction. A model can also call the rules, through a protocol of read-only operations (`query`, `explain`, `why_not`, `evidence_needs`, `validate`, `preview`).

The implementation of each pattern shows the whole chain: the rule that requests the model call, the rules that read its receipt, and the rules that decide. Each program is printed in full, with its declarations, its data, its prompts and its queries; it runs against a model or, from recorded answers, without one.

2.2 Five invariants

Every pattern and every combination of patterns keeps five invariants.

1. **Model text is data.** The engine never executes model-written rules or code unchecked; an answer is a receipt that rules read. A rule or query written by a model is a candidate that a check admits or refuses before it runs (Sections 3.2 and 3.3).

2. **Every model answer is a dated observation** with its prompt, model and inputs, so a run can be replayed and audited.

3. **Rules decide.** Whether a statement counts, which step comes next, when a task is complete: the rules derive each of these.

4. **Evidence decides acceptance.** A model's confidence is stored as a number. A statement is accepted on evidence the rules can check, such as a quote found in the source.

5. **Failures stay distinct.** A failed call, a reply in the wrong format, an *unclear* judgement and a logical *unknown* are four different results.

The division of labour follows from these invariants. The model handles perception, language and ambiguity. The rules handle whatever must be guaranteed, audited or explained. What passes between them is an explicit, validated record.

2.3 Four elementary moves

Every pattern is assembled from four moves between the rules and the outside (Fig. 1).

Request A rule whose conditions hold issues an external request to a model or a tool. What the model sees is what the rule puts into the request.

Receipt The answer returns as dated facts about the request, joined to its subject through the request's arguments. Nothing else enters the engine from outside.

Accept A rule decides whether what a receipt proposes becomes a fact, is refused with its reason, or stays *unknown*.

Ask A rule derives a question for a person, or the evidence that would settle an open goal; the answer returns as a receipt.

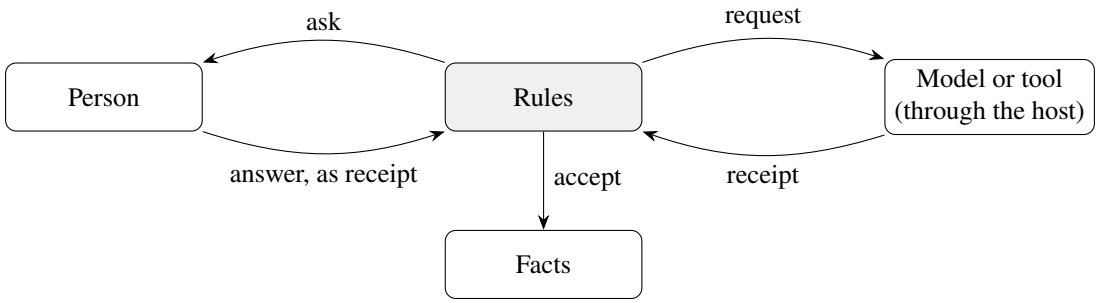

Figure 1: The four elementary moves. Every arrow that enters the rules is a receipt; only the accept move writes facts.

Table 3 shows how the thirteen patterns compose the moves. They differ in what triggers the request and in what the accept move decides; Critique and Repair is the only pattern that repeats a move, under a bound.

2.4 The catalogue form

Each pattern is described under the headings of Gamma et al. (1994, sec. 1.3), adapted to couplings. *Structure* is drawn in the boxology notation of hybrid learning and reasoning systems (van Harmelen and ten Teije 2019; van Bekkum et al. 2021), with the generative model of de Boer et al. (2025): rectangles are instances (*data* such as text, or *symbol*), hexagons are models (*model:gen* for a language model, *model:sem* for a rule base), ovals are processes (*infer*, *transform*, *generate*) and triangles are actors, here persons. The dashed arrows are ours: they mark a request, from the process that issues it to the process that answers, and so show who triggers whom where the solid arrows show what flows. A reply of the model is text; where it becomes symbols, the receipt is drawn as a transformation. Table 8 in Section A lists every symbol; Section 9.2 relates the notation to the catalogue. *Sample code* is folded into *Implementation*: each pattern prints its complete program and draws it as a *rule graph*, generated from the program by a script of the companion folder. Its boxes are the facts, the requests, the receipts and the derived slots; a small circle is one rule, labelled with the line where it starts in the listing ("l. 23" for line 23), with arrows from what it reads to what it derives (Table 9). The boxology shows the roles of a pattern, the rule graph the chain of one program. For the dialogue flow of pattern 5 the graph is a flow diagram of its steps. We add one heading, *Evidence*, for measured results where they exist, because the consequences of a coupling are empirical in a way the consequences of a class structure are not.

Patterns 1 to 11 were drawn from the applications of FrameX. Patterns 12 and 13 were added later from couplings that recurred in the submissions to a hackathon; they keep their numbers, which other documents cite, and stand in the group they belong to (Table 4).

Table 3: The patterns as compositions of the four moves.

Pattern	Moves	Request triggered by	Accept decides
1 Model as Sensor	request, receipt, accept	a pending input	whether a reading becomes a fact
2 Model as Author	request, receipt, accept, ask	a policy text or a failing case	whether a draft is offered; a person adopts it
3 Model as Translator	request, receipt, accept, request, receipt	a person's question	whether the request may run; whether the reply states the rules' decision
4 Model Proposes, Rules Dispose	request, receipt, accept	the next step of a task	admit, refuse with a reason, or narrow
5 Rules Steer the Model	request, receipt, ask	the derived step	which step follows; the judge's reading as in 1
6 Critique and Repair	accept, then request again, bounded	a refused reply	the kind of failure and when to stop
7 Division by Competence	accept, then request or ask	an *unknown* or a conflict	what the rules decide themselves
8 Model as Hypothesis Generator	request, receipt, accept in a copy, ask	an observation to explain	which candidates survive; ask names the deciding evidence
9 Model as Critic of the Rules	request, receipt, accept, ask	a change before release	which remarks become findings; a reviewer disposes
10 Rules Teach the Model	request, receipt, accept	labelling, selection, grading	the grade of a model answer
11 Rule-Governed Memory	request, receipt, accept; later request	an utterance; the next step	admission, supersession, withdrawal; what enters the prompt
12 Model Selects, Rules Write	request, receipt, accept, twice	a person's question; a value to state	whether an identifier names a passage; whether the words leave the value to the rules
13 Fail-Safe Union	request, receipt, accept, from several sources	every case, for every source	the strictest verdict in the declared order; silence as caution

Table 4: The thirteen patterns in four groups.

#	Pattern	The model	The rules
What crosses the boundary			
1	Model as Sensor	turns text into proposed facts	accept or reject each fact
2	Model as Author	drafts rules, tables, flows, tests	check them; a person adopts
3	Model as Translator	turns questions into checked requests, proofs into prose	answer and prove
12	Model Selects, Rules Write	chooses passage identifiers, phrases around values	show passages and values word for word
Who controls whom			
4	Model Proposes, Rules Dispose	proposes an answer or action	admit, refuse with a reason, narrow
5	Rules Steer the Model	answers what it is asked	decide what to ask, when, with what
6	Critique and Repair	tries again	classify the fault, state it, bound retries
7	Division by Competence	handles what the rules leave open	decide what they can, hand on the rest
13	Fail-Safe Union	gives a verdict that can only tighten	combine by a declared order; silence counts as caution
The model in the reasoning			
8	Model as Hypothesis Generator	proposes possible causes	test each, ask for deciding evidence
9	Model as Critic of the Rules	reviews rules and conclusions	record the critique, open a review
10	Rules Teach the Model	learns from labels, examples, grades	supply them with proofs
Over time			
11	Rule-Governed Memory	proposes what to remember, cites what it used	admit, supersede, withdraw, select

3 Patterns at the boundary: what crosses it

3.1 Pattern 1: Model as Sensor

Intent

Let the model read what rules cannot (free text, documents, answers in a conversation) and turn it into proposed facts; a rule decides whether each proposal becomes a fact.

Also known as

Extraction, classification as a predicate, perception, learned predicate, judge.

Motivation

An assignment review must know whether a contract names a signatory with authority, whether a consent form is signed and which party a letter addresses. Rules can decide what follows from these facts, but they cannot read the documents. A model can read them, but it also reports things the document does not say, and a fluent wrong reading looks like a right one. The pattern lets the model read and keeps the decision about what counts with the rules: the model reports each statement with a verbatim quote, the host measures whether the quote occurs, and a rule accepts a statement only if the quote was found and the kind of document can make such a statement.

Applicability

Use it when

- the inputs are unstructured (text, speech, images) but the decision over them is known and can be written as rules;
- a reading can be checked by something measurable: a quote that occurs, a value from a closed list, a hash of the source;
- the same reading is needed by several rules, so it should exist once as a fact with a provenance.

Do not use it when the decision itself is what is unknown (then see Division by Competence, Section 4.4).

Structure

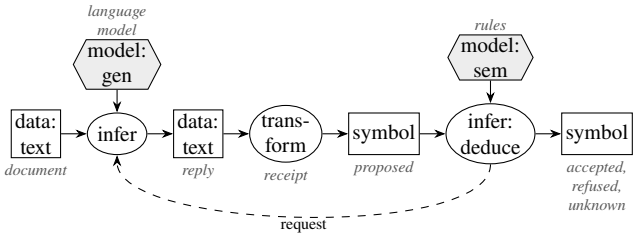

Participants

Model: reads one input and answers one question about it in a fixed format. *Host*: calls the model, measures what can be measured (format, quote occurrence), stores the answer as a receipt. *Rules*: request the reading, accept or reject each proposed fact with a reason. *Source*: the document or utterance, identified by hash.

Collaborations

A rule fires an external request for a document that is pending; the host returns a receipt with the model's answer and its measurements; an acceptance rule turns admissible statements into facts; rejected statements keep their reason. A later revocation of the source withdraws what rested on it.

Consequences

Each fact has a provenance (source, model, prompt version, date) and a reason for acceptance. An answer that cannot be grounded stays a rejected statement with its reason. A model's *unknown* leaves the criterion open. The cost is one call per input and question; a plausible but wrong value that passes every measurable check slips through this pattern and needs a second reader, a person or a cross-check.

Implementation

Ask for one thing per call: a value from a fixed list, or a statement with its quote. The request names the source text as `groundIn`, which the model does not see; the host then records for every row of the answer whether its quote occurs verbatim in that text (`quoteFound`). The program is complete: it declares its world and the slot it negates, states what a contract may say, gives one contract and the prompt, and then the steps every pattern has: the request, the rule that reads the receipt (row `?I` is one statement about document `?D`), the rules that accept or reject, and the rule for a failed call.

```
// Pattern 1, Model as Sensor: the model reads, the rules decide what counts.
world open.
closed slot allowsSlot.

// What a kind of document can state, and one document to read.
"contract"[allowsSlot -> "customer"; allowsSlot -> "provider"].
d1:Document[pending -> true; kind -> "contract";
    text -> "Service Agreement between Beta GmbH (the Provider) and Acme AG (the Customer)
        . This agreement is governed by Swiss law."].

// The instruction the model receives.
readPrompt:Prompt[text -> "Report what the document states: one row {slot, value, quote}
    per statement, slot one of customer, provider, governingLaw, the quote copied
    verbatim. Answer as JSON {\"rows\": [...]}."].

// Request: one call per pending document; groundIn makes the host check every quote.
external read(?D) using _llm with readPrompt[document -> ?T; groundIn -> ?T]
  <- ?D:Document[pending -> true; text -> ?T].

// Receipt: each answer row becomes a proposed statement about the document.
?D[slot(?I) -> ?S; value(?I) -> ?V; quoteFound(?I) -> ?Q]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?D; row(?I, "slot") -> ?S;
       row(?I, "value") -> ?V; row(?I, "quoteFound") -> ?Q].

// Accept: a quoted statement that this kind of document may make.
?D[accepted(?I) -> ?V]
  <- ?D:Document[kind -> ?K; slot(?I) -> ?S; value(?I) -> ?V;
       quoteFound(?I) -> true] AND ?K[allowsSlot -> ?S].

?D[rejected(?I) -> "quote not found in the document"]
  <- ?D[quoteFound(?I) -> false].

?D[rejected(?I) -> "not a statement this kind of document can make"]
  <- ?D:Document[kind -> ?K; slot(?I) -> ?S] AND NOT ?K[allowsSlot -> ?S].

// A failed call (no answer, or a reply that is not valid JSON) proposes nothing;
// the document stays unread, with the host's reason.
?D[notRead -> ?E]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?D; tool#status -> "failed"; error -> ?E].

?- d1[accepted(?I) -> ?V].
?- d1[rejected(?I) -> ?Why].
?- d1[notRead -> ?E].
```

Figure 2 draws the program as a rule graph. A statement the model does not report stays *unknown*: no rule accepts or rejects it. A failed call is no rejection either; `notRead` keeps the host's reason (invariant 5).

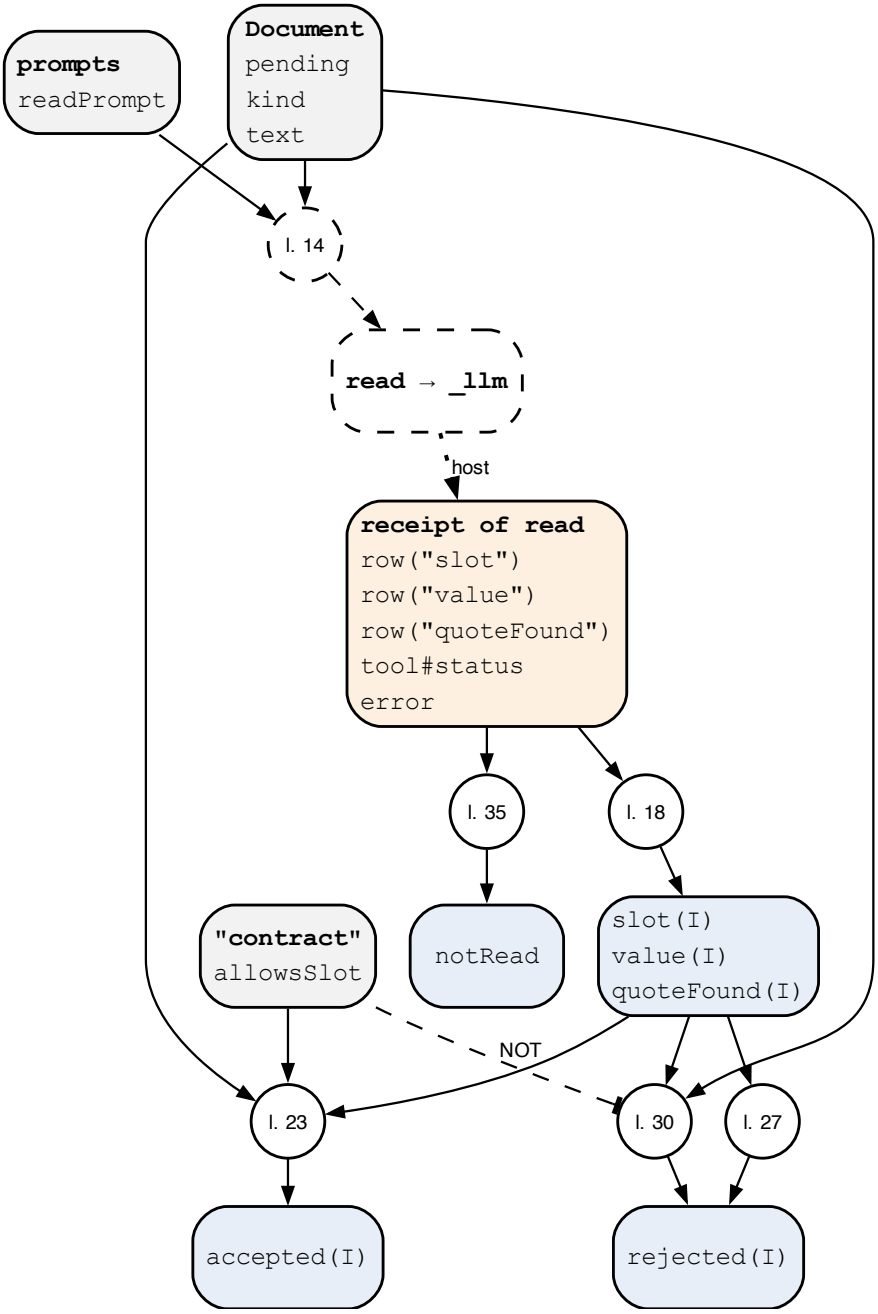

Figure 2: Rule graph of `p01-sensor.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

Library form

The rules above are the same for every source, so they can live in a library, as the dialogue flows of pattern 5 do. With the library `sensor.fx` (Section B) loaded first, the program states facts only, and the library derives `sensor#accepted`, `sensor#rejected` and `sensor#notRead` for every source:

```
// Pattern 1 in library form: load sensor.fx first. Same document, same prompt as p01-sensor.fx.
world open.

"contract"[sensor#allows -> "customer"; sensor#allows -> "provider"].

readPrompt:Prompt[text -> "Report what the document states: one row {slot, value, quote} per statement, slot one of customer, provider, governingLaw, the quote copied verbatim. Answer as JSON {\"rows\": [...]}."].

d1:sensor#Source[sensor#kind -> "contract"; sensor#prompt -> readPrompt;
    sensor#text -> "Service Agreement between Beta GmbH (the Provider) and Acme AG (the Customer). This agreement is governed by Swiss law."].

?- d1[sensor#accepted(?I) -> ?V].
?- d1[sensor#rejected(?I) -> ?Why].
?- d1[sensor#notRead -> ?E].
```

The recorded run gives the same results as the program above. Write the rules out when a case needs its own acceptance rule, and use the library when it does not. A confidence reported by the model is recorded as a number and never read as truth. Where the serving stack supports it, decode every reply under a JSON schema derived from the closed set of allowed values.

Known uses

Assignment review with quote-checked extraction; campervan search (one classification per listing, from a fixed list or *unknown*); the judge in dialogue flows (one criterion about one answer); memory claims with quotes in FrameX-Agents; model measurements per object in the Q-calculus example. In the literature: neural perception feeding a reasoner, Kautz's Neuro|Symbolic (Kautz 2022); learning from data with symbolic output (van Harmelen and ten Teije 2019; van Bekkum et al. 2021); knowledge-graph construction by a language model (S. Pan et al. 2024).

Related patterns

The judge of Rules Steer the Model (Section 4.2) is a sensor for one answer. Critique and Repair (Section 4.3) handles the replies that fail a check. A Sensor whose readings are aggregated by rules over many objects (lists, counts) is where the hybrid's accuracy advantage was largest.

Variants

Several Readers, One Rule. The same question goes to several models or samples, and rules adjudicate: agreeing readings are accepted, a majority is accepted with its count recorded, and a disagreement becomes *unclear* and is handed on (Section 4.4); it is never settled by a coin toss. Voting and parallel sampling are published agent and workflow patterns (Liu et al. 2025; Schluntz and Zhang 2024), and self-consistency beat multi-agent debate at the same number of answers (Huang et al. 2024). Asking a judge twice with the answers in swapped order and requiring agreement is the same variant applied to position bias (Zheng et al. 2023).

Evidence

In a portfolio of 100 agreements (376 documents), a model alone answered single decisions and four-step chains correctly but produced exact exhaustive lists in 0 of 6 questions at low effort; with the model as sensor and the lists derived by rules, 10 of 10 questions were correct, with a completeness check leaving no silently wrong answer. The extraction was the expensive step: 43.6 cents once against 1.2 cents per question for the model alone. With an 8-billion-parameter open model, the strict quote

rule refused correct values and made extraction worse than the model alone (45 % against 60 %): the acceptance rule must fit the reader's ability to quote.

3.2 Pattern 2: Model as Author

Intent

Let the model write formal material (rules, decision tables, schemas, tests, dialogue flows, thresholds) and check every draft offline before a person adopts it. The rules never run what the model wrote unchecked.

Also known as

Rule drafting, policy-to-rules, specification-driven authoring, knowledge acquisition with a model.

Motivation

A policy sentence ("claims above 500 francs need a receipt") must become a rule, and an amended regulation must become a changed decision table. A model drafts the rule in seconds, and the expensive part is checking it. The pattern moves the cost to a check that is complete for the faults that matter: the draft must parse, may only use allowed heads, and must pass the regression cases and expectations, and the check reports which decisions the draft would change.

Applicability

Use it when

- the target notation has a complete check for the faults that matter (syntax, allowed vocabulary, structural anomalies of a decision table, regression cases);
- regression cases or expectations exist or can be written before the draft;
- a person is available to adopt what passes.

Where no complete check exists (semantic faults in a large rule base), the draft can only open a review (see Model as Critic, Section 5.2).

Structure

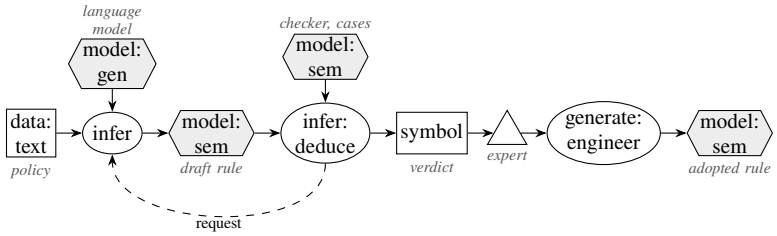

Participants

Model: drafts one artefact from a policy text, a change request or failing cases. *Checker*: parses, restricts heads, evaluates the candidate in a copy of the session against expectations and regression cases. *Rules*: request the check and derive *adoptable*, or a failure with the checker's message that drives the repair. *Expert*: adopts, edits or refuses.

Collaborations

The model's draft is stored as a receipt; a rule requests its check; the checker's receipt states regressions and changed decisions; a rule derives either a recommendation for the expert or a precise repair message for the model (Section 4.3).

Consequences

Writing becomes cheap and checking becomes the bottleneck; the quality of the engine's error messages becomes a design parameter. A draft outside the target language is refused at its first cell or token. What the organisation keeps is the adopted artefact, and the model can be exchanged. The pattern does not detect a draft that is well-formed, passes every case and is still wrong where no case covers it.

Implementation

Give the model a template to fill (a rule skeleton, the cell language of a decision table). Check in a copy of the session, and let a person adopt. Phrase engine messages for a reader who repairs: "separate the slots of a frame with ';', not ',"' led to more repairs than "expected ';' ".

The first request calls the model, the second the checker `_checkCandidate`. The operator fixes in an adapter file what the checker compares against: the vocabulary file, the hidden cases (`expect true: c3:NeedsReceipt.`) and the only head a draft may have. Its receipt carries `status` (`checked`; `rejected` before evaluation, for example for a syntax error or a head outside the allowed set; or `error`), and after evaluation `regressions` (a case that passed before fails now), `failedAfter` (cases failing after) and one line of `feedback`. The second half of the program is the repair of Critique and Repair:

```
// Pattern 2, Model as Author, with Pattern 6, Critique and Repair: the model drafts a
// rule, the checker measures it against hidden cases, and the checker's message
// drives a bounded repair. Nothing is adopted here; a person adopts.
world open.

p1:Policy[pending -> true; maxAttempts -> 3;
    text -> "A claim above 500 francs needs a receipt.";
    vocabulary -> "Class Claim with slot amount (a number). Write rules whose head is ?C:
        NeedsReceipt and nothing else. A rule is HEAD <- BODY. Slots: ?C:Claim[amount ->
        ?A]. Comparisons: > >= < <=."].

draftPrompt:Prompt[text -> "Write the FrameX rule for the policy sentence, using only the
    vocabulary. Answer as JSON {\"rule\": ...}."].
repairPrompt:Prompt[text -> "Your FrameX rule for the policy sentence failed the check.
    previousRule is your rule, problem is the checker's message. Write the corrected rule
    , using only the vocabulary. Answer as JSON {\"rule\": ...}."].

// Attempt 1: the model drafts.
external draft(?P, 1) using _llm with draftPrompt[policy -> ?T; vocabulary -> ?V;
       tool#order -> 2]
  <- ?P:Policy[pending -> true; text -> ?T; vocabulary -> ?V].

?P[rule(?N) -> ?T]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?P; tool#arg2 -> ?N; result("rule") -> ?T].

// Every draft is checked; the cases stay hidden from the model.
external check(?P, ?N) using _checkCandidate[candidate -> ?T; tool#order -> 1]
  <- ?P:Policy[rule(?N) -> ?T].

?P[passed(?N) -> true]
  <- ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P; tool#arg2 -> ?N;
       result("status") -> "checked"; result("regressions") -> false;
       result("failedAfter") -> 0].

?P[adoptable -> ?T] <- ?P[passed(?N) -> true; rule(?N) -> ?T].

// A draft fails in one of three ways; each failure keeps the checker's message.
?P[failed(?N) -> ?F]
  <- ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P; tool#arg2 -> ?N;
       result("status") -> ?St; result("feedback") -> ?F] AND ?St != "checked".

?P[failed(?N) -> ?F]
```

```
  <- ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P; tool#arg2 -> ?N;
       result("regressions") -> true; result("feedback") -> ?F].

?P[failed(?N) -> ?F]
  <- ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P; tool#arg2 -> ?N;
       result("failedAfter") -> ?K; result("feedback") -> ?F] AND ?K > 0.

// Pattern 6: ask again with the rule and the problem in the request, as a new
// request with the next number. Stop at the bound or when the message repeats.
?P[newProblem(1) -> true] <- ?P[failed(1) -> ?F].

?P[newProblem(?N) -> true]
  <- ?P[failed(?N) -> ?F] AND ?K := ?N - 1 AND ?P[failed(?K) -> ?G] AND ?F != ?G.

external draft(?P, ?M) using _llm with repairPrompt[policy -> ?T; vocabulary -> ?V;
       previousRule -> ?Old; problem -> ?F; tool#order -> 2]
  <- ?P:Policy[text -> ?T; vocabulary -> ?V; maxAttempts -> ?Max;
       rule(?N) -> ?Old; failed(?N) -> ?F; newProblem(?N) -> true]
     AND ?N < ?Max AND ?M := ?N + 1.

?P[stalled -> ?N]
  <- ?P[failed(?N) -> ?F] AND ?K := ?N - 1 AND ?P[failed(?K) -> ?F].

?P[gaveUp -> ?F] <- ?P:Policy[maxAttempts -> ?N; failed(?N) -> ?F].

?- p1[rule(?N) -> ?T].
?- p1[failed(?N) -> ?F].
?- p1[adoptable -> ?T].
?- p1[stalled -> ?N].
```

Figure 3 draws the program as a rule graph. `adoptable` is a recommendation for a person; nothing in the program adopts a rule.

Library form

With the library `author.fx` (Section B) loaded first and the adapter of the checker named to the host, the program states the policy, its vocabulary, the two prompts and the bound. The library drafts, checks, repairs and stops; the repair of Critique and Repair is part of it:

```
// Patterns 2 and 6 in library form: load author.fx first; the host needs ../
    author_adapter.py.
world open.

draftPrompt:Prompt[text -> "Write the FrameX rule for the policy sentence, using only the
    vocabulary. Answer as JSON {\"rule\": ...}."].
repairPrompt:Prompt[text -> "Your FrameX rule for the policy sentence failed the check.
    previousRule is your rule, problem is the checker's message. Write the corrected rule
    , using only the vocabulary. Answer as JSON {\"rule\": ...}."].

p1:author#Policy[author#maxAttempts -> 3;
    author#draftPrompt -> draftPrompt; author#repairPrompt -> repairPrompt;
    author#text -> "A claim above 500 francs needs a receipt.";
    author#vocabulary -> "Class Claim with slot amount (a number). Write rules whose head
        is ?C:NeedsReceipt and nothing else. A rule is HEAD <- BODY. Slots: ?C:Claim[
        amount -> ?A]. Comparisons: > >= < <=."].

?- p1[author#failed(?N) -> ?F].
?- p1[author#adoptable -> ?T].
```

Known uses

Rule drafting from one policy sentence with six hidden cases; decision-table drafts and change requests in the DMN round trip, checked against every case and reviewed; the EU Regulation 261/2004 case,

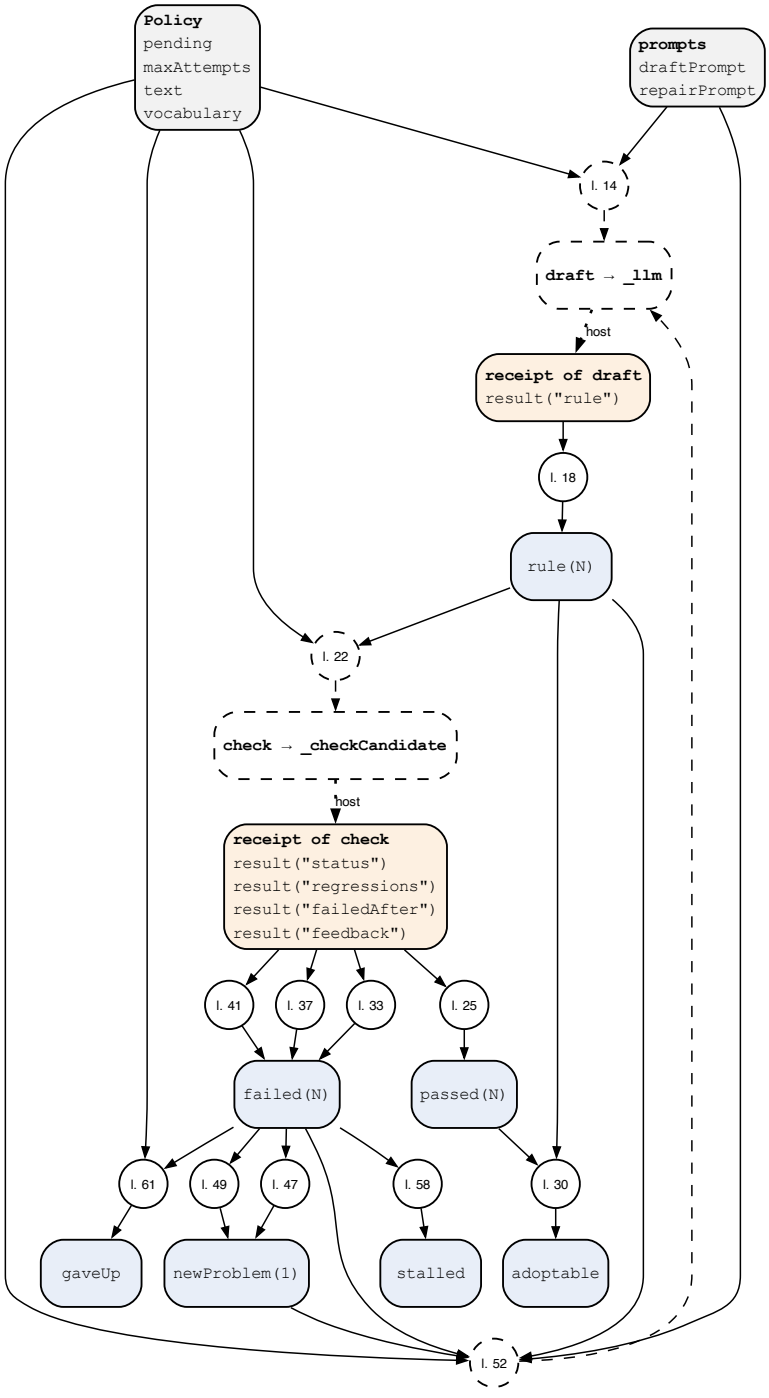

Figure 3: Rule graph of `p02-author.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

where a draft outside the cell language was refused and a change request fixed six of six failing cases, broke none and passed review; mapping suggestions between vocabularies kept as recorded proposals. In the literature: domain-model acquisition with experts in LLM-Modulo (Kambhampati et al. 2024); action-schema generation (Rani et al. 2026); the engineered model of an actor (pattern 1c) of van Bekkum et al. (2021), with the model in the engineer's role.

Related patterns

Critique and Repair (Section 4.3) closes the loop with the checker's message. Division by Competence (Section 4.4) supplies drafts from recurring model decisions. Model as Critic (Section 5.2) applies where the check is incomplete.

Evidence

In five live runs, a model drafted a rule from one policy sentence; every run ended with a rule that passed all six hidden cases, three at the first attempt, two after one repair of the same syntax slip. From the same faulty drafts, a repair succeeded at the first request in 5 of 10 runs with the message "expected ';' " and in 10 of 10 with a message naming the fix. Across 36 drafts for 12 policies, messages that name the fix raised repairs within two attempts from 8 of 15 to 14 of 15. In a checking study, static checks over rules found every seeded structural fault at every size at no model cost, while model reviews of code and prompts found an unreachable rule in only 2 or 3 of 10 artefacts at 64 rules.

3.3 Pattern 3: Model as Translator

Intent

Put the model between a person and the rules, in both directions: the person's words into a request the rules can answer, and the rules' answer and proof into words for that person.

Also known as

Natural-language interface, verbalisation, explanation, the model calls the rules as a tool.

Motivation

A clerk asks "can we pay this invoice?" and expects an answer in words, with the reason. The rules hold the answer and its proof but speak in queries and proof trees. The model speaks the clerk's language but should not decide. The pattern gives the model the conversation and the rules the decision: the model calls the rules as a tool, receives the verdict, the decisive quotes and the missing evidence, and phrases them; the host checks that the reply uses the rules' decision word.

Applicability

Use it when people ask in their own words, when answers must come with reasons a non-expert can read, or when a model-based assistant needs a verified source of decisions. Prefer fixed templates over a model where the audience and the wording are known (sentences generated from a decision table need no model).

Structure

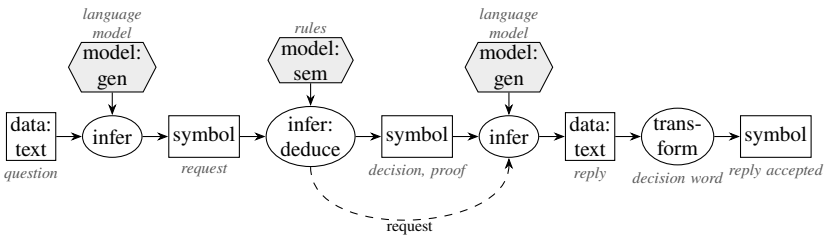

Participants

Person; *Model* (translates); *Rules* (answer, prove, name what is missing: `query`, `explain`, `why_not`, `preview`, `evidence_needs`); *Host* (checks the reply against the answer).

Collaborations

The model turns the question into a call of a read-only operation from a closed set, or, where free queries are admitted, into a candidate query that the rules validate like any draft (Section 3.2) before it runs; the rules answer with a proof; the model phrases the answer; a check binds the phrasing to the answer (decision word, cited proof nodes). In the other direction the rules can be offered read-only to any model client as a server.

Consequences

The person gets an answer in words with a proof behind it. A mistranslated question gives a correct answer to the wrong question, so a back-translation shown to the person is the main safeguard. Prose that adds facts the proof lacks must be rejected, which requires the phrasing to cite the proof. The two forms of request differ in what can be checked. A call of a closed operation set keeps invariant 1 without qualification and lets a rule refuse a request with a reason that names the field. A free query is more expressive, but a check can only establish its form, its vocabulary and that it is read-only; it inherits the consequences of Model as Author, and a query that validates can still ask the wrong question.

Implementation

Prefer a closed set of read-only operations whose arguments are identifiers; the host builds the executable query from the call. Admit a free query only through `validate` (syntax, allowed classes and slots, no assertions) and record it as a candidate with its receipt. Keep the model's view small: the verdict, the decisive quotes and the missing evidence, without the documents. Offer the rules as read-only tools with a hash of the files that answered. Bind every sentence of an explanation to a proof node, or use templates.

In rules, the model's first call turns the question into a request, the request runs only when every field names a member of a closed list, a refusal names the field, and the rules derive the decision. The second call names the decision word as `mustSay`, which the model does not see; the host records whether the reply contains it (`saysDecision`), and only such a reply is accepted. The test is a substring test that ignores case: "not admitted" contains "admitted". Where one decision word contains another, ask for the decision word in a field of its own as well and compare it exactly. The program answers admission questions; the operations *requirements* and *missing* are recognised and refused with a reason:

```
// Pattern 3, Model as Translator: a question becomes a checked request, the rules'
// decision becomes a sentence that must state it.
world open.
closed class Operation. closed class Programme. closed class Applicant.
closed slot admitted.

"admission":Operation. "requirements":Operation. "missing":Operation.
"mbi":Programme. "mgt":Programme.
"nora":Applicant. "tomas":Applicant.
"nora"[admitted -> "mbi"].
q1:Question[text -> "Hi, I am Nora. Am I admitted to the Master in Business Innovation?"].

translatePrompt:Prompt[text -> "Map the question to one call. Answer as JSON {\"op\":
    admission, requirements or missing, \"programme\": mbi or mgt, \"applicant\": the
    lowercase first name}."].
phrasePrompt:Prompt[text -> "Tell the person the decision in one friendly sentence that
    uses the decision word. Answer as JSON {\"text\": ...}."].

// From words to a request.
external translate(?Q) using _llm with translatePrompt[question -> ?T]
  <- ?Q:Question[text -> ?T].

?Q:Request[op -> ?O; programme -> ?P; applicant -> ?A]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?Q; result("op") -> ?O;
        result("programme") -> ?P; result("applicant") -> ?A].
```

```
// The request runs only if every field names something the rules know.
?Q[runs -> true] <- ?Q:Request[op -> ?O; programme -> ?P; applicant -> ?A]
   AND ?O:Operation AND ?P:Programme AND ?A:Applicant.

?Q[refused -> "unknown operation"] <- ?Q:Request[op -> ?O] AND NOT ?O:Operation.
?Q[refused -> "unknown programme"] <- ?Q:Request[programme -> ?P] AND NOT ?P:Programme.
?Q[refused -> "unknown applicant"] <- ?Q:Request[applicant -> ?A] AND NOT ?A:Applicant.
?Q[refused -> "this program answers admission questions only"]
  <- ?Q:Request[runs -> true; op -> ?O] AND ?O != "admission".

// The rules decide.
?Q[decision -> "admitted"]
  <- ?Q:Request[runs -> true; op -> "admission"; programme -> ?P; applicant -> ?A]
     AND ?A[admitted -> ?P].

?Q[decision -> "not admitted"]
  <- ?Q:Request[runs -> true; op -> "admission"; programme -> ?P; applicant -> ?A]
     AND NOT ?A[admitted -> ?P].

// From the decision to words; the reply counts only if it says the decision word.
external phrase(?Q) using _llm with phrasePrompt[decision -> ?D; mustSay -> ?D]
  <- ?Q:Request[decision -> ?D].

?Q[reply -> ?Text]
  <- ?Q:Request AND ?X[tool#name -> "_llm"; tool#arg1 -> ?Q;
       result("saysDecision") -> true; result("text") -> ?Text].

?Q[replyRefused -> "the reply does not state the decision"]
  <- ?Q:Request AND ?X[tool#name -> "_llm"; tool#arg1 -> ?Q;
       result("saysDecision") -> false].

?- ?Q:Request[decision -> ?D].
?- ?Q[reply -> ?Text].
?- ?Q[refused -> ?Why].
```

Figure 4 draws the program as a rule graph.

Library form

With `translator.fx` (Section B) the program declares the closed list of each answer field and keeps its decision rules. The library reads the fields, refuses a request with the name of the field that lies outside its list, and phrases the decision under `mustSay`. Decision rules read only valid fields (`tr#valid`). A rule that asks whether *all* fields are valid negates over the receipts, and the reply request depends on the decision; FrameX refuses that as a negative cycle, so the check is spelled out per field:

```
// Pattern 3 in library form: load translator.fx first.
world open.
closed slot admitted.

translatePrompt:Prompt[text -> "Map the question to one call. Answer as JSON {\"op\":
    admission, requirements or missing, \"programme\": mbi or mgt, \"applicant\": the
    lowercase first name}."].
phrasePrompt:Prompt[text -> "Tell the person the decision in one friendly sentence that
    uses the decision word. Answer as JSON {\"text\": ...}."].
tr#phrasing[tr#prompt -> phrasePrompt].

"op"[tr#allows -> "admission"].
"programme"[tr#allows -> "mbi"; tr#allows -> "mgt"].
"applicant"[tr#allows -> "nora"; tr#allows -> "tomas"].
"nora"[admitted -> "mbi"].

```

```
q1:tr#Question[tr#prompt -> translatePrompt;
    tr#text -> "Hi, I am Nora. Am I admitted to the Master in Business Innovation?"].

// The rules decide; this part stays with the program.
?Q[tr#decision -> "admitted"]
  <- ?Q:tr#Question[tr#valid("op") -> "admission"; tr#valid("programme") -> ?P;
       tr#valid("applicant") -> ?A] AND ?A[admitted -> ?P].
?Q[tr#decision -> "not admitted"]
  <- ?Q:tr#Question[tr#valid("op") -> "admission"; tr#valid("programme") -> ?P;
       tr#valid("applicant") -> ?A] AND NOT ?A[admitted -> ?P].

?- ?Q[tr#decision -> ?D].
?- ?Q[tr#reply -> ?Text].
?- ?Q[tr#refused -> ?F].
```

Known uses

The assignment-review chat model calls a reviewer tool that runs the rules; FrameX as a read-only server for model clients; a research agent that keeps the model's explanation beside a checked conclusion over a knowledge graph; sentences generated from decision-table rows without a model. In the literature: the model formalises, a solver answers, in Logic-LM (L. Pan et al. 2023), LINC (Olausson et al. 2023) and SatLM (Ye et al. 2023); knowledge-graph question answering and graph-to-text (S. Pan et al. 2024); solver-aided methods (Yang et al. 2025).

Related patterns

Model as Sensor (Section 3.1) translates documents where this pattern translates questions. Rules Steer the Model (Section 4.2) adds a flow to the conversation.

Variants

Read-Back Check. An independent reader that sees only the written reply reads its facts back, and rules compare them with the facts the reply was meant to state: the decision, the amounts, the dates. It is the translator's check applied to the output side, and it catches a changed value that a decision word alone does not. The reader must differ from the writer (another model or another prompt), or it repeats the writer's misreading. Reported in several Hack Apertus submissions (Sturmglocke, Lightpath, SONA with entailment in both directions); not built in FrameX. In one of them, a 70B second reader cut instructions with a changed meaning from 4.6 % to 0.2 % of 2,078 (result file committed in the submission).

Evidence

On a clean five-document case, a model alone answered the assignment-review question correctly in 12 of 12 runs at lower cost and time than the hybrid (0.28 against 0.72 cents per case). On easy cases the translator earns its cost through the proof, the stated missing evidence and the fixed decision word.

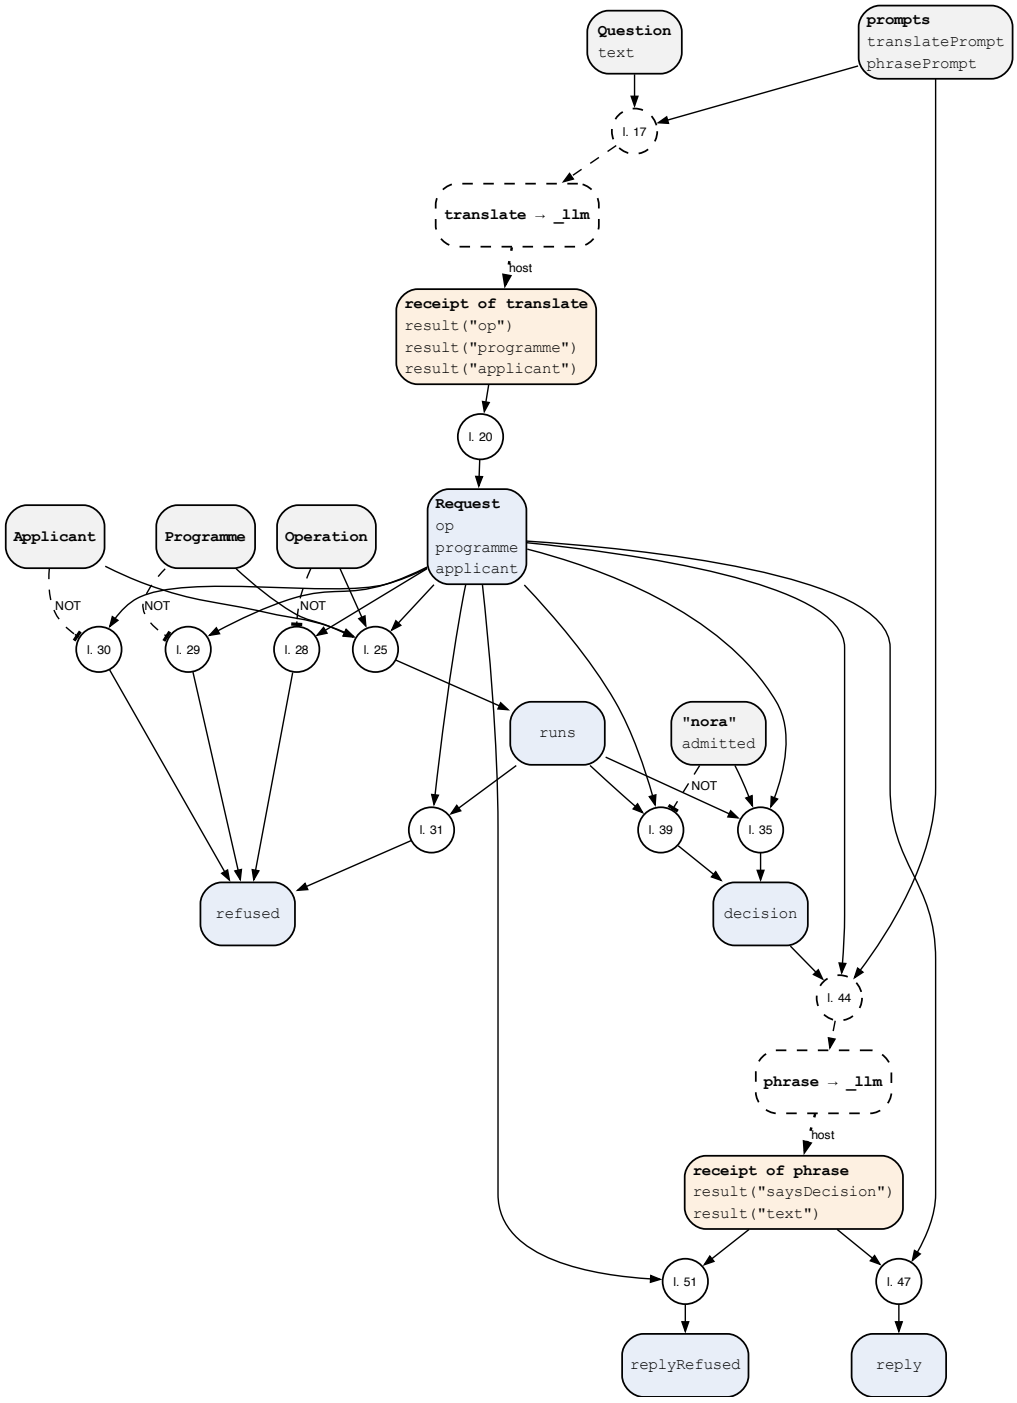

Figure 4: Rule graph of `p03-translator.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

3.4 Pattern 12: Model Selects, Rules Write

Intent

Let the model choose and phrase, but never state content: it returns identifiers of passages that the host numbered, and words around values it never sees. Every passage and every value that reaches a person comes from facts, word for word.

Also known as

Model points, code speaks; model phrases, code fills; extractive answer by identifier; slot filling; template generation with a model.

Motivation

An applicant asks which documents the Master in Business Innovation requires and what the application costs. The answer must quote the regulation exactly and state the fee exactly. A model that writes the answer can paraphrase a requirement into a different one or round a fee, and a check afterwards (a quote that must occur, a decision word that must appear) catches only what it looks for. The pattern moves the check before the writing: the model sees the passages under identifiers and answers with identifiers, and it phrases a sentence around a gap that the rules fill. A requirement the model invents has no identifier, and a fee it never saw it cannot change.

Applicability

Use it when

- the content of an answer exists as facts or source text (regulations, a catalogue, computed amounts, dates) and only its selection or wording needs a model;
- a wrong word in the content is costly, so a reply must be correct by construction rather than by check;
- the source can be split into passages that a person reads on their own.

Do not use it when the answer must combine or summarise what no single passage says; then use Model as Translator (Section 3.3) with its check of the reply.

Structure

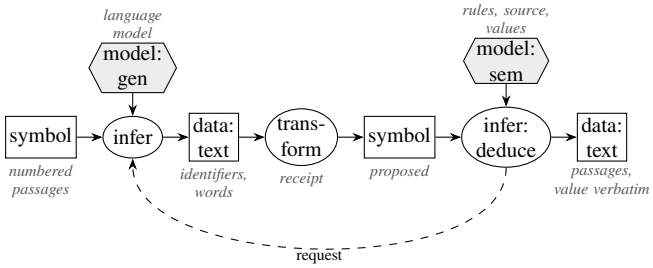

Participants

Host: splits the source into passages and numbers them, calls the model, stores the reply as a receipt. *Model*: chooses identifiers from the list it was given, and phrases words around a gap. *Rules*: request the selection and the phrasing, accept an identifier only if it names a passage of the source, and write the answer from the passages and the values. *Person*: reads passages and values as they stand in the source.

Collaborations

A rule requests a selection for a question with the numbered passages; the receipt holds identifiers; a rule shows the passage of each identifier that belongs to the source and refuses any other with its reason. A second request asks for words before and after a value without the value; a rule inserts the value between them, unless the words state a number of their own.

Consequences

No sentence of the model reaches the person as content: an invented passage has no identifier and is refused, a fee is never retyped. What the model can still get wrong is the choice: a passage that does not answer the question, or one that is missing. That is a fault of relevance, not of truth, and a person sees it in the source's own words. Answers are as fluent as the source; text that combines several passages is out of reach. The cost is one call per selection and per phrasing.

Implementation

Number the passages in the host, where the source is split; the model sees `a1: ...`, the rules know `"a1"` as a `Passage` with its source and text. Declare the class closed, so that an identifier outside it is refused rather than left *unknown*. For values, the request carries what the value is about but never the value; the rules join the words and the value with `concat`, and refuse words that contain a digit, so that a number the model made up cannot sit beside the right one. The program answers one question from three passages and states one fee:

```
// Pattern 12, Model Selects, Rules Write: the model returns identifiers and phrasing;
// every passage and every value shown to the person comes from facts.
world open.
closed class Passage. closed class Digit. closed slot hasDigit.

// The source, split into passages by the host, and the same passages numbered for the model.
"a1":Passage[source -> reg; text -> "Applicants hold a recognised Bachelor degree."].
"a2":Passage[source -> reg; text -> "Applicants prove English at level C1."].
"a3":Passage[source -> reg; text -> "The programme starts in the autumn semester."].
reg[numbered -> "a1: Applicants hold a recognised Bachelor degree. a2: Applicants prove English at level C1. a3: The programme starts in the autumn semester."].

// A question about the source, and a value an answer must state.
q1:Question[source -> reg; text -> "Which documents do I need for the MBI?"].
fee1:Amount[question -> q1; value -> "CHF 250"; about -> "the application fee"].

selectPrompt:Prompt[text -> "Choose the passages that answer the question. Answer as JSON {\"rows\": [{\"passage\": an identifier from the list}]}."].
phrasePrompt:Prompt[text -> "Write one sentence that tells the applicant the amount named in about; the system inserts the amount. Answer as JSON {\"before\": the text before the amount, \"after\": the text after it}."].

// Request 1: the model sees the numbered passages and returns identifiers only.
external select(?Q) using _llm with selectPrompt[question -> ?T; passages -> ?N]
  <- ?Q:Question[source -> ?S; text -> ?T] AND ?S[numbered -> ?N].

// Accept: an identifier shows its own passage, word for word.
?Q[shows(?I) -> ?Text]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?Q; row(?I, "passage") -> ?P]
     AND ?Q:Question[source -> ?S] AND ?P:Passage[source -> ?S; text -> ?Text].

?Q[refused(?I) -> "not a passage of this source"]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?Q; row(?I, "passage") -> ?P]
     AND ?Q:Question AND NOT ?P:Passage.

// Request 2: the model phrases around the value; the request does not carry the value.
external phrase(?V) using _llm with phrasePrompt[about -> ?A]
  <- ?V:Amount[about -> ?A].

"0":Digit. "1":Digit. "2":Digit. "3":Digit. "4":Digit.
"5":Digit. "6":Digit. "7":Digit. "8":Digit. "9":Digit.

?R[hasDigit -> true]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?V; result("before") -> ?B]
     AND ?V:Amount AND ?D:Digit
```

```
    AND ?N := occurrences(?B, ?D) AND ?N > 0.
?R[hasDigit -> true]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?V; result("after") -> ?F]
     AND ?V:Amount AND ?D:Digit
     AND ?N := occurrences(?F, ?D) AND ?N > 0.

// Write: the rules put the value between the model's words.
?Q[answer -> ?Text]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?V; result("before") -> ?B; result("after") ->
       ?F]
     AND ?V:Amount[question -> ?Q; value -> ?X] AND NOT ?R[hasDigit -> true]
     AND ?Text := concat(concat(?B, ?X), ?F).

?V[rejected -> "the phrasing states a number of its own"]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?V; hasDigit -> true] AND ?V:Amount.

// A failed call shows nothing and keeps the host's reason.
?X[notAnswered -> ?E]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?X; tool#status -> "failed"; error -> ?E].

?- q1[shows(?I) -> ?Text].
?- q1[refused(?I) -> ?Why].
?- q1[answer -> ?A].
?- fee1[rejected -> ?Why].
```

Figure 5 draws the program as a rule graph. The two requests have their own subjects, the question and the amount, so that each receipt is read only by its own rules. Where the serving stack decodes under a JSON schema, an enumeration of the passage identifiers makes an identifier outside the list impossible to produce; the rule of line 28 stays as the check that does not depend on the stack. There is no library form yet.

Known uses

In FrameX, the program above; decision-table rows become sentences from templates without any model, the limit case of this pattern. In the submissions to the Hack Apertus hackathon (October 2026), as described in the jury assessments, models chose passage or line identifiers that the application then quoted (QuietLetter, Beneficia, Comfort, StimmFolgen, Fedlex, ClearDesk, a meditation retrieval system) or wrote with placeholders that code filled with numbers and terms (Apertus QA, Bon·Appétit, Sturmglocke, Velum). In the literature: template-based generation, in which the content is fixed before the text is planned (Reiter and Dale 2000); guardrails that restrict a reply to retrieved passages (Rebedea et al. 2023). In these the selection is checked afterwards or the template is written by hand; here the model fills neither content nor value.

Related patterns

Model as Sensor (Section 3.1) reads a source and checks quotes afterwards; this pattern makes the quote the only possible answer. Model as Translator (Section 3.3) checks a written reply for the decision word; here the reply cannot contain the decision. Deriving the options first under Model Proposes, Rules Dispose (Section 4.1) is the same move for actions.

Evidence

The program is checked against recorded receipts only; FrameX has no measurement of its own. Figures from the hackathon are taken from result files committed in the submissions' repositories at a named commit; we did not rerun them, and figures that appear only in a README, report or pitch are left out. In an alert generator for civil protection, Apertus v1.5 8B filled code-built slots with a locked official hazard term; the official term and all facts appeared verbatim in 99 % of 120 alerts, against 45 % when the same model translated freely (70B: 98 % against 71 %). There the model still saw the facts and code checked them afterwards, a weaker form of the pattern. In a beneficial-ownership tool, 8B in Q4_K_M on a CPU, citing line identifiers that code verified, got 14 of 14 cases right, with no

held-out set. Selection does not close every gap: in court-decision anonymisation, 8B reached a recall of 0.764 against 0.894 for 70B on 537 decisions, and in a statistics question-answering system a rule baseline without any model matched 70B on the main set.

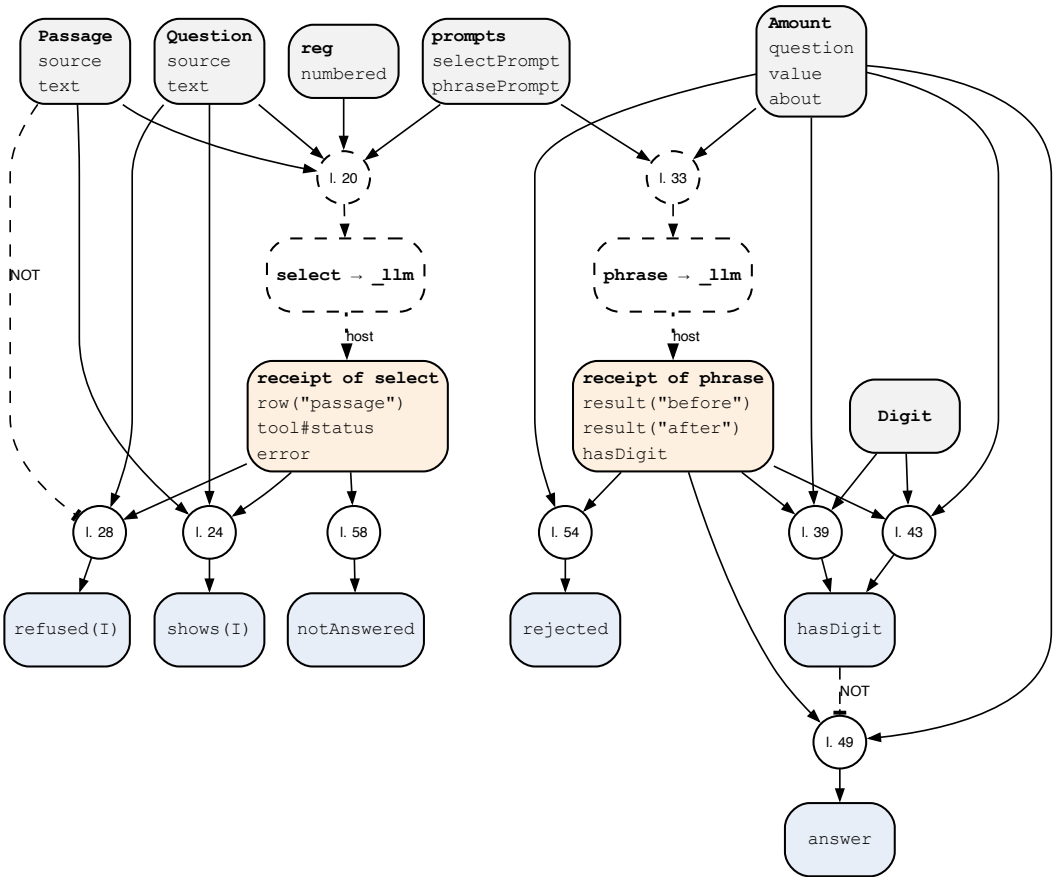

Figure 5: Rule graph of `p12-select.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

4 Patterns of control: who decides

4.1 Pattern 4: Model Proposes, Rules Dispose

Intent

Let the model propose an answer, an action or a choice, and let a rule admit it, refuse it with a reason, or narrow what may be proposed.

Also known as

Guardrail, action gating, policy enforcement point, bounded recommendation.

Motivation

An agent proposes to write a file, a model recommends a portfolio weight, an assistant proposes to send a refund. Each proposal may be reasonable, but some must never happen: a write outside the task, a weight change below a materiality threshold, a refund above an authority limit. The pattern keeps the model's initiative and puts every effect behind a rule that can say no and say why.

Applicability

Use it when the model acts (tool calls, messages, transactions), when the space of admissible outcomes is known even if the best one is not, or when a proposal must rest on evidence available at a given time.

Structure

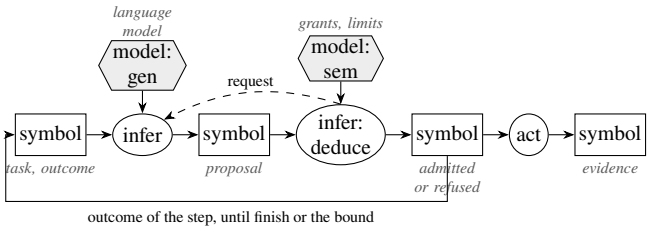

Participants

Model (proposes one action or value); *Rules* (capabilities, preconditions, limits, evidence requirements); *Host* (executes only what the rules admit); *Person* (approves what the rules require).

Collaborations

The proposal arrives as a receipt; rules derive admission or a refusal with its reason; a refusal returns to the model, which replans; the step bound stops a task whose proposals keep being refused. A task is complete when a rule derives it from recorded evidence; a model saying it is done is one more proposal.

Consequences

Effects become auditable and bounded; the model may still waste turns on refused proposals. Conflicts between constraints are reported to a person, and no priority order settles them silently. The pattern bounds what can happen; whether an admitted proposal is the best one stays open.

Implementation

Derive the options first where they are few (configurations, offers in a legal band) and let the model choose among them; a rule checks that the choice is one of them. Export the closed set of allowed values as the structured-output schema of the call, and still check the receipt.

Each step is a new request `propose(?T, ?N)` whose receipt becomes a proposal. A proposal is admitted only for an action the task grants; any other is refused with a reason. An admitted look-up runs as a query on the rules themselves (`_query`), and the next step carries the outcome of the last one, the refusal or the looked-up result. The task is complete when an admitted *finish* rests on evidence, and it stops at the step bound otherwise:

```
// Pattern 4, Model Proposes, Rules Dispose: one action per step, admitted or refused.
world open.
closed slot grants.

// The requirements, and a task with the actions it grants.
mbi[requires -> "recognised Bachelor"; requires -> "English C1"].
t1:Task[goal -> "Tell Nora which documents the Master in Business Innovation requires.";
    actions -> "lookUp: read the requirements; sendOffer: send Nora an admission offer;
        finish: answer Nora with what you found"; maxSteps -> 4].
t1[grants -> "lookUp"; grants -> "finish"].

proposePrompt:Prompt[text -> "Propose the next action for the goal. outcome says what
    became of your last action; once it answers the goal, choose finish. Answer as JSON
    {\"action\": one of the actions}."].

// Step 1: the model proposes one action; outcome is "none" before the first step.
external propose(?T, ?N) using _llm with proposePrompt[goal -> ?G; actions -> ?Acts;
       outcome -> "none"; tool#order -> 2]
  <- ?T:Task[goal -> ?G; actions -> ?Acts] AND ?N := 1.
?R:Proposal[task -> ?T; step -> ?N; action -> ?K]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?T; tool#arg2 -> ?N;
       result("action") -> ?K].

// Dispose: an action is admitted only if the task grants it.
?A[allowed -> true] <- ?A:Proposal[task -> ?T; action -> ?K] AND ?T[grants -> ?K].
?A[refused -> "action not granted for this task"]
  <- ?A:Proposal[task -> ?T; action -> ?K] AND NOT ?T[grants -> ?K].

// An admitted lookUp runs as a query on the rules themselves.
external lookUp(?T) using _query[query -> "?- mbi[requires -> ?X].";
       tool#order -> 1]
  <- ?A:Proposal[task -> ?T; action -> "lookUp"; allowed -> true].

?T[evidence -> true]
  <- ?R[tool#name -> "_query"; tool#arg1 -> ?T; tool#status -> "succeeded"].

// The next step carries the outcome of the last one; maxSteps stops the loop.
?T[next(?M) -> ?W]
  <- ?A:Proposal[task -> ?T; step -> ?N; refused -> ?W]
     AND ?T[maxSteps -> ?Max] AND ?N < ?Max AND ?M := ?N + 1.

?T[next(?M) -> ?X]
  <- ?A:Proposal[task -> ?T; step -> ?N; action -> "lookUp"; allowed -> true]
     AND ?R[tool#name -> "_query"; tool#arg1 -> ?T; result("text") -> ?X]
     AND ?T[maxSteps -> ?Max] AND ?N < ?Max AND ?M := ?N + 1.

external propose(?T, ?N) using _llm with proposePrompt[goal -> ?G; actions -> ?Acts;
       outcome -> ?W; tool#order -> 2]
  <- ?T:Task[goal -> ?G; actions -> ?Acts; next(?N) -> ?W].

// Completion only on recorded evidence.
?T[verdict -> "completed"]
  <- ?A:Proposal[task -> ?T; action -> "finish"; allowed -> true]
     AND ?T[evidence -> true].

?T[verdict -> "stopped at the step bound"]
  <- ?A:Proposal[task -> ?T; step -> ?N; action -> ?K] AND ?T[maxSteps -> ?N]
     AND ?K != "finish".

?- ?A:Proposal[step -> ?N; action -> ?K].
?- ?A[refused -> ?Why].
?- t1[verdict -> ?V].
```

Figure 6 draws the program as a rule graph.

Library form

With `agent.fx` (Section B) the program states the task, its prompt and what it grants, and keeps what each granted action does: here the query behind `lookUp`, its outcome for the next step, and the evidence that completion rests on. Step loop, admission, refusal, bound and verdict come from the library:

```
// Pattern 4 in library form: load agent.fx first.
world open.

mbi[requires -> "recognised Bachelor"; requires -> "English C1"].
proposePrompt:Prompt[text -> "Propose the next action for the goal. outcome says what
    became of your last action; once it answers the goal, choose finish. Answer as JSON
    {\"action\": one of the actions}."].

t1:ag#Task[ag#maxSteps -> 4; ag#prompt -> proposePrompt;
   ag#goal -> "Tell Nora which documents the Master in Business Innovation requires.";
   ag#actions -> "lookUp: read the requirements; sendOffer: send Nora an admission offer;
        finish: answer Nora with what you found"].
t1[ag#grants -> "lookUp"; ag#grants -> "finish"].

// What the action lookUp does, and what it reports; this stays with the program.
external lookUp(?T) using _query[query -> "?- mbi[requires -> ?X]."; tool#order -> 1]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#action -> "lookUp"; ag#allowed -> true].
?T[ag#outcome(?N) -> ?X]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#step -> ?N; ag#action -> "lookUp"; ag#allowed ->
      true]
     AND ?R[tool#name -> "_query"; tool#arg1 -> ?T; result("text") -> ?X].
?T[ag#evidence -> true]
  <- ?R[tool#name -> "_query"; tool#arg1 -> ?T; tool#status -> "succeeded"].

?- ?A:ag#Proposal[ag#step -> ?N; ag#action -> ?K].
?- ?A[ag#refused -> ?Why].
?- t1[ag#verdict -> ?V].
```

Known uses

FrameX-Agents (one action at a time, admitted against the task's capabilities, completion from evidence); allocation recommendation with point-in-time evidence and a materiality gate; casework actions that need a native authorisation proof. In the literature: rule-based safeguards around a model (Washizaki et al. 2020; Washizaki et al. 2022); hard critics in LLM-Modulo (Kambhampati et al. 2024); guardrails (Rebedea et al. 2023; Liu et al. 2025); gates in prompt chains (Schluntz and Zhang 2024).

Related patterns

Critique and Repair (Section 4.3) handles refused proposals. Rules Steer the Model (Section 4.2) decides what the model is asked to propose next.

Variants

Simulator Disposes. A simulation instead of a rule decides about the proposal: the host runs the proposed plan in a simulator, and its receipt (reached, collided, within budget) is what the rules accept or refuse. The simulator is an external tool like any other, so the receipt keeps its version and inputs, and the rules still make the decision. It is close to the preview in a copy of the session of Model as Hypothesis Generator (Section 5.1). Reported in Hack Apertus submissions (Lightpath, Sturmglocke, Rescue Briefings, Forester); not built in FrameX. In one of them, simulated violations were 16.4 % when the model chose the path and 10.5 % when a rule engine did, over three seeds (result file committed in the submission).

Evidence

One demonstration run. No measurement isolates this pattern yet. It runs in every action of the agent framework and of the allocation application (Table 5). The consequences above describe how those systems are designed and have not been measured.

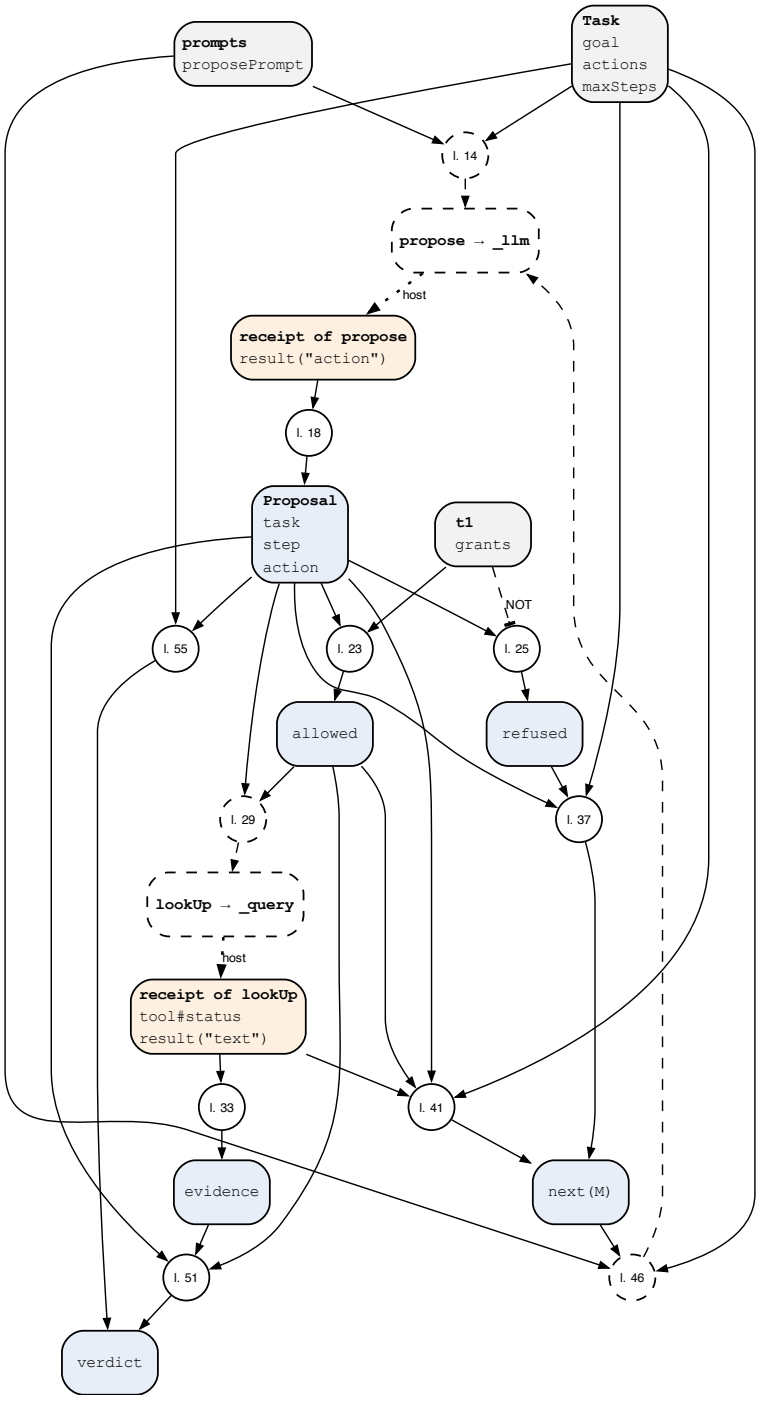

Figure 6: Rule graph of `p04-dispose.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

4.2 Pattern 5: Rules Steer the Model

Intent

Let rules decide what the model is asked, when, with which context and which tools. The model never decides the next step.

Also known as

Dialogue flow, guided conversation, symbolic orchestration, state machine with model leaves, rule-selected context.

Motivation

A triage call must ask about emergency signs first, ask one thing at a time, ask again when an answer is unclear, and hand over to a nurse when a sign is present. A model asked to run the whole call from a long prompt skips questions as the protocol grows. The pattern turns the protocol into rules: a flow derives the current step, the model phrases its question, a judge decides whether the answer satisfies the step, and rules decide branching, re-asking, hand-over and the end.

Applicability

Use it when the process is known (a protocol, an interview, an intake, an agent's phases), when the number of rules a model would have to apply by itself grows with the case, or when what the model may see must be decided by policy (permissions, point-in-time availability, jurisdiction).

Structure

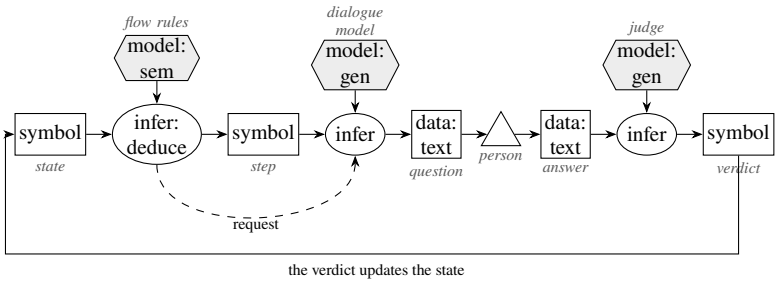

Participants

Rules (flow steps, slots, branches, phase, permitted tools, selected sources, missing evidence); *Dialogue model* (phrases the current question); *Judge* (one criterion about one answer, a verdict from a fixed set); *Person* (answers, or takes over).

Collaborations

Each turn: rules derive the step; the dialogue model asks; the person answers; the judge returns *answered*, *not answered* or *unclear* with a reason; rules decide the next step. An unanswered question never decides a branch.

Consequences

The protocol is executed completely regardless of its length; the model's work per call stays small, so a weaker model suffices. Every turn has a receipt. The cost is one or more calls per turn, and a judge's verdict that contradicts its own reason is a failure the rules must catch.

Implementation

Ask one thing per question. Name verdict labels that cannot be read as content (*answered* instead of *yes*). Let the judge read the earlier answers it needs. Derive the missing evidence for a goal and ask for exactly that. Select only the rules that apply or nearly apply to the case as context.

A flow is data that the rules of FrameX's flow library read. The library is the file `flow.fx` of FrameX, loaded before the program; the program itself contains only facts:

```
// Pattern 5, Rules Steer the Model: a dialogue flow written as facts. FrameX's flow
// library (python/framex/policies/flow.fx, loaded first) holds the rules that decide
// which step is active, request _turn (the dialogue model asks the question and returns
// the person's answer) and _judge (a judge model decides the step's criterion), ask
// again after a negative or unclear verdict, and repair a reply that breaks the format
// (Pattern 6).
world open.

advice:flow#Flow[flow#start -> programme].

// Ask which programme is meant; at most two attempts, then hand on to an officer.
programme:flow#Collect[flow#ask -> "Which Master do you mean?";
    flow#criterion -> "Does the person name Business Innovation or Management?";
    flow#slot -> programmeGiven; flow#maxAttempts -> 2;
    flow#ifSlot -> programmeGiven; flow#to -> language;
    flow#otherwise -> officer].

// Ask for the language certificate; at most three attempts (the default).
language:flow#Collect[flow#ask -> "Which English certificate do you hold?";
    flow#criterion -> "Does the person name an English certificate at level C1?";
    flow#slot -> c1Named;
    flow#ifSlot -> c1Named; flow#to -> check; flow#otherwise -> missing].

// Three ways to end.
check:flow#Say[flow#text -> "Thank you. We check your record."; flow#next -> flow#end].
missing:flow#Say[flow#text -> "An English certificate at level C1 is missing."; flow#next
    -> flow#end].
officer:flow#Say[flow#text -> "An admissions officer will contact you."; flow#next ->
    flow#end].

?- ?S:flow#Active.
?- workflow[status -> ?St].
```

Figure 7 draws the flow as a diagram. A `flow#Collect` step asks `flow#ask`, lets a judge decide `flow#criterion` about the answer and records the outcome in the slot named by `flow#slot`. `flow#ifSlot`, `flow#to` and `flow#otherwise` branch on a slot; `flow#maxAttempts` bounds the attempts (three by default); a `flow#Say` step shows a text, and `flow#end` ends the flow. The library derives which step holds the token (`flow#Active`) and which attempt is due (`flow#askNow`), and it issues two model calls per attempt. `advice:flow#Flow[flow#start -> programme]` declares the flow and its first step. A rejected attempt (`flow#rejected`) is asked again while attempts remain below the step's limit (`flow#limit`, from `flow#maxAttempts`). Four of the library's rules, shown only to make the chain visible:

```
// From FrameX's flow.fx (library; a program does not repeat these).
external turn(?S, ?A) using _turn[question -> ?Q; attempt -> ?A; tool#order -> 1]
    <- ?S[flow#askNow -> ?A; flow#ask -> ?Q].

?S[flow#answer(?A) -> ?Text]
    <- ?R[tool#arg1 -> ?S; tool#arg2 -> ?A; result("answer") -> ?Text].

external judge(?S, ?A) using _judge[criterion -> ?C; question -> ?Q;
       answer -> ?Text; tool#order -> 2]
    <- ?S[flow#answer(?A) -> ?Text; flow#criterion -> ?C; flow#ask -> ?Q].

?S[flow#askNow -> ?B]
    <- ?S[flow#rejected(?A) -> true; flow#limit -> ?M] AND ?B := ?A + 1 AND ?B <=?M.
```

`_turn` lets the dialogue model phrase the question, shows it to the person and returns the answer; `_judge` gives the judge model the criterion and that answer. For a Collect step the judge answers with

the labels *answered*, *not answered* or *unclear*; the host writes them as `"yes"`, `"no"` and `"unclear"` into `result("verdict")`, which the library reads, and keeps the label itself in `result("label")`. In the recorded run the first answer was "the Master"; the receipts read (without the line `tool#completed -> true` that the host also writes):

```
`turn(programme, 1)`[tool#status -> "succeeded"; result("answer") -> "the Master"].

`judge(programme, 1)`[tool#status -> "succeeded"; result("verdict") -> "no";
   result("label") -> "not answered"; result("reason") ->
   "The answer \"the Master\" does not mention Business Innovation or Management."
      ].
```

A *no* or *unclear* verdict rejects the attempt, so the library asks the programme question a second time, and the dialogue model rephrases it with the judge's reason. The person answers at the terminal or, for a run that can be repeated, from a file of answers.

Library form

Pattern 5 is written in library form from the start. FrameX's flow library `flow.fx` holds the rules, and the program above states facts only.

Known uses

Interview, suitability and triage flows; customer support with hand-over; FrameX-Agents phases, grants and context; point-in-time retrieval in the allocation application; the evidence-needs search for the questions that settle a goal. In the literature: programmable dialogue rails (Rebedea et al. 2023), routing and incremental querying (Schluntz and Zhang 2024; Liu et al. 2025); knowledge-enhanced inference and retrieval (S. Pan et al. 2024; de Boer et al. 2025). In these the flows are scripted or the model drives; here a rule base derives every step from its state.

Related patterns

Model as Sensor (Section 3.1) is the judge's role; Critique and Repair (Section 4.3) repairs malformed turns; Division by Competence (Section 4.4) sends *unclear* to a person.

Evidence

With synthetic intake protocols of 4 to 96 questions, the hybrid decided all 770 runs correctly; a model applying the same rules in the conversation fell to 39 of 60 at 64 questions and 31 of 60 at 96, with 45 of its 50 errors unsafe (McNemar $p < 0.001$). Without reasoning, the same model decided 10 of 30 conversations as a pure model and 30 of 30 in the hybrid. On a ladder of six local models, no model up to 70 billion parameters reached the criterion as a pure model, while a 9-billion model sufficed for the hybrid in single decisions. Cost favours neither side: with cached input priced, the two cost about the same at 64 and 96 questions. In a telephone-triage head-to-head on 14 simulated callers, a model alone was first more accurate (274 of 280 against 243 of 280). The hybrid's errors had one cause, found in its recorded verdicts and reasons and corrected in one place in the judge; afterwards it was right in 278 of 280 runs and in 199 of 200 on unseen callers, level with the model alone within noise, at lower cost per call and with every level traced to recorded verdicts.

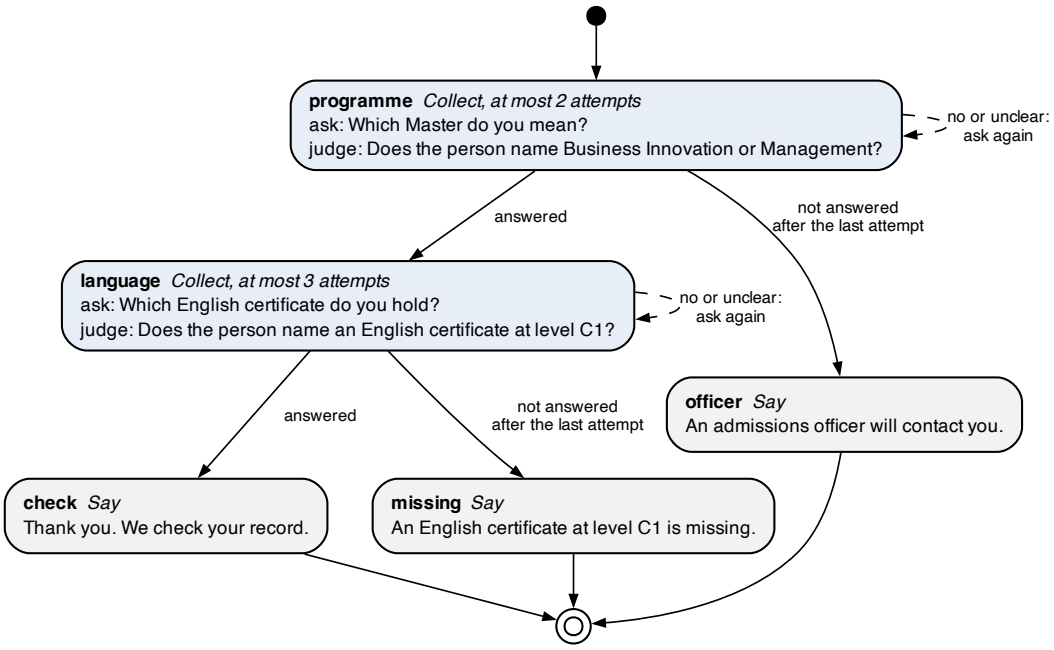

Figure 7: Flow diagram of `p05-steer.fx`: the steps of the flow and the transitions the flow library derives from them. Legend in Table 9.

4.3 Pattern 6: Critique and Repair

Intent

When a model's reply fails a check, classify the failure by rule, state it, and ask again with the problem and the faulty reply in the prompt, a bounded number of times; stop on repetition.

Also known as

Generate, verify, repair; validated retry; feedback loop with an external critic.

Motivation

At temperature zero a model gives the same reply to the same input; asking again without changing the input does not help, and a model asked to check itself without external feedback often makes things worse. The rules know what is wrong (a key misspelled, a quote not found, an action not allowed, rows missing) and can say so.

Applicability

Use it whenever another pattern checks a reply: format errors, ungrounded statements, refused proposals, incomplete extractions, failed regression cases of a draft.

Structure

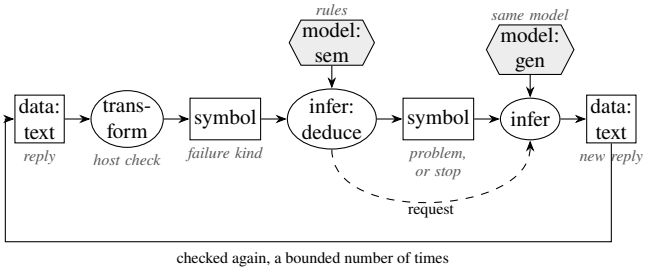

Participants

Validator in the host (measures format and occurrence); *Rules* (classify, decide, bound); *Model*; *Person* or a stronger model (escalation).

Collaborations

The rules classify a failure into one of six kinds and respond per kind:

Failure	Response
wrong format	ask again with the problem stated, up to a bound; a repeated failure stops with its reason
not grounded	reject the statement with its reason
not allowed	refuse with a reason the model receives on its next turn
incomplete	derive what is missing; the answer becomes *unknown*
failed call	stop with the reason; never treat it as an *unclear* answer
plausible but wrong	not detectable from the reply; needs grounding, a second reader or a person

Each repair is a new request with its own identity and receipt; both replies stay recorded. The first five kinds are results the rules can tell apart, as invariant 5 requires. The sixth is the class that no check can detect; we list it so that a passed check is never read as protection against it.

Consequences

Repairs succeed when the message names the fix. The bound keeps cost predictable. Silent correction is excluded: mapping "ja" to *yes* is guessing and is refused; a correction is accepted without a call only when trusted facts entail it, as a derived fact with its own proof.

Implementation

For a plain model call, the repair is part of the program. The program of Model as Author (Section 3.2) contains it: every failure keeps the checker's message, and a new request `draft(?P, ?M)` with the next number hands the model its rule and that message. It stops at `maxAttempts`, and earlier when the same message comes back twice (`stalled`):

```
// Pattern 6: ask again with the rule and the problem in the request, as a new
// request with the next number. Stop at the bound or when the message repeats.
?P[newProblem(1) -> true] <- ?P[failed(1) -> ?F].

?P[newProblem(?N) -> true]
  <- ?P[failed(?N) -> ?F] AND ?K := ?N - 1 AND ?P[failed(?K) -> ?G] AND ?F != ?G.

external draft(?P, ?M) using _llm with repairPrompt[policy -> ?T; vocabulary ->
    ?V;
       previousRule -> ?Old; problem -> ?F; tool#order -> 2]
  <- ?P:Policy[text -> ?T; vocabulary -> ?V; maxAttempts -> ?Max;
       rule(?N) -> ?Old; failed(?N) -> ?F; newProblem(?N) -> true]
    AND ?N < ?Max AND ?M := ?N + 1.

?P[stalled -> ?N]
  <- ?P[failed(?N) -> ?F] AND ?K := ?N - 1 AND ?P[failed(?K) -> ?F].

?P[gaveUp -> ?F] <- ?P:Policy[maxAttempts -> ?N; failed(?N) -> ?F].
```

In the recorded run the first draft lacked its final full stop and was rejected before evaluation; the repair request carried the message "expected '.' ", and the second draft passed all cases.

In a dialogue flow the library does the same for the dialogue model. The host validates each reply of `_turn` against the agreed format and writes what it found into the receipt; a rule of the library reads a format error and requests `_repairTurn`, which calls the same model again with the question, the problem and the faulty reply:

```
// From FrameX's flow.fx (library).
external repair(?S, ?A) using _repairTurn[question -> ?Q; attempt -> ?A;
       problem -> ?E; received -> ?X; tool#order -> 1]
   <- ?R[tool#name -> "_turn"; tool#arg1 -> ?S; tool#arg2 -> ?A;
        result("formatError") -> ?E; result("received") -> ?X]
   AND ?S[flow#ask -> ?Q].
```

The receipt that triggers it reads, for example, `result("formatError") -> "missing key: answer"`. A second broken reply ends the flow with `workflow[status -> "interrupted"]` and its reason.

Library form

The repair loop is part of two libraries: `author.fx` for drafts checked by a tool (Section 3.2) and FrameX's `flow.fx` for dialogue turns (Section 4.2). A program that loads one of them gets the bounded repair without writing a rule.

Known uses

Format repair of dialogue turns and judges; refused agent actions and premature finishes with their failed checks; completeness checks that turn a partial extraction into *unknown*; the rule-drafting loop with the checker's message. In the literature: refinement from solver error messages in Logic-LM (L. Pan et al. 2023), tool-grounded critique (Gou et al. 2024), the back-prompt controller of LLM-Modulo (Kambhampati et al. 2024); self-feedback without an external check (Madaan et al. 2023; Shinn et al. 2023), whose gains Huang et al. (2024) trace to oracle stopping signals.

Related patterns

It completes Model as Sensor, Model as Author and Model Proposes, Rules Dispose. Model as Critic (Section 5.2) is its reverse: there the model criticises the rules.

Evidence

See Model as Author: naming the fix raised first-request repairs from 5 of 10 to 10 of 10. In the literature, self-correction without an external signal lowered accuracy, and the gains reported for self-refinement rested on oracle stopping signals (Huang et al. 2024); back-prompting from a sound validator raised executable plans from about 12 % to 82 % in one planning domain (Kambhampati et al. 2024). The rules play the part of that signal and of that validator.

4.4 Pattern 7: Division by Competence

Intent

Let the rules decide the cases they cover, send what they leave open to a model or a person, and turn recurring answers into rules over time.

Also known as

Selective delegation, escalation cascade, tiered autonomy, rule distillation, human in the loop.

Motivation

Most cases of a claims process are routine and fully covered by rules; a few are ambiguous. Sending every case to a model costs money and gives up proofs where they were available; sending none leaves the ambiguous cases undecided. The pattern uses the rules' own three-valued result as the gate: *true* and *false* are decided, *unknown*, *unclear* and conflicts are handed on.

Applicability

Use it when coverage by rules is high but incomplete, when the number of model calls should follow the number of ambiguous cases, and when a person must decide the residue.

Structure

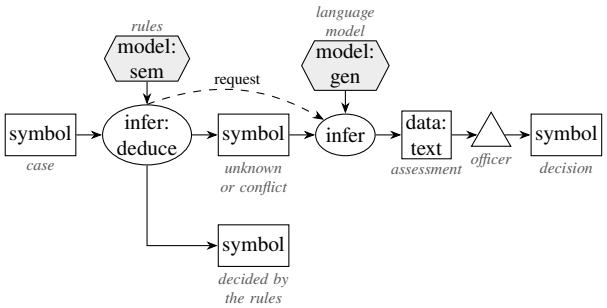

Participants

Rules (decide, mark the residue); *Model* (assesses the residue, as a recorded assessment); *Person*; *Counter* (counts recurring model decisions).

Collaborations

Rules decide; a case left *unknown* or in conflict becomes a request; the model's verdict is an assessment that travels with the case to a person and decides nothing; recurring verdicts become candidate rules for Model as Author.

Consequences

The rule part grows and the model part shrinks over time. The gate must never be the model's own confidence; use *unknown*, conflict, or a measured precision per model and task.

Implementation

Gate on the rules' own result. Under an open world a case counts as decided only when the rules derive a decision. Every other case goes to the model for an assessment and, with that assessment, to a person:

```
// Pattern 7, Division by Competence: the rules decide what they cover; the rest is
// assessed by the model and handed on to a person, with the assessment as a note.
world open.
closed slot decided. closed slot admitted. closed slot notAdmitted.

// Three files; the rules have decided two of them.
nora:Applicant[file -> "Bachelor HSG; English C1 certificate on file."].
tomas:Applicant[file -> "Bachelor ETH; English C1 recorded as failed."].
yara:Applicant[file -> "Bachelor University of Bologna; IELTS 7.0."].
nora[admitted -> mbi]. tomas[notAdmitted -> mbi].

assessPrompt:Prompt[text -> "Assess this admission file for the MBI. Answer as JSON {\"
    verdict\": likely or unlikely or unclear, \"reason\": ...}."].

?A[decided -> true] <- ?A[admitted -> ?P].
?A[decided -> true] <- ?A[notAdmitted -> ?P].

// Only an undecided file goes to the model.
external assess(?A) using _llm with assessPrompt[file -> ?T]
  <- ?A:Applicant[file -> ?T] AND NOT ?A[decided -> true].

?A[assessment -> ?V]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?A; result("verdict") -> ?V].

?A[assessmentReason -> ?W]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?A; result("reason") -> ?W].

// Every undecided file goes to an officer. The model's assessment travels with it as
// a note and decides nothing; a failed call is noted as such.
?A[handOn -> officer] <- ?A:Applicant AND NOT ?A[decided -> true].

?A[note -> ?W] <- ?A[handOn -> officer; assessmentReason -> ?W].

?A[note -> "no assessment: the model call failed"]
  <- ?A[handOn -> officer] AND ?R[tool#name -> "_llm"; tool#arg1 -> ?A;
       tool#status -> "failed"].

?- ?A[handOn -> ?W].
?- ?A[assessment -> ?V; note -> ?N].
```

Figure 8 draws the program as a rule graph. The hand-over is a derived fact, `handOn -> officer`. Showing it to a person is the host's task; FrameX's tool `_ask` puts a question to a person as written and returns the answer as a receipt.

Library form

With `competence.fx` (Section B) the program states the cases and keeps its rules for what it decides (`cp#decided`). The library assesses every other case and hands it on with the model's reason as a note:

```
// Pattern 7 in library form: load competence.fx first.
world open.

assessPrompt:Prompt[text -> "Assess this admission file for the MBI. Answer as JSON {\"
    verdict\": likely or unlikely or unclear, \"reason\": ...}."].

nora:cp#Case[cp#prompt -> assessPrompt; cp#file -> "Bachelor HSG; English C1 certificate
    on file."].
```

```
tomas:cp#Case[cp#prompt -> assessPrompt; cp#file -> "Bachelor ETH; English C1 recorded as
    failed."].
yara:cp#Case[cp#prompt -> assessPrompt; cp#file -> "Bachelor University of Bologna; IELTS
    7.0."].
nora[admitted -> mbi]. tomas[notAdmitted -> mbi].

// What the rules decide; this stays with the program.
?A[cp#decided -> true] <- ?A[admitted -> ?P].
?A[cp#decided -> true] <- ?A[notAdmitted -> ?P].

?- ?C[cp#handOn -> true; cp#assessment -> ?V; cp#note -> ?N].
```

Known uses

Three-way branches in triage that send *unclear* to a person; *awaiting user* in FrameX-Agents; required human review for evidence-gated actions in casework. In the literature: separating business logic from the model (Washizaki et al. 2020); routing (Schluntz and Zhang 2024); the inverse direction, a model that decides when to call a reasoner, is Kautz's Neuro[Symbolic] (Kautz 2022); extracting rules from a model is the "decompiler" of the neural-symbolic cycle (d'Avila Garcez and Lamb 2023).

Related patterns

Model as Author (Section 3.2) for distillation; Rules Steer the Model (Section 4.2) for the hand-over.

Evidence

One demonstration run. No measurement isolates this pattern yet. The triage flows of Section 4.2 use it to send *unclear* to a person; a measurement should report the share of cases handed on, the accuracy of those decided by rules, and the cost per decided case.

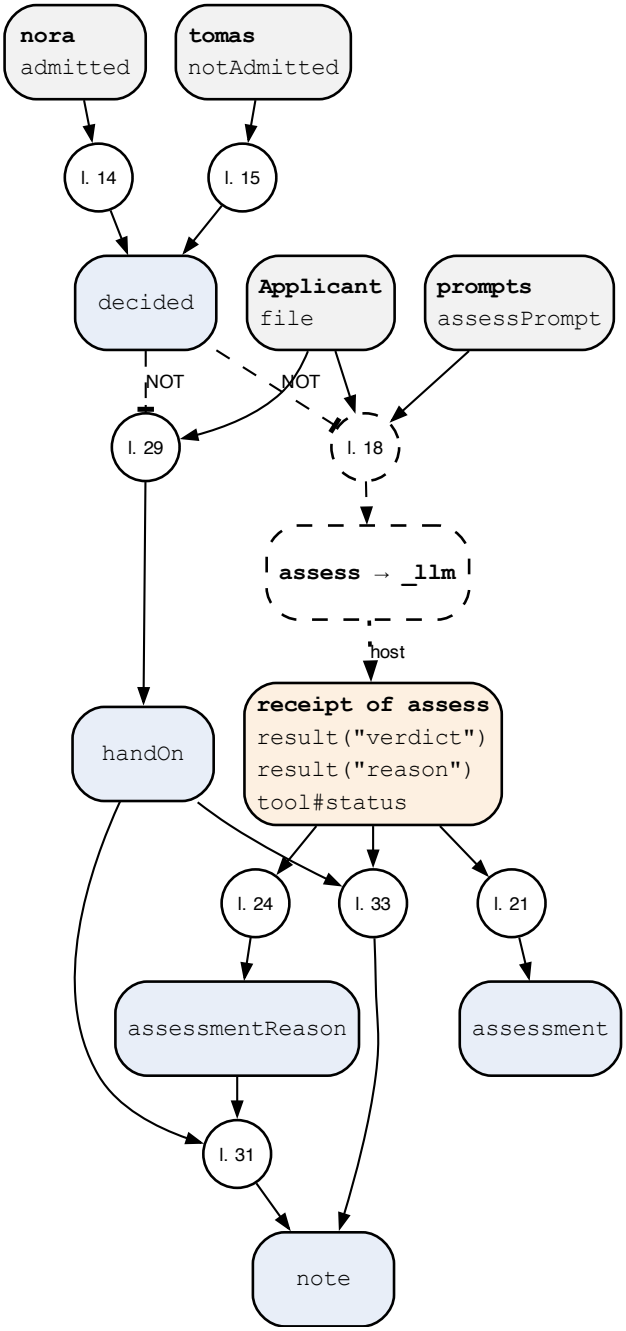

Figure 8: Rule graph of `p07-competence.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

4.5 Pattern 13: Fail-Safe Union

Intent

Combine the verdicts of rules and models under a declared order of outcomes in which each source can only make the result stricter. A source that is silent, fails or answers outside the order counts as the cautious outcome.

Also known as

Deny overrides, most restrictive wins, one-sided combination, fail closed, defence in depth.

Motivation

A payment desk releases transfers. Rules know the limits and the blocked accounts but cannot read a purpose line that pressures the clerk ("the CEO approved it by phone, do not call back"); a model can read it but misses limits and sometimes waves a fraud through. Taking a vote lets one source's leniency outvote the other's warning, and handing every disagreement to a person floods the desk. In this domain the harms are lopsided: a held payment costs a day, a released fraud costs the amount. The pattern states the order of outcomes once (release, hold, block) and lets the strictest verdict of any source decide. Neither source can loosen what the other found, and a model that did not answer is not a model that agreed.

Applicability

Use it when:

1: The outcomes can be ordered by risk, and wrongly choosing a stricter outcome costs much less than wrongly choosing a looser one;

2: The sources see different things (rules the structured fields, a model the text), so their failures are not the same;

3: A cautious outcome is not final: a hold leads to a person or a later step.

Do not use it when both directions of error cost alike, or when the stricter outcome is itself harmful (a withheld emergency call); then report the disagreement as *unclear* (Several Readers, Section 3.1) or hand it on (Section 4.4).

Structure

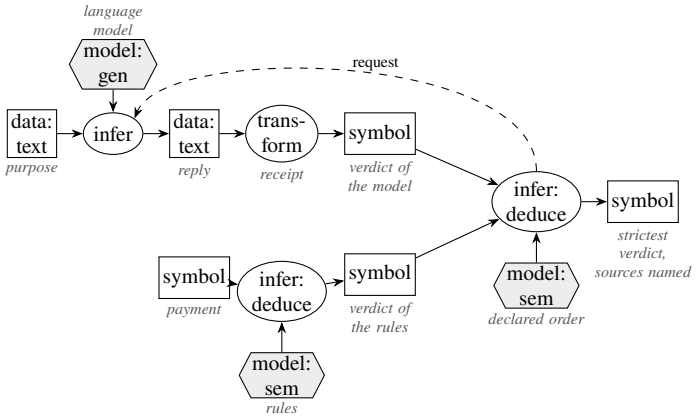

Participants

Rules: give their own verdict from the structured fields and combine all verdicts. *Model*: gives a verdict from the text, one of the ordered outcomes. *Host*: calls the model and records a failed call as such. *Order*: the outcomes with their rank, declared as facts. *Person*: takes up what is held.

Collaborations

Each source derives a verdict for each case. A source without a usable verdict (no receipt, a failed call, a word outside the order) is given the cautious verdict by a rule. The outcome is the verdict of highest rank, and the sources that gave it are named, so a person sees who held the payment and why.

Consequences

Adding a source can only make the outcome stricter, so a new check never weakens the old ones, and a model outage stops releases instead of letting them through. The price is more held cases: every source's false alarms add up, so the hold rate grows with each source and must be measured. Unlike a priority order, nothing is settled silently: the order is a fact in the program, and `decidedBy` names the source behind every outcome. The guarantee holds only if every branch moves toward the safe side, including the two that are easy to forget: a source that answers *unclear* must not let another source loosen the outcome, and a failed call must not count as agreement. One branch that does either breaks the whole union. In the hackathon below, a comparison system recorded an unreachable judge as *yes*, a submission a reply that was not JSON as *clean*.

Implementation

Declare the order as facts (`"hold":Verdict[rank -> 1]`) and the verdicts and ranks closed, so that `max` over them has a value. Write the cautious default for silence explicitly, one rule per way a source can fail to answer. The union is one aggregate; it reads the verdicts of all sources and has no rule that lowers one:

```
// Pattern 13, Fail-Safe Union: every source can only tighten the outcome;
// a source that is silent or unreadable counts as the safe side.
world open.
closed class Verdict. closed method verdict/1. closed slot rank.
closed slot blocked. closed slot screened. closed slot outcome.

// The outcomes in a declared order: a higher rank is safer.
"release":Verdict[rank -> 0]. "hold":Verdict[rank -> 1]. "block":Verdict[rank -> 2].

// The sources whose verdicts are combined, and two payments.
"rules":Source. "model":Source.
p1:Payment[amount -> 1200; payee -> "CH93 0076 2011 6238 5295 7";
   purpose -> "Invoice 2026-114, catering for the alumni evening"].
p2:Payment[amount -> 8400; payee -> "CH56 0483 5012 3456 7800 9";
   purpose -> "Urgent: transfer today, the CEO approved it by phone, do not call back"].
"CH44 3199 9123 0008 8901 2"[blocked -> true].
limit[amount -> 5000].

screenPrompt:Prompt[text -> "Read the purpose of a payment. Answer as JSON {\"verdict\":
    release or hold or block, \"reason\": ...}; hold when you are unsure."].

// Source 1, the rules: block a blocked payee, hold above the limit, release otherwise.
?P[verdict(?S) -> "block"]
  <- ?P:Payment[payee -> ?I] AND ?I[blocked -> true] AND ?S:Source AND ?S = "rules".
?P[verdict(?S) -> "hold"]
  <- ?P:Payment[amount -> ?A] AND limit[amount -> ?L] AND ?A > ?L
     AND ?S:Source AND ?S = "rules".
?P[verdict(?S) -> "release"]
  <- ?P:Payment[amount -> ?A; payee -> ?I] AND limit[amount -> ?L] AND ?A <= ?L
     AND NOT ?I[blocked -> true] AND ?S:Source AND ?S = "rules".

// Source 2, the model: one screening per payment.
external screen(?P) using _llm with screenPrompt[purpose -> ?T]
  <- ?P:Payment[purpose -> ?T].

?P[screened -> true] <- ?R[tool#name -> "_llm"; tool#arg1 -> ?P; tool#status -> "succeeded
    "].

```

```
?P[verdict(?S) -> ?V]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?P; result("verdict") -> ?V] AND ?V:Verdict
     AND ?S:Source AND ?S = "model".

// Silence and unreadable replies are not release: they count as hold.
?P[verdict(?S) -> "hold"]
  <- ?P:Payment AND NOT ?P[screened -> true] AND ?S:Source AND ?S = "model".
?P[verdict(?S) -> "hold"]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?P; result("verdict") -> ?V] AND NOT ?V:Verdict
     AND ?S:Source AND ?S = "model".

// The union: the safest verdict of any source decides, and the sources that set it are
    named.
?P[outcome -> ?V]
  <- ?P:Payment AND ?M := max{?K | ?P[verdict(?S) -> ?W] AND ?W[rank -> ?K]}
     AND ?V:Verdict[rank -> ?M].
?P[decidedBy -> ?S] <- ?P[outcome -> ?V; verdict(?S) -> ?V].

?- ?P[outcome -> ?V].
?- ?P[decidedBy -> ?S].
```

Figure 9 draws the program as a rule graph. Before any receipt the program holds both payments, because the model has not spoken (line 42). The rules give no verdict for a payment without an amount; where inputs can be incomplete, add a cautious default for the rule source as well. There is no library form yet.

Known uses

In the submissions to the Hack Apertus hackathon (October 2026), as described in the jury assessments, rules and a model could each only escalate in a retrieval system for meditation texts, Inbrev, Schild, Mandate Desk and PoesIA. In FrameX, the program above. A comparison system built for the hackathon looked like this pattern but is a Silent Rule Override (Section 8). In the literature: the deny-overrides combining algorithm of access-control policies, in which any applicable deny wins over permits (OASIS 2013); layered guardrails in which every rail can block (Rebedea et al. 2023). Here a model is one of the sources, and its silence counts as a deny.

Related patterns

Several Readers, One Rule (Section 3.1) makes a disagreement *unclear*; this pattern resolves it toward the safe side by a declared order. Model Proposes, Rules Dispose (Section 4.1) lets only the rules refuse; here the model may refuse too, but never admit. Division by Competence (Section 4.4) takes up what is held.

Variants

Strict union. A failed call counts as the cautious outcome, as in the program above. *Rule floor.* A failed call leaves only the rules' verdict, but never less than it: the model can only add findings, and an outage falls back to the rules instead of to the safe outcome. Use the floor when holding everything during an outage costs more than the rules alone miss, and say in the outcome that the model was absent.

Evidence

The program is checked against recorded receipts only; FrameX has no measurement of its own. Figures from the hackathon are taken from result files committed in the submissions' repositories at a named commit; we did not rerun them, and figures that appear only in a README, report or pitch are left out. In an invoice-payment agent (Apertus v1.5 8B, 500 synthetic invoices, 300 of which must not be paid), the model alone paid 178 of the 300, code checks alone 89, and the strict union none (95 % CI 0–1.3 %); it held 7 of 200 legitimate invoices, against 2 for the model alone. On a set run blind before later checks were added, the figures were 337, 174 and 29 of 600. In redaction of personal data (8B, 300 synthetic documents, 1,860 entities), recall was 38.7 % for rules, 57.7 % for the model and

78.5 % for the union, and 81 % of documents still leaked at least one entity: the union is stricter than each source, not safe by itself.

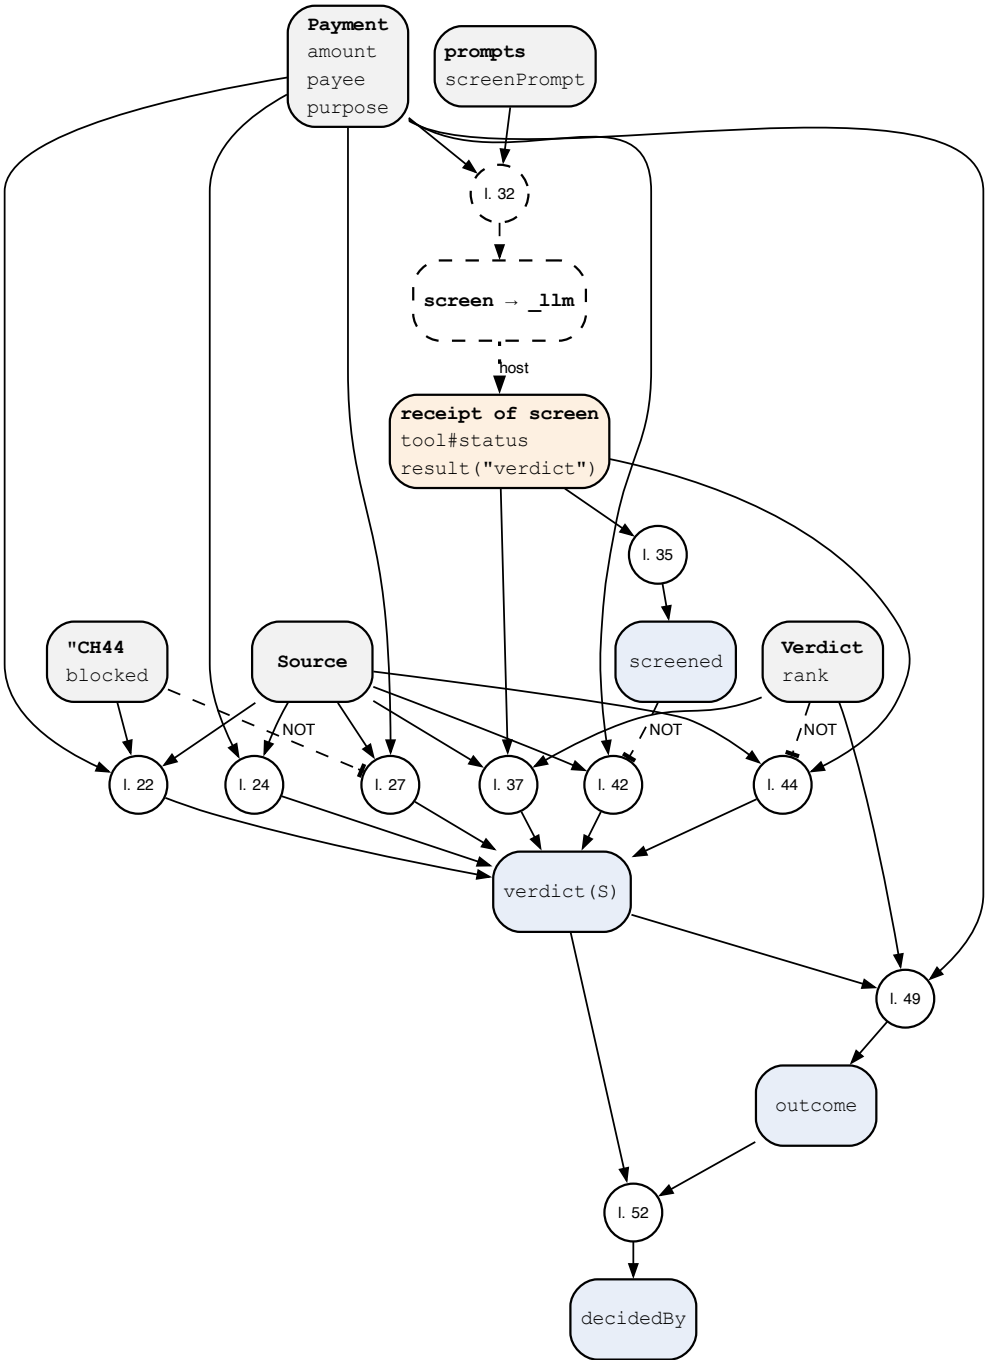

Figure 9: Rule graph of `p13-failsafe.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

5 Patterns in the reasoning itself

5.1 Pattern 8: Model as Hypothesis Generator

Intent

Let the model propose what might be the case (a cause, a missing document, an explanation) and let the rules test each proposal in a copy of the session and ask for the evidence that decides between them. A hypothesis never becomes a fact.

Also known as

Abduction with a model, diagnosis, candidate generation and test.

Motivation

A bank reconciliation shows a difference between the ledger and the bank statement. A model suggests causes (a fee, a transposed amount, a payment booked on the wrong day); each is plausible. The rules can test each: assumed in a copy, does it explain the difference without violating a constraint? The surviving candidates become questions, and only the answer decides.

Applicability

Use it for diagnosis, reconciliation, investigation and *unknown* goals, where generating candidates needs world knowledge but testing them needs exactness.

Structure

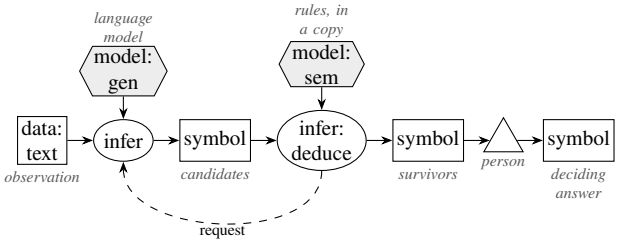

Participants

Model (proposes candidates marked hypothetical); *Rules* (test each in a copy, derive evidence questions); *Person* or a source (answers the deciding question).

Collaborations

Candidates arrive as receipts. In the full pattern each is previewed in a copy of the session; rules keep those that derive every observation and violate nothing; the evidence search names the one to four answers that would decide.

Consequences

Hypotheses are cheap to generate and never pollute the knowledge base. The search is bounded; an incomplete search is never a claim that no explanation exists.

Implementation

Ask the model for candidates as rows, keep each as a hypothesis keyed by its row number, refute it by rule, and ask about the survivors only; a refuted candidate never becomes a fact. The question itself goes to a person through a flow (Section 4.2). From the bank reconciliation:

```
// Pattern 8, Model as Hypothesis Generator: the model proposes causes, the rules test
// each against what is known and ask about the survivors.
world open.
closed method refuted/1.

// The observation to explain: ledger and bank differ by 40 francs.
recon[difference -> 40;
    ledger -> "Ledger balance 12 340.00 CHF on 30 September; bank statement 12 300.00 CHF.
         Account fee 40.00 CHF charged monthly. Payment of 90.00 CHF to Muster AG booked
        on 30 September."].

proposePrompt:Prompt[text -> "List possible causes of the difference between ledger and
    bank statement. Answer as JSON {\"rows\": [{\"entry\": ..., \"amount\": a number}]}."
    ].

// A request without arguments: the request object is propose itself.
external propose using _llm with proposePrompt[difference -> ?D; ledger -> ?L]
  <- recon[difference -> ?D; ledger -> ?L].

// Each row is a candidate cause, never a fact.
recon[cause(?I) -> ?E; amount(?I) -> ?A]
  <- propose[tool#status -> "succeeded"; row(?I, "entry") -> ?E;
       row(?I, "amount") -> ?A].

// Test: a cause whose amount differs from the difference cannot explain it.
recon[refuted(?I) -> true]
  <- recon[amount(?I) -> ?A; difference -> ?D] AND ?A != ?D.

recon[survives(?I) -> true]
  <- recon[cause(?I) -> ?E] AND NOT recon[refuted(?I) -> true].

// Ask: every surviving cause becomes a question for a person.
recon[askAbout(?I) -> ?E] <- recon[survives(?I) -> true; cause(?I) -> ?E].

?- recon[askAbout(?I) -> ?E].
?- recon[refuted(?I) -> true].
```

Figure 10 draws the program as a rule graph. This program tests every candidate with one rule, its amount. A fuller test assumes each candidate in a copy of the session (`preview`) and keeps it only if it derives the observation without violating a constraint, and the evidence search (`_evidence_needs`) names the answers that would decide. The recorded run shows why the test should be as strong as the domain allows: four of five candidates passed the amount test, among them an "unrecorded deposit" whose direction contradicts the ledger.

Library form

With `hypothesis.fx` (Section B) the program states the observation, names the row field that a person is asked about (`hy#label`) and keeps its tests (`hy#refuted`). The library keeps every field of every row as `hy#candidate(?I, field)`. It sends the observation as one text, so the model sees a different request than in the program above:

```
// Pattern 8 in library form: load hypothesis.fx first.
world open.

proposePrompt:Prompt[text -> "List possible causes of the difference between ledger and
    bank statement. Answer as JSON {\"rows\": [{\"entry\": ..., \"amount\": a number}]}."
    ].

recon:hy#Observation[hy#prompt -> proposePrompt; hy#label -> "entry"; difference -> 40;
    hy#text -> "Ledger balance 12 340.00 CHF on 30 September; bank statement 12 300.00 CHF
        . Account fee 40.00 CHF charged monthly. Payment of 90.00 CHF to Muster AG booked
        on 30 September."].
```

```
8
9  // The test stays with the program: a cause whose amount differs cannot explain the
       difference.
10 recon[hy#refuted(?I) -> true]
11   <- recon[hy#candidate(?I, "amount") -> ?A; difference -> ?D] AND ?A != ?D.
12
13 ?- recon[hy#askAbout(?I) -> ?E].
```

Known uses

Bank reconciliation with model-proposed causes; the evidence-needs search over declared open questions; the candidate checker's `facts` mode. In the literature: few-shot abduction tested against retrieved events (Shi et al. 2023); hypotheses implemented as programs and verified on examples (Wang et al. 2024). Neither asks for the evidence that decides between surviving candidates.

Related patterns

Critique and Repair (Section 4.3): a repair suggestion is a hypothesis that is a correction. Rules Steer the Model (Section 4.2) asks the deciding question.

Evidence

One demonstration run; no measurement yet. The bank reconciliation of Table 5 runs the pattern.

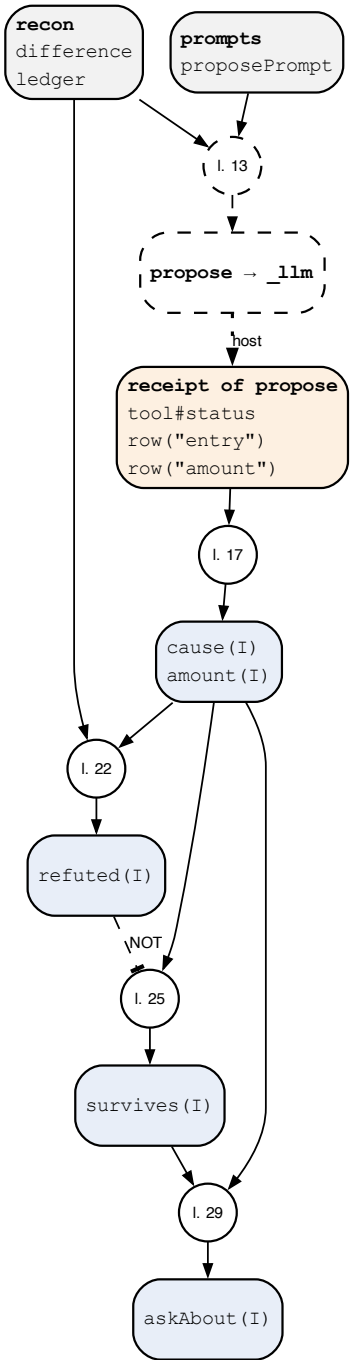

Figure 10: Rule graph of `p08-hypothesis.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

5.2 Pattern 9: Model as Critic of the Rules

Intent

Let the model review what the rules hold or conclude (a policy change, a rule against its source text, a surprising decision); its critique opens a review and never changes a rule or a conclusion.

Also known as

Policy critic, plausibility review, two-channel cross-check.

Motivation

A policy change passes every regression case and still changes the decision for a case nobody expected. The change's author wrote a rationale; a model can read the changed decisions against it and flag what looks inconsistent. Its remarks may be wrong, so they open a review that a person disposes of.

Applicability

Use it where checks are incomplete: semantic faults, fidelity of a rule to the text it cites, plausibility of changed decisions before a release.

Structure

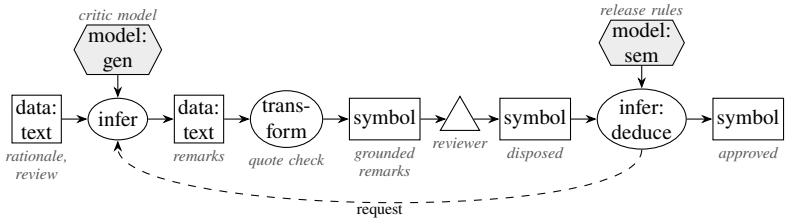

Participants

Model (critic); *Rules* (keep a remark only if it names a case and quotes the rationale verbatim; flag declared invariants without a model; gate approval); *Reviewer* (disposes of every finding).

Collaborations

A sealed before-and-after review of the policy is the critic's input; grounded remarks become findings; approval is derived only after every finding is disposed of. A corrected rule proposed by the critic is checked like any draft.

Consequences

The critic adds recall where checks are incomplete; its false remarks cost reviewer time, so grounding requirements filter them. It never weakens a complete check.

Implementation

The model receives the rationale and the sealed review and answers with rows, as in the Sensor, with the rationale as `groundIn`, the text its quotes must come from. A remark counts only if its quote was found and it names a changed decision. Approval needs a completed critique and a disposal of every remark; a failed call blocks it with the host's reason:

```
// Pattern 9, Model as Critic of the Rules: grounded remarks open a review; approval
// needs a completed critique and a reviewer's disposal of every remark.
world open.
closed class ChangedDecision. closed method disposed/1. closed slot hasOpen.
closed slot critiqued.

// A rule change with its rationale and the decisions it changes.
"c7":ChangedDecision. "c9":ChangedDecision.
chg:Change[reviewed -> true;
```

```
    rationale -> "The amendment lowers the language requirement for applicants with a
        Bachelor taught in English. Applicants with other degrees keep the C1 requirement.
        ";
    review -> "c7: Bachelor taught in German, no C1; before not admitted, after admitted.
        c9: Bachelor taught in English, no C1; before not admitted, after admitted."].

critiquePrompt:Prompt[text -> "Read the changed decisions against the rationale. Report
    one row {case, quote, remark} per decision that contradicts the rationale; quote the
    rationale verbatim. Answer as JSON {\"rows\": [...]}."].

// The quote of every remark must occur in the rationale.
external critique(?C) using _llm with critiquePrompt[rationale -> ?Ra;
       review -> ?Rv; groundIn -> ?Ra]
  <- ?C:Change[rationale -> ?Ra; review -> ?Rv].

// A remark counts only with a found quote and a case that the change really changed.
?C[remark(?I) -> ?K]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?C; row(?I, "case") -> ?K;
       row(?I, "quoteFound") -> true] AND ?K:ChangedDecision.

?C[critiqued -> true]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?C; tool#status -> "succeeded"].

?C[blocked -> ?E]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?C; tool#status -> "failed"; error -> ?E].

// A reviewer disposes of a remark by stating chg[disposed(1) -> true].
?C[hasOpen -> true] <- ?C[remark(?I) -> ?K] AND NOT ?C[disposed(?I) -> true].

?C[approved -> true]
  <- ?C:Change[reviewed -> true; critiqued -> true] AND NOT ?C[hasOpen -> true].

?- chg[remark(?I) -> ?K].
?- chg[approved -> ?A].
```

Figure 11 draws the program as a rule graph.

Library form

With `critic.fx` (Section B) the program states the change, its rationale, the review and the cases it changed; a reviewer's disposal is the fact `chg[cr#disposed(1) -> true]`:

```
// Pattern 9 in library form: load critic.fx first.
world open.

critiquePrompt:Prompt[text -> "Read the changed decisions against the rationale. Report
    one row {case, quote, remark} per decision that contradicts the rationale; quote the
    rationale verbatim. Answer as JSON {\"rows\": [...]}."].

chg:cr#Change[cr#reviewed -> true; cr#prompt -> critiquePrompt;
    cr#changed -> "c7"; cr#changed -> "c9";
    cr#rationale -> "The amendment lowers the language requirement for applicants with a
        Bachelor taught in English. Applicants with other degrees keep the C1 requirement.
        ";
    cr#review -> "c7: Bachelor taught in German, no C1; before not admitted, after
        admitted. c9: Bachelor taught in English, no C1; before not admitted, after
        admitted."].

?- chg[cr#remark(?I) -> ?K].
?- chg[cr#approved -> ?A].
```

Known uses

The policy critic in a university administration case with seeded defects. In the literature: soft critics in LLM-Modulo (Kambhampati et al. 2024) and cross-reflection (Liu et al. 2025) judge a candidate; we found no catalogue in which the model reviews the rules themselves.

Related patterns

Reverse of Critique and Repair (Section 4.3); input to Model as Author (Section 3.2) when it proposes a correction.

Evidence

In a pilot study with `gpt-oss:20b`, 20 changes over 10 rule domains were each reviewed three times. In the 10 changes with one seeded decision that contradicts the rationale, the critic found it in 29 of 30 reviews; in the 10 clean changes it made no remark in 29 of 30. All 34 of its remarks quoted the rationale verbatim, so the quote check removed nothing. The changes were short and each defect was stated plainly; harder defects and longer reviews are not yet measured. In the checking study, model reviews found contradictory rules and unknown facts almost always, but a missing rule in only 3 of 10 rule reviews at 64 rules and in none of the code or prompt reviews, so a critic belongs on top of the checks.

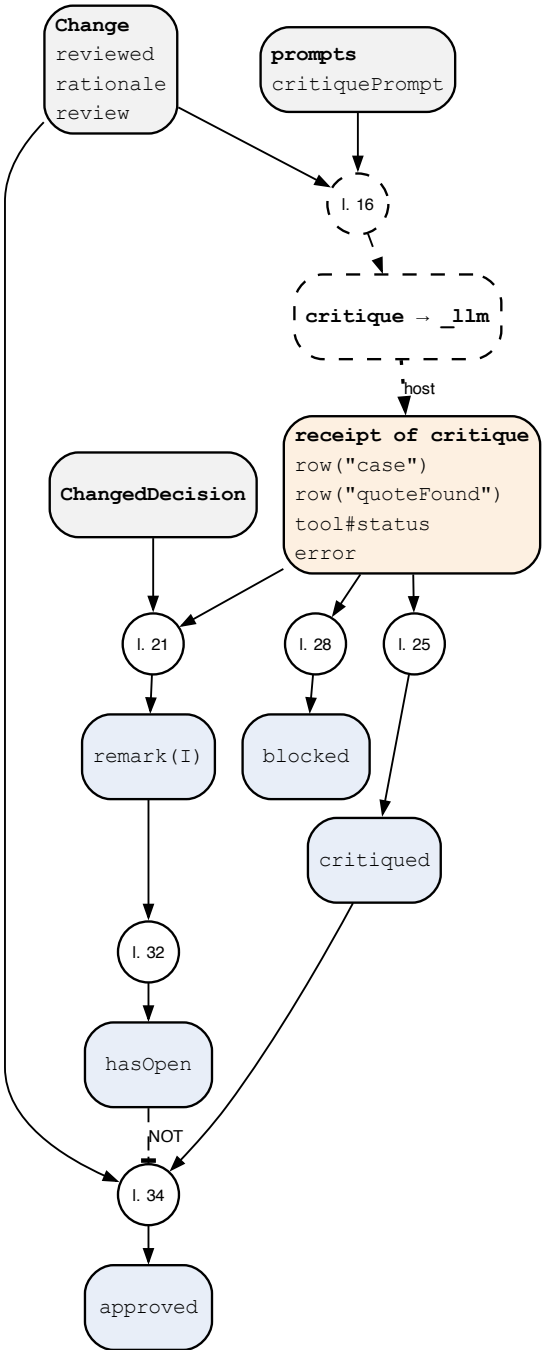

Figure 11: Rule graph of `p09-critic.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

5.3 Pattern 10: Rules Teach the Model

Intent

Let the rules supply labels, rule-selected worked examples and grades, with proofs, so that knowledge flows into the model's context or weights; the rules stay in production as the independent check.

Also known as

Rules as labels and graders, rule-selected examples, rule-informed preference data.

Motivation

A model that must answer admission questions can be shown worked cases. Which cases, and whether its answers are right, the rules know: they label generated cases with decisions and proofs, select the examples closest to the current case, and grade the answers.

Applicability

Use it when a model must work where the rules are not called (a fast path, an edge device) or must learn the vocabulary of a domain, and when training or prompt examples need trustworthy labels.

Structure

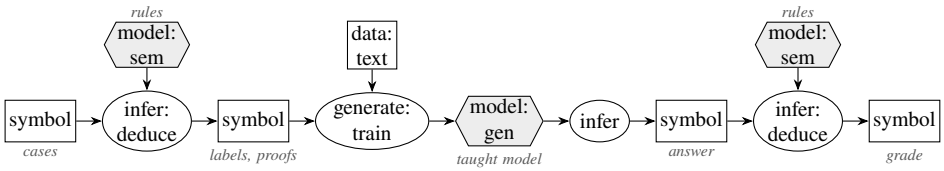

Participants

Rules (label, select, grade); *Model*; *Trainer* outside the rule engine.

Collaborations

The rules label every generated case with its decision and the missing requirements; the model answers each case; a rule grades the answer against the label and records each gap the answer missed. Labels and grades, with their proofs, leave the engine as training or evaluation data, and after training the same rules grade the taught model.

Consequences

The labels are as good as the rules and come with proofs; what the model learned is never the authority.

Implementation

Grade a model's answer against the rules' decision on generated cases with complete evidence, and record each gap the answer missed. The model answers each case; its receipt becomes the answer that the rules grade:

```
// Pattern 10, Rules Teach the Model: the rules label the cases and grade the model's
// answers. Labels and grades can become training data; training happens outside.
world open.
closed class Admitted. closed slot named.
closed slot hasBachelor. closed slot hasC1.

// Two cases: the rules read the structured fields, the model reads the description.
k1:Case[hasBachelor -> true; hasC1 -> true;
    description -> "Recognised Bachelor; English C1 certificate on file."].
k2:Case[hasBachelor -> true;
    description -> "Recognised Bachelor; no language record."].

// The rules label each case: its decision and every missing requirement.
?C:Admitted <- ?C:Case[hasBachelor -> true; hasC1 -> true].
```

```
?C[gap -> "englishC1"] <- ?C:Case AND NOT ?C[hasC1 -> true].
?C[gap -> "bachelor"] <- ?C:Case AND NOT ?C[hasBachelor -> true].

answerPrompt:Prompt[text -> "Requirements: a recognised Bachelor and English C1. Is the
    applicant admitted? Answer as JSON {\"decision\": admitted or not admitted, \"rows\":
     [{\"gap\": englishC1 or bachelor}] per missing requirement}."].

// The request field is called description, because the rules below derive a slot case.
external answer(?C) using _llm with answerPrompt[description -> ?T]
  <- ?C:Case[description -> ?T].

?R:Answer[case -> ?C; decision -> ?D]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?C; result("decision") -> ?D].

?R[named -> ?G] <- ?R:Answer AND ?R[row(?I, "gap") -> ?G].

// Grade: right decision, and every gap named.
?A[right -> true] <- ?A:Answer[case -> ?C; decision -> "admitted"] AND ?C:Admitted.

?A[right -> true]
  <- ?A:Answer[case -> ?C; decision -> "not admitted"] AND NOT ?C:Admitted.

?A[missedGap -> ?G]
  <- ?A:Answer[case -> ?C] AND ?C[gap -> ?G] AND NOT ?A[named -> ?G].

?- ?A:Answer[case -> ?C; decision -> ?D; right -> ?R].
?- ?A[missedGap -> ?G].
```

Figure 12 draws the program as a rule graph. The training step of the Structure diagram happens outside the engine and is not part of the program.

Library form

With `teach.fx` (Section B) the program states the cases and keeps its labelling rules (`te#label`, `te#gap`); the library asks the model and grades each answer against the label:

```
// Pattern 10 in library form: load teach.fx first.
world open.
closed slot hasBachelor. closed slot hasC1.

answerPrompt:Prompt[text -> "Requirements: a recognised Bachelor and English C1. Is the
    applicant admitted? Answer as JSON {\"decision\": admitted or not admitted, \"rows\":
     [{\"gap\": englishC1 or bachelor}] per missing requirement}."].

k1:te#Case[hasBachelor -> true; hasC1 -> true; te#prompt -> answerPrompt;
   te#description -> "Recognised Bachelor; English C1 certificate on file."].
k2:te#Case[hasBachelor -> true; te#prompt -> answerPrompt;
   te#description -> "Recognised Bachelor; no language record."].

// The labels come from the rules; this stays with the program.
?C[te#label -> "admitted"] <- ?C:te#Case[hasBachelor -> true; hasC1 -> true].
?C[te#label -> "not admitted"] <- ?C:te#Case AND NOT ?C[hasC1 -> true].
?C[te#label -> "not admitted"] <- ?C:te#Case AND NOT ?C[hasBachelor -> true].
?C[te#gap -> "englishC1"] <- ?C:te#Case AND NOT ?C[hasC1 -> true].
?C[te#gap -> "bachelor"] <- ?C:te#Case AND NOT ?C[hasBachelor -> true].

?- ?A:te#Answer[te#case -> ?C; te#decision -> ?D; te#right -> ?R].
?- ?A[te#missedGap -> ?G].
```

Known uses

Labels, rule-selected examples and grading for 144 generated admission cases (built, measures fixed, no model run yet); rules that govern a small network's training loop. In the literature: symbolic rules

generate training data, Kautz's Neuro:Symbolic→Neuro (Kautz 2022); informed learning with prior knowledge and meta-reasoning for control (van Bekkum et al. 2021); the "compiler" of the neural-symbolic cycle (d'Avila Garcez and Lamb 2023); natural-language principles as labels in Constitutional AI and RLAIF (Bai et al. 2022; Lee et al. 2024).

Related patterns

Rules Steer the Model (Section 4.2) for selecting context; Model as Sensor (Section 3.1) for the judge that may be taught.

Evidence

One demonstration run; no measurement yet.

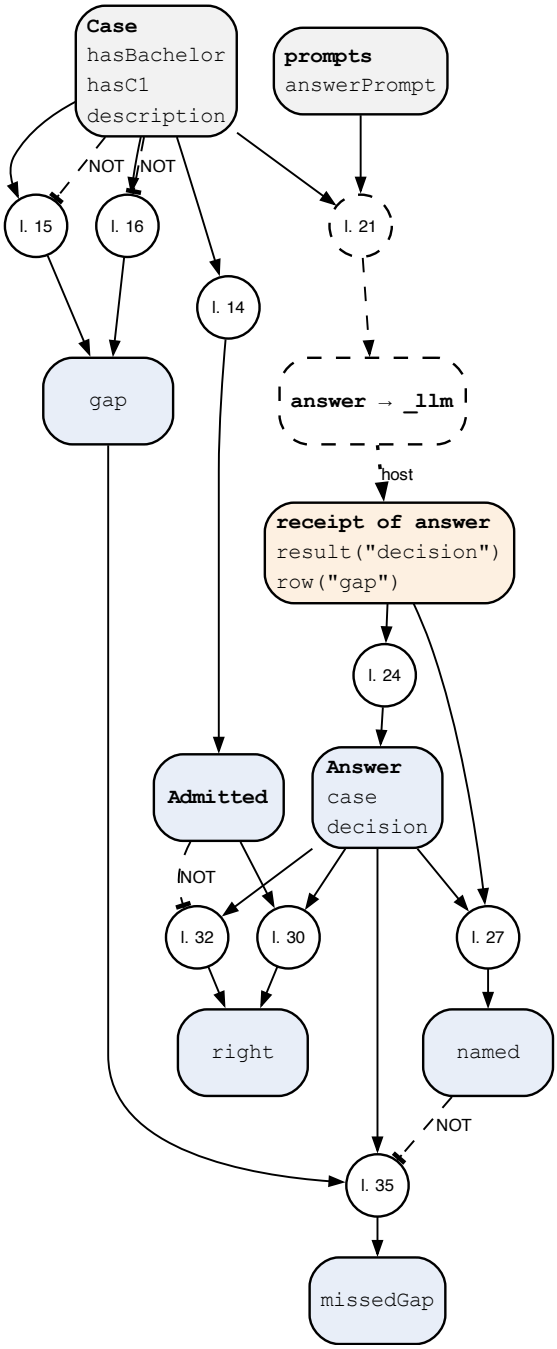

Figure 12: Rule graph of `p10-teach.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

6 Patterns over time

6.1 Pattern 11: Rule-Governed Memory

Intent

Keep what an assistant learns across conversations as statements that rules admit, supersede and withdraw, and let rules decide which remembered statements enter the next prompt.

Also known as

Agent memory, long-term memory, memory handler, belief revision with a model as source.

Motivation

An advisor learns in March that a student prefers German and in May that English is fine; in June the student asks it to forget a message. A model that keeps its own notes either keeps both values or overwrites one, cannot say which holds or why, and its notes quietly become records: a grade mentioned in a chat sits next to the grade on the transcript. A single reading (Section 3.1) decides whether a statement holds now; memory must decide whether it still holds later. The pattern keeps every remembered statement with its utterance and quote, lets a slot policy decide which sources may establish which facts, lets a newer admitted statement supersede an older one without deleting it, and recomputes what rests on an utterance when that utterance is forgotten. Rules choose what the model sees next, with a similarity score as one input, and the model cites the statements it used.

Applicability

Use it when

- an assistant serves the same person or case across sessions;
- what is known changes over time, and the history of a value matters;
- sources differ in what they may establish (a chat may set a preference, only a certificate a qualification);
- a request to forget must take effect on everything that rested on the forgotten input.

For a single session a Sensor suffices. Retrieval by similarity over documents is the task of a retriever outside the engine; its results enter as receipts like any other.

Structure

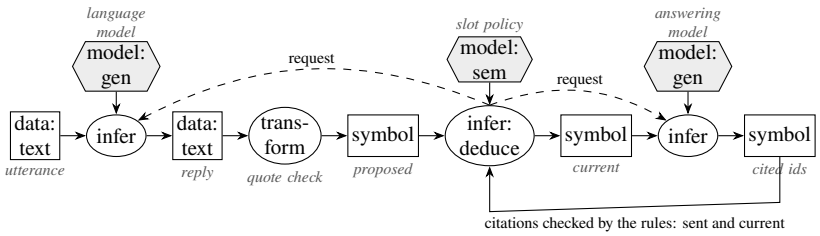

Participants

Model: proposes statements with quotes from an utterance, and later cites the statements it used. *Rules*: slot policy, admission, supersession, withdrawal, and the selection of statements for a prompt. *Host*: stores utterances and receipts, measures quotes, removes the content of a forgotten utterance. *Person*: states, corrects, asks to forget.

Collaborations

Each utterance triggers a Sensor request; the proposed statements return as receipts; rules admit them per slot; supersession follows the order of admission, and the current value is a query. A request to

forget marks the utterance; in the next derivation every statement that rested on it loses its acceptance, and an older statement can become current again. Before a model call, a steering rule (Section 4.2) selects the current statements the step needs; the reply cites their identifiers, and a rule checks that every cited statement is current.

Consequences

Every remembered value has an utterance, a quote, a date and a history. Forgetting is exact without manual clean-up, because acceptance is derived and recomputed. A chat cannot establish what its source may not, and a derived fact such as an admission is never written to memory. The cost is one Sensor call per utterance and a store that grows, since only a request to forget removes content. A statement admitted wrongly holds until it is superseded or forgotten, and a conflict between sources needs a person (Section 4.4).

Implementation

Each utterance triggers a Sensor request with the utterance as `groundIn`. Row `?I` of the answer for utterance `?U` becomes one statement object `st(?U, ?I)`, a function term in the rule head, kept only with a found quote. For recall, the rules mark the statements one step may see with `selectedFor`; the request names the step with `contextKey`, and the host sends exactly those statements and records their identifiers in the receipt (`tool#sent`). The reply cites identifiers, and rules accept it only if every cited statement was sent and is current. `tool#order` makes every `remember` run before the answer, so the selection is made over all statements:

```
// Pattern 11, Rule-Governed Memory: statements read from utterances are admitted,
// superseded and forgotten by rules; for one step the rules select what a reply may
// use, and the reply must cite only those statements, and only current ones.
world open.
closed slot mayEstablish. closed slot forgotten. closed slot hasNewer.
closed slot current. closed slot hasBad.
closed method sentId/1.

// What a chat may establish, two utterances, and one question to answer later.
chat[mayEstablish -> "preferredLanguage"; mayEstablish -> "callbackTime"].
u1:Utterance[kind -> chat; seq -> 1; text -> "Please write to me in German."].
u2:Utterance[kind -> chat; seq -> 2;
   text -> "English is fine from now on. By the way, my Bachelor grade was 5.8."].
q1:Question[step -> replyLanguage;
   text -> "In which language should we write to nora?"].

rememberPrompt:Prompt[text -> "Report what the utterance states about the speaker nora:
    one row {about, slot, value, quote}, slot one of preferredLanguage, callbackTime,
    bachelorGrade; the quote copied verbatim. Answer as JSON {\"rows\": [...]}."].
recallPrompt:Prompt[text -> "Answer the question from the statements in context only. Cite
     every statement you use by its id. Answer as JSON {\"text\": \"the answer\", \"rows
    \": [{\"cited\": \"id\"}]}."].

// Storing (Pattern 1): one statement object st(utterance, row) per quoted row.
external remember(?U) using _llm with rememberPrompt[utterance -> ?T; groundIn -> ?T;
   tool#order -> 2]
  <- ?U:Utterance[text -> ?T] AND NOT ?U[forgotten -> true].

st(?U, ?I):Statement[from -> ?U; about -> ?A; slot -> ?S; value -> ?V]
  <- ?R[tool#name -> "_llm"; tool#arg1 -> ?U; row(?I, "about") -> ?A;
       row(?I, "slot") -> ?S; row(?I, "value") -> ?V; row(?I, "quoteFound") -> true].

// Admission: the source must be allowed to establish the slot.
?St[accepted -> true]
  <- ?St:Statement[from -> ?U; slot -> ?S] AND ?U[kind -> ?K] AND ?K[mayEstablish -> ?S]
     AND NOT ?U[forgotten -> true].

?St[rejected -> "this kind of source cannot establish it"]
```

```
  <- ?St:Statement[from -> ?U; slot -> ?S] AND ?U[kind -> ?K]
     AND NOT ?K[mayEstablish -> ?S].

// Supersession: a later accepted statement on the same slot replaces an earlier one.
?St[hasNewer -> true]
  <- ?St[accepted -> true; from -> ?U; about -> ?A; slot -> ?S] AND ?U[seq -> ?N]
     AND ?Nw[accepted -> true; from -> ?W; about -> ?A; slot -> ?S] AND ?W[seq -> ?M]
     AND ?M > ?N.
?St[current -> true] <- ?St[accepted -> true] AND NOT ?St[hasNewer -> true].

// Forgetting: a person's request is the fact u2[forgotten -> true]. Its statements are
// withdrawn, and every recorded model answer that holds the utterance is removed.
external forget(?U) using _forgetAnswers[text -> ?T; tool#order -> 1]
  <- ?U:Utterance[forgotten -> true; text -> ?T].

// Recall (Pattern 5): the rules select what the step may see.
?St[selectedFor -> replyLanguage]
  <- ?St:Statement[current -> true; about -> "nora"; slot -> "preferredLanguage"].
external answer(?Q) using _llm with recallPrompt[question -> ?T; contextKey -> ?K;
    tool#order -> 3]
  <- ?Q:Question[text -> ?T; step -> ?K].

// Citations: every cited id must have been sent, and its statement must be current.
?R[sentId(?C) -> true] <- ?R[tool#sent(?C) -> ?St].
?Q[cites -> ?St]
  <- ?Q:Question AND ?R[tool#arg1 -> ?Q; row(?I, "cited") -> ?C; tool#sent(?C) -> ?St].
?Q[refused -> ?C]
  <- ?Q:Question AND ?R[tool#arg1 -> ?Q; row(?I, "cited") -> ?C]
     AND NOT ?R[sentId(?C) -> true].
?Q[refused -> ?St] <- ?Q[cites -> ?St] AND NOT ?St[current -> true].
?Q[hasBad -> true] <- ?Q[refused -> ?X].
?Q[answer -> ?Text]
  <- ?Q:Question[cites -> ?St] AND ?R[tool#arg1 -> ?Q; result("text") -> ?Text]
     AND NOT ?Q[hasBad -> true].

?- ?St[current -> true; value -> ?V].
?- ?St[rejected -> ?Why].
?- q1[answer -> ?A].
?- q1[refused -> ?X].
```

Figure 13 draws the program as a rule graph. Forgetting `u2` withdraws its statements, `st(u1, 1)` with *German* becomes current again, and `_forgetAnswers` removes every recorded model answer whose prompt or answer holds the utterance, so a replay cannot bring it back. A similarity score from a retriever can be one more condition of the selection rule.

Library form

With the library `memory.fx` (Section B) loaded first, the program states only what each kind of source may establish, the utterances, what a step needs, and the question. Selection is a fact (`mem#Step`), and a request to forget is the fact `u2[mem#forgotten -> true]`:

```
// Pattern 11 in library form: load memory.fx first. Same utterances and prompts as p11-
    memory.fx.
world open.

chat[mem#mayEstablish -> "preferredLanguage"; mem#mayEstablish -> "callbackTime"].

rememberPrompt:Prompt[text -> "Report what the utterance states about the speaker nora:
    one row {about, slot, value, quote}, slot one of preferredLanguage, callbackTime,
    bachelorGrade; the quote copied verbatim. Answer as JSON {\"rows\": [...]}."].
recallPrompt:Prompt[text -> "Answer the question from the statements in context only. Cite
     every statement you use by its id. Answer as JSON {\"text\": \"the answer\", \"rows
    \": [{\"cited\": \"id\"}]}."].
```

```
8
9  u1:mem#Utterance[mem#kind -> chat; mem#seq -> 1; mem#prompt -> rememberPrompt;
10     mem#text -> "Please write to me in German."].
11 u2:mem#Utterance[mem#kind -> chat; mem#seq -> 2; mem#prompt -> rememberPrompt;
12     mem#text -> "English is fine from now on. By the way, my Bachelor grade was 5.8."].
13
14 replyLanguage:mem#Step[mem#about -> "nora"; mem#needs -> "preferredLanguage"].
15 q1:mem#Question[mem#step -> replyLanguage; mem#prompt -> recallPrompt;
16     mem#text -> "In which language should we write to nora?"].
17
18 ?- ?St[mem#current -> true; mem#value -> ?V].
19 ?- ?St[mem#rejected -> ?Why].
20 ?- q1[mem#answer -> ?A].
21 ?- q1[mem#refused -> ?X].
```

The library derives the statements, their admission and supersession, the context of each step, the answer and every refused citation, and it removes forgotten utterances from the answer store.

Known uses

The program above, which has both halves; DreamRAG lite among the FrameX applications, the memory handler: proposal from an utterance, admission per slot, supersession, conflict and withdrawal. Memory claims with quotes in FrameX-Agents. In the literature: memory as a component of agent architectures (Liu et al. 2025); reflections kept across trials (Shinn et al. 2023); knowledge graphs as updatable memory for language models (S. Pan et al. 2024). In these the model decides what to keep and recall; here the rules do, and forgetting withdraws what rested on the forgotten input.

Related patterns

Storing is Model as Sensor (Section 3.1) applied to utterances over time; recall is Rules Steer the Model (Section 4.2); conflicts go to Division by Competence (Section 4.4); a reply that cites a superseded statement is repaired by Critique and Repair (Section 4.3).

Evidence

One demonstration run; no measurement yet.

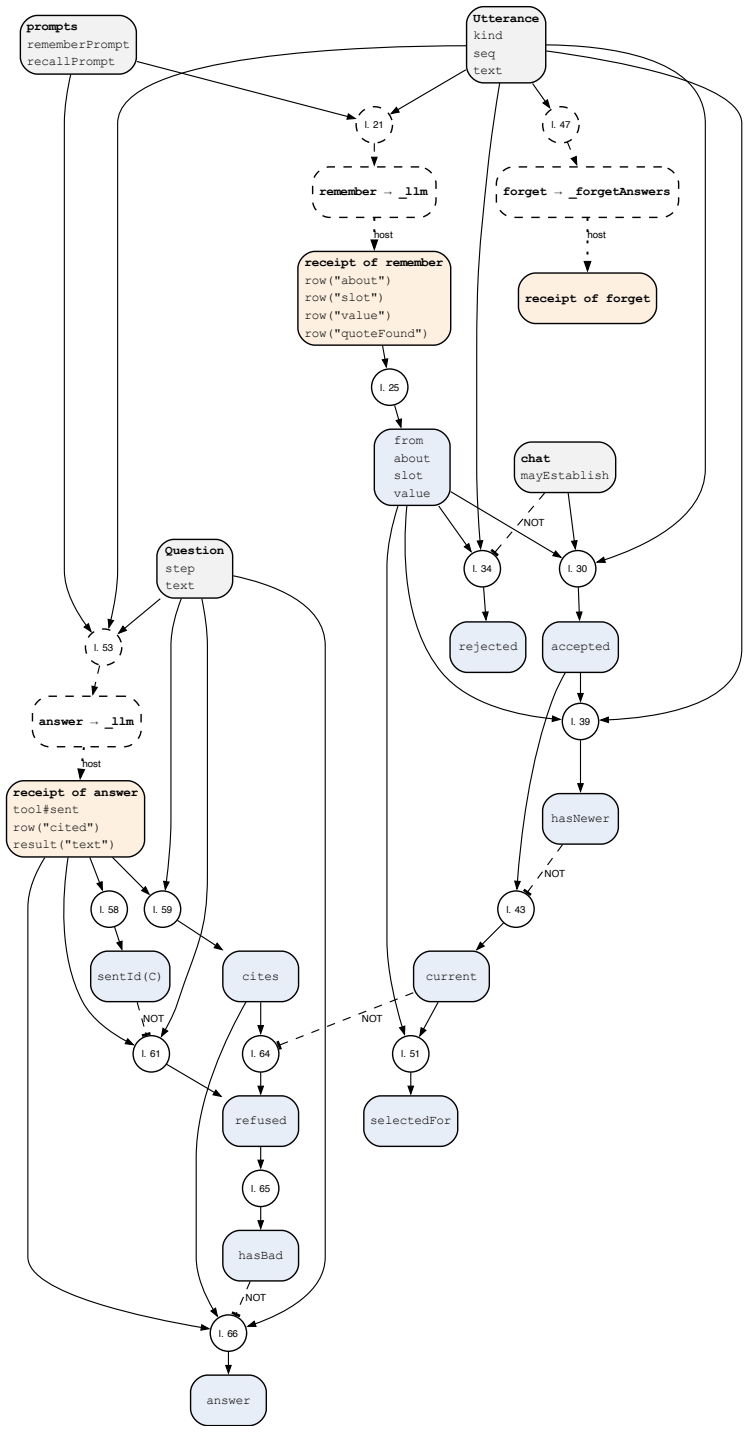

Figure 13: Rule graph of `p11-memory.fx`, generated from the program. A circle is one rule; "l. 23" means that the rule starts in line 23 of the listing. Legend in Table 9.

7 Compositions

Real systems combine patterns. Table 5 lists the compositions in the applications from which the catalogue was drawn. Three combinations recur: a flow that steers a model and uses a judge as sensor, with repair and hand-over (5 + 1 + 6 + 7); an agent whose proposals are admitted by rules that also steer its context (4 + 5 + 6); and an authoring loop in which a checker's message drives the repair (2 + 6).

Rule-Governed Memory (Section 6.1) is itself built from two of the other patterns: storing is the Sensor, recall is Rules Steer the Model.

Table 5: Patterns in the applications from which the catalogue was drawn.

Application	Patterns	How
Assignment review	1, 3, 6	quote-checked extraction; chat model calls the rules; completeness check
Interview, suitability, triage	5, 1, 6, 7	flow steers; judge as sensor; format repair; *unclear* to a person
Campervan search	1, 5	one classification per listing, requested by rule
Allocation	4, 5, 1	bounded recommendation; point-in-time evidence; grader verdicts
Agent framework	4, 5, 6, 1	admitted actions; derived phase and context; stalled repair stops; memory claims
Personal memory (DreamRAG lite)	11 (storing)	admission, supersession, conflict and withdrawal of memories as rules; this application has no recall by rules, the program of pattern 11 has
Rule drafting, decision tables	2, 6	model drafts; the checker and a bounded repair decide
Bank reconciliation	8, 5	model proposes causes; rules test them and ask for the deciding date
Policy critic	9	critic reads a sealed policy review; grounded findings gate approval
Model teaching	10	rules label cases, select examples and grade answers

8 Failure patterns

Five couplings recur in working systems and break an invariant of Section 2. Each looks like one of the patterns above, which is why it survives review. We name them so that a design review can look for them; each entry says what it looks like, what it breaks and which pattern replaces it. Most of them were named in the jury assessments of the hackathon of Sections 3.4 and 4.5; the first is the one that Critique and Repair already excludes by name.

Silent Rule Override Code outside the rules changes a verdict before the rules see it: a regular expression in the host that turns a judge's *no* into *yes*, a post-processing step that corrects a value. The outcome may be right, but the proof does not contain the step that decided it, and nobody can tell the correction from the model. It breaks invariant 3, and Critique and Repair (Section 4.3) excludes it by name. In a comparison system built for the hackathon, checked in its code, a host function rewrote the judge's reply before it became a receipt: *yes* became *no* when a pattern found the answer vague, which only tightens; but *unclear* was settled by the pattern in either direction, and an unreachable model was recorded as *yes*. Only a note in the reason text was left, and the rules saw the corrected verdict alone. Instead, write the correction as a rule that reads the receipt, so that `explain` shows it; if the correction can only make the outcome stricter, it is a source in a Fail-Safe Union (Section 4.5).

Circular Evaluation The cases or the gold labels that measure a coupling come from the model, the prompt or the developer that is being measured, or the criteria are chosen after the results are seen. The numbers then measure agreement with oneself. It breaks the claim behind every *Evidence* heading. Instead, fix the cases, the truths and the criteria before the first model call, and derive labels by rules or by construction: an injected fault, a value read from an exact source (Rules Teach the Model, Section 5.3, and C3 in Section 9).

Self-Judge The model that wrote an answer, or a model of the same family with the same prompt, judges whether the answer is right. Its errors and its judgement fail together, and self-correction without an external signal does not help (Huang et al. 2024). It breaks the independence that Several Readers, One Rule (Section 3.1) and the Read-Back Check (Section 3.3) require. Instead, let a rule check what can be checked, and use a different model or a different prompt for the rest (Sections 4.3 and 5.2).

Prose Policy The policy that decides exists only as sentences in a prompt. Nothing can test it, a change cannot be reviewed against cases, and an answer cannot show which sentence it followed. It breaks invariant 3 and leaves the guide's first question unanswered (Section 10). Instead, turn the policy into rules with Model as Author (Section 3.2) and let Rules Steer the Model (Section 4.2) put into the prompt only what one step needs.

Decorative Check A check is computed and recorded, but nothing depends on it: a `quoteFound` that no rule reads, a confidence that is logged, a validation whose failure still lets the action through. It looks like evidence in a review and changes no outcome. It breaks invariant 4. Instead, let an acceptance rule read every check, and look for receipt fields without an outgoing arrow in the rule graph (Section C).

9 Related catalogues

Two bodies of work describe couplings of learning and reasoning. Neuro-symbolic research classifies how a learned and a symbolic component are combined, from Kautz's six types (Kautz 2022) to the boxology of van Harmelen and ten Teije (2019), its modular taxonomy (van Bekkum et al. 2021) and its extension to language models (de Boer et al. 2025), and from the roadmap that unifies language models and knowledge graphs (S. Pan et al. 2024) to recent surveys (d'Avila Garcez and Lamb 2023; Yang et al. 2025; Rani et al. 2026). Software engineering for model-based systems collects architectural patterns for agents (Liu et al. 2025), practitioners' workflow patterns (Schluntz and Zhang 2024), guardrail toolkits (Rebedea et al. 2023) and design patterns for machine-learning applications (Washizaki et al. 2020; Washizaki et al. 2022). Between them lie systems in which a model translates for a solver (L. Pan et al. 2023; Olausson et al. 2023; Ye et al. 2023) or generates candidates for sound critics (Kambhampati et al. 2024).

9.1 Mapping

Table 6 places each pattern in these classifications. Kautz's types name who calls whom; the boxology names the flow of data and symbols; the roadmap names which side enhances which; the agent literature names the mechanism.

9.2 The patterns in boxology notation

The boxology began with van Harmelen and ten Teije (2019): ovals for processes, either deductive (KR) or inductive (ML), and boxes for the symbols or data that processes consume and produce. Van Bekkum et al. (2021) refined it into a vocabulary of four kinds of node, instances (rectangles, data or symbols), models (hexagons, statistical or semantic), processes (ovals: generate, transform, infer) and actors (triangles), and into elementary patterns that compose. De Boer et al. (2025) added the generative model, with an elementary pattern in which a model infers new data from data; the earlier patterns let a model infer only symbols or models, which cannot describe a language model that writes text.

Every pattern in this catalogue carries its structure in this notation, with *model:gen* for the language model and *model:sem* for the rule base, and one mark of ours, the dashed request arrow (Section 2.3 and the legend in the catalogue form).

The drawings show what the notation captures and what it leaves out.

- In the Sensor (Section 3.1) the model's reply is text, that is data; it becomes symbols only through a transformation, which is our receipt, and the rules then infer which symbols count. The notation shows this shape. It cannot show that the first symbol box holds proposals and the second accepted facts, refusals with their reason and *unknown* values; the labels below the boxes add that.

- In Model as Author (Section 3.2) the generative model produces a semantic model, the draft rule, which an inference with the checker and the cases evaluates; an expert adopts it through the elementary pattern in which an actor engineers a model.

- In Rules Steer the Model (Section 4.2) the rules derive the step, one generative model phrases the question, a person answers and a second generative model judges the answer. The closest named pattern is meta-reasoning for control, a symbolic process that directs a learned one; here it does so step by step through a conversation.

- In Rules Teach the Model (Section 5.3) the labels derived by the rules enter the training of the model, the shape of informed learning with prior knowledge.

- Model Proposes, Rules Dispose (Section 4.1), Critique and Repair (Section 4.3) and Rule-Governed Memory (Section 6.1) are loops; the notation draws the arrow back, but the bound, the order of the steps and the supersession of one value by a later one are rules that the drawing cannot show.

Table 6: The thirteen patterns in published classifications. "–" means no counterpart.

Pattern	Kautz type	Boxology, modular patterns	Language-model and agent literature	Status
1 Model as Sensor	Neuro | Symbolic	learning with symbolic output (4.1)	KG construction (S. Pan et al. 2024)	known
2 Model as Author	–	engineered model (1c)	model acquisition (Kambhampati et al. 2024)	known in part
3 Model as Translator	Symbolic[Neuro]	LLM-augmented KG (de Boer et al. 2025)	Logic-LM, LINC, SatLM	known
4 Model Proposes, Rules Dispose	Symbolic[Neuro]	–	guardrails, rule safeguards, hard critics	known
5 Rules Steer the Model	Symbolic[Neuro]	meta-reasoning for control (4.8), training only	routing, dialogue rails	new as stated
6 Critique and Repair	–	–	solver feedback, back-prompting	known
7 Division by Competence	inverse of Neuro[Symbolic]	–	routing; decompiler	new in part
8 Model as Hypothesis Generator	–	–	abduction and hypothesis search	new in part
9 Model as Critic of the Rules	–	–	soft critics review candidates only	new
10 Rules Teach the Model	Neuro:Symbolic→ Neuro	informed learning (4.4)	Constitutional AI, RLAIF	known
11 Rule-Governed Memory	Symbolic[Neuro]	–	memory component (Liu et al. 2025), Reflexion	new as stated
12 Model Selects, Rules Write	Symbolic[Neuro]	–	template-based generation (Reiter and Dale 2000), restricted retrieval answers	known in part
13 Fail-Safe Union	Symbolic[Neuro]	–	deny-overrides (OASIS 2013), layered guardrails	new as stated

Three things the boxes and arrows cannot say are the subject of this catalogue: the direction of control, which the dashed arrows add; the status of a symbol (proposed, accepted, refused, *unknown*) with its date and source; and repetition over time, a retry, a next step or a later supersession. De Boer et al. (2025) name the last gap themselves, and they observe that the notation often yields several drawings of the same system. The four elementary moves of Section 2.3 supply all three: a request carries control, a receipt carries date and source, and accept and ask carry the status.

9.3 What the catalogue adds

Direction. Most published systems let the model drive and call a solver or a tool: Kautz's Neuro[Symbolic], ReAct (Yao et al. 2023), Toolformer (Schick et al. 2023), Logic-LM and LLM-Modulo. All thirteen patterns here are Symbolic[Neuro]: the rules issue requests and the model is a subroutine whose answers return as data. We chose this direction deliberately, for accountability. The agenda of d'Avila Garcez and Lamb (2023) names trust and accountability as aims of tight integration, and loose coupling with explicit provenance reaches both without tight integration.

Time and iteration. De Boer et al. (2025) state that temporal, recurring or iterative aspects "have not yet been taken into account and can currently not be visualised adequately" in the boxology, and that the

same system can be drawn in several ways. Patterns 5, 6 and 8 are iterative by nature, and the dated receipt gives each iteration an identity, an order and a record. Pattern 11 makes time its subject: supersession, withdrawal and the current value as a query are rules over that order.

Three values and receipts. None of the sources above distinguishes a failed call, a malformed reply, an *unclear* judgement and a logical *unknown*, and none of them makes every model answer a replayable observation. These are the invariants of Section 2, and they are what lets patterns 6 and 7 route by result and leave the model's confidence aside.

Patterns without precedent. Rules that derive every step of a conversation from their state (5) are closest to programmable dialogue rails (Rebedea et al. 2023), whose flows are written by hand. Recurring model answers that become rules (7) appear only as offline rule extraction. Asking for the evidence that decides between surviving hypotheses (8) is absent from abduction with language models (Shi et al. 2023; Wang et al. 2024). A model that reviews the rules themselves (9) appears in none of the catalogues we read; in all of them the model is the one being checked. The measurements of Section 10 go beyond the validation the catalogues report, which is coverage of literature (van Harmelen and ten Teije 2019), redrawing published systems (de Boer et al. 2025), or a systematic review with known uses (Liu et al. 2025).

9.4 What the catalogue lacks

Six published patterns have no counterpart, or only an unnamed one; C1 is now named as a variant of the Sensor and C5 has become pattern 11. We rank them by how often they appear and how well they fit the invariants.

C1. **Several answers, one rule.** Ask several models or samples and let rules adjudicate: agreement, majority with a count, *unclear* on disagreement. Voting and debate are agent patterns (Liu et al. 2025), parallelisation by voting is a workflow (Schluntz and Zhang 2024), and self-consistency beat multi-agent debate at the same number of answers (Huang et al. 2024). Three-valued logic suits it: a disagreement is simply *unclear*. It is now named as a variant of the Sensor (Section 3.1); with its own measurements it may become a pattern of its own.

C2. **Gate and route before the model.** Rules decide whether a request may reach a model at all and which model or prompt it goes to: input rails (Rebedea et al. 2023), routing (Schluntz and Zhang 2024). It is implicit in the triage flow, which sends only the careful judgements of emergency signs to the stronger model; it should be named as a variant of pattern 5. Hackathon submissions add a stricter form, a *sovereign gatekeeper*: a local model and rules decide what may reach an external model, pseudonymise it on the way out and translate it back on the way in, and send nothing when the local side fails. It makes the rules the guard of the data flow between two models, and it is a candidate for a pattern of its own once it is built and measured.

C3. **Rules evaluate the model.** Rules reconstruct or audit what a model concluded: rational reconstruction (van Bekkum et al. 2021), knowledge graphs for interpretability (S. Pan et al. 2024), the agent evaluator (Liu et al. 2025). The break-even and capability studies are this pattern run offline; online it is the two-channel disagreement listed under pattern 9. It belongs with pattern 10 as grading, or as a pattern of its own. Its labels are cheapest where the truth holds by construction: a fault the code injected, a value read from an exact oracle such as a barcode or a simulator; several hackathon submissions measured this way, which also avoids Circular Evaluation (Section 8).

C4. **Pooled critics.** Several critics (rules, a model, a person) behind one controller that merges their findings into the next request (Kambhampati et al. 2024). A variant of pattern 6.

C5. **Rule-governed memory.** Retrieval and episodic memory admitted, superseded and cited by rules (Liu et al. 2025; Shinn et al. 2023). It is now Rule-Governed Memory, pattern 11 (Section 6.1), built from the Sensor for storing and Rules Steer the Model for recall; DreamRAG lite implements the storing half, the program of pattern 11 both halves.

C6. **Specification elicitation and the person as participant.** The model helps a person state a goal or a specification, and rules check that it is complete (Kambhampati et al. 2024; Liu et al. 2025); human feedback as a pattern of its own (Witschel et al. 2020). Here the person is a participant in every pattern; the elicitation is the idea of interviewing an expert under pattern 2.

Three families are outside the catalogue on purpose: the model as driver that calls tools (ReAct, Toolformer, Neuro[Symbolic]); rules compiled into the network, Kautz's $\text{Neuro}_{\text{Symbolic}}$ and the tight integration of d'Avila Garcez and Lamb (2023); and training itself.

9.5 What the catalogue learns from them

- **Correctness comes from the sound check.** In LLM-Modulo, soundness is inherited from the hard critics and completeness depends on the model; a model alone produced executable plans in about 12 % of cases, and back-prompting from a sound validator raised this to 82 % in one domain (Kambhampati et al. 2024). The same statement holds for patterns 4 and 8.

- **Self-correction needs an external signal.** Without external feedback, self-correction lowered accuracy, and earlier gains rested on oracle stopping signals (Huang et al. 2024). The rules are that signal in pattern 6; citing self-refinement (Madaan et al. 2023; Shinn et al. 2023) as precedent requires this caveat.

- **Judges have known biases, and rules can mitigate them.** A strong judge agreed with people about as often as people agreed with each other, but gave the same verdict when two answers were swapped only about two thirds of the time; references and examples reduced the failures (Zheng et al. 2023). A rule can ask twice in both orders and require agreement, and supply the reference.

- **Complementary failures justify division.** A solver-backed method and chain-of-thought failed on different cases (Olausson et al. 2023), which is the condition under which pattern 7 pays.

- **Safeguards cover known unknowns.** A rule safeguard around a model covers only the failures its authors foresaw (Washizaki et al. 2020); pair pattern 4 with 8 or 9 where unforeseen failures matter.

- **Form and selection.** Agent patterns are written with context, problem, forces, solution, consequences tied to quality attributes, known uses and related patterns, and come with a decision model that is not empirically validated (Liu et al. 2025); practitioners advise to start with the simplest option and to use evaluator loops only where the criteria are clear (Schluntz and Zhang 2024). The modular taxonomy separates elementary from composed patterns (van Bekkum et al. 2021); the thirteen patterns here are compositions of four elementary moves (request, receipt, accept, ask), made explicit in Section 2.3.

10 Choosing a pattern

A pattern is chosen by the problem it solves and by the forces of the context. For couplings the forces are measurable: how many rules a model would have to apply by itself, whether a reply can be checked, whether the model acts, who must be able to audit the result, and which model is available. This section first asks whether a hybrid is needed at all, then leads from the problem to the patterns, and closes with the measured results the guide rests on. Of the ways to select a pattern that Gamma et al. (1994, sec. 1.7) list, two fit couplings: scanning the intents (Table 4) and starting from the problem one has (Fig. 14 and Table 7). Following practitioners' advice, the guide prefers the simplest coupling that meets the requirement (Schluntz and Zhang 2024).

10.1 First question: is a hybrid needed?

A model alone was the better or equal choice in four measured situations, and the guide says so first.

- **Few rules.** In conversations of up to 48 questions, a reasoning model applying the rules in its prompt was not measurably less accurate than the hybrid (58 or more of 60 against 60 of 60); the accuracy break-even lay at 64 questions.
- **Easy documents, a strong model.** On a clean five-document review, a model alone decided 12 of 12 runs correctly, at 0.28 instead of 0.72 cents and in a third of the time.
- **Cost alone.** Once a provider's discount on cached input is priced, the hybrid's cost advantage in conversations disappears; at 64 and 96 questions the two cost about the same.
- **A sensor that cannot quote.** An 8-billion-parameter open model extracted worse in the hybrid than alone (45 % against 60 %), because a strict quote rule refused correct values.

A hybrid is needed when at least one of the following holds: the decision must come with a proof or an audit trail; the model acts and some actions must never happen; the rules are many or the process is long (from about 64 rules applied in one conversation in our runs); the answer aggregates over many objects (lists, counts, totals); a weaker, local or cheaper model must do the work; or the rules change and the change must be checked before it takes effect.

10.2 From problem to pattern

The questions are not exclusive: a triage line answers *yes* to the first, third and fourth and becomes a composition of 1, 4, 5, 6 and 7. Table 7 gives the same guide as problems, with the force that decides and the caution that the measurements add.

Table 7: Problems, patterns, deciding forces and cautions.

Problem	Patterns	Deciding force	Caution from the measurements
Facts are in documents or speech	1, 6	a reading can be checked (quote, closed list)	fit the acceptance rule to the reader: a strict quote rule hurt a small model
Answer aggregates over many objects	1, then rules	lists and counts are exact in rules	a model alone failed exhaustive lists at 100 agreements; extraction cost is paid once
A policy text must become rules	2, 6	a complete check exists (syntax, cell language, cases)	messages that name the fix doubled first-request repairs
A regulation changed	2, 9	regression cases decide; a person adopts	a draft outside the cell language is refused with its reason

Problem	Patterns	Deciding force	Caution from the measurements
People ask in their own words	3	the answer needs a proof	on easy cases a model alone was as accurate and cheaper
An agent acts	4, 5, 6	some effects must never happen	completion only from recorded evidence
A long protocol or conversation	5, 1, 6, 7	rules per conversation (break-even at 64)	a judge's verdict labels must not be readable as content
A weak or local model must suffice	5 or 1	the hybrid needs less capability	on our ladder a 9B model sufficed for the hybrid; no model up to 70B did as a pure model
Some cases are ambiguous	7	three-valued result as gate	never gate by the model's own confidence
Something must be explained	8, 5	generation needs world knowledge, testing needs exactness	incomplete search is never a "no explanation"
A soft criterion must be judged	1, candidate C1 (Section 9)	no closed check exists	judges are biased by answer order; ask in both orders and require agreement (Zheng et al. 2023)
Semantic faults matter	9	checks are incomplete	model reviews lost semantic faults as the rule base grew; a model critic never replaces a sound check (Huang et al. 2024)
A model must learn the domain	10	labels with proofs	no measured model run yet
What a person said must be remembered	11, 1, 5	sources differ in what they may establish; requests to forget	no measurement yet; similarity retrieval stays outside the rules
An answer must quote a source or state values exactly	12	the content exists as passages or facts	hackathon figures only; the model can still choose the wrong passage
Several checks, harms of unequal weight	13, 7	a wrongly loose outcome costs far more than a wrongly strict one	hackathon figures only; false alarms of every source add up

10.3 What the guide rests on

The cautions come from pre-registered studies in which designs and criteria were fixed before the runs.

Break-even. Synthetic intake protocols of 4 to 96 questions; the hybrid decided 770 of 770 runs correctly; a model applying the rules itself fell to 39 of 60 at 64 questions and 31 of 60 at 96, with 45 of its 50 errors unsafe.

Capability. Without reasoning, a model decided 10 of 30 conversations as a pure model and 30 of 30 in the hybrid. On a ladder of six local models from 1 to 70 billion parameters, the hybrid's floor in single decisions was a 9-billion-parameter model; the pure model had no floor on the ladder, and for it reasoning mattered more than size.

Checking. With one seeded fault per artefact, static checks over rules found every structural fault at every size at no model cost; model reviews of code and prompts lost unreachable rules and missing rules as size grew.

Authoring. Five of five drafting runs ended with a rule passing all hidden cases; error messages that name the fix decided how many repairs were needed.

Aggregation. At 100 agreements, a model alone produced no exact exhaustive list at low effort; the hybrid answered 10 of 10.

Decomposition. With an 8-billion-parameter open model, classifying into 77 intents by first choosing a group and then a label, as rules direct, gained 6 points over one prompt with all labels; putting the groups into one prompt instead collapsed.

The literature adds three cautions that the guide follows: correctness is inherited from the sound check (Kambhampati et al. 2024); a model's self-correction without an external signal does not help (Huang et al. 2024); and a rule safeguard covers the failures its authors foresaw (Washizaki et al. 2020).

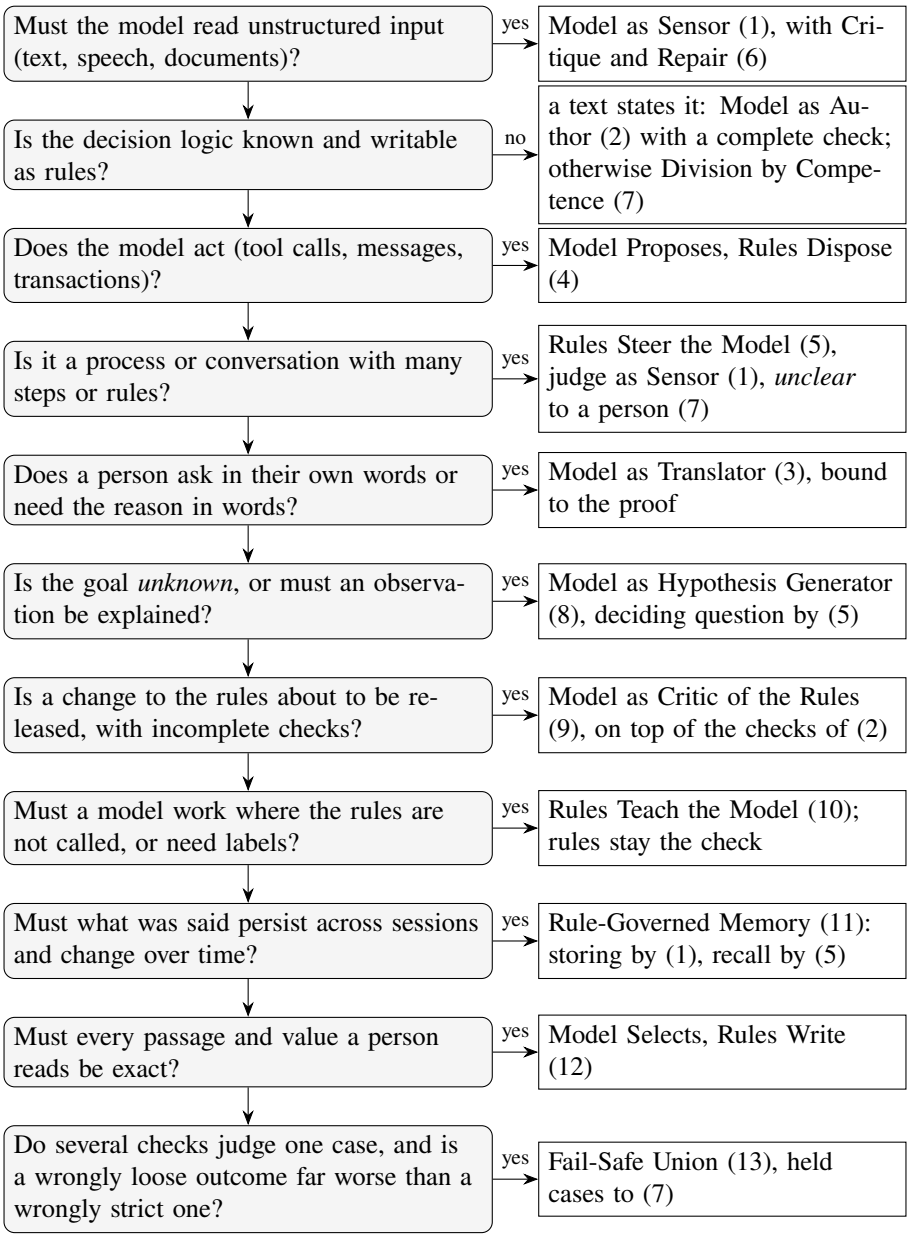

Figure 14: Decision guide. Answer every question; the patterns of all answers that lead right combine. Critique and Repair (6) accompanies every pattern whose reply is checked.

11 Limits

We drew the catalogue from one engine and its applications and then checked it against the literature. It covers every coupling we have seen so far around a symbolic, offline core, and that does not prove it complete. Patterns 12 and 13 come from hackathon submissions; their figures are taken from result files committed in those repositories, and we have not rerun them. Every pattern up to 11 has a program that ran against a model; the programs of patterns 12 and 13 are checked against recorded receipts only. Patterns 4, 7, 8, 10, 11, 12 and 13 have no FrameX measurement of their own yet, and the measurement of pattern 9 is a pilot on easy cases. Most measurements use synthetic protocols and one model family; the capability ladder mixes model families and reasoning modes. Couplings that need a different kind of core (constraint solving and plan search, probabilistic inference over model scores, differentiable logic inside the model, training itself) are outside the catalogue. The rules can check what such a core produces, as they check any external tool.

A Boxology legend

Table 8 lists every element of the Structure diagrams. The node kinds and their labels follow van Bekkum et al. (2021), the generative model follows de Boer et al. (2025); the request arrow and the notes are ours.

Table 8: Elements of the boxology diagrams.

Symbol	Name	Meaning in this catalogue
data: text	instance: data	Model-free input or output, here text: a document, a question, a model's reply before it is read as facts.
symbol	instance: symbol	Model-based structure: facts, requests, proposals, verdicts. A note below the box gives its status, such as *proposed* or *accepted*.
model: gen	model: generative (the LLM)	A language model that infers new data from data (de Boer et al. 2025); the note names its role (dialogue model, judge, critic).
model: sem	model: semantic (the FrameX rule base)	A rule base: the FrameX rules, a flow, a slot policy, a checker with its cases.
infer	process: infer	Applying a model to an instance. With a generative model above it, one model call.
infer: deduce	process: infer, deduce	Applying the rules: deriving, accepting, refusing, deciding.
trans-form	process: transform	Turning one kind of instance into another. Here mostly the receipt: the host turns a reply into symbols and adds its measurements (quote found, decision word).
generate: train	process: generate, train	Building a model from instances (pattern 10).
generate: engineer	process: generate, engineer	A person builds or adopts a model, such as a rule (pattern 2).
person	actor	A person who answers, decides or reviews: applicant, expert, officer, reviewer.
⟶	flow	What a process consumes or produces; from a model to a process, the model that the process applies.
- - - -➝	request (ours)	Control: the inference with the rules issues the call that the process at the arrow's head carries out. It marks the external requests of Section 2.1.
note	note (ours)	The role of a node in the pattern, written below or above it.

B Pattern libraries

A library holds the rules of a pattern, and a program that loads it states facts and only the rules that belong to its domain, in the way the flow library of FrameX serves pattern 5. The companion folder has one library per pattern up to 11 (`library/`); patterns 12 and 13 have none yet. Each library comes with the program of its pattern in library form, tests, and a recorded run. Load the library before the program:

```
$FRAMEX_HOME/target/release/framex test library/sensor.fx library/p01-sensor-lib.fx \
    library/tests/p01-sensor-lib.test.fx
python3 -m framex.tools library/memory.fx library/p11-memory-lib.fx
python3 -m framex.tools --adapter author_adapter.py library/author.fx library/p02-
    author-lib.fx
```

Every name a library defines carries its prefix (`sensor#`, `author#`, `tr#`, `ag#`, `cp#`, `hy#`, `cr#`, `te#`, `mem#`), so it cannot collide with the program's own names. What stays with the program is what only the domain knows: the decision of pattern 3, what the actions of pattern 4 do, which cases pattern 7 decides, the tests of pattern 8, the labels of pattern 10. The request fields keep fixed names; a library meant for other sources would make them parameters. The libraries are prototypes in the companion folder and are meant to move into FrameX beside `flow.fx`.

`sensor.fx`: pattern 1, Model as Sensor

```
// Sensor library (pattern 1, Model as Sensor). Load it before the program.
// The program states facts only:
// ?D:sensor#Source[sensor#text -> "..."; sensor#kind -> ?K; sensor#prompt -> ?P]
// ?K[sensor#allows -> "slot"] what this kind of source may state
// ?P:Prompt[text -> "..."] asks for rows {slot, value, quote}
// The library derives, for every source ?D and answer row ?I:
// ?D[sensor#accepted(?I) -> ?V] quoted, and a statement this kind may make
// ?D[sensor#rejected(?I) -> "reason"] with the reason
// ?D[sensor#notRead -> "error"] the call failed; nothing proposed
closed slot sensor#allows.

// Request: one call per source; the host checks every quote against the text.
external sense(?D) using _llm[prompt -> ?P; document -> ?T; groundIn -> ?T]
  <- ?D:sensor#Source[sensor#text -> ?T; sensor#prompt -> ?Pr] AND ?Pr[text -> ?P].

// Receipt: each answer row is a proposed statement about the source.
?D[sensor#slot(?I) -> ?S; sensor#value(?I) -> ?V; sensor#quoteFound(?I) -> ?Q]
  <- ?D:sensor#Source AND ?R[tool#name -> "_llm"; tool#arg1 -> ?D;
       row(?I, "slot") -> ?S; row(?I, "value") -> ?V; row(?I, "quoteFound") -> ?Q].

// Accept and reject.
?D[sensor#accepted(?I) -> ?V]
  <- ?D:sensor#Source[sensor#kind -> ?K; sensor#slot(?I) -> ?S; sensor#value(?I) ->
      ?V;
       sensor#quoteFound(?I) -> true] AND ?K[sensor#allows -> ?S].

?D[sensor#rejected(?I) -> "quote not found in the source"]
  <- ?D:sensor#Source[sensor#quoteFound(?I) -> false].

?D[sensor#rejected(?I) -> "not a statement this kind of source can make"]
  <- ?D:sensor#Source[sensor#kind -> ?K; sensor#slot(?I) -> ?S]
     AND NOT ?K[sensor#allows -> ?S].

// A failed call proposes nothing.
?D[sensor#notRead -> ?E]
```

```
  <- ?D:sensor#Source AND ?R[tool#name -> "_llm"; tool#arg1 -> ?D;
       tool#status -> "failed"; error -> ?E].
```

author.fx: pattern 2, Model as Author, with Critique and Repair

```
// Author library (pattern 2, Model as Author, with pattern 6, Critique and Repair)
    .
// Load it before the program; the host needs the adapter for _checkCandidate.
// The program states facts only:
// ?P:author#Policy[author#text -> "..."; author#vocabulary -> "...";
// author#draftPrompt -> ?D; author#repairPrompt -> ?R; author#maxAttempts -> N]
// The library derives author#rule(N), author#failed(N) with the checker's message,
// author#adoptable (a recommendation for a person), author#stalled and
    author#gaveUp.

// Attempt 1: the model drafts.
external draft(?P, 1) using _llm[prompt -> ?Pt; policy -> ?T; vocabulary -> ?V;
       tool#order -> 2]
  <- ?P:author#Policy[author#text -> ?T; author#vocabulary -> ?V;
      author#draftPrompt -> ?D]
     AND ?D[text -> ?Pt].

?P[author#rule(?N) -> ?T]
  <- ?P:author#Policy AND ?R[tool#name -> "_llm"; tool#arg1 -> ?P; tool#arg2 -> ?N;
       result("rule") -> ?T].

// Every draft is checked; the cases stay hidden from the model.
external check(?P, ?N) using _checkCandidate[candidate -> ?T; tool#order -> 1]
  <- ?P:author#Policy[author#rule(?N) -> ?T].

?P[author#passed(?N) -> true]
  <- ?P:author#Policy AND ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P;
      tool#arg2 -> ?N;
       result("status") -> "checked"; result("regressions") -> false;
       result("failedAfter") -> 0].

?P[author#adoptable -> ?T] <- ?P:author#Policy[author#passed(?N) -> true;
    author#rule(?N) -> ?T].

?P[author#failed(?N) -> ?F]
  <- ?P:author#Policy AND ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P;
      tool#arg2 -> ?N;
       result("status") -> ?St; result("feedback") -> ?F] AND ?St != "checked".

?P[author#failed(?N) -> ?F]
  <- ?P:author#Policy AND ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P;
      tool#arg2 -> ?N;
       result("regressions") -> true; result("feedback") -> ?F].

?P[author#failed(?N) -> ?F]
  <- ?P:author#Policy AND ?R[tool#name -> "_checkCandidate"; tool#arg1 -> ?P;
      tool#arg2 -> ?N;
       result("failedAfter") -> ?K; result("feedback") -> ?F] AND ?K > 0.

// Repair (pattern 6): ask again with the rule and the problem, as a new request.
?P[author#newProblem(1) -> true] <- ?P:author#Policy[author#failed(1) -> ?F].

```

```
?P[author#newProblem(?N) -> true]
  <- ?P:author#Policy[author#failed(?N) -> ?F] AND ?K := ?N - 1
     AND ?P[author#failed(?K) -> ?G] AND ?F != ?G.

external draft(?P, ?M) using _llm[prompt -> ?Pt; policy -> ?T; vocabulary -> ?V;
       previousRule -> ?Old; problem -> ?F; tool#order -> 2]
  <- ?P:author#Policy[author#text -> ?T; author#vocabulary -> ?V;
      author#maxAttempts -> ?Max;
       author#repairPrompt -> ?Rp; author#rule(?N) -> ?Old; author#failed(?N) -> ?F
           ;
       author#newProblem(?N) -> true]
     AND ?Rp[text -> ?Pt] AND ?N < ?Max AND ?M := ?N + 1.

?P[author#stalled -> ?N]
  <- ?P:author#Policy[author#failed(?N) -> ?F] AND ?K := ?N - 1 AND ?P[
      author#failed(?K) -> ?F].

?P[author#gaveUp -> ?F] <- ?P:author#Policy[author#maxAttempts -> ?N; author#failed
    (?N) -> ?F].
```

`translator.fx`: pattern 3, Model as Translator

```
// Translator library (pattern 3, Model as Translator). Load it before the program.
// The program states:
// ?Q:tr#Question[tr#text -> "..."; tr#prompt -> ?P] the person's words
// "field"[tr#allows -> "value"] the closed list of each field
// tr#phrasing[tr#prompt -> ?P] how a decision is put into words
// and its own decision rules over valid fields only:
// ?Q[tr#decision -> "word"] <- ?Q:tr#Question[tr#valid("f") -> ?V; ...] AND ...
// The library derives tr#field("f"), tr#valid("f"), tr#refused with the field,
// tr#reply or tr#replyRefused.
closed slot tr#allows. closed slot tr#hasBad.

external translate(?Q) using _llm[prompt -> ?P; question -> ?T]
  <- ?Q:tr#Question[tr#text -> ?T; tr#prompt -> ?Pr] AND ?Pr[text -> ?P].

// Every answer field the program declares, read from the receipt.
?Q[tr#field(?F) -> ?V]
  <- ?Q:tr#Question AND ?F[tr#allows -> ?X]
     AND ?R[tool#name -> "_llm"; tool#arg1 -> ?Q; result(?F) -> ?V].

// A field inside its closed list is valid; decision rules read only valid fields.
?Q[tr#valid(?F) -> ?V] <- ?Q:tr#Question[tr#field(?F) -> ?V] AND ?F[tr#allows -> ?V
    ].

// A field outside its closed list refuses the request and names the field.
?Q[tr#bad(?F) -> true] <- ?Q:tr#Question[tr#field(?F) -> ?V] AND NOT ?F[tr#allows
    -> ?V].
?Q[tr#refused -> ?F] <- ?Q:tr#Question[tr#bad(?F) -> true].
?Q[tr#hasBad -> true] <- ?Q:tr#Question[tr#bad(?F) -> true].

// For display only: every declared field valid. Decision rules must not read it,
    because
// it negates over receipts, and the reply request depends on the decision.
?Q[tr#runs -> true]
  <- ?Q:tr#Question AND ?R[tool#name -> "_llm"; tool#arg1 -> ?Q; result(?F) -> ?V]
     AND NOT ?Q[tr#hasBad -> true].
```

```

// From the decision to words; the reply counts only if it says the decision word.
external phrase(?Q) using _llm[prompt -> ?P; decision -> ?D; mustSay -> ?D]
  <- ?Q:tr#Question[tr#decision -> ?D] AND tr#phrasing[tr#prompt -> ?Pr] AND ?Pr[
      text -> ?P].

?Q[tr#reply -> ?Text]
  <- ?Q:tr#Question AND ?X[tool#name -> "_llm"; tool#arg1 -> ?Q;
       result("saysDecision") -> true; result("text") -> ?Text].

?Q[tr#replyRefused -> "the reply does not state the decision"]
  <- ?Q:tr#Question AND ?X[tool#name -> "_llm"; tool#arg1 -> ?Q;
       result("saysDecision") -> false].
```

agent.fx: pattern 4, Model Proposes, Rules Dispose

```
// Agent library (pattern 4, Model Proposes, Rules Dispose). Load it before the
    program.
// The program states:
// ?T:ag#Task[ag#goal -> "..."; ag#actions -> "..."; ag#maxSteps -> N; ag#prompt ->
     ?P]
// ?T[ag#grants -> "action"] what the task may do
// and, for each granted action, the rules that carry it out and report:
// ?T[ag#outcome(?N) -> "..."] what became of step ?N
// ?T[ag#evidence -> true] what completion rests on
// The library derives the proposals, admission or refusal with a reason, the next
    step,
// and ?T[ag#verdict -> "completed"] or "stopped at the step bound".
closed slot ag#grants.

external propose(?T, ?N) using _llm[prompt -> ?P; goal -> ?G; actions -> ?Acts;
       outcome -> "none"; tool#order -> 2]
  <- ?T:ag#Task[ag#goal -> ?G; ag#actions -> ?Acts; ag#prompt -> ?Pr] AND ?Pr[text
      -> ?P]
     AND ?N := 1.

?R:ag#Proposal[ag#task -> ?T; ag#step -> ?N; ag#action -> ?K]
  <- ?T:ag#Task AND ?R[tool#name -> "_llm"; tool#arg1 -> ?T; tool#arg2 -> ?N;
       result("action") -> ?K].

?A[ag#allowed -> true]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#action -> ?K] AND ?T[ag#grants -> ?K].

?A[ag#refused -> "action not granted for this task"]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#action -> ?K] AND NOT ?T[ag#grants -> ?K].

// The next step carries the outcome of the last one; maxSteps stops the loop.
?T[ag#next(?M) -> ?W]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#step -> ?N; ag#refused -> ?W]
     AND ?T[ag#maxSteps -> ?Max] AND ?N < ?Max AND ?M := ?N + 1.

?T[ag#next(?M) -> ?X]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#step -> ?N; ag#allowed -> true]
     AND ?T[ag#outcome(?N) -> ?X; ag#maxSteps -> ?Max] AND ?N < ?Max AND ?M := ?N +
         1.

external propose(?T, ?N) using _llm[prompt -> ?P; goal -> ?G; actions -> ?Acts;
```

```
      outcome -> ?W; tool#order -> 2]
  <- ?T:ag#Task[ag#goal -> ?G; ag#actions -> ?Acts; ag#prompt -> ?Pr; ag#next(?N)
      -> ?W]
     AND ?Pr[text -> ?P].

// Completion only on recorded evidence.
?T[ag#verdict -> "completed"]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#action -> "finish"; ag#allowed -> true]
     AND ?T[ag#evidence -> true].

?T[ag#verdict -> "stopped at the step bound"]
  <- ?A:ag#Proposal[ag#task -> ?T; ag#step -> ?N; ag#action -> ?K] AND ?T[
      ag#maxSteps -> ?N]
     AND ?K != "finish".
```

`competence.fx`: pattern 7, Division by Competence

```
// Competence library (pattern 7, Division by Competence). Load it before the
    program.
// The program states the cases and its own rules for what it decides:
// ?C:cp#Case[cp#file -> "..."; cp#prompt -> ?P]
// ?C[cp#decided -> true] <- ... the cases the rules cover
// The library sends every other case to the model for an assessment and hands it
    on to a
// person, ?C[cp#handOn -> true], with the assessment as ?C[cp#note -> "..."].
closed slot cp#decided.

external assess(?C) using _llm[prompt -> ?P; file -> ?T]
  <- ?C:cp#Case[cp#file -> ?T; cp#prompt -> ?Pr] AND ?Pr[text -> ?P]
     AND NOT ?C[cp#decided -> true].

?C[cp#assessment -> ?V]
  <- ?C:cp#Case AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C; result("verdict") ->
      ?V].

?C[cp#handOn -> true] <- ?C:cp#Case AND NOT ?C[cp#decided -> true].

?C[cp#note -> ?W]
  <- ?C:cp#Case[cp#handOn -> true] AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C;
       result("reason") -> ?W].

?C[cp#note -> "no assessment: the model call failed"]
  <- ?C:cp#Case[cp#handOn -> true] AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C;
       tool#status -> "failed"].
```

`hypothesis.fx`: pattern 8, Model as Hypothesis Generator

```
// Hypothesis library (pattern 8, Model as Hypothesis Generator). Load it before
    the program.
// The program states the observation and its own tests:
// ?O:hy#Observation[hy#text -> "..."; hy#prompt -> ?P; hy#label -> "key"]
// ?O[hy#refuted(?I) -> true] <- ... hy#candidate(?I, "key") ... what rules out a
    candidate
// The prompt asks for rows, one per candidate; hy#label names the row field that a
// person is asked about. The library derives the candidates, the survivors and
// ?O[hy#askAbout(?I) -> "..."]. A candidate never becomes a fact.
```

```
closed method hy#refuted/1.

external propose(?O) using _llm[prompt -> ?P; observation -> ?T]
  <- ?O:hy#Observation[hy#text -> ?T; hy#prompt -> ?Pr] AND ?Pr[text -> ?P].

// Every field of every row is kept, keyed by row and field name.
?O[hy#candidate(?I, ?K) -> ?V]
  <- ?O:hy#Observation AND ?R[tool#name -> "_llm"; tool#arg1 -> ?O; tool#status ->
      "succeeded";
       row(?I, ?K) -> ?V].

?O[hy#survives(?I) -> true]
  <- ?O:hy#Observation[hy#label -> ?L; hy#candidate(?I, ?L) -> ?E]
     AND NOT ?O[hy#refuted(?I) -> true].

?O[hy#askAbout(?I) -> ?E]
  <- ?O:hy#Observation[hy#label -> ?L; hy#survives(?I) -> true; hy#candidate(?I, ?L
      ) -> ?E].
```

`critic.fx`: pattern 9, Model as Critic of the Rules

```
// Critic library (pattern 9, Model as Critic of the Rules). Load it before the
    program.
// The program states the change:
// ?C:cr#Change[cr#rationale -> "..."; cr#review -> "..."; cr#prompt -> ?P;
// cr#reviewed -> true; cr#changed -> "case"]
// ?C[cr#disposed(?I) -> true] a reviewer has dealt with remark ?I
// The library derives grounded remarks cr#remark(?I), cr#blocked for a failed call
    , and
// cr#approved once the critique ran and every remark is disposed of.
closed method cr#disposed/1. closed slot cr#hasOpen. closed slot cr#critiqued.

external critique(?C) using _llm[prompt -> ?P; rationale -> ?Ra; review -> ?Rv;
    groundIn -> ?Ra]
  <- ?C:cr#Change[cr#rationale -> ?Ra; cr#review -> ?Rv; cr#prompt -> ?Pr] AND ?Pr[
      text -> ?P].

// A remark counts only with a found quote and a case the change really changed.
?C[cr#remark(?I) -> ?K]
  <- ?C:cr#Change[cr#changed -> ?K] AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C;
       row(?I, "case") -> ?K; row(?I, "quoteFound") -> true].

?C[cr#critiqued -> true]
  <- ?C:cr#Change AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C; tool#status -> "
      succeeded"].

?C[cr#blocked -> ?E]
  <- ?C:cr#Change AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C; tool#status -> "
      failed";
       error -> ?E].

?C[cr#hasOpen -> true] <- ?C:cr#Change[cr#remark(?I) -> ?K] AND NOT ?C[cr#disposed(
    ?I) -> true].

?C[cr#approved -> true]
  <- ?C:cr#Change[cr#reviewed -> true; cr#critiqued -> true] AND NOT ?C[cr#hasOpen
      -> true].
```

`teach.fx`: pattern 10, Rules Teach the Model

```
// Teach library (pattern 10, Rules Teach the Model). Load it before the program.
// The program states the cases and its own labelling rules:
// ?C:te#Case[te#description -> "..."; te#prompt -> ?P]
// ?C[te#label -> "decision"] <- ... ?C[te#gap -> "requirement"] <- ...
// The prompt asks for {"decision": ..., "rows": [{"gap": ...}]}. The library derives one
// te#Answer per case and grades it: te#right, te#missedGap.
closed slot te#named.

external answer(?C) using _llm[prompt -> ?P; description -> ?T]
  <- ?C:te#Case[te#description -> ?T; te#prompt -> ?Pr] AND ?Pr[text -> ?P].

?R:te#Answer[te#case -> ?C; te#decision -> ?D]
  <- ?C:te#Case AND ?R[tool#name -> "_llm"; tool#arg1 -> ?C; result("decision") -> ?D].

?R[te#named -> ?G] <- ?R:te#Answer AND ?R[row(?I, "gap") -> ?G].

?A[te#right -> true] <- ?A:te#Answer[te#case -> ?C; te#decision -> ?D] AND ?C[te#label -> ?D].

?A[te#missedGap -> ?G]
  <- ?A:te#Answer[te#case -> ?C] AND ?C[te#gap -> ?G] AND NOT ?A[te#named -> ?G].
```

`memory.fx`: pattern 11, Rule-Governed Memory

```
// Memory library (pattern 11, Rule-Governed Memory). Load it before the program.
// The program states facts only:
// ?K[mem#mayEstablish -> "slot"] what a kind of source may establish
// ?U:mem#Utterance[mem#kind -> ?K; mem#seq -> N; mem#text -> "..."; mem#prompt -> ?P]
// ?S:mem#Step[mem#about -> "who"; mem#needs -> "slot"] what one step may see
// ?Q:mem#Question[mem#step -> ?S; mem#text -> "..."; mem#prompt -> ?P]
// ?U[mem#forgotten -> true] a person's request to forget
// The library derives statements st(?U, ?I) with mem#accepted, mem#rejected, mem#current,
// the context of each step, and ?Q[mem#answer -> "..."] or ?Q[mem#refused -> ...].
closed slot mem#mayEstablish. closed slot mem#forgotten. closed slot mem#hasNewer.
closed slot mem#current. closed slot mem#hasBad.
closed method mem#sentId/1.

// Storing (pattern 1): one statement object per quoted answer row.
external remember(?U) using _llm[prompt -> ?P; utterance -> ?T; groundIn -> ?T; tool#order -> 2]
  <- ?U:mem#Utterance[mem#text -> ?T; mem#prompt -> ?Pr] AND ?Pr[text -> ?P]
     AND NOT ?U[mem#forgotten -> true].

st(?U, ?I):mem#Statement[mem#from -> ?U; mem#about -> ?A; mem#slot -> ?S; mem#value -> ?V]
  <- ?U:mem#Utterance AND ?R[tool#name -> "_llm"; tool#arg1 -> ?U; row(?I, "about") -> ?A;
       row(?I, "slot") -> ?S; row(?I, "value") -> ?V; row(?I, "quoteFound") -> true].

// Admission, rejection, supersession.
```

```
?St[mem#accepted -> true]
  <- ?St:mem#Statement[mem#from -> ?U; mem#slot -> ?S] AND ?U[mem#kind -> ?K]
     AND ?K[mem#mayEstablish -> ?S] AND NOT ?U[mem#forgotten -> true].

?St[mem#rejected -> "this kind of source cannot establish it"]
  <- ?St:mem#Statement[mem#from -> ?U; mem#slot -> ?S] AND ?U[mem#kind -> ?K]
     AND NOT ?K[mem#mayEstablish -> ?S].

?St[mem#hasNewer -> true]
  <- ?St:mem#Statement[mem#accepted -> true; mem#from -> ?U; mem#about -> ?A;
      mem#slot -> ?S]
     AND ?U[mem#seq -> ?N]
     AND ?Nw:mem#Statement[mem#accepted -> true; mem#from -> ?W; mem#about -> ?A;
         mem#slot -> ?S]
     AND ?W[mem#seq -> ?M] AND ?M > ?N.

?St[mem#current -> true] <- ?St:mem#Statement[mem#accepted -> true] AND NOT ?St[
    mem#hasNewer -> true].

// Forgetting also reaches the answer store.
external forget(?U) using _forgetAnswers[text -> ?T; tool#order -> 1]
  <- ?U:mem#Utterance[mem#forgotten -> true; mem#text -> ?T].

// Recall (pattern 5): a step sees the current statements it needs, and nothing
    else.
?St[selectedFor -> ?K]
  <- ?K:mem#Step[mem#about -> ?A; mem#needs -> ?S]
     AND ?St:mem#Statement[mem#current -> true; mem#about -> ?A; mem#slot -> ?S].

external answer(?Q) using _llm[prompt -> ?P; question -> ?T; contextKey -> ?K;
    tool#order -> 3]
  <- ?Q:mem#Question[mem#text -> ?T; mem#step -> ?K; mem#prompt -> ?Pr] AND ?Pr[
      text -> ?P].

// Citations: every cited id must have been sent, and its statement must be current
    .
?R[mem#sentId(?C) -> true] <- ?R[tool#sent(?C) -> ?St].
?Q[mem#cites -> ?St]
  <- ?Q:mem#Question AND ?R[tool#arg1 -> ?Q; row(?I, "cited") -> ?C; tool#sent(?C)
      -> ?St].
?Q[mem#refused -> ?C]
  <- ?Q:mem#Question AND ?R[tool#arg1 -> ?Q; row(?I, "cited") -> ?C]
     AND NOT ?R[mem#sentId(?C) -> true].
?Q[mem#refused -> ?St] <- ?Q:mem#Question[mem#cites -> ?St] AND NOT ?St[mem#current
     -> true].
?Q[mem#hasBad -> true] <- ?Q:mem#Question[mem#refused -> ?X].
?Q[mem#answer -> ?Text]
  <- ?Q:mem#Question[mem#cites -> ?St] AND ?R[tool#arg1 -> ?Q; result("text") ->
      ?Text]
     AND NOT ?Q[mem#hasBad -> true].
```

C Rule graph legend

Table 9 lists the elements of the rule graphs and the flow diagram. The script `rulegraph.py` of the companion folder draws them from a program with Graphviz: `python3 rulegraph.py p01-sensor.fx > p01.dot`, then `dot -Tpdf p01.dot -o p01.pdf`.

Table 9: Elements of the rule graphs and the flow diagram.

Symbol	Meaning
Document	Facts of a class or object, with their slots, and the prompt objects.
read → _llm	An external request and its tool.
receipt	The receipt the host writes, with the fields the rules read: `result("k")`, `row("k")`, `tool#status`, `error`.
accepted(I)	Slots or classes that the rules derive. A box lists the slots of one rule head.
l. 23	One rule. The label is the line of the program listing where the rule starts: "l. 23" is the rule in line 23.
l. 14	A rule that issues an external request, labelled by its line in the same way.
⟶	Into a rule: the rule reads it. Out of a rule: the rule derives it.
- - - -┤	The rule reads it under `NOT`.
- - - -➤	The rule issues the request.
·····➤	The host carries out the request and writes the receipt.
Collect	Flow diagram: a step that asks a question and lets a judge decide its criterion.
Say	Flow diagram: a step that shows a text.
• ◎	Start and end of the flow.
↻	A rejected attempt: the question is asked again while attempts remain.

References

Alexander, Christopher et al. (1977). *A Pattern Language: Towns, Buildings, Construction*. New York: Oxford University Press.

Bai, Yuntao et al. (2022). *Constitutional AI: Harmlessness from AI Feedback*. arXiv: 2212.08073.

Buschmann, Frank et al. (1996). *Pattern-Oriented Software Architecture, Volume 1: A System of Patterns*. Chichester: Wiley.

D'Avila Garcez, Artur and Luís C. Lamb (2023). "Neurosymbolic AI: The 3rd Wave". In: *Artificial Intelligence Review* 56.11, pp. 12387–12406. DOI: 10.1007/s10462-023-10448-w.

De Boer, Maaike et al. (2025). "Design Patterns for Large Language Model Based Neuro-Symbolic Systems". In: *Neurosymbolic Artificial Intelligence* 1, 29498732251377499. DOI: 10.1177/29498732251377499.

Gamma, Erich et al. (1994). *Design Patterns: Elements of Reusable Object-Oriented Software*. Reading, MA: Addison-Wesley.

Gou, Zhibin et al. (2024). "CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing". In: *International Conference on Learning Representations*. URL: https://openreview.net/forum?id=Sx038qxjek.

Huang, Jie et al. (2024). "Large Language Models Cannot Self-Correct Reasoning Yet". In: *International Conference on Learning Representations*. URL: https://openreview.net/forum?id=IkmD3fKBPQ.

Kambhampati, Subbarao et al. (2024). "Position: LLMs Can't Plan, But Can Help Planning in LLM-Modulo Frameworks". In: *Proceedings of the 41st International Conference on Machine Learning*. Vol. 235. PMLR, pp. 22895–22907.

Kautz, Henry A. (2022). "The Third AI Summer: AAAI Robert S. Engelmore Memorial Lecture". In: *AI Magazine* 43.1, pp. 105–125. DOI: 10.1002/aaai.12036.

Kifer, Michael, Georg Lausen, and James Wu (1995). "Logical Foundations of Object-Oriented and Frame-Based Languages". In: *Journal of the ACM* 42.4, pp. 741–843. DOI: 10.1145/210332.210335.

Lee, Harrison et al. (2024). "RLAIF vs. RLHF: Scaling Reinforcement Learning from Human Feedback with AI Feedback". In: *Proceedings of the 41st International Conference on Machine Learning*. Vol. 235. PMLR, pp. 26874–26901.

Liu, Yue et al. (2025). "Agent Design Pattern Catalogue: A Collection of Architectural Patterns for Foundation Model Based Agents". In: *Journal of Systems and Software* 220, 112278. DOI: 10.1016/j.jss.2024.112278.

Madaan, Aman et al. (2023). "Self-Refine: Iterative Refinement with Self-Feedback". In: *Advances in Neural Information Processing Systems*. Vol. 36, pp. 46534–46594.

OASIS (Jan. 2013). *eXtensible Access Control Markup Language (XACML) Version 3.0*. OASIS Standard. OASIS.

Olausson, Theo et al. (2023). "LINC: A Neurosymbolic Approach for Logical Reasoning by Combining Language Models with First-Order Logic Provers". In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 5153–5176. DOI: 10.18653/v1/2023.emnlp-main.313.

Pan, Liangming et al. (2023). "Logic-LM: Empowering Large Language Models with Symbolic Solvers for Faithful Logical Reasoning". In: *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 3806–3824. DOI: 10.18653/v1/2023.findings-emnlp.248.

Pan, Shirui et al. (2024). "Unifying Large Language Models and Knowledge Graphs: A Roadmap". In: *IEEE Transactions on Knowledge and Data Engineering* 36.7, pp. 3580–3599. DOI: 10.1109/TKDE.2024.3352100.

Rani, Maneeha, Bhupesh Kumar Mishra, and Dhavalkumar Thakker (2026). "Neurosymbolic Large Language Models: A Survey of Symbolic Integration, Reasoning and Explainability". In: *Information Systems Frontiers*. Online first. DOI: 10.1007/s10796-026-10794-4.

Rebedea, Traian et al. (2023). "NeMo Guardrails: A Toolkit for Controllable and Safe LLM Applications with Programmable Rails". In: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 431–445. DOI: 10.18653/v1/2023.emnlp-demo.40.

Reiter, Ehud and Robert Dale (2000). *Building Natural Language Generation Systems*. Cambridge: Cambridge University Press.

Schick, Timo et al. (2023). "Toolformer: Language Models Can Teach Themselves to Use Tools". In: *Advances in Neural Information Processing Systems*. Vol. 36, pp. 68539–68551.

Schluntz, Erik and Barry Zhang (Dec. 19, 2024). *Building Effective Agents*. Anthropic. URL: https://www.anthropic.com/research/building-effective-agents (visited on 10/12/2026).

Shi, Xiaoming et al. (2023). "Language Models Can Improve Event Prediction by Few-Shot Abductive Reasoning". In: *Advances in Neural Information Processing Systems*. Vol. 36, pp. 29532–29557.

Shinn, Noah et al. (2023). "Reflexion: Language Agents with Verbal Reinforcement Learning". In: *Advances in Neural Information Processing Systems*. Vol. 36, pp. 8634–8652.

Van Bekkum, Michael et al. (2021). "Modular Design Patterns for Hybrid Learning and Reasoning Systems: A Taxonomy, Patterns and Use Cases". In: *Applied Intelligence* 51.9, pp. 6528–6546. DOI: 10.1007/s10489-021-02394-3.

Van Harmelen, Frank and Annette ten Teije (2019). "A Boxology of Design Patterns for Hybrid Learning and Reasoning Systems". In: *Journal of Web Engineering* 18.1–3, pp. 97–124. DOI: 10.13052/jwe1540-9589.18133.

Wang, Ruocheng et al. (2024). "Hypothesis Search: Inductive Reasoning with Language Models". In: *International Conference on Learning Representations*. URL: https://openreview.net/forum?id=G7UtIGQmjm.

Washizaki, Hironori et al. (2020). "Software Engineering Patterns for Machine Learning Applications (SEP4MLA) – Part 2". In: *Proceedings of the 27th Conference on Pattern Languages of Programs (PLoP '20)*. The Hillside Group.

Washizaki, Hironori et al. (2022). "Software-Engineering Design Patterns for Machine Learning Applications". In: *Computer* 55.3, pp. 30–39. DOI: 10.1109/MC.2021.3137227.

Witschel, Hans Friedrich et al. (2020). "Visualization of Patterns for Hybrid Learning and Reasoning with Human Involvement". In: *New Trends in Business Information Systems and Technology*. Studies in Systems, Decision and Control. Cham: Springer, pp. 193–204. DOI: 10.1007/978-3-030-48332-6_13.

Yang, Xiao-Wen et al. (2025). "Neuro-Symbolic Artificial Intelligence: Towards Improving the Reasoning Abilities of Large Language Models". In: *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence (IJCAI-25), Survey Track*, pp. 10770–10778. DOI: 10.24963/ijcai.2025/1195.

Yao, Shunyu et al. (2023). "ReAct: Synergizing Reasoning and Acting in Language Models". In: *International Conference on Learning Representations*.

Ye, Xi et al. (2023). "SatLM: Satisfiability-Aided Language Models Using Declarative Prompting". In: *Advances in Neural Information Processing Systems*. Vol. 36, pp. 45548–45580.

Zheng, Lianmin et al. (2023). "Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena". In: *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*. Vol. 36, pp. 46595–46623.