arXiv:2611.06612v1 [cs.LG] 4 Nov 2026

From Reconstruction to Imputation: Test-Time Optimization for Missing Time-Series Channels

Björn Ludwig*, **Alexander Windmann**, **René Heesch**, **Oliver Niggemann**

HSU-AI – Institute for Artificial Intelligence
Helmut Schmidt University, Hamburg, Germany
bjoern.ludwig@hsu-hh.de

Abstract

Complete sensor-channel failures create a difficult form of time-series missingness: local interpolation is unavailable, and deploying a separately trained imputer for every failure pattern may be impractical. We ask whether a sequence reconstructor trained without artificial mask augmentation and then frozen can be converted into an imputer at inference time. Our method treats the missing entries of each test window as optimization variables, clamps all observations, and backpropagates through the frozen model to minimize reconstruction error only on observed coordinates. The resulting observed-self-consistency objective requires neither artificial mask augmentation nor full-channel dropout during reconstructor training. We characterize when this refinement is informative through the cross-Jacobian from missing inputs to observed residuals, and use an optional temporal-smoothness prior when that coupling is weak. Across five real-world datasets with exactly one channel synthetically masked per test window and four reconstruction architectures, no-prior test-time refinement reduces MAE relative to one-pass reconstruction in 19 of 20 dataset–backbone combinations. Compared with mask-aware imputers trained under either standard structured masks or matched full-channel dropout, the best refined reconstructor achieves the lowest MAE on two datasets and trails the strongest baseline on three. These gains require substantially more inference time and are sensitive to cross-channel coupling. Our results establish test-time self-consistency as a conditional inference-time adapter for frozen reconstruction models that complements mask-aware imputation without retraining the reconstructor.

1 Introduction

Missing values are common in multivariate time series from environmental monitoring, energy systems, and traffic sensing (De Vito et al. 2008; Zhang et al. 2017; Trindade 2015; Liu et al. 2023a). An especially hard case arises when one complete channel is unavailable. Unlike isolated points or short gaps, the missing channel provides no local observations within the affected window, so recovery must rely on temporal structure learned elsewhere and on dependencies with the remaining channels (Khayati et al. 2020; Cini, Marisca, and Alippi 2021; Mitra et al. 2023).

This paper asks a specific operational question: *Given a frozen sequence reconstructor, can it be turned into an imputer for a previously unseen full-channel mask without retraining its weights?*

*Corresponding author.

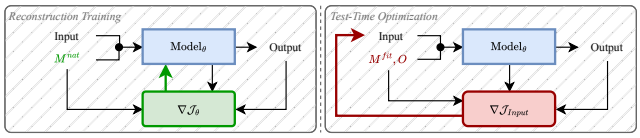

Figure 1: **High-level overview.** Left: train a reconstructor on baseline-filled windows using observed-entry loss. Right: TTO updates missing inputs using frozen model to reduce fit-observed residuals and yield the imputation.

Deep recurrent, attention-based, and diffusion-based imputers learn conditional completions or sampling procedures during training (Cao et al. 2018; Du, Côté, and Liu 2023; Tashiro et al. 2021). Mask-aware methods can explicitly condition on the observation mask and enforce consistency with available values. Instead of learning or amortizing the completion for the target mask, we optimize the missing coordinates of each test window by differentiating through a fixed reconstructor. A useful reconstructor must couple the missing coordinates to observed outputs through a bottleneck or another learned dependency, since an identity-like model whose observed outputs are independent of the missing inputs provides no gradient with which to recover the missing channel.

Figure 1 summarizes the two phases. We first train a sequence reconstructor with a reconstruction loss on observed entries and without artificial mask augmentation, then freeze its parameters. At inference time, Test-Time Optimization (TTO) treats missing entries as free variables, clamps observations, and minimizes the frozen model residual only on observed coordinates. The resulting observed-self-consistency objective can incorporate simple priors or feasibility constraints when the data term alone is underconstrained. Its informativeness is governed locally by the cross-Jacobian from optimized missing inputs to residuals on observed outputs: if this coupling is weak or rank-deficient, multiple completions can be similarly self-consistent.

Contributions. We make four contributions. (1) We formulate observed-self-consistency TTO as an inference-time adapter that turns a frozen differentiable reconstructor into a point-estimate imputer while clamping observations and fixing model parameters. (2) We characterize its local infor-

mativeness through the missing-to-observed cross-Jacobian and explain how temporal smoothness regularizes weakly constrained directions. (3) Across five real-world datasets, four architectures, and synthetic one-channel masks, TTO improves over forward-only reconstruction in 19 of 20 dataset–backbone combinations. (4) Against mask-aware imputers trained with structured masks or matched full-channel dropout, the best TTO result wins on ETTh1 and Electricity but trails on BeijingAir, ItalyAir, and PeMS, while requiring substantially more inference time.

Together, these results support TTO as a conditional adapter when a frozen reconstructor provides informative missing-to-observed coupling, not as a blanket replacement for mask-aware imputation models.

2 Related Work

Deep time series imputation: from recurrent baselines to modern generative models. Deep imputation methods have become the dominant approach for handling missing values in multivariate time series, as they can model nonlinear dynamics and cross-variable dependencies under structured missingness (Wang et al. 2025). Most deep imputers perform amortized inference, producing imputations by evaluating a trained network (or a fixed learned sampling procedure) on a partially observed window, rather than optimizing missing values per instance (Cao et al. 2018; Mattei and Frellsen 2019; Tashiro et al. 2021; Du, Côté, and Liu 2023). Within this paradigm, the community has progressed from recurrent baselines to attention-based and generative formulations, with each generation improving robustness to longer-range dependencies and more complex missingness patterns (Wen et al. 2023).

Predictive (point-estimate) imputers. Recurrent models constitute a canonical starting point due to their sequential inductive bias. Earlier RNN-based approaches such as GRU-D introduced explicit decay mechanisms to handle irregular sampling and missingness, influencing later recurrent imputers (Che et al. 2018). BRITS is a widely used recurrent baseline that performs bidirectional recurrent imputation to propagate information from both temporal directions (Cao et al. 2018). Beyond recurrence, attention mechanisms improved long-range modeling under missingness; SAITS is a representative attention-based imputer that leverages masked self-attention blocks and missingness indicators (Du, Côté, and Liu 2023).

Graph-based formulations explicitly model cross-variable structure by learning dependencies between features. GRAPE is a prominent example that addresses missingness through graph representation learning, making cross-channel coupling an explicit modeling object rather than an implicit byproduct of sequence modeling (You et al. 2020). In parallel, convolutional and frequency-inspired designs demonstrate that strong performance can be achieved by restructuring temporal variation. TimesNet is a competitive backbone that uses FFT-based reformulation and is frequently adopted as a high-performing baseline across time-series tasks (Wu et al. 2022).

Generative imputers and uncertainty. Generative approaches address the ambiguity of imputation by learning distributions over plausible completions rather than producing a single point estimate. MIWAE provides a likelihood-based framework for imputing incomplete data with deep latent-variable models (Mattei and Frellsen 2019), while GP-VAE incorporates a Gaussian-process prior to model temporal dependencies under missingness (Fortuin et al. 2020). Diffusion and score-based methods are an important generative paradigm for time-series imputation; CSDI is a representative diffusion-based method that conditions generation on observed values to impute missing segments (Tashiro et al. 2021). Overall, these generative methods provide principled ways to represent uncertainty. Although diffusion samplers are themselves iterative, their test-time iterations follow a learned conditional denoising process rather than directly optimizing missing entries for each instance through a frozen reconstructor.

Iterative imputation and test-time refinement. A complementary line of work emphasizes *iterative* imputation, where solutions are refined through repeated updates rather than a single feed-forward pass. HyperImpute formalizes generalized iterative imputation with automatic model selection (Jarrett et al. 2022), and IGRM performs imputation through iterative graph reconstruction (Zhong, Gui, and Ye 2023). Related perspectives include distribution-matching objectives for missing-value imputation (Zhao et al. 2023). Our method aligns with the broader message that per-instance refinement can be beneficial, but differs in mechanism and scope: instead of alternating model fitting and filling or employing bespoke iterative update rules, we optimize missing entries *directly* by differentiating through a *frozen* sequence reconstructor, yielding an inference-time adapter for differentiable reconstructors with informative missing-to-observed coupling. Conceptually, our method follows the general idea of optimizing inputs while keeping a differentiable model fixed, as in differentiable planning (Wu, Say, and Sanner 2020) and network inversion (Haim et al. 2022). Unlike these settings, our objective is mask-conditioned imputation as we optimize only missing entries to enforce self-consistency on observed values, rather than maximizing reward or reconstructing training data.

Takeaway. Taken together, current competitive imputers provide powerful amortized mappings across diverse architectures, but complete-channel masks remain difficult because no within-channel observations are available and recovery depends on learned cross-channel structure. This motivates our focus on test-time optimization: enforce observed self-consistency under a frozen reconstructor and regularize ill-posed cases with simple priors and feasibility constraints, conditionally bridging amortized imputation with iterative per-instance inference rather than replacing mask-aware models.

3 Problem Setting

We consider a window $X \in \mathbb{R}^{T \times D}$ with D measurement channels and length T . We use *channel* for one dimension

$d \in \{1, \dots, D\}$ and *coordinate* for one time–channel entry (t, d).

Mask contract. Let $M^{\mathrm{nat}} \in \{0,1\}^{T \times D}$ be the native-observed mask, where $M^{\mathrm{nat}}_{t,d} = 1$ exactly when $X_{t,d}$ is available before controlled evaluation masking. To evaluate complete-channel recovery, we choose one channel $d^\star$ and define the evaluation-missing mask

$$M^{\mathrm{eval}}_{t,d} := \mathbb{I}[d = d^\star], \tag{1}$$

so one entire channel is synthetically held out in each test window. The selected channel is balanced across the test set. We then define

$$\begin{aligned} M^{\mathrm{fit}} &:= M^{\mathrm{nat}} \odot (1 - M^{\mathrm{eval}}), \\ O &:= 1 - M^{\mathrm{fit}}, \\ H &:= M^{\mathrm{nat}} \odot M^{\mathrm{eval}}. \end{aligned} \tag{2}$$

Here M^{fit} marks entries available to the observed-residual objective, O marks coordinates optimized by TTO, and H marks artificial holdouts that have reference values and can therefore be scored. Thus O includes both native missing entries and artificial holdouts, whereas H excludes native missing entries. We refer to M^{nat}, M^{eval}, M^{fit}, and O as the native-observed, evaluation-missing, fit-residual, and optimization masks, respectively.

Conversion contract. The inference-time adapter receives X, M^{fit}, O, and a frozen reconstructor f_θ. It keeps θ fixed, clamps entries selected by M^{fit}, and changes only candidate values on O. If $\bar{X} = F_\theta(X, M^{\mathrm{fit}})$ denotes the reconstructor prediction after optimization, the completed window is

$$\hat{X} := M^{\mathrm{fit}} \odot X + O \odot \bar{X}. \tag{3}$$

Consequently, every value available to the adapter remains unchanged in the completed output.

Evaluator-only targets. For controlled evaluation, $\tilde{X}$ denotes reference values on the artificial holdouts H. Neither $\tilde{X}$ nor H is supplied to TTO as the adapter sees only the combined availability/optimization masks M^{fit} and O. Native missing entries are not scored. The reported mean absolute error is

$$\mathcal{E}_{\mathrm{MAE}}(\bar{X}; \tilde{X}, H) := \frac{\|H \odot (\bar{X} - \tilde{X})\|_1}{\|H\|_1}. \tag{4}$$

At deployment, M^{fit} can be supplied directly from the observed/missing pattern, the decomposition into M^{nat}, M^{eval}, and H is evaluation bookkeeping only.

4 Method

As illustrated in Fig. 2, the method converts a trained sequence reconstructor into an imputer without updating its parameters or accessing held-out values. It consists of two phases: offline reconstruction training supplies a learned sequence prior, after which post-training test-time optimization refines only the missing entries by enforcing observed self-consistency. The following subsections specify these phases and the resulting optimization procedure.

4.1 Phase 1: Offline Reconstruction Training

We train a standard autoencoder as a sequence reconstructor by reconstructing observed entries of baseline-filled windows. Let (X, M^{nat}) denote a training window and its native-observed mask (Sec. 3), and let $B(\cdot)$ be a baseline fill operator (e.g., zero or per-channel mean). The reconstructor is an encoder–decoder trained from

$$\begin{aligned} X^{\mathrm{in}} &:= M^{\mathrm{nat}} \odot X + (1 - M^{\mathrm{nat}}) \odot B(X), \\ X^{\mathrm{rec}} &:= f_\theta(X^{\mathrm{in}}). \end{aligned} \tag{5}$$

by minimizing the reconstruction loss on native-observed entries,

$$\min_\theta \ \mathbb{E}_{(X, M^{\mathrm{nat}})} \Big[\, \| M^{\mathrm{nat}} \odot (f_\theta(X^{\mathrm{in}}) - X) \|_F^2 \, \Big]. \tag{6}$$

The mask M^{nat} affects baseline filling and loss masking but is not provided as a separate network input, and training uses no artificial masking/denoising objective. After training, the reconstructor is frozen so the conversion stage below changes only missing input coordinates and does not retrain θ.

4.2 Phase 2: Post-Training Conversion by Input Optimization

At test time, we use the shared mask contract defined in section 3. The adapter receives only X, the fit-residual mask M^{fit}, and the optimization mask O. It does not receive the evaluator-only target mask H or the reference values $\tilde{X}$. We freeze θ and optimize only coordinates selected by O. We introduce an explicit optimization variable $V \in \mathbb{R}^{T \times D}$ storing candidate values at optimized coordinates and define the completed input

$$X_V := M^{\mathrm{fit}} \odot X + O \odot V, \qquad \bar{X}_V := f_\theta(X_V). \tag{7}$$

By construction, X_V agrees with X on observed indices, thus $f_\theta(X_V)$ already provides a one-shot completion, which we further refine by optimizing V.

From reconstructor to imputation model. Under the conversion contract in Sec. 3, our imputation model composes a learned reconstructor with a test-time optimization over missing entries. We choose V by enforcing *observed self-consistency* via the residual $R(V) := M^{\mathrm{fit}} \odot (\bar{X}_V - X)$ and solve

$$\begin{aligned} V^* &\in \arg\min_{V \in \mathcal{Z}} \ \mathcal{J}(V), \ \text{with} \\ \mathcal{J}(V) &:= \tfrac{1}{2} \|R(V)\|_F^2 + \lambda \, \mathcal{P}(V), \end{aligned} \tag{8}$$

where $\mathcal{Z}$ encodes optional constraints (e.g., value bounds), $\mathcal{P}(V)$ is an optional prior on missing trajectories (e.g., temporal smoothness or a proximal term), and $\lambda \geq 0$ controls the prior strength (specified in Sec. 5). An informative $\mathcal{P}(V)$ can further improve imputation when domain knowledge is available, but the method remains well-defined without it (i.e., $\lambda = 0$). Crucially, the data term depends only on M^{fit} to ensure neither evaluator-only reference values on H nor baseline-filled native missing entries provide supervision. The resulting imputation output is

$$\bar{X} = F_\theta(X, M^{\mathrm{fit}}) := \bar{X}_{V^*} = f_\theta(X_{V^*}), \tag{9}$$

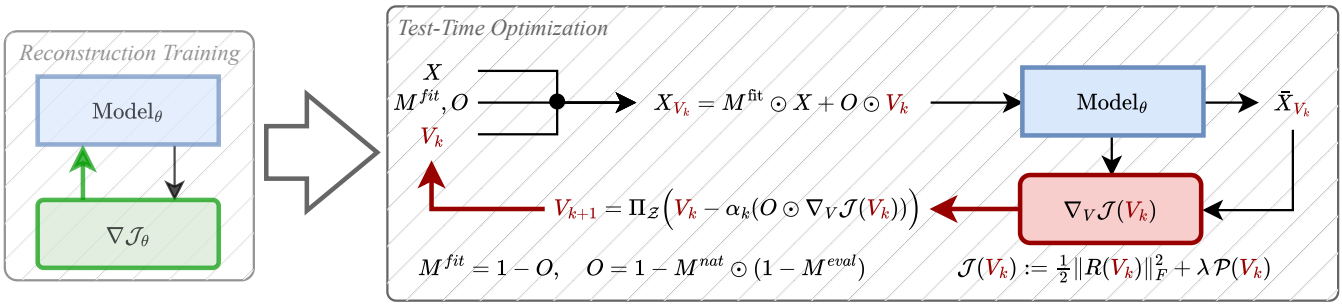

Figure 2: **Converting a Frozen Reconstructor into an Imputer.** In the prior/offline phase, an autoencoder learns to reconstruct observed entries from baseline-filled windows; the loss is evaluated only where training data are natively observed. In the proposed post-training conversion, the reconstructor is frozen and candidate missing values V are optimized through it to minimize the fit-observed reconstruction residual (optionally regularized). Observed inputs remain clamped, and gradients update only optimized entries.

and the completed series is $\hat{X} := M^{\text{fit}} \odot X + O \odot \bar{X}$ as in Eq. (3). We return the reconstructor output $f_\theta(X_{V^*})$, rather than the auxiliary variable V^* itself, because V^* is optimized only to elicit an observed-self-consistent response from the frozen model. The output is the model-consistent completion learned by the reconstruction objective and keeps forward-only and optimized inference comparable. Directly returning V^* would bypass the learned reconstruction mapping. We next discuss initialization and when the resulting refinement problem is locally well-posed.

Initialization. We initialize the optimization variable with a baseline completion,

$$V_0 := B(X), \tag{10}$$

which yields the initial completed input X_{V_0} and prediction $\bar{X}_{V_0}$. In our experiments, we use per-channel mean of the training set.

Well-posedness of the refinement problem. The refinement step is informative only if changing missing entries affects the model's predictions on observed entries, i.e., if the observed residual $R(V) = M^{\text{fit}} \odot (f_\theta(X_V) - X)$ depends on V through the optimized coordinates. Let $\Omega = \{(t,d) : M^{\text{fit}}_{t,d} = 1\}$ be the fit-observed set and $\bar{\Omega} = \{(t,d) : O_{t,d} = 1\}$ the optimized set. Let $r(V) \in \mathbb{R}^{|\Omega|}$ stack $R(V)$ on Ω, and let $V_{\bar{\Omega}} \in \mathbb{R}^{|\bar{\Omega}|}$ collect optimized entries. We define the *cross-Jacobian block*

$$C(V) := \frac{\partial r(V)}{\partial V_{\bar{\Omega}}} \in \mathbb{R}^{|\Omega| \times |\bar{\Omega}|}, \tag{11}$$

i.e., the sensitivity of the observed residual to perturbations of the missing entries. Note that $R(V)$ depends on V only through optimized coordinates because X_V clamps entries in M^{fit}.

This coupling condition is necessary for an informative refinement signal. For example, if f_θ is the identity map, its observed outputs depend only on the clamped observed inputs. Consequently, if $C(V) = 0$ the data-term gradient provides no information about the missing coordinates.

Under the differentiability and exact-self-consistency assumptions in Appendix A, if $C(V^\star)$ has full column rank at a point $V^\star$ with $R(V^\star) = 0$, then $V^\star$ is isolated among nearby self-consistent completions and the local Gauss–Newton curvature is governed by $C(V^\star)^\top C(V^\star)$. This is a local statement around an exact solution and does not establish global identifiability or rule out distant alternative completions. If $C(V^\star)$ is rank-deficient or ill-conditioned, multiple nearby completions can explain the observed entries equally well, making the refinement sensitive. The optional prior $\lambda \mathcal{P}(V)$ regularizes directions that observed self-consistency leaves weakly constrained. It can stabilize or select among such solutions, but it is not the source of the basic conversion (Appendix A).

4.3 Optimization algorithm

We minimize Eq. (8) by projected first-order methods, backpropagating through the frozen reconstructor f_θ. Let $G_k := \nabla_V \mathcal{J}(V_k)$ computed by automatic differentiation. In the first step we update only missing coordinates by using a defined step size α_k and in the second step we (optionally) use a projection $\Pi_{\mathcal{Z}}$ to $\mathcal{Z}$:

$$V_{k+\frac{1}{2}} = V_k - \alpha_k\big(O \odot G_k\big), \tag{12}$$

$$V_{k+1} = \Pi_{\mathcal{Z}}(V_{k+\frac{1}{2}}), \tag{13}$$

with early stopping based on the observed residual $\|R(V_k)\|_F$ or relative objective decrease. Algorithm 1 summarizes the resulting test-time optimization procedure.

In our experiments, the optional $\mathcal{P}(V)$ is instantiated as a first-order temporal smoothness penalty on missing indices. Full details and stopping criteria are given in Appendix B.

Runtime. Test-time optimization requires multiple forward/backward passes per window, which we implement in a batched, fully vectorized manner. Implementation details are provided in Appendix B and the full runtime comparison is stated in Appendix C.8.

5 Experiments

Our main goal is to evaluate the effectiveness and trade-offs of *test-time optimization* (Sec. 4.2) in *hard missingness*

Algorithm 1 **T**est-**T**ime **O**ptimization (TTO) of missing values (frozen reconstructor)

Require: input window X, masks M^{fit}, O, frozen reconstructor f_θ, prior weight λ, steps K

Initialize $V \leftarrow B(X)$ ▷ baseline fill for missing entries
for $k = 1, \dots, K$ **do**
  $X_V \leftarrow M^{\text{fit}} \odot X + O \odot V$ ▷ clamp fit entries
  $\bar{X}_V \leftarrow f_\theta(X_V)$
  $R(V) \leftarrow M^{\text{fit}} \odot (\bar{X}_V - X)$ ▷ observed residual
  $\mathcal{J} \leftarrow \frac{1}{2}\|R(V)\|_F^2 + \lambda \mathcal{P}(V)$
  $G \leftarrow \nabla_V \mathcal{J}$ ▷ backprop through f_θ; weights frozen
  $V \leftarrow \Pi_{\mathcal{Z}}\big(\textsc{OptStep}(V, O \odot G)\big)$ ▷ e.g., Adam; optional constraints
  if $\textsc{Converged}(\mathcal{J})$ **then**
    break
  end if
end for
return $\hat{X} \leftarrow M^{\text{fit}} \odot X + O \odot f_\theta(M^{\text{fit}} \odot X + O \odot V)$

regimes that stress cross-coordinate coupling. In particular, we use a synthetic evaluation mask that removes exactly one complete channel from each test window. Throughout, we report $\mathcal{E}_{\text{MAE}}(\bar{X}; \tilde{X}, H)$ from Eq. (4). The scoring mask H contains only natively observed entries in the synthetically held-out channel, so native missing values never enter the reported errors.

5.1 Evaluation Questions

We organize the evaluation around three questions addressing effectiveness, comparison, and optimization behavior and cost: **(EQ1)** Under an unseen full-channel mask, does TTO improve a frozen reconstructor over forward-only completion? **(EQ2)** How does the converted reconstructor compare with mask-aware imputers trained using either standard structured masks or matched full-channel masks? **(EQ3)** What optimization behavior and inference-time cost does conversion exhibit?

5.2 Datasets and preprocessing

We evaluate on five multivariate time series from different domains: *BeijingAir* (air quality), *ItalyAir* (air quality), *ETTh1* (energy/temperature), *Electricity* (electricity load), and *PeMS* (traffic). We follow the standard time-based splitting and windowing protocol used in TSI-Bench (Du et al. 2024) to avoid leakage. All datasets are standardized using statistics from the training split. For full details, see Du et al. (2024). In our experiments only the *BeijingAir* Dataset has a low rate of native missingness of 1.6%.

5.3 Missingness patterns and evaluation settings

Structured missingness for deep imputation baselines. Following the TSI-Bench protocol (Du et al. 2024), we evaluate deep imputation baselines under three structured missingness patterns: *point*, *subsequence*, and *block* missingness, each at a 10% missing rate. For these baselines, we adopt the standard practice of training and evaluating under the same pattern family (point/subsequence/block) to match the intended operating regime of mask-aware imputers.

Full-channel missing evaluation. We use the mask construction in Sec. 3. Thus, for each test window, one channel $d^\star$ is selected and $M^{\text{eval}}_{t,d} = \mathbb{I}[d = d^\star]$ for every time step. The fit-residual and optimization masks then follow from Eq. (2). We assign $d^\star$ cyclically so each channel is held out equally often across the test set.

Training-mask regimes for baseline comparison. Table 1 summarizes mask-aware baselines with two per-dataset lower envelopes: *Best OOD mask-aware* over checkpoints trained with 10% point, subsequence, or block corruption and evaluated under full-channel corruption, and *Best matched mask-aware* over checkpoints trained with matched full-channel dropout. Appendix C.4 reports every constituent checkpoint with its *point*, *subseq*, *block*, or *full feature* label. In contrast, we train and tune the reconstruction architectures for this study using only native missingness, without full-channel masks or additional artificial masking. We then freeze their weights and apply TTO in Eq. (8). The comparison therefore has a deliberate supervision boundary: mask-aware models receive explicit corruption training, whereas reconstructors do not.

5.4 Compared methods.

We compare against strong baselines covering major architecture families: BRITS (RNN) (Cao et al. 2018), SAITS (masked self-attention) (Du, Côté, and Liu 2023), iTransformer (Transformer) (Liu et al. 2023b), TimesNet (CNN/FFT-style) (Wu et al. 2022), and FreTS (MLP/spectral) (Yi et al. 2023). Several of the considered backbones were originally developed for forecasting. We use them here as *representative architectures* by following the standard adaptation procedure in the imputation literature (notably as operationalized by Du, Côté, and Liu (2023)). Model configurations for all baselines follow the fixed hyperparameter settings released with TSI-Bench (Du et al. 2024). As a non-neural sanity check, we additionally report MeanFill, which replaces missing entries by per-channel means from the corresponding training split.

5.5 Our models and hyperparameters.

Our imputer uses the reconstructor plus test-time optimization pipeline from Sec. 4. We instantiate f_θ as a standard reconstruction autoencoder, using different backbones (LSTM, MTAD-GAT (Zhao et al. 2020), TCN-AE (Thill et al. 2021), TransformerAE) that represent architectures from RNNs, CNNs, GNNs, and transformers. We define the baseline fill $B(X)$ as per-channel mean of the training set. We tune our model hyperparameters with Optuna on the validation split using a dataset-agnostic search space, see Appendix C.2.

5.6 Test-time optimization details.

Test-time input optimization runs for at most $K = 5000$ iterations with per-sample early stopping. We use Adam with learning rate 10^{-2}. Early stopping uses a bootstrap window of 100 iterations. Letting $\bar{\mathcal{L}}_{\text{init}}$ be the mean objective over this

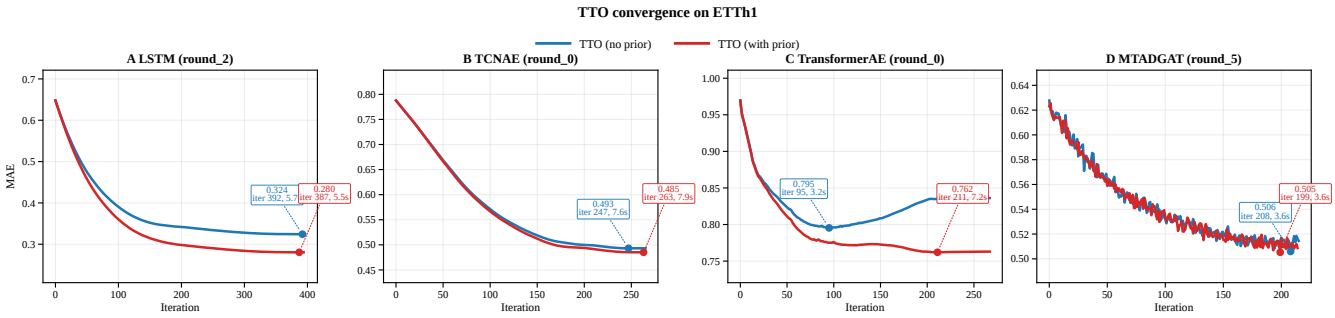

Figure 3: **Representative TTO convergence on ETTh1.** Imputation MAE over optimization iterations for four frozen reconstruction backbones, with (red) and without (blue) the smoothness prior. The trajectories are smooth in this representative case and stabilize within a small fraction of the conservative $K = 5000$ cap, using objective-based early stopping.

window, we set an improvement threshold $\Delta_{\min} = 0.01\,\bar{\mathcal{L}}_{\text{init}}$ and a plateau threshold $\tau = 10^{-4}\,\bar{\mathcal{L}}_{\text{init}}$. A sample is marked converged if it fails to improve by $\Delta_{\min}$ for 100 consecutive iterations or if the per-iteration objective change stays below τ for 200 consecutive iterations. The MAE values in Table 1 are computed from the output selected by this objective-based stopping rule, not from an oracle choice of the iteration with the best test MAE. We report three variants for our method: a pure forward pass without refinement (no optimization), test-time optimization with λ=0, and test-time optimization with a simple first-order temporal smoothness regularizer $\mathcal{P}(V)$ on missing indices (λ=10^{-3}). Reconstructor hyperparameters are chosen on validation data; the forward, no-prior, and prior variants are then all reported on the test set. The per-dataset *best* rows used for the family comparison are post-hoc descriptive lower envelopes and never select the prediction for an individual test window. Implementation and stopping details are summarized in Appendix B.

5.7 Compute and reproducibility.

Experiments were run in a containerized environment limited to 12 vCPUs and 48 GiB RAM with access to an NVIDIA A30 MIG instance (16 GiB). We report mean $\pm$ standard deviation across 6 random seeds.

5.8 Results

EQ1: Does conversion improve forward-only completion? Panel A of Table 1 makes the conversion ablation the primary result by comparing each frozen reconstructor with the same model after TTO. No-prior TTO ($\lambda = 0$) lowers MAE in 19 of 20 backbone–dataset combinations. The sole exception is TCN-AE on Electricity, where MAE shows a 0.8% relative increase. Using the reported means, averaging the pairwise relative MAE reduction across the four backbones gives 14.6% on BeijingAir, 29.1% on ETTh1, 0.4% on Electricity, 19.7% on ItalyAir, and 0.5% on PeMS, with a 12.9% macro-average over all 20 pairs.

The smoothness prior is a separate, selective effect: relative to no-prior TTO, it lowers reported MAE in 11 of 20 combinations, leaves 2 unchanged at the shown precision, and raises MAE in 7. Thus, the 19/20 conversion result does not depend on uniform gains from the optional prior.

EQ2: How does conversion compare with mask-aware imputers? Panel B reports the lower mean MAE within each method family and dataset. Each cell in a *best* row is the post-hoc minimum over the fully reported rows in that family: five matched full-channel-mask baselines, fifteen OOD point/subsequence/block checkpoints, or the eight refined reconstructor variants in Panel A, all individual baseline rows are traceable in Appendix C.4. Taking the refined-reconstructor envelope, TTO yields lower MAE than both mask-aware envelopes on ETTh1 and Electricity, but higher MAE than the strongest mask-aware result on BeijingAir, ItalyAir, and PeMS. Electricity requires a precise distinction: forward TCN-AE is the best reconstructor row, while the best *refined* row is TCN-AE with the prior. That refined result still beats the OOD and matched-mask envelopes. Matched full-channel training substantially improves some baselines, but not uniformly. This comparison should retain the supervision boundary from Sec. 5.3. The mask-aware models receive artificial corruption during training, while the reconstructors are trained only on native missingness.

Mechanism diagnostics. The cross-Jacobian analysis in Sec. 4.2 establishes the local coupling condition behind observed-self-consistent refinement. Appendix C.6 reports illustrative diagnostics for three backbones across all five datasets, including stronger and weakly coupled cases.

EQ3: What optimization behavior and inference-time cost does conversion exhibit? Figure 3 shows representative ETTh1 trajectories for all four reconstruction backbones. In this case, imputation MAE evolves smoothly rather than oscillating strongly, and the effect of the smoothness prior is most visible for LSTM and TransformerAE.

The hard cap $K = 5000$ is therefore not the typical executed budget. For the fastest TTO rows used in the runtime comparison, median stopping is 201-206 iterations and the 90th percentile is 201-209. Although early stopping limits the realized iteration count, roughly 200 forward/backward updates per window still impose substantial computational overhead.

Using the per-dataset lower-envelope comparison defined in App. C.8, TTO is 6.7—$174.7\times$ slower than one-shot inference on identical hardware. This comparison uses the fastest

	Setting	BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
Panel A: Conversion ablation for frozen reconstructors						
LSTM	forward	0.4968 ± 0.0145	0.6488 ± 0.0024	1.1901 ± 0.0389	0.5845 ± 0.0586	0.2944 ± 0.0064
	TTO, λ=0	0.3413 ± 0.0176	0.3286 ± 0.0216	1.1821 ± 0.0402	0.3806 ± 0.0365	0.2913 ± 0.0060
	TTO + prior	0.3099 ± 0.0169	0.2898 ± 0.0416	1.1821 ± 0.0402	0.3796 ± 0.0365	0.2917 ± 0.0055
MTAD-GAT	forward	0.2776 ± 0.0054	0.6227 ± 0.0203	1.5608 ± 0.0982	0.4942 ± 0.0128	0.3393 ± 0.0064
	TTO, λ=0	0.2726 ± 0.0042	0.5320 ± 0.0321	1.5580 ± 0.0990	0.4815 ± 0.0357	0.3378 ± 0.0062
	TTO + prior	0.2728 ± 0.0051	0.5323 ± 0.0311	1.5580 ± 0.0990	0.4822 ± 0.0336	0.3377 ± 0.0061
TCN-AE	forward	0.2455 ± 0.0039	0.6501 ± 0.0676	0.9853 ± 0.0287	0.7200 ± 0.0241	0.2898 ± 0.0030
	TTO, λ=0	0.2289 ± 0.0059	0.4299 ± 0.0347	0.9928 ± 0.0230	0.6557 ± 0.0268	0.2892 ± 0.0026
	TTO + prior	0.2203 ± 0.0067	0.4329 ± 0.0307	0.9926 ± 0.0229	0.6404 ± 0.0241	0.2884 ± 0.0028
TransformerAE	forward	0.2763 ± 0.0041	0.7610 ± 0.1168	1.0525 ± 0.0167	0.6440 ± 0.0399	0.2822 ± 0.0033
	TTO, λ=0	0.2255 ± 0.0043	0.6197 ± 0.1172	1.0354 ± 0.0220	0.4350 ± 0.0615	0.2811 ± 0.0051
	TTO + prior	0.2200 ± 0.0034	0.5543 ± 0.1091	1.0353 ± 0.0221	0.4353 ± 0.0619	0.2816 ± 0.0055
Panel B: Descriptive lower envelopes for external comparison						
MeanFill	train mean	0.7718	0.8043	1.9657	0.7974	0.7742
Best matched mask-aware	per-dataset envelope	**0.1383 ± 0.0028**	0.3222 ± 0.0261	1.3480 ± 0.0879	0.3588 ± 0.0244	0.2508 ± 0.0037
Best OOD mask-aware	per-dataset envelope	0.2108 ± 0.0018	0.3180 ± 0.0111	1.1452 ± 0.0348	**0.3488 ± 0.0087**	**0.2470 ± 0.0016**
Best refined reconstructor	per-dataset envelope	0.2200 ± 0.0034	**0.2898 ± 0.0416**	**0.9926 ± 0.0229**	0.3796 ± 0.0365	0.2811 ± 0.0051

Table 1: **Full-channel-missing evaluation (MAE; mean ± std where applicable).** Panel A asks whether post-training conversion improves each frozen reconstructor. Panel B compares post-hoc per-dataset lower envelopes for matched-mask baselines, OOD structured-mask checkpoints, and refined reconstructors. Bold and underline mark the best and second-best Panel-B values. Every individual baseline checkpoint is reported in Appendix C.4.

measured configuration from each family and provides an estimate of TTO's computational overhead. Complete per-dataset results are reported in Table 11, while Figs. 6-10 report convergence, active-sample counts, and runtime.

6 Conclusion and outlook

A frozen differentiable sequence reconstructor can often be converted into a point-estimate imputer without retraining its weights. Its learned temporal and cross-channel dependencies must make observed self-consistency informative about the missing inputs. TTO performs this conversion by optimizing only the missing entries of a test window through the frozen model while keeping observations clamped. We evaluate this claim in a controlled setting with five real-world datasets, four reconstruction-model families trained using only native missingness, and one complete channel synthetically masked in each test window. The experiments therefore demonstrate post-training conversion of these reconstructors.

Empirically, no-prior TTO lowers MAE relative to forward-only reconstruction in 19 of 20 backbone-dataset combinations, with a 12.9% macro-average relative reduction (Table 1). The exception is TCN-AE on Electricity, whose forward result is better than either refined variant. The optional smoothness prior is selectively beneficial as it improves 11 of 20 rows, is unchanged in 2 at reported precision, and worsens 7. Taking the best TTO-refined reconstructor per dataset, TTO yields lower MAE than the strongest mask-aware baseline on ETTh1 and Electricity, but higher MAE on BeijingAir, ItalyAir, and PeMS. This conditional accuracy benefit comes at cost of computation as stated in Sec. 5.6 and App. C.8.

Limitations. The conversion fails at its mechanism boundary. For an identity-like reconstructor whose observed outputs do not depend on missing inputs, the missing-to-observed cross-Jacobian is zero and the observed residual supplies no recovery gradient. More generally, weak, ill-conditioned, or rank-deficient coupling leaves directions ambiguous, so multiple completions can be similarly self-consistent. Priors can regularize but cannot remove this ambiguity. Our Jacobian diagnostics provide mechanism-consistent evidence for three of the four backbones. TTO also introduces iterative compute and implementation choices such as optimizer, normalization, and stopping criteria, and the evaluation does not establish performance for arbitrary reconstructors.

Outlook. Future work should evaluate broader hard-missingness regimes, including multiple missing channels, mixed block+channel dropouts, and longer horizons. Stronger domain-aware priors, uncertainty diagnostics based on residual curvature or cross-Jacobian spectra, and faster quasi-Newton or Gauss-Newton-style solvers could make TTO more predictable and scalable. Given its dependence on reconstructor coupling and its iterative cost, TTO is best understood as a conditional inference-time adapter that complements, rather than replaces, amortized mask-aware imputers.

Impact This work studies test-time conversion of compatible frozen reconstructors into point-estimate imputers. Although reconstruction objectives also appear in anomaly detection, we do not evaluate deployed-model reuse, anomaly-detection performance, or maintenance savings. For sensor dropouts, imputed values used in downstream monitoring or forecasting should be flagged rather than treated as ground truth.

References

Akiba, T.; Sano, S.; Yanase, T.; Ohta, T.; and Koyama, M. 2019. Optuna: A Next-generation Hyperparameter Optimization Framework. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, KDD '19, 2623–2631. New York, NY, USA: Association for Computing Machinery. ISBN 9781450362016.

Cao, W.; Wang, D.; Li, J.; Zhou, H.; Li, Y.; and Li, L. 2018. BRITS: bidirectional recurrent imputation for time series. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, NIPS'18, 6776–6786. Red Hook, NY, USA: Curran Associates Inc.

Che, Z.; Purushotham, S.; Cho, K.; Sontag, D.; and Liu, Y. 2018. Recurrent Neural Networks for Multivariate Time Series with Missing Values. *Scientific Reports*, 8(1): 6085.

Cini, A.; Marisca, I.; and Alippi, C. 2021. Filling the G_ap_s: Multivariate Time Series Imputation by Graph Neural Networks. In *International Conference on Learning Representations*.

De Vito, S.; Massera, E.; Piga, M.; Martinotto, L.; and Di Francia, G. 2008. On field calibration of an electronic nose for benzene estimation in an urban pollution monitoring scenario. *Sensors and Actuators B: Chemical*, 129(2): 750–757.

Du, W.; Côté, D.; and Liu, Y. 2023. SAITS: Self-attention-based imputation for time series. *Expert Systems with Applications*, 219: 119619.

Du, W.; Wang, J.; Qian, L.; Yang, Y.; Liu, F.; Wang, Z.; Ibrahim, Z.; Liu, H.; Zhao, Z.; Zhou, Y.; Wang, W.; Ding, K.; Liang, Y.; Prakash, B. A.; and Wen, Q. 2024. TSI-Bench: Benchmarking Time Series Imputation. *arXiv preprint arXiv:2406.12747*.

Fortuin, V.; Baranchuk, D.; Raetsch, G.; and Mandt, S. 2020. GP-VAE: Deep Probabilistic Time Series Imputation. In Chiappa, S.; and Calandra, R., eds., *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, volume 108 of *Proceedings of Machine Learning Research*, 1651–1661. PMLR.

Haim, N.; Vardi, G.; Yehudai, G.; Shamir, O.; and Irani, M. 2022. Reconstructing training data from trained neural networks. In *Proceedings of the 36th International Conference on Neural Information Processing Systems*, NIPS '22. Red Hook, NY, USA: Curran Associates Inc. ISBN 9781713871088.

Jarrett, D.; Cebere, B. C.; Liu, T.; Curth, A.; and van der Schaar, M. 2022. HyperImpute: Generalized Iterative Imputation with Automatic Model Selection. In Chaudhuri, K.; Jegelka, S.; Song, L.; Szepesvari, C.; Niu, G.; and Sabato, S., eds., *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, 9916–9937. PMLR.

Khayati, M.; Lerner, A.; Tymchenko, Z.; and Cudré-Mauroux, P. 2020. Mind the gap: an experimental evaluation of imputation of missing values techniques in time series. *Proc. VLDB Endow.*, 13(5): 768–782.

Liu, X.; Xia, Y.; Liang, Y.; Hu, J.; Wang, Y.; Bai, L.; Huang, C.; Liu, Z.; Hooi, B.; and Zimmermann, R. 2023a. LargeST: a benchmark dataset for large-scale traffic forecasting. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS '23. Red Hook, NY, USA: Curran Associates Inc.

Liu, Y.; Hu, T.; Zhang, H.; Wu, H.; Wang, S.; Ma, L.; and Long, M. 2023b. iTransformer: Inverted Transformers Are Effective for Time Series Forecasting. *arXiv preprint arXiv:2310.06625*.

Mattei, P.-A.; and Frellsen, J. 2019. MIWAE: Deep Generative Modelling and Imputation of Incomplete Data Sets. In *International Conference on Machine Learning*.

Mitra, R.; McGough, S. F.; Chakraborti, T.; Holmes, C.; Copping, R.; Hagenbuch, N.; Biedermann, S.; Noonan, J.; Lehmann, B.; Shenvi, A.; Doan, X. V.; Leslie, D.; Bianconi, G.; Sanchez-Garcia, R.; Davies, A.; Mackintosh, M.; Andrinopoulou, E.-R.; Basiri, A.; Harbron, C.; and MacArthur, B. D. 2023. Learning from data with structured missingness. *Nature Machine Intelligence*, 5(1): 13–23.

Tashiro, Y.; Song, J.; Song, Y.; and Ermon, S. 2021. CSDI: conditional score-based diffusion models for probabilistic time series imputation. In *Proceedings of the 35th International Conference on Neural Information Processing Systems*, NIPS '21. Red Hook, NY, USA: Curran Associates Inc. ISBN 9781713845393.

Thill, M.; Konen, W.; Wang, H.; and Bäck, T. 2021. Temporal convolutional autoencoder for unsupervised anomaly detection in time series. *Applied Soft Computing*, 112: 107751.

Trindade, A. 2015. ElectricityLoadDiagrams20112014. UCI Machine Learning Repository. DOI: https://doi.org/10.24432/C58C86.

Wang, J.; Du, W.; Yang, Y.; Qian, L.; Cao, W.; Zhang, K.; Wang, W.; Liang, Y.; and Wen, Q. 2025. Deep learning for multivariate time series imputation: a survey. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence*, IJCAI '25. ISBN 978-1-956792-06-5.

Wen, Q.; Zhou, T.; Zhang, C.; Chen, W.; Ma, Z.; Yan, J.; and Sun, L. 2023. Transformers in time series: a survey. In *Proceedings of the Thirty-Second International Joint Conference on Artificial Intelligence*, IJCAI '23. ISBN 978-1-956792-03-4.

Wu, G.; Say, B.; and Sanner, S. 2020. Scalable planning with deep neural network learned transition models. *Journal of Artificial Intelligence Research*, 68: 571–606.

Wu, H.; Hu, T.; Liu, Y.; Zhou, H.; Wang, J.-M.; and Long, M.-S. 2022. TimesNet: Temporal 2D-Variation Modeling for General Time Series Analysis. *ArXiv*, abs/2210.02186.

Yi, K.; Zhang, Q.; Fan, W.; Wang, S.; Wang, P.; He, H.; Lian, D.; An, N.; Cao, L.; and Niu, Z. 2023. Frequency-domain MLPs are more effective learners in time series forecasting. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS '23. Red Hook, NY, USA: Curran Associates Inc.

You, J.; Ma, X.; Ding, D. Y.; Kochenderfer, M.; and Leskovec, J. 2020. Handling missing data with graph representation learning. In *Proceedings of the 34th International*

Conference on Neural Information Processing Systems, NIPS '20. Red Hook, NY, USA: Curran Associates Inc. ISBN 9781713829546.

Zhang, S.; Guo, B.; Dong, A.; He, J.; Xu, Z.; and Chen, S.-X. 2017. Cautionary tales on air-quality improvement in Beijing. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 473.

Zhao, H.; Sun, K.; Dezfouli, A.; and Bonilla, E. V. 2023. Transformed distribution matching for missing value imputation. In *Proceedings of the 40th International Conference on Machine Learning*, ICML'23. JMLR.org.

Zhao, H.; Wang, Y.; Duan, J.; Huang, C.; Cao, D.; Tong, Y.; Xu, B.; Bai, J.; Tong, J.; and Zhang, Q. 2020. Multivariate Time-Series Anomaly Detection via Graph Attention Network. In *2020 IEEE International Conference on Data Mining (ICDM)*, 841–850.

Zhong, J.; Gui, N.; and Ye, W. 2023. Data imputation with iterative graph reconstruction. In *Proceedings of the Thirty-Seventh AAAI Conference on Artificial Intelligence and Thirty-Fifth Conference on Innovative Applications of Artificial Intelligence and Thirteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI'23/IAAI'23/EAAI'23. AAAI Press. ISBN 978-1-57735-880-0.

A Cross-Jacobian coupling and local well-posedness

This appendix supports the discussion in Sec. 4.2 by (i) making the dependence on the missing coordinates explicit, (ii) deriving the cross-Jacobian identity in (11), and (iii) stating a simple local well-posedness implication.

A.1 Vectorization and index operators

Let $M := M^{\text{fit}}$ be the fit-residual mask from Sec. 3, let $\Omega := \{(t, d) \mid M_{t,d} = 1\}$ be the fit-observed index set, and let $\bar{\Omega}$ denote the optimized index set. Let $\text{vec} : \mathbb{R}^{T\times D} \to \mathbb{R}^{TD}$ denote any fixed vectorization convention.

We use the standard linear maps induced by Ω:

- A *selection* matrix $S_\Omega \in \{0, 1\}^{|\Omega|\times TD}$ such that $S_\Omega \,\text{vec}(A)$ stacks the entries of A indexed by Ω.
- An *embedding* matrix $E_{\bar{\Omega}} \in \{0, 1\}^{TD\times|\bar{\Omega}|}$ such that $E_{\bar{\Omega}} v$ places $v \in \mathbb{R}^{|\bar{\Omega}|}$ in the coordinates indexed by $\bar{\Omega}$ (and zeros elsewhere).

By construction, $S_\Omega E_{\bar{\Omega}} = 0$.

Define the missing-coordinate optimization vector

$$v \;:=\; V_{\bar{\Omega}} \in \mathbb{R}^{|\bar{\Omega}|}. \tag{14}$$

Then clamping observed entries and inserting missing values (Eq. (7)) can be written as

$$\text{vec}(X_V) \;=\; \text{vec}(M \odot X) \;+\; E_{\bar{\Omega}} v. \tag{15}$$

Eq. (15) formalizes the statement in Sec. 4.2 that X_V agrees with X on Ω while the coordinates in $\bar{\Omega}$ are free.

A.2 Observed residual and cross-Jacobian identity

Recall the masked residual and its observed stacking from Sec. 4.2:

$$R(V) \;:=\; M \odot (f_\theta(X_V) - X), \qquad r(V) \;:=\; S_\Omega \,\text{vec}(R(V)) \in \mathbb{R}^{|\Omega|}. \tag{16}$$

Since S_Ω selects only observed entries, the elementwise mask is redundant after selection:

$$r(V) = S_\Omega \,\text{vec}\big(f_\theta(X_V) - X\big) = S_\Omega \,\text{vec}\big(f_\theta(X_V)\big) \;-\; S_\Omega \,\text{vec}(X). \tag{17}$$

Let $J_\theta(X_V)$ denote the Jacobian of the reconstructor output w.r.t. the vectorized input at X_V:

$$J_\theta(X_V) \;:=\; \frac{\partial \,\text{vec}(f_\theta(U))}{\partial \,\text{vec}(U)}\Big|_{U=X_V} \in \mathbb{R}^{TD\times TD}. \tag{18}$$

Applying the chain rule to (17) and the affine map (15) yields

$$C(V) \;:=\; \frac{\partial r(V)}{\partial v} \;=\; S_\Omega \, J_\theta(X_V) \, E_{\bar{\Omega}} \;\in\; \mathbb{R}^{|\Omega|\times|\bar{\Omega}|}. \tag{19}$$

Eq. (19) is the operator form of the cross-Jacobian block in (11).

A.3 Local isolation of self-consistent completions

Consider the noise-free self-consistency equation on observed entries,

$$r(V) = 0, \tag{20}$$

and let $v^\star := V^\star_{\bar{\Omega}}$ be a solution (i.e., $r(V^\star) = 0$). Assume f_θ is differentiable in a neighborhood of $X_{V^\star}$.

Proposition A.1 (Local isolation under full-column rank)**.** *If $C(V^\star)$ in* (19) *has full column rank, then $v^\star$ is an* isolated *solution of* (20)*: there exists a neighborhood $\mathcal{U}$ of $v^\star$ such that $r(V(v)) = 0$ and $v \in \mathcal{U}$ imply $v = v^\star$.*

Proof sketch. Work in the reduced variable v via (15) and define $\phi(v) := r(V(v))$. Full column rank of $C(V^\star) = \partial\phi/\partial v|_{v^\star}$ means ϕ is an immersion at $v^\star$. By the constant-rank theorem, ϕ is locally an embedding around $v^\star$, hence locally injective. Therefore $\phi(v) = \phi(v^\star) = 0$ implies $v = v^\star$ in a sufficiently small neighborhood. □

Conditioning interpretation. For the least-squares data term $\frac{1}{2}\|r(V)\|_2^2$, the Gauss–Newton approximation to the Hessian in v is $C(V)^\top C(V)$. Thus, if $C(V^\star)$ is ill-conditioned or rank-deficient, the self-consistency signal is weak or ambiguous in some missing directions, motivating regularization through $\lambda\,\mathcal{P}(V)$ as described in Sec. 4.2.

A.4 Minimal finite-dimensional example

We give a concrete linear example illustrating $C(V)$ and why full column rank corresponds to informative coupling.

Setup. Let $T = 2$ and $D = 2$. Vectorize $x = \mathrm{vec}(X) \in \mathbb{R}^4$ as $x = [x_1^{(1)}, x_2^{(1)}, x_1^{(2)}, x_2^{(2)}]^\top$. Assume channel 1 is observed at both times and channel 2 is missing at both times: $\Omega = \{(1,1),(2,1)\}$ and $\bar{\Omega} = \{(1,2),(2,2)\}$, so $|\Omega| = |\bar{\Omega}| = 2$. Let $v = [v_1, v_2]^\top$ denote the missing coordinates.

A time-/channel-coupled linear reconstructor. Define $y = f_\theta(x) \in \mathbb{R}^4$ by

$$y_1^{(1)} = x_1^{(1)} + \alpha\, x_2^{(1)}, \qquad y_2^{(1)} = x_2^{(1)}, \tag{21}$$

$$y_1^{(2)} = x_1^{(2)} + \beta\, x_2^{(1)} + \gamma\, x_2^{(2)}, \qquad y_2^{(2)} = x_2^{(2)}. \tag{22}$$

This encodes that the (missing) channel 2 influences the model's prediction on the (observed) channel 1, including across time via β.

Observed residual and cross-Jacobian. Clamping observed entries and inserting v for the missing channel yields

$$r(V) = \begin{bmatrix} y_1^{(1)} - x_1^{(1)} \\ y_1^{(2)} - x_1^{(2)} \end{bmatrix} = \begin{bmatrix} \alpha\, v_1 \\ \beta\, v_1 + \gamma\, v_2 \end{bmatrix}, \qquad C(V) = \frac{\partial r(V)}{\partial v} = \begin{bmatrix} \alpha & 0 \\ \beta & \gamma \end{bmatrix}. \tag{23}$$

Full column rank holds iff $\alpha\gamma \neq 0$, i.e., iff both missing coordinates locally affect the observed residual. If $\alpha = 0$, then v_1 is unidentifiable from observed self-consistency; if $\gamma = 0$, then v_2 is unidentifiable. This mirrors the intuition in Sec. 4.2: test-time refinement is effective only when the reconstructor induces sufficient coupling from missing to observed coordinates.

B Algorithm

In our experiments we instantiate the optional prior $\mathcal{P}(V)$ as a first-order temporal smoothness penalty on the optimized entries. We use the mask contract from Sec. 3: $M^{\mathrm{fit}} := M^{\mathrm{nat}} \odot (1 - M^{\mathrm{eval}})$ is the fit-residual mask and $O := 1 - M^{\mathrm{fit}}$ is the optimization mask. The data term uses only M^{fit}; baseline-filled native missing entries enter O, never the observations, while evaluator-only H and $\tilde{X}$ never enter optimization. Define the adjacent-validity mask $W_\Delta \in \{0,1\}^{(T-1)\times D}$ by $(W_\Delta)_{t,d} = O_{t,d}\, O_{t-1,d}$, which activates only when both endpoints of a temporal difference are optimized. We then use

$$\mathcal{P}(V) := \frac{1}{2(\|W_\Delta\|_1 + \varepsilon)} \left\| W_\Delta \odot (V_{2:T,:} - V_{1:T-1,:}) \right\|_F^2, \tag{24}$$

with $\varepsilon = 10^{-8}$ for numerical stability. In the batched implementation, W_Δ is computed per sample. Full implementation details are summarized in Algorithm 2. In our current implementation we keep the model in training mode, but disable dropout to be able to use cuDNN backward. This is the same as setting the model to evaluation mode, if no BatchNorm is used. If the model would use BatchNorm, set model to evaluation mode and set torch.backends.cudnn.flags(enabled=False). This is typically slower.

Algorithm 2 Test-time optimization of missing values (frozen reconstructor), batched (implementation-faithful)

Require: baseline-filled batch $Y \in \mathbb{R}^{B\times T\times D}$ (observed entries intact), optimize-mask $O \in \{0,1\}^{B\times T\times D}$ (1 = optimize), fit-residual mask $M^{\text{fit}} \in \{0,1\}^{B\times T\times D}$, frozen reconstructor f_θ, max iterations K, optimizer step size η, prior weight λ, bootstrap window W_0, patience p, max-consecutive q, $\varepsilon > 0$, (optional) projection scale κ with feature mean $\mu \in \mathbb{R}^D$ and std $\sigma \in \mathbb{R}^D$

$V \leftarrow Y$ ▷ optimization variable; same shape as Y
best $\leftarrow (+\infty)^B$, noImp $\leftarrow \mathbf{0}^B$, cons $\leftarrow \mathbf{0}^B$, conv $\leftarrow \mathbf{0} \in \{0,1\}^B$, $\ell_{\text{old}} \leftarrow (+\infty)^B$
$\mathcal{L}_{\text{init}} \leftarrow [\,]$, $\Delta_{\min} \leftarrow 0$, $\tau \leftarrow 0$
Initialize optimizer Opt on V (e.g., Adam with lr η)
for $k = 0, \ldots, K-1$ **do**
  act $\leftarrow \mathbf{1} -$ conv ▷ active samples
  $V[\neg O] \leftarrow Y[\neg O]$ ▷ clamp all fixed (non-optimized) entries
  $\bar{Y} \leftarrow f_\theta(V)$
  $n_F \leftarrow \|M^{\text{fit}}\|_1$ ▷ per-sample count
  $E_{\text{fit}} \leftarrow \frac{1}{2\,(n_F+\varepsilon)}\,\|M^{\text{fit}} \odot (\bar{Y} - Y)\|_F^2$
  $W_\Delta \leftarrow O_{:,2:T,:} \odot O_{:,1:T-1,:}$ ▷ valid diffs: both endpoints optimized
  $n_\Delta \leftarrow \|W_\Delta\|_1$
  $\mathcal{P} \leftarrow \frac{1}{2\,(n_\Delta+\varepsilon)}\,\|W_\Delta \odot (V[:,2:T,:] - V[:,1:T-1,:])\|_F^2$ ▷ per-sample Frobenius norm
  $\mathcal{J} \leftarrow E_{\text{fit}} + \lambda\,\mathcal{P}$ ▷ per-sample objective
  $\mathcal{J}_{\text{tot}} \leftarrow \sum_{b=1}^{B} \text{act}_b\,\mathcal{J}_b$
  Opt.zero_grad(); $\nabla_V \mathcal{J}_{\text{tot}}$ via backprop
  $G \leftarrow (\nabla_V \mathcal{J}_{\text{tot}}) \odot O$ ▷ only update optimized entries
  $G \leftarrow G \odot$ act ▷ mask converged samples
  $G \leftarrow G/(\|G\|_2 + \varepsilon)$ ▷ per-sample gradient normalization
  Set $\nabla V \leftarrow G$ and perform Opt.step()
  $V[\neg O] \leftarrow Y[\neg O]$ ▷ re-clamp fixed entries after the step
  if κ is set **then**
    lower $\leftarrow \mu - \kappa\sigma$, upper $\leftarrow \mu + \kappa\sigma$
    $V[O] \leftarrow \text{clip}(V[O], \text{lower}, \text{upper})$ ▷ projection on optimized entries only
  end if
  $\delta \leftarrow |\ell_{\text{old}} - \mathcal{J}|$, $\ell_{\text{old}} \leftarrow \mathcal{J}$
  if $k < W_0$ **then**
    Append mean($\mathcal{J}$) to $\mathcal{L}_{\text{init}}$
  else if $k = W_0$ **then**
    $\bar{\mathcal{J}}_{\text{init}} \leftarrow \text{mean}(\mathcal{L}_{\text{init}})$
    $\Delta_{\min} \leftarrow 0.01\,\bar{\mathcal{J}}_{\text{init}}$, $\tau \leftarrow 0.01\,\Delta_{\min}$ ▷ adaptive thresholds
    best $\leftarrow \mathcal{J}$
  else
    impr $\leftarrow \mathcal{J} < (\text{best} - \Delta_{\min})$
    best[impr] $\leftarrow \mathcal{J}$[impr]
    noImp[¬impr] $\leftarrow$ noImp[¬impr] $+ 1$
    good $\leftarrow \delta > \tau$
    cons[¬good] $\leftarrow$ cons[¬good] $+ 1$; cons[good] $\leftarrow 0$
    conv $\leftarrow$ conv $\vee$ (noImp $\geq p$) $\vee$ (cons $\geq q$)
  end if
end for
$V[\neg O] \leftarrow Y[\neg O]$, $\bar{Y} \leftarrow f_\theta(V)$
$\hat{X} \leftarrow Y$; $\hat{X}[O] \leftarrow \bar{Y}[O]$ ▷ replace optimized entries by model prediction
return $\hat{X}$

C Experiments

C.1 Benchmark models

We briefly summarize representative deep time-series models used as benchmarks or points of comparison. *iTransformer* (Liu et al. 2023b) adapts the Transformer to multivariate time series by *inverting* the tokenization, i.e., treating variables as tokens (instead of time steps) so that attention primarily models cross-variate dependencies while temporal information is embedded in the token representations. *SAITS* (Du, Côté, and Liu 2023) is a Transformer-based imputation model that uses diagonally masked self-attention blocks and combines multiple attention-based reconstructions to explicitly target missing-value recovery in multivariate sequences. *BRITS* (Cao et al. 2018) performs imputation with bidirectional recurrent networks, propagating information forward and backward in time and iteratively refining estimates through consistency between the two directions. *TimesNet* (Wu et al. 2022) models time series via a 2D-style representation that captures multi-periodic temporal structure, enabling a unified architecture that has been applied across forecasting and other sequence tasks, including imputation. *FreTS* (Yi et al. 2023) operates in the frequency domain by transforming the input sequence and applying MLP-based modeling on spectral components, leveraging frequency-domain structure for time-series learning.

C.2 Hyperparameter Optimization

For the backbone models, we used Optuna (Akiba et al. 2019) to sample hyperparameters from the discrete search space in Table 2. Unless stated otherwise, categorical variables were sampled uniformly over the listed choices, and continuous variables were sampled uniformly over the given interval. The learning rate was sampled log-uniformly.

C.3 Native Missingness

Table 3 reports the dimensionality, window length, and native missingness rate for each dataset before any artificial corruption is applied. This distinction matters because the frozen reconstructors are trained only on the native training data, while the benchmark evaluation masks are added synthetically at test time.

C.4 Detailed Full-Channel-Missing Results

Table 4 reports every matched-mask and OOD structured-mask checkpoint underlying the compact family-level comparison in the main paper. The individual rows are retained here so that the descriptive lower envelopes in Table 1 can be traced to their source methods and training-mask regimes.

C.5 Trainable Parameters

Table 5 compares trainable parameter counts for the mask-aware baselines and the reconstruction backbones used with test-time optimization. We retain these counts as supplementary capacity context rather than as a core contribution or scaling claim; several TTO backbones are smaller than the largest deep imputation baselines, but model size is not the focus of the conversion study.

C.6 Jacobian diagnostics for test-time optimization

The analysis in Appendix A identifies the cross-Jacobian $C(V) = \partial r(V)/\partial v$ as the local object that determines whether changes in the optimized missing coordinates are visible through the observed self-consistency residual. To check this mechanism empirically, we ran Jacobian diagnostics on the full-channel-missing setting for LSTM, TCN-AE, and TransformerAE across all five datasets, comparing TTO with λ=0 against TTO with the smoothness prior λ=10^{-3}. MTAD-GAT (Zhao et al. 2020) was omitted because of its substantially higher diagnostic runtime. For each model dataset variant pair, we used the selected checkpoint round from the corresponding TTO run and evaluated 32 test windows, recording the mean MAE, the smallest singular value of $C(V)$, the spectral condition number, and the Gauss–Newton curvature range along the optimization trajectory.

In this setting, the condition number is

$$\kappa(C(V)) = \frac{\sigma_{\max}(C(V))}{\sigma_{\min}(C(V))}, \tag{25}$$

where $C(V)$ maps perturbations of the optimized missing entries to perturbations of the observed residual. Small $\sigma_{\min}$ or large $\kappa(C(V))$ indicate directions in the missing channel that are only weakly constrained by observed self-consistency; in such directions, the least-squares curvature $C(V)^\top C(V)$ is flat or anisotropic, so optimization can be slow, sensitive, or biased toward arbitrary completions. Moderate condition numbers with non-negligible $\sigma_{\min}$ instead indicate that the reconstructor couples the missing channel to the observed channels strongly enough for TTO to receive informative gradients. The smoothness prior does not change $C(V)$ itself; rather, it adds an external regularizing term to the objective, which can stabilize underconstrained directions when the data term alone is weak.

Table 6 summarizes the paired diagnostic runs. The prior improves the final diagnostic MAE in 11 of 15 model–dataset pairs, with the largest gains on LSTM/BeijingAir, LSTM/ETTh1, TCN-AE/ItalyAir, TransformerAE/BeijingAir, and TransformerAE/ETTh1. The few negative cases are small, which is consistent with treating the prior as a stabilizer rather than as a uniformly beneficial accuracy mechanism. Figure 4 visualizes the same paired effect as a model-by-dataset heatmap. Figure 5

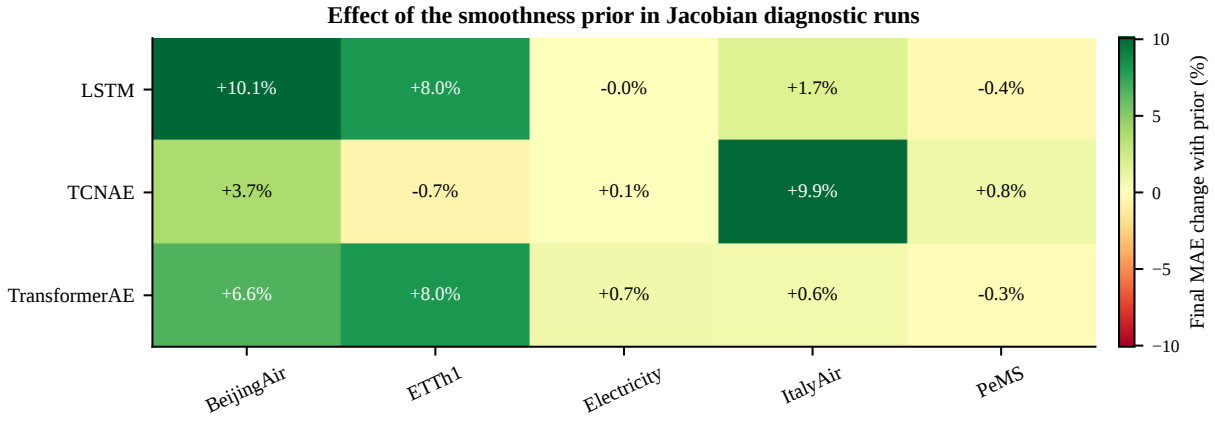

Figure 4: **Effect of the smoothness prior in the Jacobian diagnostic runs.** Each cell reports the percentage change in final diagnostic MAE when adding the smoothness prior (λ=10^{-3}) relative to TTO without the prior (λ=0), for the same model, dataset, and selected checkpoint round. Positive values indicate that the prior reduces MAE. The strongest gains occur in several BeijingAir and ETTh1 settings, while PeMS and TCN-AE/ETTh1 show small negative or near-neutral changes.

shows representative optimization trajectories: the left panels track the mean MAE over TTO iterations, while the right panels track the Jacobian condition number on a log scale. The representative plots are consistent with the mechanism that useful missing-to-observed coupling permits informative refinement, while weak or anisotropic coupling can leave directions under-constrained. They are illustrative diagnostics and do not establish a systematic association by which condition number alone predicts imputation accuracy or prior benefit.

C.7 Results on Benchmark Missingness

Tables 7, 8, and 9 report in-distribution MAE on the standard TSI-Bench point, subsequence, and block missingness settings. They complement the main full-channel-missing table by showing how the same benchmark baselines and frozen reconstructors behave when the test masks match the usual structured missingness regimes.

C.8 Clock Runtime

Table 10 reports wall-clock inference time on the full-channel test split. The benchmark rows measure one-shot imputation, while the TTO rows separately report optimization without the smoothness prior and optimization with the prior, making the accuracy–runtime tradeoff explicit. Because the test splits contain different numbers of windows, Table 11 additionally normalizes runtime to seconds per 1000 windows and reports the median/p90 stopping iterations for TTO. The normalized table is generated from the runtime summary CSVs: `benchmark_runtime_summary.csv`, `autoencoder_runtime_summary.csv`, and the corresponding batch-probe metadata in `batch_probe_summary.csv`.

C.9 Convergence

Figures 6, 7, 8, 9, and 10 visualize TTO convergence on the full-channel-missing setting for each dataset. These plots are diagnostic only: they show test MAE over optimization iterations, and the marked best-MAE iterations are not used to select the table results, which use the objective-based stopping rule described in Sec. 5.6.

Hyperparameter	Type / Range	Notes
LSTM		
hidden_size	categorical $\in \{16, 32, 64, 96, 128, 172, 256, 384, 512, 768, 1024\}$	LSTM hidden dimension
latent_size	categorical $\in \{4, 6, 8, 12, 16, 32, 64, 96, 128, 172, 256, 384, 512, 768, 1024\}$	bottleneck dimension
n_layers	integer $\in \{1, 2, 3, 4, 5\}$	number of LSTM layers
dropout	float $\in \{0.00, 0.05, 0.10\}$	LSTM dropout
batch_size	categorical $\in \{32, 64, 128, 256\}$	
lr	log-uniform in $[5 \cdot 10^{-4},\ 5 \cdot 10^{-3}]$	learning rate
proj_dim	categorical $\in \{16, 32, 64, 128, 256\}$	projection-layer width
input_proj_dim	categorical $\in \{16, 32, 64, 128, 256\}$	input projection width
post_mlp_ratio	categorical $\in \{2, 4, 8\}$	MLP expansion ratio
post_mlp_dropout	float $\in \{0.00, 0.05, 0.10\}$	MLP dropout
TCNAE		
channel_preset	categorical $\in$ {tiny, small_tight, small, medium_tight, medium_plus, large_tight, large}	sets base and latent channels
latent_pool	categorical $\in \{2, 3, 4, 6, 8, 12\}$	temporal bottleneck pooling
kernel_size	categorical $\in \{3, 5, 7, 9\}$	temporal convolution kernel
dilation_preset	categorical $\in$ {local, default, deep, extra_deep}	encoder/decoder dilation schedule
dropout	categorical $\in \{0.00, 0.05, 0.10, 0.20\}$	convolutional dropout
batch_size	categorical $\in \{1, 2, 4, 8, 16, 32, 64\}$	dataset-dependent subset
lr	log-uniform in $[10^{-4},\ 5 \cdot 10^{-3}]$	dataset-dependent upper bound
TransformerAE		
d_model	categorical $\in \{32, 64, 96, 128\}$	transformer width
n_heads	categorical $\in \{2, 4, 8\}$	attention heads
latent_dim_ratio	categorical $\in \{0.125, 0.25, 0.5\}$	determines latent dimension
ff_ratio	categorical $\in \{2, 4\}$	encoder/decoder FFN expansion
latent_ff_ratio	categorical $\in \{2, 4\}$	latent FFN expansion
encoder_depth	categorical $\in \{1, 2, 3\}$	encoder blocks
latent_depth	categorical $\in \{1, 2\}$	latent transformer blocks
decoder_depth	categorical $\in \{1, 2, 3\}$	decoder blocks
downsample_factor	categorical $\in \{2, 3, 4, 6, 8, 12\}$	temporal compression
dropout	categorical $\in \{0.00, 0.05, 0.10, 0.20\}$	transformer dropout
stem_kernel_size	categorical $\in \{1, 3, 5\}$	input convolution kernel
output_head_kernel_size	categorical $\in \{1, 3\}$	output head kernel
upsample_mode	categorical $\in$ {linear, nearest}	decoder upsampling
batch_size	categorical $\in \{1, 2, 4, 8, 16, 32, 64\}$	dataset-dependent subset
lr	log-uniform in $[10^{-4},\ 5 \cdot 10^{-3}]$	dataset-dependent upper bound
MTAD-GAT		
kernel_size	categorical $\in \{1, 3, 5, 7, 9\}$	temporal convolution kernel
feat_gat_embed_dim	categorical $\in \{2, 4, 8, 16, 24\}$	feature-attention embedding
time_gat_embed_dim	categorical $\in \{4, 8, 16, 24, 32\}$	temporal-attention embedding
gru_n_layers	categorical $\in \{1, 2\}$	GRU depth
gru_hid_dim	categorical $\in \{8, 16, 32, 48, 64, 96, 128\}$	GRU hidden dimension
latent_dim_ratio	categorical $\in \{0.125, 0.25, 0.5\}$	determines latent dimension
recon_n_layers	categorical $\in \{1, 2\}$	reconstruction MLP depth
recon_hid_dim	categorical $\in \{8, 16, 32, 48, 64, 96, 128\}$	reconstruction hidden width
dropout	categorical $\in \{0.00, 0.05, 0.10, 0.20\}$	model dropout
alpha	categorical $\in \{0.05, 0.10, 0.20, 0.30\}$	LeakyReLU slope in GAT layers
batch_size	categorical $\in \{1, 2, 4, 8, 16, 32, 64\}$	dataset-dependent subset
lr	log-uniform in $[10^{-4},\ 5 \cdot 10^{-3}]$	dataset-dependent upper bound
Fixed training settings		
epochs	fixed $= 100$	maximum training epochs
patience	fixed $= 50$	early-stopping patience
n_trials	fixed $= 20$ for backbone autoencoders	Optuna trials per model–dataset pair

Table 2: Optuna hyperparameter search spaces for the LSTM and backbone autoencoder models. For TCN-AE, TransformerAE, and MTAD-GAT, reported ranges are the union over datasets; the implementation uses dataset-specific subsets to account for sequence length and memory constraints.

Dataset	D	T (window)	Native missing rate
BeijingAir	132	24	1.60%
ItalyAir	13	12	0%
ETTh1	7	48	0%
Electricity	370	96	0%
PeMS	862	24	0%

Table 3: **Dataset characteristics and native missingness.** Native missingness rates refer to the raw series prior to any artificial corruption. D is the number of variables and T is the non-overlapping window length used to generate samples, following the preprocessing summary in TSI-Bench (Du et al. 2024).

Setting		BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
Simple statistical baseline						
MeanFill	train mean	0.7718	0.8043	1.9657	0.7974	0.7742
Matched full-channel-mask baselines						
BRITS	full feature	**0.1383 ± 0.0028**	0.9848 ± 0.1342	1.3824 ± 0.0316	0.3588 ± 0.0244	0.2690 ± 0.0032
FreTS	full feature	0.1655 ± 0.0024	0.6003 ± 0.0359	1.3480 ± 0.0879	0.4360 ± 0.0318	0.3284 ± 0.0096
SAITS	full feature	0.1575 ± 0.0053	0.3222 ± 0.0261	1.6103 ± 0.0193	0.3648 ± 0.0197	0.2508 ± 0.0037
TimesNet	full feature	0.7039 ± 0.0024	0.8059 ± 0.0021	1.9124 ± 0.0827	0.5522 ± 0.0094	0.6518 ± 0.0192
iTransformer	full feature	0.7634 ± 0.0048	0.8373 ± 0.0032	1.9117 ± 0.0480	0.7981 ± 0.0006	0.6733 ± 0.0325
OOD transfer from standard structured-mask training						
BRITS	block	0.2108 ± 0.0018	0.3180 ± 0.0111	1.2736 ± 0.0307	0.3665 ± 0.0133	0.2485 ± 0.0024
	point	0.2115 ± 0.0019	0.3227 ± 0.0129	1.2963 ± 0.0620	0.3740 ± 0.0144	0.2472 ± 0.0033
	subseq	0.2128 ± 0.0033	0.3287 ± 0.0124	1.2555 ± 0.0126	0.3797 ± 0.0184	0.2494 ± 0.0026
FreTS	block	0.2371 ± 0.0122	0.4330 ± 0.0395	1.1770 ± 0.0879	0.4054 ± 0.0397	0.3743 ± 0.0069
	point	0.2431 ± 0.0200	0.4322 ± 0.0288	1.1452 ± 0.0348	0.4137 ± 0.0338	0.3773 ± 0.0113
	subseq	0.2418 ± 0.0185	0.4302 ± 0.0343	1.1581 ± 0.0754	0.4122 ± 0.0333	0.3732 ± 0.0076
SAITS	block	0.2233 ± 0.0020	0.3546 ± 0.0109	1.6379 ± 0.0675	**0.3488 ± 0.0087**	0.2501 ± 0.0046
	point	0.2276 ± 0.0066	0.3543 ± 0.0072	1.6135 ± 0.0465	0.3513 ± 0.0105	0.2481 ± 0.0025
	subseq	0.2243 ± 0.0031	0.3545 ± 0.0066	1.6170 ± 0.0732	0.3501 ± 0.0088	**0.2470 ± 0.0016**
TimesNet	block	0.4501 ± 0.0074	0.7612 ± 0.0108	1.7867 ± 0.0142	0.5805 ± 0.0048	0.2666 ± 0.0025
	point	0.4471 ± 0.0101	0.7624 ± 0.0103	1.8161 ± 0.0526	0.5798 ± 0.0054	0.2678 ± 0.0037
	subseq	0.4582 ± 0.0124	0.7637 ± 0.0097	1.7747 ± 0.0101	0.5812 ± 0.0047	0.2663 ± 0.0019
iTransformer	block	0.7631 ± 0.0032	0.8072 ± 0.0052	1.9836 ± 0.0429	0.7980 ± 0.0014	0.5479 ± 0.0220
	point	0.7634 ± 0.0040	0.8062 ± 0.0033	2.0083 ± 0.1209	0.7980 ± 0.0013	0.5596 ± 0.0118
	subseq	0.7613 ± 0.0030	0.8047 ± 0.0023	1.9637 ± 0.0382	0.7976 ± 0.0011	0.5810 ± 0.0342
Frozen reconstructors with test-time optimization						
LSTM	forward (no opt)	0.4968 ± 0.0145	0.6488 ± 0.0024	1.1901 ± 0.0389	0.5845 ± 0.0586	0.2944 ± 0.0064
	opt, λ=0 (no prior)	0.3413 ± 0.0176	0.3286 ± 0.0216	1.1821 ± 0.0402	0.3806 ± 0.0365	0.2913 ± 0.0060
	opt + smooth prior (λ=10^{-3})	0.3099 ± 0.0169	**0.2898 ± 0.0416**	1.1821 ± 0.0402	0.3796 ± 0.0365	0.2917 ± 0.0055
MTAD-GAT	forward (no opt)	0.2776 ± 0.0054	0.6227 ± 0.0203	1.5608 ± 0.0982	0.4942 ± 0.0128	0.3393 ± 0.0064
	opt, λ=0 (no prior)	0.2726 ± 0.0042	0.5320 ± 0.0321	1.5580 ± 0.0990	0.4815 ± 0.0357	0.3378 ± 0.0062
	opt + smooth prior (λ=10^{-3})	0.2728 ± 0.0051	0.5323 ± 0.0311	1.5580 ± 0.0990	0.4822 ± 0.0336	0.3377 ± 0.0061
TCN-AE	forward (no opt)	0.2455 ± 0.0039	0.6501 ± 0.0676	**0.9853 ± 0.0287**	0.7200 ± 0.0241	0.2898 ± 0.0030
	opt, λ=0 (no prior)	0.2289 ± 0.0059	0.4299 ± 0.0347	0.9928 ± 0.0230	0.6557 ± 0.0268	0.2892 ± 0.0026
	opt + smooth prior (λ=10^{-3})	0.2203 ± 0.0067	0.4329 ± 0.0307	0.9926 ± 0.0229	0.6404 ± 0.0241	0.2884 ± 0.0028
TransformerAE	forward (no opt)	0.2763 ± 0.0041	0.7610 ± 0.1168	1.0525 ± 0.0167	0.6440 ± 0.0399	0.2822 ± 0.0033
	opt, λ=0 (no prior)	0.2255 ± 0.0043	0.6197 ± 0.1172	1.0354 ± 0.0220	0.4350 ± 0.0615	0.2811 ± 0.0051
	opt + smooth prior (λ=10^{-3})	0.2200 ± 0.0034	0.5543 ± 0.1091	1.0353 ± 0.0221	0.4353 ± 0.0619	0.2816 ± 0.0055

Table 4: **Detailed full-channel-missing evaluation (MAE; mean ± std where applicable).** One channel is fully masked at test time, in addition to native missingness, with balanced cyclic channel assignment. MeanFill is a deterministic per-channel training-mean imputer. *Full feature* denotes matched channel-drop training; *point*, *subseq*, and *block* denote OOD transfer from 10% structured-mask training. Frozen reconstructors are evaluated with forward-only inference and TTO. Best results are marked in bold; second-best results are underlined. This table expands the compact family-level comparison in the main paper and preserves every individual checkpoint row.

	BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
BRITS	3,598,496	2,178,496	17,082,800	596,912	32,012,048
FreTS	909,852	465,271	3,706,194	668,313	1,715,958
SAITS	7,153,808	88,235,470	63,624,720	16,628,642	78,229,072
TimesNet	87,063,940	5,510,663	45,569,394	22,051,853	91,622,238
iTransformer	8,286,232	23,723,056	12,989,024	18,932,236	1,854,744
TCN-AE	1,030,052	1,000,999	513,810	*39,965*	**461,822**
TransformerAE	**338,676**	1,075,975	768,322	89,997	*650,182*
MTAD-GAT	*446,624*	*106,188*	977,936	**32,156**	1,774,276
LSTM	873,492	**12,669**	**42,930**	494,781	894,142

Table 5: Trainable parameters per model and dataset. Bold indicates the smallest model for each dataset; italics indicate the second-smallest model.

Model	Dataset	Round	MAE, λ=0	MAE, prior	Δ final (%)	Δ best (%)	$\sigma_{\min}$, prior	Cond. ratio
LSTM	BeijingAir	5	0.3849	0.3461	**+10.1**	**+8.7**	0.0856	1.02
LSTM	ETTh1	2	0.3523	0.3241	**+8.0**	**+8.0**	0.1943	1.00
LSTM	Electricity	1	1.0405	1.0405	-0.0	+0.0	0.0891	1.00
LSTM	ItalyAir	0	0.3701	0.3638	**+1.7**	**+1.7**	0.2317	1.02
LSTM	PeMS	1	0.2916	0.2929	-0.4	-0.4	6.23e-03	0.89
TCNAE	BeijingAir	4	0.2992	0.2881	**+3.7**	**+1.7**	0.0145	1.04
TCNAE	ETTh1	0	0.6019	0.6059	-0.7	-0.7	0.0605	1.00
TCNAE	Electricity	0	1.0434	1.0423	**+0.1**	+0.0	1.03e-03	1.04
TCNAE	ItalyAir	0	0.6632	0.5973	**+9.9**	**+5.3**	0.0747	1.03
TCNAE	PeMS	3	0.2943	0.2919	**+0.8**	**+0.8**	1.13e-04	1.00
TransformerAE	BeijingAir	3	0.2811	0.2627	**+6.6**	**+6.6**	0.1165	0.98
TransformerAE	ETTh1	0	0.9656	0.8888	**+8.0**	**+1.5**	0.0394	0.95
TransformerAE	Electricity	5	1.0124	1.0049	**+0.7**	**+0.7**	0.0438	1.00
TransformerAE	ItalyAir	0	0.3979	0.3954	**+0.6**	**+0.6**	0.1409	0.98
TransformerAE	PeMS	4	0.3048	0.3057	-0.3	-0.3	0.0622	0.97

Table 6: **Jacobian diagnostic summary for TTO with and without the smoothness prior.** Each row compares paired diagnostic runs for the selected checkpoint round. Δ final and Δ best report the percentage MAE reduction from adding the prior (positive values are improvements). The condition ratio is the final median Jacobian condition number with prior divided by the corresponding value without prior.

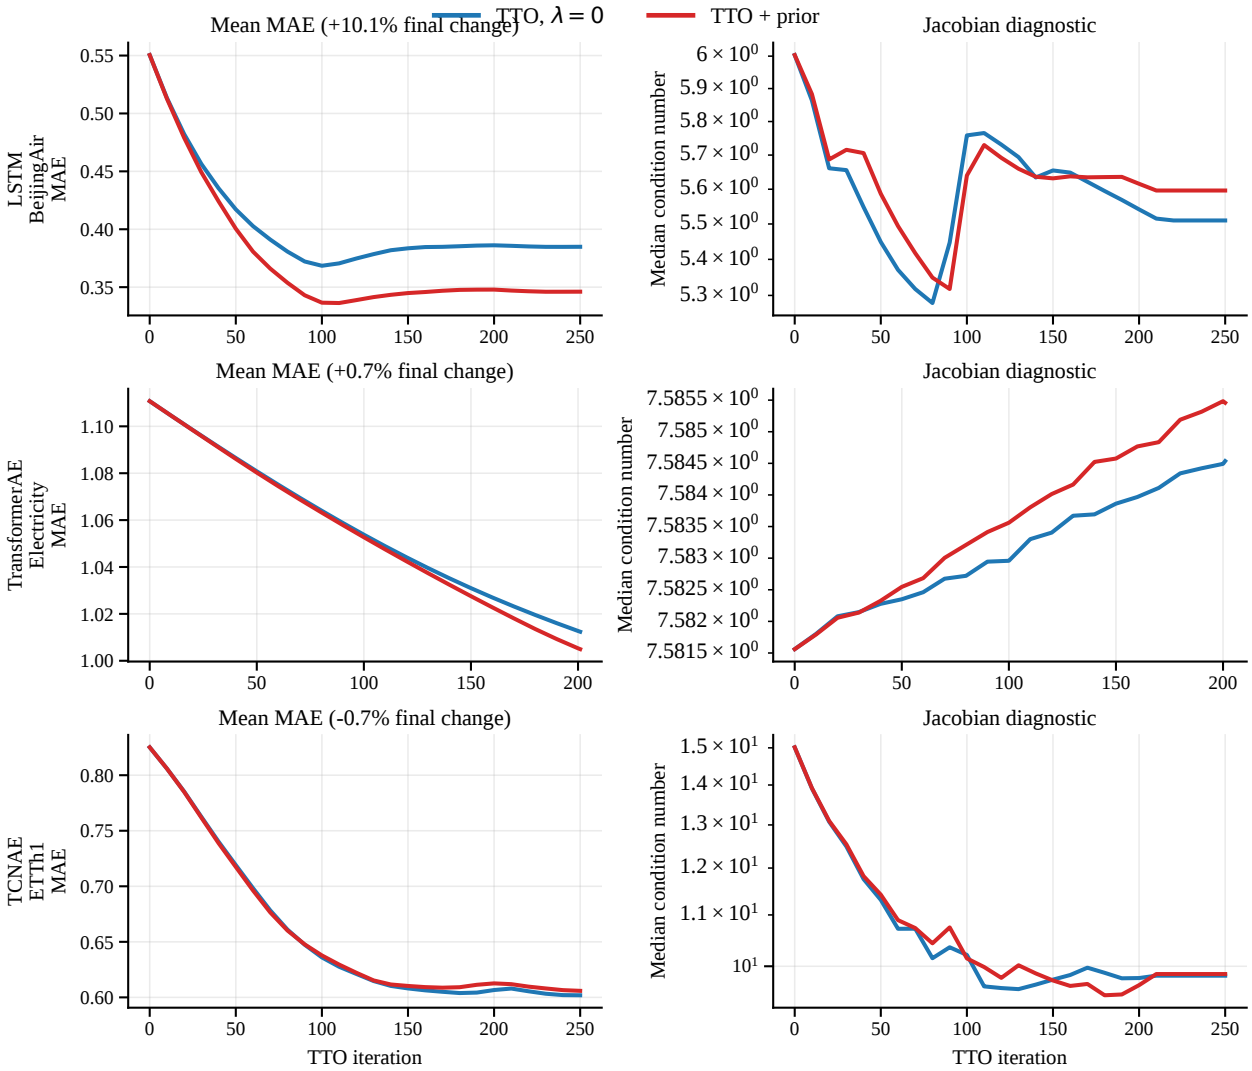

Figure 5: **Representative TTO trajectories with Jacobian diagnostics.** For three model–dataset pairs, the left column shows mean evaluation MAE over diagnostic TTO iterations and the right column shows the median cross-Jacobian condition number $\kappa(C(V))$ on a log scale. Blue curves correspond to TTO without the prior and red curves to TTO with the smoothness prior. These traces connect the accuracy changes in Table 6 to the local conditioning of the observed self-consistency map: successful refinement reduces MAE while maintaining a usable coupling between optimized missing entries and observed residuals.

	Setting	BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
MeanFill	train mean	0.7563	0.8305	1.8667	0.7716	0.8037
BRITS	point	0.1279 ± 0.0020	0.1550 ± 0.0041	0.9704 ± 0.0182	0.2269 ± 0.0059	0.2605 ± 0.0019
FreTS	point	0.2012 ± 0.0091	0.2069 ± 0.0071	**0.6853 ± 0.1011**	0.2362 ± 0.0158	0.3515 ± 0.0145
SAITS	point	0.1550 ± 0.0042	**0.1215 ± 0.0040**	1.4574 ± 0.0146	0.1707 ± 0.0083	0.2656 ± 0.0010
TimesNet	point	0.4098 ± 0.0077	0.7260 ± 0.0080	1.7587 ± 0.0131	0.5349 ± 0.0095	0.2814 ± 0.0025
iTransformer	point	**0.1046 ± 0.0024**	0.2623 ± 0.0066	0.8429 ± 0.1469	0.1806 ± 0.0032	**0.1490 ± 0.0110**
LSTM	forward (no opt)	0.5111 ± 0.0116	0.6856 ± 0.0023	1.1178 ± 0.0355	0.5474 ± 0.0387	0.3198 ± 0.0037
	opt, λ=0 (no prior)	0.4317 ± 0.0063	0.2930 ± 0.0182	1.0843 ± 0.0328	0.3103 ± 0.0470	0.3077 ± 0.0014
	opt + smooth prior (λ=10^{-3})	0.4232 ± 0.0059	0.2873 ± 0.0187	1.0842 ± 0.0328	0.3074 ± 0.0474	0.3075 ± 0.0016
MTAD-GAT	forward (no opt)	0.2917 ± 0.0045	0.4617 ± 0.0143	1.4727 ± 0.0705	0.3592 ± 0.0095	0.3593 ± 0.0024
	opt, λ=0 (no prior)	0.2699 ± 0.0014	0.4254 ± 0.0171	1.3859 ± 0.0842	0.3157 ± 0.0142	0.3418 ± 0.0083
	opt + smooth prior (λ=10^{-3})	0.2701 ± 0.0014	0.4274 ± 0.0169	1.3859 ± 0.0842	0.3157 ± 0.0141	0.3423 ± 0.0081
TCN-AE	forward (no opt)	0.2585 ± 0.0011	0.3856 ± 0.0272	0.9277 ± 0.0186	0.4285 ± 0.0126	0.3292 ± 0.0044
	opt, λ=0 (no prior)	0.2182 ± 0.0010	0.1940 ± 0.0052	0.9069 ± 0.0190	0.1822 ± 0.0167	0.3121 ± 0.0008
	opt + smooth prior (λ=10^{-3})	0.2180 ± 0.0010	0.1933 ± 0.0052	0.9068 ± 0.0190	0.1780 ± 0.0173	0.3120 ± 0.0008
TransformerAE	forward (no opt)	0.2575 ± 0.0039	0.6651 ± 0.1522	0.9027 ± 0.0239	0.4165 ± 0.1061	0.3211 ± 0.0006
	opt, λ=0 (no prior)	0.2074 ± 0.0011	0.3451 ± 0.1150	0.8650 ± 0.0232	0.1624 ± 0.0199	0.3064 ± 0.0013
	opt + smooth prior (λ=10^{-3})	0.2073 ± 0.0010	0.3367 ± 0.1167	0.8652 ± 0.0232	**0.1606 ± 0.0194**	0.3059 ± 0.0013

Table 7: **Normal test set evaluation under point missingness (MAE; mean ± std where applicable).** MeanFill is a deterministic per-channel training-mean imputer and is reported as a single value. BRITS, FreTS, SAITS, TimesNet, and iTransformer are SOTA baselines evaluated under point missingness. LSTM, MTAD-GAT, TCN-AE, and TransformerAE are frozen reconstructors trained without artificial missingness and evaluated with forward-only inference and test-time input optimization. Best results are marked in bold; second-best results are underlined.

	Setting	BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
MeanFill	train mean	0.7610	0.9757	1.8906	0.7615	0.8097
BRITS	subseq	**0.1417 ± 0.0008**	0.3876 ± 0.0223	1.2899 ± 0.0085	0.3697 ± 0.0187	0.3208 ± 0.0014
FreTS	subseq	0.2324 ± 0.0065	0.4546 ± 0.0281	1.0865 ± 0.0884	0.3765 ± 0.0103	0.3980 ± 0.0090
SAITS	subseq	0.1602 ± 0.0020	**0.3414 ± 0.0106**	1.5867 ± 0.0628	**0.2575 ± 0.0157**	**0.3184 ± 0.0011**
TimesNet	subseq	0.4642 ± 0.0074	0.9219 ± 0.0112	1.7342 ± 0.0094	0.5939 ± 0.0070	0.3440 ± 0.0033
iTransformer	subseq	0.6077 ± 0.0805	0.8846 ± 0.0043	1.6973 ± 0.0693	0.5060 ± 0.0253	0.5700 ± 0.0217
LSTM	forward (no opt)	0.5268 ± 0.0112	0.7869 ± 0.0026	1.1682 ± 0.0426	0.5679 ± 0.0494	0.3715 ± 0.0054
	opt, λ=0 (no prior)	0.4494 ± 0.0089	0.5533 ± 0.0251	1.1279 ± 0.0367	0.3755 ± 0.0405	0.3583 ± 0.0024
	opt + smooth prior (λ=10^{-3})	0.4168 ± 0.0065	0.4773 ± 0.0250	1.1279 ± 0.0367	0.3665 ± 0.0408	0.3582 ± 0.0029
MTAD-GAT	forward (no opt)	0.3059 ± 0.0059	0.6623 ± 0.0165	1.5176 ± 0.0682	0.4724 ± 0.0114	0.4166 ± 0.0031
	opt, λ=0 (no prior)	0.2775 ± 0.0022	0.5202 ± 0.0343	1.4845 ± 0.0807	0.3906 ± 0.0185	0.3953 ± 0.0097
	opt + smooth prior (λ=10^{-3})	0.2775 ± 0.0023	0.5175 ± 0.0359	1.4845 ± 0.0807	0.3917 ± 0.0178	0.3980 ± 0.0085
TCN-AE	forward (no opt)	0.3207 ± 0.0018	0.7947 ± 0.0762	1.0273 ± 0.0200	0.6806 ± 0.0196	0.3819 ± 0.0041
	opt, λ=0 (no prior)	0.2672 ± 0.0015	0.5316 ± 0.0086	1.0025 ± 0.0215	0.5988 ± 0.0388	0.3717 ± 0.0011
	opt + smooth prior (λ=10^{-3})	0.2652 ± 0.0011	0.5313 ± 0.0086	**1.0024 ± 0.0215**	0.5441 ± 0.0374	0.3713 ± 0.0011
TransformerAE	forward (no opt)	0.3502 ± 0.0052	0.9488 ± 0.1760	1.0727 ± 0.0244	0.6143 ± 0.0493	0.3738 ± 0.0013
	opt, λ=0 (no prior)	0.2804 ± 0.0017	0.7832 ± 0.0633	1.0328 ± 0.0298	0.3566 ± 0.0365	0.3638 ± 0.0011
	opt + smooth prior (λ=10^{-3})	0.2788 ± 0.0015	0.7173 ± 0.0635	1.0331 ± 0.0298	0.3335 ± 0.0265	0.3621 ± 0.0011

Table 8: **Normal test set evaluation under subsequence missingness (MAE; mean ± std where applicable).** MeanFill is a deterministic per-channel training-mean imputer and is reported as a single value. BRITS, FreTS, SAITS, TimesNet, and iTransformer are SOTA baselines evaluated under subsequence missingness. LSTM, MTAD-GAT, TCN-AE, and TransformerAE are frozen reconstructors trained without artificial missingness and evaluated with forward-only inference and test-time input optimization. Best results are marked in bold; second-best results are underlined.

	Setting	BeijingAir	ETTh1	Electricity	ItalyAir	PeMS
MeanFill	train mean	0.7534	0.8154	1.8593	0.8918	0.7972
BRITS	block	**0.1474 ± 0.0012**	0.4695 ± 0.0100	1.0111 ± 0.0197	0.3635 ± 0.0027	0.3083 ± 0.0031
FreTS	block	0.2054 ± 0.0207	0.4226 ± 0.0385	0.9126 ± 0.2514	0.4007 ± 0.0308	0.3880 ± 0.0067
SAITS	block	0.1683 ± 0.0018	**0.3448 ± 0.0806**	1.5833 ± 0.0267	**0.2476 ± 0.0112**	0.3043 ± 0.0018
TimesNet	block	0.4482 ± 0.0081	0.7725 ± 0.0083	1.7068 ± 0.0106	0.6722 ± 0.0077	0.3308 ± 0.0022
iTransformer	block	0.1951 ± 0.0066	0.3939 ± 0.0086	**0.6942 ± 0.1106**	0.3394 ± 0.0403	**0.3010 ± 0.0143**
LSTM	forward (no opt)	0.5264 ± 0.0092	0.7706 ± 0.0034	1.1792 ± 0.0376	0.7185 ± 0.0597	0.3571 ± 0.0058
	opt, λ=0 (no prior)	0.4749 ± 0.0083	0.8890 ± 0.0486	1.1350 ± 0.0353	0.4506 ± 0.0427	0.3435 ± 0.0028
	opt + smooth prior (λ=10^{-3})	0.4334 ± 0.0072	0.7275 ± 0.0306	1.1349 ± 0.0354	0.4117 ± 0.0391	0.3436 ± 0.0029
MTAD-GAT	forward (no opt)	0.2951 ± 0.0052	0.5144 ± 0.0097	1.4856 ± 0.0712	0.4730 ± 0.0040	0.4035 ± 0.0027
	opt, λ=0 (no prior)	0.2704 ± 0.0023	0.4815 ± 0.0143	1.4272 ± 0.0761	0.3865 ± 0.0248	0.3807 ± 0.0102
	opt + smooth prior (λ=10^{-3})	0.2702 ± 0.0024	0.4803 ± 0.0176	1.4272 ± 0.0761	0.3866 ± 0.0244	0.3840 ± 0.0094
TCN-AE	forward (no opt)	0.3471 ± 0.0013	0.7166 ± 0.0141	1.0702 ± 0.0174	0.7647 ± 0.0159	0.3689 ± 0.0042
	opt, λ=0 (no prior)	0.2899 ± 0.0026	0.5184 ± 0.0156	1.0230 ± 0.0197	0.6337 ± 0.0392	0.3564 ± 0.0009
	opt + smooth prior (λ=10^{-3})	0.2859 ± 0.0015	0.4837 ± 0.0079	1.0227 ± 0.0197	0.5662 ± 0.0253	0.3559 ± 0.0009
TransformerAE	forward (no opt)	0.3832 ± 0.0049	0.8001 ± 0.0416	1.0665 ± 0.0235	0.6863 ± 0.0718	0.3608 ± 0.0013
	opt, λ=0 (no prior)	0.2966 ± 0.0017	0.9486 ± 0.0600	0.9956 ± 0.0242	0.3718 ± 0.0410	0.3491 ± 0.0012
	opt + smooth prior (λ=10^{-3})	0.2930 ± 0.0018	0.7698 ± 0.0226	0.9960 ± 0.0242	0.3253 ± 0.0281	0.3475 ± 0.0012

Table 9: **Normal test set evaluation under block missingness (MAE; mean ± std where applicable).** MeanFill is a deterministic per-channel training-mean imputer and is reported as a single value. BRITS, FreTS, SAITS, TimesNet, and iTransformer are SOTA baselines evaluated under block missingness. LSTM, MTAD-GAT, TCN-AE, and TransformerAE are frozen reconstructors trained without artificial missingness and evaluated with forward-only inference and test-time input optimization. Best results are marked in bold; second-best results are underlined.

Model	BeijingAir	ETT h1	Electricity	ItalyAir	PeMS
BRITS	0.501	0.309	2.214	0.139	0.347
FreTS	0.056	0.022	0.788	0.051	0.102
SAITS	0.038	0.149	0.869	0.021	0.156
TimesNet	0.105	0.008	0.865	0.010	0.144
iTransformer	0.391	0.014	1.484	0.050	0.796
LSTM + TTO (no prior)	10.341	2.255	5.276	4.987	2.912
LSTM + TTO (with prior)	10.223	2.486	6.314	5.053	3.249
MTAD-GAT + TTO (no prior)	12.737	1.383	193.631	1.177	171.675
MTAD-GAT + TTO (with prior)	12.783	1.409	194.713	1.258	171.788
TCN-AE + TTO (no prior)	12.381	3.510	10.659	6.330	3.251
TCN-AE + TTO (with prior)	12.641	3.554	11.631	6.423	3.442
TransformerAE + TTO (no prior)	5.480	3.969	20.833	6.393	3.731
TransformerAE + TTO (with prior)	5.563	4.045	21.790	6.606	4.040

Table 10: Wall-clock inference time in seconds on the full test set on identical hardware. Benchmark rows report the fastest one-shot runtime across training missingness types (`point`/`subseq`/`block`). TTO rows report the runtime for each optimization variant separately.

Model	BeijingAir (N=304)	ETT h1 (N=72)	Electricity (N=303)	ItalyAir (N=155)	PeMS (N=153)
Fastest one-shot benchmark latency (s / 1000 windows)					
BRITS	1.65	4.29	7.31	0.90	2.27
FreTS	0.18	0.31	2.60	0.33	0.67
SAITS	0.13	2.08	2.87	0.14	1.02
TimesNet	0.34	0.11	2.85	0.07	0.94
iTransformer	1.29	0.20	4.90	0.32	5.20
TTO latency (s / 1000 windows; median/p90 stopping iterations in parentheses)					
LSTM + TTO (no prior)	34.02 (201/203)	31.32 (201/328)	17.41 (201/201)	32.18 (201/204)	19.04 (202/203)
LSTM + TTO (with prior)	33.63 (201/203)	34.53 (201/327)	20.84 (201/201)	32.60 (201/203)	21.24 (201/203)
MTAD-GAT + TTO (no prior)	41.90 (202/204)	19.22 (206/209)	639.05 (201/201)	7.59 (204/208)	1122.06 (201/201)
MTAD-GAT + TTO (with prior)	42.05 (202/204)	19.57 (205/208)	642.62 (201/201)	8.12 (204/208)	1122.80 (201/201)
TCN-AE + TTO (no prior)	40.73 (201/203)	48.75 (208/236)	35.18 (201/201)	40.84 (201/208)	21.25 (201/201)
TCN-AE + TTO (with prior)	41.58 (201/202)	49.36 (207/235)	38.39 (201/201)	41.44 (201/206)	22.50 (201/201)
TransformerAE + TTO (no prior)	18.03 (201/203)	55.13 (201/212)	68.75 (201/201)	41.25 (201/201)	24.39 (201/201)
TransformerAE + TTO (with prior)	18.30 (201/202)	56.18 (201/213)	71.91 (201/201)	42.62 (201/201)	26.41 (201/201)

Table 11: **Normalized runtime and stopping iterations on the full-channel test split.** Columns report the number of test windows used for each dataset. Benchmark entries are seconds per 1000 windows for the fastest one-shot row across point-, subsequence-, and block-trained checkpoints. TTO entries are seconds per 1000 windows; parentheses report median/p90 stopping iterations across optimized samples. Measurements use the same hardware, auto-selected batch sizes, warmup runs, and median-of-three timing protocol as Table 10.

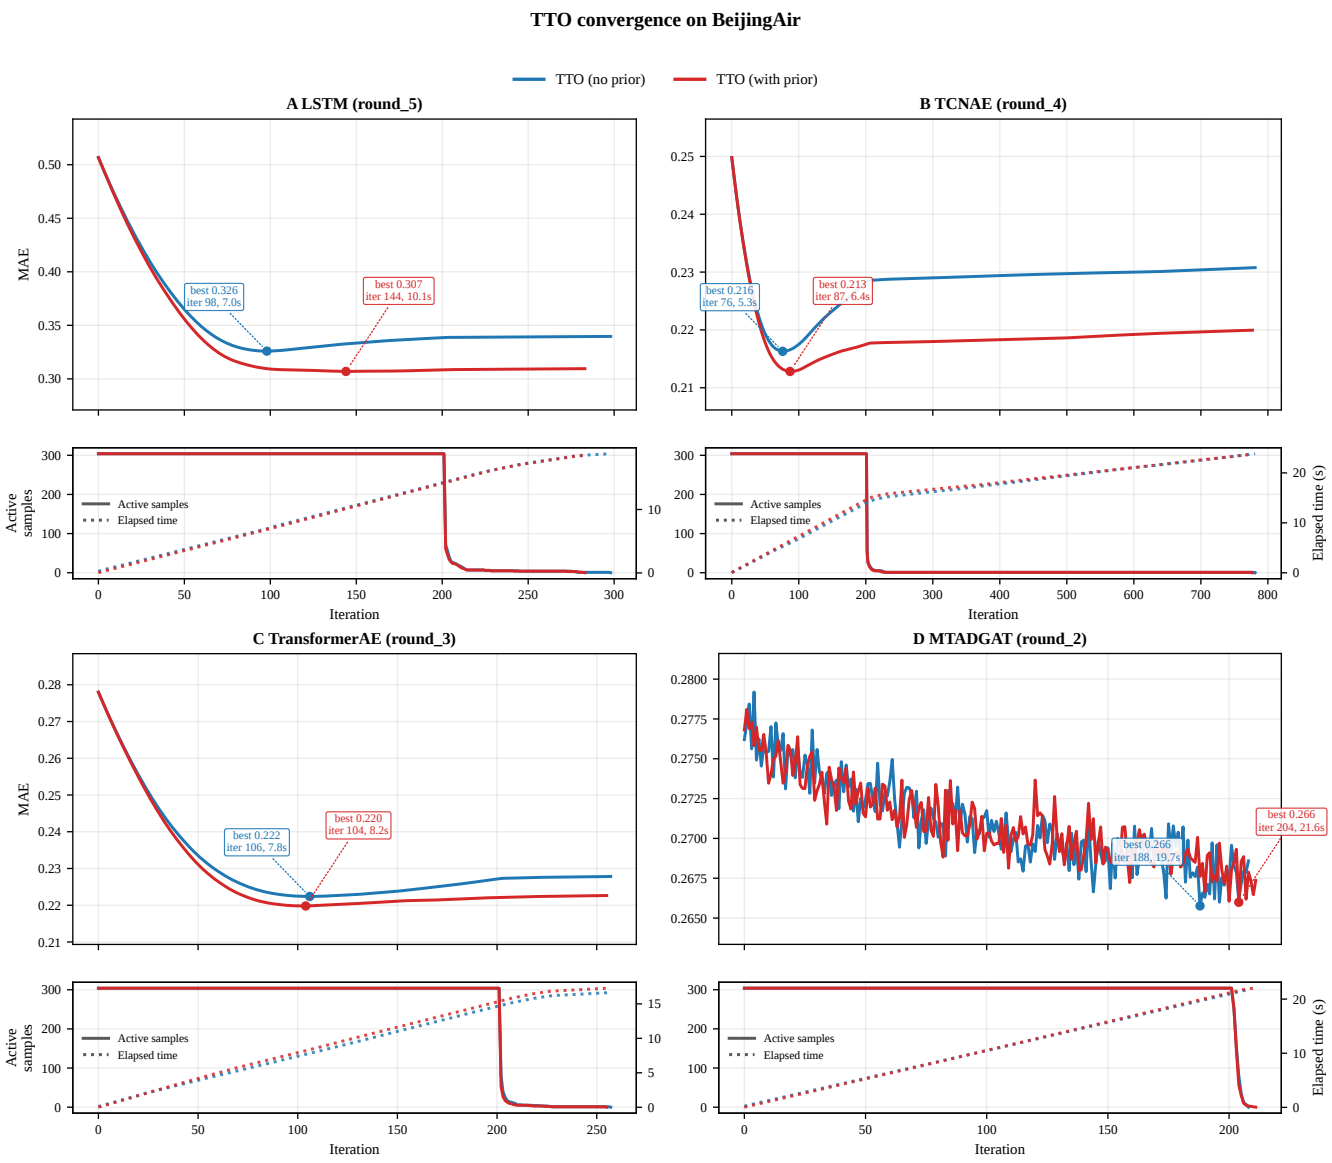

Figure 6: **Test-Time Optimization convergence on BeijingAir** Convergence behavior of test-time optimization (TTO) for four backbone architectures on the full-channel-missing setting. Each panel reports the tracked MAE over optimization iterations for TTO without prior regularization and TTO with prior regularization. Markers denote the best MAE reached by each variant for diagnostic visualization only, with annotations indicating the corresponding iteration and elapsed time. The lower subplots show the number of active samples during optimization, while dotted curves indicate elapsed wall-clock time.

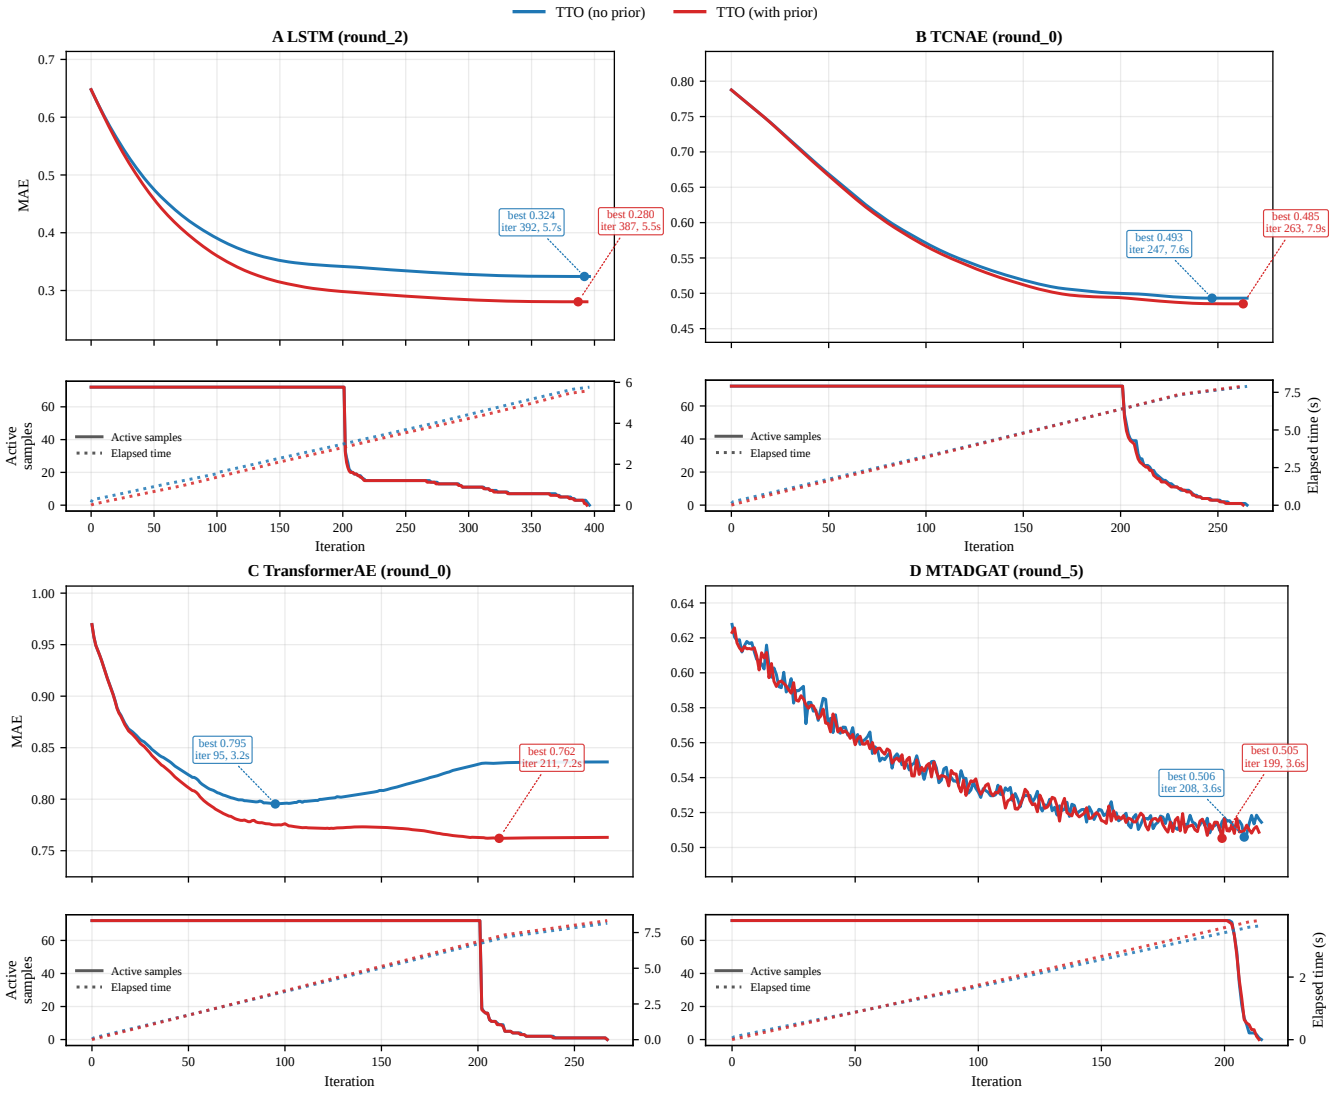

Figure 7: **Test-Time Optimization convergence on ETTh1** Convergence behavior of test-time optimization (TTO) for four backbone architectures on the full-channel-missing setting. Each panel reports the tracked MAE over optimization iterations for TTO without prior regularization and TTO with prior regularization. Markers denote the best MAE reached by each variant for diagnostic visualization only, with annotations indicating the corresponding iteration and elapsed time. The lower subplots show the number of active samples during optimization, while dotted curves indicate elapsed wall-clock time. For the two backbones LSTM and TransformerAE the prior helps convergence.

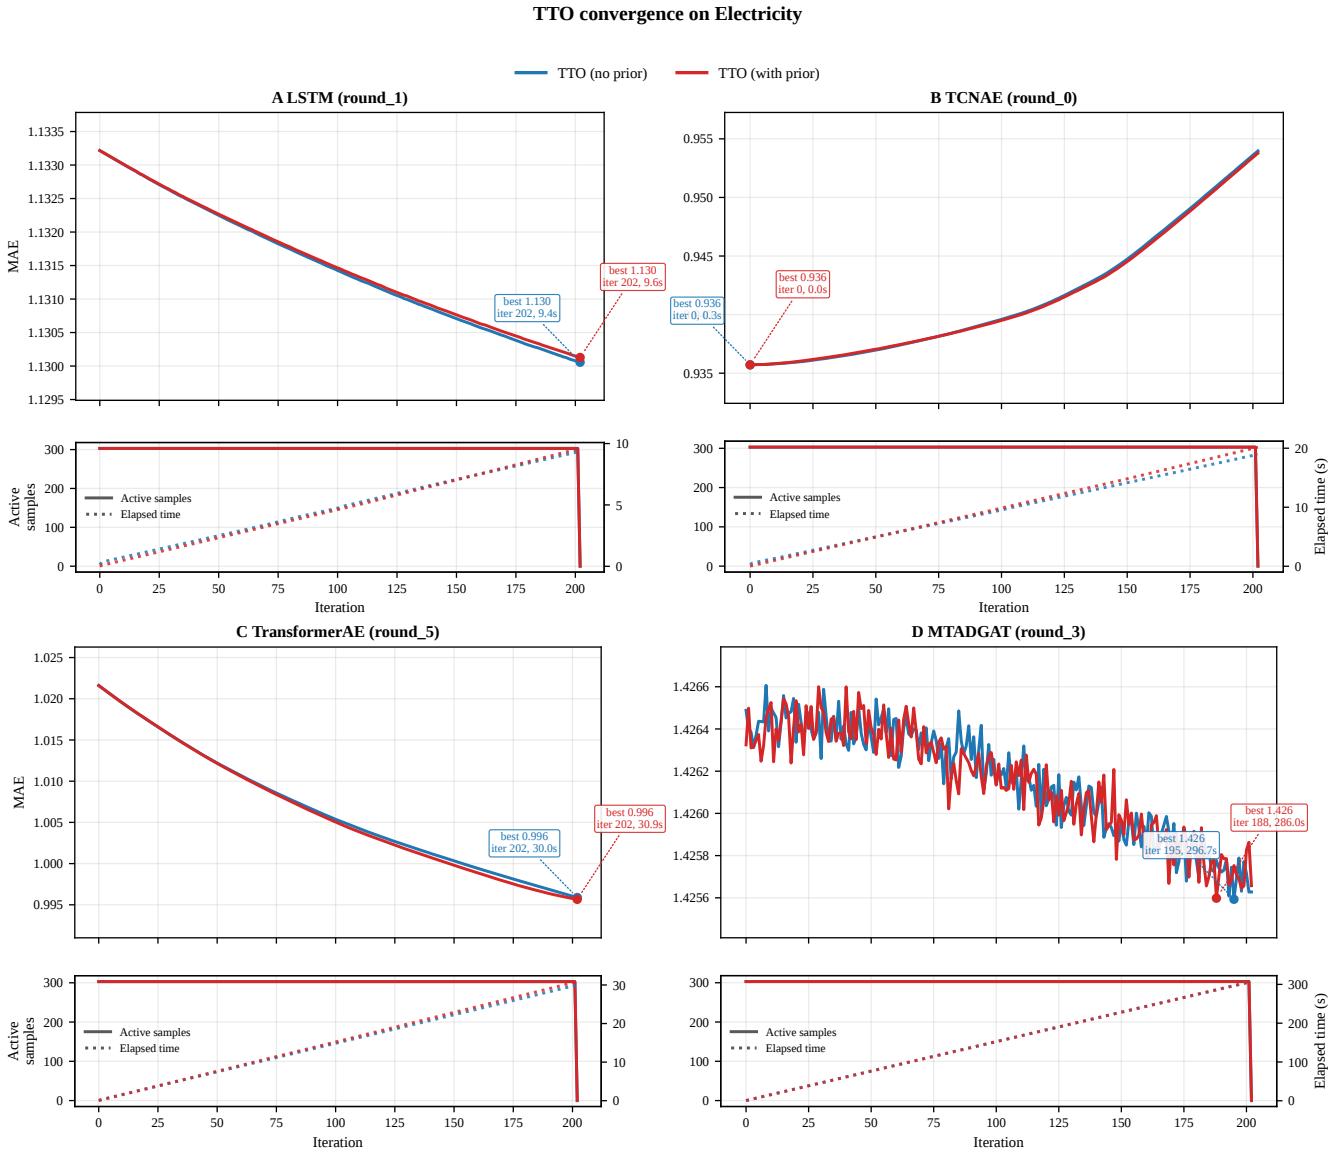

Figure 8: **Test-Time Optimization convergence on Electricity** Convergence behavior of test-time optimization (TTO) for four backbone architectures on the full-channel-missing setting. Each panel reports the tracked MAE over optimization iterations for TTO without prior regularization and TTO with prior regularization. Markers denote the best MAE reached by each variant for diagnostic visualization only, with annotations indicating the corresponding iteration and elapsed time. The lower subplots show the number of active samples during optimization, while dotted curves indicate elapsed wall-clock time. For the backbone TCN-AE TTO does not improve the imputation performance.

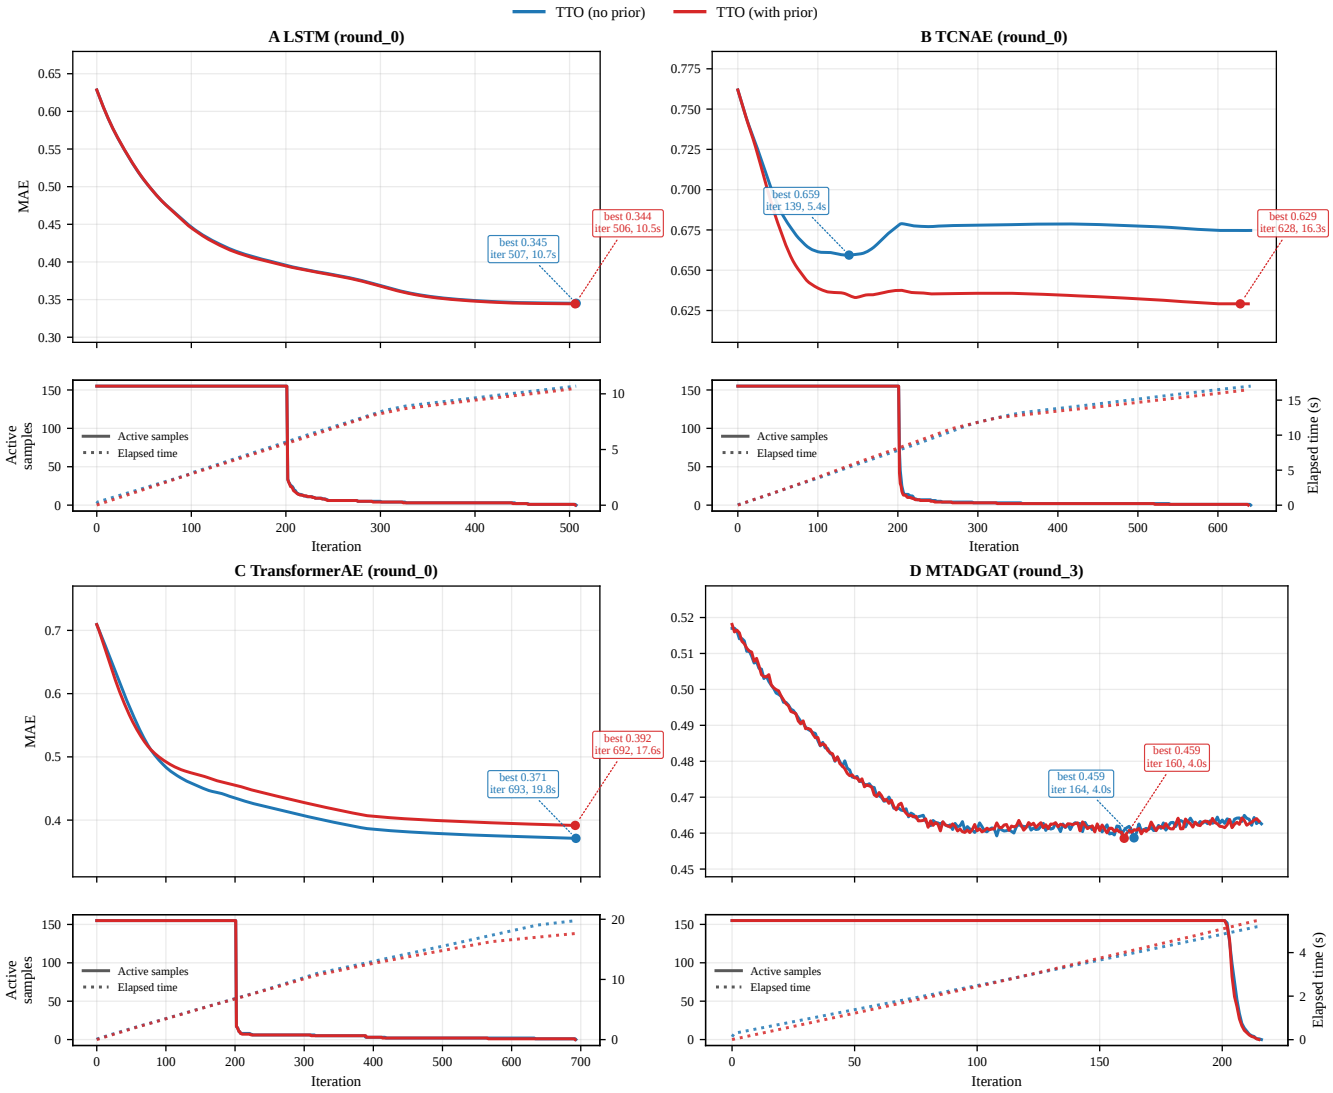

Figure 9: **Test-Time Optimization convergence on ItalyAir** Convergence behavior of test-time optimization (TTO) for four backbone architectures on the full-channel-missing setting. Each panel reports the tracked MAE over optimization iterations for TTO without prior regularization and TTO with prior regularization. Markers denote the best MAE reached by each variant for diagnostic visualization only, with annotations indicating the corresponding iteration and elapsed time. The lower subplots show the number of active samples during optimization, while dotted curves indicate elapsed wall-clock time. For the backbone TCN-AE the prior helps convergence.

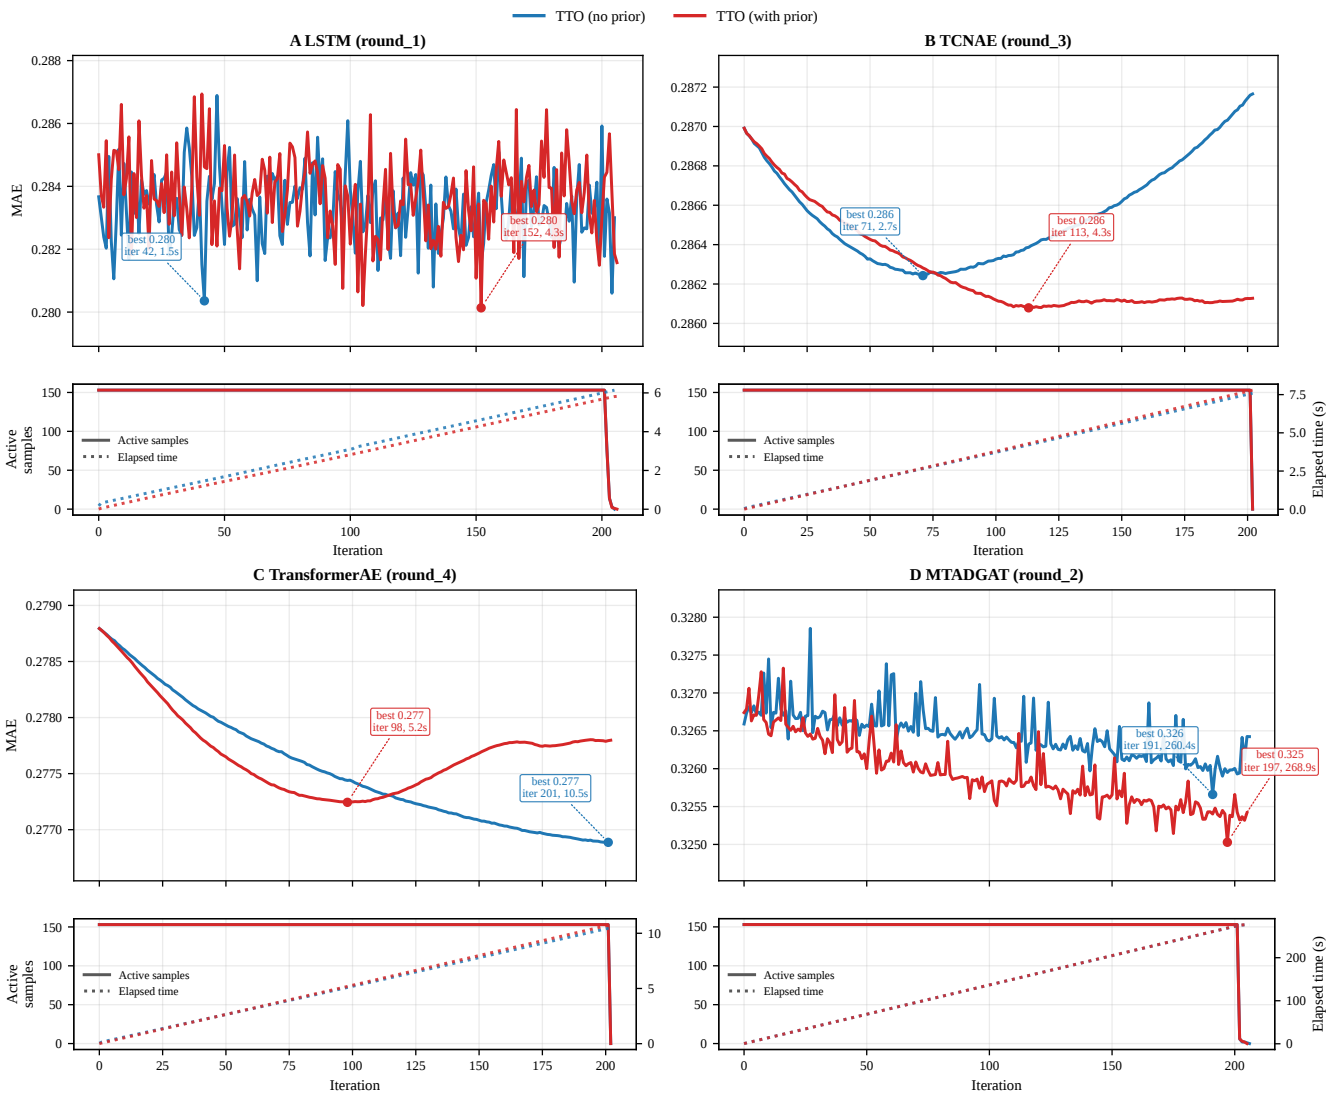

Figure 10: **Test-Time Optimization convergence on PeMS** Convergence behavior of test-time optimization (TTO) for four backbone architectures on the full-channel-missing setting. Each panel reports the tracked MAE over optimization iterations for TTO without prior regularization and TTO with prior regularization. Markers denote the best MAE reached by each variant for diagnostic visualization only, with annotations indicating the corresponding iteration and elapsed time. The lower subplots show the number of active samples during optimization, while dotted curves indicate elapsed wall-clock time. For the backbone TCN-AE the prior improves convergence, while for the backbone Transformer-AE it does not.