arXiv:2611.06910v1 [cs.AI] 4 Nov 2026

FrontierPhysics: Evaluating Agents for End-to-End Physics Research

Bingran You[1,2*], Wenjun Ke[3*], Ria Rosenauer[4*], Bo Liu[5,6*], Wenbo Chen[7†§], Min Chen[8†], Shuqi Xu[2†], Yue Ma[9], Jiechao Feng[2], Bo Chen[10], Xiangkai Sun[3], Jiarui Liu[2], Jinen Guo[10], Zeyu Zhou[11,12], Xu Huang[2], Zhuohan Li[13], Haoyu Li[14], Heguang Lin[15], Lijie Ding[16], Yadong Zeng[17], Kangqi Zhang[18], Xinge Liu[19], Hanwen Liu[20], Zhenghong Li[10], Alexander von Recum[6,21], Yuchen Zhu[12], Felix Schwarzfischer[21], Yizhou Zhao[22], Jicheng Wang[23], Binxu Li[24], Qian Yu[2], Chan Li[9], Cheng Dai[10], Di Wang[10], Haoyu Zhang[25], Jiankai Sun[6], Kaituo Feng[10], Mian Qin[26], Ruihong Yin[27], Yiqun Luo[22], Weiwei Wu[6], Shihong Pan[28], Tsung Yeh Hsieh[22], Xiao Huang[29], Yan Zhang[10], Sihan Li[30], Zirui Ren[31], Zengji Tu[31], Kyoung Whan Choe[32], Ke Sun[2], Vincenzo Galia[33], Vaisakhi Mishra[34], Jiashu He[26], Hongyang Wang[7§], Zhiheng Yang[35], Shubham Gaur[36], Yifan Gu[7§], Shuyi Wang[10], Zihao Zuo[14], Steven Dillmann[6], Yiyou Sun[2], Xinyang Han[2], Hanchung Lee[10], Molei Tao[12], Yue Zhao[20], Shilong Liu[24], Yuchen Wu[5], Xuandong Zhao[2], Xiangyi Li[1], Florian Brand[37], Gerbrand Ceder[2], Hartmut Haeffner[2]

Abstract

AI agents are playing an increasingly important role in physics research. However, existing physics benchmarks only evaluate fixed questions, isolated sub-tasks, or paper reproduction with given solutions, and rarely check research quality alongside correctness. We introduce FRONTIERPHYSICS, an end-to-end benchmark with 56 tasks made by 42 physicists based on their own research, covering nine domains across theoretical, computational, and experimental physics. Each task requires agents to review the literature, develop a research plan, conduct research and experiments, and produce structured outputs and scientific manuscripts. Combining deterministic verifiers with expert-authored rubrics, FRONTIERPHYSICS evaluates whether agents' research is correct and publishable. Across 22 models in 26 configurations, no configuration passes more than 46 of its 168 rollouts. Manual review of agent trajectories shows that frontier agents' main weaknesses are insufficiently rigorous reasoning and limited expert intuition, rather than gaps in physics knowledge or execution errors.

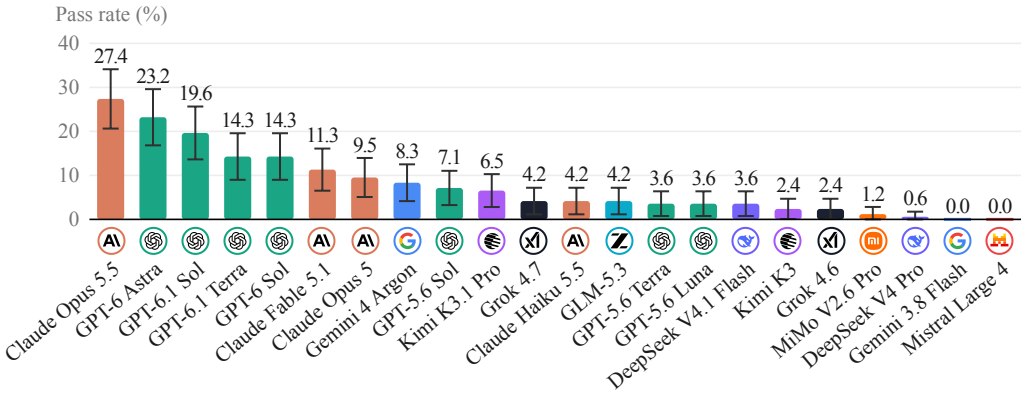

Figure 1: **FRONTIERPHYSICS LEADERBOARD.** Overall pass rate (Pass@1) under the Table 2 configurations; a pass requires passing all verifier tests and blocker rubrics. Harnesses and inference level: GPT (Codex, max), Claude (Claude Code, max), Grok (Grok Build, xhigh), Gemini (Antigravity, high), Kimi (Claude Code, max), DeepSeek (Claude Code, max), MiMo (Claude Code, max), GLM (Claude Code, max), Mistral (Claude Code, max). Error bars show approximate 95% confidence intervals.

[1]BenchFlow, [2]University of California, Berkeley, [3]California Institute of Technology, [4]Harvard University, [5]University of Washington, [6]Stanford University, [7]Amazon, [8]University of Pittsburgh, [9]University of California, San Diego, [10]Independent Researcher, [11]ZEISS, [12]Georgia Institute of Technology, [13]Lawrence Berkeley National Laboratory, [14]University of Michigan, Ann Arbor, [15]The Scripps Research Institute, [16]Oak Ridge National Laboratory, [17]University of Texas at Austin, [18]Purdue University, [19]University of Toronto, [20]University of Southern California, [21]Ludwig-Maximilians-Universität München, [22]Carnegie Mellon University, [23]University of California, Davis, [24]Princeton University, [25]Northeastern University, [26]University of Pennsylvania, [27]University of Minnesota, Twin Cities, [28]New York University, [29]Duke University, [30]Nanyang Technological University, [31]Pipeline Lab, [32]RLWRLD, [33]University of Palermo, [34]IBM, [35]University of Amsterdam, [36]University of California, Santa Cruz, [37]Prime Intellect [*]Equal contribution. [†]Core contribution. [§]Work conducted outside of the authors' roles at Amazon.

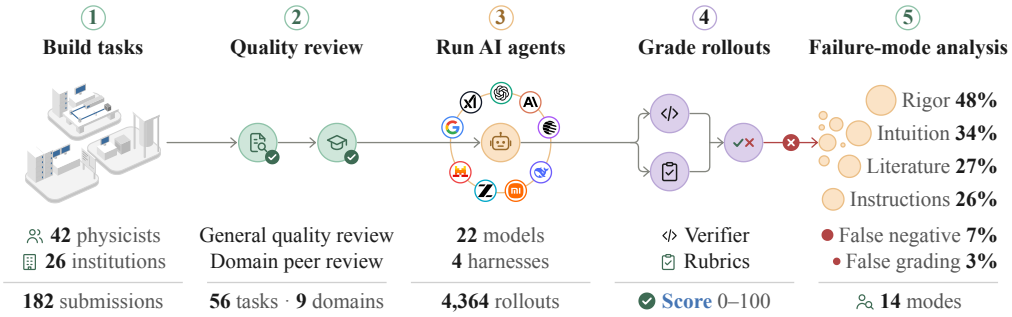

Figure 2: **Overview of FRONTIERPHYSICS.** (1) Physicists build tasks from their own research. (2) An admin team of physicists accepts those passing both reviews. (3) 4,364 rollouts cover 22 models and 4 harnesses, each is one trial of a task by a model in an agent harness. (4) A deterministic verifier and a reviewer agent with the task's rubric grade each rollout (Eq. 1). (5) The admin team and contributors manually analyze failure modes of the 168 GPT-6 Astra rollouts (Fig. 6, Table 4).

1 Introduction

Large language models (LLMs) are evolving from passive scientific assistants into agents that take on parts of the research process. When equipped with external tools, these agents can search and summarize literature, formulate hypotheses, write and debug scientific code, and produce technical reports. Recent systems have demonstrated tool-augmented reasoning in chemistry, iterative code-based experimentation, multi-agent hypothesis generation, scientific-software optimization, and even interaction with laboratory instruments [4, 6, 15, 17, 18, 20, 22, 29–31, 33, 37, 41]. These systems point toward agents that contribute to more complete scientific workflows and projects, not only isolated subtasks [9].

Evaluating progress toward this goal requires benchmarks that capture the open-ended, multi-stage nature of scientific research. Scientific benchmarks are accordingly shifting from textbook-level questions [16, 23, 35, 44, 48] toward expert-level closed-ended questions [8] and research-level problems [40]. Among existing research-level benchmarks, some assess scientific reasoning [10, 25, 34, 46, 50], while others evaluate the agents' ability to complete individual tasks within realistic scientific workflows [7, 19, 21, 28, 31, 43]. However, these benchmarks target isolated capabilities, and because many of their tasks are closed-ended with deterministic answers, models quickly saturate them, so the benchmarks lose signal. To move toward end-to-end research, some benchmarks formulate tasks as full-paper reproductions [12, 14, 36, 49]. Yet because agents are given the original publications, the research questions and solution structures are largely specified in advance, limiting their ability to evaluate open-ended scientific inquiry. More recent benchmarks address this limitation by converting published studies into research problems without revealing their solutions [24, 32, 45, 47]. Nevertheless, designing and evaluating such tasks involves many subtleties that require the expertise of the original authors: for example, how to write instructions that preserve as much open-endedness as possible while avoiding false negatives, and how to determine if an agent's failure reflects true scientific misunderstanding or a lack of expert intuition and experience. To our knowledge, no existing benchmark is fully agentic, built by physicists from their own research, graded on both paper quality and result correctness, and designed to address these subtleties through manual task review and failure-mode analysis (Table 1).

In this paper, we introduce **FRONTIERPHYSICS** (Fig. 1), a benchmark evaluating whether AI agents can conduct end-to-end research across theoretical, computational, and experimental physics. FRONTIERPHYSICS tasks are designed by active physics researchers (including Ph.D. students, postdocs, and research scientists) from research projects they have personally conducted. On average, each task author has 420 citations. To preserve open-endedness while minimizing false negatives and positives, each task goes through manual general review, domain-expert peer review, and failure-mode analysis. For each task (Fig. 2), an agent receives a high-level prompt that resembles the instructions provided by a principal investigator (PI) at the beginning of a research project. The agent must then carry out the research and produce both a scientific manuscript and numerical outputs that can be verified deterministically. Much of this process is open-ended: the agent must

investigate the topic, formulate a research plan, select appropriate methods, validate the results, and establish the significance of its findings. To evaluate these capabilities, FRONTIERPHYSICS combines a deterministic verifier (test scripts that check the numerical outputs) with rubric-based grading by reviewer agents.

2 FrontierPhysics

Figure 3 shows how a FRONTIERPHYSICS task is curated (a), structured (b), and run (c).

2.1 Task Design

Tasks are designed to reconstruct end-to-end research projects. Figure 3b shows an example task and Appendix A gives its full instruction and proposal. Some important questions to figure out or directions to explore might be given, but no explicit methodology or solution pathways will be specified. When necessary, input data and other supporting files are provided in the task environment (Fig. 3b) to enable the agent to configure and execute the relevant computational or experimental workflow.

Each task requires two types of final deliverables: deterministic artifacts, such as computed quantities, simulation outputs, or processed data, and a paper manuscript submitted as `paper.pdf`. The paper is expected to meet the standards of a peer-reviewed paper by clearly presenting the research context, methodology, results, and conclusions, while faithfully incorporating the evidence and computational artifacts produced during task execution.

The rubric-based assessment includes global rubrics and task-specific rubrics. The global rubrics evaluate behavior alignment including reward hacking [1, 38], hallucination and general research practices including literature synthesis, result validation, and the use of evidence to support scientific claims, whereas task-specific rubrics are developed by domain experts to assess the agent's understanding of the underlying physics, scientific presentation, and execution of the research task. Because rubric design requires task-specific expertise, we ask each contributor to design the rubrics

Table 1: **Comparison of scientific research benchmarks.** We compare scientific research benchmarks across the research authenticity, research components covered in the evaluation, and the evaluation methods. Meanings of labels: **Tasks/Phys.**: total/physics task counts in released task lists. **Own Research**: the tasks come from the contributors' own research. **No Reference Given**: there is no reference (e.g., the original paper) provided to the agent. **Literature Graded**: the quality of agents' literature review and synthesis is evaluated. **Write-up Graded**: agents write up scientific manuscripts and the manuscripts are graded. **Expert Rubric**: at least part of the evaluation protocol uses rubrics created by domain experts. **Deterministic Verifiers**: at least part of the evaluation protocol contains deterministic verifiers. **Failure Analysis**: domain experts manually check the trajectories and output papers of selected rollouts. ✓: documented; ○: partial; ×: no.

Benchmark	Tasks	Phys.	Research Authenticity		Research Coverage		Evaluation		
			Own Research	No Reference Given	Literature Graded	Write-up Graded	Expert Rubric	Deterministic Verifiers	Failure Analysis
Reasoning and QA									
CritPt [50]	71	71	✓	✓	×	×	×	✓	✓
FrontierScience-Research [43]	60	20	✓	✓	×	×	✓	×	×
Subtasks from scientific workflows									
Terminal-Bench-Science [42]	70	5	✓	✓	×	×	×	✓	○
Paper reproduction									
PRBench [36]	30	30	○	×	×	×	✓	×	○
ReplicationBench [49]	111	111	○	○	×	×	✓	✓	✓
QMP-Bench [12]	100	100	×	×	×	×	✓	✓	○
End-to-end research									
NatureBench [45]	90	13	×	✓	×	×	×	✓	○
PRL-Bench [32]	100	100	×	✓	×	×	✓	✓	○
ResearchClawBench [47]	40	4	×	✓	×	✓	✓	×	○
PhySciBench [24]	200	100	×	✓	○	×	✓	✓	○
FRONTIERPHYSICS	**56**	**56**	✓	✓	✓	✓	✓	✓	✓

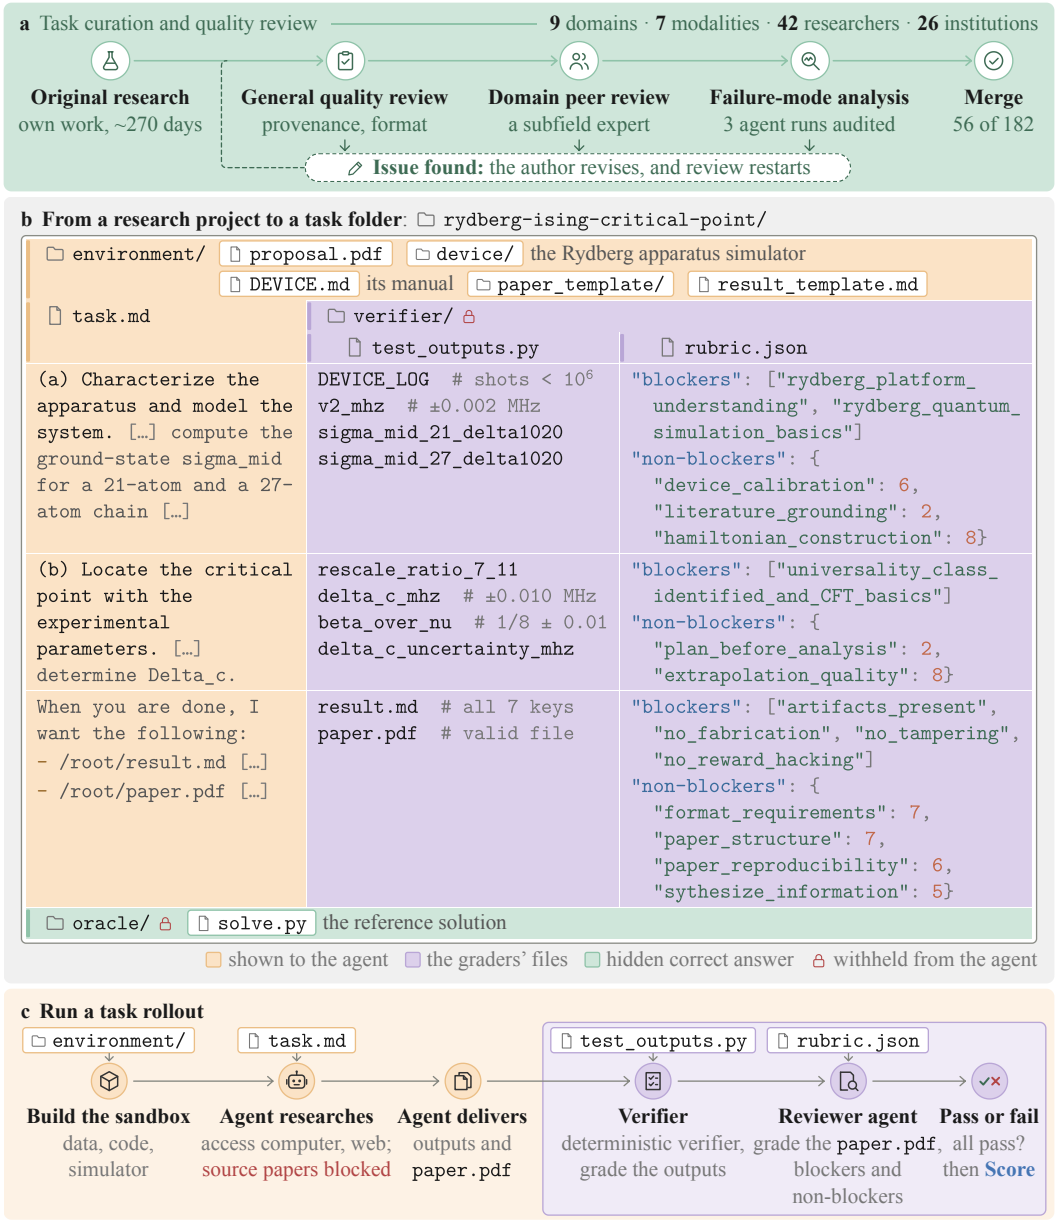

Figure 3: (a), FRONTIERPHYSICS tasks QA pipeline. (b), FRONTIERPHYSICS tasks structure, shown for the example task `rydberg-ising-critical-point`. (c), How to run a task rollout.

for their own task and research project. Depending on the task, the rubric criteria may assess whether the agent identifies and correctly uses key references, discovers the key scientific ideas through its investigation, considers realistic experimental constraints, or provided good motivations for its work in the paper.

2.2 Task Curation and Quality Review

General quality review. When a contributor creates a task submission, we first assess its authenticity and provenance (Fig. 3a): the contributor must provide supporting information (e.g., reference papers or patent links) to demonstrate that they have personally conducted the project. Contributors are asked to provide the project background, a timeline of the original research process, and the amount of time they spent completing it. We require the underlying project to have involved at least two weeks of research effort. We then check that `task.md` and `rubric.json` are human-readable and contain genuine human input, rejecting tasks that appear purely AI-generated.

Domain peer review. We aim to have every task reviewed by an expert in a closely related field. To scale this, we invite contributors to become domain reviewers once their own tasks pass the first-stage general quality check. Our domain reviewers have therefore all conducted research in their fields, and consist mostly of Ph.D. students (47%) and Ph.D. holders (37%). Each new reviewer is onboarded through a one-on-one meeting, and we audit the quality of their first review to align standards across reviewers. When a domain reviewer judges a task ready to merge, we perform a final check against our standards.

Each domain reviewer is provided with a domain-review checklist which covers three aspects: First, the reviewer verifies that `task.md` does not over-specify the problem: the prompt should not contain domain knowledge that a competent researcher is expected to know or figure out, so that the task remains as open as possible. Second, the reviewer confirms that `rubric.json` encodes real physics insight that only someone who has conducted the research would possess, rather than the low-signal criteria typical of AI-generated rubrics. Third, the reviewer examines failure modes after getting agents trajectories: contributors must perform at least three local runs with SOTA agents and an initial failure-mode analysis, and the reviewer inspects these traces for finding potential false positives and false negatives. Common issues include unjustified numerical tolerances that fail agents implementing a correct physical model or ambiguous instructions that lead to agents' misunderstanding.

When all checks pass, an admin team with six physics PhD researchers from diverse subfield backgrounds approves and merges the submission into the final task set. In total, 56 tasks were accepted from 42 physics researchers from 26 institutions worldwide in 9 different physics domains out of 182 initial task submissions. A summary can be seen in Fig. 4. For 18 tasks, the underlying research has been published in papers co-authored by the task contributors. These include four tasks based on papers published in Nature, with other journals including Nature Communications, Physical Review Letters, and Physical Review X. On average contributors report to have spent about 270 days per task for the original projects, and in total more than 8,000 hours of research.

2.3 Evaluation

In each task rollout (Fig. 3c), our evaluation scheme is designed to assess both the objective correctness of the agent's results and the overall quality of its research. To this end, we combine programmatic deterministic verifiers with an Agent-as-judge. The deterministic verifier checks structured outputs against reference values. After the agents finish running, agent trajectory files, verifier results, and output files including `paper.pdf` are recorded and saved. These files are then copied into another sandbox for a reviewer agent to review and grade the trajectory files and `paper.pdf` based on a `rubric.json`. We use an agent rather than a single LLM call as the judge because many rubric criteria, such as those on reward hacking, hallucinated results, and literature search, must be checked against the agent's trajectory, code, and output files rather than `paper.pdf` alone. All rollouts are graded by the same reviewer agent, GPT-5.6 Sol (xhigh) in Codex CLI. Section 3.3 checks its grades against an expert audit (Table 5, Fig. 6) and on the authors' own papers (Table 11). Regrading rollouts of the first evaluation round with Claude Opus 5.5 in Claude Code CLI as the reviewer agent gives the same pass/fail decision for almost all GPT-6 Astra rollouts, but the Opus 5.5 reviewer is more lenient on rollouts from its own model family (Table 6). Workspace files and trajectory files of the reviewer agents are collected the same way to HuggingFace datasets for later analysis.

Rubric-based grading has been widely adopted in benchmarks involving open-ended exploration or qualitative dimensions of agent performance [3, 11, 13, 36, 39]. Whereas deterministic verification establishes whether an output is correct, rubric-based evaluation assesses more aspects, for example whether the agent conducted and communicated the research well. Inspired by the rubric design of FrontierCode [11], we divide our rubric criteria into *blockers* and *non-blockers*. Blockers are binary

gates that represent requirements essential to the scientific validity of a task. Thus, an agent must satisfy every blocker as well as the deterministic verifier to pass. Non-blockers capture qualitative aspects of the work. Each non-blocker is scored on a three-level scale: 0 indicates that the requirement is absent or incorrect, 1 indicates partial satisfaction, and 2 indicates complete satisfaction. The precise anchors for these levels are defined separately for each criterion. The task author also gives every criterion an integer *weight* from 1 to 10 for its relative importance. Blocker weights are retained as rubric metadata but do not contribute to the numerical quality score. The final score of a task:

$$S_{\text{final}} = \begin{cases} \dfrac{\sum_{i \in \mathcal{N}} w_i s_i}{2 \sum_{i \in \mathcal{N}} w_i} \times 100, & \text{if } D = 1 \text{ and } b_j = 1 \ \ \forall j \in \mathcal{B}, \\ 0, & \text{otherwise,} \end{cases} \tag{1}$$

Here, $\mathcal{N}$ and $\mathcal{B}$ denote non-blockers and blockers respectively, $s_i \in \{0, 1, 2\}$, w_i is the criterion weight, D is the deterministic verifier result, and b_j is the blocker result.

$S_{\text{final}} = 0$ indicates that the results are not publishable, while the value of S_{final} when $S_{\text{final}} \neq 0$ measures the quality of the publication. High scores on non-blockers cannot compensate for an incorrect deterministic result or a failed blocker. Conversely, an agent may pass all deterministic checks while producing a poorly motivated, inadequately supported, or badly communicated manuscript. Such a submission passes the verifiers and blockers but receives a low S_{final}.

3 Experiments

We evaluate 22 frontier models in four agent harnesses. Each model attempts all 56 tasks three times at its highest available inference effort, and four GPT models are also run at xhigh.

3.1 Experimental Setup

As in SkillsBench [27] and ClawsBench [26] we use BenchFlow CLI [5], a research infra command line tool for creating RL environments, post-training, and evals, to run experiments as well as track agent trajectories, verifier, and rubrics results. We use Codex CLI for GPT models, Claude Code CLI for Claude models, Grok build for Grok models, and Antigravity CLI for Gemini models. The other seven models, DeepSeek V4 Pro, DeepSeek V4.1 Flash, GLM-5.3, Kimi K3, Kimi K3.1 Pro, MiMo V2.6 Pro and Mistral Large 4, also run in Claude Code CLI. Each agent runtime is started in parallel cloud sandboxes. For tasks with resource requirements under 4 vCPUs, 16GB RAM, and 10GB storage, we use Daytona as a sandbox provider. For tasks with higher resource requirements, we use Docker in Azure VMs with Lambda API access so the model can call GPU resources as needed.

All results are for FRONTIERPHYSICS v0.1, evaluated on all 56 tasks. Where a task changed after the first evaluation round in a way that could alter a result, the affected rollouts were re-run or regraded; the other changes were judged unable to alter a result.

Each rollout has a uniform 10-hour timeout cap. Endings caused by the model, such as reaching the budget or exhausting the sandbox's memory, are final and stay in the denominator; endings caused by the infrastructure, such as a lost sandbox or provider errors that outlast retries, are re-run. Table 7 counts the slots left without a valid result.

3.2 Results

The result (Fig. 1) indicates that the overall pass rates are below 30% for all agents tested. Claude Opus 5.5 has the highest pass rate, 27.4%. Table 2 provides more detailed information on the different runs and scores.

Figure 5 plots the pass rate versus the mean cost and time per rollout. Pass rate generally rises with cost. However, for some Claude models, the pass rate decreases as cost increases, in the order Opus 5.5, Fable 5.1, and Opus 5. In addition, GPT models are generally more cost-efficient. Figure 4 gives the pass rate per physics domain of Opus 5.5 and GPT-6 Astra; with few tasks per domain, the differences between domains are uncertain. Appendix D gives per-task results.

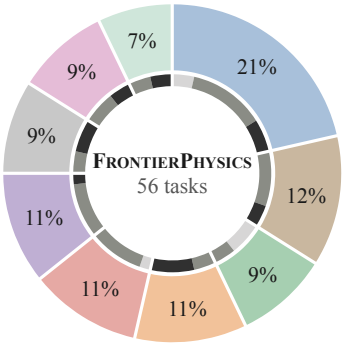

Physics domain	Tasks n	Opus 5.5	GPT-6 Astra
Condensed Matter Physics	12	33.3	25.0
AMO Physics	7	9.5	14.3
Quantum Information	5	26.7	0.0
Astrophysics	6	22.2	33.3
Fluid Dynamics	6	50.0	61.1
Statistical Physics	6	11.1	5.6
Photonics & Imaging	5	53.3	33.3
High Energy Physics	5	33.3	6.7
Chemical Physics	4	0.0	25.0
Research mode			
Theoretical	6	11.1	5.6
Computational	33	28.3	25.3
Experimental	17	31.4	25.5

Pass rate (%) 0 70

Figure 4: Task composition and overall pass rate (%) by primary physics domain across all 56 tasks. Donut labels give task shares. n gives task counts.

Figure 5 and Appendix G compare agent time, steps and output tokens with the pass rate. Table 10 gives the mean, median and 95th percentile of tool calls, agent time and cost per rollout for each configuration, and the mean and median of output tokens.

3.3 Assessing Verifier and Rubric Reliability

Verifier and rubric defects may remain after review. One major source of defects is adversarial verifier or rubric engineering designed to suppress task pass rates, such that even correct solutions produced by human experts may fail. To mitigate this, we have a three stage process:

First, we ask every task contributor to do a failure-mode report of the three GPT-6 Astra rollouts, as this is the best performing model with single agent harness when the analysis started, check the agent trajectories and output files, and explain where the agent failed and why. Second, our admin team reviewers re-read these contributors' failure-mode reports and double check them against the agent trajectories and output files. After reviewing around 60 rollouts, reviewers consolidated the recurring failures into the taxonomy in Table 4. False negatives, where agents deliver the correct results but failed by verifier or reviewer agent, are found during this process and are labeled in Appendix E. As highlighted in Table 5, rollouts with exclusively false failures lead to an underestimated pass rate. In total, nine of the 168 rollouts incorrectly failed. This coincides with the recent research [2], in which domain experts regrade some leading benchmarks and find that many failures are due to grader errors, incorrect reference solutions, and ambiguous or underspecified questions. This suggests that finding false negatives takes both domain expertise and time, but is crucial for maintaining high benchmark quality.

The non-blocking rubric scores are harder to assess because they involve more subjective judgments. Third, as a sanity check of the rubric quality we asked authors to include their own papers or manuscripts as reference materials in the oracle when possible. Thirty-three tasks included these materials. For those with available reference, we selected the manuscript-level rubric criteria used for grading these materials, including blocker and non-blocker scoring criteria against the authors' reference work. We then use the same configuration as main experiment to run reviewer agents in sandboxes to grade these files with rubrics. The authors' own materials score 96.2% on average under corresponding rubrics (Table 11), showing good alignment between rubrics grading with the human expert materials.

3.4 Failure Mode Analysis

To find why agents fail and how well the agent score approximates what we want to measure, we follow the protocol of Section 3.3. Figure 6 summarizes the earlier matched-cohort annotations for all 168 single agent GPT-6 Astra rollouts. Table 5 reports the final review-sheet labels. More detailed

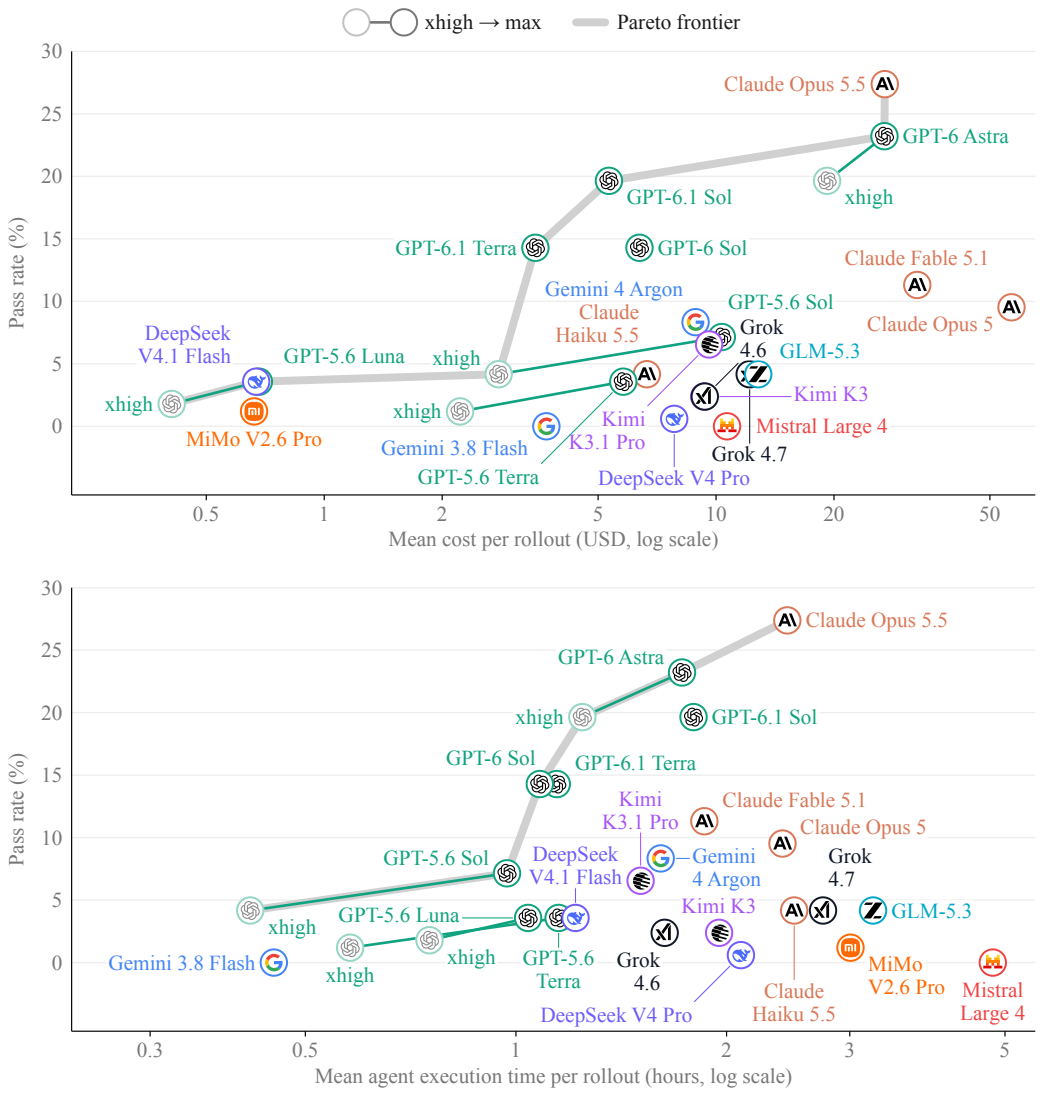

Figure 5: Mean reported inference cost (top) and agent execution time (bottom) per rollout versus overall pass rate (%). Time excludes environment setup, verification and rubric review.

information and an explanation of the failure modes can be found in Table 4. This process found design issues in some tasks. We fixed them and re-ran or regraded the affected results (Section 3.1). The pass rate of GPT-6 Astra (xhigh) after the re-runs (Table 2) agrees well with the manually corrected pass rate in Figure 6.

Figure 6 shows that agents mostly do not fail because of physics mistakes (e.g., a wrong calculation, misunderstanding of the physics concepts or lacking of physics knowledge), contrary to what one might expect from a physics benchmark. Instead, they often struggle to reason rigorously (48%), apply good physical intuition when making decisions (34%), and synthesize the literature (27%). Typical rigorous thinking failures are: circular self-verification, where the agent's own validation method depends on its wrong assumption, so it can never find out its mistakes; and edge cases, where the agent fails to consider all the edge cases of its method or code, so that it fails at some hidden tests. The third most frequent failure is in the literature. In many cases, agents fail to find the key literature for the research (as identified by the task author), struggle to synthesize information from its literature search and provide a motivation for its project, or struggle to support its argument in the manuscript with literature. Missing physics intuition is another major failure mode. Typical examples of this failure mode include: agents fail to determine the major error source of an ex-

Table 2: Detailed Frontier-model performance across all 56 tasks. For the experiment we use the highest available inference effort by default, where configurable. GPT models also appear at xhigh. The Score column reports the score defined in Equation 1, averaged over the three rollouts for each task and then across all tasks. Deterministic pass requires all verifier tests to pass and overall pass additionally requires all blocking rubric criteria to pass. Pass@1 gives the average pass rate over the three rollouts, while Pass@3 gives the percentage of tasks with at least one passing rollout. Uncertainties are approximate 95% binomial confidence intervals over all rollouts, as in Figure 1. Table 8 gives 95% confidence intervals from a bootstrap over tasks, which account for the correlation of the three trials of a task. Rows are ordered by overall Pass@1, then by deterministic Pass@1 and Score.

		Deterministic pass (%)		Overall pass (%)	
Model	**Score**	Pass@1	Pass@3	Pass@1	Pass@3
Claude Opus 5.5 Claude Code	26.0 ± 6.4	45.8 ± 7.5	62.5	**27.4** ± 6.7	44.6
GPT-6 Astra (max) Codex	21.8 ± 6.0	39.3 ± 7.4	51.8	**23.2** ± 6.4	30.4
GPT-6 Astra (xhigh) Codex	18.0 ± 5.5	38.7 ± 7.4	51.8	**19.6** ± 6.0	30.4
GPT-6.1 Sol Codex	18.4 ± 5.7	33.3 ± 7.1	50.0	**19.6** ± 6.0	30.4
GPT-6.1 Terra Codex	12.6 ± 4.7	28.6 ± 6.8	39.3	**14.3** ± 5.3	21.4
GPT-6 Sol Codex	12.5 ± 4.7	27.4 ± 6.7	44.6	**14.3** ± 5.3	28.6
Claude Fable 5.1 Claude Code	10.2 ± 4.3	31.0 ± 7.0	44.6	**11.3** ± 4.8	21.4
Claude Opus 5 Claude Code	8.6 ± 4.0	20.2 ± 6.1	30.4	**9.5** ± 4.4	17.9
Gemini 4 Argon Antigravity	6.5 ± 3.3	31.5 ± 7.0	46.4	**8.3** ± 4.2	17.9
GPT-5.6 Sol (max) Codex	6.2 ± 3.4	19.6 ± 6.0	28.6	**7.1** ± 3.9	12.5
Kimi K3.1 Pro Claude Code	6.1 ± 3.5	15.5 ± 5.5	30.4	**6.5** ± 3.7	12.5
Grok 4.7 Grok Build	3.8 ± 2.8	16.7 ± 5.6	35.7	**4.2** ± 3.0	10.7
Claude Haiku 5.5 Claude Code	3.6 ± 2.6	15.5 ± 5.5	30.4	**4.2** ± 3.0	10.7
GPT-5.6 Sol (xhigh) Codex	3.5 ± 2.6	15.5 ± 5.5	26.8	**4.2** ± 3.0	8.9
GLM-5.3 Claude Code	3.6 ± 2.6	7.1 ± 3.9	14.3	**4.2** ± 3.0	8.9
GPT-5.6 Terra (max) Codex	2.9 ± 2.3	16.1 ± 5.6	25.0	**3.6** ± 2.8	7.1
GPT-5.6 Luna (max) Codex	2.9 ± 2.3	12.5 ± 5.0	21.4	**3.6** ± 2.8	3.6
DeepSeek V4.1 Flash Claude Code	3.3 ± 2.6	10.7 ± 4.7	14.3	**3.6** ± 2.8	5.4
Kimi K3 Claude Code	1.8 ± 1.7	11.3 ± 4.8	21.4	**2.4** ± 2.3	5.4
Grok 4.6 Grok Build	2.2 ± 2.1	10.7 ± 4.7	21.4	**2.4** ± 2.3	3.6
GPT-5.6 Luna (xhigh) Codex	1.4 ± 1.5	8.9 ± 4.3	16.1	**1.8** ± 2.0	3.6
MiMo V2.6 Pro Claude Code	1.0 ± 1.4	10.1 ± 4.6	19.6	**1.2** ± 1.6	3.6
GPT-5.6 Terra (xhigh) Codex	0.9 ± 1.3	9.5 ± 4.4	16.1	**1.2** ± 1.6	1.8
DeepSeek V4 Pro Claude Code	0.3 ± 0.6	3.0 ± 2.6	5.4	**0.6** ± 1.2	1.8
Gemini 3.8 Flash Antigravity	0.0 ± 0.0	7.7 ± 4.0	14.3	**0.0** ± 0.0	0.0
Mistral Large 4 Claude Code	0.0 ± 0.0	5.4 ± 3.4	12.5	**0.0** ± 0.0	0.0

periment, which is not because they do not understand it, but because they do not have the expert intuition to identify what is important; when debugging, agents cannot realize the root of the problem while keeping tuning the wrong knob. Besides the top three failure modes, instruction following accounts for 26% of the failures. It could be agents missing some requirements in the instructions or misunderstanding the instructions. In many cases, the task authors believe the meaning should be clear to human domain experts, while the agent cannot get the explicit/implicit message.

The current failure mode analysis focuses on the GPT-6 Astra model (see Section 3.3) since it is challenging for tasks authors to manually read trajectory files with subagents involved. Other mod-

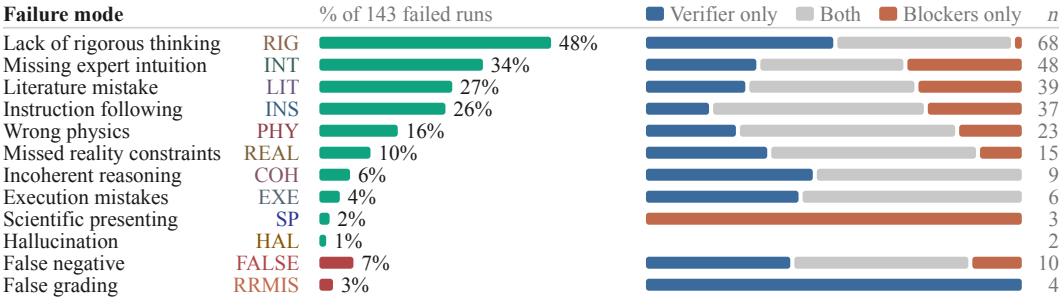

Figure 6: Failure mode analysis of GPT-6 Astra xhigh across 56 tasks (3 rollouts for each task). Counts n and percentages ($n/143$) use the final review labels in Table 5. Each rollout may have multiple tags. Verifier/rubric proportions are retained from the original recorded assessment. Correcting the eight FALSE and one RRMIS outcomes changes the pass rate from 14.9% (25/168) to 20.2% (34/168).

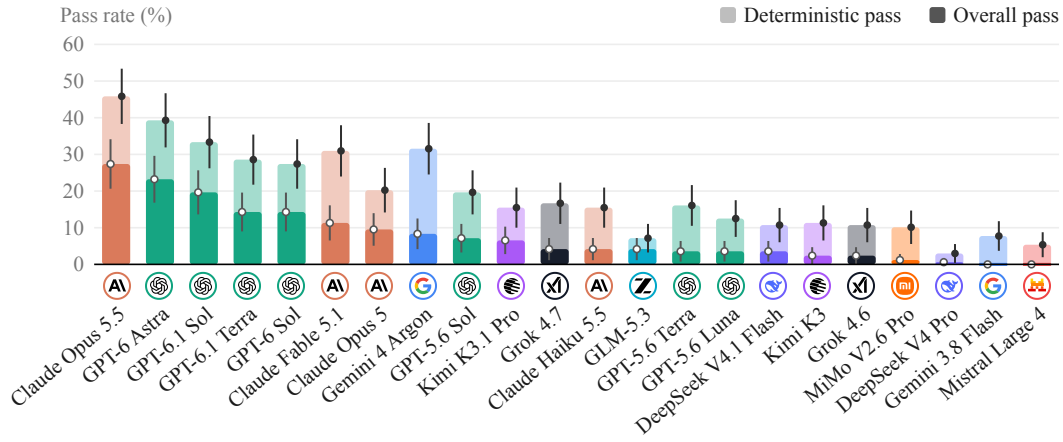

Figure 7: Deterministic and overall pass rates (Pass@1, %) of the models in Fig. 1. Column height: deterministic pass rate (all verifier tests pass). Solid part: overall pass rate (blocker rubrics also pass). Error bars: approximate 95% confidence intervals (dot: deterministic; circle: overall).

els, or other model-harness pairs, may fail in different ways, including ways outside this taxonomy. Also because of this, we can only accurately tag the false negatives of GPT-6 Astra's rollouts and therefore correct the final pass rate of GPT-6 Astra.

4 Conclusion

We have introduced FRONTIERPHYSICS, a benchmark dataset with 56 tasks evaluating the ability of AI agents to conduct physics research across nine domains. FRONTIERPHYSICS shows end-to-end physics research remains challenging to current models, given that Opus 5.5 in Claude Code CLI, the configuration with the highest pass rate, passes only 46 of its 168 rollouts. Based on the human experts auditing of 168 SOTA agents rollouts trajectory and output files, we build a failure mode taxonomy and label all these 168 rollouts manually. From the taxonomy distribution, we conclude that lack of rigorous thinking (for example, circular self-verification) followed by missing expert intuition, instead of insufficient physics knowledge or execution mistakes, are the main weaknesses of frontier LLM agents.

Capable physics research agents should go beyond specified tasks to explore scientific problems independently, from problem definition to novel findings, with only high-level guidance. Training and evaluating such agents will require more open-ended research tasks and flexible evaluation frameworks that can assess exploratory reasoning, partial progress, and unexpected findings.

References

[1] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete Problems in AI Safety. *arXiv preprint arXiv:1606.06565*, 2016.

[2] Ali Ansari, Haoran Sun, Andy Zeyi Liu, Mark Jabbour, et al. How good are frontier models at physics? expert re-grading reveals broken evaluations and near-saturation of leading benchmarks. *arXiv preprint arXiv:2609.13009*, 2026.

[3] Rahul K. Arora, Jason Wei, Rebecca Soskin Hicks, Preston Bowman, Joaquin Quiñonero-Candela, Foivos Tsimpourlas, Michael Sharman, Meghan Shah, Andrea Vallone, Alex Beutel, Johannes Heidecke, and Karan Singhal. HealthBench: Evaluating Large Language Models Towards Improved Human Health. *arXiv preprint arXiv:2505.08775*, 2025.

[4] Eser Aygün, Anastasiya Belyaeva, Gheorghe Comanici, Marc Coram, Hao Cui, Jake Garrison, Renee Johnston, Anton Kast, Cory Y. McLean, Peter Norgaard, Zahra Shamsi, David Smalling, James Thompson, Subhashini Venugopalan, Brian P. Williams, Chujun He, Sarah Martinson, Martyna Plomecka, Lai Wei, Yuchen Zhou, Qian-Ze Zhu, Matthew Abraham, Erica Brand, Anna Bulanova, Jeffrey A. Cardille, Chris Co, Scott Ellsworth, Grace Joseph, Malcolm Kane, Ryan Krueger, Johan Kartiwa, Dan Liebling, Jan-Matthis Lueckmann, Paul Raccuglia, Xuefei Julie Wang, Katherine Chou, James Manyika, Yossi Matias, John C. Platt, Lizzie Dorfman, Shibl Mourad, and Michael P. Brenner. An AI system to help scientists write expert-level empirical software. *Nature*, 654:909–916, 2026. doi: 10.1038/s41586-026-10658-6.

[5] BenchFlow team. BenchFlow: Framework for RL environments for LLM agents. Software repository, 2026. URL https://github.com/benchflow-ai/benchflow.

[6] Daniil A. Boiko, Robert MacKnight, Ben Kline, and Gabe Gomes. Autonomous chemical research with large language models. *Nature*, 624:570–578, 2023. doi: 10.1038/s41586-023-06792-0.

[7] Yadi Cao, Sicheng Lai, Jiahe Huang, Yang Zhang, Zach Lawrence, Rohan Bhakta, Izzy F. Thomas, Mingyun Cao, Chung-Hao Tsai, Zihao Zhou, Yidong Zhao, Hao Liu, Alessandro Marinoni, Alexey Arefiev, and Rose Yu. SimulCost: A Cost-Aware Benchmark and Toolkit for Automating Physics Simulations with LLMs. *arXiv preprint arXiv:2603.20253*, 2026.

[8] Center for AI Safety, Scale AI, and HLE Contributors Consortium. A benchmark of expert-level academic questions to assess AI capabilities. *Nature*, 649:1139–1146, 2026. doi: 10.1038/s41586-025-09962-4.

[9] Meng Chu, Xuan Billy Zhang, Kevin Qinghong Lin, Lingdong Kong, Jize Zhang, Teng Tu, Weijian Ma, Ziqi Huang, Senqiao Yang, Wei Huang, Yeying Jin, Zhefan Rao, Jinhui Ye, Xinyu Lin, Xichen Zhang, Qisheng Hu, Shuai Yang, Leyang Shen, Wei Chow, Yifei Dong, Fengyi Wu, Quanyu Long, Bin Xia, Shaozuo Yu, Mingkang Zhu, Wenhu Zhang, Jiehui Huang, Haokun Gui, Runyi Li, Chenyu Tang, Dong Huang, Xuhang Chen, Rui Liu, Chengzu Li, Shiyi Du, Xu Huang, Haoxuan Che, Long Chen, Qifeng Chen, Wenya Wang, Wenxuan Zhang, Xiaojuan Qi, Yang Deng, Yanwei Li, Mike Zheng Shou, Zhi-Qi Cheng, See-Kiong Ng, Ziwei Liu, Philip Torr, and Jiaya Jia. Agentic World Modeling: Foundations, Capabilities, Laws, and Beyond. *arXiv preprint arXiv:2604.22748*, 2026.

[10] Daniel J. H. Chung, Zhiqi Gao, Yurii Kvasiuk, Tianyi Li, Moritz Münchmeyer, Maja Rudolph, Frederic Sala, and Sai Chaitanya Tadepalli. Theoretical Physics Benchmark (TPBench) – a Dataset and Study of AI Reasoning Capabilities in Theoretical Physics. *arXiv preprint arXiv:2502.15815*, 2025.

[11] Cognition. FrontierCode 1.1. Official web page, 2026. URL https://cognition.com/blog/frontier-code-1.1.

[12] Ken Deng, Xiangfei Wang, Guijing Duan, Chen Mo, Junkun Huang, Runqing Zhang, Ling Qian, Zhiguo Huang, Jize Han, and Di Luo. QMP-Bench: A Research-Level Benchmark for Autonomous End-to-End Quantum Many-Body Simulations. Accepted paper under non-archival workshop policy, 2026. URL https://openreview.net/forum?id=9IdnFv2TmQ.

[13] Zhiyuan Fan, Weinong Wang, Xing Wu, and Debing Zhang. SedarEval: Automated Evaluation using Self-Adaptive Rubrics. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 16916–16930. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.findings-emnlp.984.

[14] Darius A. Faroughy, Sofia Palacios Schweitzer, Ian Pang, Siddharth Mishra-Sharma, and David Shih. Collider-Bench: Benchmarking AI Agents with Particle Physics Analysis Reproduction. *arXiv preprint arXiv:2605.13950*, 2026.

[15] Yuxing Fei, Bernardus Rendy, Xiaochen Yang, Junhee Woo, Xu Huang, Chang Li, Shilong Wang, David Milsted, Yan Zeng, and Gerbrand Ceder. Agentic LLM Reasoning in a Self-Driving Laboratory for Air-Sensitive Lithium Halide Spinel Conductors. *arXiv preprint arXiv:2604.11957*, 2026.

[16] Kaiyue Feng, Yilun Zhao, Yixin Liu, Tianyu Yang, Chen Zhao, John Sous, and Arman Cohan. Physics: Benchmarking Foundation Models on University-Level Physics Problem Solving. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 11717–11743, 2025. doi: 10.18653/v1/2025.findings-acl.610.

[17] Juraj Gottweis, Wei-Hung Weng, Alexander Daryin, Tao Tu, Petar Sirkovic, Artiom Myaskovsky, Grzegorz Glowaty, Felix Weissenberger, Alessio Orlandi, Dan Popovici, Anil Palepu, Keran Rong, Ryutaro Tanno, Khaled Saab, Fan Zhang, Jacob Blum, Andrew Carroll, Kavita Kulkarni, Nenad Tomašev, Dina Zverinski, Ivor Rendulic, Elahe Vedadi, Florian Hasler, Luka Rimanic, Marina Boia, Ivan Budiselic, Ben Feinstein, Mathias Bellaiche, Tom Sheffer, Jan Freyberg, Jeremy Ratcliff, Ottavia Bertolli, Katherine Chou, Avinatan Hassidim, Burak Gokturk, Amin Vahdat, Yuan Guan, Vikram Dhillon, Eeshit Dhaval Vaishnav, Byron Lee, Tiago R. D. Costa, José R. Penadés, Gary Peltz, Yossi Matias, James Manyika, Demis Hassabis, Yunhan Xu, Pushmeet Kohli, Annalisa Pawlosky, Alan Karthikesalingam, and Vivek Natarajan. Accelerating scientific discovery with Co-Scientist. *Nature*, 655:487–496, 2026. doi: 10.1038/s41586-026-10644-y.

[18] Tanjin He, Aikaterini Vriza, Logan Ward, Xu Huang, Yiming Chen, Anubhav Jain, Gerbrand Ceder, Rajeev S. Assary, Ian T. Foster, and Maria K. Y. Chan. Harnessing X-ray Absorption Spectroscopy Data through Multimodal Mining of Battery Literature. *arXiv preprint arXiv:2607.23886*, 2026.

[19] Peiyan Hu, Haodong Feng, Hongyuan Liu, Tongtong Yan, Wenhao Deng, Tianrun Gao, Rong Zheng, Haoren Zheng, Chenglei Yu, Chuanrui Wang, Kaiwen Li, Zhi-Ming Ma, Dezhi Zhou, Xingcai Lu, Dixia Fan, and Tailin Wu. RealPDEBench: A Benchmark for Complex Physical Systems with Real-World Data. *arXiv preprint arXiv:2601.01829*, 2026.

[20] Xu Huang, Junwu Chen, Yuxing Fei, Zhuohan Li, Philippe Schwaller, and Gerbrand Ceder. CASCADE: Cumulative Agentic Skill Creation through Autonomous Development and Evolution. *arXiv preprint arXiv:2512.23880*, 2025.

[21] Xu Huang, Junwu Chen, Yuxing Fei, Zhuohan Li, Philippe Schwaller, and Gerbrand Ceder. Sciskillbench: Cascade benchmark for evaluating llm agents on scientific tasks, 2025.

[22] Xu Huang, Junwu Chen, Philippe Schwaller, and Gerbrand Ceder. SkillPuzzler: A Self-Evolving Agentic Framework for Materials and Chemistry Research with Minimal Reliance on Predefined Tools. In *NeurIPS 2025 Workshop on AI for Science: The Reach and Limits of AI for Scientific Discovery*, 2025. URL https://neurips.cc/virtual/2025/125935.

[23] Zhen Huang, Zengzhi Wang, Shijie Xia, Xuefeng Li, Haoyang Zou, Ruijie Xu, Run-Ze Fan, Lyumanshan Ye, Ethan Chern, Yixin Ye, Yikai Zhang, Yuqing Yang, Ting Wu, Binjie Wang, Shichao Sun, Yang Xiao, Yiyuan Li, Fan Zhou, Steffi Chern, Yiwei Qin, Yan Ma, Jiadi Su, Yixiu Liu, Yuxiang Zheng, Shaoting Zhang, Dahua Lin, Yu Qiao, and Pengfei Liu. OlympicArena: Benchmarking Multi-discipline Cognitive Reasoning for Superintelligent AI. In *Advances in Neural Information Processing Systems 37*, pages 19209–19253, 2024. doi: 10.52202/079017-0607.

[24] Yigeng Jiang, Tengchao Yang, Taoyong Cui, Jiaxing Wan, Yuan Wang, Weida Wang, Zhiyu Liu, Chuyi Peng, Binzhao Luo, Maoli Gao, Huaihai Huang, Yuqianer Zeng, Ziyang Zheng, Dongchen Huang, Chao Chen, Zichao Liu, Weiping Shen, Shuchen Pu, Siyu Zhou, Runmin Ma, Yusong Hu, Fei Chao, Bo Zhang, Xiawu Zheng, Zifu Wang, Lei Bai, Yunqi Cai, and Shufei Zhang. Deep Research in Physical Sciences: A Multi-Agent Framework and Comprehensive Benchmark. *arXiv preprint arXiv:2606.18648*, 2026.

[25] Hanzheng Li, Xi Fang, Yixuan Li, Chaozheng Huang, Junjie Wang, Xi Wang, Hongzhe Bai, Bojun Hao, Shenyu Lin, Huiqi Liang, Linfeng Zhang, and Guolin Ke. RxnBench: A Multimodal Benchmark for Evaluating Large Language Models on Chemical Reaction Understanding from the Scientific Literature. *Journal of Chemical Information and Modeling*, 66: 5747–5756, 2026. doi: 10.1021/acs.jcim.6c00286.

[26] Xiangyi Li, Kyoung Whan Choe, Yimin Liu, Xiaokun Chen, Chujun Tao, Bingran You, et al. ClawsBench: Evaluating capability and safety of LLM productivity agents in simulated workspaces. *arXiv preprint arXiv:2604.05172*, 2026. Accepted to COLM 2026.

[27] Xiangyi Li, Yimin Liu, Wenbo Chen, Shenghan Zheng, Yifeng He, Xiaokun Chen, Bingran You, Haotian Shen, Yubo Li, Kyoung Whan Choe, Binxu Li, Shuyi Wang, Luyang Kong, Xinyi Liu, Jiankai Sun, Runhui Wang, Xuanqing Liu, Jiachen Li, Xin Lan, Yuanli Wang, Xuandong Zhao, Yueqian Lin, Wengao Ye, Junwei He, Di Wang, Roey Ben Chaim, Songlin Li, Hanwen Xing, Yue Zhang, Yipeng Gao, Yijiang Li, Ze Ma, Liqiang Jing, Qunhong Zeng, Tianyu Wang, Kaixin Li, Yiqi Xue, Zonglin Di, Yizhuo He, Yuchen Tian, Yixuan Gao, Hejia Geng, Bo Chen, Litong Liu, Sikai Cheng, Zelin Tan, Jiajun Bao, Junwei Zhou, Shuaicheng Tong, Shuwen Xu, Penghao Jiang, Terry Yue Zhuo, Tinghan Ye, Chujun Tao, Qi Qi, Miao Li, Longtai Liao, Zhikang Dong, Chang Shi, Steven Dillmann, and Hanchung Lee. SkillsBench: Benchmarking How Well Agent Skills Work Across Diverse Tasks. In *The Fortieth Annual Conference on Neural Information Processing Systems Evaluations and Datasets Track*, 2026. URL https://openreview.net/forum?id=Utmr8dZbrt.

[28] Xinge Liu, Terry Jingchen Zhang, Bernhard Schölkopf, Zhijing Jin, and Kristen Menou. Stargazer: A Scalable Model-Fitting Benchmark Environment for AI Agents under Astrophysical Constraints. *arXiv preprint arXiv:2604.15664*, 2026.

[29] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *arXiv preprint arXiv:2408.06292*, 2024.

[30] Andres M. Bran, Sam Cox, Oliver Schilter, Carlo Baldassari, Andrew D. White, and Philippe Schwaller. Augmenting large language models with chemistry tools. *Nature Machine Intelligence*, 6:525–535, 2024. doi: 10.1038/s42256-024-00832-8.

[31] Indrajeet Mandal, Jitendra Soni, Mohd Zaki, Morten M. Smedskjaer, Katrin Wondraczek, Lothar Wondraczek, Nitya Nand Gosvami, and N. M. Anoop Krishnan. Evaluating large language model agents for automation of atomic force microscopy. *Nature Communications*, 16, 2025. doi: 10.1038/s41467-025-64105-7.

[32] Tingjia Miao, Wenkai Jin, Muhua Zhang, Jinxin Tan, Yuelin Hu, Tu Guo, Jiejun Zhang, Yuhan Wang, Wenbo Li, Yinuo Gao, Shuo Chen, Weiqi Jiang, Yayun Hu, Zixing Lei, Xianghe Pang, Zexi Liu, Yuzhi Zhang, Linfeng Zhang, Kun Chen, Wei Wang, Weinan E, and Siheng Chen. PRL-Bench: A Comprehensive Benchmark Evaluating LLMs' Capabilities in Frontier Physics Research. *arXiv preprint arXiv:2604.15411*, 2026.

[33] Alexander Novikov, Ngân Vũ, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, Sergey Shirobokov, Borislav Kozlovskii, Francisco J. R. Ruiz, Abbas Mehrabian, M. Pawan Kumar, Abigail See, Swarat Chaudhuri, George Holland, Alex Davies, Sebastian Nowozin, Pushmeet Kohli, and Matej Balog. AlphaEvolve: A coding agent for scientific and algorithmic discovery. *arXiv preprint arXiv:2506.13131*, 2025.

[34] Haining Pan, James Roggeveen, Erez Berg, Juan Alvarez, Debanjan Chowdhury, Surya Ganguli, Federico Ghimenti, Juraj Hasik, Henry Hunt, Hong-Chen Jiang, Mason Kamb, Ying-Jer Kao, Ehsan Khatami, Michael Lawler, Di Luo, Titus Neupert, Xiaoliang Qi, Michael Brenner,

and Eun-Ah Kim. CMT-Benchmark: A Benchmark for Condensed Matter Theory Built by Expert Researchers. In *International Conference on Learning Representations*, volume 2026, pages 35522–35540, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/3bf2d6047e8599e1b7567e5d154d39e9-Abstract-Conference.html.

[35] Shi Qiu, Shaoyang Guo, Zhuo-Yang Song, Yunbo Sun, Zeyu Cai, Jiashen Wei, Tianyu Luo, Yixuan Yin, Zhang Haoxu, Yi Hu, Chenyang Wang, Chencheng Tang, Haoling Chang, Qi Liu, Ziheng Zhou, Tianyu Zhang, Jingtian Zhang, Zhangyi Liu, Minghao Li, Yuku Zhang, Boxuan Jing, Xianqi Yin, Yutong Ren, Zizhuo Fu, Jiaming Ji, Weike Wang, Xudong Tian, Anqi Lv, Laifu Man, Jianxiang Li, Feiyu Tao, Qihua Sun, Zhou Liang, Yushu Mu, Zhongxuan Li, Jing-Jun Zhang, Shutao Zhang, Xiaotian Li, Xingqi Xia, Jiawei Lin, Zheyu Shen, Jiahang Chen, Qiuhao Xiong, Binran Wang, Fengyuan Wang, Ziyang Ni, Bohan Zhang, Fan Cui, Changkun Shao, Qing-Hong Cao, Ming-xing Luo, Yaodong Yang, Muhan Zhang, and Hua Xing Zhu. PHYBench: Holistic Evaluation of Physical Perception and Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems 38*, pages 1345–1385, 2025. doi: 10.52202/085713-0044.

[36] Shi Qiu, Junyi Deng, Yiwei Deng, Haoran Dong, Jieyu Fu, Mao Li, Zeyu Li, Zhaolong Zhang, Huiwen Zheng, Leidong Bao, Anqi Lv, Zihan Mo, Yadi Niu, Yiyang Peng, Yu Tian, Yili Wang, Ziyu Wang, Zi-Yu Wang, Jiashen Wei, Liuheng Wu, Aoran Xue, Leyi Yang, Guanglu Yuan, Xiarui Zhan, Jingjun Zhang, Zifan Zheng, Pengfei Liu, Linrui Zhen, Kaiyang Li, Qichang Li, Ziheng Zhou, Guo-En Nian, Yunwei Xiao, Qing-Hong Cao, Linjie Dai, Xu Feng, Peng Gao, Ying Gu, Chang Liu, Jia Liu, Ming-xing Luo, Yan-Qing Ma, Liang-You Peng, Huichao Song, Shufeng Wang, Chenxu Wang, Tao Wang, Yi-Nan Wang, Chengyin Wu, Pengwei Zhao, and Hua Xing Zhu. PRBench: End-to-end Paper Reproduction in Physics Research. *arXiv preprint arXiv:2603.27646*, 2026.

[37] Samuel Schmidgall, Yusheng Su, Ze Wang, Ximeng Sun, Jialian Wu, Xiaodong Yu, Jiang Liu, Michael Moor, Zicheng Liu, and Emad Barsoum. Agent Laboratory: Using LLM Agents as Research Assistants. *arXiv preprint arXiv:2501.04227*, 2025.

[38] Joar Skalse, Nikolaus Howe, Dmitrii Krasheninnikov, and David Krueger. Defining and Characterizing Reward Hacking. In *Advances in Neural Information Processing Systems 35*, pages 9460–9471, 2022. doi: 10.52202/068431-0687.

[39] Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. PaperBench: Evaluating AI's Ability to Replicate AI Research. In *International Conference on Machine Learning*, pages 56843–56873, 2025. URL https://proceedings.mlr.press/v267/starace25a.html.

[40] Yiyou Sun, Xinyang Han, Weichen Zhang, Yuanbo Pang, Tianyu Wang, Yuhan Cao, Yixiao Huang, et al. Agents' Last Exam. *arXiv preprint arXiv:2606.05405*, 2026.

[41] Kyle Swanson, Wesley Wu, Nash L. Bulaong, John E. Pak, and James Zou. The Virtual Lab of AI agents designs new SARS-CoV-2 nanobodies. *Nature*, 646:716–723, 2025. doi: 10.1038/s41586-025-09442-9.

[42] Terminal-Bench-Science Team. Terminal-Bench-Science: Evaluating AI agents on research workflows across scientific domains, 2026. URL https://github.com/harbor-framework/terminal-bench-science.

[43] Miles Wang, Robi Lin, Kat Hu, Joy Jiao, Neil Chowdhury, Ethan Chang, and Tejal Patwardhan. FrontierScience: Evaluating AI's Ability to Perform Expert-Level Scientific Tasks. *arXiv preprint arXiv:2601.21165*, 2026.

[44] Weida Wang, Dongchen Huang, Jiatong LI, Tengchao Yang, Ziyang Zheng, Chuyi Peng, Di Zhang, Dong Han, Benteng Chen, Binzhao Luo, Zhiyu Liu, kunling liu, Zhiyuan Gao, Shiqigeng geng, Wei Ma, Jiaming Su, Xin Li, Shuchen Pu, Yuhan Shui, Qianjia Cheng, Zhihao Dou, Dongfei Cui, Changyong He, Jin Zeng, Zeke Xie, Mao Su, Dongzhan Zhou, Yuqiang Li, Wanli Ouyang, Yunqi Cai, Xi Dai, Shufei Zhang, LEI BAI, Jinguang Cheng, Zhong Fang, and Hongming Weng. CMPhysBench: A Benchmark for Evaluating Large Language Models

in Condensed Matter Physics. In *International Conference on Learning Representations*, volume 2026, pages 3388–3421, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/06e1da559a7227d7cce7364f92976a0e-Abstract-Conference.html.

[45] Yuru Wang, Lejun Cheng, Yuxin Zuo, Sihang Zeng, Bingxiang He, Che Jiang, Junlin Yang, Yuchong Wang, Kaikai Zhao, Weifeng Huang, Kai Tian, Zhenzhao Yuan, Jincheng Zhong, Weizhi Wang, Ning Ding, Bowen Zhou, and Kaiyan Zhang. NatureBench: Can Coding Agents Match the Published SOTA of Nature-Family Papers? *arXiv preprint arXiv:2606.24530*, 2026.

[46] Zirui Wang, Jiaqi Wang, Qinghan Wang, Yuzhi Xu, Gang Du, Tingjun Hou, and Odin Zhang. EpiBench: Can LLMs Understand Epitopes for Antibody Drug Discovery? *arXiv preprint arXiv:2608.06022*, 2026.

[47] Wanghan Xu, Shuo Li, Tianlin Ye, Qinglong Cao, Yixin Chen, Hengjian Gao, Yiheng Wang, Qi Li, Kun Li, Sheng Xu, Shengdu Chai, Fangchen Yu, Xiangyu Zhao, Zhangrui Zhao, Weijie Ma, Zijie Guo, Koutian Wu, Haoyu Zhou, Haoxiang Yin, Lixue Cheng, Chaofan Hu, Haoxuan Li, Lu Mi, Xuxuan Xie, Yifan Zhou, Ruizhe Chen, Zhiwang Zhou, Xingjian Guo, Yuhao Zhou, Xuming He, Shengyuan Xu, Xinyu Gu, Jiamin Wu, Mianxin Liu, Chunfeng Song, Fenghua Ling, Dongzhan Zhou, Shixiang Tang, Yuqiang Li, Mao Su, Peng Ye, Siqi Sun, Bin Wang, Xue Yang, Zhenfei Yin, Tianfan Fu, Guangtao Zhai, Wanli Ouyang, Bo Zhang, Lei Bai, and Wenlong Zhang. ResearchClawBench: A Benchmark for End-to-End Autonomous Scientific Research. *arXiv preprint arXiv:2606.07591*, 2026.

[48] Xin Xu, Qiyun Xu, Tong Xiao, Tianhao Chen, Yuchen Yan, Jiaxin Zhang, Shizhe Diao, Can Yang, and Yang Wang. UGPhysics: A Comprehensive Benchmark for Undergraduate Physics Reasoning with Large Language Models. *arXiv preprint arXiv:2502.00334*, 2025.

[49] Christine Ye, Sihan Yuan, Suchetha Cooray, Steven Dillmann, Ian L. V. Roque, Dalya Baron, Philipp Frank, Sergio Martin-Alvarez, Nolan Koblischke, Frank J Qu, Diyi Yang, Risa Wechsler, and Ioana Ciuca. ReplicationBench: Can AI Agents Replicate Astrophysics Research Papers? *arXiv preprint arXiv:2510.24591*, 2025.

[50] Minhui Zhu, Minyang Tian, Xiaocheng Yang, Tianci Zhou, Lifan Yuan, Penghao Zhu, Eli Chertkov, Shengyan Liu, Yufeng Du, Ziming Ji, Indranil Das, Qingzhi Chen, Junyi Cao, Yufeng Du, Jiabin Yu, Peixue Wu, Jinchen He, Yifan Su, Yikun Jiang, Yujie Zhang, Chang Liu, Ze-Min Huang, Weizhen Jia, Yunkai Wang, Farshid Jafarpour, Yong Zhao, Xinan Chen, Jessie Shelton, Aaron W. Young, John Bartolotta, Wenchao Xu, Yue Sun, Anjun Chu, Victor Colussi, Chris Akers, Nathan Brooks, Wenbo Fu, Jinchao Zhao, Marvin Qi, Anqi Mu, Yubo Yang, Allen Zang, Yang Lyu, Peizhi Mai, Christopher Wilson, Xuefei Guo, Juntai Zhou, Daniel Inafuku, Chi Xue, Luyu Gao, Ze Yang, Yaïr Hein, Yonatan Kahn, Kevin Zhou, Di Luo, John Drew Wilson, Jarrod T. Reilly, Dmytro Bandak, Ofir Press, Liang Yang, Xueying Wang, Hao Tong, Nicolas Chia, Eliu Huerta, and Hao Peng. Probing the Critical Point (CritPt) of AI Reasoning: a Frontier Physics Research Benchmark. *arXiv preprint arXiv:2509.26574*, 2025.

Appendix

A Example Task

Research task
task.md

Read /root/proposal.pdf first. Then carry out the following work.

The implementation is divided into two stages, with each stage handling one key step of this task. All numerics are checked against at least 4 significant figures.

(a) Characterize the apparatus and model the system. First, use the device to measure whatever you need to pin down the interactions of the chain, by driving small, well-chosen atom configurations and analyzing the measured occupations.

Then construct a model of the chain from your measurements (with the drive Omega = 2pi x 6 MHz) and compute the ground-state sigma_mid for a 21-atom and a 27-atom chain at Delta = 10.20 MHz. Because the device supports only a few atoms, these results must come from your own simulation, which needs to be scalable to about 30 atoms with the desired accuracy. Report the value to at least 6 decimal places.

(b) Locate the critical point with the experimental parameters. Using the model you built in (a), compute the ground-state sigma_mid across a range of system sizes and detunings yourself, then perform a finite-size scaling analysis of those curves to determine Delta_c.

(i) Report the finite-size scaling method that you used. For that, report the ratio of the finite-size rescaling factor you apply to sigma_mid for a 7-atom array to that for an 11-atom array.

(ii) Report Delta_c in MHz together with the order-parameter scaling exponent you used, and provide a full analysis of the uncertainty of the obtained Delta_c. Report the value to 6 decimals, including the uncertainty.

When you are done, I want the following:
- /root/result.md: key-value lines following /root/result_template.md.
[The seven keys and their one-line definitions.]
- /root/paper.pdf: a short paper for peer review (abstract, model and methods, results with a finite-size-scaling figure, discussion, references).

Deterministic verifier test_outputs.py	**Reviewer agent** rubric.json
☐ Device log $0 <$ total shots $< 10^6$	**Rydberg platform understanding** (blocker) blockade explained; blockade-violating shots postselected out
☐ v2_mhz ±0.002 MHz of the device's V_2	**Device calibration** (weight 6) Ω, V_1, V_2, V_3 measured; V_1 breaks r^{-6}; decay, loss and shot-to-shot noise; traces under 20 µs
	Literature grounding (weight 2) sources for the hard-boson model, second-order PXP and chain criticality
☐ sigma_mid_21_delta1020 $\pm 2.0 \times 10^{-3}$	**Rydberg quantum simulation basics** (blocker) blockade (PXP) subspace or postselection, with the reason in the paper
☐ sigma_mid_27_delta1020 $\pm 2.6 \times 10^{-3}$	**Hamiltonian construction** (weight 8) second-order PXP solved by ED or MPS; full-space ED cannot reach 30 atoms
☐ rescale_ratio_7_11 0.956346 ± 10^{-6}, so the bulk $(7/11)^{1/8} = 0.945068$ fails	**Universality class and CFT basics** (blocker) (1+1)D Ising; open-boundary $\sin(\pi/(L+2))^{-1/8}$ σ_{mid}, not the bulk $L^{1/8}$
	Plan before analysis (weight 2) σ_{mid} numerics validated before any fit
☐ delta_c_mhz ±0.010 MHz	**Extrapolation quality** (weight 8) crossings per odd-L pair (L, L+4); beyond-linear extrapolation; fit order and window varied
☐ beta_over_nu 1/8 ± 0.01	
☐ delta_c_uncertainty_mhz in (0, 0.05] MHz	
☐ result.md all seven keys	Artifacts present · No fabrication · No tampering · No reward hacking (blockers)
☐ paper.pdf 3+ pages; abstract, method, result, reference	Format requirements (weight 7) · Paper structure (7) · Paper reproducibility (6) · Synthesize information (5)

Pass all 10 tests and all 7 blockers. **Score** $100 \cdot \Sigma w_i s_i / (2\Sigma w_i)$ over the 9 weighted criteria, $s_i \in \{0, 1, 2\}$; 0 if it fails.

Figure 8: **Example task.** `rydberg-ising-critical-point`, the example of Table 4. Left: its `task.md`. The cited proposal (Figure 9) defines σ_{mid} as the order parameter on the central bond; the device simulates a Rydberg apparatus of at most three atoms. The tests that check each step in `task.md` (±: tolerance around withheld reference values) and the summarized rubric criteria that grade it.

Research proposal
proposal.pdf, which task.md says to read first

Research and plan

Quantum phase transition describes a sudden change of quantum system behavior under zero temperature. At the phase transition point, systems display universal behvaiors [sic], and the point of quantum phase transition is called the critical point. Studying universal behavior is therefore an important subject in condensed matter physics.

Recently, Rydberg atom arrays become a competitive candidate in studying such universal behaviors due to its programmability and flexibility. In these setups, an array of neutral atoms are trapped in an equally spaced lattice. Throughout this work, we consider a one-dimensional array of Strontium atoms with open boundary conditions, sitting on the lattice sites [0, 1, 2, …]. A qubit is encoded between its metastable ground state 5s5p $^3P_0$ and a Rydberg state. The system is driven by a global laser with Rabi frequency Omega = 2\pi \times 6 MHz and detuning Delta, and we work in a nearest-neighbor-Rydberg-blockade regime. As the detuning Delta is varied, the chain crosses a quantum phase transition between a disordered phase and an ordered phase at a critical detuning Delta_c.

We want to determine the critical detuning Delta_c of this disorder-to-ordered transition. A simulator of the apparatus is provided as a tool in the workspace (see /root/DEVICE.md for the interface); you can access the device to measure whatever properties of the system you need in order to model it. To locate Delta_c you will need an accurate model of the chain that you can simulate at large sizes; we characterize the transition through the charge-density-wave (CDW) order parameter at the center of the array: on bond j, sigma_j = (-1)^j (n_j - n_{j+1}) (0 indexed) with n being the Rydberg state occupation, and sigma_mid is the order parameter on the central bond, j = floor((N-1)/2).

Research how a critical point of such a chain can be determined from finite-size numerics with open boundary conditions: ❶ what universality class governs this transition and its critical exponents, ❷ how sigma_mid should behave with detuning Delta and system size N near the transition, and ❸ what algorithm and system sizes you need to reach the target precision (as detailed below). ❹ Discuss your reasoning and the relevant literature in the final paper.

	What the literature must supply	Sources the rubric names	Checked by
❶	(1+1)D Ising universality for the disorder-to-Z_2 transition of the blockaded chain; the order parameter's scaling dimension is 1/8, so β/ν = 1/8.	Fendley, Sengupta & Sachdev 2004 (the hard-boson model); Bernien et al. 2017 or Keesling et al. 2019	☐ beta_over_nu Universality class and CFT basics (blocker)
❷	With open boundaries $\sin(\pi/(L+2))^{-1/8}\ \sigma_{mid}$, not the bulk $L^{1/8}\ \sigma_{mid}$, is size-independent at Δ_c; rescaled curves cross for odd-L pairs (L, L+4), and the crossings need a beyond-linear extrapolation.	Fisher & Barber 1972; a boundary-critical-phenomena or crossing-method reference	☐ rescale_ratio_7_11 ☐ delta_c_mhz Extrapolation quality (weight 8)
❸	V_1 is not much larger than Ω, so first-order PXP falls short: a second-order PXP Hamiltonian from a Schrieffer–Wolff expansion, solved by ED in the blockade subspace or by MPS up to about 30 atoms.	Bluvstein et al. 2021	☐ sigma_mid_21_delta1020 ☐ sigma_mid_27_delta1020 Hamiltonian construction (weight 8) Rydberg quantum simulation basics (blocker)
❹	At least three such sources, found and used correctly, one of them for the effective model, and the work placed against them.	—	Literature grounding (weight 2) Synthesize information (weight 5)

Figure 9: **The proposal of the example task.** Top: `proposal.pdf`, the one-page brief that agent need to read first. It sets out the physics and the goal, and describe what are the questions should be answered after literature research. The last paragraph leaves four questions to the agent's own research (numbered). Bottom: for each question, what the literature must supply.

B Task Domain Taxonomy

The 56 tasks are assigned to nine primary domains. Each task contributes to exactly one primary-domain. The "Also spans" column in Table 3 records secondary domains for 15 tasks but these labels do not add extra observations to domain summaries.

Table 3: Primary and secondary domain labels for all 56 FRONTIERPHYSICS tasks. A dash indicates no secondary domain label.

#	Task	Domain	Also spans
1	arpes-sample-alignment	Condensed Matter Physics	Photonics & Imaging
2	baf2-ih001-strain-energy	Condensed Matter Physics	-
3	cpt-frustrated-chain-spectra	Condensed Matter Physics	-
4	eigenstate-optical-response-derivation	Condensed Matter Physics	Photonics & Imaging
5	lattice-band-geometry-discovery	Condensed Matter Physics	-
6	li-graphene-mlip-transferability	Condensed Matter Physics	Chemical Physics
7	llzo-proton-exchange-thermodynamics	Condensed Matter Physics	Chemical Physics
8	lpbf-melt-pool-window	Condensed Matter Physics	Fluid Dynamics
9	mlg-quantum-hall-dot-hartree-fock	Condensed Matter Physics	-
10	moire-continuum-model-construction	Condensed Matter Physics	-
11	rixs-entanglement-witness	Condensed Matter Physics	-
12	smt-reflow-window	Condensed Matter Physics	-
13	ion-trap-3d-vs-2d	AMO Physics	-
14	ion-trap-design	AMO Physics	-
15	multiplexing-ion-chain-qnet	AMO Physics	Quantum Information
16	phase-controlled-optical-tweezer-holography	AMO Physics	-
17	rydberg-cft-modulation-spectroscopy	AMO Physics	-
18	rydberg-cz-phase-noise-infidelity	AMO Physics	Quantum Information
19	rydberg-ising-critical-point	AMO Physics	-
20	QNTK-research	Quantum Information	-
21	cvqkd-calibrated-detector-noise	Quantum Information	-
22	ising-qml-dwave	Quantum Information	-
23	time-bin-qkd-product-povm	Quantum Information	-
24	transmon-cz-pulse-engineering	Quantum Information	-
25	desi-bao-sn-dark-energy	Astrophysics	-
26	ems-scalar-spectral-census	Astrophysics	-
27	nonlinear-boson-star-sideband-threshold	Astrophysics	-
28	ns-eos-multimessenger	Astrophysics	-
29	pulsar-timing-phase-connection	Astrophysics	-
30	solar-wind-turbulence-spectral-break	Astrophysics	-
31	consistent-two-phase-ls-vof	Fluid Dynamics	-
32	diffuse-ib-rigid-body-fsi	Fluid Dynamics	-
33	iga-branched-incompressible-flow	Fluid Dynamics	-
34	material-agnostic-visual-inverse-dynamics	Fluid Dynamics	-
35	multi-object-visual-inverse-dynamics	Fluid Dynamics	-
36	twisted-slinky-wave-coupling	Fluid Dynamics	-
37	glass-nonequilibrium-landscape-evolution	Statistical Physics	-
38	kinked-ladder-polymer-scattering	Statistical Physics	Chemical Physics
39	learning-credit-assignment	Statistical Physics	-

Continued on next page

Table 3: Task domain labels (continued).

#	Task	Domain	Also spans
40	`low-dimensional-deep-network-learning`	Statistical Physics	-
41	`metallic-glass-pel-ruggedness`	Statistical Physics	-
42	`quantized-ising-3sat-encoding`	Statistical Physics	Quantum Information
43	`calibrated-cbct-local-reconstruction`	Photonics & Imaging	-
44	`fbg-rail-axle-inversion`	Photonics & Imaging	-
45	`fbg-strain-amplifier-design`	Photonics & Imaging	-
46	`heat-assisted-detection-and-ranging`	Photonics & Imaging	-
47	`hspic-spectrometer-e2e`	Photonics & Imaging	-
48	`corsika-air-shower-inverse`	High Energy Physics	Astrophysics
49	`icecube-pointsource-unbinned`	High Energy Physics	Astrophysics
50	`o3-monodromy-line-defect-mc`	High Energy Physics	Statistical Physics
51	`o3-surface-cft-fuzzy-sphere`	High Energy Physics	Statistical Physics
52	`sipm-spe-gain-breakdown`	High Energy Physics	Photonics & Imaging
53	`agi-surface-hydrophobicity`	Chemical Physics	-
54	`hole-bifurcation-ml-landscape`	Chemical Physics	-
55	`p2d-cell-model-identification`	Chemical Physics	-
56	`smfret-research-v5`	Chemical Physics	Photonics & Imaging

C Input Modalities and Research Activities

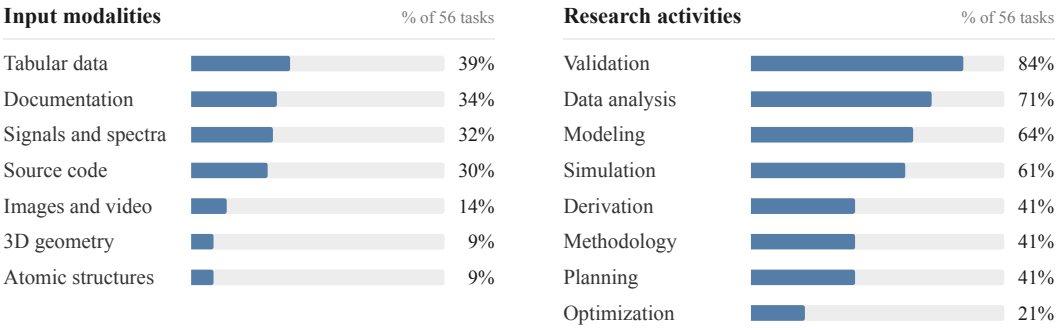

Figure 10: Input modalities and required research activities across the 56 accepted tasks. Each bar gives the percentage of tasks carrying that label. Modality and activity categories are non-exclusive: a single task can contribute to multiple categories, so the percentages within each panel need not sum to 100%.

D Per-Task Results

Figures 11–14 break the results of Table 2 down by task for all 26 model configurations. Each cell summarizes the three designated rollouts of one model on one task. As in Table 2, an unavailable rollout counts as zero. The bottom row of each figure is the model's mean over all 56 tasks. The four following figures show:

- **Verifier pass rate** (Figure 11): the percentage of rollouts in which every applicable deterministic test passes. Its bottom row is the deterministic Pass@1 of Table 2.
- **Rubric blocker pass rate** (Figure 12): the percentage of rollouts in which every rubric blocker passes.
- **Rubric score** (Figure 13): the weighted non-blocker score in Equation 1, before the verifier and blocker gates are applied.
- **Score** (Figure 14): the rubric score (Equation 1) of rollouts that pass the verifier and every blocker, and zero otherwise. Its bottom row is the Score of Table 2.

Figure 11: Verifier pass rate per task and model. The bottom row is each model's deterministic Pass@1 over all 56 tasks (Table 2). Unavailable outcomes contribute zero to the aggregate.

Figure 12: Rubric blocker pass rate per task and model. The bottom row is each model's mean over all 56 tasks. Unavailable outcomes contribute zero to the aggregate.

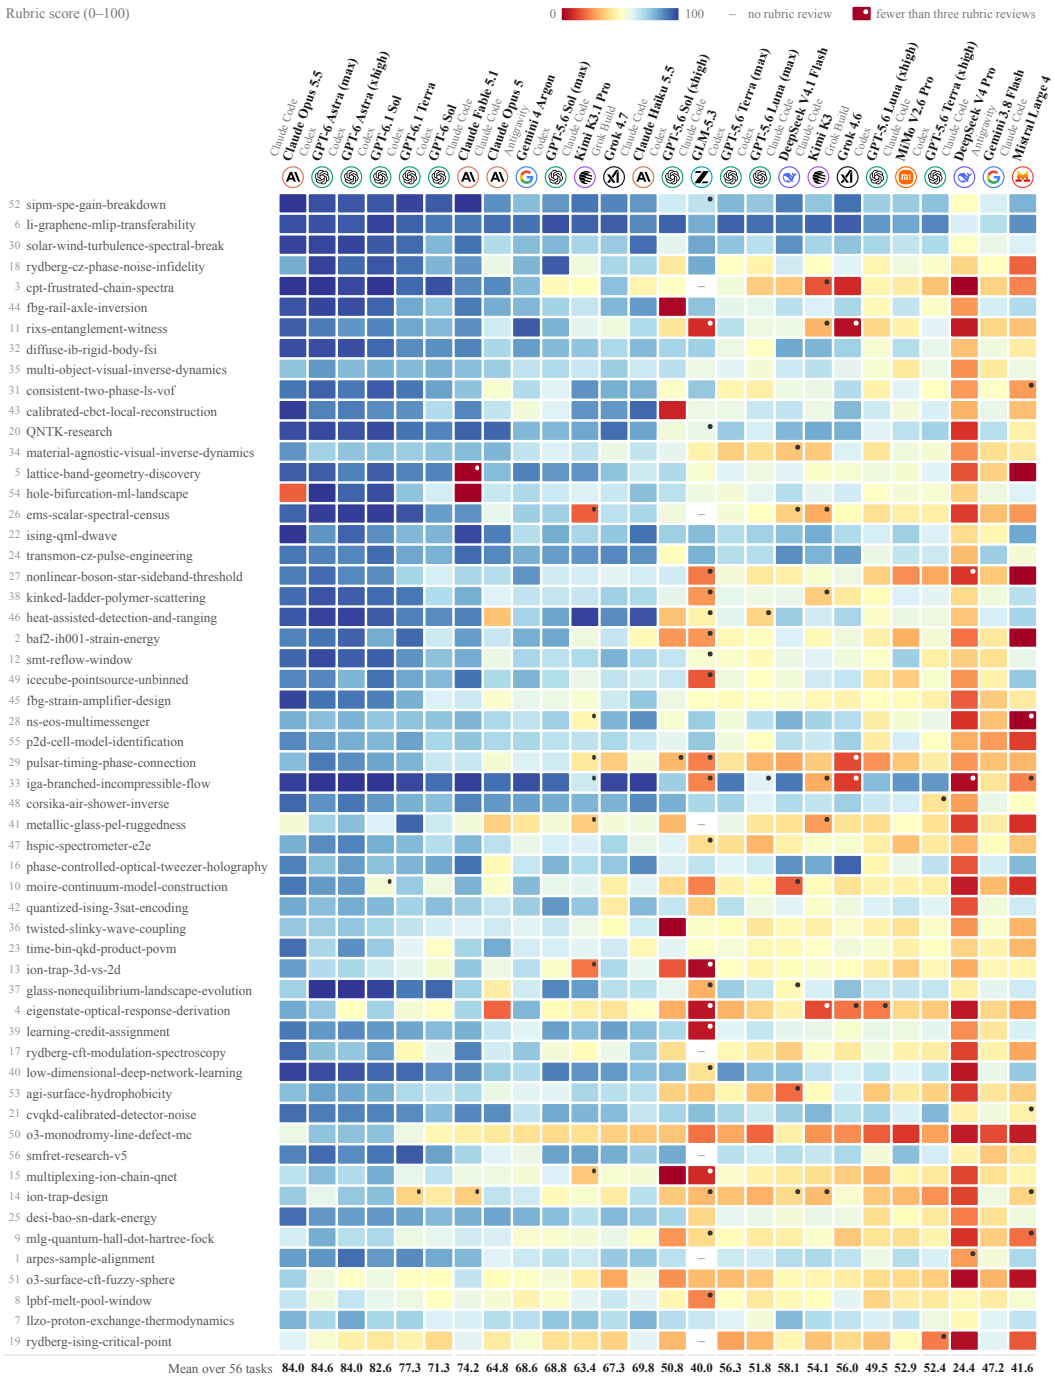

Figure 13: Rubric score per task and model. The weighted non-blocker score (Equation 1) before the verifier and blocker gates, averaged over three rollouts. The bottom row is each model's mean over all 56 tasks.

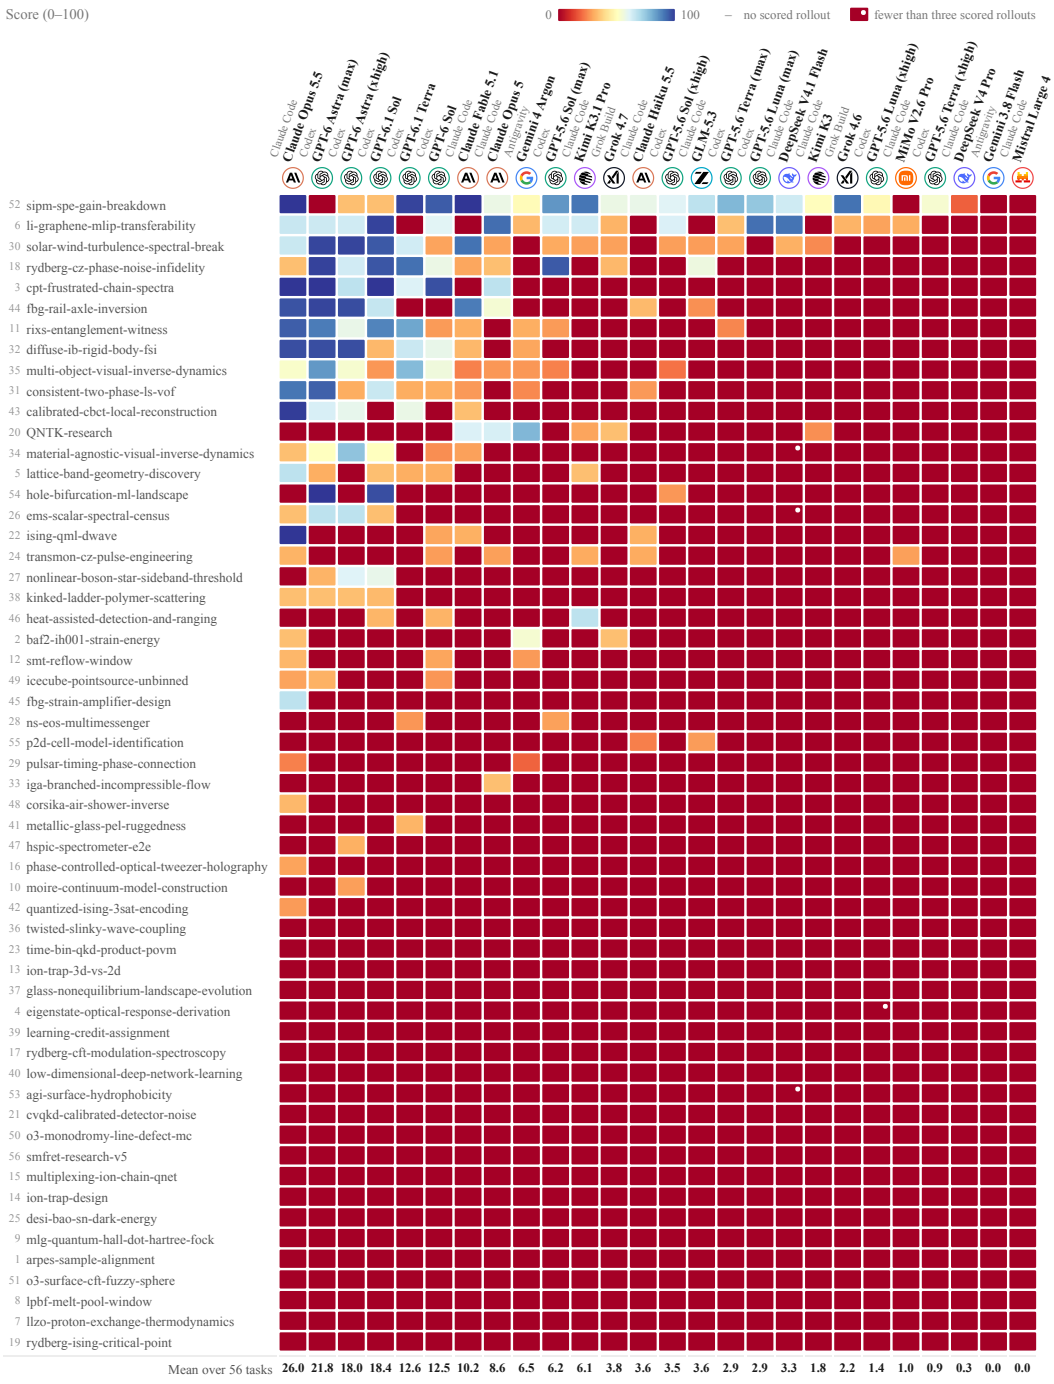

Figure 14: Score (Equation 1) per task and model. The bottom row reproduces Score in Table 2. Rows follow Figure 11. Missing rollouts count as zero.

E Failure-Mode Analysis

The following table shows a detailed explanation of the failure mode taxonomy that was extracted on the latest GPT-6 Astra (`xhigh`, `codex-acp`) cohort based on the method described in Section 3.3.

Table 4: Failure mode taxonomy.

Tag	Label	Description	Examples
Get the physics wrong	**PHY**	• Subtle mistakes in the formulas or theory applied; • Agents do not have such physics knowledge (no signal when checking agents' trajectories).	`mlg-quantum-hall-dot-hartree-fock` `material-agnostic-visual-inverse-dynamics`
Scientific presenting	**SP**	Agents have a good understanding of the task, but the paper writing is problematic and cannot present the research professionally.	`rixs-entanglement-witness`
False grading	**RRMIS**	The reviewer agent's rubric grading does not align with human researchers' grading. For example, the rubric is correct and the worker agent does the right thing, but the reviewer agent makes a mistake (often by not following the instructions properly).	`li-graphene-mlip-transferability`
Literature mistake	**LIT**	• Agents cannot find the important papers (specified by the task author in rubrics as important papers that should be found) in the research stage; • In `paper.pdf`, agents cannot use the literature to support their claims correctly.	`arpes-sample-alignment` `calibrated-cbct-local-reconstruction`

Agents do not follow `task.md` properly	**INS**	• Agents miss certain instructions in `task.md` (do not make the judgement based only on agents' artifacts; check their trajectories to see if agents ever recognize the instructions).
		• Agents don't understand or misunderstand the instructions.
		In the example task, agents do not get that "we work in a nearest-neighbor-Rydberg-blockade regime" implies they should make approximations to the Hamiltonian or perform postselection.
			`glass-nonequilibrium-landscape-evolution`
			`low-dimensional-deep-network-learning`
			`rydberg-ising-critical-point` (Trials 1–3; supporting; task-dominant: REAL)

Agents do not understand the reality constraints	**REAL**	In the example task, agents do not get the reality constraints of the simulated device.	`rydberg-ising-critical-point`

Missing physics expert intuition or experience	**INT**	In the example task rollouts, agents do not consider Rydberg decay when analyzing sources of error. An AMO physicist would naturally consider this mechanism, suggesting a gap in the agents' physical intuition.
		Also, sometimes in the thinking traces, agents are aware of the instructions but still fail in the end since they lack the expert-level physics intuition needed to apply them correctly.
			`multiplexing-ion-chain-qnet`
			`lpbf-melt-pool-window`
			`hspic-spectrometer-e2e`
			`rydberg-cft-modulation-spectroscopy` (Trial 1; supporting; task-dominant: REAL)

Incoherent reasoning	**COH**	• Contradictory logic chains or arguments. In the example task rollouts, agents realize that they are wrong, but they later forget this.	`kinked-ladder-polymer-scattering`

Continued on next page

Table 4: Failure mode taxonomy (continued).

Tag	Label	Description	Examples

Lack of rigorous thinking	RIG	• Circular self-verification.
		• Agents do not consider enough edge cases, so the program or method only works for limited cases.
		In the example task rollouts, agents' self-verification method depends on their own assumption, which is wrong, so they can never find their mistakes.
			`p2d-cell-model-identification`
			`desi-bao-sn-dark-energy`
			`time-bin-qkd-product-povm`
			`hole-bifurcation-ml-landscape`

False negative	FALSE	Agents do the task correctly but still fail because of a task-design issue (unclear `task.md`, a verifier or rubric that contradicts the instructions, an overly strict tolerance).	`phase-controlled-optical-tweezer-holography`
Misgrade	MIS	Certain test or rubric results are not correct for this trial, but they do not influence whether this trial passes or not.	`arpes-sample-alignment` (Trials 1,3; task-dominant: LIT)
Agents' execution mistakes	EXE	Agents cannot write proper code or use the tools properly, etc.	`icecube-pointsource-unbinned`
Hallucination	HAL	Agents fabricate evidence, data, experiments, literature, etc.	`nonlinear-boson-star-sideband-threshold` (Trial 1; task-dominant: RIG)
Reward hacking	HACK	Agents try to hack to get the full reward by: • Trying to change verifiers or hack rubrics to find answers; • Trying to directly find answers online and copy them to the final workspace to pretend the job has been done, etc.	-

E.1 Final Trial Labels and Main Findings

Each trial retains all labels and their order from the final review sheet, including locally rerun trials. PASS may coexist with other annotations. ^{V} preserves a verifier marker from the sheet. A dash in the finding column means no finding was supplied. For `baf2-ih001-strain-energy`, the three columns correspond to local Trials 4–6. Task numbers follow Table 3.

Table 5: Final review-sheet labels for 56 tasks. Cyan cells (eight FALSE, one RRMIS) identify the 9 trials whose final pass/fail outcomes should change; MIS denotes verifier or rubric errors that do not change the final outcome.

#	Task	Trial 1	Trial 2	Trial 3	Main finding
Condensed Matter Physics					
1	`arpes-sample-alignment`	RIG LIT MIS	LIT RIG	LIT MIS	Agents recognized tilt–translation coupling but missed pose-recovery tolerances and omitted electron-optics or inverse-optimization references. Reviewers overlooked executed one-coordinate scans in Trials 1 and 3, although both still failed pose recovery.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
2	baf2-ih001-strain-energy	INT RIG	INT RIG	RIG	Displayed Trials 1–2 insufficiently justified bulk-like ice selection and failed AgI transfer. Displayed Trial 3 passed AgI but did not justify its selected ice population despite crop-dependent results.
3	cpt-frustrated-chain-spectra	PASS	PASS	PASS	-
4	eigenstate-optical-response-derivation	INS MIS	INS MIS	INS	Agents did not start from the required current-density operator. PDF extraction and Lean environment setup also produced erroneous checks, while other verifier and blocker failures remained.
5	lattice-band-geometry-discovery	RIG LIT	FALSE	INS RIG	Trial 1 proved the scaling limit, but an early mistake spoiled its numerics; Trial 2's correct numerics hit a run-time error because the task requested too large a grid, a false negative, and Trial 3 met the same problem.
6	li-graphene-mlip-transferability	PASS RIG	RRMIS RIG	PASS	Trials 1 and 2 computed frame errors correctly but did not discuss why two of the 25 frames had larger errors; Trial 2's verifier missed a correctly printed adsorption energy after OCR dropped its minus sign. Trial 3's paper needed better writing.
7	llzo-proton-exchange-thermodynamics	RIG LIT	RIG LIT	RIG LIT	None of the agents reported the trigonal planar Li that belong to neither cavity, and two built their parent lattice on coordinates with no traceable source.
8	lpbf-melt-pool-window	RIG INT	INT	INT	Agents recommended unreliable process windows, retaining unstable candidates and marking unsafe settings as stable_screening; Trial 2 nearly passed but claimed an unstable setting stable.
9	mlg-quantum-hall-dot-hartree-fock	PHY RIG	PHY RIG	PHY RIG	Agents left the external confinement unscreened, omitting the RPA screening that this simulation setting required, among other physics mistakes; an inexperienced human might also have made this omission.
10	moire-continuum-model-construction	INS MIS	PASS INS	INS MIS	Agents passed every numerical test but mostly reached the answer through a special symmetry argument, too specialized for this model class, although the task asked for a universal low-energy reduction instead.
11	rixs-entanglement-witness	SP INT	PASS SP	PASS SP INT	Agents got the physics right but left out important assumptions a publishable paper should discuss; in Trials 1 and 3 they also simulated correctly but missed the phase transition in their results.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
12	`smt-reflow-window`	INT REAL	INT REAL	INT REAL	Agents missed real-world constraints and lacked expert intuition (e.g., that excess solder raises the risk of shorts); their bridge and cold-joint scores did not respond to clearly risky paste or thermal conditions, and Trial 2 started from random guesses.
AMO Physics					
13	`ion-trap-3d-vs-2d`	INT MIS	INT	INT	Agents solved the electrostatics correctly but failed to connect static trap depth to experimental ion loss. Trial 1 also received inconsistent grading on pseudopotential validity.
14	`ion-trap-design`	FALSE	RIG COH	INT RIG COH	Trial 1 met trap-design and operating-configuration requirements; Trial 2 met only the former. Trial 3 applied inappropriate gapless electrostatics to a slitted trap and lacked independent verification despite recognizing its importance.
15	`multiplexing-ion-chain-qnet`	INT	INT	INT	Agents could not fit Rabi-flop data with a model combining coherent excitation and thermal heating.
16	`phase-controlled-optical-tweezer-holography`	FALSE	FALSE	FALSE	-
17	`rydberg-cft-modulation-spectroscopy`	INT REAL PHY	REAL	REAL	All trials omitted device calibration. Trial 1 additionally misunderstood CFT, inadequately assessed Rydberg blockade, and omitted decay as an error source. Its rubric review also contained both false-negative and false-positive judgments.
18	`rydberg-cz-phase-noise-infidelity`	COH INT RIG EXE (FALSE)	PASS RIG FALSE	PASS RIG INT	All trials passed the verifier tests. Trial 1 failed a blocker by justifying its answer only with in-loop noise squashing, which its own data contradicted; Trials 1 and 3 did not state that the PDH servo bump causes the infidelity gap.
19	`rydberg-ising-critical-point`	REAL MIS INS LIT	REAL MIS INS LIT	REAL MIS INS LIT	Agents exceeded realistic shot budgets, omitted calibration, misunderstood open-boundary finite-size numerics, and weakly synthesized literature. Grading also relied on rigid section-word checks and an unclear nearest-neighbor blockade assumption.
Quantum Information					
20	`QNTK-research`	MIS RIG INS	MIS RIG INS	INS MIS RIG	Agents did not use the required circuit model. Their self-verification reused their own assumptions and provided weak checks.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
21	`cvqkd-calibrated-detector-noise`	PHY LIT RIG	PHY LIT RIG	PHY LIT RIG	Agents missed the intended detector model, using security assumptions that differed from the reference and giving a wrong modulation variance and key rates negative beyond 100 km; key papers appeared late, not in the introduction, and the PLOB bound went unchecked.
22	`ising-qml-dwave`	PASS	PASS LIT	PASS	The physics and optimization were correct in all trials; failures concerned the analysis and write-up.
23	`time-bin-qkd-product-povm`	RIG LIT	RIG INS LIT	RIG LIT	Agents computed the numerical results correctly, but their papers omitted the security implications and why Lavie–Lim's extra-vacuum-pulse method is needed, and missed key prior work; Trial 2 also skipped the IID discussion and once tried to access papers past the literature cutoff.
24	`transmon-cz-pulse-engineering`	INT COH	INT LIT COH	INT COH	Agents implemented the physics correctly but failed line-accuracy tests: they recognized noise in the final third of their results yet published it unsmoothed. Trial 2 used a method without properly citing its paper.
Astrophysics					
25	`desi-bao-sn-dark-energy`	RIG INT INS	RIG INT INS	RIG INT INS	The pipeline always returns a precise BAO scale, even for spectra with no acoustic feature, and its weak-signal intervals are uncalibrated; closure is run but not reported in the required form. Downstream cosmology, SN selection, degeneracies and null calibration are correct in all three trials.
26	`ems-scalar-spectral-census`	RIG PASS	PASS	PASS	Trial 1 missed hidden edge cases. Trial 2 explicitly considered the edge case; Trial 3 avoided it through a different design.
27	`nonlinear-boson-star-sideband-threshold`	HAL INT RIG PHY SP LIT	RIG INT SP LIT	PASS	Trial 1's mode classifier missed edge cases and miscounted nodes. Its grading also depended on a precision requirement absent from task.md.
28	`ns-eos-multimessenger`	FALSE	RIG	PASS	Trial 1 was penalized for an ambiguous prior convention. Trial 2 inconsistently enforced stable-branch selection, and its input guard rejected hidden data. The physics was largely correct; papers omitted equations and caveats present in code.
29	`pulsar-timing-phase-connection`	RIG MIS	RIG MIS	RIG MIS	Agents applied a standard threshold without adapting or validating it for this specific case.
30	`solar-wind-turbulence-spectral-break`	PASS	PASS	PASS	All three trials passed.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
Fluid Dynamics					
31	`consistent-two-phase-ls-vof`	FALSE	FALSE	FALSE	Trial 2's verifier failure may stem from an unspecified accuracy requirement that allowed premature stopping.
32	`diffuse-ib-rigid-body-fsi`	PASS INT MIS	PASS	PASS	Agents implemented surface versus volumetric forcing as superficial point-placement differences rather than using the correct fractional-step stage. Reviewers also penalized coding before a detailed plan, which the task did not require.
33	`iga-branched-incompressible-flow`	INT (REAL) (RIG)	INT (REAL) (RIG)	INT (REAL) (RIG)	Agents failed hidden cases with different geometries: they lacked a global understanding of the final model, whose issues a global overview in professional software would reveal, and gave edge cases too little thought.
34	`material-agnostic-visual-inverse-dynamics`	PHY	PHY	PHY	Agents failed to design sufficient supervision for stable constitutive learning.
35	`multi-object-visual-inverse-dynamics`	PHY RIG INS	PHY RIG INS	PHY RIG INS	Agents failed to distinguish constitutive errors from geometry and contact errors.
36	`twisted-slinky-wave-coupling`	INT LIT	INT LIT	INT LIT	Agents treated the spring as a whole instead of analyzing transverse motion at the cross-sectional and coil scales. They therefore failed to explain transverse–torsional coupling and provided only qualitative relations and overall trend fits.
Statistical Physics					
37	`glass-nonequilibrium-landscape-evolution`	INS	INS	INS	Agents passed all verifier tests but omitted the required comparison of factorized and dependence-preserving kinetic models and its implications for the ageing–rejuvenation crossover; Trial 2 also left activation–relaxation dependence unassessed.
38	`kinked-ladder-polymer-scattering`	INT COH RIG	INT COH RIG	INT COH RIG	Agents retained an incorrect sign rule while tuning other parameters, then claimed validation despite documenting the unresolved problem in the appendix.
39	`learning-credit-assignment`	INS	INS	INS	Agents gave results around the requested limit rather than under it and, when told to apply their model to other problems, used another method.
40	`low-dimensional-deep-network-learning`	INS LIT	INS INT LIT	INS PHY INT LIT	Trial 1 chose a wrong algorithm by ignoring conditions in task.md. Agents performed the mathematics correctly but selected the less-efficient, unshared design and did not recognize that their theory required a simplifying assumption.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
41	`metallic-glass-pel-ruggedness`	INS RIG	INS RIG	PASS	Agents recognized the described differences but did not identify a quantitative measure with physical implications; one failed trial also omitted the required raw energy and force measurements.
42	`quantized-ising-3sat-encoding`	PASS INS LIT	PASS INS LIT	PASS INS	All three trials solved the formulas; the failures concerned the papers.
Photonics & Imaging					
43	`calibrated-cbct-local-reconstruction`	INS LIT	INS LIT	RIG LIT	-
44	`fbg-rail-axle-inversion`	PASS	PASS	PASS	All three trials passed.
45	`fbg-strain-amplifier-design`	INS MIS	INS MIS	INS MIS	Agents may have misinterpreted one task requirement, leading them to violate another.
46	`heat-assisted-detection-and-ranging`	REAL INT RIG EXE	REAL INT RIG EXE	REAL INT RIG EXE	Agents overlooked emissivity structure, real sensing conditions, and range identifiability, then failed to implement stable decomposition of temperature, emissivity, texture, and distance.
47	`hspic-spectrometer-e2e`	INT RIG MIS	INT LIT RIG	PASS RIG	Agents compute the spectral metric C and the residual correlation correctly, but in Trials 1 and 2 they never explain why the residual check is needed: a periodic (comb-like) response can have a narrow first crossing and still show large later correlation peaks (INT). The same two trials fail complexity_scaling. They discuss round-trip delays and reflection-free cascades but never derive that a transmissive cascade's path order is limited by its stage count and that reflections remove this limit (INT). Trial 2 also fails positioning, because its comparison table has no filter-array entry and does not treat reconstructive spectrometers as a separate family (LIT). Trial 3 makes the periodic-response point explicit with a long-delay MZI (C = 6.630 nm^{-1}, residual correlation 0.226 > 0.15) and passes.
High Energy Physics					
48	`corsika-air-shower-inverse`	INT LIT	INT LIT	INT LIT	-
49	`icecube-pointsource-unbinned`	EXE RIG	EXE	PASS INT LIT	All three runs build the search correctly and pass every hidden scientific re-run; the failures are at the program's edge: one delivered table inconsistent with its own fit, and a cache written into the read-only code directory in –localisation-batch mode. Trial 3 passes.

Continued on next page

Table 5: Final review-sheet labels and task findings (continued).

#	Task	Trial 1	Trial 2	Trial 3	Main finding
50	o3-monodromy-line-defect-mc	PASS RIG PHY	PASS RIG PHY	PASS RIG PHY	In the line-defect Monte Carlo study, agents obtained multiple RG fixed points but could not identify which were physically relevant.
51	o3-surface-cft-fuzzy-sphere	RIG PHY	RIG PHY	RIG PHY	The task hinged on distinguishing ordinary from extraordinary boundary conditions and pairing operators correctly; agents misimplemented one boundary condition from the start and misidentified the second and third S_z=1 states.
52	sipm-spe-gain-breakdown	INS LIT	PASS	PASS	-
Chemical Physics					
53	agi-surface-hydrophobicity	INT RIG	INT RIG MIS	INT RIG	Agents avoided the intended basal interface by selecting a nonpolar prism. Cross-surface probes and system-size justification were incomplete; Trials 2–3 lacked stable estimates. Trial 3's water model conflicted with the planned study; reviewers overlooked Trial 2's conclusion.
54	hole-bifurcation-ml-landscape	PASS	RIG LIT	RIG LIT	Rerun after a task fix: Trial 1 passed. Trials 2 and 3 simulated many designs differently from the reference, changing high-yield labels and the train/test split; Trial 2 checked results with the model that produced them and omitted the PCCP study.
55	p2d-cell-model-identification	RIG RRMIS PHY LIT	RIG INS RRMIS LIT	RIG RRMIS PHY LIT	Agents meshed the hole-covered surface too coarsely, and their models, with poorly chosen transport and kinetic prefactors, did not generalize to other currents and temperatures; the rubric grader also penalized missing measured curves that the task forbade downloading.
56	smfret-research-v5	PHY RIG LIT	PHY RIG	HAL PHY	Agents used incorrect physical models and unsupported literature, and their conclusions did not match the evidence.

Table 6: Regrading with a second reviewer agent. GPT-5.6 Sol (xhigh), the default reviewer agent of all main results, and Claude Opus 5.5 grade the same first-round rollouts with the same rubrics. Rubric scores are as in Figure 13. The pass/fail columns cover the rollouts that pass the verifier, the only ones whose outcome the reviewer agent decides; κ is Cohen's kappa, and Opus only and Sol only count rollouts passed by just one of the two reviewer agents. Two Claude Opus 5.5 rollouts were not regraded.

				Pass/fail, verifier-passing rollouts			
Rollouts	**Reviewer agent**	**Overall passes**	**Rubric score (%)**	Agree	κ	Opus only	Sol only
GPT-6 Astra (xhigh) Codex · 168 rollouts	**GPT-5.6 Sol** default	25	83.4	48/53	0.81	3	2
	Claude Opus 5.5	26	84.1				
Claude Opus 5.5 Claude Code · 166 rollouts	**GPT-5.6 Sol** default	37	83.6	58/67	0.72	8	1
	Claude Opus 5.5 self	44	87.7				

F Slot Accounting

Table 7 counts the slots without a valid result, and Table 8 gives confidence intervals for the overall pass rates from a bootstrap over tasks.

Table 7: Slots without a valid result per configuration: no rollout, or a rollout without complete scoring. Each of them counts as a failure in the overall pass rate and as zero in the score. A rollout that ended at a time or memory limit, or that the model ended, is a scored failure and is not counted here. Every configuration has 168 slots.

Configuration	Slots	Configuration	Slots	Configuration	Slots
Claude Opus 5.5	0	GPT-5.6 Sol (max)	0	Grok 4.7	0
Claude Fable 5.1	0	GPT-5.6 Sol (xhigh)	0	Grok 4.6	0
Claude Opus 5	0	GPT-5.6 Luna (max)	0	Kimi K3.1 Pro	0
Claude Haiku 5.5	0	GPT-5.6 Luna (xhigh)	1	Kimi K3	0
GPT-6 Astra (max)	0	GPT-5.6 Terra (max)	0	MiMo V2.6 Pro	0
GPT-6 Astra (xhigh)	0	GPT-5.6 Terra (xhigh)	0	GLM-5.3	0
GPT-6.1 Sol	0	Gemini 4 Argon	0	Mistral Large 4	0
GPT-6 Sol	0	Gemini 3.8 Flash	0	DeepSeek V4 Pro	0
GPT-6.1 Terra	0			DeepSeek V4.1 Flash	4

Table 8: Overall pass rates with task-clustered uncertainty, over all 56 tasks and 3 trials per task; unavailable outcomes count as failures, as in Table 2. 95% confidence intervals from a bootstrap over tasks, 10,000 resamples, which accounts for the correlation of the three trials of a task. The Binomial column repeats the previous approximate 95% binomial (Wald) half-width over 168 rollouts, which treats the rollouts as independent. Pass@3 is the share of tasks with at least one passing trial. 3/3 and $\geq$1/3 count the tasks (of 56) passed in all three trials and in at least one. †Not statistically separable from Claude Opus 5.5 (two-sided paired sign-flip permutation test over tasks, unadjusted $p \geq 0.05$).

	Pass@1 (%)			Pass@3 (%)		Tasks passed	
Configuration	Est.	95% CI	Binomial	Est.	95% CI	3/3	$\geq$1/3
Claude Opus 5.5	27.4	[17.9, 36.9]	±6.7	44.6	[32.1, 57.1]	8	25
GPT-6 Astra (max)†	23.2	[13.7, 33.9]	±6.4	30.4	[17.9, 42.9]	9	17
GPT-6 Astra (xhigh)†	19.6	[11.3, 28.6]	±6.0	30.4	[17.9, 42.9]	4	17
GPT-6.1 Sol†	19.6	[11.3, 29.2]	±6.0	30.4	[19.6, 42.9]	6	17
GPT-6.1 Terra	14.3	[7.1, 22.6]	±5.3	21.4	[10.7, 32.1]	4	12
GPT-6 Sol	14.3	[7.7, 21.4]	±5.3	28.6	[17.9, 41.1]	2	16
Claude Fable 5.1	11.3	[5.4, 18.5]	±4.8	21.4	[10.7, 32.1]	3	12
Claude Opus 5	9.5	[4.2, 16.1]	±4.4	17.9	[8.9, 28.6]	1	10
Gemini 4 Argon	8.3	[3.6, 14.3]	±4.2	17.9	[8.9, 28.6]	1	10
GPT-5.6 Sol (max)	7.1	[2.4, 13.7]	±3.9	12.5	[5.4, 21.4]	2	7
Kimi K3.1 Pro	6.5	[2.4, 12.5]	±3.7	12.5	[5.4, 21.4]	1	7
Grok 4.7	4.2	[1.2, 7.7]	±3.0	10.7	[3.6, 19.6]	0	6
Claude Haiku 5.5	4.2	[1.2, 7.7]	±3.0	10.7	[3.6, 19.6]	0	6
GPT-5.6 Sol (xhigh)	4.2	[1.2, 8.3]	±3.0	8.9	[1.8, 17.9]	0	5
GLM-5.3	4.2	[1.2, 8.3]	±3.0	8.9	[1.8, 17.9]	0	5
GPT-5.6 Terra (max)	3.6	[0.6, 8.3]	±2.8	7.1	[1.8, 14.3]	1	4
GPT-5.6 Luna (max)	3.6	[0.0, 8.9]	±2.8	3.6	[0.0, 8.9]	2	2
DeepSeek V4.1 Flash	3.6	[0.0, 8.3]	±2.8	5.4	[0.0, 12.5]	1	3
Kimi K3	2.4	[0.0, 5.4]	±2.3	5.4	[0.0, 12.5]	0	3
Grok 4.6	2.4	[0.0, 6.5]	±2.3	3.6	[0.0, 8.9]	1	2
GPT-5.6 Luna (xhigh)	1.8	[0.0, 4.8]	±2.0	3.6	[0.0, 8.9]	0	2
MiMo V2.6 Pro	1.2	[0.0, 3.0]	±1.6	3.6	[0.0, 8.9]	0	2
GPT-5.6 Terra (xhigh)	1.2	[0.0, 3.6]	±1.6	1.8	[0.0, 5.4]	0	1
DeepSeek V4 Pro	0.6	[0.0, 1.8]	±1.2	1.8	[0.0, 5.4]	0	1
Gemini 3.8 Flash	0.0	[0.0, 0.0]	±0.0	0.0	[0.0, 0.0]	0	0
Mistral Large 4	0.0	[0.0, 0.0]	±0.0	0.0	[0.0, 0.0]	0	0

G Agent Resource Use

Figures 5, 15, and 16 compare overall pass rate (Table 2, overall Pass@1) with agent execution time, agent steps, and output tokens per rollout. Each resource is averaged within each task and then equally across tasks with recorded values, using the same selected rollouts as Table 2. A step denotes one tool invocation. The gray bands mark the empirical Pareto frontiers for lower resource use and higher pass rate. Figures 17–19 show how each resource is distributed over tasks and models, in the layout of Figures 11–14. Their colors follow a logarithmic scale, and the few values beyond its ends take the end colors. Appendix G.1 gives the measurement definitions and coverage.

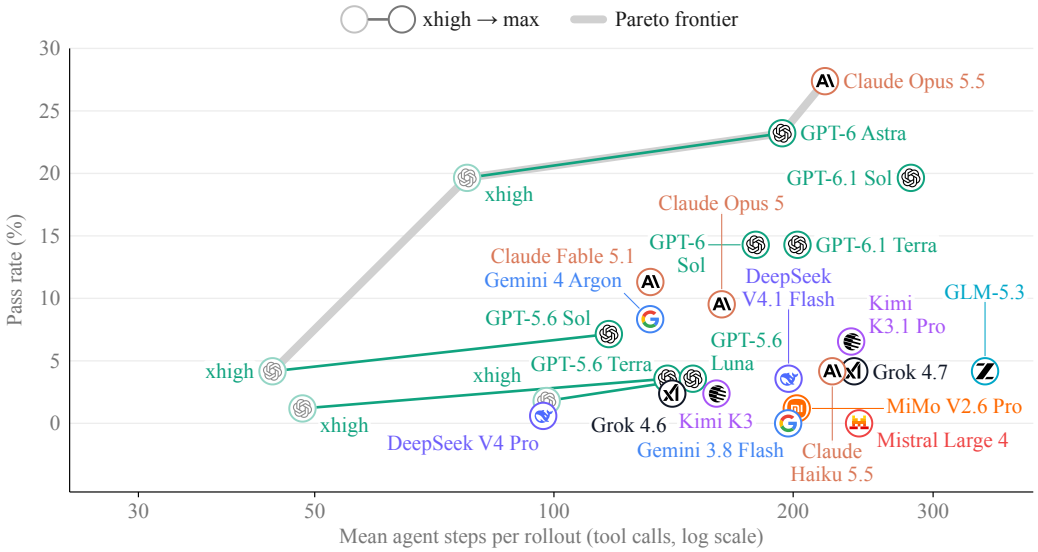

Figure 15: Mean agent steps per rollout versus overall pass rate (%). Steps count recorded tool invocations, including failed and interrupted calls.

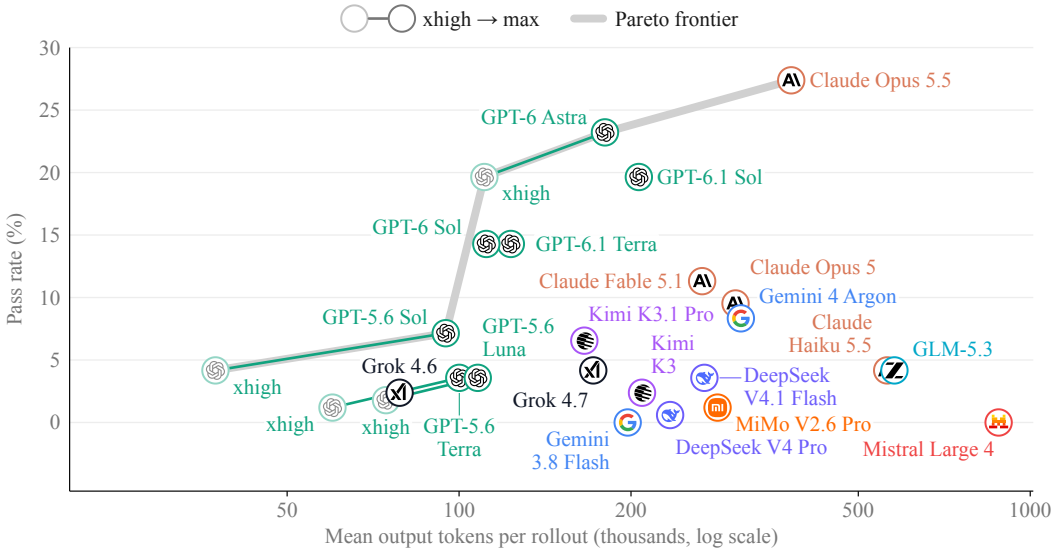

Figure 16: Mean reported output tokens per rollout, in thousands, versus overall pass rate (%). Token counts follow provider or harness usage accounting.

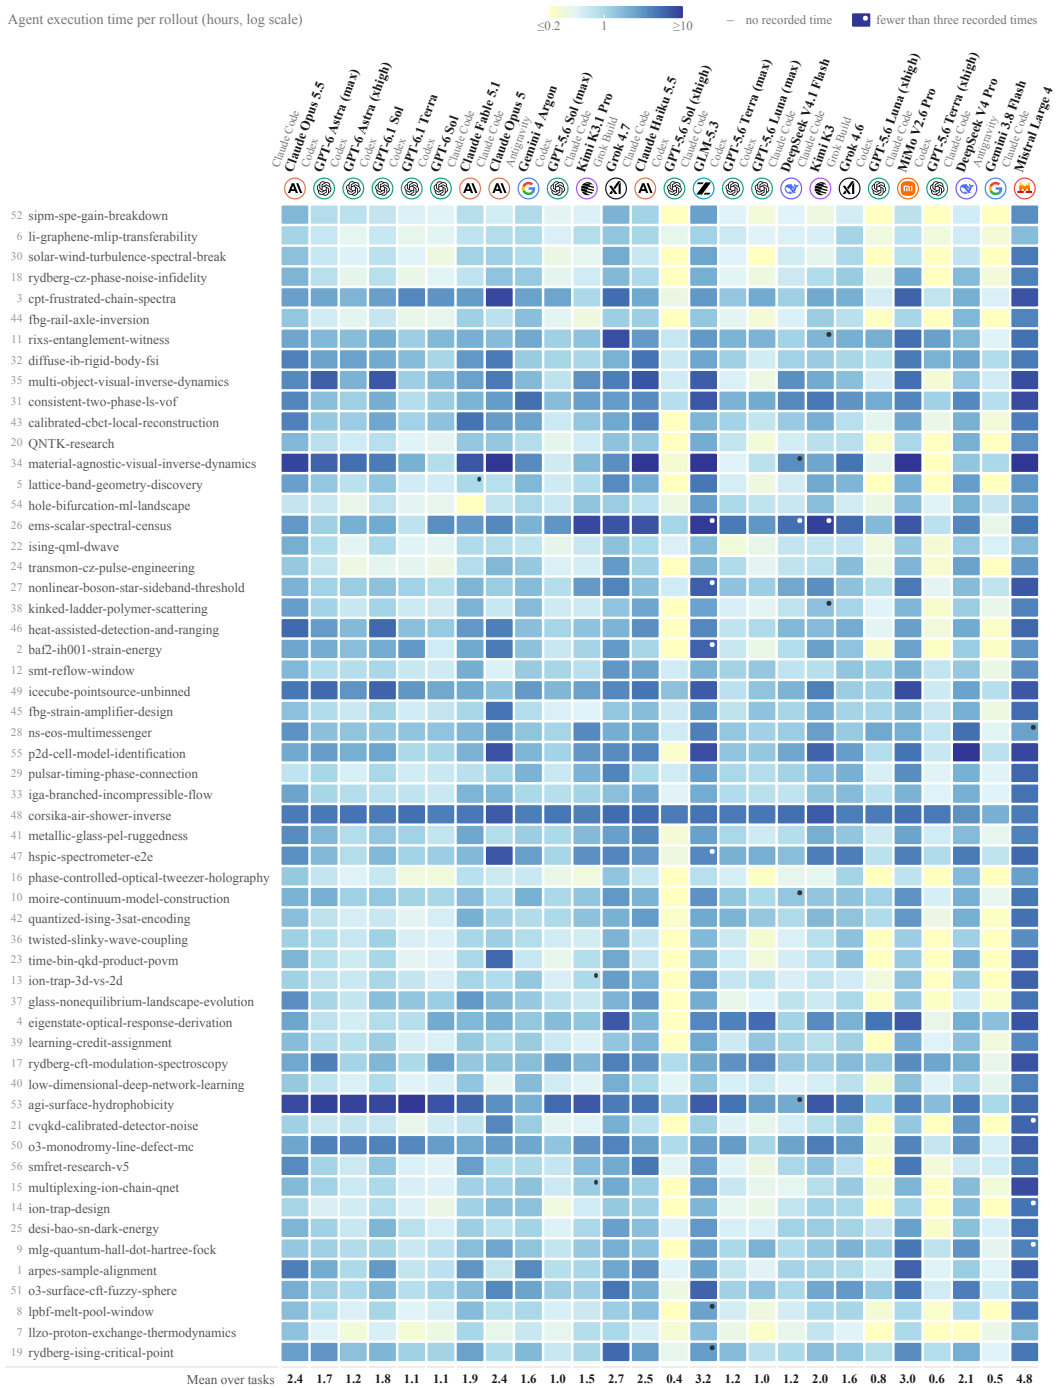

Figure 17: Agent execution time per rollout, in hours, per task and model: the mean over the recorded rollouts in each cell. The bottom row gives each model's mean over tasks with records, as plotted in Figure 5. A dash marks a task without a recorded time and a dot one with fewer than three.

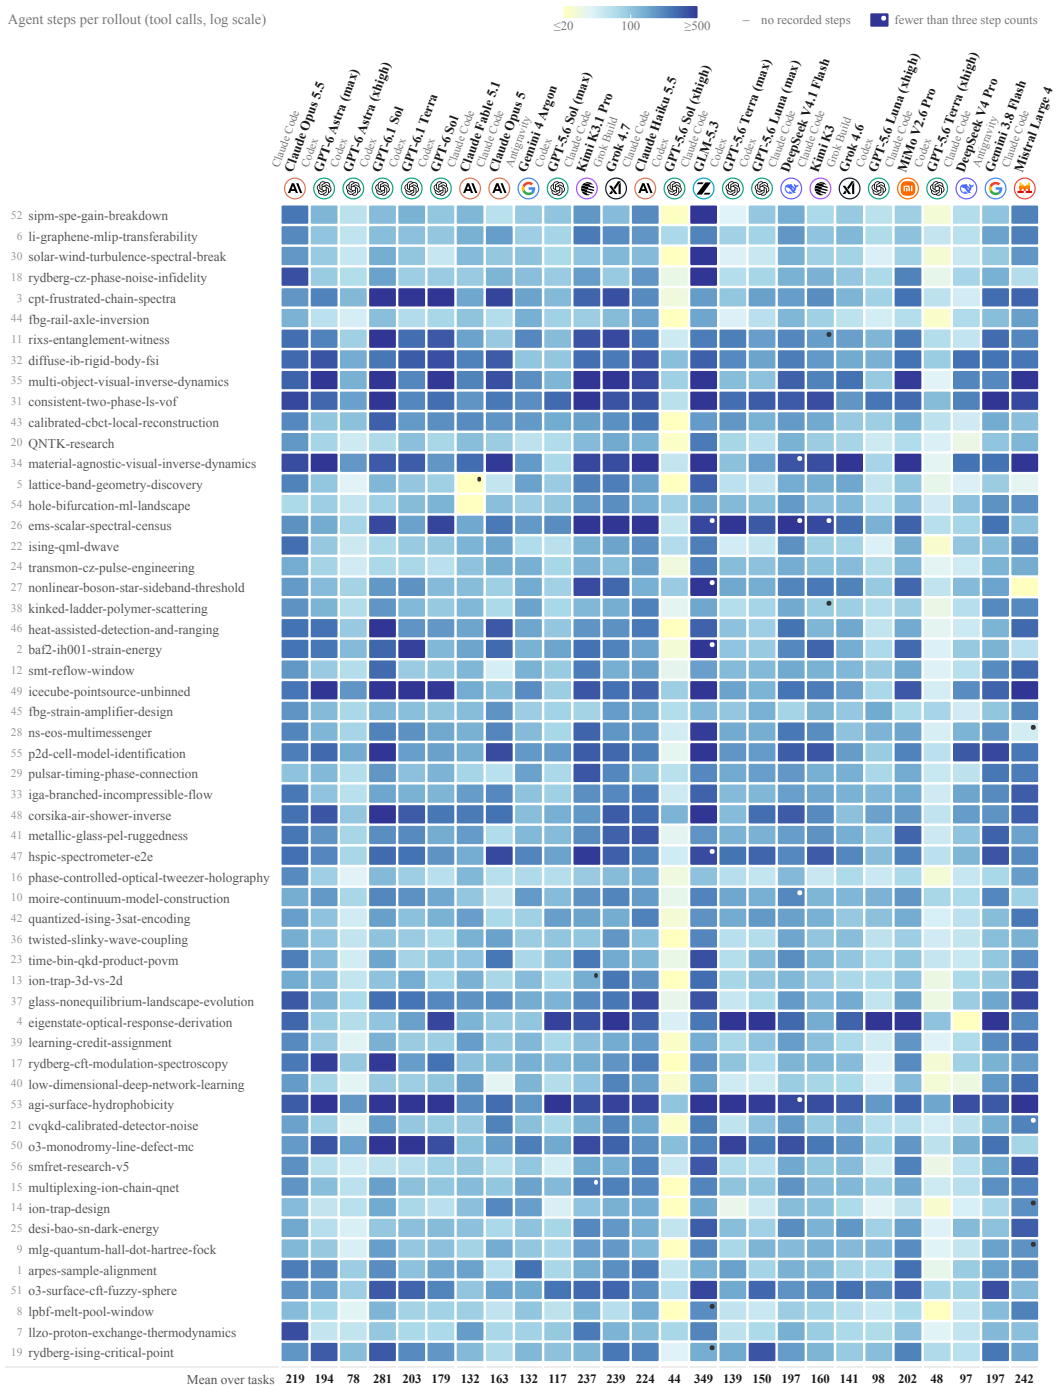

Figure 18: Agent steps (tool calls) per rollout, per task and model: the mean over the recorded rollouts in each cell. The bottom row gives each model's mean over tasks with records, as plotted in Figure 15. Rows follow Figure 11. A dash marks a task without a recorded step count and a dot one with fewer than three.

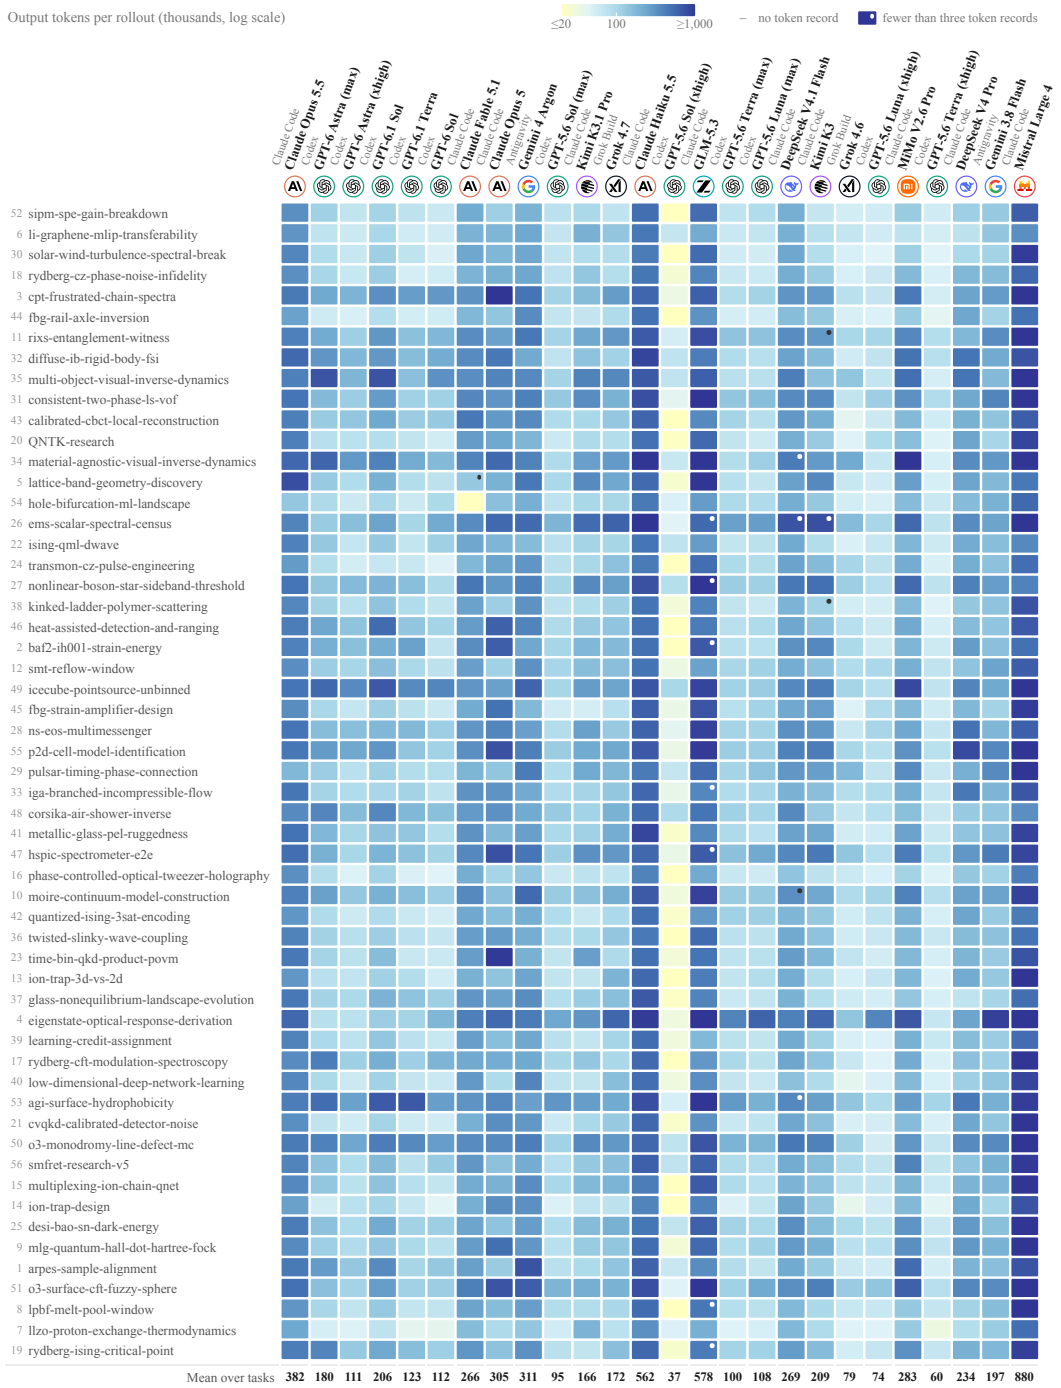

Figure 19: Output tokens per rollout, in thousands, per task and model: the mean over the recorded rollouts in each cell. The bottom row gives each model's mean over tasks with records, as plotted in Figure 16. Rows follow Figure 11. A dash marks a task without an output-token record and a dot one with fewer than three.

G.1 Agent Resource Accounting

The resource plots use the same selected rollouts and 56-task overall pass rate as the cost comparison. For each resource, available values are averaged within each task and then equally across tasks. Unavailable counters are omitted from resource means, not replaced with zero. Table 9 reports the number of rollouts and tasks contributing to each resource estimate.

Agent time is the recorded execution-stage wall-clock duration, excluding environment setup, verification and rubric review. When an idle-timeout record lacks that stage timer, its recorded elapsed execution time supplies the duration. Agent steps count recorded tool invocations, including failed and interrupted calls. A partial trace is used only when a termination monitor supplies a tool-call counter. Output tokens use the reported agent output-token count; unavailable usage records are omitted. Tokenization and output-token accounting can differ between providers. In Claude Code runs, agent time and steps include subagent work.

Table 9: Resource coverage: contributing rollouts / tasks for each model.

Model	Agent time	Agent steps	Output tokens
Claude Opus 5.5	168 / 56	168 / 56	168 / 56
GPT-6 Astra (max)	168 / 56	168 / 56	168 / 56
GPT-6 Astra (xhigh)	168 / 56	168 / 56	168 / 56
GPT-6.1 Sol	168 / 56	168 / 56	168 / 56
GPT-6.1 Terra	168 / 56	168 / 56	168 / 56
GPT-6 Sol	168 / 56	168 / 56	168 / 56
Claude Fable 5.1	166 / 56	166 / 56	166 / 56
Claude Opus 5	168 / 56	168 / 56	168 / 56
Gemini 4 Argon	168 / 56	168 / 56	168 / 56
GPT-5.6 Sol (max)	168 / 56	168 / 56	168 / 56
Kimi K3.1 Pro	165 / 56	165 / 56	168 / 56
Grok 4.7	168 / 56	168 / 56	168 / 56
Claude Haiku 5.5	168 / 56	168 / 56	168 / 56
GPT-5.6 Sol (xhigh)	168 / 56	168 / 56	168 / 56
GLM-5.3	160 / 56	160 / 56	159 / 56
GPT-5.6 Terra (max)	168 / 56	168 / 56	168 / 56
GPT-5.6 Luna (max)	168 / 56	168 / 56	168 / 56
DeepSeek V4.1 Flash	162 / 56	162 / 56	162 / 56
Kimi K3	165 / 56	165 / 56	165 / 56
Grok 4.6	168 / 56	168 / 56	168 / 56
GPT-5.6 Luna (xhigh)	168 / 56	168 / 56	168 / 56
MiMo V2.6 Pro	168 / 56	168 / 56	168 / 56
GPT-5.6 Terra (xhigh)	168 / 56	168 / 56	168 / 56
DeepSeek V4 Pro	168 / 56	168 / 56	168 / 56
Gemini 3.8 Flash	168 / 56	168 / 56	168 / 56
Mistral Large 4	164 / 56	164 / 56	168 / 56

Table 10: Resource use per rollout for each model configuration: mean, median (Med.) and 95th percentile (P95), pooled over the rollouts of Table 2 that record the value. These are the same selected rollouts as Figures 5, 15 and 16, and rows follow Table 2. Tool calls count recorded tool invocations, including failed and interrupted calls. A step is one model call of the agent that returned a response, and one step can make several tool calls or none. Time is the execution-stage wall-clock duration in hours, excluding environment setup, verification and rubric review; because it is wall-clock time, slow provider routes and rate limiting inflate it. Cost is the reported inference cost per rollout in US dollars, as in Figure 5 (metered, priced from recorded tokens at list prices, or estimated), and output tokens are in millions. Unavailable values are omitted, not set to zero; † marks a statistic over fewer than all 168 designated rollouts of the configuration (Table 9 gives the coverage of time, tool calls and output tokens). P95 and the median use linear interpolation between order statistics (NumPy's default). The figures instead average within each task and then across tasks; both means coincide wherever all 168 rollouts are recorded. In Claude Code runs, tool calls, steps and time include subagent work.

	Tool calls			Steps			Time (h)			Cost ($)			Output tokens (M)	
Model	Mean	Med.	P95	Mean	Med.	P95	Mean	Med.	P95	Mean	Med.	P95	Mean	Med.
Claude Opus 5.5	219	204	371	191	168	339	2.44	2.06	5.38	26.95	24.86	51.89	0.38	0.39
GPT-6 Astra (max)	194	111	585	108	61	335	1.73	1.05	5.28	26.91	15.24	90.77	0.18	0.13
GPT-6 Astra (xhigh)	78	65	172	86	65	197	1.24	0.75	3.72	19.21	11.60	56.87	0.11	0.09
GPT-6.1 Sol	281	170	882	158	90	583	1.79	1.17	5.21	5.33	3.20	17.51	0.21	0.15
GPT-6.1 Terra	203	124	461	203	119	553	1.14	0.65	3.90	3.46	1.23	14.25	0.12	0.09
GPT-6 Sol	179	114	501	194	110	634	1.08	0.63	3.57	6.38	3.22	22.09	0.11	0.08
Claude Fable 5.1	134†	128†	230†	81	66	190	1.87†	1.58†	4.40†	32.63	26.42	48.65	0.27†	0.27†
Claude Opus 5	163	116	425	156	105	422	2.40	1.35	7.58	56.77	32.85	163.37	0.30	0.22
Gemini 4 Argon	132	122	225	123	119	207	1.61	1.41	3.20	8.87	7.10	17.39	0.31	0.28
GPT-5.6 Sol (max)	117	89	260	154	106	401	0.97	0.62	2.77	10.35	6.93	29.21	0.09	0.09
Kimi K3.1 Pro	238†	195†	530†	187	140	450	1.52†	1.01†	4.08†	9.60	7.03	24.43	0.17	0.13
Grok 4.7	239	212	523	169	140	374	2.75	2.32	6.40	12.19	10.13	27.34	0.17	0.14
Claude Haiku 5.5	224	187	476	214	176	451	2.50	1.79	6.39	6.66†	5.13†	15.27†	0.56	0.52
GPT-5.6 Sol (xhigh)	44	34	102	66	55	168	0.42	0.27	1.33	2.79	1.77	8.43	0.04	0.03
GLM-5.3	347†	261†	1,015†	324	242	994	3.14†	2.46†	7.02†	12.81	10.05	33.40	0.57†	0.47†
GPT-5.6 Terra (max)	139	98	419	165	110	501	1.15	0.73	3.89	5.80	3.73	16.81	0.10	0.09
GPT-5.6 Luna (max)	150	103	401	197	124	525	1.04	0.57	3.44	0.68	0.42	1.78	0.11	0.09
DeepSeek V4.1 Flash	191†	166†	404†	178†	159†	394†	1.17†	0.83†	3.13†	0.65†	0.54†	1.34†	0.26†	0.24†
Kimi K3	159†	123†	358†	153	122	382	1.91†	1.24†	5.51†	9.33	6.64	22.04	0.21†	0.16†
Grok 4.6	141	122	315	83	71	179	1.63	1.27	4.02	9.21†	7.35†	19.75†	0.08	0.06
GPT-5.6 Luna (xhigh)	98	72	224	141	90	414	0.75	0.39	2.18	0.41	0.24	1.10	0.07	0.06
MiMo V2.6 Pro	202	148	497	182	133	434	3.01	2.13	7.47	0.66	0.47	1.56	0.28	0.20
GPT-5.6 Terra (xhigh)	48	42	93	70	52	168	0.58	0.31	2.02	2.22	1.53	5.92	0.06	0.06
DeepSeek V4 Pro	97	66	237	90	63	214	2.10	2.00	5.01	7.82	5.86	20.95	0.23	0.19
Gemini 3.8 Flash	197	157	433	207	165	447	0.45	0.36	1.06	3.69†	2.73†	8.38†	0.20	0.16
Mistral Large 4	244†	205†	634†	205	173	562	4.82†	4.02†	9.25†	10.67	9.29	24.39	0.88	0.69

H Reference-material Evaluation

Table 11: Evaluating rubric scores for all 33 tasks where the authors provided a manuscript or paper as a reference. Only the rubrics, that actually evaluated/ needed the present manuscript were evaluated. If no blocking rubric applied to the manuscript, it is shown as a dash. If no scoring rubric applied to the manuscript it is shown as NA.

Task / reference	Weighted score	Blockers
QNTK-research	100%	2/2
consistent-two-phase-ls-vof	100%	1/1
cpt-frustrated-chain-spectra	N/A	**0/1 - fail**
cvqkd-calibrated-detector-noise	100%	**2/3 - fail**
diffuse-ib-rigid-body-fsi	100%	-
eigenstate-optical-response-derivation	100%	5/5
fbg-rail-axle-inversion	100%	-
fbg-strain-amplifier-design	100%	2/2
glass-nonequilibrium-landscape-evolution - bundled report	**91.94%**	2/2
glass-nonequilibrium-landscape-evolution - original paper	**50%**	-
heat-assisted-detection-and-ranging	100%	-
hole-bifurcation-ml-landscape	**81.58%**	-
hspic-spectrometer-e2e	100%	3/3
icecube-pointsource-unbinned	100%	-
ion-trap-3d-vs-2d - bundled report	100%	4/4
ion-trap-3d-vs-2d - original paper	100%	1/1
ion-trap-design - bundled report	100%	-
ion-trap-design - original paper	100%	-
lattice-band-geometry-discovery	100%	1/1
learning-credit-assignment	100%	4/4
li-graphene-mlip-transferability	**93.52%**	2/2
llzo-proton-exchange-thermodynamics	100%	-
low-dimensional-deep-network-learning	100%	7/7
material-agnostic-visual-inverse-dynamics	100%	-
metallic-glass-pel-ruggedness	**72.73%**	1/1
moire-continuum-model-construction	**80%**	-
multi-object-visual-inverse-dynamics	100%	-
multiplexing-ion-chain-qnet	100%	-
o3-surface-cft-fuzzy-sphere	100%	-
phase-controlled-optical-tweezer-holography	N/A	**0/1 - fail**
rixs-entanglement-witness	100%	-
rydberg-cft-modulation-spectroscopy	100%	4/4
rydberg-cz-phase-noise-infidelity	100%	6/6
rydberg-ising-critical-point	100%	3/3
sipm-spe-gain-breakdown	100%	-
time-bin-qkd-product-povm	100%	5/5