arXiv:2611.07948v1 [cs.AI] 5 Nov 2026

Circe: Training-Free LLM Self-Evaluation via Component Interaction Consistency

Shuyi Miao
Beihang University
Beijing, China

Zheng Li
Nanjing University of Science and Technology
Nanjing, China

Xu Xie
Independent Researcher
China

Linrui Xu
Central South University
Changsha, China

Fei Shen*
National University of Singapore
Singapore, Singapore

Xiaoyu Du
Nanjing University of Science and Technology
Nanjing, China

Tat-Seng Chua
National University of Singapore
Singapore, Singapore

Abstract

As large language models (LLMs) are increasingly deployed in real-world applications, evaluating the correctness of their responses is crucial for reliable use. However, such evaluation is challenging because reference answers are rarely available in practice. Existing reference-free methods mainly rely on output signals, which may assign high scores to incorrect responses and thus do not necessarily reflect correctness. This motivates exploring the model's internal computation for complementary correctness signals. Recent studies exploit internal states, but they mainly analyze hidden-state trajectories across layers, which record only the combined updates of attention and feed-forward networks (FFNs) and overlook how these components interact. We therefore propose Circe, a training-free self-evaluation method that assesses a response using both internal and external signals, without reference answers or repeated sampling. For internal signals, we first introduce component interaction consistency (CIC), which measures the directional alignment between attention updates in earlier layers and attention-induced FFN output changes in later layers. We further incorporate representation geometry (RG) to characterize normalized displacement and angular changes along layerwise hidden-state trajectories. For the external signal, we convert the mean negative log-likelihood of generated tokens into a bounded likelihood confidence (LC) score. Finally, we fuse CIC, RG, and LC into a unified correctness score. Experiments on five LLMs and six benchmarks across four domains show that our method distinguishes correct from incorrect responses more effectively than existing baselines, and generalizes to multilingual, multimodal, and open-world settings.

CCS Concepts

• **General and reference** → **Evaluation**; • **Computing methodologies** → *Artificial intelligence.*

Keywords

Large Language Models, Self-Evaluation, Response Correctness, Latent Knowledge, Interpretability

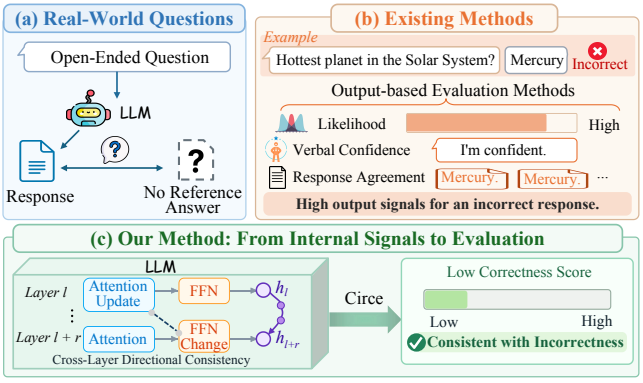

Figure 1: Evaluating LLM response correctness with internal signals. (a) Real-world questions often lack reference answers. (b) Output likelihood, verbal confidence, and response agreement can remain high for incorrect answers. (c) Circe incorporates internal computational signals, yielding a low correctness score for the illustrated incorrect answer.

1 Introduction

Large language models (LLMs) have demonstrated strong capabilities in mathematical reasoning [1], knowledge question answering [42], language understanding [22, 30, 35], and code generation [11, 34]. However, their responses can be fluent and plausible while still being incorrect, making correctness evaluation important for reliable deployment. Reference answers provide a basis for verification, but are often unavailable for questions encountered in real-world applications, as illustrated in Figure 1(a). This motivates reference-free evaluation methods that estimate response correctness from information available during or after generation.

Existing methods for reference-free correctness evaluation [19] commonly derive signals from model outputs. These include verbal confidence elicited through prompting [26, 41], probability-based measures such as entropy [18] and perplexity [38], and agreement across responses obtained through prompt perturbation [13] or repeated sampling [43]. However, these signals do not necessarily reflect response correctness: an incorrect answer may receive high probability, be accompanied by a confident self-assessment, or recur consistently across generations, as illustrated in Figure 1(b).

*Corresponding author.

Methods requiring multiple responses also introduce additional inference costs. These limitations motivate us to investigate internal computational signals that can complement output-level evidence.

Recent studies have explored correctness signals in internal representations [8], including layerwise hidden-state trajectories [44]. Such trajectories characterize the overall evolution of representations but leave the interactions between the underlying components implicit. Within a Transformer block, attention aggregates contextual information, and the FFN transforms the updated representation. Their contributions are combined in the residual stream, so similar aggregate changes can arise from different patterns of component interaction. This raises a complementary question: can the relationship between attention updates and subsequent FFN responses provide additional evidence of response correctness? We examine this relationship and observe that correct responses tend to exhibit higher directional alignment between earlier attention updates and attention-induced FFN output changes in later layers. This observation motivates an interaction-based correctness signal alongside aggregate representation trajectories.

Building on this observation, we propose Circe, a training-free self-evaluation metric that combines component interactions, representation geometry, and output confidence. Specifically, we first introduce *Component Interaction Consistency (CIC)* to quantify the cross-layer alignment described above. At each block, we compute the difference between the FFN outputs with and without the local attention update, and compare its direction with attention updates from earlier blocks. Second, we incorporate *Representation Geometry (RG)* to characterize normalized displacement and angular changes along layerwise hidden-state trajectories, providing complementary evidence from overall representation evolution. Finally, *Likelihood Confidence (LC)* transforms the mean negative log-likelihood of generated tokens into a bounded score that captures output-level confidence. Circe first combines CIC and RG into an internal score, which is then fused with LC to produce a unified response-level correctness score, as illustrated in Figure 1(c). The domain-specific fusion weights are selected on development data and fixed at test time. Circe evaluates each test response using internal activations and token probabilities, without reference answers or repeated response sampling. Our main contributions are as follows:

- We identify cross-layer directional consistency between earlier attention updates and attention-induced FFN output changes in later layers as a correctness-related signal, providing component-level evidence that complements aggregate representation trajectories.
- We propose Circe (Component Interaction, Representation, and Confidence for Evaluation), a training-free self-evaluation metric that integrates component interactions, representation geometry, and output confidence to score individual responses without reference answers at inference time or repeated sampling.
- Extensive experiments on five LLMs and six benchmarks spanning four domains demonstrate competitive performance in response correctness evaluation, with ablation studies supporting the complementary roles of the three signals. Additional evaluations examine its applicability to multilingual and multimodal tasks and mixed benchmark distributions.

2 Related Work

Existing approaches to evaluating LLM response correctness can be broadly divided according to whether they rely on reference answers. Supervised methods learn to evaluate correctness from responses annotated against reference answers, whereas reference-free methods derive correctness signals without such references.

Supervised Correctness Evaluation. These methods learn to evaluate correctness from responses labelled as correct or incorrect according to reference answers. Hidden-state probing methods [3, 36] train lightweight classifiers on internal representations, while external verifiers [10, 25] are trained with outcome-level or process-level annotations. These studies show that model representations encode information about response correctness beyond output likelihood or agreement. However, their reliance on reference-based annotations limits their applicability when reference answers are unavailable, and their performance may degrade under distribution shift between training and evaluation. These limitations motivate reference-free approaches to correctness evaluation.

Reference-free Evaluation from Model Outputs. These methods evaluate response correctness using generated text and output distributions. Probability-based methods [18, 38] compute confidence measures such as negative log-likelihood, perplexity, and predictive entropy from token probabilities during generation. Verbal-confidence methods [21, 26, 41] elicit an explicit self-assessment from the model, such as a stated confidence level or the probability it assigns to the correctness of its own answer. Another line of work estimates confidence from consistency across multiple responses. Self-Consistency [43] derives confidence from answer agreement, Semantic Entropy [23] quantifies uncertainty over semantically equivalent responses, SelfCheckGPT [32] checks individual statements against sampled responses, and SPUQ [13] aggregates outputs under prompt and temperature perturbations. Although these require no reference answers, responses that are probable, confident, or mutually consistent can still be incorrect, and methods that require multiple generations incur additional inference cost.

Reference-free Evaluation from Internal States. To exploit the correctness information encoded in internal representations without reference answers, recent work derives correctness scores directly from hidden states. EigenScore [8] measures semantic consistency through the eigenvalues of the covariance matrix of hidden representations, but still requires multiple generations. CoE [44] treats the hidden representations across layers as a Chain-of-Embedding and evaluates correctness from the magnitude and angular changes between adjacent layers, requiring neither reference answers nor repeated sampling. However, such trajectory statistics summarize the aggregate evolution of representations, leaving the interaction between attention and FFNs within each layer implicit. In contrast, our method explicitly measures the directional consistency between attention updates and the FFN output changes they induce, and combines it with representation geometry and output confidence to evaluate the response.

3 Method

We estimate response correctness from internal and external perspectives. The internal signals consist of CIC, which captures

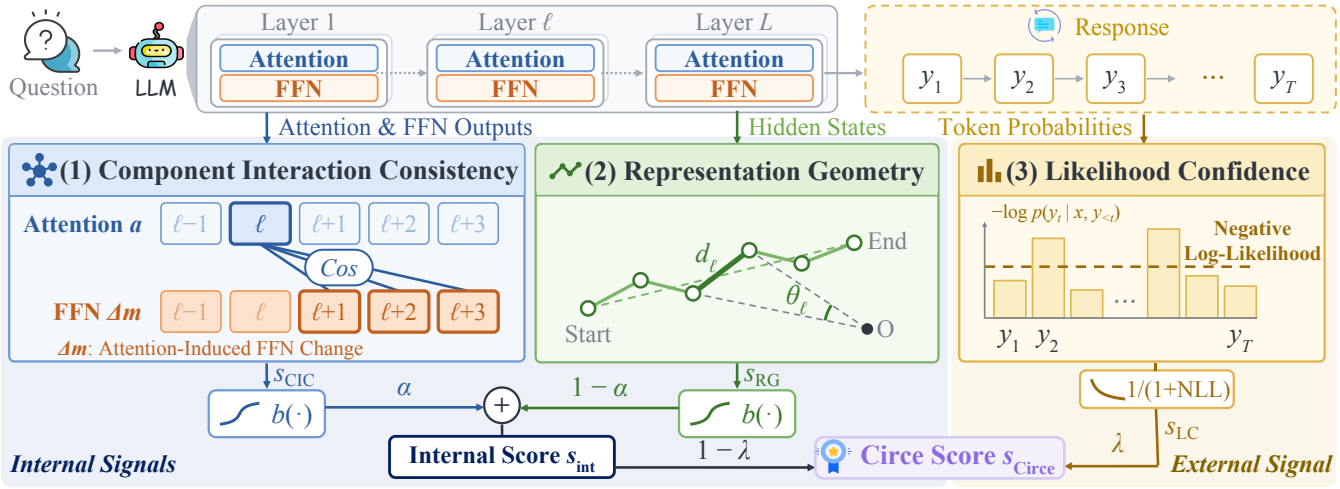

Figure 2: Overview of Circe. (1) Component interaction consistency: cross-layer directional alignment between attention updates and attention-induced FFN changes. (2) Representation geometry: normalized displacement and angular change along the layerwise hidden state trajectory. (3) Likelihood confidence: a bounded transformation of the mean NLL of the generated tokens. Internal signals are fused into an internal score, which is combined with the output signal to yield the final Circe score.

component-level interactions across layers, and RG, which describes the overall evolution of residual representations. The external signal is LC, which reflects the probabilities assigned to generated tokens. As illustrated in Figure 2, CIC and RG are first fused into an internal score, which is then combined with LC to obtain the final Circe score.

3.1 Internal Evaluation Signals

The internal signals examine the model from two complementary levels: CIC analyzes the directional consistency between internal components within each layer, whereas RG analyzes the geometry of the hidden states that each layer finally outputs.

Component Interaction Consistency Score. Attention and FFNs jointly shape information processing across Transformer layers. We examine whether attention updates in earlier blocks align with attention-related changes in FFN outputs in later blocks. We think that stronger alignment reflects more coherent information processing and is associated with a higher likelihood of a correct response. To characterize this interaction, we first quantify how an attention update changes the FFN output within the same block. Consider a Transformer with L blocks and a response of T tokens. At token position $t \in \{1, \dots, T\}$ and block $\ell \in \{1, \dots, L\}$, let $h_{t,\ell-1}$ denote the incoming residual state. Let $a_{t,\ell}$ denote the attention update added to this state, and F_ℓ the FFN computation, including its input normalization. We compare the FFN output with and without the current attention update:

$$\Delta m_{t,\ell} = F_\ell(h_{t,\ell-1} + a_{t,\ell}) - F_\ell(h_{t,\ell-1}). \tag{1}$$

Both evaluations use the same incoming residual state and model parameters. This difference captures the change in FFN output induced by adding the current attention update. Its direction therefore characterizes how attention modifies the FFN contribution to the residual stream within that block. To obtain response-level vectors for each block, we average the attention updates and attention-induced FFN output changes over all response tokens:

$$\bar{a}_\ell = \frac{1}{T}\sum_{t=1}^{T} a_{t,\ell}, \qquad \overline{\Delta m}_\ell = \frac{1}{T}\sum_{t=1}^{T} \Delta m_{t,\ell}. \tag{2}$$

To characterize the persistence of directional alignment across layers, we compare earlier attention updates with locally defined attention-induced FFN output changes in subsequent blocks. We aggregate these cosine similarities to obtain the CIC score:

$$s_{\mathrm{CIC}} = \frac{1}{R}\sum_{r=1}^{R} \frac{1}{L-r}\sum_{\ell=1}^{L-r} \cos\left(\bar{a}_\ell, \overline{\Delta m}_{\ell+r}\right), \tag{3}$$

where r denotes the layer separation, R its maximum value ($1 \le R < L$), and $\cos(\cdot,\cdot)$ cosine similarity between nonzero vectors in the shared residual-stream space. Higher s_{CIC} indicates stronger average alignment between earlier attention directions and attention-induced FFN output changes across subsequent layers, as shown by the CIC scores for Qwen2-7B-Instruct on MATH in Figure 3.

Representation Geometry Score. Following [44], we construct a representation trajectory by averaging the residual states over all response tokens at each layer:

$$\bar{h}_\ell = \frac{1}{T}\sum_{t=1}^{T} h_{t,\ell}, \qquad \ell \in \{0, \dots, L\}. \tag{4}$$

The sequence $(\bar{h}_0, \dots, \bar{h}_L)$ describes how the response-level representation evolves across layers. $h_{t,0}$ is the residual input to the first block, and $h_{t,L}$ is the output of the final block.

We characterize each transition by its displacement and angular change, normalized by the corresponding changes between the

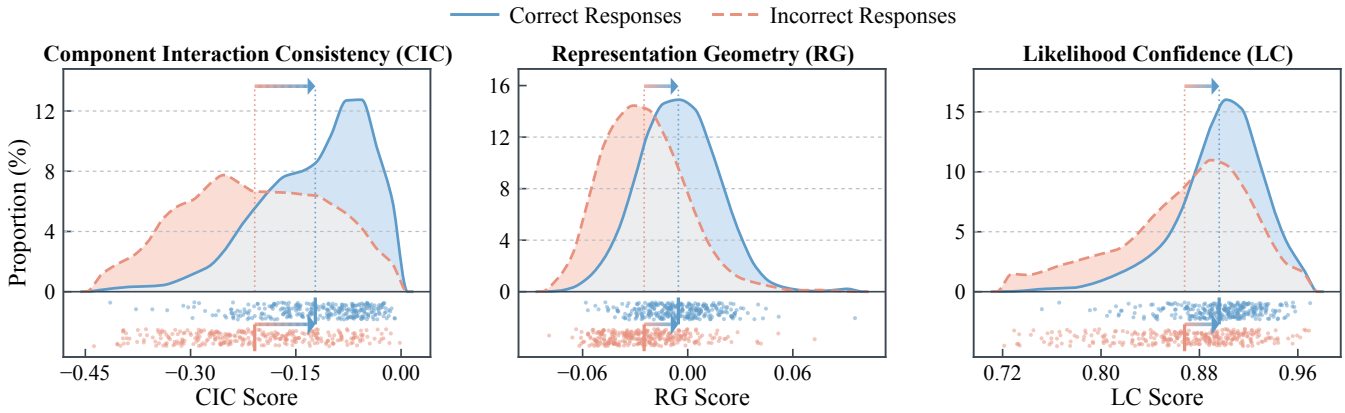

Figure 3: Correct responses have higher mean CIC, RG, and LC scores. Curves show smoothed score distributions for Qwen2-7B-Instruct on MATH, and dots below the curves represent individual responses. Vertical dotted lines and short bars mark the class means, with arrows pointing from the incorrect-response mean to the correct-response mean. Grey shading indicates overlap between the distributions of the two classes.

initial and final states:

$$d_\ell = \frac{\|\bar{h}_{\ell+1} - \bar{h}_\ell\|_2}{\|\bar{h}_L - \bar{h}_0\|_2 + \epsilon}, \tag{5}$$

$$\theta_\ell = \frac{\arccos\big(\cos(\bar{h}_\ell, \bar{h}_{\ell+1})\big)}{\arccos\big(\cos(\bar{h}_0, \bar{h}_L)\big) + \epsilon}, \tag{6}$$

where $\ell \in \{0, \ldots, L-1\}$, $\|\cdot\|_2$ denotes the Euclidean norm, and $\epsilon > 0$ ensures numerical stability in the denominators. Cosine values are clipped to $[-1, 1]$ before applying arccos. We combine the two quantities to obtain the RG score:

$$s_{\text{RG}} = \frac{1}{L}\sum_{\ell=0}^{L-1}(d_\ell - \theta_\ell)\,. \tag{7}$$

This score increases with normalized displacement and decreases with normalized angular change, summarizing representation evolution alongside the component consistency measured by CIC.

3.2 External Evaluation Signal

To complement the two internal signals, we examine the probabilities assigned to the tokens of the generated response $y = (y_1, \ldots, y_T)$ for a prompt x, and summarize them into the LC score. Motivated by the likelihood-based uncertainty term, we compute the mean token negative log-likelihood (NLL):

$$\text{NLL} = -\frac{1}{T}\sum_{t=1}^{T}\log p(y_t \mid x, y_{<t}), \tag{8}$$

where $p(y_t \mid x, y_{<t})$ is the probability assigned to the generated token y_t given the prompt and preceding response tokens. Lower NLL indicates a higher geometric mean of the probabilities assigned to the generated tokens. We transform NLL into the LC score:

$$s_{\text{LC}} = \frac{1}{1 + \text{NLL}}. \tag{9}$$

Higher s_{LC} therefore indicates greater likelihood-based confidence in the generated response. For Qwen2-7B-Instruct on MATH, correct responses have higher mean RG and LC scores than incorrect responses, as shown in Figure 3.

3.3 Constructing the Circe Score

We combine component interaction consistency, representation geometry, and output confidence into the Circe score. Before fusion, we map the two internal scores to $(0, 1)$:

$$b(z) = \frac{1}{2}\left(1 + \frac{z}{1 + |z|}\right), \tag{10}$$

where z is a scalar score. This strictly increasing transformation preserves the ordering of each internal score. We first combine the transformed CIC and RG scores:

$$s_{\text{int}} = \alpha b(s_{\text{CIC}}) + (1 - \alpha) b(s_{\text{RG}}), \tag{11}$$

where $\alpha \in [0, 1]$ balances the contributions of the two internal signals. We then incorporate LC to obtain the final Circe score:

$$s_{\text{Circe}} = (1 - \lambda) s_{\text{int}} + \lambda s_{\text{LC}}, \tag{12}$$

where $\lambda \in [0, 1]$ controls the contribution of output-level confidence. Higher s_{Circe} indicates stronger combined evidence of response correctness. Appendix A provides a theoretical analysis of the proposed scores and the rationale for their fusion.

4 Experiments

4.1 Setup

Datasets. We evaluate the discriminative effectiveness of correctness metrics on six public benchmarks, covering four different task types: mathematics, reasoning, knowledge, and understanding. Specifically, these include mathematics: GSM8K [10] and MATH [16], reasoning: CommonsenseQA [40] and TheoremQA [9], knowledge: MMLU [15], and understanding: Belebele [5].

Models. We select five open-source models with diverse architectures and training origins: Qwen2-1.5B-Instruct [46], Qwen2-7B-Instruct [46], Llama3-8B-Instruct [2], Mistral-7B-Instruct [20], and Qwen2.5-14B-Instruct [47]. The instructions used by all models to generate responses are listed in Appendix D.

Baselines. We compare against 12 reference-free evaluation baselines, grouped by the source of their correctness signals. Output-based methods derive correctness scores from generated text, output probabilities, or logits. Among single-pass methods, we select

Table 1: Evaluation performance across five models and four domains. Results are reported in AUROC (%, ↑), FPR95 (%, ↓), and AUPR (%, ↑). Average denotes the mean over the five models. † denotes methods that require multiple stochastic inferences.

Method	Model					Average
	Qwen2-1.5B-Instruct[46]	Qwen2-7B-Instruct[46]	Llama3-8B-Instruct[2]	Mistral-7B-Instruct[20]	Qwen2.5-14B-Instruct[47]	
Domain I: Mathematics						
Verbal Conf. [21]	52.10 / 88.70 / 38.05	56.21 / 91.42 / 68.23	44.19 / 98.94 / 43.56	43.26 / 97.16 / 27.85	30.06 / 98.47 / 71.97	45.16 / 94.94 / 49.93
PSA pipeline† [45]	69.58 / 84.75 / 52.87	62.28 / 85.53 / 71.15	61.34 / 86.24 / 51.24	75.68 / 79.26 / 44.98	66.57 / 77.37 / 87.69	67.09 / 82.63 / 61.59
Max Prob. [17]	63.27 / 82.28 / 44.11	68.40 / 75.67 / 73.82	58.47 / 84.60 / 54.81	65.53 / 81.09 / 37.33	62.03 / 79.27 / 82.95	63.54 / 80.58 / 58.60
Perplexity [38]	63.78 / 82.07 / 44.48	68.83 / 73.09 / 74.23	58.32 / 84.71 / 54.72	66.37 / 79.40 / 37.93	62.47 / 79.11 / 83.21	63.95 / 79.68 / 58.91
Entropy [18]	65.41 / 80.92 / 46.00	70.25 / 73.45 / 74.97	60.17 / 84.02 / 55.57	66.55 / 82.02 / 38.27	63.34 / 78.43 / 83.46	65.14 / 79.77 / 59.65
Temp. Scaling [14]	62.39 / 83.24 / 43.46	69.11 / 76.98 / 73.12	58.72 / 84.33 / 54.77	64.98 / 82.48 / 37.03	61.55 / 79.85 / 82.90	63.35 / 81.38 / 58.26
Energy [28]	74.80 / 71.62 / 61.11	56.10 / 88.52 / 70.76	53.26 / 92.64 / 50.32	49.07 / 96.57 / 26.94	71.98 / 76.94 / 88.50	61.04 / 85.26 / 59.53
MC Dropout† [12]	46.45 / 95.95 / 36.08	55.64 / 90.53 / 68.58	48.75 / 96.25 / 47.18	49.09 / 94.12 / 27.26	49.82 / 94.95 / 79.99	49.95 / 94.36 / 51.82
LN-Entropy† [31]	70.92 / 78.06 / 54.45	66.63 / 71.09 / 73.84	61.95 / 85.61 / 56.63	67.71 / 80.16 / 39.25	66.28 / 74.25 / 85.18	66.70 / 77.83 / 61.87
EigenScore† [8]	63.76 / 90.09 / 50.22	62.56 / 88.89 / 67.10	52.77 / 95.45 / 43.00	41.65 / 96.77 / 24.11	62.87 / 83.10 / 84.88	56.72 / 90.86 / 53.86
CoE-R [44]	71.85 / 78.86 / 54.34	75.70 / 69.30 / 83.37	73.19 / 72.30 / 66.84	68.38 / 82.83 / 38.41	60.46 / 86.57 / 83.05	69.91 / 77.97 / 65.20
CoE-C [44]	69.66 / 81.29 / 53.59	73.63 / 72.53 / 81.87	74.03 / 77.52 / 69.19	71.00 / 79.57 / 40.20	61.97 / 88.74 / 85.23	70.06 / 79.93 / 66.02
Ours	**79.04** / **63.63** / **65.48**	**77.68** / **64.91** / **84.69**	**79.21** / **62.86** / **73.72**	**75.84** / **66.18** / **47.01**	**75.40** / **63.17** / **90.56**	**77.43** / **64.15** / **72.29**
Domain II: Reasoning						
Verbal Conf. [21]	50.96 / 93.09 / 37.44	43.25 / 99.15 / 42.73	50.01 / 97.74 / 45.62	44.83 / 96.71 / 36.51	40.57 / 96.90 / 54.94	45.92 / 96.72 / 43.45
PSA pipeline† [45]	58.57 / 97.30 / 38.75	48.98 / 97.59 / 49.53	51.25 / 94.76 / 47.08	46.07 / 97.30 / 38.84	60.46 / 91.43 / 65.24	53.07 / 95.68 / 47.89
Max Prob. [17]	63.84 / 83.03 / 39.63	54.57 / 91.77 / 54.21	48.14 / 95.10 / 44.43	53.90 / 91.50 / 39.53	60.23 / 85.08 / 63.18	56.14 / 89.30 / 48.20
Perplexity [38]	64.05 / 84.26 / 39.91	55.39 / 93.33 / 54.78	48.40 / 95.62 / 44.53	54.70 / 90.84 / 40.02	60.74 / 84.75 / 63.54	56.66 / 89.76 / 48.56
Entropy [18]	63.84 / 79.51 / 39.71	55.97 / 92.69 / 54.80	48.56 / 95.55 / 44.64	54.83 / 90.04 / 40.12	60.92 / 83.51 / 63.50	56.82 / 88.26 / 48.55
Temp. Scaling [14]	63.55 / 83.40 / 39.75	54.06 / 94.04 / 53.87	48.12 / 95.47 / 44.39	53.43 / 92.23 / 39.23	60.11 / 85.27 / 63.28	55.85 / 90.08 / 48.10
Energy [28]	61.84 / 77.87 / 41.74	47.52 / 95.16 / 50.07	43.69 / 97.02 / 41.29	48.34 / 95.61 / 36.84	67.55 / 82.72 / 70.76	53.79 / 89.68 / 48.14
MC Dropout† [12]	45.51 / 95.44 / 31.62	50.14 / 97.17 / 49.86	46.21 / 96.60 / 42.07	51.29 / 95.09 / 37.44	49.33 / 94.05 / 57.89	48.50 / 95.67 / 43.78
LN-Entropy† [31]	67.32 / 74.00 / 43.75	56.07 / 94.25 / 54.26	52.88 / 95.36 / 40.92	59.74 / 92.26 / 44.77	61.24 / 87.15 / 64.14	59.45 / 88.60 / 49.57
EigenScore† [8]	61.53 / 88.55 / 40.38	53.39 / 94.11 / 53.39	53.58 / 95.94 / 45.80	60.70 / 91.12 / 48.29	62.32 / 87.78 / 67.71	58.30 / 91.50 / 51.11
CoE-R [44]	69.30 / 80.81 / 48.80	67.18 / 83.75 / 64.88	63.08 / 90.10 / 54.96	72.22 / 79.18 / 58.09	62.27 / 86.72 / 67.70	66.81 / 84.11 / 58.89
CoE-C [44]	69.88 / 76.38 / 48.81	62.36 / 90.33 / 63.76	55.85 / 90.12 / 50.48	69.39 / 82.93 / 57.59	64.22 / 84.68 / 69.41	64.34 / 84.89 / 58.01
Ours	**74.62** / **68.25** / **52.98**	**67.94** / **80.45** / **66.09**	**63.48** / **86.23** / **55.12**	**74.22** / **74.84** / **59.35**	**69.46** / **79.73** / **74.35**	**69.94** / **77.90** / **61.58**
Domain III: Knowledge						
Verbal Conf. [21]	57.49 / 97.28 / 56.95	51.25 / 96.61 / 49.62	46.37 / 99.56 / 59.34	53.52 / 92.78 / 51.75	35.38 / 95.79 / 72.13	48.80 / 96.40 / 57.96
PSA pipeline† [45]	49.79 / 94.68 / 50.22	**60.25** / 88.81 / 70.10	54.95 / 86.34 / 69.64	56.67 / 92.43 / 55.45	63.67 / 87.72 / 87.14	57.07 / 90.00 / 66.51
Max Prob. [17]	60.54 / 90.83 / 57.47	57.09 / 95.31 / 71.76	49.92 / 93.99 / 66.13	53.15 / 96.28 / 52.42	67.18 / 85.88 / 87.99	57.58 / 92.46 / 67.15
Perplexity [38]	61.13 / 90.22 / 58.04	57.26 / 94.79 / 72.04	50.50 / 92.90 / 66.59	53.41 / 95.95 / 52.95	67.50 / 85.66 / 88.15	57.96 / 91.90 / 67.55
Entropy [18]	60.77 / 90.34 / 57.29	57.80 / 93.75 / 71.84	50.12 / 92.35 / 66.25	54.39 / 94.59 / 53.42	68.08 / 84.81 / 88.36	58.23 / 91.17 / 67.43
Temp. Scaling [14]	60.30 / 90.81 / 58.13	56.81 / 93.75 / 71.40	49.52 / 93.99 / 65.86	53.18 / 96.96 / 52.14	66.78 / 86.12 / 87.86	57.32 / 92.33 / 67.08
Energy [28]	56.98 / 92.98 / 53.30	48.73 / 97.72 / 64.79	46.90 / 95.77 / 61.74	46.58 / 99.08 / 45.73	70.85 / 84.25 / 90.20	54.01 / 93.96 / 63.15
MC Dropout† [12]	52.62 / 94.29 / 51.82	55.64 / 95.88 / 68.78	49.87 / 93.16 / 65.98	50.24 / 98.43 / 50.69	48.25 / 95.40 / 79.07	51.32 / 95.43 / 63.27
LN-Entropy† [31]	60.63 / 90.09 / 59.22	58.01 / 92.68 / 71.99	55.36 / 91.14 / 68.90	60.42 / 90.36 / 57.63	69.68 / 83.79 / 89.04	60.82 / 89.61 / 69.36
EigenScore† [8]	50.82 / 95.04 / 50.80	50.36 / 96.35 / 67.53	58.30 / 91.80 / 69.73	61.68 / 89.53 / 62.75	64.15 / 86.51 / 86.69	57.06 / 91.85 / 67.50
CoE-R [44]	49.33 / 93.36 / 50.45	57.98 / 87.50 / 71.54	60.19 / 88.13 / 72.19	62.88 / 88.16 / 61.04	64.57 / 90.80 / 87.60	58.99 / 89.59 / 68.56
CoE-C [44]	49.57 / 93.81 / 50.83	57.56 / 88.14 / 71.45	55.68 / 89.87 / 68.27	63.73 / 88.73 / 62.53	66.03 / 92.11 / 88.67	58.51 / 90.53 / 68.35
Ours	**63.70** / **87.00** / **61.23**	59.79 / **86.22** / **75.14**	**66.60** / **85.37** / **77.44**	**67.20** / **86.24** / **65.26**	**72.38** / **83.29** / **91.01**	**65.93** / **85.62** / **74.02**
Domain IV: Understanding						
Verbal Conf. [21]	49.83 / 96.83 / 71.11	51.18 / 95.67 / 79.65	47.42 / 99.10 / 83.64	50.21 / 97.33 / 45.24	40.23 / 83.33 / 93.06	47.77 / 94.45 / 74.54
PSA pipeline† [45]	57.31 / 91.55 / 74.44	62.07 / 93.20 / 92.15	52.07 / 89.64 / 89.37	54.19 / 94.88 / 50.82	63.48 / 80.95 / 97.04	57.82 / 90.04 / 80.76
Max Prob. [17]	60.02 / **87.32** / 74.38	60.52 / 88.35 / 91.80	56.64 / 95.19 / 90.42	60.82 / 92.32 / 56.70	64.96 / 83.33 / 97.03	60.59 / 89.30 / 82.07
Perplexity [38]	60.78 / 92.25 / 75.46	60.93 / 87.38 / 91.80	56.64 / 94.23 / 90.44	61.38 / 91.92 / 57.50	65.34 / 83.33 / 96.99	61.01 / 89.82 / 82.44
Entropy [18]	59.80 / 88.03 / 73.15	62.65 / 86.41 / 92.13	56.78 / 92.31 / 90.43	61.87 / 91.11 / 58.11	66.40 / 85.71 / 97.15	61.50 / 88.71 / 82.19
Temp. Scaling [14]	60.30 / 89.44 / 76.55	60.45 / 88.35 / 91.84	56.63 / 95.19 / 90.38	60.35 / 92.12 / 56.07	64.10 / 83.33 / 96.94	60.37 / 89.69 / 82.36
Energy [28]	56.43 / 91.90 / 70.16	53.26 / 97.38 / 81.67	44.78 / 98.06 / 80.14	50.78 / 95.64 / 50.22	68.98 / 76.19 / 97.14	54.85 / 91.83 / 75.87
MC Dropout† [12]	50.49 / 94.01 / 67.82	58.23 / 91.42 / 87.09	51.20 / 97.89 / 85.68	53.97 / 95.36 / 52.55	43.48 / 100.00 / 94.03	51.47 / 95.74 / 77.43
LN-Entropy† [31]	61.76 / 91.55 / 76.55	59.65 / 90.50 / 88.61	55.42 / 92.68 / 88.79	58.71 / 89.13 / 49.54	**71.31** / 80.95 / **97.72**	61.37 / 88.96 / 80.24
EigenScore† [8]	52.34 / 94.72 / 70.27	63.59 / 90.29 / 92.79	56.32 / 91.35 / 90.19	60.32 / 90.71 / 53.58	59.25 / 95.24 / 96.82	58.36 / 92.46 / 80.73
CoE-R [44]	53.61 / 92.96 / 72.84	73.55 / 72.92 / 95.24	68.55 / 86.36 / 94.57	62.99 / **82.09** / 85.65	63.98 / 85.71 / 96.46	64.54 / 84.01 / 88.95
CoE-C [44]	53.88 / 92.96 / 73.15	72.47 / 78.12 / 95.00	61.09 / 90.91 / 92.99	63.36 / 88.56 / 85.87	65.48 / 85.71 / 96.74	63.25 / 87.25 / 88.75
Ours	**63.42** / 87.68 / **79.57**	**75.10** / **66.67** / **95.27**	**73.61** / **78.41** / **95.62**	**66.89** / 83.08 / **87.18**	70.21 / **66.67** / 97.16	**69.85** / **76.50** / **90.96**

P(True) [21] for verbal confidence (verbal conf.) because of its versatility. Maximum softmax probability (max prob.) [17], perplexity [38], entropy [18], temperature scaling (temp. scaling) [14], and energy [28] compute confidence or uncertainty measures from output probabilities or logits. Among methods requiring multiple inferences, the PSA pipeline [45] aggregates confidence over prompted and sampled responses, Monte Carlo dropout (MC dropout) [12]

Table 2: Component ablation of our method across five models. AUROC (%, ↑) is averaged over six benchmarks for each model. "only" uses a single component, and "w/o" removes the component from the full method.

	Components			AUROC					
Variant	CIC	RG	LC	Qwen2-1.5B-Instruct[46]	Qwen2-7B-Instruct[46]	Llama3-8B-Instruct[2]	Mistral-7B-Instruct[20]	Qwen2.5-14B-Instruct[47]	Average
CIC only	✓	✗	✗	59.22	68.72	66.21	66.33	69.01	65.90
LC only	✗	✗	✓	62.93	60.28	53.77	58.19	62.44	59.52
RG only	✗	✓	✗	64.21	69.55	66.88	67.84	62.33	66.16
w/o LC	✓	✓	✗	65.32	70.36	68.88	70.14	69.68	68.87
w/o CIC	✗	✓	✓	69.38	70.43	69.46	70.43	68.83	69.71
w/o RG	✓	✗	✓	70.64	69.34	67.38	68.65	70.02	69.21
Ours	✓	✓	✓	**72.41**	**71.02**	**70.93**	**72.37**	**72.05**	**71.76**

aggregates predictions across stochastic forward passes, and length-normalized entropy (LN-Entropy) [31] estimates uncertainty over multiple sampled responses. Internal-state methods derive correctness scores from internal representations without training on correctness-labelled data. EigenScore [8] measures consistency among the hidden representations of multiple generated responses, while CoE-R and CoE-C [44] assess correctness from the magnitude and angular changes of hidden-state trajectories within a response.

Evaluation Metrics. We adopt AUROC [7], FPR95, and AUPR [33] to evaluate how well each method distinguishes correct from incorrect responses. We treat correct responses as the positive class. AUROC measures the overall ranking quality between correct and incorrect responses across all decision thresholds. FPR95 denotes the proportion of incorrect responses that are wrongly accepted as correct when the recall of correct responses reaches 95%. AUPR reflects the precision–recall trade-off for correct responses and remains informative under class imbalance. For the mathematics and reasoning domains, we report the macro average over their two benchmarks so that each benchmark contributes equally.

Hyperparameters. Our method has two fusion hyperparameters: α balances CIC and RG, while λ controls the contribution of LC. In addition, we set the maximum layer separation in CIC to $R = 4$. For each benchmark, we randomly split the questions into disjoint development and test sets with a ratio of 1:4 using a fixed random seed, and the same split is shared by all five models. Only the development sets are used for hyperparameter selection, and all reported results are computed on the test sets. Since the relative reliability of the signals varies across domains, we select one (α, λ) pair for each domain and share it across all five models and all three metrics. The selected pairs are $(0.30, 0.07)$ for mathematics, $(0.07, 0.05)$ for reasoning, $(0.96, 0.21)$ for knowledge, and $(0.33, 0.06)$ for understanding. Details of the search procedure are provided in Appendix B.

4.2 Main Results

Table 1 reports the evaluation performance of our method and the baselines across five models and four domains. Among the baselines, the strongest method varies across domains: CoE-C performs best in the mathematics domain, CoE-R in the reasoning and understanding domains, and LN-Entropy in the knowledge domain, indicating that no single baseline consistently reflects response correctness. In contrast, our method achieves the best average AUROC, FPR95, and AUPR in all four domains, improving AUROC over the strongest baseline by 7.37%, 3.13%, 5.11%, and 5.31% in the mathematics, reasoning, knowledge, and understanding domains, respectively. Notably, our method requires only a single generated response, whereas sampling-based baselines rely on multiple stochastic inferences. However, the improvement is not uniform across settings, as our method achieves the highest AUROC in 18 of the 20 model–domain pairs, with the exceptions of Qwen2-7B-Instruct in the knowledge domain and Qwen2.5-14B-Instruct in the understanding domain, where it is slightly lower than the PSA pipeline and LN-Entropy, respectively. Overall, these results suggest that combining internal computational signals with output confidence provides more reliable evidence for self-evaluation. We further compare our method with supervised methods that train classifiers on the hidden states of LLMs in Appendix C.

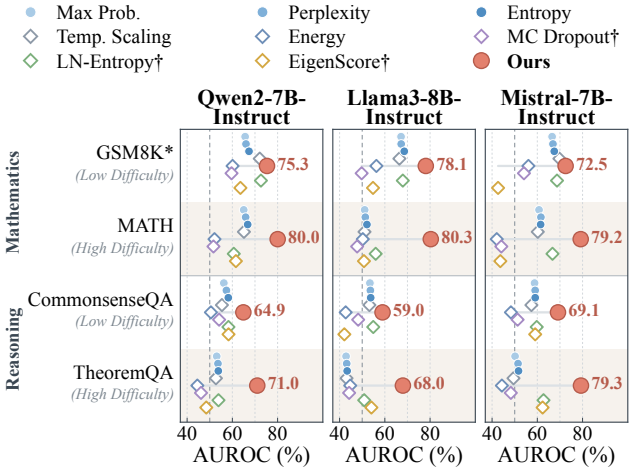

Figure 4: Self-evaluation across task difficulties. AUROC (%) on mathematics and reasoning benchmarks of low and high difficulty with three models. Our method (red circles) achieves the highest AUROC on all benchmarks, and its margin over the strongest baseline is larger on high-difficulty benchmarks. Dashed lines indicate chance-level performance (50%).

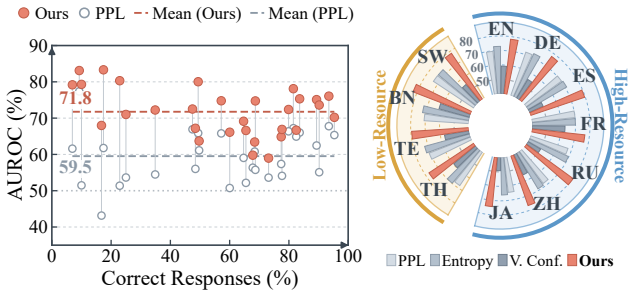

(a) Data ratio robustness. (b) Multilingual generalization.

Figure 5: Robustness to data ratio and generalization across languages. (a) AUROC (%) of Circe and Perplexity on all 30 combinations of the five models and six benchmarks, plotted against the proportion of correct responses. Dashed lines indicate the mean AUROC. (b) AUROC (%) of our method and baselines on the 11 language versions of MGSM with Llama3-8B-Instruct.

4.3 Ablation Study

To assess each component's contribution, we evaluate CIC, RG, and LC individually, their pairwise combinations, and the complete score. Table 2 reports the AUROC over six benchmarks for each model. The complete score achieves the best AUROC on all five models, with an average of 71.76%. Removing any component consistently degrades performance, reducing the average AUROC by 2.05%, 2.55%, and 2.89% without CIC, RG, and LC, respectively. Notably, the contribution of a component is not determined by its standalone performance. LC alone yields the lowest average AUROC of 59.52%, yet removing it from the complete score causes a drop comparable to that of the other two components. Similarly, CIC and RG achieve comparable average AUROC when used alone, at 65.90% and 66.16% respectively, but their relative strength varies across models, with RG outperforming CIC on Qwen2-1.5B-Instruct and CIC outperforming RG by 6.68% on Qwen2.5-14B-Instruct. Consistent with this complementarity, every pairwise combination outperforms each component on average. Overall, these results indicate that CIC, RG, and LC capture complementary correctness evidence, supporting their integration into a unified score.

4.4 Extended Analysis

Effect of Task Difficulty. Tasks within the same domain can differ substantially in difficulty, which may affect the reliability of self-evaluation signals. We therefore compare a lower-difficulty and a higher-difficulty benchmark in each of the mathematics and reasoning domains, namely GSM8K and MATH for mathematics, and CommonsenseQA and TheoremQA for reasoning. As shown in Figure 4, our method achieves the highest AUROC on all four benchmarks across the three models, and its margin over the strongest baseline is consistently larger on the higher-difficulty benchmark in each domain. On MATH, our method achieves an AUROC of 79.2%–80.3%. On TheoremQA, our method achieves an AUROC of 68.0%–79.3%. Moreover, our method achieves a higher AUROC on the higher-difficulty benchmark in every model–domain pair, whereas most baselines show no such improvement. One possible explanation is that more difficult questions elicit more complex internal computation, so that attention–FFN interactions and representation changes reveal correctness-related patterns that output confidence alone fails to capture.

Table 3: Multimodal generalization on Geometry3K with Qwen2.5-VL-7B. Our method outperforms baselines on AUROC (%, ↑), AUPR (%, ↑), and FPR95 (%, ↓).

Method	AUROC	AUPR	FPR95
Max Prob. [17]	61.60	64.28	88.84
Perplexity [38]	61.91	64.80	88.84
Entropy [18]	61.71	64.68	91.32
CoE-R [44]	65.86	72.44	80.58
CoE-C [44]	64.08	71.71	90.08
Ours	**71.25**	**77.67**	**77.27**

Data Ratio Robustness. The proportion of correct responses depends on model accuracy and varies substantially across tasks, so a reliable correctness metric should remain discriminative regardless of whether correct responses form the majority or the minority. Figure 5 (a) plots the AUROC of our method and perplexity against the proportion of correct responses for all 30 combinations of the five models and six benchmarks, where this proportion ranges from 6.9% for Mistral-7B-Instruct on MATH to 95.3% for Qwen2.5-14B-Instruct on Belebele. Across this range, our method remains above the random level of 50% in all combinations and shows no systematic decline toward either extreme, whereas perplexity falls below 50% in several low-accuracy combinations. Overall, our method achieves higher average AUROC than perplexity across all four domains. The three exceptions are Qwen2-1.5B-Instruct on MMLU and Belebele, and Qwen2.5-14B-Instruct on CommonsenseQA, where the gap for the latter is only 1.3%. These results indicate that our method is robust to varying class proportions, which is important in practice because the prevalence of incorrect responses changes with model capability and task difficulty.

Multilingual Generalization. To assess the applicability of Circe in multilingual settings, we further examine its performance on MGSM, a multilingual benchmark released by Shi et al. [37], in which GSM8K problems are manually translated into ten typologically diverse languages, yielding 11 language versions together with English. Using Llama3-8B-Instruct, we evaluate each language version separately and compute AUROC within each language. We adopt the language grouping of Shi et al. [37], which divides the languages according to their frequency in pre-training data into seven high-resource languages, namely English (EN), German (DE), Spanish (ES), French (FR), Russian (RU), Chinese (ZH), and Japanese (JA), and four low-resource languages, namely Thai (TH), Telugu (TE), Bengali (BN), and Swahili (SW). As shown in Figure 5 (b), our method achieves higher AUROC than perplexity, entropy, and verbal conf. in all 11 languages. Moreover, its advantage is not confined to high-resource languages and remains evident in the low-resource group. These results indicate that the correctness signals captured by our method generalize across languages, including those that are underrepresented in pre-training data.

Multimodal Generalization. To evaluate whether our method generalizes to tasks beyond text, we further apply it to Geometry3K [29], a geometry problem-solving benchmark in which each

Table 4: Robustness under open-world distribution shift. AUROC (%, ↑) on mixed math and mixed reasoning with Llama3-8B-Instruct and Mistral-7B-Instruct.

Model	Method	Mixed Math	Mixed Reasoning
Llama3-8B-Instruct[2]	Verbal Conf. [21]	79.26	60.03
	Max Prob. [17]	47.42	48.29
	Entropy [18]	48.37	49.65
	CoE-R [44]	78.47	60.31
	CoE-C [44]	79.89	54.81
	Ours	**83.17**	**62.04**
Mistral-7B-Instruct[20]	Verbal Conf. [21]	57.93	54.36
	Max Prob. [17]	64.81	49.49
	Entropy [18]	66.40	49.44
	CoE-R [44]	62.77	61.70
	CoE-C [44]	72.36	58.88
	Ours	**72.87**	**62.65**

problem is paired with a diagram, so that visual information is involved in response generation. Following the evaluation protocol of GeoR [6], we use Qwen2.5-VL-7B [4], an open-source vision–language model, to generate responses. In addition to output-based baselines, we compare with CoE-R and CoE-C, which estimate correctness from hidden-state trajectories. As shown in Table 3, our method achieves an AUROC of 71.25%, exceeding the strongest baseline, CoE-R, by 5.39%, and also achieves the highest AUPR and the lowest FPR95 among all methods. Notably, these results are obtained with the same fusion weights as in the text-only experiments and without additional training, indicating that our method transfers to visual–textual settings and captures correctness evidence beyond hidden-state trajectories alone.

Robustness under Open-World Distribution Shift. In practical deployment, questions often come from heterogeneous sources with different characteristics and difficulty levels. To evaluate whether our method remains effective under such distribution shifts, we construct two mixed settings that simulate open-world scenarios, with equal contributions from each source: mixed math, which combines GSM8K and TheoremQA, and mixed reasoning, which combines CommonsenseQA and MMLU. As shown in Table 4, our method achieves the highest AUROC on both mixtures with both Llama3-8B-Instruct and Mistral-7B-Instruct. For instance, on mixed math with Llama3-8B-Instruct, our method achieves an AUROC of 83.17%, surpassing the strongest baseline, CoE-C at 79.89%, by 3.28%. These results indicate that our method remains robust when responses from different sources are evaluated together.

Computational Efficiency. To evaluate the computational cost of our method, we report the mean execution time per response on GSM8K with Llama3-8B-Instruct. Table 5 reports the costs of base LLM inference, softmax computation, and the additional computation of our method. All methods require base LLM inference to generate a response, while softmax computation converts output logits into probability distributions over the vocabulary. The latter is required by baselines such as perplexity and entropy. For our method, the reported additional computation covers all computations introduced by our method during and after response generation. It takes 0.67s per response, corresponding to approximately 18.0% of the base inference time and giving a reconstructed total execution time of 4.36s per response. In contrast, methods based on sampling require multiple stochastic inferences, so their computational cost increases with the number of sampled responses. Overall, these results show that our method combines signals from internal states and model outputs while keeping the additional computational cost below that of base inference in this setting.

Table 5: Computational efficiency on GSM8K with Llama3-8B-Instruct. Values are mean ± standard deviation. For our method, the reported time covers all additional computations during and after response generation.

Operation	Time (s)
Base LLM inference	3.69 ±1.27
Softmax computation	0.07 ±0.21
Circe computation	**0.67** ±0.43

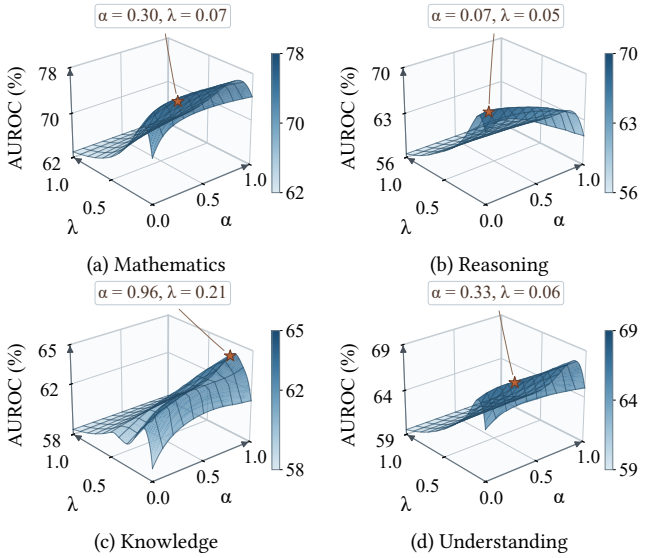

Figure 6: Hyperparameter sensitivity across four domains. Surface height and color indicate the AUROC (%) of our method averaged over the benchmarks in each domain and the five models. Stars mark the selected setting in each domain. The AUROC range is set separately for each domain.

Hyperparameter Sensitivity Across Domains. Figure 6 shows the AUROC of our method over the search grid in each domain, averaged over the test sets of the benchmarks in the domain and the five models. Overall, AUROC varies smoothly with the hyperparameters, and each domain contains a broad high-AUROC region rather than an isolated peak. In the mathematics, reasoning, and understanding domains, AUROC remains close to its maximum over a wide range of α when λ is small, so the results do not rely on a precise choice of either hyperparameter. In the knowledge domain, AUROC increases steadily with α and peaks near $\alpha = 1$, indicating that CIC is the more informative internal signal for knowledge questions. In all four domains, AUROC drops markedly as λ approaches 1, where the fused score reduces to LC alone, which confirms the benefit of combining internal signals with LC.

5 Conclusion

In this paper, we study how to evaluate the correctness of LLM responses without reference answers or repeated sampling. We observe that correct responses tend to exhibit higher cross-layer directional consistency between attention updates and the FFN output changes, revealing a component-level correctness signal that is not captured by output confidence or hidden-state trajectories alone. Building on this observation, we propose Circe, a self-evaluation method that assesses a response by combining three complementary signals: component interaction consistency, representation geometry, and likelihood confidence. Extensive experiments on five LLMs and six benchmarks spanning four domains show that our method achieves the best average performance across all evaluation metrics. Further analyses demonstrate that our method remains effective on more difficult tasks, under varying proportions of correct responses, under open-world distribution shift, and in multilingual and multimodal settings, while adding minimal scoring overhead.

References

[1] Janice Ahn, Rishu Verma, Renze Lou, Di Liu, Rui Zhang, and Wenpeng Yin. 2024. Large Language Models for Mathematical Reasoning: Progresses and Challenges. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: Student Research Workshop*. Association for Computational Linguistics, St. Julian's, Malta, 225–237. doi:10.18653/v1/2024.eacl-srw.17

[2] AI@Meta. 2024. Llama 3 Model Card. https://github.com/meta-llama/llama3/blob/main/MODEL_CARD.md

[3] Amos Azaria and Tom Mitchell. 2023. The internal state of an LLM knows when it's lying. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. 967–976.

[4] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibo Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. 2025. Qwen2.5-VL Technical Report. *arXiv preprint arXiv:2502.13923* (2025). arXiv:2502.13923 [cs.CV]

[5] Lucas Bandarkar, Davis Liang, Benjamin Muller, Mikel Artetxe, Satya Narayan Shukla, Donald Husa, Naman Goyal, Abhinandan Krishnan, Luke Zettlemoyer, and Madian Khabsa. 2024. The Belebele Benchmark: a Parallel Reading Comprehension Dataset in 122 Language Variants. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Bangkok, Thailand, 749–775. doi:10.18653/v1/2024.acl-long.44

[6] Jinhe Bi, Danqi Yan, Yifan Wang, Wenke Huang, Haokun Chen, Guancheng Wan, Mang Ye, Xun Xiao, Hinrich Schuetze, Volker Tresp, and Yunpu Ma. 2026. The Geometry of Reasoning: Self-Evaluation via Layerwise Trajectory Evolution. In *Proceedings of the 43rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 306)*. PMLR, 8076–8099. https://proceedings.mlr.press/v306/bi26e.html

[7] Kendrick Boyd, Kevin H Eng, and C David Page. 2013. Area under the precision-recall curve: point estimates and confidence intervals. In *Joint European conference on machine learning and knowledge discovery in databases*. Springer, 451–466.

[8] Chao Chen, Kai Liu, Ze Chen, Yi Gu, Yue Wu, Mingyuan Tao, Zhihang Fu, and Jieping Ye. 2024. INSIDE: LLMs' Internal States Retain the Power of Hallucination Detection. *arXiv preprint arXiv:2402.03744* (2024).

[9] Wenhu Chen, Ming Yin, Max Ku, Pan Lu, Yixin Wan, Xueguang Ma, Jianyu Xu, Xinyi Wang, and Tony Xia. 2023. TheoremQA: A Theorem-driven Question Answering Dataset. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, 7889–7901. doi:10.18653/v1/2023.emnlp-main.489

[10] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. 2021. Training Verifiers to Solve Math Word Problems. *arXiv preprint arXiv:2110.14168* (2021).

[11] Giuseppe Crupi, Rosalia Tufano, Alejandro Velasco, Antonio Mastropaolo, Denys Poshyvanyk, and Gabriele Bavota. 2025. On the effectiveness of llm-as-a-judge for code generation and summarization. *IEEE Transactions on Software Engineering* 51, 8 (2025), 2329–2345.

[12] Yarin Gal and Zoubin Ghahramani. 2016. Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. In *Proceedings of The 33rd International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 48)*. PMLR, 1050–1059.

[13] Xiang Gao, Jiaxin Zhang, Lalla Mouatadid, and Kamalika Das. 2024. Spuq: Perturbation-based uncertainty quantification for large language models. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2336–2346.

[14] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On Calibration of Modern Neural Networks. In *Proceedings of the 34th International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 70)*. PMLR, 1321–1330. https://proceedings.mlr.press/v70/guo17a.html

[15] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring Massive Multitask Language Understanding. In *International Conference on Learning Representations*.

[16] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021. Measuring Mathematical Problem Solving With the MATH Dataset. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, Vol. 1. https://datasets-benchmarks-proceedings.neurips.cc/paper_files/paper/2021/file/be83ab3ecd0db773eb2dc1b0a17836a1-Paper-round2.pdf

[17] Dan Hendrycks and Kevin Gimpel. 2017. A Baseline for Detecting Misclassified and Out-of-Distribution Examples in Neural Networks. In *International Conference on Learning Representations*. https://arxiv.org/abs/1610.02136

[18] Yuheng Huang, Jiayang Song, Zhijie Wang, Shengming Zhao, Huaming Chen, Felix Juefei-Xu, and Lei Ma. 2025. Look before you leap: An exploratory study of uncertainty analysis for large language models. *IEEE Transactions on Software Engineering* 51, 2 (2025), 413–429.

[19] Takumi Ito, Kees van Deemter, and Jun Suzuki. 2025. Reference-free evaluation metrics for text generation: A survey. *arXiv preprint arXiv:2501.12011* (2025).

[20] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, et al. 2023. Mistral 7B. *arXiv preprint arXiv:2310.06825* (2023). https://arxiv.org/abs/2310.06825

[21] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. 2022. Language Models (Mostly) Know What They Know. *arXiv preprint arXiv:2207.05221* (2022).

[22] Jiayi Kuang, Ying Shen, Jingyou Xie, Haohao Luo, Zhe Xu, Ronghao Li, Yinghui Li, Xianfeng Cheng, Xika Lin, and Yu Han. 2025. Natural language understanding and inference with mllm in visual question answering: A survey. *Comput. Surveys* 57, 8 (2025), 1–36.

[23] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. 2023. Semantic Uncertainty: Linguistic Invariances for Uncertainty Estimation in Natural Language Generation. In *International Conference on Learning Representations*.

[24] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. 2023. Inference-Time Intervention: Eliciting Truthful Answers from a Language Model. In *Advances in Neural Information Processing Systems*.

[25] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2024. Let's verify step by step. In *International Conference on Learning Representations*, Vol. 2024. 39578–39601.

[26] Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. Teaching Models to Express Their Uncertainty in Words. *Transactions on Machine Learning Research* (2022). https://openreview.net/forum?id=8s8K2UZGTZ

[27] Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring How Models Mimic Human Falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, 3214–3252. doi:10.18653/v1/2022.acl-long.229

[28] Weitang Liu, Xiaoyun Wang, John Owens, and Yixuan Li. 2020. Energy-based Out-of-Distribution Detection. *Advances in Neural Information Processing Systems* 33 (2020), 21464–21475.

[29] Pan Lu, Ran Gong, Shibiao Jiang, Liang Qiu, Siyuan Huang, Xiaodan Liang, and Song-Chun Zhu. 2021. Inter-gps: Interpretable geometry problem solving with formal language and symbolic reasoning. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. 6774–6786.

[30] Jiani Luo, Xiaoyan Zhao, Yang Zhang, Shuyi Miao, Bingbing Xu, Stefan Konigorski, and Tat-Seng Chua. 2026. Know You Before You Speak: User-State Modeling for LLM Personalization in Multi-Turn Conversation. *arXiv preprint arXiv:2605.24647* (2026).

[31] Andrey Malinin and Mark Gales. 2020. Uncertainty Estimation in Autoregressive Structured Prediction. In *International Conference on Learning Representations*.

[32] Potsawee Manakul, Adian Liusie, and Mark Gales. 2023. SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural*

Language Processing. Association for Computational Linguistics, Singapore, 9004–9017. doi:10.18653/v1/2023.emnlp-main.557

[33] Christopher D Manning, Hinrich Schütze, and Hinrich Schütze. 1999. *Foundations of statistical natural language processing*. Vol. 999. MIT press, Cambridge.

[34] Shuyi Miao, Xinyi Huang, Wangjie Qiu, Fei Shen, Jinchun He, Zhiming Zheng, and Tat-Seng Chua. 2026. ContractLens: Latent Security Knowledge for Malicious Smart Contract Detection. *arXiv preprint arXiv:2610.04992* (2026).

[35] Shuyi Miao, Wangjie Qiu, Pengyang Shao, Canran Xiao, Fei Shen, Zhiming Zheng, and Tat-Seng Chua. 2026. Who Bridges Safety? Identifying and Targeting Cross-Lingual Shared Safety Pathways. *arXiv preprint arXiv:2608.09095* (2026).

[36] Hadas Orgad, Michael Toker, Zorik Gekhman, Roi Reichart, Idan Szpektor, Hadas Kotek, and Yonatan Belinkov. 2025. Llms know more than they show: On the intrinsic representation of llm hallucinations. In *International Conference on Learning Representations*, Vol. 2025. 66880–66913.

[37] Freda Shi, Mirac Suzgun, Markus Freitag, Xuezhi Wang, Suraj Srivats, Soroush Vosoughi, Hyung Won Chung, Yi Tay, Sebastian Ruder, Denny Zhou, et al. 2022. Language models are multilingual chain-of-thought reasoners. *arXiv preprint arXiv:2210.03057* (2022).

[38] Chengcheng Si, Zhe Gan, Zhengyuan Yang, Shuohang Wang, Jianfeng Wang, Jordan Boyd-Graber, and Lijuan Wang. 2023. Prompting GPT-3 to Be Reliable. In *International Conference on Learning Representations*.

[39] Weihang Su, Changyue Wang, Qingyao Ai, Yiran Hu, Zhijing Wu, Yujia Zhou, and Yiqun Liu. 2024. Unsupervised Real-Time Hallucination Detection based on the Internal States of Large Language Models. In *Findings of the Association for Computational Linguistics: ACL 2024*. 14379–14391.

[40] Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. CommonsenseQA: A Question Answering Challenge Targeting Commonsense Knowledge. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, 4149–4158. doi:10.18653/v1/N19-1421

[41] Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D Manning. 2023. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. 5433–5442.

[42] Mingyang Wang, Alisa Stoll, Lukas Lange, Heike Adel, Hinrich Schütze, and Jannik Strötgen. 2025. Bring your own knowledge: A survey of methods for LLM knowledge expansion. In *Proceedings of the First Workshop on Large Language Model Memorization (L2M2)*. 150–168.

[43] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *International Conference on Learning Representations*.

[44] Yiming Wang, Pei Zhang, Baosong Yang, Derek F. Wong, and Rui Wang. 2025. Latent Space Chain-of-Embedding Enables Output-Free LLM Self-Evaluation. In *International Conference on Learning Representations*.

[45] Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. In *International Conference on Learning Representations*, Vol. 2024. 23650–23678.

[46] An Yang, Baosong Yang, Binyuan Hui, et al. 2024. Qwen2 Technical Report. *arXiv preprint arXiv:2407.10671* (2024). https://arxiv.org/abs/2407.10671

[47] An Yang, Baosong Yang, Beichen Zhang, et al. 2024. Qwen2.5 Technical Report. *arXiv preprint arXiv:2412.15115* (2024). https://arxiv.org/abs/2412.15115

Appendix

Appendix A gives a theoretical analysis of our method. It explains what CIC measures and derives when fusing the scores improves AUROC. Appendix B describes how the fusion hyperparameters are selected. Appendix C compares our method with label-based methods that train classifiers on hidden states, and Appendix D lists the instructions that all models follow when generating responses to the benchmark questions. Appendix E discusses whether CIC captures component interactions rather than vector similarity, why our method remains training-free despite using development data, and what CIC contributes beyond representation geometry and likelihood.

A Theoretical Analysis

This section analyzes the three scores of our method and their fusion. Appendix A.1 shows that CIC measures whether later FFNs respond to attention updates in directions consistent with earlier attention updates. Appendix A.2 derives the condition under which fusion improves AUROC, which applies both to the internal score s_{int} that fuses CIC and RG and to the final score s_{Circe} that fuses s_{int} with LC. Throughout, $\|\cdot\|$ denotes the Euclidean norm and $\cos(u,v) = u^\top v/(\|u\|\,\|v\|)$ for nonzero u, v.

A.1 What CIC Measures

We assume that each FFN map F_ℓ, including its input normalization, is twice continuously differentiable on the segments $\{h_{t,\ell-1} + \tau a_{t,\ell} : \tau \in [0,1]\}$, and that the operator norm of its second derivative is bounded by β_ℓ on these segments. This holds for smooth activations such as GELU and SiLU and for normalization layers away from the zero vector. Let $J_{t,\ell}$ denote the Jacobian of F_ℓ at $h_{t,\ell-1}$, so that $J_{t,\ell}\, a_{t,\ell}$ is the linear response of the FFN to the attention update.

Lemma A.1 (Linearization of the attention-induced FFN output change). *For every token t and block ℓ,*

$$\Delta m_{t,\ell} = J_{t,\ell}\, a_{t,\ell} + e_{t,\ell}, \qquad \|e_{t,\ell}\| \le \frac{\beta_\ell}{2}\,\|a_{t,\ell}\|^2. \tag{13}$$

Consequently, $\overline{\Delta m}_\ell = v_\ell + \bar{e}_\ell$ with $v_\ell = \frac{1}{T}\sum_{t=1}^{T} J_{t,\ell}\, a_{t,\ell}$ and $\|\bar{e}_\ell\| \le \frac{\beta_\ell}{2T}\sum_{t=1}^{T}\|a_{t,\ell}\|^2$.

Proof. Write $h = h_{t,\ell-1}$ and $a = a_{t,\ell}$. By Taylor's theorem with integral remainder, $F_\ell(h+a) - F_\ell(h) = J_{t,\ell}\, a + \int_0^1 (1-\tau)\, D^2 F_\ell(h+\tau a)[a,a]\, d\tau$, and the norm of the integral is at most $\beta_\ell\|a\|^2\int_0^1(1-\tau)\,d\tau = \beta_\ell\|a\|^2/2$. Averaging over t and applying the triangle inequality gives the second statement. □

Lemma A.2 (Stability of cosine similarity). *For nonzero $u, v \in \mathbb{R}^d$ and any e with $\|e\| < \|v\|$,*

$$\left|\cos(u, v+e) - \cos(u,v)\right| \le \frac{2\|e\|}{\|v\|}. \tag{14}$$

Proof. Let $\hat{x} = x/\|x\|$. By the Cauchy–Schwarz inequality, $|\cos(u,v+e) - \cos(u,v)| = |\hat{u}^\top(\widehat{v+e} - \hat{v})| \le \|\widehat{v+e} - \hat{v}\|$. Moreover,

$$\begin{aligned}\|\widehat{v+e} - \hat{v}\| &\le \frac{\|(v+e) - v\|}{\|v\|} + \|v+e\|\left|\frac{1}{\|v+e\|} - \frac{1}{\|v\|}\right| \\ &= \frac{\|e\| + \big|\|v\| - \|v+e\|\big|}{\|v\|} \le \frac{2\|e\|}{\|v\|}.\end{aligned}$$

□

Proposition A.3 (CIC as Jacobian-mediated alignment). *Let $\tilde{s}_{\text{CIC}}$ be defined as s_{CIC} with $\overline{\Delta m}_{\ell+r}$ replaced by $v_{\ell+r} = \frac{1}{T}\sum_t J_{t,\ell+r}\, a_{t,\ell+r}$. If $\|\bar{e}_\ell\| < \|v_\ell\|$ for every block ℓ that enters the score, then*

$$\left|s_{\text{CIC}} - \tilde{s}_{\text{CIC}}\right| \le \max_\ell \frac{2\|\bar{e}_\ell\|}{\|v_\ell\|} \le \max_\ell \frac{\beta_\ell\,\frac{1}{T}\sum_t \|a_{t,\ell}\|^2}{\|v_\ell\|}. \tag{15}$$

Proof. s_{CIC} and $\tilde{s}_{\text{CIC}}$ are the same convex combination of cosine terms, so their difference is bounded by the largest difference between corresponding terms. Lemma A.2 with $u = \bar{a}_\ell$, $v = v_{\ell+r}$, and $e = \bar{e}_{\ell+r}$ bounds each term, and Lemma A.1 bounds $\|\bar{e}_{\ell+r}\|$. □

Interpretation. When the FFN maps are nearly linear over the range of the attention updates, that is, when the right-hand side of the bound above is small, CIC is close to $\tilde{s}_{\mathrm{CIC}}$. This approximation compares the response-level attention update of block ℓ with the attention update of block $\ell + r$ after it is mapped through the local Jacobians of the FFN in that block. A high CIC score therefore indicates that later FFNs turn their attention input into output changes whose directions agree with earlier attention updates, which is the cross-layer consistency that CIC is designed to capture.

A.2 Justification for Score Fusion

Since our method fuses its scores with nonnegative weights, we analyze linear fusion under a Gaussian model with a common covariance. Let $S \in \mathbb{R}^k$ collect k scores and assume $S \mid Y = y \sim \mathcal{N}(\mu_y, \Sigma)$ for $y \in \{0, 1\}$, with $\Sigma \succ 0$, $Y = 1$ for correct responses, and $\Delta = \mu_1 - \mu_0$. Although this model simplifies the actual score distributions, it makes explicit how the correlation between scores determines the benefit of fusion.

Lemma A.4 (AUROC of a linear score). *For any $w \neq 0$, the AUROC of $w^\top S$ is $\Phi\big(w^\top\Delta/\sqrt{2\, w^\top\Sigma w}\big)$, where Φ is the cumulative distribution function of the standard normal distribution. It is maximized by $w^\star = c\,\Sigma^{-1}\Delta$ for any $c > 0$, with value $\Phi\big(\sqrt{\Delta^\top\Sigma^{-1}\Delta/2}\big)$.*

Proof. For independent S_1 and S_0 drawn from the two classes, $w^\top(S_1 - S_0) \sim \mathcal{N}(w^\top\Delta,\, 2w^\top\Sigma w)$, so $\Pr(w^\top S_1 > w^\top S_0) = \Phi\big(w^\top\Delta/\sqrt{2w^\top\Sigma w}\big)$. By the Cauchy–Schwarz inequality in the inner product induced by Σ, $(w^\top\Delta)^2 \le (w^\top\Sigma w)(\Delta^\top\Sigma^{-1}\Delta)$, with equality if and only if $w \propto \Sigma^{-1}\Delta$. □

Our method fuses its scores in two steps. The internal score s_{int} combines CIC and RG, and the final score s_{Circe} combines s_{int} with LC, so each step fuses a pair of scores. For such a pair, orient each score so that correct responses score higher on average, and standardize it by its within-class standard deviation, which does not change its AUROC. Let $\delta_1 \ge \delta_2 > 0$ be the resulting standardized mean differences, so that score 1 is the stronger one, and let $\rho \in (-1, 1)$ be their within-class correlation. Since the overall scale of w does not affect AUROC, nonnegative fusion amounts to $w = (1, \gamma)$ with $\gamma \ge 0$.

Proposition A.5 (Condition for a fusion gain). *Under the Gaussian model, the following hold.*

(1) Some $\gamma > 0$ yields a higher AUROC than the stronger score alone if and only if $\rho < \delta_2/\delta_1$.

(2) If $\rho < \delta_2/\delta_1$, the best nonnegative fusion attains AUROC $\Phi(D/\sqrt{2})$ with

$$D^2 = \delta_1^2 + \frac{(\delta_2 - \rho\,\delta_1)^2}{1 - \rho^2}. \tag{16}$$

Proof. By Lemma A.4, the AUROC of $S_1 + \gamma S_2$ is increasing in $g(\gamma) = (\delta_1 + \gamma\delta_2)/\sqrt{1 + 2\gamma\rho + \gamma^2}$, and $g(0) = \delta_1$. For $\gamma > 0$, $g(\gamma) > \delta_1$ is equivalent to $(\delta_1 + \gamma\delta_2)^2 > \delta_1^2(1 + 2\gamma\rho + \gamma^2)$, that is, to $2\delta_1(\delta_2 - \rho\delta_1) + \gamma(\delta_2^2 - \delta_1^2) > 0$. If $\rho \ge \delta_2/\delta_1$, both terms are nonpositive because $\delta_2 \le \delta_1$, so no $\gamma > 0$ improves on the stronger score. If $\rho < \delta_2/\delta_1$, the inequality holds for all sufficiently small $\gamma > 0$, which proves item 1. For item 2, $w^\star \propto \Sigma^{-1}\Delta = \frac{1}{1-\rho^2}(\delta_1 - \rho\delta_2,\ \delta_2 - \rho\delta_1)$, whose components are both positive when $\rho < \delta_2/\delta_1$, so the unconstrained optimum is feasible. Its separation is $\Delta^\top\Sigma^{-1}\Delta = \frac{\delta_1^2 - 2\rho\delta_1\delta_2 + \delta_2^2}{1-\rho^2} = \delta_1^2 + \frac{(\delta_2 - \rho\delta_1)^2}{1-\rho^2}$. □

Complementarity of the Three Scores. Proposition A.5 makes the role of complementarity explicit. A weaker score improves a stronger one only when their within-class correlation is below the ratio δ_2/δ_1 of their separations, and the attainable gain grows with $(\delta_2 - \rho\delta_1)^2/(1 - \rho^2)$. A weak score can therefore still be useful if it is weakly correlated with the stronger one, whereas a strong score that is highly correlated with it adds little. The ablation in Table 2 is consistent with this analysis. LC has the lowest average AUROC among the single scores at 59.52%, yet fusing it with the internal score s_{int} raises the average AUROC from 68.87% to 71.76%. Under the Gaussian model, such a gain requires the likelihood signal to carry information about correctness that is not already captured by the internal signals CIC and RG.

B Hyperparameter Selection and Sensitivity

Circe involves two groups of hyperparameters: the fusion hyperparameters α and λ, which weight the three component scores, and the maximum layer separation R, which determines the layer pairs used in CIC. We describe how each is set below.

B.1 Fusion Weights

Selection Procedure. The two fusion hyperparameters α and λ assign the weights $w_{\mathrm{CIC}} = (1 - \lambda)\alpha$, $w_{\mathrm{RG}} = (1 - \lambda)(1 - \alpha)$, and $w_{\mathrm{LC}} = \lambda$ to CIC, RG, and LC, respectively, which sum to one. We search all such weight vectors on a grid with a step size of 0.01, namely $\{(i/100,\ j/100,\ k/100) \mid i, j, k \in \mathbb{Z}_{\ge 0},\ i + j + k = 100\}$, which contains 5,151 combinations. All selection is performed on the development sets, which are obtained by randomly splitting the questions of each benchmark into development and test sets with a ratio of 1:4. For each domain and each candidate weight, we compute 15 results, one for each of the five models under each of the three metrics, where each result is averaged over the development sets of the benchmarks in the domain. We compare each result with the strongest baseline under the same model and metric, and select the weight that outperforms the strongest baseline in the largest number of the 15 results. The selected weight vector is shared across all five models and all three metrics within the domain and is fixed for all test-set evaluations. The selected weights $(w_{\mathrm{CIC}}, w_{\mathrm{RG}}, w_{\mathrm{LC}})$ are $(0.28, 0.65, 0.07)$ for mathematics, $(0.07, 0.88, 0.05)$ for reasoning, $(0.76, 0.03, 0.21)$ for knowledge, and $(0.31, 0.63, 0.06)$ for understanding, corresponding to $(\alpha, \lambda) \approx (0.30, 0.07)$, $(0.07, 0.05)$, $(0.96, 0.21)$, and $(0.33, 0.06)$, respectively.

B.2 Maximum Layer Separation

The maximum layer separation R controls the range of layer separations included in CIC, and we fix $R = 4$ in all experiments. Table 6 examines this choice across the four domains with Qwen2-7B-Instruct. At $R = 4$, AUROC remains close to the best displayed result in every domain. Increasing R further produces only small, domain-dependent changes rather than a consistent improvement, while the number of layer-pair terms in the score grows with R. $R = 4$ therefore provides a compact multi-lag setting without sacrificing detection performance.

Table 6: Sensitivity to the maximum layer separation. AUROC (%, ↑) on Qwen2-7B-Instruct, averaged over the benchmarks in each domain.

Domain	$R = 1$	$R = 4$	$R = 8$	$R = 16$	$R = 27$
Mathematics	77.37	**77.68**	77.65	77.64	77.67
Reasoning	67.93	**67.94**	67.93	67.87	67.89
Knowledge	59.60	59.79	**59.81**	59.74	59.62
Understanding	74.85	**75.10**	75.06	74.93	74.72

Table 7: Comparison with label-based methods. AUROC (%, ↑) of the label-based methods ITI and MIND and the label-free methods CoE-R, CoE-C, and ours, evaluated on responses generated by Llama3-8B-Instruct.

Method	TruthfulQA	GSM8K	MATH	MMLU
ITI [24]	48.06	48.36	57.51	53.09
MIND [39]	51.95	53.03	56.89	45.83
CoE-R [44]	57.15	71.15	75.23	60.19
CoE-C [44]	60.92	73.11	74.97	55.67
Ours	**62.85**	**76.14**	**79.84**	**66.31**

C Label-Based Evaluation with Hidden States

ITI [24] and MIND [39] are supervised methods that train classifiers on LLM hidden states. ITI trains linear probes on the labelled candidate answers of TruthfulQA [27], and MIND trains a classifier on hallucination labels automatically constructed from Wikipedia. Since these methods also rely on internal states, we compare them with our method on TruthfulQA, GSM8K, MATH, and MMLU with Llama3-8B-Instruct. On all four benchmarks, we evaluate the correctness of responses generated by the LLM itself rather than the labelled candidate answers, so the evaluated responses differ from the training data of both methods in task type, data source, or the distribution of responses.

As shown in Table 7, the supervised methods ITI and MIND remain close to chance on all four benchmarks, with AUROC between 45.83% and 57.51%. Even on TruthfulQA, whose questions are used to train ITI, its probe does not transfer from labelled candidate answers to LLM-generated responses, suggesting that supervised classifiers capture features specific to their training data. In contrast, all three label-free methods exceed 55% AUROC on every benchmark, and our method achieves the highest AUROC on all four benchmarks. These results indicate that fitting a classifier with labels does not guarantee reliable correctness evaluation on new tasks, whereas label-free signals derived from internal states transfer more consistently across different task types.

D Response Generation Instructions

All models generate responses with their official chat templates and the dataset-specific instructions listed below. Each instruction describes the task, asks for step-by-step reasoning, and fixes the format of the final answer so that the answer can be extracted and compared with the reference. The placeholder `{input_data}` is replaced by the test input, and `{answer_type}` is replaced by the answer type provided by TheoremQA.

GSM8K and MGSM

Solve this math problem. Give the reasoning steps before giving the final answer on the last line by itself in the format of "`Answer:`". Do not add anything other than the integer answer after "`Answer:`".
Question: `{input_data}`

MATH

Please reason step by step, and put your final answer within `\boxed{{}}`
Question: `{input_data}`

CommonsenseQA

Answer the following multiple choice common-sense reasoning question. The last line of your response should be of the following format: '`Answer: $LETTER`' (without quotes) where LETTER is one of ABCDE. Think step by step and output the reasoning process before answering.
Question: `{input_data}`

TheoremQA

Below is an instruction that describes a task, paired with an input that provides further context. Write a response that appropriately completes the request.
Instruction:
Please read a math problem, and then think step by step to derive the answer. The answer is decided by Answer Type.
If the Answer type in [bool], the answer needs to be True or False.
Else if the Answer type in [integer, float] , The answer needs to be in numerical form.
Else if the Answer type in [list of integer, list of float] , the answer needs to be a list of number like [2, 3, 4].
Else if the Answer type in [option], the answer needs to be an option like (a), (b), (c), (d).
You need to output the answer in your final sentence like 'Therefore, the answer is ...'.
Question: `{input_data}`
Answer_type: `{answer_type}`

MMLU

Answer the following multiple choice question. The last line of your response should be of the following format: '`Answer: $LETTER`' (without quotes) where LETTER is one of ABCD. Think step by step before answering.
Question: `{input_data}`

Belebele

Answer the following multiple choice reading-comprehension question. The last line of your response should be of the following format: '`Answer: $LETTER`' (without quotes) where LETTER is one of ABCD. Please fully understand the passage and give explanations step by step before answering.
Question: `{input_data}`

E More Discussions

▹ ***Does CIC characterize component interactions or merely vector similarity?***

CIC characterizes component interactions. It does not compare earlier attention updates with raw FFN outputs, but with attention-induced FFN changes. Specifically, Eq. (1) evaluates the same FFN with and without its current attention update while holding the incoming residual state and model parameters fixed, which isolates how attention modifies the FFN contribution within each block. CIC then measures whether these interaction-derived changes in later blocks align with earlier attention directions. Appendix A further shows that, under a local linear approximation, these changes correspond to attention updates transformed by the FFN Jacobian. The compared vectors therefore have a specific computational interpretation, and their alignment reflects cross-layer consistency in attention–FFN processing.

▹ ***Why is Circe training-free despite using development data?***

Training-free refers to the absence of parameter optimization: Circe neither updates the backbone nor trains an additional correctness detector, and its component scores are computed directly from

model activations and token probabilities. Correctness-labelled development data are used only to select the fusion weights, with a single configuration shared across all five models and all three metrics within each domain. Development and test questions are disjoint, and the selected weights remain fixed during test evaluation. Thus, the scoring itself requires no training, and labelled data enter only through the selection of a few hyperparameters.

▹ ***What does CIC contribute beyond representation geometry and likelihood?***

CIC contributes component-level evidence that neither RG nor LC captures. RG summarizes the overall hidden-state trajectory and LC measures output likelihood, but neither separates how attention changes FFN processing. Table 2 confirms this incremental contribution: adding CIC to RG and LC raises the average AUROC from 69.71% to 71.76%, with improvements on all five models. Moreover, although CIC and RG perform comparably alone, combining them outperforms either one for every model. These results indicate that CIC provides correctness evidence complementary to global representation geometry and output confidence.