arXiv:2611.07541v1 [cs.AI] 5 Nov 2026

CST-BENCH: BENCHMARKING COUNTERFACTUAL SPATIAL TRANSITION REASONING WITH CAUSAL TRANSITION ON-POLICY SELF-DISTILLATION

Ying Li
School of Artificial Intelligence and Computer Science
North China University of Technology
Beijing, China
ying.li@ncut.edu.cn

Yongle Cheng
School of Artificial Intelligence and Computer Science
North China University of Technology
Beijing, China
chenyongle071@gmail.com

Jing Li
School of Artificial Intelligence
Jilin University
Changchun, China
lij25@mails.jlu.edu.cn

Xingquan Cai
School of Artificial Intelligence and Computer Science
North China University of Technology
Beijing, China
caixingquan@ncut.edu.cn

Yawei Luo
ReLER Lab
Zhejiang University
Hangzhou, China
yaweiluo@zju.edu.cn

November 6, 2026

ABSTRACT

Spatial reasoning is a fundamental capability of multimodal large language models (MLLMs), yet existing benchmarks mainly focus on static scene understanding. When objects are moved or swapped, models must infer not only spatial relations but also how interventions change the scene state. To address this problem, we introduce CST-Bench, a benchmark for counterfactual spatial reasoning. CST-Bench contains 7,082 examples covering 6 intervention types and 12 transformation operators, with dense causal annotations including intervention sources, affected relations, expected effects, invariant relations, and metric changes. We benchmark general-purpose MLLMs on CST-Bench and further compare post-training strategies on CST-Bench and standard spatial benchmarks. Results show that methods effective for ordinary spatial reasoning can still fail under counterfactual interventions. We further formulate counterfactual spatial reasoning as causal state-transition modeling and propose CT-OPSD, which uses CST-Bench annotations to guide a teacher model in learning counterfactual state transitions and causal reasoning chains. These chains are then distilled into a student model that reasons only from real observations. Overall, experiment results establish CST-Bench as a challenging benchmark and reusable supervision resource for counterfactual spatial reasoning. CT-OPSD further provides an initial demonstration of how its annotations can support model improvement.

1 Introduction

Spatial reasoning benchmarks test whether multimodal large language models (MLLMs) understand geometric structure rather than merely recognizing objects. Existing benchmarks cover 2D relations, depth and distance estimation, and object orientation [Fu et al., 2024, Tong et al., 2024, Ma et al., 2025], but mainly focus on static spatial reasoning.

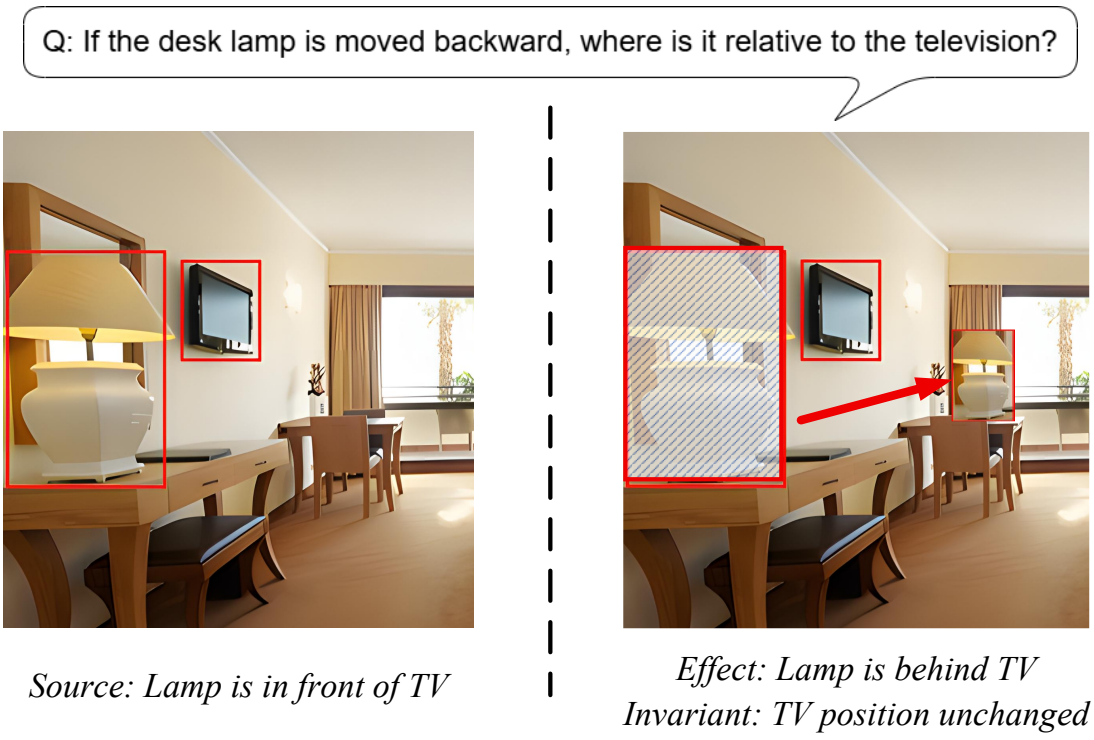

Figure 1: Counterfactual spatial transition reasoning. Moving the lamp backward changes its relation to the TV from *in front of* to *behind*, while unaffected scene properties remain invariant.

However, real-world spatial reasoning also requires predicting how relations change beyond an observed scene. Recent studies extend spatial evaluation to large-scale urban scenes with aerial viewpoints [Zhan et al., 2025], while counterfactual VQA modifies critical image content to test whether models revise their predictions [Niu et al., 2021, Chen et al., 2020]. SiT-Bench [Guo et al., 2026] separates spatial cognition from visual perception, whereas MindEdit-Bench [Yu et al., 2026a] examines object-level counterfactual spatial reasoning in natural images. Nevertheless, existing benchmarks either broaden scene coverage, separate perception from reasoning, or examine isolated counterfactual edits, providing limited supervision for complete intervention-induced spatial transitions.

We introduce CST-Bench, a benchmark centered on causal transitions induced by spatial interventions. It contains 7,082 image-grounded samples from distinct natural scenes, spanning 6 intervention families and 12 fine-grained transition operators. Each sample pairs a single query and target answer with dense annotations of source and target states, affected relations, expected effects, preserved invariants, and applicable metric changes. Final-answer accuracy is the primary metric, while these annotations support diagnostic analysis and structured causal-transition supervision. Figure 1 illustrates one complete transition from an observed state to its counterfactual state.

We benchmark general-purpose MLLMs on CST-Bench and compare several post-training strategies using a common backbone. Results reveal that existing models often fail to update affected relations while preserving invariants. Because this task requires tracing the consequences of an intervention, we formulate it as causal state-transition modeling [Urp'i et al., 2024, Peters, Janzing, and Schölkopf, 2017, Pearl, 2009] and introduce CT-OPSD, a benchmark-aligned causal transition on-policy self-distillation method. Distinct from existing OPSD approaches, CT-OPSD explicitly trains its teacher with structured source states, intervention effects, invariants, and target states, without exposing these privileged annotations to the student. On-policy alignment instead transfers the teacher's causal-transition behavior to a student that uses only the image and question. Overall, CST-Bench provides a fine-grained evaluation and supervision resource, while CT-OPSD demonstrates how its annotations can improve counterfactual spatial reasoning while retaining strong performance on standard spatial benchmarks.

Our main contributions are summarized as follows:

- We formulate counterfactual spatial transition reasoning as a standalone evaluation task requiring MLLMs to predict intervention-induced relational changes and preserved spatial invariants.
- We introduce CST-Bench, which grounds six interven- tion families and twelve transition operators in 7,082 dis- tinct natural scenes, providing systematic coverage of both changed and invariant spatial relations.
- We benchmark general-purpose MLLMs and systematically compare post-training strategies, revealing a significant gap between static spatial understanding and counterfactual transition reasoning.
- We propose CT-OPSD, a benchmark-aligned on-policy self-distillation method that leverages privileged causal-transition supervision to improve counterfactual spatial reasoning while preserving the resulting student model's general spatial capabilities.

2 Related Work

2.1 Spatial Reasoning Benchmarks

Spatial reasoning benchmarks have expanded beyond traditional visual question answering. Large-scale spatial VQA supervision has been used for qualitative relations and metric estimation [Chen et al., 2024], while multifaceted reasoning has been evaluated under common and uncommon camera viewpoints [Ma et al., 2025]. Evaluations now cover large-scale scenes [Zhan et al., 2025] and object, camera, and regional relations across multiple images [Yang et al., 2026]. Recent work also unifies diverse spatial tasks within broader evaluation frameworks [Wu et al., 2026, Cai et al., 2025].

Collectively, these benchmarks broaden spatial evaluation from pairwise relations to scene-level and multi-view reasoning, but mainly evaluate relations within a fixed observed state rather than how that state changes under an intervention. Their results also show that MLLMs perform unevenly across spatial domains, viewpoints, modalities, and reasoning skills. Spatially enhanced supervision may improve specific tasks, but does not guarantee a consistent understanding of relational changes. CST-Bench complements these benchmarks by testing whether models update affected relations after an intervention while preserving unaffected relations.

2.2 Counterfactual Spatial Reasoning

Spatial reasoning benchmarks have expanded from situated reasoning under imagined egocentric perspectives [Ma et al., 2023], ground-truth 3D annotations across indoor, outdoor, and simulated scenes [Cheng et al., 2024], and camera- and human-centered localization [Li et al., 2025], to video-based observation [Yang et al., 2025] and large-scale 3D spatial QA derived from raw video [Gao et al., 2026]. However, these benchmarks primarily recover relations from observed or reconstructed scenes while keeping the underlying scene state unchanged. SiT-Bench separates spatial reasoning from visual perception [Guo et al., 2026], while MindCube evaluates occluded objects and unseen space from limited multi-view observations [Wang et al., 2026]. These settings assess world modeling and reference-frame transformation, but mainly alter the observer rather than the scene state.

Counterfactual reasoning has also been explored in visual QA and event-centric video understanding [Niu et al., 2021, Chen et al., 2020, Yi et al., 2020]. MindEdit-Bench further evaluates object-level counterfactual editing [Yu et al., 2026a]. Nevertheless, existing benchmarks provide limited supervision for diagnosing the complete causal transition induced by an intervention. CST-Bench addresses this gap by formulating counterfactual spatial reasoning as an intervention-induced state transition with dense causal annotations, supporting both final-answer evaluation and diagnosis of transition failures.

2.3 On-Policy Self-Distillation

Traditional knowledge distillation typically relies on fixed teacher trajectories [Hinton, Vinyals, and Dean, 2015, Kim and Rush, 2016, Bengio et al., 2015], which can create a mismatch between training and student inference. Generalized distillation connects knowledge distillation with learning using privileged information [Lopez-Paz et al., 2016, Vapnik and Vashist, 2009]. More recently, on-policy distillation aligns teacher and student over student-generated trajectories [Agarwal et al., 2024], including teacher and student policies constructed from the same model but with different contexts [Zhao et al., 2026]. This paradigm has also been extended to multimodal tasks. Vision-OPD uses local crops as privileged teacher inputs and transfers fine-grained perception to a student that observes only the full image [Yuan et al., 2026]. However, privileged conditions may create advantages that the student cannot learn, resulting in the "privilege illusion" [Yu et al., 2026b]. CT-OPSD instead uses CST-Bench annotations to train a causal teacher and distills its transition reasoning into a student policy; the annotations are used only during training, while the student relies solely on the image and question at inference time.

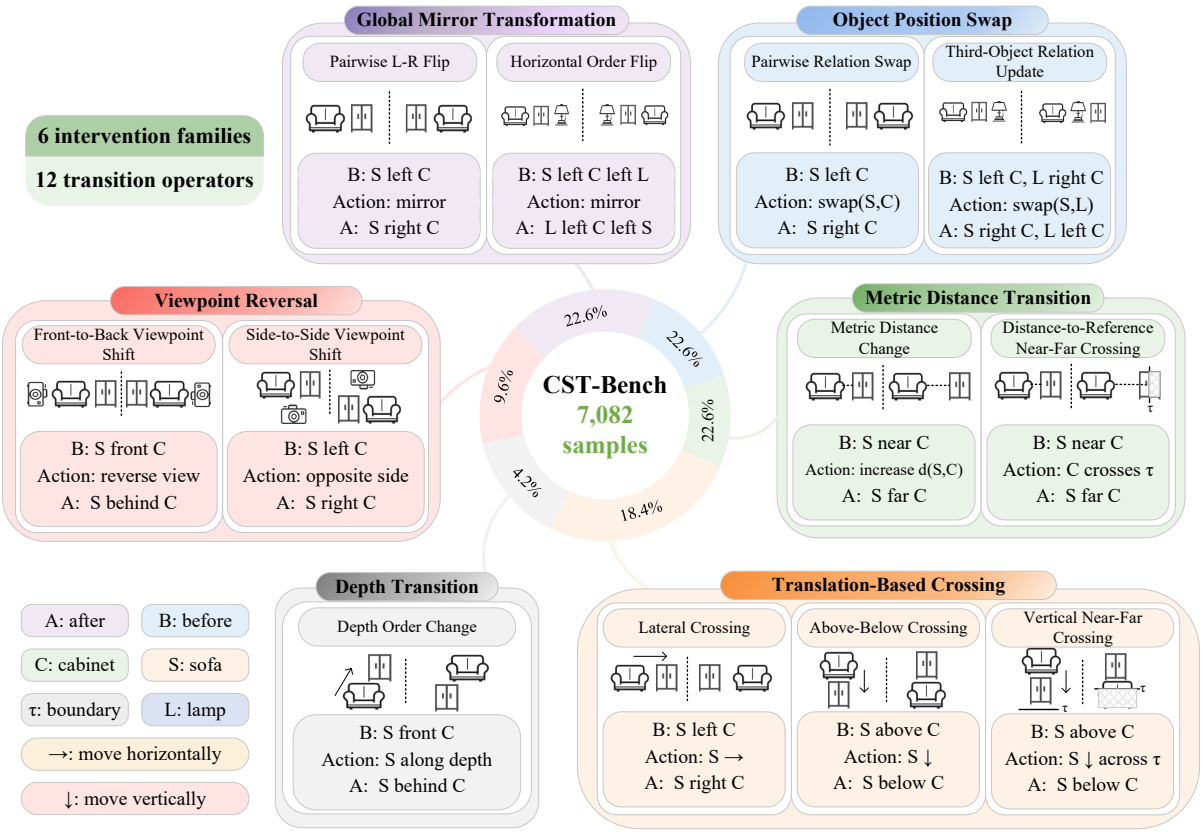

Figure 2: Overview of the CST-Bench transition taxonomy. CST-Bench contains 7,082 counterfactual spatial QA samples across six intervention families and twelve transition operators. The operators cover relational state updates, metric distance changes, and qualitative relation-threshold crossings, with dense causal transition annotations.

3 CST-Bench: A Counterfactual Spatial Reasoning Benchmark

CST-Bench frames counterfactual spatial reasoning as updating a source state after an explicit intervention. It contains six intervention families and twelve fine-grained operators (Figure 2), constructed through deterministic generation, VLM consensus validation, and human verification. Its dense annotations support diagnostic evaluation and enable CT-OPSD, an auxiliary approach that transfers transition reasoning to an observation-only student.

3.1 Task Taxonomy

CST-Bench organizes twelve transition operators into six intervention families covering global, local, metric, translation, depth, and viewpoint changes. Across 7,082 image-grounded samples, it recasts static observations as intervention-conditioned transitions, jointly testing whether models update affected relations and preserve invariants.

Global Mirror Transformation. This family tests globally consistent horizontal updates [Ma et al., 2025]. Pairwise Left–Right Flip reverses a single relation, whereas Multi-object Horizontal Order Flip updates multiple objects consistently. Both preserve object identities, object count, and vertical relations.

Object Position Swap. This family applies identity-preserving local interventions [Yu et al., 2026a]. Pairwise Relation Swap tests the direct effect between exchanged objects, while Third-object Relation Update requires propagating the change to an unchanged reference object.

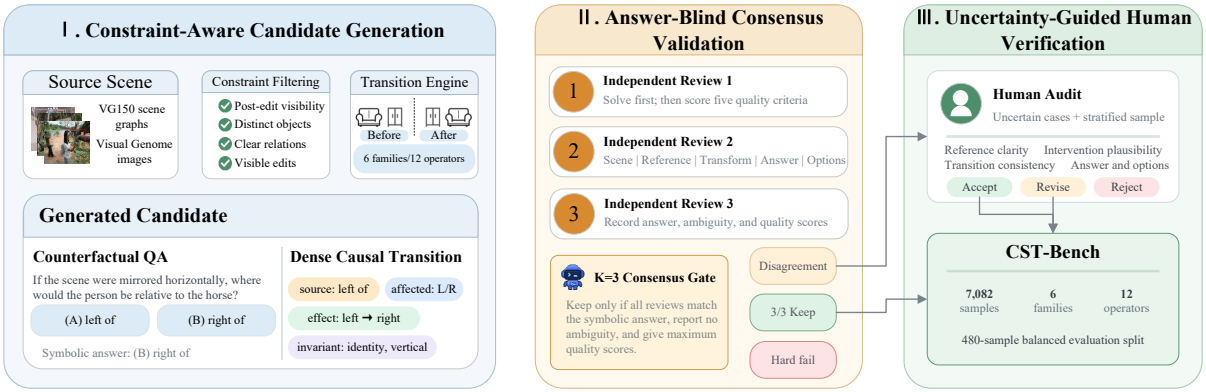

Figure 3: Overview of CST-Bench construction. Constraint-aware generation jointly derives counterfactual states and annotations, answer-blind VLM consensus filters ambiguous candidates, and uncertainty-guided human verification resolves uncertain cases and audits all transformation operators.

Metric Distance Transition. This family distinguishes continuous distance sensitivity from categorical relation updates [Chen et al., 2024, Ma et al., 2025]. Metric Distance Change (Relation Preserved) changes distance while retaining the qualitative relation, whereas Distance-to-Reference Near–Far Crossing moves it across the *near–far* threshold [Cohn and Renz, 2008].

Translation-Based Crossing. This family uses object translation to induce spatial relation changes [Yu et al., 2026a]. Lateral Relation Crossing and Above–Below Crossing test horizontal and vertical updates, while Vertical Near–Far Crossing changes the near–far relation through vertical displacement. Unaffected relations must remain unchanged.

Depth Transition. Depth Order Change moves an object across a reference along the depth axis, extending observed front–behind evaluation to intervention-based state updates [Fu et al., 2024, Ma et al., 2025]. The depth ordering changes while image-plane relations are preserved.

Viewpoint Reversal. This family changes the observer while preserving the physical scene. Front-to-Back and Side-to-Side Viewpoint Shifts require recomputing observer-centered relations without altering the allocentric layout [Li et al., 2025, Wang et al., 2026, Klatzky, 1998].

3.2 Benchmark Construction

As shown in Fig. 3, unlike standard spatial QA, the post-intervention state is not directly observed and must be inferred from a hypothetical intervention [Zhang et al., 2024]. A valid sample must realize the intended effect, propagate dependent relation changes, preserve invariants, and remain visually unambiguous. We address this challenge through three stages.

Stage I: Constraint-Aware Candidate Generation. We curate scenes from the VG150 split of Visual Genome [Xu et al., 2017, Krishna et al., 2016] with valid object boxes, sufficient distinct objects, and geometrically unambiguous relations. Pairwise operators require two usable objects, whereas Third-object Relation Update requires three objects from distinct categories. Relations are computed from normalized bounding boxes, with positional references used when necessary.

Prior work shows that automatically generated spatial QA may contain ambiguous references and shortcut-prone or unreliable labels [Chen et al., 2024, Ma et al., 2025]. To address these risks, a deterministic, operator-aware editor applies each intervention, and a symbolic solver jointly derives the pre- and post-intervention states, counterfactual QA pair, affected relations, expected effects, invariants, and applicable metric changes. Deriving these fields from the same edited state prevents contradictions between the transition and its dense annotations. We reject candidates that violate operator-specific relation constraints, contain invalid boxes or ambiguous relations, provide insufficient relation separation, or move objects outside the image boundary.

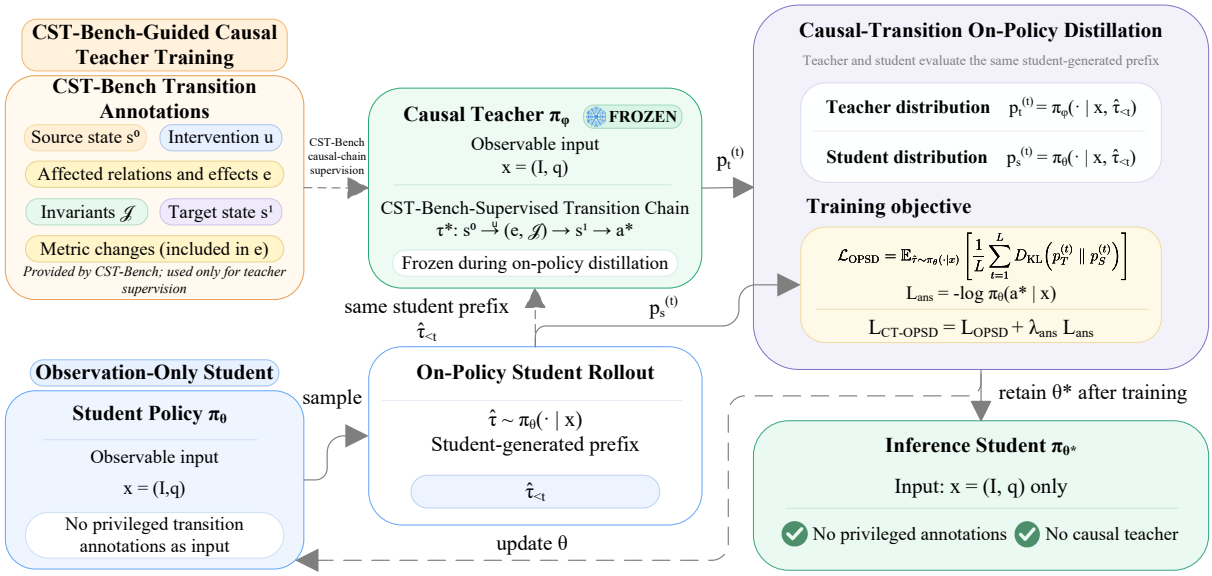

Figure 4: Overview of CT-OPSD with CST-Bench supervision. Privileged transition annotations from CST-Bench train a causal teacher but are never provided as student inputs. The frozen teacher aligns with the observation-only student on student-generated prefixes, and only the updated student is retained for inference.

We score each remaining candidate by

$$s(c) = s_{\text{ref}} + s_{\text{rel}} + s_{\text{edit}} + s_{\text{vis}} + s_{\text{nat}}, \tag{1}$$

where the terms measure reference clarity, relation clarity, edit effectiveness, post-edit visibility, and transformation naturalness. Candidates with $s(c) < 8$ are removed. The remaining candidates are ranked by

$$r(c) = \frac{1}{2}\hat{s}_{\text{op}}(c) + \frac{1}{2}\hat{s}_{\text{global}}(c), \tag{2}$$

where the normalized scores are computed within an operator and globally, respectively. We retain at most one highest-ranked transformation per source image.

Stage II: Answer-Blind Consensus Validation. Although LLM/VLM-based judging provides scalable validation, prior work has identified systematic biases and limited reliability in single-pass evaluation [Zheng et al., 2023, Wang et al., 2023a]. Drawing on multi-sample consistency [Wang et al., 2023b], each candidate is independently reviewed $K = 3$ times by a VLM. To prevent the symbolic label from influencing the initial solution, the VLM first solves the problem without accessing the generated answer and only then compares its prediction with the symbolic label. It also scores scene–image consistency, reference clarity, transformation fidelity, answer correctness, and option quality from 0 to 2. A candidate passes only under unanimous agreement:

$$\text{Keep}(c) = \bigwedge_{k=1}^{K} \left[\hat{a}_k = a^* \wedge b_k = 0 \wedge \min_j q_{k,j} = 2 \right], \tag{3}$$

where $\hat{a}_k$ is the independently predicted answer, a^* is the symbolic answer, b_k indicates ambiguity, and $q_{k,j}$ is the score for criterion j. Candidates with repeated hard failures are removed, while unresolved cases are routed to human verification.

Stage III: Uncertainty-Guided Human Verification. Human reviewers inspect unresolved candidates and a stratified sample covering all transformation operators. They verify reference clarity, intervention plausibility, transition consistency, answer correctness, and option quality. Ambiguous or inconsistent samples are corrected when possible and otherwise excluded.

Table 1: Performance of MLLMs on CST-Bench (%) under the answer-only protocol on the balanced 480-example evaluation split. Best and second-best results are bolded and underlined, respectively.

Model	Transition Families						Aggregate Results	
	Mirror	Swap	Metric	Translation	Depth	Viewpoint	Macro Avg.	Micro Acc.
General-purpose Open-source MLLMs								
Qwen2.5-VL-3B	15.5	15.1	17.8	20.7	11.6	17.0	16.3	17.0
Qwen3-VL-4B	42.5	25.0	35.6	**96.3**	**37.5**	30.0	<u>44.5</u>	**49.4**
Qwen2.5-VL-7B	27.6	19.2	22.3	36.3	12.5	20.6	23.1	25.1
LLaVA-OneVision-7B	26.3	35.0	39.4	<u>90.0</u>	32.5	27.5	41.8	46.6
InternVL3.5-4B	32.5	25.0	26.2	78.3	32.5	30.0	37.4	41.2
Phi-3.5-Vision-4B	33.8	22.5	33.8	85.0	17.5	31.2	37.3	42.9
Gemma-3-4B-it	**52.5**	**42.5**	<u>42.5</u>	60.0	<u>35.0</u>	**40.0**	**45.4**	<u>47.5</u>
Spatial- and Embodied-oriented Open-source MLLMs								
SpatialThinker-3B	27.5	31.2	37.5	70.0	22.5	26.2	35.8	39.8
RoboBrain2.0-7B	<u>50.0</u>	<u>37.5</u>	21.3	33.8	5.0	27.5	29.2	31.6
SpaceR-SFT-7B	41.3	20.0	**51.3**	70.8	32.5	<u>32.5</u>	41.4	44.6

4 Causal Transition On-Policy Self-Distillation

As an application of CST-Bench's dense annotations, CT-OPSD models counterfactual spatial reasoning as causal state-transition prediction. Distinct from standard OPSD, it explicitly trains a teacher to organize source states, intervention effects, invariants, and target states into a structured transition chain. The resulting behavior is distilled through on-policy distribution matching, while the privileged annotations are never provided as student inputs [Agarwal et al., 2024, Zhao et al., 2026].

Causally Structured Teacher. Let $x = (I, q)$ denote the observable input, where I is the image and q contains the question, intervention description, and answer options. CST-Bench additionally provides privileged transition annotations

$$z_c = (s^0, u, e, \mathcal{I}, s^1), \tag{4}$$

where s^0 and s^1 are the source and target states, u is the structured intervention, e records affected relations and expected effects, and $\mathcal{I}$ contains preserved invariants. When applicable, e also includes the corresponding metric change. Together with the standard answer label a^*, these fields form

$$\tau^* : s^0 \xrightarrow{u} (e, \mathcal{I}) \rightarrow s^1 \rightarrow a^*. \tag{5}$$

The teacher π_ϕ learns to generate this chain from x:

$$\mathcal{L}_{\text{teacher}} = -\sum_{t=1}^{|\tau^*|} \log \pi_\phi(\tau_t^* \mid x, \tau_{<t}^*). \tag{6}$$

The teacher is frozen after this stage.

Student Rollout and Causal Bridge. The student π_θ is initialized from the same base model and uses only x as input. It generates an on-policy trajectory

$$\hat{\tau} \sim \pi_\theta(\cdot \mid x) \tag{7}$$

[Ross, Gordon, and Bagnell, 2011]. At each step, the frozen teacher and student evaluate the same student-generated prefix:

$$p_T^{(t)} = \pi_\phi(\cdot \mid x, \hat{\tau}_{<t}), \qquad p_S^{(t)} = \pi_\theta(\cdot \mid x, \hat{\tau}_{<t}). \tag{8}$$

Thus, the causal bridge transfers the teacher's structured transition behavior on states encountered by the student, without exposing z_c or τ^* to the student.

Training Objective. We align the two distributions along student rollouts using forward KL divergence:

$$\mathcal{L}_{\text{OPSD}} = \mathbb{E}_{\hat{\tau} \sim \pi_\theta(\cdot|x)} \left[\frac{1}{L} \sum_{t=1}^{L} D_{\text{KL}} \left(p_T^{(t)} \parallel p_S^{(t)} \right) \right]. \tag{9}$$

Table 2: Post-training results on CST-Bench and standard spatial benchmarks (%). All variants use Qwen2.5-VL-3B as the base model. Best and second-best results are bolded and underlined, respectively.

Model	CST-Bench	3DSRBench	CV-Bench			$\text{BLINK}_{\text{val}}$		
	Micro Acc.	Acc.	2D	3D	Avg.	Spatial Relation	Relative Depth	Avg.
Post-training Strategies with Qwen2.5-VL-3B								
Qwen2.5-VL-3B	17.0	46.0	64.0	66.0	65.0	<u>83.0</u>	64.0	73.5
Qwen2.5-VL-3B + SFT	50.3	38.0	63.0	68.0	65.5	70.0	67.0	68.5
Qwen2.5-VL-3B + GRPO	46.7	**52.9**	71.0	76.3	73.7	81.8	66.9	74.4
Qwen2.5-VL-3B + OPSD	<u>53.3</u>	<u>51.8</u>	<u>73.8</u>	<u>78.3</u>	<u>76.1</u>	82.5	<u>67.2</u>	<u>74.9</u>
Qwen2.5-VL-3B + CT-OPSD	**55.2**	49.5	**75.1**	**78.6**	**76.9**	**84.0**	**69.2**	**76.6**

The complete objective additionally includes standard answer supervision:

$$\mathcal{L}_{\text{CT-OPSD}} = \mathcal{L}_{\text{OPSD}} + \lambda_{\text{ans}} \left[-\log \pi_\theta(a^* \mid x) \right]. \tag{10}$$

At inference, the teacher and all privileged transition annotations are discarded, and the student predicts from $x = (I, q)$ alone.

5 Experiments

5.1 Implementation Details

We evaluate ten open-source multimodal large language models (MLLMs): Qwen2.5-VL-3B and Qwen2.5-VL-7B, Qwen3-VL-4B [Bai et al., 2025], LLaVA-OneVision-7B [Li et al., 2024], InternVL3.5-4B [Wang et al., 2025], Phi-3.5-Vision-4B [Abdin et al., 2024], Gemma-3-4B-it [Team et al., 2025b], RoboBrain2.0-7B [Team et al., 2025a], SpatialThinker-3B [Batra et al., 2026], and SpaceR-SFT-7B [Ouyang et al., 2025]. All foundation models are evaluated without fine-tuning on CST-Bench.

We use the balanced 480-example evaluation split, which contains 40 examples for each of the twelve fine-grained operators. These operators are grouped into six transition families: Global Mirror Transformation, Object Position Swap, Metric Distance Transition, Translation-Based Crossing, Depth Transition, and Viewpoint Reversal. All retained examples keep the referenced objects visible before and after the intervention; distance-related operators test continuous metric changes or qualitative near–far threshold crossings rather than visibility changes.

At inference time, every model receives only the image, the counterfactual question, and the answer options. Models do not access scene states, intervention effects, invariant relations, or other privileged causal annotations. We use a common answer-only prompt, greedy decoding, and a maximum of 64 newly generated tokens. Each response is parsed as a multiple-choice answer.

5.2 Baselines and Evaluation Metrics

Using Qwen2.5-VL-3B as the common base, we compare the unadapted model, answer-supervised SFT, GRPO [Shao et al., 2024], OPSD [Zhao et al., 2026], and CT-OPSD. All variants use the same observable input (I, q); privileged annotations are available only when training the CT-OPSD teacher.

We report multiple-choice accuracy. Foundation models are evaluated by family accuracy, Macro Avg. over the six families, and Micro Acc. over all 480 examples. Post-trained models are additionally evaluated on 3DSRBench [Ma et al., 2025], CV-Bench [Tong et al., 2024], and $\text{BLINK}_{\text{val}}$ [Fu et al., 2024].

For component ablation, CST W-Avg. denotes the predefined weighted aggregate of causal-transition diagnostics, distinct from the family-level Macro Avg. Its definition and weights are provided in the appendix.

5.3 Main Results

Table 1 compares ten open-source MLLMs on the balanced CST-Bench evaluation split. Overall performance remains limited: no foundation model is consistently strong across all six transition families. Gemma-3-4B-it achieves the best Macro Avg. of 45.4, indicating the most balanced family-level performance, whereas Qwen3-VL-4B obtains the highest Micro Acc. of 49.4.

Table 3: Component ablation of CT-OPSD. CST W-Avg. is the weighted diagnostic aggregate. Best and second-best results are bolded and underlined, respectively.

Variant	CST W-Avg.	Ans. Acc.	Src.	Eff.	Inv.	Cons.
OPSD	24.0	<u>53.3</u>	34.1	0.7	0.3	22.5
CT w/o Aux.	<u>24.7</u>	52.7	40.1	0.9	0.4	30.0
CT w/o Bridge	24.5	52.1	40.0	1.0	<u>0.7</u>	<u>30.6</u>
CT w/o Eff.	24.0	50.4	40.0	1.2	**1.0**	28.5
CT w/o Inv.	24.6	51.7	<u>41.2</u>	<u>1.3</u>	0.5	**31.2**
CT w/o Metric	24.1	50.4	**41.3**	1.2	0.5	28.5
CT-OPSD	**25.9**	**55.2**	**41.3**	**1.5**	0.6	30.4

The difference between these rankings is expected rather than contradictory. Macro Avg. weights each transition family equally, whereas Micro Acc. weights families according to their number of examples. In particular, Translation-Based Crossing contains three operators and thus contributes more to Micro Acc. than the single-operator Depth Transition. We therefore report both metrics to distinguish balanced transition reasoning from aggregate accuracy.

5.4 Family-wise Analysis

The leading model varies substantially by transition family. Gemma-3-4B-it performs best on Global Mirror Transformation (52.5), Object Position Swap (42.5), and Viewpoint Reversal (40.0). SpaceR-SFT-7B achieves the best result on Metric Distance Transition (51.3), while Qwen3-VL-4B leads Translation-Based Crossing (96.3) and Depth Transition (37.5). This distribution indicates that competence in one type of spatial update does not reliably transfer to other transformations.

Depth Transition and Viewpoint Reversal are particularly challenging: their best scores are only 37.5 and 40.0, respectively. Qwen3-VL-4B illustrates the resulting imbalance, achieving 96.3 on translation but only 25.0 on swapping and 30.0 on viewpoint reversal. Thus, strong performance on image-plane translation alone is insufficient for reliable counterfactual spatial reasoning across relational, metric, depth, and reference-frame changes.

5.5 Post-training Comparison

Table 3 reports source-state understanding, effect reasoning, invariance identification, and transition consistency, abbreviated as Src., Eff., Inv., and Cons. Full CT-OPSD achieves the best CST W-Avg. (25.9) and answer accuracy (55.2), while all component removals reduce both metrics. Although individual diagnostic scores vary across variants, the full method provides the strongest overall balance.

Table 2 compares post-training strategies using Qwen2.5-VL-3B. CT-OPSD achieves the best CST-Bench Micro Acc. (55.2), outperforming standard OPSD by 1.9 points, and obtains the best CV-Bench and BLINK averages. Its 3DSRBench score remains below GRPO and OPSD but exceeds the unadapted base model, indicating improved counterfactual reasoning without uniformly maximizing every external benchmark.

6 Analysis

The fundamental distinction of CST-Bench lies not merely in task difficulty, but in the object of inference. Static benchmarks ask a model to recognize a spatial state already present in an image, whereas CST-Bench requires it to intervene on that state and predict a valid successor under conservation constraints. It therefore tests whether the model's spatial representation is coherent and manipulable, rather than a collection of independently recognized relation labels. The uneven family-level performance is consistent with current MLLMs lacking such a unified transition model.

Dense transition annotations are important because counterfactual correctness has two inseparable sides: applying the intended change and avoiding unintended changes. A model may fail by ignoring the intervention, but it may also fail by updating too much. Final-answer supervision alone does not distinguish these errors. CT-OPSD is effective because its supervision reflects this structure: the teacher learns selective state updating, and on-policy alignment transfers that behavior at states encountered by the student, without providing privileged evidence at inference.

The ablations reinforce this interpretation. Improvements on an isolated diagnostic dimension do not necessarily improve overall performance because effects and invariants must be satisfied jointly. The key connection is therefore structural: CST-Bench defines valid spatial transitions, its annotations expose their internal constraints, and CT-OPSD converts those constraints into a learnable training signal.

7 Conclusion

We introduced CST-Bench, a counterfactual spatial reasoning benchmark with 7,082 examples across six families and twelve operators. It evaluates intervention-induced relational, metric, depth, and viewpoint changes with dense causal transition annotations. We also proposed CT-OPSD, which distills a privileged causal teacher into an observation-only student. Experiments reveal uneven counterfactual capabilities among existing MLLMs, while CT-OPSD improves CST-Bench accuracy and retains strong performance on standard spatial benchmarks.

References

Abdin, M.; Aneja, J.; Awadalla, H.; et al. 2024. Phi-3 Technical Report: A Highly Capable Language Model Locally on Your Phone. arXiv:2404.14219.

Agarwal, R.; Vieillard, N.; Zhou, Y.; Stanczyk, P.; Ramos, S.; Geist, M.; and Bachem, O. 2024. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. arXiv:2306.13649.

Bai, S.; Chen, K.; Liu, X.; Wang, J.; Ge, W.; Song, S.; Dang, K.; Wang, P.; Wang, S.; Tang, J.; Zhong, H.; Zhu, Y.; Yang, M.; Li, Z.; Wan, J.; Wang, P.; Ding, W.; Fu, Z.; Xu, Y.; Ye, J.; Zhang, X.; Xie, T.; Cheng, Z.; Zhang, H.; Yang, Z.; Xu, H.; and Lin, J. 2025. Qwen2.5-VL Technical Report. arXiv:2502.13923.

Batra, H.; Tu, H.; Chen, H.; Lin, Y.; Xie, C.; and Clark, R. 2026. SpatialThinker: Reinforcing Scene Graph-Grounded Spatial Reasoning via Dense Rewards. arXiv:2511.07403.

Bengio, S.; Vinyals, O.; Jaitly, N.; and Shazeer, N. 2015. Scheduled sampling for sequence prediction with recurrent neural networks. *Advances in neural information processing systems*, 28.

Cai, Z.; Wang, Y.; Sun, Q.; et al. 2025. Holistic Evaluation of Multimodal LLMs on Spatial Intelligence. arXiv:2508.13142.

Chen, B.; Xu, Z.; Kirmani, S.; Ichter, B.; Driess, D.; Florence, P.; Sadigh, D.; Guibas, L.; and Xia, F. 2024. SpatialVLM: Endowing Vision-Language Models with Spatial Reasoning Capabilities. arXiv:2401.12168.

Chen, L.; Yan, X.; Xiao, J.; Zhang, H.; Pu, S.; and Zhuang, Y. 2020. Counterfactual Samples Synthesizing for Robust Visual Question Answering. arXiv:2003.06576.

Cheng, A.-C.; Yin, H.; Fu, Y.; Guo, Q.; Yang, R.; Kautz, J.; Wang, X.; and Liu, S. 2024. SpatialRGPT: Grounded Spatial Reasoning in Vision Language Models. arXiv:2406.01584.

Cohn, A. G.; and Renz, J. 2008. Qualitative spatial representation and reasoning. *Foundations of artificial intelligence*, 3: 551–596.

Fu, X.; Hu, Y.; Li, B.; Feng, Y.; Wang, H.; Lin, X.; Roth, D.; Smith, N. A.; Ma, W.-C.; and Krishna, R. 2024. BLINK: Multimodal Large Language Models Can See but Not Perceive. *ArXiv*, abs/2404.12390.

Gao, Y.; Li, H.; Liu, Y.; et al. 2026. Holi-Spatial: Evolving Video Streams into Holistic 3D Spatial Intelligence. arXiv:2603.07660.

Guo, Z.; Yang, Z.; Li, Y.; Zhang, X.; Gao, W.; Wang, J.; Li, C.; Liu, X.-Y.; and Jian, P. 2026. Can LLMs See Without Pixels? Benchmarking Spatial Intelligence from Textual Descriptions. *ArXiv*, abs/2601.03590.

Hinton, G.; Vinyals, O.; and Dean, J. 2015. Distilling the Knowledge in a Neural Network. arXiv:1503.02531.

Kim, Y.; and Rush, A. M. 2016. Sequence-Level Knowledge Distillation. arXiv:1606.07947.

Klatzky, R. L. 1998. Allocentric and egocentric spatial representations: Definitions, distinctions, and interconnections. In *Spatial cognition: An interdisciplinary approach to representing and processing spatial knowledge*, 1–17. Springer.

Krishna, R.; Zhu, Y.; Groth, O.; et al. 2016. Visual Genome: Connecting Language and Vision Using Crowdsourced Dense Image Annotations. arXiv:1602.07332.

Li, B.; Zhang, Y.; Guo, D.; et al. 2024. LLaVA-OneVision: Easy Visual Task Transfer. arXiv:2408.03326.

Li, D.; Li, H.; Wang, Z.; et al. 2025. ViewSpatial-Bench: Evaluating Multi-perspective Spatial Localization in Vision-Language Models. arXiv:2505.21500.

Lopez-Paz, D.; Bottou, L.; Schölkopf, B.; and Vapnik, V. 2016. Unifying distillation and privileged information. arXiv:1511.03643.

Ma, W.; Chen, H.; Zhang, G.; Chou, Y.-C.; Chen, J.; de Melo, C. M.; and Yuille, A. 2025. 3DSRBench: A Comprehensive 3D Spatial Reasoning Benchmark. arXiv:2412.07825.

Ma, X.; Yong, S.; Zheng, Z.; Li, Q.; Liang, Y.; Zhu, S.-C.; and Huang, S. 2023. SQA3D: Situated Question Answering in 3D Scenes. arXiv:2210.07474.

Niu, Y.; Tang, K.; Zhang, H.; Lu, Z.; Hua, X.-S.; and Wen, J.-R. 2021. Counterfactual VQA: A Cause-Effect Look at Language Bias. arXiv:2006.04315.

Ouyang, K.; Liu, Y.; Wu, H.; Liu, Y.; Zhou, H.; Zhou, J.; Meng, F.; and Sun, X. 2025. SpaceR: Reinforcing MLLMs in Video Spatial Reasoning. arXiv:2504.01805.

Pearl, J. 2009. *Causality*. Cambridge university press.

Peters, J.; Janzing, D.; and Schölkopf, B. 2017. Elements of Causal Inference: Foundations and Learning Algorithms.

Ross, S.; Gordon, G. J.; and Bagnell, J. A. 2011. A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning. arXiv:1011.0686.

Shao, Z.; Wang, P.; Zhu, Q.; Xu, R.; Song, J.; Bi, X.; Zhang, H.; Zhang, M.; Li, Y. K.; Wu, Y.; and Guo, D. 2024. DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models. arXiv:2402.03300.

Team, B. R.; Cao, M.; Tan, H.; et al. 2025a. RoboBrain 2.0 Technical Report. arXiv:2507.02029.

Team, G.; Kamath, A.; Ferret, J.; et al. 2025b. Gemma 3 Technical Report. arXiv:2503.19786.

Tong, S.; Brown, E.; Wu, P.; et al. 2024. Cambrian-1: A Fully Open, Vision-Centric Exploration of Multimodal LLMs. *ArXiv*, abs/2406.16860.

Urp'i, N. A.; Bagatella, M.; Vlastelica, M.; and Martius, G. 2024. Causal Action Influence Aware Counterfactual Data Augmentation. *ArXiv*, abs/2405.18917.

Vapnik, V.; and Vashist, A. 2009. A new learning paradigm: Learning using privileged information. *Neural networks*, 22(5-6): 544–557.

Wang, P.; Li, L.; Chen, L.; Cai, Z.; Zhu, D.; Lin, B.; Cao, Y.; Liu, Q.; Liu, T.; and Sui, Z. 2023a. Large Language Models are not Fair Evaluators. arXiv:2305.17926.

Wang, Q.; Yin, B.; Zhang, P.; et al. 2026. MindCube: Spatial Mental Modeling from Limited Views. arXiv:2506.21458.

Wang, W.; Gao, Z.; Gu, L.; et al. 2025. InternVL3.5: Advancing Open-Source Multimodal Models in Versatility, Reasoning, and Efficiency. arXiv:2508.18265.

Wang, X.; Wei, J.; Schuurmans, D.; Le, Q.; Chi, E.; Narang, S.; Chowdhery, A.; and Zhou, D. 2023b. Self-Consistency Improves Chain of Thought Reasoning in Language Models. arXiv:2203.11171.

Wu, H.; Huang, X.; Chen, Y.; Zhang, Y.; Wang, Y.; and Xie, W. 2026. SpatialScore: Towards Comprehensive Evaluation for Spatial Intelligence. arXiv:2505.17012.

Xu, D.; Zhu, Y.; Choy, C. B.; and Fei-Fei, L. 2017. Scene Graph Generation by Iterative Message Passing. arXiv:1701.02426.

Yang, J.; Yang, S.; Gupta, A. W.; Han, R.; Fei-Fei, L.; and Xie, S. 2025. Thinking in Space: How Multimodal Large Language Models See, Remember, and Recall Spaces. arXiv:2412.14171.

Yang, S.; Xu, R.; Xie, Y.; et al. 2026. MMSI-Bench: A Benchmark for Multi-Image Spatial Intelligence. arXiv:2505.23764.

Yi, K.; Gan, C.; Li, Y.; Kohli, P.; Wu, J.; Torralba, A.; and Tenenbaum, J. B. 2020. CLEVRER: CoLlision Events for Video REpresentation and Reasoning. arXiv:1910.01442.

Yu, L.; Tang, X.; Liu, M.; Li, X.; Bai, X.; Zhou, S.; Lin, Q.; Xuan, W.; and Yokoya, N. 2026a. MindEdit-Bench: Benchmarking Object-Level Counterfactual Spatial Reasoning in VLMs from In-the-Wild Photos. arXiv:2607.00491.

Yu, X.; Li, G.; Si, Q.; et al. 2026b. DOPD: Dual On-policy Distillation. arXiv:2606.30626.

Yuan, Q.; Lou, J.; Yu, X.; Lin, H.; Sun, L.; Han, X.; and Lu, Y. 2026. Vision-OPD: Learning to See Fine Details for Multimodal LLMs via On-Policy Self-Distillation. arXiv:2605.18740.

Zhan, W.; Zhou, Z.; Zheng, Z.; Gao, C.; Cui, J.; Li, Y.; Chen, X.; and Zhang, X.-P. 2025. Open3DVQA: A Benchmark for Comprehensive Spatial Reasoning with Multimodal Large Language Model in Open Space. *ArXiv*, abs/2503.11094.

Zhang, L.; Zhai, X.; Zhao, Z.; Zong, Y.; Wen, X.; and Zhao, B. 2024. What If the TV Was Off? Examining Counterfactual Reasoning Abilities of Multi-modal Language Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 21853–21862.

Zhao, S.; Xie, Z.; Liu, M.; Huang, J.; Pang, G.; Chen, F.; and Grover, A. 2026. Self-Distilled Reasoner: On-Policy Self-Distillation for Large Language Models. arXiv:2601.18734.

Zheng, L.; Chiang, W.-L.; Sheng, Y.; Zhuang, S.; Wu, Z.; Zhuang, Y.; Lin, Z.; Li, Z.; Li, D.; Xing, E. P.; Zhang, H.; Gonzalez, J. E.; and Stoica, I. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. arXiv:2306.05685.

(a) Global mirror

(b) Object swap

(c) Metric distance

(d) Translation crossing

(e) Depth transition

(f) Viewpoint reversal

Figure 5: Representative source images from the six CST-Bench intervention families. Each example is paired with a single intervention-conditioned query and a symbolic target transition.

Technical Supplement

This technical supplement provides additional details for CST-Bench and CT-OPSD. It follows the task definition, construction procedure, and observation-only inference setting described in the main paper. The main paper is self-contained; the material below specifies the benchmark records, operator semantics, validation protocol, and experimental diagnostics.

A Representative CST-Bench Examples

Figure 5 presents one source image from each intervention family. Each sample contains one observed image, one intervention-conditioned multiple-choice question, one target answer, and one causal-transition record. The post-intervention image is not rendered. Instead, the target state is defined symbolically from the observed scene and the hypothetical intervention.

Global Mirror Transformation. The query asks how an object is related to a reference after a horizontal mirror transformation. A source relation of *left of* becomes *right of*; object identity, object count, and vertical relations are preserved.

Object Position Swap. Two objects exchange positions, while the query may concern one exchanged object and a third, unchanged reference object. The example therefore tests whether the direct edit is propagated to the relevant third-object relation.

Metric Distance Transition. The intervention changes normalized distance while holding direction fixed. The qualitative relation may remain unchanged or cross the *near/ far* boundary, separating metric sensitivity from categorical updates.

Translation-Based Crossing. An object is translated across a reference horizontally or vertically. The affected relation changes, while the orthogonal relation and object properties remain invariant.

Depth Transition. An object moves across a reference along the depth axis. The *behind/ in front of* relation changes while image-plane left–right and above–below relations are preserved.

Viewpoint Reversal. The observer changes viewpoint while the physical scene remains fixed. The task recomputes observer-centered relations without treating the viewpoint change as an edit to the allocentric scene layout.

Table 4: Full-dataset distribution by intervention family.

Family	Samples	Share
Global Mirror Transformation	1,600	22.6%
Object Position Swap	1,600	22.6%
Metric Distance Transition	1,600	22.6%
Translation-Based Crossing	1,300	18.4%
Depth Transition	300	4.2%
Viewpoint Reversal	682	9.6%
Total	7,082	100.0%

Table 5: Sample counts by transition operator. The official evaluation split contains 40 examples per operator.

Transition operator	Samples	Transition operator	Samples
Pairwise Left–Right Flip	1,300	Lateral Relation Crossing	386
Multi-object Horizontal Order Flip	300	Above–Below Crossing	614
Pairwise Relation Swap	300	Vertical Near–Far Crossing	300
Third-object Relation Update	1,300	Depth Order Change	300
Metric Distance Change (Relation Preserved)	1,300	Front-to-Back Viewpoint Shift	382
Distance-to-Reference Near–Far Crossing	300	Side-to-Side Viewpoint Shift	300
		Total	**7,082**

B CST-Bench Specification and Construction

B.1 Data Source and Evaluation Split

CST-Bench curates scenes from the VG150 split of Visual Genome with valid object boxes, sufficient distinct objects, and geometrically unambiguous relations. The source annotations supply object categories, bounding boxes, and scene relations used to identify valid spatial configurations. Pairwise operators require two usable objects; Third-object Relation Update and multi-object order operators require three distinct object categories.

The benchmark contains 7,082 image-grounded samples across six intervention families and twelve transition operators. The balanced evaluation split contains 480 examples, with 40 examples for each operator. Samples are partitioned by source image so that source images in the evaluation split are excluded from post-training data. This grouped split prevents a model from observing another counterfactual variant of an evaluation image during training.

B.2 Spatial Relation Definitions

For a box $b_i = [x_i^1, y_i^1, x_i^2, y_i^2]$ in an image of width W and height H, the normalized center is

$$(\tilde{x}_i, \tilde{y}_i) = \left(\frac{x_i^1 + x_i^2}{2W}, \frac{y_i^1 + y_i^2}{2H} \right). \tag{11}$$

For an ordered pair (i, j), *left of*/*right of* are determined by the sign of $\tilde{x}_i - \tilde{x}_j$, and *above*/*below* are determined by the sign of $\tilde{y}_i - \tilde{y}_j$. Candidates with insufficient center separation along the queried axis are rejected.

Normalized center distance is

$$d(i, j) = \frac{\|c_i - c_j\|_2}{\sqrt{W^2 + H^2}}. \tag{12}$$

Near–far operators retain only pairs outside an ambiguity band around the qualitative boundary. A relation-preserved sample changes $d(i, j)$ without changing its qualitative label, whereas a threshold-crossing sample changes both the metric distance and the qualitative relation.

All boxes must have positive extent and remain within the image boundary before and after the symbolic edit. Thus, every object needed to answer a retained question remains visible. Front–behind labels are taken from available scene annotations or from the operator-defined symbolic update; they are not inferred from two-dimensional box centers.

Table 6: Complete CST-Bench operator specification.

Family	Operator	Precondition	Intervention	Affected / expected	Principal invariants
Global mirror	Pairwise Left–Right Flip	Clear horizontal pair	Mirror image coordinates	left ↔ right	identity, count, vertical relation
Global mirror	Multi-object Horizontal Order Flip	Three ordered, distinct objects	Mirror all selected objects	reverse all horizontal pair orders	identity, count, vertical relations
Object swap	Pairwise Relation Swap	Two distinct objects	Exchange their positions	queried-pair relation reverses	identity, count, attributes
Object swap	Third-object Relation Update	Swapped pair plus unchanged reference	Exchange two positions	update relation to third object	reference object, identity, count
Metric distance	Metric Distance Change (Relation Preserved)	Clear near/far pair	Change subject–reference distance	metric distance changes; label preserved	direction, identity, image-plane order
Metric distance	Distance-to-Reference Near–Far Crossing	Clear pair with valid offset	Change distance without crossing reference object	near ↔ far	side of reference, identity, orthogonal relation
Translation	Lateral Relation Crossing	Clear horizontal pair	Translate subject across reference	left ↔ right	identity, count, vertical relation
Translation	Above–Below Crossing	Clear vertical pair	Translate subject across reference	above ↔ below	identity, count, horizontal relation
Translation	Vertical Near–Far Crossing	Clear pair with valid vertical offset	Translate vertically without crossing reference	near ↔ far	identity, count, horizontal relation
Depth	Depth Order Change	Source front/behind label	Move subject across reference along depth	front ↔ behind	identity, count, left/right, above/below
Viewpoint	Front-to-Back Viewpoint Shift	Observer-centered depth relation	Reverse viewing direction	front ↔ behind	physical scene and allocentric layout
Viewpoint	Side-to-Side Viewpoint Shift	Observer-centered lateral relation	View scene from opposite side	left ↔ right	physical scene and allocentric layout

B.3 Complete Operator Specification

Table 6 defines the twelve operators. "Affected" lists the relation updated by the symbolic editor; "invariants" lists the properties that must remain unchanged for a valid transition.

B.4 Constraint-Aware Candidate Generation

The principal construction difficulty is that the post-intervention state is not directly visible in the source image. CST-Bench addresses this issue with a deterministic, operator-aware editor and symbolic solver. Starting from a curated source state, the editor applies an intervention and the solver jointly derives the source state, target state, affected relations, expected effects, invariants, applicable metric changes, and final answer. Generating these fields from one edited state prevents contradictions between a transition and its dense annotations.

Candidates are discarded if they violate an operator-specific constraint, contain invalid boxes or ambiguous references, lack sufficient relation separation, fail to realize the intended effect, or move a referenced object outside the image. Remaining candidates are scored by

$$s(c) = s_{\text{ref}} + s_{\text{rel}} + s_{\text{edit}} + s_{\text{vis}} + s_{\text{nat}}, \tag{13}$$

where the five terms respectively measure reference clarity, relation clarity, edit effectiveness, post-edit visibility, and transformation naturalness. Candidates with $s(c) < 8$ are removed. The remaining candidates are ranked by

$$r(c) = \frac{1}{2}\hat{s}_{\text{op}}(c) + \frac{1}{2}\hat{s}_{\text{global}}(c), \tag{14}$$

where the two normalized scores are computed within an operator and globally, respectively. At most one highest-ranked transformation is retained for each source image.

B.5 Answer-Blind Consensus Validation and Human Verification

The symbolic label can be internally consistent while the image, reference, or question remains ambiguous. CST-Bench therefore uses answer-blind consensus validation before human verification. For each candidate, a VLM independently solves the question $K = 3$ times without seeing the symbolic answer. It then compares its independent prediction with the symbolic label and scores five criteria from 0 to 2: scene–image consistency, reference clarity, transformation fidelity, answer correctness, and option quality.

A candidate is retained only when all three reviews recover the symbolic answer, report no ambiguity, and assign the maximum score to every criterion:

$$\text{Keep}(c) = \bigwedge_{k=1}^{K} \left[\hat{a}_k = a^* \wedge b_k = 0 \wedge \min_j q_{k,j} = 2 \right], \tag{15}$$

where $\hat{a}_k$ is the independent prediction, a^* is the symbolic answer, b_k is the ambiguity indicator, and $q_{k,j}$ is the score of criterion j. Candidates with repeated hard failures are excluded. Candidates without unanimous agreement are routed to human verification rather than retained by a majority vote.

Human verification follows a two-part protocol. First, two reviewers inspect each unresolved candidate independently. They assess reference clarity, intervention plausibility, source-to-target transition consistency, target answer correctness, and distractor quality. A reviewer marks a sample as *retain*, *correct*, or *exclude*, with a brief rationale for every correction or exclusion. Second, a stratified audit samples retained examples from every operator to check that automated filtering does not conceal operator-specific errors. When the two reviewers disagree, a third reviewer adjudicates using the source image, the symbolic transition record, and the operator definition in Table 6. Only adjudicated retain or corrected decisions enter the final benchmark; ambiguous or inconsistent samples are excluded.

B.6 Released Data Fields

Table 7 lists the principal fields in each released sample. Structured fields can be distributed as JSON or as serialized values in a CSV representation.

Table 7: Principal CST-Bench fields.

Field	Meaning
`sample_uid`	Stable sample identifier.
`image_id`	Source-image identifier used for grouped splitting.
`operator`	Intervention family and fine-grained transition operator.
`question`	Intervention-conditioned natural-language query.
`options`, `answer`	Answer choices and target choice.
`source_state`	Relevant source objects and spatial relations.
`target_state`	Symbolically updated post-intervention state.
`intervention`	Operator and its object-specific parameters.
`affected_relations`	Relations changed by the intervention.
`expected_effects`	Required pre-to-post relational or metric changes.
`invariants`	Properties that must remain unchanged.
`metric_change`	Metric-distance information, when applicable.

C CT-OPSD Implementation Details

C.1 Teacher Supervision and Student Inputs

Let $x = (I, q)$ denote the observable image-question input. For training, each CST-Bench example additionally provides the causal transition annotations

$$z_c = (s^0, u, e, \mathcal{I}, s^1), \tag{16}$$

where s^0 and s^1 are the source and target states, u is the intervention, e records affected relations and expected effects, and $\mathcal{I}$ is the set of preserved invariants. Metric information is included in e when applicable. Together with the answer a^*, these labels define the reference transition

$$\tau^* : s^0 \xrightarrow{u} (e, \mathcal{I}) \to s^1 \to a^*. \tag{17}$$

The causal teacher is supervised to generate this structured transition from x. Its teacher-training loss is

$$\mathcal{L}_{\text{teacher}} = -\sum_{t=1}^{|\tau^*|} \log \pi_\phi(\tau_t^* \mid x, \tau_{<t}^*). \tag{18}$$

The privileged labels specify the teacher target; they are not supplied as inputs to the student. The teacher is frozen before student training.

The student is initialized from the same base model and receives only $x = (I, q)$ throughout rollout training and inference. It generates an on-policy trajectory $\hat{\tau} \sim \pi_\theta(\cdot \mid x)$. At step t, the frozen teacher and student evaluate the same student-generated prefix:

$$p_T^{(t)} = \pi_\phi(\cdot \mid x, \hat{\tau}_{<t}), \qquad p_S^{(t)} = \pi_\theta(\cdot \mid x, \hat{\tau}_{<t}). \tag{19}$$

The resulting causal bridge transfers teacher behavior at states visited by the student rather than only along reference trajectories.

Table 8: CT-OPSD components and their roles.

Component	Role
Observable input x	Image, question, intervention description, and answer options.
Causal teacher	Predicts source state, intervention consequences, invariants, target state, and answer.
Student rollout	Samples a trajectory from the observation-only student policy.
On-policy alignment	Matches the teacher and student distributions on the shared student prefix.
Answer supervision	Anchors the final response to the gold answer a^*.
Inference student	Uses only the image and question after teacher removal.

C.2 Training Objective and Inference

CT-OPSD minimizes forward KL divergence along student rollouts together with standard answer supervision:

$$\mathcal{L}_{\text{OPSD}} = \mathbb{E}_{\hat{\tau}\sim\pi_\theta(\cdot|x)}\left[\frac{1}{L}\sum_{t=1}^{L} D_{\text{KL}}\left(p_T^{(t)} \parallel p_S^{(t)}\right)\right], \tag{20}$$

$$\mathcal{L}_{\text{CT-OPSD}} = \mathcal{L}_{\text{OPSD}} + \lambda_{\text{ans}}\left[-\log \pi_\theta(a^* \mid x)\right]. \tag{21}$$

At inference, the teacher and all privileged causal-transition annotations are discarded. The retained student predicts from the image and question alone.

C.3 Prompt and Output Format

The student uses the same answer-only input as the evaluated foundation models: the image, intervention-conditioned question, and answer options. It is not given source-state labels, affected relations, expected effects, invariants, target states, or metric annotations. The teacher is trained to emit a compact causal transition chain followed by the answer; the student is distilled on its own generated continuations and is evaluated by its final multiple-choice answer.

C.4 Ablation Definitions

The component ablations in the main paper remove one source of causal supervision while keeping the base model, observable input, evaluation split, and answer format fixed. **CT w/o Aux.** removes auxiliary structured supervision beyond the final answer. **CT w/o Bridge** removes the causal-transition bridge between teacher and student rollouts. **CT w/o Eff.**, **CT w/o Inv.**, and **CT w/o Metric** remove the corresponding effect, invariant, and metric information from the teacher's causal-transition supervision. Standard **OPSD** does not use CST-Bench's structured causal-transition annotations. All variants are evaluated observation-only, using the image and question without privileged annotations.

D Evaluation Protocol and Additional Results

D.1 Answer-Only Evaluation

At inference, every foundation model and post-trained student receives the source image, counterfactual question, and answer options. Models do not access source states, intervention effects, invariant relations, target states, or other privileged annotations. Decoding uses a common answer-only prompt and the response is parsed as a multiple-choice answer; invalid or unmatched responses receive zero credit.

Let $a_i \in \{0, 1\}$ indicate whether prediction i is correct. For family f with index set D_f,

$$\text{Acc}_f = \frac{1}{|D_f|}\sum_{i\in D_f} a_i, \qquad \text{Macro} = \frac{1}{6}\sum_{f=1}^{6}\text{Acc}_f, \tag{22}$$

and

$$\text{Micro} = \frac{1}{480}\sum_{i=1}^{480} a_i. \tag{23}$$

Macro Avg. weights each family equally, whereas Micro Acc. weights every example in the balanced 480-example evaluation split equally.

D.2 Causal-Transition Diagnostics

The diagnostic evaluation measures source-state understanding (Src.), effect reasoning (Eff.), invariance identification (Inv.), and transition consistency (Cons.). Source-state understanding assesses whether the model recovers the relevant pre-intervention relation. Effect reasoning assesses whether it identifies the intervention-induced change. Invariance identification checks whether unaffected properties are preserved, and transition consistency checks whether the predicted relation update agrees with the final answer.

For ablations, CST W-Avg. aggregates answer accuracy and the structured diagnostics according to the predefined benchmark weights:

$$\text{CST W-Avg.} = \frac{\sum_m w_m S_m}{\sum_m w_m}, \tag{24}$$

where S_m is the score of diagnostic dimension m and w_m is its fixed weight. This aggregate is distinct from family-level Macro Avg. and is used to compare whether a method jointly captures source states, effects, invariants, and consistent target transitions.

D.3 External Spatial Benchmarks

Post-trained variants are also evaluated on 3DSRBench, CV-Bench, and the spatial-relation and relative-depth subsets of $\text{BLINK}_{\text{val}}$. All post-training methods use the same base model and the same evaluation protocol within each external benchmark. CST-Bench metrics are not combined with these external benchmark scores; the reported external columns retain the standard task-level accuracy or average for each benchmark.

E Limitations, Ethics, and Release Scope

Geometric approximation. Bounding boxes support controlled image-plane relations and normalized distance changes, but they are not metric three-dimensional reconstructions. Apparent distance depends on projection, object scale, and depth. Depth transitions therefore rely on source annotations and symbolic state updates rather than two-dimensional box geometry alone.

Symbolic interventions. The observed image is real, but the post-intervention state is symbolic rather than a photorealistically rendered edit. CST-Bench tests state updating under a hypothetical intervention; it does not evaluate image-edit realism, physical simulation, collision, deformation, or long-horizon dynamics.

Validation limits. Answer-blind consensus validation and human verification reduce ambiguity but cannot remove all systematic errors. VLM reviewers can share model-specific biases, and human judgments can remain uncertain for borderline spatial cases. The released validation records should therefore be retained and extended when the benchmark is revised.

Data and task coverage. The benchmark is English-language, multiple-choice, and limited to twelve operators. Its visual, object, language, geographic, cultural, and demographic coverage inherits that of its source data. Performance may not transfer to open-ended responses, videos, robotics, safety-critical systems, or physical settings with richer dynamics.

Responsible use. The source images may depict people and everyday environments. CST-Bench does not add identity labels or infer sensitive attributes. A high benchmark score does not establish that a model is safe for autonomous navigation or robotic control, where an incorrect spatial update can cause unsafe actions.