arXiv:2611.12324v1 [cs.LG] 9 Nov 2026

How Many Exposures Does a World Model Need? Learning Unseen Game Mechanics at Test Time

Ku Onoda[1,2] Haruto Kawano[1,2] Shota Takashiro[1,2] Kohei Hayashi[1,2]
[1]Third Intelligence, Inc. [2]The University of Tokyo

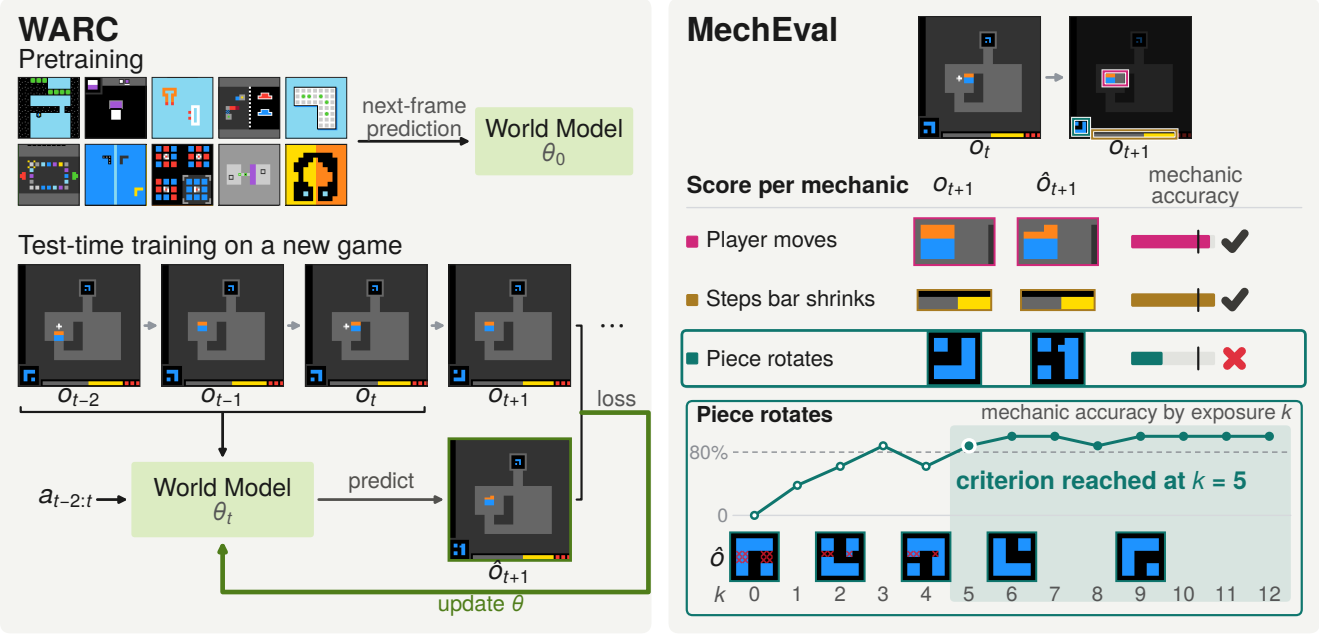

Figure 1. **WARC and MechEval.** Left: WARC is pretrained on recordings of many games (top). In a new game (bottom), it predicts the next frame from the last three frames and actions, then updates its weights. Right: For each mechanic that fires, MechEval scores the prediction made before the update on the cells it changes, as given by the game program. The exposure k counts the mechanic's earlier occurrences. It reaches the criterion at the first k from which its accuracy stays at or above 80%. The accuracies shown are illustrative.

Abstract

World models enable agents to plan and learn from imagined experience, but in an unfamiliar environment they must learn its mechanics from limited interaction. A mechanic is a rule for how an action or an interaction changes part of the environment. Each occurrence of a mechanic that the model observes is one exposure. We ask how many exposures a world model needs to predict the effects of a new environment's mechanics. We address this question on ARC-AGI-3, a benchmark of interactive games defined by their programs. We introduce MechEval, which labels 69 recurring mechanics across 20 games from these programs. It scores each prediction on the cells the mechanic changes at each exposure. As a test bed, we build WARC (World model for ARC-AGI-3), a 35M-parameter vision transformer that is pretrained on other games and updates all of its parameters after each transition. On held-out games, WARC learns to predict four in five mechanics with at least 80% accuracy on the cells they change and keeps that accuracy to the end of the recordings. Half of them reach that accuracy within about 80 exposures, in a single play of the game, and without updates only one in seven does at all. In closed-loop play on the public games, WARC-Agent, which pairs WARC with a fixed graph explorer, triples the explorer's official efficiency score from 0.15 to 0.46, mostly through test-time training. On the hidden games of the Kaggle leaderboard, our entry scores 0.20.

1. Introduction

An agent that enters a new environment must learn how it works from the little experience it gathers there. How efficiently it learns is central to intelligence, which has been defined as the efficiency of acquiring new skills [15]. ARC-AGI-3 measures this efficiency directly, comparing agents and humans by the actions they need to clear each level of interactive games they have never seen [4, 16]. To clear a level, an agent must discover the rules of the game, explore it, and find its goal. Exploring and finding the goal both depend on predicting what its actions will do. A world model provides these predictions [32, 33], but it is usually learned from 10^4 to 10^8 interactions with the environment it models [33]. In a new game, as in any new environment, it must instead learn the rules from a few interactions and predict their effects in situations it has not seen [62].

Two gaps make this regime hard to study. First, methods differ in what they predict and how they use new evidence, from an inferred context to gradient steps on latent dynamics [62, 92], and have not been compared under a common protocol. Second, there is no established way to measure the adaptation itself. Ying et al. [96] call for assessing world models by how efficiently they adapt, but the usual measures do not isolate it. An agent's return confounds model learning with exploration and planning [33, 50].

We propose to measure adaptation by how efficiently a world model learns the *mechanics* of the new environment. A mechanic is a conditional rule for how an action or an interaction changes part of the environment. For example, a tile that rotates whatever is placed on it does so wherever it appears. A model that has learned the rule can predict its effect in new places, so a few transitions inform many predictions [45, 67, 85]. Each occurrence of a mechanic adds evidence for one rule, so we count experience in *exposures*, the occurrences observed so far. Exposure does not depend on the model, so methods can be compared at equal exposure, and accuracy on occurrences in new configurations separates learning the rule from memorizing transitions. We ask: *How many exposures to a new mechanic does a world model need to predict its effects?*

Each game of ARC-AGI-3 is a program, so the rules that act at each step can be read from it. **MechEval**, an annotator built into the game programs and shown in Fig. 1, applies this measure to ARC-AGI-3. For any transition, it reports which mechanics fire and which cells each one changes. Each prediction is scored on those cells and tracked against the mechanic's exposure.

With this protocol, we study what makes a world model learn fast. Our test bed is **WARC** (World model for ARC-AGI-3), a 35M-parameter vision transformer, pretrained on other games, that predicts the next frame in pixel space and updates all of its parameters after each transition. With WARC, we compare two design choices: gradient updates versus a recurrent state, and full-parameter updates versus LoRA [38]. Finally, **WARC-Agent**, WARC inside a fixed graph explorer that ranks actions by the novelty of its predicted outcomes, tests whether faster learning carries over to closed-loop play. Our contributions are:

- **MechEval**: an annotator built into the game programs that measures per-mechanic prediction accuracy for 69 mechanics in 20 games.
- **WARC**: a simple vision transformer that learns to predict mechanics reliably in fewer exposures than the baselines. Four in five mechanics reach criterion, half of them within about 80 exposures.
- **WARC-Agent**: a simplified graph explorer [73] guided by WARC that triples the explorer's official score, Relative Human Action Efficiency (RHAE), from 0.15 to 0.46.

2. Related Work

Adapting a world model to a new environment. A world model is usually learned from the environment it models. Model-based agents such as DreamerV3 [33] learn a world model from 10^4 to 10^8 interactions with each game [37, 60, 98]. Pretrained world models transfer across games and tasks but are still finetuned on thousands of transitions or more of each new one [14, 34, 35, 94]. AdaWorld adapts from few interactions by pretraining with latent actions [29]. Test-time adaptation of world models has been studied for changes in a known quantity, such as mass, appearance, or layout. In-context methods infer a context from recent transitions [52, 62, 68, 74], amortize the update into a hypernetwork or a mixture of experts [40, 64], or learn from a synthetic prior [24]. Gradient-based methods update the model at test time. AdaJEPA does so inside the planning loop [92], others on explored transitions or in residual modules [81, 87], and CAVIA only on a context vector [100]. On static ARC puzzles, test-time training on each puzzle's demonstrations is the standard recipe [1, 39], and WARC applies it along the stream of a game. We compare both forms on the mechanics of unseen games and score the world model's predictions rather than an agent's return.

Learning the rules of a game. Game description languages such as VGDL [76] and Autumn [19] write a game as a set of rules, like our mechanics. Theory-based agents infer such a rule set from play as a program [85], and language-model agents, including several published agents on ARC-AGI-3, write and repair it as code [71, 79, 84]. These methods recover the rules explicitly, whereas we ask how many exposures a neural world model needs to predict them.

Evaluating world models. Evaluation of world models has been extended from visual fidelity to physical understanding [61] and the consistency of game mechanics [55]. ScratchWorld credits a prediction only when it computes

the fields that an action changes [54], and AutumnBench evaluates prediction and planning after a reward-free interaction phase [93]. Ying et al. [96] propose new games under a trial budget to assess how efficiently a world model adapts. In contrast, MechEval scores each mechanic at each exposure while the model learns, with ground truth from the game programs. Appendix J covers other work.

3. Preliminaries

Games. Each ARC-AGI-3 game is a turn-based environment governed by a deterministic program with a hidden state s_t [4]. At every turn, the agent observes a frame o_t and takes an action a_t. A frame is a 64×64 grid of cells in 16 colors, and an action is one of up to five keys or a click on a cell, as in Fig. 1. The state changes only when the agent acts, $s_{t+1} = T(s_t, a_t)$, where T is the game's transition function. The frame is rendered from the state, $o_t = \rho(s_t)$. Each game consists of a sequence of levels. A level ends once its hidden goal is reached, and the next level applies the same dynamics to a new layout. The setting is discrete and deterministic, and since the frame shows the whole board, it is nearly fully observable.

Mechanics. As in game description languages [19, 76, 85], a *mechanic* m pairs a binary condition c_m on the state and action with a change e_m to the objects it acts on. It *fires* whenever $c_m = 1$. The transition function is then the combined effect of the mechanics that fire,

$$T(s_t, a_t) = s_t \oplus \bigcup_{m \in M_t} e_m(s_t, a_t), \tag{1}$$

where $M_t = \{m : c_m(s_t, a_t) = 1\}$ is the set of mechanics that fire and $\oplus$ applies their changes to the state in the order the program defines. Learning a game's dynamics therefore means learning its mechanics. A mechanic is a rule as the program implements it, and we measure learning per mechanic. Each mechanic changes only a few cells, so most of the frame stays the same. Mechanics also recur during play, at other positions and in later levels.

World models and test-time training. A world model with parameters θ predicts the next frame from the last h frames and actions, $\hat{o}_{t+1} = f_\theta(o_{t-h+1:t},\, a_{t-h+1:t})$. Such a model is usually trained on the environment it models and then kept fixed. Test-time training instead updates θ with a self-supervised loss on the test input before predicting [82].

VARC. VARC [39] applies test-time training to the static puzzles of ARC-AGI-1 and ARC-AGI-2, and WARC builds on it (Sec. 5). It treats a puzzle as an image-to-image task, in which a vision transformer predicts the color of each cell of the output grid from the input grid. Each puzzle has its own learned task token that conditions the network. For a new puzzle, VARC trains the network and a new task token on the two to four demonstration pairs and their rotated, flipped, and recolored copies, then predicts the test output. A puzzle is static, with all its demonstrations given at once.

Learning along a stream. In a game, by contrast, the input arrives as a stream. A model pretrained on other games observes the transitions of an unseen game one at a time and can learn from each transition only after predicting it [20],

$$\theta_{t+1} = \theta_t - \eta\, \nabla_\theta\, \mathcal{L}\big(\theta_t;\, o_{t-h+1:t+1},\, a_{t-h+1:t}\big), \tag{2}$$

where $\mathcal{L}$ is the next-frame loss and η the learning rate. The parameters persist along the stream, as in online test-time training [82, 90]. A model can instead keep θ fixed and condition on a context of recent transitions. In both cases, the model carries what it has learned along the stream.

4. MechEval

Our question is how many exposures a model needs before it predicts the effect of a mechanic in a new game. Accuracy over all cells of the frame cannot answer this question for two reasons. First, copying the last frame already predicts 98.8% of the cells correctly, because most cells do not change. Second, an average over a game is dominated by its most frequent mechanics. MechEval therefore follows each mechanic separately along the stream of Sec. 3. It scores every *occurrence* of a mechanic, a step at which the mechanic fires, on the cells it changes. It reports this accuracy against the model's exposures to the mechanic, as in Fig. 1.

Annotation. MechEval covers 20 of the 25 public games, with 69 mechanics and 279 human recordings [43], each truncated when level 2 is first completed. Each mechanic produces one type of change, such as the rotation of the key piece when the player steps on a tile, or a counter that decreases with every move. We first described the mechanics of each game in natural language, from the game code and the screen transitions, in a codebook of 107 entries. We checked every entry against the game programs and the recordings, and instrumented the programs to log which mechanics fire at every step and on which objects. Occurrences and exposures are thus read from the program rather than inferred from the frames (Appendix B). For every step of every recording, the ground truth gives the mechanics M_t that fire and, for each, the cells of its objects that change color, with rules for cells that two mechanics share. These are all that MechEval needs from an environment. Other environments can be scored in the same way if their simulator can log them. A change that happens only once or twice, such as picking up the only key of a level, leaves no later occurrence to predict. We therefore keep a mechanic only if it fires at least three times in the median recording that contains it.

Protocol. A model follows the recorded actions and, at every step, predicts the next frame $\hat{o}_{t+1}$ before any update on

that transition (Eq. (2)). Every model observes the same occurrences in the same order, so dynamics learning is measured separately from exploration and goal inference.

Accuracy by exposure. The *exposure* of mechanic m at step t is the number of earlier steps at which it fired,

$$k_m(t) = \big|\{u < t : \ m \in M_u\}\big|, \tag{3}$$

including those whose change is not visible, which have no cells to score. It does not depend on the model, so methods are compared at equal k. For an occurrence of m at step t, $C_m(t)$ denotes the cells of its objects that change color, as given by the ground truth. The occurrence is scored by its *mechanic accuracy*,

$$\mathrm{acc}_m(t) = \frac{1}{|C_m(t)|} \sum_{g \in C_m(t)} \mathbf{1}\big[\hat{o}_{t+1}(g) = o_{t+1}(g)\big], \tag{4}$$

and its average over a mechanic's occurrences with the same k gives its *learning curve*. In Fig. 1 (right), the player and the steps bar are predicted with at least 80% accuracy at step t, and the piece rotation is not. A mechanic is fully learned only when the model also predicts its effect in a position, object, or level it has not observed.

Exposures to criterion. To summarize each learning curve by one number, we follow trials to criterion from studies of human and animal learning [78]: the number of trials a learner needs before its responses meet a criterion that the study sets, such as ten correct answers in a row. It marks when the predictions of a mechanic become reliable rather than how accurate they are on average. We count exposures instead of trials. We define the *exposures to criterion* (EtC) of a mechanic as the first k from which its learning curve stays at or above 80% for every later exposure. A mechanic whose accuracy is below 80% at its last exposure has not reached the criterion (n.r.), and its EtC is that last exposure.

5. WARC

WARC, the base model of our study, builds on VARC. It keeps the vision transformer that predicts the color of each cell, and it learns each new task at test time, as VARC does. The main change is in how it learns. VARC trains on a few demonstrations and then answers, whereas WARC predicts each transition of a game and then trains on it (Sec. 5.2). WARC also has no task token, so everything it learns about a new game is held in its weights. The other changes adapt the model and its data to games (Sec. 5.1). It takes the last h frames and actions instead of one grid, patchifies the frames jointly, and is pretrained on human recordings of other games. Tab. 6 lists every difference, including augmentations that rotate the actions with the frames.

5.1. Model

WARC is a vision transformer [22] that maps the window $x_t = (o_{t-h+1:t}, a_{t-h+1:t})$ of the last h frames and actions to a distribution $p_\theta(o_{t+1} \mid x_t)$ over the next frame. The distribution factorizes over cells, each with a distribution over the 16 colors, and the predicted frame $\hat{o}_{t+1}$ takes the most probable color at each cell. Both pretraining and test-time training minimize the negative log-likelihood of the observed next frame, which is the sum of the cross-entropies of its cells,

$$\mathcal{L}(\theta) = -\log p_\theta(o_{t+1} \mid x_t). \tag{5}$$

The model reads several frames because a few state variables are not shown in a single frame. The $h = 3$ frames are patchified jointly into 3D patches of 2×2 cells and three frames, so each token covers the same cells across time. The actions condition every block. No token identifies the game, since a new game comes with no identity and the frames show which game it is.

5.2. Learning on the Stream

The input of WARC covers only the last h frames and actions, so a transition outside this window affects a prediction only through the updates it has caused. The weights thus play the role of a recurrent state, as the weights of the inner model do in test-time training layers [83]. Section 6.4 compares these updates with a recurrent state that reads the same transitions. To learn a new mechanic from its first few exposures, WARC updates after every transition. To keep the mechanics it has already learned, it also replays earlier transitions. After each transition, it takes 20 gradient steps of Eq. (2) with the loss of Eq. (5), on minibatches of eight drawn from the ten most recent transitions, twenty older ones sampled at random, and their augmented copies. When a mechanic fires for the $(k{+}1)$-th time, WARC has already trained on all k earlier occurrences. All parameters are updated, so a new mechanic can change any part of the model. Each transition is augmented with its three rotations, with the click positions rotated with the frames and the four movement keys relabeled, which assumes that these keys move in the four directions. The updates use the pretraining optimizer, Muon with AdamW, at a constant learning rate, keeping its state across updates. Appendix C gives the architecture and hyperparameters.

6. Experiments on MechEval

We address three questions on MechEval. How quickly does WARC learn the mechanics of an unseen game, and how do the alternatives compare (Sec. 6.2)? Does it learn the mechanics or memorize the transitions it has seen (Sec. 6.3)? Which design choices in the model and the update determine its speed (Sec. 6.4)?

6.1. Setup

Evaluation. Learning curves are computed on the 20 games that MechEval covers. Each recording is processed independently. The model starts from its pretrained checkpoint, learns along the recording, and is then discarded. The 25 public games are split into three folds, and each is evaluated with the model pretrained on the human recordings of the other two folds, with augmented copies (Appendix C).

Metrics. We report EtC with the fraction of mechanics not reached (n.r., Sec. 4). The *false-edit rate* (FE) is the fraction of unchanged cells predicted to change, including steps where a move is blocked (Appendix F).

Kinds of mechanics. To break down the results, we label each mechanic by how it follows from an action (Appendix B). A direct mechanic is the effect of a key press (15 mechanics) or a click (16) on the object it acts on. A linked mechanic follows another change of the state, such as a marker that moves to the selected piece (10). A contact mechanic happens only when the action brings two objects together, such as a block that slides when pushed (6). A counter is a step or budget bar (21), and one autonomous mechanic is an agent that moves on its own.

Model selection. The test-time recipe was selected on the five public games not covered by MechEval, so that no evaluated game influenced the choice.

Comparisons. All methods receive the same frames and actions and are scored identically. Those that learn by gradient descent also receive WARC's update budget (Sec. 5.2). We compare WARC with the following baselines (Appendix D).

- **Episodic retrieval** stores the transitions observed so far and predicts each cell from the color that followed the same 3×3 neighborhood under the same action in an earlier transition. It tests whether the mechanics can be looked up rather than learned.
- **WARC (frozen)** is the pretrained model without updates and measures what pretraining alone transfers.
- **WARC (scratch)** starts from random initialization and measures what the updates learn without pretraining.
- **WARC (GRU)** reads the recording through a convolutional GRU instead of three frames and is not updated. It tests whether a recurrent state can replace the updates.
- **DreamerV3** [33] is a latent world model whose recurrent state is recomputed over the last 64 frames, so it adapts in context even without updates. It is pretrained on the same folds, updated under a recipe tuned on the same five games as WARC's, and scored on its decoded frames.

Statistics. Results are for test-time seed 0 unless stated otherwise. Over all exposures, the other two seeds of WARC are within 0.3 points of it. Curves average over occurrences within a mechanic and then over mechanics with equal weight. Intervals are 95% bootstrap intervals over games and then recordings, with 2,000 resamples.

Method	k=0↑ %	All↑ %	EtC↓ exposures	n.r.↓ %	FE↓ %
Episodic retrieval	0.1	32.9	109.0	76.8	**0.08**
WARC (frozen)	1.8	18.4	114.0	85.5	0.11
WARC (scratch)	3.9	35.0	105.0	51.2	1.15
WARC (GRU)	0.8	24.0	87.0	84.1	0.24
DreamerV3	**14.2**	25.9	114.0	90.8	3.07
WARC	3.7	**69.1**	**74.0**	**18.8**	0.11

Table 1. Learning the mechanics of unseen games (69 mechanics). k=0: mechanic accuracy before any exposure; All: over all exposures; EtC: median over all mechanics, with the last exposure for those not reached; n.r.: the fraction of mechanics not reached; FE: false-edit rate. WARC, WARC (scratch), and DreamerV3 average three test-time seeds; the other rows do not update, so their predictions are the same under every seed. Bold: best per column.

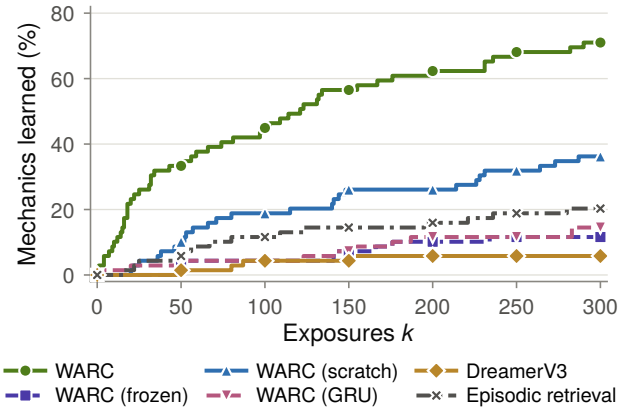

Figure 2. **Mechanics learned.** Share of the 69 mechanics whose accuracy stays at or above 80% from exposure k to their last exposure. Solid lines update the weights at test time and dashed lines do not.

6.2. Learning the Mechanics of Unseen Games

In Tab. 1 and Fig. 2, the lead of WARC over WARC (frozen) grows with exposure. Over all exposures its mechanic accuracy is 69.1%, against 18.4% without updates, and the gain holds in every fold and seed. Four in five mechanics reach criterion, half of them within about 80 exposures. WARC (frozen) reaches 10 mechanics. The gain comes from each mechanic's own transitions rather than from adapting to the game as a whole. A mechanic's accuracy stays near WARC (frozen) when the transitions in which it fires are kept out of the update buffer, with the number of updates unchanged (Appendix E). Episodic retrieval reaches 32.9% and brings 16 mechanics to criterion, mostly

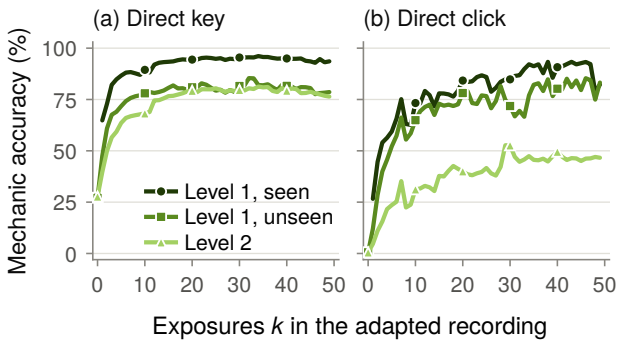

Figure 3. **Generalization to other recordings.** WARC adapts on one recording and, after k exposures to a mechanic, predicts the mechanic in the game's other recordings without updates, for (a) 14 of the 15 direct key and (b) 15 of the 16 direct click mechanics (14 in level 1), those with occurrences to predict. Level-1 predictions are split by whether the adapted recording has shown the same key acting at the same place or the same click. Beyond $k = 20$, only a few direct click mechanics have predictions.

counters and linked mechanics. Most mechanics cannot be predicted by looking up earlier transitions.

6.3. Memorization or Generalization

High accuracy could also come from memorizing transitions the model has seen. We test this on the direct mechanics, the 15 key and 16 click mechanics, where what is new can be read from the action. WARC adapts on one recording and, after k exposures to a mechanic, predicts the same mechanic in the game's other recordings without further updates. A prediction in level 1 is seen if the adapted recording has shown the same key acting at the same place or a click on the same object at the same cell, and unseen otherwise (Appendix F). Predictions in level 2, which the adapted recording has not reached, form a third group. If WARC memorized transitions, only seen predictions would improve with exposure. Instead, Fig. 3 shows that unseen predictions also rise, to about 80% for keys and 70% for clicks at ten to nineteen exposures. Seen key predictions stay easier, so part of what WARC learns is specific to them. For keys, level 2 approaches the unseen predictions, but for clicks it reaches only about 40%. In the adapted recording, moves generalize to new directions and positions. In Fig. 4, after one move up, WARC predicts the first move to the left.

A side effect. Once WARC has learned how an object moves, it also predicts the move the first time a wall or another object blocks it, so the object passes through the obstacle in the prediction. At these first blocked moves it keeps 58% of the cells that should stay unchanged, against 89% for WARC (frozen).

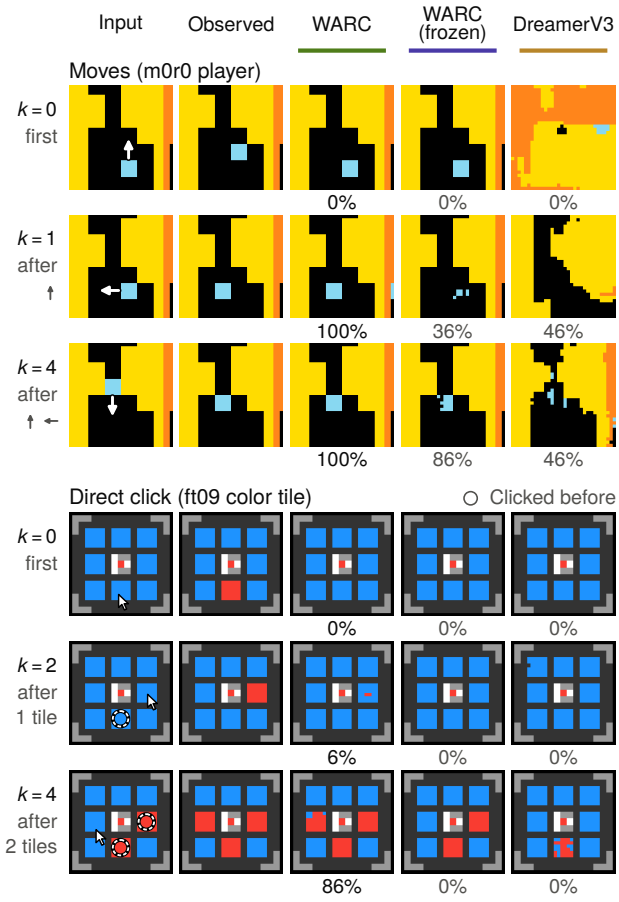

Figure 4. **Generalization to new situations.** Each row is the first time the mechanic fires ($k = 0$) or a later step new to the recording ($k \leq 5$), with the input frame and action, the observed next frame, and the predictions with their mechanic accuracy. No model predicts the first time, but WARC then predicts the first move to the left after one move up, and a click on a new tile after seeing two, which WARC (frozen) and DreamerV3 miss.

6.4. What Makes Test-Time Training Work

Pretraining. Pretraining sets how fast a new mechanic is learned. Without it, WARC (scratch) in Tab. 1 reaches 35.0%, against 69.1%, and half of the mechanics do not reach criterion, against 19%. Fig. 5 shows how pretraining scales with the number of games and of recordings per game. This experiment alone uses a variant of WARC and a different setup (Appendix H). The exact-frame error after test-time training falls steadily as games are added, whereas more recordings per game help little beyond two. Beyond the smallest size, more games are better at equal data.

Gradient updates versus a recurrent state. Fast learning needs a gradient update after every transition. In Tab. 2, updating only every five transitions leaves the latest exposures unlearned and costs 10.5 points over the first four ex-

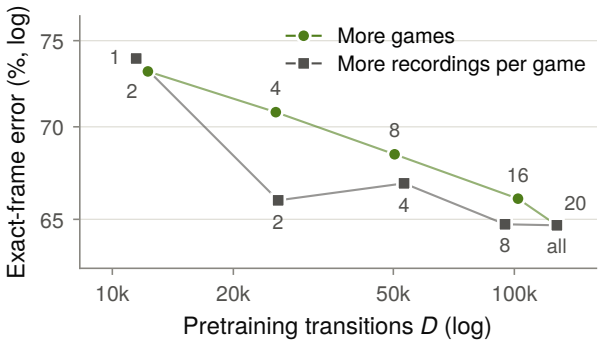

Figure 5. **Scaling pretraining data.** A variant of WARC with a token per game, pretrained on one five-fold split (Appendix H). Exact-frame error, the fraction of next frames predicted with any wrong cell, on its five held-out games after test-time training on 1,000 of their transitions, against the number of pretraining transitions D. Labels give the number of games, or of recordings per game over all 20 games.

		Accuracy (%)	
Choice	Change	All	k=1–4
WARC		69.1	41.9
Pretraining	none	−34.1	−28.2
Update	every 5 transitions	−1.7	−10.5
Updated part	LoRA, rank 64	−1.4	+3.8
	last two blocks	−4.6	−6.7
	modulation layers	−9.4	−7.2

Table 2. Ablations of WARC. Entries give the change from WARC in points when one choice changes, averaged over three test-time seeds. All, k=1–4: mechanic accuracy over all exposures and the first four. FE is 0.11–0.12% except without pretraining (1.15%). Paired intervals and the augmentation rows are in Tab. 11.

posures. The update budget and the augmentation change accuracy over all exposures by at most 1.5 points (Appendix G). The recurrent state of WARC (GRU), which reads the same transitions, does not replace these updates. Without updates, WARC (GRU) brings only 16% of the mechanics to criterion in Tab. 1, and once updated with the same recipe it matches WARC (Appendix D). DreamerV3, which also updates its weights at test time, brings 9% of the mechanics to criterion and makes many false edits. In-context adaptation emerges with pretraining on many environments [51, 69, 89], so these results concern a recurrent state pretrained on 16 or 17 games.

Full-parameter updates versus LoRA. The update need not change every parameter, since in Tab. 2 a low-rank update with LoRA [38] matches the full one. Updating only the modulation layers or only the last two blocks falls short.

Agent	L1@1000↑	L2@1000↑	RHAE@1000↑
Random	10.4	0.0	0.045
From-scratch CNN [80]	18.4	0.0	0.099
Graph	31.6	5.6	0.149
Graph + WARC (frozen)	34.4	4.4	0.181
WARC-Agent	**46.0**	**8.4**	**0.455**
Graph + oracle world model	61.6	26.8	1.392

Table 3. Closed-loop play on the 25 public games with equal weight per game, over 10 seeds. L1@N, L2@N: fraction of runs that clear level 1 or 2 within N actions; RHAE@N: the official action-efficiency score over the first N actions, on its 0–100 scale [3] (Appendix I). Graph: our explorer, alone or with the world model ranking its actions. The from-scratch CNN is StochasticGoose [80], reinitialized at every level. The oracle world model returns the true next frame, an upper bound rather than an agent. Bold: best agent.

7. From Prediction to Play

Accurate predictions alone do not make an agent act better [50]. We therefore place WARC in a fixed explorer, which proposes actions that WARC ranks by the novelty of their predicted outcomes, and call this agent WARC-Agent. We ask whether learning the mechanics in play makes exploration more efficient.

7.1. Setup

Agent. The explorer, a simplified version of a published graph explorer [73], hashes frames into states, records the actions tried in each state, and returns along known paths to states with untried actions. Its actions are the available keys, a click on each connected component of the frame, and a grid of 64 clicks, tried in priority groups. Without a model, the explorer picks an untried action of the current group at random. With WARC, it ranks the untried actions of the current group, as in model-based exploration [77]. WARC predicts the next frame for each of these actions and extends the most novel ones by a second step. The explorer takes the action whose predicted frame differs most from the frames seen in the level.

Protocol. Each agent plays each of the 25 public games with seeds 0 to 9, and every measure is read from its first 1,000 actions. Unlike in Sec. 6, WARC learns from the agent's own transitions.

7.2. Results

In Tab. 3, WARC-Agent triples the efficiency score of the explorer alone, from 0.15 to 0.46. WARC (frozen) reaches only 0.18, so test-time training raises the score 2.5 times and the share of runs that clear level 1 from 34% to 46%. The efficiency gain holds across games, with a paired interval of +0.27 [0.10, 0.48] over WARC (frozen), whereas the inter-

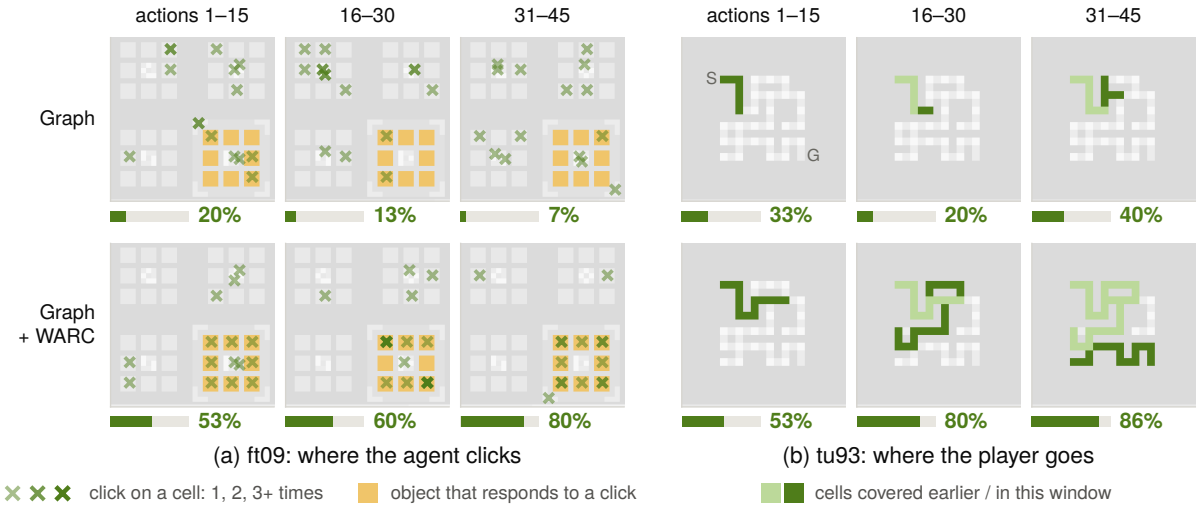

Figure 6. **Where the explorer acts in ft09 and tu93.** The first 45 actions of level 1 in three windows of 15, for the explorer alone (Graph) and for WARC-Agent (Graph + WARC). The bar under each board is the share of the window's actions that reached a game state not seen before in the level. These runs were repeated with logging and differ from those in Tab. 3, as test-time training on the GPU is not deterministic. For tu93, we show the run with the median level-1 clear among seeds 0 to 9. S and G mark the start and the goal.

val for level-1 clears includes zero. Fig. 6 shows the difference in play. In ft09, the explorer alone keeps clicking tiles that do nothing, while with WARC its clicks move to the tiles that respond. In tu93, the explorer alone keeps walking the same corridors near the start, while with WARC the player moves on to unvisited parts of the maze and reaches the goal. The explorer takes the key predicted to lead to a new state rather than one that returns or is blocked. With the explorer fixed, the score rises with the accuracy of the world model. Mechanic accuracy in Tab. 1 is 18.4% for WARC (frozen) and 69.1% for WARC, and with an oracle world model the explorer reaches an RHAE of 1.39. A world model that learns faster could therefore still raise the score. Our Kaggle entry, WARC-Agent, scores 0.20 on the public leaderboard (Appendix A).

8. Limitations

MechEval covers levels 1 and 2 of the 20 public games whose changes can be attributed to single objects. It scores a mechanic only in the situations that the human recordings reach, so reaching the criterion is weaker than learning the rule in general. This work does not model the goal of a game, which changes once per level and so falls outside the recurring mechanics we evaluate. In closed-loop play, the action selector is hand-designed and shared across conditions, so our claims concern the world model rather than the agent. Even with an oracle world model, the explorer clears no level in five games. All results come from one benchmark of deterministic, discrete games on a 64×64 grid. Whether updating the weights after each transition works as well with continuous or stochastic dynamics is untested.

9. Conclusion

We asked how many exposures a world model needs to predict the mechanics of a new game. Within a single play of an unseen game, WARC brings four in five mechanics to criterion, half of them within about 80 exposures, and without updates only one in seven reaches it. One play has far fewer interactions than the 10^4 to 10^8 from which world models usually learn their environment. What WARC learns comes from each mechanic's own transitions and carries over to other recordings of the game, including moves it has not seen. The speed comes from pretraining on other games and from a gradient update after each transition. The recurrent state of WARC (GRU), which reads the same transitions, does not replace these updates at this pretraining scale. With the action selector fixed, WARC-Agent triples the efficiency score of the explorer alone, mostly through test-time training, and an oracle world model triples it again. Higher scores need both a world model that learns faster and an action selector that aims at the goal of each game. Unlike the score of a level, MechEval shows which mechanics a model learns and after how many exposures. It can score any model that predicts the next frame, including the programs that language-model agents write to simulate a game [17, 53, 71]. These agents could in turn use a world model learned in play as a simulator, to test the actions they propose or to roll out plans toward a goal. Other environments, such as robotics simulators, can be scored in the same way if they log which mechanics fire and what each one changes, which we leave to future work.

References

[1] Ekin Akyürek, Mehul Damani, Adam Zweiger, Linlu Qiu, Han Guo, Jyothish Pari, Yoon Kim, and Jacob Andreas. The surprising effectiveness of test-time training for few-shot learning. *arXiv preprint arXiv:2411.07279*, 2024. 2

[2] Eloi Alonso, Adam Jelley, Vincent Micheli, Anssi Kanervisto, Amos Storkey, Tim Pearce, and François Fleuret. Diffusion for world modeling: Visual details matter in Atari. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 31

[3] ARC Prize Foundation. ARC-AGI-3 scoring methodology. https://docs.arcprize.org/methodology, 2026. ARC-AGI-3 Docs. Accessed 2026-10-05. 7, 29

[4] ARC Prize Foundation. ARC-AGI-3: A new challenge for frontier agentic intelligence. *arXiv preprint arXiv:2603.24621*, 2026. 2, 3

[5] Mohsen Arjmandi. Sensi: Learn one thing at a time – curriculum-based test-time learning for LLM game agents. *arXiv preprint arXiv:2603.17683*, 2026. 31

[6] Mido Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Mojtaba Komeili, Matthew Muckley, Ammar Rizvi, Claire Roberts, Koustuv Sinha, Artem Zholus, Sergio Arnaud, Abha Gejji, Ada Martin, Francois Robert Hogan, Daniel Dugas, Piotr Bojanowski, Vasil Khalidov, Patrick Labatut, Francisco Massa, Marc Szafraniec, Kapil Krishnakumar, Yong Li, Xiaodong Ma, Sarath Chandar, Franziska Meier, Yann LeCun, Michael Rabbat, and Nicolas Ballas. V-JEPA 2: Self-supervised video models enable understanding, prediction and planning. *arXiv preprint arXiv:2506.09985*, 2025. 31

[7] Jonathan Balloch, Zhiyu Lin, Mustafa Hussain, Aarun Srinivas, Robert Wright, Xiangyu Peng, Julia Kim, and Mark Riedl. NovGrid: A flexible grid world for evaluating agent response to novelty. In *AAAI Spring Symposium on Designing Artificial Intelligence for Open Worlds*, 2022. 31

[8] Amir Bar, Gaoyue Zhou, Danny Tran, Trevor Darrell, and Yann LeCun. Navigation world models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. 31

[9] Ruairidh M. Battleday, Kai Sandbrink, Jimi Cullen-Drohan, Zihan Yan, Timothy Muller, Clare Maguire, Ales Kubicek, Fraser Greenlee-Scott, Sukrit Sumant, Tri Dao, Jürgen Schmidhuber, Michal Valko, Joshua Tenenbaum, Thomas L. Griffiths, Zeb Kurth-Nelson, and James C. R. Whittington. DiG-bench: Discovery in games. *arXiv preprint arXiv:2608.12593*, 2026. 31

[10] Terrance Boult, Przemyslaw Grabowicz, Derek Prijatelj, Roni Stern, Lawrence Holder, Joshua Alspector, Mohsen M. Jafarzadeh, Toqueer Ahmad, Akshay Dhamija, Chunchun Li, Steve Cruz, Abhinav Shrivastava, Carl Vondrick, and Walter Scheirer. Towards a unifying framework for formal theories of novelty. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, pages 15047–15052, 2021. 31

[11] Jake Bruce, Michael Dennis, Ashley Edwards, Jack Parker-Holder, Yuge Shi, Edward Hughes, Matthew Lai, Aditi Mavalankar, Richie Steigerwald, Chris Apps, Yusuf Aytar, Sarah Bechtle, Feryal Behbahani, Stephanie Chan, Nicolas Heess, Lucy Gonzalez, Simon Osindero, Sherjil Ozair, Scott Reed, Jingwei Zhang, Konrad Zolna, Jeff Clune, Nando de Freitas, Satinder Singh, and Tim Rocktäschel. Genie: Generative interactive environments. *arXiv preprint arXiv:2402.15391*, 2024. 31

[12] Yuri Burda, Harrison Edwards, Amos Storkey, and Oleg Klimov. Exploration by random network distillation. *arXiv preprint arXiv:1810.12894*, 2018. 31

[13] Terry Chen, Yeyin (Eva) Zhu, Zhifan Ye, Jean-Francois Puget, and Humphrey Shi. NVIDIA AVO reaches 100% on ARC-AGI-3, demonstrating a frontier-level general-purpose architecture for long-horizon autonomous agents. https://developer.nvidia.com/blog/nvidia-avo-reaches-100-on-arc-agi-3-demonstrating-a-frontier-level-general-purpose-architecture-for-long-horizon-autonomous-agents/, 2026. NVIDIA Technical Blog, published 21 Aug 2026. Accessed 2026-10-05. 31

[14] Jie Cheng, Ruixi Qiao, Yingwei Ma, Binhua Li, Gang Xiong, Qinghai Miao, Yongbin Li, and Yisheng Lv. Scaling offline model-based RL via jointly-optimized world-action model pretraining. In *International Conference on Learning Representations (ICLR)*, 2025. 2

[15] François Chollet. On the measure of intelligence. *arXiv preprint arXiv:1911.01547*, 2019. 2

[16] François Chollet, Mike Knoop, Gregory Kamradt, and Bryan Landers. ARC prize 2025: Technical report. *arXiv preprint arXiv:2601.10904*, 2026. 2

[17] David Courtis, Wenhao Li, and Scott Sanner. OPINE-World: Programmatic world modeling with ontology-error-prioritized interactive exploration for ARC-AGI-3. *arXiv preprint arXiv:2607.01531*, 2026. 8, 31

[18] Karan Dalal, Daniel Koceja, Gashon Hussein, Jiarui Xu, Yue Zhao, Youjin Song, Shihao Han, Ka Chun Cheung, Jan Kautz, Carlos Guestrin, Tatsunori Hashimoto, Sanmi Koyejo, Yejin Choi, Yu Sun, and Xiaolong Wang. One-minute video generation with test-time training. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 17702–17711, 2025. 31

[19] Ria Das, Joshua B. Tenenbaum, Armando Solar-Lezama, and Zenna Tavares. Combining functional and automata synthesis to discover causal reactive programs. *Proceedings of the ACM on Programming Languages*, 7(POPL):56, 2023. 2, 3

[20] A. P. Dawid. Present position and potential developments: Some personal views: Statistical theory: The prequential approach. *Journal of the Royal Statistical Society. Series A (General)*, 147(2):278–292, 1984. 3

[21] Shibhansh Dohare, J. Fernando Hernandez-Garcia, Qingfeng Lan, Parash Rahman, A. Rupam Mahmood, and Richard S. Sutton. Loss of plasticity in deep continual learning. *Nature*, 632(8026):768–774, 2024. 31

[22] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner,

Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations (ICLR)*, 2021. 4

[23] Adrien Ecoffet, Joost Huizinga, Joel Lehman, Kenneth O. Stanley, and Jeff Clune. First return, then explore. *Nature*, 590:580–586, 2021. 31

[24] Fabio Ferreira, Moreno Schlageter, Raghu Rajan, Andre Biedenkapp, and Frank Hutter. One-shot world models using a Transformer trained on a synthetic prior. *arXiv preprint arXiv:2409.14084*, 2024. 2

[25] Alexis Fox, Junlin Wang, Paul Rosu, and Bhuwan Dhingra. PRO-LONG: Programmatic memory enables long-horizon reasoning. *arXiv preprint arXiv:2607.20064*, 2026. 31

[26] Paul Furgale, Severin Klingler, James Nolan, Matt Staats, Gaia Di Lorenzo, Elisa Martinez Abad, Christian Schüller, Razvan Dinu, Alessio Devoto, Pascal Berard, Gal Kaplun, Elad Sarafian, Riccardo Roveri, Leon Derczynski, and Ricardo Silveira Cabral. NVIDIA-labs OO agents: Native Python object-oriented agents. *arXiv preprint arXiv:2607.20709*, 2026. 31

[27] Chathura Gamage, Vimukthini Pinto, Cheng Xue, Peng Zhang, Ekaterina Nikonova, Matthew Stephenson, and Jochen Renz. NovPhy: A testbed for physical reasoning in open-world environments. *arXiv preprint arXiv:2303.01711*, 2023. 31

[28] Gege Gao, Bernhard Schölkopf, and Andreas Geiger. Slots, transitions, loops: Learning composable world models for ARC. *arXiv preprint arXiv:2606.12316*, 2026. 31

[29] Shenyuan Gao, Siyuan Zhou, Yilun Du, Jun Zhang, and Chuang Gan. AdaWorld: Learning adaptable world models with latent actions. In *International Conference on Machine Learning (ICML)*, 2025. 2

[30] Shivam Goel, Gyan Tatiya, Matthias Scheutz, and Jivko Sinapov. NovelGridworlds: A benchmark environment for detecting and adapting to novelties in open worlds. In *AAMAS Workshop on Adaptive Learning Agents (ALA)*, 2021. 31

[31] Shivam Goel, Yichen Wei, Panagiotis Lymperopoulos, Klára Churá, Matthias Scheutz, and Jivko Sinapov. NovelGym: A flexible ecosystem for hybrid planning and learning agents designed for open worlds. In *Proceedings of the 23rd International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pages 688–696. International Foundation for Autonomous Agents and Multiagent Systems, 2024. 31

[32] David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018. 2

[33] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse control tasks through world models. *Nature*, 640:647–653, 2025. 2, 5, 19

[34] Nicklas Hansen and Xiaolong Wang. Hallucination in world models is predictable and preventable. *arXiv preprint arXiv:2606.27326*, 2026. 2

[35] Nicklas Hansen, Hao Su, and Xiaolong Wang. Learning massively multitask world models for continuous control. *arXiv preprint arXiv:2511.19584*, 2025. 2

[36] Xianglong He, Chunli Peng, Zexiang Liu, Boyang Wang, Yifan Zhang, Qi Cui, Fei Kang, Biao Jiang, Mengyin An, Yangyang Ren, Baixin Xu, Hao-Xiang Guo, Kaixiong Gong, Size Wu, Wei Li, Xuchen Song, Yang Liu, Yangguang Li, and Yahui Zhou. Matrix-game 2.0: An open-source, real-time, and streaming interactive world model. *arXiv preprint arXiv:2508.13009*, 2025. 31

[37] Conor Heins, Toon Van de Maele, Alexander Tschantz, Hampus Linander, Dimitrije Markovic, Tommaso Salvatori, Corrado Pezzato, Ozan Catal, Ran Wei, Magnus Koudahl, Marco Perin, Karl Friston, Tim Verbelen, and Christopher Buckley. AXIOM: Learning to play games in minutes with expanding object-centric models. *arXiv preprint arXiv:2505.24784*, 2025. 2

[38] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022. 2, 7, 27

[39] Keya Hu, Ali Cy, Linlu Qiu, Xiaoman Delores Ding, Runqian Wang, Yeyin Eva Zhu, Jacob Andreas, and Kaiming He. ARC is a vision problem! *arXiv preprint arXiv:2511.14761*, 2025. 2, 3, 15, 31

[40] Jinwoo Jang, Minjong Yoo, Sihyung Yoon, and Honguk Woo. Test-time mixture of world models for embodied agents in dynamic environments. In *International Conference on Learning Representations (ICLR)*, 2026. 2

[41] Keller Jordan, Yuchen Jin, Vlado Boza, Jiacheng You, Franz Cesista, Laker Newhouse, and Jeremy Bernstein. Muon: An optimizer for hidden layers in neural networks. https://kellerjordan.github.io/posts/muon/, 2024. Blog post. 19

[42] Greg Kamradt. ARC-AGI-3 preview: 30-day learnings. https://arcprize.org/blog/arc-agi-3-preview-30-day-learnings, 2025. ARC Prize blog, published 19 Aug 2025. Accessed 2026-10-05. 31

[43] Greg Kamradt. Measuring human performance on ARC-AGI-3. https://arcprize.org/blog/arc-agi-3-human-dataset, 2026. ARC Prize blog, published 14 Apr 2026. Accessed 2026-10-05. 3

[44] Anssi Kanervisto, Dave Bignell, Linda Yilin Wen, Martin Grayson, Raluca Georgescu, Sergio Valcarcel Macua, Shan Zheng Tan, Tabish Rashid, Tim Pearce, Yuhan Cao, Abdelhak Lemkhenter, Chentian Jiang, Gavin Costello, Gunshi Gupta, Marko Tot, Shu Ishida, Tarun Gupta, Udit Arora, Ryen W. White, Sam Devlin, Cecily Morrison, and Katja Hofmann. World and human action models towards gameplay ideation. *Nature*, 638(8051):656–663, 2025. 31

[45] Ken Kansky, Tom Silver, David A. Mély, Mohamed Eldawy, Miguel Lázaro-Gredilla, Xinghua Lou, Nimrod Dorfman, Szymon Sidor, Scott Phoenix, and Dileep George. Schema networks: Zero-shot transfer with a generative causal model of intuitive physics. In *Proceedings of the 34th International Conference on Machine Learning*, pages 1809–1818, 2017. 2

[46] Seth Karten, Alex L. Zhang, Kevin Thomas, Sebastian Müller, Elie Bakouch, Daniel Auras, Mika Senghaas, Fares

Obeid, Konstantin Dunas, Johannes Hagemann, and Sami Jaghouar. Prime agent: A self-improving RLM harness. *arXiv preprint arXiv:2608.23552*, 2026. 31

[47] Seung Wook Kim, Yuhao Zhou, Jonah Philion, Antonio Torralba, and Sanja Fidler. Learning to simulate dynamic environments with GameGAN. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1231–1240, 2020. 31

[48] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017. 31

[49] Ben Krause, Emmanuel Kahembwe, Iain Murray, and Steve Renals. Dynamic evaluation of neural sequence models. In *International Conference on Machine Learning (ICML)*, pages 2766–2775, 2018. 31

[50] Nathan Lambert, Brandon Amos, Omry Yadan, and Roberto Calandra. Objective mismatch in model-based reinforcement learning. In *Learning for Dynamics and Control (L4DC)*, 2020. 2, 7

[51] Michael Laskin, Luyu Wang, Junhyuk Oh, Emilio Parisotto, Stephen Spencer, Richie Steigerwald, DJ Strouse, Steven Hansen, Angelos Filos, Ethan Brooks, Maxime Gazeau, Himanshu Sahni, Satinder Singh, and Volodymyr Mnih. In-context reinforcement learning with algorithm distillation. In *International Conference on Learning Representations (ICLR)*, 2023. 7

[52] Kimin Lee, Younggyo Seo, Seunghyun Lee, Honglak Lee, and Jinwoo Shin. Context-aware dynamics model for generalization in model-based reinforcement learning. In *International Conference on Machine Learning (ICML)*, 2020. 2

[53] Jens Lehmann, Andrei Aioanei, and Sahar Vahdati. Tycho: Active abstraction with programmatic world models for ARC-AGI-3. *arXiv preprint arXiv:2607.28287*, 2026. 8, 31

[54] Yufeng Lin and Jialu Zhang. ScratchWorld: Evaluating if world models compute executable consequences. *arXiv preprint arXiv:2606.31689*, 2026. 3

[55] Zijun Lin, Zeqing Wang, Cheston Tan, Bihan Wen, and Yeying Jin. StatePlay: State-aware game world models for mechanics-consistent generation. *arXiv preprint arXiv:2607.26754*, 2026. 2

[56] Binnan Liu, Yechi Ma, Tian Xie, and Wei Hua. TraceViT: Grounded trace supervision for visual abstract reasoning. *arXiv preprint arXiv:2607.29586*, 2026. 31

[57] Zhiguang Liu and Yi Shang. RoSeViT: Role-separated vision transformers for ARC visual reasoning. *arXiv preprint arXiv:2601.13562*, 2026. 31

[58] Ilya Loshchilov and Frank Hutter. SGDR: Stochastic gradient descent with warm restarts. In *International Conference on Learning Representations (ICLR)*, 2017. 19

[59] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019. 19

[60] Vincent Micheli, Eloi Alonso, and François Fleuret. Transformers are sample-efficient world models. In *International Conference on Learning Representations (ICLR)*, 2023. 2

[61] Saman Motamed, Laura Culp, Kevin Swersky, Priyank Jaini, and Robert Geirhos. Do generative video models understand physical principles? *arXiv preprint arXiv:2501.09038*, 2025. 2

[62] Anusha Nagabandi, Ignasi Clavera, Simin Liu, Ronald S. Fearing, Pieter Abbeel, Sergey Levine, and Chelsea Finn. Learning to adapt in dynamic, real-world environments through meta-reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2019. 2

[63] NVIDIA, Niket Agarwal, Arslan Ali, Maciej Bala, Yogesh Balaji, Erik Barker, Tiffany Cai, Prithvijit Chattopadhyay, Yongxin Chen, Yin Cui, Yifan Ding, Daniel Dworakowski, Jiaojiao Fan, Michele Fenzi, Francesco Ferroni, Sanja Fidler, Dieter Fox, Songwei Ge, Yunhao Ge, Jinwei Gu, Siddharth Gururani, Ethan He, Jiahui Huang, Jacob Huffman, Pooya Jannaty, Jingyi Jin, Seung Wook Kim, Gergely Klár, Grace Lam, Shiyi Lan, Laura Leal-Taixe, Anqi Li, Zhaoshuo Li, Chen-Hsuan Lin, Tsung-Yi Lin, Huan Ling, Ming-Yu Liu, Xian Liu, Alice Luo, Qianli Ma, Hanzi Mao, Kaichun Mo, Arsalan Mousavian, Seungjun Nah, Sriharsha Niverty, David Page, Despoina Paschalidou, Zeeshan Patel, Lindsey Pavao, Morteza Ramezanali, Fitsum Reda, Xiaowei Ren, Vasanth Rao Naik Sabavat, Ed Schmerling, Stella Shi, Bartosz Stefaniak, Shitao Tang, Lyne Tchapmi, Przemek Tredak, Wei-Cheng Tseng, Jibin Varghese, Hao Wang, Haoxiang Wang, Heng Wang, Ting-Chun Wang, Fangyin Wei, Xinyue Wei, Jay Zhangjie Wu, Jiashu Xu, Wei Yang, Lin Yen-Chen, Xiaohui Zeng, Yu Zeng, Jing Zhang, Qinsheng Zhang, Yuxuan Zhang, Qingqing Zhao, and Artur Zolkowski. Cosmos world foundation model platform for physical AI. *arXiv preprint arXiv:2501.03575*, 2025. 31

[64] Fernando Palafox and David Fridovich-Keil. Amortized low-rank adaptation for model-based reinforcement learning. *arXiv preprint arXiv:2609.12278*, 2026. 2

[65] Deepak Pathak, Pulkit Agrawal, Alexei A. Efros, and Trevor Darrell. Curiosity-driven exploration by self-supervised prediction. In *International Conference on Machine Learning (ICML)*, 2017. 31

[66] William Peebles and Saining Xie. Scalable diffusion models with transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4195–4205, 2023. 15

[67] Thomas Pouncy, Pedro Tsividis, and Samuel J. Gershman. What is the model in model-based planning? *Cognitive Science*, 45(1):e12928, 2021. 2

[68] Sai Prasanna, Karim Farid, Raghu Rajan, and André Biedenkapp. Dreaming of many worlds: Learning contextual world models aids zero-shot generalization. In *Reinforcement Learning Conference (RLC)*, 2024. 2

[69] Allan Raventós, Mansheej Paul, Feng Chen, and Surya Ganguli. Pretraining task diversity and the emergence of non-Bayesian in-context learning for regression. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 7

[70] Sergey Rodionov. Do coding agents need executable world models, simplification, and verification to solve ARC-AGI-3? *arXiv preprint arXiv:2607.15439*, 2026. 31

[71] Sergey Rodionov. Executable world models for ARC-AGI-3 in the era of coding agents. In *Artificial General Intelligence (AGI 2026), Lecture Notes in Computer Science*, pages 198–210. Springer Nature Switzerland, 2026. 2, 8

[72] David Rolnick, Arun Ahuja, Jonathan Schwarz, Timothy Lillicrap, and Gregory Wayne. Experience replay for continual learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019. 31

[73] Evgenii Rudakov, Jonathan Shock, and Benjamin Ultan Cowley. Graph-based exploration for ARC-AGI-3 interactive reasoning tasks. In *AAAI 2026 Workshop on AI for Scientific Research*, 2026. 2, 7, 28, 31

[74] Frank Röder, Jan Benad, Manfred Eppe, and Pradeep Kr. Banerjee. Dynamics-aligned latent imagination in contextual world models for zero-shot generalization. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. 2

[75] Elad Sarafian, Gal Kaplun, Ron Banner, Daniel Soudry, and Boris Ginsburg. Workspace optimization: How to train your agent. *arXiv preprint arXiv:2605.09650*, 2026. 31

[76] Tom Schaul. A video game description language for model-based or interactive learning. In *IEEE Conference on Computational Intelligence in Games (CIG)*, 2013. 2, 3

[77] Ramanan Sekar, Oleh Rybkin, Kostas Daniilidis, Pieter Abbeel, Danijar Hafner, and Deepak Pathak. Planning to explore via self-supervised world models. In *International Conference on Machine Learning (ICML)*, 2020. 7, 31

[78] Roger N. Shepard, Carl I. Hovland, and Herbert M. Jenkins. Learning and memorization of classifications. *Psychological Monographs: General and Applied*, 75(13):1–42, 1961. 4

[79] Alexy Skoutnev, Kirill Acharya, Gaston Longhitano, Madeleine Udell, Kevin Ellis, and Iddo Drori. Twin: Playing an unknown game with a test-time digital twin. *arXiv preprint arXiv:2608.14490*, 2026. 2

[80] Dries Smit. ARC-AGI challenge 3 — StochasticGoose. https://github.com/DriesSmit/ARC3-solution, 2025. Tufa Labs. Submission to the ARC-AGI-3 Developer Preview Agent Competition. Accessed 2026-10-05. 7, 31

[81] Krishnam Soni, Aditya Sehgal, Vedant Dave, and Elmar Rueckert. Sandwich-residuals: Parameter-efficient test-time adaptation of world models. *arXiv preprint arXiv:2609.21740*, 2026. 2

[82] Yu Sun, Xiaolong Wang, Zhuang Liu, John Miller, Alexei A. Efros, and Moritz Hardt. Test-time training with self-supervision for generalization under distribution shifts. In *International Conference on Machine Learning (ICML)*, 2020. 3, 31

[83] Yu Sun, Xinhao Li, Karan Dalal, Jiarui Xu, Arjun Vikram, Genghan Zhang, Yann Dubois, Xinlei Chen, Xiaolong Wang, Sanmi Koyejo, Tatsunori Hashimoto, and Carlos Guestrin. Learning to (Learn at test time): RNNs with expressive hidden states. In *International Conference on Machine Learning (ICML)*, pages 57503–57522, 2025. 4, 31

[84] Hao Tang, Darren Key, and Kevin Ellis. WorldCoder, a model-based LLM agent: Building world models by writing code and interacting with the environment. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 2

[85] Pedro A. Tsividis, Joao Loula, Jake Burga, Nathan Foss, Andres Campero, Thomas Pouncy, Samuel J. Gershman, and Joshua B. Tenenbaum. Human-level reinforcement learning through theory-based modeling, exploration, and planning. *arXiv preprint arXiv:2107.12544*, 2021. 2, 3

[86] Dani Valevski, Yaniv Leviathan, Moab Arar, and Shlomi Fruchter. Diffusion models are real-time game engines. In *International Conference on Learning Representations (ICLR)*, 2025. 31

[87] Chi Wan, Kangrui Wang, Yuan Si, Pingyue Zhang, and Manling Li. WorldAgen: Unified state-action prediction with test-time world model training. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2026. 2

[88] Dequan Wang, Evan Shelhamer, Shaoteng Liu, Bruno Olshausen, and Trevor Darrell. Tent: Fully test-time adaptation by entropy minimization. In *International Conference on Learning Representations (ICLR)*, 2021. 31

[89] Fan Wang, Zhiyuan Chen, Yuxuan Zhong, Sunjian Zheng, Pengtao Shao, Bo Yu, Shaoshan Liu, Jianan Wang, Ning Ding, Yang Cao, and Yu Kang. Context and diversity matter: The emergence of in-context learning in world models. *arXiv preprint arXiv:2509.22353*, 2025. 7

[90] Qin Wang, Olga Fink, Luc Van Gool, and Dengxin Dai. Continual test-time domain adaptation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7201–7211, 2022. 3, 31

[91] Renhao Wang, Yu Sun, Arnuv Tandon, Yossi Gandelsman, Xinlei Chen, Alexei A. Efros, and Xiaolong Wang. Test-time training on video streams. *Journal of Machine Learning Research*, 26(9):1–29, 2025. 31

[92] Ying Wang, Oumayma Bounou, Yann LeCun, and Mengye Ren. AdaJEPA: An adaptive latent world model. *arXiv preprint arXiv:2606.32026*, 2026. 2

[93] Archana Warrier, Dat Nguyen, Michelangelo Naim, Moksh Jain, Yichao Liang, Karen Schroeder, Cambridge Yang, Joshua B. Tenenbaum, Sebastian Vollmer, Kevin Ellis, and Zenna Tavares. Benchmarking world-model learning with environment-level queries. In *International Conference on Machine Learning (ICML)*, 2026. 3

[94] Yifan Xu, Nicklas Hansen, Zirui Wang, Yung-Chieh Chan, Hao Su, and Zhuowen Tu. On the feasibility of cross-task transfer with model-based reinforcement learning. *arXiv preprint arXiv:2210.10763*, 2022. 2

[95] Sherry Yang, Yilun Du, Kamyar Ghasemipour, Jonathan Tompson, Leslie Kaelbling, Dale Schuurmans, and Pieter Abbeel. Learning interactive real-world simulators. In *International Conference on Learning Representations (ICLR)*, 2024. 31

[96] Lance Ying, Katherine M. Collins, Prafull Sharma, Cedric Colas, Kaiya Ivy Zhao, Adrian Weller, Zenna Tavares,

Phillip Isola, Samuel J. Gershman, Jacob D. Andreas, Thomas L. Griffiths, Francois Chollet, Kelsey R. Allen, and Joshua B. Tenenbaum. Assessing adaptive world models in machines with novel games. *arXiv preprint arXiv:2507.12821*, 2025. 2, 3

[97] Tianyuan Zhang, Sai Bi, Yicong Hong, Kai Zhang, Fujun Luan, Songlin Yang, Kalyan Sunkavalli, William T. Freeman, and Hao Tan. Test-time training done right. *arXiv preprint arXiv:2505.23884*, 2025. 31

[98] Weipu Zhang, Gang Wang, Jian Sun, Yetian Yuan, and Gao Huang. STORM: Efficient stochastic Transformer based world models for reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. 2

[99] Gaoyue Zhou, Hengkai Pan, Yann LeCun, and Lerrel Pinto. DINO-WM: World models on pre-trained visual features enable zero-shot planning. In *International Conference on Machine Learning (ICML)*, 2025. 31

[100] Luisa M. Zintgraf, Kyriacos Shiarlis, Vitaly Kurin, Katja Hofmann, and Shimon Whiteson. Fast context adaptation via meta-learning. In *International Conference on Machine Learning (ICML)*, 2019. 2

Appendix

A. Kaggle Submission

Submission. Our entry to the ARC-AGI-3 competition on Kaggle runs WARC-Agent (Appendix I), in which WARC ranks the explorer's actions and updates with the recipe of Sec. 5.2 after every valid transition. The competition scores agents on 110 hidden games, none of which is among the 25 public games. Since pretraining on more games helps (Appendix H), the submitted model is pretrained on all 25 public games with the recipe of Appendix C. The rest of the paper holds out games by fold, so that every evaluated game is unseen in pretraining, as the hidden games are. The notebook, the agent, the pretraining and test-time training code, and the weights are at `https://github.com/third-intelligence/WARC`. We will release the annotations and the scorer of MechEval. The submission scores 0.20 on the public leaderboard, which uses half of the hidden games. The hidden games are not public, so we cannot tell how they differ from the public games, and the score on them may differ widely from our results on the public games.

Time budget. The notebook plays the hidden games one at a time on one RTX PRO 6000 Server Edition GPU without network access and must finish within nine hours. We gave the games 8.6 hours in total, so each game gets about 280 s. We timed the notebook on the same GPU type in two games played with keys. In these games, an action takes 0.74 s, most of it for the 20 gradient steps after each transition, so a game gets about 380 actions. Games with clicks get fewer actions, because the model predicts the next frame for more candidates at each step. The runs of Tab. 3 play 1,000 actions per game with no time limit. Cut at 380 actions, they clear level 1 in 13% of runs with Graph, 19% with Graph + WARC (frozen), and 40% with WARC-Agent, against 32%, 34%, and 46% at 1,000 actions. Their RHAE hardly changes (0.144, 0.177, and 0.451, against 0.149, 0.181, and 0.455), and is 0.143, 0.175, and 0.448 when cut at 250 actions.

B. MechEval Details

Games and recordings. We use 20 of the 25 public games. In the other five, changes cannot be assigned to single objects, for example because joints overlap in cn04 or the whole level is drawn as one sprite in bp35 and lf52. We replay each recording through the game with logging added and keep only the segments whose frames match the recording exactly. Tab. 4 gives the resulting data per game.

Mechanics. A mechanic is a branch of the game program that updates the state, defined by its trigger, the object it changes, and its effect. Parts that always change together form one object, and a change to another object is a separate mechanic. We started from the 107 entries of the codebook (Sec. 4), checked each against the programs and the recordings, and had two raters check the result. This gives 81 mechanics. Three of them never fire in the recordings, two never change a visible cell, and seven fire fewer than three times in the median recording that contains them (Sec. 4). The seven are the key pickup and the switch highlight in dc22, the return to the start in ls20, the goal flash in r11l, the resizing of the player in sc25, the selector button in tn36, and the enemy ball in tu93. The remaining 69 mechanics are scored. Figs. 7 to 9 show every mechanic at a representative occurrence.

Kinds. Tab. 5 defines the kinds. The moves of Sec. 6.3 are the 13 direct key mechanics that move, tilt, or extend an object. Two raters who did not see our kinds sorted the mechanics into direct, counter, and other. Our kinds agree with each rater on 86% of the mechanics (Cohen's κ 0.77), and the raters agree with each other on 87% (κ 0.78).

Occurrences and scored cells. Mechanic m fires at step t if its branch runs and the state of its object changes. We compare the states after any animation has ended. The scored cells $C_m(t)$ of m are the changed cells inside the bounding boxes of its objects, taken before and after the step. When the program's log does not name every changed object, or the boxes of two mechanics overlap, rules derived only from the programs and the frames assign the cells. For example, a selection marker keeps only its own cells and leaves the selected object to that object's mechanic. After these rules, 0.03% of the changed cells lie outside every region. Steps where a mechanic's condition fails, such as a move blocked by a wall, are not occurrences and do not count as exposures. They are recorded separately (Appendix F).

	VARC	WARC
Task	output grid of a static puzzle	next frame of an interactive game
Input	one grid on a 64×64 canvas	the last three 64×64 frames and actions
Transformer	width 512, 10 blocks, 8 heads, 2×2 patches	the same, with 3D patches over the three frames
Output	a color for every cell	the same
Conditioning	one learned token per puzzle	the actions through adaptive layer normalization and a learned vector at a clicked patch, no task token
Parameters	18M	35.3M, of which 15.8M modulate the blocks by the actions and 3.1M embed the 3D patches
Pretraining data	the 400 ARC training puzzles and RE-ARC, about 400k pairs	human recordings of 16 or 17 games, 90k to 123k transitions
Pretraining augmentation	scale and translation on the canvas	three rotations and two flips with ten color sets, 51 copies, actions transformed with the frames
Test-time data	the 2 to 4 demonstration pairs of the puzzle	the stream so far, recent and older transitions
Test-time schedule	100 epochs, then the prediction	20 steps after every transition, each transition predicted before training on it
Updated at test time	the network and a new task token	all parameters
Test-time augmentation	flips, rotations, and color permutations, 51 views, each with its own token	three rotations, 31 copies, actions and clicks rotated with the frames
Test-time optimizer	Adam, cosine schedule	Muon with AdamW, constant learning rate
Inference	majority vote over 510 views	one forward pass per step

Table 6. What WARC takes from VARC and what it changes (Sec. 5). The VARC column follows Hu et al. [39] for their default model.

Game	Rec.	Trans.	Mech.	Occ.	Med.	Game	Rec.	Trans.	Mech.	Occ.	Med.
ar25	10	1,002	4	1,570	32	s5i5	11	1,942	2	2,460	62
cd82	11	654	5	824	7	sb26	12	824	3	973	22.5
dc22	11	1,632	4	2,064	22.5	sc25	15	694	3	865	13
ft09	10	547	2	636	28	sk48	14	3,730	5	6,280	66
ka59	10	1,163	5	1,611	16	su15	13	844	3	1,262	27
lp85	54	4,266	3	6,834	24	tn36	14	1,462	3	2,351	34
ls20	13	2,411	4	4,745	35	tr87	12	1,696	3	2,539	64.5
m0r0	11	1,748	4	4,169	88	tu93	13	492	2	938	32.5
r11l	10	533	3	985	27	vc33	10	325	3	659	21
re86	11	1,008	3	1,653	41	wa30	14	2,920	5	6,178	62
All	279	29,893	69	49,596	30						

Table 4. MechEval data per game, up to the first completion of level 2: recordings (Rec.), transitions (Trans.), scored mechanics (Mech.), their scored occurrences (Occ.), and the median number of occurrences of a mechanic in a recording in which it fires (Med.).

Kind	Definition	Mechanics	Example
Direct key	A key moves, tilts, extends, or changes an object	15	ls20 player moves
Direct click	A click changes the clicked object or the object that a button drives	16	lp85 tile row shifts
Linked	Copies or marks another change	10	re86 piece selection
Contact	An object reaches a place or touches another object	6	ka59 pushed block slides
Counter	A step or budget bar changes	21	ls20 steps bar
Autonomous	An object moves by its own rule	1	wa30 orange agent

Table 5. The kinds of the 69 mechanics.

C. WARC Details

Architecture. Each cell takes one of 16 colors, or a padding symbol before the start of an episode or level. A 3D convolution with kernel and stride $3\times2\times2$ turns the three embedded frames into 32×32 tokens with rotary position embeddings, followed by the transformer of VARC (Tab. 6). The current action is embedded by its type and, for a click, by its x and y coordinates, and conditions every block through adaptive layer normalization initialized at zero [66]. A click also adds a learned vector to the token of the clicked patch, one for each of the four cells in the patch. The two earlier actions are embedded in the same way and added to every token through a linear layer initialized at zero, and the output head predicts the 2×2 cells of each token.

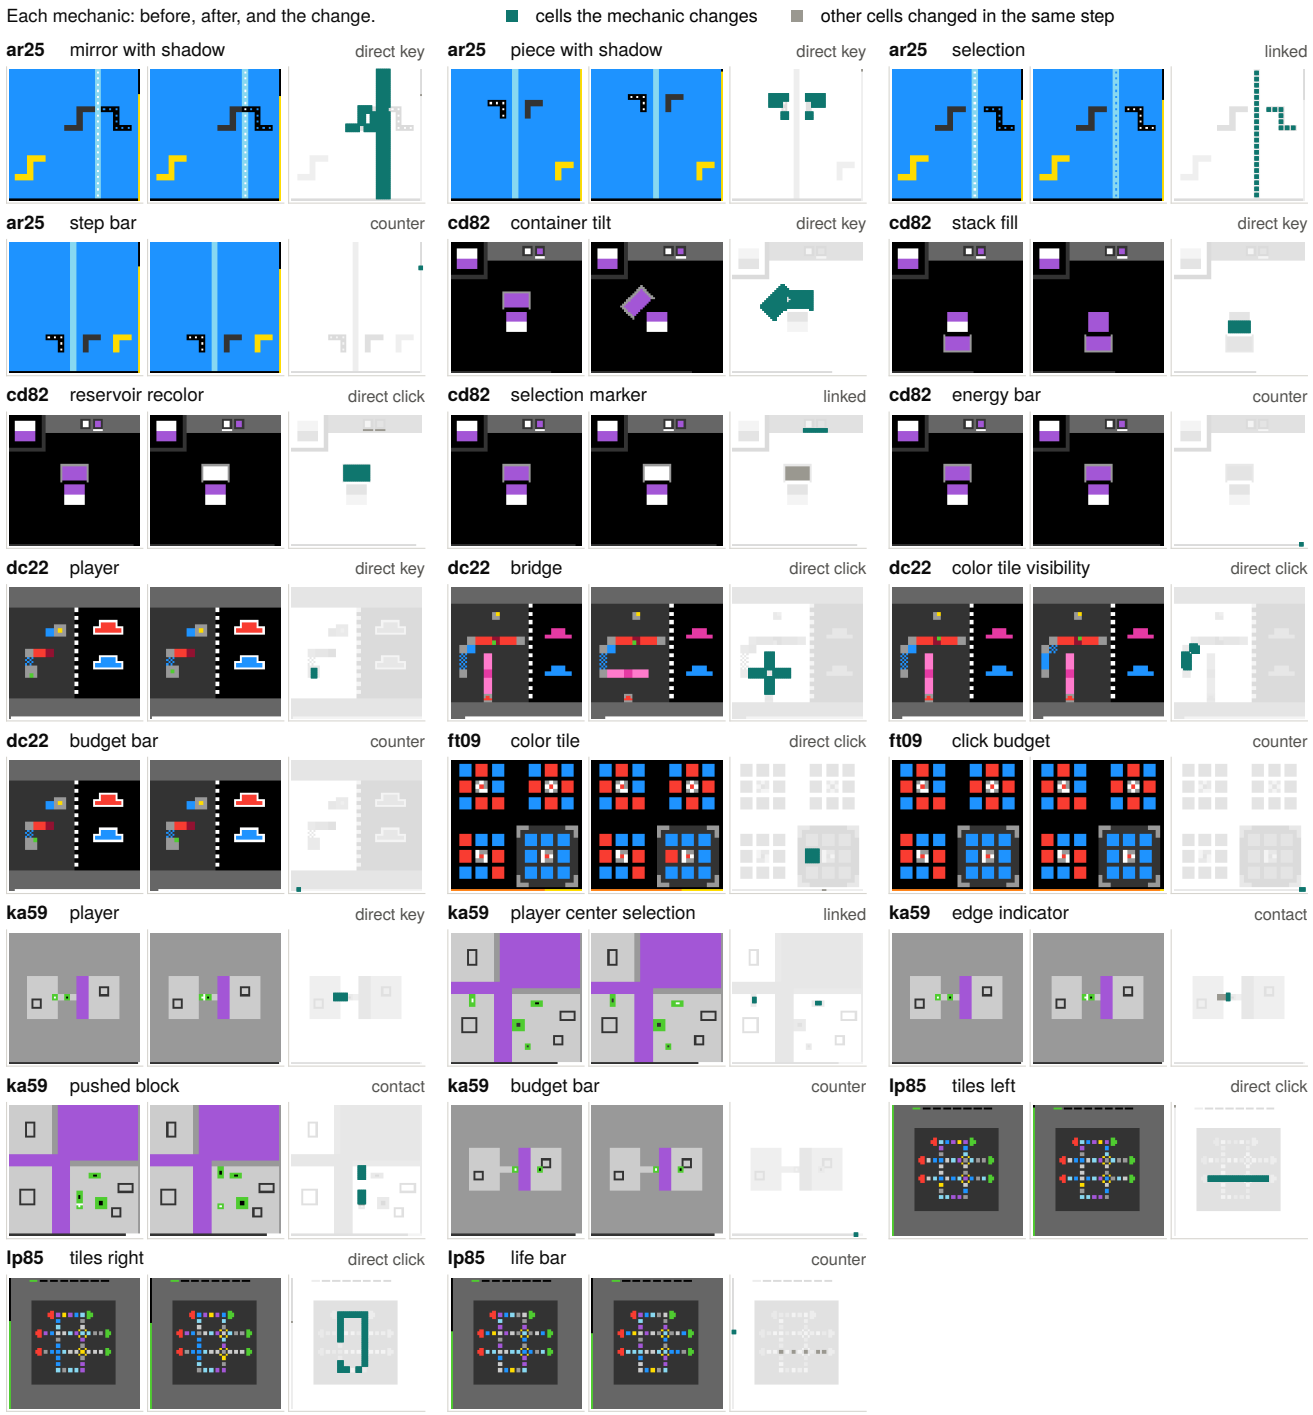

Figure 7. **The mechanics of MechEval, game by game: ar25 to lp85** (ls20 to wa30 in Figs. 8 and 9). Each mechanic is shown at a representative occurrence: among its scored occurrences in all recordings, one whose number of scored cells is closest to the mechanic's median, with the fewest other cells changed in the same step. The panels show the frame before, the frame after, and the change, with the scored cells $C_m(t)$ in teal and the other cells changed in the same step in gray. Within a game, the mechanics are ordered by kind, which the label above each panel names (Tab. 5).

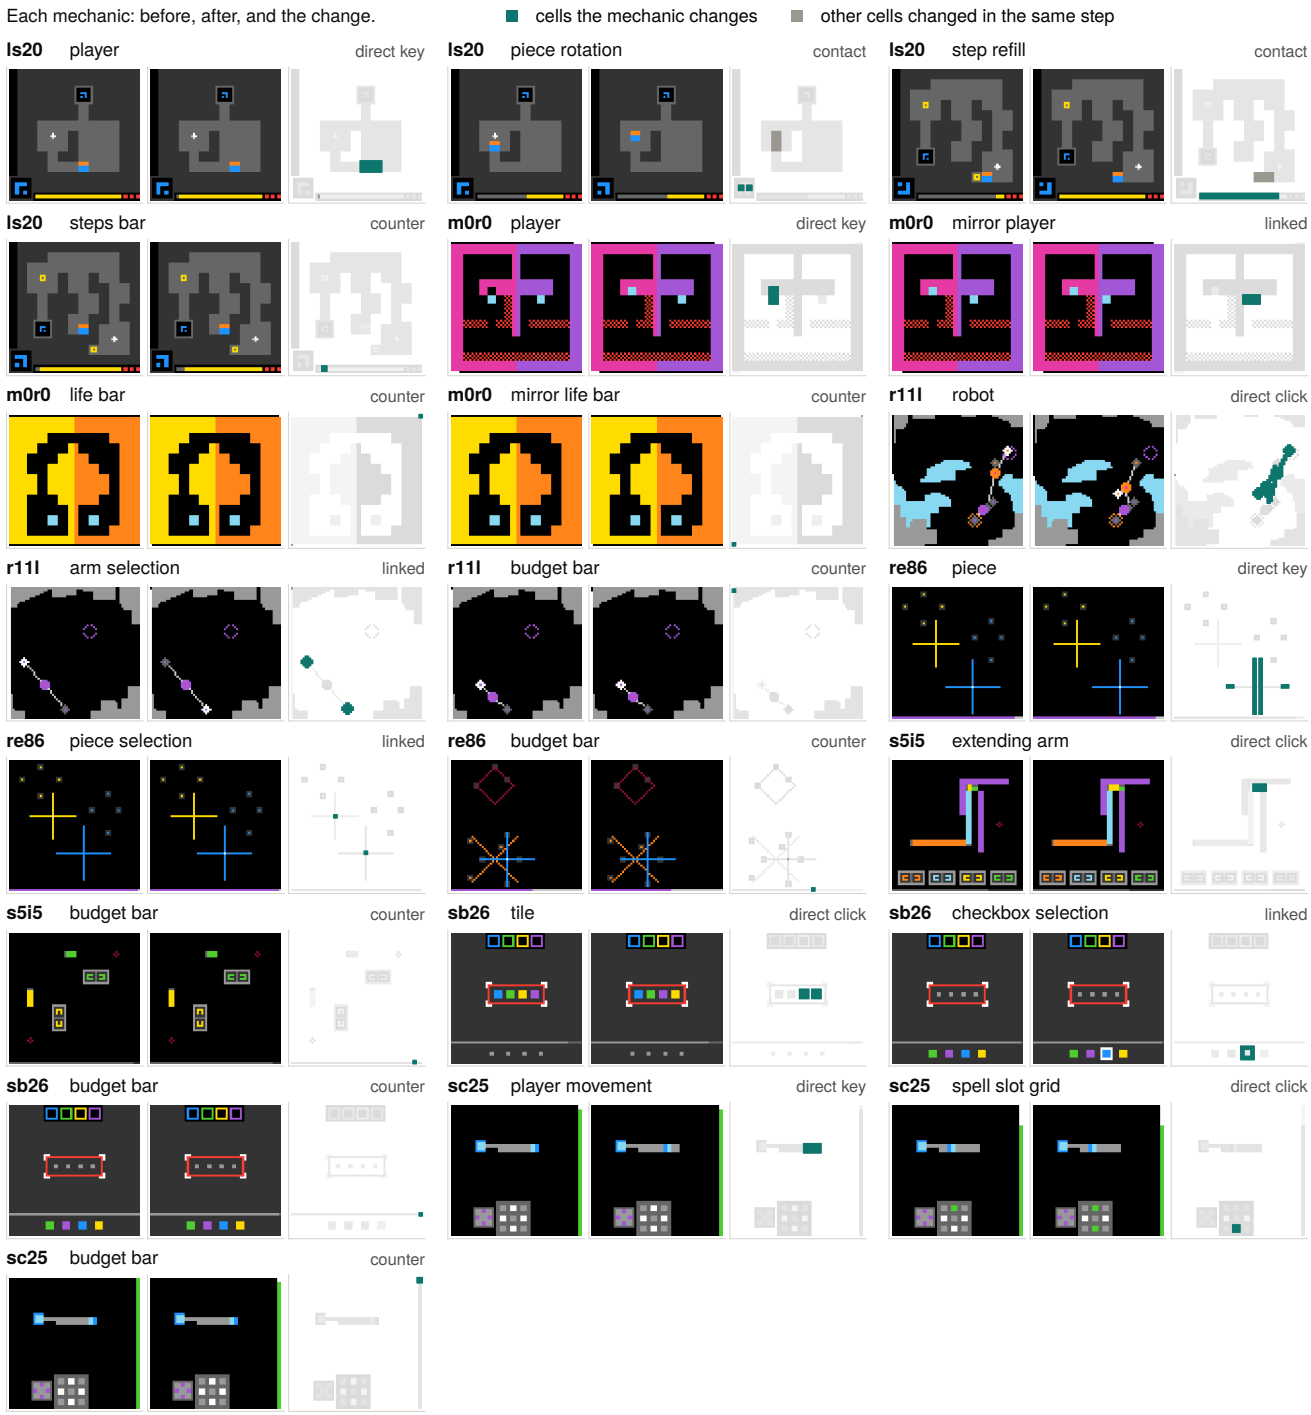

Figure 8. **The mechanics of MechEval, game by game: ls20 to sc25**, shown as in Fig. 7. The ls20 step refill is shown at a refill of the steps bar, since half of its occurrences take the item while the bar is full.

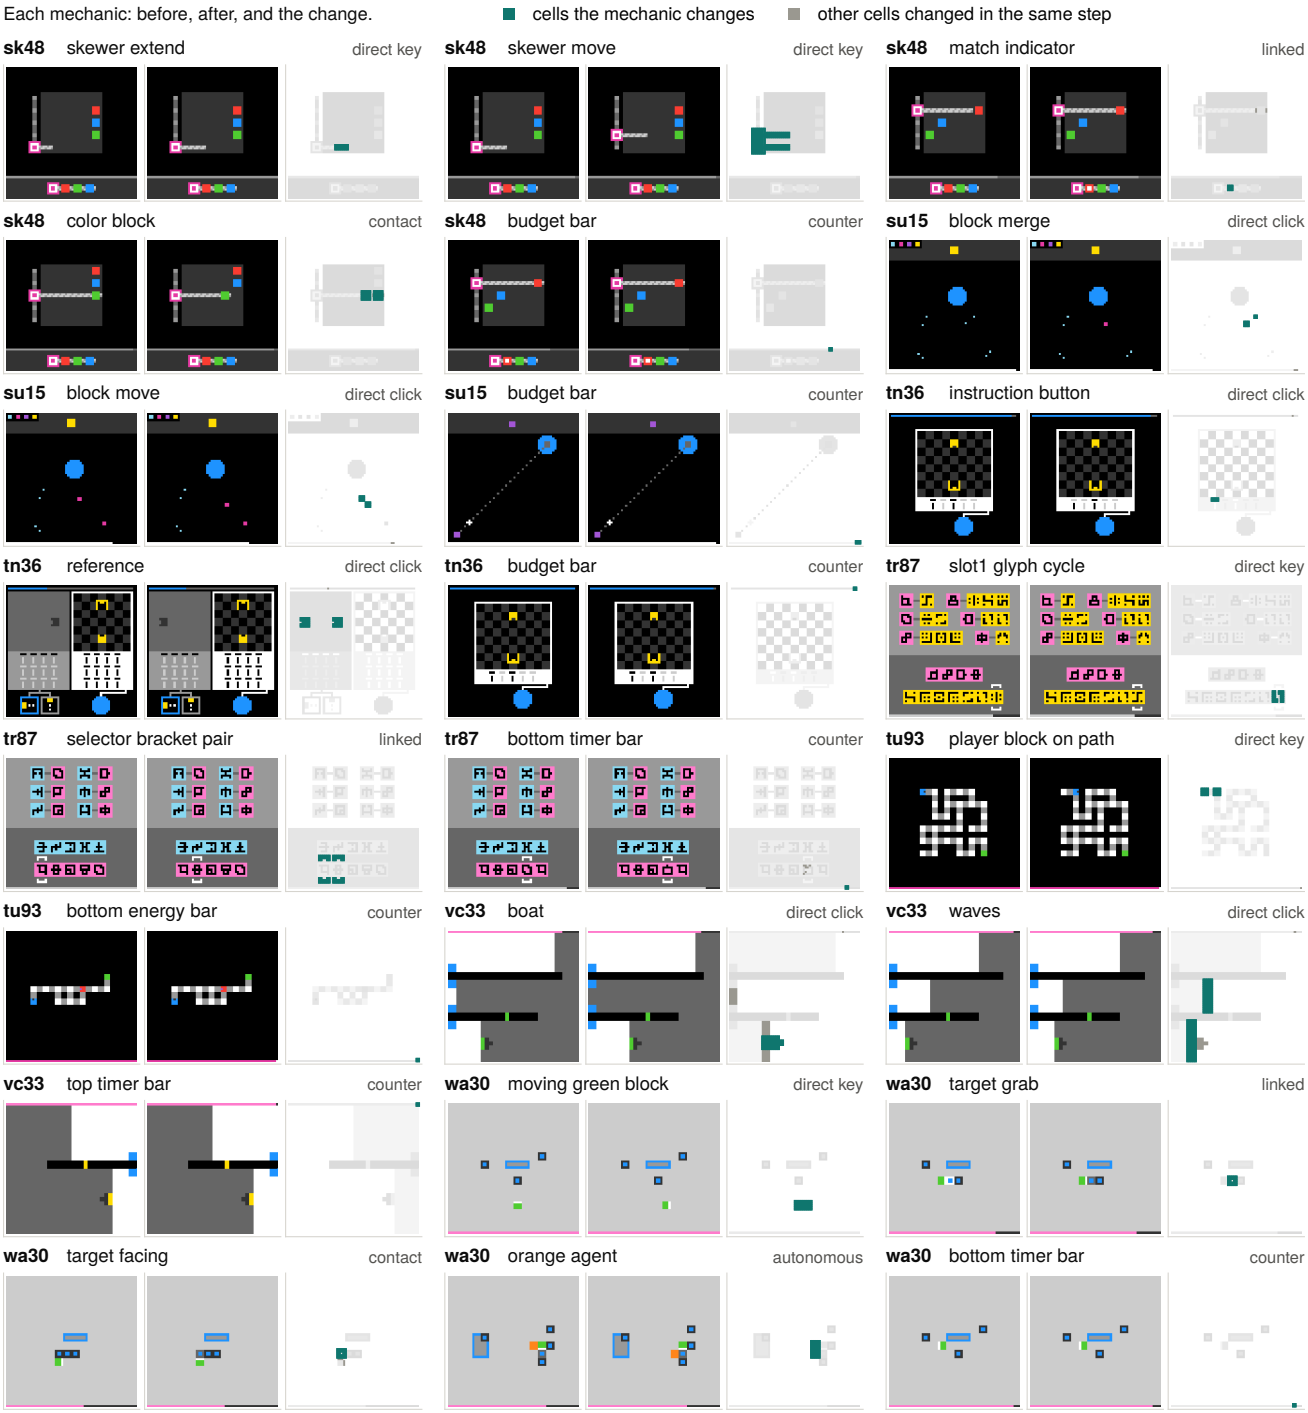

Figure 9. **The mechanics of MechEval, game by game: sk48 to wa30**, shown as in Fig. 7.

	Pretraining	Test-time training
Optimizer	Muon + AdamW	Muon + AdamW
Learning rate	3×10^{-4}	3×10^{-4}
Weight decay	0	0.01
Schedule	warmup, cosine	constant
Batch	516	8
Length	30 epochs	20 steps per transition
Augmented copies	51	31 (rotations)
Precision	bfloat16	bfloat16
Gradient clipping	1.0	1.0
Compute	13–20 h per fold on six GPUs	one GPU (Appendix A)

Table 7. Optimization settings. In both stages, Muon [41] updates the two-dimensional weights of the transformer blocks except the modulation layers, with momentum 0.95, and AdamW [59] updates the remaining parameters. Pretraining warms up the learning rate linearly over six epochs and then decays it with a cosine schedule [58]. The GPUs are RTX PRO 6000, and test-time training runs on one.

Pretraining. The 25 public games are split into three folds at random. Each fold holds two of the six games without clicks, and the five games outside MechEval are spread over the folds as two, two, and one. We train each fold model on the human recordings of its 16 or 17 training games, 90k to 123k transitions after removing resets. For each training game, one recording that reaches the last level is held out to select the checkpoint by exact-frame accuracy, the fraction of next frames predicted with every cell correct. Each recording is augmented into 51 copies, the original and the 50 combinations of five rotations and flips with the original colors or one of nine color permutations, with the actions transformed with the frames.

Test-time training. A transition is valid unless it crosses a reset, a level change, or the end of an episode, and the history window restarts at these boundaries. After every valid transition, the model takes 20 gradient steps on the ten most recent valid transitions and twenty older ones drawn at random, with a new draw at every update. Each drawn transition enters the minibatches as the original and its three rotations, each rotation repeated ten times (31 copies). The optimizer state persists across updates, and every parameter is updated. Tab. 7 lists the optimization settings of both stages.

Choosing the recipe. We chose the recipe on the five public games outside MechEval (61 recordings), each with the fold model that holds it out, by a rule fixed before any run. The score is the accuracy on the changed cells of every valid transition that changes the frame, predicted before the update on that transition. Candidates whose false-edit rate exceeds twice the lowest are discarded. Starting from AdamW at 10^{-4} with five recent and ten older transitions and rotations with color permutations, we changed one setting at a time. This gave Muon at 3×10^{-4}, ten recent and twenty older transitions, and rotations only. The final recipe reaches 82.8% over three seeds, against 82.0% for the starting point.

D. Baselines

The learned baselines, DreamerV3 and WARC (GRU), are pretrained on the same folds as WARC, and when they are updated, they take 20 gradient steps on minibatches of eight after every valid transition.

Episodic retrieval. Episodic retrieval stores every valid transition of the recording so far and has no parameters. For each cell, it predicts the most frequent next color among earlier cells with the same 3×3 neighborhood under the same action type, breaking ties by recency. For a click, the earlier cell must also lie at the same offset from the clicked cell, up to two cells away. A cell without a match is predicted unchanged. A 5×5 neighborhood scores 34.5% against 32.9%, and copying the most similar earlier transition scores at most 19.1%.

DreamerV3. We use a PyTorch port of the DreamerV3 world model [33] that matches the official JAX implementation. We take its 50M configuration with a CNN depth of 64 (45.9M parameters). A categorical decoder over the 16 colors lets it be scored cell by cell like WARC, and the action is factored into its type and click coordinates. It is pretrained with only the world-model loss for 50 epochs on batches of 16 sequences of 64 steps, with the depth and the learning rate chosen on fold 1. We keep the checkpoint with the best accuracy over all cells of the next frame on held-out recordings. At test time, it predicts from a state recomputed over the last 64 frames. Every part of the model is updated, and each transition is trained as the last step of a sequence of up to nine frames. Its update was tuned on the five games outside MechEval by the rule of Appendix C. The tuning gave AdamW at a constant 3×10^{-4}, five recent and ten older transitions, and rotations only. Muon was not tried, since the model was not pretrained with it.

WARC (GRU). WARC (GRU) keeps the blocks, the click vectors, and the head of WARC and replaces the three-frame

Method	k=0↑ %	All↑ %	EtC↓ exposures	n.r.↓ %	FE↓ %
WARC (frozen)	1.8	18.4	114.0	85.5	0.11
WARC	3.3	69.1	74.0	18.8	0.11
WARC (GRU)	0.8	24.0	87.0	84.1	0.24
WARC (GRU), updated	2.5	70.0	74.0	18.8	0.11

Table 8. **Gradient updates versus a recurrent state.** WARC and WARC (GRU) with and without updates on all three folds (69 mechanics), with the columns of Tab. 1, in test-time seed 0.

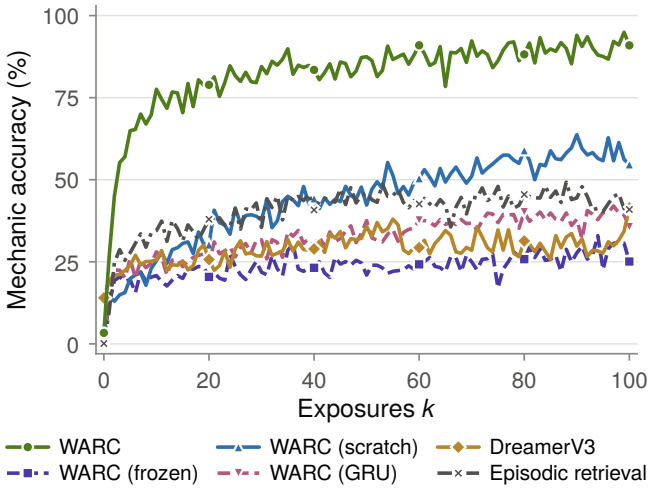

Figure 10. **Learning curves of every method.** Mechanic accuracy against exposure k, averaged over the mechanics with an occurrence at k, up to $k = 100$, where about half of them remain. Solid lines update the weights at test time and dashed lines do not.

Method	70%	80%	90%
Episodic retrieval	109.0 (76.8)	109.0 (76.8)	109.0 (76.8)
WARC (frozen)	114.0 (85.5)	114.0 (85.5)	114.0 (87.0)
WARC (scratch)	94.7 (44.9)	105.0 (51.2)	114.0 (56.5)
WARC (GRU)	87.0 (82.6)	87.0 (84.1)	87.0 (85.5)
DreamerV3	114.0 (87.4)	114.0 (90.8)	114.0 (91.8)
WARC	**61.0** (**16.9**)	**74.0** (**18.8**)	**77.0** (**20.3**)

Table 9. EtC, with n.r. (%) in parentheses, for accuracy thresholds of 70%, 80% (Tab. 1), and 90%. WARC, WARC (scratch), and DreamerV3 average three test-time seeds.

window by a convolutional GRU over the 32×32 patch grid. Each transition is encoded into a 64-channel map from its frames, its action, and a reset flag, and the GRU reads these maps in order from a zero state at the start of the recording, ending with the current frame. The state, mapped to the token width, forms the input tokens, so the model sees the frames only through the state. The action also enters as a prefix token and through adaptive layer normalization. The model has 32.2M parameters. We pretrain one model per fold, reading each recording in chunks of 64 transitions with the state carried across chunks. In half of the recordings, the five keys are permuted in the action labels, so that what a key does must be read from the state. Test-time training uses the recipe of WARC and keeps the encoder and the GRU fixed, since they produce the stored states. Tab. 8 gives it with and without updates, next to WARC, in test-time seed 0. Updated, it matches WARC over all exposures, with a paired difference of $+0.9$ points $[-0.1, +1.9]$.

E. Learning Curves

All methods. Fig. 10 shows the mechanic accuracy of every method of Tab. 1 at each exposure. No model predicts more than a fifth of the changed cells at $k = 0$, and from $k = 1$ to $k = 100$ WARC is the most accurate model at every exposure.

Choice of threshold. The criterion of 80% is a convention, so Tab. 9 also gives EtC and n.r. at 70% and 90%. WARC has the lowest EtC and n.r. at every threshold. WARC (scratch) reaches fewer mechanics as the threshold rises, with n.r. from 44.9% to 56.5%, and the other models change little, since most of their mechanics stay below 70%.

Every mechanic. Figs. 11 to 13 show the curve of every mechanic. Drops in the per-exposure values mark occurrences new to the recording, such as a new form or a new level (Appendix F), and accuracy recovers within a few occurrences. Some counters drop because their bar moves on an irregular schedule that three frames cannot always resolve.

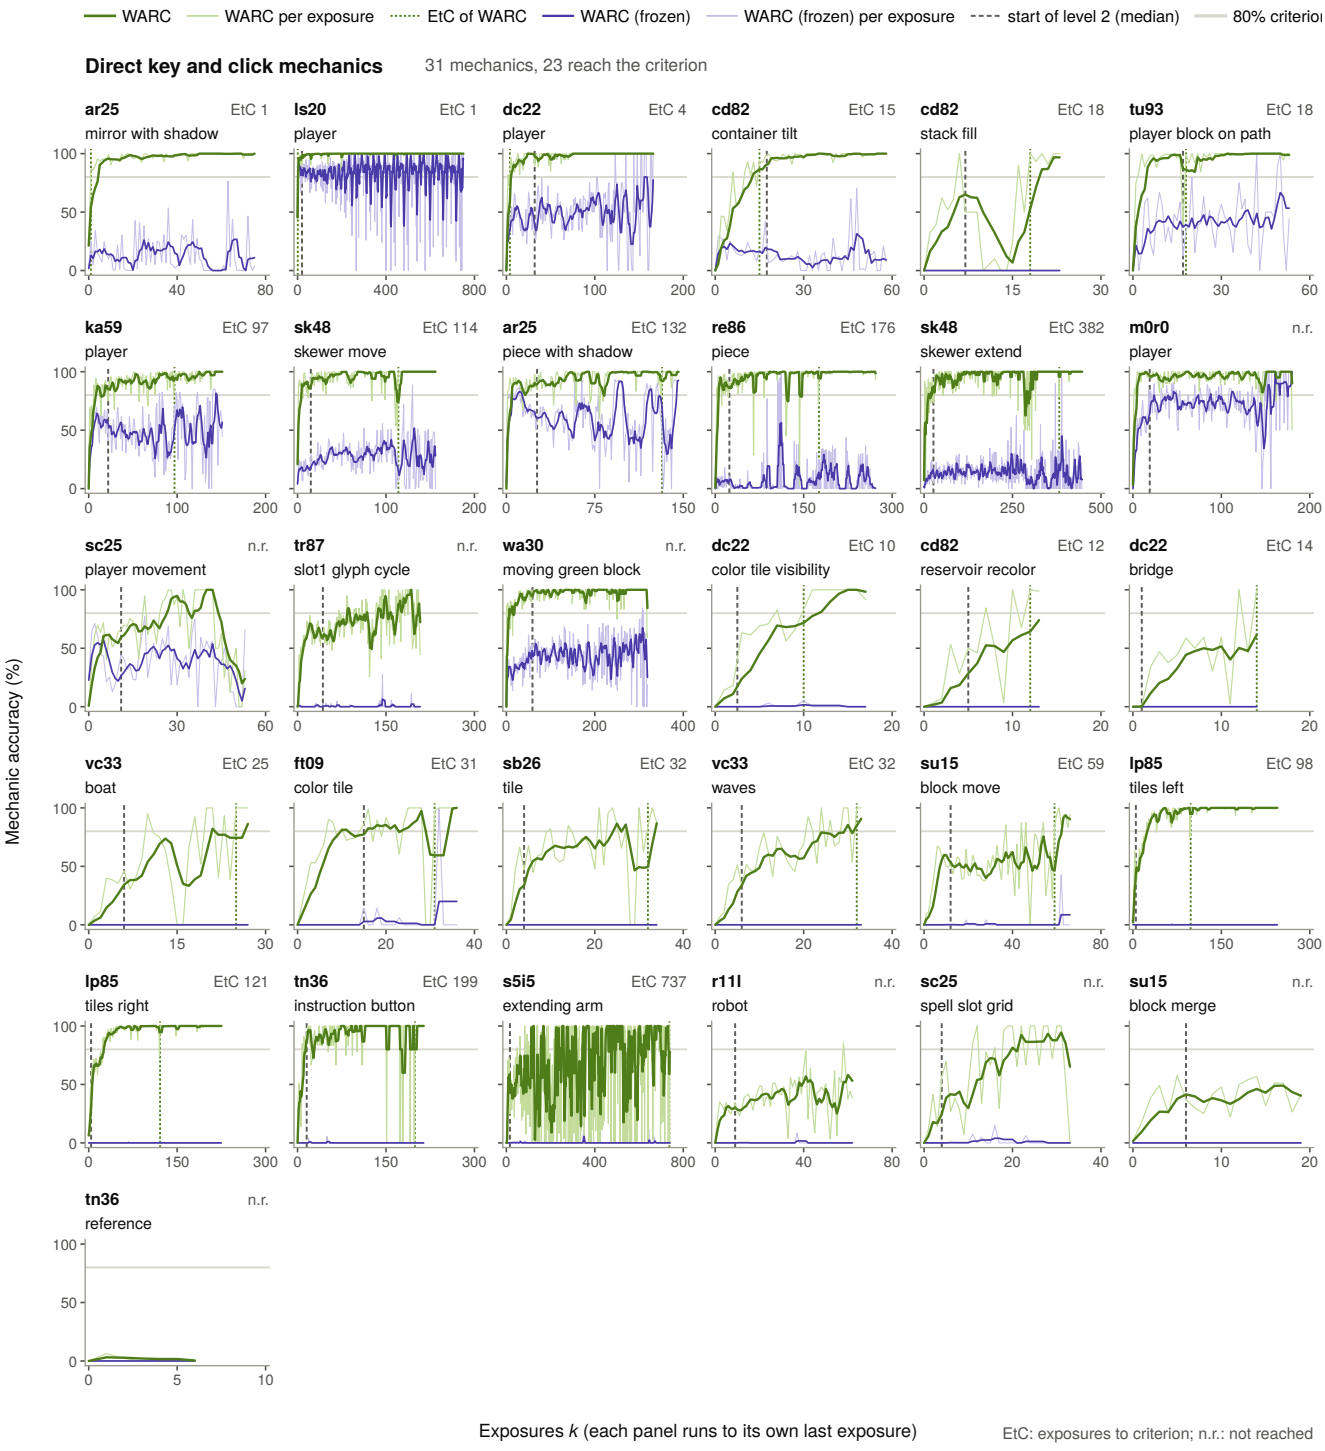

Figure 11. **The learning curve of every mechanic: direct key and direct click mechanics** (the other kinds in Figs. 12 and 13). Mechanic accuracy of WARC and WARC (frozen) against exposure k, averaged over each mechanic's occurrences with k earlier occurrences in all recordings, at every k, for WARC in test-time seed 0. Lines are the mean over the last five exposures, and the faint lines the per-exposure values, on which EtC is computed. Every test-time seed brings 56 mechanics to criterion, as in Tab. 1. Mechanics are grouped by kind (Tab. 5) and ordered by EtC within each kind, the exposures to criterion of WARC (n.r.: not reached). The dashed line marks the median exposure at the mechanic's first occurrence in level 2, over the recordings in which it fires in both levels, and the gray line the 80% criterion.

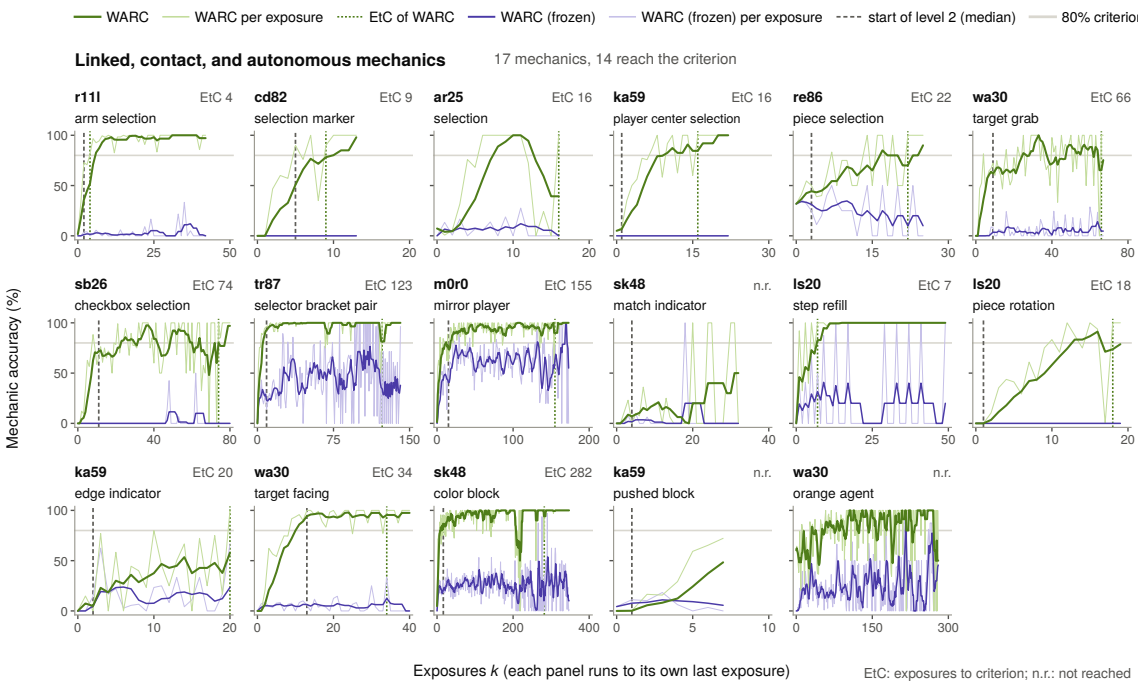

Figure 12. **The learning curve of every mechanic: linked, contact, and autonomous mechanics**, shown as in Fig. 11.

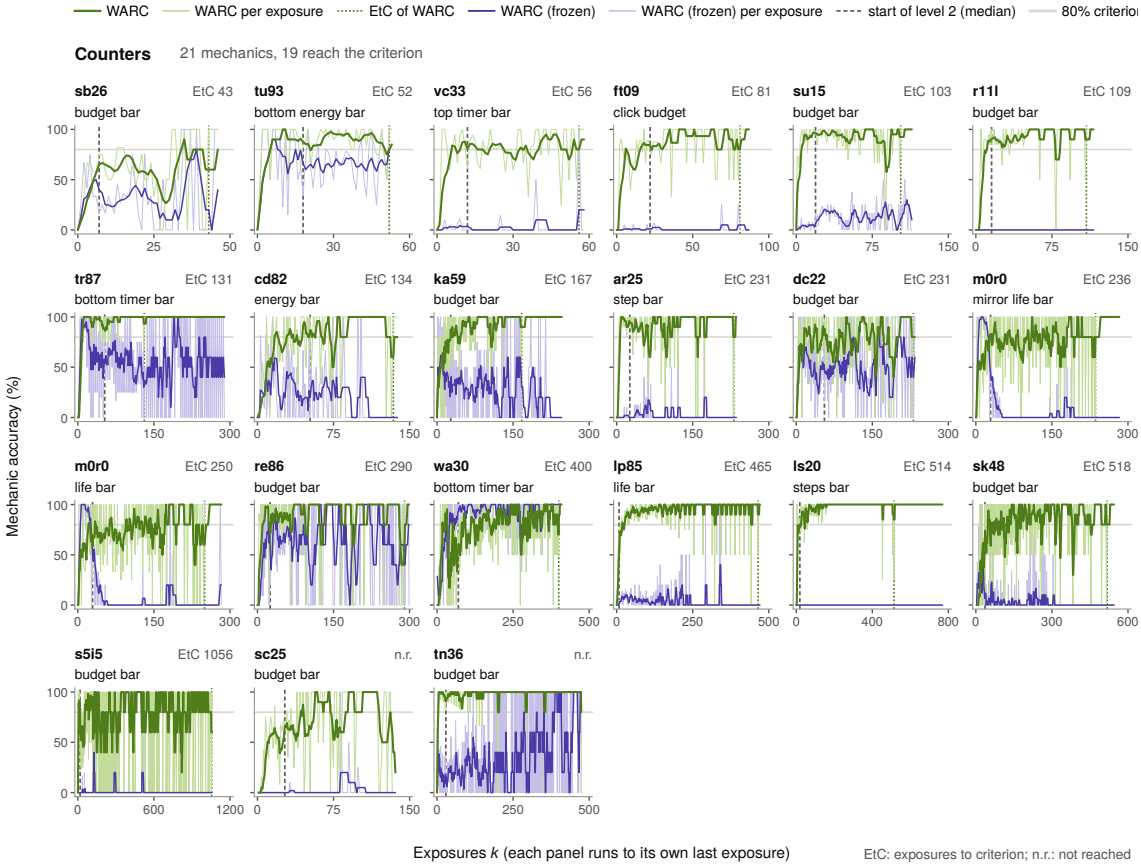

Figure 13. **The learning curve of every mechanic: counters**, shown as in Fig. 11.

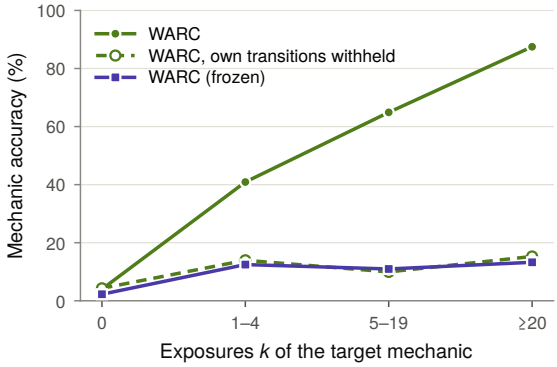

Figure 14. **Withheld transitions.** Mechanic accuracy on 32 target mechanics for WARC with their transitions in the buffer or withheld from it, and for WARC (frozen), which is not updated.

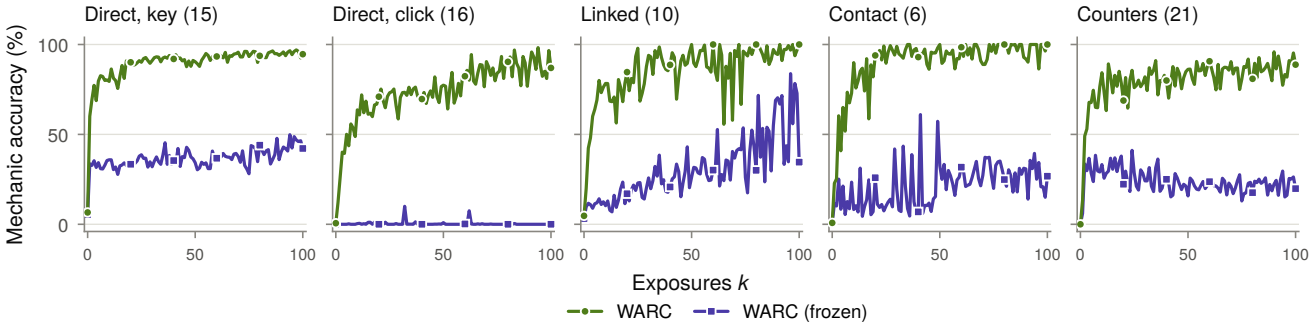

Figure 15. **Learning curves by kind of mechanic.** Mechanic accuracy against exposure k, at every exposure up to $k = 100$, with equal weight per mechanic within each kind of Tab. 5, without the single autonomous mechanic.

What pretraining alone predicts. Of the 17 mechanics that WARC (frozen) predicts at 30% or more, 15 are counters or direct key mechanics, and both kinds recur across games. For counters, it reaches a median of 32% on the 12 whose bar has the same form in a training game and 2% on the other nine.

Withheld transitions. In each recording, we choose a target mechanic from occurrence counts, before looking at any prediction. It must fire at least five times and in at most a quarter of the transitions, and the targets of a game rotate across its recordings. This gives 32 targets in 144 recordings. The transitions in which the target fires are predicted and scored but never enter the update buffer, while the updates and the input frames are unchanged. On the targets, WARC then stays close to WARC (frozen) in Fig. 14, 1.2 points $[-5.9, 5.3]$ below it and 42.9 points $[-48.3, -34.0]$ below WARC with the transitions used. The other mechanics of the same recordings lose 3.4 points, since the withheld transitions contain some of their occurrences.

Kinds of mechanics. Fig. 15 groups the mechanics by kind. Pretraining alone predicts part of the direct key mechanics and counters, but none of the direct click mechanics. Test-time training learns every kind and learns direct key mechanics fastest. Most of the mechanics that do not reach criterion are direct key or direct click mechanics.

F. Memorization or Generalization

Other recordings. At each occurrence of a mechanic in level 1 of the adapted recording A, the model predicts the other recordings before it updates on that occurrence. Each predicted occurrence uses the three frames and actions before it in its own recording. We predict up to 48 occurrences in level 1, taken from the other recordings in turn, and every occurrence in level 2. For level 2, adaptation stops at the end of level 1 of A. WARC (frozen) predicts the same occurrences. The same place of a key means the same start and end positions of the object it acts on. A mechanic is included when it fires in level 1 of A and in the predicted level of another recording. This holds for 14 direct key and 14 direct click mechanics in level 1,

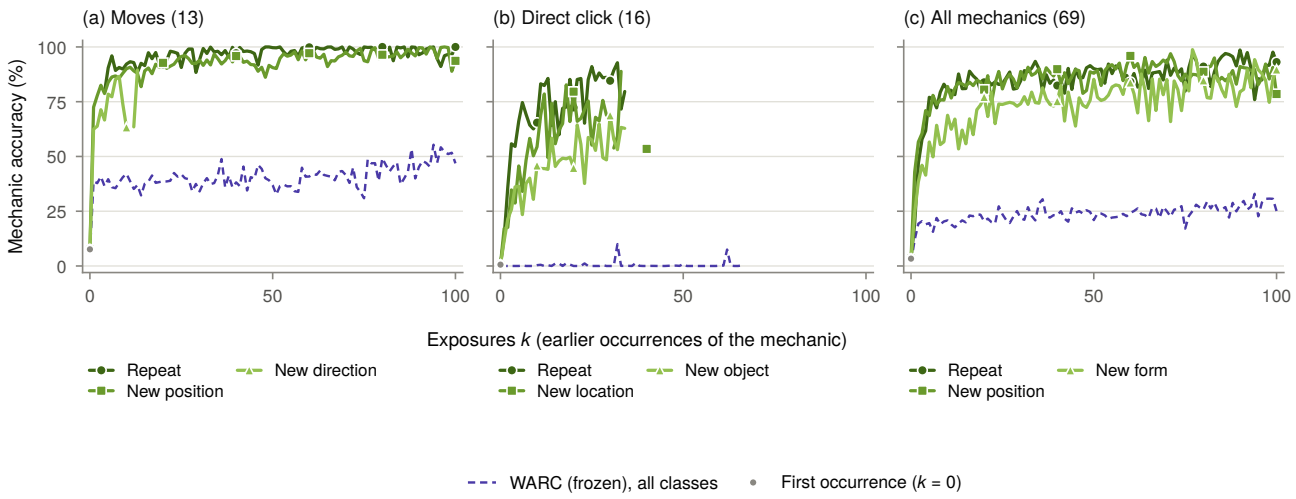

Figure 16. **Memorization or generalization.** Mechanic accuracy against exposure k, at every exposure up to $k = 100$, split by what is new about each occurrence within its recording: (a) the 13 moves, (b) the 16 direct click mechanics, and (c) all 69 mechanics by the pattern of changed cells. Each occurrence belongs to one class, so a class can be absent at small k. For example, a move at $k = 1$ is never a repeat. A point is drawn where at least five mechanics have an occurrence, and the line breaks elsewhere.

and for 14 and 15 in level 2. Means give equal weight per mechanic after averaging over the adapted recordings.

By class. Fig. 17 splits the unseen occurrences by what is new about them. At five to nineteen exposures, a move from a new position trails a seen move by 12.3 points [−16.9, −7.1] and a move in a new direction by 15.2 [−22.1, −8.4]. A known object clicked at a new location is predicted as well as a seen click (−1.9 [−8.6, 5.8]), whereas a new object trails by 15.2 [−28.9, −0.2]. WARC (frozen) dips at $k = 1$ for moves. Most seen occurrences there are the first move of their recording, which WARC (frozen) predicts at 11% because its input frames show no earlier move.

Classes. Within a recording, every scored occurrence with $k \geq 1$ is assigned one class by what is new about it, with

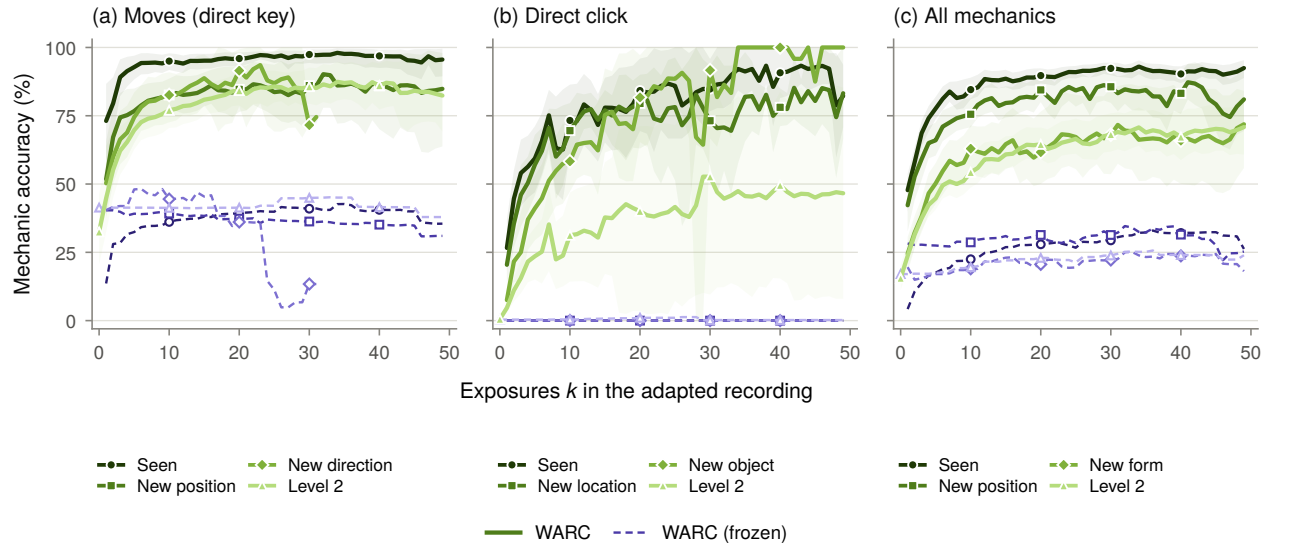

Figure 17. **Generalization to other recordings, by class.** Fig. 3 with the unseen occurrences split by what is new about them, as in Fig. 16, for (a) the 12 moves and (b) the 15 direct click mechanics with occurrences to predict (14 in level 1), and (c) all 64 mechanics with occurrences in other recordings (63 in level 1). Bands show 95% intervals for WARC, and dashed lines show WARC (frozen) on the same occurrences. Every exposure up to $k = 49$ is drawn. Beyond $k = 20$, only a few direct click mechanics have occurrences.

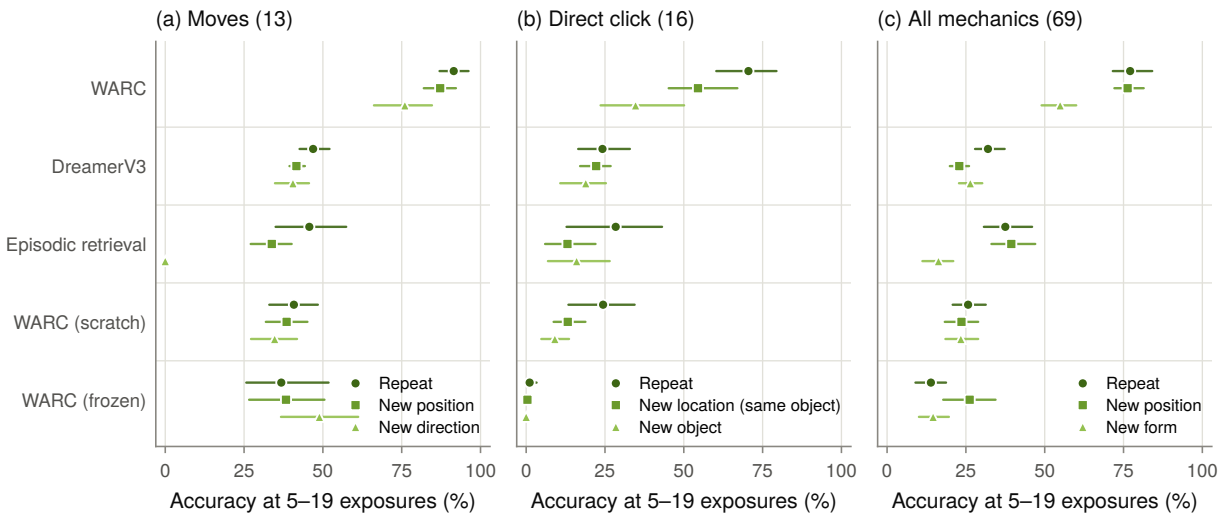

Figure 18. **The split for other world models.** Mechanic accuracy at five to nineteen exposures on repeats, new positions, and new directions, objects, or forms, for WARC, WARC (frozen), WARC (scratch), episodic retrieval, and DreamerV3. DreamerV3 uses its recipe tuned on the five games outside MechEval.

a definition for each kind of mechanic. The history of an occurrence is every earlier occurrence of the mechanic in the recording, scored or not. The first occurrence of a mechanic, and its first occurrence in level 2, are not classed. The classes are heuristic and the mechanics behind each class differ, so classes are compared within a mechanic before averaging.

Moves (direct key). For the 13 moves, what is new is the direction and the position of a move. A move is a *repeat* if the mechanic has fired earlier in the recording under the same key at the same place, that is, with the same start and end positions of the object. It is a *new position* if the mechanic has fired under the key only at other places, and a *new direction* if it has not fired under the key. A move in a new direction is predicted less well than a repeat but better than WARC (frozen) on the same occurrences (Fig. 18a), and a move from a new position nearly as well as a repeat (Tab. 10).

Direct click. For the 16 direct click mechanics, what is new is the clicked object and the cell. The clicked object is the sprite under the clicked cell before the click, identified by its name and position. A click on the background or on a sprite larger than 1,024 cells is identified by the cell itself. A click is a *repeat* if the same object was clicked at the same cell earlier in the recording, a *new location* if the same object was clicked at another cell, and a *new object* otherwise. A click on a known object at a new location is predicted less well than a repeat, and a click on a new object less well still.

All mechanics. For all 69 mechanics (Fig. 16c), we compare the pattern of changed cells, taken relative to their bounding box and including the colors. An occurrence is a repeat if the same cells changed in the same way before, a new position if the same pattern appeared at another place, and a new form otherwise. A new form can be a new effect or the same effect against different surroundings, so this split is coarser than those of the direct mechanics. A new position is predicted nearly as well as a repeat, and a new form clearly lower. Excluding the 21 counters gives the same result.

Other models. The classes depend only on the recording, so Fig. 18 applies them to every model. Episodic retrieval predicts a move from a new position less well than a repeat and never a move in a new direction. WARC (scratch) shows the same ordering as WARC. DreamerV3 is below WARC in every class.

Robustness. The ordering of the classes holds in each of the three test-time seeds, over coarser ranges of exposure, and when restricted to level 1.

Blocked moves. Blocked moves are steps at which a move fails because a wall, a boundary, or another object stops it. The object should stay in place, so a blocked move is scored on the unchanged cells of its region, as the fraction that the model also predicts unchanged, and is classed as moves are (Fig. 19). The gap to WARC (frozen) at the first blocked move is -31.5 points $[-40.8, -23.9]$. After the object has been blocked before, WARC predicts blocked moves as well as WARC (frozen) from five exposures on, with intervals that include zero. Over the first four, a blocked move at a new position trails by 9.5

Group	Class	Seed 0	Seed 1	Seed 2	Mean	vs. repeat
Moves	repeat	91.5	91.9	91.8	91.7	–
	new position	87.2	86.8	86.6	86.9	−4.9
	new direction	76.0	77.3	78.5	77.3	−14.5
Direct click	repeat	70.5	73.0	74.3	72.6	–
	new location	54.5	54.4	54.9	54.6	−18.0
	new object	34.7	36.6	35.6	35.6	−37.0
All	repeat	77.1	79.0	79.5	78.5	–
	new position	76.3	74.0	75.2	75.2	−3.4
	new form	54.9	55.8	54.8	55.2	−23.4

Table 10. The split of Fig. 16 over three test-time seeds: mechanic accuracy (%) at five to nineteen exposures, with equal weight per mechanic, and the mean difference from repeats in the same seed. The differences for new directions, locations, objects, and forms exclude zero in every seed, and the intervals for new positions reach zero in every seed.

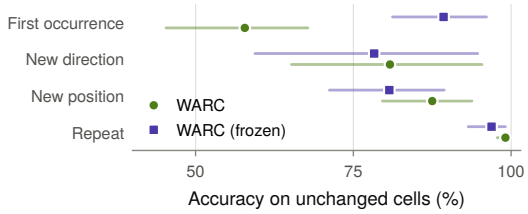

Figure 19. **Blocked moves.** The split of moves in Fig. 16a for blocked moves (19 types in 12 games), scored on the region's unchanged cells (predicting no change scores 100), at the first blocked move and at five to nineteen exposures, with 95% intervals.

Choice	Change	All	k=1–4
Pretraining	none	−34.1 [−37.4, −31.3]	−28.2 [−32.2, −24.2]
Update	every 5 transitions	−1.7 [−2.6, −0.9]	−10.5 [−14.2, −7.0]
Updated part	LoRA, rank 64	−1.4 [−2.6, +0.2]	+3.8 [+2.1, +5.6]
	last two blocks	−4.6 [−5.6, −3.7]	−6.7 [−9.3, −4.1]
	modulation layers	−9.4 [−11.6, −7.4]	−7.2 [−10.0, −4.5]
Augmentation	none	−1.1 [−1.9, −0.2]	−1.8 [−3.6, +0.1]
	+ flips and colors	+0.3 [−0.9, +1.3]	+1.4 [−0.3, +3.1]
	colors, not rotations	−0.1 [−1.5, +1.3]	+0.7 [−1.3, +2.9]

Table 11. Paired 95% bootstrap intervals for the ablations: the change from WARC in mechanic accuracy over all exposures and the first four, averaged over three test-time seeds, resampling games and then recordings 2,000 times with draws shared by the seeds and the rows.

points [−17.4, −1.1].

More examples. Fig. 20 shows six more mechanics in the format of Fig. 4, without the limit $k \leq 5$. WARC does not always predict a move in a new direction. In cd82 it leaves the container tilted at both new directions, where WARC (frozen) and DreamerV3 predict more of the changed cells, and in sc25 it misses the first move to the right, which WARC (frozen) predicts. WARC (frozen) and DreamerV3 miss every click on a new object shown, whereas WARC predicts 31 to 100% of their changed cells.

G. Ablation Details

All ablations change one choice of WARC.

Intervals. Tab. 11 gives paired intervals for Tab. 2. Removing pretraining, updating less often, and restricting the update all lower accuracy, with intervals that exclude zero. LoRA is the exception and is ahead over the first four exposures.

Update cadence. Updating every five valid transitions takes 100 gradient steps at a time on minibatches from one draw of the buffer, so the steps per transition and the minibatches match WARC.

Updated parameters. The rows that restrict the update keep the recipe of WARC and change only which parameters are

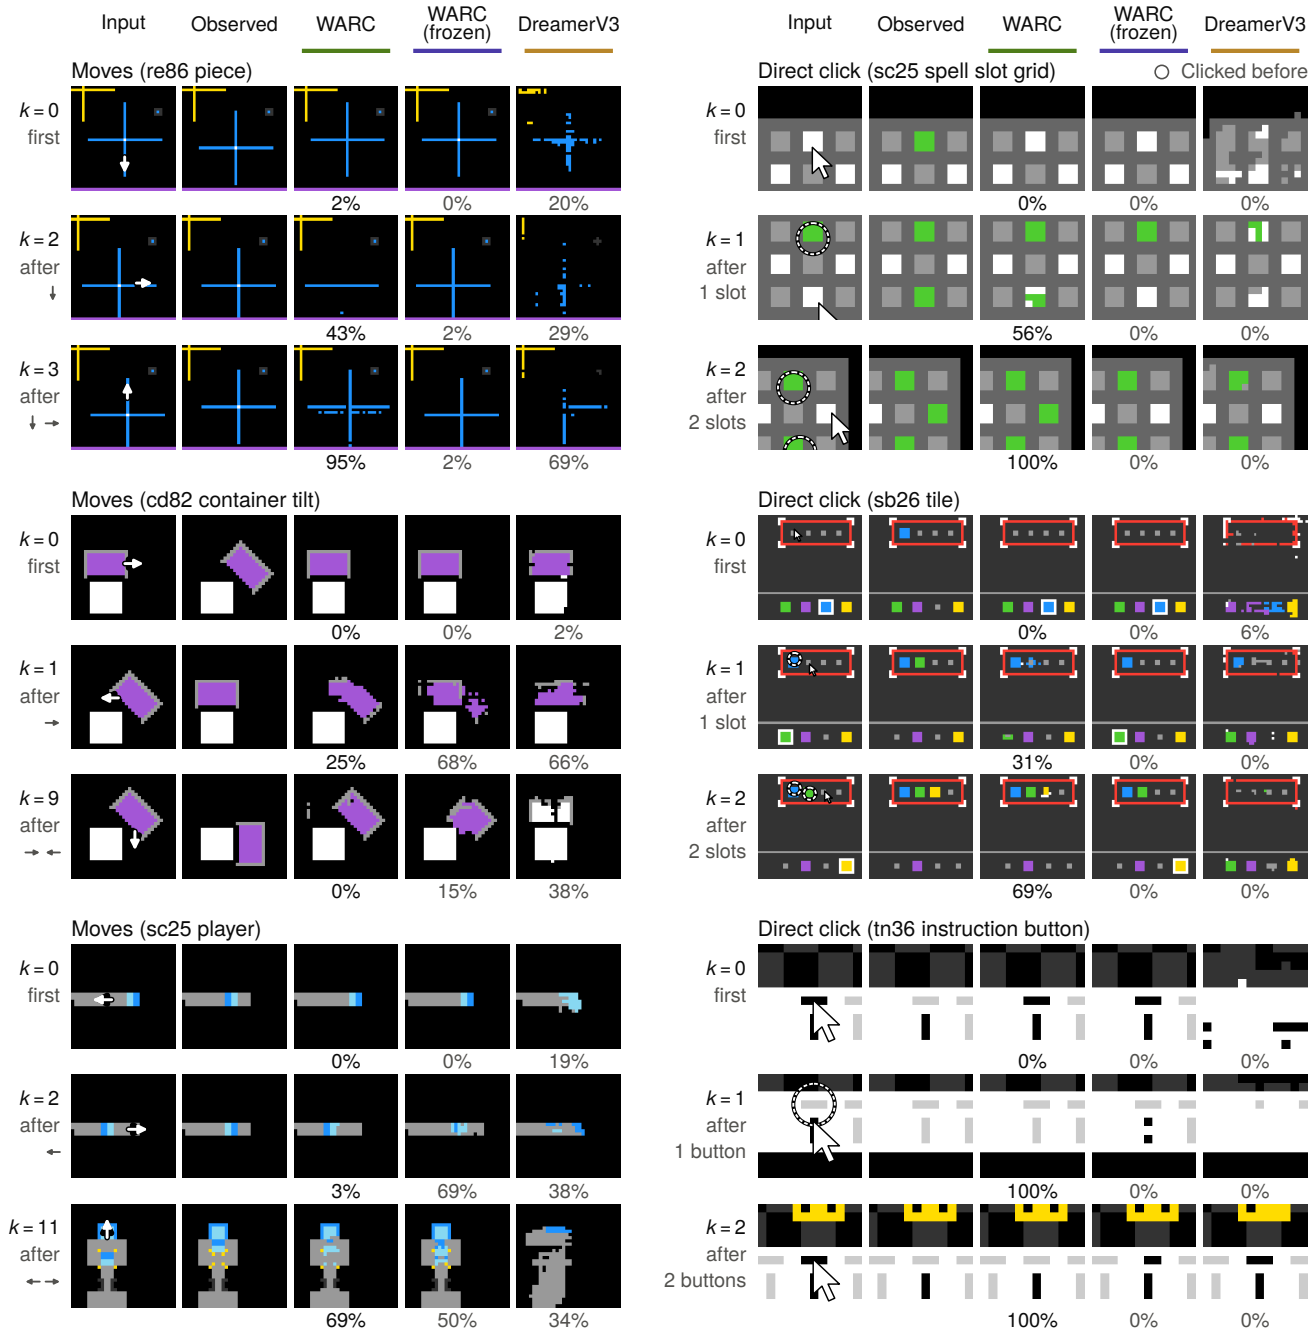

Figure 20. **More examples of generalization to new occurrences.** Fig. 4 for three moves (left) and three direct click mechanics (right). The last row of sc25 player is in level 2.

updated. The modulation layers are the adaptive layer norms of every block, and the last two blocks include the final norm and the head. LoRA [38] adds rank-64 adapters, 3.3M parameters, to the attention and MLP linear layers of every block. The learning rate of LoRA (10^{-3}) was chosen on the five games outside MechEval by the rule of Appendix C. Ranks 8 and 32 also stay within the interval of the full update over all exposures in test-time seed 0.

Augmentation and gradient steps. Without augmentation, accuracy over all exposures is 1.1 points lower, and the intervals of the other two augmentation rows in Tab. 11 include zero. The rotations matter most for moves in a new direction. At five

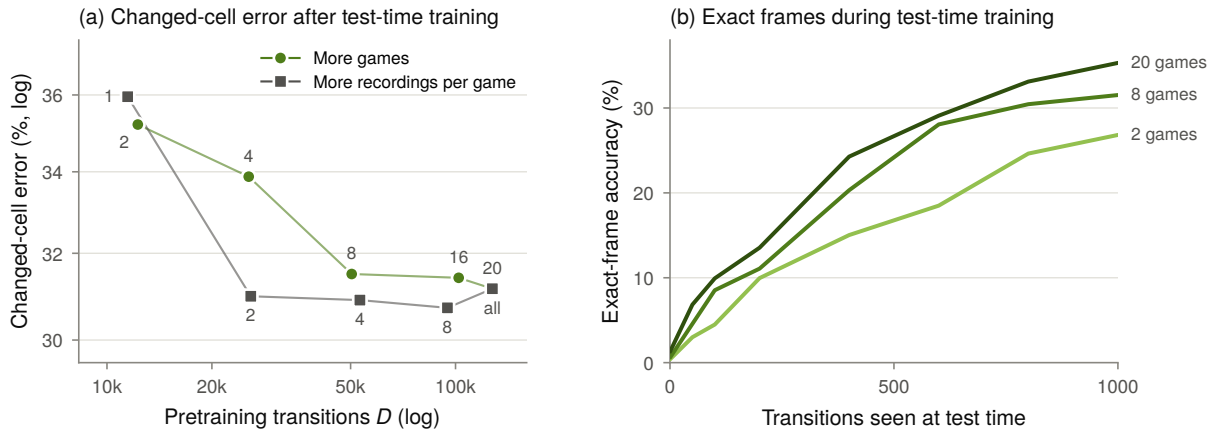

Figure 21. **Pretraining scale, continued.** (a) Changed-cell error after test-time training against pretraining transitions, for the runs of Fig. 5. Both series flatten. (b) Exact-frame accuracy during test-time training for models pretrained on 2, 8, and 20 games, evaluated at 0, 50, 100, and every 200 transitions. Means over five held-out games, one pretraining seed.

to nineteen exposures, averaged over the three test-time seeds as in Tab. 10, a new direction trails a repeat by 28.0 points without augmentation and by 14.5 with it. For new positions and for clicks on new objects, the gap changes by less than three points. Without augmentation, WARC still predicts the first move to the left in Fig. 4, after one move up. Across the augmentation rows, the false-edit rate stays at 0.11 to 0.12%. We keep rotations, chosen on the five games outside MechEval (Appendix C). Doubling the gradient steps per transition does not raise accuracy, and halving them lowers it by 1.5 points.

H. Pretraining Scale

Setup. This study uses a variant of WARC with a seven-block encoder that does not see the action, a three-block dynamics stack that does, and a token per game. Nine models are pretrained on nested subsets of the 20 training games of one five-fold split, and each is adapted by test-time training on 1,000 transitions of human play in the five held-out games. One axis keeps all 20 games and adds recordings per game (1, 2, 4, 8, all), and the other keeps all recordings and adds games (2, 4, 8, 16, 20), so the two meet at the full set. We evaluate on up to 1,024 other held-out transitions by exact-frame accuracy and by changed-cell accuracy, the accuracy on all cells that change.

Results. Exact-frame accuracy after test-time training rises from 26.8% with 2 games to 35.3% with 20 (Fig. 5). Changed-cell accuracy rises from 64.8% to 68.5% with 8 games and then flattens (Fig. 21a). At similar data, 20 games with two recordings each reach 34.0% exact-frame accuracy, against 29.2% for 4 games with all their recordings. Before test-time training, no model predicts more than 3% of frames exactly, so the gain from broader pretraining appears only after test-time training (Fig. 21b). These results come from one pretraining seed and five held-out games, and the test-time training of this study used a recipe developed while looking at these games.

I. Closed-Loop Details

I.1. Agent

Explorer. The explorer is a simplified version of a published graph explorer [73]. It keeps a graph of the states reached in the current level and the actions tried in each. A state is keyed by a hash of the frame, with the status bars masked as in the first frame of the level, and by the set of available actions. The actions are the available keys, one click at the center of each connected component of the frame, and a fixed grid of 64 clicks. They are tried in priority order, first the keys and the components of at most 1,024 cells, then the grid, then the larger components. In a state with untried actions, the explorer chooses among the untried actions of the highest open priority group, uniformly when it has no world model. Otherwise it moves along known edges to the nearest untried action, by priority and then by distance. If none is reachable, it retries the least-tried actions of the current state. The published explorer adds hand-set rules on the frame. It detects status bars by their shape and position, sorts the clicks into five groups by color and size, with segments of non-gray colors and medium size

first, and opens the groups one at a time. Our explorer instead orders the clicks by size and masks only thin bars along the top and bottom edges. It clicks at the center of each component and on a fixed grid rather than at a random cell of each segment, adds the available actions to the state key, and never rebuilds its graph within a level.

Ranking actions with the world model. For a frame o, let $\phi(o)$ be the color histograms of its 4×4-cell blocks, with the status bars and the action-independent cells (defined below) masked, and let $d(\phi, \phi') = \frac{1}{2}\|\phi - \phi'\|_1$. With $\mathcal{O}_\ell$ the distinct frames of the current level and r the median of $d\big(\phi(\hat{o}_{t+1}), \phi(o_{t+1})\big)$ over the last 20 transitions, the novelty of a predicted frame $\hat{o}$ at depth j is

$$\nu_j(\hat{o}) = \max\Big(0,\ \min_{o\in\mathcal{O}_\ell} d\big(\phi(\hat{o}), \phi(o)\big) - j\,r\Big). \tag{6}$$

Subtracting r keeps the prediction errors of a still inaccurate model from counting as novelty. WARC predicts the next frame $\hat{o}_a$ of every untried action a of the open group, or of a random sample of 256 when there are more. Candidates with the same predicted frame are merged into one. The 32 actions with the highest ν_1 are extended by a second step over a itself and three other candidates b drawn at random, and are scored by

$$s(a) = \max\Big(\nu_1(\hat{o}_a),\ \tfrac{1}{2}\max_b \nu_2(\hat{o}_{a,b})\Big), \tag{7}$$

while every other action keeps $s(a) = \nu_1(\hat{o}_a)$. The explorer takes the action with the highest score, or a uniformly random candidate with probability $\epsilon = 0.2$. Until five prediction errors have been recorded, or when every score is zero, the choice is uniform. $\mathcal{O}_\ell$ and the errors are reset at each new level, and WARC updates with the recipe of Sec. 5.2 after every valid transition (Algorithm 1). WARC (frozen) runs Algorithm 1 without UPDATE, and Graph runs it without a world model, choosing uniformly among U.

Action-independent cells. A cell is action-independent if every candidate action leads to the same predicted color for it and this color differs from the current frame, as for a timer or a step counter. The mask extends to the connected components that contain such cells, excluding components of more than 1,024 cells and cells on which the predictions disagree. It is recomputed at every step, skipped when all actions lead to the same predicted frame, and applies only to the novelty score, not to the state keys.

I.2. Protocol and Measures

Protocol. Each run plays one of the 25 public games with environment seed 0 and agent seeds 0 to 9, and every measure is read from its first 1,000 actions, including resets. Random, which picks a uniformly random available action, and the from-scratch CNN play 4,000 actions under a limit of 600 s per run. Each of these runs uses up its actions within 246 s, so the time limit never binds. Graph plays 1,000 actions on a CPU with no time limit, in 98 s on average. Runs with a world model also play 1,000 actions with no time limit. Every run with these seeds is included.

RHAE. The official score, Relative Human Action Efficiency [3], rates each cleared level by the squared ratio of the human players' actions to the agent's, weights later levels more heavily, and normalizes over all levels of a game on a 0–100 scale. An agent that clears only the first levels, using many more actions than humans, therefore scores below 1.

I.3. Results

Results per game. In six games (ar25, bp35, ft09, m0r0, s5i5, tu93), WARC raises the level-1 clears over WARC (frozen) by 40 to 100 points. In four (cn04, ls20, sk48, sp80) it lowers them by 10 to 70 points, and in the other 15 the two agree. Paired over games and seeds, the efficiency score rises by 0.31 [0.12, 0.52] over the explorer alone, and the level-1 clears by 12 points [−3, 27] over WARC (frozen).

Explorer ablations. The explorer was fixed before the runs of Tab. 12, which remove one component at a time on seeds 0 to 4, paired with the first row. The random sample of 32 actions also runs without the mask, so it is compared with the row without the mask. We read the table mainly by L1@1000, which 42 to 63% of the runs reach. No variant differs from the full explorer beyond its paired interval in level-1 clears or in the efficiency score. With 25 games, the level-1 intervals span 14 to 21 points, so the table rules out only large effects. Ranking up to 256 candidates rather than a random 32, and planning two steps rather than one, each add a few level-1 clears (-4.8 and -3.2 points without them), within the intervals. Without the mask, level-1 clears are unchanged, and without the random action they drop by 2.4 points. We keep both for games that differ from the public ones. The random action keeps the explorer moving when the model's predictions are wrong, and the

Algorithm 1 WARC-Agent (Graph + WARC) in one game.

Require: world model f_θ pretrained on other games, reading the last h frames and actions of H; $\epsilon = 0.2$; at most 256 candidates, 32 of them extended by 4 second steps

$o \leftarrow$ RESET; $H \leftarrow (o)$; $\mathcal{B} \leftarrow \emptyset$; new graph, $\mathcal{O}_\ell \leftarrow \{o\}$, no errors
for $t = 1, \dots, 1000$ **do** ▷ a reset after a game over counts as an action
  $U \leftarrow$ untried actions of the open group in state KEY(o), else the next step toward one, else the least-tried
  $C \leftarrow U$, or 256 actions of U drawn at random
  $\hat{o}_a \leftarrow f_\theta(H, a)$ for every $a \in C$ ▷ one batch
  recompute the mask of action-independent cells from $\{\hat{o}_a\}$
  keep one action of C per distinct $\hat{o}_a$
  if fewer than five errors are recorded **then**
    $s(a) \leftarrow 0$ for every $a \in C$
  else
    $s(a) \leftarrow \nu_1(\hat{o}_a)$ for every $a \in C$ ▷ Eq. (6)
    for the 32 actions a with the highest $s(a)$ **do**
      $C_2 \leftarrow \{a\} \cup \{3 \text{ other actions of } C \text{ drawn at random}\}$
      $s(a) \leftarrow \max\left(s(a),\ \frac{1}{2} \max_{b \in C_2} \nu_2(f_\theta(H + (a, \hat{o}_a), b))\right)$ ▷ Eq. (7)
    end for
  end if
  $a \leftarrow$ uniform over C with probability ϵ or if $\max_a s(a) = 0$, else $\arg\max_a s(a)$
  $o' \leftarrow$ STEP(a); add (KEY$(o), a,$ KEY$(o'))$ to the graph and o' to $\mathcal{O}_\ell$
  if the transition is valid **then**
    record the error $d\big(\phi(\hat{o}_a), \phi(o')\big)$ ▷ the last 20 set r
    $\mathcal{B} \leftarrow \mathcal{B} \cup \{(H, a, o')\}$; UPDATE$(\mathcal{B})$; $H \leftarrow H + (a, o')$
  else
    $H \leftarrow (o')$; at a new level, a new graph, $\mathcal{O}_\ell \leftarrow \{o'\}$, no errors
  end if
  $o \leftarrow o'$
end for
return the levels cleared

function UPDATE$(\mathcal{B})$ ▷ Sec. 5.2
  $\mathcal{D} \leftarrow$ the 10 newest and 20 random older transitions of $\mathcal{B}$, each as the original and its three rotations, each rotation repeated ten times
  for 20 steps **do**
    take a gradient step on Eq. (5) for a minibatch of 8 from $\mathcal{D}$
  end for
end function

Explorer	L1@1000↑	L2@1000↑	RHAE@1000↑
Graph + WARC (WARC-Agent)	47.2	8.8	0.456
without the mask of action-independent cells	47.2	9.6	0.513
a random sample of 32 actions, no mask	42.4	10.4	0.346
one planning step	44.0	8.0	0.473
no random action	44.8	10.4	0.520
Graph + oracle world model	63.2	27.2	1.326

Table 12. The explorer with one component removed, and with an oracle world model, on the 25 public games over seeds 0 to 4, with the measures of Tab. 3. Every row but the last uses WARC with test-time training, and the last row is the oracle world model of Tab. 3, on these seeds.

mask keeps cells that change under every action, such as a timer, from counting as novelty. We do not choose among the variants by their scores, since that would tune the explorer on the evaluated games.

Oracle world model. The oracle row of Tab. 3 runs the same explorer, actions, and random choices with a perfect world model. A copy of the game program follows the run, and each candidate action is tried on a further copy, so each predicted

frame is the true next frame. An agent cannot copy its environment, so the oracle row is for analysis only. Paired over games and seeds, the oracle clears level 1 in 16 points more runs than WARC [7, 26], and it clears level 2 in 27% of runs against 8%. Of the ten games that WARC never clears, the oracle clears five. Five games are never cleared even with the oracle, so there the explorer is the limit. In 20% of the decisions at which some action would complete the level, the explorer took another, since it ranks actions by novelty alone and acts at random with probability 0.2.

J. Additional Related Work

Video world models. Trained on large video and interaction datasets, video world models simulate real-world interaction [95], physical scenes [63], and navigation [8]. In games, GameGAN [47], Genie [11], GameNGen [86], DIAMOND [2], Muse [44], and Matrix-Game [36] simulate the games they were trained on. Latent world models such as DINO-WM [99] and V-JEPA 2 [6] predict in the feature space of a pretrained encoder and plan with them. These models are finetuned for a new domain or applied to it without updates. We instead study how a world model adapts online to an unseen game.

Test-time training and continual learning. Test-time training updates a model on each test input with a self-supervised loss [82], and dynamic evaluation adapts a sequence model to the sequence it is predicting [49]. TENT updates only the normalization parameters on unlabeled test data [88]. Test-time training layers make the hidden state of a sequence model the weights of an inner model, updated by a gradient step on each token [83]. LaCT updates them on large chunks [97]. Such layers also generate minute-long videos [18]. On a video stream, a model can also be trained online on each frame [91]. Over a stream of updates, continual learning limits forgetting by replaying earlier data [72] or by protecting important weights [48]. Continual test-time adaptation restores part of the source weights [90]. Plasticity can also be lost as training continues [21]. WARC replays earlier transitions in every update (Sec. 5.2).

Exploration. Curiosity-driven agents reward transitions that their own model predicts poorly [12, 65], and Plan2Explore plans toward states where an ensemble of world models disagrees [77]. Go-Explore keeps an archive of visited states and returns to promising ones before exploring from them [23]. Our agent combines a graph of visited states with a ranking of untried actions by the novelty of their predicted outcomes (Appendix I).

ARC puzzles. Vision transformers such as VARC [39], RoSeViT [57], and TraceViT [56] cast a puzzle as image-to-image translation. Loop-OWM instead learns a composable, object-centric world model of a puzzle's rule [28].

Agents on ARC-AGI-3. Besides the agents cited in Sec. 2, published agents have a frontier language model write and repair a programmatic world model [70], as in OPINE-World [17] and Tycho [53]. Others run a language model within a harness with memory, feedback, or a curriculum [5, 13, 25, 26, 46, 75]. Agents without a language model explore with a graph over states [73] or train a predictor from scratch on each game, like StochasticGoose [42, 80]. We compare with both in closed-loop play (Sec. 7).

Rule changes and rule discovery. Boult et al. [10] formalize novelty, and NovGrid [7], NovelGridworlds [30], NovelGym [31], and NovPhy [27] inject controlled rule changes into grid worlds and physics puzzles. DiG-bench makes rule discovery the task itself [9]. MechEval instead follows each mechanic of an existing game across its occurrences.