arXiv:2611.11995v1 [cs.AI] 9 Nov 2026

Infrathon: Benchmarking Language Agents for Real-World, Long-Horizon AI Infra Optimization

Zhengyu Hu[1], Zheyuan Xiao[1], Yike Wang[1,2], Hamish Ivison[1], Haoyu Wang[3], Yige Yuan[1], Rulin Shao[1], Yuetai Li[1], Niqi Lyu[4], Yicheng Ding[4*], Shu Yang[5], Zhichen Zeng[1,6], Banghua Zhu[1,6], Jiaxin Pei[5,7], Kaize Ding[8], Pradeep Dasigi[2], Dan Roth[3], Hannaneh Hajishirzi[1], Noah A. Smith[1,2], Teng Xiao[1,2*†], Radha Poovendran[1,*]

[1]University of Washington, [2]Allen Institute for AI, [3]University of Pennsylvania, [4]Pyromind Dynamics, [5]Stanford University, [6]RadixArk, [7]The University of Texas at Austin, [8]Northwestern University

Abstract. In this work, we study whether frontier agent can continually improve real-world AI infrastructure over long horizons. We introduce `Infrathon`, a benchmark of diverse, long-horizon AI infrastructure optimization tasks across different levels. We evaluate six models across 150 runs with 24-hour budgets, totaling over 10,000 GPU-hours. Our results show that model rankings vary with the evaluation horizon, while 135 of the 150 runs achieve further improvements during the second half of their budgets. These findings indicate that shorter evaluations can overlook both optimization gains and changes in relative model performance. We also find agents can discover solutions that outperform human-designed baselines: Claude Opus 5.5 increases the mean KV-cache compression ratio from 6.82× with the human-designed Quant-VideoGen codec to 19.86×. Analysis of agent trajectories identifies task violations, unreliable testing, and premature completion as barriers to sustained improvement. Together, Our `Infrathon` framework provide resources for constructing long-horizon tasks, evaluating frontier agents, and discovering good infrastructure solutions.

Code: https://github.com/InfrathonX/Infrathon
Docker Environment: https://hub.docker.com/r/infrathon/infrathon
HuggingFace: https://huggingface.co/datasets/Infrathon/Infrathon
Website: https://InfrathonX.github.io

.

1 Introduction

AI infrastructure, the hardware and software used to train and serve AI models, affects their computational efficiency. Advances such as FlashAttention [1], SGLang [2], and Ansor [3] demonstrate the benefits of optimizing attention kernels, model serving, and tensor programs. However, developing such optimizations requires repeated implementation, profiling, and validation while preserving correctness under numerical and hardware constraints. Changes to one component can expose new bottlenecks or introduce regressions elsewhere, requiring engineers to use performance feedback to guide further improvements [4, 5]. AI agents could automate this process, reducing engineering effort and improving model training and serving.

Existing benchmarks provide only a partial evaluation of this capability, with limitations in task coverage, execution budgets, or publicly available GPU tasks (Table 1). KernelBench [6] focuses on operators and modules, while InferenceBench [7] evaluates complete inference services. Although Φ-Bench [8] spans all three levels, most tasks are limited to 1.5-hour runs. EdgeBench [9] uses 12-hour budgets in its GPU tasks, while `Infrathon` provides 12–24-hour budgets for GPU-based AI infrastructure optimization. These limitations leave open whether frontier agents can repeatedly identify bottlenecks, revise solutions, and validate improvements

[†]Project leader, [*]Corresponding authors. Contact emails: zhu7@uw.edu and tengxiao01@gmail.com

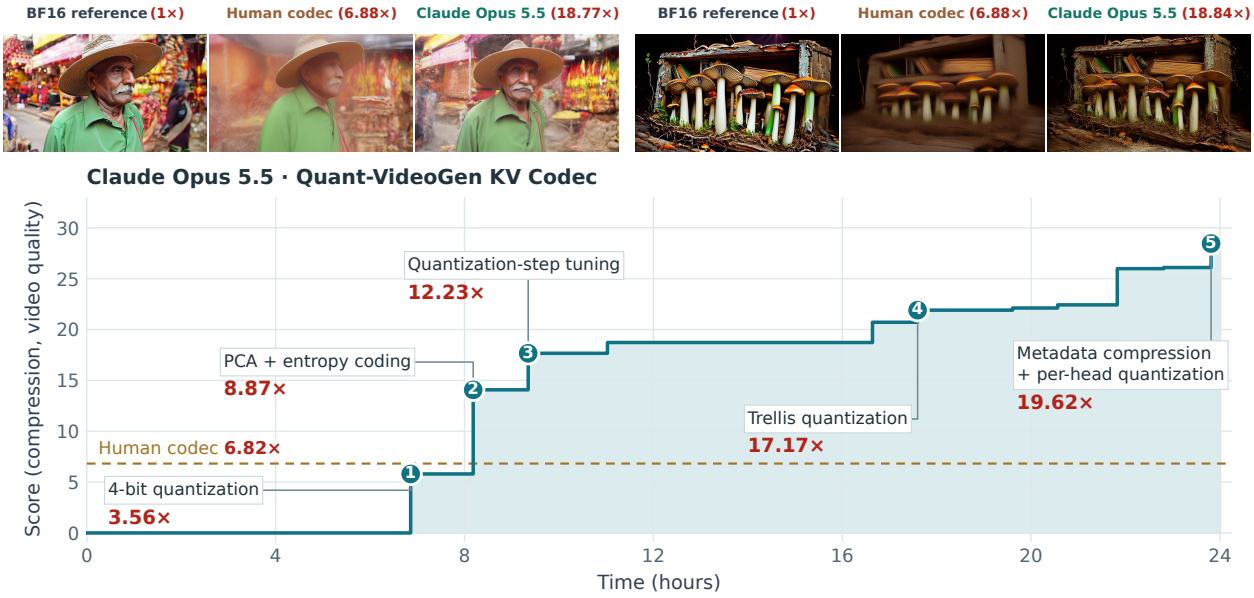

Figure 1: **An optimization example with Infrathon.** On Quant-VideoGen KV Codec, Claude Opus 5.5 develops a codec that reduces memory use during video generation, with higher KV-cache compression ratios and better video quality than the human codec. Top: two video examples; compression ratios are red. Bottom: Claude's scores over time combine compression ratio and video quality; dashed line: Human Solution. Numbered points mark five selected submissions; boxes summarize their code changes (Appendix E.3).

across diverse AI infrastructure tasks over extended runs. This raises the following research question:

Can Frontier Agents Solve Diverse, Long-Horizon AI Infrastructure Tasks?

To answer this question, we introduce `Infrathon`, a benchmark of *25 real-world, long-horizon AI infrastructure optimization tasks* across three levels. We construct these tasks with human experts using `Infrathon-build`, a tool that helps humans efficiently build long-horizon tasks. They draw on open-source projects (e.g., DeepGEMM [10], FlashMLA [11], and Liger-Kernel [12]) and research projects (e.g., FlashAttention [13], SageAttention [14], and Flash-KMeans [15]). *11 operator-level tasks* optimize computational primitives with *12-hour* budgets; *5 module-level tasks* optimize interacting components with *24-hour* budgets; and *9 system-level tasks* optimize end-to-end training or serving workflows with *24-hour* budgets. `Infrathon` is the first benchmark for long-horizon AI infrastructure optimization across three levels, with 12–24-hour budgets and metrics like TFLOP/s without fixed upper bounds. The framework couples *Task Definition* (Section 3.1) with *Task Execution* (Section 3.2): agents develop solutions and run public tests in the *Work* environment, submit solutions to the *Controller* for code review, and receive private test results from the *Judge*.

We then evaluate frontier models with a total budget of *over 10,000 GPU-hours*. Across 150 runs covering six models and all 25 tasks (Section 4), Claude Opus 5.5 leads on 17 tasks and GPT-6 Astra on six. Rankings change with the time budget: Astra leads overall at 25% of the budget, while Claude leads at 50%. Of the 150 runs, 135 improved their best score in the second half. Infrathon also supports knowledge discovery in AI infrastructure (Section 4.2). On Quant-VideoGen KV Codec, Claude Opus 5.5 increases mean KV-cache compression from 6.75× to 20.89× on WorldScore and from 6.88× to 18.83× on MovieGen, compared with the Quant-VideoGen codec (Human). Both evaluations show better video quality. With `Infrathon-build` and `Infrathon`, users can turn their own infrastructure problems into tasks, compare agent solutions with existing implementations, and iterate using test feedback (Figure 1). Our trajectory analysis shows that agents sometimes violate task instructions, test the wrong code, or declare completion before finding further improvements (Section 4.4). For example, the Controller rejects an early Quant-VideoGen KV Codec [16]

submission in which Claude Opus 5.5 submits human codec weights downloaded from the internet as its own.

Together, `Infrathon` and `Infrathon-build` help humans build long-horizon tasks, evaluate LLMs, and discover new solutions and knowledge. We hope they help develop agents for efficient training and inference.

Table 1: **Coverage of related benchmarks.** *Infra only*: exclusively AI infrastructure; *Long horizon*: support for continuous execution of every task for at least 12 hours; *GPU tasks*: at least one GPU task in the open-source release. *Operator*, *Module*, and *System* denote computational primitives, components, and complete workflows. ✔ and ✖ indicate whether coverage is documented (Appendices A.2 and D).

Benchmark	Infra only	Long horizon	GPU tasks	Operator level	Module level	System level
KernelBench [6]	✔	✖	✔	✔	✔	✖
KernelBenchX [17]	✔	✖	✔	✔	✖	✖
FlashInfer-Bench [18]	✔	✖	✔	✔	✖	✖
RE-Bench [19]	✖	✖	✔	✔	✖	✔
InferenceBench [7]	✔	✖	✔	✖	✖	✔
Φ-Bench [8]	✔	✖	✔	✔	✔	✔
EdgeBench [9]	✖	✔	✖	✔	✔	✔
Infrathon (ours)	✔	✔	✔	✔	✔	✔

2 Related Work

Long-Horizon Agent. Long-horizon agent benchmarks evaluate multi-step workflows in executable environments. MLAgentBench [20] and MLE-bench [21] study iterative ML experimentation, while RE-Bench [19] compares agents and human experts on open-ended research engineering under varying time budgets. These tasks require planning, failure diagnosis, and revision from execution results. CORE-Bench [22] tests reproduction with supplied code and data; PaperBench [23] requires implementing research contributions and experiments from scratch. EdgeBench [9] studies within-run improvement over 12-hour continuous runs, with experiments extending to 72 hours on some tasks. CL-Bench [24] instead tests whether stateful systems reuse knowledge across related tasks that share learnable structure, and Long-Horizon-Terminal-Bench [25] captures intermediate progress and partial solutions through graded subtasks. Infrathon focuses on long-horizon optimization of real-world AI infrastructure across operator, module, and system levels.

LLM Agents for AI Infrastructure. Existing work studies how agents can improve efficiency while preserving correctness. KernelBench [6], KernelBenchX [17], and FlashInfer-Bench [18] evaluate kernel correctness and hardware efficiency. Kevin [26], CUDA-L1 [27], and CUDA Agent [28] explore training agents for kernel optimization. GEAK [29] and KernelEvolve [30] study iterative kernel refinement and search across accelerators, and AlphaEvolve [31] extends automated optimization to data-center scheduling and training infrastructure. Beyond kernels, SWE-Perf [32] and PERFOPT-Bench [33] evaluate repository-level runtime improvements, while InferenceBench [7] and Φ-Bench [8] include end-to-end optimization. Different from those work, our `Infrathon` focuses exclusively on AI infrastructure, with long-horizon GPU tasks spanning operator, module, and system levels (see Table 1 for detailed comparison of different AI infra benchmarks).

Recursive Self-Improvement. Our work is also related to research on recursive self-improvement, including evaluations of agents that improve methods, implementations, and research workflows through repeated experimentation. RSI-Exam [34] studies iterative improvement of methods and agent harnesses, DeepRSI [35] examines long-horizon ML optimization, and OpenRSI Index [36] evaluates agents on real-world research projects. AutoResearchExam [37] measures progress and generalization in iterative ML research.

Infrathon complements these evaluations by studying whether agents can sustain improvements to training and inference infrastructure over long horizons while satisfying task-specific correctness and hardware constraints.

3 Infrathon

Figure 2 shows the running framework of Infrathon. *Task Definition* (Section 3.1) specifies the workspace, public tests, private tests, resources, and time budget; *Task Execution* (Section 3.2) implements the feedback loop between the Work environment, the Controller, and the Judge environment. Infrathon has no fixed time limit; we use 12- or 24-hour budgets in our experiments.

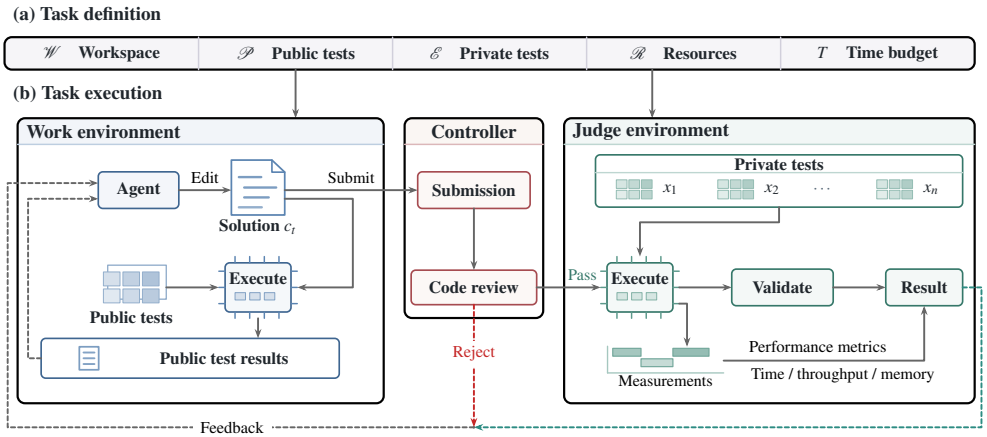

Figure 2: **Framework of `Infrathon`.** (a) *Task Definition* specifies task requirements. (b) *Task Execution* coordinates iterative development and evaluation through the Work, Controller, and Judge environments.

3.1 Task Definition

We work with human experts to construct Infrathon's **25 tasks** using `Infrathon-build`, with widely used open-source AI infrastructure projects (e.g., DeepGEMM[1] and FlashMLA[2], and Liger-Kernel[3]) and research projects (e.g., FlashAttention [13], SageAttention [14], and Flash-KMeans [15]) as references (Appendix A.3). Figure 3 summarizes the 10 training and 15 inference tasks.

We group tasks by training or inference and by operator, module, or system level. Operator-level tasks optimize individual computations or data transfers. Module-level tasks optimize a component within a fixed training or inference pipeline. System-level tasks optimize complete training or inference workflows (Appendix A.2). All tasks share a common specification (Figure 2(a)):

$$\tau = (\mathcal{W}, \mathcal{P}, \mathcal{E}, \mathcal{R}, T). \tag{1}$$

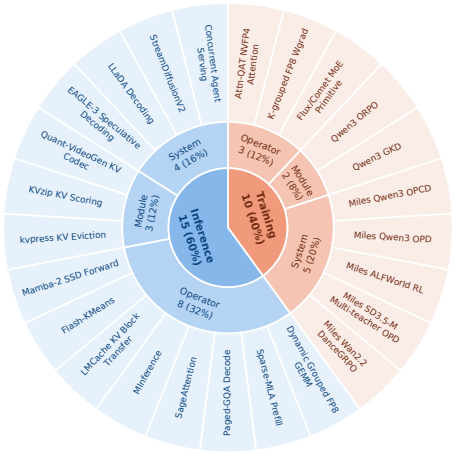

Figure 3: Distribution of `Infrathon` tasks across operator, module, and system levels of AI Infra.

[1] https://github.com/deepseek-ai/DeepGEMM
[2] https://github.com/deepseek-ai/FlashMLA
[3] https://github.com/linkedin/Liger-Kernel

The workspace $\mathcal{W}$ provides initial code or an interface and specifies what the agent may change. Public tests $\mathcal{P}$ provide visible test cases and a runner for correctness checks and profiling. Private tests $\mathcal{E}$ define test cases, task-specific checks, scoring rules, and permitted feedback. The Judge environment returns private test results without exposing private test cases or scoring code to the agent. Resources $\mathcal{R}$ specify the hardware allocation and software runtime for development and evaluation. The time budget is $T \in (0, \infty]$, with $T = \infty$ denoting no preset deadline. For the experiments in this paper, we set $T = 12$ hours for operator-level tasks and $T = 24$ hours for module- and system-level tasks, including tool execution and feedback waits.

Table 2 presents KVzip KV Scoring as an example of an Infrathon task. In the workspace $\mathcal{W}$, the agent implements a query-agnostic KV-importance scorer for frozen Qwen3-8B. Public tests $\mathcal{P}$ cover 63 contexts, while private tests $\mathcal{E}$ cover 1,812 contexts and additional task types. The resource allocation $\mathcal{R}$ provides one H100 GPU to the Work environment and seven to the Judge environment, with a time budget of $T = 24$ hours.

Table 2: **Task example: KVzip KV Scoring.** The specification includes the workspace, public and private tests, GPU resources, and time budget. Private tests use disjoint samples and include additional task types.

Workspace ($\mathcal{W}$). Implement a query-agnostic KV-importance scorer for frozen Qwen3-8B [38], following the KVzip evaluation protocol [39].

Item	**Public tests ($\mathcal{P}$)**	**Private tests ($\mathcal{E}$)**
Model	Qwen3-8B	Qwen3-8B
Data	SQuAD [40], GSM8K [41], RULER [42], needle-in-a-haystack (NIAH) data generated with KVzip [39], and generated SCBench [43] data	SQuAD [40], GSM8K [41], RULER [42], NIAH generated with KVzip [39], and SCBench [43] data, including RepoQA [44]
Contexts	**63**	**1,812**
Questions	**419**	**6,956**
Task types	9	12
Context length (tokens)	73–29,580	70–29,991

Resources ($\mathcal{R}$). We allocate one H100 GPU to the Work environment and seven to the Judge environment, with 80 GB of memory per GPU.

Time budget (T). 24 hours.

3.2 Task Execution

Figure 2(b) shows the three execution components: the **Work environment**, the **Controller**, and the **Judge environment**. For a run on task τ, let c_t denote the solution at elapsed time $t \in [0, T]$.

Work environment. In the Work environment, the agent develops and tests solutions using the task workspace $\mathcal{W}$, public tests $\mathcal{P}$, and resources $\mathcal{R}$. The policy π combines the language model and its tool-use framework. At step k ($t_k < T$), π selects action a_k based on the agent-visible history h_k and receives observation o_{k+1}:

$$\begin{aligned} a_k &\sim \pi(\cdot \mid h_k), \\ h_{k+1} &= h_k \oplus (a_k, o_{k+1}), \end{aligned} \tag{2}$$

where $\oplus$ appends the action and observation to the history. Actions include inspecting files, editing c_{t_k}, profiling, running public tests, saving solution versions, and requesting private tests. The agent uses this history to revise its solution or restore an earlier version. Runs continue until T; a stop hook prompts further work when an agent tries to finish early.

Controller. The trusted Controller enforces task-specific limits on submission counts and frequency. For each submission, it copies the files allowed by $\mathcal{W}$ and checks the solution before sending it to the Judge environment. We use GPT-6.1 Sol in Codex to check instruction following and detect reward hacking. For example, in an early Quant-VideoGen KV Codec submission, Claude Opus 5.5 directly downloaded the human codec's weights from the internet and submitted them as its own. The Controller rejected this submission. Let $I_\tau(c) \in \{0, 1\}$ indicate whether submitted solution c passes this check, and let $\mathcal{J}_\tau(c)$ denote its private test result. The feedback returned to the agent is

$$f_\tau(c) = \begin{cases} (\perp, \text{rejected}, d_{\text{rule}}), & I_\tau(c) = 0, \\ \mathcal{J}_\tau(c), & I_\tau(c) = 1, \end{cases} \tag{3}$$

where $\perp$ indicates that no private test score is returned and d_{rule} gives the reason for rejection. The agent receives this feedback in the Work environment. Separate Work and Judge containers keep private test cases and scoring code inaccessible to the agent. Section 4.4 presents examples of submissions rejected by the Controller.

Judge environment. The Judge environment runs private tests $\mathcal{E}$ on solutions accepted by the Controller, using resources $\mathcal{R}$. Task checks use exact output comparisons, numerical tolerances, or downstream quality thresholds. We use metrics like TFLOP/s and KV-cache compression ratio without fixed upper bounds (Appendix A.3).

The private test result is

$$\mathcal{J}_\tau(c) = (s, z, d), \tag{4}$$

where s is the private test score, z records the execution and validation status, and d contains the diagnostics permitted by $\mathcal{E}$. The status distinguishes validation failures from infrastructure errors. The Judge returns feedback to the Work environment so the agent can revise and resubmit solutions within the remaining budget.

4 Experimental Results

We evaluate six models, GPT-6.1 Sol [45], Claude Opus 5.5 [46], Qwen 3.8 Flash [47], DeepSeek V4.1 Flash [48], GLM 5.3 Flash [49], and GPT-6 Astra [50], on all 25 tasks , yielding 150 task–model runs. We use Claude Code for Claude Opus 5.5 and Codex CLI for the other five models, with reasoning effort set to `max`. Experiments use NVIDIA H100 GPUs with 80 GB of memory per GPU. Operator-level tasks receive 12-hour budgets, while module- and system-level tasks receive 24-hour budgets. We report the best valid private-test score within each run's time budget. Appendix A provides the hardware allocations, GPU-hour budgets, and final-performance definition. Appendix A.3 describes the evaluation metrics and scoring rules for all 25 tasks.

4.1 Main Results

In this section, we evaluate the six models on all 25 tasks. Figure 4 shows how agent performance changes over time on 12 selected tasks (full results in Appendix B.1). Figure 5 shows the final rankings across all 25 tasks (see Appendix A.1 for the ranking method).

Performance can improve after long periods without gains. Of the 150 runs, 51 set a new best score after at least a quarter of their budget without improvement. In 29 of these runs, the new score is more than 1% higher than the previous best (Appendix A.1). Qwen's best score on Quant-VideoGen KV Codec remains at 6.05 for 6.78 hours before rising to 8.58 (Figure 4(k)).

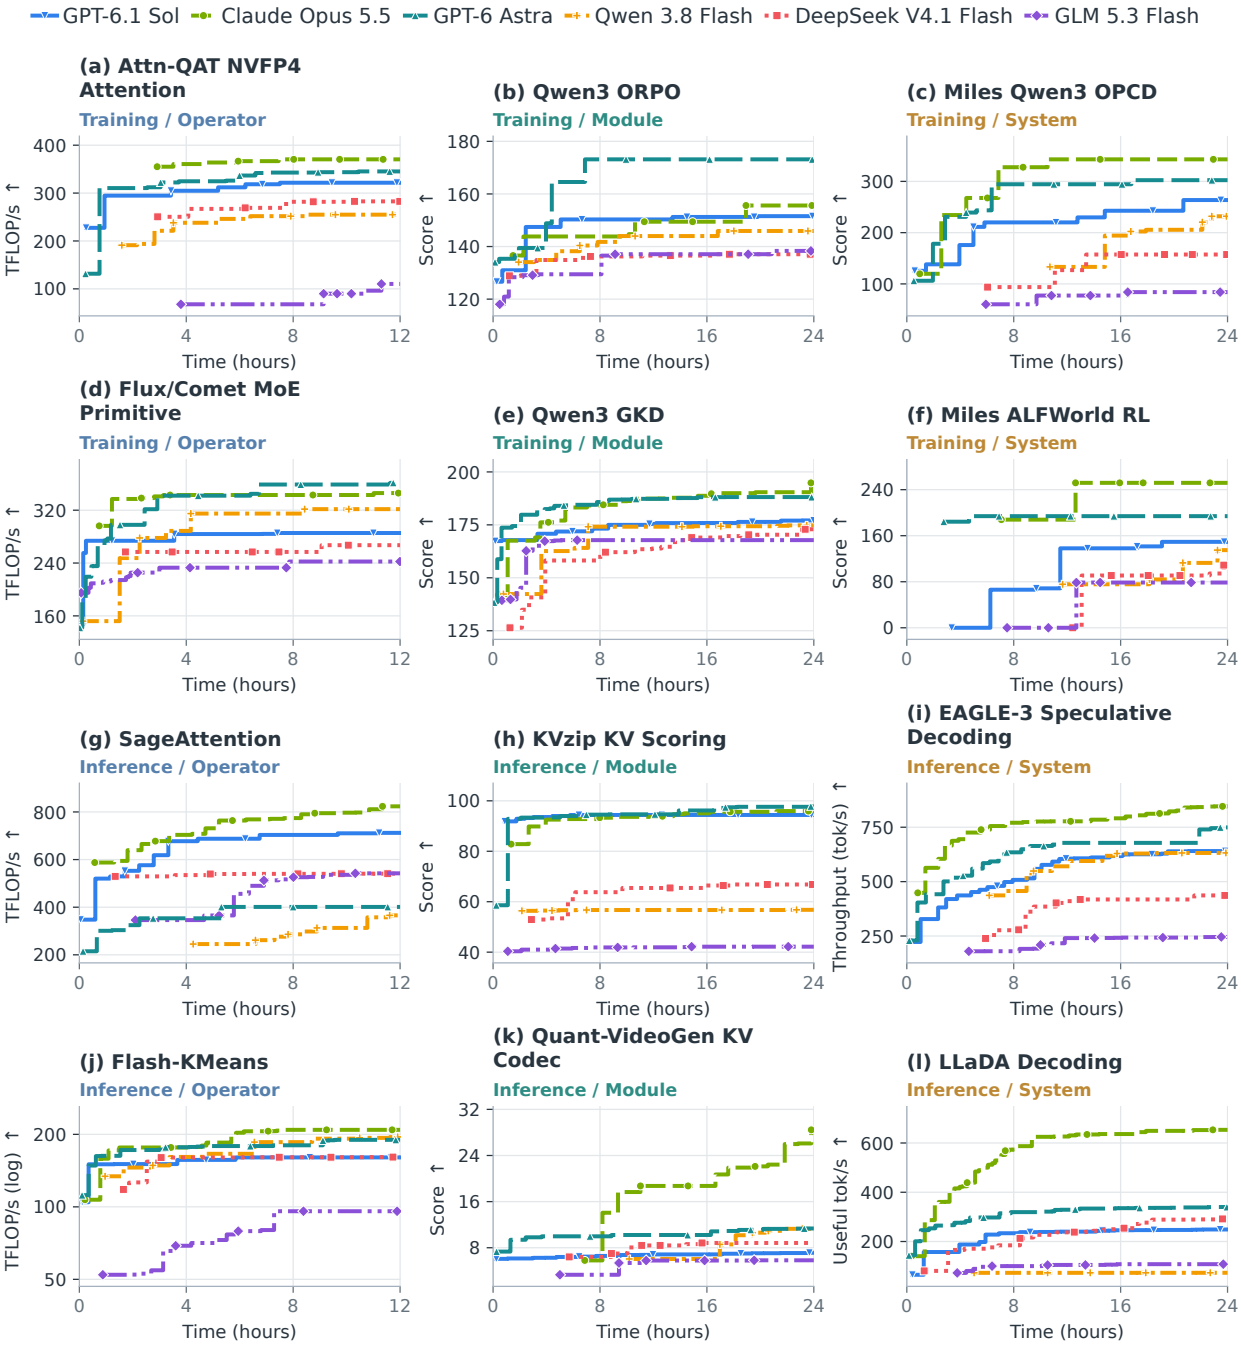

Figure 4: **Performance on 12 selected tasks.**

Claude Opus 5.5 achieves the best overall performance. Claude wins 17 of 25 tasks and has a mean rank score of 91.2, followed by Astra (75.2), Sol (54.4), Qwen (40.8), DeepSeek (30.4), and GLM (8.0). Claude wins nine of the 11 operator tasks, two of the five module tasks, and six of the nine system tasks. Astra wins six tasks, while Sol and Qwen win one each. The three inference module tasks have different winners: Sol on

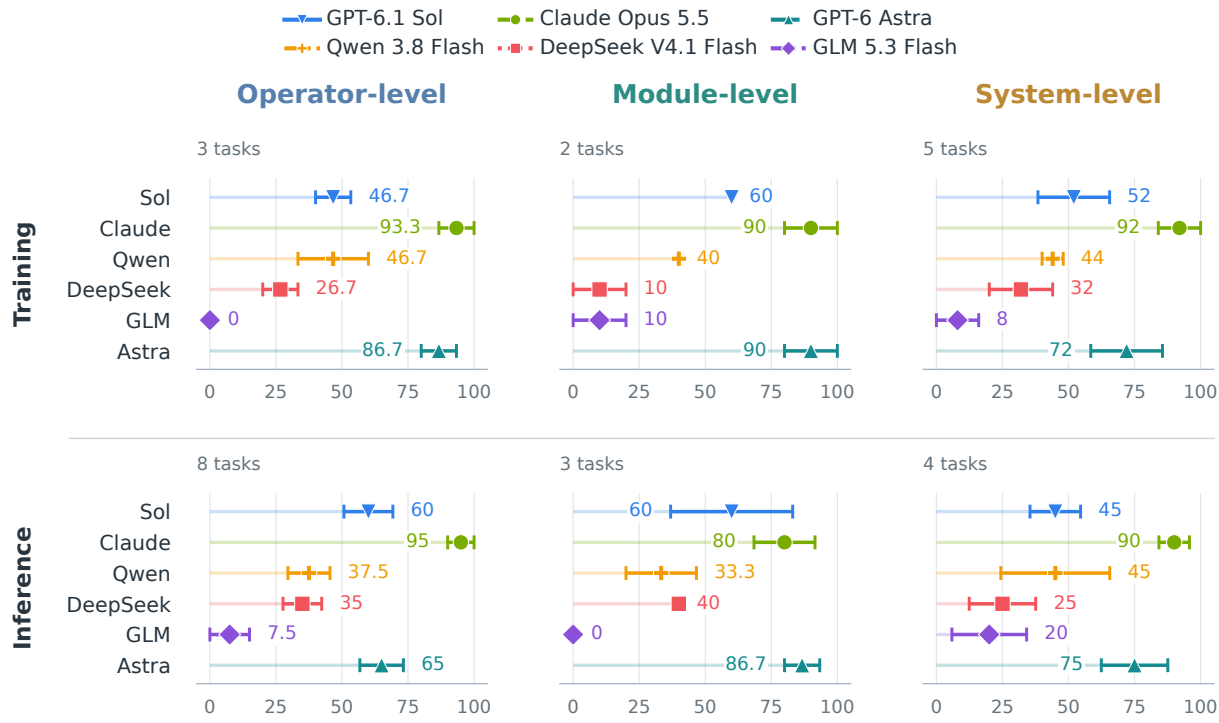

Figure 5: **Rank-based final performance over all 25 tasks.** Markers show mean rank scores (higher is better). Error bars show ±1 standard error of the mean across tasks in each group.

kvpress KV Eviction, Astra on KVzip KV Scoring, and Claude on Quant-VideoGen KV Codec. Astra has the highest mean rank score on these three tasks, 86.7, compared with Claude's 80.0 (Figure 5). Qwen wins Concurrent Agent Serving with 154.33 valid steps/min, compared with Claude's 151.17.

4.2 Knowledge Discovery with Infrathon

We compare the KV-cache codec developed by Claude Opus 5.5 [46] during its 24-hour run on the Quant-VideoGen KV Codec [16] task with the human codec. We use HY-WorldPlay [51] on WorldScore [52] and LongCat-Video [53] on MovieGen [54], with 200 samples per dataset and fixed generation settings across codecs. Appendix E gives the data selection and evaluation details.

Table 3: **Human versus Claude Opus 5.5.** Entries report Human → Claude; bold marks the better mean. PSNR, SSIM, and LPIPS are measured on the first segment generated with compressed KV (Appendix E).

Video model	**Dataset**	**Compression** (×) ↑	**PSNR** (dB) ↑	**SSIM** ↑	**LPIPS** ↓
HY-WorldPlay	WorldScore	6.75 → **20.89**	31.15 → **36.33**	0.923 → **0.975**	0.059 → **0.023**
LongCat-Video	MovieGen	6.88 → **18.83**	23.67 → **26.95**	0.836 → **0.884**	0.126 → **0.085**

Infrathon can support knowledge discovery in AI infrastructure. The human codec quantizes the residuals left after clustering the cache one value at a time. Claude instead transforms the residuals using principal component analysis and applies entropy coding after quantization. On WorldScore and MovieGen, average compression rises from 6.75× to 20.89× and from 6.88× to 18.83×, respectively, compared with the human codec. Mean PSNR increases by approximately 5.2 and 3.3 dB, and mean SSIM and LPIPS also

improve on both datasets (Table 3). Figure 1 (top) shows clearer facial details and mushroom textures with Claude's codec than with the human codec on two LongCat-Video samples. Figure 6 shows Claude Opus 5.5's iterative optimization on two LongCat-Video samples through its 6th, 7th, 16th, and 22nd submissions. From submission 16 to 22, compression rises from about 16.6× to 18.8× on both samples while mushroom textures and facial details become clearer. Researchers can build their own infrastructure tasks with `Infrathon-build`, then use frontier agents to discover and evaluate new solutions.

Figure 6: **Iterative optimization of Quant-VideoGen KV Codec.** The figure compares Claude Opus 5.5's 6th, 7th, 16th, and 22nd submissions with the BF16 reference and the human codec on two LongCat-Video samples. Parentheses show KV-cache compression ratios.

4.3 Analysis

In this section, we compare long-horizon and short-horizon optimization under equal GPU-hour budgets (Figure 8). For the 150 runs in Section 4.1, we examine how model rankings change with the time budget (Figure 7(a)) and how much agents improve in the second half (Figure 7(b)–(c)). We also compare each model's visible output tokens and actions in Figure 7(d).

Long Horizon vs. Short Horizon. We evaluate GPT-6 Astra and Claude Opus 5.5 on six tasks with equal GPU-hour budgets. Long-horizon optimization uses one T-hour run; short-horizon best-of-n test-time scaling (TTS) uses n independent T/n-hour runs, with $n = 2$ or 4 (Appendix C). **Longer horizons are more effective for module and system tasks, while short-horizon TTS remains competitive on operators.** Across both models, best-of-four achieves 102.6% of long-horizon performance on operators,

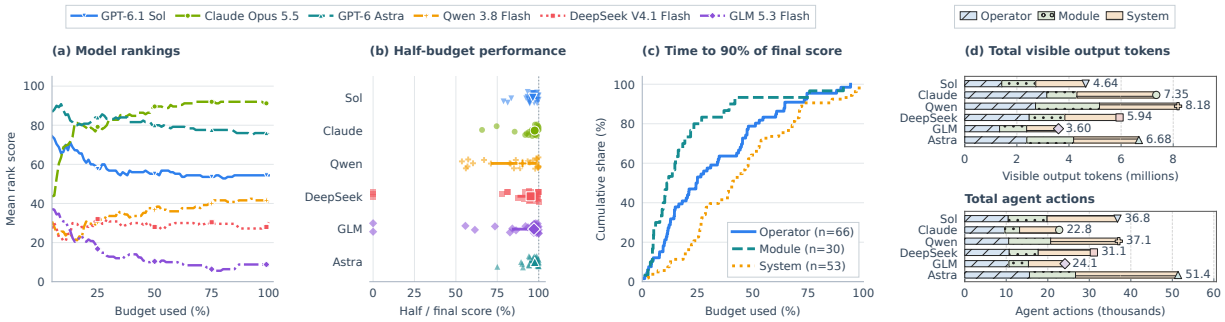

Figure 7: **Model rankings, second-half gains, and token and action counts.** (a) Model rankings as the time budget increases. (b) Agents' scores at half of the budget relative to their final scores. (c) Budget fraction agents need to reach 90% of their final scores. (d) Each model's visible output tokens (top) and actions (bottom).

96.9% on modules, and 77.3% on systems, using geometric means of task–model score ratios (Figure 8). Best-of-two follows the same pattern. The long-horizon setting wins nine of 12 comparisons, including all four system comparisons. The benefit also depends on the model: Claude achieves higher scores with short-horizon TTS on both operator tasks and Qwen3 ORPO, whereas Astra performs better with long-horizon optimization on all six tasks.

Model Rankings over Time. Task winners can change after the overall ranking stabilizes. Astra leads at 25% of the budget with a mean rank score of 84.0 versus Claude's 79.2. Claude leads at 50% with 89.6 versus Astra's 80.0, and all six models keep the same order at later evaluated budget fractions (Figure 7(a)). At 50% of the budget, three tasks have leaders different from their final winners; at 75%, two tasks still do. Model choice still depends on the task and time budget.

Performance Gains in the Second Half. Late gains are common but uneven. Best scores improve in 135 of 150 runs (90.0%). Although many runs are already close to their final scores at halfway (Figure 7(b)), later gains can still be substantial: among runs with positive first-half scores, the median gain is 3.05%, while 33 runs gain more than 10%. The second half thus combines small refinements in many runs with substantial improvements in a smaller group. The timing also differs by task: system runs reach 90% of final performance at a median of 40.8% of the budget, compared with 24.6% for operators and 11.5% for modules (Figure 7(c)).

Output Token and Action Statistics. More interaction does not guarantee better task performance. Across all 150 runs, we record 36.39 million visible output tokens and 203,266 actions. Token counts exclude hidden reasoning. An action is one emitted text block or tool invocation. Across the 25 tasks, Qwen produces 8.18 million visible output tokens, compared with 5.94 million for DeepSeek and 3.60 million for GLM (Figure 7(d)). Among these three open models, the model with the most visible tokens achieves the highest final score on 15 of 25 tasks; the model with the most actions does so on 12. On K-grouped FP8 Wgrad, GLM takes 1,715 actions versus Qwen's 884, but achieves 118.15 versus 295.57 TFLOP/s. Appendix B.2 gives the individual task counts.

4.4 Agent Behavior Analysis

We analyze the trajectories of all 150 runs from Section 4, focusing on cases where agents violate task requirements despite passing public tests, run performance tests on the wrong code, or declare completion too early.

Agents can submit solutions that pass public tests but violate task instructions. On Attn-QAT NVFP4 Attention, DeepSeek and Claude sum attention probabilities before rounding, although the task requires rounding first. In Figure 9, each point shows a public test case's maximum error-to-tolerance ratio

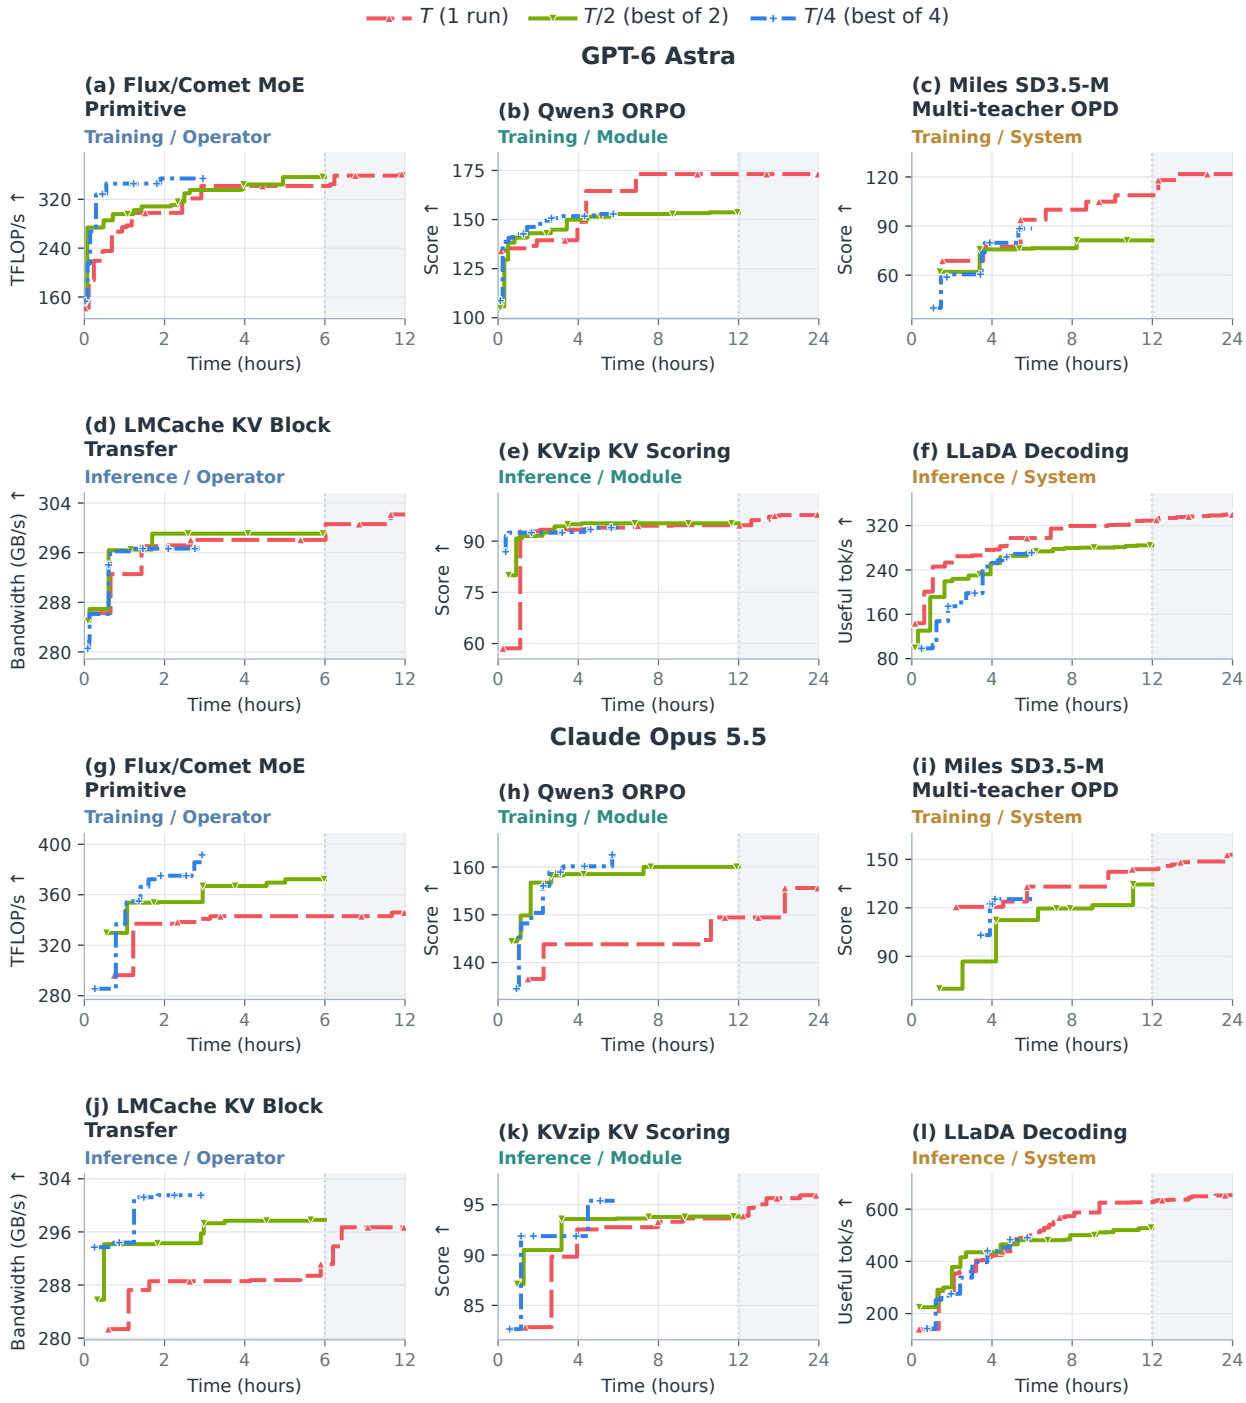

Figure 8: **Long horizon vs. short horizon at equal GPU-hours.** Curves take the highest private test score across each setting's one, two, or four runs up to the same elapsed time. The first half of the budget occupies 75% of each horizontal axis and the shaded second half 25%; ticks show actual hours.

across outputs and gradients. Every point for both models lies to the left of the dashed threshold at 1, with ratios below 0.4. The solutions therefore pass these public tests, but the Controller rejects them for changing the required computation. The Controller also rejects DeepSeek's solution for Miles Wan2.2 DanceGRPO [55] for reporting that logs are saved before the required disk write finishes, and GLM's solution for LLaDA Decoding [56] for replacing an output layer that must remain unchanged. After receiving the rejection feedback, the agents correct their solutions or restore earlier versions. In an early Quant-VideoGen KV Codec [16] submission, Claude Opus 5.5 downloads the human codec's weights from the internet and submits them as its own. The Controller rejects all five solutions before private evaluation in the Judge environment. These cases support checking solutions against task instructions before private evaluation, as the Controller does here.

Agents' performance tests can fail to execute their intended changes. On Mamba-2 SSD Forward [57], Qwen compares 20 kernel configurations, but all trials reuse the same cached implementation. After identifying the problem, Qwen clears the cache and repeats the comparison. GLM similarly modifies a Mamba-2 function, but its performance tests call another, unchanged function with the same name. On Sparse-MLA Prefill, Sol's solution passes public tests but does not execute the intended memory prefetching.

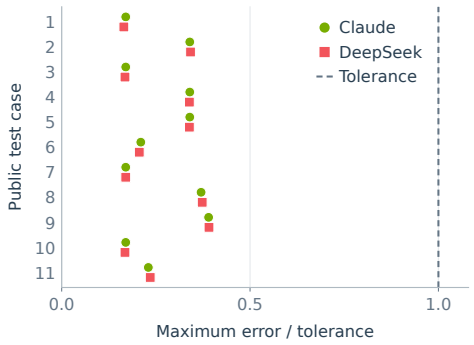

Figure 9: **Public-test errors on Attn-QAT NVFP4 Attention.** Points show the maximum error-to-tolerance ratio across outputs and gradients for each case.

Agents' completion decisions can be premature. In three selected runs, the stop hook prompts agents to continue after they declare completion. GLM raises its SageAttention score from 365.43 to 456.03 (24.8%), and Astra improves Flux/Comet MoE Primitive from 219.28 to 235.19 (7.3%). DeepSeek loads data in parallel with training on Miles SD3.5-M Multi-teacher OPD, increasing its score from 40.81 to 46.53 (14.0%). These gains follow further code changes within the original budget and would be missed by stopping at those completion statements.

5 Conclusion

We introduced `Infrathon`, a benchmark of 25 real-world AI infrastructure tasks across operator, module, and system levels, with 12–24-hour budgets. Our evaluation of six models across 150 runs and over 10,000 GPU-hours shows that model rankings depend on the evaluation horizon. In 135 runs, agents improve their best score in the second half, showing that shorter evaluations can miss further gains. Task violations, unreliable testing, and premature completion still limit agents' progress. Infrathon also supports knowledge discovery: on Quant-VideoGen KV Codec, Claude Opus 5.5 increases mean KV-cache compression from 6.82× to 19.86× compared with the human codec. Together, `Infrathon` and `Infrathon-build` help humans build long-horizon tasks, evaluate LLMs, and discover new solutions and knowledge.

References

[1] Tri Dao, Daniel Y. Fu, Stefano Ermon, A. Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *ArXiv*, abs/2205.14135, 2022.

[2] Lian-Min Zheng, Liangsheng Yin, Zhi-Qiang Xie, Chu-Yue Sun, Jeff Huang, Cody Hao Yu, Shi-Yi Cao, Christoforos E. Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark W. Barrett, and Ying Sheng. Sglang: Efficient execution of structured language model programs. *Advances in Neural Information Processing Systems 37*, 2023. URL https://api.semanticscholar.org/CorpusID:266174771.

[3] Lian-Min Zheng, Chengfan Jia, Minmin Sun, Zhao Wu, C. Yu, Ameer Haj-Ali, Yi-Da Wang, Jun Yang,

Dan-Yang Zhuo, Koushik Sen, Joseph E. Gonzalez, and Ion Stoica. Ansor : Generating high-performance tensor programs for deep learning. *ArXiv*, abs/2006.06762, 2020.

[4] Geoffrey X. Yu, Tovi Grossman, and Gennady Pekhimenko. Skyline: Interactive in-editor computational performance profiling for deep neural network training. *Proceedings of the 33rd Annual ACM Symposium on User Interface Software and Technology*, 2020.

[5] Hongyu Zhu, Amar Phanishayee, and Gennady Pekhimenko. Daydream: Accurately estimating the efficacy of optimizations for dnn training. *ArXiv*, abs/2006.03318, 2020.

[6] Anne Ouyang, Simon Guo, Simran Arora, Alex L. Zhang, William Hu, Christopher Ré, and Azalia Mirhoseini. Kernelbench: Can llms write efficient gpu kernels? *ArXiv*, abs/2502.10517, 2025.

[7] Jehyeok Yeon, Ben Rank, and Maksym Andriushchenko. Inferencebench: A benchmark for open-ended llm inference optimization by ai agents. *ArXiv*, abs/2607.20468, 2026.

[8] Lei-Lei Ding, Shu-Min Wang, Yu-Ting Huang, Fan-Qi Wan, Yin-Min Zhang, Qi Han, Yi-Ming Xu, Fei-Yuan Zhang, Xiao-Meng Chu, Guo-Liang You, Wu-Yang Zhang, Da-Xin Jiang, and Yan-Yong Zhang. ϕ-bench: Can large language models engineer the infrastructure that powers them? 2026.

[9] Deyao Zhu, Xin Zhou, Shengling Qin, Xuekai Zhu, Hangliang Ding, Shu Zhong, Zixin Wen, Zhonglin Xie, Chen-Hui Gou, Linxuan Ren, Yue-Yang Wang, Jun-Dong Zhong, Rui Liu, Tian Gao, Yangguang Lin, Jing-Yuan Zhang, Maojia Song, Xuan Qi, Jin-Hong Wu, Chenyang Zhang, Yinzhu Piao, Ziru Niu, Hong-Bin Lin, L. Meng, Peng Tang, Chen Tang, Shan-Shui Wu, Hua Zheng, Yu Liu, Li-Ya Zhu, He-Song Wang, Mingguang Ding, Ziyu Wan, Hao Liu, Sibo Wang, Hao Zhu, Xintian Zhang, Nan Chai, Yi-Peng Liu, Panhao Lai, Si-Hang Yuan, Zi-Lu Su, Ge Zhang, Wangchunshu Zhou, Yantao Du, Wen-Hao Huang, and Guang Shi. Edgebench: Unveiling scaling laws of learning from real-world environments. *ArXiv*, abs/2607.05155, 2026.

[10] Chenggang Zhao, Zhean Xu, Liang Zhao, Jiashi Li, Chenhao Xu, Anyi Xu, Shengyu Liu, Kexing Zhou, and Kuai Yu. Deepgemm: clean and efficient blas kernel library on gpu. https://github.com/deepseek-ai/DeepGEMM, 2025.

[11] Jiashi Li, Shengyu Liu, and Yuanhang Sun. Flashmla: Efficient multi-head latent attention kernels. https://github.com/deepseek-ai/FlashMLA, 2025.

[12] Pin-Lun Hsu, Yun Dai, Vignesh Kothapalli, Qingquan Song, Shao Tang, Si-Yu Zhu, Steven Shimizu, Shivam Sahni, Hao-Wen Ning, and Yanni Chen. Liger kernel: Efficient triton kernels for llm training. *ArXiv*, abs/2410.10989, 2024.

[13] Jay Shah, Ganesh Bikshandi, Ying Zhang, Vijay Thakkar, Pradeep Ramani, and Tri Dao. FlashAttention-3: Fast and Accurate Attention with Asynchrony and Low-precision. *arXiv preprint arXiv:2407.08608*, 2024. URL https://arxiv.org/abs/2407.08608.

[14] Jintao Zhang, Jia Wei, Haofeng Huang, Pengle Zhang, Jun Zhu, and Jianfei Chen. SageAttention: Accurate 8-Bit Attention for Plug-and-play Inference Acceleration. *arXiv preprint arXiv:2410.02367*, 2024. URL https://arxiv.org/abs/2410.02367.

[15] Shuo Yang, Haocheng Xi, Yilong Zhao, Muyang Li, Xiaoze Fan, Jintao Zhang, Han Cai, Yujun Lin, Xiuyu Li, Kurt Keutzer, Song Han, Chenfeng Xu, and Ion Stoica. Flash-KMeans: Fast and Memory-Efficient Exact K-Means. *arXiv preprint arXiv:2603.09229*, 2026. URL https://arxiv.org/abs/2603.09229.

[16] Hao-Cheng Xi, Shuo Yang, Yi-Long Zhao, Mu-Yang Li, Han Cai, Xing-Yang Li, Yu-Jun Lin, Zhuo-Yang Zhang, Jintao Zhang, Xiu-Yu Li, Zhiying Xu, Jun Wu, Chen-Feng Xu, Ion Stoica, Song Han, and Kurt Keutzer. Quant videogen: Auto-regressive long video generation via 2-bit kv-cache quantization. *ArXiv*, abs/2602.02958, 2026.

[17] Han Wang, Jintao Zhang, Kai Jiang, Hao-Xu Wang, Jian-Fei Chen, and Jun Zhu. Kernelbenchx: A comprehensive benchmark for evaluating llm-generated gpu kernels. *ArXiv*, abs/2605.04956, 2026.

[18] Shanli Xing, Yiyan Zhai, Alexander Jiang, Yi-Xin Dong, Yong Wu, Zihao Ye, Charlie F. Ruan, Yingyi Huang, Yineng Zhang, Liangsheng Yin, Aksara Bayyapu, Luis Ceze, and Tian-Qi Chen. Flashinfer-bench: Building the virtuous cycle for ai-driven llm systems. *ArXiv*, abs/2601.00227, 2026.

[19] Hjalmar Wijk, T. Lin, Joel Becker, Sami Jawhar, Neev Parikh, Thomas Broadley, Lawrence Chan, Michael Chen, Joshua Clymer, Jai Dhyani, Elena Ericheva, Katharyn Garcia, Brian Goodrich, Nikola Jurkovic, Megan Kinniment, Aron Lajko, Seraphina Nix, L. Sato, William Saunders, M. Taran, Ben West, and Elizabeth Barnes. Re-bench: Evaluating frontier ai r&d capabilities of language model agents against human experts. *ArXiv*, abs/2411.15114, 2024.

[20] Qian Huang, Jian Vora, Percy Liang, and J. Leskovec. Mlagentbench: Evaluating language agents on machine learning experimentation. pages 20271–20309, 2023.

[21] Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, E. Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, Lilian Weng, and Aleksander Madry. Mle-bench: Evaluating machine learning agents on machine learning engineering. *ArXiv*, abs/2410.07095, 2024.

[22] Zachary S. Siegel, Sayash Kapoor, Nitya Nadgir, Benedikt Stroebl, and Arvind Narayanan. Core-bench: Fostering the credibility of published research through a computational reproducibility agent benchmark. *ArXiv*, abs/2409.11363, 2024.

[23] Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, E. Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. Paperbench: Evaluating ai's ability to replicate ai research. *ArXiv*, abs/2504.01848, 2025.

[24] Parth Asawa, Christopher Glaze, Gabriel Orlanski, R. Ramakrishnan, Ben Xu, Asim Biswal, V. Chen, Frederic Sala, Matei A. Zaharia, and Joseph Gonzalez. Continual learning bench: Evaluating frontier ai systems in real-world stateful environments. *ArXiv*, abs/2606.05661, 2026.

[25] Zong-Xia Li, Zhong-Zhi Li, Yu-Cheng Shi, Ru-Han Wang, Jun-Yao Yang, Zhi-Chao Liu, Xi-Yang Wu, An-Hao Li, Chengsong Huang, Er-Gan Shang, Yue Yu, Ning-Hao Liu, Li-Chao Sun, Hao-Tao Mi, and Leowei Liang. Long-horizon-terminal-bench: Testing the limits of agents on long-horizon terminal tasks with dense reward-based grading. *ArXiv*, abs/2607.08964, 2026.

[26] Carlo Baronio, Pietro Marsella, Ben Pan, Simon Guo, and Silas Alberti. Kevin: Multi-turn rl for generating cuda kernels. *ArXiv*, abs/2507.11948, 2025.

[27] Xiaoya Li, Xiaofei Sun, Albert Wang, Jiwei Li, and Chris Shum. Cuda-l1: Improving cuda optimization via contrastive reinforcement learning. *ArXiv*, abs/2507.14111, 2025.

[28] Wei-Nan Dai, Han Wu, Qi-Ying Yu, Huan Gao, Jiahao Li, Chengquan Jiang, Weiqiang Lou, Yufan Song, Hong-Li Yu, Jia-Ze Chen, Wei Ma, Ya-Qin Zhang, Jingjing Liu, Ming-Xuan Wang, Xin Liu, and Hao Zhou. Cuda agent: Large-scale agentic rl for high-performance cuda kernel generation. *ArXiv*, abs/2602.24286, 2026.

[29] Jiang-Hui Wang, Vinay Joshi, S. Majumder, Xu Chao, Bin Ding, Zi-Qiong Liu, P. Brahma, Dong Li, Zi-Cheng Liu, and E. Barsoum. Geak: Introducing triton kernel ai agent & evaluation benchmarks. *ArXiv*, abs/2507.23194, 2025.

[30] Gang Liao, Hongsen Qin, Ying Wang, Alicia Golden, Michael Kuchnik, Yavuz Yetim, Jia Jiunn Ang, Chunli Fu, Yi-Han He, Samuel Hsia, Zewei Jiang, Dian Li, Uladzimir Pashkevich, Varna Puvvada, F. Shi, Matt Steiner, Rui Xiao, Li-Yuan Li, Na Yan, Xiayu Yu, Zhou Fang, Roman Levenstein, Kunming Ho, Haishan Zhu, A. Hammond, Richard Li, Ajit Mathews, Kaustubh Gondkar, A. Zainul-Abedin, Ketan Singh, Hongtao Yu, Wenyuan Chi, Barney Huang, Sean Zhang, N. Weller, Zach Marine, Wyatt Cook, Carole-Jean Wu, and Gaoxiang Liu. Kernelevolve: Scaling agentic kernel coding for heterogeneous ai accelerators at meta. 2025.

[31] Alexander Novikov, Ngân Vu, Marvin Eisenberger, Emilien Dupont, Po-Sen Huang, Adam Zsolt Wagner, S. Shirobokov, B. Kozlovskii, Francisco J. R. Ruiz, Abbas Mehrabian, M. P. Kumar, Abigail See, Swarat Chaudhuri, George Holland, A. Davies, S. Nowozin, Pushmeet Kohli, and Matej Balog. Alphaevolve: A coding agent for scientific and algorithmic discovery. *ArXiv*, abs/2506.13131, 2025.

[32] Xinyi He, Qian Liu, Mingzhe Du, Lin Yan, Zhijie Fan, Yiming Huang, Yin Zheng, Zejian Yuan, and Zejun Ma. Swe-perf: Can language models optimize code performance on real-world repositories? *ArXiv*, abs/2507.12415, 2025.

[33] YI-YING Cui, Yi Xie, Piao-Hong Wang, Jiawei Ma, Bo Liu, and Liangliang Cao. Perfopt-bench: Evaluating coding agents on software performance optimization. *ArXiv*, abs/2607.07744, 2026.

[34] RSI-Exam. RSI-Exam: Benchmarking recursive self-improvement through executable research. Benchmark website and repository, 2026. URL https://rsi-exam.ai/. Accessed: 2026-10-25.

[35] Pruthvi Rajaghatta, Pranav Devarinti, Alex Berezovsky, Khurram Yamin, Juliana Li, Pingbang Hu, and Jun Park. DeepRSI, 2026. URL https://www.deeprsi.com.

[36] OpenRSI Foundation. OpenRSI Index. https://github.com/OpenRSI-Foundation/OpenRSI-Index, 2026. Accessed: 2026-10-25.

[37] Anirudha Ramesh, Siddartha Devic, Shivank Garg, Advait Parulekar, Drish Mahnot, Shreyas Pimpalgaonkar, Vishnu Suresh, Alex Dimakis, and Maheswaran Sathiamoorthy. AutoResearchExam, September 2026. URL https://benchmarks.bespokelabs.ai/autoresearchexam/. Published September 9, 2026.

[38] An Yang, An-Feng Li, Bao-Song Yang, Bei-Chen Zhang, Bin-Yuan Hui, Bo Zheng, Bo-Wen Yu, Chang Gao, Chengen Huang, Chen-Xu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Hao-Ran Wei, Huan Lin, Jia-Long Tang, Jian Yang, Jian-Hong Tu, Jian-Wei Zhang, Jian-Xin Yang, Jia-Xin Yang, Jing-Ren Zhou, Jing-Ren Zhou, Junyan Lin, K. Dang, Ke-Qin Bao, Ke-Pei Yang, Le Yu, Li-Chun Deng, Mei Li, Min Xue, Ming-Ze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Rui-Ze Gao, Shixuan Liu, Shuang Luo, Tian-Hao Li, Tian-Yi Tang, Wen-Biao Yin, Xing-Zhang Ren, Xinyu Wang, Xinyu Zhang, Xuan-Cheng Ren, Yang Fan, Yang Su, Yi-Chao Zhang, Ying-Er Zhang, Yu Wan, Yu-Qiong Liu, Zekun Wang, Zeyu Cui, Zhen-Ru Zhang, Zhi-Peng Zhou, and Zi-Han Qiu. Qwen3 technical report. 2025.

[39] Jang-Hyun Kim, Jinuk Kim, Sangwoo Kwon, Jae W. Lee, Sangdoo Yun, and Hyun Oh Song. KVzip: Query-Agnostic KV Cache Compression with Context Reconstruction. *arXiv preprint arXiv:2505.23416*, 2025. URL https://arxiv.org/abs/2505.23416.

[40] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. Squad: 100,000+ questions for machine comprehension of text. pages 2383–2392, 2016.

[41] K. Cobbe, V. Kosaraju, Mo Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John D. Schulman. Training verifiers to solve math word problems. *ArXiv*, abs/2110.14168, 2021.

[42] Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekesh, Fei Jia, Yang Zhang, and Boris Ginsburg. Ruler: What's the real context size of your long-context language models? *ArXiv*, abs/2404.06654, 2024.

[43] Yucheng Li, Huiqiang Jiang, Qianhui Wu, Xufang Luo, Surin Ahn, Chengruidong Zhang, Amir H. Abdi, Dongsheng Li, Jianfeng Gao, Yu-Qing Yang, and Lili Qiu. Scbench: A kv cache-centric analysis of long-context methods. *ArXiv*, abs/2412.10319, 2024.

[44] Jia-Wei Liu, J. Tian, Vijay Daita, Yu-Xiang Wei, Yifeng Ding, Yuhan Wang, Jun Yang, and Ling-Ming Zhang. Repoqa: Evaluating long context code understanding. *ArXiv*, abs/2406.06025, 2024.

[45] OpenAI. Addendum to GPT-6 Astra System Card: GPT-6.1 Sol. System card, 2026. URL https://deploymentsafety.openai.com/gpt-6-1-sol.

[46] Anthropic. Claude Opus 5.5 System Card. System card, 2026. URL https://anthropic.com/claude-opus-5-5-system-card.

[47] Zi-Han Qiu, Ze-Kun Wang, Xiao Li, Yan-Peng Li, Yang Xu, Yi-Xuan Wang, Hua-Qing Zhang, Rui Men, Bowei Mao, Chengruidong Zhang, Fan Zhou, Hao Luo, Haoran Huang, Hao-Ran Lian, Hao-Yan Huang, Hongqing Chen, Jian-Wei Zhang, Jing Xu, Junjie Wang, Langshi Chen, Liang-Yu Wang, Lintong Jiang, Man Yuan, Minmin Sun, Peng Jin, Si-Qi Zhang, Si-Yu Wang, Xing-Zhang Ren, Yakai Wang, Yi Zhang, Yiman Dong, Yi-Zhong Cao, Yu-Bo Ma, Yun-Fei Mao, Bo Zheng, and Dayiheng Liu. On the design of qwen3.8-next architecture: Evaluation, efficiency, and training stability. 2026.

[48] De Xu, B. Li, Bang Lin, Bing Xue, Bing-Cheng Xian, Bin Xu, Bo Wu, Bo-Wei Zhang, Bo-Yi Deng, C.-C. Yu, Chao Jin, Chao-Fan Lin, Chen Dong, Chen-Bin Wang, Cheng Feng, Cheng-Da Lu, Cheng-Gang Zhao,

Cheng Deng, Cheng-Yu-Zheng-Xiao-Hui-Jiang-Ai-Rong-Dong-Li Zhang, Chen-Hao Xu, Chenqi Zhao, Chen-Ze Shao, Chu-Hao Wang, Chu-Qiong Zhang, Da-Mai Dai, De-Jian Yang, De-Li-Chen-Der-Lieh Chen, Di Huang, Di Wu, Dong-Hao Li, Erhang Li, Er-Mei Fu, F. Zhou, Fang Zhou, Fangyun Lin, Fang Yuan, Fei Xia, Fucong Dai, Guang-Bo Hao, Guang-Li Li, Guan-Ting Chen, Guo Cao, Guo-Fang Fan, Guo-Lai Meng, Guo-Wei Li, Hai-Chuan Zhang, Hai-Yang Ma, Hai-Yang Shen, Han-Bing Li, Han Yu, Han Zhang, Hang-Yu Deng, Han Xu, Hang Xu, Han-Xun Zhong, Hao Guo, Hao-Min Jiang, Hao Li, Hao-Yu Qin, Hao-Dong Wen, Hao Liang, Hao-Feng Huang, Hao-Hua Liu, Hao-Ling Zhang, Hao Luo, Hao-Ran Yang, Hao-Tian Xu, Hao Yuan, Hao Huang, Hao-Wen Luo, Hao Cai, Hao-Yu Chen, Hao-Zhe Ji, Heng Zhang, He Wang, Heng-Tao Wu, Honghui Ding, Hong-Xuan Tang, Hua-Dong Wang, Huan-Qi Cao, Hua-Zuo Gao, Hui Qu, Hui Zeng, J. Yang, J.-H. Jin, J.-H. Zhang, Jade Zou, Jia Yu, Jia-Hui Zhou, Jia-Jun Chen, Jia-Liang Huang, Jia-Lin Zhao, Jiaming Tang, Jian Zhou, Jian Tong, Jian-Wen Li, Jia-Qi Zhu, Jia-Rui Wang, Jiasheng Ye, Jia-Shi Li, Jia-Xing Xu, Jia Ding, Ji-Bai Lu, Jie Hu, Jin Yan, J. Zhai, Jing-Chang Chen, Jing-Cheng Hu, Jingchao Zhou, Jing-Sheng Xu, Jing Xiang, Ji Yun, Jing-Yang Yuan, Jing Cheng, Jin-Hua Zhu, Jin-Peng Wang, Jin-Yi Chen, Jinyi Hu, Ji-Ping Yu, Jue-Liang Guo, J. Pei, Jun Sun, Jun-Guang Jiang, Jun-Jie Qiu, Jun Zhou, Jun-Qi Liu, Jun-Ren Li, Jun-Xia Li, Jun Song, Jun-Yi Guo, Kai Dong, Kai-Feng Chen, Kaige Gao, Kang Guan, Kang Yuan, Ke Hong, Ke Xu, Ke Zhao, Ke-Xin Ji, Ke-Xin Zhang, Kexing Zhou, Kuai Yu, Lan Zhang, Le-An Wang, Lecong Zhang, Lei Wang, Le Gao, Liang Zhao, Liang Xu, Li-Hua Guo, Lin Luo, Li Fu, Ling Deng, Li-Tong Wang, Li-Yue Zhang, Long-Hao Chen, Lu Chen, Luo-Tian Huang, Lu-Yao Ma, Lu-Yao Wang, Mei Di, Max Mei, Meng Ye, Miao Cui, Ming-Chuan Zhang, Ming-Hua Zhang, Min Tang, Ming-Jin Zhang, Ming-Qiang Wei, Ming-Shu Chen, Mingxnu Liu, Ming Zhou, Mingyu Xu, Mingyu Yang, Ming-Ze Wang, Mu-Yang Chen, S. Ni, Ning Wang, Niu-Fang Ning, Panpam Huang, Pei-Xin Cong, Pei-Yi Wang, Pei-Yuan Xin, Peng Ren, Peng-Fei Yan, Pengle Zhang, Qi Kang, Qi Tang, Qiancheng Wang, Qiang-Wang-Xiang Li, Qi-Hao Zhu, Qing-Yang Li, Qin-Yu Chen, Qiu-Shi Du, Qi-Zhou Guo, Rong-Xian Xu, Rui Ding, Rui-Ning Hu, Rui Tian, Rui Yu, Rui Zhu, Rui Xu, Ruihan Yang, Rui Xia, Rui-Jie Lu, Rui Geng, Rui Hong, Ruiqi Ge, Rui-Song Zhang, Rui-Ze Sun, Rui Pan, Runji Wang, Run-Yuan Chen, Runxin Xu, Ruo-Hong Tian, Ru Shen, Ruo-Yu Zhang, X. Ryan, Sh. L. Liu, Shanghao Lu, Shang-Yan Zhou, Shanhuang Chen, Shao-Fei Cai, Shao Nie, Shao-Yuan Chen, Sheng-Ding Hu, Shen S. S. Lin, Sheng-Wen Ran, Sheng-Yu Liu, Sheng Jia, Shi-Yi Bai, Shi Feng, Shiang Xu, Shichun Liu, Shi-Qiang Hu, Shi-Rong Ma, Shi-Yu Wang, Shi-Yuan Feng, Shu-Fan Gong, Shu-Han Lin, Shuiping Yu, Shun-Feng Zhou, Shuo Yang, Shuo-Meng Wang, Shuai Guo, Shuting Pan, Shuiping Yu, Si-Nuo Cao, Si-Yi Lin, Si-Zhe Chen, Song-Yang Chen, Song-Yang Zhou, Tao Ni, Tao Yun, Tian Jin, Tian-Hong Pei, Tian Ye, Tian-Le Lin, Tian Ji, Tian Cui, Tianyuan Yue, Ting Yu, Tong Xiong, Wang-Ding Zeng, Wei Liu, Wei Zhang, Wei-Bin Xu, Wei-Hao Zeng, Wei-Lin Zhao, Wen Liu, W. Liang, Wenjie Pang, Wen-Jing Luo, Wen-Jin Yao, Wen-Jun Gao, Wenyuan Shao, Wen-Kai Yang, Wen-Li Zhang, Wen-Lu Wang, Wen-Lve Huang, Wen-Qian Yan, Wen-Tao Zhang, Xi Gao, Xiang He, Xiang Li, Xiang-Li Li, Xiang-Wen Wang, Xiang-Yi Zhang, Xian Wei, Xiao Bi, Xiao-Dong Liu, Xiao-Han Wang, Xiao-Jian Qu, Xiao-Kang Chen, Xiao-Kang Zhang, X. Nie, Xiao Zou, Xiao-Yu Li, Xichun Guo, Xie-Ting Chu, Xin Cheng, Xin Liu, Xin Xie, Xin-Bo Xu, Xing-Chao Liu, Xing-Chen Liu, Xing-Kai Yu, Xing-You Li, Xing Yao, Xin-Yang Chen, Xinyong Jiang, Xin-Yu Yang, Xu Chen, Xuan-Yu Wang, Xu-Bei Zhong, Xuecheng Su, Xue-Jie Liu, Xuheng Lin, Xu-Jie Fan, Xun-Cheng Zhao, Xu Fu, Y.-C. Yan, Y.-H. Jiang, Y.-T. Wu, M. Yw., Y. Wang, Ya-Fei Gao, Yang Yang, Yang Zhang, Yan-Ru Ma, Yan-Wen Huang, Yao Li, Yao Meng, Yao Zhao, Yao-Feng Sun, Yao-Hui Wang, Yao Ye, Ye-Hang Yin, Ye Wu, Yi Qian, Ying-Da Tao, Yi Yu, Yi-Chao Zhang, Yi-Chen Jiang, Yi-Cheng Wang, Yi-Fan Ding, Yi-Fan Shi, Yihui Peng, Yi-Feng Zhai, Yi-Jia Wu, Yi-Yu Xiong, Yi-Lun Wang, Ying He, Ying Zhou, Yi Luo, Yin-Min Zhong, Yi-Ping Wang, Yi-Song Wang, Yi-Xiang Zhang, Yi-Xiao Chen, Yi-Xuan Tan, Yi-Xuan Wei, Ying Ma, Yi-Yao Yang, Yi-Yuan Liu, Yi Cai, Yi Wei, Yi-Zhi Wang, Yong-Lu Yang, Yong-Qi Zhuo, Yong-Qiang Guo, Yong-Tong Wu, Yu Wu, Yu Zhang, Yu-Zhi Bian, Yuan Cheng, Yuan Ou, Yuan Sun, Yuan-Fan Xu, Yuan Sun, Yuan-Hao Li, Yu-Chen Liu, Yua Yao, Yu-Dong Han, Yuduan Wang, Yu-Han Wu, Yu-Hao Meng, Yuwei Zou, Yu-Kun Li, Yun-Chuan Wang, Yun-Fan Xiao, Yu Xiong, Yu-Peng Chen, Yu Cao, Yu-Qian Wang, Yu-Qing Chen, Yu-Shun Zhang, Yu-Tong Lin, Yue-Tu Xiao, Yu-Xian Gu, Yu-Xiang Chen, Yu-Xiang Huang, Yu-Xiang Luo, Yu-Mei You, Yu-Xin Chen, Yu-Xin Xiang,

Yuxuan Liu, Yuxuan Zhou, Yu-Yang Zhou, Yu-Zhe Guo, Yuzhi Huang, Yu-Zhuo Bai, Z. Zy, Zan-Lin Ni, Ze-Hao Wang, Ze-Hua Zhao, Zehui Ren, Ze Zhao, Zhangli Sha, Zhan Wang, Zhao Zhang, Zhao Du, Zhe Fu, Zhe-An Xu, Zhen-Da Xie, Zheng-Guo Liu, Zhengyan Zhang, Zhen-Hua Dong, Zhe-Wen Hao, Zhibo Wang, Zhibin Gou, Zhi Ma, Zhi-Hao Li, Zhi-Hong Shao, Zhi-Hua Huang, Zhi-Jie Li, Zhi Lu, Zhi-Xiang Huang, Zhi-Xuan Chen, Z. Pan, Zhi-Yu Wu, Zhi-Zhou Ren, Zhuan-Di He, Zhu Li, Zhu-Ping Zhang, Zian Xu, Zi-Hao Wang, Zi-Hui Gu, Zijia Zhu, Zi-Li Zhang, Zi-Lin Li, Z. Hou, Z. Lyu, Zi-Qiao Wang, Zi-Wei Xie, Zi-Yao Zhang, Ziyun Gao, Zi-Zheng Pan, Zong-Lin Li, Zong-Xiang Yao, Zui Chen, Zuo-Fan Wu, Chen Ling, Cheng-Yu Hou, Chong Chen, D. Li, Di Qi, Dong-Li Ji, Fang-Chen Wei, Fang Xia, Fei Xie, Fei-Yi Tan, Hai-Long Guo, Hai-Yan Zhai, Hui Zhou, Hui-Hui Tan, Hui-Jie Li, Jia Luo, Jia Song, Jia-Lu Cai, Jian Liang, Jiang-Ting Zhou, Jia-Qi Gao, Jia-Yi Shao, Jie Chen, Jie Yang, Jin Chen, Jing-De Zhang, Jing-Zi Zhou, Jin-Qiang Wang, Jin-Ya Liu, Jin-Zhao Sun, Jun-Hua Ling, Jun-Min Zheng, Kai-Cheng Yang, Le Su, Lei Xia, Li Ding, Lin Zhuo, Lin Ma, Lin Zhu, Li-Yu Cai, Lu Yao, Min Zhang, Meng Li, Miao-Fa Lin, Miao Wang, Min Zhang, Ming-Ming Li, Ming-Ming Wang, Ming-Ze Yin, Min-Min Han, Nan Cao, Ning-Xin Ma, Pan Wang, P. Lin, Peng Sun, Peng Zhang, Qian Ying, Qian-Jiang Xiang, Qiao Wang, Qin Mao, Qian Jiang, Rong Jin, Ru-Yi Chen, Shan Tao, Shang-Mian Sun, Shao-Qing Wu, Shi-Chao Zou, Si-Xuan Lei, Tian-Yang Zhang, Tian-Yu Sun, Ting-Ting Yin, Wang Xiao, Wei An, Wei Li, Wei Wang, Wei-Wei Lin, Wenqing Hou, X. Lin, Xiang-Fei Meng, Xiang Huang, Xiao-Jue-Peng-Xiao-Jue Peng, Xiao-Qian Li, Xiao-Ting Zhang, Xiao-Wen Sun, Xiaoxiang Wang, Xiaochang Ye, Xin-Rou Zhang, Xin-Yu Zhang, Xue-Gang Cao, Xue-Yin Chen, Ya-Nan Zhou, Yan-Hong Xu, Yao-Hua Xia, Yao Xu, Yi Shao, Yi-Hong Zhang, Yi-Lin Ma, Ying Tang, Yi Lou, Yi-Rui Chen, Y. Piao, Yi-Xuan Chen, Yong-Ju Xiong, Yue-Han Xuan, Yue Yang, Yue-Qiang Xu, Yukun Zha, Yun-Xiang Ma, Yu-Ping Lin, Yu-Ting Yan, Yu Xie, Yu Sheng, Yuxuan Zhu, Ze-Kai Zhang, Zhe Ju, Zheng Lin, Zhe Gao, Zhe Sun, Zhi-Gang Yan, ZhongChen Wu, Zi Wang, Z. Qu, Ziaho Yan, and Zi-Yi Wan. Deepseek-v4.1-flash: Pushing the limits of kv cache compression. 2026.

[49] GLM-4.5 Team Aohan Zeng, Xin Lv, Zhen-Yu Hou, Zheng-Xiao Du, Qinkai Zheng, Bin Chen, Da Yin, Chen-Di Ge, Cheng-Xing Xie, Cun-Xiang Wang, Gengzheng Pan, Hao Zeng, Hao-Ke Zhang, Haoran Wang, Huilong Chen, Jia-Jie Zhang, Jian Jiao, Jia-Qi Guo, Jingsen Wang, Jin-Cheng Du, Jinzhu Wu, Kedong Wang, Lei Li, L. Fan, L. Zhong, Ming-Dao Liu, Mingming Zhao, Pengfan Du, Qian Dong, Rui Lu, Shuang-li, S. Cao, Song Liu, Ting Jiang, Xiaodong Chen, Xiaohan Zhang, Xuancheng Huang, Xue Dong, Yabo Xu, Yao Wei, Yi-Fan An, Yi-Lin Niu, Yitong Zhu, Yuan Wen, Yukuo Cen, Yu Bai, Zhongpei Qiao, Zihan Wang, Zi-Kang Wang, Zilin Zhu, Ziqian Liu, Zixuan Li, Bojie Wang, Bo-Si Wen, Can Huang, Changpeng Cai, Chao Yu, Chen Li, Chen Huang, C. Hu, Chenhui Zhang, Chenzheng Zhu, Congfeng Yin, D. Lin, Da-Wei Yang, Dingkun Wang, Ding Ai, Er-Le Zhu, Fang-Yu Yi, Feiyu Chen, Guohong Wen, Hai-Long Sun, Hai Zhao, Haiyi Hu, Hanchen Zhang, Han Liu, Hanyu Zhang, Hao Peng, Hao Tai, Haobo Zhang, He Liu, Hongwei Wang, Hong Yan, Hong-Yu Ge, Huan Liu, Hua Chu, Jia-Ni Zhao, Jiachen Wang, Jiajing Zhao, Jia-Ming Ren, Jiapeng Wang, Jiaxin Zhang, Jia-Yi Gui, Jia yin Zhao, Jijie Li, Jing An, Jing Li, Jing Yuan, Jinhua Du, Jinxin Liu, J. Zhi, Junwen Duan, Kai Zhou, Kang Wei, Kedong Wang, Keyu Luo, Laiqiang Zhang, Lei Sha, Liang Xu, Lindong Wu, Li Ding, Lu Chen, Minghao Li, Nianyi Lin, Pankova Ta, Qiang Zou, Rong-Huang Song, Rui-Qi Yang, Shangqing Tu, Shangtong Yang, Shao-Xiang Wu, Sheng-Yan Zhang, Shijie Li, Shuang Li, Shuying Fan, Wei Qin, Wei Tian, Wei-Ning Zhang, Wen-Bo Yu, Wen-Jie Liang, Xiang-Tian Kuang, Xiang Cheng, Xiangyang Li, Xiao Yan, Xiaowei Hu, Xiaoying Ling, Xing Fan, Xi Xia, Xinyuan Zhang, Xin-Ze Zhang, Xirui Pan, Xun-Kai Zhang, Yan-Dong Wu, Yanfu Li, Yangyi Fang, Yidong Wang, Yi fan Zhu, Yi-Shuai Tan, Yi-Lin Zhou, Yiming Pan, Ying Zhang, Yinpei Su, Yi xin Geng, Yong Yan, Yonglin Tan, Yue Bi, Yu-Han Shen, Yu-Hao Yang, Yujiang Li, Yunan Liu, Yunqing Wang, Yuntao Li, Yurong Wu, Yutao Zhang, Yu Duan, Yuxuan Zhang, Ze-Xian Liu, Zhen Jiang, Zhenhe Yan, Zhe-Yu Zhang, Zhi-Xiang Wei, Zhuo-Fu Chen, Zhuo-Er Feng, Zi-Jun Yao, Ziwei Chai, Ziyuan Wang, Zuzhou Zhang, Bin Xu, Minlie Huang, Hong-Ning Wang, Juan-Zi Li, Yu-Chao Dong, and Jie Tang. Glm-5: from vibe coding to agentic engineering. *ArXiv*, abs/2602.15763, 2026.

[50] OpenAI. GPT-6 Astra System Card. System card, 2026. URL https://deploymentsafety.openai.com/gpt-6-astra.

[51] Wen-Qiang Sun, Hai-Yu Zhang, Haoyuan Wang, Jun-Tao Wu, Ze-Han Wang, Zhen-Wei Wang, Yun-Hong

Wang, Jun Zhang, Teng-Fei Wang, and Chun-Chao Guo. Worldplay: Towards long-term geometric consistency for real-time interactive world modeling. *ArXiv*, abs/2512.14614, 2025.

[52] Hao-Yi Duan, Hong-Xing Yu, Si-Ru Chen, Fei-Fei Li, and Jia-Jun Wu. Worldscore: A unified evaluation benchmark for world generation. *2025 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 01–12, 2025.

[53] Xunliang Cai, Qi-Long Huang, Zhuo-Liang Kang, Hong-Yu Li, Shi-Jun Liang, Li-Ya Ma, Si-Yu Ren, Xiao-Ming Wei, Ri-Chu Xie, and Tong Zhang. Longcat-video technical report. *ArXiv*, abs/2510.22200, 2025.

[54] Adam Polyak, Amit Zohar, Andrew Brown, Andros Tjandra, Animesh Sinha, Ann Lee, Apoorv Vyas, Bo-Wen Shi, Chih-Yao Ma, Ching-Yao Chuang, David Yan, Dhruv Choudhary, Ding-Kang Wang, Geet Sethi, Guan Pang, Haoyu Ma, Ishan Misra, Ji Hou, Jialiang Wang, K. Jagadeesh, Kunpeng Li, Luxin Zhang, Mannat Singh, Mary Williamson, Matt Le, Matthew Yu, Mitesh Kumar Singh, Peizhao Zhang, Peter Vajda, Quentin Duval, Rohit Girdhar, Roshan Sumbaly, Sai Saketh Rambhatla, Sam S. Tsai, S. Azadi, Samyak Datta, Sanyuan Chen, Sean Bell, Sharadh Ramaswamy, Shelly Sheynin, Siddharth Bhattacharya, Simran Motwani, Tao Xu, Tianhe Li, Tingbo Hou, Wei-Ning Hsu, Xi Yin, Xiaoliang Dai, Y. Taigman, Yaqiao Luo, Yen-Cheng Liu, Yi-Chiao Wu, Yue Zhao, Yuval Kirstain, Zecheng He, Zi-Jian He, Albert Pumarola, Ali K. Thabet, Artsiom Sanakoyeu, Arun Mallya, Baishan Guo, Boris Araya, B. Kerr, Carleigh Wood, Ce Liu, Cen Peng, Dimitry Vengertsev, E. Schonfeld, Elliot Blanchard, Felix Juefei-Xu, F. Nord, Jeff Liang, John Hoffman, J. Kohler, Kaolin Fire, Karthik Sivakumar, Lawrence Chen, Li-Cheng Yu, Lu-Ya Gao, Markos Georgopoulos, Rashel Moritz, S. Sampson, Shikai Li, Simone Parmeggiani, S. Fine, Tara Fowler, Vladan Petrovic, and Yuming Du. Movie gen: A cast of media foundation models. 2024.

[55] Ang Wang, Bao-Le Ai, Bin Wen, Chao-Jie Mao, Chen-Wei Xie, Di Chen, Feiwu Yu, Hai-Ming Zhao, Jian-Xiao Yang, Jianyuan Zeng, Jiayu Wang, Jing-Feng Zhang, Jing-Ren Zhou, Jinkai Wang, Ji-Xuan Chen, Kai Zhu, Kang Zhao, Keyu Yan, Liang-Hua Huang, Xiaofeng Meng, Ningying Zhang, Pandeng Li, Ping Wu, Rui-Hang Chu, Rui Feng, Shiwei Zhang, Si-Yang Sun, Tao Fang, Tianxing Wang, T. Gui, Ting-Yu Weng, Tong Shen, Wei Lin, Wei Wang, Wei Wang, Wen-Chao Zhou, Wente Wang, Wen Shen, Wenyuan Yu, Xianzhong Shi, Xiaomin Huang, Xin Xu, Yan Kou, Yan-Mei Lv, Yifei Li, Yi-Jing Liu, Yiming Wang, Yingya Zhang, Yi-Tong Huang, Yong Li, You Wu, Yu Liu, Yulin Pan, Yun Zheng, Yuntao Hong, Yupeng Shi, Yutong Feng, Zeyinzi Jiang, Zhen Han, Zhi-Fan Wu, and Ziyu Liu. Wan: Open and advanced large-scale video generative models. *ArXiv*, abs/2503.20314, 2025.

[56] Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-Rong Wen, and Chongxuan Li. Large language diffusion models. *ArXiv*, abs/2502.09992, 2025.

[57] Tri Dao and Albert Gu. Transformers are ssms: Generalized models and efficient algorithms through structured state space duality. *ArXiv*, abs/2405.21060, 2024.

[58] Qwen Team. Qwen3.6-27B: Flagship-level coding in a 27B dense model, April 2026. URL https://qwen.ai/blog?id=qwen3.6-27b.

[59] Peiyuan Zhang, Matthew Noto, Wenxuan Tan, Chengquan Jiang, Will Lin, Wei Zhou, and Hao Zhang. Attn-QAT: 4-Bit Attention With Quantization-Aware Training. *arXiv preprint arXiv:2603.00040*, 2026. URL https://arxiv.org/abs/2603.00040.

[60] Li-Wen Chang, Wenlei Bao, Qi Hou, Chengquan Jiang, Ningxin Zheng, Yinmin Zhong, Xuanrun Zhang, Zuquan Song, Chengji Yao, Ziheng Jiang, Haibin Lin, Xin Jin, and Xin Liu. FLUX: Fast Software-based Communication Overlap On GPUs Through Kernel Fusion. *arXiv preprint arXiv:2406.06858*, 2024. URL https://arxiv.org/abs/2406.06858.

[61] Shulai Zhang, Ningxin Zheng, Haibin Lin, Ziheng Jiang, Wenlei Bao, Chengquan Jiang, Qi Hou, Weihao Cui, Size Zheng, Li-Wen Chang, Quan Chen, and Xin Liu. Comet: Fine-grained Computation-communication Overlapping for Mixture-of-Experts. *arXiv preprint arXiv:2502.19811*, 2025. URL https://arxiv.org/abs/2502.19811.

[62] Jiwoo Hong, Noah Lee, and James Thorne. ORPO: Monolithic Preference Optimization without Reference Model. *arXiv preprint arXiv:2403.07691*, 2024. URL https://arxiv.org/abs/2403.07691.

[63] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. *arXiv preprint arXiv:2306.13649*, 2023. URL https://arxiv.org/abs/2306.13649.

[64] Tianzhu Ye, Li Dong, Xun Wu, Shaohan Huang, and Furu Wei. On-Policy Context Distillation for Language Models. *arXiv preprint arXiv:2602.12275*, 2026. URL https://arxiv.org/abs/2602.12275.

[65] Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. ALFWorld: Aligning Text and Embodied Environments for Interactive Learning. *arXiv preprint arXiv:2010.03768*, 2020. URL https://arxiv.org/abs/2010.03768.

[66] Quanhao Li, Junqiu Yu, Kaixun Jiang, Yujie Wei, Zhen Xing, Pandeng Li, Ruihang Chu, Shiwei Zhang, Yu Liu, and Zuxuan Wu. DiffusionOPD: A Unified Perspective of On-Policy Distillation in Diffusion Models. *arXiv preprint arXiv:2605.15055*, 2026. URL https://arxiv.org/abs/2605.15055.

[67] Patrick Esser, S. Kulal, A. Blattmann, Rahim Entezari, Jonas Muller, Harry Saini, Yam Levi, Dominik Lorenz, Axel Sauer, Frederic Boesel, Dustin Podell, Tim Dockhorn, Zion English, Kyle Lacey, Alex Goodwin, Yannik Marek, and Robin Rombach. Scaling rectified flow transformers for high-resolution image synthesis. pages 12606–12633, 2024.

[68] Yuval Kirstain, Adam Polyak, Uriel Singer, Shahbuland Matiana, Joe Penna, and Omer Levy. Pick-a-pic: An open dataset of user preferences for text-to-image generation. *ArXiv*, abs/2305.01569, 2023.

[69] Zeyue Xue, Jie Wu, Yu Gao, Fangyuan Kong, Lingting Zhu, Mengzhao Chen, Zhiheng Liu, Wei Liu, Qiushan Guo, Weilin Huang, and Ping Luo. DanceGRPO: Unleashing GRPO on Visual Generation. *arXiv preprint arXiv:2505.07818*, 2025. URL https://arxiv.org/abs/2505.07818.

[70] Yuhang Ma, Yunhao Shui, Xiaoshi Wu, Keqiang Sun, and Hongsheng Li. HPSv3: Towards Wide-Spectrum Human Preference Score. *arXiv preprint arXiv:2508.03789*, 2025. URL https://arxiv.org/abs/2508.03789.

[71] DeepSeek-AI. Deepseek-v3.2: Pushing the frontier of open large language models. *ArXiv*, abs/2512.02556, 2025.

[72] DeepSeek-AI, Anyi Xu, Bang Lin, Bing Xue, Bing-Li Wang, Bin Xu, Bo Wu, Bowei Zhang, Chao Lin, Chengyao Dong, Chen Ling, Cheng-Da Lu, Chen Zhao, Chengqi Deng, C. Hou, Chen Xu, Chen-Ze Shao, Chong Ruan, Conner Sun, Da-Mai Dai, Daya Guo, De-Bin Yang, De-Li Chen, Dong-Hui Li, Dong-Li Ji, Erhang Li, Fang-Chen Wei, Fangyun Lin, Fang Yuan, Fei Xia, Fucong Dai, Guangbo Hao, Guan-Ting Chen, Guo Cao, Guo-Hui Meng, Guo-Wei Li, Han Yu, Han Zhang, Han-Wei Xu, Hao Li, Hao Liang, Hao-Ling Zhang, Hao-Ming Luo, Hao-Ran Wei, Hao Yuan, Hao-Wei Zhang, Hao-Wen Luo, Hao Chen, Hao-Zhe Ji, Heng Zhang, Honghui Ding, Hong-Xuan Tang, Huan-Qi Cao, Hua-Zuo Gao, Hui Qu, Hui Zeng, J. Yang, J.-Q. Zhu, Jian-Ming Luo, Jia Song, Jia Yu, Jialiang Huang, Jia-Lu Cai, Jian Liang, Jiang-Ting Zhou, Jiasheng Ye, Jia-Shi Li, Jia-Xin Xu, Jie-Wen Hu, Jie Yang, Jin Chen, Jin Yan, Jing-Chang Chen, Jing-Li Zhou, Jing Xiang, Jing-Yang Yuan, Jing Cheng, Jing-Zi Zhou, Jin-Hua Zhu, Ji-Ping Yu, Joseph Sun, Junliang Ran, Jun Jiang, Jun-Jie Qiu, Jun-Long Li, Junming Zheng, Jun-Mei Song, Kai Dong, Kaige Gao, Kang Guan, Kexing Zhou, Ke-Zhao Huang, Kuai Yu, Le-An Wang, Lecong Zhang, Lei Wang, Leyi Xia, L. Zhang, Liang Zhao, Li-Hua Guo, Lin Luo, Lin Ma, Linyan Zhu, Li-Tong Wang, Li-Yu Cai, Li-Yue Zhang, Long-Hao Chen, Mingxi Di, M.-Y. Xu, Max Mei, Miaojun Wang, Ming-Chuan Zhang, Ming-Hua Zhang, Minghui Tang, Ming-Ming Li, Mingxu Zhou, Min Han, Ning Wang, Pan Huang, Pan-Pan Wang, Pei-Xin Cong, Pei-Yi Wang, Peng Zhang, Qiancheng Wang, Qi-Hao Zhu, Qing-Yang Li, Qin-Yu Chen, Qiu-Shi Du, Qi Jiang, Rui Tian, Rui-Fan Xu, Rui-Jie Lu, Rui-Ling Xu, Ruiqi Ge, Rui-Song Zhang, Rui-Zhe Pan, Runji Wang, Run-Yuan Chen, Run-Qiu Yin, Run-Xin Xu, Ruo-Han Shen, Ruo-Yu Zhang, Ru-Yi Chen, Sh. L. Liu, Shanghao Lu, Shang-Mian Sun, Shang-Yan Zhou, Shanhuang Chen, Shaofei Cai, Shao Nie, Shao-Ping Wu, Shao-Yuan Chen, Sheng-Ding Hu, Sheng-Yu Liu, Shi-Qiang Hu, Shi-Rong Ma, Shi-Yu Wang, Shuiping Yu, Shun-Feng Zhou, Shuting Pan, Shuang Yu, Song-Yang Zhou, Tao Ni, Tao Yun, Tian Jin, Tian-Hong Pei, Tian Ye, Tian-Le Lin, Tian Ji, Tianyi Cui, Tianyuan Yue, Tingting Yu, Tu-Tu Wang, W. Zhang, WL Xiao, Wang-Ding Zeng, Wei An, Weilin Zhao, Wen Liu, W. Liang, Wenjie Pang, Wenjing Luo, Wen-Jin Yao, Wen-Jun Gao, Wenkai Yang, Wenlve Huang, Wenqing Hou, Wen-Tao Zhang, Wen-Ting Ma, Xi Gao, Xiang He,

Xiang-Wen Wang, Xianzu Wang, Xiao Bi, Xiao-Dong Liu, Xiao-Han Wang, Xiao-Kang Chen, Xiao-Kang Zhang, X. Nie, Xiaowen Sun, Xiaoxiang Wang, Xin Cheng, Xin Liu, Xin Xie, Xing-Chao Liu, Xing-Chen Liu, Xing-Kai Yu, Xing-You Li, Xin-Yu Yang, Xin-Yu Zhang, Xu Chen, Xuan-Yu Wang, Xuecheng Su, Xue-Yin Chen, Xuheng Lin, Xu Fu, Y.-C. Yan, Y.-Q. Wang, Y.-W. Ma, Yan-Fen Luo, Yang Zhang, Yan-Hong Xu, Yan-Ru Ma, Yan-Wen Huang, Yao Li, Yao Xu, Yao Zhao, Yao-Feng Sun, Yao-Hui Wang, Yi Qian, Yingxia Shao, Yi Yu, Yi-Chao Zhang, Yi-Fan Ding, Yi-Fan Shi, Yi-Jia Wu, Yi-Yu Xiong, Yi-Lin Ma, Ying He, Ying Tang, Ying Zhou, Yi Luo, Yin-Min Zhong, Y. Piao, Yi-Song Wang, Yi-Xiang Zhang, Yi-Xiao Chen, Yi-Xuan Tan, Yixuan Wei, Yi-Yang Ma, Yi-Yuan Liu, Yong-Lu Yang, Yong-Qiang Guo, Yong-Tong Wu, Yu Wu, Yu-Kun Li, Yuan Cheng, Yuan Ou, Yuan-Fan Xu, Yuan-Hao Li, Yuduan Wang, Yue-Han Yang, Yue-Qiang Xu, Yu-Han Wu, Yu Meng, Yu Zou, Yukun Zha, Yu Xiong, Yu-Peng Chen, Yu-Ping Lin, Yu Cao, Yu-Qian Wang, Yu-Shun Zhang, Yu-Ting Yan, Yu-Tong Lin, Yu-Xian Gu, Yu-Xiang Luo, Yu-Mei You, Yuxuan Liu, Yuxuan Zhou, Yu-Yang Zhou, Yu-Zhen Huang, Z.-F. Wu, Ze-Hao Wang, Ze Zhao, Zehui Ren, Ze-Kai Zhang, Zhangli Sha, Zhe Fu, Zhe Ju, Zhe-An Xu, Zhen-Da Xie, Zheng-Yan Zhang, Zhe Gao, Zhe-Wen Hao, Zhibin Gou, Zhi-Cheng Ma, Zhi-Gang Yan, Zhi-Hong Shao, Zhi-Xiang Huang, Zhi-Xuan Chen, Zhi-Yu Wu, Zhi-Zhou Ren, Zhong-Yuan Wu, Zhu Li, Zhu-Ping Zhang, Zian Xu, Zi-Hao Wang, Z. Qu, Zihui Gu, Zijia Zhu, Zi-Lin Li, Zi-Peng Zhang, Zi-Wei Xie, Zi-Yi Gao, Zi-Yi Wan, Zi-Zheng Pan, and Zong-Xiang Yao. Deepseek-v4: Towards highly efficient million-token context intelligence. 2026.

[73] Woosuk Kwon, Zhuo-Han Li, Si-Yuan Zhuang, Ying Sheng, Lian-Min Zheng, C. Yu, Joseph E. Gonzalez, Haotong Zhang, and Ion Stoica. *Efficient Memory Management for Large Language Model Serving with PagedAttention*. 2023.

[74] Huiqiang Jiang, Yucheng Li, Chengruidong Zhang, Qianhui Wu, Xufang Luo, Surin Ahn, Zhenhua Han, Amir H. Abdi, Dongsheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. MInference 1.0: Accelerating Pre-filling for Long-Context LLMs via Dynamic Sparse Attention. *arXiv preprint arXiv:2407.02490*, 2024. URL https://arxiv.org/abs/2407.02490.

[75] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, A. Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, A. Hartshorn, Ao-Bo Yang, Archi Mitra, A. Sravankumar, A. Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurélien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Rozière, B. Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, C. Touret, Chunyang Wu, C. Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, D. Song, Danielle Pintz, D. Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab Albadawy, E. Lobanova, Emily Dinan, E. Smith, Filip Radenovic, Francisco (Paco) Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, G. Nail, G. Mialon, Guanglong Pang, Guillem Cu-curell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Iliyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, J. van der Linde, Jennifer Billock, Jenny Hong, Jen-Ya Lee, Jeremy Fu, Jianfeng Chi, Jian-Yu Huang, Jia-Wen Liu, Jie Wang, Jie-Cao Yu, Joanna Bitton, J. Spisak, Jongsoo Park, Joseph Rocca, Joshua John-stun, Joshua Saxe, Ju-Qing Jia, Kalyan Vasuden Alwala, Karthik Prasad, K. Upasani, Kate Plawiak, Ke Li, K. Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuen ley Chiu, Kunal Bhalla, Kushal Lakhotia, Lauren Rantala-Yeary, L. van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, L. Malo, Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, M. Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri S. Chatterji, Ning Zhang, Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Peng-Chuan Zhang, Peng-Wei Li, P. Vasić, Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Pu-Xin Xu, Qing He, Qing-Xiao Dong, Ragavan Srinivasan, R. Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic, R. Raileanu, R. Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ron nie Polidoro, Roshan Sumbaly, Ross

Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, S. Kim, Sergey Edunov, Shao-Liang Nie, Sharan Narang, S. Raparthy, Sheng Shen, Sheng-Ye Wan, Shruti Bhosale, Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, S. Collot, Suchin Gururangan, S. Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Vogeti, Vítor Albiero, Vladan Petrovic, Wei-Wei Chu, Wen-Han Xiong, Wen-Yin Fu, Whit ney Meers, X. Martinet, Xiao-Dong Wang, Xiao-Fang Wang, X. Tan, Xi-De Xia, Xinfeng Xie, Xuchao Jia, Xue-Wei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yiqian Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yu-Ning Mao, Zacharie Delpierre Coudert, Zhe Yan, Zheng-Xing Chen, Zoe Papakipos, Aaditya K. Singh, Aayushi Srivas-tava, Abha Jain, A. Kelsey, Adam Shajnfeld, Adi Gangidi, Adolfo Victoria, Ahuva Goldstand, A. Menon, Ajay Sharma, Alex Boesen-berg, Alexei Baevski, Allie Fein-stein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, A. Goyal, Aparajita Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin R. Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Po-Yao (Bernie) Huang, Beth Loyd, B. de Paola, Bhargavi Paranjape, Bing Liu, Bo Wu, B. Ni, Braden Hancock, Bram Wasti, Brandon Spence, B. Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Chang-Han Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer, Cynthia Gao, D. Civin, D. Beaty, Daniel Kreymer, Shang-Wen Li, David Adkins, David Xu, D. Testuggine, Delia David, D. Parikh, Diana Liskovich, Didem Foss, Dingkang Wang, D. Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Lê, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smoth-ers, Fei Sun, F. Kreuk, Feng Tian, Filippos Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank J. Kanayet, F. Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, G. Zhang, G. Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Han Zha, Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry As-pegren, Hunter Goldman, Hong-Yuan Zhan, Ibrahim Damlaj, Igor Molybog, I. Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, J. Weissman, James Geboski, J. Kohli, Janice Lam, J. Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, J. Chan, Jenny Zhen, J. Reizenstein, Jeremy Teboul, J. Zhong, Jian Jin, Jing-Yi Yang, Joe Cummings, J. Carvill, Jon Shepard, J. McPhie, J. Torres, Josh Ginsburg, Junjie Wang, Kaixing(Kai) Wu, U. KamHou, Karan Saxena, Kartikay Khandelwal, Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Ke-Qian Li, K. Jagadeesh, Kun Huang, Kunal Chawla, Kyle Huang, Lai-Lin Chen, Lakshya Garg, A. Lavender, Leandro Silva, Lee Bell, Lei Zhang, Liang-Peng Guo, Li-Cheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus, Matan Hasson, Matthew Lennie, Matthias Reso, M. Groshev, M. Naumov, Maya Lathi, Meghan Keneally, Miao Liu, M. Seltzer, Michal Valko, Michelle Re-strepo, Mihir M. Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark, M. Macey, M. Wang, Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, N. Santhanam, Natascha Parks, Natasha White, Navy ata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, O. Hart, Omkar Salpekar, Ozlem Kalinli, Parkin Kent, Parth Parekh, P. Saab, Pavan Balaji, Pe dro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollár, Polina Zvyagina, Prashant Ratanchandani, P. Yuvraj, Qian Liang, Rachad Alao, R. Rodriguez, R. Ayub, Raghotham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara Chugh, Sara Hunt, Sargun Dhillon, S. Sidorov, S. Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto, Sharadh Ramaswamy, Shaun Lind-say, Sheng Feng, Shenghao Lin, Shengxin Zha, Shishir G. Patil, S. Shankar, Shu-Qiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala, Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, S. Gupta, Summer Deng, Sung-Bae Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara Best, Thilo Koehler, Thomas S. Robinson, Tianhe Li,

Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked, Varun Vontimitta, Victoria O. Ajayi, Victoria Montanez, Vijai Mohan, Vinay Kumar, Vishal Mangla, Vlad Ionescu, V. Poenaru, Vlad T. Mihailescu, V. Ivanov, Wei Li, Wenchen Wang, Wen-Wen Jiang, W. Bouaziz, Will Constable, Xia Tang, Xiao-Jian Wu, Xiao-Lan Wang, Xi-Lun Wu, Xin-Bo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu, Ye Jia, Ye Qi, Ye-Rong Li, Yi-Lin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu Wang, Yu Zhao, Yuchen Hao, Yun-Di Qian, Yun-Lu Li, Yuzi He, Zach Rait, Z. DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhen-Yu Yang, Zhi-Wei Zhao, and Zhi-Yu Ma. The llama 3 herd of models. 2024.

[76] Yuhan Liu, Yihua Cheng, Jiayi Yao, Yuwei An, Xiaokun Chen, Shaoting Feng, Yuyang Huang, Samuel Shen, Rui Zhang, Kuntai Du, and Junchen Jiang. LMCache: An Efficient KV Cache Layer for Enterprise-Scale LLM Inference. *arXiv preprint arXiv:2510.09665*, 2025. URL https://arxiv.org/abs/2510.09665.

[77] Alessio Devoto, Maximilian Jeblick, and Simon Jégou. Expected Attention: KV Cache Compression by Estimating Attention from Future Queries Distribution. *arXiv preprint arXiv:2510.00636*, 2025. URL https://arxiv.org/abs/2510.00636.

[78] Jang-Hyun Kim, Dongyoon Han, and Sangdoo Yun. Fast KVzip: Efficient and Accurate LLM Inference with Gated KV Eviction. *arXiv preprint arXiv:2601.17668*, 2026. URL https://arxiv.org/abs/2601.17668.

[79] Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE-3: Scaling up Inference Acceleration of Large Language Models via Training-Time Test. *arXiv preprint arXiv:2503.01840*, 2025. URL https://arxiv.org/abs/2503.01840.

[80] Chengyue Wu, Hao Zhang, Shuchen Xue, Zhijian Liu, Shizhe Diao, Ligeng Zhu, Ping Luo, Song Han, and Enze Xie. Fast-dLLM: Training-free Acceleration of Diffusion LLM by Enabling KV Cache and Parallel Decoding. *arXiv preprint arXiv:2505.22618*, 2025. URL https://arxiv.org/abs/2505.22618.

[81] Tianrui Feng, Zhi Li, Shuo Yang, Haocheng Xi, Muyang Li, Xiuyu Li, Lvmin Zhang, Keting Yang, Kelly Peng, Song Han, Maneesh Agrawala, Kurt Keutzer, Akio Kodaira, and Chenfeng Xu. StreamDiffusionV2: A Streaming System for Dynamic and Interactive Video Generation. *arXiv preprint arXiv:2511.07399*, 2025. URL https://arxiv.org/abs/2511.07399.

[82] Hao Kang, Ziyang Li, Weili Xu, Xinyu Yang, Yinfang Chen, Junxiong Wang, Beidi Chen, Tushar Krishna, Chenfeng Xu, and Simran Arora. ThunderAgent: A Simple, Fast and Program-Aware Agentic Inference System. *arXiv preprint arXiv:2602.13692*, 2026. URL https://arxiv.org/abs/2602.13692.

[83] Zhou Wang, A. Bovik, H. Sheikh, and Eero P. Simoncelli. Image quality assessment: from error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13:600–612, 2004.

[84] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 586–595, 2018.

[85] Zi-Qi Huang, Yi-Nan He, Jia-Shuo Yu, Fan Zhang, Chenyang Si, Yu-Ming Jiang, Yuan-Han Zhang, Tianxing Wu, Qin Jin, Nattapol Chanpaisit, Yaohui Wang, Xin-Yuan Chen, Li-Min Wang, Da-Hua Lin, Yu Qiao, and Ziwei Liu. Vbench: Comprehensive benchmark suite for video generative models. *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 21807–21818, 2023.

A Implementation Details

A.1 Experimental Setup

We follow the experimental setup in Section 4, using Claude Code version 2.1.287 and Codex CLI version 0.159.3. All models are accessed through the OpenRouter API. Each model is evaluated on all 25 tasks, with 12-hour budgets for operators and 24-hour budgets for modules and systems. Public tests support development; the Judge environment scores submitted solutions on private tests after checking task correctness and quality requirements.

Hardware. We deploy the Work environment and the Judge environment in separate containers and use NVIDIA H100 GPUs with 80 GB of memory each. Each environment receives one GPU by default, two for training modules and systems and StreamDiffusionV2, and eight for Flux/Comet MoE Primitive. For inference-module tasks, the Work environment receives one GPU, while the Judge environment receives seven GPUs for kvpress KV Eviction and KVzip KV Scoring and six for Quant-VideoGen KV Codec. For Concurrent Agent Serving, the Work environment and the Judge environment run in CPU-only containers connected to a dedicated two-GPU vLLM backend serving Qwen3.6-27B [58] with tensor parallelism.

GPU budgets. GPU budgets include the Work environment, the continuously allocated Judge environment, and any dedicated serving backend. For a run with a budget of T_{h} hours, the GPU-hour budget is

$$B_{\mathrm{GPU}} = T_{\mathrm{h}} \left(N_{\mathrm{Work}} + N_{\mathrm{Judge}} + N_{\mathrm{backend}}\right). \tag{5}$$

Here N_{Work}, N_{Judge}, and N_{backend} are the numbers of GPUs allocated to the Work environment, the Judge environment, and a dedicated serving backend, respectively.

Across the 150 main-experiment runs, the total budget is 11,376 GPU-hours: 4,320 for the Work environment, 6,768 for the Judge environment, and 288 for the serving backend.

The long-horizon vs. short-horizon comparison (Section 4.3) reuses 12 main-experiment runs and adds 72 short-horizon runs across Astra and Claude, requiring 2,592 additional GPU-hours.

Final performance. Most Infrathon tasks use performance metrics without fixed upper bounds.

For task τ and model m, let $\mathcal{M}_{\tau,m}(t)$ index submissions with a valid private test result, a finite private test score, and elapsed time $t_i \le \min(t, T)$. Write $u_{\tau,m,i}$ for the private test score oriented so that higher is better. The best-so-far private test score is

$$s^{\star}_{\tau,m}(t) = \max\big(\{0\} \cup \{u_{\tau,m,i} : i \in \mathcal{M}_{\tau,m}(t)\}\big). \tag{6}$$

Final performance is the best private test score reached within the task budget:

$$s^{\mathrm{final}}_{\tau,m} = s^{\star}_{\tau,m}(T). \tag{7}$$

All task scores are nonnegative. Runs without a valid scored submission remain at zero. Curves show valid scored submissions and carry the best score forward through the budget.

Rank scores. Tasks use different metrics and scales, so directly averaging their raw scores is not meaningful. We therefore obtain an overall ordering by aggregating within-task ranks with equal task weights.

Let $\rho_{\tau,m}$ be model m's rank on task τ among the $K = 6$ models, with rank 1 assigned to the highest final performance and average ranks for exact ties. For a task group G, we define

$$\mathrm{RankScore}_m(G) = \frac{100}{|G|(K-1)} \sum_{\tau \in G} (K - \rho_{\tau,m}). \tag{8}$$

This rescales the mean rank to $[0, 100]$, with 100 and 0 corresponding to first and last on every task in the group.

Run selection. Among the main-experiment runs, we select the most recently started run for each task–model pair.

Trajectory and interaction statistics. We analyze private test scores and submission times from all 150 main-experiment runs. At each budget fraction, ranking uses the best private test score submitted by that time, with average ranks for ties. Second-half improvement means that the final best-so-far score exceeds its value at the end of the first half; the score is initialized to zero. Percentage increases require a positive first-half score. Figure 7(b)–(c) uses only runs with positive final scores. In panel (b), dots represent individual runs; large markers and horizontal bars show each model's median and interquartile range, respectively.

For the plateau analysis, a record is a submission that raises the best-so-far score. Across all 150 runs, we count those with a pair of consecutive records separated by one quarter of the task budget or more. We also count runs in which the later record exceeds the preceding best by more than 1%; each run is counted once per statistic.

We count visible output tokens in agent text and tool calls with `o200k_base`, removing duplicate displays. One text block or tool invocation counts as one action.

A.2 Optimization Levels

We classify tasks by their optimization target and the code the agent is allowed to change. The level describes the task as a whole, including any kernels optimized within it.

Operator level. Operator-level tasks optimize a specified computation or data transfer with fixed input–output behavior. Typical tasks implement or accelerate GPU kernels, including prescribed fused operations. Infrathon examples include Dynamic Grouped FP8 GEMM, SageAttention, and LMCache KV Block Transfer. A solution may use one or multiple kernels.

Module level. Module-level tasks optimize a functional component or a training step within a fixed workflow. The task exposes the component's implementation or policy while fixing the surrounding training or inference pipeline. Infrathon examples include Qwen3 ORPO, Qwen3 GKD, KVzip KV Scoring, and Quant-VideoGen KV Codec.

System level. System-level tasks optimize a complete training or generation workflow, or an end-to-end service. The editable scope covers its training or decoding loop, pipeline stages, or serving orchestration. Infrathon examples include Miles Qwen3 OPD, EAGLE-3 Speculative Decoding, and Concurrent Agent Serving.

A.3 Task Descriptions

The 25 entries below summarize each task's objective, reference project or method, and scoring rule. Scores describe valid private-test results, with higher values indicating better performance. Final run performance is the best score reached within the budget (Equation 7).

Throughout, GM and AM denote geometric and arithmetic means:

$$\begin{aligned} \mathrm{GM}(x_1, \ldots, x_n) &= \left(\prod_{i=1}^{n} x_i \right)^{1/n}, \\ \mathrm{AM}(x_1, \ldots, x_n) &= \frac{1}{n} \sum_{i=1}^{n} x_i. \end{aligned} \tag{9}$$

WGM denotes a geometric mean with task-defined workload weights. For compute tasks, q_i is case i's effective throughput in TFLOP/s, and G_f is the geometric-mean case throughput within workload family f.

Attn-QAT NVFP4 Attention. Using Attn-QAT [59] in FastVideo as the reference, this task accelerates quantization-aware forward and backward attention for video models. The score is effective training throughput in TFLOP/s, accounting for activation recomputation and aggregating geometrically within and across model families.

$$\text{score} = \text{GM}_f(G_f).$$

K-grouped FP8 Wgrad. Using DeepGEMM[4] as the reference, this task accelerates grouped FP8 weight-gradient matrix multiplication for MoE training. The score measures useful matrix-multiplication throughput in TFLOP/s, aggregated geometrically within and across grouping regimes.

$$\text{score} = \text{GM}_f(G_f).$$

Flux/Comet MoE Primitive. Following FLUX [60] and Comet [61], this task accelerates token all-gather, expert scatter, and grouped matrix multiplication across GPUs. The score measures useful per-rank throughput in TFLOP/s using the slowest rank's completion time, then aggregates geometrically across communication workloads.

$$\text{score} = \text{GM}_f(G_f).$$

Qwen3 ORPO. Using TRL's ORPO implementation [62] as the reference, this task accelerates complete Qwen3-1.7B [38] preference-optimization steps, including forward/backward computation, communication, and parameter updates. We measure training throughput in completion tokens used in the loss per second. Let r_i be the median ratio of the submitted solution's throughput to the reference throughput for workload i; a score of 100 denotes equal throughput.

$$\text{score} = 100\ \text{GM}_i(r_i).$$

Qwen3 GKD. Using TRL's GKD implementation [63] as the reference, this task accelerates off-policy distillation from a frozen Qwen3-4B teacher into Qwen3-1.7B. Timing covers teacher and student computation, the distillation objective, and student updates. As in Qwen3 ORPO, let r_i be the median ratio of the submitted solution's training throughput to the reference throughput for workload i.

$$\text{score} = 100\ \text{GM}_i(r_i).$$

Miles Qwen3 OPCD. This task implements on-policy context distillation [64] in Miles. A Qwen3-1.7B student learns from a same-size teacher with privileged medical context, then answers questions without that context. Let Δq be held-out MedMCQA accuracy gained over the original model, in percentage points, and h the complete training time in hours.

$$\text{score} = \frac{\Delta q}{h}.$$

Miles Qwen3 OPD. This task implements sampled-token on-policy distillation in Miles: Qwen3-1.7B generates mathematical reasoning trajectories and learns from a frozen Qwen3-4B teacher's token scores. Let Δq be the held-out DAPO-Math accuracy gain in percentage points and h the complete training time in hours.

$$\text{score} = \frac{\Delta q}{h}.$$

Miles ALFWorld RL. This task integrates ALFWorld [65] into Miles and trains Qwen3-4B with multi-turn group-relative policy optimization on household tasks. Let Δq be the held-out game-success gain over the original policy, in percentage points, and h the complete training time in hours.

$$\text{score} = \frac{\Delta q}{h}.$$

[4] https://github.com/deepseek-ai/DeepGEMM

Miles SD3.5-M Multi-teacher OPD. Following DiffusionOPD [66], this task trains a Stable Diffusion 3.5 Medium [67] student LoRA in Miles with OCR and aesthetics teachers. Let c_r be the fraction of the original-to-teacher quality gap closed on route r (text-recognition accuracy or PickScore [68]), and h the median training-replay duration in hours. The score measures mean gap closure per replay hour.

$$\text{score} = \frac{\text{AM}_r(c_r)}{h}.$$

Miles Wan2.2 DanceGRPO. This task implements DanceGRPO [69] in Miles for Wan2.2 [55] text-to-video training with HPSv3 rewards [70], optimizing sampling, reward computation, policy updates, and weight publication. Let h be the median fixed-workload replay duration in hours. For runs meeting the held-out reward requirement, the score is replays per hour.

$$\text{score} = \frac{1}{h}.$$

Dynamic Grouped FP8 GEMM. Using DeepGEMM[5] as the reference, this task accelerates block-scaled FP8 matrix multiplication for MoE inference with varying token counts per expert. The score is useful throughput in TFLOP/s, excluding padded work and aggregating geometrically within and across storage layouts.

$$\text{score} = \text{GM}_f(G_f).$$

Sparse-MLA Prefill. Using FlashMLA[6] as the reference, this task accelerates sparse multi-head latent attention, where each query attends to selected keys. The score measures throughput over the selected attention interactions in TFLOP/s, aggregated geometrically within and across DeepSeek model configurations [71, 72].

$$\text{score} = \text{GM}_f(G_f).$$

Paged-GQA Decode. Using FlashAttention [13] and PagedAttention [73] as references, this task accelerates grouped-query attention over paged KV caches for decoding and short-query prefill. The score measures throughput over causally visible attention pairs in TFLOP/s, aggregated geometrically within and across query-length groups.

$$\text{score} = \text{GM}_f(G_f).$$

SageAttention. Using SageAttention [14] as the reference, this task accelerates quantized attention, including smoothing, quantization, softmax, and value aggregation. The score is dense-equivalent throughput in TFLOP/s, aggregated geometrically within and across video-model and causal-language-model workloads.

$$\text{score} = \text{GM}_f(G_f).$$

MInference. Using MInference [74] as the reference, this task accelerates long-context attention on Llama-3.1-8B [75] with sparse patterns while retaining attention accuracy. The score is the geometric mean of dense-equivalent throughput in TFLOP/s across workloads and model layers. Timing includes both pattern selection and attention computation.

$$\text{score} = \text{GM}_i(q_i).$$

LMCache KV Block Transfer. Using LMCache [76] as the reference, this task accelerates KV-block transfers between paged serving buffers and contiguous device or host cache storage. Let b_g be the geometric-mean bandwidth in GB/s for workload group g, measured as logical payload transferred per second. The score weights device, host, and trace-replay groups.

$$\text{score} = \text{WGM}_g(b_g).$$

[5] https://github.com/deepseek-ai/DeepGEMM
[6] https://github.com/deepseek-ai/FlashMLA

Flash-KMeans. Using Flash-KMeans [15] as the reference, this task accelerates an exact Lloyd iteration: assigning points to their nearest centroids and updating centroid means. The native metric is weighted geometric-mean latency across assignment, update, and full-iteration workloads. Let F_{eff} be the fixed weighted geometric mean of effective operation counts and t_{eff} the weighted latency in seconds. We report equivalent throughput in TFLOP/s:

$$\text{score} = \frac{F_{\text{eff}}}{10^{12} t_{\text{eff}}}.$$

Mamba-2 SSD Forward. This task accelerates the state-space-dual forward computation of Mamba-2 [57], with TileLang[7] as an implementation reference. The score is effective recurrent-work throughput in TFLOP/s, aggregated geometrically within and across model configurations.

$$\text{score} = \text{GM}_f(G_f).$$

kvpress KV Eviction. Using kvpress [77] as the reference, this task selects KV-cache positions to retain in Qwen3-8B while preserving its next-token distributions. Item fidelity is $f = \exp(-\overline{D}_{\text{KL}})$, where $\overline{D}_{\text{KL}}$ is mean KL divergence from full-cache to compressed-cache distributions on the same continuation. Let $\bar{f}_{g,\rho}$ be mean item fidelity in workload family g at eviction level ρ. The score averages across families and eviction levels:

$$\text{score} = 100\ \text{AM}_\rho\left[\text{GM}_g(\bar{f}_{g,\rho})\right].$$

KVzip KV Scoring. Using KVzip [39] and Fast KVzip [78] as references, this task scores Qwen3-8B KV entries before seeing downstream questions and retains the highest-scoring entries. Let $r_{g,\kappa}$ be the mean dataset-level compressed/full-cache answer-quality ratio in workload group g at cache budget κ. The score averages these ratios across groups and budgets:

$$\text{score} = 100\ \text{AM}_{g,\kappa}(r_{g,\kappa}).$$

Quant-VideoGen KV Codec. Using Quant-VideoGen [16] as the reference, this task optimizes a shared KV-cache codec for LongCat-Video [53] and HY-WorldPlay [51]. The score combines compression and video quality:

$$\text{score} = \text{GM}_m[R_m(1 + q_m)]. \tag{10}$$

Here R_m is video model m's compression ratio and q_m measures video quality relative to the human codec. Solutions failing correctness, quality, or runtime checks score zero. Appendix E.4 gives the metric definitions.

EAGLE-3 Speculative Decoding. Using EAGLE-3 [79] as the reference, this task accelerates Qwen3-8B speculative decoding while preserving target-model decoding behavior. Let u_c be generated-token throughput for workload c, in tokens/s. Timing includes prefill, drafting, verification, and KV-cache updates.

$$\text{score} = \text{GM}_c(u_c).$$

LLaDA Decoding. Using Fast-dLLM [80] as the reference, this task accelerates LLaDA's iterative decoding [56] while retaining math and code answer quality. Let u_c be useful completion-token throughput for workload c, in tokens/s. Useful tokens exclude EOS and padding; timing covers the complete decoding request.

$$\text{score} = \text{GM}_c(u_c).$$

[7] https://github.com/tile-ai/tilelang

StreamDiffusionV2. Using StreamDiffusionV2 [81] as the reference, this task accelerates live video-to-video streaming through scheduling, communication, and cache management. For stream i, let v_i be steady-state FPS and a_i its latency-adjustment factor, derived from the ratio of the reference's 99th-percentile input-to-output delay to the submitted solution's. The score is latency-adjusted FPS:

$$\text{score} = \text{GM}_i(v_i a_i).$$

Concurrent Agent Serving. Using ThunderAgent [82] as the reference, this task optimizes admission and scheduling for concurrent coding agents connected to a fixed Qwen3.6-27B backend. Let N_{valid} be executable, protocol-valid agent steps completed during a scored window of d minutes. The score is valid steps/min:

$$\text{score} = \frac{N_{\text{valid}}}{d}.$$

B Complete Results and Trajectories

B.1 Optimization Trajectories

Figures 10 and 11 show performance over time for all 150 runs across the 25 tasks.

Among runs with positive first-half scores, Qwen has the largest model-level median second-half gain, 10.81%, compared with 2.43% for Astra and 2.47% for Claude. Among runs with positive final scores, 64.2% of system runs reach 90% of final performance by half of the budget, compared with 78.8% of operator runs and 93.3% of module runs.

B.2 Output Token and Action Statistics by Task

Figure 12 shows visible output tokens and agent actions for each of the 150 runs.

C Long Horizon vs. Short Horizon

We compare long-horizon optimization with best-of-n test-time scaling (TTS) over independent short-horizon runs, using GPT-6 Astra and Claude Opus 5.5 on six tasks: one training and one inference task each for operators, modules, and systems. Each task–model pair receives equal GPU-hours as one T-hour run, two $T/2$-hour runs, or four $T/4$-hour runs. Here T is 12 hours for operators and 24 hours otherwise. Each setting reports the best private test score across its run histories.

Figures 13–18 show all 84 runs. Each panel shows one full-budget run, two half-budget runs, and four quarter-budget runs, aligned by elapsed time from each run's start. Runs with the same budget share their color, line style, and marker. Panel letters match Figure 8.

Claude achieves higher scores with short-horizon TTS on both operator tasks and Qwen3 ORPO, whereas Astra performs better with long-horizon optimization on all six tasks. In four task–model comparisons, short-horizon TTS leads at the end of its budget but is later overtaken by long-horizon optimization. On KVzip KV Scoring, both models trail their best two-run result at 12 hours, then surpass it at 13.90 hours for Astra and 12.67 hours for Claude.

D Benchmark Coverage and Horizon Comparison

This appendix documents the coverage and execution settings underlying Table 1. A check denotes coverage or support documented in the inspected paper or repository; a cross denotes no documented coverage. Infra only requires an exclusive focus on AI infrastructure. GPU tasks requires at least one task in the open-source

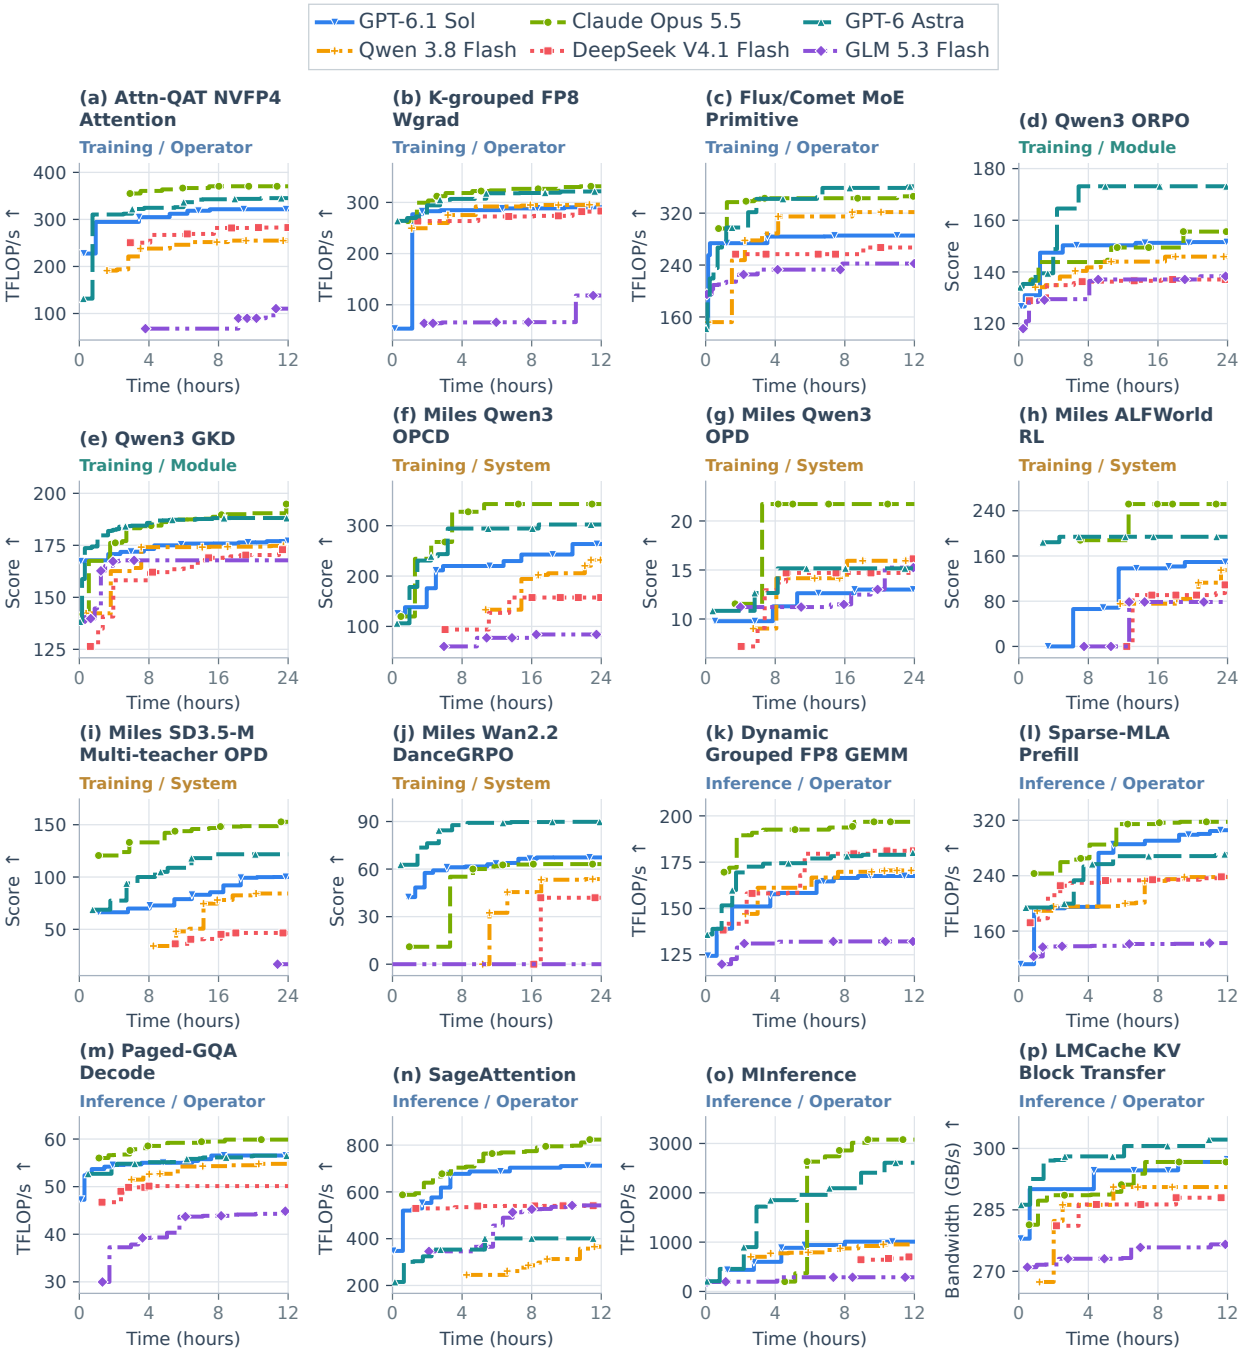

Figure 10: **Complete optimization trajectories for training tasks and inference operators.** Panels (a)–(p) cover the first 16 tasks in Appendix A.3: all 10 training tasks and the first six inference operators. Curves show performance over time.

release that explicitly uses a GPU. Long horizon requires the framework to support at least 12 hours of continuous execution on every task, with mechanisms that prevent voluntary early completion and recover

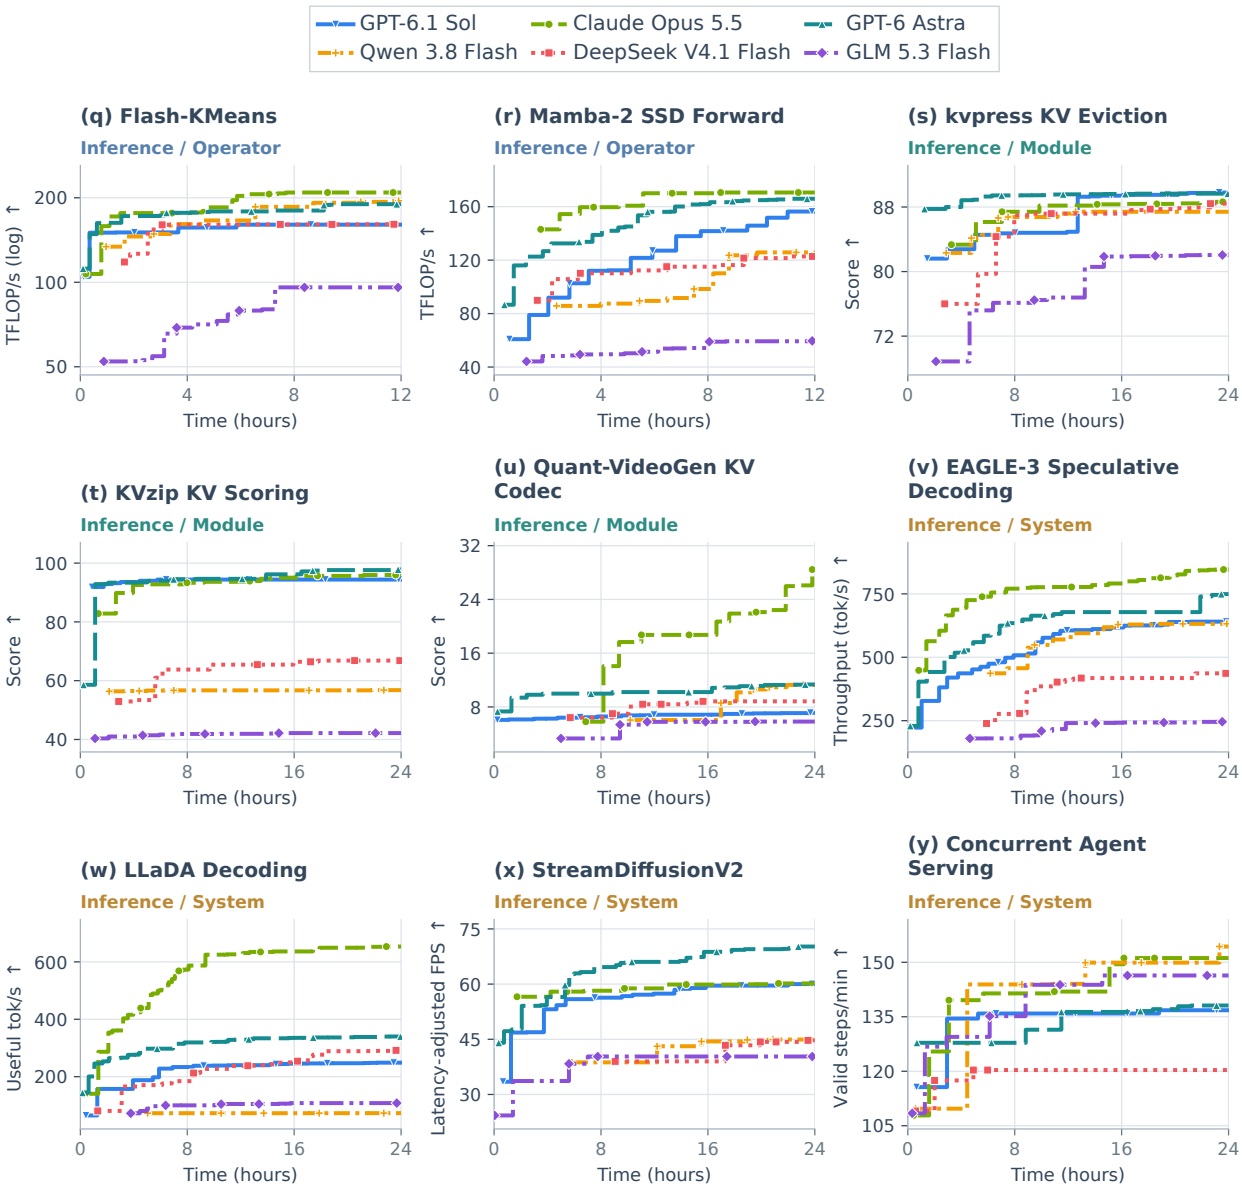

Figure 11: **Complete optimization trajectories for the remaining inference tasks.** Panels (q)–(y) continue the task order in Appendix A.3 with two operators, three modules, and four systems. Curves show performance over time.

interrupted sessions within a run.

We use the definitions in Section A.2 to assign task levels in Table 1. A check requires at least one independently specified task at that level. Components used inside a higher-level task do not count as separate tasks. The same definitions apply to non-AI tasks in general-purpose benchmarks.

Figure 12: **Visible output tokens and actions for every task–model run.** Rows follow the benchmark task order. Cell values show (a) visible output tokens in thousands and (b) agent actions. Color intensity is logarithmic within each panel.

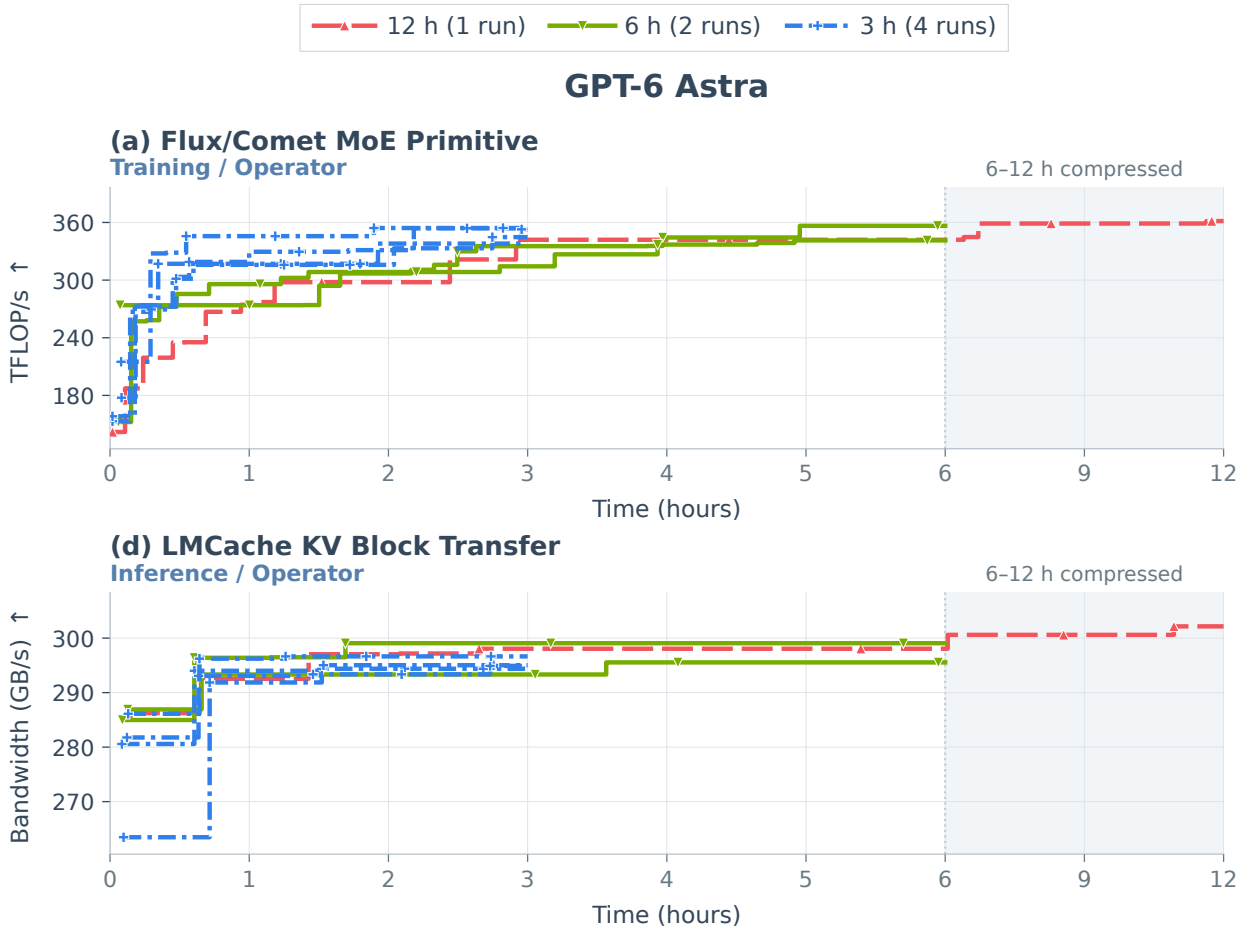

Figure 13: **Individual operator trajectories with GPT-6 Astra.** Flux/Comet MoE Primitive and LMCache KV Block Transfer each use one 12-hour run, two 6-hour runs, and four 3-hour runs. Curves show performance over time. The 6–12-hour interval is compressed threefold; ticks show actual hours.

D.1 Benchmarks in the Main Comparison

KernelBench [6]. KernelBench evaluates generated GPU implementations of PyTorch computations using correctness and speedup over PyTorch Eager. Its individual-operation and simple fusion tasks establish operator coverage. Its Level 3 tasks include independent Inception, ResNet, and Transformer blocks, with interacting branches, attention, normalization, and feed-forward components; the benchmark permits the agent to choose the kernel granularity. These establish module coverage. Level 4 tasks evaluate Hugging Face model forward calls that return logits. Its refinement study uses 10 model calls with code and execution feedback; no numerical wall-clock cap is reported.

KernelBenchX [17]. KernelBenchX evaluates correctness, speedup, and hardware efficiency across operators and GPUs. Its tasks specify individual GPU kernels, including fused computations. Its KernelAgent setting uses three workers with at most five refinement rounds each; its GEAK setting uses three iterations with four candidate solutions per round. The core KernelAgent implementation stops when any worker passes its tests or all workers exhaust their rounds. Test and seed generation also use model calls.

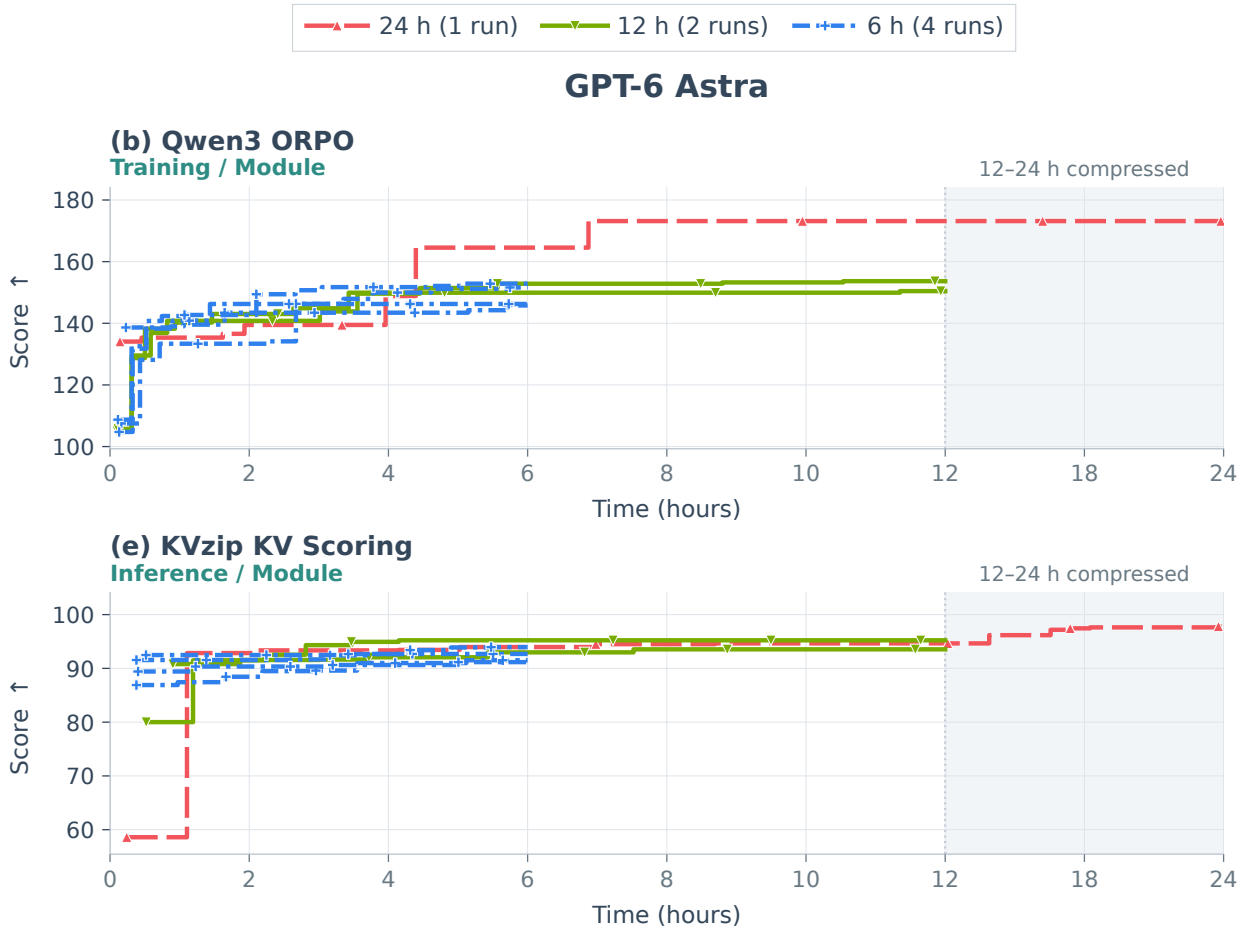

Figure 14: **Individual module trajectories with GPT-6 Astra.** Qwen3 ORPO and KVzip KV Scoring each use one 24-hour run, two 12-hour runs, and four 6-hour runs. Curves show performance over time. The 12–24-hour interval is compressed threefold; ticks show actual hours.

FlashInfer-Bench [18]. FlashInfer-Bench evaluates inference kernels on workloads derived from serving traces. The optimization targets are individual GPU kernels, including fused MoE kernels. The sequential setting uses 10 model calls with execution feedback and returns the best solution that passes its tests. No numerical wall-clock cap is reported.

RE-Bench [19]. RE-Bench compares agents and human experts on seven research-engineering tasks, with up to 8 hours per attempt. Its scope includes ML research beyond infrastructure. Its prefix-sum GPU kernel establishes operator coverage, and optimization of the complete LLM Foundry finetuning script establishes system coverage. The restricted-architecture task allows designing and training an entire masked language model, while Fix Embedding scores a repaired model checkpoint and permits further training. Runs end when the agent finishes or reaches the time limit. The paper excludes pauses for API issues from agents' time usage.

InferenceBench [7]. InferenceBench requires a running OpenAI-compatible inference service on an H100 GPU. All four scenarios evaluate complete services. It uses 2-hour main budgets and a 1/2/4/8-hour time ablation. Its runners resume sessions using the remaining time, with a 300-second grace period added to the

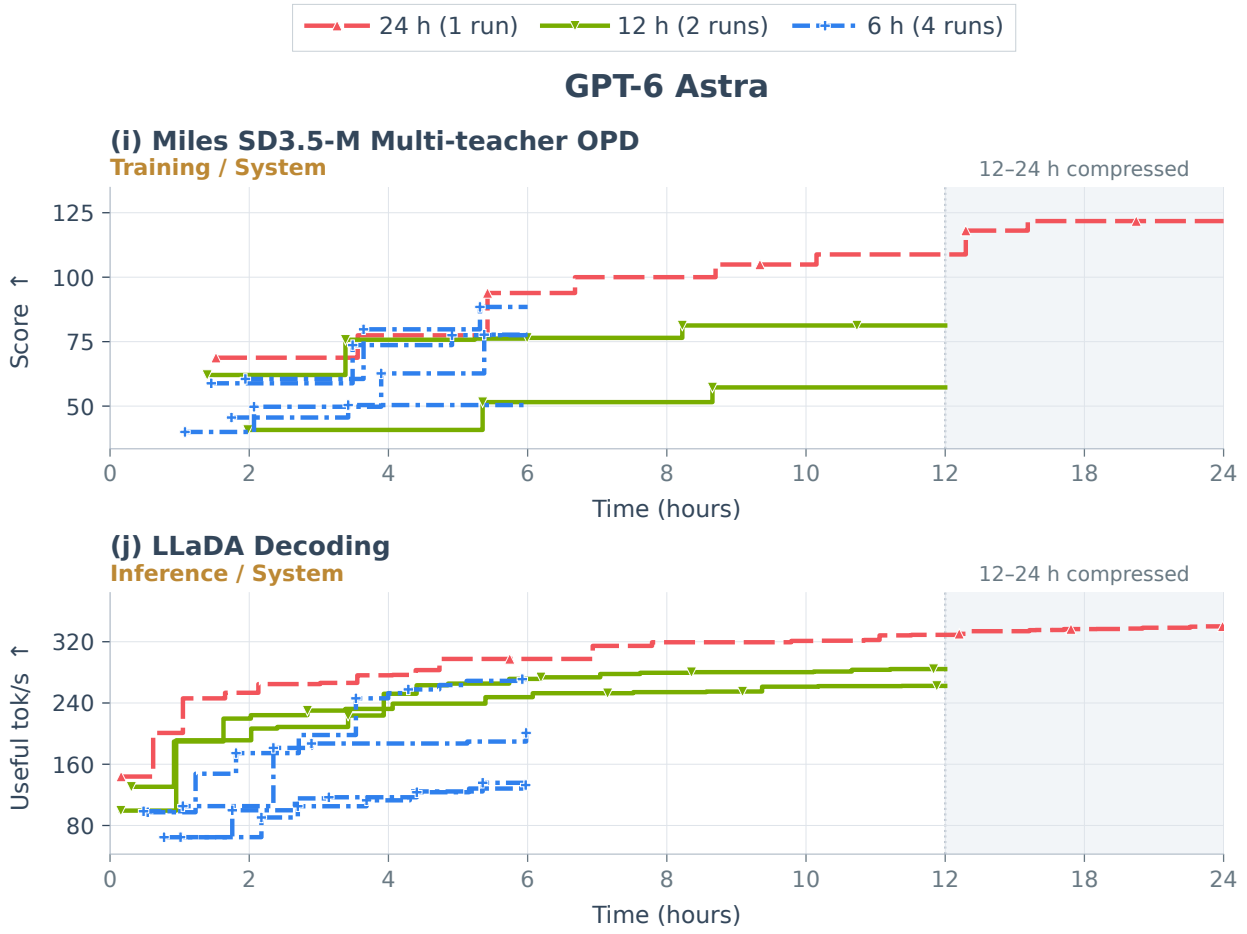

Figure 15: **Individual system trajectories with GPT-6 Astra.** Miles SD3.5-M Multi-teacher OPD and LLaDA Decoding each use one 24-hour run, two 12-hour runs, and four 6-hour runs. Curves show performance over time. The 12–24-hour interval is compressed threefold; ticks show actual hours.

configured budget. Errors or exhausted retries can end a run early.

Φ-Bench [8]. The benchmark covers kernel completion, long-horizon implementation, and end-to-end optimization. An audit of all 85 task manifests finds 43 tasks declaring GPU resources. Kernel implementations establish operator coverage. The priority-request-queue task permits changes only to the queue implementation, with exact ordering behavior and the surrounding vLLM code fixed; the memory-accounting task restricts edits to its accountant and scheduler. These establish module coverage. Complete nanoGPT training and vLLM serving tasks establish system coverage. The 85 tasks use caps of 1.5, 4, 5, or 6 hours. All 20 LHI (long-horizon implementation) tasks belong to the 1.5-hour group. The runner applies a timeout to a single agent invocation and ends the attempt when it returns.

EdgeBench [9]. EdgeBench covers science, ML, software, systems, math optimization, knowledge work, mathematics, and games. The paper's task descriptions establish operator coverage through a VLIW/SIMD kernel on a custom architecture simulator, module coverage through the EnTT graph component, and system coverage through a PostgreSQL-compatible server. These include non-AI tasks. The released official experiment configuration and container backends allocate CPU and memory resources without GPU device

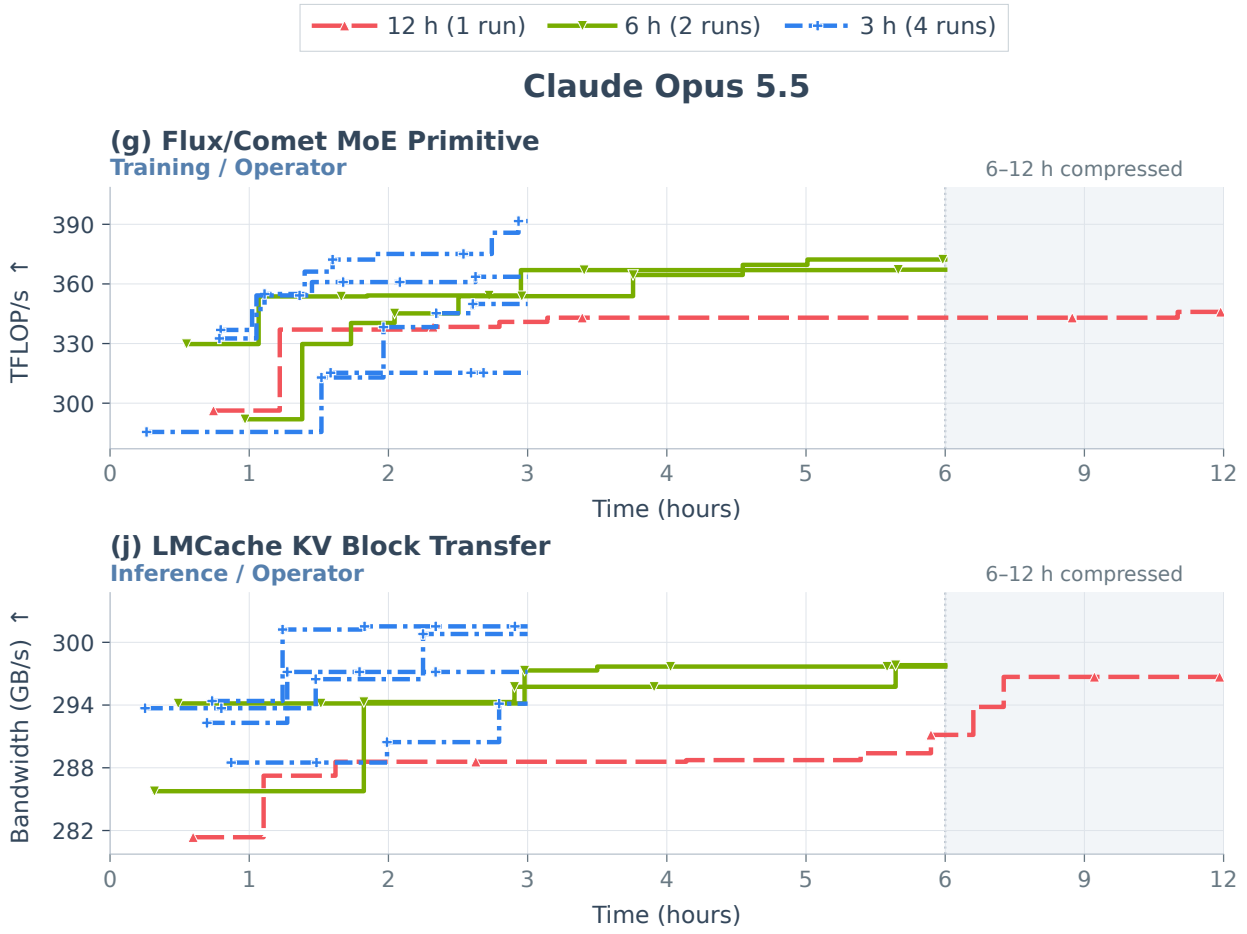

Figure 16: **Individual operator trajectories with Claude Opus 5.5.** Flux/Comet MoE Primitive and LMCache KV Block Transfer each use one 12-hour run, two 6-hour runs, and four 3-hour runs. Curves show performance over time. The 6–12-hour interval is compressed threefold; ticks show actual hours.

requests; several ML tasks also explicitly require CPU-only execution. Its main runs use 12 hours; extended experiments reach 28 and 72 hours. The released harness blocks voluntary stops and resumes eligible exits using the remaining execution budget, establishing framework support for long-horizon execution. Runs may end at the time or resume limit, or after exits that do not qualify for recovery.

E Quant-VideoGen KV Codec Case Study Details

This appendix gives the experimental settings, codec designs, optimization steps, and metric definitions for Section 4.2.

E.1 Experimental Setup

We compare the human codec with Claude Opus 5.5's codec on two video models, using 200 samples per dataset. Both codecs and the uncompressed BF16 reference use the same inputs and generation settings (Table 4).

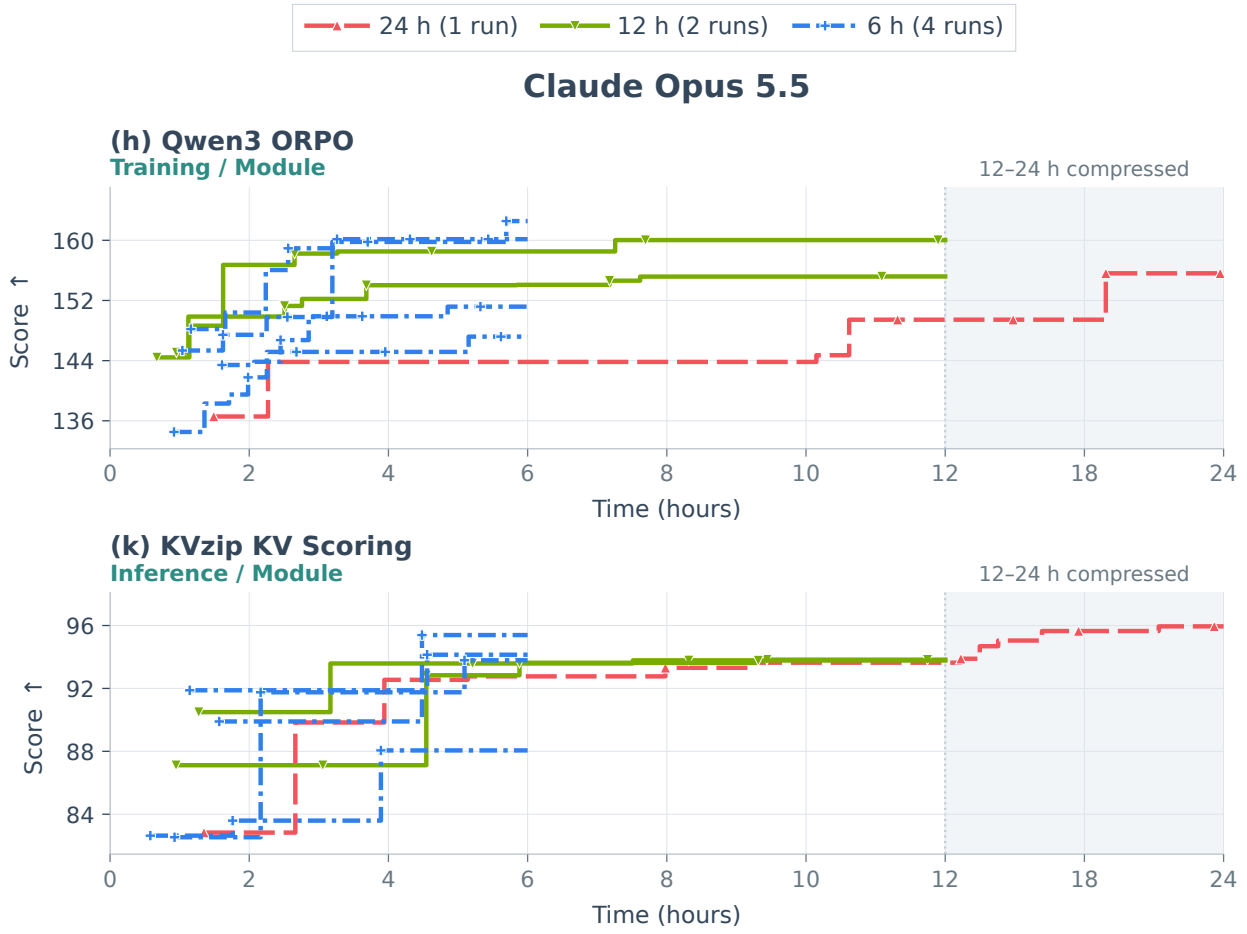

Figure 17: **Individual module trajectories with Claude Opus 5.5.** Qwen3 ORPO and KVzip KV Scoring each use one 24-hour run, two 12-hour runs, and four 6-hour runs. Curves show performance over time. The 12–24-hour interval is compressed threefold; ticks show actual hours.

Table 4: **Video generation and evaluation settings.** Frame indices start at zero.

Setting	**HY-WorldPlay**	**LongCat-Video**
Dataset	WorldScore [52]	MovieGen [54]
Input	Image, prompt, and camera actions	Prompt and fixed 73-frame prefix
Resolution; seed	832 × 480; 0	832 × 480; 0
Output frames	285	273: 73 fixed and ten 20-frame segments
Frame rate	24 fps	15 fps
Denoising	Four steps	50 steps; guidance scale 4.0
PSNR/SSIM/LPIPS frames	29–44	73–92

WorldScore selection. We map camera actions from the static WorldScore dataset to HY-WorldPlay controls and extend generation to 285 frames. Screening for severe collapse in uncompressed videos retains 200 of 289 reviewed samples. The screen uses only uncompressed references, except for eight pilot samples whose codec results had already been inspected; two of these were excluded. The retained set contains 131

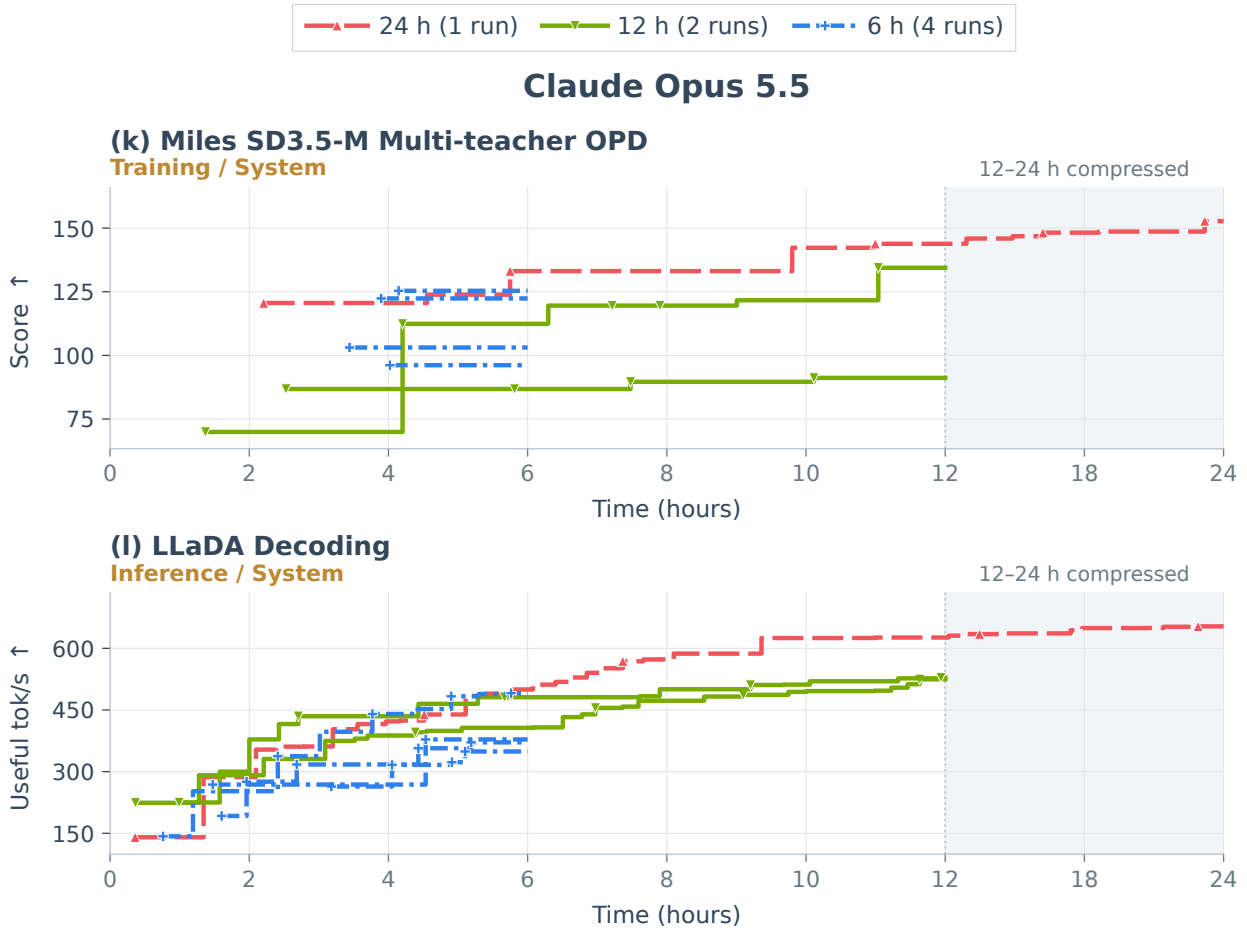

Figure 18: **Individual system trajectories with Claude Opus 5.5.** Miles SD3.5-M Multi-teacher OPD and LLaDA Decoding each use one 24-hour run, two 12-hour runs, and four 6-hour runs. Curves show performance over time. The 12–24-hour interval is compressed threefold; ticks show actual hours.

outdoor and 69 indoor samples, and 115 photorealistic and 85 stylized samples. Its prompts do not overlap with the 11 original HY-WorldPlay task prompts. HY-WorldPlay uses our 8B integration with the official quantization code.

MovieGen selection. We sample 200 prompts from the 1,003-entry list released with Self-Forcing, excluding the six private-test and two public-test prompts, without screening generation quality.

E.2 Codec Designs

Table 5 compares the two codecs. For one attention head, $X \in \mathbb{R}^{S\times D}$ denotes either the cached keys (K) or values (V), with S tokens and head dimension D. Let x_i be token i's vector, $\widehat{x}_i$ its reconstruction, C the number of clusters, and a_i its assigned cluster.

Human codec. Given cluster centers c_j, the codec [16] assigns each token to its nearest center and quantizes the residual:

$$a_i = \arg\min_j \|x_i - c_j\|_2^2, \qquad \widehat{x}_i = c_{a_i} + Q_3(x_i - c_{a_i}). \tag{11}$$

Here Q_3 rounds each residual value to $\{-s, 0, s\}$, with FP8 scale s.

Table 5: **Codec designs for Human and Claude Opus 5.5.**

Component	Human	Claude Opus 5.5
K/V clustering	Separate assignments	Shared assignments; separate when $2D > 256$
Cluster count	256	Adaptive (Equation 12)
Lloyd iterations	Two for HY-WorldPlay; up to 100 for LongCat-Video	Ten
Residual transform	None	PCA, retaining all dimensions
Quantization	Ternary values; one FP8 scale per 64 elements	Eight-state trellis quantization
Residual storage	Two bits per symbol	rANS entropy coding

Claude's codec. Let $\bar{k}_i$ and $\bar{v}_i$ be the mean-centered, root-mean-square-normalized K and V vectors, and d_j a joint cluster center. The cluster count and shared assignment are

$$\begin{aligned} C &= \max(1, \min(256, \lfloor S/128 \rfloor)), \\ u_i &= [1.3\bar{k}_i, \bar{v}_i], \qquad a_i = \arg\min_j \|u_i - d_j\|_2^2. \end{aligned} \tag{12}$$

For each of K and V, let μ be the mean token vector, $\widetilde{c}_j$ a stored quantized center, and $B \in \mathbb{R}^{D\times D}$ the PCA residual transform. With transformed residual z_i and trellis quantizer Q_{TCQ} of step size Δ, reconstruction is

$$\begin{aligned} z_i &= (x_i - \mu - \widetilde{c}_{a_i})B^{-1}, \\ \widehat{x}_i &= \mu + \widetilde{c}_{a_i} + Q_{\mathrm{TCQ}}(z_i; \Delta)B. \end{aligned} \tag{13}$$

B is stored as an INT8 basis with scales. Both B and Δ include the scaling factors for K, V, and each head. The trellis minimizes squared error after clipping, without a rate penalty. Range asymmetric numeral systems (rANS) coding losslessly encodes the resulting indices using their empirical frequencies. Stored data include centers, transforms, scales, shared assignments, and frequency tables.

Task scores. The human codec has a quality-weighted task score of 6.82, compared with 28.45 for Claude's final submission.

E.3 Optimization Steps

Table 6 maps the five numbered points in Figure 1 to Claude Opus 5.5's submissions and explains the changes leading to each point. Scores are recomputed using the quality-weighted metric in Equation 10.

E.4 Evaluation Metrics

Table 7 summarizes the four metrics reported in Section 4.2. The results tables report arithmetic means over samples. PSNR, SSIM, and LPIPS compare raw RGB frames with the corresponding uncompressed BF16 frames, using the windows in Table 4; frame scores are averaged within each window.

Compression ratio. For encoding call j, let b_j be the original BF16 KV bytes and $\widehat{b}_j$ the stored bytes, including all reconstruction metadata. A video's compression ratio is

$$R = \frac{\sum_j b_j}{\sum_j \widehat{b}_j}. \tag{14}$$

Shared storage is counted once within each call. Only KV tensors passed to the codec enter this ratio. The task score combines compression and video quality for each video model (Appendix A.3).

Table 6: **Claude's optimization steps on Quant-VideoGen KV Codec.** Point numbers match Figure 1.

Point	Submission	Optimization	Score
1	6th	**4-bit quantization.** Claude tests a simple 4-bit codec and obtains its first valid private test score.	5.80
2	7th	**PCA + entropy coding.** Claude returns to clustering and PCA residual quantization with rANS entropy coding. Video tests in the Work environment guide changes to the quantization step and encoded symbol set.	14.07
3	8th	**Quantization-step tuning.** Claude uses video tests in the Work environment to select larger quantization steps for K and V, increasing compression while passing private tests.	17.66
4	16th	**Trellis quantization.** Claude jointly quantizes PCA coefficients and uses finer steps for high-variance key directions.	21.90
5	22nd	**Metadata compression + per-head quantization.** Claude losslessly compresses metadata and shares K/V cluster assignments. It also refines clustering and uses finer quantization steps for attention heads with larger residual magnitudes.	28.45

Table 7: **Metrics for the codec comparison.** Arrows indicate the preferred direction.

Metric	What it measures	Computation or model
Compression ↑	KV storage savings	Ratio of original to stored bytes (Equation 14)
PSNR ↑	Pixel reconstruction accuracy	Mean squared error, expressed in dB (Equation 15)
SSIM ↑	Local luminance, contrast, and structural similarity	Uniform 11×11 windows on each RGB channel (Equation 16)
LPIPS ↓	Perceptual difference from the BF16 reference	Weighted VGG feature distances; LPIPS version 0.1

PSNR. Let $I_t, \widehat{I}_t \in [0,1]^{H\times W\times 3}$ be reference and reconstructed RGB frames at time t, with height H and width W. Peak signal-to-noise ratio is

$$\mathrm{MSE}_t = \frac{\|I_t - \widehat{I}_t\|_F^2}{3HW}, \qquad \mathrm{PSNR}_t = -10\log_{10}\big(\max(\mathrm{MSE}_t, 10^{-10})\big). \tag{15}$$

SSIM. Structural similarity [83] compares corresponding local windows x and y. Let μ_x, μ_y be their means, σ_x^2, σ_y^2 their variances, and σ_{xy} their covariance. The local score is

$$\mathrm{SSIM}(x,y) = \frac{(2\mu_x\mu_y + C_1)(2\sigma_{xy} + C_2)}{(\mu_x^2 + \mu_y^2 + C_1)(\sigma_x^2 + \sigma_y^2 + C_2)}, \tag{16}$$

with $C_1 = 0.01^2$ and $C_2 = 0.03^2$. We average the local scores over all pixels and RGB channels, using 11×11 uniform windows with stride 1 and zero padding of 5 pixels.

LPIPS. Learned perceptual image patch similarity [84] measures weighted distances between VGG features. We use version 0.1 and rescale RGB inputs from $[0,1]$ to $[-1,1]$.

Quality-weighted task score. We rescore submitted solutions using compression and video quality. For each video model m, R_m is the ratio of total BF16 to stored cache bytes, and quality metrics are averaged over test cases.

For frame similarity, we measure PSNR, SSIM, and LPIPS on the first segment generated with compressed KV. Let $\mathbf{x}_m$ contain their mean values and $\mathbf{x}_m^{\mathrm{Human}}$ the corresponding values for the human codec. We normalize each metric toward $\mathbf{b} = (40, 1, 0)$ and cap it at 1:

$$f_m = \frac{1}{3}\sum_{j=1}^{3} \min\left(\frac{x_{m,j} - x_{m,j}^{\mathrm{Human}}}{b_j - x_{m,j}^{\mathrm{Human}}},\, 1\right). \tag{17}$$

VBench [85] measures full-video quality. For dimension k, let $V_{m,k}$ and $V_{m,k}^{\mathrm{BF16}}$ be the mean scores with and without compression, and $\Delta_{m,k}$ the allowed decrease. We use the smallest normalized margin across its four dimensions:

$$g_{m,k} = 1 - \frac{\max(V_{m,k}^{\mathrm{BF16}} - V_{m,k},\, 0)}{\Delta_{m,k}}, \qquad v_m = \min_k g_{m,k} - G_m^{\mathrm{Human}}, \tag{18}$$

where G_m^{Human} is the corresponding minimum margin for the human codec. We set $q_m = (f_m + v_m)/2$ in Equation 10.

Table 8: **Human codec values and quality thresholds.** VBench tuples follow aesthetic quality, background consistency, imaging quality, and subject consistency.

Quantity	LongCat-Video	HY-WorldPlay
$\mathrm{PSNR}^{\mathrm{Human}}$ (dB)	23.6403	31.0633
$\mathrm{SSIM}^{\mathrm{Human}}$	0.8704	0.9483
$\mathrm{LPIPS}^{\mathrm{Human}}$	0.0964	0.0435
G^{Human}	0.4060	0.3939
Δ	(1.873, 2.685, 6.174, 4.196)	(2.055, 1.511, 2.415, 1.984)

E.5 Additional Results

Table 9 gives all reported metrics; Table 10 reports how often Claude improves quality and separates encoding from decoding time.

Table 9: **Human and Claude on WorldScore and MovieGen.** Entries are dataset means. Time is in seconds; MovieGen timing uses 188 matched pairs. VBench [85] columns report aesthetic quality, background consistency, imaging quality, and subject consistency.

			Efficiency		Frame similarity			Full-video VBench ↑			
Video model	Dataset	Codec	Comp. ↑	Time ↓	PSNR ↑	SSIM ↑	LPIPS ↓	Aes.	Bg.	Img.	Subj.
HY-WorldPlay	WorldScore	Human	6.75	305	31.15	0.923	0.059	59.8	93.2	69.9	88.2
		Claude	20.89	262	36.33	0.975	0.023	61.8	94.0	70.9	89.6
LongCat-Video	MovieGen	Human	6.88	1734	23.67	0.836	0.126	63.8	93.0	65.8	89.1
		Claude	18.83	1760	26.95	0.884	0.085	66.0	94.2	68.4	91.2

Table 10: **Per-sample improvements and codec timing.** Counts compare Claude with the human codec. Times show Human → Claude in seconds per video.

Measurement	**WorldScore**	**MovieGen**
Higher compression	200/200	200/200
Higher PSNR	200/200	170/200
Higher SSIM	200/200	169/200
Lower LPIPS	200/200	176/200
Higher VBench scores (four dimensions)	615/800	659/800
Mean total encoding time (s)	67.3 → 41.4	103.2 → 49.5
Mean total decoding time (s)	25.2 → 8.6	18.5 → 102.4
Decodes per compressed cache	4	≈ 56