arXiv:2611.13755v1 [quant-ph] 10 Nov 2026

A Hardware-Trained Generative Quantum Eigensolver for Sample-Based Quantum Diagonalization of Organotin Hydroxides up to 70 Qubits

Kunal Kumar*[1], Kimberlee Keithley[2,4], Toshiya Iwamori[6], Shu Kanno[2,4], Ikko Hamamura[5], Takao Kobayashi[2], Akira Koseki[6], Qi Gao†[2,4], Naoki Yamamoto[3,4], Hiroshi Horii[1], Yuri Kobayashi[1], Borja Peropadre[7], Kouhei Nakaji[8], Toru Imai[1], and Takayuki Osogami[6]

[1]IBM Quantum, IBM Japan, Tokyo 103-8510, Japan
[2]Mitsubishi Chemical Corporation, Science & Innovation Center, 1000 Kamoshida-cho, Aoba-ku, Yokohama 227-8502, Japan
[3]Department of Applied Physics and Physico-Informatics, Keio University, Hiyoshi 3-14-1, Kohoku, Yokohama 223-8522, Japan
[4]Quantum Computing Center, Keio University, Hiyoshi 3-14-1, Kohoku, Yokohama 223-8522, Japan
[5]NVIDIA, ATT EAST 12F, 2-11-7 Akasaka, Minato-ku, Tokyo 107-0052, Japan
[6]IBM Research, IBM Japan, Tokyo 103-8510, Japan
[7]IBM Quantum, IBM Research Cambridge, Cambridge, Massachusetts, USA
[8]NVIDIA, Santa Clara, California, USA

Abstract

Quantum hardware can provide feedback for training circuit generators beyond the reach of exact classical simulation. We train generative quantum eigensolver (GQE) policies, transformers conditioned on molecular descriptors, to generate the state-preparation circuits for sample-based quantum diagonalization (SQD). Training runs first on a noiseless simulator and then with the IBM Heron processor `ibm_kobe` inside the loop, at 46 and 70 qubits (1.2×10^8 shots for the four main training runs). A classical CCSD(T) reference energy of the training molecules enters only an admission gate and two auxiliary reward terms. During hardware training, a reward penalty prevents collapse to circuits that sample only the Hartree–Fock determinant. Policies trained on transition-metal complexes are applied without retraining to $Sn(OH)_3$-Me and its hydrates, minimal models of organotin photoresists for extreme-ultraviolet lithography, from 32 to 70 qubits. After one round of hardware sampling, followed by classical ranking and recovery of the cached counts with the molecular Hamiltonian and no further QPU calls, 27 of 28 runs at 46–70 qubits are within 15 mHa of the density-matrix renormalization group (DMRG) reference (median 4.5 mHa), and in all eight runs with an untrained control the trained policy is more accurate. At 56–70 qubits the selected circuits use 34–72 times fewer two-qubit gates than the local unitary cluster Jastrow (LUCJ) circuit transpiled for the same device. These results demonstrate that circuit generators can be trained with quantum hardware at 70 qubits.

1 Introduction

Strongly correlated electrons, as in transition-metal complexes and organometallic compounds, are a main benchmark for quantum computational chemistry [1]. Exact classical diagonalization becomes impossible when the active space grows, because its cost increases exponentially. Density-matrix renormalization group (DMRG) methods [2] work well for low-dimensional systems, but the growth of entanglement limits them in complex active spaces. Quantum computers are therefore expected to help beyond the classically solvable regime. Early work reduced molecular problems to a few qubits and ran them on small superconducting devices [3, 4], and processors now have more than a hundred qubits and run circuits beyond exact classical simulation [5]. Present quantum processors are noisy and allow only shallow circuits. Quantum-selected configuration interaction (QSCI) [6] and sample-based quantum diagonalization (SQD) [7] suit this situation: the quantum processor only samples important electronic configurations from a prepared

*Corresponding author: kunalkumar@ibm.com
†Corresponding author: qi.gao.mb@mcgc.com

state, and a classical supercomputer diagonalizes the Hamiltonian in the subspace that these configurations define. This hybrid scheme has made active spaces of several tens of qubits accessible on superconducting processors [7], and it has been extended to molecular vibrations [8, 9]. Recent work has also explored broken-symmetry reference states for SQD and reported near-DMRG accuracy [10].

The quality of an SQD calculation depends on the circuit that prepares the state. Fixed ansätze such as the local unitary cluster Jastrow (LUCJ) ansatz [11] often give deep circuits, which must then be adapted to the connectivity of the device. The generative quantum eigensolver (GQE) offers another route: an autoregressive model, first formulated as a transformer ("GPT-QE"), generates a sequence of operations from a predefined operator pool, so the circuit is learnt and not designed by hand [12]. Transformer-based GQE has been used to numerically design circuits for QSCI, with a substantial reduction in two-qubit gates compared with Trotterized baselines, for active spaces of up to 32 qubits [13]. A GQE and QSCI pipeline has also been applied to tin photoresist models of 14 to 44 qubits, including methyltin trihydroxide, with hardware runs on an IQM processor for smaller tin systems at 14 and 22 qubits [14].

Training one model on many molecules (amortization) would let the same model propose circuits for a new molecule, without designing a new ansatz for each active space. Training such a model requires an energy estimate for every circuit it proposes. In this work, the noiseless training stage uses molecules of up to 32 qubits, and exact classical simulation becomes intractable from about 46 qubits (Section 2.7); for the 46- and 70-qubit targets we obtain the energy feedback from the quantum processor. As a result, the policy is trained on hardware-measured (noisy) outcomes, and circuits are scored according to their performance under the device noise and transpilation constraints. Here we extend amortized circuit design to transition-metal active spaces of 46–70 qubits and train the policy directly on superconducting hardware.

Sustained semiconductor scaling demands the transition of lithographic patterning from deep ultraviolet toward extreme ultraviolet lithography (EUVL). In this context, organotin materials have emerged as leading photoresist candidates owing to the high absorption cross-section of tin at the EUV excitation energy of 92 eV [15–17]. Consequently, diverse organometallic and hybrid tin clusters have been synthesized, including butyltin oxo hydroxo complexes [18], sodium-templated alkyltin Keggin clusters [19, 20], and organotin carboxylates [21]. Photolysis of the Sn–C bond is believed to drive the EUV chemistry of these materials, and its description requires multireference methods [22]. Such systems therefore need careful treatment when the size of the problem is chosen for a quantum algorithm. Here, we employ the tin-organometallic family as held-out targets for an amortized generative policy trained on benchmark transition-metal complexes, using $Sn(OH)_3$-Me (methyltin trihydroxide), the smallest model of a monoorganotin hydroxide, and its mono- and di-hydrates.

To make policy training feasible under device noise, we use a two-stage curriculum. Stage 1 uses a noiseless simulator and small molecules, while Stage 2 continues the training on the IBM processor `ibm_kobe` using transition-metal complexes with 46- and 70-qubit active spaces. The resulting policies are then evaluated zero-shot on the held-out $Sn(OH)_3$-Me targets introduced above. This progression from easier simulated problems to larger hardware-based problems is intended to stabilize training under device noise [23]. In Stage 2, we compare two rewards: one incorporating a Hartree–Fock penalty, which prevents the policy from producing only the reference determinant, and one tuned jointly on a noisy simulator (Section 2.7.1). Our evidence that the policy transfers to new chemistry is qualitative (Section 4).

The trained policies give energies close to DMRG with short circuits. For each target, the policy generates ten circuits, each circuit is sampled once on `ibm_kobe`, and the circuit with the lowest energy computed from its measured shots is used for configuration recovery. At 46 to 70 qubits, 27 of 28 runs are within 15 mHa of the DMRG reference (median 4.5 mHa), and in all eight runs with an untrained control the trained policy is more accurate (Section 3.3). At 56 to 70 qubits the selected circuits use 64 to 160 two-qubit gates. This is 34 to 72 times fewer than the LUCJ circuit transpiled for the same device (Qiskit optimization level 1), which needs 4,248 to 6,952 two-qubit gates (Section 3.4). Sampling ten circuits takes about 11 times the QPU time per job of sampling one LUCJ circuit (Section 3.9).

The primary contributions of this work are as follows:

(i) We develop amortized GQE policies trained with a two-stage curriculum (a noiseless-simulator warm-up followed by continuation on `ibm_kobe`) for 46- and 70-qubit active spaces, with a Hartree–Fock penalty in the hardware-stage reward (Section 2.7.1), and compare them with a jointly tuned reward (Section 3.5).

(ii) We perform a zero-shot evaluation across the eleven active-space tiers of the $Sn(OH)_3$-Me family (32–70 qubits) against DMRG references, comparing five distinct circuit selection and pooling strategies, with matched untrained controls.

(iii) We present a circuit-level analysis (depth, valid-shot yield, spin-resolved excitation hierarchy and operator statistics) of the sampled quantum circuits (Sections 3.4 and 3.8, Appendix F).

(iv) We report the quantum and classical cost of training and inference: shots, QPU time and classical wall

time (Section 3.9). To our knowledge, no GQE policy has been trained on quantum hardware at this size before, so these costs are a first reference point for future studies.

(v) We open-source the related code, the generated sequence tokens, and the analyses of their excitation content, its correlation with the spin sectors, and the correctness of the populated orbitals.

The remainder of the paper is organized as follows. Section 2 describes the method: the pipeline, sample-based quantum diagonalization, the policy and its architecture, the molecules, the training procedure, the hyperparameters and the computing resources. Section 3 reports the results: the two training stages (Sections 3.1 and 3.2), the zero-shot inference on unseen molecules (Section 3.3), the circuit cost against LUCJ (Section 3.4), the comparison of the two rewards (Section 3.5), the divergence from the LUCJ circuit (Section 3.6), the comparison with DMRG (Section 3.7), the determinants of recovery accuracy (Section 3.8) and the computational cost (Section 3.9). Section 4 discusses the results, and Section 5 concludes.

2 Methods

2.1 Pipeline overview

Figures 1 and 2 summarize the full training and inference control flow, elaborated in the remainder of this section. Training proceeds per molecule-selection iteration: the size curriculum (Section 2.7) picks a molecule, its descriptor conditions the policy (Section 2.5, Fig. 3), the policy emits a batch of candidate gate sequences, and—depending on which stage is active for that lineage—the resulting circuits are either simulated via statevectors (Stage 1, Section 2.7) or executed on physical superconducting hardware (Stage 2). Every sampled circuit is transpiled, its subspace is built via recovery and global merge (Section 2.4), and the resulting raw energy error is checked, together with the circuit depth and non-triviality criteria, against a hard admission gate; only circuits that pass the gate contribute a policy-gradient update (Section 2.3). Inference reuses the identical policy forward pass and recovery machinery with frozen weights, with one key addition relative to training: a single hardware sampling round trip's raw counts are cached, enabling subsequent ranking, selection, and subspace recovery to execute locally without invoking additional quantum processor queries.

2.2 Sample-based quantum diagonalization

Let an ansatz circuit $U(\boldsymbol{\theta})$ act on $n_q = 2n_o$ qubits, where n_o is the number of active spatial orbitals (Jordan–Wigner-mapped, one qubit per spin-orbital, spin-separated so that an n_o-orbital active space on (n_α, n_β) electrons occupies n_q qubits in total). We write $|\psi_{\boldsymbol{\theta}}\rangle = U(\boldsymbol{\theta})\,|\mathrm{HF}\rangle$ for the state prepared by the circuit. Measuring $|\psi_{\boldsymbol{\theta}}\rangle$ in the computational basis yields a set of bitstrings, each splittable into an α-string and a β-string of Hamming weight n_α and n_β respectively. Let $\mathcal{A}, \mathcal{B}$ denote the sets of distinct α- and β-strings observed over N_{shots} measurements, with empirical counts $c(a, b)$ for each observed pair $(a, b) \in \mathcal{A} \times \mathcal{B}$; the observed determinants form the set S_{raw}.

A classical configuration-interaction subspace S is constructed from the samples, but is not restricted to S_{raw}. In our recovery procedure (Section 2.4), the observed strings are repaired, the Hartree–Fock determinant and carryover strings are added, and the determinant set is taken as the product $S_{\text{rec}} = \mathcal{A}' \times \mathcal{B}'$, which in general contains determinants that were not directly measured. For a circuit's own recovered subspace we set $S = S_{\text{rec}}$, while the pooled result uses a merged set S_{merge}.

The molecular Hamiltonian $\hat{H}$ is projected onto S and diagonalized exactly. Here $|v\rangle$ is a trial vector in the span of the determinants of S; it is not the state $|\psi_{\boldsymbol{\theta}}\rangle$ prepared by the circuit, and the minimum can lie below $\langle\psi_{\boldsymbol{\theta}}|\hat{H}|\psi_{\boldsymbol{\theta}}\rangle$:

$$E_S = \min_{|v\rangle \in \mathrm{span}(S)} \frac{\langle v|\hat{H}|v\rangle}{\langle v|v\rangle}. \tag{1}$$

The subspace S comes from the samples of one circuit $U_{\mathbf{g}}$ (Section 2.3), so E_S is the SQD energy of that circuit; we also write it $E_{\text{SQD}}(\mathbf{g})$. We report the signed deviation $\Delta E = E_{\text{SQD}} - E_{\text{DMRG}}$ from the DMRG reference energy (in mHa); a positive value means that the SQD energy lies above the reference. Because the DMRG energy is itself an approximation, the sign is informative. All deviations in this study are positive. Because S is built from sampled determinants rather than an a priori truncation, the quality of E_S depends both on the circuit's ability to place amplitude on determinants near the true ground state and on how S is selected and enlarged from the raw sample (Section 2.4).

2.3 Amortized generative policy

Rather than fixing $U(\boldsymbol{\theta})$ to a hand-designed form, a policy π_ϕ generates the circuit itself, token by token. At each of L autoregressive decoding steps, π_ϕ samples a gate token g_t from a problem-specific operator pool $\mathcal{P}$ (Section 2.5), conditioned on the tokens emitted so far and on a molecule-specific descriptor $\mathbf{d}$. Here L is the number of gate tokens per circuit: $L = 28$ for the policies trained for 46 qubits and $L = 44$ for those trained for 70 qubits. The sampling is defined as follows:

$$g_t \sim \pi_\phi(\cdot \mid g_{<t}, \mathbf{d}), \qquad t = 1, \ldots, L, \tag{2}$$

yielding a full gate sequence $\mathbf{g} = (g_1, \ldots, g_L)$ and hence a circuit $U_{\mathbf{g}}$. A batch of N_{samples} sequences is drawn per

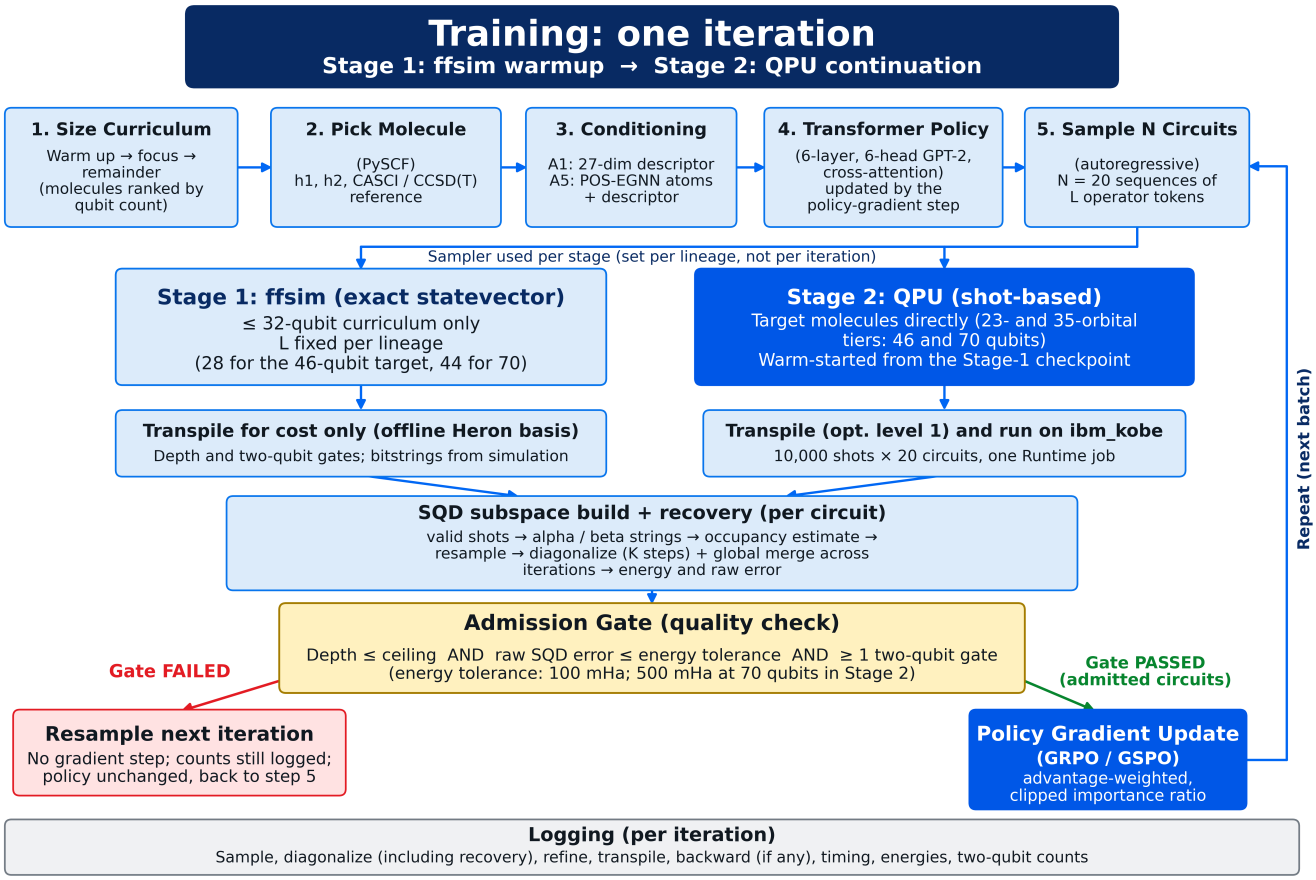

Figure 1: Training control flow for one iteration. The size curriculum picks a molecule; its conditioning (A1 or A5) drives the transformer policy, which samples $N = 20$ circuits of L operator tokens each. In Stage 1 the circuits are evaluated with an exact statevector simulator (curriculum of at most 32 qubits; the transpilation only measures depth and two-qubit-gate count); in Stage 2 they are transpiled and run on `ibm_kobe` (10,000 shots for each of 20 circuits in one Runtime job) on the 23- and 35-orbital target molecules, initialized from the full Stage-1 policy checkpoint (every weight is carried over; Section 2.7). The SQD subspace build with self-consistent recovery and global merge gives each circuit's energy and raw error; an admission gate then checks the best circuit of the batch (depth within the ceiling, raw error within the energy tolerance, at least one two-qubit gate). Admitted iterations drive a GRPO/GSPO policy-gradient update over all 20 circuits; for rejected iterations the counts are logged but nothing is back-propagated and the next iteration resamples.

training iteration; each is scored by its resulting SQD energy $E_{\mathrm{SQD}}(\mathbf{g})$ (Eq. 1), and the policy is updated to favor sequences with lower energy via a group-relative policy-gradient objective (GRPO/GSPO; [24, 25]): with the batch energies $E_i = E_{\mathrm{SQD}}(\mathbf{g}_i)$, $i = 1, \ldots, N_{\mathrm{samples}}$,

$$A_i = \frac{\bar{E} - E_i}{\mathrm{std}(\{E_j\}) + \varepsilon},$$

$$\mathcal{L}(\phi) = -\min\Big(r_i(\phi)\, A_i,\ \mathrm{clip}\left(r_i(\phi), 1 - \epsilon_{\mathrm{lo}}, 1 + \epsilon_{\mathrm{hi}}\right) A_i\Big), \tag{3}$$

where $\bar{E}$ is the batch mean energy, $r_i(\phi)$ is the (token- or sequence-level) importance ratio between the current and the behavior policy, and the lowest-energy sequence in the batch additionally receives a direct imitation term. A_i is undefined (and the ratio term dropped) when the batch has zero energy variance, which the training logs record as a 'reward-degenerate' iteration.

In each training iteration, the policy samples $N = 20$ circuits and we compute a raw SQD energy E_i for each. The admission gate is evaluated once per iteration, on the best circuit (lowest E_i): its transpiled depth must be within the ceiling, its raw SQD error within the energy tolerance, and it must contain at least one two-qubit gate (ruling out the trivial empty-circuit optimum). If these criteria are met, we compute advantages over all 20 circuits in the batch and apply one GRPO/GSPO policy-gradient update. If the gate fails, the iteration is logged but no gradient is back-propagated and the next iteration resamples.

2.4 Recovery and global-merge subspace construction

Sampled bitstrings obtained from physical quantum processors are typically sparse and corrupted by device noise, precluding their direct use as S in Eq. 1. Two

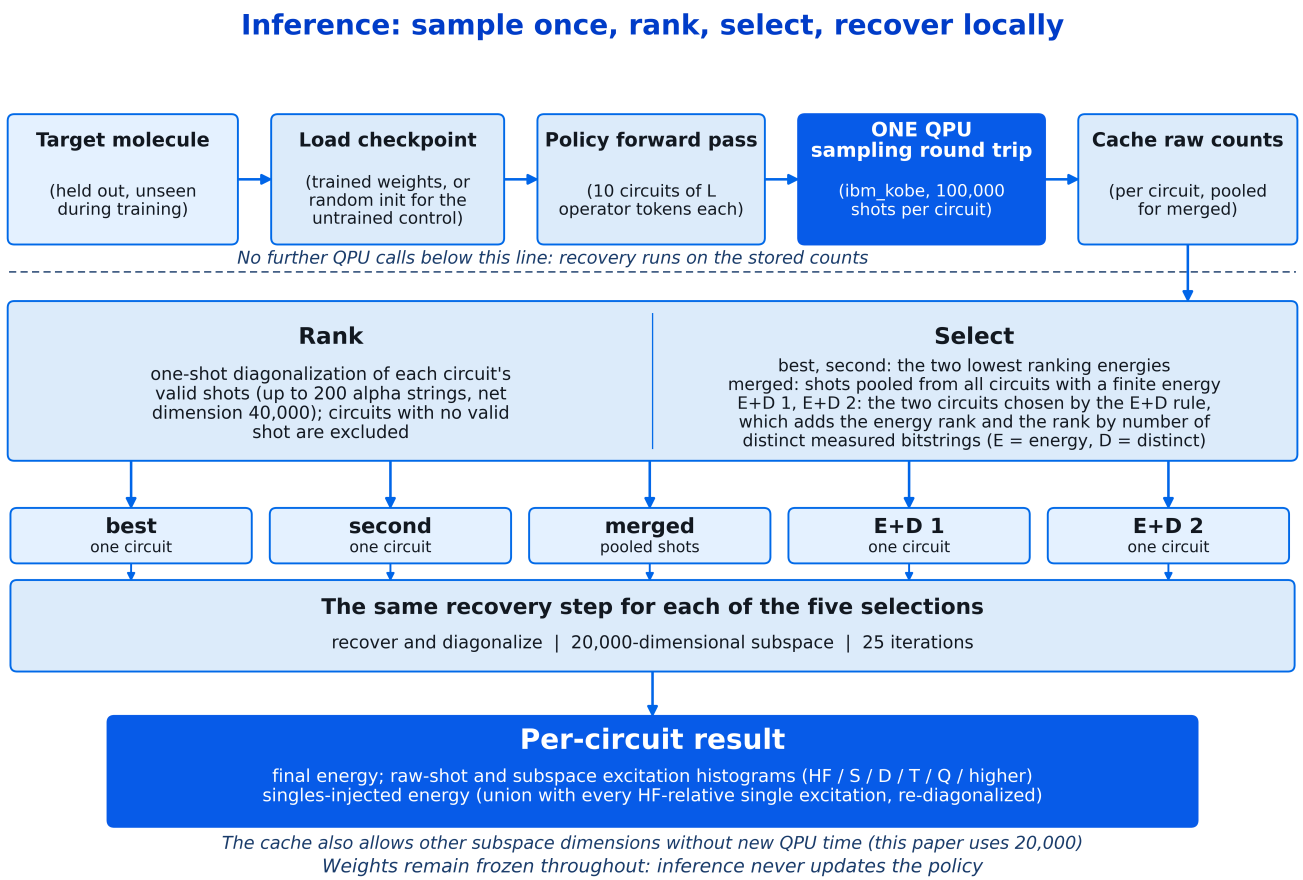

Figure 2: Inference control flow. A frozen policy (trained, or randomly initialized for the untrained control) emits 10 circuits for an unseen molecule; they are sampled once on `ibm_kobe` (100,000 shots each) and the raw counts are cached. Below the dashed line everything runs locally on the cached counts: the circuits are ranked by a one-shot diagonalization (up to 200 alpha strings, net dimension 40,000; circuits without a valid shot are excluded); five results are selected: best and second (the two lowest ranking energies), merged (the shots pooled from all circuits with a finite energy), and E+D 1 and E+D 2 (the two circuits chosen by the E+D rule, which combines the rank by energy, E, with the rank by the number of distinct measured bitstrings, D); the same recovery step is then applied to each of the five, with $D = 20{,}000$ strings per spin sector and 25 iterations. The result per circuit is the energy, the raw-shot and subspace excitation histograms and the singles-injected energy.

complementary subspace-construction mechanisms are employed to resolve this limitation.

Self-consistent configuration recovery. Given raw sampled counts $c(a, b)$, an iterative recovery loop (following the configuration-recovery procedure of the Qiskit SQD addon [7, 26]; code labels in Table A3) alternates between (i) estimating a running per-orbital occupancy vector $\boldsymbol{\mu}^{(k)}$ for each spin from the current best determinant set, (ii) repairing the observed bitstrings, in both the α and the β half, against the occupancy of each spin and post-selecting those whose halves have the correct electron numbers n_α and n_β (Hamming weight), (iii) subsampling down to a fixed target number D of strings per spin sector (the recovery subspace size)—where the Hartree–Fock determinant and carryover determinants from the prior step are enforced, with the remainder drawn from the recovered distribution; for the closed-shell molecules of this work the repaired α and β halves are first folded into one list of distinct strings, used for both spin sectors—and (iv) diagonalizing the determinant subspace $\mathcal{A}_k \times \mathcal{B}_k$ with $|\mathcal{A}_k| = |\mathcal{B}_k| = D$, whose net dimension is D^2 (Eq. 1), on GPU. Here D counts strings per spin sector: $D = 200$ at the ranking stage (40,000 determinants) and $D = 20{,}000$ at the final recovery (4×10^8 determinants, capped by the number of strings that exist). The lowest-energy determinant subspace over K recovery iterations is retained, with each step seeding the subsequent occupancy estimate and carryover list. In a variant called K-batch SQD [7], B independent subspaces are drawn at each step from the same recovered distribution and only the one with the lowest energy is kept; this costs B diagonalizations per step and gives a steadier occupancy estimate for the next step.

Global merge. Across training iterations (as opposed to within one circuit's own recovery loop), two global-refinement policies determine how much of the recovery history is retained and merged with the current batch's local result: a bounded mode keeps the union of the previous retained state and the current local state, capped at a fixed refine dimension, while a cumulative mode keeps the union of the previous retained state, the current local state, and every per-circuit state sampled this iteration, capped only by an overall cumulative dimension. The merged state does not enter the training reward. The reward of each circuit is its own recovered energy, computed from that circuit's samples alone (Eq. 1), before any merging. The merge is applied afterwards to the batch, and its energy is recorded as the best energy of the iteration, which is reported but not used as a reward or in the admission gate (which tests the best circuit's own raw error). At inference, the merge pools the shots of several circuits into the reported "merged" result (Section 3.3).

2.5 Model architecture

The policy π_ϕ is a GPT-2-style autoregressive transformer decoder with 6 layers, 6 attention heads and about 56.8 million parameters. The decoder follows the GPT-2 architecture [27] that GPT-QE uses for GQE [12]; the cross-attention to a molecule prompt is our addition. We did not tune the model size. Its cross-attention layers let the decoder attend to a molecule-specific prompt of variable length, instead of a single conditioning vector. Its position embeddings cover the $L+1$ positions of the gate-token sequence (a start token and L gate tokens), where L is defined in Section 2.3. The prompt enters only through cross-attention and carries no position embedding. The decoder is shared by the two conditioning schemes of this work, called A1 and A5. They differ only in how the prompt is built from the molecule; both are described in detail below and shown in Figure 3.

A1: scalar cross-attention conditioning. A fixed 27-dimensional hand-built active-space descriptor $\mathbf{d} \in \mathbb{R}^{27}$ (per-system scalars: orbital/electron counts, HF/CASCI-adjacent energetics, and related active-space statistics) is mapped by a three-layer MLP (hidden width 1500, GELU activations, 3.45 M parameters) to a single prompt token of the decoder's embedding dimension $E = 768$, with the output bounded element-wise by a scaled hyperbolic tangent (bound $b = 3$):

$$\mathbf{p}_{\mathrm{A1}} = b\,\tanh\!\big(\mathrm{MLP}_{27\to E}(\mathbf{d})/b\big) \in \mathbb{R}^{1\times E}. \qquad (4)$$

The bound caps the scale of the cross-attention keys and values, which would otherwise grow without limit as the encoder weights train. This one-token prompt is computed once per molecule and cached and cross-attended at every one of the L autoregressive decoding steps.

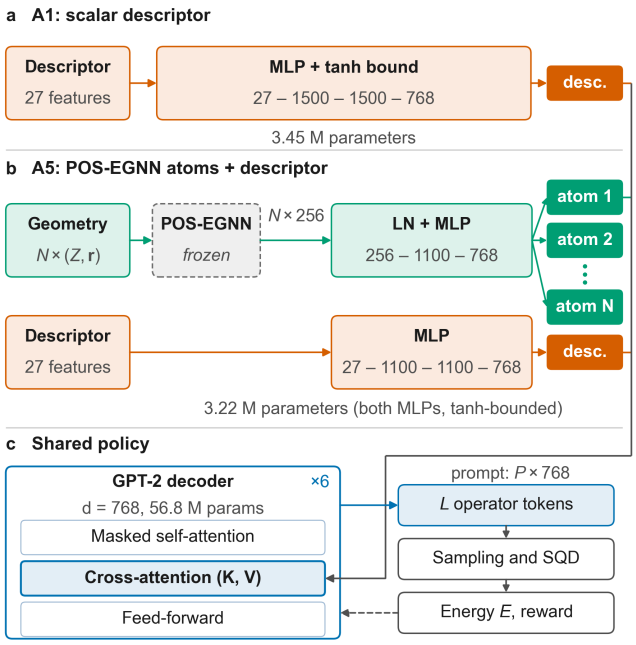

Figure 3: Conditioning encoders and shared policy. **(a) A1.** The 27-dimensional active-space descriptor is mapped by a three-layer MLP (27–1500–1500–768, GELU, 3.45 M parameters) to a single 768-dimensional prompt token; the output is bounded element-wise as $3\tanh(x/3)$. **(b) A5.** The N atoms of the molecule (atomic numbers Z and positions $\mathbf{r}$) are embedded by the pre-trained POS-EGNN [28, 29] into $N \times 256$ invariant per-atom features. This network is frozen: its output is computed once per geometry and stored. The features are layer-normalised and mapped by a two-layer MLP (256–1100–768) to N atom tokens. Independently, the descriptor (the A1 features, with the HOMO–LUMO gap and the mean and standard deviation of the orbital energies recomputed from the eigenvalues of the active-space one-body Hamiltonian) is mapped by a separate MLP (27–1100–1100–768) to one further token. The $N+1$ tokens are bounded as in A1 and kept as separate prompt tokens; the two MLPs have 3.22 M parameters in total. **(c) Shared policy.** Both arms feed the same GPT-2 decoder (6 layers, 6 heads, embedding dimension 768, ≈56.8 M parameters). In every layer, masked self-attention over the gate tokens generated so far is followed by a cross-attention block whose keys and values (K, V) are the prompt tokens ($P = 1$ for A1, $P = N+1$ for A5) and a feed-forward block. The decoder samples L operator tokens ($L = 28$ for the 46-qubit and 44 for the 70-qubit lineages), which are assembled into a circuit, sampled and processed by SQD; the resulting energy defines the reward. The dashed arrow marks the policy-gradient update of the decoder and of the MLP encoders. Vermillion: descriptor stream; green: atom stream; blue: policy; grey dashed: frozen network.

A5: POS-EGNN structural + scalar cross-attention conditioning. POS-EGNN (position-based equivariant graph neural network) is a model released by IBM Research within its Foundation Models

for Materials collection [28]. It was pre-trained as a machine-learning potential, predicting energies, forces and stresses, on about 1.4 million structures of the Materials Project trajectory data set MPtrj [29], and its invariant per-atom features can be extracted for downstream tasks. We use it as a frozen feature extractor: for a molecule with N atoms and geometry $\{(Z_i, \mathbf{r}_i)\}$ it yields 256-dimensional invariant embeddings $\mathbf{e}_i \in \mathbb{R}^{256}$, $i = 1, \ldots, N$. Because the geometry is fixed per molecule, these embeddings are computed once, in a separate environment (POS-EGNN pins a different PyTorch version than the training environment), and stored; they are neither recomputed nor updated during training. The embeddings are layer-normalized and projected by a two-layer MLP (256–1100–768) to N prompt tokens, and the descriptor, independently projected with separate weights by a three-layer MLP (27–1100–1100–768), gives one additional token. The descriptor $\mathbf{d}_{\mathrm{A5}}$ contains the same 27 features as in A1 except that the three orbital-energy entries are recomputed from the eigenvalues of the active-space one-body Hamiltonian, which repairs a defect of the original entries for systems loaded from integral files. The two MLPs have 3.22 M parameters in total.

$$\mathbf{p}_{\mathrm{A5}} = b\,\tanh\big([\,\mathbf{q}_1; \ldots; \mathbf{q}_N; \mathbf{q}_d\,]/b\big) \in \mathbb{R}^{(N+1)\times E},$$
$$\mathbf{q}_i = \mathrm{MLP}_{256\to E}\big(\mathrm{LN}(\mathbf{e}_i)\big), \qquad \mathbf{q}_d = \mathrm{MLP}_{27\to E}(\mathbf{d}_{\mathrm{A5}}), \tag{5}$$

i.e. a prompt of N+1 tokens, kept as separate tokens (rather than pooled into one vector) so that the decoder's cross-attention, not a fixed pooling operator, determines which atoms matter for a given decoding step. Both token streams pass through the same bound $b = 3$ as in A1, so neither stream can dominate cross-attention through an unbounded activation scale. The prompt tokens carry no position or token-type embedding: the atom tokens are an unordered set, and the descriptor token differs from them only through its separate projection.

2.5.1 Data flow through the policy

Input → output, training. For a given molecule, both arms' encoders and the shared decoder are run forward together: build the prompt ($\mathbf{p}_{\mathrm{A1}}$ or $\mathbf{p}_{\mathrm{A5}}$) from the molecule's descriptor (and geometry, for A5) → cache it → autoregressively sample N_{samples} gate-token sequences of length L, cross-attending the cached prompt at every step → assemble each sequence into a circuit via the operator pool $\mathcal{P}$ → sample/simulate each circuit (Stage 1 or Stage 2, Section 2.7) → build and recover a subspace per circuit (Section 2.4) → admission gate → for admitted iterations (those whose best circuit passes), backpropagate the GRPO/GSPO loss (Eq. 3) using all $N = 20$ circuits in the batch through both the decoder and the active encoder (A1's MLP or A5's two MLPs), since the prompt is kept attached to its autograd graph rather than detached. Both the decoder and the conditioning encoder are therefore trained jointly, end to end, from the SQD energy signal alone.

Table 1: Token map of the operator pool (`FfsimLUCJPool`) for a molecule with $n = n_{\mathrm{orb}}$ spatial orbitals. Token ids above $3n+4$ are identity padding to reach the fixed vocabulary size 110.

Token id(s)	Operator
0	identity
1	LUCJ layer
2–5	orbital rotations (four fixed seeds)
6–$2n+3$	Givens rotations on neighbouring pairs $(p, p{+}1)$ (two angles per pair)
$2n+4$, $2n+5$	diagonal-Coulomb layers (two time values)
$2n+6$–$3n+4$	number–number phase gates on neighbouring pairs $(p, p{+}1)$ (one gate per pair)
Total	$3n+5$ operators

Input → output, inference. Identical forward path through the encoder and decoder, with every weight frozen (no backward pass, no optimizer step): molecule descriptor (and geometry, for A5) → prompt → sampled circuits → one QPU sampling round trip (cached, Fig. 2) → ranking and selection of the best, second, merged and E+D circuits → recovery/diagonalization at the reported subspace size → final energy and excitation histograms. The ten circuits are ranked by a one-shot diagonalization of their valid shots; best and second are the two circuits with the lowest ranking energy, merged pools the shots of all circuits with a finite ranking energy, and E+D 1 and E+D 2 are the two circuits that combine the rank by energy (E) with the rank by the number of distinct measured bitstrings (D). An untrained control uses the same path with randomly initialized decoder and encoder weights instead of a checkpoint, so that the difference to a trained policy shows what the training contributed beyond the architecture and the ranking and recovery steps. We ran eight such controls, on three molecules (SnOH3-me L, SnOH3-me-h2o M and SnOH3-me LL).

The operator pool $\mathcal{P}$ comprises local two-orbital gates (Givens rotations, which have single-excitation character, and number–number phase gates, each acting on one spatial-orbital pair, i.e. four qubits including both spin sectors) together with full-coverage operators (four fixed orbital rotations, two diagonal-Coulomb layers and one LUCJ layer) that touch every qubit in the active space. For a molecule with n_{orb} spatial orbitals, the pool has $3\,n_{\mathrm{orb}} + 5$ operators: 74 for the 23-orbital (46-qubit) and 110 for the 35-orbital (70-qubit) molecules. The vocabulary of the transformer is fixed at 110 for all molecules, with the remaining token ids acting as identity padding when $3\,n_{\mathrm{orb}} + 5 < 110$ (token map in Table 1).

2.6 Molecules and active spaces

Training curriculum. A size-ordered curriculum of small and medium molecules (H_2, LiH, H_2O, BeH_2, NH_3, CH_4, N_2, CH_2O, HCN, C_2H_4, CO_2) together with three transition-metal complexes — $Fe(CN)_6$, $Co(NH_3)_6$, and ferrocene — at a small, $\leq$16-orbital active-space tier for Stage-1 pretraining, and at a 23-orbital (46-qubit) or 35-orbital (70-qubit) tier for the target of Stage-2 continuation (Section 2.7). The molecules are ordered by qubit count, which sets the length of the generated sequence and the size of the operator pool and the cost of the simulation; we did not order them by the difficulty of sampling their ground states.

Zero-shot/held-out target. The SnOH3-me family (and its mono- and di-hydrate adducts) is used as a held-out target, at active-space tiers S/M/L/LL spanning 24–50 electrons and up to 35 orbitals (70 qubits), with DMRG references available at every tier (Table A1).

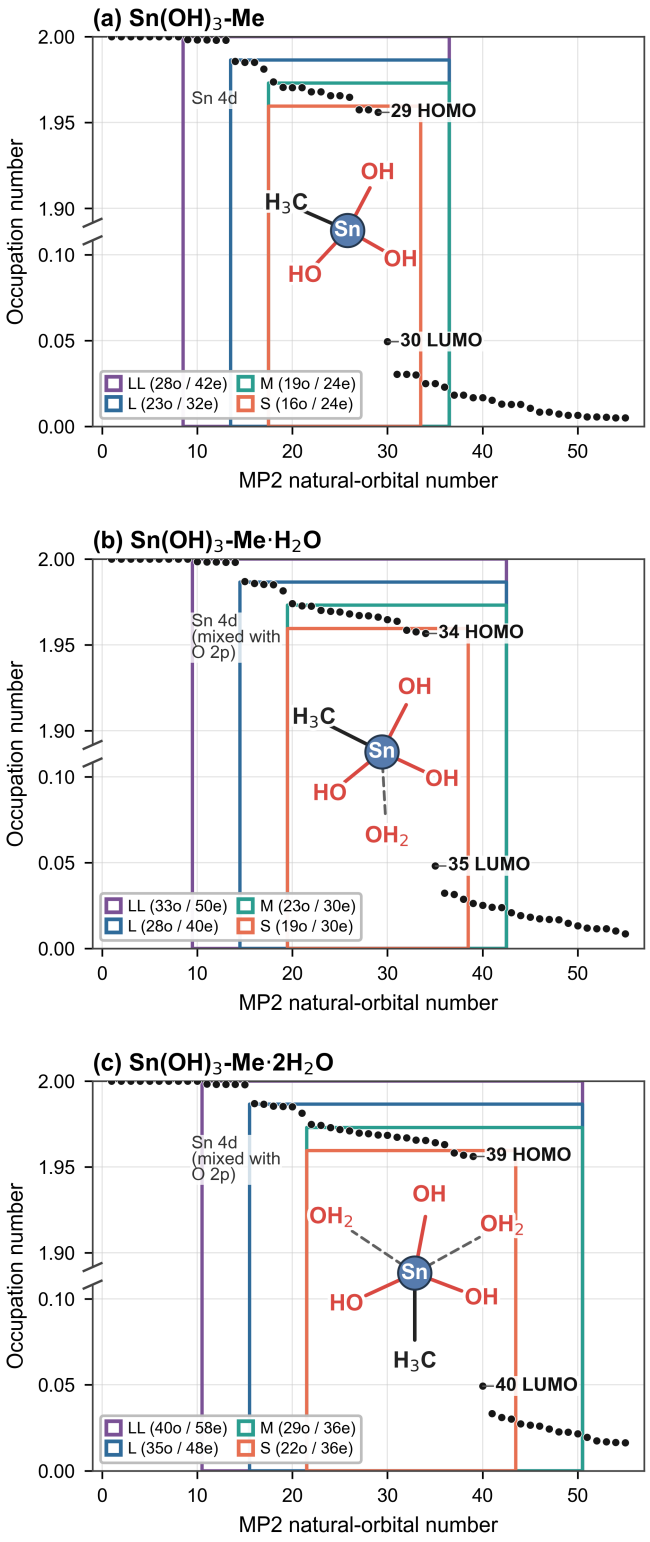

Figure 4: MP2 natural-orbital occupation numbers (MP2/def2-SVP, frozen core) of (a) $Sn(OH)_3$-Me, (b) $Sn(OH)_3$-Me·H_2O and (c) $Sn(OH)_3$-Me·$2H_2O$, with the S, M, L and LL active spaces boxed (orbitals, electrons; the natural-orbital ranges are in Table A2). The orbitals labelled Sn 4d are mixed with O 2p. The occupation axis is compressed between 0.1 and 1.9. The sketch in each panel is the structure of the molecule (Sn blue, O red, C black; dashed: water bound to tin). In the optimized geometry used for the calculations the first water of the dihydrate is coordinated to tin (Sn–O 2.4 Å) and the second lies 3.5 Å from tin, hydrogen-bonded to two hydroxyl groups.

Held-out molecules and how their active spaces were chosen. The held-out targets are $Sn(OH)_3$-Me and its mono- and di-hydrate adducts, $Sn(OH)_3$-Me·H_2O and $Sn(OH)_3$-Me·$2H_2O$. $Sn(OH)_3$-Me is the smallest model of a monoorganotin hydroxide, the Sn–C and Sn–O motif of the organotin oxo-hydroxo clusters discussed in the introduction (Section 1); in the optimized structures one water is coordinated to tin in each hydrate (Sn–O 2.45 Å in the monohydrate, 2.41 Å in the dihydrate; the Sn–OH distances are 2.0 Å) and is hydrogen-bonded to two hydroxyl groups, while the second water of the dihydrate lies in the second shell (Sn–O 3.53 Å) and is hydrogen-bonded to two hydroxyl groups, so that the water changes the environment of tin while keeping the chemistry the same, and together the three molecules give a ladder of active spaces from 32 to 70 qubits with DMRG references at every system size (Table A1). The structures were optimized and shown, by vibrational-frequency analysis, not to be saddle points; the MP2 natural orbitals were computed with Gaussian 16 [30] at the frozen-core MP2 level [31] in the def2-SVP basis [32], and the active spaces are built on these natural orbitals [33], not on Hartree–Fock canonical orbitals. The natural orbitals are ordered by occupation number (Figure 4). An active space is a contiguous range of them: it starts at the lowest occupied natural orbital of the range, runs through the HOMO and includes the first virtual natural orbitals. Eleven tiers were defined (Table A2): S, M, L and LL for $Sn(OH)_3$-Me and for the monohydrate, and S, M and L for the dihydrate (there is no LL tier for the dihydrate). S and M share their lowest orbital (occupation 1.974–1.975) and differ in the number of virtual orbitals; L and LL add occupied orbitals whose occupation is closer to 2 (lowest orbital of L: 1.986–1.987; of LL: 1.998), and only LL reaches down to the orbitals labelled Sn 4d, which mix with O 2p (Figure 4). The occupation-number threshold that ends the virtual range is not well defined; the tiers use the first 4 (S) or 7 to 11 (M, L, LL, depending on the molecule) virtual natural orbitals. Whether the Sn 4d orbitals must be in the active space is not settled by this work. The natural orbitals themselves are shown in Figures A1, A2 and A3 and discussed in Appendix B. Natural orbitals were preferred to selections by Hartree–Fock orbital energy because they are more localized on bonds and atoms and so have more distinct characters (an assessment made when the active spaces were chosen; the two selections were not benchmarked against each other here).

2.7 Training procedure

Training proceeds in two stages per (architecture, qubit-count-target) pair.

Stage 1 — noiseless-simulator warmup. The policy is trained from scratch, that is, from random weights, with an exact-statevector (`ffsim`) sampler on the ≤32-qubit curriculum (Section 2.6), using $L = 28$ gate tokens per circuit for the 46-qubit-target lineage and $L = 44$ for the 70-qubit-target lineage (gate count cannot change once a lineage is fixed, since it resizes the transformer's position-embedding table). In both stages every generated circuit acts on the Hartree–Fock reference state, so the policy only chooses the operators applied to it. The policy's vocabulary is 110 operators from the first iteration, the size of the pool of the 35-orbital targets, so that the output head and token embedding keep their shape in Stage 2. Stage 1 runs for a fixed 150 iterations, without early stopping. This stage is run once per architecture (A1 or A5) and per circuit-length lineage ($L = 28$ or 44), not per target molecule, since the exact-statevector sampler's memory footprint scales as $\binom{n_o}{n_\alpha}\binom{n_o}{n_\beta}$ and becomes intractable above $n_o \approx 23$ — the curriculum therefore never includes the actual 46-/70-qubit target molecules in Stage 1. This Stage 1 is shared by the Stage-2 configurations described below.

Stage 2 — physical hardware continuation. The Stage-1 checkpoint initialises the policy for 150 further iterations on the target molecules (the 23- or 35-orbital tier of $Fe(CN)_6$/$Co(NH_3)_6$/ferrocene) with a shot-based sampler executing on the IBM quantum processing unit (QPU) `ibm_kobe` (Heron generation), rather than the statevector simulator, since Stage-1's qubit-count memory ceiling does not constrain shot-based hardware sampling. Circuits were submitted at 10,000 shots per circuit, batched 20 circuits per Qiskit Runtime job, in standard job execution mode rather than a persistent Runtime session to prevent prolonged device locking on shared multi-user allocations. The Stage-2 reward adds the Hartree–Fock penalty (Section 2.8.1).

What the warm start carries over. The vocabulary is 110 in both stages, so every policy tensor keeps its shape: the transformer layers, the position embedding, the cross-attention, the token embedding, the output head and the prompt encoder are all carried from Stage 1 into Stage 2. The price of the larger vocabulary is that operator ids above the $3\,n_{\mathrm{orb}} + 5$ real operators of a molecule's pool act as identity padding; on a 16-orbital Stage-1 molecule, $3 \cdot 16 + 5 = 53$ ids are real and 57 of 110 ids are padding (52%). Inference restores the whole checkpoint, including the prompt encoder. The Stage-1 and Stage-2 histories are single runs of 150 iterations each (Appendix G).

2.7.1 The two reward configurations

Stage 2 is run with two different rewards, which define the two training configurations (Table 2). Everything else is the same: the architecture, the optimiser (learning rate 5×10^{-6}, gradient-norm clipping at 8.0), the size curriculum, the seed (1234), the vocabulary of 110

Table 2: The two training configurations (Section 2.7.1); everything not listed is identical. The meaning of each reward weight is in Table 4. A variant without a penalty ($\lambda_{\mathrm{HF}} = 0$, $\lambda_{\mathrm{exc}} = 2.0$, $\lambda_{\mathrm{cov}} = 10.0$) was completed for A1-46q only (the other lineages were stopped after 17 or 18 iterations).

	Hartree–Fock penalty (main)	**Tuned reward**
Policy vocabulary, Stage 1 / Stage 2	110 / 110	110 / 110
Iterations, Stage 1 / Stage 2	150 / 150	150 / 150
λ_{HF}	0.05	0.07
λ_{spin}, c_H	0, 0	0.77, 0.26
λ_{exc}, λ_{cov}	2.0, 10.0	3.9, 7.2
Checkpoint used at inference	final	final (A1-70q: iteration 129)

operators, the fixed budget of 150 iterations per stage without early stopping, and the sampler and recovery settings. Both configurations are trained for the four lineages (A1 and A5 conditioning, each for 46 and for 70 qubits; Section 2.5).

Hartree–Fock penalty (the main configuration). The Stage-2 reward contains a penalty $\lambda_{\mathrm{HF}} = 0.05$ on the share of right-spin-sector shots that is the Hartree–Fock determinant. It lowers the reward of circuits whose valid shots are mostly the reference determinant, which a policy without this term converges to on the hardware (Appendix F). The final checkpoint is used at inference. The Stage-1 policy is shared with the second configuration.

Tuned reward (a comparison). The Stage-2 reward weights were chosen by a joint search on a noisy 24-qubit simulator (Section 2.8.1): $\lambda_{\mathrm{HF}} = 0.07$, a spin-balance term $\lambda_{\mathrm{spin}} = 0.77$, an entropy bonus $c_H = 0.26$, $\lambda_{\mathrm{exc}} = 3.9$ and $\lambda_{\mathrm{cov}} = 7.2$. For A1-70q the policy used at inference is the one saved at iteration 129, because the run ended by a segmentation fault at iteration 137.

The Hartree–Fock penalty is the main configuration because it is the simpler reward and, on the 43 test cases that both have, it gives the more reliable result at inference (Section 3.5). A third variant without any penalty was completed for one lineage only and is used in the training comparisons of the appendices, not at inference.

2.8 Hyperparameters and tuning

Settings. Table 3 lists the settings of the four lineages, Table 4 the Stage-2 reward weights of the two configurations (Table 2), and Table 5 the settings used at inference. Appendix C defines every setting by its name in the code (Table A3), and Appendix D describes the preliminary runs that fixed them, including the depth sweep that set the admission gate. The optimiser, the size curriculum, the seed and the recovery settings are the same for all runs.

2.8.1 The reward and its weights

The adjusted energy. For each of the $N = 20$ circuits sampled at an iteration, the policy is trained on an adjusted energy

$$E_i^{\mathrm{adj}} = E_i + g_i \left(\lambda_{\mathrm{exc}}\, x_i^{\mathrm{exc}} + \lambda_{\mathrm{cov}}\, x_i^{\mathrm{cov}}\right) + \lambda_{\mathrm{over}}\, x_i^{\mathrm{over}} + \lambda_{\mathrm{HF}}\, x_i^{\mathrm{HF}} + \lambda_{\mathrm{spin}}\, x_i^{\mathrm{spin}}, \tag{6}$$

where E_i is the SQD energy of the circuit (Hartree) and $g_i \in \{0, 1\}$ is one when E_i is within the auxiliary-term threshold (50 mHa) of the reference energy, distinct from the gradient-admission threshold of 100 or 500 mHa (Table 3), and zero otherwise. The group-relative objective then subtracts the batch mean and divides by the batch standard deviation of E^{adj}, so only the differences between circuits of one batch matter, not the absolute scale. Table 4 lists the five weights and the entropy coefficient. The terms are:

- x_i^{exc}: the share of sampled determinants with excitation rank above 3 relative to the Hartree–Fock determinant, from the raw counts. It favours a wavefunction dominated by singles, doubles and triples. Gated, because for a circuit far from the reference energy the rank distribution is not informative.
- x_i^{cov}: the share of qubits that the transpiled circuit leaves idle. It pushes the policy to use the qubits of the molecule. Gated.
- $x_i^{\mathrm{over}} = \log_{10}(d_i/d_{\max})$ when the transpiled depth d_i exceeds the ceiling $d_{\max}$ (Table 3), zero otherwise. It is the one depth term that is not gated, because on a hard molecule no circuit reaches the gate.
- x_i^{HF}: the share of right-spin-sector shots (those with the target alpha and beta electron counts) that are exactly the Hartree–Fock determinant. It penalizes a circuit that leaves the reference state unchanged, which the policy converged to in the first Stage-2 run. Ungated.
- $x_i^{\mathrm{spin}} = |w_\alpha - w_\beta|$, where w_α (w_β) is the share of right-spin-sector shots with an excitation in the alpha (beta) channel only. It is zero for a balanced state. Ungated.
- The entropy coefficient c_H acts on the loss and not on the reward, $\mathcal{L} \leftarrow \mathcal{L} - c_H\, \overline{H}$, with $\overline{H}$ the mean token entropy of the sampled sequences. It counters the collapse of the policy onto one sequence.

The reasons for the Hartree–Fock, spin and entropy terms are given in Appendix D. The three optional penal-

ties on a proxy gate count, on the diagonalisation subspace size and on the transpiled depth (Table A3) are zero throughout.

Reference energy in training. Each training molecule is associated with an approximate classical reference energy (CCSD(T) for the Stage-2 molecules). This reference enters only through the admission gate and two auxiliary reward terms (Section 2.7); the policy-gradient update depends on it only through these binary decisions, and not through backpropagation of the reference energy itself. At 70 qubits the admission gate passed 99–100% of iterations once the tolerance was loosened (Section 3.2), so in that regime it constrained the updates only weakly. We do not estimate an uncertainty of the reference energies. For systems where a suitable classical reference is unavailable, a different gating strategy would be required; we do not test such a reference-free protocol here. Inference uses no reference energy.

How the weights were chosen. The depth and coverage weights were fixed in preliminary runs, where the depth ceiling was the dominant lever of the admission gate. The Hartree–Fock weight was set to 0.05 in a test on a 24-qubit simulator with readout noise: it lowered the Hartree–Fock share of the valid shots without raising the raw error, whereas larger values removed that share only by moving the policy onto a single excited determinant. Because the five weights interact, we then searched them jointly (42 configurations, two seeds each) on a noisy 24-qubit simulator, and took the tuned reward from the configurations that met our criteria on raw error, Hartree–Fock share, policy diversity and circuit content. Two seeds cannot separate the passing configurations, so the search supports a region of the weight space and not one point in it, and it does not show that the tuned weights improve the hardware result, which is why the tuned reward is a comparison and the Hartree–Fock penalty is the main configuration (Section 3.5). The search, its criteria and its limits are given in Appendix D.

Comparison protocol. To compare the two rewards at inference, we ran the same zero-shot protocol with the policies of each configuration: the final policy of each run (for A1-70q with the tuned reward, the policy saved at iteration 129 of 137), the same lineages and molecules, one batch of ten circuits sampled at 100,000 shots, ranking at net dimension 40,000, and the same recoveries. A variant without any penalty is not evaluated, because three of its four Stage-2 runs were stopped after at most 18 iterations.

The remaining optimizer and curriculum settings (the learning rate, the gradient-norm clipping threshold, the fraction of iterations spent on the largest molecules, and the fraction of molecules counted as "largest"; exact definitions and code labels are in Table A3, Appendix C) were tuned to varying degrees of rigor, distinguished in Table A4. The fraction of molecules counted as "largest" was set by a controlled pair of runs with a clear, measured winner. The gradient-clipping threshold was recalibrated by measuring the policy's gradient-norm distribution rather than by a grid sweep. The learning rate's existing default was validated against, and beat, one higher alternative, and was not derived from a wider search. The fraction of iterations spent on the largest molecules was changed together with the fraction of molecules counted as "largest" and was never isolated in its own test.

The full per-parameter tuning narrative (what was varied, what the sweep/ablation found, and the setting this motivated) is reported in Table A4, Appendix E.

2.9 Computing resources

All classical computation (policy training, circuit generation, and the SQD ranking and recovery) ran on the roquo cluster at the RIKEN Center for Computational Science (R-CCS), a GPU supercomputer coupled to IBM quantum processors through the JHPC-quantum project. It has 135 nodes, each with 144 Arm (Grace) cores, about 1.7 TB of memory and four NVIDIA GB200 GPUs. Jobs are scheduled with Slurm; a typical job used a quarter of a node (one GPU, 36 cores), and the recoveries used two GPUs. The quantum circuits were executed on `ibm_kobe` through IBM's Qiskit Runtime service.

3 Results

3.1 Stage-1 (simulator) training

The Stage-1 policy is the first stage of both training configurations (Section 2.7); all numbers in this subsection are from it. It ran for a fixed 150 iterations per lineage.

Loss on admitted iterations (Figure 5 for the 46-qubit lineages, Figure A8 in Appendix G for the 70-qubit ones) starts at 2.2–2.5 (mean of the first three iterations) for all four lineages and falls below 0.1 after 58 (A1-46q), 60 (A5-46q), 92 (A1-70q) and 90 (A5-70q) iterations, ending at 0.002–0.010. The raw energy error of the same iterations (Figures A9 and A10) does not decay monotonically because the size curriculum cycles through the training molecules: it returns to small values on the small warm-up molecules and rises to tens of mHa on the larger ones in every cycle. Averaged over the admitted iterations of the last 30, the error against the reference is 36.1, 28.1, 27.6 and 21.1 mHa for A1-46q, A5-46q, A1-70q and A5-70q (Table A5). These errors are against a coupled-cluster or exact reference on the training molecules, which are not the inference molecules, so they bound the accuracy of the warm start but do not predict the zero-shot error (Section 3.5). The admission gate passes in 82–86% of the Stage-1 iterations (Figure A11, Appendix G).

Table 3: Final hyperparameters, per lineage. All four lineages share the same optimizer settings (base learning rate 5×10^{-6}, encoder LR multiplier 1.0, weight decay 0.01, grad clip 8.0) and the same size-curriculum schedule (warmup/focus/remainder fractions 0.25/0.30/0.45, top-40% molecules by qubit count treated as the curriculum's "large set"); only the values shown here differ across lineages. The recovery size D is the number of strings per spin sector kept during training (same quantity as at inference time, Table 5); the resulting determinant subspace contains up to D^2 determinants (Section 2.4). Code labels for every row are in Table A3.

Setting (code labels: Table A3)	A1-46q	A1-70q	A5-46q	A5-70q
Conditioning	scalar27	scalar27	POS-EGNN+s27	POS-EGNN+s27
Gate tokens per circuit, L	28	44	28	44
Policy params (M)	56.79	56.83	56.79	56.83
Encoder params (M)	3.45	3.45	3.22	3.22
Depth ceiling (Stage 1)	10000	20000	10000	20000
Depth ceiling (Stage 2)	10000	50000	10000	50000
Coverage penalty weight λ_{cov}	10.0	10.0	10.0	10.0
Depth-overshoot penalty weight λ_{over}	2.0	2.0	2.0	2.0
Excitation penalty weight λ_{exc}	2.0	2.0	2.0	2.0
Energy admission tolerance, mHa (Stage 1)	100	100	100	100
Energy admission tolerance, mHa (Stage 2)	100	500	100	500
Recovery size D (strings per spin; Stage 1)	20000	20000	20000	20000
Recovery size D (strings per spin; Stage 2)	40000	40000	40000	40000
Shots per circuit (Stage 2)	10000	10000	10000	10000
Stage-2 backend	ibm_kobe	ibm_kobe	ibm_kobe	ibm_kobe

Table 4: Reward and loss weights of the Stage-2 runs; Section 2.8.1 defines every term, and Appendix D describes how the weights were tuned. The first five terms add to the per-circuit reward energy E_i^{adj} (Hartree); the entropy coefficient enters the loss. A "gated" term is applied only when the circuit's energy is within the auxiliary-term threshold (50 mHa; Table A3), distinct from the gradient-admission threshold of 100 or 500 mHa (Table 3), of the reference energy; the other terms always apply. "Reference" is the setting of the preliminary runs and of the runs without a penalty, "HF penalty" changes only λ_{HF}, and "Tuned" is the configuration selected by the joint search of Section 2.8.1 (rounded to the values submitted). The search range is the interval explored. The optional penalties on a proxy gate count, on the diagonalisation subspace size and on the transpiled depth (Table A3) are zero in every run. The code label of each weight is the entry in the "Flag" column; all labels are defined in Table A3.

Term	**Flag**	**Quantity x_i of circuit i**	**Gated**	**Range**	**Ref.**	**HF pen.**	**Tuned**
λ_{exc}	`cost-lambda-excitation`	weight of excitation ranks >3 among the sampled determinants	yes	0–4	2.0	2.0	3.9
λ_{cov}	`cost-lambda-coverage`	idle-qubit fraction of the transpiled circuit	yes	0–20	10.0	10.0	7.2
λ_{over}	`cost-lambda-overshoot`	$\log_{10}(d_i/d_{\max})$ when the depth d_i exceeds the ceiling $d_{\max}$	no	fixed	2.0	2.0	2.0
λ_{HF}	`cost-lambda-hf`	share of right-spin-sector shots that are the Hartree–Fock determinant	no	0–0.3	0	0.05	0.07
λ_{spin}	`cost-lambda-spin`	|alpha-only − beta-only| share of right-spin-sector shots	no	0–1	0	0	0.77
c_H	`entropy-coef`	mean token entropy of the sampled sequences (nats), subtracted from the loss	—	0–0.3	0	0	0.26

The depth of the best of the 20 circuits sampled in each iteration (Figures A12 and A14; 70-qubit lineages in Figures A13 and A15, Appendix G) grows during Stage 1: its median over the first ten iterations is 436–1158 and over the last ten 3997 (A1-46q), 6634 (A5-46q), 8384 (A1-70q) and 10,454 (A5-70q), as the policy learns to put more gates on the 16-orbital simulator molecules, where depth is not penalised by a device. The Stage-2 values (right of the dashed line in the same figures) are 50–150, but the two stages use different molecules and samplers and are not a like-for-like comparison (Section 4).

Sampling dominates Stage-1 wall time (Figure 6, left; Table A7): 284–539 s per iteration out of 496–756 s in total (57–71%), even though Stage 1 uses an exact-statevector `ffsim` sampler rather than hardware: the

Table 5: Inference-time hyperparameters (zero-shot campaign on unseen molecules, Section 3.3), unique to inference and not already listed in Table 3. All values are uniform across every lineage and every target molecule in this campaign – unlike training (Table 3), nothing here varies by architecture or qubit-count tier. The recovery size D is the same quantity as in training (Table 3, Section 2.4): the final-recovery number of strings per spin sector, which can be changed with zero further QPU calls once a lineage's circuits are sampled and cached (Fig. 2). The ranking cap (distinct from the training recovery size, although the code flag has the same name) is a separate, smaller cap used only to rank the 10 sampled circuits against each other before the final recovery – not swept, intentionally cheap: 200 strings per spin sector (net dimension 40,000) in the fresh-shot campaign. The E+D rule uses the same ranking.

Parameter (code labels in Table A3)	Value
Shots per circuit	100000
Recovery size D (strings per spin sector; final recovery; the cache allows other values)	20000
Final recovery iterations K	25
Ranking cap (pre-recovery ranking only; fresh-shot campaign)	200 strings per spin sector (net dimension 40,000)
Depth limit for diagonalisation (70q lineages, matches Stage-2 training ceiling)	50000

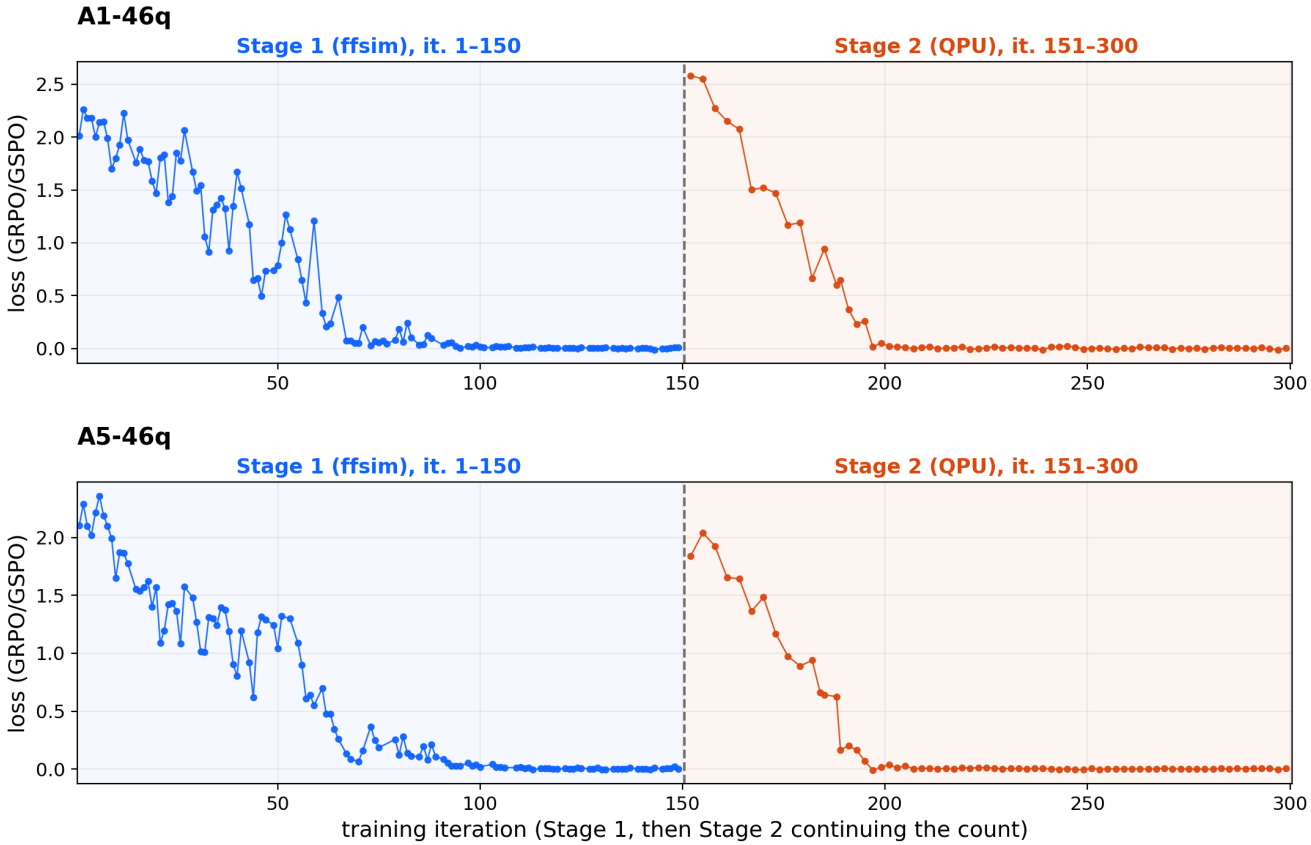

Figure 5: Training loss (GRPO/GSPO) of the 46-qubit lineages, A1-46q (top) and A5-46q (bottom): Stage 1 (ffsim simulator, iterations 1–150, left of the dashed line) followed by Stage 2 with the Hartree–Fock penalty (`ibm_kobe` QPU, iterations 151–300, right of the line), with the iteration count continuing across the two stages. Points are admitted iterations only; an iteration that fails the admission gate triggers no update and has no loss. The 70-qubit lineages are in Figure A8 (Appendix G).

cost is the policy's autoregressive generation of L gate tokens for 20 circuits per iteration and their simulation, not a hardware round-trip, and it is largest for the 70-qubit lineages (427–539 s) whose sequences have 44 gates. Diagonalization (including recovery) is the second-largest cost at 190–196 s; refine, transpile, and backward are each under 15 s. A Stage-1 run takes 20.7–31.5 h.

The gradient-norm breakdown by model component (Figure 7, left of each dashed line) shows the policy transformer's gradient norm exceeding the prompt encoder's, by a median factor of 7.3–10.8 across all four lineages in Stage 1 – consistent with the parameter-count asymmetry between the two components (a 56.8M-parameter policy backbone vs. a 3.2–3.4M-parameter prompt encoder, Section 2.5). Both norms decay as training con-

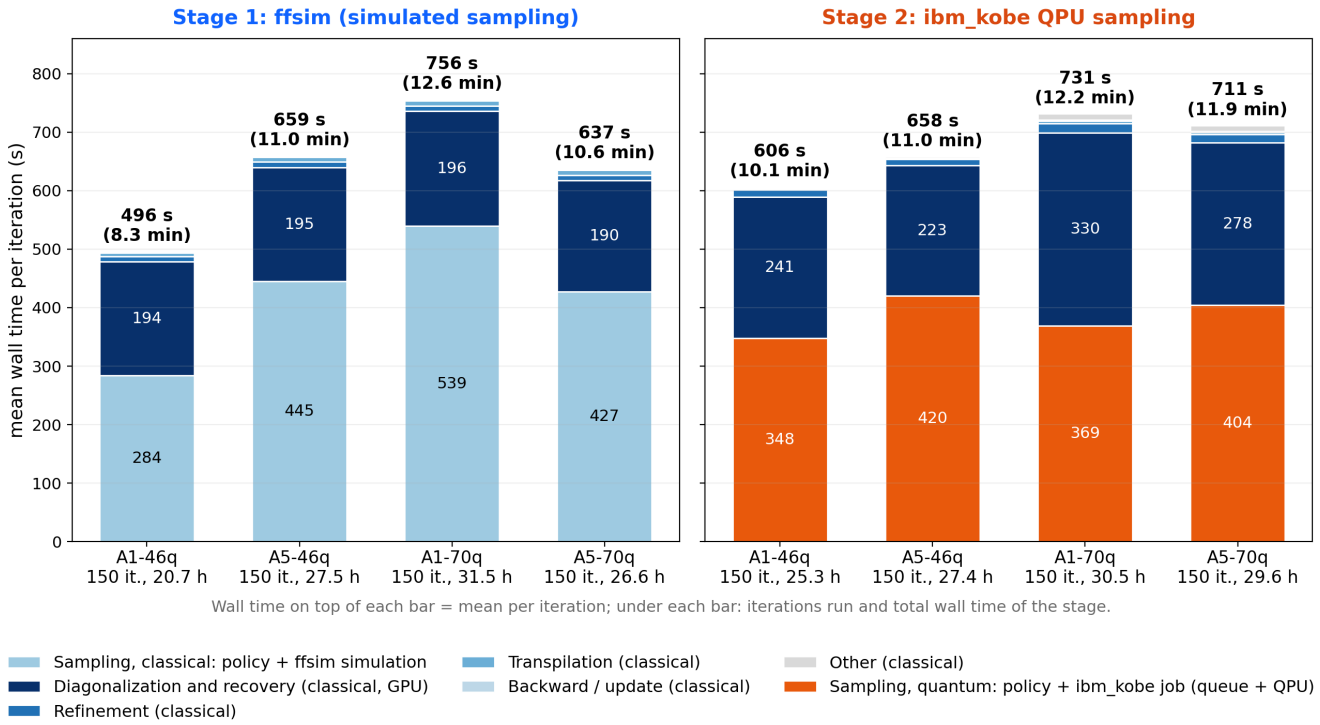

Figure 6: Mean wall time per training iteration by pipeline step, for Stage 1 (left) and Stage 2 with the Hartree–Fock penalty (right) and the four lineages. The sampling step (policy generation of the gate sequences, then an `ffsim` simulation in Stage 1 or an IBM Runtime job on `ibm_kobe`, including queue time and QPU execution, in Stage 2) is the only step that involves the quantum computer (orange); every other step is classical (blue shades): diagonalization and recovery on the GPUs, refinement, transpilation, and the backward pass. Bold numbers on top are the mean wall time per iteration; below each bar are the iterations run and the total wall time of the stage. One A5-46q, A1-70q and A5-70q iteration whose sampling exceeded one hour (device-queue stalls) is excluded from the Stage-2 means.

verges and the encoder's never reaches the policy's. This is the closest per-component analogue this architecture logs to a separate "encoder loss / decoder loss" decomposition, since both components are optimized jointly against a single scalar GRPO/GSPO objective (Section 2.3).

3.2 Stage-2 (hardware) training with the Hartree–Fock penalty

On hardware, all four lineages run for 150 iterations and settle on shallow circuits, but the sampled circuits narrow to a small family. The loss on admitted iterations (Figures 5 and A8) restarts at 1.9–2.9 and falls below 0.1 after 18 (A1-46q and A5-46q), 48 (A1-70q) and 51 (A5-70q) iterations, ending at 0.001–0.004. The raw error on the hardware targets (Figures A9 and A10) is nearly flat, with a mean over the stage of 46.2 and 46.5 mHa for the 46-qubit lineages and 300.6 and 294.1 mHa for the 70-qubit lineages (Table A5).

The admission gate passes 100% and 99% of the iterations for A1-70q and A5-70q, and 46% and 47% for A1-46q and A5-46q (Figure A11). The depth of the best of the 20 circuits is low from the start and stays so; its median over the last ten iterations is 56 (A1-46q), 54 (A5-46q), 76 (A1-70q) and 152 (A5-70q), against 44–66 in the first ten (Figures A12–A15).

The chemistry of the sampled circuits (Appendix F) shows the narrowing. By iteration 50–55 the effective vocabulary of the policy is about one token. In the best circuit, 6–21% of the shots are in the right spin sector, and 0–4% of those are the Hartree–Fock determinant (60–76% in Stage 1). The 46-qubit circuits are strongly beta-only, with a beta-only minus alpha-only weight difference of 0.56 (A1-46q) and 0.67 (A5-46q). The tuned reward keeps an effective vocabulary of 21–27 (Table A5). Per-iteration times and billed QPU time are given in Tables A7 and A8 (Appendix E) and the total cost in Section 3.9. Each result is from one seed per lineage.

3.3 Zero-shot inference on unseen molecules

Every lineage's trained policy of the HF-penalty recipe (A1/A5, 46q-trained and 70q-trained) was run zero-shot on the unseen molecules of the $Sn(OH)_3$-Me family – the small, medium, large and extra-large active-space tiers of the bare molecule and its mono- and di-hydrate adducts, none of which appear in any lineage's own

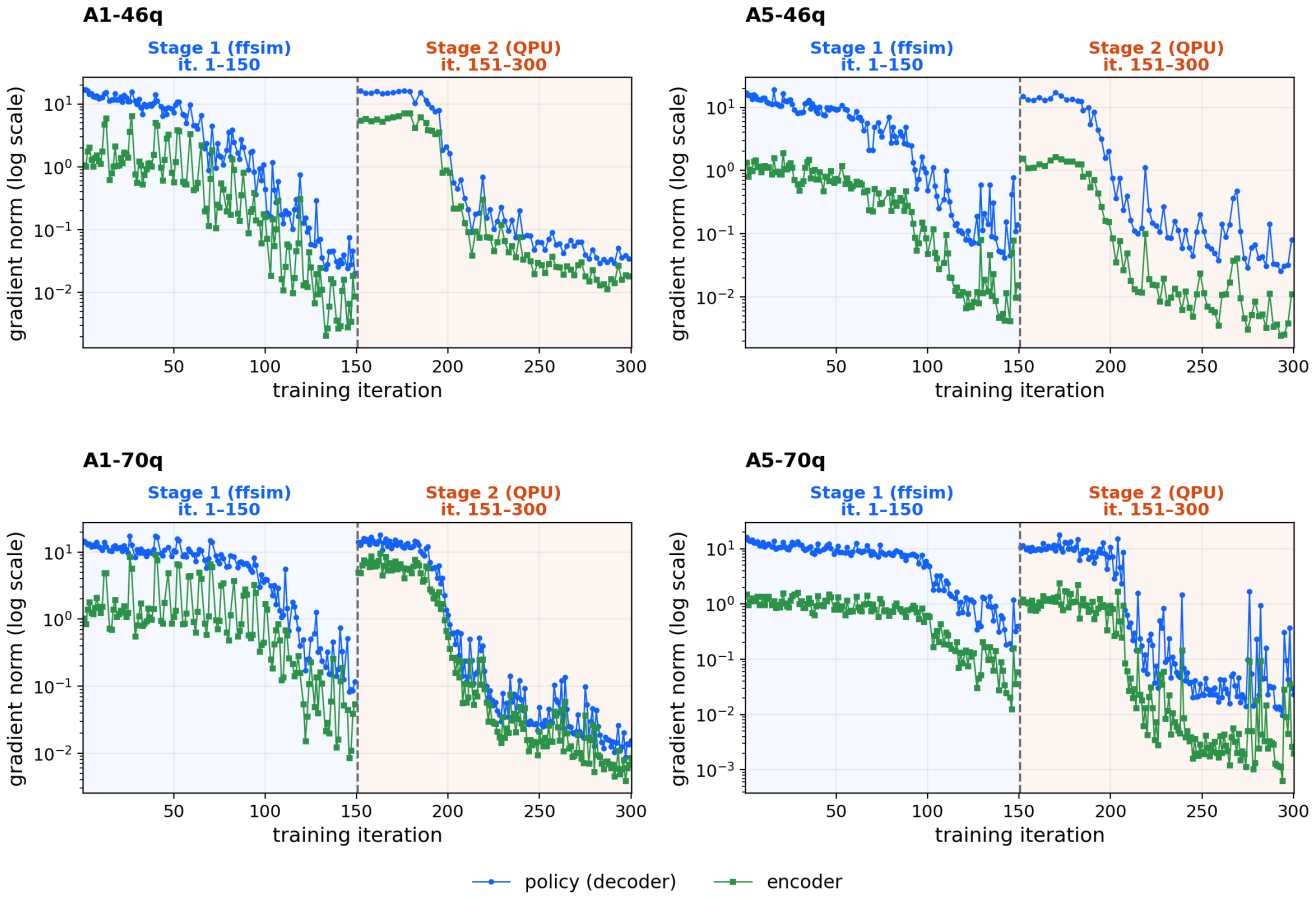

Figure 7: Gradient norm of the policy transformer (decoder) and of the prompt encoder on admitted iterations, for the four lineages (log scale). In each panel Stage 1 (ffsim) is left and Stage 2 (`ibm_kobe` QPU, HF penalty) right of the dashed line, with the iteration count continuing across the stages.

training curriculum. All four lineages are run on all eleven tiers (44 runs, all measured on `ibm_kobe`; 43 of the Hartree–Fock-penalty runs have a result in the table, see below). Table 6 is the comprehensive summary, trained policies only; the matched untrained controls (eight runs on three molecules; the controls are randomly initialized and so do not depend on the training recipe) and the full per-molecule detail – excitation-level composition, circuit-cost/coverage breakdowns, per-stage timing – are in Appendix J. Section J.1 explains why some lineage/molecule combinations are a matched-norb case and others a cross-norb vocab-mismatch ablation (the lineage's own training norb differs from the target's). Figures 8 and A16 show, for the A1 policy trained at 46 and at 70 qubits, the raw excitation composition and as-submitted depth of all ten sampled circuits of every compound; they make visible that the selected circuits are the ones with weight at low excitation ranks and small depth, which Section 3.8 quantifies.

Summary of the HF-penalty runs. Of the 44 runs, 43 have a recovered best circuit; the other (SnOH3-me LL with A5-46q) produced valid shots for one circuit only, so the ranking found no best and second pair (the selections best, second, merged and E+D are defined in Section 2.1 and Figure 2). Sampling the ten circuits again with a different seed gave eight circuits with valid shots but a recovered error of about 200 mHa (Appendix N), so this test case is not reported. For the tuned reward, the one run with a single usable circuit (SnOH3-me-2h2o S with A5-70q) was sampled again with a different seed, which gave five usable circuits and a best-circuit error of 0. 07 mHa; this second draw is the one reported (Appendix N). The median error of the best circuit against DMRG is 2.8 mHa; 13 runs are within 1 mHa, 28 within 5 mHa and 42 within 15 mHa. The eight untrained controls have a median error of 106 mHa (one within 15 mHa, at the 56-qubit tier, where the untrained policy happened to pick a depth-89 circuit). The comparison is paired by molecule and lineage and uses the same sampling, ranking and recovery. The depth of the selected circuit is 30–168 (median 55) and its two-qubit-gate count 36–160 (median 84), against a median depth of 12,300 for the untrained controls' selected circuits. The error grows with the size of the active space, as expected for a fixed recovery size of 20,000 strings per spin sector; the exceptions at small size are the SnOH3-

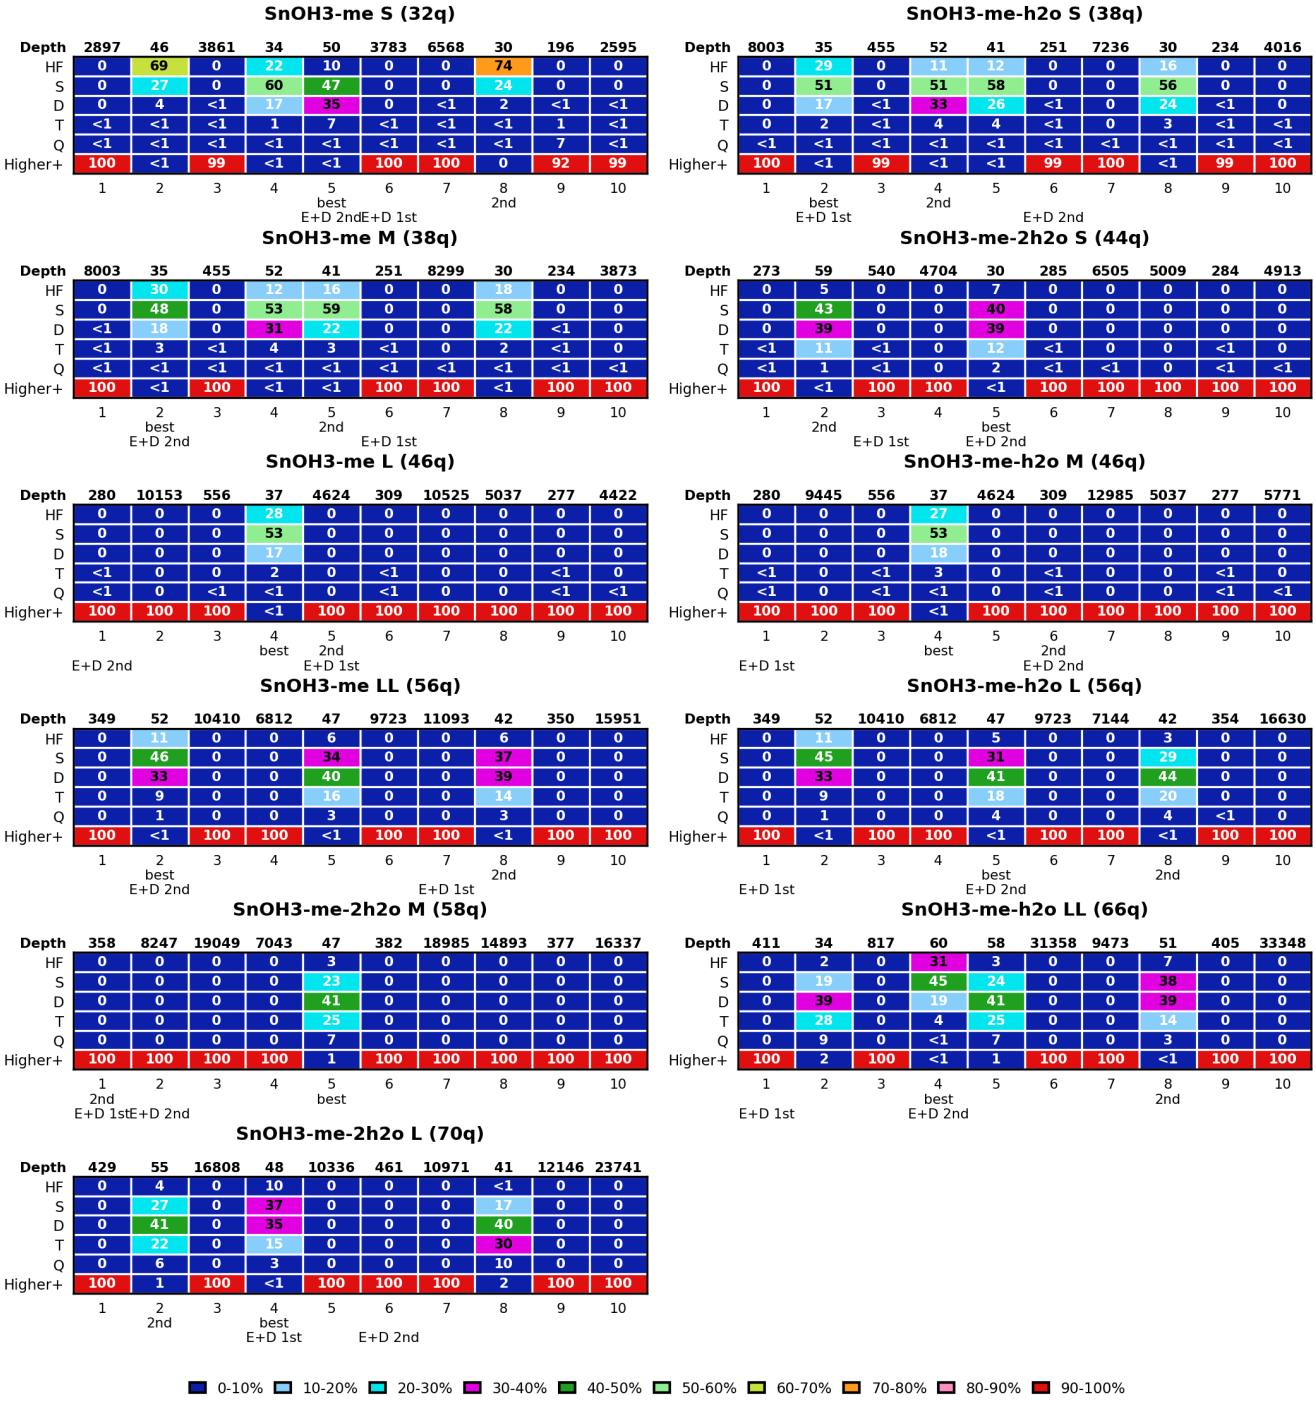

Figure 8: Raw-shot excitation composition (HF, S, D, T, Q, higher; percent of 100,000 shots) of the ten circuits sampled by the A1 policy trained at 46 qubits with the Hartree–Fock penalty, for each of the eleven compounds, ordered from the smallest to the largest active space. Numbers above each column are the circuit depth as-submitted to `ibm_kobe`; "best" and "2nd" mark the two circuits with the lowest ranking energy, and "E+D 1st" and "E+D 2nd" the two circuits chosen by the combined energy and distinct-bitstring rank (a circuit can carry two labels). In every compound the best circuit puts at most 1.2% of its shots in "higher" and is among the four shallowest of the ten (the shallowest in four of the eleven panels); the second circuit does so in eight of the eleven panels, and in the SnOH3-me L, SnOH3-me-h2o M and SnOH3-me-2h2o M panels it has essentially all its weight in "higher". Most of the other circuits are one to three orders of magnitude deeper with all their weight in "higher".

Table 6: Zero-shot inference comprehensive summary, trained policies only, every lineage×molecule combination. Each table entry is energy (E_h)/error vs. DMRG (mHa). Depth/2q are the depth and two-qubit-gate count of the best circuit as transpiled and submitted to `ibm_kobe` (live transpilation at Qiskit optimization level 1); the offline estimate on the `ibm_heron_r2` basis is not used. The distance of the generated circuits from the idealised LUCJ circuit is in Section 3.6. Matched untrained controls are in Appendix J. All values come from the fresh-shot re-ranking campaign: every run's ten sampled circuits were sampled on `ibm_kobe` with 100,000 new shots each (no error mitigation), and the raw counts and transpiled circuits were stored. Circuits are ranked by a one-shot diagonalization of the electron-number-conserving shots in a subspace of up to 200 most-measured alpha strings (net CI dimension 40,000); circuits with no valid shot have no ranking energy and are excluded. "Best"/"second" are the two lowest-ranking-energy circuits, "merged" pools the shots of every circuit with a finite ranking energy, and "E+D 1"/"E+D 2" are the two circuits with the lowest sum of two ranks, the rank of the ranking energy (the same 40,000-determinant ranking) and the rank by number of distinct measured bitstrings, recovered on the same fresh shots. Every recovery uses $D = 20{,}000$ strings per spin sector and 25 iterations. Errors are relative to the DMRG reference.

Molecule (tier, qubits)	Arch. (lineage)	Best E_h/mHa	Second E_h/mHa	Merged E_h/mHa	E+D 1/2 (mHa)	Depth	2q
SnOH3-me S (32q)	A1 (A1-46q)	−479.193266/0.12	−479.192672/0.71	−479.193374/0.01	0.01/0.12	50	56
SnOH3-me S (32q)	A1 (A1-70q)	−479.193164/0.22	−479.193270/0.11	−479.193374/0.01	0.01/0.01	41	68
SnOH3-me S (32q)	A5 (A5-46q)	−479.192727/0.65	−479.193081/0.30	−479.193374/0.01	0.30/0.01	31	36
SnOH3-me S (32q)	A5 (A5-70q)	−479.192924/0.46	−479.193134/0.25	−479.193374/0.01	0.01/0.46	69	60
SnOH3-me-h2o S (38q)	A1 (A1-46q)	−555.108538/0.29	−555.108684/0.15	−555.108826/0.01	0.29/0.01	35	52
SnOH3-me-h2o S (38q)	A1 (A1-70q)	−555.108766/0.07	−555.108751/0.08	−555.108826/0.01	0.01/0.01	57	92
SnOH3-me-h2o S (38q)	A5 (A5-46q)	−555.108713/0.12	−555.108826/0.01	−555.108826/0.01	0.12/0.01	34	40
SnOH3-me-h2o S (38q)	A5 (A5-70q)	−555.108740/0.09	−555.108764/0.07	−555.108826/0.01	0.09/0.19	110	84
SnOH3-me M (38q)	A1 (A1-46q)	−479.292683/9.6	−479.296591/5.7	−479.286434/15.8	152.0/9.6	35	52
SnOH3-me M (38q)	A1 (A1-70q)	−479.298463/3.8	−479.295672/6.6	−479.282569/19.7	144.4/139.1	57	92
SnOH3-me M (38q)	A5 (A5-46q)	−479.285058/17.2	−479.282828/19.5	−479.271416/30.9	143.9/130.6	32	52
SnOH3-me M (38q)	A5 (A5-70q)	−479.288699/13.6	−479.152486/149.8	−479.265913/36.4	131.8/149.1	152	104
SnOH3-me-2h2o S (44q)	A1 (A1-46q)	−631.083103/0.21	−631.083239/0.07	−631.083269/0.04	0.04/0.21	30	60
SnOH3-me-2h2o S (44q)	A1 (A1-70q)	−631.083257/0.06	−631.083084/0.23	−631.083269/0.04	0.06/0.04	60	100
SnOH3-me-2h2o S (44q)	A5 (A5-46q)	−631.083185/0.13	−631.083207/0.11	−631.083267/0.04	0.11/0.06	49	64
SnOH3-me-2h2o S (44q)	A5 (A5-70q)	−631.083092/0.22	−631.083044/0.27	−631.083269/0.04	0.22/0.27	168	104
SnOH3-me L (46q)	A1 (A1-46q)	−479.299926/14.3	−479.121483/192.7	−479.258786/55.4	192.7/198.1	37	64
SnOH3-me L (46q)	A1 (A1-70q)	−479.306409/7.8	−479.116356/197.9	−479.272504/41.7	193.1/198.6	69	100
SnOH3-me L (46q)	A5 (A5-46q)	−479.301748/12.5	−479.306572/7.6	−479.284378/29.8	199.6/12.5	55	80
SnOH3-me L (46q)	A5 (A5-70q)	−479.308601/5.6	−479.297326/16.9	−479.284631/29.6	199.1/5.6	144	124
SnOH3-me-h2o M (46q)	A1 (A1-46q)	−555.123748/1.4	−555.102442/22.7	−555.118018/7.1	22.7/22.7	37	64
SnOH3-me-h2o M (46q)	A1 (A1-70q)	−555.124450/0.71	−555.102207/23.0	−555.118773/6.4	23.0/22.6	69	100
SnOH3-me-h2o M (46q)	A5 (A5-46q)	−555.123672/1.5	−555.123151/2.0	−555.123131/2.0	1.4/2.0	58	64
SnOH3-me-h2o M (46q)	A5 (A5-70q)	−555.123764/1.4	−555.123482/1.7	−555.124246/0.91	1.4/1.7	108	100
SnOH3-me LL (56q)	A1 (A1-46q)	−479.300183/14.9	−479.299292/15.8	−479.286822/28.2	203.8/14.9	52	80
SnOH3-me LL (56q)	A1 (A1-70q)	−479.304625/10.4	−479.112860/202.2	−479.279821/35.2	202.2/202.1	62	124
SnOH3-me LL (56q)	A5 (A5-70q)	−479.301896/13.2	−479.111752/203.3	−479.261236/53.8	203.3/204.3	79	104
SnOH3-me-h2o L (56q)	A1 (A1-46q)	−555.132208/2.8	−555.132421/2.6	−555.129520/5.5	33.1/2.8	47	80
SnOH3-me-h2o L (56q)	A1 (A1-70q)	−555.133193/1.9	−555.132552/2.5	−555.127221/7.8	33.1/33.0	65	124
SnOH3-me-h2o L (56q)	A5 (A5-46q)	−555.131417/3.6	−555.132660/2.4	−555.129768/5.3	33.1/3.6	53	64
SnOH3-me-h2o L (56q)	A5 (A5-70q)	−555.132405/2.7	−555.132408/2.6	−555.130714/4.3	2.7/33.1	44	76
SnOH3-me-2h2o M (58q)	A1 (A1-46q)	−631.112084/6.5	−631.077480/41.1	−631.095105/23.5	41.1/41.1	47	88
SnOH3-me-2h2o M (58q)	A1 (A1-70q)	−631.114317/4.3	−631.113058/5.6	−631.106177/12.4	41.1/41.1	65	132
SnOH3-me-2h2o M (58q)	A5 (A5-46q)	−631.114080/4.5	−631.113620/5.0	−631.105052/13.6	41.1/41.1	45	72
SnOH3-me-2h2o M (58q)	A5 (A5-70q)	−631.113705/4.9	−631.111954/6.7	−631.110745/7.9	41.1/41.1	62	68
SnOH3-me-h2o LL (66q)	A1 (A1-46q)	−555.135328/2.1	−555.133503/3.9	−555.131333/6.1	35.4/2.1	60	92
SnOH3-me-h2o LL (66q)	A1 (A1-70q)	−555.134846/2.6	−555.102002/35.4	−555.131599/5.8	2.6/35.4	51	156
SnOH3-me-h2o LL (66q)	A5 (A5-46q)	−555.133877/3.5	−555.134389/3.0	−555.130850/6.5	35.4/3.5	44	84
SnOH3-me-h2o LL (66q)	A5 (A5-70q)	−555.135231/2.2	−555.134346/3.1	−555.133797/3.6	3.1/2.2	57	84
SnOH3-me-2h2o L (70q)	A1 (A1-46q)	−631.123614/8.6	−631.123272/9.0	−631.119423/12.8	8.6/54.8	48	100
SnOH3-me-2h2o L (70q)	A1 (A1-70q)	−631.124552/7.7	−631.077480/54.8	−631.099821/32.4	54.8/7.7	60	160
SnOH3-me-2h2o L (70q)	A5 (A5-46q)	−631.124857/7.4	−631.123201/9.0	−631.122669/9.6	54.8/7.2	44	96
SnOH3-me-2h2o L (70q)	A5 (A5-70q)	−631.120256/12.0	−631.124322/7.9	−631.110701/21.5	12.0/54.8	134	128

me M (38q) and SnOH3-me L (46q) tiers (best-circuit errors 3.8–17.2 and 5.6–14.3 mHa), where the non-hydrate molecule is harder than its size suggests. At the three S tiers, the recovery cap of 20,000 strings exceeds the number of α strings (1,820–7,315; Table A2), so the cap does not restrict the strings that can be retained, although

the basis still depends on the sampled configurations. The near-exact merged energies at these tiers therefore give limited evidence that the circuits preferentially sample important configurations. Larger spaces, where the cap is smaller than the full α-string space, are a more demanding test. For SnOH3-me M (38 qubits), the cap covers at most 39.7% of the 50,388 α strings, which may contribute to its larger errors.

The matched untrained controls are compared with the trained policies in Appendix J (Table A19). Paired by lineage, the trained policy is more accurate in all eight comparisons. At SnOH3-me L the four controls give 188–200 mHa against 5.6–14.3 mHa for the four trained policies, and at SnOH3-me-h2o M the three controls give 22.6–22.9 mHa against 0.71–1.5 mHa. At SnOH3-me LL the margin is small (10.4 against 12.3 mHa for A1-70q), and the control lies, inside the 10.4–14.9 mHa range of the three trained policies, so the size of the separation depends on the molecule.

3.4 Circuit cost: trained policy vs. pure LUCJ ansatz

The transpilation costs and the hardware results of the LUCJ benchmark are collected in Appendix K.

The relevant LUCJ comparator is the circuit that was benchmarked on hardware: the spin-balanced LUCJ operator of the `ffsim` library [34] built from CCSD amplitudes, with one ansatz repetition (the shallowest variant). For the generated circuits, depths and two-qubit-gate counts are those of the circuits as transpiled at Qiskit optimization level 1 [26] and submitted to `ibm_kobe`; for LUCJ, we report our own transpilation of the circuit that was run. The LUCJ circuit has depth 356 and 3040 two-qubit gates at 46 qubits, and depth 428–536 and 4248–6952 two-qubit gates at 56–70 qubits (Tables A22 and A23).

Selected circuits. At 56–70 qubits the best circuits of the 19 HF-penalty runs (SnOH3-me LL, SnOH3-me-h2o L, SnOH3-me-2h2o M, SnOH3-me-h2o LL and SnOH3-me-2h2o L, with the lineages that were run) have depth 44–134 and 64–160 two-qubit gates, i.e. 4–12× shallower and 34–72× fewer two-qubit gates than the LUCJ circuit of the same molecule (428–536 deep with 4,248–6,952 two-qubit gates). At the 46-qubit hydrate (SnOH3-me-h2o M) the four trained lineages' best circuits have depth 37–108 and 64–100 two-qubit gates against 356 and 3040 for LUCJ. No run selected a very deep circuit: the largest selected depth is 168. The shallowest selected circuits are not always the most accurate: the deepest of the 56–70-qubit selections (A5-70q at SnOH3-me-2h2o L, depth 134) has an error of 12.0 mHa, and the A5-70q SnOH3-me LL circuit (depth 79) 13.2 mHa; the two largest errors of the 19 runs (14.9 and 13.2 mHa) are at SnOH3-me LL.

Typical circuits. The selected circuits are not typical of what the policy samples. Over all ten circuits of each HF-penalty run at 56–70 qubits, the median as-submitted depth is between 58 and 12,424 (median over the 20 runs, 5,960; 20 rather than 19 because the run without a recovered best circuit has sampled circuits), and only 1–9 of the ten circuits per run are shallower than the LUCJ circuit; the A1-70q runs are the deepest (6,515–12,424) and the A5-70q runs the shallowest (58–6,733). The cost advantage therefore depends on selecting the right circuit among the ten; the selection (lowest ranking energy) is made after sampling, so the advantage belongs to the sample-and-select procedure and not to every circuit the policy emits. A smaller-scale version of this comparison (16 qubits, $N_2/CH_4/H_2O$/benzene) previously found the trained policy's circuits 60–180× shallower in two-qubit-gate count at matched accuracy. The LUCJ depths in Table A23 are from our own transpilation (not re-extracted from the submitted LUCJ job), so the depth ratios should be read as indicative.

The LUCJ circuit was also run on hardware with the pipeline of the other campaigns of this study: 100,000 shots on `ibm_kobe`, a 25-iteration GPU recovery with $D = 20{,}000$ strings per spin sector, and no error mitigation, dynamical decoupling or twirling. Unlike the trained GQE policy, LUCJ is a single fixed ansatz and not a sampled distribution: there is one circuit per molecule, sampled once and recovered once, so no best/second-best selection applies. Table A22 reports all completed results across all 9 target molecules in the experimental campaign; Table A23 reports the corresponding raw-shot excitation distributions evaluated across the 100,000 physical shots (following the same convention as Table A21) alongside transpiled circuit metrics.

The LUCJ results show a trend across the nine target molecules. At the three smallest active spaces (SnOH3-me S at 32 qubits, SnOH3-me-h2o S at 38 qubits and SnOH3-me-2h2o S at 44 qubits), LUCJ reaches the DMRG reference energy to within 0.04 mHa, as do the GQE lineages; there the recovery cap of 20,000 strings already covers the whole α-string space (Table A2). From 46 qubits onward, the raw error of LUCJ, that is, the error of the energy recovered from the sampled bitstrings, ranges from 22.7 to 204.1 mHa; it is 22.7–54.8 mHa for all targets except SnOH3-me LL (56 qubits), an outlier at 204.1 mHa. We also report the error after singles injection: the single excitations of the Hartree–Fock determinant are added to the recovered set of configurations, and the Hamiltonian is diagonalized again in the enlarged basis (Appendix J). This lowers all errors from 46 qubits onward to 5.2–12.6 mHa, except SnOH3-me LL, which falls from 204.1 to 41.2 mHa. At SnOH3-me-h2o M (46 qubits), the one target where LUCJ and the untrained GQE policies can be compared on identical footing, the singles-injected error of LUCJ (5.2 mHa) is close to that of the three untrained GQE lineages

Table 7: Lowest recovered SQD energy vs. DMRG reference for every SnOH3-me-family target, with the signed deviation $\Delta E = E_{\mathrm{SQD}} - E_{\mathrm{DMRG}}$ (all positive), taken over all lineages, policies and result types (best, second, merged, E+D 1/2) in Tables A15 and A19; the lineage and type that achieved it are noted. DMRG references are from Table A1 (maximum bond dimension 4000, 17 sweeps). Selecting the minimum over many results is a best-of-k statistic, not a single pre-declared run.

Target	DMRG (E_h)	Best SQD (E_h)	ΔE (mHa)	Lineage	Type
SnOH3-me S (32q)	−479.193382	−479.193374	<0.01	A1-46q trained	merged
SnOH3-me-h2o S (38q)	−555.108833	−555.108826	<0.01	A1-46q trained	merged
SnOH3-me M (38q)	−479.302278	−479.298463	3.8	A1-70q trained	best
SnOH3-me-2h2o S (44q)	−631.083312	−631.083269	0.04	A1-46q trained	merged
SnOH3-me L (46q)	−479.314216	−479.308601	5.6	A5-70q trained	best
SnOH3-me-h2o M (46q)	−555.125158	−555.124450	0.71	A1-70q trained	best
SnOH3-me LL (56q)	−479.315057	−479.304625	10.4	A1-70q trained	best
SnOH3-me-h2o L (56q)	−555.135058	−555.133193	1.9	A1-70q trained	best
SnOH3-me-2h2o M (58q)	−631.118616	−631.114317	4.3	A1-70q trained	best
SnOH3-me-h2o LL (66q)	−555.137396	−555.135328	2.1	A1-46q trained	best
SnOH3-me-2h2o L (70q)	−631.132238	−631.125022	7.2	A5-46q trained	E+D 2

(4.9–5.1 mHa, Table A19), although the LUCJ circuit is much deeper (depth 356 against 37–108 for the trained GQE circuits). The best circuits selected by the ranking for the four trained Hartree–Fock-penalty lineages reach 0.71–1.5 mHa at this target, with as-submitted depths of 37–108 and 64–100 two-qubit gates, so at SnOH3-me-h2o M the trained policy is better than LUCJ in accuracy and in circuit cost. This rests on one molecule and one fresh hardware sample per circuit. The GQE circuit was selected from ten by ranking energy, not by comparison with DMRG, whereas LUCJ is a single circuit with no such selection step.

3.5 Hartree–Fock penalty and tuned reward

We compare the two training configurations at inference, with the protocol of Section 2.8.1, on the molecule–lineage pairs for which both have a best and a second circuit. A lineage is one trained policy (A1 or A5, trained at 46 or 70 qubits), and each molecule–lineage pair is one test case: that policy run once on one of the eleven $Sn(OH)_3$-Me active spaces, giving $4 \times 11 = 44$ test cases per configuration. The per-test-case values are in Appendix H, and the per-run results of each configuration are in the reports described in Appendix P. Table 8 gives the 43 of 44 test cases for which both have a best and a second circuit; for one Hartree–Fock-penalty run no pair exists (Appendix N). All entries are medians over these test cases.

The two sets are close, and which one is better depends on the statistic. The tuned reward has the lower median error of the best circuit (2.4 mHa against 2.8 mHa for the Hartree–Fock penalty), of the better of the best and second circuit (2.4 against 2.6 mHa), and is the lower of the two in 26 of the 43 test cases. The Hartree–Fock penalty has the lower mean (4.8 mHa against 10.5) and the more reliable tail: none of its best circuits is above 30 mHa, against three for the tuned reward (41, 41 and 198 mHa, at SnOH3-me-2h2o M with the 46-qubit lineages and at SnOH3-me L with A5-70q). The two sets are within 5 mHa in 28 and 24 of the 43 test cases, and within 1 mHa in 13 and 12. The second circuit is less accurate for the tuned reward (median 7.3 mHa against 5.0 mHa), and the selected circuits have the same depth (median 55 and 56).

The circuits of the two sets differ in the chemistry of their shots. The best circuits of the Hartree–Fock penalty have the larger share of shots in the right spin sector (8.0% against 4.4%), a larger Hartree–Fock share of the valid shots (50% against 24%), fewer alpha-only shots (10.7% against 14.5%) and a smaller both-spin share (4.4% against 14.8%), with similar beta-only shares (21% and 19%).

3.6 Divergence of the generated circuits from the idealised LUCJ circuit

The depth and error comparisons above do not say how different a generated circuit is, as a physical object, from the idealised circuit it is meant to replace. We therefore compare each generated circuit U_g (the product of the pool operators selected by the token sequence) with the LUCJ circuit U_{LUCJ} of the same molecule (the spin-balanced LUCJ operator of `ffsim` [34] built from CCSD amplitudes, one ansatz repetition, as in the hardware baseline), in an exact, noiseless calculation with state vectors in the full N-electron space. This is possible for the three tiers of at most 44 qubits whose determinant space fits in memory (SnOH3-me S, SnOH3-me-h2o S and SnOH3-me-2h2o S; 16, 19 and 22 orbitals); all tiers of 46 qubits and above cannot be computed exactly: their full state vectors need 1 to 4 terabytes at 46 qubits, 22 and 155 terabytes at 56 qubits and 3×10^9 gigabytes at 70 qubits (16 bytes per amplitude). SnOH3-me M (38 qubits, 2.5×10^9 determinants, a 41-gigabyte state vector) was not computed. The comparison is therefore limited to the three tiers that were computed, and

Table 8: The two training configurations on the 43 test cases (molecule–lineage pairs) for which both have a best and a second circuit (of 44; the Hartree–Fock-penalty run of SnOH3-me LL with A5-46q has none, see Appendix N). Medians over test cases, by lineage (last block: all cells). Deviation columns (mHa, $\Delta E = E_{\mathrm{SQD}} - E_{\mathrm{DMRG}}$): "best ΔE" and "2nd ΔE" are the deviations of the recovered energy of the best and second circuit, and "min(best,2nd)" is the smaller of the two. "depth" is the as-submitted depth of the best circuit. The remaining columns give the shot composition of the best circuit, in percent: "valid" is the share of all shots with the correct electron numbers in each spin sector; the next columns are shares of these valid shots. "HF" is the share equal to the Hartree–Fock reference determinant, "HL" the share that contains a HOMO→LUMO excitation (the highest occupied to the lowest unoccupied orbital), and "α only", "β only" and "both" the shares that excite only the spin-up channel, only the spin-down channel, or both.

Lineage	Model set	test cases	best ΔE	2nd ΔE	min(best,2nd)	depth	valid	HF	HL	α only	β only	both
A1-46q	HF penalty	11	2.85	5.69	2.64	47	5.0	53.5	0.8	11.8	18.6	3.8
	tuned reward	11	2.30	22.35	2.30	40	6.0	26.5	0.2	23.8	19.6	6.5
A5-46q	HF penalty	10	3.58	2.70	2.70	44	10.6	54.7	1.6	4.7	22.4	2.9
	tuned reward	10	2.37	3.03	2.31	49	10.8	39.9	0.6	8.4	14.8	7.0
A1-70q	HF penalty	11	2.55	6.61	2.55	60	7.3	49.3	2.5	7.6	26.4	4.4
	tuned reward	11	2.55	22.90	2.55	63	2.9	21.1	0.8	22.8	5.0	14.3
A5-70q	HF penalty	11	2.65	3.05	2.65	108	7.5	38.6	2.3	13.8	16.7	6.5
	tuned reward	11	2.62	4.46	2.62	61	2.5	17.0	0.4	9.6	23.2	22.9
all	HF penalty	43	2.85	5.00	2.64	55	8.4	49.3	1.8	10.3	21.7	4.4
	tuned reward	43	2.40	7.33	2.39	56	4.3	25.0	0.7	14.4	19.0	14.3

the larger tiers are compared with LUCJ through circuit cost and recovered energy (Sections 3.4 and 3.7). All circuits of both training configurations on these tiers are included: 12 runs of ten circuits each for each of the two configurations (Table 9); the tuned-reward run at SnOH3-me-2h2o S with A5-70q is the repeated draw of Appendix N.

Measures. With $|\psi_g\rangle = U_g|\mathrm{HF}\rangle$ and $|\phi\rangle = U_{\mathrm{LUCJ}}|\mathrm{HF}\rangle$ and the determinant distributions $p = |\psi_g|^2$, $q = |\phi|^2$ we report (i) the state fidelity $F = |\langle\psi_g|\phi\rangle|^2$ [35, 36]; (ii) the Hellinger distance $H = \sqrt{1 - \mathrm{BC}}$ [37] with the Bhattacharyya coefficient $\mathrm{BC} = \sum_x \sqrt{p_x q_x}$ [38], which is 0 for identical and 1 for disjoint distributions; (iii) the Jensen–Shannon divergence in bits [39], bounded by 1; and (iv) $d_{\mathrm{HF}} = \sqrt{2 - 2|\langle\psi_g|\phi\rangle|}$, the distance between the two states up to a global phase, which is a lower bound of the phase-minimised operator distance $\min_\alpha \|U_g - e^{i\alpha}U_{\mathrm{LUCJ}}\|$ because the two operators are compared on one input state.

We also report two operator-level quantities, both without phase minimisation. First, for the best and second circuits of SnOH3-me S we estimate the operator norm $\|U_g - U_{\mathrm{LUCJ}}\|$ (maximum 2) by a matrix-free Lanczos iteration [40] on $(U_g - U_{\mathrm{LUCJ}})^\dagger(U_g - U_{\mathrm{LUCJ}})$ with a stopping criterion of top-Ritz residual below 10^{-4}; 12 of 16 circuits met this criterion, while four reached the 60-step cap without convergence. Second, we estimate the normalised Hilbert–Schmidt distance as $\|U_g - U_{\mathrm{LUCJ}}\|_F/\sqrt{D}$ (with D the Hilbert-space dimension) using a 3-vector Hutchinson trace estimate. This quantity is 0 for identical unitaries and $\sqrt{2}$ for unrelated Haar-random unitaries.

Result. Table 9 gives the medians; the per-circuit values are in Appendix I and in the released reports. The circuits that the ranking selects (best and second best) have high fidelity with the LUCJ state in the Hartree–Fock-penalty set without exception: all 24 selected circuits have $F > 0.95$ (median $F = 0.998$, Hellinger distance 0.035, $d_{\mathrm{HF}} = 0.05$). However, the LUCJ state itself has fidelity 0.9976–0.9997 with the Hartree–Fock determinant, so the unchanged reference would also score comparably high; the selected circuits exceed the Hartree–Fock fidelity by at most 0.0004. For the tuned reward 19 of the 24 selected circuits are near the LUCJ state ($F > 0.5$; 17 have $F > 0.95$): 11 of the 12 best circuits, but only 8 of the 12 second circuits. The other five selected circuits are orthogonal to it (fidelity 0), among them one best circuit. The circuits that are not selected are mostly the opposite: their median fidelity is 0 and their median Hellinger distance 1.0 (essentially orthogonal states, $d_{\mathrm{HF}} \approx \sqrt{2}$), although about half of them (46 of 96 for the Hartree–Fock penalty) are also near the LUCJ state. Divergence is related to depth: 98% of the circuits of the Hartree–Fock penalty that are shallower than 200 are near the LUCJ state, against 36% of the deeper ones (95% against 27% for the tuned reward). Closeness to the LUCJ state does not translate into a lower recovered error for the second circuit: on the three tiers where the divergence is computed, its median error is 0.13 mHa for the Hartree–Fock penalty, where all second circuits are near the LUCJ state, and 0.07 mHa for the tuned reward, where four of twelve are not.

Operator-level distances are not used. The operator-level quantities above are reported without

Table 9: Ideal (noiseless) distance of the generated circuits from the idealised LUCJ circuit (Section 3.6), medians over the circuits of the three small tiers (SnOH3-me S, SnOH3-me-h2o S, SnOH3-me-2h2o S; ten circuits per run; 12 runs for each of the two configurations). F: state fidelity; H: Hellinger distance between the determinant distributions; d_{HF}: distance of the two states up to a global phase (a lower bound of the operator distance); "near": circuits with $F > 0.5$. Selected = best and second circuit of the ranking. The per-circuit values are in Appendix I.

Set	Circuits	n	F	H	d_{HF}	near	depth
HF penalty	selected	24	0.998	0.035	0.049	24/24	57
	not selected	96	0.000	0.999	1.413	46/96	3372
tuned reward	selected	24	0.998	0.035	0.049	19/24	61
	not selected	96	0.000	1.000	1.414	28/96	3830

phase minimisation, whereas d_{HF} minimises over a global phase on the Hartree–Fock input state. Because the global phase of `ffsim`'s LUCJ circuit is arbitrary, the unphased operator-level distances can substantially overstate the discrepancy. For the selected circuits, the unphased state distance $\sqrt{2 - 2\mathrm{Re}\langle\psi_g|\phi\rangle}$ ranges from 0.17 to 0.59, while d_{HF} ranges from 0.05 to 0.22; for one tuned circuit it is 2.00 while d_{HF} is 0.05. We therefore do not use the operator-level values in the analysis; a phase-minimised operator distance was not computed here.

What the ideal state does not show. The reference state is almost the Hartree–Fock determinant (99.8–99.97% of its probability, because the CCSD amplitudes are small), so a high fidelity to it mostly says that the generated state is also close to the Hartree–Fock determinant (99.5%, 100% and 98% for the best A1-46q circuit at the three tiers). The hardware does not sample that state: the measured shots of the same circuits are 7–29% Hartree–Fock, 40–51% single and 17–39% double excitations (Table A21), so device noise, not the ideal circuit, produces the excitation content that the recovery step exploits. The divergence is thus a statement about the ideal circuits, not about what the device returns; the two should be read together. Noiseless sampling with the matrix-product-state (MPS) simulator of cuQuantum [41] is consistent with this reading (Appendix O). The simulation used 8-token checkpoints, so its circuits differ from those run on `ibm_kobe`. Without noise, circuits without a global orbital-rotation operator (`orb_rot` in Table A3) returned the Hartree–Fock determinant in at least 90% of their shots, and 4–7 circuits per test case returned only that determinant in 10^5 shots. Circuits containing a global `orb_rot` operator instead produced many distinct configurations.

Relation to the recovered error. Among the recovered circuits of the three small tiers, the fidelity is only weakly related to the recovered error (Spearman correlation 0.31 for the HF-penalty set), because every error on these tiers is within a fraction of a mHa (median 0.1 mHa both for circuits near the LUCJ state and for those far from it).

As noted in Section 3.3, the recovery cap exceeds the full alpha string dimension at these tiers. In the noiseless MPS sampling of Appendix O, the merged recovered bases contained more than 99.5 % of all alpha strings. The best individual circuits recovered to 0.25–0.51 mHa and had broad alpha-string coverage. The comparison cannot say whether closeness to LUCJ helps at the larger tiers, where it cannot be computed.

3.7 Comparison against DMRG

Table 7 gives, for each SnOH3-me-family target, the lowest recovered energy over every lineage, policy and result type in the fresh-shot campaign, with the DMRG reference (Table A1) and the signed deviation ΔE; the lineage and result type that achieved it are noted. This is a best-of-k statistic over up to eight lineage/policy combinations and five result types, so it is more favourable than any single pre-declared run; per-lineage values are in Tables A15 and A19.

The three small active spaces (SnOH3-me S (32q), SnOH3-me-h2o S (38q), SnOH3-me-2h2o S (44q)) are recovered to within 0.1 mHa of DMRG by the best result of each (0.01, 0.01 and 0.04 mHa). On the larger tiers the lowest errors are 0.71 mHa (SnOH3-me-h2o M (46q)), 2.1 mHa (SnOH3-me-h2o LL (66q)), 1.9 mHa (SnOH3-me-h2o L (56q)), 3.8 mHa (SnOH3-me M (38q)), 4.3 mHa (SnOH3-me-2h2o M (58q)), 5.6 mHa (SnOH3-me L (46q)), 7.2 mHa (SnOH3-me-2h2o L (70q)) and 10.4 mHa (SnOH3-me LL (56q)). The SnOH3-me (non-hydrate) series is the hardest of the family for the policy: its 38-, 46- and 56-qubit tiers have the largest errors, and the SnOH3-me M (38q) tier, which has fewer qubits than the 46-qubit hydrate, is 3.8–17.2 mHa from DMRG where the hydrate is within 1.5 mHa.

The choice of circuit matters more than shot noise. Selecting the two circuits with the lowest one-shot ranking energy at net CI dimension 40,000 gives errors of 0.71–1.5 mHa for Sn(OH)3-h2o M for the best circuit of each trained lineage, whereas the E+D circuits (Section 2.1), recovered on the same fresh shots, remain at 22.6–23.0 mHa in two of the four lineages (A1-46q and A1-70q), as do all untrained controls (22.6–22.9 mHa). Each circuit was sampled once, so this campaign cannot separate the choice of circuit from shot noise. The best circuits have 5.3×10^3–2.4×10^4 valid shots of 100,000 at these test cases; the E+D picks that stay near 23 mHa have 13–2457 valid shots, but an E+D pick of A5-46q with 1992 valid shots reaches 1.4 mHa, so the valid-shot count alone does not separate them. We did not run the ablation that would separate valid-shot count from ranking energy or depth, so we do not attribute the im-

provement to either.

3.8 What distinguishes a good recovery: valid shots and selection

The stored per-circuit data (released with the paper, see Appendix L) allow a few observational checks. They are not controlled experiments: valid-shot count, circuit depth and connectivity are correlated with one another, and the runs are few. We use the 94 distinct individual-circuit recoveries of the HF-penalty trained runs at 44 qubits and above (best, second and the two E+D picks, duplicates removed).

Valid shots. The recovered error falls steeply with the number of valid shots of the recovered circuit: median 41.1 mHa for circuits with fewer than 100 valid shots (25 recoveries; 4 below 25 mHa), 41.1 mHa for 100–1000 (11; 3 below 25 mHa), 5.3 mHa for 1000–10,000 (42; 39 below 25 mHa) and 1.8 mHa above 10,000 (16; all below 25 mHa). Within a single molecule the Spearman correlation between $\log_{10}$ valid shots and error is negative at 7 of the 8 molecules with enough recoveries (-0.65 to -0.89); the one positive value ($+0.57$) occurs at SnOH3-me-2h2o S, where every recovery is within a fraction of a mHa.

Ranking energy does not order recoveries. Of the 43 HF-penalty runs in which both the best and the second best circuit were recovered, the best circuit (lowest ranking energy) recovered to an energy no higher than the second in 28, i.e. 65%. The ranking energy is a useful screen (it discards circuits with too few valid shots) but its order among the leading circuits is only weakly informative about the recovered order.

3.9 Computational cost: training and inference

This section reports the quantum and classical resources of the training and inference runs. Wall-clock figures are Slurm elapsed times, shot counts are circuits times shots per circuit, and QPU time is IBM's billed quantum time per job. The same numbers, with the per-run breakdown, are in Appendix L and in the released HTML. The `ibm_kobe` calibration in force during the experiments (T_1, T_2, readout, single- and two-qubit gate errors) is given in Section L.1, Tables A24 and A25.

3.9.1 Training

Stage 1 runs on an exact-statevector simulator and uses no QPU time. Stage 2 samples on `ibm_kobe` at 10,000 shots per circuit, 20 circuits per iteration batched in one Qiskit Runtime job, so each iteration is 200,000 shots and one job. For the Hartree–Fock-penalty configuration the four lineages ran 150 iterations of Stage 1 and 150 of Stage 2 each, i.e. 600 Stage-2 iterations, 12,000 circuits and 120,000,000 shots; Stage 1 took 106.3 h and Stage 2 112.8 h of wall time in total (the lineages ran concurrently). IBM billed 648 QPU minutes (10.8 h) for the 600 jobs, 65 s per job on average: 59.8 and 59.4 s per job for the 46-qubit lineages and 69.9 and 70.0 s for the 70-qubit lineages (Table A8), so the 70-qubit iterations cost 1.2 times as much QPU time as the 46-qubit ones. The tuned reward costs 1175 QPU minutes (120 s per job) because its 70-qubit circuits are deeper. Table 10 gives the quantum resources and Table A26 (Appendix M) the budget per lineage. Iterations lost to crashed Runtime jobs are not counted, so the shot totals are lower bounds.

The mean iteration time by pipeline step is in Figure 6 (Table A7, Appendix E). In Stage 1 the circuit sampling (exact simulation) accounts for the largest share (57–71% of an iteration); in Stage 2 of the HF-penalty recipe the sampling share is 47–61% and includes the IBM device queue, with the SQD diagonalization/recovery taking most of the remainder. The refinement, transpilation and backward steps together are under 5% of an iteration in every lineage.

3.9.2 Inference

Inference runs once on hardware per run (a batch of 10 circuits at 100,000 shots each, one Runtime job) and then entirely on classical hardware. Each of the two training configurations has 44 runs (440 circuits, 4.4×10^7 shots, 44 Runtime jobs); the billed QPU time and the classical wall time are in Table 11. The LUCJ baseline used 18 Runtime jobs (1,800,000 shots, 11 QPU minutes; one per molecule at each of the two transpilation levels) and the eight untrained controls 8 jobs (Table 10).

Table 11 gives the inference timing and the billed QPU time of the two configurations side by side, stage by stage.

In both configurations the classical recovery dominates the wall time, far more than the hardware round trip (mostly IBM queue time and local transpilation) or the ranking (Table 11); merged recoveries are the most expensive per job because they pool up to ten times the shots. Failed, cancelled and timed-out jobs are not counted in the table, and the per-run times are in the released reports.

Table 10: Quantum resources on `ibm_kobe`: circuits, shots, Qiskit Runtime jobs and IBM-billed QPU time (minutes) from the account's own job records. A training job is one iteration (20 circuits at 10,000 shots); an inference job is one run (10 circuits at 100,000 shots). Stage 1 of the training is a simulation and uses no QPU time. The inference rows are the sampling jobs of the 44 runs of each configuration; the untrained controls are the 8 control runs, and the LUCJ baseline one run per molecule at each of two transpilation levels. Training without a penalty is the A1-46q run and the stopped runs of the other lineages.

Phase	Circuits	Shots	Jobs	QPU time (min)	mean QPU s/job
HF penalty, training Stage 2, 46-qubit models	6,000	60,000,000	300	298	60
HF penalty, training Stage 2, 70-qubit models	6,000	60,000,000	300	350	70
HF penalty, training total	12,000	120,000,000	600	648	65
HF penalty, inference (44 runs)	440	44,000,000	44	315	429
Tuned reward, training Stage 2	11,740	117,400,000	587	1175	120
Tuned reward, inference (44 runs)	440	44,000,000	44	366	500
No penalty, training Stage 2	3,700	37,000,000	185	205	67
Untrained controls, inference (8 runs)	80	8,000,000	8	85	639
LUCJ baseline, inference	18	1,800,000	18	11	38

Table 11: Inference timing and cost of the two training configurations (one zero-shot run = ten circuits sampled once on `ibm_kobe` at 100,000 shots, ranked, then up to five recoveries). Wall times are Slurm elapsed times of the finished jobs on the roquo cluster; "jobs" counts finished jobs, failed, cancelled and timed-out jobs are listed separately and not counted. Recovery jobs use two GPUs each. QPU time is the time IBM billed for the sampling jobs (job records); the untrained controls are not included. Recoveries still running when the table was made are listed in the row "Still running" and are not counted.

Set	Stage	jobs	wall (h)	mean per job (min)	GPU-hours
HF penalty	Sampling and transpilation (IBM round trip)	44	128.1	175	128.1
	Ranking diagonalisation	44	1.0	1	2.0
	Recovery: best / second	86	196.5	137	392.9
	Recovery: E+D 1 / 2	48	306.9	384	613.8
	Recovery: merged	43	191.5	267	383.0
	All recoveries	177	694.9	236	1389.8
	QPU time billed by IBM (sampling)	44/44 jobs: 18,877 s (315 min), 429 s per job			
	Not counted: failed, cancelled or timed out	1	0.00		
tuned reward	Sampling and transpilation (IBM round trip)	44	191.5	261	191.5
	Ranking diagonalisation	44	1.2	2	2.4
	Recovery: best / second	86	254.1	177	508.3
	Recovery: E+D 1 / 2	50	332.9	399	665.8
	Recovery: merged	43	228.2	318	456.3
	All recoveries	179	815.2	273	1630.4
	QPU time billed by IBM (sampling)	44/44 jobs: 21,980 s (366 min), 500 s per job			

4 Discussion

The two rewards give results that are equivalent in accuracy at the resolution of this study (Section 3.5): on the 43 test cases that both have, the median error of the best circuit is 2.8 mHa for the Hartree–Fock penalty and 2.4 mHa for the tuned reward, with one training seed per lineage. We centre the paper on the Hartree–Fock penalty because it is the simpler reward, with one added term, and because it has no best-circuit error above 30 mHa, against three for the tuned reward. The more diverse policy of the tuned reward (an effective vocabulary of 21–27 tokens against one; Appendix F) does not translate into more accurate circuits on unseen molecules in this setting. At inference, the penalty does not remove the Hartree–Fock weight of the shots as it does in training; it changes the balance between the spin channels (Section 3.5). We have not tested whether these differences cause the differences in accuracy. What both configurations share, and what the comparison with the untrained controls supports, is that a trained policy recovers well below an untrained one at the tiers where untrained policies fail (a median of 106 mHa for the eight untrained controls against 2.8 mHa for the 43 runs of the Hartree–Fock penalty).

The HF-penalty policy has collapsed onto about one

operator sequence by iteration 50–55 of Stage 2, and its 46-qubit circuits excite one spin channel only (Section 3.2). Both statements describe the policy as trained, and the zero-shot circuits are drawn from it with the molecule-conditioned input, so they are not a fixed circuit; but ten samples are far from ten independent circuits (Section 3.3 shows that some runs' ten circuits are nearly identical). The practical consequence is that a large share of the accuracy comes from the ranking and recovery steps that follow the sampling, and from a modest number of well-yielding circuits, not from a rich learnt distribution. The tuned reward gives a slightly lower median error (2.4 against 2.8 mHa on the 43 paired test cases) but a worse tail, at the price of deeper 70-qubit circuits and 1.8 times the QPU time per job. Whether a more diverse policy would help at larger scale, or with more circuits per run, is open question.

The dominant lever found during hyperparameter tuning was the circuit-depth admission ceiling, not the magnitude of the coverage/overshoot reward terms: once the ceiling is set generously relative to the transpiled depths produced by a given circuit length L and qubit count, a wide range of reward-weight combinations learn comparably, whereas an overly tight ceiling produces sustained zero-gradient stalls regardless of reward shaping. The admission gate's role as a hard constraint on which samples ever reach the policy-gradient update is a more consequential design choice, for this architecture and reward formulation, than the relative weighting of the soft reward terms. The pass rate of the Stage-2 gate differs by qubit count and not by conditioning scheme (100% and 99% at 70 qubits against 46% and 47% at 46 qubits, Section 3.2); because A1 and A5 agree, we attribute it to the energy-admission tolerance of the 46-qubit lineage and not to either architecture.

A related, more acute version of the same structural issue appeared at the 70-qubit hardware tier in the preliminary runs: the energy admission tolerance that worked at 46 qubits produced zero admitted iterations over many consecutive hardware iterations, a genuine deadlock rather than slow convergence, since without an admitted iteration no gradient exists. Loosening the tolerance by a fixed factor broke the deadlock, and the 70-qubit lineages of the final runs, which start from that loosened tolerance, admit 99–100% of their iterations (Table A5). The 46-qubit lineages never needed the loosening, yet it is they that show the weaker admission outcome: both settle near a 46–47% pass rate (Figure A11), two independent conditioning schemes giving the same outcome, which points at the 46-qubit lineage's energy tolerance itself being comparatively tight for its own curriculum, not at either architecture. It is the mirror image of the 70-qubit finding: there the tolerance blocked learning until loosened; here learning proceeds, but roughly half of all hardware circuits fail the gate throughout training. We did not test a different tolerance for the 46-qubit lineages.

The depth and two-qubit-gate curves cannot be read as evidence of transfer. The best circuit of Stage 2 is shallow from the first iterations (Section 3.2), whereas Stage 1 ends with depths of several thousand (Section 3.1), but the two stages differ in the molecules (the Stage-1 curriculum against the 23- and 35-orbital targets), the sampler, the best-of-20 selection by energy and the depth ceiling, so the change cannot be attributed to the carried policy. We did not establish that the Stage-1 warm start improves Stage 2 over a random initialization under the same budget. A short test on a small molecule (ethylene) on a noisy simulator showed no clear difference between warm-started and randomly initialized policies; however, this test was not designed as a controlled comparison of the hardware stage, so no quantitative conclusion about warm-start benefits can be drawn from it.

The divergence measures (Section 3.6) are computed for the ideal, noiseless circuit on the Hartree–Fock input, and only for the three tiers of at most 44 qubits. They support three statements. First, the circuits that the ranking selects have a fidelity above 0.95 with the LUCJ state, but no higher than that of the Hartree–Fock determinant itself (0.9976–0.9997), so this fidelity does not show that the circuits reproduce the correlated part of LUCJ. Second, the deep circuits that the policy also samples are orthogonal to it (fidelity 0, Hellinger distance 1), which matches their absence of valid shots on the device; the same circuits carry the diagonal-Coulomb layers that dominate their depth. Third, closeness is not implied by shallowness: in the SnOH3-me S run of the A1-46q policy, a circuit of depth 46 has fidelity 0.954 and one of depth 196 has fidelity 0, whereas the circuits of depth 30–57 reach 0.993–0.998. Operator-level distances are reported without phase minimisation and are not used in the analysis (Section 3.6). The accuracy of the method is therefore established by the recovered energies against DMRG, and not by the closeness to LUCJ.

We looked for two trends and found neither clearly. (i) Does a policy predict better on active spaces close to the size it was trained for? For the HF-penalty set, the median error of the selected circuit at the spaces of at most 46 qubits is 0.85 and 1.07 mHa for the A1-46q and A5-46q policies, and 0.46 and 0.93 mHa for the A1-70q and A5-70q policies; at 56–70 qubits it is 6.5 and 4.1 mHa (46-qubit policies) against 4.3 and 4.9 mHa (70-qubit policies). The 70-qubit policy is better in 5 of 6 test cases (A1) and 5 of 6 (A5) at the small spaces and in 4 of 5 (A1) and 2 of 4 (A5) at the large ones, so the larger policy is slightly better throughout, including on the spaces the smaller policy was trained for, and the gap does not grow with the size of the target. The medians over all eleven spaces (2.6–3.6 mHa) do not separate the four lineages with one seed each, so we find no specialisation of a policy to the size it was trained for; the largest errors go with few usable circuits in the draw (Appendix N), an observation on one draw per

test case and not a test. (ii) Does the heavier A5 encoder generate deeper circuits? No. Over all 220 sampled circuits per encoder, the A1 circuits are deeper (median as-submitted depth 3,427 against 528 for A5; mean 5,693 against 3,550). This is carried by the A1-70q policy, whose circuits are the deepest of the four lineages (median 5,581, 65% deeper than 1,000), whereas A5-70q has the shallowest and the highest valid-shot yield. The extra structural information of the A5 encoder is therefore not paid for in circuit depth, and it does not give a consistently lower error either (median over all spaces 3.6 and 2.7 mHa for A5, 2.9 and 2.6 mHa for A1). All comparisons rest on one training seed per lineage.

The depth advantage of Section 3.4 is paid for with shots. A fresh-shot run samples ten circuits at 100,000 shots each (10^6 shots, one Runtime job) to obtain one or two circuits for recovery, and ranks all ten by a diagonalization; eight or nine of the ten are discarded, and in most runs the typical sampled circuit is deeper than LUCJ. A conventional LUCJ-based SQD run uses one fixed circuit and 100,000 shots, needs no training or selection step, uses 11 to 13 times less QPU time (38 s per job against 429 and 500 s per job for the Hartree–Fock-penalty and tuned-reward runs, Table 11), and its depth can be tuned directly through the number of ansatz repetitions and the transpiler optimization level. Whether the policy is worth its additional shots and its one-time training (Table 10) depends on the size of the active space. At the three smallest spaces (32, 38 and 44 qubits) the recovery subspace cap alone reached the DMRG answer from the LUCJ samples, as it did from every GQE circuit, so there the additional shots and the policy buy nothing. From 46 qubits up, LUCJ-SQD's raw errors are 22.7–204.1 mHa and its singles-injected errors 5.2–41.2 mHa (5.2–12.6 mHa except at SnOH3-me LL). The policy's accuracy gain over LUCJ is established only at SnOH3-me-h2o M, with one fresh sample per circuit; at the other tiers its errors (Table 7) are comparable to or only modestly better than LUCJ's singles-injected errors, and the benefit there is the shallower circuit (Section 3.4).

5 Conclusion

We trained amortized GQE policies in two stages, a noiseless-simulator warm-up and 150 iterations with the IBM Heron processor `ibm_kobe` in the training loop, at 46 and 70 qubits, with a Hartree–Fock penalty in the hardware reward. Applied without retraining to the eleven active spaces of the $Sn(OH)_3$-Me family, 32 to 70 qubits, the policies gave, after one round of hardware sampling, followed by classical ranking and recovery of the cached counts with the molecular Hamiltonian and no further QPU calls, a median error of 2.8 mHa against DMRG over the 43 trained runs (13 within 1 mHa, 42 within 15 mHa). At 46 to 70 qubits, 27 of 28 runs are within 15 mHa (median 4.5 mHa); the errors are 0.06–0.65 mHa at the 32–44-qubit spaces (3.8–17 mHa at SnOH3-me M, 38 qubits) and 0.7–1.5 mHa at the 46-qubit hydrate. The trained policy is more accurate than its untrained control in all eight runs that have one (median error of the controls 106 mHa). At 56–70 qubits the selected circuits used 34–72× fewer two-qubit gates and were 4–12× shallower than the LUCJ circuit transpiled for the same device, at the price of about 11 times the QPU time per job, because ten circuits are sampled instead of one. These results show that training a circuit generator with the quantum hardware in the loop is feasible at 70 qubits. The tuned reward gives a slightly lower median error, 2.4 against 2.8 mHa on the 43 test cases that both have (lower in 26), but a worse tail: three of its best circuits exceed 30 mHa, none does for the Hartree–Fock penalty. We therefore keep the simpler Hartree–Fock penalty as the main configuration, for its robustness in these single-seed runs.

The evidence has clear limits. The cost and accuracy advantages are those of the selected circuit, not of the typical sampled one; the ranking energy screens out circuits with too few valid shots but does not order the leading circuits; the trained policy collapses onto essentially one operator sequence in training, and its 46-qubit circuits excite one spin channel only; each circuit was sampled once and each configuration trained with one seed; training uses an approximate classical reference energy (CCSD(T) for the Stage-2 molecules) in the admission gate and in two reward terms, so the protocol applies to systems for which such a reference exists; the effect of model size was not studied; the evidence that the policy transfers to new chemical systems remain qualitative; the untrained controls cover three molecules; and no ablation separates valid-shot count, depth and connectivity as the cause of the gains, or tests whether idle qubits matter (Appendix J). The four main training runs required 1.2×10^8 hardware shots and 648 QPU minutes. The code and the per-run results are released (Appendix P).

Several follow-up studies are suggested by these limits. Controlled hardware comparisons of warm-started and randomly initialized policies, multiple training seeds per lineage, and smaller or parameter-efficient transformers would help quantify how sensitive the results are to the training recipe and model size. Whether a more diverse policy improves performance at larger scale, or with more circuits per run, remains open. Because several runs rely on modest valid-shot yields, a natural next step is to combine the approach with error mitigation (not used here) and to order training molecules by a measure of how difficult their ground states are to sample. Finally, the cost accounting reported here can guide the choice of the number of circuits and shots when extending the approach to larger active spaces.

Acknowledgements

We thank the RIKEN JHPC-quantum operations team for providing access to the Fugaku supercomputers' IBM Quantum allocation (`ibm_kobe` based on the IBM Heron processor) and for promptly flagging and help- ing resolve a technical issue during the hardware training campaign. Part of this work was carried out on the Miyabi supercomputer system (Miyabi-G) of the Joint Center for Advanced High Performance Computing (JCAHPC), through its "Large-scale HPC Challenge", a call-for-proposals program that gives exclusive use of the accelerated GPU nodes of Miyabi-G, the general-purpose CPU nodes of Miyabi-C, or both. All main-text training, inference and LUCJ results were obtained on `ibm_kobe`. Additional preliminary experiments that are not part of the reported analyses were performed on `ibm_kawasaki` managed by the University of Tokyo. A part of this work was performed for Council for Science, Technology and Innovation (CSTI), Cross-ministerial Strategic Innovation Promotion Program (SIP), "Promoting the application of advanced quantum technology platforms to social issues" (Funding agency: QST).

We used Claude Sonnet 5.5 (Anthropic) to assist with drafting and editing the manuscript and the analysis code. The authors reviewed the output and are responsible for the content of this paper.

The sampled circuits, raw counts and result files behind every table are described in Appendix P.

A Classical reference benchmarks

Table A1: DMRG reference energies, SnOH3-me family, all active-space tiers (maximum bond dimension 4000, 17 sweeps, energy-convergence threshold $10^{-5}\,E_\mathrm{h}$). The truncation error is the discarded weight of the last sweep; it is not an uncertainty estimate. No sweep-by-sweep energies or lower-bond-dimension calculations are reported, so differences below about 10^{-2} mHa should not be over-interpreted.

Molecule	Tier (active space)	DMRG energy (E_h)	Truncation error
SnOH3-me	S (16o, 24e)	−479.193382	2.87×10^{-12}
SnOH3-me	M (19o, 24e)	−479.302278	4.73×10^{-7}
SnOH3-me	L (23o, 32e)	−479.314216	1.45×10^{-6}
SnOH3-me	LL (28o, 42e)	−479.315057	2.04×10^{-7}
SnOH3-me-h2o	S (19o, 30e)	−555.108833	2.53×10^{-12}
SnOH3-me-h2o	M (23o, 30e)	−555.125158	2.33×10^{-9}
SnOH3-me-h2o	L (28o, 40e)	−555.135058	4.70×10^{-9}
SnOH3-me-h2o	LL (33o, 50e)	−555.137396	1.64×10^{-8}
SnOH3-me-2h2o	S (22o, 36e)	−631.083312	3.04×10^{-12}
SnOH3-me-2h2o	M (29o, 36e)	−631.118616	2.80×10^{-7}
SnOH3-me-2h2o	L (35o, 48e)	−631.132238	5.76×10^{-7}

B MP2 natural orbitals and active spaces of the target molecules

This appendix documents the natural orbitals behind the active spaces of the held-out molecules (Section 2.6): the definition of the tiers (Table A2) and the orbitals of all three molecules (Figures A1–A3), followed by comments on the orbital shapes. The occupation numbers are those of the Gaussian 16 MP2 densities (for $Sn(OH)_3$-Me as stored with the natural orbitals, for the hydrates from the MP2 density and the def2-SVP overlap matrix) [30–32]. The comments are our reading of the plotted isosurfaces and occupation numbers; they are qualitative and have not been checked by a population analysis.

Table A2: Active spaces built on the MP2 natural orbitals (NOs) of the three target molecules (MP2/def2-SVP, frozen core, Gaussian 16). A tier is a contiguous range of NOs ordered by occupation number, from the lowest NO of the range through the HOMO and the first virtual NOs; orbitals and electrons give the active space, "occupied" and "virtual" the NOs of the range below and above the HOMO. Alpha strings is the number of alpha-spin occupation strings, $\binom{n_\mathrm{orb}}{n_\mathrm{elec}/2}$, which equals the number of beta strings for these closed-shell active spaces. The last column says whether the tier was run in this paper (all eleven tiers of 32 to 70 qubits were run for the four lineages of each configuration).

Molecule	Tier	(orbitals, electrons)	NO range	Occupied	Virtual	Qubits	Alpha strings	HOMO/LUMO NO	Run here
$Sn(OH)_3$-Me	LL	28, 42	9–36	21	7	56	1,184,040	29 / 30	yes
$Sn(OH)_3$-Me	L	23, 32	14–36	16	7	46	245,157	29 / 30	yes
$Sn(OH)_3$-Me	M	19, 24	18–36	12	7	38	50,388	29 / 30	yes
$Sn(OH)_3$-Me	S	16, 24	18–33	12	4	32	1,820	29 / 30	yes
$Sn(OH)_3$-Me·H_2O	LL	33, 50	10–42	25	8	66	13,884,156	34 / 35	yes
$Sn(OH)_3$-Me·H_2O	L	28, 40	15–42	20	8	56	3,108,105	34 / 35	yes
$Sn(OH)_3$-Me·H_2O	M	23, 30	20–42	15	8	46	490,314	34 / 35	yes
$Sn(OH)_3$-Me·H_2O	S	19, 30	20–38	15	4	38	3,876	34 / 35	yes
$Sn(OH)_3$-Me·$2H_2O$	L	35, 48	16–50	24	11	70	417,225,900	39 / 40	yes
$Sn(OH)_3$-Me·$2H_2O$	M	29, 36	22–50	18	11	58	34,597,290	39 / 40	yes
$Sn(OH)_3$-Me·$2H_2O$	S	22, 36	22–43	18	4	44	7,315	39 / 40	yes

Occupation numbers. In all three molecules the natural-orbital occupations stay close to 2 or 0 (Figure 4): no orbital has an occupation between about 0.1 and 1.9. The largest deviations are at the frontier: the HOMO has an occupation of 1.956–1.957 and the LUMO of 0.048–0.049, in each molecule. The occupied orbitals fall off smoothly from 2.00 to the HOMO, so there is no gap in the occupied part that would define a natural boundary for an active space; the virtual orbitals drop from the LUMO (0.048–0.049) to 0.031–0.033 and decay slowly. The orbitals nearest the frontier, which are the most fractionally occupied, are the first that every tier includes. This is a limitation, not

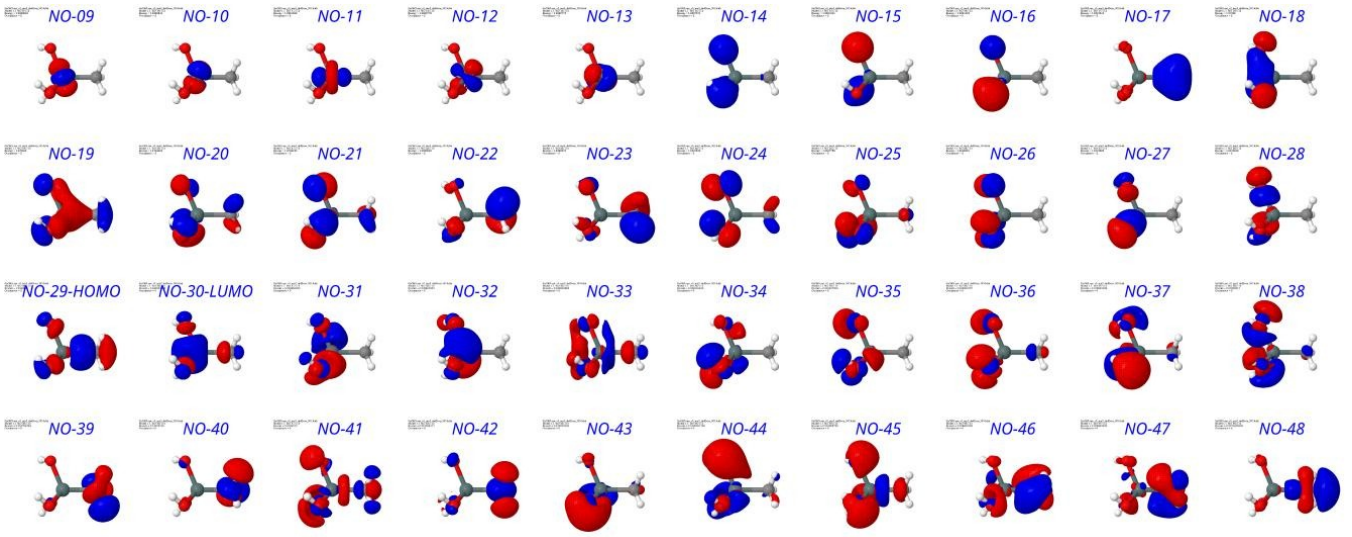

Figure A1: MP2 natural orbitals NO-9 to NO-48 of $Sn(OH)_3$-Me (MP2/def2-SVP, frozen core, Gaussian 16; isosurfaces, the two colours are the two signs of the orbital). The orbitals are ordered by occupation number; the HOMO and LUMO are labelled. The active spaces of Table A2 are ranges of these orbitals (Figure 4). The first eight natural orbitals are not shown.

a finding: MP2 occupations stay close to 0 and 2 for weakly correlated systems and do not show whether strong static correlation exists. The comparison with DMRG validates the SQD energies within each selected active space, not the physical completeness of that space.

Orbitals near the frontier. In $Sn(OH)_3$-Me the HOMO (NO-29) is concentrated along the Sn–CH_3 axis, with lobes of opposite sign on the tin and on the methyl side and small amplitude on the oxygens, which looks like a Sn–C σ-type orbital. The LUMO (NO-30) has a large lobe on the tin and hydroxyl side and a smaller lobe on the methyl carbon. In the hydrates the frontier orbitals are less symmetric because the water ligands break the three-fold arrangement of the hydroxyl groups. The HOMO of the monohydrate (NO-34) has its main lobe above tin towards the methyl group and small lobes on the hydroxyl oxygens, and its LUMO (NO-35) consists of small lobes spread over the hydroxyl region. In the dihydrate the HOMO (NO-39) has lobes on the methyl side and on one oxygen and the LUMO (NO-40) lobes on tin and neighbouring oxygens. If, as these shapes suggest, the Sn–C bond region contributes to the frontier orbitals, the active spaces are consistent with the Sn–C bond cleavage that matters for organotin photoresists, but this is a reading of the pictures, not a result.

Orbitals below the frontier and the Sn 4d orbitals. The occupied natural orbitals at the start of the LL range (NO-9 to about NO-15 in $Sn(OH)_3$-Me, with the corresponding ranges in the hydrates) have compact, multi-lobed shapes around tin with small contributions from the oxygens, which is the mixing of Sn 4d with O 2p noted in Figure 4; their occupations are 1.998 or higher. These orbitals are only in the LL tier. The orbitals between them and the HOMO are mostly lone-pair-like on single hydroxyl or water oxygens, or localized along Sn–O bonds, and the hydrates add orbitals that extend over the water ligands.

Higher virtual orbitals. The virtual orbitals above the LUMO have more nodes, and several of them are large diffuse lobes that extend outside the molecular frame. Such orbitals are probably correlating or diffuse functions rather than bonding orbitals. NO-44 and NO-45 of the monohydrate lie above its active spaces (they end at NO-42), so they are in no tier. NO-44 of the dihydrate is inside its M and L tiers (ending at NO-50) and outside S (ending at NO-43). Whether adding a few higher natural orbitals would improve the active spaces, as the selection of these spaces suggested, and whether the Sn 4d orbitals are needed, were not tested here.

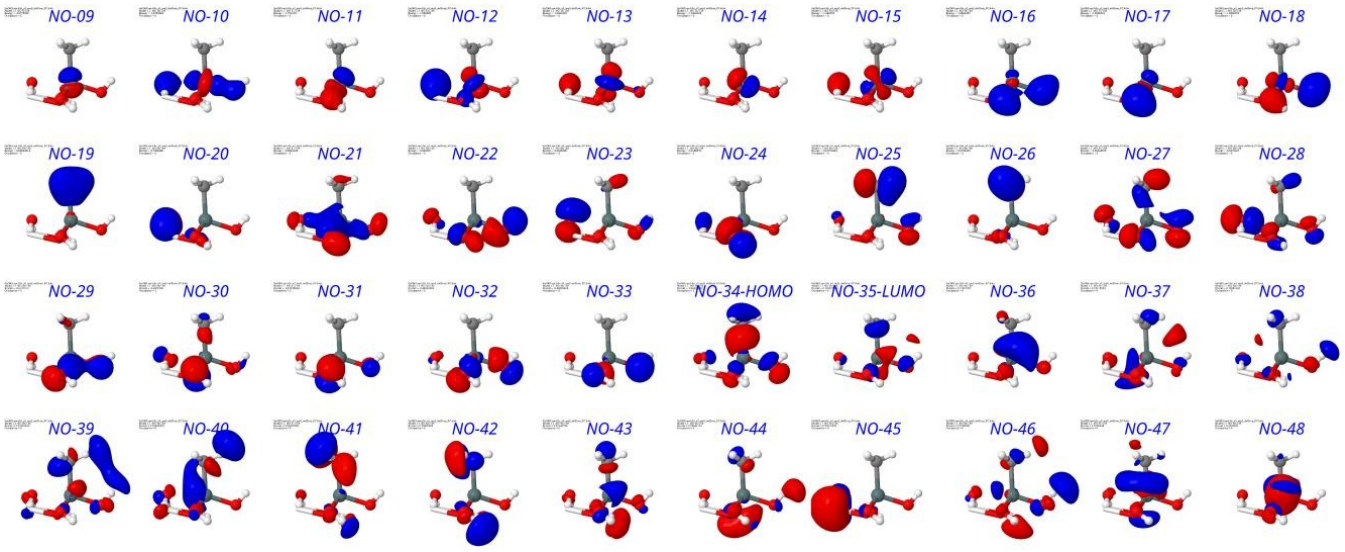

Figure A2: MP2 natural orbitals NO-9 to NO-48 of $Sn(OH)_3$-Me·H_2O (MP2/def2-SVP, frozen core, Gaussian 16; isosurfaces, the two colours are the two signs of the orbital). The orbitals are ordered by occupation number; the HOMO and LUMO are labelled. The active spaces of Table A2 are ranges of these orbitals (Figure 4). The first eight natural orbitals are not shown.

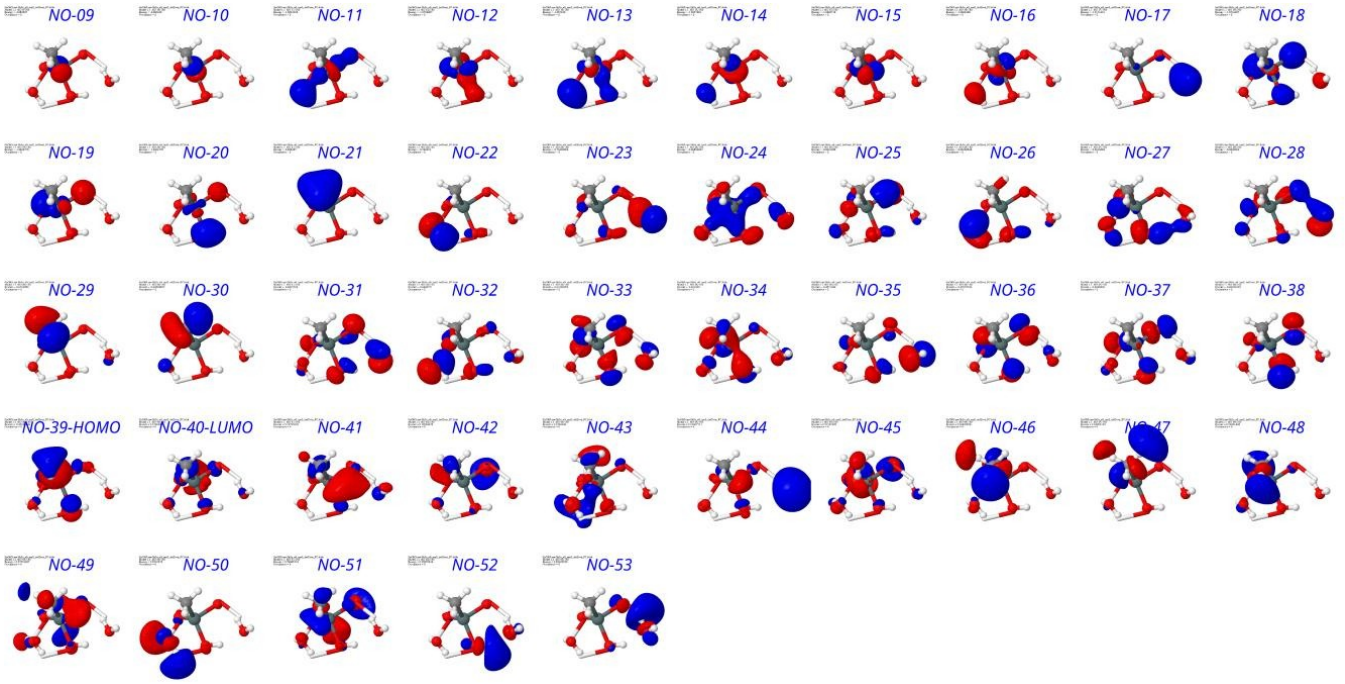

Figure A3: MP2 natural orbitals NO-9 to NO-53 of $Sn(OH)_3$-Me·$2H_2O$ (MP2/def2-SVP, frozen core, Gaussian 16; isosurfaces, the two colours are the two signs of the orbital). The orbitals are ordered by occupation number; the HOMO and LUMO are labelled. The active spaces of Table A2 are ranges of these orbitals (Figure 4). The first eight natural orbitals are not shown.

C Glossary of code-level hyperparameters and labels

The main text uses descriptive names for every setting; this appendix gives, for each, the flag or identifier it carries in the released code, its exact meaning, and the value used. Units: energies in mHa (10^{-3} Hartree) unless stated; "iteration" is one policy update from a batch of sampled circuits; "admitted" means that the iteration passed the admission gate (the best circuit of the batch of 20) and so contributed a gradient step.

Table A3: Glossary. "Name in the text" is the descriptive name used in the main text; "Code label" is the flag or identifier in the released code.

Name in the text	Code label	Definition
Policy and circuit		
Number of gate tokens per circuit, L	`ngates`	Number of operator tokens the policy generates for one circuit (28 for the 46-qubit lineages, 44 for the 70-qubit lineages). It fixes the size of the policy's position-embedding table and therefore cannot change between training stages.
Operator pool; vocabulary	`FfsimLUCJPool`	The set of unitary operators that the tokens select. For a molecule with n_{orb} spatial orbitals it has $3\,n_{\mathrm{orb}} + 5$ operators; token ids above that are identity padding so that one output layer serves every molecule.
Local excitation gates	`givens`, `num_num`	Givens rotation between two neighbouring spatial orbitals (single-excitation character) and number–number (density–density) phase gate between two orbitals.
Global operators	`orb_rot`, `diag_coulomb`, `LUCJ`	Operators acting on every qubit: four fixed (seeded) orbital rotations, a diagonal-Coulomb layer with two time values, and one LUCJ layer (Table 1).
Molecule conditioning input	`set_prompt`	Computes the molecule descriptor (and, for the A5 encoder, the molecular graph embedding) once and stores it, so that the decoder attends to it at every step.
Optimiser and curriculum		
Learning rate	`base-lr`	Step size of the AdamW optimiser, 5×10^{-6} for all lineages.
Prompt-encoder learning-rate multiplier	`encoder-lr-mult`	The prompt encoder's learning rate is the learning rate times this factor; 1.0 here.
Gradient-norm clipping threshold	`grad-clip`	Maximum L_2 norm of the combined policy and encoder gradient; larger gradients are rescaled to this norm before the update. 8.0 here.
Warm-up fraction	`warmup-frac`	Fraction of the iterations in the first phase of the three-phase size curriculum, in which the training molecules are visited in plain round-robin order (0.25).
Focus fraction	`focus-frac`	Fraction of the iterations in the second phase, which visits only the largest molecules, without noise (0.30 here). The remaining 0.45 form the third phase, with the same restricted set and Gaussian noise (standard deviation 0.01 Hartree) added to the reward.
Large-molecule fraction	`curriculum-top-frac`	Fraction of the training molecules, ordered by qubit count, that count as the "largest" set in the second and third phases (0.4 here; rounded up to at least one molecule).
Admission gate and reward		
Depth ceiling	`depth-ceiling-max`	Largest transpiled circuit depth that is still admitted: a batch whose best circuit is deeper triggers no gradient update. 10,000 (46 qubits) and 20,000 (70 qubits) in Stage 1; 10,000 (46 qubits) and 50,000 (70 qubits) in Stage 2.

Name in the text	Code label	Definition
Energy admission tolerance (gradient-admission threshold)	`gate-energy-mha`	A batch is admitted only if the magnitude $\|\Delta E\|$ of the raw deviation of its best circuit from the reference energy is at most this value (mHa): 100 in Stage 1, and 100 (46 qubits) or 500 (70 qubits) in Stage 2.
Auxiliary-term threshold	`gate-loose-mha`	Accuracy threshold (50 mHa) inside which the gated reward terms (excitation and coverage penalties) apply; wider than chemical accuracy (1.6 mHa) because the training molecules are far from it.
Excitation penalty weight, λ_{exc}	`cost-lambda-excitation`	Weight of the penalty on the share of sampled determinants of excitation rank above three.
Coverage penalty weight, λ_{cov}	`cost-lambda-coverage`	Weight of the penalty on the fraction of qubits that carry no gate in the transpiled circuit.
Depth-overshoot penalty weight, λ_{over}	`cost-lambda-overshoot`	Weight of the penalty on $\log_{10}(d/d_{\max})$ when the depth d exceeds the ceiling $d_{\max}$.
Hartree–Fock penalty weight, λ_{HF}	`cost-lambda-hf`	Weight of the penalty on the share of right-spin-sector shots that is the Hartree–Fock determinant.
Spin-balance penalty weight, λ_{spin}	`cost-lambda-spin`	Weight of the penalty on the imbalance between alpha-only and beta-only excitations.
Entropy coefficient, c_H	`entropy-coef`	Weight of the mean token entropy of the sampled sequences, subtracted from the loss to keep the policy from collapsing.
Gate-count, subspace-size and real-depth weights	`cost-lambda`, `cost-lambda-dim`, `cost-lambda-real-depth`	Optional accuracy-gated penalties on a proxy gate count, on $\log_{10}$ of the diagonalisation subspace size, and on $\log_{10}$ of the backend-transpiled depth; all zero in every run reported here.
Subspace recovery and diagonalisation		
Recovery size, D	`max-dim` (training), `-subspace-dim` (inference), `subspace_dim`	Target number of strings per spin sector kept in each configuration-recovery step; the resulting determinant subspace contains up to D^2 determinants (closed-shell product subspace). $D = 20{,}000$ in the final recoveries; $D = 40{,}000$ as the training cap in Stage 2.
Number of recovery iterations, K	`recover-max-iterations`, `FINAL_RECOVER_ITERS`	Number of configuration-recovery iterations; 25 in the final recoveries.
Independent subspaces per step, B	`num_batches`	Number of independent subspaces drawn from the recovered distribution in one step, of which the lowest-energy one is kept.
Ranking cap	`-max-dim` (inference)	Smaller cap used only to rank the ten sampled circuits before the final recovery: 200 strings per spin sector (net dimension 40,000); the E+D rule of the trained policies uses the same ranking. Only the E+D picks of the eight untrained controls, sampled in the first campaign, come from the earlier ranking with a cap of 2,000 determinants (`RANKING_PAIR_CAP`).
Depth limit for diagonalisation	`-max-depth-for-diag`	Circuits deeper than this are not diagonalised at inference (50,000 for the 70-qubit lineages, matching the training ceiling).
Software and devices		
Recovery routine	`recover_and_diagonalize`, `qiskit_addon_sqd`	Our GPU implementation of the configuration-recovery loop of SQD, built on the Qiskit SQD addon.

Name in the text	Code label	Definition
Local unitary cluster Jastrow ansatz	`UCJOpSpinBalanced`, `n_reps`	ffsim's spin-balanced LUCJ operator built from CCSD amplitudes; `n_reps` is the number of ansatz repetitions (1 in the LUCJ baseline).
Hardware backend	`ibm_kobe`, `ibm_kawasaki`, `ibm_heron_r2`	IBM superconducting backends (Heron generation). `ibm_kobe` is the physical processor used for all reported sampling. `ibm_kawasaki` was used only in preliminary runs not included in the reported analyses. `ibm_heron_r2` is an offline backend model whose coupling map is used for the per-iteration training transpilation; depths reported for training are on this model (not necessarily the circuits submitted). Inference depths are live-transpiled for `ibm_kobe` and submitted.

D Hyperparameter tuning: preliminary runs and reward search

This appendix describes the preliminary runs and the joint search behind the settings of Section 2.8 and the tuned reward of Section 2.8.1. The tables of the tuning (Tables A4 and A6) are in Appendix E.

Settings. Table 3 lists the settings of the four lineages; the Stage-2 reward weights of the two configurations are in Table 2 and Section 2.8.1, and Table A3 defines every setting by its name in the code. The optimiser, the size curriculum, the seed and the recovery settings are the same for all runs.

Most of these settings were fixed in preliminary Stage-2 runs, which also determined the admission gate. Two findings from them are used here:

1. **The depth ceiling is the dominant lever of the admission gate.** In a sweep over the depth ceiling and the coverage and overshoot weights, the ceiling decided whether training stalled with zero gradient: with a loose ceiling a large share of the sampled circuits cleared the admission gate and contributed a gradient update, and the particular coverage and overshoot weights mattered much less; a tight ceiling stalled regardless of the reward shaping.

2. **The energy admission tolerance must be loosened for the 70-qubit targets.** With the tolerance that worked at 46 qubits, no circuit cleared the gate over many consecutive iterations at 70 qubits. Without an admitted iteration there is no gradient, so the policy distribution cannot change (the next stochastic batch can still differ), a deadlock and not slow learning. Loosening the tolerance by a factor of five (100 to 500 mHa) restored a non-trivial pass rate.

Why the Hartree–Fock, spin and entropy terms were added. In the first Stage-2 run on `ibm_kobe`, the policy converged to a circuit with no gate able to change the Hartree–Fock determinant, and 94% of its valid shots were the Hartree–Fock determinant itself; this collapse happens before the energy gate opens, which is why the Hartree–Fock term is ungated. With the Hartree–Fock penalty alone, the 46-qubit lineages settled on circuits whose valid shots were mostly beta-only (about 60%, with almost no alpha-only weight), which motivated the spin term. Without the entropy term, the effective vocabulary of the policy, $\exp \overline{H}$, fell from about 100 to 1 in all lineages, which removes the exploration that the group-relative advantage needs.

How the weights were chosen. Three stages, in order, each informed by the previous one. (i) Preliminary runs and their depth sweeps (Table A4). $\lambda_{\text{cov}} = 10$, $\lambda_{\text{over}} = 2$ and $\lambda_{\text{exc}} = 2$ were fixed in the 36-job sweep over the depth ceiling and the coverage and overshoot weights: the depth ceiling was the dominant lever, and several weight combinations performed comparably. λ_{exc} was set based on an observation in an earlier zero-shot run of a Stage-1 checkpoint on $Sn(OH)_3$-Me L (accepted circuit with 97.5% high-rank weight); the organotin targets were therefore not untouched when this weight was set, and λ_{exc} was not swept. (ii) A one-term test of λ_{HF}. On a 24-qubit CPU test (C_2H_4, 12 orbitals, 3 seeds, 60 iterations, symmetric 5% readout flips) we compared $\lambda_{\text{HF}} \in \{0, 0.05, 0.2, 1.0\}$ and a noise-free control. At 0.05 the Hartree–Fock share of the valid shots fell from 57% to 37%, the number of HOMO–LUMO rotations per circuit doubled and the raw error stayed within 1 mHa of the unpenalized run (10.7

against 11.8 mHa). At 0.2 and 1.0 the share fell to 1% and 0% but the raw error rose to 44 and 75 mHa: the policy escaped Hartree–Fock by landing on one excited determinant, not on a correlated state. We therefore started the hardware runs with $\lambda_{\mathrm{HF}} = 0.05$. (iii) A joint search. With the penalty on hardware, the policy still collapsed to one sequence, and in the 46-qubit lineages onto a spin-asymmetric one. A grid of the spin weight ($\lambda_{\mathrm{spin}} \in \{0, 0.2, 0.6\}$) and the entropy coefficient ($c_H \in \{0, 0.05, 0.2\}$) at fixed $\lambda_{\mathrm{HF}} = 0.05$, with symmetric and with asymmetric readout noise, showed that $c_H = 0.2$ was the only setting that kept the effective vocabulary above one, that the spin weight changed nothing under symmetric noise (alpha and beta weights agreed to 0.001 in every arm), that $\lambda_{\mathrm{spin}} = 0.2$ barely reduced the imbalance under asymmetric noise (0.114 against 0.128), and that $\lambda_{\mathrm{spin}} = 0.6$ removed it only by degenerating (Hartree–Fock share 0.01 and raw error 62 mHa). Since the five weights interact (a larger λ_{cov} or λ_{exc} changes which circuits the entropy bonus protects), we then searched them together: 40 configurations drawn as a Latin hypercube over $\lambda_{\mathrm{HF}} \in [0, 0.3]$, $\lambda_{\mathrm{spin}} \in [0, 1]$, $c_H \in [0, 0.3]$, $\lambda_{\mathrm{exc}} \in [0, 4]$ and $\lambda_{\mathrm{cov}} \in [0, 20]$, plus the two reference settings ($\lambda_{\mathrm{HF}} = 0$ and 0.05, the rest at their reference values), two seeds each, 80 iterations each, on the same 24-qubit test. This search used a noise model fitted to the `ibm_kobe` counts of the first Stage-2 run: 3,000 circuits at 23 orbitals show that the shots in the right spin sector fall from 49% at depth 45–60 to 12% at depth 60–100 and to 2–3% above depth 300, equivalent to independent per-bit flip probabilities of 0.017, 0.07 and 0.3–0.45. The simulator therefore flips each bit with probability $p(d) = 0.5\,[1 - \exp(-(d - 45)/200)]$ (plus 1% readout noise) from the transpiled depth d, which reproduces 0.017, 0.08 and 0.42 at depths 52, 80 and 400, and applies it to the alpha and beta registers separately. The model is crude: independent flips cannot reproduce which bitstrings the device returns.

Selection. Each configuration was scored on the last 20 of its 80 iterations by the raw error of the best circuit, the share of the valid shots that is Hartree–Fock, the number of gates that can change that determinant, the effective vocabulary and the depth. We required the raw error to be at most 20 mHa (the best unpenalized runs reach 11 mHa), the Hartree–Fock share at most 0.30, an effective vocabulary of at least 3 tokens, and at least 1.5 gates per circuit able to change the determinant. Table A6 lists all 42 configurations, 80 runs in all. The hardware runs had to start before the search had finished, so the configuration used for them (number 3: $\lambda_{\mathrm{HF}} = 0.07$, $\lambda_{\mathrm{spin}} = 0.77$, $c_H = 0.26$, $\lambda_{\mathrm{exc}} = 3.9$, $\lambda_{\mathrm{cov}} = 7.2$) was chosen on the 26 configurations that had a finished run at that point, not on the final ranking.

Of the 42 configurations, five meet the criteria: numbers 24, 2, 16, 32 and 3, with mean raw errors of 11.9, 12.4, 12.5, 14.5 and 14.8 mHa. The selected configuration is the fifth of these and tenth overall by raw error. Over its two seeds it gave 16.8 and 12.7 mHa, a Hartree–Fock share of 0.17–0.21, an effective vocabulary of 5.8–7.5 tokens, 2.9–3.8 Hartree–Fock-changing gates per circuit and a depth of 63–68. The unpenalized reference (configuration 0) gave 104 and 11.1 mHa over its two seeds: one seed collapsed at depth 7,056 and the other onto the Hartree–Fock circuit, with 63% of its valid shots Hartree–Fock. $\lambda_{\mathrm{HF}} = 0.05$ alone (configuration 1) gave 123 and 10.8 mHa. Two seeds per configuration cannot separate the five passing configurations from each other: their errors differ by less than the difference between the two seeds of the selected one. The ranking therefore supports the choice of a region of the weight space (a Hartree–Fock weight of 0.04–0.15, a spin weight above 0.2, an entropy coefficient above 0.1), not of one point in it. Configurations 7, 13 and 22 reach 11–13 mHa at shallow depth, but their policy has collapsed to 1.4–2.2 effective tokens, and configuration 37 (11.4 mHa, 3.5 tokens) fails on its Hartree–Fock share of 0.34; the five passing configurations keep 3.7–9.6 effective tokens. The spin weight cannot be judged in this simulator, because its readout noise is the same for the alpha and beta registers in the joint search (the imbalance is below 0.002 in every row); its value (0.77) was therefore chosen heuristically for the hardware runs, because the hardware showed one-sided excitation, and it has not been tested against that failure.

Limits. The search used one molecule, 24 qubits and 80 iterations on a noise model that is fitted to one hardware run, with two seeds per configuration. It shows which settings keep the policy diverse and the raw error moderate under that model. It does not show that the selected weights improve the hardware result; the hardware runs with this setting are reported next to the runs with $\lambda_{\mathrm{HF}} = 0.05$ and without a penalty (Section 3 and Appendix F).

E Training detail: hyperparameter tuning, comparison of the runs and timing

This appendix collects the tables of the hyperparameter tuning (Appendix D): the preliminary runs (Table A4), the comparison of the training runs (Table A5), the reward search (Table A6), and the training time and billed QPU time (Tables A7 and A8). It is referenced from Section 2.8 and Section 3.

Table A4: Hyperparameter tuning in the preliminary runs: what was varied, what the sweep or ablation found, and the setting this motivated for the final runs. The names are the descriptive names of the code labels defined in Table A3.

Setting	What was varied	What was found	Setting used
Depth ceiling	Swept jointly with the two reward weights below (grid over ceiling × coverage × overshoot weight) on the ≤32q curriculum.	The ceiling, not the reward-weight combination, was the dominant lever for avoiding sustained zero-gradient stalls: a loose ceiling let a wide range of reward weights learn comparably; a tight one stalled regardless of reward shaping.	10000 (46q), 20000 (70q, Stage 1); loosened to 50000 for 70q Stage 2 once hardware depths ran higher than in simulation.
Coverage and overshoot weights	Same grid as above.	Secondary to the ceiling: multiple weight combinations performed comparably once the ceiling itself was set generously; no single combination was uniquely best across all swept ceilings.	10.0 / 2.0, all four lineages, both stages.
Energy admission tolerance	Observed directly on the first Stage-2 attempts at the 70q target: with the Stage-1 value (100), the energy tolerance never admitted a single iteration across many consecutive hardware iterations.	A structural deadlock, not slow learning: with zero admitted iterations the policy cannot update, so the policy distribution cannot change (the next stochastic batch can still differ). Loosening by 5× broke the deadlock immediately; both 70q lineages then admitted 99–100% of their iterations for the rest of training.	100 (46q, both stages; never hit this deadlock), 500 (70q, Stage 2 only).
Gate tokens per circuit	Fixed per (encoder, qubit-target) lineage before training begins, not swept within a lineage.	Cannot change across a warmstart transfer: it sizes the transformer's position-embedding table ($n_{\text{positions}}$ = ngates+1), so Stage 1 and Stage 2 of one lineage must agree on it by construction.	28 (46q lineage), 44 (70q lineage).
Recovery subspace size	Fixed at 20000 for every Stage-1 run (ffsim); raised to 40000 for every Stage-2 run (hardware).	Not independently swept here; chosen to give the hardware stage, which collects far fewer effective samples per circuit than ffsim's exact simulation, more room in the per-iteration subspace cap.	20000 (Stage 1), 40000 (Stage 2), uniformly.

Fraction of molecules counted as largest	Controlled pair of runs on a 5-molecule subset, 0.2 (the original default) vs. 0.4, holding every other setting fixed.	0.4 won on every hard molecule measured: co2 21.2 vs. 112.4 mHa, fe_cn6 132.7 vs. 240.8 mHa, co_nh3_6 294.1 vs. 459.1 mHa (for ferrocene the difference was smaller: 173.5 vs. 181.1 mHa). A wider "hard" set (top 40% of the molecule pool by qubit count, rather than top 20%) gave the curriculum's focus phase more molecules to distribute attention across.	0.4, all four lineages.
Fraction of iterations on the largest molecules	Changed alongside curriculum-top-frac adjusted in the same tuning round (0.5→0.3), but not isolated in its own A/B test the way curriculum-top-frac was.	The value 0.3 was carried over from the preceding tuning round on the rationale that a shorter focus phase leaves more of the run for the final consolidation phase; this rationale was not tested in isolation.	0.3, all four lineages (a fixed choice, not swept).
Learning rate	Controlled pair of runs on a 5-molecule subset, the existing default 5×10^{-6} vs. a 4× higher 2×10^{-5}.	The higher rate reached a numerically higher early gate-pass rate (41.7%, 10/24 admitted iterations) but produced the worst final accuracy of any hyperparameter test run in this comparison (fe_cn6 312 mHa, ferrocene 478 mHa) – confirming the pre-existing default was not improved on by this specific alternative, rather than establishing 5×10^{-6} as independently optimal from a wider sweep.	5×10^{-6}, all four lineages, both stages.
Gradient-norm clipping threshold	Not swept over a grid; recalibrated by measuring the policy's gradient-norm distribution across 24 training runs, after the originally-used value of 1.0 was found to fire on essentially every accepted step.	Raw gradient norms' 90th percentile measured at ≈8.2 across those runs; clip threshold set to 8.0 (its measured 90th percentile) so it acts as an occasional safety net rather than a de facto fixed step size – confirmed to then fire on only 3.5% of accepted steps (90/2561) in the recalibrated runs. This recalibration showed no measurable accuracy or gate-pass-rate benefit over no clipping at all (medians differed by only 1–3 percentage points across 8 seeds); it was adopted as zero-cost insurance against rare large updates, not because it measurably improved training.	8.0, all four lineages, both stages.

Table A5: The training runs (single seed per run), by stage and lineage. n is the number of iterations of the run; "gate" the share of iterations that passed the admission gate; "error" the mean raw SQD error (mHa) of the best circuit over the admitted iterations of the last 30 iterations ("last") and of the whole run ("all"); "depth" the median depth of the best circuit as transpiled for the offline backend model `ibm_heron_r2` at Qiskit optimization level 1 (not the circuit submitted) and ">1000" the share of iterations whose best circuit is deeper than 1,000; "eff. voc." the policy's effective vocabulary defined as $\exp(\bar{H})$ with $\bar{H}$ the token entropy in nats over all 110 token ids (including padding); "valid" the share of shots in the right spin sector, "HF" the share of the valid shots that is the Hartree–Fock determinant, "HL" the share that is a HOMO→LUMO excitation and "$|\alpha - \beta|$" the difference of the alpha-only and beta-only weights among the valid shots, all tracked over the last 30 iterations. Stage 1 is shared by the three Stage-2 variants. Runs that were stopped (Stage 2 without penalty, A5-46q and A5-70q) are shown with the iterations they have; the tuned run of A1-70q ended by a segmentation fault at iteration 137.

Lineage	Run	n	gate (%)	error last	error all	depth	>1000 (%)	eff. voc.	valid	HF	HL	$\lvert\alpha - \beta\rvert$
A1-46q	Stage 1	150	84	36.1	26.5	4394	92	1.0	1.00	0.76	0.16	0.000
A1-46q	Stage 2, no penalty	150	51	26.2	37.3	48	0	1.0	0.51	0.95	0.00	0.008
A1-46q	Stage 2, HF penalty	150	46	48.1	46.2	56	0	1.0	0.10	0.02	0.07	0.558
A1-46q	Stage 2, tuned	150	59	48.0	46.7	44	5	21.5	0.03	0.03	0.00	0.018
A5-46q	Stage 1	150	82	28.1	22.3	6008	95	1.0	1.00	0.70	0.15	0.000
A5-46q	Stage 2, no penalty	17	47	43.8	43.8	39	0	61.2	0.05	0.06	0.05	0.064
A5-46q	Stage 2, HF penalty	150	47	45.0	46.5	53	0	1.0	0.21	0.00	0.00	0.666
A5-46q	Stage 2, tuned	150	63	56.7	52.5	43	7	20.9	0.04	0.06	0.00	0.043
A1-70q	Stage 1	150	86	27.6	24.0	7008	93	1.1	1.00	0.72	0.14	0.000
A1-70q	Stage 2, HF penalty	150	100	287.9	300.6	76	0	1.0	0.07	0.01	0.03	0.045
A1-70q	Stage 2, tuned	137	99	284.7	286.6	79	39	26.1	0.02	0.00	0.00	0.014
A5-70q	Stage 1	150	86	21.1	20.9	6832	91	1.2	1.00	0.60	0.20	0.000
A5-70q	Stage 2, no penalty	18	94	348.4	348.4	66	6	56.5	0.02	0.02	0.00	0.020
A5-70q	Stage 2, HF penalty	150	99	276.0	294.1	152	1	1.0	0.06	0.04	0.15	0.073
A5-70q	Stage 2, tuned	150	100	269.6	285.4	79	39	27.0	0.02	0.01	0.00	0.015

The training runs. Tables A5 and A7 compare the Stage-1 run with the three Stage-2 variants, which differ only in the reward. No penalty uses the excitation, coverage and overshoot terms only. HF penalty, the main configuration, adds the Hartree–Fock term ($\lambda_{\mathrm{HF}} = 0.05$). Tuned uses the jointly tuned reward of Table 4. Stage 1 is shared by the three variants (one Stage-1 run per lineage). The no-penalty Stage-2 run finished only for A1-46q; the other three lineages were stopped after 17 or 18 iterations (A1-70q never started its first hardware batch) so that the device could be used for the tuned runs. The tuned runs reached 150 iterations, except A1-70q, which ended by a segmentation fault in the transpiler at iteration 137; its latest saved policy is that of iteration 129, and no later checkpoint exists.

Why the wall times differ. The per-iteration times of Table A7 differ between runs for reasons that we can separate only in part. (i) Stage-1 iterations are slower at 70 qubits (539 and 427 s of sampling against 284 and 445 s at 46 qubits for A1 and A5), because the policy generates $L = 44$ gate tokens per circuit instead of 28 and the simulator applies them; the longest single sampling step was 57 min (the ferrocene iteration of each 14-molecule cycle). (ii) In Stage 2, the sampling column is the time to get a batch of 20 circuits back from `ibm_kobe`, which includes the device queue, not the QPU time. It grows from the no-penalty run (212 to 267 s) to the HF-penalty runs (315 to 368 s) and the tuned runs (523 to 569 s). Two effects add up. The load on the queue was not constant: between one and eight of our own Stage-2 jobs were submitted at the same time, on top of other tenants' jobs, and the sampling time of the same lineage is not reproducible across these periods; we did not control for this. The circuits themselves also differ. A collapsed policy emits the same short circuit twenty times, whereas the tuned policy, which does not collapse, emits twenty different ones and, at 70 qubits, many deep ones. The billed QPU time per job, which excludes the queue, shows this: it is 59 to 60 s for the 46-qubit HF-penalty runs and 83 to 85 s for the tuned ones, and 70 s against 159 to 167 s at 70 qubits (Table A8). Diagonalization, which does not depend on the device, is similar in all Stage-2 regimes (215 to 330 s). (iii) The stalled iterations (Table A7, last column) are excluded from the Stage-2 means. There are six, one in each of six runs. Four (no-penalty A1-46q iteration 102, and the first iteration of three HF-penalty runs) are attributed to leftover Runtime jobs, which we did not verify run by run; two (tuned 70-qubit, iteration 128) are unexplained. One iteration of the no-penalty A1-46q run (8.8 h of sampling) waited behind Qiskit Runtime jobs that cancelled Slurm jobs had left queued on `ibm_kobe`, and the first iteration of each HF-penalty run of A5-46q, A1-70q and A5-70q (2.3 to 2.4 h) waited behind the same leftover jobs from the runs we had just replaced. Two were not: iteration 128 of both 70-qubit tuned runs (1.8 h each) fell into a period of about three hours in which `ibm_kobe` started none of our jobs. We do not know the reason; the device reported itself operational throughout.

The chemistry that each regime produced is in Appendix F. The comparison of the energies in Table A5 shows no regime that is better than the others on hardware, and the trade-offs between them are discussed there.

Table A6: Joint reward search on the noisy 24-qubit simulator: all 42 configurations, ranked by the mean raw error of the best circuit over the last 20 of 80 iterations (mean over the finished seeds, two for all but configuration 34, which has one, 80 runs in all). Configurations 0 and 1 are the two references (no Hartree–Fock penalty, and $\lambda_{\mathrm{HF}} = 0.05$ alone, the other weights at their reference values); 2–41 are Latin-hypercube draws over the ranges of Table 4. Columns: the five weights and the entropy coefficient c_H; the share of the valid (right-spin-sector) shots that is the Hartree–Fock determinant; the policy's effective vocabulary; the number of gates per circuit that can change the Hartree–Fock determinant; the share of valid shots on HOMO→LUMO determinants; the share of all shots that is valid; the median depth. A check mark means: raw error $\leq$ 20 mHa, Hartree–Fock share ≤ 0.30, effective vocabulary ≥ 3 and ≥ 1.5 gates able to change the determinant (5 of 42). The configuration used for the hardware runs (3) is in bold. The alpha/beta imbalance is below 0.002 in every row, because the noise of this simulator is the same for both registers.

rank	cfg	λ_{HF}	λ_{spin}	c_H	λ_{exc}	λ_{cov}	error (mHa)	HF share	eff. voc.	HF-changing	HOMO→LUMO	valid	depth	ok
1	7	0.01	0.86	0.14	3.2	10.7	11.0	0.52	1.4	1.07	0.17	0.31	66	
2	37	0.06	0.90	0.18	3.0	19.1	11.4	0.34	3.5	1.82	0.20	0.26	70	
3	24	0.07	0.26	0.23	3.2	7.7	11.9	0.28	3.8	2.28	0.22	0.24	70	✓
4	22	0.05	0.38	0.18	0.3	2.8	12.3	0.21	2.2	1.33	0.06	0.13	99	
5	2	0.15	0.72	0.19	0.2	0.8	12.4	0.10	3.7	3.49	0.08	0.11	104	✓
6	16	0.04	0.68	0.28	2.3	13.7	12.5	0.23	6.3	2.74	0.13	0.18	69	✓
7	19	0.02	0.21	0.11	0.8	17.5	12.7	0.39	1.1	0.53	0.08	0.18	87	
8	13	0.02	0.25	0.16	3.7	16.0	12.9	0.37	1.9	2.14	0.26	0.38	63	
9	32	0.04	0.53	0.29	1.2	2.3	14.5	0.20	9.6	3.95	0.06	0.14	76	✓
10	**3**	**0.07**	**0.77**	**0.26**	**3.9**	**7.2**	**14.8**	**0.19**	**6.7**	**3.36**	**0.17**	**0.22**	**66**	✓
11	29	0.14	0.67	0.03	0.5	14.4	28.0	0.10	1.0	1.50	0.44	0.27	88	
12	17	0.25	0.09	0.21	1.5	15.4	29.4	0.03	3.6	3.71	0.31	0.16	84	
13	12	0.12	0.34	0.24	4.0	4.3	38.1	0.04	5.9	3.69	0.11	0.10	97	
14	21	0.18	0.28	0.03	1.7	15.7	44.1	0.02	1.0	3.50	0.78	0.48	63	
15	0	0.00	0.00	0.00	2.0	10.0	57.6	0.31	1.0	4.00	0.10	0.23	3560	
16	26	0.10	0.99	0.11	2.0	9.2	57.7	0.07	1.1	4.25	0.39	0.31	2715	
17	33	0.22	0.32	0.04	0.9	14.5	66.5	0.02	1.0	3.50	0.11	0.09	1478	
18	1	0.05	0.00	0.00	2.0	10.0	66.8	0.18	1.0	3.50	0.13	0.14	2221	
19	36	0.18	0.15	0.20	1.1	5.8	68.2	0.03	5.8	2.85	0.13	0.10	1952	
20	39	0.26	0.41	0.17	1.5	13.3	69.5	0.01	2.7	4.07	0.13	0.09	1918	
21	20	0.27	0.02	0.07	2.6	17.4	73.2	0.01	1.0	5.67	0.06	0.08	3752	
22	6	0.28	0.19	0.10	1.0	19.8	73.3	0.01	1.1	6.46	0.38	0.21	549	
23	5	0.27	0.48	0.15	2.7	11.0	86.9	0.02	2.4	5.39	0.29	0.24	3775	
24	23	0.09	0.15	0.10	2.1	10.3	89.1	0.16	1.1	2.68	0.01	0.10	3578	
25	8	0.10	0.91	0.05	0.9	8.7	89.3	0.16	1.0	1.50	0.06	0.11	2370	
26	11	0.30	0.64	0.08	3.3	1.6	90.2	0.01	1.0	4.51	0.07	0.09	2050	
27	9	0.14	0.81	0.27	2.2	16.9	95.6	0.00	7.4	7.98	0.00	0.05	6364	
28	10	0.07	0.45	0.02	2.5	12.6	100.6	0.23	1.0	3.00	0.08	0.14	1737	
29	34	0.19	0.12	0.22	1.8	1.1	104.7	0.00	5.5	4.19	0.02	0.06	2420	
30	25	0.13	0.74	0.21	2.9	3.3	109.2	0.01	6.6	5.45	0.02	0.06	1414	
31	41	0.17	0.56	0.15	3.5	7.0	116.7	0.00	2.8	2.54	0.00	0.05	4786	
32	30	0.21	0.59	0.29	2.0	18.4	118.0	0.00	10.2	5.08	0.00	0.05	2318	
33	15	0.23	0.38	0.26	1.8	12.3	119.8	0.01	10.7	5.04	0.00	0.05	4318	
34	40	0.00	0.61	0.00	3.4	0.4	122.6	0.00	1.0	3.01	0.00	0.05	5265	
35	27	0.20	0.83	0.05	3.0	6.4	125.9	0.00	1.0	3.50	0.23	0.11	296	
36	31	0.23	0.06	0.24	1.3	8.2	135.5	0.00	7.2	5.38	0.00	0.05	3157	
37	38	0.24	0.51	0.14	0.4	18.6	138.2	0.00	2.3	3.96	0.00	0.05	5354	
38	14	0.29	0.94	0.01	3.8	3.6	139.4	0.00	1.0	7.50	0.00	0.05	6970	
39	35	0.09	0.78	0.12	0.4	5.1	139.5	0.00	1.4	2.55	0.00	0.05	7718	
40	28	0.20	0.96	0.25	2.8	11.6	146.4	0.00	7.1	4.09	0.00	0.05	3961	
41	18	0.16	0.47	0.06	0.1	9.5	155.6	0.00	1.0	2.99	0.00	0.05	4039	
42	4	0.12	0.04	0.08	0.6	4.5	159.0	0.00	1.0	3.23	0.00	0.05	2786	

Table A7: Time per iteration (seconds, mean over the iterations of each run) by pipeline step, Stage 1 on the simulator and Stage 2 on `ibm_kobe`; the last column is the analysed elapsed time of the run with the stalled Stage-2 iterations removed. In the Stage-2 rows, iterations whose circuit sampling took longer than one hour (device-queue stalls caused by Runtime jobs left queued by cancelled Slurm jobs) are excluded from the means and from the analysed elapsed time, and counted in the column "stalled"; Stage-1 iterations are all kept, including the ferrocene iterations, whose simulation can take up to an hour. The sampling column of Stage 2 includes the IBM device queue. n is the number of iterations the run has.

Run	Lineage	n	wall	sampling	diagonalize	refine	transpile	backward	stalled	analysed (h)
Stage 1	A1-46q	150	496	284	194	9	6	1.2	0	20.7
Stage 1	A5-46q	150	659	445	195	10	8	1.3	0	27.5
Stage 1	A1-70q	150	756	539	196	10	8	1.2	0	31.5
Stage 1	A5-70q	150	637	427	190	9	8	1.1	0	26.6
Stage 2, no penalty	A1-46q	150	494	265	215	11	1	0.8	1	20.5
Stage 2, no penalty	A5-46q	17	474	212	239	13	4	0.8	0	2.2
Stage 2, no penalty	A5-70q	18	621	267	284	15	16	1.5	0	3.1
Stage 2, HF penalty	A1-46q	150	606	348	241	13	2	0.8	0	25.3
Stage 2, HF penalty	A5-46q	150	606	368	223	11	1	0.7	1	25.1
Stage 2, HF penalty	A1-70q	150	677	315	330	16	5	1.5	1	28.0
Stage 2, HF penalty	A5-70q	150	654	348	278	14	4	1.3	1	27.1
Stage 2, tuned	A1-46q	150	791	525	241	14	5	0.8	0	33.0
Stage 2, tuned	A5-46q	150	790	527	237	15	5	0.9	0	32.9
Stage 2, tuned	A1-70q	137	938	563	298	20	23	1.4	1	35.4
Stage 2, tuned	A5-70q	150	919	546	294	20	22	1.4	1	38.1

Table A8: IBM-billed QPU time of the Stage-2 runs, from the account's job records (one Runtime job per iteration, 20 circuits of 10,000 shots). "Jobs" counts the iterations that were submitted; "mean" is the billed seconds per job. "Stopped": cancelled after 17 or 18 iterations; "crashed": A1-70q tuned ended by a segmentation fault in the transpiler at iteration 137. The 1372 jobs total 121,661 s (2,028 min, 33.8 h). These jobs are listed separately in Table 10.

Regime	Lineage	Jobs	QPU time (min)	mean (s/job)
No penalty	A1-46q	150	146.7	58.7
No penalty	A5-46q (stopped)	17	21.2	74.9
No penalty	A5-70q (stopped)	18	37.2	124.2
HF penalty	A1-46q	150	149.4	59.8
HF penalty	A5-46q	150	148.5	59.4
HF penalty	A1-70q	150	174.8	69.9
HF penalty	A5-70q	150	175.1	70.0
Tuned	A1-46q	150	211.3	84.5
Tuned	A5-46q	150	206.3	82.5
Tuned	A1-70q (crashed)	137	365.7	160.1
Tuned	A5-70q	150	391.5	156.6

F Chemistry learnt during training

This appendix shows what the policy generated in terms of the two quantities that connect a circuit to the chemistry of a closed-shell molecule, using the per-iteration tracking recorded during training. HOMO–LUMO content. The Hartree–Fock determinant fills the lowest orbitals. A rotation between two neighbouring occupied orbitals, or two neighbouring virtual ones, leaves it unchanged, whereas a rotation between the highest occupied orbital (HOMO) and the lowest virtual one (LUMO) is the only neighbouring rotation that moves an electron from an occupied to a virtual orbital. We therefore count the HOMO–LUMO rotations per circuit, the share of the right-spin-sector shots that are exactly a HOMO→LUMO excitation, and the share that is the Hartree–Fock determinant itself. Alpha/beta balance. For a closed-shell molecule the alpha and beta electrons are equivalent, so excitations that move only an alpha electron (alpha-only) and only a beta electron (beta-only) should have equal weight. A large imbalance means that the circuit, or the noise acting on it, favours one spin channel. The right-spin-sector shots (called valid shots) are those with the right number of alpha and beta electrons; all other shots (wrong electron number, or an electron moved between the two channels) are noise, because every operator in the pool conserves both numbers. Unless stated otherwise, the values reported below are means over the last 30 iterations (or over all iterations where noted), computed from per-iteration batch means over the $N = 20$ circuits. The HF, alpha-only, beta-only and both-channel shares form a partition of the valid shots, while the HOMO→LUMO share overlaps them.

Stage 1 (Figure A4). On the simulator the circuits keep a small excitation content and the two spin channels stay exactly equal (alpha-only and beta-only both 0.06 to 0.13 over the last 30 iterations; the simulator and all gates are spin-symmetric, so this is expected). The policy generates 0.5 (A1-46q) to 1.6 HOMO–LUMO rotations per circuit, 14 to 20% of the sampled weight is on HOMO→LUMO determinants and 60 to 76% is the Hartree–Fock determinant. The HOMO–LUMO content rises during Stage 1 in all lineages except A1-46q, where it falls after iteration 80 before the policy collapses onto one sequence.

Stage 2 (Figures A5, A6 and A7). Table A9 gives the last 30 iterations of each run.

The three regimes on A1-46q, compared with Stage 1. A1-46q is the one lineage for which all three Stage-2 regimes ran 150 iterations, so it is the clean comparison. Table A10 puts them next to the Stage-1 simulator run that all three started from. Two cautions apply to reading it. First, the Stage-1 shares are shares of all shots, because on the simulator every shot is valid, whereas the Stage-2 shares are shares of the valid shots, which are 51, 10 and 2.7% of all shots for the three regimes; the Stage-2 columns therefore describe few shots, and the tuned regime's the fewest. Second, Stage 1 and Stage 2 use different molecules (14 small and medium molecules against the three transition-metal complexes), so the comparison is of the same quantities, not of the same circuits.

No penalty. Stage 1 left 76% of the weight on the Hartree–Fock determinant and 17% on HOMO→LUMO, with two or three gates per circuit able to change that determinant. The unpenalized Stage 2 pushes this to its limit: no gate can change the determinant (0.0 against 2.5 per circuit), no HOMO–LUMO rotation remains (0.50 to 0.00), and 95% of the valid shots are the Hartree–Fock determinant. Of the three regimes it keeps the Stage-1 spin symmetry (0.022 and 0.030), because it excites almost nothing. HF penalty. It reverses the collapse onto Hartree–Fock (2% of the valid shots) and has more HOMO–LUMO rotations than Stage 1 (3.0 against 0.5) and a similar share of HOMO→LUMO weight (7% against 17%), but breaks the Stage-1 spin symmetry, with 57% of the valid shots beta-only and 2% alpha-only. Tuned. It removes the Hartree–Fock circuit (3% of the valid shots) and keeps the alpha and beta weights within a factor of 2.1 of each other (0.017 and 0.035), closer to the Stage-1 symmetry than the HF penalty, but carries almost no HOMO→LUMO weight (0.3% against 17%) and fewer HOMO–LUMO rotations than Stage 1 (0.23 against 0.50). Most of its valid shots are excited in both channels at once (92%), against 11% in Stage 1.

The three regimes on the other lineages. The same comparison for A5-46q and the two 70-qubit lineages leaves out the no-penalty regime, whose runs were stopped after 17 or 18 iterations (A1-70q has none). For the HF penalty and the tuned reward the picture of A1-46q repeats. The HF penalty gives HOMO–LUMO rotation counts of 2.0 (A5-46q, A5-70q) and 1.0 (A1-70q) per circuit and beta-only weights of 0.67 (A5-46q), 0.12 (A1-70q) and 0.20 (A5-70q) against alpha-only weights of 0.003, 0.072 and 0.127. The tuned reward gives 0.25 to 1.06 HOMO–LUMO rotations, a beta-only to alpha-only ratio of 2.7 (A5-46q), 3.1 (A1-70q) and 1.8 (A5-70q), and HOMO→LUMO weight of at most 0.5%, with 0.4 to 6% of the valid shots on the Hartree–Fock determinant. Stage 1 of these lineages has 1.5 to 1.6 HOMO–LUMO rotations per circuit and 14 to 20% of the weight on HOMO→LUMO determinants; neither Stage-2 regime reaches the second figure except the HF penalty at A5-70q (15%).

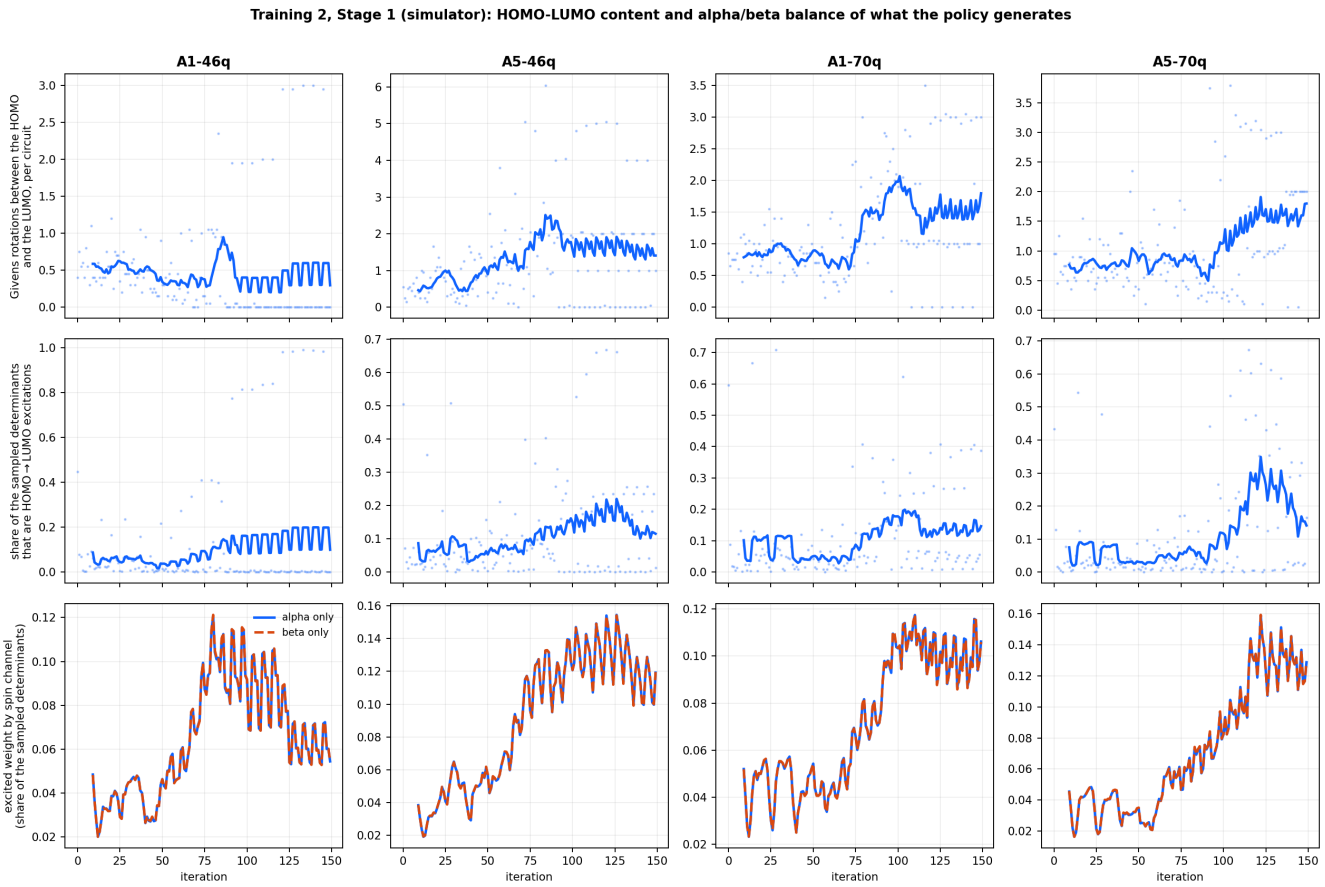

Figure A4: Stage 1 (simulator), per lineage (columns). Top: HOMO–LUMO rotations per circuit. Middle: share of the sampled determinants that are HOMO→LUMO excitations. Bottom: weight of the alpha-only (solid) and beta-only (dashed) excitations. Dots are single iterations, lines are rolling means over 10 iterations. On the simulator all shots are in the right spin sector.

Table A9: Chemistry of the last 30 iterations of the Stage-2 runs. "Valid": share of all shots in the right spin sector. The other columns are shares of the valid shots: Hartree–Fock determinant (HF), HOMO→LUMO excitation (HL), alpha-only (a) and beta-only (b) excitation, and excitation in both channels at once. b/a is the ratio of beta-only to alpha-only weight: for example, alpha-only weights 0.016 and 0.003 correspond to about 16 and 7 valid shots per circuit (at 10000 shots), so ratios such as 36 and 207 are uncertain by a large factor. HOMO–LUMO rotations (rot.) are per circuit. The tuned runs are complete (150 iterations) except A1-70q (137); the no-penalty run is complete for A1-46q only.

Lineage	Regime	valid	HF	HL	a	b	both	b/a	HL rot.
A1-46q	no penalty	0.51	0.95	0.00	0.022	0.030	0.00	1.4	0.00
A1-46q	HF penalty	0.10	0.02	0.07	0.016	0.574	0.39	36	3.00
A1-46q	tuned	0.03	0.03	0.00	0.017	0.035	0.92	2.1	0.23
A5-46q	HF penalty	0.21	0.00	0.00	0.003	0.669	0.32	207	2.00
A5-46q	tuned	0.04	0.06	0.00	0.025	0.068	0.85	2.7	0.25
A1-70q	HF penalty	0.07	0.01	0.03	0.072	0.117	0.80	1.6	1.00
A1-70q	tuned	0.02	0.01	0.00	0.024	0.040	0.93	1.7	0.57
A5-70q	HF penalty	0.06	0.04	0.15	0.127	0.200	0.63	1.6	2.00
A5-70q	tuned	0.02	0.00	0.00	0.008	0.023	0.96	2.9	1.06

Trade-offs. Each regime gives up something that another keeps.

- No penalty: usable shots, no HOMO→LUMO excitation. The only regime with a high share of valid shots (0.51 in A1-46q), because its circuit does nothing to the Hartree–Fock determinant: 95% of the valid shots are that determinant and none is a HOMO→LUMO excitation. The shots are clean because there is nothing to disturb, and the recovered energy then comes from the configuration-recovery step, not from the circuit.

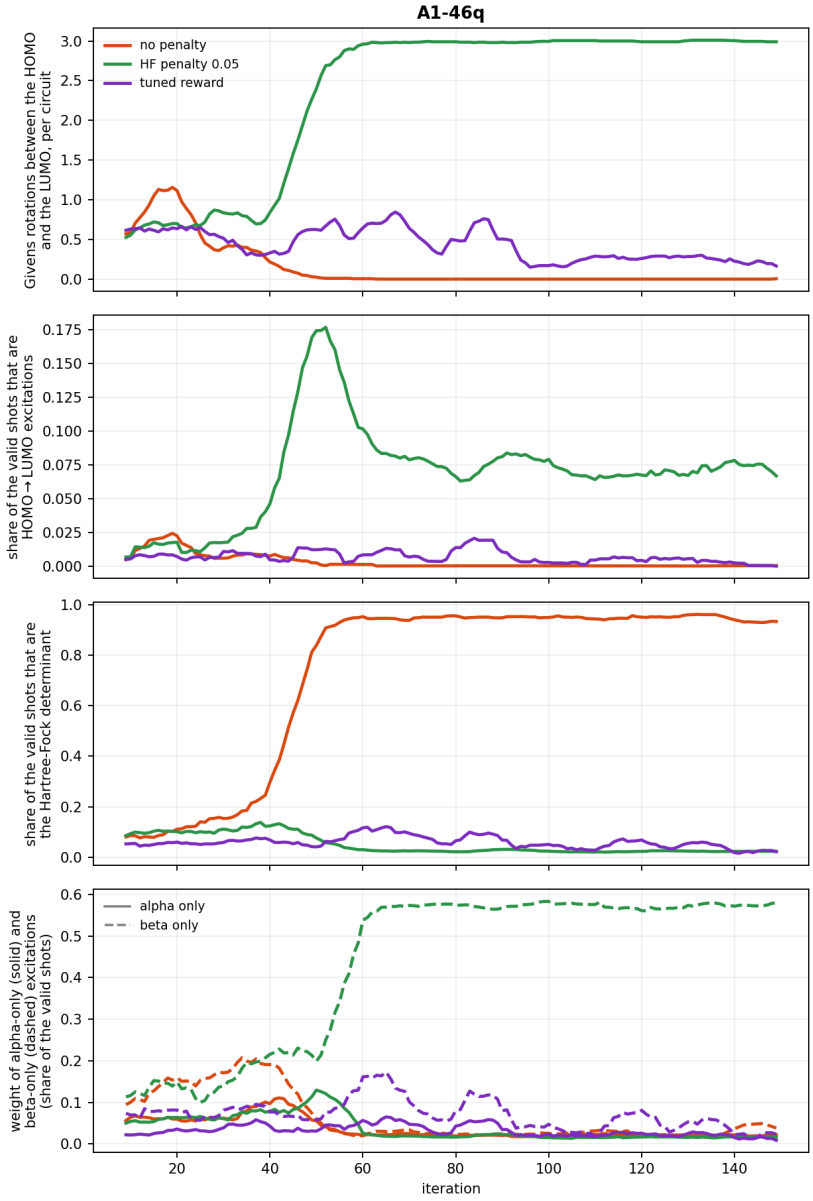

Figure A5: Stage 2 (`ibm_kobe`), A1-46q, the only lineage with all three reward regimes complete (150 iterations each): no penalty (orange), HF penalty 0.05 (green) and tuned reward (purple, 150 iterations). Rows: HOMO–LUMO rotations per circuit; share of the valid shots that are HOMO→LUMO excitations; share of the valid shots that are the Hartree–Fock determinant; weight of alpha-only (solid) and beta-only (dashed) excitations among the valid shots. Rolling means over 10 iterations. A closed-shell molecule should give equal alpha-only and beta-only weight.

- HF penalty: chemistry content, but spin-asymmetric and a collapsed policy. It removes the Hartree–Fock circuit (HF share 0.00 to 0.04) and has the most HOMO→LUMO weight (0.07 in A1-46q, 0.15 in A5-70q), but the 46-qubit circuits excite the beta channel almost alone (beta-only to alpha-only weight 36 in A1-46q and 207 in A5-46q), which a closed-shell molecule should not do, and only 7 to 21% of the shots are valid. The policy still collapses onto one sequence (effective vocabulary 1.0 in all lineages), so the penalty changes which circuit is repeated, not whether one is.
- Tuned reward: balance and diversity, but a low HOMO→LUMO content and few valid shots. It keeps the policy from collapsing (21 to 27 effective tokens in the last 30 iterations, against 1.0), keeps the alpha and beta weights within a factor of 1.7 to 3.3 of each other, and removes the Hartree–Fock circuit (HF share 0.00 to 0.06). In

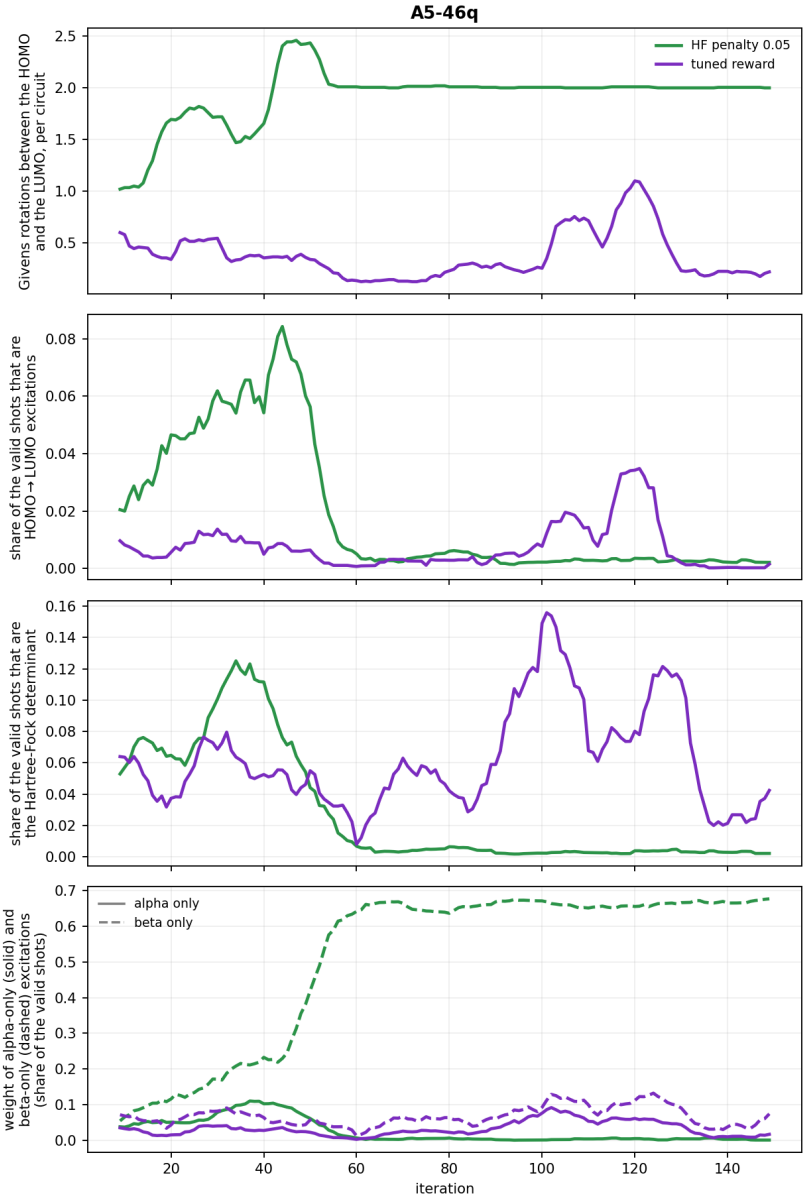

Figure A6: As Figure A5 for A5-46q, HF penalty 0.05 (green, 150 iterations) and tuned reward (purple, 150 iterations). The no-penalty run of this lineage was stopped after 17 iterations and is not shown.

return the HOMO→LUMO weight is at most 0.005, the number of HOMO–LUMO rotations is 0.23 and 0.25 in the 46-qubit tuned runs (against 3.0 and 2.0 for the HF penalty), and only 2 to 4% of the shots are valid. The valid shots are mostly excited in both channels at once (0.85 to 0.97), and with so few of them we cannot tell designed double excitations from noise that happens to land in the right sector. At 70 qubits it also emits very deep circuits in 35 to 44% of the iterations (Table A5); for circuits deeper than 1,000, the median valid-shot share is 10^{-5} and the maximum is 0.6% (214 circuits of the HF-penalty set).

- Cost. The tuned regime is the slowest per iteration and uses the most QPU time per job (Tables A7 and A8), as discussed above.
- Energy. None of this shows up as a better energy on hardware. For $Co(NH_3)_6$, the molecule whose iterations pass the gate in every run, the mean raw error over the last 45 admitted iterations is 26.2, 47.1 and 33.9 mHa for A1-46q (no penalty, HF penalty, tuned), 46.1 and 36.2 mHa for A5-46q (HF penalty, tuned; the no-penalty run

Table A10: A1-46q: the chemistry of what the policy generates in the last 30 iterations of the simulator stage and of each of the three hardware regimes. Shares are of all shots in Stage 1 (all valid) and of the valid shots in Stage 2. "HL" is the HOMO→LUMO excitation, "a" and "b" the alpha-only and beta-only excitation, "both" an excitation in both channels at once, "rot." the HOMO–LUMO rotations per circuit, "HF gates" the gates per circuit that can change the Hartree–Fock determinant, "ident." the share of identity tokens and "depth" the median depth of the best circuit.

Run	valid	HF	HL	a	b	both	HL rot.	HF gates	ident.	depth
Stage 1, simulator	1.00	0.76	0.17	0.062	0.062	0.11	0.50	2.5	0.60	3478
Stage 2, no penalty	0.51	0.95	0.00	0.022	0.030	0.00	0.00	0.0	0.36	48
Stage 2, HF penalty	0.10	0.02	0.07	0.016	0.574	0.39	3.00	3.0	0.25	56
Stage 2, tuned	0.03	0.03	0.00	0.017	0.035	0.92	0.23	1.8	0.34	48

was stopped), 225.5 and 224.1 mHa for A1-70q (HF penalty, tuned) and 225.2 and 217.5 mHa for A5-70q (HF penalty, tuned). These are single runs and the differences are of the size of the run-to-run noise.

In short, the no-penalty conclusions rest on A1-46q, the only lineage for which that regime completed 150 Stage-2 iterations. The penalties and the tuned reward change what the policy generates in the way they were designed to (no Hartree–Fock-only circuit, no collapse, balanced spin channels) but, within the noise of this device, they have not yet produced more HOMO→LUMO weight or lower energies than the plain reward. All statements rest on one run per regime and lineage.

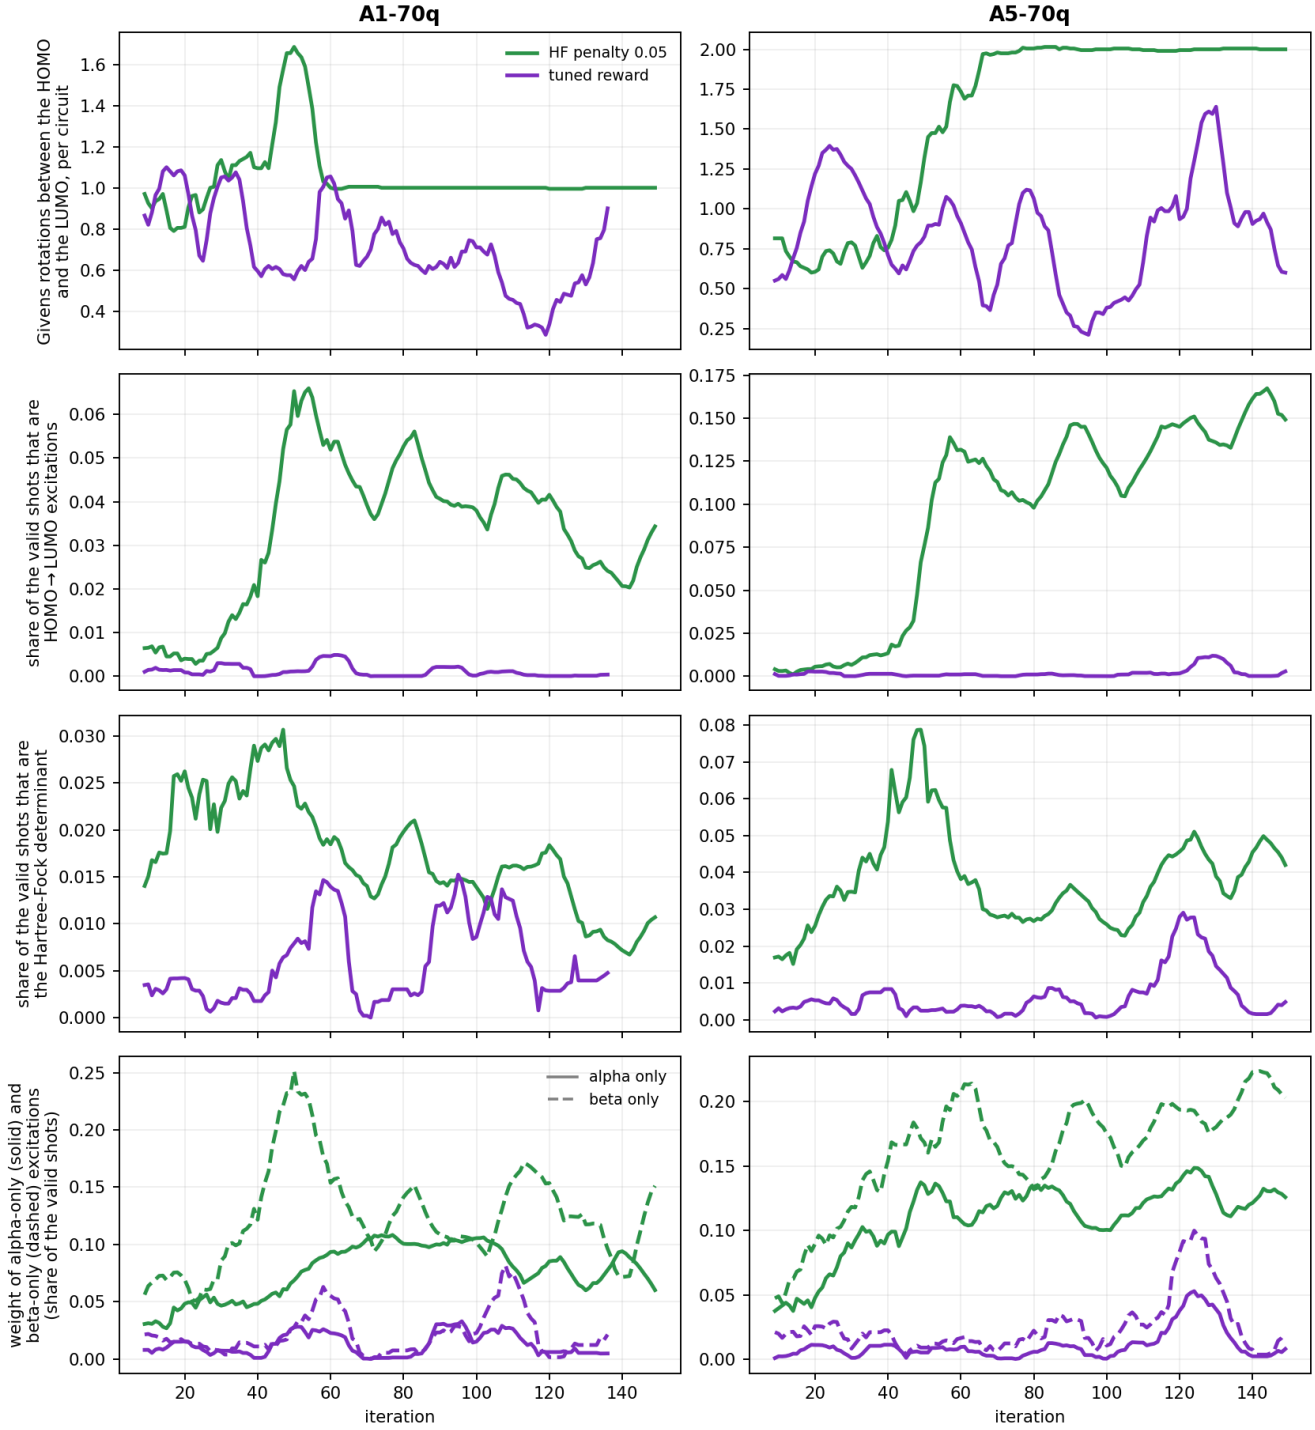

Figure A7: As Figure A5 for the 70-qubit lineages A1-70q (left) and A5-70q (right), HF penalty 0.05 (green, 150 iterations) and tuned reward (purple; 137 iterations for A1-70q, 150 for A5-70q). The no-penalty runs of these lineages were stopped after at most 18 iterations (A1-70q never started) and are not shown.

G Additional training curves

The curves below complete the main-text Figures 5, 7 and 6. All use the run histories of the shared Stage 1 and the HF-penalty Stage 2; the training histories of all runs are part of the released data.

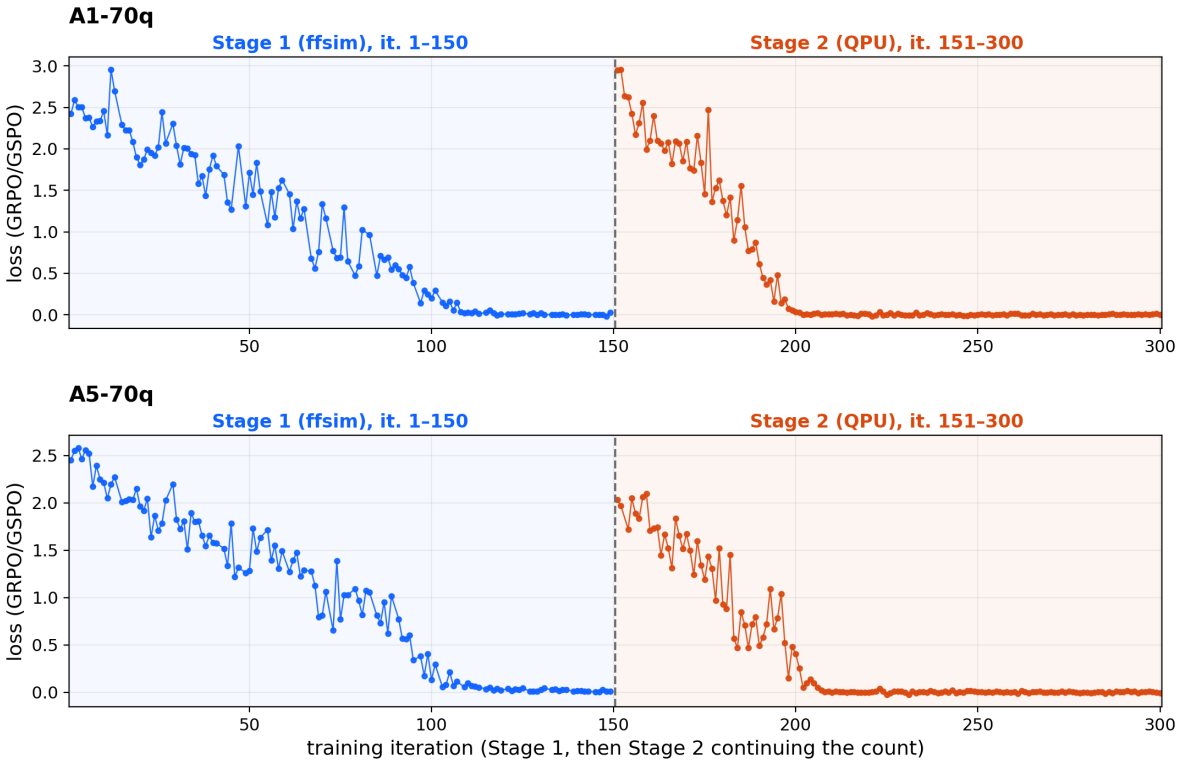

Figure A8: Training loss (GRPO/GSPO) of the 70-qubit lineages, A1-70q (top) and A5-70q (bottom), on admitted iterations. Stage 1 is the shared Stage 1 (ffsim simulator, iterations 1–150, left of the dashed line) and Stage 2 the HF-penalty run (`ibm_kobe` QPU, iterations 151–300, right of it).

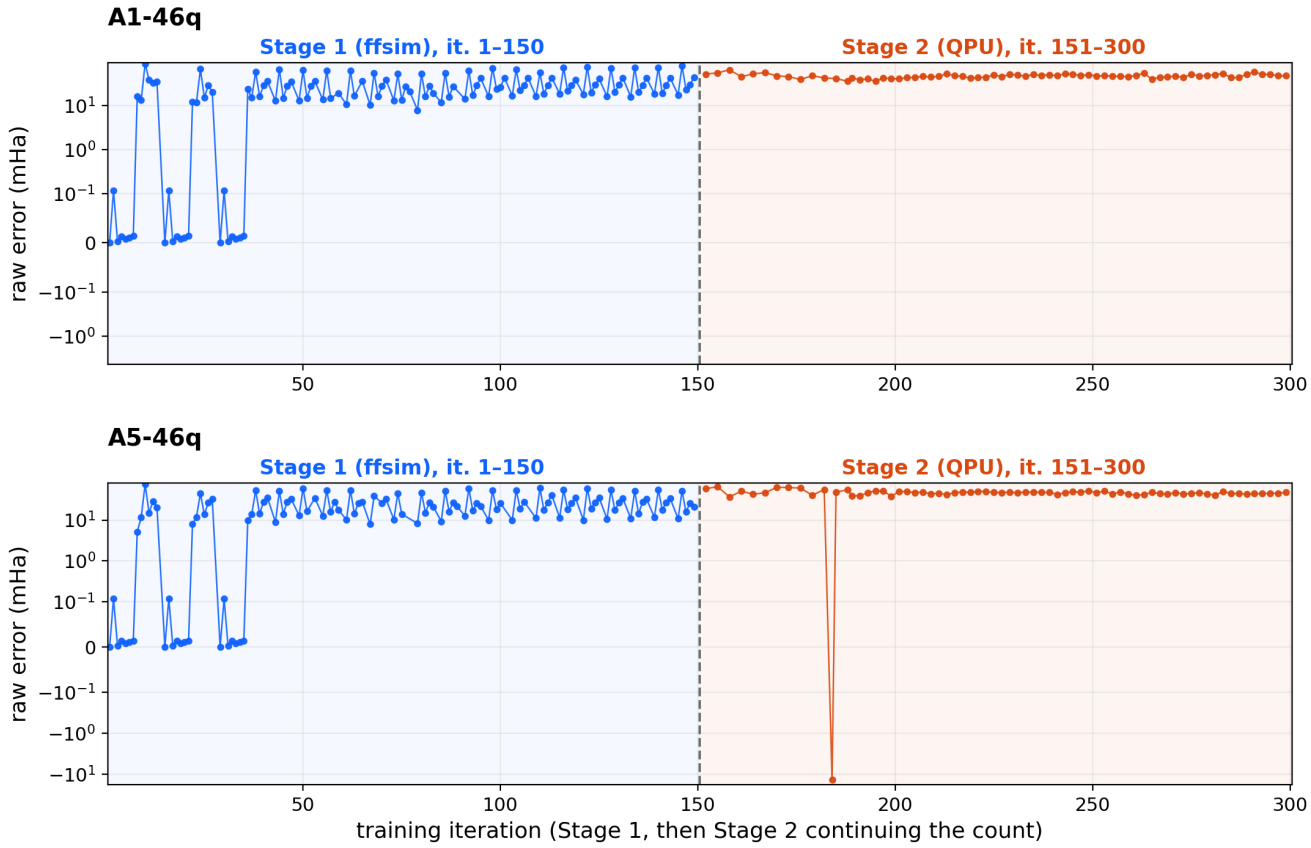

Figure A9: Raw SQD energy error (mHa; symmetric logarithmic axis) of the 46-qubit lineages on admitted iterations. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line. In Stage 1 the error cycles with the size curriculum (small warm-up molecules near zero, larger ones tens of mHa). Negative values are iterations whose recovered energy lies below the CCSD(T) reference, which is not variational.

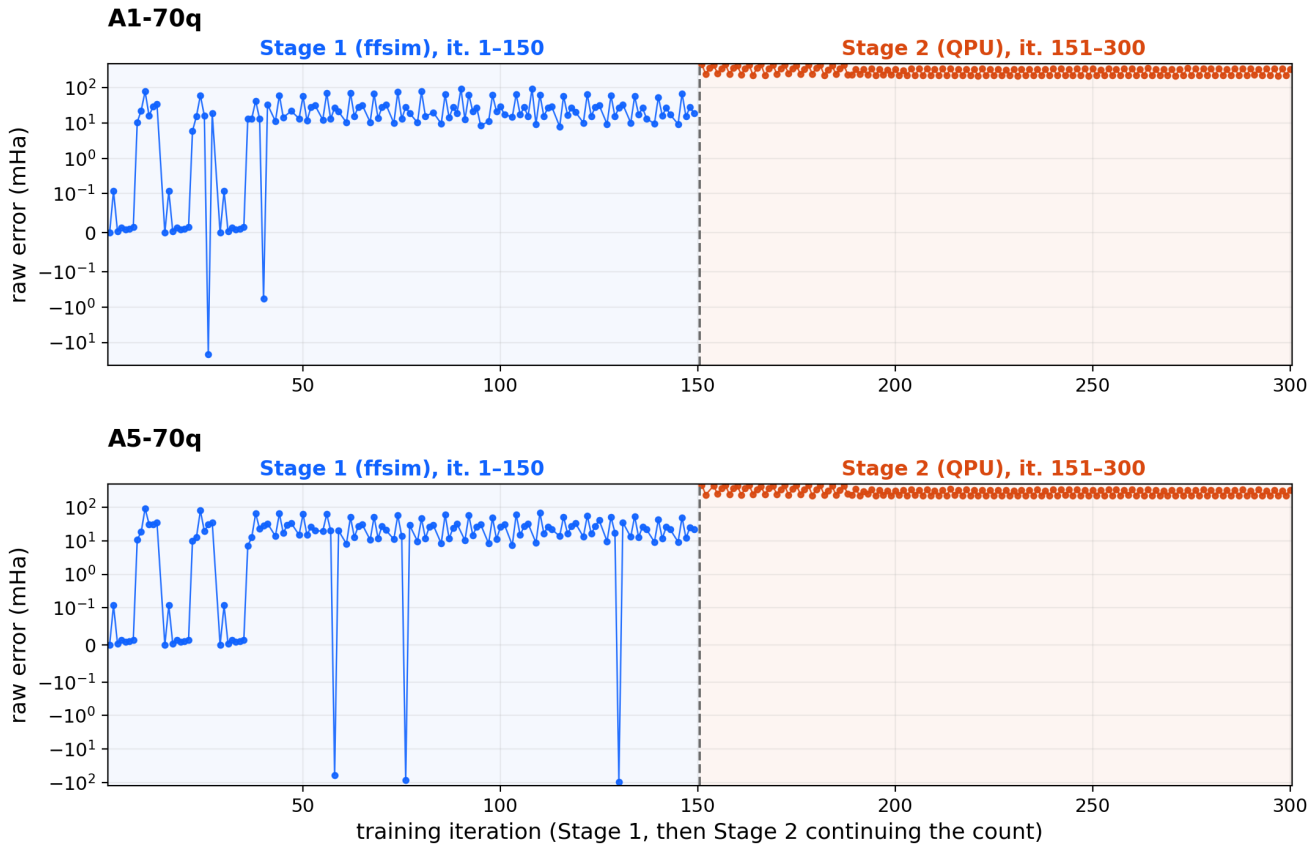

Figure A10: As Figure A9 for the 70-qubit lineages. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line.

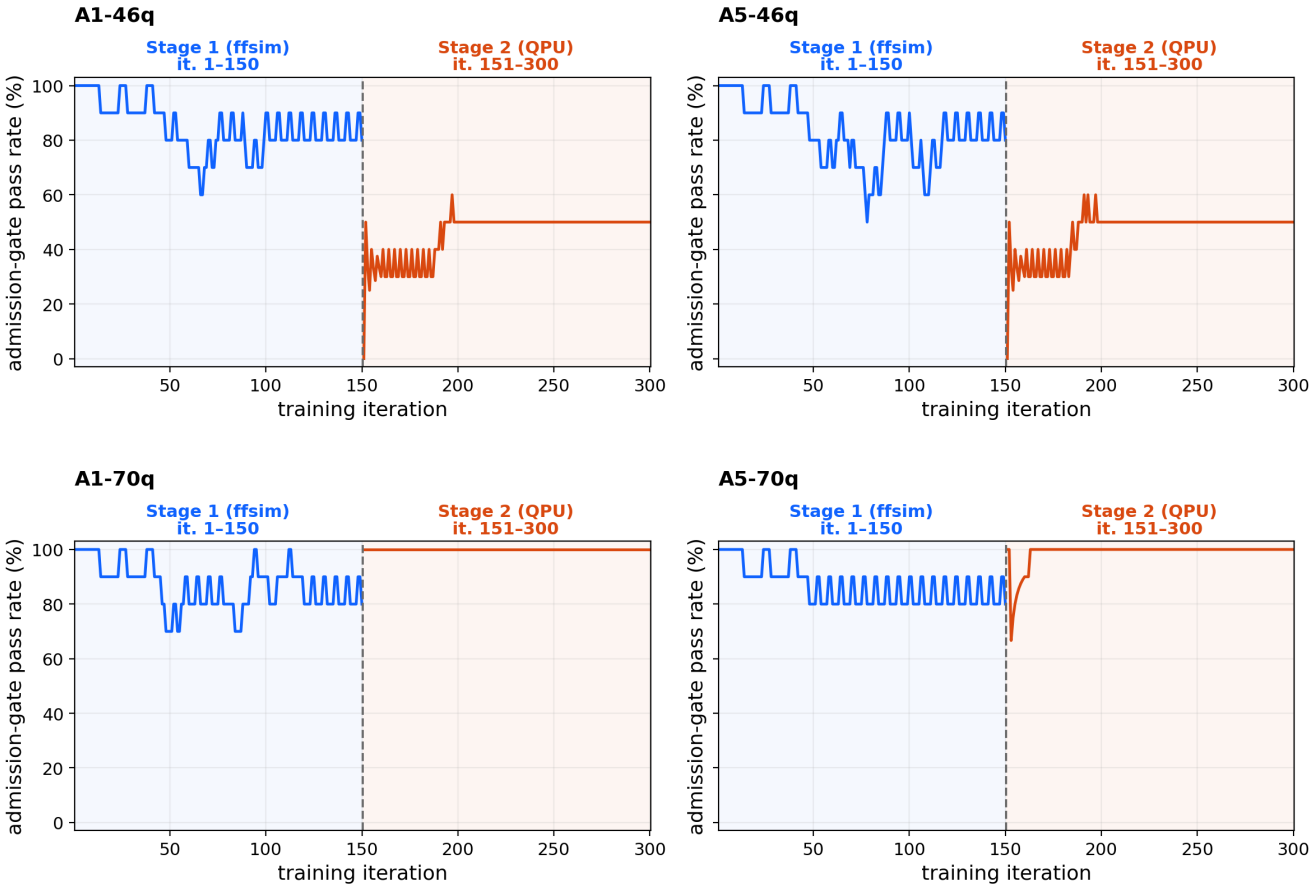

Figure A11: Rolling admission-gate pass rate (window of 10 iterations) of the four lineages. Stage 1 (ffsim simulator, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty) right of the dashed line; the iteration count continues across the stages.

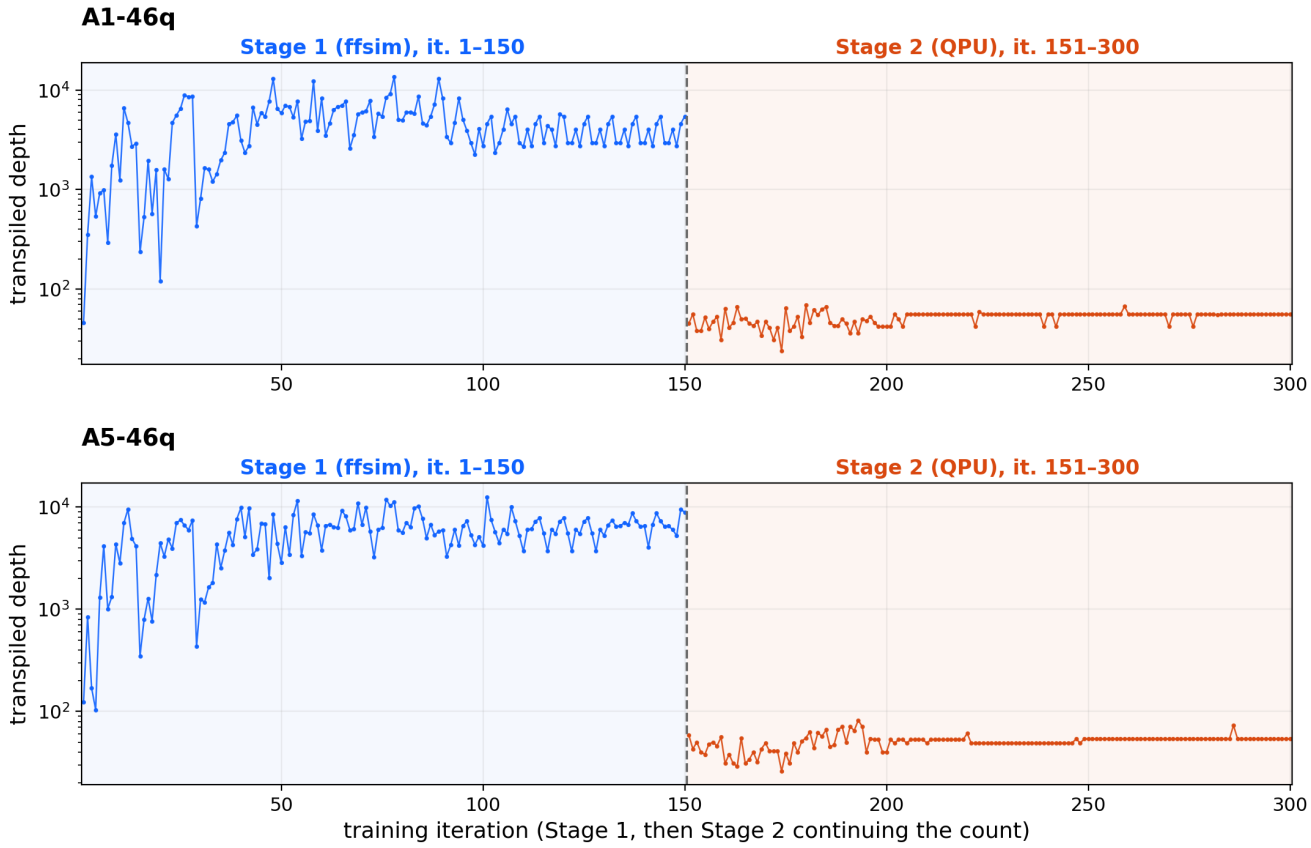

Figure A12: Transpiled depth of the best circuit (lowest energy) of the 20 sampled in every iteration, admitted or not (logarithmic axis; transpiled at optimisation level 1 for the `ibm_heron_r2` coupling map). 46-qubit lineages. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line.

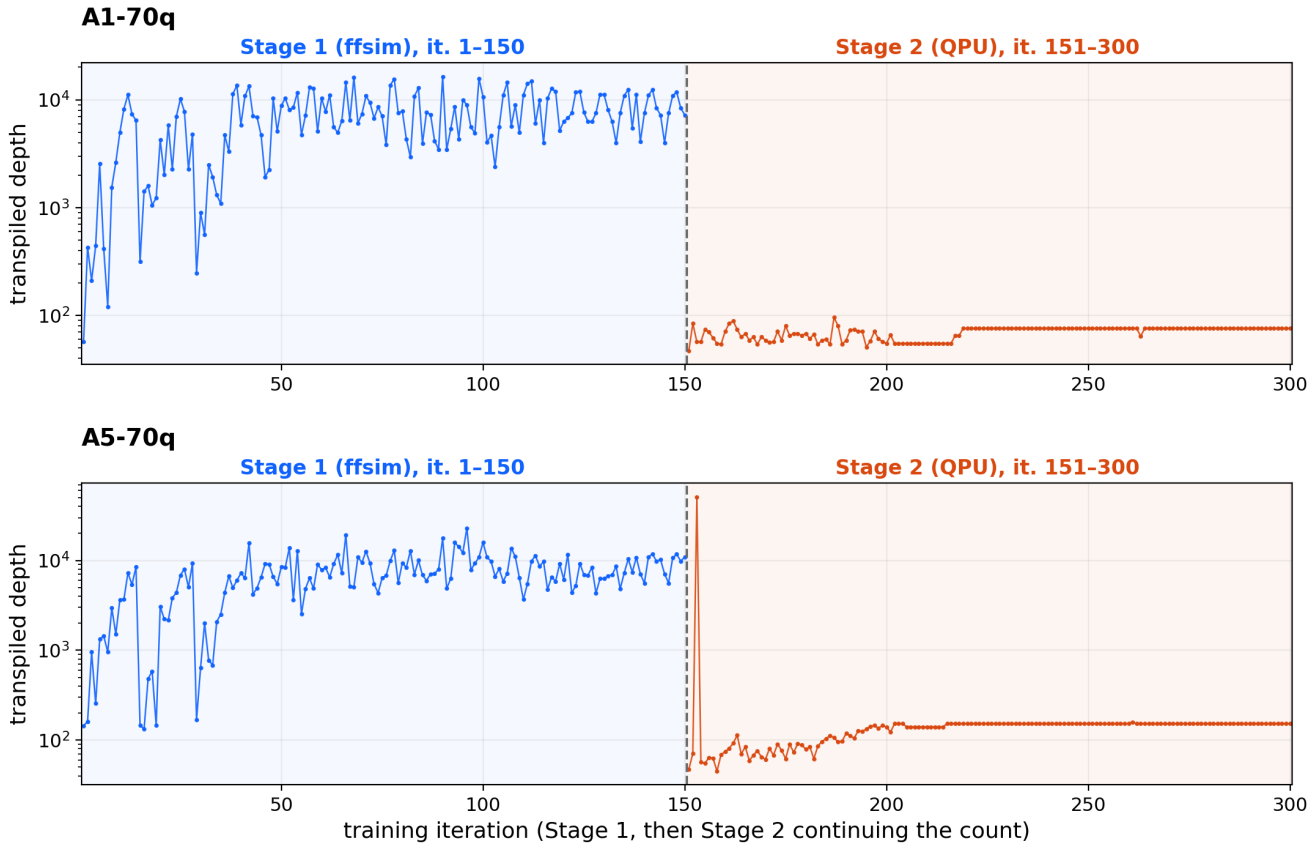

Figure A13: Transpiled depth of the best circuit (lowest energy) of the 20 sampled in every iteration, admitted or not (logarithmic axis; transpiled at optimisation level 1 for the `ibm_heron_r2` coupling map). 70-qubit lineages. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line.

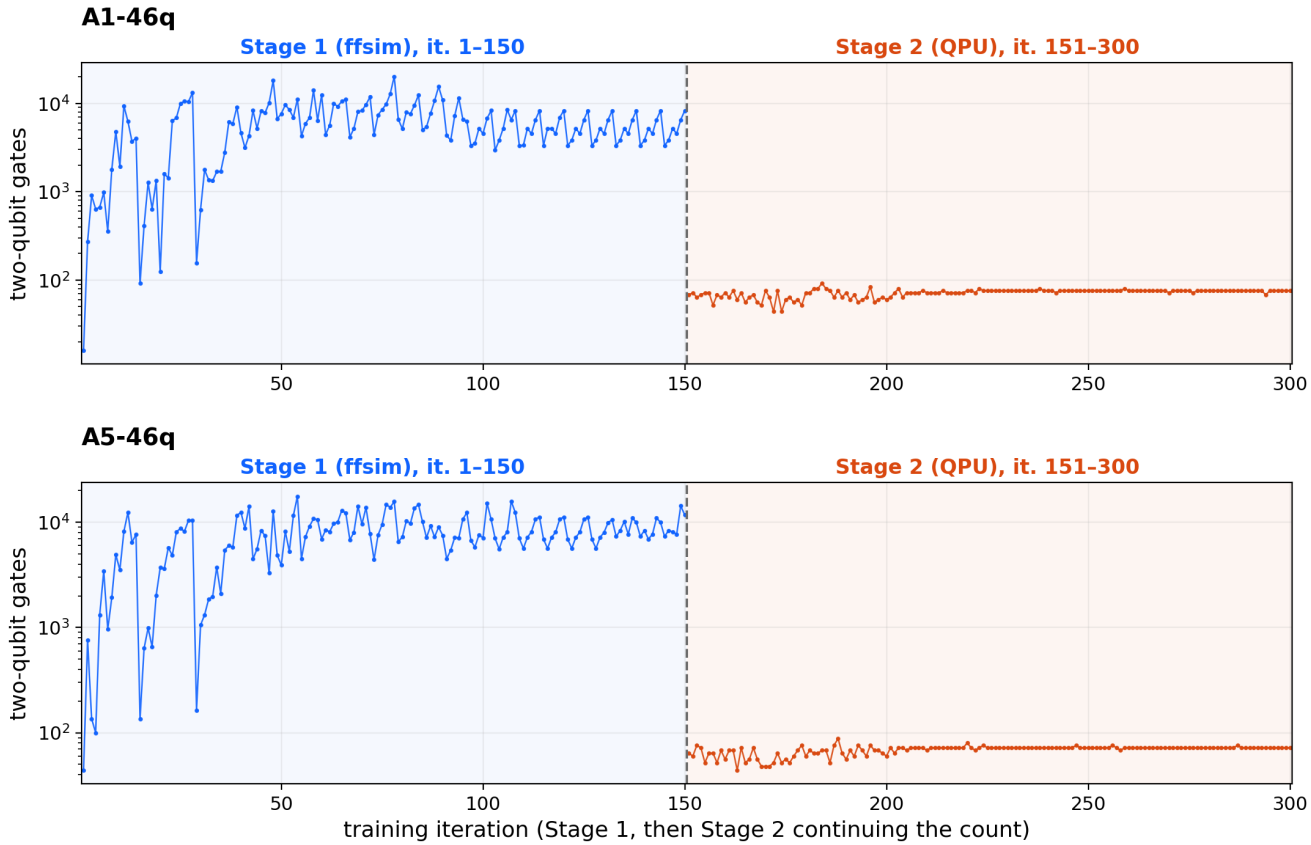

Figure A14: Two-qubit-gate count of the best circuit (lowest energy) of the 20 sampled in every iteration, admitted or not (logarithmic axis). 46-qubit lineages. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line.

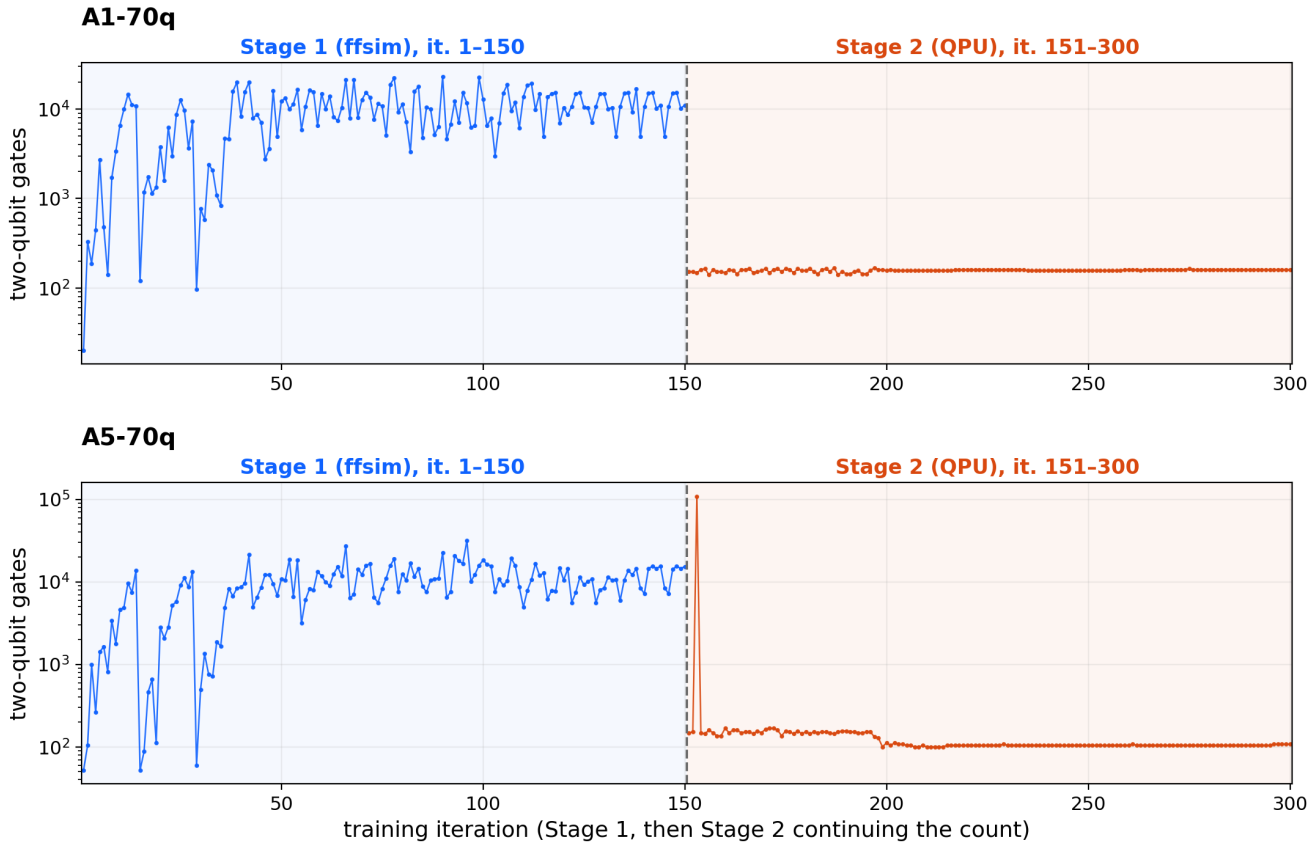

Figure A15: Two-qubit-gate count of the best circuit (lowest energy) of the 20 sampled in every iteration, admitted or not (logarithmic axis). 70-qubit lineages. Stage 1 (ffsim simulator, iterations 1–150, shared Stage 1) is left and Stage 2 (`ibm_kobe` QPU, HF penalty, iterations 151–300) right of the dashed line.

H Comparison of the two training configurations, per test case

Best and second circuit of every test case (molecule–lineage pair) in which both training configurations have both circuits recovered (43 of 44 test cases): recovered energy E (E_h), signed deviation ΔE from the DMRG reference (mHa), as-submitted depth of the best circuit, and the spin-resolved shot analysis of the best circuit. "valid" is the share of all shots with the correct electron numbers in each spin sector; "HF" is the share of the valid shots equal to the Hartree–Fock determinant, "HL" the share containing a HOMO→LUMO excitation (highest occupied to lowest unoccupied orbital), and α/β/both the shares that excite only the spin-up channel, only the spin-down channel, or both. Two repeated energies in the table (−631.077480 and −555.102002) are Hartree–Fock determinant energies computed from the integrals, not converged SQD or DMRG results.

Target	Lineage	Set	best E (E_h)	ΔE (mHa)	2nd E (E_h)	ΔE (mHa)	depth	valid (%)	HF (%)	HL (%)	α/β/both (%)
S 32q	A1-46q	HF penalty	-479.193266	0.12	-479.192672	0.71	50	2.7	58.1	0.8	24.8/10.9/6.3
		tuned reward	-479.193089	0.29	-479.193374	0.01	27	6.4	56.5	0.0	35.1/5.0/3.4
S 32q	A1-70q	HF penalty	-479.193164	0.22	-479.193270	0.11	41	21.3	12.6	0.0	0.6/82.4/4.4
		tuned reward	-479.193259	0.12	-479.193255	0.13	34	5.8	21.1	28.6	4.2/63.2/11.5
S 32q	A5-46q	HF penalty	-479.192727	0.65	-479.193081	0.30	31	40.9	94.6	1.4	3.3/2.0/0.1
		tuned reward	-479.193374	0.01	-479.193374	0.01	2770	0.0 (9)	0.0	0.0	0.0/0.0/100.0
S 32q	A5-70q	HF penalty	-479.192924	0.46	-479.193134	0.25	69	33.1	31.0	51.9	1.8/63.1/4.0
		tuned reward	-479.191604	1.78	-479.193374	0.01	92	3.0	17.0	4.4	5.3/54.8/22.9
h2o S 38q	A1-46q	HF penalty	-555.108538	0.29	-555.108684	0.15	35	9.0	79.9	0.0	7.3/11.5/1.3
		tuned reward	-555.108769	0.06	-555.108664	0.17	40	4.3	40.7	3.1	6.6/46.2/6.5
h2o S 38q	A1-70q	HF penalty	-555.108766	0.07	-555.108751	0.08	57	14.7	76.7	0.5	20.6/2.1/0.6
		tuned reward	-555.108816	0.02	-555.108680	0.15	68	2.9	35.5	0.8	6.0/48.1/10.4
h2o S 38q	A5-46q	HF penalty	-555.108713	0.12	-555.108826	0.01	34	40.4	93.1	0.0	0.6/6.2/0.1
		tuned reward	-555.108740	0.09	-555.108540	0.29	64	2.8	64.8	0.4	8.7/23.3/3.1
h2o S 38q	A5-70q	HF penalty	-555.108740	0.09	-555.108764	0.07	110	12.9	61.5	6.6	21.7/12.4/4.4
		tuned reward	-555.108752	0.08	-555.108824	0.01	45	22.4	3.5	0.1	0.0/95.0/1.4
M 38q	A1-46q	HF penalty	-479.292683	9.60	-479.296591	5.69	35	12.0	64.2	11.1	11.8/20.2/3.8
		tuned reward	-479.297136	5.14	-479.290408	11.87	40	8.4	48.9	8.5	23.2/18.2/9.7
M 38q	A1-70q	HF penalty	-479.298463	3.82	-479.295672	6.61	57	14.8	52.1	9.1	35.2/7.5/5.2
		tuned reward	-479.295479	6.80	-479.295776	6.50	68	12.5	2.5	32.8	8.8/20.3/68.4
M 38q	A5-46q	HF penalty	-479.285058	17.22	-479.282828	19.45	32	20.7	3.8	4.0	0.1/91.5/4.6
		tuned reward	-479.290816	11.46	-479.288843	13.43	49	21.7	7.8	0.4	0.3/89.5/2.4
M 38q	A5-70q	HF penalty	-479.288699	13.58	-479.152486	149.79	152	7.0	45.0	0.4	11.6/33.9/9.5
		tuned reward	-479.292726	9.55	-479.150765	151.51	56	4.0	25.0	0.4	32.6/19.0/23.4
2h2o S 44q	A1-46q	HF penalty	-631.083103	0.21	-631.083239	0.07	30	1.5	64.1	2.7	8.0/25.3/2.6
		tuned reward	-631.083257	0.05	-631.083235	0.08	59	5.7	25.3	1.0	31.8/19.6/23.2
2h2o S 44q	A1-70q	HF penalty	-631.083257	0.06	-631.083084	0.23	60	1.9	61.7	2.9	7.6/26.4/4.3
		tuned reward	-631.083269	0.04	-631.083269	0.04	298	0.0 (7)	0.0	0.0	0.0/0.0/100.0
2h2o S 44q	A5-46q	HF penalty	-631.083185	0.13	-631.083207	0.11	49	11.9	41.4	34.3	53.3/2.2/3.1
		tuned reward	-631.083122	0.19	-631.083109	0.20	33	14.1	45.6	0.7	38.3/8.2/7.8
2h2o S 44q	A5-70q	HF penalty	-631.083092	0.22	-631.083044	0.27	168	14.9	17.1	29.0	0.6/79.7/2.6
		tuned reward	-631.083237	0.07	-631.083255	0.06	66	2.5	75.7	0.4	3.9/19.4/1.0
L 46q	A1-46q	HF penalty	-479.299926	14.29	-479.121483	192.73	37	11.1	53.5	0.8	1.6/43.7/1.3
		tuned reward	-479.301231	12.98	-479.115176	199.04	36	19.0	14.7	0.2	0.7/80.7/4.0
L 46q	A1-70q	HF penalty	-479.306409	7.81	-479.116356	197.86	69	9.0	59.2	0.5	3.5/35.3/2.0
		tuned reward	-479.300283	13.93	-479.116713	197.50	69	8.1	29.1	0.0	64.6/2.0/4.3
L 46q	A5-46q	HF penalty	-479.301748	12.47	-479.306572	7.64	55	2.4	74.1	1.8	2.2/23.1/0.6
		tuned reward	-479.301200	13.02	-479.120938	193.28	41	13.7	13.7	22.6	32.0/16.4/37.8
L 46q	A5-70q	HF penalty	-479.308601	5.62	-479.297326	16.89	144	4.3	38.6	0.8	44.8/8.0/8.7
		tuned reward	-479.115938	198.28	-479.115299	198.92	17042	0.0 (47)	0.0	0.0	0.0/0.0/100.0
h2o M 46q	A1-46q	HF penalty	-555.123748	1.41	-555.102442	22.72	37	10.3	52.9	0.7	3.0/41.8/2.4
		tuned reward	-555.123722	1.44	-555.102808	22.35	36	22.6	8.5	0.1	0.4/87.6/3.5
h2o M 46q	A1-70q	HF penalty	-555.124450	0.71	-555.102207	22.95	69	10.2	49.3	1.6	3.2/44.3/3.1
		tuned reward	-555.124065	1.09	-555.102254	22.90	69	7.7	29.0	0.4	64.1/1.9/5.0
h2o M 46q	A5-46q	HF penalty	-555.123672	1.49	-555.123151	2.01	58	23.9	23.4	1.5	0.9/72.8/2.8
		tuned reward	-555.122758	2.40	-555.122769	2.39	45	9.5	66.1	0.0	14.6/13.1/6.1
h2o M 46q	A5-70q	HF penalty	-555.123764	1.39	-555.123482	1.68	108	5.3	57.9	2.2	13.8/21.8/6.5
		tuned reward	-555.124431	0.73	-555.124133	1.02	44	1.8	22.2	0.9	10.9/47.2/19.7
LL 56q	A1-46q	HF penalty	-479.300183	14.87	-479.299292	15.76	52	3.8	41.6	1.0	8.7/41.7/7.9
		tuned reward	-479.305726	9.33	-479.109979	205.08	52	2.2	26.5	0.7	29.2/20.5/23.8
LL 56q	A1-70q	HF penalty	-479.304625	10.43	-479.112860	202.20	62	1.5	38.0	0.8	2.2/56.3/3.6
		tuned reward	-479.306343	8.71	-479.110787	204.27	59	0.4	19.2	0.7	60.5/5.0/15.3

Target	Lineage	Set	best E (E_h)	ΔE (mHa)	2nd E (E_h)	ΔE (mHa)	depth	valid (%)	HF (%)	HL (%)	α/β/both (%)
LL 56q	A5-70q	HF penalty	-479.301896	13.16	-479.111752	203.30	79	1.8	24.0	2.3	30.5/16.7/28.9
		tuned reward	-479.302562	12.50	-479.297283	17.77	63	14.9	56.5	3.7	14.4/23.2/5.8
h2o L 56q	A1-46q	HF penalty	-555.132208	2.85	-555.132421	2.64	47	5.0	15.0	1.0	28.8/18.6/37.6
		tuned reward	-555.132762	2.30	-555.102002	33.06	52	2.4	26.3	0.0	29.3/21.4/23.0
h2o L 56q	A1-70q	HF penalty	-555.133193	1.86	-555.132552	2.51	65	7.3	47.9	2.5	39.6/6.7/5.8
		tuned reward	-555.133432	1.63	-555.102199	32.86	59	0.9	16.3	1.1	67.6/2.9/13.2
h2o L 56q	A5-46q	HF penalty	-555.131417	3.64	-555.132660	2.40	53	8.4	72.3	4.9	12.4/13.1/2.2
		tuned reward	-555.132718	2.34	-555.132834	2.22	49	1.7	43.5	1.1	5.5/44.8/6.2
h2o L 56q	A5-70q	HF penalty	-555.132405	2.65	-555.132408	2.65	44	14.2	6.7	44.7	14.6/24.2/54.5
		tuned reward	-555.132949	2.11	-555.133484	1.57	60	2.1	7.8	14.3	9.6/35.7/46.9
2h2o M 58q	A1-46q	HF penalty	-631.112084	6.53	-631.077480	41.14	47	2.3	18.4	4.8	39.3/11.9/30.4
		tuned reward	-631.077480	41.14	-631.077480	41.14	10985	0.1	0.0	0.0	0.0/0.0/100.0
2h2o M 58q	A1-70q	HF penalty	-631.114317	4.30	-631.113058	5.56	65	5.6	50.9	2.5	38.6/5.9/4.7
		tuned reward	-631.113121	5.49	-631.077480	41.14	49	2.1	13.1	1.0	62.3/4.0/20.5
2h2o M 58q	A5-46q	HF penalty	-631.114080	4.54	-631.113620	5.00	45	7.1	7.4	14.4	6.1/45.7/40.7
		tuned reward	-631.077480	41.14	-631.077480	41.14	368	0.6	0.0	0.0	0.0/0.0/100.0
2h2o M 58q	A5-70q	HF penalty	-631.113705	4.91	-631.111954	6.66	62	19.8	59.2	2.2	31.0/6.5/3.4
		tuned reward	-631.113566	5.05	-631.077480	41.14	56	6.6	22.1	0.0	54.1/5.9/17.9
h2o LL 66q	A1-46q	HF penalty	-555.135328	2.07	-555.133503	3.89	60	12.1	80.1	0.0	14.6/4.4/0.8
		tuned reward	-555.135415	1.98	-555.133657	3.74	34	8.1	56.1	6.5	31.2/8.1/4.6
h2o LL 66q	A1-70q	HF penalty	-555.134846	2.55	-555.102002	35.39	51	2.2	8.2	3.2	1.5/70.4/19.8
		tuned reward	-555.134843	2.55	-555.134225	3.17	58	4.4	37.0	17.4	16.1/32.6/14.3
h2o LL 66q	A5-46q	HF penalty	-555.133877	3.52	-555.134389	3.01	44	9.3	46.5	0.1	8.5/37.9/7.0
		tuned reward	-555.135580	1.82	-555.133728	3.67	49	12.1	80.5	4.5	8.1/10.4/1.0
h2o LL 66q	A5-70q	HF penalty	-555.135231	2.17	-555.134346	3.05	57	7.5	79.9	1.5	11.2/6.8/2.1
		tuned reward	-555.134772	2.62	-555.132932	4.46	69	2.5	8.1	9.7	2.3/66.7/22.9
2h2o L 70q	A1-46q	HF penalty	-631.123614	8.62	-631.123272	8.97	48	3.7	48.4	0.0	38.7/7.2/5.6
		tuned reward	-631.124462	7.78	-631.077480	54.76	48	6.0	59.5	0.0	23.8/11.4/5.4
2h2o L 70q	A1-70q	HF penalty	-631.124552	7.69	-631.077480	54.76	60	0.4	14.5	3.1	19.1/25.2/41.2
		tuned reward	-631.122201	10.04	-631.077480	54.76	63	0.4	31.7	0.7	22.8/23.5/22.1
2h2o L 70q	A5-46q	HF penalty	-631.124857	7.38	-631.123201	9.04	44	5.7	63.0	0.0	10.3/21.7/5.0
		tuned reward	-631.126224	6.01	-631.124909	7.33	40	14.1	36.2	10.8	30.4/18.0/15.5
2h2o L 70q	A5-70q	HF penalty	-631.120256	11.98	-631.124322	7.92	134	4.9	0.5	49.4	6.1/6.2/87.1
		tuned reward	-631.124065	8.17	-631.124165	8.07	61	2.5	16.0	0.1	41.8/10.7/31.6

I Per-circuit divergence from the idealised LUCJ circuit

Fidelity, Hellinger distance, Jensen–Shannon divergence, d_{HF} and (where computed) the operator-norm lower bound of every circuit of the three small tiers, for the two training configurations; definitions and references are in Section 3.6. The tuned run of SnOH3-me-2h2o S with A5-70q is the repeated draw of Appendix N. Larger tiers cannot be computed exactly, for the memory reason given there.

Table A12: Per-circuit ideal divergence from the LUCJ circuit, HF-penalty set (Section 3.6). F, H, JS (bits) and d_{HF} as in Table 9. "op. norm*" is a lower-bound estimate of $\|U_g - U_{\mathrm{LUCJ}}\|$ (maximum 2) computed without phase minimisation, and only for the eight selected SnOH3-me S circuits of each set; 2.00 indicates a non-converged value (Lanczos stopped at the step cap). $F = 0.000$ means $F < 10^{-5}$, and H or JS values above 0.998 are rounded to 1.000. Depth is as-submitted. "sel.": best or second circuit.

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
SnOH3-me S (32q), A1-46q	1		2897	0.000	0.999	1.000	1.414	–
	2		46	0.954	0.152	0.024	0.216	–
	3		3861	0.000	1.000	1.000	1.414	–
	4		34	0.998	0.035	0.001	0.049	–
	5	best	50	0.993	0.059	0.004	0.084	1.71
	6		3783	0.000	1.000	1.000	1.414	–
	7		6568	0.000	1.000	1.000	1.414	–
	8	2nd	30	0.998	0.035	0.001	0.049	1.33
	9		196	0.000	1.000	1.000	1.414	–
	10		2595	0.954	0.152	0.024	0.215	–

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
SnOH3-me S (32q), A1-70q	1		6798	0.000	1.000	1.000	1.414	–
	2		2698	0.923	0.197	0.040	0.279	–
	3		3356	0.000	1.000	1.000	1.414	–
	4	best	41	0.998	0.035	0.001	0.049	1.79
	5	2nd	57	0.993	0.059	0.004	0.084	1.94
	6		8368	0.000	1.000	1.000	1.414	–
	7		6606	0.000	1.000	1.000	1.414	–
	8		2760	0.999	0.019	0.000	0.027	–
	9		3460	0.000	0.999	0.999	1.413	–
	10		3214	0.000	0.999	1.000	1.413	–
SnOH3-me S (32q), A5-46q	1		395	0.000	1.000	1.000	1.414	–
	2		218	0.000	1.000	1.000	1.414	–
	3	best	31	0.998	0.035	0.001	0.049	1.69
	4		37	0.998	0.035	0.001	0.049	–
	5		222	0.000	1.000	1.000	1.414	–
	6		48	0.998	0.035	0.001	0.049	–
	7		37	0.998	0.035	0.001	0.049	–
	8		41	0.998	0.035	0.001	0.049	–
	9	2nd	30	0.998	0.035	0.001	0.049	1.96
	10		69	0.998	0.035	0.001	0.049	–
SnOH3-me S (32q), A5-70q	1		7403	0.000	1.000	1.000	1.414	–
	2		6470	0.994	0.053	0.003	0.075	–
	3		83	0.998	0.035	0.001	0.049	–
	4		230	0.000	0.999	1.000	1.414	–
	5		224	0.000	1.000	1.000	1.414	–
	6	2nd	64	0.993	0.059	0.004	0.084	1.72
	7		2587	0.998	0.035	0.001	0.049	–
	8	best	69	0.954	0.152	0.024	0.216	1.58
	9		272	0.000	1.000	1.000	1.414	–
	10		59	0.978	0.104	0.011	0.148	–
SnOH3-me-h2o S (38q), A1-46q	1		8003	0.000	1.000	1.000	1.414	–
	2	best	35	1.000	0.015	0.000	0.022	–
	3		455	0.000	1.000	1.000	1.414	–
	4	2nd	52	0.999	0.022	0.001	0.031	–
	5		41	0.999	0.022	0.001	0.031	–
	6		251	0.000	1.000	1.000	1.414	–
	7		7236	0.000	1.000	1.000	1.414	–
	8		30	1.000	0.015	0.000	0.022	–
	9		234	0.000	1.000	1.000	1.414	–
	10		4016	1.000	0.006	0.000	0.008	–
SnOH3-me-h2o S (38q), A1-70q	1		5630	0.000	1.000	1.000	1.414	–
	2	best	57	1.000	0.015	0.000	0.022	–
	3		485	0.000	1.000	1.000	1.414	–
	4	2nd	51	0.999	0.022	0.001	0.031	–
	5		72	0.991	0.069	0.005	0.098	–
	6		5581	0.000	1.000	1.000	1.414	–
	7		2835	0.000	1.000	1.000	1.414	–
	8		2805	1.000	0.006	0.000	0.008	–
	9		473	0.000	1.000	1.000	1.414	–
	10		5322	0.000	1.000	1.000	1.414	–
SnOH3-me-h2o S (38q), A5-46q	1		11431	0.995	0.052	0.003	0.074	–
	2		3284	1.000	0.015	0.000	0.022	–
	3		3113	0.979	0.103	0.011	0.145	–

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
	4	best	34	1.000	0.015	0.000	0.022	–
	5		3389	1.000	0.015	0.000	0.022	–
	6		3808	1.000	0.015	0.000	0.022	–
	7		8494	0.000	1.000	1.000	1.414	–
	8		11595	1.000	0.015	0.000	0.022	–
	9		4075	0.000	1.000	1.000	1.414	–
	10	2nd	3867	1.000	0.006	0.000	0.008	–
SnOH3-me-h2o S (38q), A5-70q	1		7577	1.000	0.015	0.000	0.022	–
	2		145	1.000	0.015	0.000	0.022	–
	3		168	1.000	0.015	0.000	0.022	–
	4		130	1.000	0.015	0.000	0.022	–
	5		7292	1.000	0.015	0.000	0.022	–
	6		6542	1.000	0.015	0.000	0.022	–
	7		4397	0.000	1.000	1.000	1.414	–
	8	best	110	1.000	0.015	0.000	0.022	–
	9	2nd	161	1.000	0.015	0.000	0.022	–
	10		3802	0.994	0.052	0.003	0.074	–
SnOH3-me-2h2o S (44q), A1-46q	1		273	0.000	1.000	1.000	1.414	–
	2	2nd	59	0.969	0.126	0.016	0.178	–
	3		540	0.000	1.000	1.000	1.414	–
	4		4704	1.000	0.013	0.000	0.018	–
	5	best	30	0.982	0.095	0.009	0.134	–
	6		285	0.000	1.000	1.000	1.414	–
	7		6505	0.000	1.000	1.000	1.414	–
	8		5009	1.000	0.013	0.000	0.018	–
	9		284	0.000	1.000	1.000	1.414	–
	10		4913	1.000	0.005	0.000	0.007	–
SnOH3-me-2h2o S (44q), A1-70q	1		6251	0.000	1.000	1.000	1.414	–
	2	2nd	59	0.998	0.034	0.001	0.048	–
	3		569	0.000	1.000	1.000	1.414	–
	4		56	1.000	0.013	0.000	0.018	–
	5	best	60	0.995	0.048	0.002	0.069	–
	6		7112	0.000	1.000	1.000	1.414	–
	7		4460	0.000	1.000	1.000	1.414	–
	8		9153	1.000	0.005	0.000	0.016	–
	9		6979	0.000	1.000	1.000	1.414	–
	10		6506	0.000	1.000	1.000	1.414	–
SnOH3-me-2h2o S (44q), A5-46q	1		44	0.999	0.020	0.000	0.028	–
	2		11122	0.000	1.000	1.000	1.414	–
	3	best	49	1.000	0.013	0.000	0.018	–
	4		37	1.000	0.013	0.000	0.018	–
	5	2nd	62	1.000	0.013	0.000	0.018	–
	6		5757	0.000	1.000	1.000	1.414	–
	7		32	1.000	0.013	0.000	0.018	–
	8		5062	1.000	0.013	0.000	0.018	–
	9		4180	1.000	0.005	0.000	0.007	–
	10		4886	1.000	0.013	0.000	0.018	–
SnOH3-me-2h2o S (44q), A5-70q	1	2nd	142	0.992	0.064	0.004	0.090	–
	2		10847	0.000	1.000	1.000	1.414	–
	3	best	168	0.992	0.064	0.004	0.090	–
	4		410	0.000	1.000	1.000	1.414	–
	5		5564	0.000	1.000	1.000	1.414	–
	6		6509	0.000	1.000	1.000	1.414	–

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
	7		4229	1.000	0.013	0.000	0.018	–
	8		4403	0.995	0.048	0.002	0.069	–
	9		5665	0.000	1.000	1.000	1.414	–
	10		4271	0.999	0.020	0.000	0.028	–

Table A13: Per-circuit ideal divergence from the LUCJ circuit, tuned reward set (Section 3.6). F, H, JS (bits) and d_{HF} as in Table 9. "op. norm*" is a lower-bound estimate of $\|U_g - U_{\mathrm{LUCJ}}\|$ (maximum 2) computed without phase minimisation, and only for the best and second circuit of SnOH3-me S; 2.00 indicates a non-converged value (Lanczos stopped at the step cap). $F = 0.000$ means $F < 10^{-5}$, and H or JS values above 0.998 are rounded to 1.000. Depth is as-submitted. "sel.": best or second circuit.

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
SnOH3-me S (32q), A1-46q	1		5561	0.000	0.999	1.000	1.414	–
	2		2789	0.954	0.152	0.024	0.216	–
	3		3873	0.000	1.000	1.000	1.414	–
	4	best	27	0.998	0.035	0.001	0.049	1.36
	5		219	0.000	0.999	1.000	1.414	–
	6	2nd	4954	0.000	1.000	1.000	1.414	2.00
	7		7665	0.000	0.999	1.000	1.413	–
	8		2564	0.999	0.017	0.000	0.024	–
	9		196	0.000	1.000	1.000	1.414	–
	10		6806	0.999	0.017	0.000	0.024	–
SnOH3-me S (32q), A1-70q	1		6932	0.000	1.000	1.000	1.414	–
	2	2nd	63	0.921	0.201	0.042	0.284	1.61
	3		8003	0.000	1.000	1.000	1.414	–
	4	best	34	0.954	0.152	0.024	0.216	1.79
	5		237	0.000	1.000	1.000	1.414	–
	6		5057	0.000	1.000	1.000	1.414	–
	7		7134	0.000	1.000	1.000	1.414	–
	8		2283	0.000	0.999	1.000	1.414	–
	9		7522	0.000	0.999	0.999	1.413	–
	10		2877	0.000	0.999	1.000	1.413	–
SnOH3-me S (32q), A5-46q	1		763	0.000	1.000	1.000	1.414	–
	2	2nd	5551	0.000	1.000	1.000	1.414	2.00
	3		7077	0.000	1.000	1.000	1.414	–
	4		411	0.000	1.000	1.000	1.414	–
	5		418	0.000	1.000	1.000	1.414	–
	6		2888	0.000	0.999	1.000	1.414	–
	7		2713	0.993	0.059	0.004	0.084	–
	8		2757	0.000	1.000	1.000	1.414	–
	9		235	0.000	1.000	1.000	1.414	–
	10	best	2770	0.998	0.035	0.001	0.049	2.00
SnOH3-me S (32q), A5-70q	1		15128	0.000	1.000	1.000	1.414	–
	2		6375	0.776	0.345	0.124	0.488	–
	3		3124	0.000	1.000	1.000	1.414	–
	4	best	92	0.776	0.345	0.124	0.488	1.99
	5		6838	0.000	1.000	1.000	1.414	–
	6		2771	0.993	0.059	0.004	0.084	–
	7		3516	0.000	1.000	1.000	1.414	–
	8		3413	0.000	1.000	1.000	1.414	–
	9		3214	0.000	0.999	1.000	1.414	–
	10	2nd	3041	0.000	1.000	1.000	1.414	2.00

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
SnOH3-me-h2o S (38q), A1-46q	1		3627	0.000	1.000	1.000	1.414	–
	2	best	40	1.000	0.015	0.000	0.022	–
	3		453	0.000	1.000	1.000	1.414	–
	4	2nd	31	0.999	0.022	0.001	0.031	–
	5		253	0.000	1.000	1.000	1.414	–
	6		470	0.000	1.000	1.000	1.414	–
	7		8879	0.000	1.000	1.000	1.414	–
	8		3381	1.000	0.006	0.000	0.008	–
	9		233	0.000	1.000	1.000	1.414	–
	10		2933	1.000	0.006	0.000	0.008	–
SnOH3-me-h2o S (38q), A1-70q	1		5592	0.000	1.000	1.000	1.414	–
	2	2nd	51	0.999	0.022	0.001	0.031	–
	3		4446	0.000	1.000	1.000	1.414	–
	4	best	68	0.999	0.022	0.001	0.031	–
	5		292	0.000	1.000	1.000	1.414	–
	6		6360	0.000	1.000	1.000	1.414	–
	7		8075	0.000	1.000	1.000	1.414	–
	8		9337	0.000	1.000	1.000	1.414	–
	9		5594	0.000	1.000	1.000	1.414	–
	10		3326	0.000	1.000	1.000	1.414	–
SnOH3-me-h2o S (38q), A5-46q	1		7517	0.000	1.000	1.000	1.414	–
	2		35	1.000	0.015	0.000	0.022	–
	3		3428	1.000	0.015	0.000	0.022	–
	4		32	1.000	0.015	0.000	0.022	–
	5		4205	0.000	1.000	1.000	1.414	–
	6		267	0.000	1.000	1.000	1.414	–
	7		670	0.000	1.000	1.000	1.414	–
	8	2nd	53	1.000	0.015	0.000	0.022	–
	9		253	0.000	1.000	1.000	1.414	–
	10	best	64	1.000	0.015	0.000	0.022	–
SnOH3-me-h2o S (38q), A5-70q	1		7960	1.000	0.015	0.000	0.022	–
	2		3786	0.995	0.052	0.003	0.074	–
	3		8351	0.000	1.000	1.000	1.414	–
	4	best	45	1.000	0.015	0.000	0.022	–
	5		4040	0.994	0.052	0.003	0.074	–
	6		3726	0.994	0.052	0.003	0.074	–
	7		492	0.000	1.000	1.000	1.414	–
	8		7283	0.995	0.052	0.003	0.074	–
	9	2nd	77	0.979	0.103	0.011	0.145	–
	10		3189	0.994	0.052	0.003	0.074	–
SnOH3-me-2h2o S (44q), A1-46q	1		283	0.000	1.000	1.000	1.414	–
	2	best	59	0.987	0.079	0.006	0.112	–
	3		530	0.000	1.000	1.000	1.414	–
	4	2nd	30	0.999	0.020	0.000	0.028	–
	5		279	0.000	1.000	1.000	1.414	–
	6		527	0.000	1.000	1.000	1.414	–
	7		5410	0.000	1.000	1.000	1.414	–
	8		11178	1.000	0.005	0.000	0.016	–
	9		277	0.000	1.000	1.000	1.414	–
	10		4300	1.000	0.005	0.000	0.007	–
SnOH3-me-2h2o S (44q), A1-70q	1		6947	0.000	1.000	1.000	1.414	–
	2		5317	0.988	0.079	0.006	0.112	–
	3		6533	0.000	1.000	1.000	1.414	–

Run	#	sel.	depth	F	H	JS	d_{HF}	op. norm*
	4		4338	1.000	0.013	0.000	0.018	–
	5	best	298	0.000	1.000	1.000	1.414	–
	6	2nd	806	0.000	1.000	1.000	1.414	–
	7		5296	0.000	1.000	1.000	1.414	–
	8		9736	0.000	1.000	1.000	1.414	–
	9		6734	0.000	1.000	1.000	1.414	–
	10		6418	0.000	1.000	1.000	1.414	–
SnOH3-me-2h2o S (44q), A5-46q	1		4696	0.995	0.048	0.002	0.069	–
	2		11682	0.000	1.000	1.000	1.414	–
	3	2nd	45	1.000	0.013	0.000	0.018	–
	4		66	1.000	0.013	0.000	0.018	–
	5	best	33	1.000	0.013	0.000	0.018	–
	6		11243	0.000	1.000	1.000	1.414	–
	7		4257	1.000	0.013	0.000	0.018	–
	8		4855	1.000	0.013	0.000	0.018	–
	9		6320	0.995	0.047	0.002	0.067	–
	10		288	0.000	1.000	1.000	1.414	–
SnOH3-me-2h2o S (44q), A5-70q	1		570	0.000	1.000	1.000	1.414	–
	2		9938	0.000	1.000	1.000	1.414	–
	3		29113	0.000	1.000	1.000	1.414	–
	4		294	0.000	1.000	1.000	1.414	–
	5	best	66	0.999	0.020	0.000	0.028	–
	6		329	0.000	1.000	1.000	1.414	–
	7		21016	0.996	0.047	0.002	0.067	–
	8	2nd	58	0.995	0.048	0.002	0.069	–
	9		11617	0.000	1.000	1.000	1.414	–
	10		5285	1.000	0.013	0.000	0.018	–

J Zero-shot inference: full detail

Table A14: Which (target, policy) runs have fresh-shot results. "done": sampled, ranked and recovered; "not run": not part of the design (untrained controls were run for a subset of targets); "excluded†": the batch of ten circuits had a single circuit with valid shots and a second draw gave a recovered error of about 200 mHa, so the run is not part of the statistics (Appendix N). Columns are the lineage (A1/A5, trained at 46 or 70 qubits).

	Trained				Untrained control			
Target	A1-46q	A1-70q	A5-46q	A5-70q	A1-46q	A1-70q	A5-46q	A5-70q
SnOH3-me S (32q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me-h2o S (38q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me M (38q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me-2h2o S (44q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me L (46q)	done	done	done	done	done	done	done	done
SnOH3-me-h2o M (46q)	done	done	done	done	done	done	not run	done
SnOH3-me LL (56q)	done	done	excluded†	done	not run	done	not run	not run
SnOH3-me-h2o L (56q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me-2h2o M (58q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me-h2o LL (66q)	done	done	done	done	not run	not run	not run	not run
SnOH3-me-2h2o L (70q)	done	done	done	done	not run	not run	not run	not run

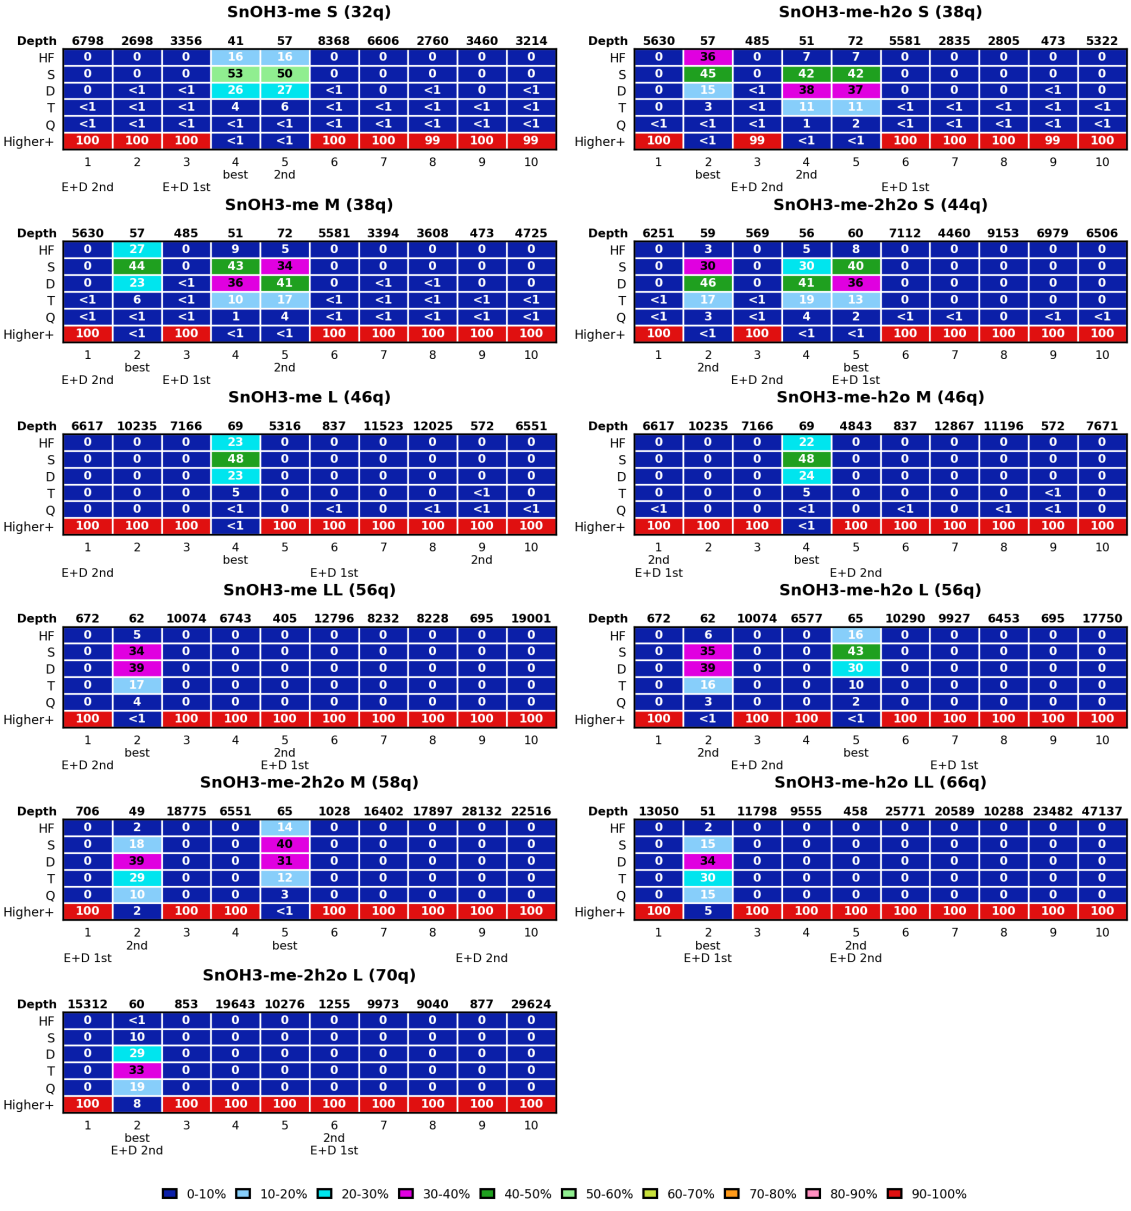

Figure A16: As Fig. 8 for the A1 policy trained at 70 qubits. The best circuit is the shallowest of the ten in six of the eleven panels and among the three shallowest in all of them, and it has at most 7.9% of its weight in "higher" (at most 0.6% in nine of the panels). The second circuit has little weight in "higher" in six panels; in SnOH3-me L, SnOH3-me-h2o M, SnOH3-me LL, SnOH3-me-h2o LL and SnOH3-me-2h2o L it has essentially all its weight there. The other circuits are one to three orders of magnitude deeper with essentially all weight in "higher".

Table A15: Zero-shot inference on the SnOH3-me L tier (46 qubits), per lineage and policy. Depth and 2q are the as-submitted (live) cost of the best/second circuit; "valid shots" is the number of the 100,000 shots with the correct electron numbers, i.e. the data behind each pick. "SI" (singles-injected) adds the single-excitation configurations to the best circuit's recovered determinant set and re-diagonalizes. DMRG reference −479.314216 E_h. Ranking, selection, shot counts and recovery settings are as in Table 6, except for the E+D picks of the untrained controls: these eight runs were sampled in the first campaign, and their E+D circuits were chosen with the energy rank of the earlier ranking, which used a cap of 2,000 determinants instead of 40,000, added to the rank by distinct bitstrings. Their best, second and merged results use the 40,000-determinant ranking.

Lineage/Policy	Depth (best/2nd)	2q (best/2nd)	Valid shots (best/2nd)	ΔE best/2nd/merged (mHa)	ΔE E+D 1/2 (mHa)	ΔE best-SI (mHa)
Trained policies						
A1-46q trained	37/4624	64/9866	11097/3	14.3/192.7/55.4	192.7/198.1	14.3
A1-70q trained	69/572	100/2104	9034/355	7.8/197.9/41.7	193.1/198.6	7.8
A5-46q trained	55/61	80/88	2445/5055	12.5/7.6/29.8	199.6/12.5	12.5
A5-70q trained	144/142	124/120	4324/15345	5.6/16.9/29.6	199.1/5.6	5.6
Untrained controls						
A1-46q untrained	14298/334	25505/1120	21/1859	200.2/196.2/188.7	196.2/200.0	39.7
A1-70q untrained	340/7267	1168/13540	605/62	193.7/194.7/192.2	196.9/199.9	38.2
A5-46q untrained	12650/24025	24649/44951	48/3	192.9/193.2/194.3	199.2/185.6	38.8
A5-70q untrained	10282/24119	17269/45351	8/26	188.2/200.9/198.9	198.1/200.9	39.2

Table A16: Measured breakdown of the first-campaign sample-job elapsed time for the A1-70q jobs (not repeated in Table A15 for the fresh-shot campaign), queried directly from the IBM Quantum Runtime job metadata (the job-metrics interface) via saved job IDs, not estimated from local logs. Both batch jobs started at the moment of submission (confirmed from the scheduler accounting records), so none of the batch elapsed time is scheduler queue wait; the principal breakdown is between IBM's device queue, device execution, and local overhead (Python startup, pipeline construction, the circuit-depth transpilation transpile loop over all 10 sampled circuits). "Billed QPU" is the quantum time IBM billed for the job (561 and 624 s here); "per circuit" is this billed time divided by the 10 circuits in the batched Runtime job (56.1 and 62.4 s). The IBM-reported device-execution time is not the billed time.

	sbatch total	IBM queue	IBM device exec.	IBM total	Local overhead	Billed QPU / per circuit
A1-70q trained	12m18s	5.3s	10m04s	10m09s	2m09s	561 s / 56.1 s
A1-70q untrained	39m46s	4m40s	32m33s	37m13s	2m34s	624 s / 62.4 s

Table A17: Recovered energy and excitation-level composition of the recovered determinant set for the SnOH3-me L tier, per lineage/policy and circuit type. **Caution:** S/D/T/Q/higher are counts of determinant pairs (a_i, b_j) in the Cartesian product of the recovered string list with itself (the same list serves the α and the β sector of these closed-shell systems), not counts out of the 100,000 shots; the per-shot picture is Table A21. With n recovered strings the product has n^2 pairs, and the five columns sum to n^2-1 in every row (the Hartree–Fock pair is not listed); n = 20,000, i.e. 4×10^8 pairs, in the rows that reached the cap. "Higher" absorbs every order above quadruple. Ranking, selection, shot counts and recovery settings are as in Table 6.

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
A1-46q trained	best	−479.299926	14.3	224	16206	414092	0	18885502
A1-46q trained	second	−479.121483	192.7	28	616	9224	0	399990131
A1-46q trained	merged	−479.258786	55.4	202	11075	91470	0	399897252
A1-46q untrained	best	−479.114005	200.2	16	392	5434	0	399994157
A1-46q untrained	second	−479.117972	196.2	16	488	6624	0	399992871
A1-46q untrained	merged	−479.125474	188.7	26	661	9786	0	399989526
A1-70q trained	best	−479.306409	7.8	224	17064	515550	0	93984445
A1-70q trained	second	−479.116356	197.9	24	624	9540	0	399989811
A1-70q trained	merged	−479.272504	41.7	216	12692	113924	0	399873167
A1-70q untrained	best	−479.120472	193.7	30	751	11662	0	399987556
A1-70q untrained	second	−479.119544	194.7	24	528	7576	0	399991871
A1-70q untrained	merged	−479.122028	192.2	18	497	6830	0	399992654
A5-46q trained	best	−479.301748	12.5	224	16092	402204	0	36706128
A5-46q trained	second	−479.306572	7.6	224	16448	444030	0	43615618
A5-46q trained	merged	−479.284378	29.8	222	13875	174854	0	399811048
A5-46q untrained	best	−479.121282	192.9	14	473	6342	0	399993170
A5-46q untrained	second	−479.120983	193.2	20	516	7320	0	399992143
A5-46q untrained	merged	−479.119918	194.3	18	495	7024	0	399992462
A5-70q trained	best	−479.308601	5.6	224	17250	537532	0	97573829

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
A5-70q trained	second	−479.297326	16.9	224	16526	452264	0	44945106
A5-70q trained	merged	−479.284631	29.6	224	14468	218526	0	399766781
A5-70q untrained	best	−479.126065	188.2	34	739	11126	0	399988100
A5-70q untrained	second	−479.113307	200.9	16	416	5772	0	399993795
A5-70q untrained	merged	−479.115344	198.9	12	416	5374	0	399994197

Several patterns are visible from the eight L-tier lineages (Table A15).

First, trained policies separate from untrained controls. Every untrained lineage's best, second, merged and E+D result lies between 186 and 201 mHa, whereas each trained lineage has at least one individual-circuit result at 5.6–14.3 mHa (A1-46q: 14.3 best; A5-46q: 12.5 best / 7.6 second; A1-70q: 7.8 best; A5-70q: 5.6 best).

Second, the 70q-trained lineages are as good as the 46q-trained ones on this target even though it is a cross-norb target for them: A1-70q and A5-70q reach 7.8 and 5.6 mHa, against 14.3 and 12.5 mHa for the matched-norb A1-46q and A5-46q. With one molecule and one fresh sample per lineage we do not rank the lineages.

Third, a second circuit with few valid shots recovers poorly. The A1-46q and A1-70q second circuits have 3 and 355 valid shots and recover to 192.7 and 197.9 mHa, while the A5-46q second circuit (5055 valid shots) recovers to 7.6 mHa. A finite ranking energy from a few valid shots is not a reliable ranking signal.

Fourth, merging pools weaker circuits with better ones: the merged result (29.6–55.4 mHa for the trained lineages) is worse than the best individual circuit in every trained lineage. Finally, ranking-energy order is not recovered-energy order: A5-46q trained's second circuit (7.6 mHa) beats its best (12.5 mHa) here.

Across the 43 HF-penalty runs and all 430 individual sampled circuits (ten per run, not only the best and second), 329 circuits carry at least one gate on every one of the target molecule's qubits; 101 circuits in 37 runs leave between 2 and 14 qubits with no gate after translation to {rz, sx, cx} (Table A18), and 128 circuits have at least one qubit that is in no two-qubit gate. Idle-qubit circuits include many of the selected ones: 30 of the 43 best circuits and 21 of the 43 second circuits leave qubits idle, for example the best circuit of SnOH3-me-2h2o L with A5-70q (14 idle of 70). A coarser count stored with each circuit counts qubits named by any instruction of the pre-transpile circuit, including composite gates that translate to nothing on some of their qubits, so it does not detect these cases. Whether the idle qubits matter for accuracy was not tested: SnOH3-me-h2o M with A1-70q, whose best circuit has 2 idle qubits, gives the lowest error of that target (0.71 mHa). Operators that act on the full register (an orbital rotation, a diagonal-Coulomb layer, the LUCJ layer) link all qubits in one block, whereas circuits built only from local gates split into two spin blocks (Section 2.5).

Representative circuit diagrams are not shown, since no single-page rendering stays legible at these qubit counts and depths. The per-circuit depth and two-qubit-gate counts are in Table A15: the best circuits of the four trained L-tier lineages have 64–124 two-qubit gates (as-submitted depth 37–144), against depth 334–24,119 and 1,120–45,351 two-qubit gates for the untrained lineages' best and second circuits.

The gap between the trained and untrained A1-70q sample jobs' total elapsed time (12m18s vs. 39m46s) is not Slurm queue wait or local overhead – both are a few minutes at most and comparable between the two jobs (Table A16) – it is almost entirely time spent on IBM's own side: the 27.5-minute gap lies about 22.5 min in device execution (32m33s vs. 10m04s) and about 4.6 min in the IBM device queue (4m40s vs. 5.3s). Despite the first-campaign untrained policy's circuits being 3.0–3.2× deeper/costlier (those sample jobs' depths, not repeated in Table A15), the billed QPU time differs by only 11% (624 vs. 561 s; Table A16). The metadata do not show the cause of the extra device-execution time.

J.1 Cross-norb vocab-mismatch ablation

A separate campaign ran the 70q-trained A1/A5 checkpoints (and, for one target, the 46q-trained A5 checkpoint) against three further unseen molecules whose active-space orbital count (norb) differs from that lineage's own training norb – confirmed in direct code inspection before running (Section 3.4) to be a vocab-mismatch condition: the operator pool's gate vocabulary is built fresh per molecule from that molecule's own norb, so a token index that meant one specific gate during training can mean a different gate (or an identity pad) at the target's different norb. The numbering depends on the electron count as well. Table 1 gives the token ordering. In particular, the two Givens angles on neighbouring orbitals $(p, p{+}1)$ occupy tokens $6{+}2p$ and $6{+}2p{+}1$, so the HOMO–LUMO pair (the only neighbouring pair that changes the Hartree–Fock determinant) appears at different token ids in different molecules: token 28 for SnOH3-me S (16 orbitals, 24 electrons), 34 for the hydrate M tier (23 orbitals, 30 electrons), 36 for SnOH3-me L (23 orbitals, 32 electrons) and 52 for the dihydrate L tier (35 orbitals, 48 electrons). Whether the policies learned such a preference was not checked. Results below are reported as an ablation of the vocabulary mismatch, not as clean zero-shot transfer. Targets: SnOH3-me S (norb=16, 32 qubits, trained checkpoints only –

Table A18: Qubit coverage of the fresh-shot sampled circuits, per run, over all ten circuits of each of the 43 runs (430 circuits). Each generated circuit was translated to {rz, sx, cx} without a coupling map and the qubits carrying at least one gate (or at least one two-qubit gate) were counted; "idle" means no gate at all after translation. 101 circuits leave at least one qubit idle. The earlier `active_qubits` field stored with each circuit counts the qubits named by any instruction of the pre-transpile circuit, including composite gates that translate to nothing on some qubits; it equals the register width for every circuit and so does not detect these idle qubits. Targets of a different norb than the lineage's training norb are the cross-norb ablation of Section J.1.

Target	Lineage/Policy	Target qubits	Qubits with a gate (min / max)	Qubits in a 2q gate (min / max)	Max idle qubits	Circuits with idle qubits
SnOH3-me S (32q)	A1-46q trained	32	30 / 32	20 / 32	2	2/10
SnOH3-me S (32q)	A1-70q trained	32	32 / 32	28 / 32	0	0/10
SnOH3-me S (32q)	A5-46q trained	32	28 / 32	22 / 32	4	4/10
SnOH3-me S (32q)	A5-70q trained	32	28 / 32	22 / 32	4	1/10
addlinespace						
addlinespace SnOH3-me-h2o S (38q)	A1-46q trained	38	34 / 38	26 / 38	4	4/10
SnOH3-me-h2o S (38q)	A1-70q trained	38	36 / 38	32 / 38	2	1/10
SnOH3-me-h2o S (38q)	A5-46q trained	38	30 / 38	20 / 38	8	1/10
SnOH3-me-h2o S (38q)	A5-70q trained	38	30 / 38	26 / 38	8	3/10
addlinespace						
addlinespace SnOH3-me M (38q)	A1-46q trained	38	32 / 38	26 / 38	6	4/10
SnOH3-me M (38q)	A1-70q trained	38	36 / 38	32 / 38	2	1/10
SnOH3-me M (38q)	A5-46q trained	38	38 / 38	30 / 38	0	0/10
SnOH3-me M (38q)	A5-70q trained	38	38 / 38	30 / 38	0	0/10
addlinespace						
addlinespace SnOH3-me-2h2o S (44q)	A1-46q trained	44	42 / 44	40 / 44	2	1/10
SnOH3-me-2h2o S (44q)	A1-70q trained	44	42 / 44	40 / 44	2	2/10
SnOH3-me-2h2o S (44q)	A5-46q trained	44	38 / 44	32 / 44	6	2/10
SnOH3-me-2h2o S (44q)	A5-70q trained	44	42 / 44	34 / 44	2	1/10
addlinespace						
addlinespace SnOH3-me L (46q)	A1-46q trained	46	42 / 46	34 / 46	4	1/10
SnOH3-me L (46q)	A1-70q trained	46	44 / 46	42 / 46	2	1/10
SnOH3-me L (46q)	A5-46q trained	46	38 / 46	34 / 46	8	3/10
SnOH3-me L (46q)	A5-70q trained	46	42 / 46	40 / 46	4	2/10
addlinespace						
addlinespace SnOH3-me-h2o M (46q)	A1-46q trained	46	42 / 46	34 / 46	4	1/10
SnOH3-me-h2o M (46q)	A1-70q trained	46	44 / 46	42 / 46	2	1/10
SnOH3-me-h2o M (46q)	A5-46q trained	46	40 / 46	30 / 46	6	7/10
SnOH3-me-h2o M (46q)	A5-70q trained	46	38 / 46	32 / 46	8	7/10
addlinespace						
addlinespace SnOH3-me LL (56q)	A1-46q trained	56	52 / 56	46 / 56	4	3/10
SnOH3-me LL (56q)	A1-70q trained	56	56 / 56	52 / 56	0	0/10
SnOH3-me LL (56q)	A5-70q trained	56	54 / 56	42 / 56	2	1/10
addlinespace						
addlinespace SnOH3-me-h2o L (56q)	A1-46q trained	56	52 / 56	46 / 56	4	3/10
SnOH3-me-h2o L (56q)	A1-70q trained	56	56 / 56	52 / 56	0	0/10
SnOH3-me-h2o L (56q)	A5-46q trained	56	44 / 56	38 / 56	12	4/10
SnOH3-me-h2o L (56q)	A5-70q trained	56	50 / 56	36 / 56	6	3/10
addlinespace						
addlinespace SnOH3-me-2h2o M (58q)	A1-46q trained	58	54 / 58	48 / 58	4	1/10
SnOH3-me-2h2o M (58q)	A1-70q trained	58	54 / 58	52 / 58	4	2/10
SnOH3-me-2h2o M (58q)	A5-46q trained	58	54 / 58	44 / 58	4	2/10
SnOH3-me-2h2o M (58q)	A5-70q trained	58	52 / 58	32 / 58	6	5/10
addlinespace						
addlinespace SnOH3-me-h2o LL (66q)	A1-46q trained	66	54 / 66	50 / 66	12	3/10
SnOH3-me-h2o LL (66q)	A1-70q trained	66	66 / 66	64 / 66	0	0/10
SnOH3-me-h2o LL (66q)	A5-46q trained	66	60 / 66	48 / 66	6	3/10
SnOH3-me-h2o LL (66q)	A5-70q trained	66	56 / 66	42 / 66	10	6/10
addlinespace						
addlinespace SnOH3-me-2h2o L (70q)	A1-46q trained	70	62 / 70	52 / 70	8	3/10
SnOH3-me-2h2o L (70q)	A1-70q trained	70	64 / 70	60 / 70	6	1/10
SnOH3-me-2h2o L (70q)	A5-46q trained	70	62 / 70	50 / 70	8	6/10
SnOH3-me-2h2o L (70q)	A5-70q trained	70	56 / 70	42 / 70	14	5/10

DMRG reference −479.193382 E_h), SnOH3-me-h2o M (norb=23, 46 qubits, trained and untrained – DMRG reference −555.125158 E_h), and SnOH3-me LL (norb=28, 56 qubits – DMRG reference −479.315057 E_h, Table A1).

Table A20: Recovered energy and pair-counted excitation composition for the targets of Table A19; same pair-counting convention and caution as Table A17. Ranking, selection, shot counts and recovery settings are as in Table 6.

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
SnOH3-me S (32q) A1-46q trained	best	−479.193266	0.12	96	3088	38790	0	1248521
SnOH3-me S (32q) A1-46q trained	second	−479.192672	0.71	96	2996	33466	0	239066
SnOH3-me S (32q) A1-46q trained	merged	−479.193374	0.01	96	3096	39776	0	3269431
SnOH3-me S (32q) A1-70q trained	best	−479.193164	0.22	96	3096	39054	0	1053962
SnOH3-me S (32q) A1-70q trained	second	−479.193270	0.11	96	3086	38722	0	1310664
SnOH3-me S (32q) A1-70q trained	merged	−479.193374	0.01	96	3096	39776	0	3269431
SnOH3-me S (32q) A5-46q trained	best	−479.192727	0.65	96	3040	35626	0	296478
SnOH3-me S (32q) A5-46q trained	second	−479.193081	0.30	96	3092	38574	0	690973
SnOH3-me S (32q) A5-46q trained	merged	−479.193374	0.01	96	3096	39776	0	3269431

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
SnOH3-me S (32q) A5-70q trained	best	−479.192924	0.46	96	3074	37530	0	527815
SnOH3-me S (32q) A5-70q trained	second	−479.193134	0.25	96	3072	37644	0	709143
SnOH3-me S (32q) A5-70q trained	merged	−479.193374	0.01	96	3096	39776	0	3269431
SnOH3-me-h2o S (38q) A1-46q trained	best	−555.108538	0.29	120	4658	64666	0	1992651
SnOH3-me-h2o S (38q) A1-46q trained	second	−555.108684	0.15	120	4730	68696	0	1151902
SnOH3-me-h2o S (38q) A1-46q trained	merged	−555.108826	0.01	120	4860	79268	0	18466000
SnOH3-me-h2o S (38q) A1-70q trained	best	−555.108766	0.07	120	4834	75660	0	3293954
SnOH3-me-h2o S (38q) A1-70q trained	second	−555.108751	0.08	120	4808	74266	0	3546021
SnOH3-me-h2o S (38q) A1-70q trained	merged	−555.108826	0.01	120	4860	79268	0	18466000
SnOH3-me-h2o S (38q) A5-46q trained	best	−555.108713	0.12	120	4720	67806	0	1029853
SnOH3-me-h2o S (38q) A5-46q trained	second	−555.108826	0.01	120	4860	79278	0	18873057
SnOH3-me-h2o S (38q) A5-46q trained	merged	−555.108826	0.01	120	4860	79268	0	18466000
SnOH3-me-h2o S (38q) A5-70q trained	best	−555.108740	0.09	120	4774	71788	0	2276473
SnOH3-me-h2o S (38q) A5-70q trained	second	−555.108764	0.07	120	4834	75254	0	2130960
SnOH3-me-h2o S (38q) A5-70q trained	merged	−555.108826	0.01	120	4860	79274	0	18777394
SnOH3-me M (38q) A1-46q trained	best	−479.292683	9.6	168	9310	193124	0	19930566
SnOH3-me M (38q) A1-46q trained	second	−479.296591	5.7	168	9408	201334	0	18955973
SnOH3-me M (38q) A1-46q trained	merged	−479.286434	15.8	168	8936	162134	0	399828761
SnOH3-me M (38q) A1-70q trained	best	−479.298463	3.8	168	9688	226642	0	41404710
SnOH3-me M (38q) A1-70q trained	second	−479.295672	6.6	168	9488	210426	0	61355326
SnOH3-me M (38q) A1-70q trained	merged	−479.282569	19.7	168	8712	142664	0	399848455
SnOH3-me M (38q) A5-46q trained	best	−479.285058	17.2	168	8972	163332	0	9920856
SnOH3-me M (38q) A5-46q trained	second	−479.282828	19.5	168	8982	164246	0	11557228
SnOH3-me M (38q) A5-46q trained	merged	−479.271416	30.9	168	7994	81576	0	399910261
SnOH3-me M (38q) A5-70q trained	best	−479.288699	13.6	168	9142	178184	0	8526809
SnOH3-me M (38q) A5-70q trained	second	−479.152486	149.8	56	1688	30954	0	399967301
SnOH3-me M (38q) A5-70q trained	merged	−479.265913	36.4	162	7809	106826	0	399885202
SnOH3-me-2h2o S (44q) A1-46q trained	best	−631.083103	0.21	144	6886	125446	0	7640467
SnOH3-me-2h2o S (44q) A1-46q trained	second	−631.083239	0.07	144	6904	126698	0	6798942
SnOH3-me-2h2o S (44q) A1-46q trained	merged	−631.083269	0.04	144	7020	138720	0	53363340
SnOH3-me-2h2o S (44q) A1-70q trained	best	−631.083257	0.06	144	6964	131694	0	10685297
SnOH3-me-2h2o S (44q) A1-70q trained	second	−631.083084	0.23	144	6818	120528	0	7380109
SnOH3-me-2h2o S (44q) A1-70q trained	merged	−631.083269	0.04	144	7020	138720	0	53363340
SnOH3-me-2h2o S (44q) A5-46q trained	best	−631.083185	0.13	144	6836	120756	0	3444363
SnOH3-me-2h2o S (44q) A5-46q trained	second	−631.083207	0.11	144	6876	124062	0	4841817
SnOH3-me-2h2o S (44q) A5-46q trained	merged	−631.083267	0.04	144	7020	137030	0	15292846
SnOH3-me-2h2o S (44q) A5-70q trained	best	−631.083092	0.22	144	6834	120950	0	4690096
SnOH3-me-2h2o S (44q) A5-70q trained	second	−631.083044	0.27	144	6798	118334	0	4309959
SnOH3-me-2h2o S (44q) A5-70q trained	merged	−631.083269	0.04	144	7020	138720	0	53348711
SnOH3-me-h2o M (46q) A1-46q trained	best	−555.123748	1.4	240	18500	496376	0	24205667
SnOH3-me-h2o M (46q) A1-46q trained	second	−555.102442	22.7	14	411	5278	0	399994296
SnOH3-me-h2o M (46q) A1-46q trained	merged	−555.118018	7.1	210	11639	66550	0	399921600
SnOH3-me-h2o M (46q) A1-46q untrained	best	−555.102302	22.9	8	294	3508	0	399996189
SnOH3-me-h2o M (46q) A1-46q untrained	second	−555.102364	22.8	10	339	3992	0	399995658
SnOH3-me-h2o M (46q) A1-46q untrained	merged	−555.102310	22.8	6	295	3022	0	399996676
SnOH3-me-h2o M (46q) A1-70q trained	best	−555.124450	0.71	240	19496	621246	0	105552042
SnOH3-me-h2o M (46q) A1-70q trained	second	−555.102207	23.0	4	212	2308	0	399997475
SnOH3-me-h2o M (46q) A1-70q trained	merged	−555.118773	6.4	230	14021	93558	0	399892190
SnOH3-me-h2o M (46q) A1-70q untrained	best	−555.102456	22.7	10	249	3314	0	399996426
SnOH3-me-h2o M (46q) A1-70q untrained	second	−555.102468	22.7	14	181	1768	0	399998036
SnOH3-me-h2o M (46q) A1-70q untrained	merged	−555.102255	22.9	14	289	3702	0	399995994
SnOH3-me-h2o M (46q) A5-46q trained	best	−555.123672	1.5	240	18354	478028	0	16700986
SnOH3-me-h2o M (46q) A5-46q trained	second	−555.123151	2.0	240	17760	408178	0	27918797
SnOH3-me-h2o M (46q) A5-46q trained	merged	−555.123131	2.0	240	18042	441442	0	399540275
SnOH3-me-h2o M (46q) A5-70q trained	best	−555.123764	1.4	240	19000	561904	0	96224776
SnOH3-me-h2o M (46q) A5-70q trained	second	−555.123482	1.7	240	18406	488464	0	59137618
SnOH3-me-h2o M (46q) A5-70q trained	merged	−555.124246	0.91	240	19488	619226	0	399361045
SnOH3-me-h2o M (46q) A5-70q untrained	best	−555.102543	22.6	8	242	2946	0	399996803
SnOH3-me-h2o M (46q) A5-70q untrained	second	−555.102455	22.7	6	43	580	0	399999370

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
SnOH3-me-h2o M (46q) A5-70q untrained	merged	−555.102186	23.0	8	238	2950	0	399996803
SnOH3-me LL (56q) A1-46q trained	best	−479.300183	14.9	294	27785	919854	0	237940002
SnOH3-me LL (56q) A1-46q trained	second	−479.299292	15.8	294	27607	892998	0	218178304
SnOH3-me LL (56q) A1-46q trained	merged	−479.286822	28.2	294	25185	528534	0	399445986
SnOH3-me LL (56q) A1-70q trained	best	−479.304625	10.4	294	28161	975858	0	398995686
SnOH3-me LL (56q) A1-70q trained	second	−479.112860	202.2	10	173	2370	0	399997446
SnOH3-me LL (56q) A1-70q trained	merged	−479.279821	35.2	294	22871	187598	0	399789236
SnOH3-me LL (56q) A1-70q untrained	best	−479.302734	12.3	294	29007	1103188	0	379663706
SnOH3-me LL (56q) A1-70q untrained	second	−479.110195	204.9	0	16	160	0	399999823
SnOH3-me LL (56q) A1-70q untrained	merged	−479.279604	35.5	294	23847	331504	0	399644354
SnOH3-me LL (56q) A5-70q trained	best	−479.301896	13.2	294	27359	859128	0	303022707
SnOH3-me LL (56q) A5-70q trained	second	−479.111752	203.3	10	241	3240	0	399996508
SnOH3-me LL (56q) A5-70q trained	merged	−479.261236	53.8	254	17777	211854	0	399770114
SnOH3-me-h2o L (56q) A1-46q trained	best	−555.132208	2.8	320	32232	1073160	0	288166351
SnOH3-me-h2o L (56q) A1-46q trained	second	−555.132421	2.6	320	31772	999304	0	272871103
SnOH3-me-h2o L (56q) A1-46q trained	merged	−555.129520	5.5	314	26771	334486	0	399638428
SnOH3-me-h2o L (56q) A1-70q trained	best	−555.133193	1.9	320	33296	1245064	0	398721319
SnOH3-me-h2o L (56q) A1-70q trained	second	−555.132552	2.5	320	32806	1166154	0	398800719
SnOH3-me-h2o L (56q) A1-70q trained	merged	−555.127221	7.8	314	26311	262464	0	399710910
SnOH3-me-h2o L (56q) A5-46q trained	best	−555.131417	3.6	320	31506	954976	0	204677478
SnOH3-me-h2o L (56q) A5-46q trained	second	−555.132660	2.4	320	32916	1182632	0	391101380
SnOH3-me-h2o L (56q) A5-46q trained	merged	−555.129768	5.3	320	29162	573050	0	399397467
SnOH3-me-h2o L (56q) A5-70q trained	best	−555.132405	2.7	320	31550	962282	0	206625128
SnOH3-me-h2o L (56q) A5-70q trained	second	−555.132408	2.6	320	31702	986450	0	186370248
SnOH3-me-h2o L (56q) A5-70q trained	merged	−555.130714	4.3	320	29542	634626	0	399335511
SnOH3-me-2h2o M (58q) A1-46q trained	best	−631.112084	6.5	394	45927	1418666	0	398535012
SnOH3-me-2h2o M (58q) A1-46q trained	second	−631.077480	41.1	0	14	190	0	399999795
SnOH3-me-2h2o M (58q) A1-46q trained	merged	−631.095105	23.5	232	14422	112804	0	399872541
SnOH3-me-2h2o M (58q) A1-70q trained	best	−631.114317	4.3	396	49464	2047290	0	397902849
SnOH3-me-2h2o M (58q) A1-70q trained	second	−631.113058	5.6	396	47314	1623144	0	398329145
SnOH3-me-2h2o M (58q) A1-70q trained	merged	−631.106177	12.4	374	36859	354898	0	399607868
SnOH3-me-2h2o M (58q) A5-46q trained	best	−631.114080	4.5	396	48848	1927832	0	358908932
SnOH3-me-2h2o M (58q) A5-46q trained	second	−631.113620	5.0	396	48692	1893402	0	227398245
SnOH3-me-2h2o M (58q) A5-46q trained	merged	−631.105052	13.6	346	31705	308132	0	399659816
SnOH3-me-2h2o M (58q) A5-70q trained	best	−631.113705	4.9	396	47538	1658058	0	86973896
SnOH3-me-2h2o M (58q) A5-70q trained	second	−631.111954	6.7	396	45066	1170682	0	115812979
SnOH3-me-2h2o M (58q) A5-70q trained	merged	−631.110745	7.9	394	42011	632994	0	399324600
SnOH3-me-h2o LL (66q) A1-46q trained	best	−555.135328	2.1	400	51898	2393474	0	284585591
SnOH3-me-h2o LL (66q) A1-46q trained	second	−555.133503	3.9	400	48660	1745814	0	236845829
SnOH3-me-h2o LL (66q) A1-46q trained	merged	−555.131333	6.1	400	44596	923320	0	399031683
SnOH3-me-h2o LL (66q) A1-70q trained	best	−555.134846	2.6	400	49602	1935652	0	398014345
SnOH3-me-h2o LL (66q) A1-70q trained	second	−555.102002	35.4	0	36	544	0	399999419
SnOH3-me-h2o LL (66q) A1-70q trained	merged	−555.131599	5.8	398	42647	608990	0	399347964
SnOH3-me-h2o LL (66q) A5-46q trained	best	−555.133877	3.5	400	49706	1953242	0	197173420
SnOH3-me-h2o LL (66q) A5-46q trained	second	−555.134389	3.0	400	51528	2322426	0	380884574
SnOH3-me-h2o LL (66q) A5-46q trained	merged	−555.130850	6.5	400	42928	587058	0	399369613
SnOH3-me-h2o LL (66q) A5-70q trained	best	−555.135231	2.2	400	51502	2316104	0	318256829
SnOH3-me-h2o LL (66q) A5-70q trained	second	−555.134346	3.1	400	49142	1846532	0	398103925
SnOH3-me-h2o LL (66q) A5-70q trained	merged	−555.133797	3.6	400	49072	1824750	0	398125777
SnOH3-me-2h2o L (70q) A1-46q trained	best	−631.123614	8.6	528	81784	3207164	0	396710523
SnOH3-me-2h2o L (70q) A1-46q trained	second	−631.123272	9.0	528	79590	2629546	0	397290335
SnOH3-me-2h2o L (70q) A1-46q trained	merged	−631.119423	12.8	516	68948	616726	0	399313809
SnOH3-me-2h2o L (70q) A1-70q trained	best	−631.124552	7.7	528	79440	2588922	0	397331109
SnOH3-me-2h2o L (70q) A1-70q trained	second	−631.077480	54.8	0	0	16	0	399999983
SnOH3-me-2h2o L (70q) A1-70q trained	merged	−631.099821	32.4	330	28945	285278	0	399685446
SnOH3-me-2h2o L (70q) A5-46q trained	best	−631.124857	7.4	528	83086	3552970	0	396363415
SnOH3-me-2h2o L (70q) A5-46q trained	second	−631.123201	9.0	526	79153	2643128	0	397277192
SnOH3-me-2h2o L (70q) A5-46q trained	merged	−631.122669	9.6	526	76467	1925512	0	397997494
SnOH3-me-2h2o L (70q) A5-70q trained	best	−631.120256	12.0	500	67066	1152464	0	398779969

Target/Lineage	Type	Energy (E_h)	ΔE (mHa)	S	D	T	Q	Higher
SnOH3-me-2h2o L (70q) A5-70q trained	second	−631.124322	7.9	526	77617	2239082	0	334368575
SnOH3-me-2h2o L (70q) A5-70q trained	merged	−631.110701	21.5	410	44667	546210	0	399408712

Table A21: Raw-shot excitation-level composition (HF/S/D/T/Q/higher) of the sampled bitstrings, per lineage/policy, for the best and second circuit and for the merged pool. Unlike Tables A17 and A20, counts are over the measured shots themselves (`raw_shot_excitation_histogram`, weighted by multiplicity), so each best/second row sums to 100,000 ("n"); the merged row sums to the pooled shots of all finite-energy circuits. Ranking, selection, shot counts and recovery settings are as in Table 6.

Target/Lineage	Type	HF	S	D	T	Q	Higher	n
SnOH3-me S (32q) A1-46q trained	best	9563	47174	35302	7217	704	40	100000
SnOH3-me S (32q) A1-46q trained	second	73648	24464	1823	64	1	0	100000
SnOH3-me S (32q) A1-46q trained	merged	174718	158128	57496	10431	9313	589914	1000000
SnOH3-me S (32q) A1-70q trained	best	15864	53088	26321	4322	395	10	100000
SnOH3-me S (32q) A1-70q trained	second	15619	50395	27497	5836	601	52	100000
SnOH3-me S (32q) A1-70q trained	merged	31483	103483	53831	10447	4126	796630	1000000
SnOH3-me S (32q) A5-46q trained	best	75767	22037	2076	119	1	0	100000
SnOH3-me S (32q) A5-46q trained	second	37986	49858	11077	1010	68	1	100000
SnOH3-me S (32q) A5-46q trained	merged	343611	271049	76571	9812	9776	289181	1000000
SnOH3-me S (32q) A5-70q trained	best	38407	50890	9743	917	41	2	100000
SnOH3-me S (32q) A5-70q trained	second	41440	47694	9839	961	64	2	100000
SnOH3-me S (32q) A5-70q trained	merged	233308	140523	24451	6476	19835	575407	1000000
SnOH3-me-h2o S (38q) A1-46q trained	best	29168	50965	17380	2308	170	9	100000
SnOH3-me-h2o S (38q) A1-46q trained	second	11483	51150	32605	4475	279	8	100000
SnOH3-me-h2o S (38q) A1-46q trained	merged	68750	216647	100105	13754	2961	397783	800000
SnOH3-me-h2o S (38q) A1-70q trained	best	36208	45204	15453	2803	297	35	100000
SnOH3-me-h2o S (38q) A1-70q trained	second	7494	42005	37942	10962	1496	101	100000
SnOH3-me-h2o S (38q) A1-70q trained	merged	50894	129503	90324	25433	5245	598601	900000
SnOH3-me-h2o S (38q) A5-46q trained	best	71348	25611	2864	171	6	0	100000
SnOH3-me-h2o S (38q) A5-46q trained	second	0	0	0	5	29	99966	100000
SnOH3-me-h2o S (38q) A5-46q trained	merged	71348	25611	2864	176	46	199955	300000
SnOH3-me-h2o S (38q) A5-70q trained	best	33180	49434	15023	2157	193	13	100000
SnOH3-me-h2o S (38q) A5-70q trained	second	47581	42541	8936	900	41	1	100000
SnOH3-me-h2o S (38q) A5-70q trained	merged	141386	251716	92217	13551	1081	100049	600000
SnOH3-me M (38q) A1-46q trained	best	29986	48487	18233	2995	284	15	100000
SnOH3-me M (38q) A1-46q trained	second	15637	59066	21921	3089	270	17	100000
SnOH3-me M (38q) A1-46q trained	merged	75820	217967	92926	12312	1482	599493	1000000
SnOH3-me M (38q) A1-70q trained	best	26574	43964	22656	5803	911	92	100000
SnOH3-me M (38q) A1-70q trained	second	4831	33697	40786	16732	3508	446	100000
SnOH3-me M (38q) A1-70q trained	merged	40414	120949	99670	32619	6377	699971	1000000
SnOH3-me M (38q) A5-46q trained	best	5815	40511	44079	8934	649	12	100000
SnOH3-me M (38q) A5-46q trained	second	9128	53451	31943	5040	412	26	100000
SnOH3-me M (38q) A5-46q trained	merged	14943	93962	76022	13993	1320	799760	1000000
SnOH3-me M (38q) A5-70q trained	best	15975	51746	27391	4502	370	16	100000
SnOH3-me M (38q) A5-70q trained	second	0	0	1	4	93	99902	100000
SnOH3-me M (38q) A5-70q trained	merged	15975	51746	27392	4522	768	899597	1000000
SnOH3-me-2h2o S (44q) A1-46q trained	best	6777	39820	38991	12317	1883	212	100000
SnOH3-me-2h2o S (44q) A1-46q trained	second	5397	43261	39345	10603	1297	97	100000
SnOH3-me-2h2o S (44q) A1-46q trained	merged	12174	83081	78336	22952	3611	399846	600000
SnOH3-me-2h2o S (44q) A1-70q trained	best	8424	40114	35955	12733	2458	316	100000
SnOH3-me-2h2o S (44q) A1-70q trained	second	2848	30352	46181	17456	2875	288	100000
SnOH3-me-2h2o S (44q) A1-70q trained	merged	16026	100280	123367	49556	9698	301073	600000
SnOH3-me-2h2o S (44q) A5-46q trained	best	23489	53166	20330	2792	213	10	100000
SnOH3-me-2h2o S (44q) A5-46q trained	second	13090	47446	30633	7606	1112	113	100000
SnOH3-me-2h2o S (44q) A5-46q trained	merged	60325	226041	164735	42601	5784	514	500000
SnOH3-me-2h2o S (44q) A5-70q trained	best	13935	49122	30029	6236	631	47	100000
SnOH3-me-2h2o S (44q) A5-70q trained	second	14446	49445	29194	6252	631	32	100000
SnOH3-me-2h2o S (44q) A5-70q trained	merged	28381	98567	59231	12577	2021	99223	300000
SnOH3-me L (46q) A1-46q trained	best	27683	52820	17094	2250	145	8	100000

Target/Lineage	Type	HF	S	D	T	Q	Higher	n
SnOH3-me L (46q) A1-46q trained	second	0	0	0	0	0	100000	100000
SnOH3-me L (46q) A1-46q trained	merged	27683	52820	17094	2256	229	799918	900000
SnOH3-me L (46q) A1-46q untrained	best	0	0	0	0	0	100000	100000
SnOH3-me L (46q) A1-46q untrained	second	0	0	0	2	51	99947	100000
SnOH3-me L (46q) A1-46q untrained	merged	0	0	0	4	80	899916	900000
SnOH3-me L (46q) A1-70q trained	best	22910	47962	23493	5013	575	47	100000
SnOH3-me L (46q) A1-70q trained	second	0	0	0	1	6	99993	100000
SnOH3-me L (46q) A1-70q trained	merged	22910	47962	23493	5014	590	700031	800000
SnOH3-me L (46q) A1-70q untrained	best	0	0	0	1	37	99962	100000
SnOH3-me L (46q) A1-70q untrained	second	0	0	0	0	0	100000	100000
SnOH3-me L (46q) A1-70q untrained	merged	0	0	0	1	44	999955	1000000
SnOH3-me L (46q) A5-46q trained	best	11558	47087	32743	7713	844	55	100000
SnOH3-me L (46q) A5-46q trained	second	14844	50116	28378	5962	647	53	100000
SnOH3-me L (46q) A5-46q trained	merged	56270	147149	78341	16361	1774	500105	800000
SnOH3-me L (46q) A5-46q untrained	best	0	0	0	0	4	99996	100000
SnOH3-me L (46q) A5-46q untrained	second	0	0	0	0	1	99999	100000
SnOH3-me L (46q) A5-46q untrained	merged	0	0	0	1	22	999977	1000000
SnOH3-me L (46q) A5-70q trained	best	10334	44027	33879	10033	1563	164	100000
SnOH3-me L (46q) A5-70q trained	second	3866	32527	44595	16051	2690	271	100000
SnOH3-me L (46q) A5-70q trained	merged	21220	116464	117086	38289	6359	700582	1000000
SnOH3-me L (46q) A5-70q untrained	best	0	0	0	0	0	100000	100000
SnOH3-me L (46q) A5-70q untrained	second	0	0	0	0	1	99999	100000
SnOH3-me L (46q) A5-70q untrained	merged	0	0	0	1	11	999988	1000000
SnOH3-me-h2o M (46q) A1-46q trained	best	26532	52848	17834	2599	178	9	100000
SnOH3-me-h2o M (46q) A1-46q trained	second	0	0	0	1	12	99987	100000
SnOH3-me-h2o M (46q) A1-46q trained	merged	26532	52848	17834	2605	234	899947	1000000
SnOH3-me-h2o M (46q) A1-46q untrained	best	0	0	0	0	4	99996	100000
SnOH3-me-h2o M (46q) A1-46q untrained	second	0	0	0	1	5	99994	100000
SnOH3-me-h2o M (46q) A1-46q untrained	merged	0	0	0	3	49	999948	1000000
SnOH3-me-h2o M (46q) A1-70q trained	best	21761	47739	24387	5377	682	54	100000
SnOH3-me-h2o M (46q) A1-70q trained	second	0	0	0	0	1	99999	100000
SnOH3-me-h2o M (46q) A1-70q trained	merged	21761	47739	24387	5378	694	900041	1000000
SnOH3-me-h2o M (46q) A1-70q untrained	best	0	0	0	0	2	99998	100000
SnOH3-me-h2o M (46q) A1-70q untrained	second	0	0	0	0	1	99999	100000
SnOH3-me-h2o M (46q) A1-70q untrained	merged	0	0	0	1	8	999991	1000000
SnOH3-me-h2o M (46q) A5-46q trained	best	26761	54384	16890	1849	112	4	100000
SnOH3-me-h2o M (46q) A5-46q trained	second	6148	42569	40210	9857	1127	89	100000
SnOH3-me-h2o M (46q) A5-46q trained	merged	83510	271220	240736	87508	15304	301722	1000000
SnOH3-me-h2o M (46q) A5-70q trained	best	15641	46468	29123	7493	1163	112	100000
SnOH3-me-h2o M (46q) A5-70q trained	second	14142	49156	29314	6458	860	70	100000
SnOH3-me-h2o M (46q) A5-70q trained	merged	107078	398102	297743	83360	12527	101190	1000000
SnOH3-me-h2o M (46q) A5-70q untrained	best	0	0	0	0	3	99997	100000
SnOH3-me-h2o M (46q) A5-70q untrained	second	0	0	0	0	2	99998	100000
SnOH3-me-h2o M (46q) A5-70q untrained	merged	0	0	0	0	11	999989	1000000
SnOH3-me LL (56q) A1-46q trained	best	11109	46363	32695	8558	1155	120	100000
SnOH3-me LL (56q) A1-46q trained	second	6209	37216	39430	14250	2589	306	100000
SnOH3-me LL (56q) A1-46q trained	merged	23464	117381	112204	39228	6898	300825	600000
SnOH3-me LL (56q) A1-70q trained	best	5291	34388	39206	16886	3640	589	100000
SnOH3-me LL (56q) A1-70q trained	second	0	0	0	0	0	100000	100000
SnOH3-me LL (56q) A1-70q trained	merged	5291	34388	39206	16886	3640	400589	500000
SnOH3-me LL (56q) A1-70q untrained	best	8196	33661	35363	17227	4661	892	100000
SnOH3-me LL (56q) A1-70q untrained	second	0	0	0	0	0	100000	100000
SnOH3-me LL (56q) A1-70q untrained	merged	8196	33661	35363	17227	4661	200892	300000
SnOH3-me LL (56q) A5-70q trained	best	3320	26316	41043	21969	6207	1145	100000
SnOH3-me LL (56q) A5-70q trained	second	0	0	0	0	1	99999	100000
SnOH3-me LL (56q) A5-70q trained	merged	3320	26316	41043	21969	6208	301144	400000
SnOH3-me-h2o L (56q) A1-46q trained	best	5390	31433	40536	18063	4021	557	100000
SnOH3-me-h2o L (56q) A1-46q trained	second	3280	28637	43522	19745	4254	562	100000
SnOH3-me-h2o L (56q) A1-46q trained	merged	19279	105553	117429	46890	9615	401234	700000

Target/Lineage	Type	HF	S	D	T	Q	Higher	n
SnOH3-me-h2o L (56q) A1-70q trained	best	15894	42646	29701	9609	1862	288	100000
SnOH3-me-h2o L (56q) A1-70q trained	second	5851	35489	39026	15758	3374	502	100000
SnOH3-me-h2o L (56q) A1-70q trained	merged	21745	78135	68727	25367	5236	700790	900000
SnOH3-me-h2o L (56q) A5-46q trained	best	25345	48690	21049	4359	517	40	100000
SnOH3-me-h2o L (56q) A5-46q trained	second	10940	42947	34403	10013	1560	137	100000
SnOH3-me-h2o L (56q) A5-46q trained	merged	49353	168718	133148	41605	6515	200661	600000
SnOH3-me-h2o L (56q) A5-70q trained	best	5621	34344	41636	15295	2757	347	100000
SnOH3-me-h2o L (56q) A5-70q trained	second	9317	43133	35791	10169	1464	126	100000
SnOH3-me-h2o L (56q) A5-70q trained	merged	19891	112942	120891	39422	6202	200652	500000
SnOH3-me-2h2o M (58q) A1-46q trained	best	3041	23236	40636	24716	7127	1244	100000
SnOH3-me-2h2o M (58q) A1-46q trained	second	0	0	0	0	0	100000	100000
SnOH3-me-2h2o M (58q) A1-46q trained	merged	3041	23236	40636	24716	7127	901244	1000000
SnOH3-me-2h2o M (58q) A1-70q trained	best	13830	39660	31266	11948	2812	484	100000
SnOH3-me-2h2o M (58q) A1-70q trained	second	1644	18231	38792	28591	10270	2472	100000
SnOH3-me-2h2o M (58q) A1-70q trained	merged	15474	57891	70058	40539	13082	802956	1000000
SnOH3-me-2h2o M (58q) A5-46q trained	best	4001	28326	41146	20746	4958	823	100000
SnOH3-me-2h2o M (58q) A5-46q trained	second	7976	44070	35710	10444	1616	184	100000
SnOH3-me-2h2o M (58q) A5-46q trained	merged	11977	72396	76856	31190	6574	801007	1000000
SnOH3-me-2h2o M (58q) A5-70q trained	best	37557	44812	14856	2536	218	21	100000
SnOH3-me-2h2o M (58q) A5-70q trained	second	9683	45606	34012	9236	1327	136	100000
SnOH3-me-2h2o M (58q) A5-70q trained	merged	72528	197532	160181	56849	11329	501581	1000000
SnOH3-me-h2o LL (66q) A1-46q trained	best	31411	45412	18803	3853	476	45	100000
SnOH3-me-h2o LL (66q) A1-46q trained	second	7197	37715	38763	13528	2531	266	100000
SnOH3-me-h2o LL (66q) A1-46q trained	merged	43401	126104	137371	69934	19321	303869	700000
SnOH3-me-h2o LL (66q) A1-70q trained	best	1510	15119	33817	29825	14518	5211	100000
SnOH3-me-h2o LL (66q) A1-70q trained	second	0	0	0	0	0	100000	100000
SnOH3-me-h2o LL (66q) A1-70q trained	merged	1510	15119	33817	29825	14518	105211	200000
SnOH3-me-h2o LL (66q) A5-46q trained	best	18906	46401	26915	6685	1000	93	100000
SnOH3-me-h2o LL (66q) A5-46q trained	second	9291	39559	35828	12677	2316	329	100000
SnOH3-me-h2o LL (66q) A5-46q trained	merged	32446	115285	103553	39417	8154	301145	600000
SnOH3-me-h2o LL (66q) A5-70q trained	best	23092	45098	24376	6353	967	114	100000
SnOH3-me-h2o LL (66q) A5-70q trained	second	2779	24312	41636	23602	6452	1219	100000
SnOH3-me-h2o LL (66q) A5-70q trained	merged	41304	180903	226667	115109	30330	105687	700000
SnOH3-me-2h2o L (70q) A1-46q trained	best	9888	36840	34706	14553	3453	560	100000
SnOH3-me-2h2o L (70q) A1-46q trained	second	3686	27231	40572	21692	5757	1062	100000
SnOH3-me-2h2o L (70q) A1-46q trained	merged	14444	80586	115765	66438	19022	403745	700000
SnOH3-me-2h2o L (70q) A1-70q trained	best	648	9799	29489	33364	18814	7886	100000
SnOH3-me-2h2o L (70q) A1-70q trained	second	0	0	0	0	0	100000	100000
SnOH3-me-2h2o L (70q) A1-70q trained	merged	648	9799	29489	33364	18814	607886	700000
SnOH3-me-2h2o L (70q) A5-46q trained	best	16872	44324	28844	8366	1404	190	100000
SnOH3-me-2h2o L (70q) A5-46q trained	second	7448	36091	38624	14732	2770	335	100000
SnOH3-me-2h2o L (70q) A5-46q trained	merged	38431	185481	223634	114076	32098	306280	900000
SnOH3-me-2h2o L (70q) A5-70q trained	best	145	3260	22593	43300	24325	6377	100000
SnOH3-me-2h2o L (70q) A5-70q trained	second	612	8134	32663	39988	15446	3157	100000
SnOH3-me-2h2o L (70q) A5-70q trained	merged	944	16631	98517	194045	135955	553908	1000000

Is the amount of training what matters? A trained policy is not always better than a random one, and the separation depends on the target. Three patterns stand out. First, at the h2o-M target all four trained lineages (A1-46q, A1-70q, A5-46q, A5-70q) reach 0.71–1.5 mHa with the circuit chosen by lowest ranking energy, while for A1-46q and A1-70q the E+D 1/2 circuits (22.6–23.0 mHa) and every untrained control (22.6–22.9 mHa) stay at the level of the untrained policies (for the two A5 lineages an E+D pick coincides with the best or second circuit and is also at 1.4–2.0 mHa); the cause, e.g. valid-shot count versus ranking energy versus depth, was not isolated. Merged recoveries are 0.91–7.1 mHa for the trained lineages and 22.9–23.0 mHa for the untrained ones. The singles-injected error of the untrained lineages' best circuit is 4.9–5.1 mHa, while for the trained lineages' best it equals the plain best deviation. Second, at the 32q target the merged result of all four trained checkpoints is exact to 0.01 mHa and the best circuit's is within 0.12–0.65 mHa; this active space (norb=16) is small enough that the result says little about whether the vocab mismatch is harmless at larger norb. Third, at the largest tier where an untrained control was run (LL, norb=28, 56 qubits) the trained A1-70q best circuit reaches 10.4 mHa (as-submitted depth 62, 124 two-qubit

Table A19: Cross-norb ablation and remaining matched-norb targets, per target, lineage and policy: SnOH3-me S (32q), LL (56q), SnOH3-me-h2o S/M/LL (38/46/66q) and SnOH3-me-2h2o S/M/L (44/58/70q), with DMRG references from Table A1. Same columns as Table A15. Where valid shots are few (tens to hundreds of 100,000) the ranking energy and the pick rest on very little data. Wall times of the first campaign are not repeated here. Ranking, selection, shot counts and recovery settings are as in Table 6.

Target	Lineage/Policy	Depth (best/2nd)	2q (best/2nd)	Valid shots (best/2nd)	ΔE best/2nd/merged (mHa)	ΔE E+D 1/2 (mHa)	ΔE best-SI (mHa)
Trained policies							
SnOH3-me S (32q)	A1-46q trained	50/30	56/36	2652/34212	0.12/0.71/0.01	0.01/0.12	0.1
SnOH3-me S (32q)	A1-70q trained	41/57	68/80	21300/4078	0.22/0.11/0.01	0.01/0.01	0.2
SnOH3-me S (32q)	A5-46q trained	31/30	36/36	40942/11383	0.65/0.30/0.01	0.30/0.01	0.7
SnOH3-me S (32q)	A5-70q trained	69/64	60/76	33072/14648	0.46/0.25/0.01	0.01/0.46	0.5
SnOH3-me-h2o S (38q)	A1-46q trained	35/52	52/48	8991/13563	0.29/0.15/0.01	0.29/0.01	0.3
SnOH3-me-h2o S (38q)	A1-70q trained	57/51	92/80	14708/16736	0.07/0.08/0.01	0.01/0.01	0.1
SnOH3-me-h2o S (38q)	A5-46q trained	34/3867	40/5856	40442/3	0.12/0.01/0.01	0.12/0.01	0.1
SnOH3-me-h2o S (38q)	A5-70q trained	110/161	84/104	12917/23131	0.09/0.07/0.01	0.09/0.19	0.1
SnOH3-me M (38q)	A1-46q trained	35/41	52/60	12049/3485	9.6/5.7/15.8	152.0/9.6	9.6
SnOH3-me M (38q)	A1-70q trained	57/72	92/100	14809/2736	3.8/6.6/19.7	144.4/139.1	3.8
SnOH3-me M (38q)	A5-46q trained	32/51	52/68	20660/25340	17.2/19.5/30.9	143.9/130.6	17.2
SnOH3-me M (38q)	A5-70q trained	152/401	104/788	6989/1005	13.6/149.8/36.4	131.8/149.1	13.6
SnOH3-me-2h2o S (44q)	A1-46q trained	30/59	60/76	1544/768	0.21/0.07/0.04	0.04/0.21	0.2
SnOH3-me-2h2o S (44q)	A1-70q trained	60/59	100/108	1898/1854	0.06/0.23/0.04	0.06/0.04	0.1
SnOH3-me-2h2o S (44q)	A5-46q trained	49/62	64/80	11891/10298	0.13/0.11/0.04	0.11/0.06	0.1
SnOH3-me-2h2o S (44q)	A5-70q trained	168/142	104/112	14875/19156	0.22/0.27/0.04	0.22/0.27	0.2
SnOH3-me-h2o M (46q)	A1-46q trained	37/309	64/1068	10254/771	1.4/22.7/7.1	22.7/22.7	1.4
SnOH3-me-h2o M (46q)	A1-70q trained	69/6617	100/13041	10230/13	0.71/23.0/6.4	23.0/22.6	0.7
SnOH3-me-h2o M (46q)	A5-46q trained	58/51	64/72	23929/3130	1.5/2.0/2.0	1.4/2.0	1.5
SnOH3-me-h2o M (46q)	A5-70q trained	108/139	100/104	5346/6485	1.4/1.7/0.91	1.4/1.7	1.4
SnOH3-me LL (56q)	A1-46q trained	52/42	80/76	3788/3000	14.9/15.8/28.2	203.8/14.9	14.9
SnOH3-me LL (56q)	A1-70q trained	62/405	124/1636	1525/14	10.4/202.2/35.2	202.2/202.1	10.4
SnOH3-me LL (56q)	A5-70q trained	79/363	104/1588	1770/52	13.2/203.3/53.8	203.3/204.3	13.2
SnOH3-me-h2o L (56q)	A1-46q trained	47/42	80/76	5017/12757	2.8/2.6/5.5	33.1/2.8	2.8
SnOH3-me-h2o L (56q)	A1-70q trained	65/62	124/124	7270/1656	1.9/2.5/7.8	33.1/33.0	1.9
SnOH3-me-h2o L (56q)	A5-46q trained	53/41	64/84	8388/2750	3.6/2.4/5.3	33.1/3.6	3.6
SnOH3-me-h2o L (56q)	A5-70q trained	44/69	76/80	14182/2605	2.7/2.6/4.3	2.7/33.1	2.7
SnOH3-me-2h2o M (58q)	A1-46q trained	47/358	88/1692	2263/1196	6.5/41.1/23.5	41.1/41.1	6.0
SnOH3-me-2h2o M (58q)	A1-70q trained	65/49	132/124	5605/1202	4.3/5.6/12.4	41.1/41.1	4.3
SnOH3-me-2h2o M (58q)	A5-46q trained	45/46	72/88	7066/4184	4.5/5.0/13.6	41.1/41.1	4.5
SnOH3-me-2h2o M (58q)	A5-70q trained	62/41	68/64	19811/3396	4.9/6.7/7.9	41.1/41.1	4.9
SnOH3-me-h2o LL (66q)	A1-46q trained	60/51	92/96	12073/3357	2.1/3.9/6.1	35.4/2.1	2.1
SnOH3-me-h2o LL (66q)	A1-70q trained	51/458	156/2256	2232/3	2.6/35.4/5.8	2.6/35.4	2.6
SnOH3-me-h2o LL (66q)	A5-46q trained	44/32	84/84	9310/10483	3.5/3.0/6.5	35.4/3.5	3.5
SnOH3-me-h2o LL (66q)	A5-70q trained	57/58	84/80	7521/1020	2.2/3.1/3.6	3.1/2.2	2.2
SnOH3-me-2h2o L (70q)	A1-46q trained	48/55	100/100	3740/1157	8.6/9.0/12.8	8.6/54.8	8.6
SnOH3-me-2h2o L (70q)	A1-70q trained	60/1255	160/7296	393/29	7.7/54.8/32.4	54.8/7.7	7.7
SnOH3-me-2h2o L (70q)	A5-46q trained	44/68	96/84	5651/2579	7.4/9.0/9.6	54.8/7.2	7.4
SnOH3-me-2h2o L (70q)	A5-70q trained	134/160	128/132	4941/11792	12.0/7.9/21.5	12.0/54.8	9.6
Untrained controls							
SnOH3-me-h2o M (46q)	A1-46q untrained	11979/10835	19458/17473	177/55	22.9/22.8/22.8	22.5/22.8	5.0
SnOH3-me-h2o M (46q)	A1-70q untrained	14901/6614	27111/12183	181/123	22.7/22.7/22.9	22.9/22.9	4.9
SnOH3-me-h2o M (46q)	A5-70q untrained	13777/349	24044/1180	148/263	22.6/22.7/23.0	23.0/22.8	5.1
SnOH3-me LL (56q)	A1-70q untrained	89/12313	152/23965	12522/1	12.3/204.9/35.5	12.3/204.4	12.3

gates, 1,525 valid shots) and A5-70q 13.2 mHa (depth 79), but the A1-70q untrained control's best circuit reaches 12.3 mHa (as-submitted depth 89, 12,522 valid shots), so the trained-versus-untrained separation seen at h2o-M and L is not present at this target; this is one pair of runs. Coverage of these circuits is in Table A18.

K LUCJ baseline: transpilation and hardware benchmark

Circuit cost and hardware results of the fixed LUCJ ansatz that serves as the classical-ansatz comparator of Section 3.4.

L Dataset-level numbers

The deposit accompanying this paper (Zenodo, Section P) consists of two self-contained HTML reports, one for each training configuration. For every run they give the ten sampled circuits as gate-token sequences with their depth and two-qubit-gate count, the valid-shot, Hartree–Fock, HOMO→LUMO and spin-resolved shares of the measured shots, and the five recovery energies, together with the timing of the campaign. The raw counts and the transpiled circuits are not part of the deposit.

Table A22: LUCJ-ansatz hardware benchmark, `optimization_level=1`, energy and pair-counted subspace excitation (same convention as Table A17/Table A20: counts of determinant pairs in the recovered $n \times n$ Cartesian product, not out of the 100,000 physical shots—see Table A23 for that). "SI" = singles-injected error. All 9 target molecules across the hardware benchmark campaign.

Target	Energy (E_h) / ΔE (mHa)	ΔE SI (mHa)	S	D	T	Higher
S (32q)	−479.193374 / 0.0076	0.0076	96	3096	39776	3269431
h2o-S (38q)	−555.108826 / 0.0069	0.0069	120	4860	79278	18864350
h2o-M (46q)	−555.102419 / 22.7	5.2	4	42	410	399999543
2h2o-S (44q)	−631.083269 / 0.04	0.04	144	7020	138720	53363340
LL (56q)	−479.110941 / 204.1	41.2	10	145	2152	399997692
2h2o-M (58q)	−631.077480 / 41.1	9.6	2	21	220	399999756
h2o-LL (66q)	−555.102002 / 35.4	8.0	2	25	388	399999584
2h2o-L (70q)	−631.077480 / 54.8	12.6	0	0	26	399999973
h2o-L (56q)	−555.102002 / 33.1	7.3	2	45	862	399999090

Table A23: LUCJ-ansatz hardware benchmark, `optimization_level=1`: raw-shot excitation distributions (evaluated directly out of the 100,000 physical sampled shots from each circuit's measured counts, with the same convention as Table A21) alongside transpiled circuit metrics (transpiled at optimisation level 1). "n" is the total valid sampled-shot count attributed to a named excitation rank; S (32q) has n<100,000 because 13 sampled shots were non-particle-number- conserving readout errors, excluded from every named bucket (confirmed directly in the raw histogram). All 9 target molecules across the benchmark campaign.

Target	HF	S	D	T	Q	Higher	Depth	2q
S (32q)	6	97	672	2694	7922	88609	251	1368
h2o-S (38q)	0	0	3	5	124	99868	293	1872
h2o-M (46q)	0	0	0	0	1	99999	356	3040
2h2o-S (44q)	0	0	0	0	4	99996	337	2448
LL (56q)	0	0	0	0	0	100000	428	4248
2h2o-M (58q)	0	0	0	0	0	100000	446	4912
h2o-LL (66q)	0	0	0	0	0	100000	506	5880
2h2o-L (70q)	0	0	0	0	0	100000	536	6952
h2o-L (56q)	0	0	0	0	1	99999	430	4360

The tables of this appendix are the numbers most often needed to reuse the data: Tables A15–A21 (energies, depth, valid shots, excitation), Table A18 (coverage), the cost tables of Section 3.9 and, once fetched, the `ibm_kobe` calibration tables. Circuits were generated and transpiled with Qiskit [26].

L.1 `ibm_kobe` calibration during the experiments

Median over the device's qubits (or couplers) of the calibration IBM reported as in force at the stated time (UTC), from the IBM Quantum backend properties (fetched 2026-10-09 09:04 UTC). T_1/T_2 are medians; the last column is the worst two-qubit error on the device. The first table samples the device every few hours over the last days of training and inference; the second gives the calibration in force when each fresh-shot hardware sampling job was created. IBM recalibrates several times a day, so a job's actual calibration is the snapshot in force at its execution time. A worst two-qubit error of 100% means that IBM lists some couplers with error 1 (unusable; 10 of the 352 listed two-qubit gate entries, both directions, in the first snapshot); the median is unaffected.

Table A25: `ibm_kobe` calibration in force at the creation of each fresh-shot sampling job (run, job creation time UTC). RO: readout error; 1q: single-qubit `sx` error; 2q: median two-qubit `cz` error; worst 2q: largest two-qubit error on the device.

Run, job created (UTC)	T_1 (μs)	T_2 (μs)	RO (%)	1q (%)	2q (%)	worst 2q (%)
h2o m no23 a5-46q trained, 2026-10-08 11:01	218	105	0.62	0.030	0.20	100.0
h2o m no23 a5-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
sno32 a5-46q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
sno32 a1-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
2h2o s22 a5-46q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
sno32 a5-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o s19 a1-46q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o m no23 a1-46q untrained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0

Run, job created (UTC)	T_1 (μs)	T_2 (μs)	RO (%)	1q (%)	2q (%)	worst 2q (%)
2h2o s22 a1-46q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o m no23 a1-46q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o m no23 a1-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o ll33 a5-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
2h2o l35 a5-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o s19 a1-46q untrained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o m no23 a5-70q untrained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
llno28 a5-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o m no23 a1-70q untrained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
h2o ll33 a1-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
2h2o l35 a1-70q trained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
llno28 a1-70q untrained, 2026-10-08 11:02	218	105	0.62	0.030	0.20	100.0
2h2o m no29 a1-70q trained, 2026-10-08 11:03	218	105	0.62	0.030	0.20	100.0
L46 a1-46q untrained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a1-46q trained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a5-46q trained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a1-70q trained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a1-70q untrained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a5-70q trained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a5-46q untrained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0
L46 a5-70q untrained, 2026-10-08 12:34	218	105	0.62	0.030	0.20	100.0

M Computational cost: complementary figures

The main-text cost section (Section 3.9, Table 10) gives the quantum resources. The training budget per lineage of the main (HF-penalty) recipe is collected here; the inference timing is in Table 11.

Table A24: `ibm_kobe` calibration snapshots over the last days of training and inference.

Snapshot (UTC)	T_1 (μs)	T_2 (μs)	readout error (%)	1q (`sx`) error (%)	2q (`cz`) error (%)	worst 2q error (%)
2026-10-05 09:04	226	120	0.61	0.026	0.21	100.0
2026-10-05 15:04	226	120	0.66	0.026	0.21	100.0
2026-10-05 21:04	226	120	0.59	0.026	0.21	100.0
2026-10-06 03:04	188	113	0.63	0.029	0.21	100.0
2026-10-06 09:04	188	113	0.66	0.029	0.22	100.0
2026-10-06 15:04	188	113	0.66	0.029	0.22	100.0
2026-10-06 21:04	188	113	0.62	0.029	0.22	100.0
2026-10-07 03:04	188	113	0.62	0.029	0.22	100.0
2026-10-07 09:04	183	104	0.61	0.028	0.20	100.0
2026-10-07 15:04	183	104	0.68	0.028	0.20	100.0
2026-10-07 21:04	183	104	0.62	0.028	0.20	100.0
2026-10-08 03:04	183	104	0.62	0.028	0.20	100.0
2026-10-08 09:04	218	105	0.62	0.030	0.20	100.0
2026-10-08 15:04	218	105	0.60	0.030	0.20	100.0
2026-10-08 21:04	218	105	0.63	0.030	0.20	100.0
2026-10-09 03:04	218	105	0.64	0.030	0.20	100.0
2026-10-09 09:04	184	101	0.59	0.029	0.20	100.0

Table A26: Stage-2 training budget per lineage of the HF-penalty recipe (150 iterations each; shots = iterations × 20 circuits × 10,000 shots, one Runtime job per iteration; wall times per lineage are in Table A7, IBM-billed QPU time in Table A8). The inference timing and cost of the two configurations is in Table 11.

Lineage	Stage-2 iterations	Circuits	Shots	Runtime jobs
A1-46q	150	3,000	30,000,000	150
A5-46q	150	3,000	30,000,000	150
A1-70q	150	3,000	30,000,000	150
A5-70q	150	3,000	30,000,000	150
Total	600	12,000	120,000,000	600

N Repeated sampling of two test cases

In two of the 88 runs the ten sampled circuits contained a single circuit with valid shots, so the ranking found no best and second pair: the Hartree–Fock-penalty run of SnOH3-me LL (56 qubits) with A5-46q (one circuit, 5,772 valid shots) and the tuned-reward run of SnOH3-me-2h2o S (44 qubits) with A5-70q (one circuit, 28,055 valid shots). In both, the other nine circuits carried a diagonal-Coulomb layer and had a transpiled depth of several thousand; their valid-shot shares are near zero (Section 3.8). The number of usable circuits depends on the draw of the policy: across the 88 runs, circuits without a diagonal-Coulomb operator are shallow in every case, and the share of such circuits differs between draws. We therefore repeated the draw of the ten gate sequences once for each of the two test cases with a different random seed (4321, 5678 and 8765 were drawn without hardware; the draw with the most circuits free of diagonal-Coulomb operators, seed 8765, was measured). Sampling, ranking and recovery are otherwise identical to the main campaign.

Table A27: Repeated draws for the two test cases with a single usable circuit. Errors are those of the recovered energy against DMRG (mHa). Usable: circuits with valid shots at the ranking step.

Test case	Draw	Usable	Best	Second	Merged	Reported
Tuned reward, SnOH3-me-2h2o S, A5-70q	seed 1234	1 of 10	–	–	–	no
	seed 8765	5 of 10	0.07	0.06	0.04	yes
HF penalty, SnOH3-me LL, A5-46q	seed 1234	1 of 10	–	–	–	no
	seed 8765	8 of 10	203	199	205	no

For the tuned-reward test case the second draw gives five usable circuits, with a best-circuit error of 0.07 mHa, and this draw is used in all tables of the main text. For the Hartree–Fock-penalty test case the second draw gives eight usable circuits but a recovered error of about 200 mHa for every selection (best, second, merged and the combined energy and distinct-bitstring picks), at no more than 188 valid shots per circuit. The same plateau of about 200 mHa appears at this molecule for the second circuits of the 70-qubit policies, and the 46-qubit A1 policy reaches 15 mHa. A policy trained at 46 qubits applied to a 56-qubit active space is a possible reason, but one further draw does not establish it, and the test case is excluded from the Hartree–Fock-penalty statistics (43 of 44 runs). Both outcomes show that the result is sensitive to the draw of the sequences. They do not establish that a policy has to be applied at the active-space size it was trained for (Section 4 finds no clear size dependence).

O Noiseless matrix-product-state sampling of three earlier checkpoints

This appendix reports a noiseless simulation of circuits generated by three 46-qubit checkpoints, labelled A1-v5, A5-v5 and A5-v6. These checkpoints generate eight operator tokens per circuit, whereas the 46-qubit lineages of the main text use $L = 28$. The simulated circuits are therefore different from those run on `ibm_kobe`. The results describe these checkpoints and are not a noiseless counterpart of the hardware runs.

Protocol. We evaluated the three checkpoints on five active spaces: the three S tiers (32, 38 and 44 qubits), $Sn(OH)_3$-Me L (46 qubits) and $Sn(OH)_3$-Me·H_2O M (46 qubits). For each checkpoint and active space we generated ten circuits with $\beta = 0.2$ and sampled each circuit with 10^5 shots. The circuits were simulated as matrix product states with cuTensorNet, part of the cuQuantum SDK [41]. We used a maximum bond dimension of $\chi = 512$, complex128 precision and a relative singular-value cutoff of 10^{-8}. No noise was added. The generation and recovery seed was 1234, and the measurement seeds were 1234–1243. Ranking and recovery used the budgets of Table 5: at most 200 strings per spin for ranking, and a target of 20,000 strings per spin with 25 iterations for recovery. Unlike on the hardware, we recovered all ten circuits of each test case, so that the lowest error among the ten circuits after recovery is known. Merged pools the shots of all ten circuits (10^6 shots).

Two kinds of circuits. Table A28 lists the errors against DMRG. Of the 150 circuits, 30 contain a global orbital-rotation operator, two in every test case. These returned the Hartree–Fock determinant at most once in 10^5 shots and produced 73,000–100,000 distinct configurations. The other 120 circuits returned the Hartree–Fock determinant in at least 90% of their shots and produced at most 329 distinct configurations. Of these, 4–7 per test case returned only the Hartree–Fock determinant.

Small active spaces. At the three S tiers, each orbital-rotation circuit alone covers 92–98 % of the α-strings and recovers to 0.25–1.97 mHa. The near-Hartree–Fock circuits recover only to 3.6–84 mHa. Merging the ten circuits gives errors of 0.008–0.049 mHa and covers 99.58–99.95 % of the α-strings. This agrees with the counting in Section 3.3: at these tiers the 20,000-string cap exceeds the full α-string dimension (1,820, 3,876 and 7,315 strings). The merged energies also agree with the hardware merged energies in Table 6 to within a few μHa, although the circuits are different.

46-qubit active spaces. At the two 46-qubit tiers, almost every shot of an orbital-rotation circuit is a distinct configuration. The recovered subspaces reach the cap of 20,000 strings, which is 8.2 % of the α-strings for $Sn(OH)_3$-Me L and 4.1 % for $Sn(OH)_3$-Me·H_2O M. The lowest error among the ten circuits is 73–173 mHa at $Sn(OH)_3$-Me L and 14–22 mHa at $Sn(OH)_3$-Me·H_2O M. The merged results are 43–149 mHa at $Sn(OH)_3$-Me L and 13–21 mHa at $Sn(OH)_3$-Me·H_2O M. Merging improves on the lowest of the ten circuits in five of the six cases, but not for A5-v5 at $Sn(OH)_3$-Me·H_2O M (14.5 against 14.0 mHa), where the cap already binds.

Comparison with the hardware runs. On the same two active spaces the hardware runs of the main text are much more accurate: 5.6–14.3 mHa at $Sn(OH)_3$-Me L and 0.71–1.5 mHa at $Sn(OH)_3$-Me·H_2O M (Table 6), against 73–173 and 14–22 mHa for the lowest of the ten noiseless circuits, and 30–41 and 4.9–5.2 mHa even after singles injection. The circuits differ (eight tokens against 28), so this is not a controlled comparison. It does show that sampling a comparable trained policy without noise does not give the accuracy of the hardware campaign. The noiseless circuits with a global orbital rotation spread their shots over so many configurations that the recovered subspace reaches only 4–8 % of the α-strings, and the other circuits return the Hartree–Fock determinant. On the device the shots of the selected shallow circuits are spread by noise into single and double excitations (Table A21), which the recovery exploits. The accuracy of the hardware runs is therefore not an artefact of the ranking and recovery steps, which were identical here, and the device noise contributes to the result instead of only degrading it.

Ranking. The ranking placed the two orbital-rotation circuits last in 13 of the 15 test cases. With at most 200 strings per spin, their ranking energies at the S tiers lie 1.0–2.3 E_h above DMRG, while those of the near-Hartree–Fock circuits lie only 3–84 mHa above it. The ranking therefore favours circuits whose samples are concentrated on a few configurations. As a result, the best circuit by ranking was not the lowest of the ten after recovery in 11 of the 15 test cases. On the hardware, noise spreads the samples of every circuit, so this result may not carry over to the runs in the main text, where the selected circuit recovered to within 15 mHa in 42 of 43 runs (Section 3.3).

Table A28: Deviation $\Delta E = E_{\mathrm{SQD}} - E_{\mathrm{DMRG}}$ (mHa) of the noiseless matrix-product-state sampling of the three earlier checkpoints. Best and Second are the recovered deviations of the first- and second-ranked circuits by the pre-recovery ranking energy (one-shot diagonalization with at most 200 strings per spin sector, as in Section 2.1). Where the two entries are equal ($Sn(OH)_3$-Me S, A5-v5 and A5-v6), both circuits recovered to the same energy. Lowest of ten is the lowest recovered deviation among all ten circuits, before SI; it cannot be known without recovering every circuit and is not a selection. Merged pools all ten circuits. SI augments the recovered α and β string sets with strings associated with single excitations of the Hartree–Fock determinant and diagonalizes again. The 20,000-string cap applies before SI. HF-only is the number of circuits, out of ten, whose shots contain only the Hartree–Fock determinant. Coverage is the share of the α-strings in the merged recovered set. The circuits differ from those run on `ibm_kobe`.

Checkpoint	Best	Second	Lowest of ten	Merged	Best + SI	Merged + SI	HF-only	Coverage (%)
$Sn(OH)_3$-Me S (32 qubits, 16o24e)								
A1-v5	28.956	78.702	0.505	0.008	16.050	0.008	4	99.89
A5-v5	82.595	82.595	0.398	0.008	19.121	0.008	5	99.95
A5-v6	82.595	82.595	0.501	0.008	19.121	0.008	6	99.89
$Sn(OH)_3$-Me·H_2O S (38 qubits, 19o30e)								
A1-v5	3.629	6.425	0.248	0.008	1.226	0.008	6	99.85
A5-v5	3.629	6.425	0.255	0.008	1.226	0.008	5	99.85
A5-v6	6.357	6.425	0.259	0.008	1.476	0.008	6	99.85
$Sn(OH)_3$-Me·$2H_2O$ S (44 qubits, 22o36e)								
A1-v5	4.243	5.563	0.449	0.049	1.063	0.049	6	99.58
A5-v5	4.243	5.563	0.380	0.045	1.063	0.045	6	99.60
A5-v6	5.563	5.832	0.379	0.045	1.204	0.045	7	99.60
$Sn(OH)_3$-Me L (46 qubits, 23o32e)								
A1-v5	73.440	199.536	73.440	43.222	38.537	20.266	6	8.16
A5-v5	199.536	203.430	173.048	148.614	41.332	26.471	6	8.16
A5-v6	169.534	199.536	169.534	142.858	30.841	27.022	6	8.16
$Sn(OH)_3$-Me·H_2O M (46 qubits, 23o30e)								
A1-v5	14.048	20.022	14.048	12.987	4.932	3.746	6	4.08
A5-v5	14.048	22.750	14.048	14.488	4.932	4.147	5	4.08
A5-v6	22.682	22.750	21.771	21.199	5.235	4.374	6	4.08

Singles injection. Adding the single excitations of the Hartree–Fock determinant (SI) to the recovered set of the best circuit by ranking and diagonalizing again lowers its error, for example from 29.0 to 16.1 mHa at $Sn(OH)_3$-Me S and from 73.4 to 38.5 mHa at $Sn(OH)_3$-Me L for A1-v5. No noise is present, so this gain comes from the larger basis and not from the removal of noise.

Limitations. Only one generation seed was used, and convergence beyond 25 recovery iterations was not checked. The MPS truncation error was not recorded. One of the 1.5×10^7 shots had a wrong number of alpha or beta electrons. The circuits conserve both numbers, so this cannot occur in an exact simulation. The shot came from an orbital-rotation circuit at 44 qubits and may reflect MPS truncation.

Ranking and recovery build their bases differently. Ranking performs spin-symmetry completion. In eight test cases, the recovered basis of the best circuit by ranking omitted strings added during ranking, and its energy was 0.6–12 mHa above the ranking energy. Independent PySCF calculations [42] on the same bases reproduced these increases. In three of the eight affected test cases, all at 46 qubits, the affected circuit was also the lowest of the ten after recovery.

P Data and code availability

The data of this work are available on Zenodo at `https://zenodo.org/records/24069330`. The record contains two self-contained HTML reports, one for each training configuration: the Hartree–Fock penalty (the main configuration) and the tuned reward. Each report gives the training of the configuration, the table of every run, the chemistry of the generated circuits on unseen molecules (valid share, Hartree–Fock, HOMO→LUMO and spin-resolved shares, operator-family statistics) set against the chemistry during training, the comparison of the two configurations, the

timing and cost, and, for every run, the ten generated gate-token sequences with their depth, two-qubit-gate count and recovered energies. The code is available on GitHub at https://github.com/kunalroy332/sqd-gqe.

References

[1] S. McArdle, S. Endo, A. Aspuru-Guzik, S. C. Benjamin, and X. Yuan, Reviews of Modern Physics **92**, 015003 (2020).

[2] S. R. White, Physical Review Letters **69**, 2863 (1992).

[3] S. Bravyi, J. M. Gambetta, A. Mezzacapo, and K. Temme, Tapering off qubits to simulate fermionic Hamiltonians (2017), arXiv:1701.08213 .

[4] A. Kandala, A. Mezzacapo, K. Temme, M. Takita, M. Brink, J. M. Chow, and J. M. Gambetta, Nature **549**, 242 (2017).

[5] Y. Kim, A. Eddins, S. Anand, K. X. Wei, E. van den Berg, S. Rosenblatt, H. Nayfeh, Y. Wu, M. Zaletel, K. Temme, and A. Kandala, Nature **618**, 500 (2023).

[6] K. Kanno et al., Quantum-selected configuration interaction: classical diagonalization of Hamiltonians in subspaces selected by quantum computers (2023), arXiv:2302.11320 .

[7] J. Robledo-Moreno, M. Motta, H. Haas, A. Javadi-Abhari, P. Jurcevic, W. Kirby, S. Martiel, K. Sharma, S. Sharma, T. Shirakawa, I. Sitdikov, R.-Y. Sun, K. J. Sung, M. Takita, M. C. Tran, S. Yunoki, and A. Mezzacapo, Science Advances **11**, eadu9991 (2025), arXiv:2405.05068 .

[8] K. Kumar, E. Lötstedt, Y. Kobayashi, and S. Yunoki, Vibrational sample-based quantum diagonalization with unitary-cluster-Jastrow ansätze on IBM QPUs (2026), arXiv:2610.09532 .

[9] P. J. Ollitrault, A. Kandala, C.-F. Chen, P. K. Barkoutsos, N. M. Meana, M. Pistoia, S. Blanco, K. M. Eckstein, and I. Tavernelli, Chemical Science **11**, 10642 (2020).

[10] K. Kumar, T. Shirakawa, B. Kim, R. Wakizaka, I. Hamamura, M. Sato, S. Yunoki, M. Motta, H. Horii, and Y. Kobayashi, Broken-symmetry reference states for sample-based quantum diagonalization with near-DMRG accuracy (2026), arXiv:2611.04496 .

[11] M. Motta et al., Chemical Science **14**, 11213 (2023).

[12] K. Nakaji, L. B. Kristensen, R. Kemmoku, J. A. Campos-Gonzalez-Angulo, M. G. Vakili, H. Huang, M. Bagherimehrab, C. Gorgulla, F. Wong, A. McCaskey, J.-S. Kim, T. Nguyen, P. Rao, Q. Gao, M. Sugawara, N. Yamamoto, and A. Aspuru-Guzik, The generative quantum eigensolver (GQE) and its application for ground state search (2024), arXiv:2401.09253 .

[13] R. Kemmoku, Q. Gao, S. Kanno, K. Keithley, I. Hamamura, N. Yamamoto, and K. Nakaji, Generative circuit design for quantum-selected configuration interaction (2026), arXiv:2604.09756 .

[14] K. Elgammal and M. Maußner, Scaling a CUDA-Q GQE + QSCI pipeline to 40 qubits for EUV photoresist chemistry (2026), arXiv:2607.23988 [quant-ph] .

[15] J. H. Ma, C. Needham, H. Wang, A. Neureuther, D. Prendergast, and P. Naulleau, ACS Applied Materials & Interfaces **14**, 5514 (2022).

[16] T. Du, X. Yang, Y. Zhao, P. Han, J. Zhao, and S. Zhou, Journal of Chemical Physics **160**, 154307 (2024).

[17] J. Haitjema, Y. Zhang, M. Vockenhuber, D. Kazazis, Y. Ekinci, and A. M. Brouwer, Journal of Micro/Nanolithography, MEMS, and MOEMS **16**, 033510 (2017).

[18] R. T. Frederick, S. Saha, J. T. Diulus, F. Luo, J. M. Amador, M. Li, D.-H. Park, E. L. Garfunkel, D. A. Keszler, and G. S. Herman, Microelectronic Engineering **205**, 26 (2019).

[19] J. T. Diulus, R. T. Frederick, D. C. Hutchison, I. Lyubinetsky, R. Addou, M. Nyman, and G. S. Herman, ACS Applied Nano Materials **3**, 2266 (2020).

[20] S. Saha, D.-H. Park, D. C. Hutchison, M. R. Olsen, L. N. Zakharov, D. Marsh, S. Goberna-Ferrón, R. T. Frederick, J. T. Diulus, N. Kenane, et al., Angewandte Chemie International Edition **56**, 10274 (2017).

[21] M. C. Sharps, D. A. Marsh, L. N. Zakharov, J. E. Hutchison, and D. W. Johnson, Crystal Research and Technology **52**, 1700081 (2017).

[22] J. Chalabala, O. Dvořák, and P. Slavíček, Chemical Physics **515**, 221 (2018).

[23] Y. J. Patel, A. Kundu, M. Ostaszewski, X. Bonet-Monroig, V. Dunjko, and O. Danaci, in International Conference on Learning Representations (ICLR) (2024) arXiv:2402.03500 .

[24] Z. Shao, P. Wang, Q. Zhu, R. Guo, J. Ni, H. Li, M. Lu, Y. Wu, W. Xiong, H. Liu, and D. Xie, DeepSeekMath: Pushing the limits of mathematical reasoning in open language models (2024), arXiv:2402.03300 .

[25] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, Proximal policy optimization algorithms (2017), arXiv:1707.06347 .

[26] A. Javadi-Abhari et al., Quantum computing with Qiskit (2024), arXiv:2405.08810 .

[27] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, Language Models are Unsupervised Multitask Learners, Technical Report (OpenAI, 2019).

[28] IBM Research, pos-egnn: position-based equivariant graph neural network foundation model for chemistry and materials, Part of IBM Foundation Models for Materials (FM4M). Model release (Apache-2.0), `https://huggingface.co/ibm-research/materials.pos-egnn`, source `https://github.com/IBM/materials/tree/main/models/pos_egnn` (2025), no separate publication accompanies the release.

[29] B. Deng et al., Nature Machine Intelligence **5**, 1031 (2023), source of the MPtrj data set.

[30] M. J. Frisch et al., Gaussian 16, Gaussian, Inc., Wallingford, CT (2016).

[31] C. Møller and M. S. Plesset, Physical Review **46**, 618 (1934).

[32] F. Weigend and R. Ahlrichs, Physical Chemistry Chemical Physics **7**, 3297 (2005).

[33] P.-O. Löwdin, Physical Review **97**, 1474 (1955).

[34] K. J. Sung, I. Choi, M. Amico, B. Andrews, E. Ayantuna, Y. Kawashima, W.-H. Lin, D. Omanovic, S. Piccinelli, J. R. Moreno, A. A. Saki, J. Shee, S. Shin, M. C. Tran, K. Ueda, H. Zhang, and M. Motta, ffsim: Faster simulation of fermionic quantum circuits (2026), arXiv:2605.03123 [quant-ph] .

[35] R. Jozsa, Journal of Modern Optics **41**, 2315 (1994).

[36] M. A. Nielsen and I. L. Chuang, Quantum Computation and Quantum Information, 10th ed. (Cambridge University Press, Cambridge, 2010).

[37] E. Hellinger, Journal für die reine und angewandte Mathematik **136**, 210 (1909).

[38] A. Bhattacharyya, Bulletin of the Calcutta Mathematical Society **35**, 99 (1943).

[39] J. Lin, IEEE Transactions on Information Theory **37**, 145 (1991).

[40] C. Lanczos, Journal of Research of the National Bureau of Standards **45**, 255 (1950).

[41] The cuQuantum Development Team, NVIDIA cuQuantum SDK, Version v24.11.0, `https://github.com/nvidia/cuquantum` (2024).

[42] Q. Sun, T. C. Berkelbach, N. S. Blunt, G. H. Booth, S. Guo, Z. Li, J. Liu, J. D. McClain, E. R. Sayfutyarova, S. Sharma, S. Wouters, and G. K.-L. Chan, WIREs Computational Molecular Science **8**, e1340 (2018).