arXiv:2611.14174v1 [cs.AI] 10 Nov 2026

Cross-Border Fulfillment Center Selection via Doubly Robust Estimation and Linear Programming

Kriti Mahajan [1] **Patrick Bloebaum** [2]

Abstract

We introduce the cross-border (XB) fulfillment center (FC) Selection Model: a multi-stage optimization framework combining forecasting, doubly robust estimation, and linear programming to select optimal source FCs. For each source–destination shipping arc the objective is to choose a subset of source-country FCs that minimize delivery time and cost while maximizing profit and shipped units under a fixed domestic network. We evaluate the model on three production-scale international shipping corridors for a large e-commerce company using 12 months of shipment data, employing both aggregate offline replay and shipment-level counterfactual nearest-neighbor matching. Across all shipping corridors, the XB FC Selection Model achieves 8.5–51.1% reductions in delivery time and 4.4–48.3% reductions in logistics cost relative to human expert-defined FC selection policies, with up to 6.5 percentage point improvements in 5/5 target compliance (5 days, $5 USD). Pareto analysis shows the model consistently dominates single-objective ranking strategies and unconstrained LP on the efficiency–reallocation frontier, with all results statistically significant.

1. Introduction

Cross-border (XB) e-commerce has grown rapidly, but unlike domestic operations, XB fulfillment must contend with customs regulations, volatile international freight conditions, and longer transit times while meeting rising customer expectations. Delivery speed and cost efficiency are critical to customer satisfaction and competitive performance, as long transit times and high delivery costs reduce conversion, retention, and repeat purchases (Deshpande & Pendem, 2020) and inhibit cross-border growth (Li, 2022). Therefore, effective cost control and lead-time reduction through well-designed logistics networks are essential for sustaining growth and competitiveness in XB e-commerce.

Given the above, a core challenge in XB order fulfillment is selecting source fulfillment centers (FCs)—origin warehouses that store, pick, pack, and dispatch orders for international customers. This FC selection is a key component of regionalization, which structures a fulfillment network into geographic regions so that customer orders are primarily served by nearby FCs, reducing distance, cost, and delivery time. Unlike domestic fulfillment, where FCs can be newly sited or networks redesigned, XB operations typically operate within existing source-country infrastructure and transportation networks; the XB division cannot create new FCs or routes and must use the network as given. FC selection for each cross-border corridor therefore becomes a constrained subset selection problem. Moreover, the optimal FC subset for one source–destination corridor (e.g., US>Australia (AU)) may not generalize to another (e.g., US>Singapore (SG)) due to differing geographic, regulatory, and demand patterns. Selecting an inappropriate subset can consequently increase transit times and logistics costs thereby degrading customer experience.

To address this challenge, we introduce the XB FC Selection Model, a data-driven framework integrating short-horizon demand forecasting, doubly robust estimation (DRE), and constrained linear programming to select source FC subsets for each XB corridor. Forecasting provides next-period operational metrics, DRE corrects selection bias in historical FC assignments and generates candidate FCs, and a multi-objective LP selects the final subset while controlling shipped-volume reallocation. The model minimizes delivery time and logistics cost while maximizing profit and shipped volume, implementing soft regionalization: orders are preferentially fulfilled by selected (in-region) FCs, while non-selected (out-of-region) FCs remain available for residual demand. Although expert-based FC selection performs competitively (Section 4), its manual process is difficult to scale across growing FC networks, corridors, and decision periods. The proposed ML framework by contrast enables scalable automated, forward-looking, multi-objective optimization

[1]Cross Border Tech & Science, Amazon [2]Agentic AI DevOps, Amazon. Correspondence to: Kriti Mahajan <kritimhj@amazon.com>.

International Conference on Machine Learning (Expo), Seoul, South Korea. PMLR 306, 2026. Copyright 2026 by the author(s).

beyond rank-based approaches. Existing ML-based regionalization methods for domestic fulfillment do not directly apply because XB operations must use a fixed source-country network without creating new FCs or routes. Offline evaluation on production-scale data shows 8.5–51.1% lower delivery time, 4.4–48.3% lower logistics cost, and up to 6.5 percentage-point higher 5/5 compliance versus expert policies. Ablation analysis (Sections 3–4) confirms that forecasting, DRE-based FC selection, and constrained LP optimization are all necessary for the model's observed Pareto dominance over comparable baselines.

Related Work: Our work sits at the intersection of supply chain optimization (SCO) and FC location choice, extending the limited literature on cross-border (XB) e-commerce logistics. E-commerce fulfillment has been extensively studied (Acimovic & Graves, 2015),(Jasin et al., 2019). Prior work has studied large-scale e-commerce fulfillment optimization, including dynamic order fulfillment, transportation-network design, and domestic regionalization (Sinha et al., 2026), (O'Neill, 2023), (Xu et al., 2009), (Lara et al., 2023), (Amazon Science, 2022),(Kassmann & Allgor, 2006),(Allgor et al., 2023). Cross-border logistics introduces additional regulatory and transportation complexity, and prior XB studies have considered overseas warehouse planning, logistics-mode choice, and network design (Cecil, 2024), (Guan et al., 2021).However, these approaches do not address the problem studied here: selecting a subset of existing source-country FCs for individual XB corridors under a fixed transportation network. Our approach combines causal estimation with forecasting and constrained optimization to address this production-scale FC subset-selection problem. The closest work in the XB FC selection literature is (Guan et al., 2021) proposes a game-theoretic framework for overseas FC siting; our paper differs in both theory and setup, using causal estimation and LP to select optimal FCs within the existing source-country network rather than creating new overseas facilities.

2. XB FC Selection Model Architecture

The XB FC Selection Model consists of three components: Short-horizon forecasting, Causal FC candidate generation, Constrained optimization.

2.1 Problem Set-Up and Notation : Let $\mathcal{A} = \{(s, d)\}$ denote a set of source country s to destination country d shipping corridors; $\mathcal{W}_s$ denote the set of fulfillment centers located in source country s ; $t \in \mathbb{N}$ index discrete decision periods (here days). Let $i \in \mathcal{W}_s$ index fulfillment centers. For each arc a, FC i, and time period t, we observe: $T_{i,t,a} \in \{0, 1\}$ - indicator if FC i was active for a on t, $Y_{i,t,a} \in \mathbb{R}$ - observed mean FC profit, and $X_{i,t,a} \in \mathbb{R}^d$ - the observed FC covariates.

Given that cross-border FC selection must balance multiple competing objectives, the observed covariates are: $\tau_{i,t,a}$:= Mean Delivery time from FC i to destination d, $c_{i,t,a}$:= Mean Shipping Cost from FC i to destination d, $u_{i,t,a}$:= Total Shipped Units from FC i to destination d. $\{\tau, c, u, Y\}$ define the origin-destination (OD) flow values for FC i on arc a for each time period t. Note that delivery time is measured from source FC to destination country delivery pin code. Ship cost includes all costs across all types of transportation , taxes etc. Shipped Units is included to make sure that sufficient shipment volume is handled by the chosen FCs and is a proxy for product mix diversity. [1] For each arc $a = (s, d) \in \mathcal{A}$ and time t, the XB FC selection model's task is to select a subset of FCs in s at time $t + 1$ that will minimize delivery time and cost while maximizing shipped units and profit i.e. the model's task is to select $\mathcal{W}_{a,t} \subseteq \mathcal{W}_s$ of fulfillment centers. The FC selection has to be optimized in a forward looking fashion because due to technical constraints of the deployment platform, real (or near-real) time deployment of selected FCs is not possible. Thus, the model has to make a choice that is optimal for the next period $t+1$, given the information it has in the current period t.

The XB FC Selection Model is designed to achieve "5/5" through soft regionalization, optimizing FC-level profit in both the DRE and LP stages. Profit serves as a composite objective that balances cost efficiency and demand coverage: optimizing time or cost alone can concentrate volume in a few fast or cheap FCs, causing severe shipped-unit reallocation (up to 86%, Section 4) and reducing in-region coverage. By contrast, profit penalizes such concentration because lower volume reduces total profit, acting as an endogenous regularizer that favors FC subsets maximizing in-region assignment while remaining cost- and time-efficient. DRE further benefits because profit captures each FC's net causal contribution across operational dimensions.

2.2 Training Phase - Origin-Destination (OD) Forecasting: Since FC selection at time t affects outcomes at time $t + 1$, we require forward-looking estimates. The forecasts are multivariate, jointly incorporating the temporal dynamics of the OD metrics τ, c, and u. Specifically, for each OD metric $z \in \{\tau, c, u\}$, the target is the value of z at time t, while the values of all three OD metrics at time $t - 1$ are used jointly as predictors. Formally, for each arc a, we fit a multivariate forecasting model of the form $z_{a,t} = \beta_{z,\tau}\tau_{a,t-1} + \beta_{z,c}c_{a,t-1} + \beta_{z,u}u_{a,t-1} + \varepsilon_{a,t-1,z}$ where $z \in \{\tau, c, u\}$.

2.3 Training Phase - Candidate FC generation using

[1] This design is due to unavailability of FC capacity and XB inventory data, preventing direct inventory modeling. Based on stakeholder alignment, the LP assumes (1) all source-country inventory is XB-eligible and (2) XB inventory equals 2% of total FC inventory, ensuring capacity for increased volumes in selected FCs. Shipped Units serves as a proxy for product-mix diversity.

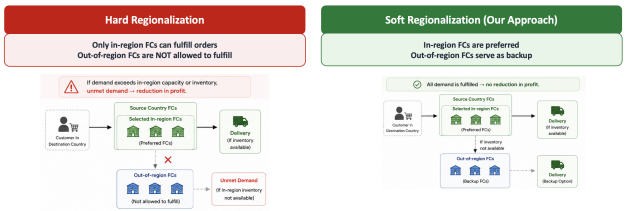

Figure 1. Hard vs. Soft Regionalization: Most businesses choose Soft Regionalization as it ensures zero demand loss while still capturing efficiency gains from better FC targeting

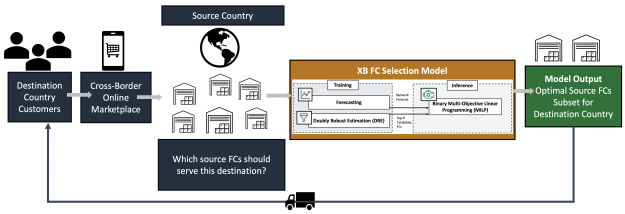

Figure 2. Cross-Border (XB) Fulfillment Center (FC) Selection Model

Algorithm 1 XB FC Selection Model

Require: Historical shipment data $\mathcal{D}$; shipping corridors $\mathcal{A}$; FC sets $\{\mathcal{W}_s\}$; hyperparameters $(P, K_{\min}, K_{\max}, \alpha)$
Ensure: Daily FC selection policy $\Pi(a, t)$
for all $a = (s, d) \in \mathcal{A}$ **do**
  Aggregate shipment data into OD metrics $\mathcal{Z}_{i,t}$ for all $i \in \mathcal{W}_s$
  Forecast next-day metrics $\widehat{\mathcal{Z}}_{i,t+1}$ using short-horizon models
  Estimate per-FC treatment effects $\widehat{\text{ATE}}_i$ via doubly robust learning
  $\mathcal{W}^{\text{cand}} \leftarrow$ top-P FCs with significant $\widehat{\text{ATE}}_i$
  Solve binary LP over $\mathcal{W}^{\text{cand}}$ using $\widehat{\mathcal{Z}}_{i,t+1}$
  $\Pi(a, t+1) \leftarrow \{i \in \mathcal{W}^{\text{cand}} : x_i = 1\}$
end for

Doubly Robust Estimation (DRE): As mentioned previously, XB orders use shared domestic infrastructure, so FC selection is influenced by prior domestic-optimized policies, most notably domestic regionalization since 2022. Historical assignments $T_{i,t}$ are therefore non-random, reflecting heuristics, constraints, and capacity limits, making naive rankings by observed speed or profitability biased. We address this using Doubly Robust Estimation (DRE) to narrow the FC candidate set to 45, the business-defined maximum number of in-region FCs under soft regionalization. DRE estimates each FC's causal impact on serving an arc despite biased and incomplete historical data, where only realized outcomes—not counterfactuals—are observed and past assignments are confounded by manual rules. This avoids systematically overvaluing frequently selected FCs and provides a more robust estimate of each FC's true contribution. DRE is used to compute the Average Treatment Effect (ATE) for each origin FC i as follows: For each arc a, we model FC selection as a treatment where: $T_{i,t} = \mathbb{I}\{\text{FC } i \text{ selected}\}$, $Y_{i,t} = \pi_{i,t}(\text{Observed profit})$ and $X_{i,t-1} =$ observed lagged confounders. The confounders are lagged at time t-1 to prevent bi-directionality and endogeneity. The observed lagged confounders are: $\{\texttt{Distance}_{t-1}, \texttt{time}_{t-1}, \texttt{price}_{t-1}, \texttt{cost}_{t-1}, \texttt{shipped_units}_{t-1}, \texttt{Regionalization}$ where Regionalization denotes the start of domestic regionalizations dummy indicator. We seek the ATE for each FC i: $\text{ATE}_i = \mathbb{E}[Y_{i,t}(1) - Y_{i,t}(0)]$. To generate the list of candidate FCs, for each arc we retain the top P fulfillment centers with statistically significant positive ATEs: $\mathcal{W}_a^{\text{cand}} = \{i : \text{ATE}_i > 0,\ p\text{-value} < \alpha\}$.

2.4 Inference Phase - Optimal top-k FC Selection via Multi-Objective Linear Programming: We model the source FC selection problem as a binary linear program, optimizing multiple objectives: profit, shipped units, cost and time. We run the below optimization routine at each time period t (for the purpose of this paper, the model is run daily) using the OD forecasted values for time $t+1$ to determine the optimal FC selection at time $t+1$. The resulting problem is a binary linear program solved daily using a MILP solver. More formally: **2.4.1 Decision Variables:** For arc a and time t, define: $x_i \in \{0, 1\}, \quad i \in \mathcal{W}_a^{\text{cand}}$ where $x_i = 1$ indicates FC i is selected. **2.4.2. Parameters**: Weights for profit, cost, and time, subject to $w_\pi + w_c + w_\tau = 1$, are tuned using Bayesian optimization (BO). The Bayesian optimization objective function maximizes a composite score combining delivery-time and cost reduction while penalizing profit reduction. For each candidate weight vector (w_π, w_c, w_τ), the resulting FC selections from the optimization formulation in Sections 2.4.3-2.4.4 are evaluated via aggregate offline replay (Section 3.3.1) on a held-out validation window preceding the evaluation period. Weights are re-estimated monthly using a rolling validation window of the most recent month's data. **2.4.3. Objective:** We optimize a weighted combination of normalized forecasted metrics with a small lexicographic bias:

$$\max_x \sum_i x_i \left(w_\pi \tilde{\pi}_{i,t+1} - w_c \tilde{c}_{i,t+1} - w_\tau \tilde{\tau}_{i,t+1} + \varepsilon b_i\right),$$

where $\tilde{\cdot}$ denotes z-score normalization, and $\varepsilon = 10^{-6} \ll 1$ is a small coefficient used to enforce deterministic tie-breaking (e.g., based on the hash or ID of FC i). **2.4.4. Constraints:**The optimization is subject to two constraints. The first limits the number of selected origin fulfillment centers (FCs):

$$K_{\min} \leq \sum_i x_i \leq K_{\max}, \tag{1}$$

where $K_{\min}$ and $K_{\max}$ denote the lower and upper bounds, respectively, on the number of selected origin FCs. The second enforces minimum shipment coverage:

$$\sum_i x_i \hat{u}_{i,t+1} \geq \alpha \sum_i \hat{u}_{i,t+1}, \tag{2}$$

where $\alpha = 0.1$ specifies the minimum fraction of total shipped units that must be covered by the selected origin FCs. The number of selected FCs (top-k) varies by shipping corridor (arc), since different destinations may require different numbers of FCs to serve demand, depending on local demand–supply dynamics.

3. Experimental Methodology

3.1 Dataset Description: We evaluate the XB FC Selection Model on three shipping corridors — US>SG (Singapore), US>AU (Australia), and US>DE (Germany) — spanning 12 months of production data. The underlying dataset consists of anonymized cross-border shipment records aggregated from order-level data to the FC i × arc a × day t level. **OD Flow Dataset Construction:** For each FC i × arc a × day t observation, we compute origin-destination (OD) flow metrics - mean delivery time (including in-transit and customs processing time), mean shipping cost (across all transport modes and duties), total shipped units, and mean FC profit.

3.2 Baselines: We compare the XB FC Selection Model against nine baselines across four categories: **Rule-based rankings**: (1) *Profit Rank*: top-k FCs by prior-period profit; (2) *Time Rank*: top-k FCs by prior-period delivery speed. **ATE-augmented rankings**: (3) *ATE-Profit Rank*: DRE candidate generation (top 45 FCs with significant positive ATEs), followed by profit-based top-k selection; (4) *ATE-Time Rank*: the same candidate generation followed by time-based top-k selection. **LP-based methods**: (5) *Vanilla LP*: top-k selection using the multi-objective LP (Section 2.4) over all FCs; (6) *Profit Rank + LP*: top 45 FCs by prior-period profit followed by LP-based top-k selection; (7) *Time Rank + LP*: top 45 FCs by prior-period delivery speed followed by LP-based top-k selection. **Composite-score heuristics**: (8) *Observed Score Top-k*: FCs are scored using previous-day observed metrics, $s_i = w_\pi \hat{\pi}_i - w_c \hat{c}i - w\tau \hat{\tau}_i + w_u \hat{u}_i$, where $\hat{\pi}_i, \hat{c}_i, \hat{\tau}i, \hat{u}i$ are within-day standardized (zero-mean, unit-variance) profit, shipping cost, transit time, and shipped volume, respectively, with fixed weights $(w\pi, w_c, w\tau, w_u) = (1, 1, 1, 0.5)$, chosen a priori and not tuned on the evaluation period. The top-k FCs are selected, followed by greedy volume-floor repair: high-volume FCs replace the lowest-scored selected FC only if selected volume strictly increases, until 10% shipment-volume coverage is reached or no improving swap exists. (9) *Forecasted Score Top-k*: identical to (8), but using forecasted z values instead of lagged metrics. Since profit is not forecast by the existing predictive-model block, we fit a one-step linear regression, $\hat{\pi}_i^f = \beta_0 + \beta_1 \pi_i^{\text{lag-1}}$, on the pre-cutoff training split and apply it at inference; all other metrics use forecasts from the main model pipeline. All strategies are benchmarked against two reference policies: (i) **the Expert baseline** — FC selection policy defined by human supply chain experts using inventory stock based pairwise FC evaluation, distance filtering, and sort-point availability checks — and (ii) a **No Regionalization baseline** in which all source-country FCs are eligible for all shipping corridors; the Expert baseline is also measured against this policy.

3.3 Evaluation Framework: XB fulfillment's order-to-FC assignment system supports only one active FC selection policy per arc, precluding within-arc randomization and conventional A/B testing. Because alternative policies require a full-volume cutover rather than controlled diversion, strong offline performance guarantees are needed before any online pilot. We therefore use a two-stage evaluation protocol: (i) aggregate offline replay on historical data, followed by (ii) shadow deployment, where candidate FC selection strategies f make daily FC selections using live data in parallel with the no-regionalization baseline, without affecting live routing. We use two offline evaluation frameworks: **3.3.1 Aggregate Offline Replay:** Each strategy f from Section 3.2, including the XB FC Selection model, makes daily FC selections using only information available at time t. We simulate the network assuming only f's selected FCs are active and compute outcomes from production data at $t + 1$. Thus, if f selects 10 of 100 FCs on a corridor, we characterize the resulting network using the aggregate mean time, cost, and shipped volume of those 10 FCs. Performance is aggregated monthly and reported relative to baseline as $\%$ change $=$ $(\text{Baseline Value} - \text{Strategy Value})/\text{Baseline Value} \times 100$. We report Time (days; larger reduction is better), Cost (thousands USD; larger reduction is better), and Shipped Units (thousands; smaller reduction is better) for both pre- and post-shadow deployment windows to assess temporal stability. Under soft regionalization, all demand is fulfilled; thus, "shipped unit reduction" reflects volume handled by non-optimal (out-of-region) FCs, not unserved demand.**3.3.2 Counterfactual Matching (only for post-shadow deployment)**: Offline evaluation cannot observe what would have occurred if a different FC had fulfilled an XB order. We therefore partition shipments into *optimal* FC orders (fulfilled by an FC in strategy f's recommended set) and *non-optimal* FC orders (fulfilled by an FC outside the set), and construct shipment-level counterfactuals using two complementary matching approaches. **Same-ASIN Matching:** For each non-optimal shipment, we match historical shipments of the same ASIN fulfilled by an optimal FC and use their realized delivery time and shipping cost as counterfactuals. This is highly sparse, however, with fewer than 1% of

daily demand having the same ASIN assigned to both optimal and non-optimal FCs.**Nearest-Neighbor Matching:** We therefore use nearest-neighbor matching as a complementary approach for at-scale counterfactual evaluation. For each non-optimal order i with standardized package dimensions $\mathbf{x}i \in \mathbb{R}^3$ and GL g_i, we find the nearest optimal order $j^* = \arg\min j \in \mathcal{O}_g^{\text{in}} |\mathbf{x}_i - \mathbf{x}_j|_2$ using a GL-stratified KD-Tree. The matched order's realized delivery time and shipping cost provide counterfactual estimates for i. These estimates are then used to compute Mean Delivery Time, Mean Shipping Cost, and % Meeting the 5/5 Target (delivery ≤ 5 days and cost $\leq \$5$ USD) for each strategy.

3.4 Implementation Details: AutoML fits the forecasting functions, with each variable/arc trained through a period ending one month before forecasting period t to prevent look-ahead bias. OD forecasting and DRE-based candidate generation run in parallel. DRE uses EconML's DR-Learner, with propensity and outcome models automatically selected from linear and gradient-boosted forest models; the final CATE model is estimated by OLS. LP optimization uses CBC via PuLP. The selected FC count ($topk$) is 11/6 for US>SG (pre/post-shadow), 8 for US>DE, and 6 for US>AU.

4. Results

4.1 Aggregate Performance: 4.1.1 Against No Regionalization Baseline. All strategies achieve ¿95% time reduction, confirming the fundamental value of FC selection (Appendix Table 3). Strategies separate sharply on cost versus volume: LP maximizes cost reductions (88–91%) but at near-total volume loss (87–90%), while the Expert policy achieves the lowest volume reduction (36–61%) with modest cost gains (46–63%). The XB FC Model balances these extremes: 51–80% cost reduction at 41–79% volume reduction, consistently improving upon Expert in cost efficiency without proportional volume loss.

4.1.2 Against Expert Baseline: The trend observed in Section 4.1.1 continues here. Relative to Expert selection, the XB FC Selection Model provides the best cross-corridor balance between efficiency and volume retention (Table 1; Figure 3). The main baselines exhibit sharper trade-offs: Single-objective rankings show the sharpest trade-offs—Profit Rank preserves volume but increases delivery time in every corridor (up to $+40.5\%$), while Time Rank achieves large time reductions for US>SG and US>DE but increases delivery time for US>AU and incurs severe volume loss (46–87%). ATE-based candidate generation yields marginal improvements over single-objective ranking counterparts but does not resolve these trade-offs. All three LP variants produce substantial time reductions for US>SG and US>DE but diverge sharply on volume retention, with Vanilla LP incurring 83–94% volume loss for US>DE and US>AU. The composite-score heuristics occupy a moderate band but increase delivery time for US>AU by 14.8–21.9%, demonstrating that static scoring rules cannot capture corridor-specific structure. By contrast, the XB FC Model is one of only three strategies that achieve delivery-time reductions across all three corridors—and it does so at substantially lower volume cost than either alternative. Specifically, the model achieves 24.4–41.8% / 27.1–51.1% / 8.5–13.3% time reductions and 4.4–8.6% / 45.9–48.3% / 32.6–36.4% cost reductions for US>SG / US>DE / US>AU, respectively, while limiting shipped-unit reallocation to -7.7% to $+2.5\%$ / 48.5–48.9% / 30.1–35.9%. These results are stable across pre- and post-shadow-deployment windows and statistically significant (Wilcoxon $p < 0.02$; paired t-test $p < 0.005$; Cohen's d: -1.03 to -3.89; see Appendix Tables 4–9).**Geographic Distribution of Selected FCs:** Expert and XB FC Model selections broadly overlap in coastal orientation for each arc (see Appendix Figure 6), but Figure 5 reveals where they diverge.

4.1.3 Pareto Trade-Off Analysis (see Figure 4). To characterize multi-objective efficiency, we visualize the trade-off between operational efficiency gains and shipped volume retention relative to the no-regionalization baseline via Pareto frontier plots. Strategies separate into four structural clusters: (1) *High-efficiency, high-volume-loss* (upper-left): Time Rank, LP variants, and ATE-Time Rank achieve the largest time and cost reductions at steep volume penalties. (2) *Low-efficiency, low-volume-loss* (lower-right): Profit Rank and ATE-Profit Rank preserve volume but offer minimal or negative time improvement. (3) *Moderate* (center): Composite-score heuristics achieve intermediate trade-offs for US>SG and US>DE but drift into the low-efficiency cluster for US>AU. (4) *Pareto-optimal* (frontier): The XB FC Model consistently occupies the Pareto frontier — the region of maximum efficiency gain per unit of volume reduction — across all three corridors. No other strategy achieves this cross-corridor consistency.

4.2 Counterfactual Evaluation: Shipment-Level Performance. Shipment-level counterfactual matching corroborates the aggregate results (Tables 2). Under both same product and nearest neighbor settings, the XB FC Model achieves the lowest or near-lowest mean delivery time and shipping cost across all three corridors, and the highest 5/5 target compliance for US>SG and US>DE. For US>AU, the results diverge: Expert leads under nearest-neighbor matching (0.93 vs. 1.13 days), but the same-ASIN estimator — which controls for product-level heterogeneity — favors XB FC on every metric (1.84 vs. 2.31 days, 8.21% vs. 6.47% 5/5 compliance). Across both matching methods, the XB FC Model consistently outperforms LP-based selection on all metrics and corridors, demonstrating the incremental value of DRE-based candidate generation alongside the multi-objective LP. Paired t-tests on matched samples (Ap-

Table 1. FC Selection Strategies Comparison Across Regions – Historical Performance Relative to Expert FC Selection (% Values).

	US>SG						US>DE						US>AU					
	Pre-Deployment			Post-Deployment			Pre-Deployment			Post-Deployment			Pre-Deployment			Post-Deployment		
Strategy	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units
Profit Rank	0.092	-0.104	-0.050	0.080	-0.132	-0.120	0.405	-0.230	-0.273	0.110	-0.197	-0.206	0.316	-0.067	-0.068	0.193	-0.043	-0.040
Time Rank	-0.477	-0.452	-0.463	-0.113	-0.627	-0.628	-0.157	-0.726	-0.735	-0.640	-0.850	-0.865	0.378	-0.851	-0.862	0.476	-0.857	-0.869
ATE-Profit Rank	-0.015	-0.111	-0.048	0.013	-0.131	-0.126	0.274	-0.460	-0.471	0.010	-0.436	-0.439	0.626	-0.375	-0.374	0.419	-0.316	-0.292
ATE-Time Rank	-0.412	-0.162	-0.142	-0.274	-0.414	-0.421	-0.379	-0.691	-0.667	-0.627	-0.706	-0.690	-0.187	-0.582	-0.589	-0.094	-0.596	-0.584
LP	-0.453	0.001	0.070	-0.187	-0.056	0.040	-0.223	-0.822	-0.834	-0.613	-0.883	-0.893	0.397	-0.939	-0.940	0.382	-0.938	-0.942
XB FC Model (ATE-LP)	-0.418	-0.044	0.025	-0.244	-0.086	-0.077	-0.271	-0.459	-0.489	-0.511	-0.483	-0.485	-0.085	-0.364	-0.359	-0.133	-0.326	-0.301
Profit Rank + LP	-0.486	-0.758	-0.789	-0.130	-0.834	-0.786	-0.263	-0.612	-0.625	-0.615	-0.671	-0.668	-0.057	-0.790	-0.804	0.026	-0.784	-0.803
Time Rank + LP	-0.574	-0.067	0.002	-0.321	-0.119	-0.021	-0.323	-0.765	-0.774	-0.704	-0.823	-0.831	-0.383	-0.820	-0.833	-0.367	-0.813	-0.826
Observed Score Top-k	-0.603	-0.305	-0.238	-0.203	-0.282	-0.176	-0.178	-0.489	-0.479	-0.344	-0.385	-0.373	0.148	-0.135	-0.127	0.219	-0.152	-0.129
Forecasted Score Top-k	-0.612	-0.370	-0.315	-0.196	-0.280	-0.188	-0.181	-0.487	-0.499	-0.338	-0.386	-0.383	0.190	-0.128	-0.126	0.175	-0.101	-0.090

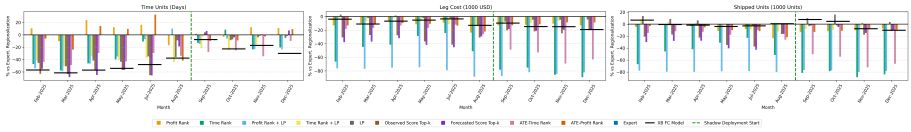

(a) US>SG

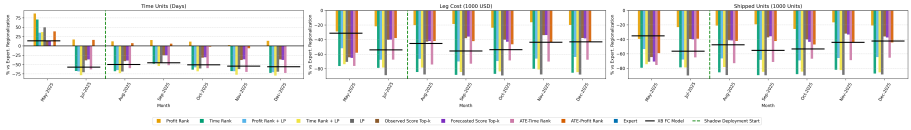

(b) US>DE

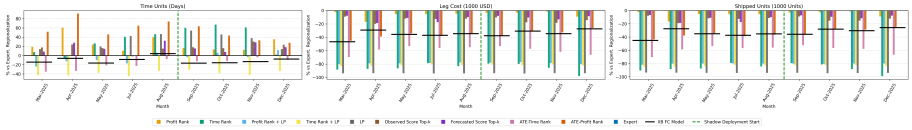

(c) US>AU

Figure 3. Aggregate Performance - Expert vs. FC Selection Strategies: Percentage Difference in Operational Metrics

Table 2. Shipment-Level Counterfactual Evaluation: Same Product and Nearest Neighbor from Optimal vs. Non-Optimal FC

		SG				DE				AU			
Evaluation	Metric	Actual	Expert	LP	XB FC	Actual	Expert	LP	XB FC	Actual	Expert	LP	XB FC
Same Product	Mean Time (Days)	3.21	2.74	1.36	1.18	5.42	5.87	3.21	2.76	4.68	2.31	2.47	1.84
	Mean Cost (USD)	8.47	8.31	7.14	6.52	13.76	13.21	11.84	10.27	14.63	12.46	14.17	11.92
	% Meeting 5/5	54.72	47.35	51.86	59.41	11.64	6.27	9.18	12.83	6.02	6.47	4.16	8.21
	Sample Size (N)	1,842	1,842	1,842	1,842	1,176	1,176	1,176	1,176	963	963	963	963
Nearest Neighbor	Mean Time (Days)	1.38	1.23	0.72	0.70	3.03	3.23	2.00	1.85	2.53	0.93	1.14	1.13
	Mean Cost (USD)	5.86	5.87	5.77	5.55	11.08	10.68	10.03	9.36	12.02	11.13	12.28	10.99
	% Meeting 5/5	41.48	38.10	40.64	44.55	6.83	3.18	5.61	6.85	3.88	4.66	2.46	3.57
	Sample Size (N)	94,216	94,216	94,216	94,216	87,543	87,543	87,543	87,543	81,769	81,769	81,769	81,769

pendix Tables 10–12) confirm all comparative differences are statistically significant at the 1% level.

5. Conclusion and Limitations

We presented the XB FC Selection Model, which combines forecasting, doubly robust candidate generation, and multi-objective LP for FC selection under soft regionalization. Across three production corridors, the model consistently provides a stronger efficiency–volume trade-off than expert, ranking, and LP baselines, reducing delivery time by 8.5–51.1% and cost by up to 48.3% relative to Expert while improving 5/5 compliance by up to 6.5 percentage points. **Limitations**: The main limitations are the lack of direct FC-level XB inventory/capacity observations and the inability to conduct online A/B tests because only one policy can be active per corridor. Furthermore, optimal FC selection requires substantial inventory reallocation (ex: 48% for US>DE), highlighting a gap with current placement. To address these limitations, future work will incorporate explicit inventory data, jointly optimize FC selection and physical inventory reallocation via reinforcement learning, and report post-deployment results.

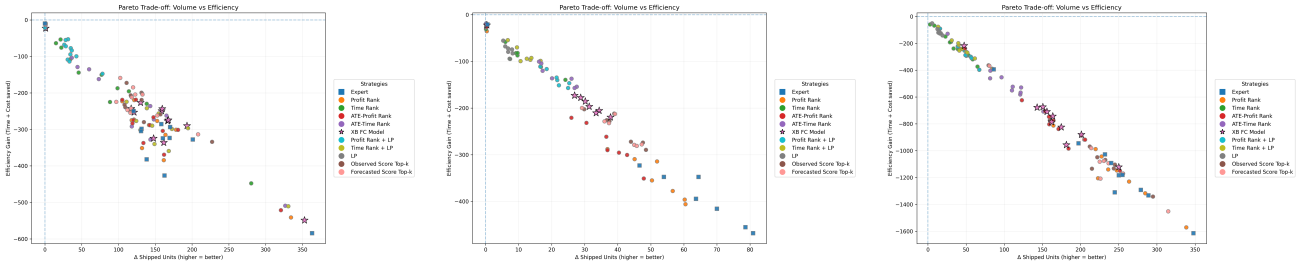

(a) US>SG (b) US>DE (c) US>AU

Figure 4. Pareto Trade-off Between Throughput and Efficiency Across Fulfillment Strategies (Monthly Absolute Values). Each point corresponds to a strategy for a month, visualizing the Pareto frontier between shipment volume and operational efficiency across FC selection strategies versus the no-regionalization baseline.

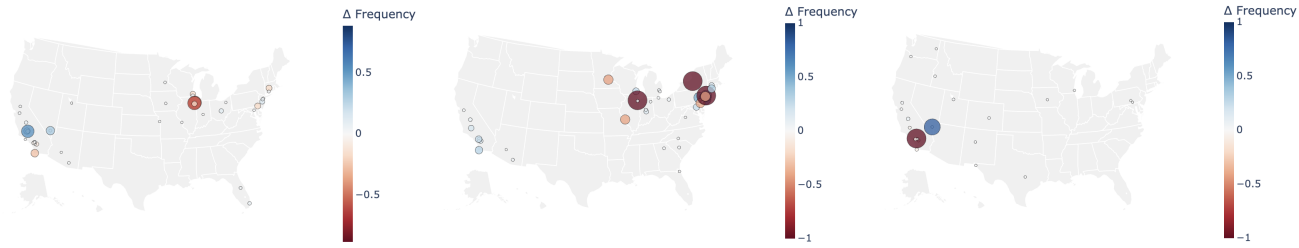

(a) US>SG (b) US>DE (c) US>AU

Figure 5. Spatial comparison of XB FC model selection frequency relative to the Expert baseline. Each point represents a warehouse location; marker size is proportional to deviation magnitude; red indicates higher-than-expert selection, while blue indicates lower-than-expert selection.

Appendix

Table 3. FC Selection Strategies Comparison Across Regions – Historical Performance Relative to No Regionalization (% Values).

	US>SG						US>DE						US>AU					
	Pre-Deployment			Post-Deployment			Pre-Deployment			Post-Deployment			Pre-Deployment			Post-Deployment		
Strategy	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units	Time	Cost	Units
Expert	-0.959	-0.452	-0.331	-0.983	-0.462	-0.328	-0.955	-0.641	-0.582	-0.957	-0.628	-0.609	-0.991	-0.519	-0.441	-0.992	-0.456	-0.378
Profit Rank	-0.957	-0.507	-0.350	-0.981	-0.531	-0.408	-0.947	-0.713	-0.670	-0.953	-0.703	-0.693	-0.988	-0.553	-0.480	-0.991	-0.479	-0.402
Time Rank	-0.980	-0.731	-0.672	-0.986	-0.765	-0.702	-0.980	-0.924	-0.912	-0.986	-0.941	-0.945	-0.988	-0.925	-0.919	-0.989	-0.919	-0.912
ATE-Profit Rank	-0.962	-0.514	-0.352	-0.982	-0.532	-0.411	-0.949	-0.788	-0.755	-0.958	-0.794	-0.786	-0.985	-0.695	-0.645	-0.989	-0.627	-0.560
ATE-Time Rank	-0.976	-0.540	-0.418	-0.988	-0.659	-0.573	-0.980	-0.890	-0.861	-0.986	-0.884	-0.871	-0.993	-0.792	-0.762	-0.993	-0.778	-0.738
LP	-0.978	-0.448	-0.269	-0.987	-0.490	-0.309	-0.980	-0.958	-0.956	-0.984	-0.957	-0.958	-0.988	-0.971	-0.967	-0.989	-0.967	-0.964
XB FC Model (ATE-LP)	-0.983	-0.492	-0.322	-0.989	-0.520	-0.345	-0.976	-0.829	-0.808	-0.980	-0.804	-0.794	-0.992	-0.689	-0.637	-0.993	-0.632	-0.565
Profit Rank + LP	-0.981	-0.869	-0.859	-0.986	-0.903	-0.857	-0.980	-0.882	-0.863	-0.985	-0.876	-0.870	-0.992	-0.898	-0.889	-0.992	-0.882	-0.878
Time Rank + LP	-0.977	-0.476	-0.303	-0.988	-0.507	-0.377	-0.983	-0.927	-0.917	-0.989	-0.937	-0.938	-0.995	-0.912	-0.905	-0.995	-0.898	-0.891
Observed Score Top-k	-0.984	-0.608	-0.467	-0.988	-0.624	-0.472	-0.970	-0.787	-0.747	-0.972	-0.767	-0.751	-0.990	-0.586	-0.514	-0.991	-0.539	-0.460
Forecasted Score Top-k	-0.985	-0.646	-0.522	-0.988	-0.628	-0.486	-0.969	-0.784	-0.746	-0.972	-0.772	-0.760	-0.989	-0.584	-0.515	-0.991	-0.510	-0.434

Table 4. US>SG: Statistical Comparison of FC Selection Strategies vs. Expert Baseline

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	Expert	ATE-Profit Rank	0.717	0.569	0.888	0.042
Leg Cost (1000 USD)	Expert	ATE-Profit Rank	-66.354	0.001	0.000	-1.775
Shipped Units (1000 Units)	Expert	ATE-Profit Rank	-9.257	0.052	0.053	-0.626
Time Units (Days)	Expert	ATE-Time Rank	-26.363	0.001	0.005	-1.021
Leg Cost (1000 USD)	Expert	ATE-Time Rank	-162.802	0.001	0.000	-1.426
Shipped Units (1000 Units)	Expert	ATE-Time Rank	-36.753	0.001	0.001	-1.391
Time Units (Days)	Expert	Forecasted Score Top-k	-33.457	0.042	0.017	-0.814
Leg Cost (1000 USD)	Expert	Forecasted Score Top-k	-181.416	0.001	0.000	-1.454
Shipped Units (1000 Units)	Expert	Forecasted Score Top-k	-42.215	0.001	0.007	-0.963
Time Units (Days)	Expert	Observed Score Top-k	-34.390	0.003	0.012	-0.868
Leg Cost (1000 USD)	Expert	Observed Score Top-k	-166.891	0.001	0.000	-1.549
Shipped Units (1000 Units)	Expert	Observed Score Top-k	-35.885	0.001	0.008	-0.943
Time Units (Days)	Expert	Profit Rank	9.849	0.027	0.021	0.777
Leg Cost (1000 USD)	Expert	Profit Rank	-63.126	0.001	0.000	-2.100
Shipped Units (1000 Units)	Expert	Profit Rank	-8.334	0.034	0.034	-0.699
Time Units (Days)	Expert	Profit Rank + LP	-25.284	0.021	0.034	-0.697
Leg Cost (1000 USD)	Expert	Profit Rank + LP	-444.192	0.001	0.000	-2.351
Shipped Units (1000 Units)	Expert	Profit Rank + LP	-125.011	0.001	0.000	-1.959
Time Units (Days)	Expert	Time Rank	-22.943	0.042	0.044	-0.655
Leg Cost (1000 USD)	Expert	Time Rank	-293.775	0.001	0.000	-2.078
Shipped Units (1000 Units)	Expert	Time Rank	-77.631	0.001	0.000	-2.447
Time Units (Days)	Expert	Time Rank + LP	-34.626	0.007	0.007	-0.946
Leg Cost (1000 USD)	Expert	Time Rank + LP	-52.239	0.001	0.000	-1.707
Shipped Units (1000 Units)	Expert	Time Rank + LP	-2.404	0.266	0.338	-0.289
Time Units (Days)	Expert	LP	-24.414	0.007	0.021	-0.774
Leg Cost (1000 USD)	Expert	LP	-18.058	0.012	0.013	-0.850
Shipped Units (1000 Units)	Expert	LP	6.539	0.064	0.041	0.668
Time Units (Days)	Expert	XB FC Model	-25.758	0.001	0.006	-0.969
Leg Cost (1000 USD)	Expert	XB FC Model	-37.739	0.001	0.002	-1.149
Shipped Units (1000 Units)	Expert	XB FC Model	0.321	0.622	0.939	0.023

Table 5. US>DE: Statistical Comparison of FC Selection Strategies vs. Expert Baseline

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	Expert	ATE-Profit Rank	14.461	0.109	0.092	0.690
Leg Cost (1000 USD)	Expert	ATE-Profit Rank	-213.611	0.008	0.000	-2.286
Shipped Units (1000 Units)	Expert	ATE-Profit Rank	-24.708	0.008	0.000	-2.219
Time Units (Days)	Expert	ATE-Time Rank	-113.836	0.008	0.000	-2.245
Leg Cost (1000 USD)	Expert	ATE-Time Rank	-347.407	0.008	0.000	-2.303
Shipped Units (1000 Units)	Expert	ATE-Time Rank	-38.476	0.008	0.000	-2.218
Time Units (Days)	Expert	Forecasted Score Top-k	-62.847	0.016	0.001	-2.010
Leg Cost (1000 USD)	Expert	Forecasted Score Top-k	-197.450	0.008	0.000	-2.204
Shipped Units (1000 Units)	Expert	Forecasted Score Top-k	-22.380	0.008	0.001	-2.090
Time Units (Days)	Expert	Observed Score Top-k	-64.021	0.016	0.001	-1.930
Leg Cost (1000 USD)	Expert	Observed Score Top-k	-197.967	0.008	0.000	-2.257
Shipped Units (1000 Units)	Expert	Observed Score Top-k	-21.960	0.008	0.001	-2.115
Time Units (Days)	Expert	Profit Rank	27.974	0.008	0.000	2.443
Leg Cost (1000 USD)	Expert	Profit Rank	-100.810	0.008	0.001	-2.053
Shipped Units (1000 Units)	Expert	Profit Rank	-12.122	0.008	0.001	-1.948
Time Units (Days)	Expert	Profit Rank + LP	-111.600	0.016	0.001	-1.924
Leg Cost (1000 USD)	Expert	Profit Rank + LP	-336.185	0.008	0.000	-2.245
Shipped Units (1000 Units)	Expert	Profit Rank + LP	-38.491	0.008	0.000	-2.179
Time Units (Days)	Expert	Time Rank	-114.044	0.016	0.001	-1.817
Leg Cost (1000 USD)	Expert	Time Rank	-407.670	0.008	0.000	-2.184
Shipped Units (1000 Units)	Expert	Time Rank	-46.993	0.008	0.001	-2.141
Time Units (Days)	Expert	Time Rank + LP	-128.606	0.016	0.001	-1.962
Leg Cost (1000 USD)	Expert	Time Rank + LP	-407.038	0.008	0.000	-2.295
Shipped Units (1000 Units)	Expert	Time Rank + LP	-46.922	0.008	0.000	-2.233
Time Units (Days)	Expert	LP	-111.850	0.016	0.001	-1.820
Leg Cost (1000 USD)	Expert	LP	-443.032	0.008	0.000	-2.298
Shipped Units (1000 Units)	Expert	LP	-51.293	0.008	0.000	-2.219
Time Units (Days)	Expert	XB FC Model	-93.582	0.016	0.001	-2.052
Leg Cost (1000 USD)	Expert	XB FC Model	-251.973	0.008	0.001	-2.090
Shipped Units (1000 Units)	Expert	XB FC Model	-29.458	0.008	0.001	-2.003

Table 6. US>AU: Statistical Comparison of FC Selection Strategies vs. Expert Baseline

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	Expert	ATE-Profit Rank	0.000	0.002	0.000	2.359
Leg Cost (1000 USD)	Expert	ATE-Profit Rank	-0.753	0.002	0.000	-4.143
Shipped Units (1000 Units)	Expert	ATE-Profit Rank	-0.080	0.002	0.000	-4.061
Time Units (Days)	Expert	ATE-Time Rank	-0.000	0.006	0.005	-1.170
Leg Cost (1000 USD)	Expert	ATE-Time Rank	-1.284	0.002	0.000	-3.453
Shipped Units (1000 Units)	Expert	ATE-Time Rank	-0.140	0.002	0.000	-3.847
Time Units (Days)	Expert	Forecasted Score Top-k	0.000	0.004	0.001	1.607
Leg Cost (1000 USD)	Expert	Forecasted Score Top-k	-0.261	0.002	0.000	-2.476
Shipped Units (1000 Units)	Expert	Forecasted Score Top-k	-0.027	0.002	0.000	-2.342
Time Units (Days)	Expert	Observed Score Top-k	0.000	0.004	0.000	2.007
Leg Cost (1000 USD)	Expert	Observed Score Top-k	-0.322	0.002	0.000	-2.399
Shipped Units (1000 Units)	Expert	Observed Score Top-k	-0.032	0.002	0.000	-2.563
Time Units (Days)	Expert	Profit Rank	0.000	0.002	0.001	1.524
Leg Cost (1000 USD)	Expert	Profit Rank	-0.124	0.002	0.002	-1.417
Shipped Units (1000 Units)	Expert	Profit Rank	-0.013	0.002	0.001	-1.481
Time Units (Days)	Expert	Profit Rank + LP	-0.000	0.695	0.610	-0.167
Leg Cost (1000 USD)	Expert	Profit Rank + LP	-1.736	0.002	0.000	-3.542
Shipped Units (1000 Units)	Expert	Profit Rank + LP	-0.194	0.002	0.000	-3.580
Time Units (Days)	Expert	Time Rank	0.000	0.010	0.004	1.224
Leg Cost (1000 USD)	Expert	Time Rank	-1.876	0.002	0.000	-3.528
Shipped Units (1000 Units)	Expert	Time Rank	-0.208	0.002	0.000	-3.777
Time Units (Days)	Expert	Time Rank + LP	-0.000	0.002	0.000	-4.845
Leg Cost (1000 USD)	Expert	Time Rank + LP	-1.801	0.002	0.000	-3.522
Shipped Units (1000 Units)	Expert	Time Rank + LP	-0.200	0.002	0.000	-3.646
Time Units (Days)	Expert	LP	0.000	0.002	0.002	1.321
Leg Cost (1000 USD)	Expert	LP	-2.074	0.002	0.000	-3.466
Shipped Units (1000 Units)	Expert	LP	-0.228	0.002	0.000	-3.548
Time Units (Days)	Expert	XB FC Model	-0.000	0.004	0.001	-1.596
Leg Cost (1000 USD)	Expert	XB FC Model	-0.750	0.002	0.000	-3.888
Shipped Units (1000 Units)	Expert	XB FC Model	-0.079	0.002	0.000	-3.847

Table 7. US>SG: Statistical Comparison of Algorithmic FC Selection Strategies vs. XB FC Selection Model

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	XB FC Model	ATE-Profit Rank	26.475	0.000	0.001	1.216
Leg Cost (1000 USD)	XB FC Model	ATE-Profit Rank	-28.615	0.001	0.000	-1.705
Shipped Units (1000 Units)	XB FC Model	ATE-Profit Rank	-9.578	0.000	0.000	-2.104
Time Units (Days)	XB FC Model	ATE-Time Rank	-0.606	0.910	0.710	-0.110
Leg Cost (1000 USD)	XB FC Model	ATE-Time Rank	-125.063	0.000	0.004	-1.033
Shipped Units (1000 Units)	XB FC Model	ATE-Time Rank	-37.074	0.000	0.001	-1.226
Time Units (Days)	XB FC Model	Forecasted Score Top-k	-7.699	0.176	0.146	-0.452
Leg Cost (1000 USD)	XB FC Model	Forecasted Score Top-k	-143.677	0.000	0.000	-1.441
Shipped Units (1000 Units)	XB FC Model	Forecasted Score Top-k	-42.536	0.000	0.001	-1.286
Time Units (Days)	XB FC Model	Observed Score Top-k	-8.633	0.129	0.087	-0.543
Leg Cost (1000 USD)	XB FC Model	Observed Score Top-k	-129.152	0.000	0.000	-1.596
Shipped Units (1000 Units)	XB FC Model	Observed Score Top-k	-36.207	0.000	0.001	-1.355
Time Units (Days)	XB FC Model	Profit Rank	35.607	0.000	0.000	1.431
Leg Cost (1000 USD)	XB FC Model	Profit Rank	-25.387	0.012	0.006	-0.974
Shipped Units (1000 Units)	XB FC Model	Profit Rank	-8.655	0.002	0.001	-1.260
Time Units (Days)	XB FC Model	Profit Rank + LP	0.474	0.970	0.917	0.031
Leg Cost (1000 USD)	XB FC Model	Profit Rank + LP	-406.453	0.000	0.000	-2.480
Shipped Units (1000 Units)	XB FC Model	Profit Rank + LP	-125.332	0.000	0.000	-2.235
Time Units (Days)	XB FC Model	Time Rank	2.814	0.519	0.494	0.204
Leg Cost (1000 USD)	XB FC Model	Time Rank	-256.035	0.001	0.000	-1.758
Shipped Units (1000 Units)	XB FC Model	Time Rank	-77.952	0.000	0.000	-2.047
Time Units (Days)	XB FC Model	Time Rank + LP	-8.869	0.064	0.060	-0.604
Leg Cost (1000 USD)	XB FC Model	Time Rank + LP	-14.500	0.110	0.272	-0.333
Shipped Units (1000 Units)	XB FC Model	Time Rank + LP	-2.725	0.129	0.356	-0.278
Time Units (Days)	XB FC Model	LP	1.343	0.970	0.630	0.143
Leg Cost (1000 USD)	XB FC Model	LP	19.681	0.007	0.038	0.680
Shipped Units (1000 Units)	XB FC Model	LP	6.218	0.000	0.012	0.862

Table 8. US>DE: Statistical Comparison of Algorithmic FC Selection Strategies vs. XB FC Selection Model

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	XB FC Model	ATE-Profit Rank	108.042	0.008	0.000	2.182
Leg Cost (1000 USD)	XB FC Model	ATE-Profit Rank	38.362	0.109	0.035	0.923
Shipped Units (1000 Units)	XB FC Model	ATE-Profit Rank	4.750	0.055	0.034	0.933
Time Units (Days)	XB FC Model	ATE-Time Rank	-20.254	0.008	0.000	-2.397
Leg Cost (1000 USD)	XB FC Model	ATE-Time Rank	-95.434	0.008	0.001	-2.029
Shipped Units (1000 Units)	XB FC Model	ATE-Time Rank	-9.019	0.008	0.002	-1.675
Time Units (Days)	XB FC Model	Forecasted Score Top-k	30.735	0.008	0.002	1.709
Leg Cost (1000 USD)	XB FC Model	Forecasted Score Top-k	54.523	0.016	0.007	1.326
Shipped Units (1000 Units)	XB FC Model	Forecasted Score Top-k	7.077	0.016	0.005	1.409
Time Units (Days)	XB FC Model	Observed Score Top-k	29.560	0.008	0.001	1.819
Leg Cost (1000 USD)	XB FC Model	Observed Score Top-k	54.006	0.016	0.007	1.328
Shipped Units (1000 Units)	XB FC Model	Observed Score Top-k	7.497	0.016	0.003	1.527
Time Units (Days)	XB FC Model	Profit Rank	121.555	0.008	0.000	2.554
Leg Cost (1000 USD)	XB FC Model	Profit Rank	151.163	0.008	0.001	2.032
Shipped Units (1000 Units)	XB FC Model	Profit Rank	17.335	0.016	0.001	1.937
Time Units (Days)	XB FC Model	Profit Rank + LP	-18.018	0.023	0.011	-1.222
Leg Cost (1000 USD)	XB FC Model	Profit Rank + LP	-84.212	0.008	0.000	-2.236
Shipped Units (1000 Units)	XB FC Model	Profit Rank + LP	-9.034	0.008	0.000	-2.225
Time Units (Days)	XB FC Model	Time Rank	-20.462	0.023	0.021	-1.045
Leg Cost (1000 USD)	XB FC Model	Time Rank	-155.697	0.008	0.001	-1.890
Shipped Units (1000 Units)	XB FC Model	Time Rank	-17.535	0.008	0.002	-1.772
Time Units (Days)	XB FC Model	Time Rank + LP	-35.024	0.016	0.002	-1.644
Leg Cost (1000 USD)	XB FC Model	Time Rank + LP	-155.065	0.008	0.000	-2.347
Shipped Units (1000 Units)	XB FC Model	Time Rank + LP	-17.465	0.008	0.000	-2.319
Time Units (Days)	XB FC Model	LP	-18.269	0.023	0.026	-0.993
Leg Cost (1000 USD)	XB FC Model	LP	-191.059	0.008	0.000	-2.422
Shipped Units (1000 Units)	XB FC Model	LP	-21.836	0.008	0.000	-2.391

Table 9. US>AU: Statistical Comparison of Algorithmic FC Selection Strategies vs. XB FC Selection Model

Metric	Baseline	Strategy	Mean Difference	Wilcoxon p-value	T-test p-value	Cohen's d
Time Units (Days)	XB FC Model	ATE-Profit Rank	0.000	0.002	0.000	2.975
Leg Cost (1000 USD)	XB FC Model	ATE-Profit Rank	-0.003	0.375	0.873	-0.052
Shipped Units (1000 Units)	XB FC Model	ATE-Profit Rank	-0.001	0.846	0.678	-0.136
Time Units (Days)	XB FC Model	ATE-Time Rank	-0.000	0.432	0.328	-0.327
Leg Cost (1000 USD)	XB FC Model	ATE-Time Rank	-0.534	0.002	0.000	-2.197
Shipped Units (1000 Units)	XB FC Model	ATE-Time Rank	-0.061	0.002	0.000	-2.429
Time Units (Days)	XB FC Model	Forecasted Score Top-k	0.000	0.002	0.000	2.993
Leg Cost (1000 USD)	XB FC Model	Forecasted Score Top-k	0.489	0.002	0.000	3.272
Shipped Units (1000 Units)	XB FC Model	Forecasted Score Top-k	0.052	0.002	0.000	3.169
Time Units (Days)	XB FC Model	Observed Score Top-k	0.000	0.002	0.000	2.557
Leg Cost (1000 USD)	XB FC Model	Observed Score Top-k	0.428	0.002	0.000	2.990
Shipped Units (1000 Units)	XB FC Model	Observed Score Top-k	0.047	0.002	0.000	2.822
Time Units (Days)	XB FC Model	Profit Rank	0.000	0.002	0.000	2.495
Leg Cost (1000 USD)	XB FC Model	Profit Rank	0.626	0.002	0.000	2.790
Shipped Units (1000 Units)	XB FC Model	Profit Rank	0.066	0.002	0.000	2.738
Time Units (Days)	XB FC Model	Profit Rank + LP	0.000	0.105	0.087	0.607
Leg Cost (1000 USD)	XB FC Model	Profit Rank + LP	-0.986	0.002	0.000	-3.025
Shipped Units (1000 Units)	XB FC Model	Profit Rank + LP	-0.115	0.002	0.000	-3.032
Time Units (Days)	XB FC Model	Time Rank	0.000	0.002	0.001	1.450
Leg Cost (1000 USD)	XB FC Model	Time Rank	-1.127	0.002	0.000	-2.864
Shipped Units (1000 Units)	XB FC Model	Time Rank	-0.129	0.002	0.000	-3.067
Time Units (Days)	XB FC Model	Time Rank + LP	-0.000	0.002	0.000	-2.903
Leg Cost (1000 USD)	XB FC Model	Time Rank + LP	-1.051	0.002	0.000	-3.036
Shipped Units (1000 Units)	XB FC Model	Time Rank + LP	-0.121	0.002	0.000	-3.125
Time Units (Days)	XB FC Model	LP	0.000	0.002	0.001	1.621
Leg Cost (1000 USD)	XB FC Model	LP	-1.324	0.002	0.000	-3.088
Shipped Units (1000 Units)	XB FC Model	LP	-0.149	0.002	0.000	-3.125

Metric	**Comparison**	**t-stat**	**p-value**
Time (Days)			
	ATE vs Actual	-72.74	$< 1e^{-300}$
	LP vs Actual	-70.26	$< 1e^{-300}$
	Expert vs Actual	-72.74	$< 1e^{-300}$
	ATE vs Expert	-78.79	$< 1e^{-300}$
	LP vs Expert	-75.21	$< 1e^{-300}$
	ATE vs LP	-15.73	1.36×10^{-55}
Cost (USD)			
	ATE vs Actual	-14.31	2.40×10^{-46}
	LP vs Actual	-4.43	9.48×10^{-6}
	Expert vs Actual	-14.31	2.40×10^{-46}
	ATE vs Expert	-16.15	1.57×10^{-58}
	LP vs Expert	-5.16	2.46×10^{-7}
	ATE vs LP	-41.71	$< 1e^{-300}$

Table 10. US¿SG: Paired t-test results comparing delivery time and cost across strategies using matched samples.

Metric	**Comparison**	**t-stat**	**p-value**
Time (Days)			
	ATE vs Actual	-62.50	$< 1e^{-300}$
	LP vs Actual	-53.28	$< 1e^{-300}$
	Expert vs Actual	-62.50	$< 1e^{-300}$
	ATE vs Expert	-83.94	$< 1e^{-300}$
	LP vs Expert	-72.97	$< 1e^{-300}$
	ATE vs LP	-24.64	3.25×10^{-133}
Cost (USD)			
	ATE vs Actual	-58.91	$< 1e^{-300}$
	LP vs Actual	-35.07	4.97×10^{-266}
	Expert vs Actual	-58.91	$< 1e^{-300}$
	ATE vs Expert	-45.90	$< 1e^{-300}$
	LP vs Expert	-22.01	7.26×10^{-107}
	ATE vs LP	-68.76	$< 1e^{-300}$

Table 11. US>DE: Paired t-test results comparing delivery time and cost across strategies using matched samples.

Metric	**Comparison**	**t-stat**	**p-value**
Time			
	ATE vs Actual	-514.73	$< 1e^{-300}$
	LP vs Actual	-507.15	$< 1e^{-300}$
	Expert vs Actual	-514.73	$< 1e^{-300}$
	ATE vs Expert	174.52	$< 1e^{-300}$
	LP vs Expert	177.90	$< 1e^{-300}$
	ATE vs LP	-24.13	1.28×10^{-128}
Cost			
	ATE vs Actual	-159.44	$< 1e^{-300}$
	LP vs Actual	35.90	3.58×10^{-282}
	Expert vs Actual	-159.44	$< 1e^{-300}$
	ATE vs Expert	-23.74	1.50×10^{-124}
	LP vs Expert	172.24	$< 1e^{-300}$
	ATE vs LP	-407.83	$< 1e^{-300}$

Table 12. US>AU: Paired t-test results comparing delivery time and cost across strategies using matched samples.

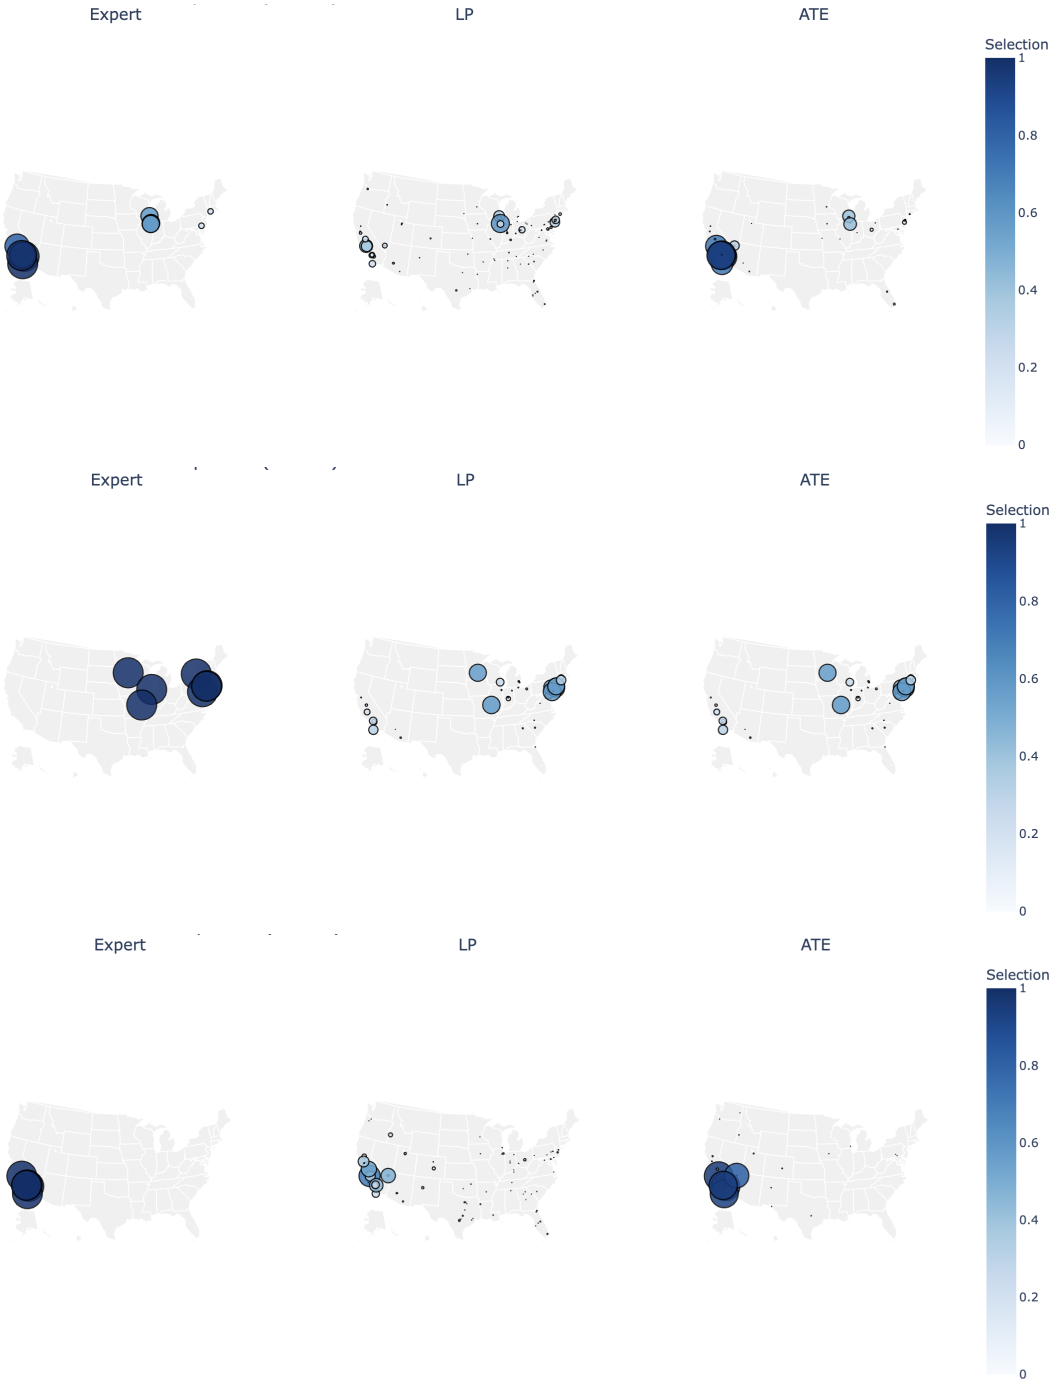

Figure 6. Spatial distribution of fulfillment center selection frequency under alternative policies. Each row shows the spatial distribution across U.S. warehouses for a distinct policy (Expert, LP, and XB FC Selection Model (ATE)). Marker size and color are proportional to normalized selection frequency, with larger and darker markers indicating higher FC utilization. The figure highlights how alternative policies redistribute fulfillment demand relative to the expert baseline.

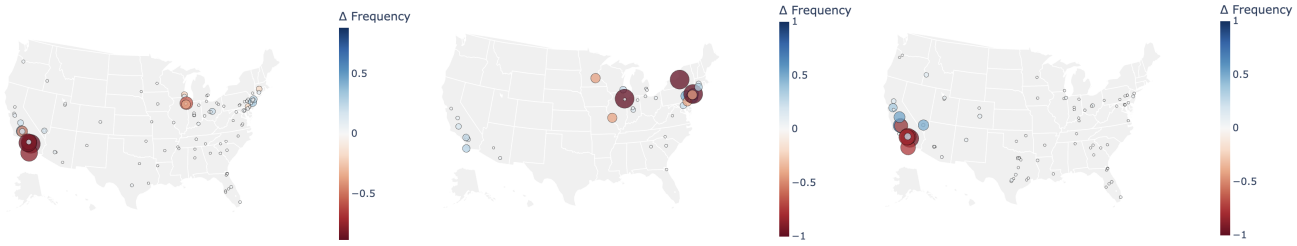

(a) US>SG *(b)* US>DE *(c)* US>AU

Figure 7. Spatial Comparison of LP FC Selection Policies Relative to Expert Baseline This figure compares fulfillment center (FC) selection frequency differences relative to an expert-defined policy across U.S. warehouses. Each point represents a warehouse location, with marker size proportional to the magnitude of the deviation and color indicating the direction and magnitude of change in selection frequency (red = higher-than-expert selection, blue = lower-than-expert selection). The visualization highlights spatial heterogeneity in how alternative policies redistribute FC utilization compared to the baseline expert strategy.

References

Acimovic, J. and Graves, S. C. Making better fulfillment decisions on the fly in an online retail environment. *Manufacturing & Service Operations Management*, 17(1): 34–51, 2015. doi: 10.1287/msom.2014.0505.

Allgor, R., Cezik, T., and Chen, D. Algorithm for robotic picking in amazon fulfillment centers enables humans and robots to work together effectively. *INFORMS Journal on Applied Analytics*, 53(4):266–282, 2023. doi: 10.1287/inte.2022.1143.

Amazon Science. Maximizing the efficiency of amazon's own delivery networks. https://www.amazon.science/blog/maximizing-the-efficiency-of-amazons-own-delivery-networks, 2022. Accessed: 2026-09-09.

Cecil, P. Cross-border supply chain optimization: Strategies for managing international operations while maintaining speed and cost efficiency. *International Journal of Scientific Research and Management*, 12(5):6565–6588, 2024. doi: 10.18535/ijsrm/v12i05.em23.

Deshpande, V. and Pendem, P. Logistics performance, ratings, and its impact on customer purchasing behavior and sales in e-commerce platforms, 2020. Forthcoming.

Guan, X., Zhou, H. Y., Zhang, L. Y., Guo, X. M., and Xie, S. F. Model design of cross-border e-commerce overseas warehouse location selection based on game theory. In *Proceedings of the 2021 International Conference on Bioinformatics and Intelligent Computing*, pp. 226–231, 2021. doi: 10.1145/3448748.3448784.

Jasin, S., Sinha, A., and Uichanco, J. Omnichannel operations: Challenges, opportunities, and models. In Gallino, S. and Moreno, A. (eds.), *Operations in an Omnichannel World*, pp. 15–34. Springer, Cham, 2019. doi: 10.1007/978-3-030-20119-7_2.

Kassmann, D. and Allgor, R. Supply chain design, management and optimization. In Marquardt, W. and Pantelides, C. (eds.), *16th European Symposium on Computer Aided Process Engineering and 9th International Symposium on Process Systems Engineering*, volume 21, pp. 101–106. Elsevier B.V., 2006. doi: 10.1016/S1570-7946(06)80031-3.

Lara, C. L., Koenemann, J., Nie, Y., and de Souza, C. C. Scalable timing-aware network design via lagrangian decomposition. *European Journal of Operational Research*, 309(1):152–169, 2023. doi: 10.1016/j.ejor.2023.01.018.

Li, Z. The impact of logistics service quality of cross border e-commerce business on customer satisfaction. *BCP Business & Management*, 29:130–134, 2022. doi: 10.54691/bcpbm.v29i.2203.

O'Neill, S. Sizing down to scale up: How amazon reworked its fulfillment network to meet customer demand. https://www.amazon.science/news-and-features/how-amazon-reworked-its-fulfillment-network-to-meet-customer-demand, July 2023. Accessed: 2026-09-09.

Sinha, A., Agte, J., Allgor, R., Lara, C. L., Agiwal, A., Atakan, S., Campo, L., Cezik, T., Chen, D., Chen, Q., Fischer, J., Gor, E., Kabbani, N., Kennedy, R., Krishnan, K., Li, Y., Mamadapur, S. M., Mathur, T., McCabe, N., Mildebrath, D., Powell, E., Qualizza, A., Schroeder, D., Seth, N., Si, X., Sinha, K., Stegner, D., Xiao, J., Zhang, L., Zhang, S., and Zou, J. Regionalize and scale: Amazon's fulfillment network design for faster and cheaper delivery. *INFORMS Journal on Applied Analytics*, 56(1):23–41, 2026. doi: 10.1287/inte.2025.0295.

Xu, P. J., Allgor, R., and Graves, S. C. Benefits of reevaluating real-time order fulfillment decisions. *Manufacturing & Service Operations Management*, 11(2):340–355, 2009. doi: 10.1287/msom.1080.0222.