Dark Matters: VAE-Aware HDR Encoding for SDR-to-HDR Video Reconstruction

Donghoon Cha[1], **Bomseok Park**[1], **Hansl von Kwon**[1], **Ku DoHyung**[1], **Hanwool Sul**[2*]

[1]STUDIO FREEWILLUSION Inc.
[2]Seoul National University

arXiv:2611.13504v1 [cs.CV] 10 Nov 2026

Abstract

SDR-to-HDR reconstruction is an ill-posed problem in which clipped highlights, crushed shadows, and quantization artifacts must be recovered while preserving the original content. Latent diffusion models provide powerful generative priors for restoring missing detail, but their effectiveness is constrained by how HDR signals pass through the pretrained VAE. We find that common log-like encodings suppress low-luminance chromatic differences before generation, producing a dark-side collapse in the frozen-VAE latent. We introduce EquivFlow, a compact, invertible exposure-equivariant encoding that maps exposure changes to additive shifts while retaining a scale-independent chromatic component, without modifying the VAE. Encoding-level tests across pretrained image and video VAEs show that its exposure structure and deep-dark chroma transport persist across the tested hosts. Integrated into SDR-to-HDR video diffusion, EquivFlow delivers consistent chromatic-fidelity gains across the tested settings while retaining competitive reconstruction fidelity.

Introduction

The growing adoption of high-dynamic-range (HDR) displays and production workflows has increased the demand for converting existing standard-dynamic-range (SDR) content into HDR. This demand extends beyond conventional film and video remastering to a range of applications, including AI-assisted content creation workflows in which image or video assets are produced in SDR or otherwise limited-dynamic-range form and must be integrated into HDR production pipelines. SDR-to-HDR reconstruction is therefore becoming increasingly relevant across these application areas.

Dark-region fidelity is also a practical requirement in HDR mastering. Production quality control treats crushed blacks as a defect when underexposure removes intended shadow detail, texture, and depth (Netflix n.d.). At the signal level, perceptual HDR coding allocates precision according to luminance-dependent visibility thresholds so that quantization artifacts remain uniformly small (Miller, Nezamabadi, and Daly 2013). These practices motivate preserving not only shadow luminance but also low-luminance chromatic evidence before it enters the pretrained VAE.

Inverse tone mapping is an inherently ill-posed, one-to-many problem. A single SDR observation may correspond to multiple plausible HDR signals because highlight clipping removes intensity variation above the SDR range, while shadow crushing, quantization, and banding suppress subtle luminance and chromatic differences in dark regions. Deterministic tone expansion alone cannot uniquely recover this missing information. A successful SDR-to-HDR method must instead infer plausible detail in ambiguous regions while preserving the spatial and chromatic evidence that remains in the input.

Recent diffusion-based SDR-to-HDR methods address this ambiguity using pretrained generative models. Before generation, latent implementations encode HDR signals with a pretrained variational autoencoder (VAE). Recent systems first compress linear HDR with prescribed log-like transforms—LogC3 in LumiVid and Log-Gamma in DiffHDR—so that it lies within the pretrained VAE's operating range (Korem et al. 2026; Yu et al. 2026). The representation passed through the VAE can therefore affect which luminance and chromatic distinctions remain available to the reconstruction model.

Observation 1: latent chromatic structure collapses at low luminance. Linear HDR signals are poorly matched to the bounded input range of pretrained VAEs; Figure 1 qualitatively illustrates how log-like chromatic trajectories contract in the frozen-VAE latent space at low exposure. Table 1 quantifies this behavior: LogC3 and Log-Gamma match or exceed EquivFlow in frozen-VAE round-trip reconstruction, yet retain less linearly recoverable chroma at -8 EV, with LogC3 unrecoverable there. These frozen-VAE representation and round-trip measurements diagnose what reaches the generative model; they are not downstream SDR-to-HDR reconstruction results.

encoding	recon	$R^2_{u'v'}(0)$	$R^2_{u'v'}(-8)$	surv.$_{-8}$	surv.$_{-10}$
LogC3	35.00	.481	-1.133	-2.354	-3.410
Log-Gamma	**36.07**	.738	.421	.571	.391
EquivFlow	34.38	**.801**	**.883**	**1.102**	**1.105**

Table 1: Frozen-VAE chroma retention for published encodings (no fitted head). $R^2_{u'v'}$ measures linear chroma recoverability; surv.$_{-k}$ is retention after k stops of darkening (1.0 = unchanged).

*Corresponding author.

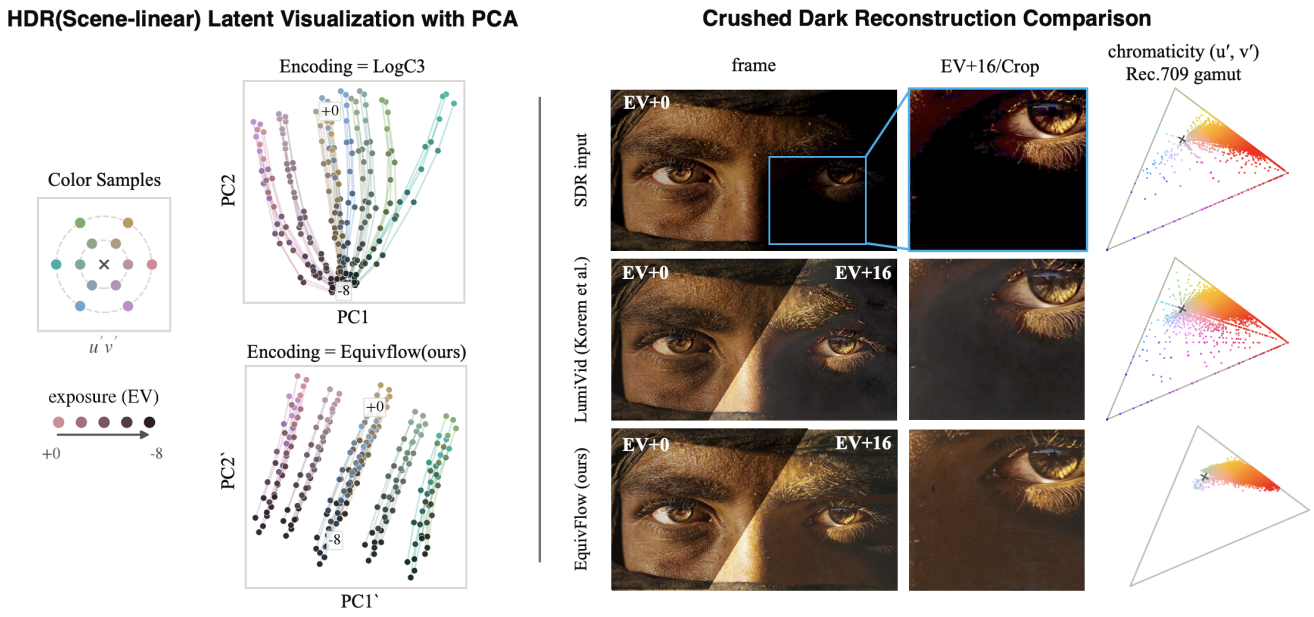

Figure 1: Motivation. **Left:** Frozen-VAE latent PCA for 12 chromaticities over 0 to -8 EV. LogC3 trajectories collapse to one point at -8 EV, while the proposed exposure-equivariant encoding keeps them parallel and separated. The frozen autoencoder and color samples are identical, isolating the encoding as the cause. **Right:** The same failure in a close-up: the EV$+0$ frame, the crushed region lifted to EV$+16$, and its CIE 1976 $u'v'$ pixel distribution relative to the Rec.709 gamut.

Observation 2: exposure changes scale, not chromaticity. In linear HDR, multiplicative exposure changes overall scale while leaving chromaticity unchanged. This suggests an encoding in which exposure becomes translation along a fixed code-space direction and a transverse chromatic component remains scale-independent. In our construction, this relation is exact above the luminance toe when both exposure endpoints remain above it.

These observations motivate EquivFlow, an invertible exposure-equivariant encoding optimized through the frozen target VAE. It adapts the representation to the VAE's operating geometry without modifying the pretrained generative backbone.

Our contributions are: *(i)* We identify and measure the loss of linearly recoverable chromatic structure in dark regions after frozen-VAE encoding, and show that the failure is asymmetric across exposure. *(ii)* We derive EquivFlow, an invertible, 13-parameter exposure-equivariant encoding with an analytic inverse, fitted through frozen-VAE round-trip measurements, and examine its structural behavior across multiple pretrained video VAEs. *(iii)* We integrate EquivFlow into SDR-to-HDR video diffusion and evaluate it against established baselines, obtaining consistent chromatic-fidelity gains across the tested settings while retaining competitive reconstruction fidelity.

Related Work

Inverse tone mapping and learning-based SDR-to-HDR. Classical iTM expands observed SDR luminance using analytical or spatially adaptive operators (Banterle et al. 2006; Rempel et al. 2007), but cannot model content missing from the input. Learning-based methods instead predict saturated radiance with HDRCNN (Eilertsen et al. 2017), regress HDR through multi-scale features with ExpandNet (Marnerides et al. 2018), or decompose global mapping, local enhancement, and highlight generation with HDRTVNet (Chen et al. 2021). These supervised feed-forward models motivate generative priors for severely underdetermined regions.

Diffusion-based SDR-to-HDR reconstruction. Diffusion priors offer a way to synthesize plausible content where the SDR input carries no usable evidence. Existing methods incorporate such priors at different stages of reconstruction: Semantic-Aware Diffusion iTM completes saturated regions before merging estimated exposures (Goswami et al. 2025); LEDiff generates and fuses exposure-dependent latents within an LDR latent coordinate system (Wang et al. 2025); and gain-map diffusion predicts an SDR base and a gain map (Guan et al. 2025).

Many of these diffusion-based approaches build on latent diffusion models, which perform denoising in the latent space of a pretrained autoencoder rather than directly in pixel space (Rombach et al. 2022). This formulation is particularly attractive for SDR-to-HDR reconstruction, as it enables the reuse of strong pretrained image or video priors while operating on compact perceptual representations. Adapting an SDR-pretrained latent diffusion model to HDR, however, also requires mapping HDR signals into a form compatible with its pretrained autoencoder. LumiVid addresses distribution alignment by applying a fixed LogC3 (Brendel 2017) trans-

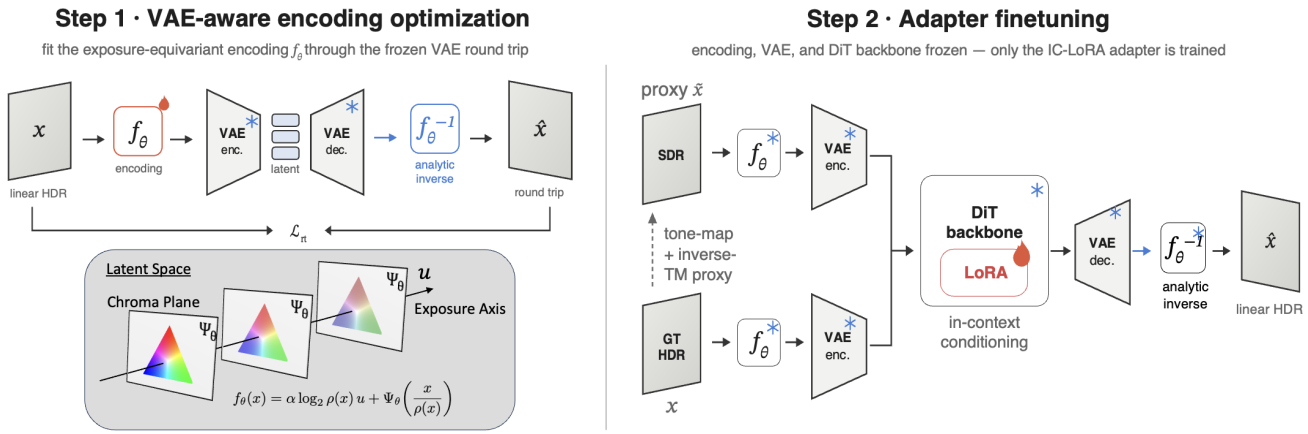

Figure 2: Method overview. The display-referred SDR observation is mapped to a scene-linear proxy by a fixed preprocessing convention, so observation and HDR target share one scene-linear coordinate system; the invertible exposure-equivariant encoding f_θ then carries that space into the pretrained VAE, and is what the generative model actually sees. The encoding is fitted first against the frozen VAE, and the adapter is trained afterwards in the resulting code space. The VAE and the diffusion backbone stay frozen throughout.

form before a frozen video VAE, thereby mapping scene-linear HDR signals closer to the pretrained latent distribution (Korem et al. 2026). DiffHDR follows a related design using a fixed Log-Gamma transform with a frozen video VAE (Yu et al. 2026). These designs highlight HDR encoding as an important component of latent diffusion pipelines for SDR-to-HDR reconstruction.

Luminance and chrominance. Low-light enhancement methods have also emphasized how luminance interacts with color. HVI introduces a trainable color space that separates intensity from chromatic components and compresses low-intensity coordinates to suppress noise (Yan et al. 2025). ICLR further improves luminance–chrominance interaction through dual-stream feature exchange and a covariance correction loss (Xu et al. 2026). These methods show that restoration quality depends not only on network capacity, but also on how luminance and chromatic information are organized before feature extraction. This consideration becomes particularly important when signals undergo strong dynamic-range compression, since subtle chromatic distinctions in dark regions may be attenuated before later restoration stages can exploit them.

Method

Overview

Our method reconstructs scene-linear HDR video from a display-referred SDR observation through a frozen latent video diffusion model. It is defined by three choices. The SDR observation is first mapped to a scene-linear proxy by a fixed preprocessing convention, so that the observation and HDR target can be expressed in one scene-linear coordinate system. That space is then carried into the pretrained VAE by an invertible exposure-equivariant color encoding f_θ, which is what the generative model actually sees. Here, "encoding" denotes an invertible, pointwise reparameterization of scene-linear color rather than a transformation of scene content. The encoding is learned first, against the frozen VAE, so that subsequent adapter training operates in a code space already adapted to the bottleneck it must pass through.

Exposure Equivariance

The constraint is motivated by what the frozen VAE receives. Log-like encodings can shrink dark chromatic differences before the VAE sees them (Figure 1), leaving the generative prior conditioned on an already attenuated signal. We therefore seek an encoding in which exposure changes move codes predictably without reducing chromatic separations.

Exposure acts on scene-linear color by a positive scalar, $x \mapsto sx$. We require the encoding to turn this multiplicative action into a translation along a fixed unit direction u of code space:

$$f_\theta(sx) = f_\theta(x) + \alpha \log_2(s)\, u, \qquad \|u\|_2 = 1, \tag{1}$$

with α the code displacement per exposure stop.

Equation (1) gives exposure and chromaticity separate roles in the code. Along u the code travels with exposure; transverse to u it is untouched by any global scaling and therefore carries chromatic information. In particular, code differences are exposure-invariant, $f_\theta(sx) - f_\theta(sy) = f_\theta(x) - f_\theta(y)$, so chromatic amplitude does not contract merely because luminance falls.

The requirement is restrictive enough to be characterized. A map satisfies Eq. (1) for all $s > 0$ if and only if it decomposes as

$$f_\theta(x) = \alpha \log_2 \rho(x)\, u + \Psi_\theta\left(\frac{x}{\rho(x)}\right), \tag{2}$$

where ρ is positive and 1-homogeneous ($\rho(sx) = s\rho(x)$) and Ψ_θ is defined on the chromaticity slice $\{x : \rho(x) = 1\}$,

with its image taken transverse to the exposure direction u. Taking Ψ_θ to be the restriction of f_θ to that slice and applying Eq. (1) with $s = \rho(x)$ gives one direction; 1-homogeneity of ρ gives the other. The supplement states and proves this formally.

Equation (2) gives a practical template for the encoding. The function ρ chooses the brightness coordinate, Ψ_θ encodes the color remaining after brightness normalization, and α sets how far the code moves per exposure stop. Since the VAE accepts only a bounded input range, these choices are fixed as a measured operating point for the target VAE rather than set from color-science principles alone. The construction gives an invertible representation in which exposure and color have explicit roles before the representation is fitted to the target VAE.

Encoding Parameterization

Our encoding, EquivFlow, realizes Eq. (2) in closed form with 13 fitted parameters. We take $\rho = Y$, the Rec.709 luminance (ITU-R 2015), which is positive and 1-homogeneous, and coordinatize the chromaticity slice by the CIE 1976 coordinates (u', v') (CIE 2024), which are invariant under $x \mapsto sx$ by construction. Both roles of Eq. (2) are then linear, and the encoding is a single affine map of the feature triple $\phi(x) = \big(\ell_\tau(Y),\, u',\, v'\big)$:

$$f_\theta(x) = B\,\phi(x) + b, \tag{3}$$

whose intensity coordinate is a toe-limited logarithm,

$$\ell_\tau(Y) = \begin{cases} \log_2 Y, & Y \geq \tau, \\ \log_2 \tau + \dfrac{Y/\tau - 1}{\ln 2}, & Y < \tau. \end{cases} \tag{4}$$

The fitted toe τ is required because $\log_2$ diverges as $Y \to 0$ while the code range is bounded; below it the intensity coordinate continues linearly with matched value and slope. This keeps the luminance coordinate finite near black without introducing a discontinuity at the toe.

Above the toe the encoding satisfies Eq. (1) exactly rather than approximately: u' and v' are unchanged by scaling and $\ell_\tau(sY) = \ell_\tau(Y) + \log_2 s$, so $f_\theta(sx) = f_\theta(x) + \log_2(s)\, Be_1$. The flow direction and the per-stop budget are therefore read off the first column of B,

$$u = \frac{Be_1}{\|Be_1\|_2}, \qquad \alpha = \|Be_1\|_2, \tag{5}$$

so the exposure direction and per-stop displacement are determined by the first column of B. Inversion needs no learned decoder: solving the affine system gives $\phi = B^{-1}(z - b)$, from which $Y = \ell_\tau^{-1}(\phi_1)$ and (Y, u', v') maps back to linear Rec.709 in closed form, followed by the nonnegativity projection. The 13 parameters are the 9 entries of B, the 3 offsets b, and the toe τ.

VAE-Aware Optimization

Analytic invertibility does not imply survival through the pretrained bottleneck. The encoder is frozen and has its own operating geometry, so theclosed-form parameters of Eq. (3) are fitted against that geometry rather than from color-science principles alone. For the encoding f_θ, the relevant measurement is the frozen VAE round trip

$$x \to f_\theta(x) \to D\left(E(f_\theta(x))\right) \to f_\theta^{-1}\left(D(E(f_\theta(x)))\right), \tag{6}$$

where E and D are the frozen encoder and decoder and the final inverse is followed by the nonnegativity projection.

Let $\hat{x}_\theta$ denote the output of this round trip. For EquivFlow, we score VAE compatibility by the round-trip reconstruction loss

$$\begin{aligned} \mathcal{L}_{\mathrm{rt}} = &\|\log(1 + \hat{x}_\theta) - \log(1 + x)\|_2^2 \\ &+ \lambda_{uv} \left\| m(x)\big(q(\hat{x}_\theta) - q(x)\big) \right\|_2^2 . \end{aligned} \tag{7}$$

Here $q(x) = (u'(x), v'(x))$ and $m(x)$ masks pixels above the dark-noise floor. The encoding is also required to keep codes in the VAE input range before hard clipping.

The EquivFlow encoding is an exactly exposure-equivariant closed form with an analytic inverse, fitted by frozen-VAE round-trip measurements before any task-level diffusion evaluation.

The fitted parameters belong to the VAE they were fitted against; how far one parameter set transfers to other frozen bottlenecks is measured in the experiments rather than assumed.

SDR-to-HDR Reconstruction

After fitting, the encoding is fixed. For supervised training and evaluation, each scene-linear HDR target x is rendered to a display-referred SDR observation by a fixed recipe: normalization, extended Reinhard tone mapping, per-channel clipping, the sRGB OETF (IEC 1999), and 8-bit quantization. The SDR observation is then mapped to a scene-linear proxy $\tilde{x}$ using a fixed Reinhard inverse-tone-mapping preprocessing convention. This proxy defines the conditioning coordinate used by the model, while clipped, crushed, and quantized content remains underdetermined. Because the preprocessing contract may affect performance, the experiments additionally evaluate tone-mapper variants rendered with Reinhard (Reinhard et al. 2002) and ACES-like (Narkowicz 2016) curves.

Both x and $\tilde{x}$ are mapped through the same encoding f_θ and encoded online by the frozen VAE. The target latent serves as the diffusion target, while the proxy latent is supplied as in-context conditioning.

At inference, the conditioned diffusion model predicts an HDR encoding latent. The frozen VAE decoder maps this latent to encoding space, after which the analytic inverse of f_θ maps the result to linear HDR:

$$\hat{x} = f_\theta^{-1}\left(D(\hat{h})\right). \tag{8}$$

Diffusion backbone and IC-LoRA adaptation. The reconstruction model uses an in-context LoRA conditioning scheme (Huang et al. 2024). In our implementation, the SDR condition latent is formed from the scene-linear proxy after the same color encoding f_θ, rather than from raw SDR. It is prepended as clean reference tokens to the noised HDR target-token sequence; reference tokens receive zero timestep and are excluded from the denoising loss, which is evaluated

method	ARRI (40)				HDRTV4K (40)				UPIQ (22)			
	PU21-PSNR ↑		$\Delta C_T C_P$ ↓		PU21-PSNR ↑		$\Delta C_T C_P$ ↓		PU21-PSNR ↑		$\Delta C_T C_P$ ↓	
TM method	rein.	aces	rein.	aces	rein.	aces	rein.	aces	rein.	aces	rein.	aces
EQUIVFLOW	**39.99**	26.11	**5.51**	**5.97**	**33.12**	26.16	**8.67**	**10.23**	**41.77**	27.61	**3.02**	**4.27**
LogC3	39.70	**26.33**	5.52	6.23	31.43	**27.44**	9.73	11.95	39.28	**28.63**	3.93	5.59
LumiVid	36.26	25.98	6.87	8.18	32.36	27.25	10.86	14.54	36.01	26.94	6.18	10.30
HDRTVNet++	26.32	20.61	6.05	6.23	—	—	—	—	28.07	22.48	4.96	6.19
LEDiff	18.37	18.74	12.25	12.29	18.88	19.12	18.87	19.41	20.10	20.64	10.29	10.74

Table 2: PU21-PSNR and the chromatic ΔE_{ITP} term $\Delta C_T C_P$ on a frozen 102-clip test set — the first 40 ARRI crops (10 parents), the first 40 HDRTV4K clips (19 parents), and all 22 UPIQ parents — frozen before any aggregation was computed, under the training curve (*rein.*) and one unseen curve (*aces*); the HDR ground truth, clip list, per-clip normalization and crop are held fixed while the SDR tone curve moves. EQUIVFLOW holds the lowest $\Delta C_T C_P$ in all six chromatic columns and the best PU21-PSNR under the training curve on all three sets; under the unseen curve the aggregate lead passes to LogC3. $\Delta C_T C_P$ is defined on PQ-encoded ICtCp and is not exposure-invariant; the exposure-invariant $\Delta u'v'$ reading is Table 3. Parent-weighted means; clip counts in parentheses; best per column in bold. Upper block: the matched internal pair; lower block: external systems. —: withheld, same-domain with HDRTVNet++'s training corpus and never allowed to win a column.

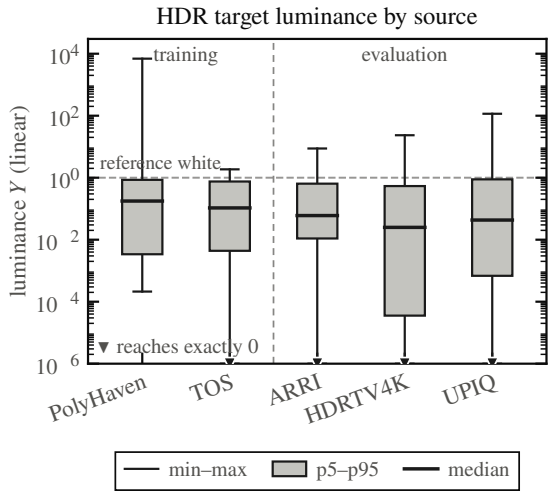

Figure 3: HDR target luminance by source (log axis, pooled over sampled pixels): line and caps span minimum to maximum, the bar the 5th–95th percentile, the dot the median, and the dashed line marks reference white. Target, not the SDR the model receives.

only on the target portion. Only the low-rank adapter (Hu et al. 2022) is optimized; the VAE, diffusion backbone, and text encoder remain frozen.

Experiments

Datasets and Setup

Training corpus. All internal arms are trained on one corpus of 5,780 crops at 960×544 and 49 frames, drawn from 1,015 parent sequences: 215 windows from 94 *Tears of Steel* (Blender Foundation 2012) shots (2,580 crops) and 800 clips rendered from PolyHaven (Poly Haven n.d.) HDRI panoramas (3,200 crops). Splits are parent-disjoint, shot-level for *Tears of Steel* and asset-level for PolyHaven.

These are fixed derivatives of the cited public sources, not a newly introduced dataset or benchmark.

Training inputs. Training SDR inputs use the fixed Reinhard-based rendering protocol described in the Method section, followed by the same camera-style degradations for every internal arm: contrast scaling, band-selective blur, 8-bit quantization, and codec compression. Evaluation additionally includes ACES-like renderings.

Evaluation sets. We evaluate on ARRI's official ALEXA Mini ProRes 4444 XQ sample reel M001R00H (ARRI n.d.), HDRTV4K (Guo et al. 2023), and UPIQ (Mikhailiuk et al. 2021), decoding every prediction to common linear Rec. 709 (ITU-R 2015) before scoring. ARRI contains 176 crops from 44 cinema-camera parents and is the brightest set. HDRTV4K contains 78 clips from 36 parents of mixed video and static content, including material from the source used by HDRTVNet++. UPIQ contains one 49-frame static clip for each of 22 calibrated image parents, so temporal metrics are undefined there. The main table reports a 102-clip test subset frozen before any aggregation was computed: the first 40 ARRI crops (10 parents), the first 40 HDRTV4K clips (19 parents), and all 22 UPIQ parents. The chromatic-decomposition panel is computed on the full HDRTV4K–UPIQ pool.

Baselines. The internal rows share the training split above, the frozen LTX-2.3 22B backbone and VAE (HaCohen et al. 2024; Lightricks 2026), IC-LoRA adapter (Huang et al. 2024), inference settings, and scorer. We report the protocol from the Method section: canonical Reinhard inputs and one unseen tone mapper, ACES-like. LogC3 is the matched internal baseline: it is retrained on the same corpus and degradation protocol with the same backbone, conditioning method, and scorer. The two internal arms differ in the HDR representation and in its paired reference contract: EQUIVFLOW conditions on the inverse-tone-mapped scene-linear proxy of the Method section, which is part of the contribution, while LogC3 is trained under raw-SDR conditioning, the official

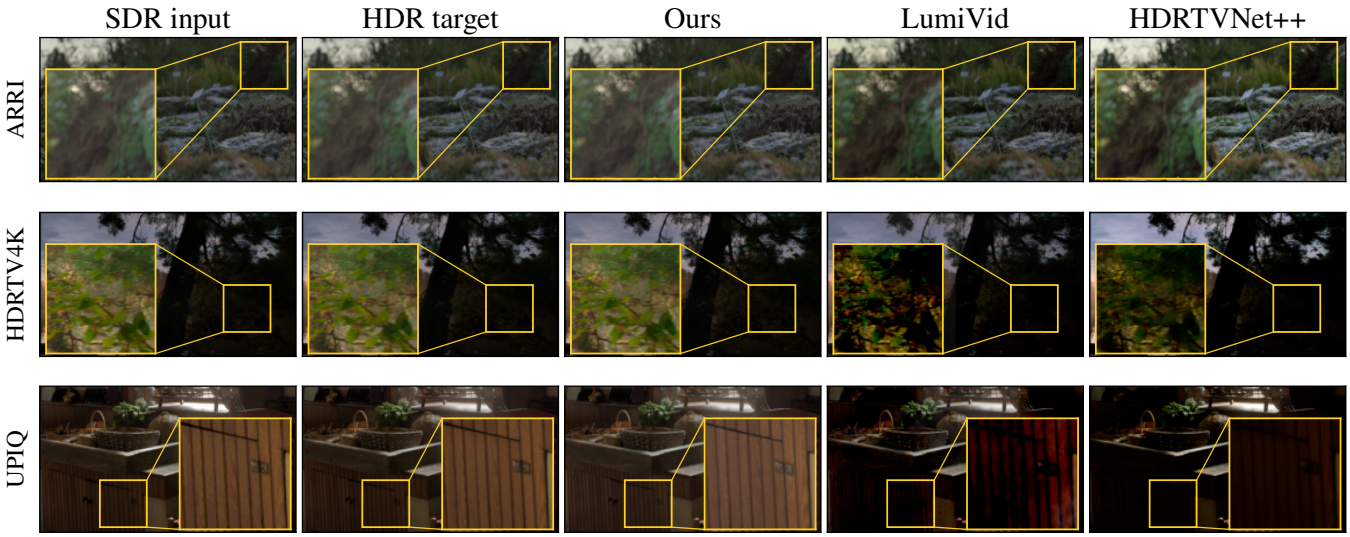

Figure 4: Qualitative comparison at frame 24. The inset covers the darkest chroma-carrying ground-truth region; each row carries a display gain applied identically across columns.

contract of LumiVid.

LumiVid is reported separately because it is the closest architectural predecessor: a pretrained video diffusion model adapted for SDR-to-HDR through a fixed LogC3 interface. HDRTVNet++ (Chen et al. 2023) and LEDiff (Wang et al. 2025) are external system baselines that differ in architecture, training data, and budget. LEDiff uses its official preprocessing and one method-level radiometric gain fitted on a frozen 24-clip training subset, without access to test images or per-clip ground truth.

Main Results

Quantitative results. Table 2 evaluates the frozen test subset under the training Reinhard curve and an unseen ACES-like curve, with the target and evaluated clips held fixed. PU21-PSNR (Mantiuk and Azimi 2021) measures aggregate HDR reconstruction, while $\Delta C_T C_P$, the chromatic component of ΔE_{ITP} (ITU-R 2019), measures perceptual chromatic error at absolute luminance. Under the training curve, EquivFlow leads the matched LogC3 control on both quantities across all three sets. Under the unseen curve, it retains the chromatic lead throughout, while LogC3 takes the aggregate PU21-PSNR lead. The consistent result is therefore improved chromatic fidelity across tone curves; aggregate reconstruction depends more strongly on the SDR rendering. Parent-level tests support the matched-pair chromatic difference on UPIQ under the training curve and on HDRTV4K and UPIQ under the unseen curve; the smaller ARRI subset does not separate under the training curve. Remaining preregistered metrics and full statistical details are reported in the supplement.

Baseline comparison. LumiVid is the closest architectural predecessor and therefore the most direct external comparison. EquivFlow has lower chromatic error in every column and higher aggregate fidelity in all but the HDRTV4K ACES-like case. It also leads LEDiff and HDRTVNet++ on both reported quantities wherever a comparison is available;

	reinhard		aces-like	
metric ↓	Ours	LogC3	Ours	LogC3
chroma L1	**.0431**	.0539	**.0525**	.0656
$\Delta u'v'$	**.0140**	.0173	**.0170**	.0247
$\Delta C_T C_P$	**6.37**	7.37	**7.98**	9.47

Table 3: Chromatic error of the internal pair on the pooled HDRTV4K–UPIQ panel (58 parents), training curve and one unseen curve. Lower is better; parent-weighted; best per pair in bold. All six paired differences favor Ours and survive Holm correction (largest corrected $p = 4.8 \times 10^{-6}$). ARRI appears in Table 2 only; nested $\Delta u'v'$ variants and per-metric intervals: supplement.

HDRTVNet++ is withheld on HDRTV4K because that set is same-domain. These remain system-level comparisons because architecture, training data, radiometric convention, and budget differ.

Chromatic decomposition. To test whether the main-table chromatic trend depends on one metric, Table 3 evaluates the matched pair using raw chroma amplitude, exposure-invariant CIE 1976 chromaticity, and PQ-encoded ICtCp (ITU-R 2018). Every metric favors EquivFlow under both tone curves, with a larger margin under the unseen curve. This consistency supports the chromatic-fidelity reading across the tested metrics and curves.

Qualitative results. Figure 4 shows the highest-margin EquivFlow–LumiVid example from each evaluation set, so the rows show selected successes rather than typical cases. In the magnified dark regions, EquivFlow remains closer to the target's color and local structure, while the comparison outputs show stronger saturation, hue, or shadow-level deviations. These examples complement, but do not replace, the dataset-level quantitative comparisons above.

VAE	$R^2_\beta \uparrow$	grid $0 \to -8 \uparrow$	$\Delta u'v'_{-8} \downarrow$	/LTX
LTX	0.9990	$0.979 \to 0.979$	0.01536	1.000
SD 2.1	0.9927	$0.667 \to 0.812$	0.00756	0.492
Wan 2.1	0.9977	$0.958 \to 0.917$	0.00824	0.536
Wan 2.2	0.9976	$0.979 \to 0.854$	0.00793	0.516

Table 4: Training-free portability of EquivFlow, measured on this draft's carrier. Grid accuracy uses a host- and exposure-specific readout; $\Delta u'v'$ is frozen-VAE round-trip error and /LTX is that error relative to the LTX host.

Representation Analysis

Table 1 tests representation behavior before diffusion by passing each published encoding through the same frozen VAE without a fitted head. The log-like encodings reconstruct well but lose linearly recoverable chroma as exposure falls, whereas EquivFlow preserves it; reconstruction quality is therefore not a proxy for dark-exposure chroma retention. This divergence is consistent with exposure–chroma coupling and matches the downstream widening of the chromatic margin under the unseen tone curve (Tables 2 and 3). We next examine the supporting design choices.

Cross-VAE transfer test. EquivFlow's exposure structure is tested on frozen LTX, Stable Diffusion 2.1 (Rombach et al. 2022; Stability AI 2022), Wan 2.1 T2V-1.3B (Team Wan et al. 2025), and Wan 2.2 TI2V-5B (Wan Team 2025) VAEs; only LTX was used during fitting. On the three foreign hosts, the exposure-linked latent displacement remains nearly linear and chromaticity readout stays above chance at deep exposure. All three foreign hosts pass the frozen roundtrip non-collapse rule in that regime.

Portability here concerns the exposure structure, not the latent coordinates themselves: an EV-0-frozen readout degrades along each host's exposure trajectory, and absolute round-trip quality varies substantially across hosts.

Conclusion

We introduced EquivFlow, a VAE-aware HDR encoding motivated by the loss of recoverable chromatic structure under darkening in common log-like representations. Its compact, invertible, pointwise construction maps exposure changes to code-space translations while preserving a scale-independent chromatic component, without modifying the pretrained VAE or diffusion backbone. Representation analyses across pretrained image and video VAEs confirm the persistence of this structure, and integration into SDR-to-HDR video diffusion yields consistent chromatic-fidelity gains across the tested datasets and tone curves while retaining competitive reconstruction fidelity. Together, these results establish VAE-aware encoding as an effective interface for adapting pretrained latent generative models to HDR reconstruction. Future work will examine this design across broader VAE families and reconstruction settings.

References

ARRI. n.d. ARRI Camera Sample Footage and Reference Images. Official sample-footage collection. ALEXA Mini ProRes 4444 XQ sample reel M001R00H.

Banterle, F.; Ledda, P.; Debattista, K.; and Chalmers, A. 2006. Inverse Tone Mapping. In *Proceedings of GRAPHITE*.

Blender Foundation. 2012. Tears of Steel. Mango Open Movie Project, https://mango.blender.org/.

Brendel, H. 2017. ALEXA Log C Curve: Usage in VFX. White paper, ARRI. LogC3 definition.

Chen, X.; Li, Z.; Zhang, Z.; Ren, J. S.; Liu, Y.; He, J.; Qiao, Y.; Zhou, J.; and Dong, C. 2023. Towards Efficient SDRTV-to-HDRTV by Learning from Image Formation. *arXiv preprint arXiv:2309.04084*. HDRTVNet++.

Chen, X.; Zhang, Z.; Ren, J. S.; Tian, L.; Qiao, Y.; and Dong, C. 2021. A New Journey from SDRTV to HDRTV. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*.

CIE. 2024. ISO/CIE 11664-5:2024: Colorimetry—Part 5: CIE 1976 $L^*u^*v^*$ Colour Space and u', v' Uniform Chromaticity Scale Diagram. Technical report, International Commission on Illumination.

Eilertsen, G.; Kronander, J.; Denes, G.; Mantiuk, R. K.; and Unger, J. 2017. HDR Image Reconstruction from a Single Exposure Using Deep CNNs. *ACM Transactions on Graphics*, 36(6).

Goswami, A.; Singh, A. R.; Banterle, F.; Debattista, K.; and Bashford-Rogers, T. 2025. Semantic Aware Diffusion Inverse Tone Mapping. *Journal of Physics: Conference Series*, 3128(1): 012009. ArXiv:2405.15468.

Guan, Y.; Xu, R.; Liao, Y.; Yao, M.; Wang, L.; and Xiong, Z. 2025. HDR Image Generation via Gain Map Decomposed Diffusion. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 17536–17545.

Guo, C.; Fan, L.; Xue, Z.; and Jiang, X. 2023. Learning a Practical SDR-to-HDRTV Up-Conversion Using New Dataset and Degradation Models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 22231–22241.

HaCohen, Y.; Chiprut, N.; Brazowski, B.; Shalem, D.; Moshe, D.; Richardson, E.; Levin, E.; Shiran, G.; Zabari, N.; Gordon, O.; Panet, P.; Weissbuch, S.; Kulikov, V.; Bitterman, Y.; Melumian, Z.; and Bibi, O. 2024. LTX-Video: Realtime Video Latent Diffusion. *arXiv preprint arXiv:2501.00103*.

Hu, E. J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; and Chen, W. 2022. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations (ICLR)*.

Huang, L.; Wang, W.; Wu, Z.-F.; Shi, Y.; Dou, H.; Liang, C.; Feng, Y.; Liu, Y.; and Zhou, J. 2024. In-Context LoRA for Diffusion Transformers. *arXiv preprint arXiv:2410.23775*.

IEC. 1999. IEC 61966-2-1:1999: Multimedia Systems and Equipment—Colour Measurement and Management—Part 2-1: Colour Management—Default RGB Colour Space—sRGB. Technical report, International Electrotechnical Commission.

ITU-R. 2015. Recommendation ITU-R BT.709-6: Parameter Values for the HDTV Standards for Production and International Programme Exchange. Technical report, International Telecommunication Union.

ITU-R. 2018. Recommendation ITU-R BT.2100-2: Image Parameter Values for High Dynamic Range Television for Use in Production and International Programme Exchange. Technical report, International Telecommunication Union. ICtCp definition.

ITU-R. 2019. Recommendation ITU-R BT.2124-0: Objective Metric for the Assessment of the Potential Visibility of Colour Differences in Television. Technical report, International Telecommunication Union. Delta E ITP definition.

Korem, N. K.; Oumoumad, M.; Cain, H.; Ben Yosef, M.; Jelercic, U.; Bibi, O.; Inger, Y.; Patashnik, O.; and Cohen-Or, D. 2026. HDR Video Generation via Latent Alignment with Logarithmic Encoding. *arXiv preprint arXiv:2604.11788*. LumiVid.

Lightricks. 2026. LTX-2.3 Model and Inference Repository. GitHub repository. LTX-2.3 22B development checkpoint.

Mantiuk, R. K.; and Azimi, M. 2021. PU21: A Novel Perceptually Uniform Encoding for Adapting Existing Quality Metrics for HDR. In *Picture Coding Symposium (PCS)*.

Marnerides, D.; Bashford-Rogers, T.; Hatchett, J.; and Debattista, K. 2018. ExpandNet: A Deep Convolutional Neural Network for High Dynamic Range Expansion from Low Dynamic Range Content. *Computer Graphics Forum*, 37(2): 37–49.

Mikhailiuk, A.; Pérez-Ortiz, M.; Yue, D.; Suen, W.; and Mantiuk, R. K. 2021. Consolidated Dataset and Metrics for High-Dynamic-Range Image Quality. *IEEE Transactions on Multimedia*, 24: 2125–2138.

Miller, S.; Nezamabadi, M.; and Daly, S. 2013. Perceptual Signal Coding for More Efficient Usage of Bit Codes. *SMPTE Motion Imaging Journal*, 122(4): 52–59.

Narkowicz, K. 2016. ACES Filmic Tone Mapping Curve. Technical blog. Fitted ACES-like approximation, not the normative ACES rendering transform.

Netflix. n.d. Crushed Blacks. Netflix Partner Help Center.

Poly Haven. n.d. Poly Haven: HDRI Library. https://polyhaven.com/hdris.

Reinhard, E.; Stark, M.; Shirley, P.; and Ferwerda, J. 2002. Photographic Tone Reproduction for Digital Images. *ACM Transactions on Graphics*, 21(3): 267–276.

Rempel, A. G.; Trentacoste, M.; Seetzen, H.; Young, H. D.; Heidrich, W.; Whitehead, L.; and Ward, G. 2007. Ldr2Hdr: On-the-Fly Reverse Tone Mapping of Legacy Video and Photographs. In *ACM SIGGRAPH 2007 Papers*.

Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; and Ommer, B. 2022. High-Resolution Image Synthesis with Latent Diffusion Models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.

Stability AI. 2022. Stable Diffusion v2-1 Model Card. Hugging Face model card; accessed through the sd2-community mirror.

Team Wan; et al. 2025. Wan: Open and Advanced Large-Scale Video Generative Models. *arXiv preprint arXiv:2503.20314*.

Wan Team. 2025. Wan2.2: Open and Advanced Large-Scale Video Generative Models. GitHub repository and model release.

Wang, C.; Xia, Z.; Leimkühler, T.; Myszkowski, K.; and Zhang, X. 2025. LEDiff: Latent Exposure Diffusion for HDR Generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 453–464. ArXiv:2412.14456.

Xu, X.; Liu, H.; Liu, W.; Wang, W.; Wu, J.; and Jiang, K. 2026. ICLR: Inter-Chrominance and Luminance Interaction for Natural Color Restoration in Low-Light Image Enhancement. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, 11415–11423.

Yan, Q.; Feng, Y.; Zhang, C.; Pang, G.; Shi, K.; Wu, P.; Dong, W.; Sun, J.; and Zhang, Y. 2025. HVI: A New Color Space for Low-light Image Enhancement. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 5678–5687.

Yu, Z.; Ma, L.; He, M.; Isikdogan, L.; Xu, Y.; Smirnov, D.; Salamanca, P.; Mi, D.; Delgado, P.; Yu, N.; Philip, J.; Li, X.; Wang, W.; and Debevec, P. 2026. DiffHDR: Re-Exposing LDR Videos with Video Diffusion Models. *arXiv preprint arXiv:2604.06161*.

Supplementary Material

Exposure-Equivariant Encodings

This section states the general characterization used in the main paper and then verifies the properties of the realized encoding. Let $f : \mathbb{R}^3_{>0} \to \mathbb{R}^d$ be a code map, let $u \in \mathbb{R}^d$ be a unit vector, and let $\alpha > 0$. We say that f is exposure-equivariant along u if, for every color $x \in \mathbb{R}^3_{>0}$ and every exposure scale $s > 0$,

$$f(sx) = f(x) + \alpha \log_2(s) u. \tag{1}$$

Proposition 1 (characterization). Let $\rho : \mathbb{R}^3_{>0} \to \mathbb{R}_{>0}$ be any positive 1-homogeneous function, $\rho(sx) = s\rho(x)$, and define

$$\mathcal{C}_\rho = \{c \in \mathbb{R}^3_{>0} : \rho(c) = 1\}.$$

A map f satisfies Eq. (1) if and only if there exists a map $\Psi : \mathcal{C}_\rho \to \mathbb{R}^d$ such that

$$f(x) = \alpha \log_2 \rho(x) u + \Psi\left(\frac{x}{\rho(x)}\right). \tag{2}$$

This characterization alone does not assert injectivity; transversality and injectivity of the chromatic component are additional requirements.

Proof. Assume Eq. (1). For $c \in \mathcal{C}_\rho$, define $\Psi(c) = f(c)$. For arbitrary x, set $c = x/\rho(x)$. Homogeneity gives $\rho(c) = 1$ and $x = \rho(x)c$. Hence

$$f(x) = f(c) + \alpha \log_2 \rho(x) u = \Psi(c) + \alpha \log_2 \rho(x) u,$$

which is Eq. (2). Conversely, homogeneity gives $\rho(sx) = s\rho(x)$ and $sx/\rho(sx) = x/\rho(x)$. Substitution into Eq. (2) yields

$$f(sx) = \alpha \log_2(s\rho(x))u + \Psi\left(\frac{x}{\rho(x)}\right) = f(x) + \alpha \log_2(s) u.$$

□

The Realized Encoding

The deployed EquivFlow encoding is

$$f(x) = B \begin{bmatrix} \ell_\tau(Y(x)) \\ u'(x) \\ v'(x) \end{bmatrix} + b, \tag{3}$$

where Y, u', v' are Rec. 709 luminance and CIE 1976 chromaticity. The toe-limited log luminance is

$$\ell_\tau(Y) = \begin{cases} \log_2 Y, & Y \geq \tau, \\ \log_2 \tau + (Y/\tau - 1)/\ln 2, & 0 \leq Y < \tau. \end{cases} \tag{4}$$

The frozen numerical values are

$$B = \begin{bmatrix} 0.0512156526 & 1.0503072513 & 0.1687983249 \\ 0.0553585156 & -0.4965997474 & 0.3491496748 \\ 0.0499266056 & 0.0399719486 & -1.0148571856 \end{bmatrix},$$

$$b = [0.4599701209 \quad 0.6956425436 \quad 1.2171179097]^\mathsf{T}, \quad \tau = 0.0010543414867.$$

Analytic inverse and transversality. The scalar function ℓ_τ is strictly increasing on $Y \geq 0$, with range $[t_{\min}, \infty)$, where $t_{\min} = \log_2 \tau - 1/\ln 2$. Its inverse is

$$\ell_\tau^{-1}(t) = \begin{cases} 2^t, & t \geq \log_2 \tau, \\ \tau[1 + \ln 2\,(t - \log_2 \tau)], & t_{\min} \leq t < \log_2 \tau. \end{cases} \tag{5}$$

Moreover, the selected B is nonsingular: it has rank 3, $\det B = 0.10697$, and spectral condition number 13.77. Thus, for a code z in the image of Eq. (3), one first computes $q = B^{-1}(z - b)$, recovers $Y = \ell_\tau^{-1}(q_1)$, and then converts $(Y, q_2, q_3) = (Y, u', v')$ uniquely back to linear Rec. 709 RGB. The standard $Yu'v'$ inverse is unique on the admissible positive-color domain ($Y > 0, v' > 0$), and the Rec. 709 RGB–XYZ matrix is nonsingular. Therefore Eq. (3) is injective on that domain.

This argument does not require the exposure and chromatic code directions to be orthogonal. Define the dual functional

$$\lambda^\mathsf{T} = e_1^\mathsf{T} B^{-1}.$$

Then $\lambda^\mathsf{T} B e_1 = 1$ and $\lambda^\mathsf{T} B e_2 = \lambda^\mathsf{T} B e_3 = 0$. Consequently, $\lambda^\mathsf{T}(z - b)$ recovers the luminance coordinate, while the chromatic plane $\operatorname{span}(Be_2, Be_3)$ lies in $\ker \lambda^\mathsf{T}$. This is the non-orthogonal direct-sum transversality realized by the fitted matrix.

Exactness domain. Above the toe, scaling RGB by s leaves (u', v') unchanged and increments the first coordinate by $\log_2 s$. Hence Eq. (1) holds whenever both Y and sY are at least τ, with $\alpha = \|Be_1\|_2$ and $u = Be_1/\|Be_1\|_2$. The linear toe is deliberately finite at black and is not exactly log-equivariant when either endpoint lies below τ.

The implementation floors RGB components at 10^{-8} before computing $u'v'$, clamps inverse RGB to the nonnegative orthant, and clamps encoded codes to $[0, 1]$ at the VAE interface before mapping them to $[-1, 1]$. Accordingly, analytic injectivity and equivariance refer to the mathematical encoding on the interior where these guards and code clipping are inactive; they are not claims of global invertibility or exact equivariance for the clipped VAE pipeline.

Encoding Fitting and Selection

Relation to the main paper. The main paper describes encoding fitting at the interface level: the closed-form parameters of its Eq. (3) are fitted by frozen-VAE round-trip measurements, and Figure 2 (Step 1) depicts that fitting as a single block. Concretely, this block is realized in two sub-stages. Stage 1 optimizes an auxiliary intermediate map through the frozen VAE by gradient descent; its round-trip component corresponds to main-paper Eq. (7) (see the normalization note below), and the main paper's code-range requirement enters as the range penalty of the stage-1 objective, Eq. (6) below. Stage 2 produces the deployed encoding in closed form: the stage-1 targets are projected into the affine family of main-paper Eq. (3) by least squares (affine distillation), and one displacement scale is then chosen by the prespecified selection gates. The deployed parameters are therefore obtained as follows: B and b from the closed-form projection, the toe τ carried over from the stage-1 map and frozen, and

the first-column scale from the frozen selection rule. Gradient optimization occurs only inside stage 1; the stage-1 map is scaffolding and is discarded after distillation.

Calibration data. The encoding was fitted once through the frozen LTX video VAE. The fixed calibration pool contains 768 crops: four crops from each of 192 temporal windows. Each crop has 9 frames at 256×256. With seed 0, saturation is sampled uniformly from $[.55, .80]$ and, with probability .5, hue is perturbed uniformly within $[-40°, 40°]$.

Stage 1: VAE-aware fitting targets. Stage 1 optimizes an intermediate map g only to provide targets for affine distillation. Concretely, g is the fixed-form analytic map

$$g(x) = s_{\text{out}} \odot M_2\{a \odot \operatorname{asinh}(b \odot M_1 x) + d\} + b_{\text{out}},$$

where $M_i = I + \Delta M_i$; its active tensors comprise the two 3×3 matrices; the three-channel $a, b, d, s_{\text{out}}, b_{\text{out}}$; an unconstrained three-vector normalized to the unit gauge u; and one log-scale whose exponential is α — 37 scalar parameters, of which u and α serve only the structural terms below. Its internal parameterization is not part of the deployed encoding. The final affine map instead obtains its realized exposure flow directly from Be_1.

The frozen VAE is kept in the computation graph: codes are clamped to $[0, 1]$, mapped to $[-1, 1]$, encoded and decoded in 32-bit floating point, and mapped back before analytic inversion. VAE parameters receive no gradient. One exposure from $\mathcal{E}_{\text{fit}} = \{0, -4, -2, +2, +4\}$ EV is sampled uniformly at each step.

For a VAE-round-tripped reconstruction $\hat{x}$, the intermediate fitting objective is

$$\mathcal{L}_{\text{fit}} = \underbrace{\mathcal{L}_{\log} + \mathcal{L}_{uv}}_{\mathcal{L}_{\text{rt}}} + \underbrace{\mathcal{L}_{\text{eq}} + 0.25\mathcal{L}_{\text{amp}}}_{\mathcal{L}_{\text{struct}}} + 10\mathcal{L}_{\text{range}}. \quad (6)$$

The three groups respectively measure VAE round-trip fidelity, regularize the exposure–chroma structure, and keep codes inside the useful VAE input range. The weights of $\mathcal{L}_{\log}$, $\mathcal{L}_{uv}$, and $\mathcal{L}_{\text{eq}}$ are one. Specifically,

$$\mathcal{L}_{\log} = \text{mean}\|\log(1 + \hat{x}) - \log(1 + x)\|_2^2,$$

and $\mathcal{L}_{uv}$ is the mean Euclidean $u'v'$ error over pixels with ground-truth $Y \geq 10^{-3}$. Thus the compact notation in the main paper denotes a masked per-pixel mean, rather than an unnormalized tensor norm. For $e \in \mathcal{E}_{\text{fit}}$, $\mathcal{L}_{\text{eq}}$ penalizes

$$\text{mean}_m\|g(2^e x) - g(x) - e\alpha u\|_2^2,$$

where g is the intermediate fitting map, u is its learned unit gauge, and $\alpha > 0$ is its learned displacement scale. The term $\mathcal{L}_{\text{amp}}$ penalizes the change with exposure of the code distance between an image and its chromatically perturbed counterpart. The same valid-luminance mask $m = \mathbf{1}[Y \geq 10^{-3}]$ is used. Finally,

$$\mathcal{L}_{\text{range}} = \text{mean}\{\,\text{ReLU}(g(x) - 0.95)^2 + \text{ReLU}(0.05 - g(x))^2\}.$$

We use AdamW for 15,000 steps with learning rate 2×10^{-3}, weight decay 10^{-4}, batch size 4, cosine decay to zero, unit gradient clipping, and seed 0.

Stage 2: affine distillation (closed form). The distillation targets are the code vectors $g(x)$ of the converged stage-1 map on 120,000 RGB samples at $\{0, -2, -4, -6, -8\}$ EV (600,000 fitting rows); Eq. (3) is fitted to these targets by least squares. The -6 and -8 EV rows evaluate the fixed target map outside its optimization ladder; they do not extend the first-stage optimization. The toe is frozen, the first column of B is swept over $c \in \{.5, .75, 1, 1.5, 2, 2.5, 3\}$, and the bias is recentered using 40,000 additional samples.

Since exposure moves the code along Be_1 with $\alpha = \|Be_1\|_2$ (Eq. (1)), scaling that column by c rescales the per-stop displacement relative to the fixed chromatic columns, trading round-trip reconstruction against deep-exposure chromatic separation; Table 1 lists the gate thresholds and the two passing candidates of this sweep. The selected c is not a separate deployed parameter: it is absorbed into the frozen matrix of Eq. (3), whose printed first column already contains it.

Selection is performed before downstream task evaluation. Reconstruction uses a fixed 16-clip encoding-fit validation panel. OOD survival is the ratio of held-out chromaticity R^2 at -8 EV to that at EV 0 on a fixed 12-clip ARRI–UPIQ OOD survival panel. Grid identification is the four-fold texture-held-out classifier defined below. The projection score is a group-held-out linear probe on absolute top-24-PCA perturbation coordinates; the separation AUC uses the fraction of perturbation energy assigned to independently estimated luminance and chroma subspaces.

No downstream SDR-to-HDR prediction or task metric enters encoding selection; the selection and main-paper reconstruction values estimate different panels.

Encoding selection predates the downstream subset aggregation, but two of the 12 OOD-survival-panel clips also occur in the frozen ARRI subset of the test set. Thus no downstream output or metric is used for selection, but encoding selection and downstream evaluation are not strictly image-disjoint. We treat this gate as design selection and do not claim an untouched-image test for representation selection.

Chromatic-Loss Ablation

We rerun both fitting stages (stage-1 optimization and stage-2 affine distillation), the scale sweep, the frozen selection rule, and 6,500-step adapter training with the chromatic round-trip term $\mathcal{L}_{uv}$ removed. All other loss terms and settings are fixed. This fit produces no candidate that passes the selection rule; for diagnostic downstream readout only, its best-reconstruction candidate is exported with the gate explicitly waived.

Removing $\mathcal{L}_{uv}$ costs 1.73 dB PU21 and increases mean $\Delta u'v'$ from .0140 to .0219. The ablation identifies $\mathcal{L}_{uv}$ as a load-bearing term in the EquivFlow fitting objective.

SDR-to-HDR System Details

Data Construction

HDR targets and split. The training corpus contains 5,780 crops from 1,015 parent sequences: 2,580 crops from 215 temporal windows extracted from 94 *Tears of Steel* shots,

candidate	recon. ↑	OOD surv. ↑	grid$_{-8}$ ↑	exposure R^2 ↑	projection ↑	separation AUC ↑
gate	35.99	$0.8\times$ fit ref.	.90	.80	.89	.85
$c = 2.5$	36.08	1.092	.979	.869	.985	1.000
$c = 3$	**36.28**	1.102	.979	.864	.999	1.000

Table 1: Encoding-selection gates and the two passing candidates. Candidates differ only in the exposure-column scale c (Stage 2); the predefined highest-reconstruction rule selects $c = 3$. These are selection-panel scores; the separate main-paper probe panel gives 34.38 dB for the selected encoding.

$\mathcal{L}_{uv}$	PU21-PSNR ↑	$\Delta u'v'$ ↓	chroma L1 ↓	$\Delta C_T C_P$ ↓	selection
on	**35.95**	**.0140**	**.0431**	**6.37**	pass; EQUIVFLOW
off	34.22	.0219	.0582	10.28	fail; waived

Table 2: Chromatic-loss ablation under the training Reinhard curve on the full HDRTV4K–UPIQ evaluation pool (100 clips, 58 parents). Entries are parent-weighted descriptive means from separately fitted encodings and separately trained adapters. The failed encoding is not admissible; its downstream result is reported only to diagnose the removed term.

and 3,200 crops from 800 PolyHaven clips. The parent-disjoint split contains 5,388 training, 324 validation, and 68 test crops. Each clip is normalized by mapping the 75th percentile of its per-frame 99th-percentile luma to unity.

The two sources differ in provenance. PolyHaven renders are scene-referred and retain highlights above reference white. *Tears of Steel* is a finished 16-bit display grade that we linearize with the inverse sRGB transfer function; its grade has already compressed highlight information. We therefore refer to the combined data as linear Rec. 709 targets, not uniformly as scene-linear radiance.

SDR generation. Each target has four offline SDR variants. Variant 0 is a clean lossless FFV1 render. Variants 1–3 use independently sampled degradation parameters. The loader selects the clean variant with probability 0.25 and otherwise selects uniformly among the three degraded variants, so each stored variant has marginal probability 0.25. The variant generator uses seed $42 +$ sample index. No online exposure, contrast, blur, or codec augmentation is applied during adapter training.

All training variants use extended Reinhard with white point $w = 4$ and no additional exposure offset. For linear luminance L,

$$L_t = \frac{L(1 + L/w^2)}{1 + L},$$
$$L = \frac{w^2}{2}\left[(L_t - 1) + \sqrt{(1 - L_t)^2 + \frac{4L_t}{w^2}}\right]. \quad (7)$$

The luminance ratio is applied to RGB, followed by per-channel clipping, sRGB OETF, and 8-bit quantization. The analytic inverse in Eq. (7) inverts the unclipped continuous luminance curve; clipping and quantization remain information-losing.

For degraded variants, sRGB contrast about 0.5 is sampled uniformly from $[0.85, 1.25]$. Gaussian blur with $\sigma \sim U[1, 3]$ is applied selectively to highlights above 0.98 and/or shadows below 0.02, each independently with probability 0.7. H.264 encoding uses `yuv420p`, integer CRF sampled from 18–32, and the `medium` preset.

Unseen tone curve. The ACES-like evaluation condition applies

$$A(y) = \mathrm{clip}\left(\frac{y(2.51y + 0.03)}{y(2.43y + 0.59) + 0.14}, 0, 1\right) \quad (8)$$

as a luminance-preserving RGB ratio, followed by the same sRGB transfer and 8-bit quantization. This curve is not used to train either internal arm.

Adapter Training and Inference

The base model is the LTX-2.3 22B checkpoint `ltx-2.3-22b-dev.safetensors`; its transformer, VAE, and text encoder remain frozen. IC-LoRA reference tokens use clean ($t = 0$) latents, participate in bidirectional attention, and are excluded from the denoising loss. Rank-32 modules use scale 32 and dropout 0. They are inserted into the q, k, v, and output projections of `attn1` and `attn2`, and into `ff.net.0.proj` and `ff.net.2`.

Each arm is trained once for a fixed 6,500-step comparison budget, rather than a validation-selected checkpoint. We use AdamW at learning rate 2×10^{-4} with unit gradient clipping and a linear factor $1 \to .1$ over a 100,000-step horizon. Training uses four NVIDIA H100 80GB GPUs, per-GPU batch size 1, accumulation 1 (effective batch 4), BF16, gradient checkpointing, and seed 42. Noise follows a shifted logit-normal distribution with standard deviation 1, $\epsilon = 10^{-3}$, a .1 uniform mixture, and the model's sequence-dependent shift.

The positive prompt is "A high dynamic range cinematic scene with bright highlights and deep shadows."; the negative prompt is "worst quality, inconsistent motion, blurry, jittery, distorted." Inference uses seed 42, 30 steps, guidance 4, and spatiotemporal guidance 1 at block 28 in `stg_v` mode. All settings other than the conditioning contract and encoding are shared: EquivFlow conditions on the analytically inverse-tone-mapped proxy, whereas LogC3 uses the display-referred SDR directly.

Evaluation Definitions

Common output domain. All outputs are decoded to linear Rec. 709 before evaluation. PU21 uses 100 nits

per normalized linear unit, clamps absolute luminance to $[0.005, 10, 000]$ nits, applies the banding–glare PU21 encoding normalized to unit peak, and reports PSNR with peak 1.

For chromatic $\Delta C_T C_P$, Rec. 709 is converted to Rec. 2020 and then to PQ ICtCp using the 10,000-nit convention. We use the ITP scaling $(I, T, P) = (I, 0.5C_T, C_P)$ and report the mean

$$720\sqrt{(0.5\Delta C_T)^2 + (\Delta C_P)^2}.$$

This omits the intensity difference by construction.

Luminance-independent chroma metrics. Let $r(x) = x/(\sum_c x_c + \epsilon)$, with $\epsilon = 10^{-4}$ used only to avoid division by zero. Chroma L1 is $\mathrm{mean}\,|r(\hat{x}) - r(x)|$. Mean $\Delta u'v'$ is the Euclidean CIE 1976 chromaticity error over ground-truth pixels with $Y \geq 10^{-3}$. LF drift first applies a spatial Gaussian with $\sigma = 8$ to the chromaticity fields and measures the same masked error. Range drift further restricts evaluation to pixels whose SDR input is neither clipped nor null: all channels are below $254/255$ and the maximum channel exceeds $3/255$. The null region is replaced by the target field before Gaussian filtering to prevent invalid boundaries from leaking into the range mask.

Main-paper frozen-VAE probe. Table 1 in the main paper uses no diffusion training. Its reconstruction column is the mean per-clip log-RGB PSNR after a frozen-VAE round trip on 16 deterministically sampled held-out validation clips (a panel distinct from the encoding-selection panel of Table 1). For the chromatic columns, 12 deterministic center-cropped clips from the ARRI–UPIQ OOD pool are evaluated at $\{0, -2, -4, -6, -8, -10\}$ EV. Only latent tokens whose EV-0 ground-truth luminance is at least 10^{-3} are used. At every exposure and for each encoding separately, a `StandardScaler` followed by ridge regression with penalty 1 is refitted to predict u' and v'. Five-fold `GroupKFold` by parent clip supplies held-out predictions; the reported $R^2_{u'v'}$ is the mean of the two coordinate scores and $\mathrm{surv}_{-k} = R^2(-k)/R^2(0)$. Because held-out R^2 is unbounded below and the numerator can exceed its EV-0 value, survival may be negative or exceed one. It measures relative linear recoverability, not chroma-energy amplification.

Cross-VAE probe. The structural-portability test applies the same fixed encoding to each frozen host VAE over the exposure ladder $\{0, -2, -4, -6, -8\}$. Its chroma grid comprises 12 chromaticities (two radii, 0.03 and 0.06, at six equally spaced hue angles about D65) rendered with four luminance textures at base luminance 0.45. Grid identification is 12-way accuracy; chance is $1/12$. At each host and exposure, a `StandardScaler` plus logistic regression ($C = 1$) is evaluated by four-fold `GroupKFold`, holding out one luminance texture per fold. Thus the table's grid readout is host- and exposure-specific.

For the exposure readout, a host-specific direction is the first right singular vector of all nonzero-EV latent differences. The scalar $\beta(e)$ is their mean projection on that direction, and R^2_β is the no-intercept linear fit of $\beta(e)$ to EV. The reported $\Delta u'v'_{-8}$ is the encoding round-trip chromaticity error at -8 EV. The prespecified non-collapse check uses 16 fixed clips (15 with a nonempty EV-0 luminance mask) and requires each foreign host's round-trip $\Delta u'v'$ to be no more than twice the LTX value at both -6 and -8 EV. This is a within-subject collapse check, not a comparative superiority test.

EV-0-frozen readout. The main paper notes that an EV-0-frozen readout degrades along each host's exposure trajectory. This diagnostic fits the chroma-grid classifier at EV 0 and applies it unchanged at deeper exposures, instead of refitting per exposure as in the grid column of main-paper Table 4. Under the frozen readout, -8 EV grid accuracy falls to .208 (LTX), .083 (SD 2.1, chance level), .167 (Wan 2.1), and .250 (Wan 2.2) from EV-0 values of .979, .667, .958, and .979, while the per-exposure refitted readout remains high (main-paper Table 4). The latent coordinates therefore rotate or rescale along each host's exposure trajectory even though exposure structure and chroma separability persist; the portability claim concerns the latter, not shared latent coordinates.

LEDiff radiometric alignment. LEDiff is run with its official preprocessing and checkpoint. A single method-level scalar gain is fitted once on 24 fixed, parent-distinct training clips and then held fixed for every evaluation set. With luminance y, the gain is

$$g_{\mathrm{LEDiff}} = \exp\big(\mathrm{mean}_{y_{\mathrm{gt}} > 10^{-4}}\,\delta\big) = 0.0294942,$$
$$\delta = \log(y_{\mathrm{gt}} + 10^{-4}) - \log(y_{\mathrm{pred}} + 10^{-4}).$$

No test image, test ground truth, or per-clip gain enters this fit.

Qualitative selection and display. For each evaluation set, the main-paper figure selects the evaluated crop with the largest EqᴜivFlow–LumiVid ΔE_{ITP} margin; it therefore shows selected successes, not representative or median cases. ARRI and HDRTV4K are displayed in the full frame containing that crop, whereas UPIQ contains one evaluated crop. The inset is chosen as the darkest chroma-carrying target region that remains inside one reconstructed crop. A row-specific display gain is applied identically to every column and affects no reported metric. For HDRTV4K, only the evaluated quadrant contributes to quantitative results; the other quadrants are display context.

Statistical Protocol and Full Results

Frozen evaluation subset. The frozen 102-clip test subset of main-paper Table 2 was fixed before any subset aggregation was computed. Table 4 records the per-set selection rule.

Aggregation and inference. Metrics are first averaged within each parent and then averaged with equal weight across parents. The chroma-decomposition panel uses the full HDRTV4K–UPIQ pool of 58 parents (36 HDRTV4K and 22 UPIQ). Confidence intervals are nonparametric parent-cluster bootstrap intervals from 10,000 resamples with seed 42. Pairwise tests are two-sided, parent-level paired Wilcoxon signed-rank tests. For the dataset-wise table, Holm correction is applied within each method pair and curve across three datasets and two primary metrics. For the chroma

	Reinhard			ACES-like		
metric	Δ	95% CI	p value	Δ	95% CI	p value
chroma L1	-0.01083	$[-0.01565, -0.00598]$	1.61×10^{-6}	-0.01317	$[-0.01834, -0.00799]$	2.47×10^{-7}
$\Delta u'v'$	-0.00330	$[-0.00451, -0.00209]$	1.61×10^{-6}	-0.00769	$[-0.00976, -0.00577]$	6.04×10^{-10}
$\Delta C_T C_P$	-0.998	$[-1.388, -0.614]$	4.83×10^{-6}	-1.489	$[-1.926, -1.065]$	1.65×10^{-7}
LF drift†	-0.00417	$[-0.00572, -0.00265]$	4.20×10^{-7}	-0.00768	$[-0.00968, -0.00583]$	1.49×10^{-10}
range drift†	-0.00407	$[-0.00557, -0.00266]$	1.37×10^{-7}	-0.00807	$[-0.01010, -0.00613]$	6.55×10^{-11}

Table 3: Parent-paired chroma differences (EquivFlow−LogC3) on the pooled HDRTV4K–UPIQ panel (58 parent scenes); negative values favor EquivFlow, and every bootstrap interval excludes zero. The three primary rows report Holm-corrected p_{H} within each curve's three-metric family. †Nested $\Delta u'v'$ variants are descriptive and report uncorrected p.

set	selection	clips	parents
ARRI	first 40 clips	40	10
HDRTV4K	first 40 clips	40	19
UPIQ	all parents	22	22

Table 4: Frozen evaluation subset. "First" refers to each set's fixed evaluation ordering; the ACES-like condition uses the same clips.

curve	set	ΔPU21	p_{H}	$\Delta C_T C_P$	p_{H}
reinhard	ARRI	$+0.294$	.465	-0.013	.492
	HDRTV4K	$+1.684$	.312	-1.066	.197
	UPIQ	$+2.485$	.00128	-0.917	.00128
aces-like	ARRI	-0.216	.465	-0.259	.465
	HDRTV4K	-1.282	.00134	-1.718	.00508
	UPIQ	-1.024	.00134	-1.327	.000157

Table 5: Internal pair on the frozen test subset. p_{H} is the two-sided paired Wilcoxon value after Holm correction within the stated six-test family.

panel, correction is applied within each curve across its three primary metrics. There is one trained model per internal arm; bootstrap intervals characterize parent-level sampling uncertainty rather than variation across training runs.

Dataset-wise internal comparison. Table 5 supplies the parent-level tests behind the dataset-wise statements in the main paper. Deltas are EquivFlow−LogC3, so positive PU21-PSNR and negative $\Delta C_T C_P$ favor EquivFlow. The UPIQ chromatic difference is significant under the training curve; HDRTV4K and UPIQ are significant under the unseen curve. ARRI does not separate under either curve, and HDRTV4K does not separate under the training curve.

Chroma-panel statistics. Table 3 gives the paired differences underlying the main paper's chroma panel for the fixed 6,500-step EquivFlow checkpoint. Units are parent scenes ($n = 58$); each entry is EquivFlow−LogC3 with a 10,000-sample bootstrap 95% interval. The three primary metrics use Holm-corrected p_{H}. The two nested filtered forms of $\Delta u'v'$ are descriptive, use uncorrected p, and lie outside the correction family.

Interpretive boundary. Under matched Reinhard, EquivFlow has higher aggregate-fidelity point estimates on the frozen subset. That aggregate advantage does not persist under the ACES-like swap. The result that remains consistent across the two tested curves is narrower: the full EquivFlow system has lower luminance-independent chromaticity error than the same-backbone LogC3 system. Because the arms also differ in reference reconstruction, these results attribute the effect to the complete conditioning system, not to the encoding alone.