arXiv:2611.14285v1 [cs.AI] 10 Nov 2026

LM-SimBench: A Unified Framework, Dataset, and Benchmark for Large Language Model Training Simulation

Yu Dong[1,†] Hong Li[1,†] Qiang Sun[1] Xu Zhou[1]
Jinbiao Huo[2] Qixiu Cheng[1,*] Zhiyuan Liu[1,*]

[1]Southeast University [2]The Hong Kong Polytechnic University

dongyu_007@seu.edu.cn hongli09@seu.edu.cn 220243655@seu.edu.cn xuzhou@seu.edu.cn
jinbiao.huo@polyu.edu.hk qixiu.cheng@seu.edu.cn zhiyuanl@seu.edu.cn

[†]These authors contributed equally.

[*]Corresponding authors: Qixiu Cheng and Zhiyuan Liu.

Abstract

Evaluating candidate parallelization strategies through distributed training runs makes strategy selection prohibitively expensive at scale. Performance simulation offers a lower-cost alternative, yet controlled comparisons among existing simulators remain difficult because they differ in the training stages they model and the workloads on which they are evaluated. We introduce LM-SimBench[1], a unified framework, profiling dataset, and benchmark for evaluating LLM training simulators. The framework decomposes simulation into functional stages from model compilation to schedule-level aggregation, connecting execution representations and cost estimators to measurable evaluation targets. Following this framework, we profile 23 model settings from 10 dense and Mixture-of-Experts (MoE) LLM families on an Ascend Neural Processing Unit (NPU) cluster. The resulting dataset contains 348 feasible configurations that jointly vary five parallelism dimensions and activation recomputation, with six views spanning configuration, rank, operator, event, and kernel granularity. Evaluations of representative training simulators demonstrate how LM-SimBench supports comparisons of prediction accuracy, assessment of configuration selection, and diagnosis of prediction errors. Matched GPU–NPU experiments validate the cross-platform applicability of the collection and evaluation protocol and demonstrate NPU-guided GPU configuration selection.

1. Introduction

Distributed LLM training requires selecting parallel configurations that fit device memory and minimize iteration time. Performance depends on coupled choices in data, tensor, pipeline, context, and expert parallelism (DP/TP/PP/CP/EP), together with activation recomputation, scheduling, and network topology [15, 27]. Evaluating these choices through full training runs is costly. For example, DeepSeek-V3 reports 2.788 million H800 GPU-hours for its official training, corresponding to an estimated US$5.576 million at US$2 per GPU-hour [11]. OpenAI likewise notes that extensive model-specific tuning is infeasible for training runs at GPT-4's scale [28]. Given a model, training strategy, and hardware platform, performance simulation estimates runtime and memory requirements to guide configuration selection before deployment (Figure 1).

Existing simulators are evaluated using different hardware platforms, software stacks, workloads, and validation targets. For example, vTrain validates training iteration times on A100 systems using Megatron-DeepSpeed [6]; ASTRA-sim2.0 validates its analytical network backend against NCCL All-Reduce measurements on V100 GPUs [41]; and SimAI evaluates GPT-3 and LLaMA workloads on A100/H100 clusters with RoCEv2 networks [39]. These differences make reported prediction errors difficult to compare directly across studies, motivating shared evaluation tasks and matched execution conditions. Evaluating training simulators involves iteration-time accuracy, internal computation and communication costs, and configuration-selection quality. These complementary targets call for end-to-end measurements linked to intermediate profiles from the same training runs.

To address these evaluation challenges, we introduce LM-SimBench, which brings together a unified simulation framework, a profiling dataset collected with MindSpeed-LLM [4] on Ascend NPUs [24], and a hierarchical benchmark. The framework identifies simulator inputs and outputs, the dataset supplies corresponding measurements, and the benchmark evaluates prediction accuracy, error diagnosis,

[1]Code: GitHub; dataset: Hugging Face.

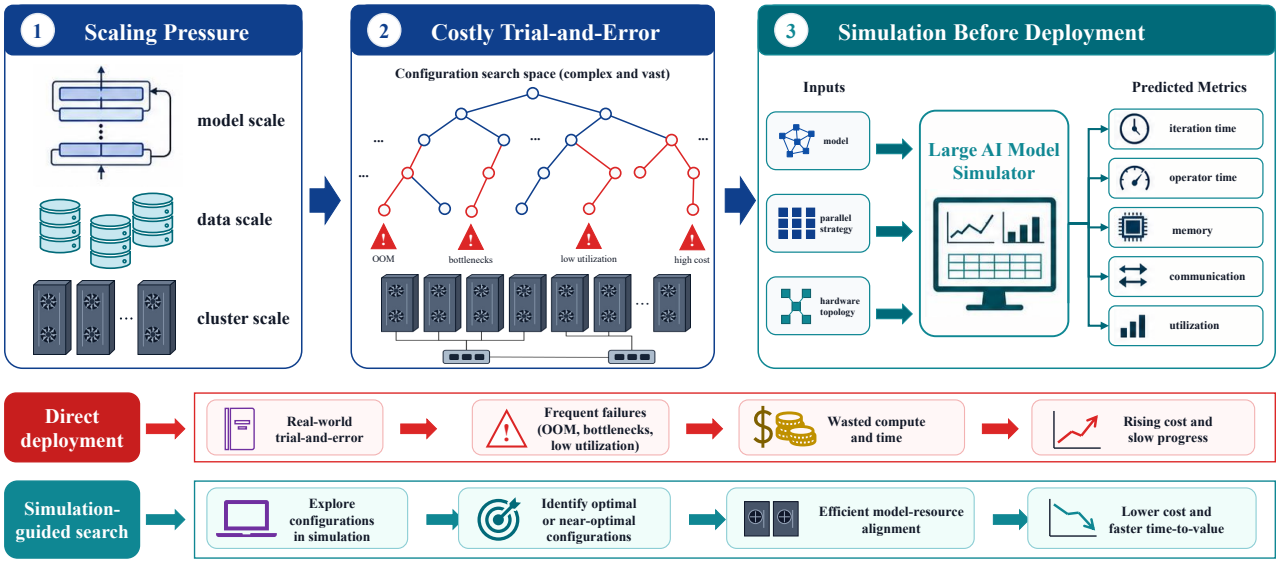

Figure 1. Simulation for model–resource alignment. Effective alignment improves utilization and reduces cost; poor alignment creates bottlenecks and resource waste. Simulation estimates these effects before deployment, reducing the need for repeated full-scale trials.

and configuration selection.

Our contributions are:

- **A unified simulation framework.** We connect model compilation and device binding with execution modeling, computation and communication cost estimation, and schedule-level aggregation, relating simulator design choices to evaluation targets.
- **A multi-granularity profiling dataset.** We release 348 feasible configurations across 23 model settings from 10 dense and MoE LLM families, covering five parallelism dimensions and recomputation.
- **A hierarchical evaluation benchmark.** System-, component-, and strategy-level tasks connect prediction accuracy with error diagnosis and configuration selection. Calculon and vTrain demonstrate comparisons of prediction accuracy, simulation cost, and diagnostic outputs.
- **Cross-platform protocol validation.** Matched measurements for Qwen3-0.6B and Qwen3-1.7B on NVIDIA V100 GPUs and Ascend NPUs demonstrate cross-platform evaluation and NPU-guided GPU configuration selection.

2. Related Work

Distributed LLM training simulators. Existing simulators estimate training time and, in some cases, device-memory requirements under different parallel configurations [18, 30]. Table 1 summarizes representative tools by their operator-time estimation approach, prediction outputs, and availability of parameter-calibration tools. Analytical estimators apply cost equations with hardware parameters [18, 26]; profiling-based estimators use measured operator or kernel latencies [6, 12, 30]; hybrid estimators combine measured costs with analytical or mechanistic models [39]. Cost estimation and execution representation are separate modeling dimensions. For example, vTrain [6] combines profiled kernel times with analytical inter-node communication estimates, while SimAI [39] combines kernel profiling with detailed communication simulation. Lumos [23] develops trace-driven performance modeling for large-scale LLM training.

Profiling data and execution traces. Existing resources provide measurements and workload descriptions at different granularities (Table 2). NeuSight [21] provides operator and end-to-end latency measurements across Transformer models and GPU platforms. AcmeTrace [17] and Alibaba GPU traces [1] record production workloads and resource usage for cluster-level analysis. Chakra [36] provides a standardized execution-trace representation, while AICB [2] generates parameterized workloads for communication simulation. LM-SimBench links end-to-end measurements to per-rank computation, communication, and memory profiles from the same training runs across controlled parallel configurations, providing a common measurement basis for simulator calibration and evaluation.

Training benchmarks and simulator evaluation. MLPerf Training [25] standardizes workloads, quality targets, and reporting for evaluating training systems. Simulator evaluation addresses a different question: how

Table 1. Representative simulators for distributed LLM training. Cost estimation refers to operator computation time: Analytical uses hardware-based cost equations; Profiling-based directly uses measured operator or kernel latencies; Hybrid uses analytical models calibrated with operator profiling data. Memory denotes device-memory usage; Calibration tool indicates the availability of parameter-fitting tools.

Simulator	Cost estimation	Time	Memory	Calibration tool
Calculon [18]	Analytical	✓	✓	×
SimuMax [26]	Analytical	✓	✓	✓
Proteus [12]	Profiling-based	✓	✓	×
vTrain [6]	Profiling-based	✓	×	×
Phantora [30]	Profiling-based	✓	✓	×
ASTRA-sim 2.0 [5, 41]	Hybrid	✓	✓	×
SimAI [39]	Hybrid	✓	×	×
Echo [14]	Hybrid	✓	×	✓

ASTRA-sim's memory tracking requires tensor metadata; Echo's calibration targets its slowdown predictor.

Table 2. Related public artifacts and the primary LM-SimBench collection. The comparison distinguishes measured profiling datasets, cluster traces, and tools for representing or generating workloads.

Artifact	Type / workloads	Released signals or specifications	Parallelism coverage	Hardware
NeuSight [21]	Profiling dataset; 5 dense + 1 MoE Transformer	Operator and end-to-end latency measurements	DP/TP/PP sweeps	NVIDIA / AMD GPUs
AcmeTrace [17]	Cluster trace; production LLM jobs	Scheduler logs, resource monitoring, timing, and failures	Job-level observations	A100 GPUs
Alibaba GPU traces [1]	Cluster trace; production GPU jobs	Job-level resource and timing records	Job-level observations	GPUs
Chakra [36]	Execution-trace format	Graph schema and example traces; collection and replay tools	Representation schema	Hardware-agnostic
Lumos [23]	Modeling toolkit; GPT-3 family	Kineto and execution-trace tooling	Modeling methodology	H100 GPUs
AICB [2]	Workload generator; dense and MoE LLMs	Parameterized workload specifications	DP/TP/PP/SP/EP specifications	GPU-agnostic
LM-SimBench	**Profiling dataset**; 23 settings, 10 dense/MoE families	**Six linked views**: iteration time, peak memory, time breakdown, operator latency, communication events, and matrix kernels	**348 configurations**; DP/TP/PP/CP/EP and recomputation	Ascend NPUs

accurately predictions reproduce measured system behavior and how effectively they guide configuration selection. Existing simulators validate their predictions using different workloads, platforms, and performance targets, including training iteration time and collective-communication latency [6, 39, 41]. LM-SimBench organizes evaluation into system-, component-, and strategy-level tasks. These tasks compare overall time and memory predictions with measurements, assess internal computation and communication costs, and evaluate the quality of recommended parallel configurations.

3. A Unified Framework for Large-Model Training Simulation

We describe training simulators through common inputs, outputs, and functional stages. This framework connects their execution and cost-modeling choices to measurable targets for overall accuracy, internal cost estimates, and configuration selection.

3.1. Pipeline and Formal Problem

A simulation instance is $x = (M, P, N)$: M specifies model architecture and workload; P specifies parallelism, micro-batch size, recomputation, precision, fusion, and overlap; and N specifies compute, memory, and interconnect characteristics. RC and noRC denote configurations with and

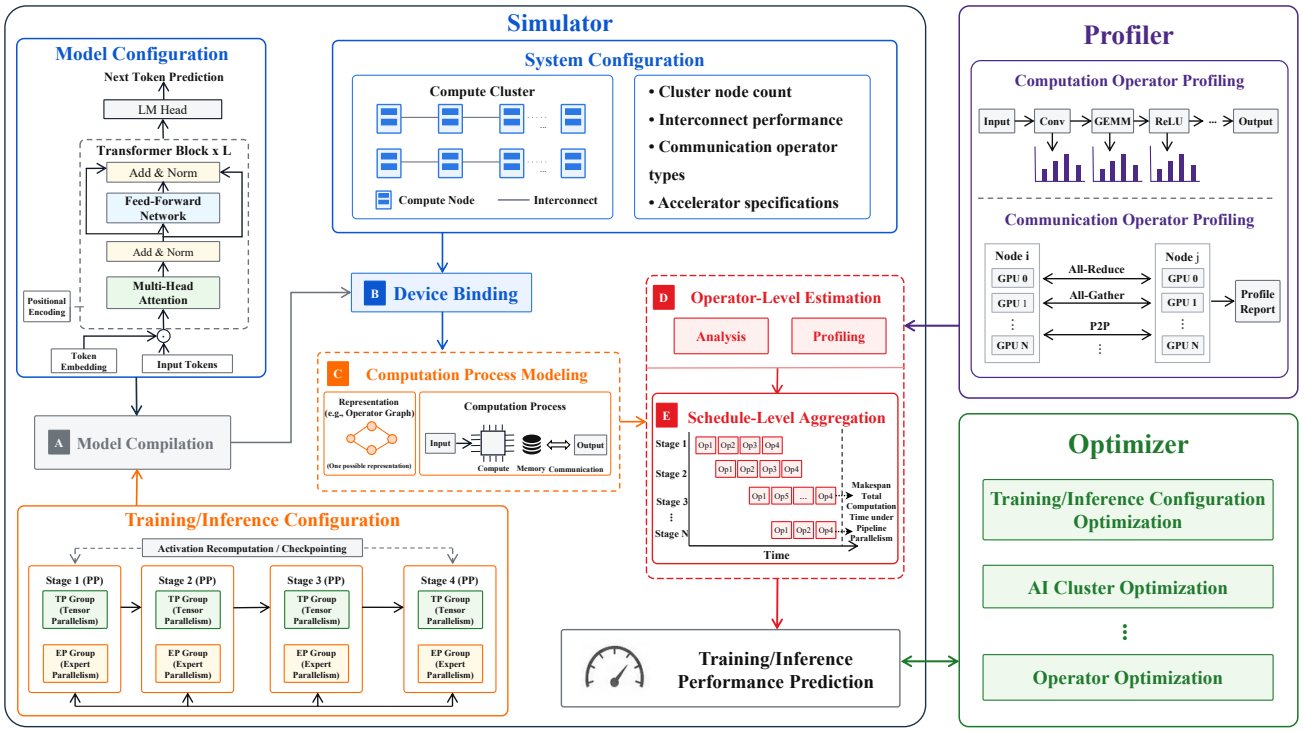

Figure 2. Unified pipeline for large-model training simulation. Model, training, and system configurations drive compilation, device binding, and computation-process modeling. Computation and communication costs from operator-level estimation feed schedule-level aggregation to predict iteration time. Profiling supplies measured costs; the optimizer uses predictions for configuration search.

without activation recomputation, respectively. The simulator predicts

$$\mathcal{S}_\theta(x) = (\widehat{T}, \widehat{\mathbf{V}}, \widehat{C}, \widehat{z}), \tag{1}$$

where θ denotes simulator parameters. Outputs comprise iteration time $\widehat{T}$, per-device peak memory $\widehat{\mathbf{V}} = (\widehat{V}_d)_{d \in D(N)}$, configuration-selection objective $\widehat{C}$, and optional diagnostics $\widehat{z}$. $D(N)$ is the device set.

Figure 2 defines five functional stages. Stage A compiles the model under P, Stage B assigns devices and parallel groups, and Stage C constructs the execution representation:

$$\begin{aligned} \bar{M} &= \mathcal{C}(M, P), \\ \Pi &= \mathcal{B}(\bar{M}, P, N), \\ R &= \mathcal{R}(\bar{M}, \Pi, P, N). \end{aligned} \tag{2}$$

Here, $\bar{M}$ is the compiled model and Π records placement and communication groups. R represents execution through analytical phases, an operator-level directed acyclic graph (DAG), a profiled trace, or a hybrid schedule.

Stage D estimates the costs of computation and communication units in R. Let $\mathcal{U}(R) = \mathcal{U}_{\text{op}}(R) \cup \mathcal{U}_{\text{comm}}(R)$ denote these units. Their predicted latencies are

$$\widehat{t}_u = \begin{cases} \mathcal{F}_{\text{op}}(\phi_u, \Pi, N), & u \in \mathcal{U}_{\text{op}}(R), \\ \mathcal{F}_{\text{comm}}(\phi_u, \Pi, N), & u \in \mathcal{U}_{\text{comm}}(R), \end{cases} \tag{3}$$

where ϕ_u describes tensor shape, FLOPs, transferred bytes, primitive, and group size. Estimators may be analytical, structure-aware, profiling-based, or learned. The profiler supplies computation and communication measurements for cost estimation.

Stage E aggregates unit costs under execution dependencies, resource constraints, and computation–communication overlap to predict iteration time. Time and memory predictions are expressed as

$$\begin{aligned} \widehat{T} &= \mathcal{A}\big(R, \Pi, \{\widehat{t}_u\}_{u \in \mathcal{U}(R)}, N\big), \\ \widehat{\mathbf{V}} &= \mathcal{V}(R, \Pi, N). \end{aligned} \tag{4}$$

The aggregation function $\mathcal{A}$ implements phase rules, constrained graph scheduling, or trace replay, determining which scheduling and overlap effects are represented. The memory estimator $\mathcal{V}$ uses tensor layouts and lifetime assumptions in R and Π, together with N.

Table 3 compares execution representations (Stage C) and cost estimators (Stage D), linking simulator design to calibration requirements and intermediate evaluation targets.

3.2. Connecting Simulation to Evaluation

LM-SimBench evaluates overall predictions, internal cost estimates, and configuration selection using linked measure-

Table 3. Representative simulators in the unified framework: execution representations in Stage C and computation and communication cost estimators in Stage D.

Simulator	Execution	Computation	Communication
Calculon [18]	Analytical phases	Analytical	Mechanistic
SimAI [39]	Workload graph	Profiled kernels / analytical scaling	Mechanistic / packet-level
Optimus [20]	Operator graph	Analytical / structure-aware	Mechanistic
ASTRA-sim 2.0 [41]	Workload graph	Profiled / roofline	Mechanistic / packet-level
vTrain [6]	Execution graph	Profiled CUDA kernels	Profiled / analytical
LLMEmu [43]	Hybrid	Profiling-based	Trace-driven

ments from the same training runs:

$$\mathcal{D} = \{(x_i, y_i^\star, z_i^\star)\}_{i=1}^{n}, \tag{5}$$

where $x_i = (M_i, P_i, N_i)$, $y_i^\star = (T_i^\star, \mathbf{V}_i^\star)$ contains measured time and per-device peak memory, and $z_i^\star$ contains rank-level breakdowns, operator latencies, communication events, and matrix-kernel records. Configuration and rank identifiers link these views.

System-level tasks compare time and memory predictions with $y_i^\star$. Component-level tasks compare intermediate estimates with $z_i^\star$ at matching units and aggregation granularity: computation and communication profiles assess Stage D, while time breakdowns help assess scheduling and overlap in Stage E. Strategy-level tasks compare candidate configurations under a fixed model and system, assessing selection quality against measured iteration times (Section 5).

Configuration selection retains strategies that are valid and predicted to fit on every device:

$$I_{\text{feas}}(M, P, N) = \mathbf{1}\{P \in \mathcal{K}(M, N)\} \cdot \prod_{d \in D(N)} \mathbf{1}\{\widehat{V}_d \le \text{cap}_d\}, \tag{6}$$

where $\mathcal{K}(M, N)$ is the valid strategy space and cap_d is device capacity. Scalar peak memory is $\widehat{V} = \max_{d \in D(N)} \widehat{V}_d$. Ranking retains predicted-feasible candidates; measured memory supplies the feasibility evaluation labels.

4. LM-SimBench Dataset

LM-SimBench contains distributed-training profiles for 23 model settings from 10 dense and MoE LLM families, covering 348 feasible parallel configurations. A model setting specifies the model and training sequence length; its configurations vary DP/TP/PP/CP/EP and recomputation settings. The primary collection uses Ascend 910B3 NPUs with 64 GiB per device and MindSpeed-LLM [4].

Training logs and per-rank Ascend Profiler outputs are processed into six measurement views (Table 4). Shared model-setting, configuration, and rank identifiers connect end-to-end labels to intermediate records from the same run. Processing checks configuration names, rank coverage, units, and schemas. The resulting 26.2 GB of CSV tables and 550 GB of raw profiles support both direct evaluation and trace-level analysis.

Table 4. Six linked measurement views and their evaluation uses.

View	Granularity	Evaluation use
Iteration time	Configuration	Runtime prediction and ranking
Peak memory	Rank	Memory prediction and feasibility
Time breakdown	Rank	Bottleneck diagnosis
Operator latency	Operator/rank	Computation cost modeling
Communication	Event/rank	Communication cost modeling
Matrix kernels	Kernel/rank	Matrix-operator prediction

The dataset also includes paired NVIDIA V100 and Ascend 910B3 measurements for Qwen3-0.6B and Qwen3-1.7B under the same schema. The dataset comprises system telemetry collected from controlled training runs.

5. Benchmark Design

LM-SimBench evaluates system- and component-level predictions, configuration-selection quality, and simulation cost. Cross-platform tasks compare paired measurements and evaluate NPU-guided GPU configuration selection.

5.1. Evaluation Protocol

For each $x = (M, P, N)$, predictions are matched to measured time, memory, and component costs through configuration, rank, and component identifiers. Shared-task comparisons use the same configurations and measurement labels, with sample counts reported for each experiment.

Calibration measurements supply simulator parameters or profiled costs. Simulator comparisons use disjoint calibration and test configurations; calibration-coverage experiments vary the calibration set while fixing the test set. Platform-calibration experiments fit parameters to each

stack's measurements and report the resulting calibration errors.

5.2. System- and Component-Level Prediction

Time and memory. Iteration-time predictions are compared with configuration-level measurements. Peak-memory error is computed per rank for platform calibration and on the maximum across ranks for configuration-level diagnostics. Configuration maxima also determine memory rankings and capacity decisions. Feasibility evaluation reports agreement with measured memory and retention of measured-feasible configurations.

Component costs. Computation targets include matrix-kernel latency per call, stage computation time, and recomputation-stage time. Predictions are paired with measurements at the corresponding operator or stage granularity. Communication targets include event durations and non-overlapped communication time from time breakdowns. Aggregated estimates are compared at matching configuration and rank granularity, with durations expressed in common units.

Prediction metrics are mean absolute error (MAE), symmetric mean absolute percentage error (sMAPE), and R^2. For matched predictions $\widehat{y}_i$ and measurements y_i, sMAPE averages $200|\widehat{y}_i - y_i|/(|\widehat{y}_i| + |y_i| + \epsilon)$, with $\epsilon = 10^{-12}$. Platform calibration reports mean absolute percentage error (MAPE) for iteration time and per-rank memory, with memory ranking evaluated on configuration maxima.

5.3. Configuration Selection and Simulation Cost

For a fixed model and system, configuration selection applies the capacity criterion in Eq. 6 and ranks candidates by predicted iteration time. Spearman ρ and Kendall τ compare predicted and measured orderings. Top-k hit records whether the measured optimum enters the recommended set; Top-k set overlap is the fraction shared with the measured top-k set.

Let $\mathcal{C}$ be the evaluation candidate set, $\mathcal{S}_k$ the recommended k configurations, and $T^\star(c)$ the measured iteration time. Shortlist regret is

$$\text{Regret}\,@k = 100\left(\frac{\min_{c\in\mathcal{S}_k} T^\star(c)}{\min_{c\in\mathcal{C}} T^\star(c)} - 1\right)\%, \qquad (7)$$

which evaluates the fastest measured member of the shortlist. At $k = 1$, it gives the selected configuration's time penalty. Simulation cost is evaluated using the median simulation time per configuration, measured as software wall-clock time in the same CPU and software environment.

5.4. Cross-Platform Evaluation

Paired comparisons match model settings and configuration identifiers across GPU and NPU stacks and report the matched sample count. Time and memory rankings are compared using Spearman ρ, Kendall τ_b, Top-1 agreement, and Top-k set overlap, both pooled and stratified by RC/noRC. Matched RC/noRC pairs assess the direction of time and memory changes.

Matrix problems are matched by batch size, dimensions (m, n, k), and recorded input types. Per-call latencies are pooled across paired configurations and ranks to obtain a median per problem and stack. These medians define latency rankings, compared by Spearman ρ and slowest-problem Top-k overlap. Operator-category and communication-primitive durations are summed across ranks per configuration and normalized into time-share vectors. The vectors are aligned over the union of categories, with zero entries for absent categories. Spearman correlation and Jensen–Shannon divergence (JSD) are computed per configuration and summarized by their medians.

NPU-guided GPU selection ranks the paired candidates by measured NPU time and retains the fastest k. GPU measurements evaluate optimum inclusion and shortlist regret (Eq. 7), using the GPU-fastest configuration in the paired candidate set as the reference.

6. Results

We evaluate simulator accuracy, configuration selection, and simulation cost, examine computation, memory, and communication through intermediate profiles, and assess platform calibration and NPU-guided GPU configuration selection.

6.1. Simulator Accuracy, Selection, and Cost

Setup. We evaluate Calculon and vTrain on Qwen1.5-1.8B at 8K tokens with CP=1 on the NPU stack. Calculon uses analytical cost models, while vTrain combines an execution graph with LM-SimBench profiles supplied through an adapter. Ten configurations calibrate vTrain, and ten disjoint configurations form the test set. We compare iteration-time prediction, configuration ranking, and simulation cost on the six test configurations for which both simulators produce valid predictions.

Accuracy, selection, and cost. The two simulators exhibit distinct accuracy and execution-cost characteristics (Table 5). vTrain achieves 1.86% iteration-time sMAPE and Spearman $\rho = 1.000$, compared with 8.88% and 0.943 for Calculon. Both identify the measured fastest configuration and recover the measured top-three set. Calculon's median simulation time is 1.78 ms per configuration, compared with 25.46 ms for vTrain, a $14.3\times$ speedup.

Calibration coverage. Broader configuration coverage improves vTrain's prediction accuracy. On a fixed set of

Table 5. Calculon and vTrain on six common test configurations of Qwen1.5-1.8B (8K tokens). Simulation time is median wall-clock time per configuration.

Metric	Calculon	vTrain
Iteration-time MAE (s) ↓	1.059	**0.207**
Iteration-time sMAPE (%) ↓	8.88	**1.86**
Iteration-time R^2 ↑	0.630	**0.990**
Ranking Spearman ρ ↑	0.943	**1.000**
Ranking Kendall τ ↑	0.867	**1.000**
Top-1 hit / regret	1 / 0%	1 / 0%
Top-3 set overlap	3/3	3/3
Simulation time (ms) ↓	**1.78**	25.46

Table 6. Calculon on platform calibration configurations. Entries report MAPE (%) / Spearman ρ. Memory error uses per-rank peaks; ranking uses the maximum peak across ranks.

Qwen3 / stack	Time	Memory
0.6B / V100	10.28 / 0.861	15.99 / 0.993
0.6B / 910B3	10.19 / 0.950	13.20 / 0.986
1.7B / V100	11.22 / 0.853	12.56 / 0.993
1.7B / 910B3	7.97 / 0.935	12.55 / 0.994

eight test configurations, calibration with ten TP/DP/PP profiles yields 2.08% sMAPE, compared with 11.77% using six TP-endpoint profiles. The 9.69-percentage-point reduction demonstrates the value of sampling multiple parallelism dimensions for profiling-based simulation.

6.2. Diagnosis with Intermediate Profiles

Computation. Operator- and stage-level profiles characterize computation predictions at their respective granularities. Calculon's matrix-operator predictions achieve $R^2 = 0.817$, with 0.109 ms MAE and 43.34% sMAPE. For vTrain, stage-computation sMAPE is 11.60%, and recomputation-stage sMAPE is 2.75%.

Memory. Calculon's memory model retains 13 of the 20 measured-feasible configurations, with a peak-memory MAE of 6.92 GiB on the retained set. Paired predictions and measurements reveal overestimates of 13–29 GiB in several noRC configurations, linking memory accounting to candidate retention.

Communication. Time breakdowns connect communication estimates with the overlap observed during training. vTrain's non-overlapped communication prediction has 136.20% sMAPE. The associated communication profiles expose the difference between serialized communication estimates and observed overlap, providing a concrete target for refining communication scheduling.

6.3. Simulator Calibration across Platforms

Setup. We collect paired profiles for Qwen3-0.6B and Qwen3-1.7B at 4K tokens on 8×NVIDIA V100 GPUs (32 GiB, Megatron-LM) and 8×Ascend 910B3 NPUs (64 GiB, MindSpeed-LLM). Each model uses 18 DP/TP/PP/recomputation candidates with CP=1, micro/global batch sizes 1/32, and a distributed optimizer. We profile the sixth iteration after five warm-up steps. V100 uses FP16; 910B3 uses BF16, fused kernels, and sequence parallelism. The paired evaluation contains 15 complete configurations for Qwen3-0.6B and 11 for Qwen3-1.7B.

Platform calibration. We apply a common Calculon implementation incorporating Qwen3 architecture, pipeline scheduling, kernel-launch latency, and rank-level memory accounting. For each platform, a single parameter set for compute efficiency, launch latency, communication, and runtime memory is calibrated jointly across both models and their candidate configurations.

On the calibration configurations, iteration-time MAPE ranges from 7.97% to 11.22%, with ranking $\rho = 0.853$–0.950; peak-memory ranking reaches $\rho = 0.986$–0.994 (Table 6). Within each of the four model×RC/noRC groups, GPU measurements, NPU measurements, and both calibrated simulators identify the same Top-1 configuration.

6.4. Cross-Platform Configuration Selection

NPU-guided GPU selection. Ranking the paired configurations by measured NPU iteration time produces a three-candidate shortlist containing the GPU-fastest configuration for both models (Table 7; Figure 3). This reduces the candidate sets from 15 to 3 and from 11 to 3. Selecting the fastest shortlisted configuration using GPU measurements yields zero Top-3 shortlist regret in both cases. Direct selection of the NPU-fastest configuration gives GPU regret of 0.0% and 6.9%, respectively.

Memory and recomputation trends. Cross-platform peak-memory rankings achieve $\rho = 0.982$ and 0.991, with identical three lowest-memory configurations. Across all nine matched RC/noRC pairs, recomputation increases iteration time and reduces peak memory on both platforms. These consistent trends support memory-oriented candidate screening across the two stacks.

Component-level execution characteristics. Matrix-problem latency rankings agree across platforms, with $\rho = 0.983$ and 0.987 and slowest-problem Top-10 overlaps of 8/10 and 10/10. Communication-primitive time shares have median rank correlations of 0.949 and 0.872.

Table 7. Cross-platform transfer from Ascend 910B3 to NVIDIA V100. Shortlist regret uses GPU time. Correlations compare measured profiles on the two stacks.

Level	Metric	Qwen3-0.6B	Qwen3-1.7B
Coverage	Paired complete configurations	15	11
Strategy	Iteration-time rank ρ	0.718	0.664
	NPU Top-1 selection: GPU regret ↓	0.0%	6.9%
	GPU optimum in NPU Top-3	**Yes**	**Yes**
	NPU Top-3 shortlist: GPU regret ↓	**0.0%**	**0.0%**
Memory	Peak-memory rank ρ	**0.982**	**0.991**
	Lowest-memory Top-3 overlap	3/3	3/3
	RC time/memory direction agreement	6/6 pairs	3/3 pairs
Component	Matrix median-latency rank ρ	**0.983**	**0.987**
	Slowest matrix-problem Top-10 overlap	8/10	10/10
	Communication-primitive median ρ	0.949	0.872
	Operator-category median ρ	0.503	0.531
	Non-overlapped communication share, GPU/NPU	17.0% / 28.5%	10.9% / 30.0%

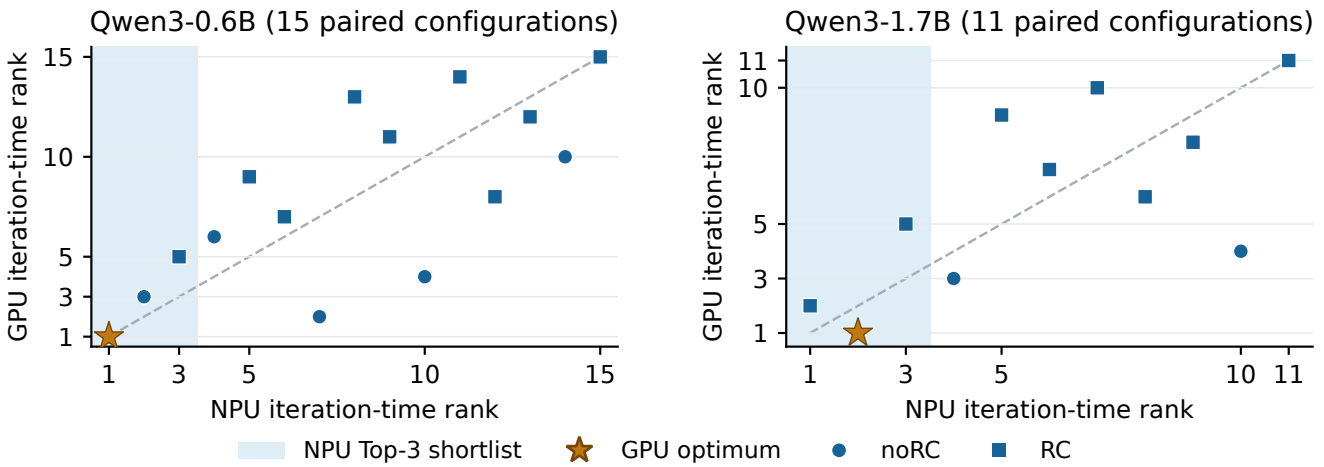

Figure 3. NPU profiles guide GPU configuration selection. Each point is a paired configuration ranked by measured iteration time (rank 1 is fastest). The shaded NPU Top-3 shortlist contains the GPU optimum (star) for both models. The dashed line denotes identical ranks.

Operator-category correlations are 0.503 and 0.531, and non-overlapped communication accounts for 17.0%/10.9% of runtime on V100 and 28.5%/30.0% on 910B3. The profiles identify expensive shared matrix problems and characterize each stack's operator composition and communication overlap.

7. Conclusion

LM-SimBench connects a unified simulation framework, a multi-granularity profiling dataset, and a hierarchical evaluation benchmark for distributed LLM training. Its linked measurements support simulator calibration, prediction assessment, and diagnosis of computation, memory, and communication costs. Evaluations of Calculon and vTrain demonstrate how the same benchmark captures prediction accuracy, configuration-selection utility, and simulation cost. GPU–NPU experiments show that NPU profiles preserve useful memory and matrix-operator rankings and identify shortlists containing the GPU-optimal configurations. Extending workload and accelerator coverage will broaden this common basis for training-system simulation and optimization.

LLM Usage Statement

We used generative AI tools to assist with the writing and presentation of this manuscript: improving its organization, clarity, and academic English, and checking internal consistency and layout. Generative AI tools were not used to formulate the research hypotheses, to design the methodology or the experiments, to implement the released code or the evaluated simulators, to generate or process the measured profiling data, or to interpret the results. All AI-assisted text was reviewed by the authors, who verified every technical claim and reported value and take responsibility for the final content of this work.

References

[1] Alibaba. Alibaba cluster trace program: GPU cluster traces, 2023. Production GPU cluster traces (`cluster-trace-gpu-v2020` and `cluster-trace-gpu-v2023`). Accessed: November 6, 2026.

[2] Alibaba. AICB: An open-source AI communication benchmark suite, 2024. Open-source software repository. Accessed: November 6, 2026.

[3] Alibaba Group. InferSim: A lightweight LLM inference performance simulator, 2025. Open-source software repository. Accessed: April 27, 2026.

[4] Ascend Community. MindSpeed-LLM: A distributed training suite for large language models on Ascend NPUs, 2024. Open-source software repository. Accessed: April 27, 2026.

[5] ASTRA-sim Developers. ASTRA-sim 2.0: Open-source implementation, 2026. Software, revision 518bd51. Accessed: November 6, 2026.

[6] Jehyeon Bang, Yujeong Choi, Myeongwoo Kim, Yongdeok Kim, and Minsoo Rhu. vTrain: A simulation framework for evaluating cost-effective and compute-optimal large language model training, 2024.

[7] James Bergstra and Yoshua Bengio. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(10):281–305, 2012.

[8] Bernd Bischl et al. Hyperparameter optimization: Foundations, algorithms, best practices, and open challenges. *WIREs Data Mining and Knowledge Discovery*, 13(2):e1484, 2023.

[9] Ermao Cai, Da-Cheng Juan, Dimitrios Stamoulis, and Diana Marculescu. NeuralPower: Predict and deploy energy-efficient convolutional neural networks. In *Proceedings of the Ninth Asian Conference on Machine Learning (ACML)*, pages 622–637. PMLR, 2017.

[10] Jaehong Cho, Minsu Kim, Hyunmin Choi, Guseul Heo, and Jongse Park. LLMServingSim: A HW/SW co-simulation infrastructure for LLM inference serving at scale. In *2024 IEEE International Symposium on Workload Characterization (IISWC)*, pages 15–29. IEEE, 2024.

[11] DeepSeek-AI et al. DeepSeek-V3 technical report, 2025.

[12] Jiangfei Duan, Xiuhong Li, Ping Xu, Xingcheng Zhang, Shengen Yan, Yun Liang, and Dahua Lin. Proteus: Simulating the performance of distributed DNN training. *IEEE Transactions on Parallel and Distributed Systems*, 35(10):1867–1878, 2024.

[13] Stefan Falkner, Aaron Klein, and Frank Hutter. BOHB: Robust and efficient hyperparameter optimization at scale. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, pages 1437–1446. PMLR, 2018.

[14] Yicheng Feng, Yuetao Chen, Kaiwen Chen, Jingzong Li, Tianyuan Wu, Peng Cheng, Chuan Wu, Wei Wang, Tsung-Yi Ho, and Hong Xu. Echo: Simulating distributed training at scale. *arXiv preprint arXiv:2412.12487*, 2024.

[15] Aaron Grattafiori et al. The Llama 3 herd of models, 2024.

[16] Zhenhua Guo et al. A survey on performance modeling and prediction for distributed DNN training. *IEEE Transactions on Parallel and Distributed Systems*, 35(12):2463–2478, 2024.

[17] Qinghao Hu et al. Characterization of large language model development in the datacenter. In *21st USENIX Symposium on Networked Systems Design and Implementation (NSDI '24)*, pages 709–729, 2024.

[18] Mikhail Isaev, Nic McDonald, Larry Dennison, and Richard W. Vuduc. Calculon: A methodology and tool for high-level co-design of systems and large language models. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '23)*. ACM, 2023.

[19] Mahmoud Khairy, Zhensheng Shen, Tor M. Aamodt, and Timothy G. Rogers. Accel-Sim: An extensible simulation framework for validated GPU modeling. In *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*, pages 473–486, 2020.

[20] Joyjit Kundu et al. Performance modeling and workload analysis of distributed large language model training and inference. In *2024 IEEE International Symposium on Workload Characterization (IISWC)*, pages 57–67, 2024.

[21] Seonho Lee, Amar Phanishayee, and Divya Mahajan. Forecasting GPU performance for deep learning training and inference. In *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS '25)*, 2025.

[22] Lisha Li, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. Hyperband: A novel bandit-based approach to hyperparameter optimization. *Journal of Machine Learning Research*, 18(185):1–52, 2018.

[23] Mingyu Liang et al. Lumos: Efficient performance modeling and estimation for large-scale LLM training. In *Proceedings of Machine Learning and Systems (MLSys)*, 2025.

[24] Heng Liao et al. Ascend: A scalable and unified architecture for ubiquitous deep neural network computing: Industry track paper. In *2021 IEEE International Symposium on High-Performance Computer Architecture (HPCA)*, pages 789–801, 2021.

[25] Peter Mattson et al. MLPerf training benchmark. In *Proceedings of Machine Learning and Systems (MLSys)*, pages 336–349, 2020.

[26] Moore Threads. SimuMax: A static analytical model for LLM distributed training, 2026. Software, revision 4dcaae5. Accessed: November 6, 2026.

[27] Deepak Narayanan, Mohammad Shoeybi, Jared Casper, Patrick LeGresley, Mostofa Patwary, Vijay Korthikanti,

Dmitri Vainbrand, Prethvi Kashinkunti, Julie Bernauer, Bryan Catanzaro, et al. Efficient large-scale language model training on gpu clusters using megatron-lm. In *Proceedings of the international conference for high performance computing, networking, storage and analysis*, pages 1–15, 2021.

[28] OpenAI. GPT-4 technical report, 2023.

[29] Hang Qi, Evan R. Sparks, and Ameet Talwalkar. Paleo: A performance model for deep neural networks. In *International Conference on Learning Representations (ICLR)*, 2017.

[30] Jianxing Qin, Jingrong Chen, Xinhao Kong, Yongji Wu, Tianjun Yuan, Liang Luo, Zhaodong Wang, Ying Zhang, Tingjun Chen, Alvin R. Lebeck, and Danyang Zhuo. Phantora: Maximizing code reuse in simulation-based machine learning system performance estimation. In *23rd USENIX Symposium on Networked Systems Design and Implementation (NSDI '26)*, pages 1809–1825, 2026.

[31] Saeed Rashidi, Srinivas Sridharan, Sudarshan Srinivasan, and Tushar Krishna. ASTRA-SIM: Enabling SW/HW co-design exploration for distributed DL training platforms. In *2020 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, pages 81–92. IEEE, 2020.

[32] Vijay Janapa Reddi et al. MLPerf inference benchmark. In *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*, pages 446–459, 2020.

[33] Ananda Samajdar, Yuhao Zhu, Paul Whatmough, Matthew Mattina, and Tushar Krishna. SCALE-Sim: Systolic CNN accelerator simulator, 2019.

[34] Ananda Samajdar, Jan Moritz Joseph, Yuhao Zhu, Paul N. Whatmough, Matthew Mattina, and Tushar Krishna. A systematic methodology for characterizing scalability of DNN accelerators using SCALE-Sim. In *2020 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, pages 58–68, 2020.

[35] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019.

[36] Srinivas Sridharan et al. Chakra: Advancing performance benchmarking and co-design using standardized execution traces. *arXiv preprint arXiv:2305.14516*, 2023.

[37] Jonas Svedas et al. A survey of end-to-end modeling for distributed DNN training: Workloads, simulators, and TCO. *arXiv preprint arXiv:2506.09275*, 2025.

[38] Xiaokai Wang, Shaoyuan Huang, Yuting Li, and Xiaofei Wang. ScaleDL: Towards scalable and efficient runtime prediction for distributed deep learning workloads. *arXiv preprint arXiv:2511.04162*, 2025.

[39] Xizheng Wang et al. SimAI: Unifying architecture design and performance tuning for Large-Scale large language model training with scalability and precision. In *22nd USENIX Symposium on Networked Systems Design and Implementation (NSDI '25)*, pages 541–558, 2025.

[40] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009.

[41] William Won, Taekyung Heo, Saeed Rashidi, Srinivas Sridharan, Sudarshan Srinivasan, and Tushar Krishna. ASTRA-sim2.0: Modeling hierarchical networks and disaggregated systems for large-model training at scale. In *2023 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, pages 283–294. IEEE, 2023.

[42] Li Yang and Abdallah Shami. On hyperparameter optimization of machine learning algorithms: Theory and practice. *Neurocomputing*, 415:295–316, 2020.

[43] Siyuan Yang, Enda Yu, Pingjing Lu, and Dezun Dong. LLMEmu: A lightweight performance emulator for high-fidelity distributed LLM training. *Performance Evaluation*, 170:102526, 2025.

[44] Yang You, Aydın Buluç, and James Demmel. Scaling deep learning on GPU and Knights Landing clusters. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '17)*, pages 1–12. ACM, 2017.

[45] Geoffrey X. Yu, Yubo Gao, Pavel Golikov, and Gennady Pekhimenko. Habitat: A runtime-based computational performance predictor for deep neural network training. In *Proceedings of the 2021 USENIX Annual Technical Conference (USENIX ATC '21)*, pages 503–521. USENIX Association, 2021.

[46] Hengrui Zhang, August Ning, Rohan Prabhakar, and David Wentzlaff. A hardware evaluation framework for large language model inference, 2023.

Appendices

Overview and Resources

These appendices follow the main paper from framework and dataset design to simulator evaluation and cross-platform application. Appendix A places LM-SimBench among profiling resources and performance-modeling tools. Appendix B explains the simulation pipeline and modeling choices, and Appendix C documents the collection, processing, and file schema of the six linked measurement views.

Appendix D reports simulator accuracy, configuration selection, simulation cost, and calibration coverage. Appendix E uses single-node and multi-node cases to connect predictions with rank-, communication-, and operator-level measurements. Appendix F evaluates platform-specific calibration and NPU-guided GPU configuration selection using paired measurements.

Data and code. The resources comprise:

- **Processed dataset:** the six linked measurement views, manifests, and field documentation.
- **Raw profiles:** the original profiling traces for detailed execution analysis.
- **Code:** benchmark evaluation, preprocessing scripts, and usage examples.

Dataset organization, documentation, and licenses are described in Appendix C.

A. Additional Related Work

This section expands Section 2 by relating profiling resources, performance predictors, and training simulators to the evaluation tasks in LM-SimBench. The organizing distinction is between the measurements an artifact provides and the modeling choices a simulator uses: linked measurements allow those choices to be assessed at both the output and component levels.

A.1. Profiling Datasets and Trace Releases

Table 2 compares workload coverage, released signals, parallelism dimensions, and hardware across related artifacts based on their released content. The LM-SimBench entry covers the primary NPU collection; Appendix F reports the matched GPU supplement.

Performance-prediction datasets. NeuSight [21] is the closest artifact in purpose: it releases operator-level and end-to-end latencies for six Transformer models across seven NVIDIA and AMD GPUs to support training-performance prediction. LM-SimBench provides per-rank memory and communication-event measurements, with a configuration sweep covering context and expert parallelism and activation recomputation across dense and MoE models.

Cluster operational traces. AcmeTrace [17] and the Alibaba GPU cluster traces [1] expose production-job behavior through scheduler logs, resource monitoring, timing, and failure records. Their job-level observations support scheduling and cluster-management studies. LM-SimBench instead associates individual training configurations with rank-level and operator-level execution measurements.

Trace formats and modeling toolkits. Chakra [36] standardizes graph-based execution traces and provides collection, generation, and replay tools. Lumos [23] develops trace-driven performance modeling, while AICB [2], the workload generator accompanying SimAI, provides parameterized LLM workload specifications. These tools address representation and modeling; LM-SimBench supplies measured profiles and labels for evaluating predictions at multiple granularities.

A.2. Neural-Network Performance Prediction

Analytical predictors estimate execution costs from model structure, operation counts, and hardware characteristics. Roofline relates arithmetic intensity to compute and memory limits [40]; SCALE-Sim and Accel-Sim model systolic-array and GPU execution [19, 33]. Profiling-based approaches measure representative operators and combine their costs into workload-level estimates. Paleo combines analytical models with calibrated measurements [29], and NeuralPower predicts CNN-layer runtime, power, and energy [9]. Subsequent predictors extend these ideas across accelerators and distributed settings [34, 44, 45]. LM-SimBench provides distributed LLM training measurements for evaluating both aggregate predictions and their intermediate operator costs.

A.3. Large-Model System Simulators

Training simulators combine analytical, graph-based, and measurement-based modeling [16, 37]. Calculon estimates iteration time and memory under Megatron-style parallelism [18, 35], while SimAI connects framework and network modeling for large-scale LLM training [39]. Graph- and task-based systems represent dependencies and scheduling explicitly: ASTRA-sim emphasizes distributed execution, collective communication, and network topology [31, 41], and Optimus uses operator graphs for workload and accelerator analysis [20].

Measurement-based systems use profiles to estimate execution costs. vTrain combines an execution graph with profiled kernel times and profiled or analytical communication estimates [6]; LLMEmu uses GPU-kernel profiles to emulate distributed training on CPU clusters [43]. Habitat and ScaleDL learn performance relationships across hardware or distributed settings [38, 45]. Related serving tools, including LLMCompass, LLMServingSim, and InferSim,

address deployment performance and request-level behavior [3, 10, 46].

LM-SimBench provides common measurements and evaluation tasks for these different modeling choices. Appendix B maps representative training simulators by execution representation, computation estimator, and communication model.

A.4. System Benchmarks and Configuration Search

MLPerf Training and Inference standardize workloads, quality targets, and reporting procedures for system-level comparisons [25, 32]. LM-SimBench complements this role by providing reusable training profiles for time and memory prediction, component-level diagnostics, and configuration ranking.

Configuration search turns performance estimates into training-system decisions. General search methods use random sampling, multi-fidelity allocation, and Bayesian optimization to allocate evaluation budgets [7, 8, 13, 22, 42]. In distributed training, parallelism, recomputation, scheduling, and placement determine runtime and memory requirements. LM-SimBench evaluates this use of simulation through configuration rankings, optimum inclusion, and selection regret, connecting prediction quality to the configurations a user would choose.

B. Unified Simulation Framework

This section expands the five-stage pipeline in Section 3. It follows a simulation instance from its input configurations to execution representation, unit-cost estimates, and aggregated predictions, then relates these choices to representative simulators and configuration selection.

B.1. Simulation Inputs and Execution Pipeline

Simulation inputs. We describe the target scenario with three inputs. The model configuration M records architectural and workload parameters such as layer types, hidden dimensions, sequence length, global batch size, and parameter count. The system configuration N specifies the target hardware, including compute capability, memory capacity, inter-device bandwidth, and topology. The training configuration P specifies data, tensor, pipeline, context, and expert parallelism; micro-batch size; recomputation; offloading; numerical precision; operator fusion; and computation–communication overlap.

Performance prediction problem. Given (M, P, N), the simulator $\mathcal{S}_\theta$ predicts the performance of one training iteration. Following Eq. 1, its outputs are iteration time $\widehat{T}$, per-device peak memory $\widehat{\mathbf{V}} = (\widehat{V}_d)_{d \in D(N)}$, a configuration-selection objective $\widehat{C}$, and optional diagnostics $\widehat{z}$. Below, $\widehat{T}$ is the representative metric, and $\widehat{V} = \max_{d \in D(N)} \widehat{V}_d$ denotes scalar peak memory when required.

Stages A–C: constructing the execution representation. For a training configuration P, model compilation transforms the abstract model into

$$\bar{M} = \mathcal{C}(M, P), \tag{8}$$

where $\mathcal{C}(\cdot)$ partitions or replicates components according to the selected parallel strategy. Tensor parallelism partitions parameters within operators, pipeline parallelism assigns blocks to stages, and data parallelism replicates the model across groups. Device binding then assigns each component of $\bar{M}$ to devices and communication groups:

$$\Pi = \mathcal{B}(\bar{M}, P, N), \tag{9}$$

where Π records placement and parallel-group membership, allowing the simulator to distinguish intra-node from inter-node communication.

The bound model is then transformed into an execution representation:

$$R = \mathcal{R}(\bar{M}, \Pi, P, N), \tag{10}$$

where R may be an analytical phase decomposition, an operator-level dependency graph, a profiled trace, or a hybrid schedule. It captures the execution order, synchronization constraints, and computation–communication structure induced by pipeline scheduling, tensor partitioning, recomputation, and overlap.

Stages D–E: estimating and aggregating costs. Stage D estimates the latency of each computation or communication unit in $\mathcal{U}(R) = \mathcal{U}_{\text{op}}(R) \cup \mathcal{U}_{\text{comm}}(R)$. Stage E aggregates these estimates under the ordering and resource constraints represented by R and Π. As in Eq. 3, the predicted latency of unit u is

$$\widehat{t}_u = \begin{cases} \mathcal{F}_{\text{op}}(\phi_u, \Pi, N), & u \in \mathcal{U}_{\text{op}}(R), \\ \mathcal{F}_{\text{comm}}(\phi_u, \Pi, N), & u \in \mathcal{U}_{\text{comm}}(R), \end{cases} \tag{11}$$

where ϕ_u records features such as tensor shape, FLOPs, bytes transferred, primitive type, group size, and placement; Π provides placement and group membership; and N provides the required hardware parameters. The estimator may be analytical, structure-aware, profiling-based, or learned. This cost-estimation choice is separate from the execution representation: for example, an operator graph can use either analytical or measured latencies.

Stage E combines these unit costs with dependencies, resource constraints, and overlap to obtain iteration time:

$$\widehat{T} = \mathcal{A}\big(R, \Pi, \{\widehat{t}_u\}_{u \in \mathcal{U}(R)}, N\big)\,. \tag{12}$$

Peak memory is estimated from the execution representation and device placement, as in Eq. 4:

$$\widehat{\mathbf{V}} = \mathcal{V}(R, \Pi, N), \tag{13}$$

Here, $\mathcal{A}(\cdot)$ aggregates time over the execution representation, and $\mathcal{V}(\cdot)$ estimates peak memory from the tensor-layout and lifetime assumptions in R and Π, together with system characteristics in N. The indicator $I_{\text{feas}}(M, P, N)$ from Eq. 6 checks strategy validity and predicted memory against device capacity. Measured memory provides the reference for evaluating this predicted feasibility.

Memory components. The principal tensor-memory components are parameters, activations, gradients, and optimizer states. Their contribution on device d is

$$\widehat{V}_{\text{tensor},d} = \widehat{V}_{\text{param},d} + \widehat{V}_{\text{act},d} + \widehat{V}_{\text{grad},d} + \widehat{V}_{\text{opt},d}, \quad (14)$$

where $\widehat{V}_{\text{param},d}$, $\widehat{V}_{\text{act},d}$, $\widehat{V}_{\text{grad},d}$, and $\widehat{V}_{\text{opt},d}$ denote the memory used by parameters, activations, gradients, and optimizer states on device d, respectively. Their sizes and lifetimes depend on architecture, sequence length, batch size, precision, parallelism, and recomputation or offloading. The memory estimator combines these terms with its runtime-memory accounting to obtain $\widehat{V}_d$; rank-level measurements provide the corresponding evaluation targets.

B.2. Modeling Choices and Simulator Mapping

The following choices determine the execution detail and cost of a simulator. Table 8 maps representative systems by execution representation (Stage C) and computation and communication estimators (Stage D), together with artifact availability and modeling focus.

Large-model abstraction. For the decoder-only Transformers in LM-SimBench, M records the layer count, hidden and feedforward dimensions, attention heads, sequence length, and batch size. These parameters determine the sizes of computation and memory operations; MoE models additionally specify routing and sparsity. Appendix C.1 lists the profiled workloads.

Cluster and hardware abstraction. Cluster modeling reduces the physical environment to three performance-relevant components: compute, memory, and network (Figure 4).

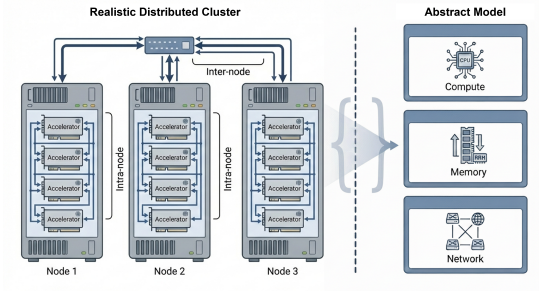

Figure 4. Simplified cluster abstraction. The physical cluster is represented by compute, memory, and network components.

Compute is characterized by data type, peak FLOPs, and empirical efficiency; memory by capacity, bandwidth, and efficiency; and the network by bandwidth, latency, topology, and communication efficiency. The simulator combines peak capabilities with efficiency factors to obtain effective performance.

Measured operator and communication latencies can calibrate or replace analytical cost estimates for the target platform.

Training workflow modeling. Workflow modeling specifies the temporal order of computation and communication after compilation and device binding. In an operator-level graph, nodes represent computation or communication and edges encode data dependencies. A training graph typically includes the forward pass, backward pass, gradient synchronization, and parameter update.

Workflow granularity determines which execution effects can be represented. Graph-based models construct an operator-level DAG for fine-grained timing and dependency analysis. Analytical models reduce execution to phases such as forward, backward, communication, and optimizer update; they are faster but may miss overlap and synchronization. Hybrid models represent critical or communication-heavy segments explicitly and estimate the remainder analytically.

Computation cost modeling. Computation models estimate the latency of operators or phases using analytical formulas, structure-aware simulation, profiles, learned predictors, or combinations thereof. Roofline-style models estimate latency from FLOPs, throughput, and efficiency. Structure-aware models additionally represent hardware hierarchy, tiling, memory movement, and scheduling. Profiling-based methods reuse measured operator latency. Analytical models can thus provide broad coverage while profiles calibrate important or hardware-specific kernels.

Communication cost modeling. Communication models estimate the latency introduced by distributed parallelism. Data, tensor, pipeline, context, and expert parallelism induce collective or point-to-point operations such as All-Reduce, All-Gather, Reduce-Scatter, Send, and Receive. Their cost depends on message size, primitive, group size, bandwidth, latency, topology, and overlap with computation.

Mechanistic models estimate collectives from topology, bandwidth, latency, and algorithmic decomposition. Trace-driven models replay measured events and preserve their dependencies; packet-level simulators represent congestion and routing at higher computational cost; and data-driven predictors learn latency from message, topology, and configuration features. The appropriate choice depends on whether the goal is rapid search, workload-specific optimization, or high-fidelity architectural analysis.

B.3. Simulation Outputs and Strategy Search

The predicted iteration time and per-device peak memory provide the efficiency and feasibility criteria for comparing candidate configurations. The diagnostic outputs connect

Table 8. Representative training simulators characterized by execution representation (Stage C), computation and communication estimators (Stage D), artifact availability, and modeling focus.

Simulator	Workflow	Comp.	Comm.	Public artifact	Primary focus
Calculon [18]	A	An	Mech	✓	Analytical iteration-time and memory estimation for LLM training.
SimAI [39]	G	Pr/An	Mech/Pkt	✓	Full-stack large-scale training simulation across framework and network layers.
Optimus [20]	G	An/St	Mech	×	Operator-graph modeling for workload and accelerator microarchitecture exploration.
ASTRA-sim 1.0 [31]	H	An	Mech/Pkt	✓	Network-topology and collective-communication simulation.
ASTRA-sim 2.0 [41]	G	Pr/An	Mech/Pkt	✓	Detailed workload graphs combined with network-level simulation.
vTrain [6]	G	Pr	Pr/An	✓	Graph-based training simulation calibrated by real-system measurements.
LLMEmu [43]	H	Pr	Trace	✓	Kernel-profiled GPU emulation of distributed LLM training on CPU clusters.
Habitat [45]	A	Pr/Data	Data	✓	Machine learning (ML)-based runtime prediction across unseen GPU generations.
ScaleDL [38]	A	Data	Data	×	Data-driven performance prediction for distributed training.
Paleo [29]	A	An	Mech	✓	Early CNN training-time estimator for GPU clusters.

G: graph-based; A/An: analytical; H: hybrid; St: structure-aware; Pr: profiling-based; Data: data-driven; Mech: white-box mechanistic; Trace: trace-driven; Pkt: packet-level.

these predictions to the operator and communication measurements used in component-level evaluation.

Configuration-selection objective. Iteration time and per-device memory can be combined in a task-specific cost function. Consistent with the mapping $\mathcal{S}_\theta(M, P, N) \to (\widehat{T}, \widehat{\mathbf{V}}, \widehat{C}, \widehat{z})$, we write

$$\widehat{C} = C(\widehat{T}, \widehat{\mathbf{V}}, N), \tag{15}$$

where N contains the system characteristics. In the benchmark of Section 5, $\widehat{C} = \widehat{T}$: configurations are ranked by predicted iteration time after the per-device capacity check. The general formulation also accommodates resource cost or other objectives when required by an application.

Strategy search. Configuration search minimizes predicted cost over valid strategies that pass the predicted feasibility check:

$$\min_{\substack{P \in \mathcal{K}(M,N) \\ I_{\text{feas}}(M,P,N)=1}} \widehat{C}. \tag{16}$$

For the DP/TP/PP setting with CP=1 and no expert parallelism, resource allocation and recomputation satisfy

$$P_{tp} \cdot P_{pp} \cdot P_{dp} \le N_{\text{proc}}, \tag{17}$$

$$P_{rc} \in \{0, 1\}, \tag{18}$$

Here P_{tp}, P_{pp}, and P_{dp} are the tensor-, pipeline-, and data-parallel degrees; P_{rc} is the recomputation switch; and N_{proc} is the number of available processors. With M and N fixed, these choices define the candidate training configurations. Table 9 groups the inputs that specify a simulation instance; the benchmark varies the training dimensions within each model setting and system.

Table 9. Configuration inputs and their roles in simulation and search.

Input	Representative parameters	Role
Model and workload M	Architecture, layer dimensions, attention heads, sequence length, global batch size	Defines the training workload.
Training configuration P	DP/TP/PP/CP/EP, micro-batch size, recomputation, precision, optimizer-state partitioning, fusion, and overlap	Defines candidate execution strategies for a fixed workload and system.
System configuration N	Device count, compute capability, memory capacity and bandwidth, interconnect, and topology	Defines resources and capacity constraints.

C. LM-SimBench Dataset Details

This section provides the collection and storage details for the dataset introduced in Section 4. Shared model-setting, configuration, and rank identifiers connect the six measurement views from each training run. We describe their coverage, collection environment, processing pipeline, and file schema so that the same records can be used for simulator evaluation and execution analysis.

C.1. Workload and Configuration Coverage

Table 10 summarizes the primary NPU collection, and Table 11 gives the per-model coverage. For each model setting, the latter lists the number of feasible configurations and the observed degrees of each parallelism dimension. A model setting specifies a model and its training sequence length; a parallel configuration combines parallelism degrees and a recomputation setting. Dataset size excludes raw profiler dumps.

Feasible coverage varies with architecture, parameter count, context length, memory capacity, and cluster size. Smaller settings therefore admit broader DP/TP/PP/CP combinations than some larger or more constrained settings. For MoE workloads, we additionally report expert parallelism because it changes both computation placement and communication.

The release identifiers in Table 11 encode the model and training sequence length. In the case studies, `qwen15_1.8b_8k` denotes Qwen1.5-1.8B at 8K tokens, and `qwen25_32b_32k` denotes Qwen2.5-32B at 32K tokens. The cross-platform identifiers `qwen3_0.6b_4k` and `qwen3_1.7b_4k` denote Qwen3-0.6B and Qwen3-1.7B, respectively, both at 4K tokens.

The cross-platform supplement pairs V100/Megatron-LM measurements with a newly collected 910B3/MindSpeed-LLM reference for two Qwen3 workloads. Its candidate names match a CP=1 subset of the primary collection; its paired measurements and comparisons are documented in Appendix F.

C.2. Profiling Platform and Software Stack

The primary collection uses Ascend 910B3 NPUs with MindSpeed-LLM. Table 12 records the host, accelerator, interconnect, and software versions used for these runs.

C.3. Processing Pipeline and Quality Control

Algorithm 1 converts per-rank Ascend Profiler outputs into the six views in Table 4. Peak memory, time breakdowns, operator latencies, communication events, and matrix-kernel records are extracted from the corresponding profiler files. Configuration-level iteration time is derived by summing the profiler-reported `Stage` and `Bubble` durations for each rank and taking the maximum across ranks, with microseconds converted to milliseconds. Shared model-setting, configuration, and rank identifiers link the configuration-level labels to their underlying execution records.

Algorithm 1 Data processing pipeline for LM-SimBench

Require: Raw profiler directory $\mathcal{P}_{\text{raw}}$
Ensure: Structured dataset directory D
Initialize the structured output directory D
for each model m in $\mathcal{P}_{\text{raw}}$ **do**
  Identify feasible parallel configurations of m
  for each configuration c of m **do**
    Parse DP, TP, PP, CP, EP, and recomputation settings from c
    for each rank directory r under c **do**
      Read profiler metadata and obtain the rank identifier
      Extract memory, operator, time-breakdown, and communication records
      Extract matrix-kernel features from kernel-level profiler files
    end for
    Sum `Stage` and `Bubble` for each rank; take the maximum across ranks and convert to milliseconds
    Aggregate rank-level records into structured CSV files
  end for
end for
Normalize field names, rank identifiers, units, and directory layout
Generate manifest records for coverage and size checks
return Structured dataset directory D

Table 10. Scale and coverage of the processed LM-SimBench dataset.

Item	Value
Model settings	23
Parallel configurations	348
Data format	CSV
Data modality	Tabular profiling records
Model coverage	Dense and MoE models from 10 LLM families
Configuration dimensions	DP, TP, PP, CP, EP, recomputation
Profiling views	Iteration time, peak memory, time breakdown, operator latency, communication, matrix kernels
Processed dataset size	26.2 GB
Raw profiles	550 GB

Table 11. Observed parallelism degrees and configuration counts for each model setting in LM-SimBench.

Model-setting identifier	Configs	DP	TP	PP	CP	EP	Recompute
gemma2_27b_8k	3	1/2	4/8	2	1/2	–	RC
gemma2_9b_8k	5	1/2	2/4	2	1/2	–	RC/noRC
glm4_9b_8k	6	1/2	2	2/4	1/2	–	RC/noRC
internlm25_1point8b_32k	16	1/2/4	1/2/4	1/2/4	2/4/8	–	RC/noRC
internlm25_20b_32k_16p	9	1/2	2/4/8	2/4/8	1/2/4	–	RC
internlm25_7b_32k	5	1/2	2/4	2/4	1/2	–	RC/noRC
internlm3_8b_8k	7	1/2	1/2	2/4	2/4	–	RC/noRC
mistral_7b_32k	11	1/2/4	1/2/4/8	1/2/4/8	1	–	RC/noRC
phi35_mini_4k	38	1/2/4/8	1/2/4/8	1/2/4/8	1/2/4	–	RC/noRC
phi35_moe_4k	10	1	1/2/4/8	1/2/4/8	1	1/2/4/8	RC/noRC
qwen15_1.8b_8k	40	1/2/4/8	1/2/4/8	1/2/4/8	1/2/4/8	–	RC/noRC
qwen15_4b_8k	40	1/2/4/8	1/2/4/8	1/2/4/8	1/2/4/8	–	RC/noRC
qwen25_0point5_32k	6	1/2	2	2/4	1/2	–	RC/noRC
qwen25_14b_8k	1	1	4	8	1	–	RC
qwen25_32b_32k	4	1/2	4/8	4	1/2	–	RC/noRC
qwen25_3b_32k	6	1	1/2	1/2	2/4	–	RC/noRC
qwen25_72b_8k	3	1	4/8	4/8	1	–	RC/noRC
qwen25_7b_8k	6	1/2	1/2	4	1/2	–	RC/noRC
qwen2_1.5b_4k	26	1/2/4/8	1/2	1/2/4	1/2/4	–	RC/noRC
qwen3_0.6b_4k	38	1/2/4/8	1/2/4/8	1/2/4	1/2/4/8	–	RC/noRC
qwen3_1.7b_4k	38	1/2/4/8	1/2/4/8	1/2/4	1/2/4/8	–	RC/noRC
qwen3_30b_a3b_4k	2	1	4	2	2	1	RC/noRC
qwen3_8b_4k	28	1/2/4	1/2/4/8	1/2/4	1/2/4	–	RC/noRC

Quality checks cover configuration parsing, rank identifiers, field names, units, and file coverage. The manifest records configuration and file counts and storage sizes for consistency checks.

C.4. File Schema and Field Definitions

The directory hierarchy preserves the associations established during processing. The root manifest indexes model settings; summary files within each model-setting directory contain configuration- and rank-level measurements; and configuration/rank directories retain the detailed communication and matrix-kernel records:

```
LM-SimBench/
  data_manifest.csv
  <model_name>/
    step_time_summary.csv
    peak_memory.csv
    step_trace_time_summary.csv
    operator_latency_by_category.csv
    comm_summary.csv
    <parallel_config>/
      rank_<id>/
        comm_time.csv
        matrix_operator_latency.csv
```

Table 13 describes each file's location, granularity, and benchmark role. Table 14 lists representative join keys, tar-

Table 12. Hardware and software environment used to collect the primary LM-SimBench profiles.

Category	Item	Value
Host processor	CPU architecture	aarch64
	CPU model	Kunpeng-920
	CPU cores	192
Accelerator	NPU model	Ascend 910B3
	HBM per NPU	65536 MiB
	NPU interconnect	HCCS
Ascend software stack	CANN toolkit	8.0.0
	CANN kernels	910B, 8.0.0
	CANN Neural Network Acceleration Library (NNAL)	8.0.0
Deep learning runtime	Python	3.9.11
	PyTorch	2.1.0
	torch-npu	2.1.0.post3
	Transformers	4.43.2
Training framework	MindSpeed	0.8.0
	MindSpeed-LLM	branch 2.1.0, commit `bf1e61f4`

gets, and features. The dataset repository provides complete column definitions.

Within a model setting, `optimization_config` joins configuration-level records and `rank_id` identifies the corresponding rank-level views. Dense-model configuration names follow `dp<dp>-tp<tp>-pp<pp>-cp<cp>-<RC|noRC>`; MoE names may additionally contain `ep<ep>`. The suffixes `RC` and `noRC` specify the recomputation setting. These keys connect the iteration-time label to the rank-level summaries and the event- and kernel-level files for the same run.

C.5. Documentation, Licensing, and Maintenance

Documentation. The resources listed at the beginning of this supplement include dataset cards, field descriptions, usage examples, and Croissant metadata. Table 15 summarizes the data's use and provenance.

License. The datasets, metadata, and documentation are released under CC BY 4.0; the preprocessing and benchmark code is released under Apache-2.0.

Maintenance. The release will be versioned, with previous versions retained for reproducibility and major updates recorded in a changelog. File, schema, preprocessing, and data-loading issues will be tracked in the code repository.

Table 13. Storage location, profiling granularity, and benchmark role of the main CSV files.

File	Location	Granularity	Benchmark role
`data_manifest.csv`	Dataset root	Dataset / model level	Provides the global index of model directories, configuration counts, rank-directory counts, file coverage, and dataset-scale statistics.
`step_time_summary.csv`	Model directory	Configuration level	Provides configuration-level iteration-time labels derived from profiler time breakdowns for runtime prediction and configuration ranking.
`peak_memory.csv`	Model directory	Rank level	Provides rank-level peak memory labels for memory modeling, feasibility analysis, and imbalance analysis across ranks.
`step_trace_time_summary.csv`	Model directory	Rank level	Decomposes iteration time into computation, communication, idle time, and pipeline-related components for bottleneck diagnosis.
`operator_latency_by_category.csv`	Model directory	Operator-category / rank level	Aggregates profiler operators into semantic categories for operator-latency modeling and module-level performance analysis.
`comm_summary.csv`	Model directory	Communication-group / rank level	Provides aggregated communication statistics for comparing collective and point-to-point communication costs across configurations.
`comm_time.csv`	Configuration / rank directory	Communication-event / rank level	Preserves event-level communication behavior for analyzing waiting time, transfer size, link bandwidth, and communication bottlenecks.
`matrix_operator_latency.csv`	Configuration / rank directory	Matrix-kernel / rank level	Provides matrix-operator shapes, latency, data types, and hardware-counter features for GEMM modeling and hardware-aware analysis.

Table 14. Representative join keys, targets, and features in the main CSV files.

File	Representative fields
data_manifest.csv	model_name, parallel_config_count, rank_directory_count, size_bytes, size_human.
step_time_summary.csv	optimization_config, iteration_time_ms.
peak_memory.csv	optimization_config, rank_id, peak_memory_mb.
step_trace_time_summary.csv	optimization_config, rank_id, computing, communication, communication_not_overlapped, overlapped, free, stage, bubble.
operator_latency_by_category.csv	optimization_config, rank_id, op_category, op_type, core_type, count, total_time_us, avg_time_us, ratio_percent.
comm_summary.csv	comm_scope, comm_type, comm_group, op_count, elapsed_time_sum_ms, wait_time_sum_ms, link_bandwidth.
comm_time.csv	comm_scope, comm_type, comm_group, op_name, elapsed_time_ms, wait_time_ms, link_type, link_transit_size_mb, link_bandwidth_gbps.
matrix_operator_latency.csv	type, duration_us, input_shapes, output_shapes, a_shape, b_shape, c_shape, batch, m, n, k, aic_mac_ratio, cube_utilization_percent.

Table 15. Use, coverage, and provenance of LM-SimBench.

Topic	Summary
Intended use	Training-simulator evaluation and development: time and memory prediction, configuration ranking, operator and communication analysis, and calibration.
Coverage and evaluation scope	Measurements describe the sampled models, feasible configurations, and software/hardware stacks in Section C.1 and Appendix F. Coverage depends on model size, memory capacity, and platform implementations. Serving behavior and model-quality, safety, or fairness assessment require separate data.
Provenance and record type	Controlled distributed training runs produce the logs and profiler outputs; the processed tables are normalized derivatives of these measured records, without synthetic profiling samples or human annotations.
Personal or sensitive information	The release contains system telemetry, with no intentional personal records, user-generated text, or private user data.

D. Simulator Evaluation and Comparison

This section provides the full results supporting Sections 6.1 and 6.2. Calculon and vTrain [6] illustrate how LM-SimBench evaluates analytical and profiling-based simulators through common time and ranking targets. Intermediate profiles then connect their overall predictions to computation, memory, and communication behavior.

D.1. Evaluation Setup

We use Qwen1.5-1.8B at an 8K-token sequence length with CP=1 on the primary NPU stack. Calculon uses analytical cost models; vTrain combines an execution graph with LM-SimBench profiles supplied through an adapter. Ten configurations calibrate vTrain, and ten disjoint configurations form the test set. The shared-task comparison uses the six test configurations for which both simulators produce valid predictions.

Predicted and measured iteration times are matched by configuration. Accuracy, ranking, and shortlist metrics follow Section 5; simulation cost is the median software wall-clock time per configuration in the same CPU and software environment. The component and memory analyses below use their corresponding profiling targets, with coverage specified alongside the results.

D.2. Prediction, Selection, and Simulation Cost

Table 16 expands the comparison in Table 5. vTrain achieves 1.86% iteration-time sMAPE and Spearman $\rho = 1.000$, compared with 8.88% and 0.943 for Calculon. Its absolute time error is lower on five of the six configurations. Both simulators identify the measured fastest configuration and recover the measured top-three set, giving zero Top-1 selection regret.

Calculon's median simulation time is 1.78 ms per configuration, compared with 25.46 ms for vTrain, a $14.3\times$ speedup. Together, the prediction, selection, and timing metrics characterize how the two modeling approaches support configuration search: prediction accuracy describes the estimated runtime, while ranking and regret describe the resulting choice.

D.3. Calibration Coverage

A separate analysis evaluates how the configuration coverage of the calibration profiles affects vTrain's predictions. The test set is fixed at eight configurations. Six TP-endpoint calibration profiles yield 11.77% sMAPE; ten calibration profiles spanning TP/DP/PP reduce it to 2.08%. The 9.69-percentage-point reduction illustrates the value of covering multiple parallelism dimensions when supplying measured costs to a profiling-based simulator.

D.4. Diagnosis with Intermediate Profiles

Calculon: memory and matrix operators. Table 17 summarizes Calculon's diagnostic results. Its memory model retains 13 of the 20 measured-feasible configurations. Time and peak-memory errors and iteration-time rankings are evaluated on those 13 configurations; memory uses the maximum across ranks. Matrix-operator predictions are evaluated against per-call kernel latencies at matching matrix dimensions.

The 6.92 GiB peak-memory MAE and 43.34% matrix-operator sMAPE connect the end-to-end assessment to specific modeling components. Appendix E.1 shows the corresponding time and memory predictions, and Appendix E.3 reports matrix errors by configuration and shape. These linked views allow a reader to follow an aggregate metric to the measurements that explain it.

vTrain: computation and communication stages. Stage-level profiles provide evaluation targets for the intermediate costs emitted by vTrain (Table 18). Computation and recomputation achieve 11.60% and 2.75% sMAPE, respectively. Non-overlapped communication has 136.20% sMAPE: vTrain's serialized communication estimates differ from the overlap observed in the training profiles. The stage breakdown thus identifies communication scheduling as a concrete component to refine, alongside the accurate recomputation estimates.

Table 16. Shared-task evaluation of Calculon and vTrain on six common test configurations of Qwen1.5-1.8B (8K tokens, CP=1).

Level	Task and metric	Calculon	vTrain
System	Iteration time, MAE / sMAPE / R^2	1.059 s / 8.88% / 0.630	**0.207 s / 1.86% / 0.990**
Strategy	Iteration-time ranking, Spearman ρ	0.943	**1.000**
	Iteration-time ranking, Kendall τ	0.867	**1.000**
	Top-1 hit / selection regret	1 / 0%	1 / 0%
	Top-3 set overlap	3/3	3/3
Efficiency	Median wall-clock per configuration	**1.78 ms**	25.46 ms

Table 17. Calculon diagnostics on Qwen1.5-1.8B (8K tokens, CP=1). Time, memory, and ranking use the 13 configurations retained from 20 measured-feasible candidates.

Task	**Level**	**Metrics**	**Results**
Iteration time prediction	System	MAE / sMAPE / R^2	1.97 s / 15.96% / 0.358
Peak memory prediction	System	MAE / sMAPE / R^2	6.92 GiB / 39.34% / -3.087
Matrix-operator latency	Component	MAE / sMAPE / R^2	0.109 ms / 43.34% / 0.817
Feasible-configuration retention	System	Retention rate	65.0% (13/20)
Iteration time ranking	Strategy	Spearman ρ / Top-1 hit	0.753 / 1

Iteration-time and peak-memory prediction and ranking use the 13 configurations retained by Calculon's memory model; the retention rate uses all 20 measured-feasible configurations. Results follow the order of the listed metrics.

Table 18. vTrain component diagnostics on Qwen1.5-1.8B (8K tokens). Predictions are compared with stage-level profiling targets.

Target	**MAE**	**sMAPE (%)**	$\boldsymbol{R^2}$
Stage computation time	716.10 ms	11.60	0.462
Recomputation-stage time	37.99 ms	2.75	0.801
Non-overlapped communication time	8.710 s	136.20	-0.643

E. Profiling Cases and Simulator Diagnostics

The cases in this section expand the intermediate-profile analysis in Section 6.2. Qwen1.5-1.8B connects configuration selection to time and memory predictions, while Qwen2.5-32B shows how rank-level profiles reveal the computation and communication behind a multi-node training iteration. In both cases, the linked views connect overall behavior to specific ranks, communication primitives, and matrix operators.

E.1. Single-Node Case: Qwen1.5-1.8B

Configuration-level time and rank-level memory. For Qwen1.5-1.8B at a training sequence length of 8K tokens, varying DP, TP, PP, CP, and recomputation produces substantial differences in iteration time (Figure 5). Rank-level memory measurements expose how parallel partitioning and recomputation change the memory distribution within the same workload (Figure 6).

Connecting memory screening to prediction accuracy. On the CP=1 subset, Calculon's memory model accepts 13 of the 20 measured-feasible configurations and rejects seven. Its largest memory overestimates occur without recomputation. Figures 7 and 8 compare time and memory predictions for the accepted configurations. Together, these measurements distinguish runtime prediction errors from the loss of feasible candidates during memory screening. The aggregate metrics are reported in Table 17; detailed matrix-operator errors for this workload are collected in Appendix E.3.

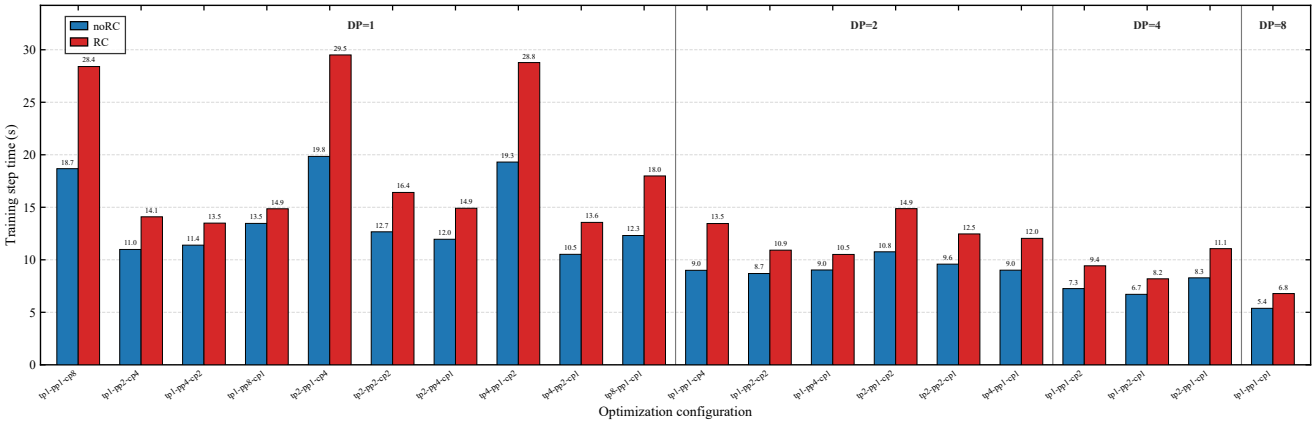

Figure 5. **Iteration time across parallel configurations.** Mean iteration time of Qwen1.5-1.8B (8K tokens) under different DP/TP/PP/CP and recomputation settings.

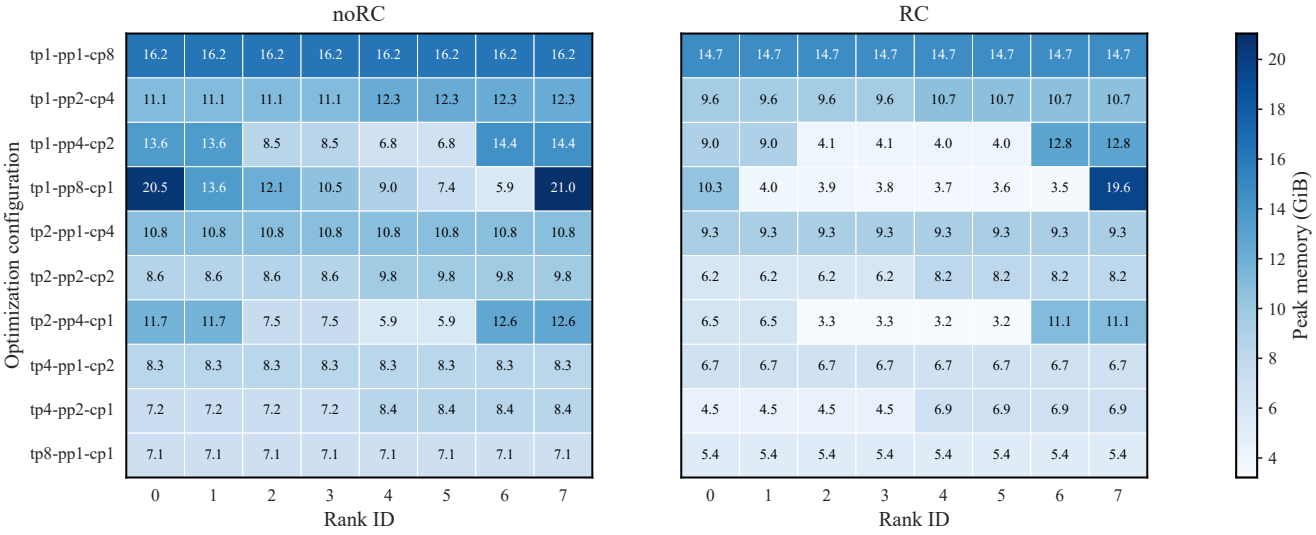

Figure 6. **Rank-level peak memory under DP=1 configurations.** Peak memory of Qwen1.5-1.8B (8K tokens) by rank across TP/PP/CP configurations.

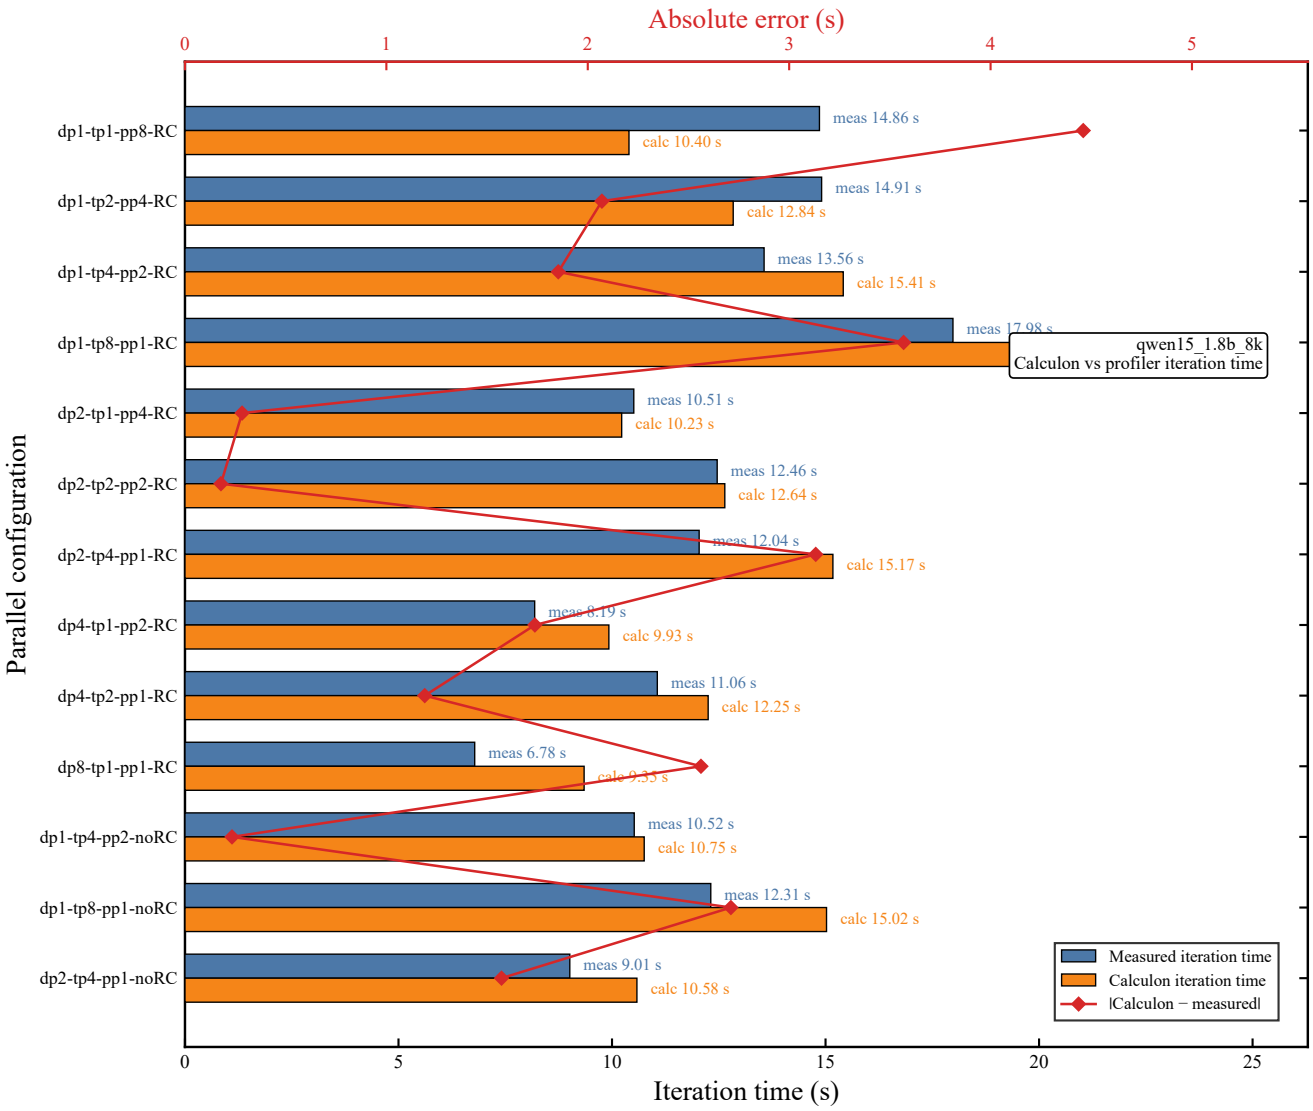

Figure 7. **Calculon iteration-time prediction on Qwen1.5-1.8B (8K tokens).** Measured and Calculon-predicted iteration time for the 13 configurations accepted by Calculon's memory model.

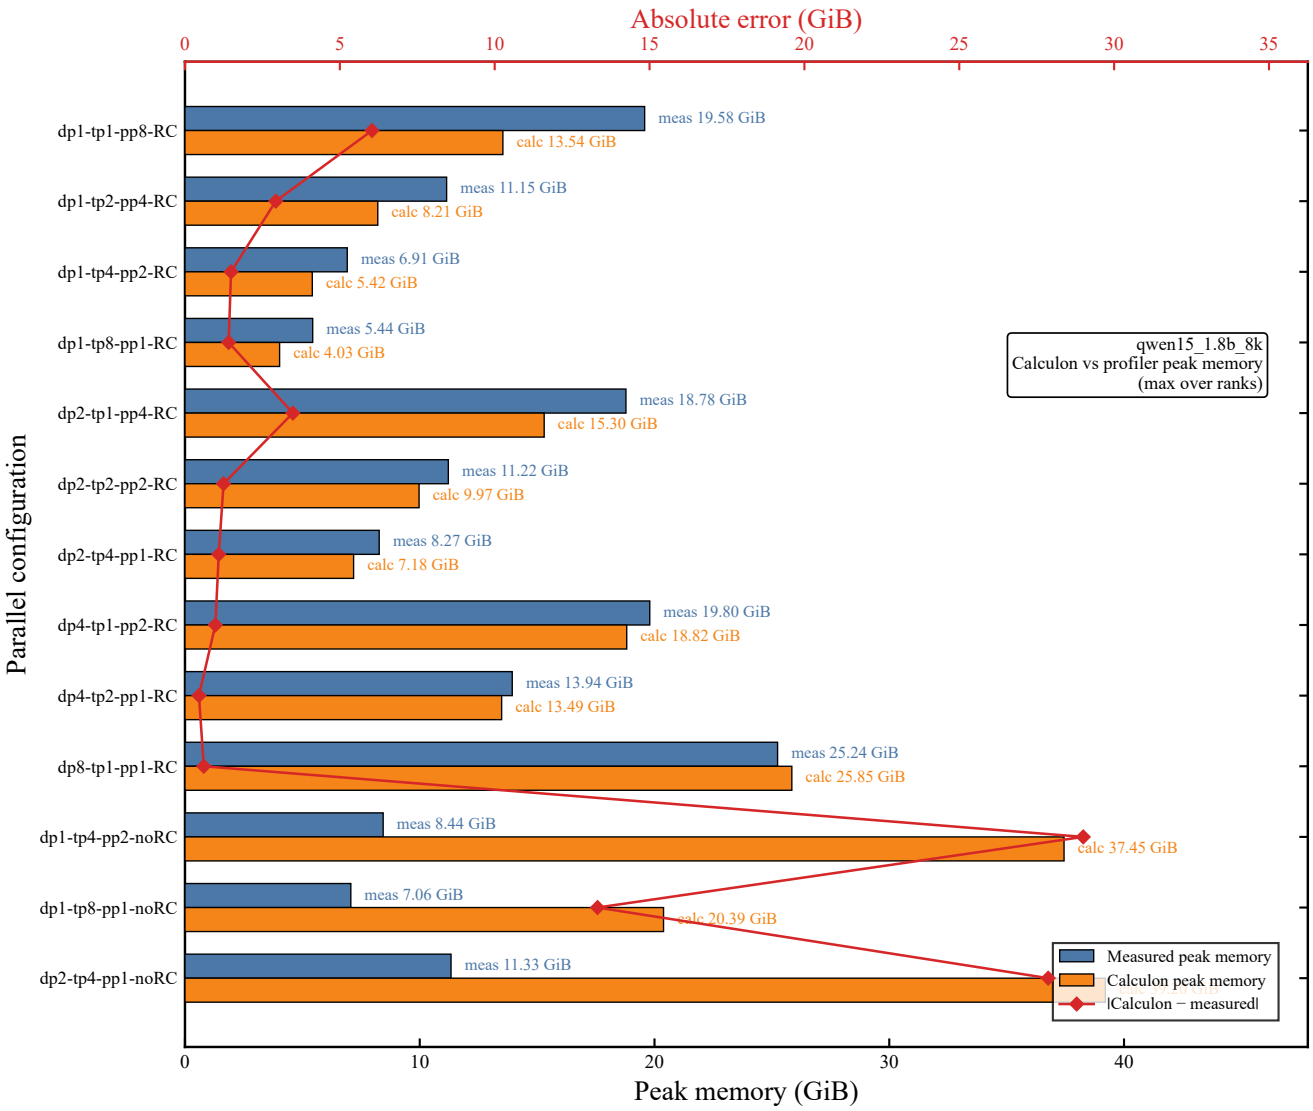

Figure 8. **Calculon peak-memory prediction on Qwen1.5-1.8B (8K tokens).** Measured maximum rank-level peak memory and Calculon predictions for 13 configurations.

E.2. Multi-Node Case: Qwen2.5-32B

Locating costs across ranks. For Qwen2.5-32B at a training sequence length of 32K tokens, rank-level profiles reveal nonuniform memory, computation, and communication across the distributed execution. Memory curves vary with parallel partitioning and pipeline-stage assignment (Figure 9); time breakdowns identify the computation and nonoverlapped communication associated with each rank (Figure 10).

Tracing dominant communication and matrix costs. The lower-level records identify the dominant costs. On rank 0 under `dp1-tp8-pp4-cp1-RC`, `batchSendRecv` accounts for most communication time (Figure 11), while a small set of matrix shapes accounts for most cumulative matrix-operator latency (Figure 12). These cumulative records locate the primitives and shapes that account for device work; the rank-level time breakdowns show how computation and communication contribute to the execution schedule.

Relating operator estimates to iteration time. The Calculon baseline uses a compatibility setup with multi-head attention in place of grouped-query attention and activation recomputation enabled; the baseline implementation supports neither grouped-query attention nor context parallelism. It predicts 73.76 s per iteration against the measured 112.22 s, a 34.3% underestimate. The operator-level comparison in Figure 13 shows that Calculon recovers the dominant matrix shapes and principal `MatMulV3` hotspots, but underestimates their cumulative cost by 27–35%. Thus, the component profiles expose underestimated operator costs alongside the end-to-end runtime error.

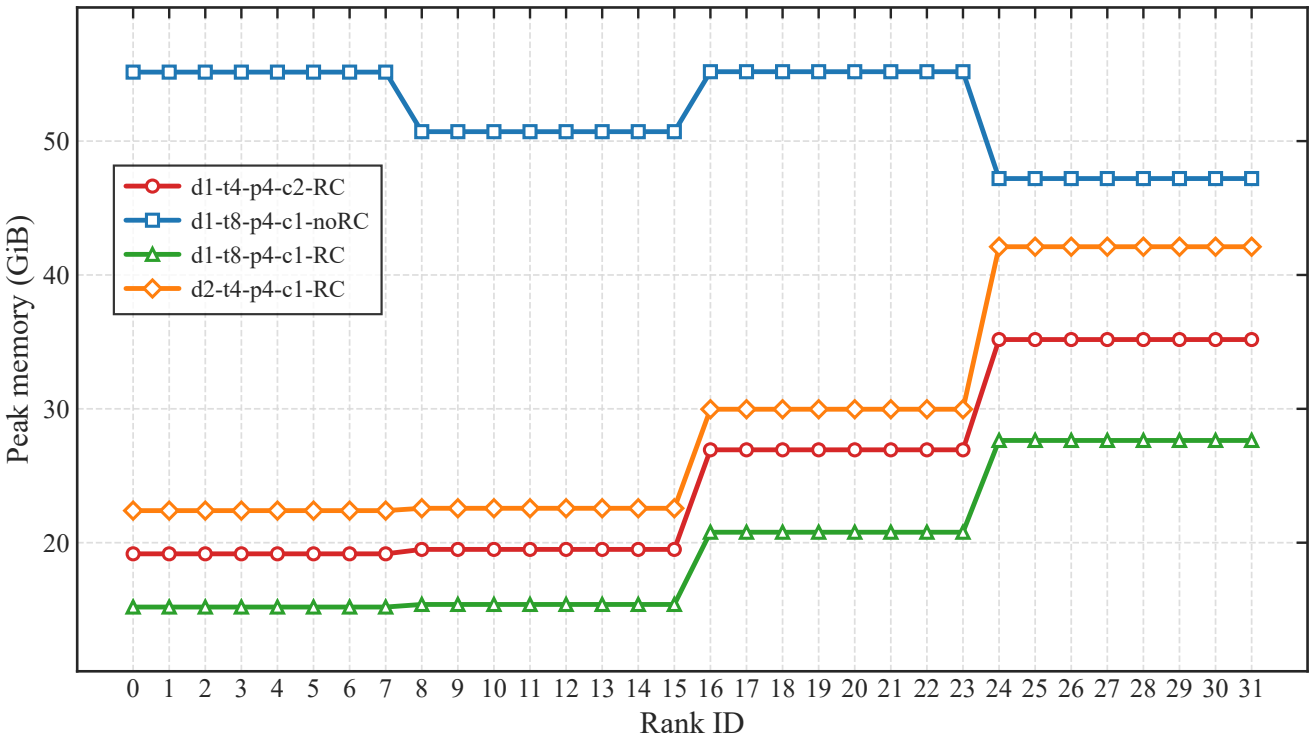

Figure 9. **Rank-level peak memory across configurations.** Peak memory by rank for Qwen2.5-32B (32K tokens) under four parallel configurations.

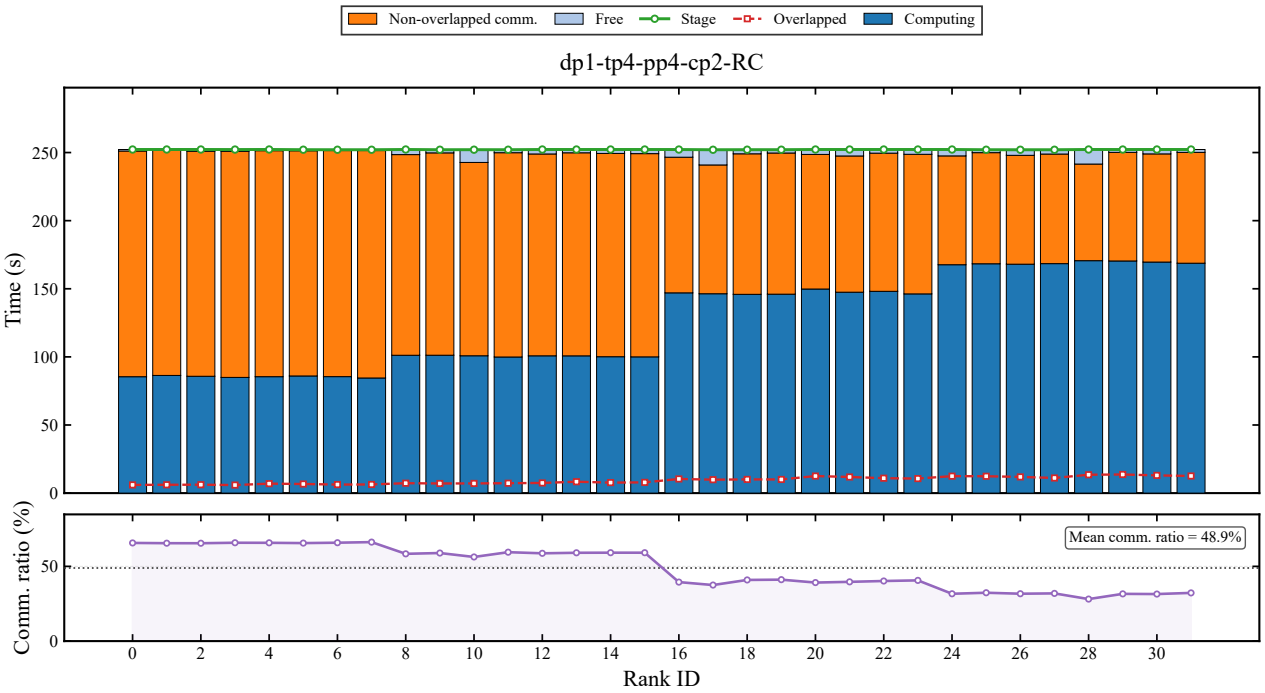

(a) `dp1-tp4-pp4-cp2-RC`.

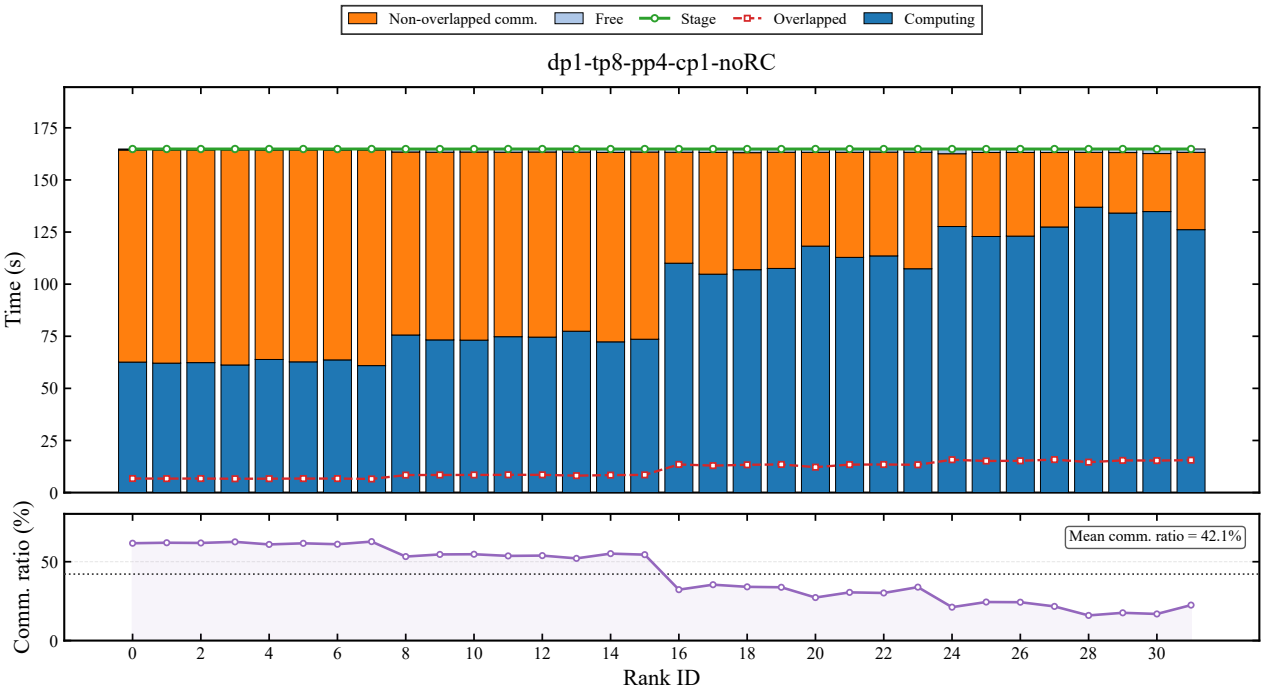

(b) `dp1-tp8-pp4-cp1-noRC`.

Figure 10. **Rank-level iteration-time breakdowns.** Each panel breaks down iteration time for Qwen2.5-32B (32K tokens) into computation, non-overlapped communication, free, stage, and overlapped time and reports the communication ratio for one configuration.

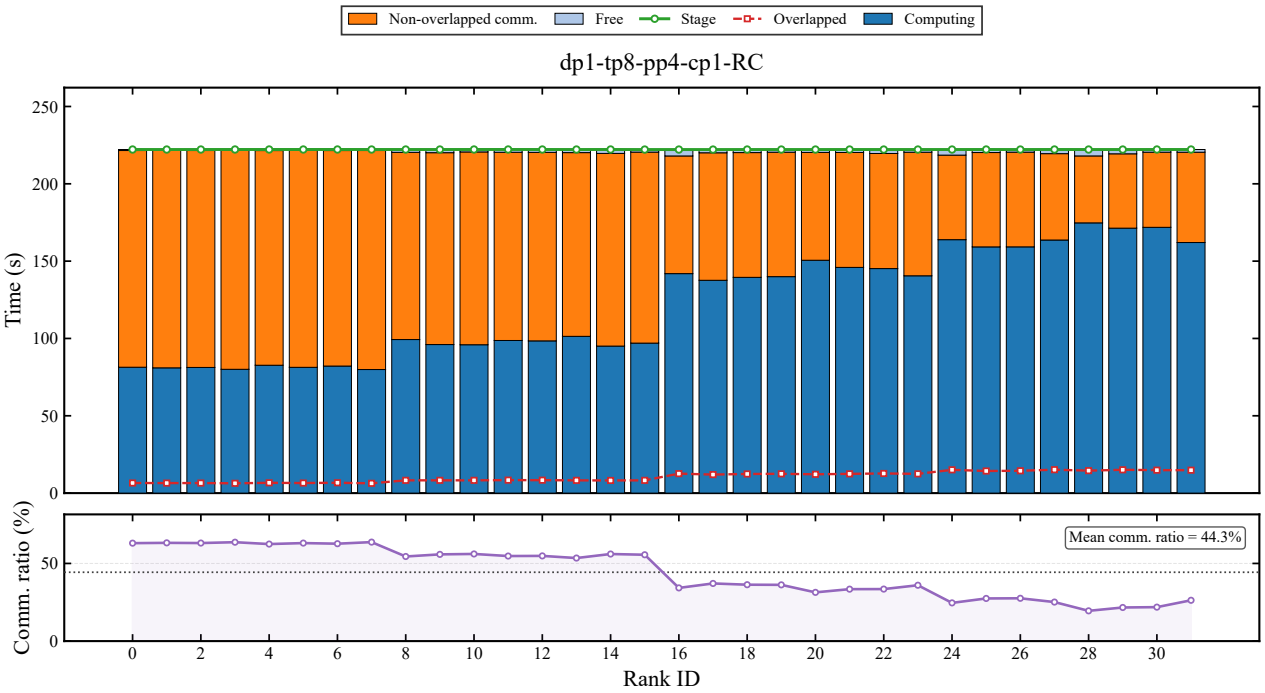

(c) `dp1-tp8-pp4-cp1-RC`.

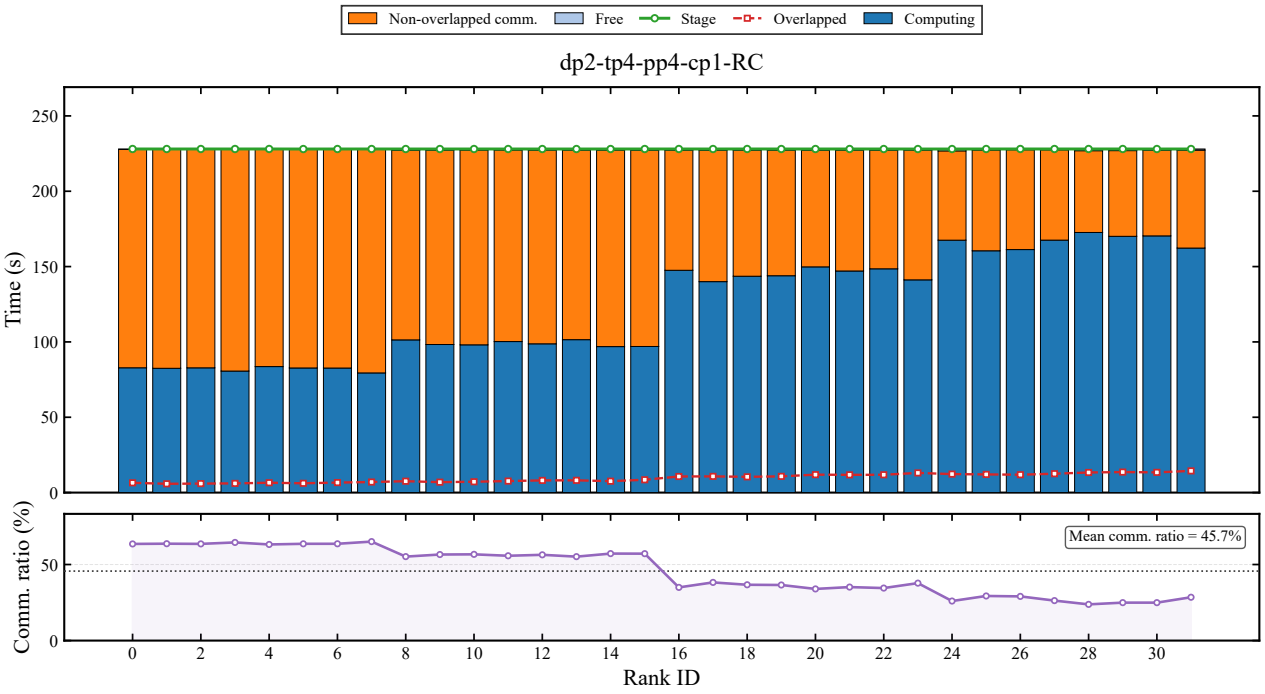

(d) `dp2-tp4-pp4-cp1-RC`.

Figure 10. **Rank-level iteration-time breakdowns.** Each panel breaks down iteration time for Qwen2.5-32B (32K tokens) into computation, non-overlapped communication, free, stage, and overlapped time and reports the communication ratio for one configuration.

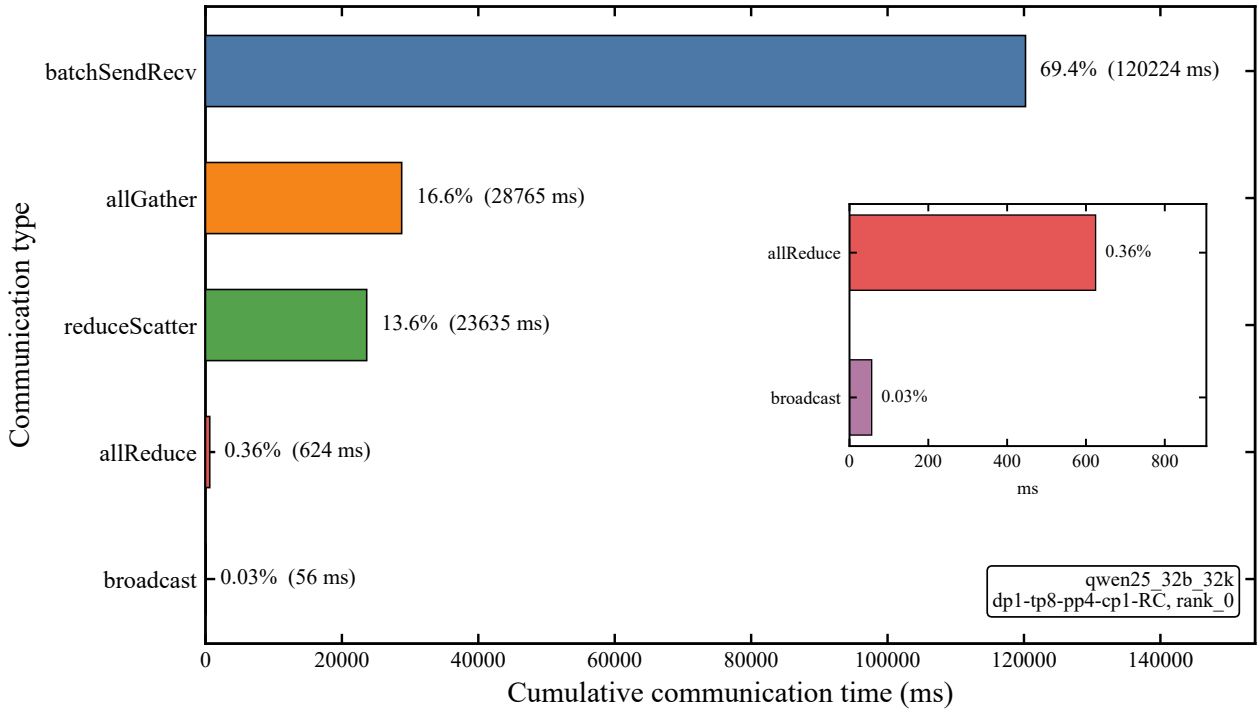

Figure 11. **Communication time breakdown.** Cumulative time by communication primitive for rank 0 of Qwen2.5-32B (32K tokens) under `dp1-tp8-pp4-cp1-RC`. The inset enlarges primitives with small shares.

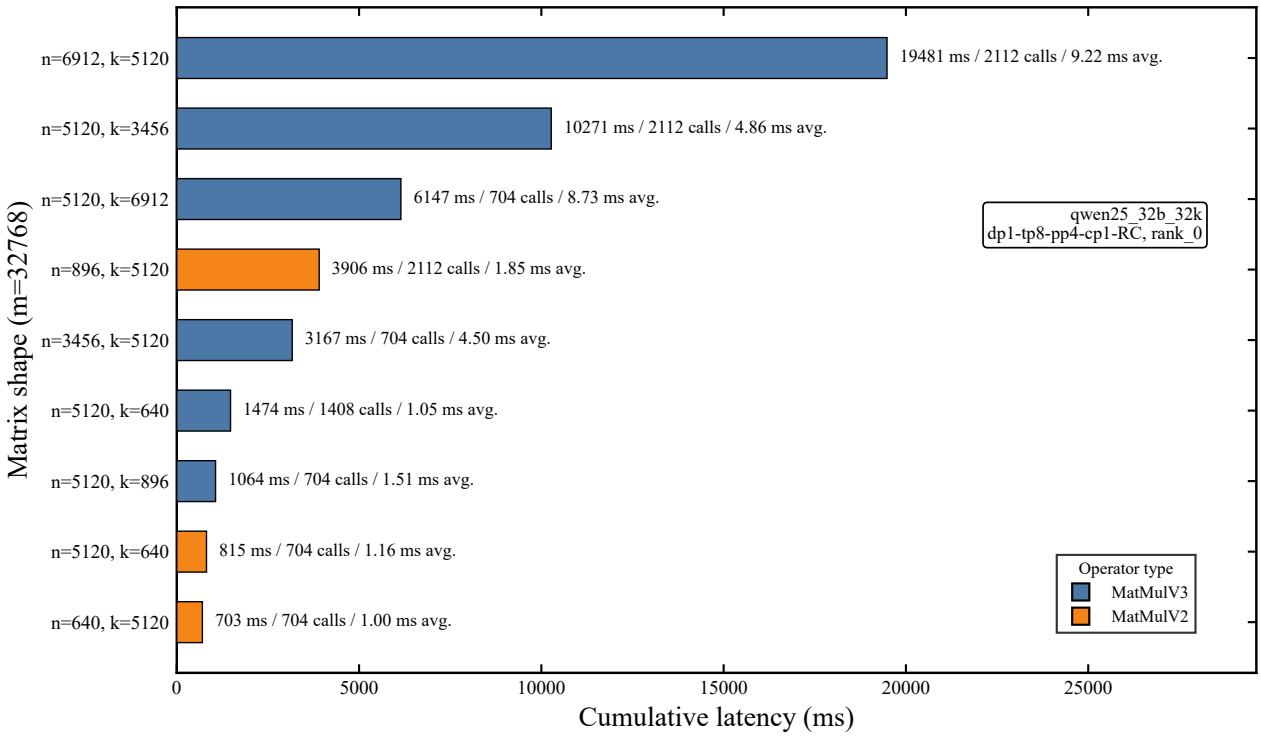

Figure 12. **Top matrix-operator shapes by cumulative latency.** Matrix operators on rank 0 are grouped by (n, k) with $m = 32768$ for Qwen2.5-32B (32K tokens) under `dp1-tp8-pp4-cp1-RC`. Each bar reports cumulative latency, call count, and mean latency.

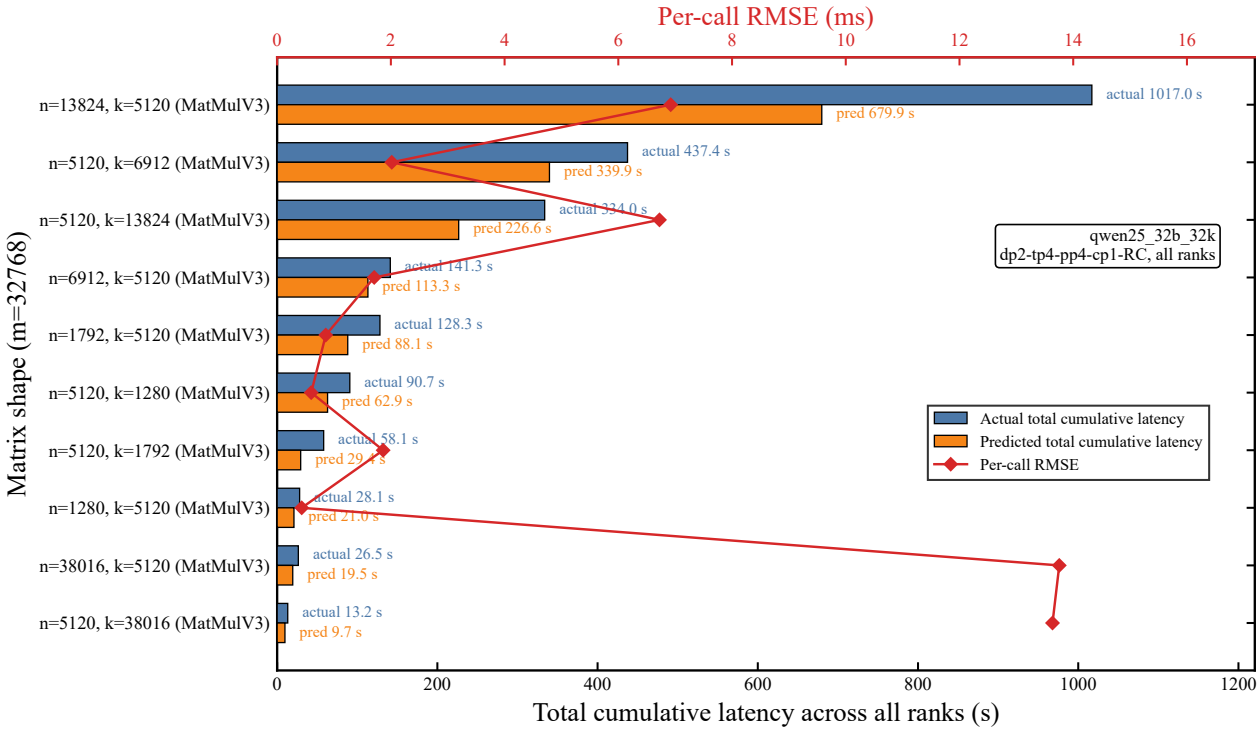

(a) `dp2-tp4-pp4-cp1-RC`.

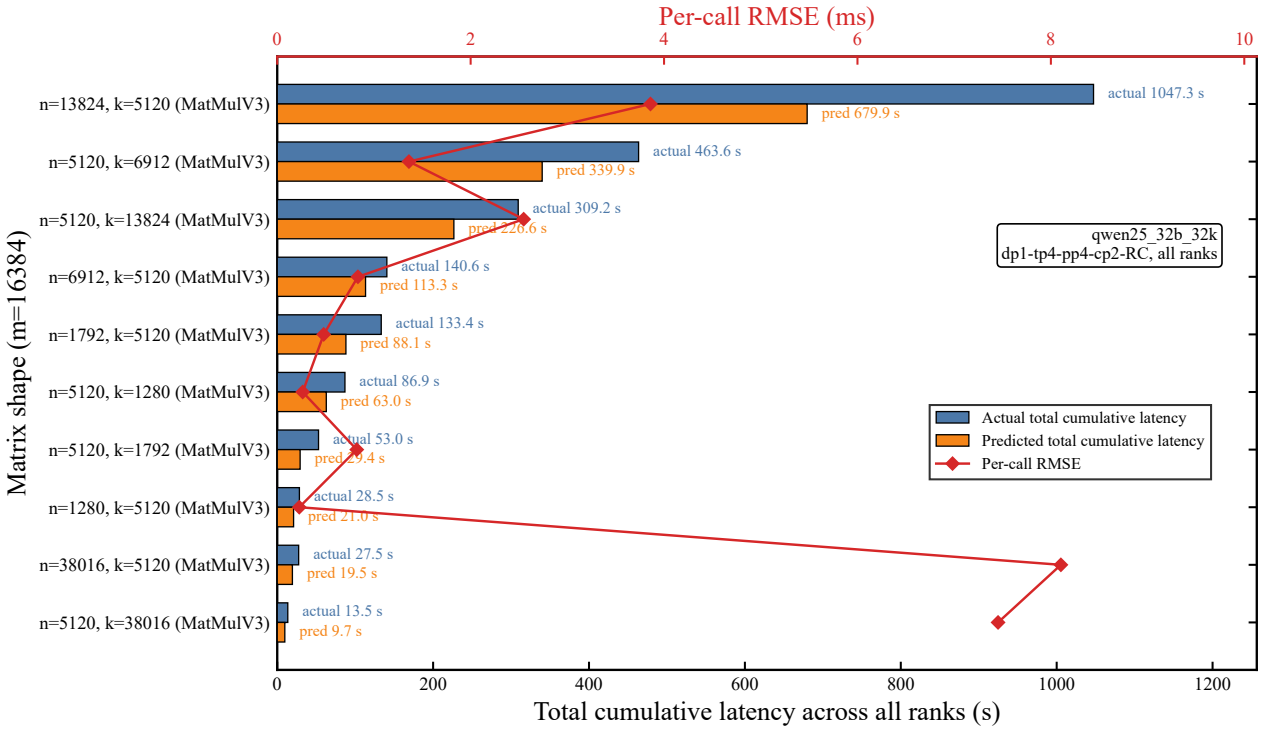

(b) `dp1-tp4-pp4-cp2-RC`.

Figure 13. Matrix-latency predictions for Qwen2.5-32B (32K tokens) under four parallel configurations. Bars compare measured and predicted cumulative latency across ranks for dominant shapes; the red curve shows per-call RMSE for each shape.

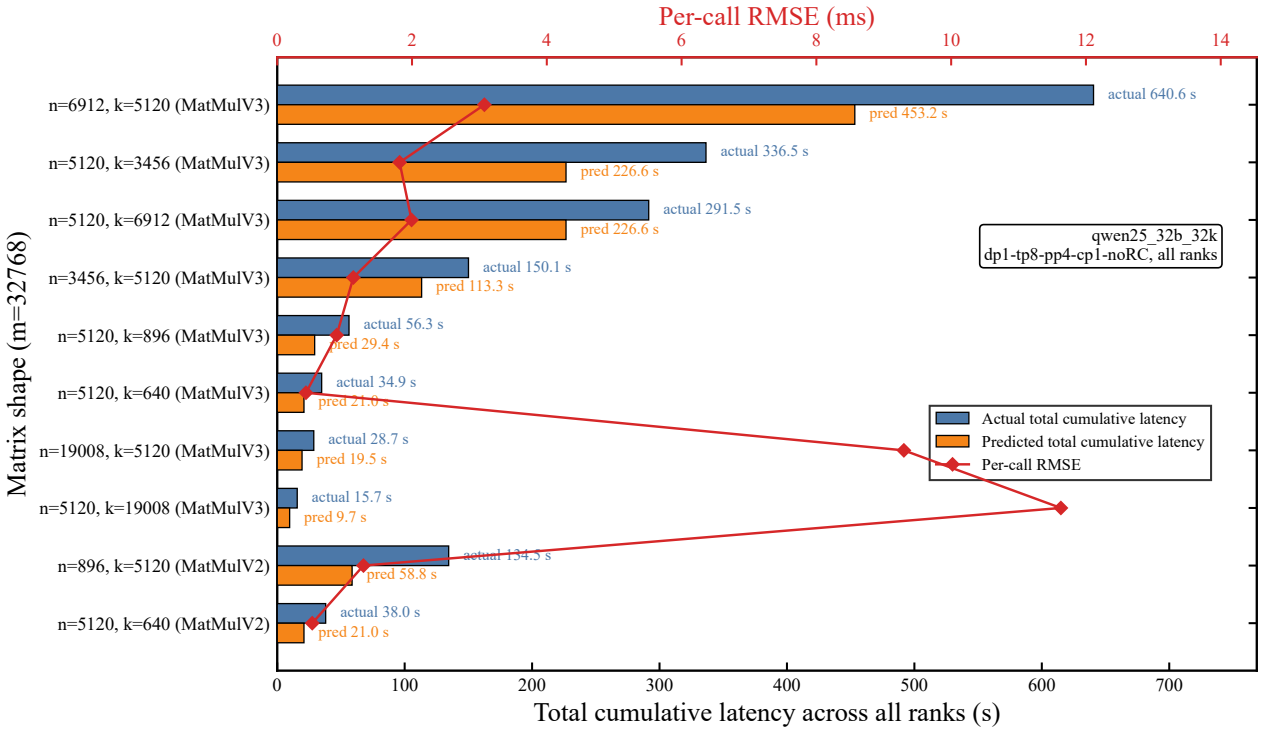

(c) `dp1-tp8-pp4-cp1-noRC`.

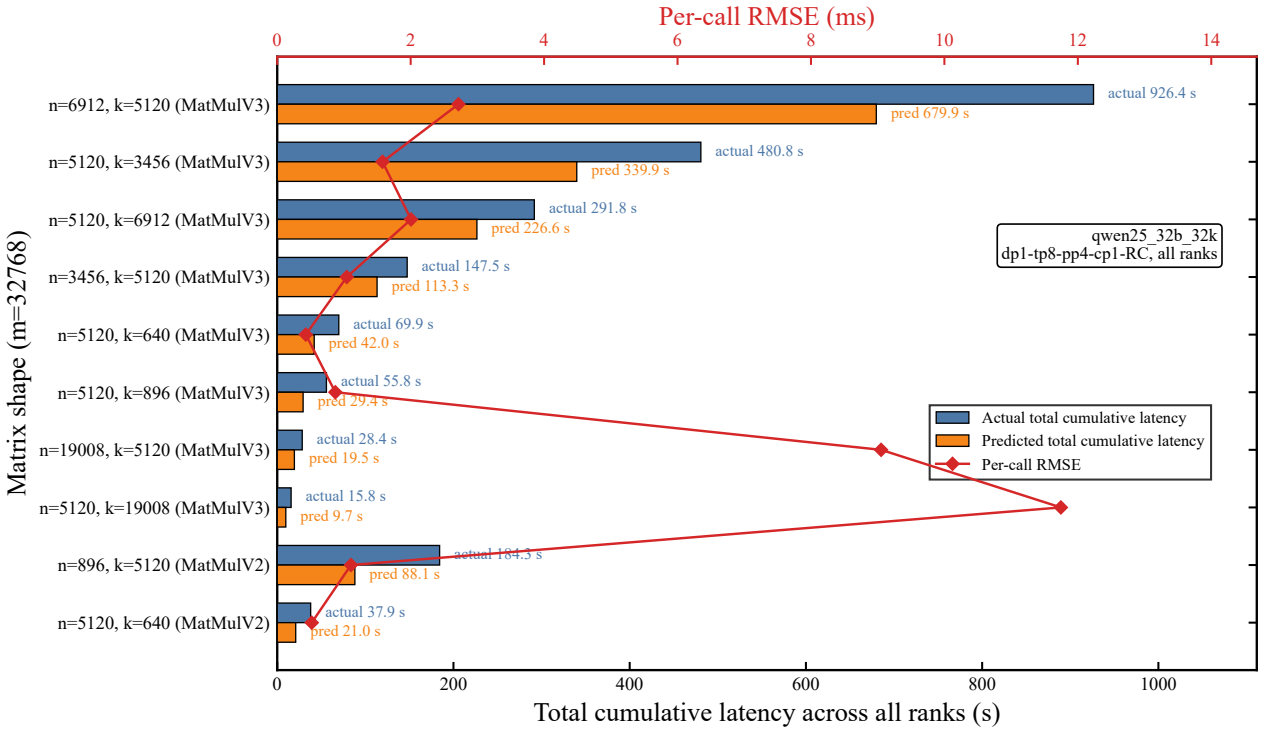

(d) `dp1-tp8-pp4-cp1-RC`.

Figure 13. Matrix-latency predictions for Qwen2.5-32B (32K tokens) under four parallel configurations. Bars compare measured and predicted cumulative latency across ranks for dominant shapes; the red curve shows per-call RMSE for each shape.

E.3. Qwen1.5-1.8B Matrix-Operator Errors

Tables 19 and 20 expand the aggregate matrix-operator metrics in Table 17 for Qwen1.5-1.8B (8K tokens). Grouping by parallel configuration shows how estimation quality varies with the training strategy; grouping by matrix dimensions identifies the operator shapes associated with those errors. Together, the two tables support examining a computation estimator at both configuration and kernel granularity.

Table 19. Matrix-operator prediction metrics by parallel configuration for Qwen1.5-1.8B (8K tokens). Configurations continue from the left panel to the right.

Parallel configuration	MAE (ms)	sMAPE (%)	R^2	Parallel configuration	MAE (ms)	sMAPE (%)	R^2
dp1-tp1-pp1-cp8-RC	0.048	40.505	0.779	dp2-tp1-pp1-cp4-RC	0.072	36.826	0.864
dp1-tp1-pp1-cp8-noRC	0.053	42.190	0.778	dp2-tp1-pp1-cp4-noRC	0.076	38.311	0.873
dp1-tp1-pp2-cp4-RC	0.072	36.794	0.865	dp2-tp1-pp2-cp2-RC	0.120	30.403	0.890
dp1-tp1-pp2-cp4-noRC	0.076	38.126	0.874	dp2-tp1-pp2-cp2-noRC	0.127	31.581	0.897
dp1-tp1-pp4-cp2-RC	0.119	30.313	0.889	dp2-tp1-pp4-cp1-RC	0.252	31.580	0.874
dp1-tp1-pp4-cp2-noRC	0.128	31.596	0.895	dp2-tp1-pp4-cp1-noRC	0.268	32.662	0.880
dp1-tp1-pp8-cp1-RC	0.251	31.436	0.875	dp2-tp2-pp1-cp2-RC	0.101	45.635	0.758
dp1-tp1-pp8-cp1-noRC	0.267	32.491	0.881	dp2-tp2-pp1-cp2-noRC	0.118	50.418	0.748
dp1-tp2-pp1-cp4-RC	0.056	42.201	0.725	dp2-tp2-pp2-cp1-RC	0.203	42.782	0.741
dp1-tp2-pp1-cp4-noRC	0.061	45.811	0.741	dp2-tp2-pp2-cp1-noRC	0.245	48.749	0.722
dp1-tp2-pp2-cp2-RC	0.093	43.991	0.792	dp2-tp4-pp1-cp1-RC	0.117	43.081	0.696
dp1-tp2-pp2-cp2-noRC	0.113	49.147	0.769	dp2-tp4-pp1-cp1-noRC	0.141	47.530	0.690
dp1-tp2-pp4-cp1-RC	0.203	42.833	0.741	dp4-tp1-pp1-cp2-RC	0.120	30.453	0.888
dp1-tp2-pp4-cp1-noRC	0.245	48.783	0.721	dp4-tp1-pp1-cp2-noRC	0.127	31.512	0.894
dp1-tp4-pp1-cp2-RC	0.060	39.834	0.731	dp4-tp1-pp2-cp1-RC	0.262	32.662	0.869
dp1-tp4-pp1-cp2-noRC	0.069	43.331	0.731	dp4-tp1-pp2-cp1-noRC	0.271	32.963	0.879
dp1-tp4-pp2-cp1-RC	0.112	41.730	0.700	dp4-tp2-pp1-cp1-RC	0.210	43.702	0.727
dp1-tp4-pp2-cp1-noRC	0.144	48.204	0.690	dp4-tp2-pp1-cp1-noRC	0.245	48.707	0.721
dp1-tp8-pp1-cp1-RC	0.107	60.527	0.608	dp8-tp1-pp1-cp1-RC	0.279	34.472	0.859
dp1-tp8-pp1-cp1-noRC	0.109	61.202	0.655	dp8-tp1-pp1-cp1-noRC	0.292	34.973	0.859

Table 20. Matrix-operator prediction metrics by matrix dimensions for Qwen1.5-1.8B (8K tokens). Shapes continue from the left panel to the right.

m	n	k	MAE (ms)	sMAPE (%)	m	n	k	MAE (ms)	sMAPE (%)
1024	2048	2048	0.011	24.151	4096	3072	2048	0.078	38.374
1024	2048	5504	0.056	55.184	4096	5504	2048	0.114	32.488
1024	2048	6144	0.068	58.875	4096	6144	2048	0.108	28.405
1024	2048	11008	0.089	46.131	4096	11008	2048	0.193	29.637
1024	5504	2048	0.046	47.035	4096	37984	2048	1.271	48.828
1024	6144	2048	0.048	45.422	4096	75968	2048	1.350	30.161
1024	11008	2048	0.064	35.649	4096	151936	2048	2.326	26.538
1024	151936	2048	0.968	40.604	8192	256	2048	0.018	36.910
2048	1024	2048	0.009	19.450	8192	512	2048	0.032	45.419
2048	2048	1024	0.002	4.102	8192	688	2048	0.191	113.572
2048	2048	2048	0.031	44.121	8192	768	2048	0.078	64.092
2048	2048	2752	0.077	68.929	8192	1024	2048	0.042	32.793
2048	2048	3072	0.082	63.776	8192	1376	2048	0.091	47.508
2048	2048	5504	0.096	48.069	8192	1536	2048	0.078	38.603
2048	2048	6144	0.075	37.723	8192	2048	256	0.009	20.407
2048	2048	11008	0.153	41.648	8192	2048	512	0.030	43.394
2048	2752	2048	0.062	59.668	8192	2048	688	0.234	120.810
2048	3072	2048	0.050	47.278	8192	2048	768	0.032	33.151
2048	5504	2048	0.070	38.430	8192	2048	1024	0.062	42.996
2048	6144	2048	0.073	36.576	8192	2048	1376	0.092	47.680
2048	11008	2048	0.093	27.610	8192	2048	1536	0.121	48.883
2048	75968	2048	0.736	32.456	8192	2048	2048	0.074	28.819
2048	151936	2048	1.269	28.616	8192	2048	2752	0.177	45.856
4096	512	2048	0.009	20.546	8192	2048	3072	0.545	88.956
4096	1024	2048	0.035	48.722	8192	2048	5504	0.435	54.585
4096	1376	2048	0.066	62.631	8192	2048	6144	0.189	26.503
4096	1536	2048	0.047	44.944	8192	2048	11008	0.387	29.873
4096	2048	512	0.005	12.478	8192	2048	18992	1.663	60.863
4096	2048	1024	0.039	51.339	8192	2048	37984	5.344	82.560
4096	2048	1376	0.065	61.595	8192	2752	2048	0.139	38.224
4096	2048	1536	0.038	37.719	8192	3072	2048	0.137	34.575
4096	2048	2048	0.044	33.438	8192	5504	2048	0.268	38.426
4096	2048	2752	0.099	50.468	8192	6144	2048	0.259	34.041
4096	2048	3072	0.192	68.259	8192	11008	2048	0.440	32.862
4096	2048	5504	0.170	43.264	8192	18992	2048	1.423	53.131
4096	2048	6144	0.079	21.620	8192	37984	2048	2.795	52.109
4096	2048	11008	0.263	38.596	8192	75968	2048	2.778	30.899
4096	2048	37984	2.144	72.136	8192	151936	2048	5.433	30.293
4096	2752	2048	0.081	43.180					

F. Cross-Platform Evaluation

This section expands the platform-calibration and configuration-selection results in Sections 6.3 and 6.4. Paired V100/910B3 profiles support two uses of LM-SimBench: calibrating a common simulator on each platform and using NPU measurements to select promising GPU configurations. Detailed time, memory, and component records connect these uses to the execution characteristics of the two stacks.

F.1. Experimental Setup and Coverage

We collect new paired measurements on 8×V100 (32 GiB per device) and 8×910B3 (64 GiB) for Qwen3-0.6B and Qwen3-1.7B at a training sequence length of 4K tokens. Each model uses the same 18 DP/TP/PP/recomputation candidates on both stacks, matching the CP=1 subset of its 38 configurations in the primary release. The NPU reference in this experiment was collected separately from that release. Both stacks use a distributed optimizer and profile the sixth iteration after five warm-up steps. The supplementary data directories use the suffix `_tp_pp_sweep_cp1_npu` for NPU runs and `_tp_pp_sweep_cp1_gpu` under `processed_gpu/` for GPU runs, prefixed by the workload name.

Table 21 lists the matched workload settings and the platform-specific precision, fusion, sequence-parallel, and profiler settings. Ascend-specific `AICore`/`MTE`/`cube` and `HCCS`/`RDMA`/`SDMA` fields are left empty in GPU records. Driver, CUDA/CANN, PyTorch, collective-library, and framework revisions are recorded in the artifact's environment metadata.

Complete GPU profiles are available for 15 of 18 candidates for Qwen3-0.6B and 12 of 18 for Qwen3-1.7B; the NPU side is complete for 18 and 16 candidates, respectively. Paired statistics therefore use 15 and 11 same-named configurations with complete profiles on both stacks, 26 pairs in total. Table 22 lists every candidate that is not complete on both stacks. Empty GPU profile directories are marked `missing_or_incomplete` in the manifest; their failure causes are unknown. Table 23 summarizes the paired comparisons.

F.2. Simulator Calibration across Platforms

Calculon provides a concrete application of the profiling data to simulator calibration and evaluation. We use a common modeling implementation on both platforms, calibrate its platform parameters against the corresponding profiles, and compare its time, memory, and ranking outputs with the measured targets. The implementation includes Qwen3 architecture and memory accounting, kernel launch latency, pipeline scheduling, and rank-level memory output. Each platform uses one parameter set for compute efficiency, launch latency, communication, and runtime memory across both models and the full candidate grid. Simulation outputs are stored under `calculon/outputs/qwen3_sweeps_final/`.

Table 24 expands Table 6 with rank-level memory errors. On the platform-calibration configurations, iteration-time MAPE is 7.97%–11.22% and time-ranking ρ is 0.853–0.950; rank-level peak-memory MAPE is 12.55%–15.99% and configuration-level peak-memory ranking ρ is 0.986–0.994. Across the four model×recomputation settings, GPU measurement, NPU measurement, Calculon-GPU, and Calculon-NPU select the same Top-1 configuration (4/4).

F.3. NPU-Guided GPU Configuration Selection

NPU measurements provide a shortlist for GPU configuration selection, following the procedure in Section 5.4. For each model, we retrospectively rank the paired-complete candidates by measured NPU iteration time, retain the fastest three, and identify the fastest GPU configuration within that shortlist using its measured GPU time. The comparison target is the fastest GPU configuration over the entire paired candidate set. Top-1 selection regret and Top-3 shortlist regret measure the GPU runtime penalty relative to this optimum when selecting the NPU-fastest candidate or the GPU-fastest member of the NPU-derived top-three shortlist, respectively (Eq. 7). The values are recomputed from the rounded times in Table 27. In both workloads the NPU-derived shortlist contains this GPU optimum, giving zero shortlist selection regret (Table 25). The candidate counts decrease from 15 to 3 and from 11 to 3, respectively.

Pooled over both recomputation settings, iteration-time rank agreement between the two stacks is moderate ($\rho = 0.718$ and 0.664, $\tau_b = 0.562$ and 0.491 for the two models). Directly choosing the NPU-fastest configuration finds the GPU optimum for Qwen3-0.6B and incurs a 6.9% GPU runtime penalty for Qwen3-1.7B. Table 26 additionally compares configurations with recomputation fixed: within each of the four model×recomputation strata, the two stacks select the same Top-1 configuration. This conditional agreement and the pooled shortlist results show how NPU observations can narrow GPU search even when the complete rankings differ.

Table 27 reports the per-configuration iteration times, with ranks computed over all paired configurations within each model, pooling RC and noRC. A leave-one-configuration-out diagnostic fits a median GPU/NPU time ratio on the other paired configurations and applies it to the held-out NPU time. It yields 22.4% and 22.9% sMAPE with $R^2 = -8.33$ and -4.74, respectively, showing that the runtime difference varies with configuration.

F.4. Memory and Recomputation

Configuration-level peak memory is the maximum across ranks, matching the memory-ranking task in Section 5. Its ordering is strongly aligned across platforms: rank correla-

Table 21. Matched workload settings and platform-specific choices in the cross-platform experiment.

Aspect	**Setting**
Matched	Model architecture, layer counts, hidden size, attention heads, and grouped-query attention; sequence length 4096; micro/global batch 1/32; eight devices; CP=1; candidate configuration names; optimizer hyperparameters and seed; profiled iteration index.
Stack-specific	Data type (FP16 on V100 vs. BF16 on 910B3); fused kernels (FlashAttention and fused RoPE/SwiGLU/RMSNorm disabled in the GPU sweep, enabled on the Ascend stack); sequence parallelism (enabled on the NPU side); profiler-visible hardware counters and communication-link fields; input data path.

Table 22. Candidates with missing or incomplete profiles on either stack, excluded from paired comparisons.

Model	Configuration	GPU status	NPU status
Qwen3-0.6B	`dp2-tp1-pp4-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-0.6B	`dp4-tp1-pp2-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-0.6B	`dp8-tp1-pp1-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-1.7B	`dp1-tp2-pp4-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-1.7B	`dp2-tp1-pp4-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-1.7B	`dp2-tp2-pp2-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-1.7B	`dp4-tp1-pp2-cp1-RC`	Complete	Missing
Qwen3-1.7B	`dp4-tp1-pp2-cp1-noRC`	Missing/incomplete (empty directory)	Missing
Qwen3-1.7B	`dp4-tp2-pp1-cp1-noRC`	Missing/incomplete (empty directory)	Complete
Qwen3-1.7B	`dp8-tp1-pp1-cp1-noRC`	Missing/incomplete (empty directory)	Complete

Table 23. Cross-platform agreement in configuration rankings, memory behavior, and operator performance between V100 GPUs and Ascend 910B3 NPUs. *Three paired configurations.

Level	Metric	Qwen3-0.6B	Qwen3-1.7B
Coverage	Paired complete configurations	15	11
Strategy	Iteration rank ρ / τ_b, pooled	0.718 / 0.562	0.664 / 0.491
	Iteration rank ρ, noRC / RC	0.714 / 0.683	1.000* / 0.690
System	Peak-memory rank ρ / τ_b	0.982 / 0.924	0.991 / 0.964
	Lowest-memory Top-3 overlap	3/3	3/3
	RC time/memory direction agreement	6 / 6	3 / 3
Component	Matrix median-latency ρ / τ_b	0.983 / 0.892	0.987 / 0.909
	Slowest matrix-problem Top-10 overlap	8/10	10/10
	Operator-category median ρ / JSD	0.503 / 0.436	0.531 / 0.356
	Communication-primitive median ρ / JSD	0.949 / 0.019	0.872 / 0.035
	Non-overlapped comm. share, GPU / NPU	17.0% / 28.5%	10.9% / 30.0%

tion is $\rho = 0.982$ and 0.991 ($\tau_b = 0.924$ and 0.964), and the three lowest-memory configurations coincide on both stacks for both models. The median GPU/NPU memory ratios are 0.77 and 0.70. Table 28 lists the per-configuration peak memory.

This agreement also holds when recomputation is fixed: the RC/noRC peak-memory rank correlations are 0.983/0.943 for Qwen3-0.6B and 0.976/1.000 for Qwen3-1.7B (the latter noRC subset contains three candidates). The agreement across these subsets supports using NPU profiles to identify relatively memory-efficient GPU strategies under either recomputation setting.

Recomputation behaves consistently across stacks. Across six matched base configurations for Qwen3-0.6B and three for Qwen3-1.7B, enabling recomputation increases iteration time and decreases peak memory in the same direction in every pair (100% directional agreement). The median time overhead is 32.9% (GPU) versus 44.0% (NPU) for the smaller model and 32.6% versus 38.8% for the larger one; the median memory saving is 85.6% versus 77.2% and 77.5% versus 69.7%, respectively.

Table 24. Calculon time, memory, and ranking accuracy on the configurations used for platform-specific calibration.

Model	Platform	Time MAPE (%)	Time ρ	Rank-mem. MAPE (%)	Config-mem. ρ
Qwen3-0.6B	V100	10.28	0.861	15.99	0.993
Qwen3-0.6B	910B3	10.19	0.950	13.20	0.986
Qwen3-1.7B	V100	11.22	0.853	12.56	0.993
Qwen3-1.7B	910B3	7.97	0.935	12.55	0.994

Table 25. GPU configuration selection using NPU-derived rankings and top-three shortlists.

Model	Paired candidates	Top-1 regret	GPU optimum in Top-3	Top-3 regret
Qwen3-0.6B	15	0.0%	Yes	0.0%
Qwen3-1.7B	11	6.9%	Yes	0.0%

Table 26. Cross-platform iteration-time rank agreement by recomputation setting.

Model	Stratum	n	ρ	τ_b	Top-3	Top-1 agreement
Qwen3-0.6B	all	15	0.718	0.562	2/3	yes
Qwen3-0.6B	noRC	6	0.714	0.600	2/3	yes
Qwen3-0.6B	RC	9	0.683	0.556	2/3	yes
Qwen3-1.7B	all	11	0.664	0.491	2/3	**no**
Qwen3-1.7B	noRC	3	1.000	1.000	3/3	yes
Qwen3-1.7B	RC	8	0.690	0.571	2/3	yes

F.5. Operator and Communication Behavior

Matrix-problem latency. Matrix operators are matched by batch size, dimensions (m, n, k), and recorded input-data-type labels. Type labels are normalized for capitalization, quotation marks, and whitespace before matching. Per-call latencies are aggregated across all paired configurations and all eight ranks to obtain a median for each problem on each stack. Each model has 27 shared problem types, and those types account for 72.8% and 76.7% of GPU matrix calls. Ranking the shared problems by median latency gives $\rho = 0.983$ and 0.987 ($\tau_b = 0.892$ and 0.909), with slowest-shape Top-10 overlap of $8/10$ and $10/10$. The relative latency ordering of the shared matrix problems is therefore preserved, whereas absolute speed differs by a median factor of 2.79 and 2.84 (Table 29).

Operator and communication composition. For each paired configuration, durations are summed across ranks by operator category or canonical communication primitive. Primitive names are normalized, with send/receive variants grouped as point-to-point communication. Vectors are aligned over the union of categories, filling absent entries with zero, and normalized separately for operators and communication. Spearman correlation and JSD are computed per configuration and summarized by their medians.

Communication-primitive time shares have median rank correlations of 0.949 and 0.872 and median JSD of 0.019 and 0.035; the dominant primitive matches in 86.7% and 72.7% of configurations. Operator-category correlations are 0.503 and 0.531, with median JSD of 0.436 and 0.356. Table 30 reports time shares pooled over the paired configurations within each model and stack. The category breakdown makes the implementation differences visible: on 910B3, *other* and *shape_transform* account for a large share of device time, while *attention* is a separate category only on the GPU stack. These categories follow each profiler's classification and reflect each stack's kernel composition.

Communication overlap. Non-overlapped communication share is computed per rank as non-overlapped communication time divided by stage time. Rank-level shares are summarized by their median within each configuration, followed by the median over paired configurations. The median non-overlapped communication share is 17.0% and 10.9% on V100 versus 28.5% and 30.0% on 910B3, and the per-configuration correlation of this share across stacks is only $\rho = 0.179$ and 0.364. Communication primitive ordering is more consistent across stacks than the degree of communication overlap.

Table 27. GPU and NPU iteration times by configuration, ordered by GPU runtime within each model.

Configuration	GPU (s)	NPU (s)	GPU/NPU	GPU rank	NPU rank
Qwen3-0.6B					
dp4-tp2-pp1-cp1-noRC	6.40	3.93	1.63	1	1
dp2-tp4-pp1-cp1-noRC	6.69	6.11	1.09	2	7
dp2-tp2-pp2-cp1-noRC	7.20	4.48	1.61	3	2
dp1-tp4-pp2-cp1-noRC	7.33	6.74	1.09	4	10
dp8-tp1-pp1-cp1-RC	7.53	4.69	1.61	5	3
dp1-tp2-pp4-cp1-noRC	8.28	5.10	1.62	6	4
dp4-tp2-pp1-cp1-RC	8.60	5.63	1.53	7	6
dp2-tp4-pp1-cp1-RC	8.93	9.48	0.94	8	12
dp4-tp1-pp2-cp1-RC	8.95	5.50	1.63	9	5
dp1-tp8-pp1-cp1-noRC	9.06	11.39	0.80	10	14
dp2-tp2-pp2-cp1-RC	9.46	6.27	1.51	11	9
dp1-tp4-pp2-cp1-RC	9.69	9.75	0.99	12	13
dp2-tp1-pp4-cp1-RC	10.25	6.25	1.64	13	8
dp1-tp2-pp4-cp1-RC	10.65	6.98	1.53	14	11
dp1-tp8-pp1-cp1-RC	12.86	18.36	0.70	15	15
Qwen3-1.7B					
dp2-tp4-pp1-cp1-noRC	8.85	6.42	1.38	1	2
dp8-tp1-pp1-cp1-RC	9.46	5.49	1.73	2	1
dp1-tp4-pp2-cp1-noRC	9.82	7.43	1.32	3	4
dp1-tp8-pp1-cp1-noRC	10.75	12.09	0.89	4	10
dp4-tp2-pp1-cp1-RC	11.20	7.04	1.59	5	3
dp2-tp4-pp1-cp1-RC	11.74	9.26	1.27	6	8
dp2-tp2-pp2-cp1-RC	12.34	7.83	1.58	7	6
dp1-tp4-pp2-cp1-RC	12.82	10.14	1.26	8	9
dp2-tp1-pp4-cp1-RC	13.28	7.46	1.78	9	5
dp1-tp2-pp4-cp1-RC	14.10	8.69	1.62	10	7
dp1-tp8-pp1-cp1-RC	14.28	16.78	0.85	11	11

Table 28. GPU and NPU peak memory by configuration, measured as the maximum across ranks.

Configuration	GPU (GiB)	NPU (GiB)	GPU/NPU	GPU rank	NPU rank
Qwen3-0.6B					
dp1-tp8-pp1-cp1-RC	1.47	2.18	0.67	1	1
dp1-tp4-pp2-cp1-RC	2.18	3.16	0.69	2	2
dp2-tp4-pp1-cp1-RC	2.37	3.25	0.73	3	3
dp1-tp2-pp4-cp1-RC	3.64	5.26	0.69	4	5
dp2-tp2-pp2-cp1-RC	3.67	5.09	0.72	5	4
dp4-tp2-pp1-cp1-RC	4.14	5.37	0.77	6	6
dp2-tp1-pp4-cp1-RC	6.44	8.91	0.72	7	7
dp4-tp1-pp2-cp1-RC	6.68	8.94	0.75	8	9
dp8-tp1-pp1-cp1-RC	7.70	9.63	0.80	9	10
dp1-tp8-pp1-cp1-noRC	8.55	8.92	0.96	10	8
dp1-tp4-pp2-cp1-noRC	14.99	13.23	1.13	11	11
dp2-tp4-pp1-cp1-noRC	15.41	14.25	1.08	12	12
dp1-tp2-pp4-cp1-noRC	28.11	22.97	1.22	13	14
dp2-tp2-pp2-cp1-noRC	28.13	22.74	1.24	14	13
dp4-tp2-pp1-cp1-noRC	29.03	24.91	1.17	15	15
Qwen3-1.7B					
dp1-tp8-pp1-cp1-RC	2.81	4.54	0.62	1	1
dp1-tp4-pp2-cp1-RC	3.61	5.72	0.63	2	2
dp2-tp4-pp1-cp1-RC	4.26	6.13	0.69	3	3
dp1-tp2-pp4-cp1-RC	5.38	8.34	0.64	4	5
dp2-tp2-pp2-cp1-RC	5.71	8.24	0.69	5	4
dp4-tp2-pp1-cp1-RC	7.16	9.31	0.77	6	6
dp2-tp1-pp4-cp1-RC	8.92	12.73	0.70	7	7
dp1-tp8-pp1-cp1-noRC	11.24	14.02	0.80	8	8
dp8-tp1-pp1-cp1-RC	12.96	15.66	0.83	9	9
dp1-tp4-pp2-cp1-noRC	18.10	18.85	0.96	10	10
dp2-tp4-pp1-cp1-noRC	18.96	20.20	0.94	11	11

Table 29. Median per-call latency of representative matched matrix operators on GPUs and NPUs.

Model	Matrix problem	GPU (µs)	NPU (µs)	GPU/NPU
Qwen3-0.6B	b8: 4096x4096 @ 4096x128	654.5	388.4	1.69
Qwen3-0.6B	b8: 4096x128 @ 128x4096	693.3	234.0	2.96
Qwen3-0.6B	b16: 4096x4096 @ 4096x128	1187.3	697.2	1.70
Qwen3-0.6B	b4: 4096x4096 @ 4096x128	278.4	133.5	2.09
Qwen3-0.6B	b4: 4096x128 @ 128x4096	352.4	97.0	3.63
Qwen3-0.6B	b16: 4096x128 @ 128x4096	1371.0	484.1	2.83
Qwen3-0.6B	b1: 4096x1024 @ 1024x1024	125.5	41.9	2.99
Qwen3-0.6B	b1: 4096x1024 @ 1024x3072	320.2	111.4	2.87
Qwen3-0.6B	b2: 4096x4096 @ 4096x128	177.6	53.7	3.31
Qwen3-0.6B	b1: 4096x1024 @ 1024x1536	168.8	63.7	2.65
Qwen3-1.7B	b1: 4096x2048 @ 2048x3072	603.4	212.4	2.84
Qwen3-1.7B	b1: 4096x2048 @ 2048x6144	1191.1	402.7	2.96
Qwen3-1.7B	b1: 4096x3072 @ 3072x2048	617.2	208.8	2.96
Qwen3-1.7B	b8: 4096x4096 @ 4096x128	710.1	387.5	1.83
Qwen3-1.7B	b1: 4096x2048 @ 2048x1536	309.2	118.0	2.62
Qwen3-1.7B	b1: 4096x1536 @ 1536x2048	374.8	111.9	3.35
Qwen3-1.7B	b4: 4096x4096 @ 4096x128	284.5	134.9	2.11
Qwen3-1.7B	b1: 4096x2048 @ 2048x12288	2395.5	795.2	3.01
Qwen3-1.7B	b1: 4096x6144 @ 6144x2048	1111.7	412.1	2.70
Qwen3-1.7B	b1: 4096x2048 @ 2048x2048	470.8	147.4	3.19

Table 30. Operator-category and communication-primitive time shares on GPUs and NPUs.

Component	0.6B GPU	0.6B NPU	1.7B GPU	1.7B NPU
Operator categories				
matmul	28.6%	14.9%	38.9%	22.0%
memory_movement	20.7%	4.2%	16.2%	3.3%
elementwise	18.6%	17.2%	19.1%	19.9%
attention	14.7%	0.0%	11.2%	0.0%
memory_init	13.0%	0.2%	10.1%	0.2%
other	2.3%	40.3%	2.2%	34.1%
shape_transform	1.6%	22.8%	1.4%	20.0%
activation	0.5%	0.2%	0.7%	0.3%
embedding	0.0%	0.0%	0.0%	0.0%
optimizer	0.0%	0.1%	0.1%	0.1%
indexing_or_embedding	0.0%	0.0%	0.0%	0.0%
normalization	0.0%	0.0%	0.0%	0.0%
Communication primitives				
all_reduce	54.4%	82.9%	54.3%	84.1%
point_to_point	42.3%	15.3%	41.3%	14.0%
broadcast	1.8%	0.9%	1.8%	0.7%
reduce_scatter	1.5%	0.5%	2.7%	0.6%
all_gather	0.0%	0.4%	0.0%	0.5%