arXiv:2611.13566v1 [cs.LG] 10 Nov 2026

Metamorphic testing of learned AC optimal power flow surrogates reveals representation dependence

Stephen Peng* **Vineet Jagadeesan Nair***
AI Science Institute, Boston Consulting Group
Peng.Stephen@bcg.com, JagadeesanNair.Vineet@bcg.com

Abstract

Evaluating a learned AC optimal power flow (AC-OPF) surrogate on one grid serialization cannot show whether its accuracy depends on that encoding. We test this dependence with thirteen metamorphic relations, each specifying when a transformation is valid and how to compare its outputs. Our audit covers 17 model and pipeline configurations: two released trained checkpoints and 15 random-weight architecture or mechanism controls. Every configuration violates at least six relations, even at reporting tolerances ten times tighter or looser. For `GridSFM-Open v1.1`, source inspection and intermediate-tensor interventions point to an indexing bug and three noncanonical or numerically unstable preprocessing choices; a repaired checkpoint remains untested. In a separate case500 training experiment, augmentation lowers transformed-view error in three of four models and median discrepancy and constraint-violation P95 in all four. These benefits depend on the training protocol: some relations worsen, and every evaluated prediction remains AC infeasible. A GPU/CPU eigensolver mismatch accounts for most of the large raw-Laplacian-baseline gain. We recommend reporting representation consistency alongside accuracy, feasibility, and cost. These results establish neither deployment readiness nor universal invariance.

1 Introduction

AC-OPF chooses generator dispatch and network states to minimize cost while satisfying physical equations and operating limits. Learned surrogates approximate this calculation. Published evaluations report accuracy, feasibility, and speed on held-out scenarios, each stored in one serialization [1–3]. A bus can move to another row without changing the grid. Eligible branches can be reversed, and each island can use a different angle reference. Component order and model-visible base-voltage metadata can also change without changing the declared problem, under the conditions specified below. Once predictions are returned to their physical devices and a common angle reference, a representation-consistent deterministic surrogate should give the same answer. Testing one encoding leaves this property unmeasured. Error on one serialization alone cannot tell whether accuracy reflects the physical problem or a convention used to encode it. This gap between the measured quantity and the intended claim is a construct-validity problem in the evaluation protocol.

Metamorphic testing checks necessary relationships between outputs for linked inputs without requiring a unique reference output [4, 5]. We define thirteen AC-OPF transformations in two tiers. Representation relations re-encode the same grid. Electrical-equivalence relations change component or graph structure while preserving the aggregate electrical problem. Each relation specifies an output map that restores the original identities and angle reference, aggregating predictions where needed. Algebraic conditions establish when the relation is valid, and sampled solver controls check

*Equal contribution.

40th Conference on Neural Information Processing Systems (NeurIPS 2026). Workshop: TAE (Trust-AI-Eval): Can We Trust AI Evaluation?.

its implementation (Sec. C.1 gives the archived admission rules and coverage limits). One admitted violation refutes invariance for that pipeline and relation. A finite set of passes shows only that the tests found no violation.

Our main contributions are:

- We define thirteen metamorphic relations with explicit output maps, numerical admission criteria, and evidence tiers that separate open-weight checkpoints, random initializations, and reimplementations of common AC-OPF neural architecture components (Secs. A and 3).
- We audit 17 model and pipeline configurations, including two open-weight checkpoints. Every configuration fails at least six relations, and for both released checkpoints the trained and random-weight versions agree on all 13 aggregate threshold classifications (Secs. 4.1 and 4.3).
- We localize the first detected discrepancies in the grid foundation model `GridSFM-Open v1.1` to derived preprocessing features, implicating one indexing bug and three noncanonical or numerically unstable feature choices; we have not validated a repaired checkpoint (Sec. 4.4).
- We test the relations as training augmentations under an AC constraint-penalty objective. Transformed-view error falls in three of four models, and median discrepancy and AC constraint-violation P95 fall for all four. Every prediction remains AC infeasible, and an eigensolver mismatch accounts for most of the largest gain (Sec. 4.5).
- We propose reporting practices and a signed-manifest scheme through which closed-weight groups can publish auditable consistency evidence without exposing parameters (Sec. 5).

Code is available in an anonymous repository.

2 Related work

2.1 Learned AC-OPF surrogates

Learned AC-OPF models map demands, generator limits, costs, constraints, and network parameters to voltage magnitudes, angles, and dispatch. They may replace a nonlinear optimizer, provide its starting point, or work with a correction step. Fast predictions with low supervised error can still violate power balance or operating limits and can have high dispatch cost. Physics-informed losses and differentiable completion or correction layers address some of these problems [6, 7].

Early surrogates often used fully connected networks tied to a fixed topology and variable order. Graph neural networks use bus–branch connectivity and shared message-passing rules, allowing permutation-equivariant computation [8, 9]. Homogeneous and heterogeneous AC-OPF models use this structure [1, 2, 10, 11]. Distinguishing buses, generators, loads, shunts, lines, and transformers preserves useful physical information, but also requires consistent handling of each family's order, identity, and aggregation.

Local message passing can struggle to represent long-range electrical coupling. Global attention, graph-transformer blocks, and positional encodings extend its reach [3, 12, 13]. Laplacian eigenvectors, effective resistance, random-walk statistics, and landmark summaries supply global information, but some also depend on arbitrary node indices, eigenvector signs or bases, selected reference nodes, or row order. These features can break equivariance in the full feature-to-prediction pipeline even when the graph layers themselves are equivariant.

2.2 Metamorphic testing and equivariance

Metamorphic testing addresses the test-oracle problem: how to test software when the correct output is unavailable or costly to obtain. It checks necessary relationships between outputs for linked inputs [4, 5]. If an input transformation has a known effect on a correct output, a violated relationship can reveal a defect without a unique reference label. For example, symmetry-based relations have been used to test human-trajectory prediction models [14].

The AC-OPF relations have different mathematical forms. A bus permutation requires the same permutation of bus outputs; after mapping them back to physical identities, predictions should be

unchanged. A reference-angle change is a gauge transformation: it shifts absolute angles while preserving angle differences, flows, and dispatch after canonicalization. Component and parallel-branch splits can change output dimension and require many-to-one aggregation. We reserve *built-in equivariance* for a property that holds at every parameter setting. Consistency observed in one fitted checkpoint is an empirical result that depends on the tested input distribution.

Learned-OPF studies typically report accuracy, optimality, feasibility, speed, and generalization to changed loads or contingencies [1–3]. An outage changes the physical optimization problem; an admitted relation in our audit preserves it after a known output map. In the literature we reviewed, we found no AC-OPF study combining a relation catalog with solver controls, trained-checkpoint tests, architecture-level failure localization, and an augmentation experiment. Our contribution is this OPF-specific construction and audit, building on established metamorphic-testing and graph-equivariance methods.

Relation to formal guarantees and invariant spectral features. Worst-case verification asks whether a predictor meets a constraint or quality bound throughout an input domain; Venzke et al. [15], for example, study neural DC-OPF guarantees. Our tests find counterexamples on sampled, conditionally equivalent AC-OPF encodings and cannot certify a whole domain. The raw Laplacian coordinates in our diagnostic baseline have known sign and eigenspace-basis ambiguities. SignNet and BasisNet address these symmetries explicitly [16]; RWSE avoids these ambiguities through a different construction. We do not benchmark trained SignNet or BasisNet, so our LapPE findings should not be generalized to all spectral graph models.

Related scientific-surrogate metamorphic testing. Li et al. [17] use metamorphic tests for scientific ML surrogates with explicit applicability conditions and numerical thresholds. We share their requirement to check that a relation applies before interpreting a violation. Here we develop the AC-OPF relations and output maps, distinguish trained-checkpoint tests from random-weight controls, and audit complete grid-model pipelines. Validity-gated metamorphic testing and its use beyond power systems precede this work; our experiments cover power systems only.

2.3 Evaluation validity beyond power systems

Evaluation choices can affect conclusions in other domains as well. Meaning-preserving prompt-format changes can shift measured accuracy by tens of points [18]; model rankings from one prompt may not hold across prompts [19]; and undetected data contamination can inflate scores [20]. Systematic reviews also find that many benchmarks leave construct validity and reporting practices unexamined [21, 22].

A heuristic perturbation may change the task it is meant to test. Our relations preserve the declared AC-OPF problem under explicit conditions, and sampled numerical controls check their implementation before we interpret the model comparisons. AC-OPF supplies known physical invariances that can be checked numerically. The same audit structure could test consistency in other domains where task-preserving transformations can be verified.

3 Metamorphic testing for AC-OPF

For a deterministic surrogate f, relation r transforms serialized case d into an equivalent input $T_r d$. Its output map P_r returns predictions to their physical devices, aligns angles independently within each connected island, and aggregates split-device outputs where needed. We measure departure from $P_r f(T_r d) = f(d)$ by

$$\Delta_r(d) = \max_{c \in \mathcal{C}} \frac{\mathrm{MAE}_c(P_r f(T_r d), f(d))}{\tau_c}, \qquad \mathcal{C} = \{V_m, V_a, P_g, Q_g\}. \tag{1}$$

Here V_m, V_a denote voltage magnitude and angle, and P_g, Q_g denote active and reactive generation. Each channel's mean absolute error (MAE) averages absolute differences over its elements. We set $\tau_c = 10^{-6}$ in native units: per unit for V_m, P_g, Q_g and radians for V_a. Angles are re-anchored without further modulo wrapping. A value of $\Delta_r = 1$ reaches the reporting threshold; it is neither a percentage error nor an AC-feasibility limit.

Verdicts and summaries. A cell in the 17-configuration screen passes only when every executed, applicable draw satisfies $\Delta_r \leq 1$. We record partial coverage where controls are unavailable and check maximum component errors separately at 10^{-5}. The five-grid trained-versus-random comparison summarizes magnitude differently: it averages draws within each source scenario, takes the median over sources, and then takes the worst grid, keeping weight seeds separate. A below-threshold aggregate does not mean that every draw passed. Sec. C.1 specifies the numerical controls, and Sec. C.2 reports threshold sensitivity and component tails.

Validity and scope. Eight representation-equivalence (`REP`) relations change labels, orientation, angle-reference conventions, or a field used by a particular pipeline. Five electrical-equivalence (`ELC`) relations change component or graph bookkeeping while preserving the declared aggregate problem. Algebraic conditions establish applicability; the saved solver controls check implementation on a sample of evaluation views. Sec. A specifies the transformations, applicability conditions, and output maps. Secs. B.1 and B.2 discuss non-uniqueness and the map from serialized data to the physical problem.

Worked output map example. If `bus-relabel` stores old buses as $(3, 1, 2)$, restore voltage predictions as $(\widehat{V}_2', \widehat{V}_3', \widehat{V}_1')$ and match generator outputs by physical device identity. If the reference bus also changes, subtract the restored reference angle from every angle in that island before comparing. A discrepancy measures deterministic representation dependence. Either prediction could still be feasible, because a nonconvex AC-OPF problem can admit several physically valid solutions.

Table 1: The thirteen relations at a glance. The two tiers describe the output-map contract; neither tier asserts that every relation is a bijective group action.

Tier	Relations	What is held fixed / restored
`REP`	`component-relabel`, `bus-relabel`, `row-order`	Physical identities; inverse permutations
`REP`	`branch-reversal`, `tap-side-flip`	Branch two-port equations; endpoint/attribute transport
`REP`	`ref-bus-move`, `island-ref-mix`	Per-island angle gauge after mapping back
`REP`	`basekv-scale`	Exported per-unit target; conditional on schema semantics
`ELC`	`component-split`, `parallel-split`	Aggregate injection, costs or admittance; sum split outputs
`ELC`	`load-as-gen`, `null-pad`	Net bus injection; undo device recoding/padding
`ELC`	`island-union`	Original disconnected subproblem; restrict to its outputs

A relation tests the full adapter, preprocessing, prediction, and output-restoration path. For example, `basekv-scale` is valid here only because the exported target equations are already in per unit and the changed field is model-visible metadata; the relation does not extend to rescaling physical voltages. We check conditional applicability before interpreting a discrepancy. A violation outside a model's documented input contract diagnoses interface compatibility and may not indicate a defect in its intended use.

We use random initialization as a counterexample screen for built-in consistency. One violation refutes architectural invariance to that relation, whereas a pass means only that the screen found no violation. The DC3 controls, for example, test only the partition-and-completion choice with a stand-in; they do not run DC3's Newton power-flow completion or correction loop. We separately compare released (open-weight) and random-weight versions of `GridSFM-Open v1.1` and `LUMINA-2M v0.1.3` to test whether these training states change the aggregate threshold classification [23–27]. Scenarios come from the N-1 topology-perturbed `OPFData` release, with `PGLib` networks and `PowerModels.jl` used for numerical solver controls [28–30]. The full model inventory, sampling scheme, and solver admission rules appear in Secs. C and C.1.

Numerical convergence alone does not admit a relation. All 162 pairs of solves in the expanded-relation controls converge, but only 17 of 18 grid–relation pairs meet the admission criteria. Case14 `component-split` is withheld because its relative objective difference, 7.26×10^{-9}, exceeds the 10^{-9} limit. The screen reuses these sampled controls without solving every model draw; fourteen `component-split` cells therefore retain only case500 evidence.

4 Audit results

4.1 Built-in consistency screen

Fig. 1 summarizes 17 model and pipeline configurations, of which two are trained open-weight checkpoints and 15 are random-weight architecture or mechanism controls (Table 6). No row passes all thirteen relations, and every row fails at least six. The six shared failures are `basekv-scale`, `ref-bus-move`, `island-ref-mix`, `component-split`, `load-as-gen`, and `null-pad`. At a relaxed native-channel threshold of 10^{-3}, all 17 configurations still fail at least four relations; excluding the pipeline-specific `basekv-scale` relation leaves at least three failures per configuration. At the primary 10^{-6} threshold, the same exclusion leaves 140 of 204 cells failed and at least five failed relations per configuration. This reaggregation of saved results tests threshold sensitivity only and does not calibrate numerical noise (Sec. C.2). Sec. 4.2 traces the pattern to specific design mechanisms.

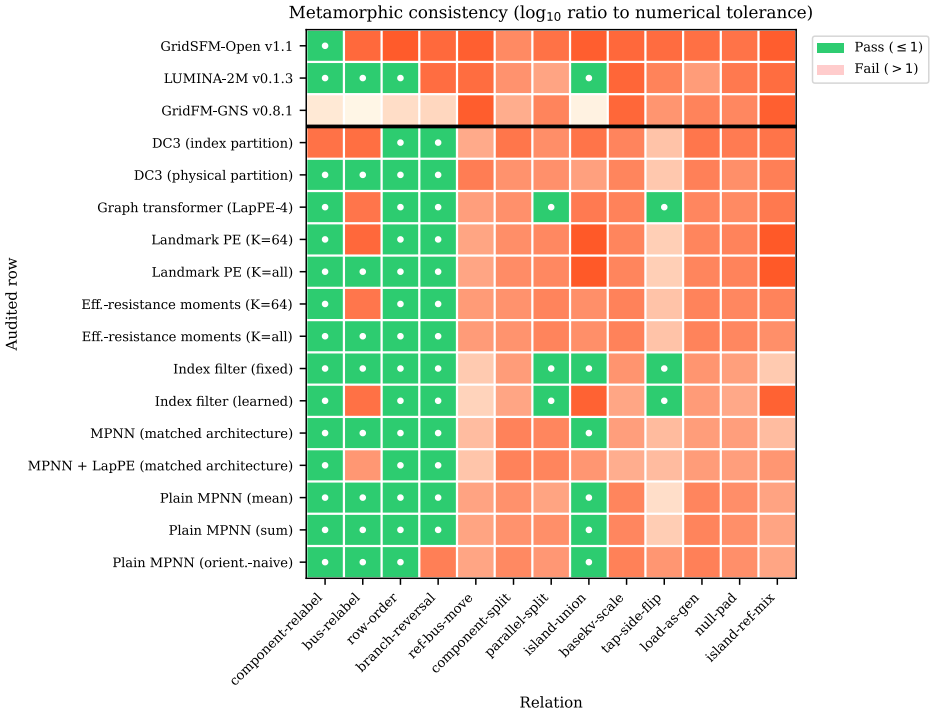

Figure 1: Metamorphic consistency for 17 model and pipeline configurations (two trained checkpoints and 15 random-weight architecture or mechanism controls) across 13 relations on case14/500. Each case uses 3 source scenarios, 3 sampled transformations per admitted case–relation pair, and 3 random initializations where applicable. Green denotes a pass at the prespecified 10^{-6} tolerance.

4.2 Architecture-level failure patterns

Relabeling gives the clearest distinction among designs. Those that tie features to physical components rather than array position pass `component-relabel`, `bus-relabel`, and `row-order`. The `row-order` relation fails only in the GridSFM-Open and GridFM processing paths, which use row position directly.

The `bus-relabel` relation fails in 8 of 17 rows. Each failing row uses index-derived or spectral coordinates through Laplacian or landmark positional encodings, index-partitioned heads, or index filters. The matched MPNN and LapPE+MPNN rows isolate one such choice: adding raw Laplacian coordinates turns `bus-relabel` and `island-union` from passes into failures, because eigenvector signs and eigenspace bases depend on ordering. Index-selected landmarks are buses chosen by array

position, so relabeling changes which physical buses are selected. Using all buses removes this selection dependence in the matched control; we have not validated a repaired trained checkpoint.

The `island-union` relation distinguishes local from global computation. Models with purely local, permutation-safe aggregation are unaffected by adding a disconnected island, while rows with global positional summaries or spectral coordinates fail. The `branch-reversal` relation fails only in rows that consume directed edge features without symmetrizing them. LUMINA shows the same pattern in the trained/random comparison (Sec. 4.3).

The six relations that fail in every configuration expose two kinds of preprocessing dependence. First, `ref-bus-move` and `island-ref-mix` change the angle reference, while `basekv-scale` changes model-visible metadata that does not affect the standard per-unit target after `OPFData` export. These failures show that the audited pipelines do not consistently canonicalize angle references or discard target-irrelevant fields. Second, `component-split`, `load-as-gen`, and `null-pad` change how device records represent a fixed bus-level injection. Their failures show dependence on component count or type beyond the electrical aggregates in bus balance. Standardizing inputs may address the first kind. The second may require a common representation for electrically equivalent components and aggregation of their contributions at each bus (Sec. F.1).

4.3 Trained versus random-weight consistency

This comparison uses three source scenarios from each of case14, case118, case500, case2000, and case4661, with three transformation draws per source–relation pair: $3 \times 13 \times 3 = 117$ comparisons per grid and model/seed run. Each released checkpoint and its three random initializations receive the same scenarios and draws. Trained runs stored under three seed labels repeat the same checkpoint and draws; the independent units remain three sources per grid. The grids differ in their relation to each checkpoint's documented training coverage, so this is not a common unseen-topology benchmark (Sec. C.1).

Trained and random-weight versions of both open-weight backbones agree on all 13 threshold classifications of the source-median, worst-grid aggregate (Fig. 4). The aggregate is below threshold for GridSFM only on `component-relabel`; for LUMINA it is additionally below threshold on `bus-relabel`, `row-order`, and `island-union`. These below-threshold aggregates agree with the finite permutation and aggregation checks and do not establish an all-input guarantee. Both trained and random LUMINA variants fail `branch-reversal`, consistent with a directed message-passing path that does not transport the branch attributes needed to compensate for endpoint exchange. Passing label permutations therefore does not imply consistency under every `REP` relation.

Training changes most discrepancy magnitudes without consistently reducing them. We exclude `component-relabel` from ratio comparisons because its trained and random discrepancies are all below 10^{-6}, making ratios between near-zero values uninformative. Among the remaining finite case–relation cells, trained discrepancy exceeds the random control in 55 of 60 GridSFM cells and 50 of 58 LUMINA cells. The median ratios are $7.50\times$ and $6.23\times$, respectively, with substantial variation across grids. Several LUMINA reference-related cells have ratios below one.

A large ratio can result from a near-zero random denominator, while a ratio below one can still leave a discrepancy far above tolerance. Both absolute and relative values are therefore needed. Random networks are inaccurate and can appear consistent because their outputs are nearly constant; these ratios do not compare model quality. For the tested checkpoints, training changed magnitudes but eliminated no above-threshold aggregate discrepancy. Sec. D.1 reports per-relation values in both views, and Fig. 5 shows how training changes discrepancies on each grid.

4.4 Localization in released pipelines

For GridSFM, the bus-to-bus edge index is permutation-equivariant, and shuffling edge rows after feature construction leaves network predictions unchanged. The first detected discrepancy appears in derived features before the prediction network. Four candidate mechanisms explain where it can enter: effective-resistance summaries select up to 64 landmarks partly by bus index; the combined line/transformer offsets count appended arrays rather than edges; applying `sign` to a near-zero DC flow turns numerical noise into a categorical feature; and the cycle basis uses a spanning tree weighted by edge-row order. We reproduced these observations against pinned commit `1ca775f` [25].

We diagnose these mechanisms by substituting aligned intermediate tensors; this procedure does not produce a repaired-and-retrained checkpoint. For case500 source 9245 under bus relabeling (draw 0), the saved cumulative intervention changes the legacy equal-channel discrepancy score from 30.4769 unpatched to 30.4712 after node/cycle feature substitution and to 5.2535×10^{-5} after cycle-incidence substitution. This score averages channel errors divided by 10^{-3}, unlike the paper's maximum-channel Δ_r at 10^{-6}. The final-to-initial ratio, 1.724×10^{-6}, measures the relative residual and is not an absolute pass threshold.

Other relations retain substantial residuals under the same interventions. In a separate network-only probe, edge-row shuffling gives zero movement, while branch-node relabeling gives a nonzero 3.51×10^{-5} discrepancy on the legacy scale. These tests locate substantial sensitivity before the prediction network. They neither prove that the entire network is invariant nor validate four independent repairs.

The candidate mechanisms suggest four repairs: select landmarks independently of arbitrary indices or summarize all buses; track cumulative edge counts when joining component families; apply a physically justified dead band or smooth mapping to near-zero flows; and construct cycle information independently of row order. A patch-and-retest experiment is still needed to establish whether these changes restore consistency. A checkpoint may have adapted to its existing features, so each repair should be followed by retraining or fine-tuning and an accuracy check.

LUMINA localizes differently from GridSFM. Its below-threshold results under component relabeling, bus relabeling, row ordering, and disconnected-island composition hold for both trained and random weights (Sec. 4.3), and its branch-reversal failure is consistent with orientation-sensitive directed processing. The random-weight comparison and network-only tests thus separate preprocessing defects, architectural transport choices, and changes in magnitude acquired during training.

Representation changes also alter the cost implied by GridSFM's active-power predictions. For the trained checkpoint, the median of the absolute cost change normalized by each scenario's solver reference cost is 23.4% under `row-order`, 50.2% under `island-union`, and 51.3% under `island-ref-mix` (Table 9). These predictions are AC infeasible: the values measure encoding sensitivity of implied cost, not realized dispatch losses or optimality gaps. Sec. D.2 defines the normalization and reports all relations.

4.5 Relations as data augmentation

We test whether training on the relations reduces representation dependence. We train baseline and augmented versions of an MPNN, an MPNN with raw Laplacian positional encodings (LapPE+MPNN) [12, 31], an MPNN with random-walk structural encodings that avoid sign and eigenvector-basis ambiguities (RWSE+MPNN) [12, 32], and LUMINA. All models minimize equal-channel MSE against the solver solution plus an AC constraint-violation penalty. For each training instance, the augmented variant applies one relation with probability 0.5. It draws from a frozen architecture-specific set chosen from earlier trained-baseline measurements, independently of the random-weight screen.

We evaluate 50 held-out scenarios per seed in canonical and transformed forms under all thirteen relations. Accuracy is the unweighted mean of native-unit MAEs across V_m, V_a, P_g, Q_g. This convention combines p.u. and radians, so it should be read alongside the per-channel values in Table 10. Sec. E.1 gives the training objective, relation sets, hyperparameters, and metrics.

Augmentation lowers mean MAE on transformed scenarios by 87.0% for LapPE+MPNN, 41.9% for LUMINA, and 1.1% for RWSE+MPNN, and raises it by 3.2% for MPNN (Table 2). Canonical MAE improves by 0.4% for RWSE+MPNN and worsens by 3.9% for MPNN and 3.0% for LUMINA. Median normalized discrepancy Δ_r falls for all four models. Some individual relations worsen, out-of-distribution grid generalization is mixed, and no prediction from our runs satisfies both the inequality and equality tolerances. Per-seed and relation-level effects appear in Sec. E.2, with feasibility and out-of-distribution results in Sec. E.3.

The relation-level results show what the aggregate obscures. Each arm's per-relation summary takes the largest source median across seeds. LapPE+MPNN worsens on `null-pad`, `component-split`, and `load-as-gen`; RWSE+MPNN worsens on `tap-side-flip`, `ref-bus-move`, and `island-ref-mix`. For example, their discrepancies increase by 29.4% on

Table 2: Aggregate augmentation outcomes under the AC constraint-penalty objective on held-out case500 scenarios (baseline → augmented; three seeds). Median Δ_r summarizes each architecture's own trained relation set and supports within-model comparisons. Constraint columns report mean ± sample standard deviation. Within each seed, we take each source scenario's worst violation over its 53 canonical and transformed views and then P95, the linearly interpolated 95th percentile, across 50 source scenarios. Every arm has 0% AC-feasible source scenarios at 10^{-6}. Equality is the larger of the maximum nodal active- and reactive-power balance violations in Eqs. (3) and (4), expressed in per unit. The pooled inequality metric is dimensionless: each family is divided by its 10^{-6} tolerance before taking the maximum, and displayed values are in millions of tolerances. Units, metric definitions, and full results appear in Sec. E.1 and Tables 10 and 11.

Model	Canonical MAE	Transformed MAE	Median Δ_r	Eq. violation P95 (p.u.)	Ineq. violation P95 ($10^6\times$ tol.)
MPNN	$0.01295 \to 0.01345$	$0.01312 \to 0.01353$	$852 \to 595$	$58.1\pm18.7 \to 42.2\pm2.3$	$19.3\pm4.2 \to 11.9\pm1.6$
LapPE+MPNN	$0.09983 \to 0.01309$	$0.10211 \to 0.01318$	$1{,}000 \to 805$	$113.9\pm42.1 \to 39.5\pm5.2$	$42.1\pm4.6 \to 12.7\pm3.6$
RWSE+MPNN	$0.01255 \to 0.01250$	$0.01271 \to 0.01257$	$883 \to 563$	$51.8\pm10.9 \to 33.5\pm2.4$	$18.8\pm2.8 \to 10.2\pm0.8$
LUMINA	$0.01834 \to 0.01889$	$0.03594 \to 0.02089$	$4{,}630 \to 1{,}260$	$871\pm342 \to 164\pm11$	$276\pm91 \to 20.1\pm2.1$

`null-pad` and 21.2% on `tap-side-flip`, respectively. In these per-relation summaries, MPNN improves on all seven relations used in its training and LUMINA on all nine (Table 11). A lower median therefore does not establish improvement on every relation.

Most of LapPE+MPNN's large improvement in Table 2 comes from a mismatch between the GPU eigensolver used in training and the CPU eigensolver used in this audit, which may also inflate its baseline violation P95 values (Sec. E.4). Both eigensolver paths leave eigenvector signs and bases unfixed. With only the backend changed, the same checkpoints' clean-cohort baseline MAE rises from 0.0125 to 0.0885, while the augmented arm changes by 3%. Median augmented discrepancies remain 563 to $1{,}260$ times the reporting threshold. Neither these reductions nor the lower constraint P95 values establish exact consistency, AC feasibility, or deployment readiness.

Scenario-clustered constraint-violation P95 also falls under augmentation for all four models under the same loss function. We have not tested whether training on invariance-preserving transformations causes this reduction, and the reduction does not show improved physical understanding.

5 Discussion

Representation consistency asks whether a deterministic pipeline returns the same physically matched prediction for equivalent encodings. It is separate from accuracy, feasibility, and optimality. Aggregation keyed to physical components, permutation-safe features, symmetric handling of branch attributes, and a common angle reference can guarantee some relations by construction (Sec. F.1). Augmentation instead encourages consistency on transformations represented in the training data.

The four augmentation models fail different relations. LUMINA's largest baseline failures are `branch-reversal` and `basekv-scale`; LapPE+MPNN's are `bus-relabel`, `island-union`, `island-ref-mix`, and `parallel-split`. A single-relation benchmark would therefore miss some models' largest failures. Orbit-averaging theory motivates augmentation under compact group actions, but does not directly cover every conditional, pipeline-specific, or dimension-changing relation in this catalog [33–35] (Sec. B.4).

5.1 Toward benchmarks that test invariance

Current learned-OPF benchmarks generally sample physical scenarios but evaluate each under one serialization. We propose testing a small set of metamorphic variants as repeated views of each physical case. Developers can run this audit before release; operators and planners can use documented inference access. For each held-out scenario, check which relations apply, generate K valid encodings, restore predictions to physical identities and a common angle reference, and inspect the worst discrepancy alongside per-channel accuracy, feasibility, optimality, and dispatch cost. Keep failed admissions and infeasible predictions in the record. The source scenario remains the independent sampling unit: permutations and splits are repeated measurements, not extra data points. Table 3 lists the minimum reporting fields.

Metamorphic variants serve two different roles. As an *evaluation augmentation*, they test whether predictions and reported conclusions hold across valid encodings. As a *training augmentation*, they change the training objective and may improve learned invariance. A finite pass after augmented training is evidence about the tested checkpoint and inputs; it certifies neither architectural equivariance nor operational use. A claim of built-in equivariance still needs an analytical argument. Benchmarks should retain random-initialization or architecture-level screens to search for counterexamples to that claim.

Table 3: Minimum reproducibility fields for a metamorphic learned-OPF benchmark.

Record	Minimum contents
Data and labels	Release, split, topology membership, scenario identifier, units, angle-reference convention, and label-generation pipeline.
Relation	Name, tier, parameters, applicability decision, transformed-case hash, and output map.
Solver control	Status, objective difference, equality and inequality constraint violations, and admission decision for both views.
Model	Code and checkpoint version, checksums, preprocessing contract, feature order, normalization, and output semantics.
Sampling	Source-scenario count, sampled transformations, weight seeds, and the clustering unit used for uncertainty.
Metrics	Per-channel median, tail, and worst discrepancy plus accuracy, AC constraint violations, bounds, feasibility, and cost.
Negative records	Unsupported relations, failed solver admission checks, missing joins, infeasible predictions, and threshold saturation retained rather than dropped.
Artifact	Case–relation–seed records, configurations, environment, and one command or script that regenerates tables and figures.

5.2 Open weights, transparency, and black-box auditing

Source code and a training recipe let an auditor test an architecture at random initialization. Testing learned behavior also requires the trained state, exact preprocessing contract, normalization statistics, feature-column order, and a runnable path from a documented case file to canonical outputs. GridSFM-Open and LUMINA-2M supplied these resources, making the audit and localization findings possible [23–27]. The releases let us compare trained and random versions of the same backbone, trace failures to input features or edge conventions, measure whether training reduced or amplified discrepancies, and evaluate augmentation without recreating an undocumented baseline. These findings, including the defects, depend on that openness. Sec. F.2 lists the minimum resources needed. Closed-weight groups can also participate through inference access: they can run the public transformations internally and publish signed or hashed manifests with applicability decisions, solver-control status, per-channel discrepancies, and failure tails. A shared format would make these consistency reports comparable without disclosing proprietary parameters.

5.3 Limitations

The thirteen relations cover the fibers enumerated in Sec. B.2, but do not generate every equivalent AC-OPF encoding (Sec. B.3). A pass means that this catalog found no violation in the tested cases, not that the pipeline is invariant to all equivalent inputs. Only two topology-varying surrogates supplied both a compatible checkpoint and a runnable inference path, which limits the breadth of our trained-weight study. Their agreement with the matched random-weight controls on all thirteen aggregate threshold classifications shows that these particular checkpoints did not eliminate the detected failures; it does not show that ordinary training can never learn a missing relation.

The augmentation study covers one training grid family, three seeds, one compute budget, and model-specific eligible relation sets. Some relations worsen, transfer to unseen grid sizes is mixed, and no arm produces an AC-feasible prediction. These results do not establish deployment readiness. PowerModels.jl and IPOPT provide high-accuracy numerical controls, not global-optimality certificates for the nonconvex problem. Distinct local optima can differ in weakly identified channels such as reactive power dispatch. A metamorphic discrepancy therefore establishes deterministic representation dependence of the surrogate without showing that either output is physically invalid.

5.4 Future work

Future benchmark versions can add transformations such as unequal component splits, inactive-bound changes, cost rescalings that preserve the optimal dispatch, or schema-level bus–branch conversions. Further work can determine which relation families admit generators or analytical guarantees, compare finite catalogs with property-based sampling over continuous fibers, and develop tolerance rules tied to solver conditioning. Model work should compare augmentation with mechanisms that enforce invariance by construction. Candidates include canonical unit and angle-reference preprocessing, frame averaging or learned canonicalization applied to the serialization group [36, 37], sign- and basis-invariant or provably stable spectral features in place of raw Laplacian coordinates [16, 38], output heads keyed by physical role and component identity rather than array position [11], and permutation-invariant set aggregation over co-located devices [39]. Feasibility layers and correction methods should be tested jointly with these representation fixes: post-processing cannot be assumed to preserve consistency. Larger training studies should add grid families, compute budgets, and checkpoint initializations while retaining negative relation-level outcomes.

6 Conclusion

A single-serialization evaluation cannot separate a learned AC-OPF surrogate's accuracy on the physical problem from its dependence on the encoding. Our thirteen transformations test that dependence with explicit output maps, validity conditions, and sampled solver admission checks. Each of the 17 audited configurations violates at least six relations. Six fail in every configuration: two change angle-reference conventions, one changes a model-visible base-voltage field outside the per-unit target, and three change how device records represent a fixed bus injection. For both released checkpoints, training changes discrepancy magnitudes, usually upward, without changing the reported aggregate threshold classifications. These finite observations do not establish which relations training can or cannot learn. In `GridSFM-Open v1.1`, intermediate-tensor interventions point to one indexing bug and three noncanonical or numerically unstable preprocessing choices. A small network-only residual remains, and a repaired checkpoint is still untested.

Training with these transformations lowers transformed-view error in three of four models and median discrepancy and AC constraint-violation P95 in all four. Individual relations and cross-grid transfer show mixed results, an eigensolver mismatch accounts for most of the largest gain, and every evaluated prediction remains AC infeasible. Representation consistency should therefore be reported alongside accuracy, optimality, and feasibility. Once a relation is admitted, measuring consistency requires model inference on transformed inputs. Learned-OPF benchmarks should evaluate several serializations of each physical scenario and compare predictions after restoring their physical identities and angle reference. The same audit structure could apply wherever task-preserving transformations can be checked independently of the model; a second domain remains untested.

Acknowledgments and Disclosure of Funding

The authors received no funding for this work and declare no competing interests.

References

[1] Luis Piloto, Sofia Liguori, Sephora Madjiheurem, Miha Zgubič, Sean Lovett, Hamish Tomlinson, Sophie Elster, Chris Apps, and Sims Witherspoon. CANOS: A Fast and Scalable Neural AC-OPF Solver Robust To N-1 Perturbations, March 2024. URL http://arxiv.org/abs/2403.17660. arXiv:2403.17660 [cs.LG].

[2] Shaohui Liu, Chengyang Wu, and Hao Zhu. Topology-Aware Graph Neural Networks for Learning Feasible and Adaptive AC-OPF Solutions. *IEEE Transactions on Power Systems*, 38 (6):5660–5670, November 2023. ISSN 1558-0679. doi: 10.1109/TPWRS.2022.3230555. URL https://ieeexplore.ieee.org/document/9992121.

[3] Olayiwola Arowolo and Jochen L. Cremer. Towards generalization of graph neural networks for AC optimal power flow. *Energy and AI*, 25:100842, September 2026. ISSN 2666-5468. doi: 10.1016/j.egyai.2026.100842. URL https://www.sciencedirect.com/science/article/pii/S2666546826001680.

[4] Tsong Yueh Chen, Fei-Ching Kuo, Huai Liu, Pak-Lok Poon, Dave Towey, T. H. Tse, and Zhi Quan Zhou. Metamorphic Testing: A Review of Challenges and Opportunities. *ACM Computing Surveys*, 51(1):4:1–4:27, January 2018. doi: 10.1145/3143561. URL https://dl.acm.org/doi/abs/10.1145/3143561.

[5] Sergio Segura, Gordon Fraser, Ana B. Sanchez, and Antonio Ruiz-Cortés. A Survey on Metamorphic Testing. *IEEE Transactions on Software Engineering*, 42(9):805–824, September 2016. ISSN 1939-3520. doi: 10.1109/TSE.2016.2532875. URL https://ieeexplore.ieee.org/abstract/document/7422146.

[6] Rahul Nellikkath and Spyros Chatzivasileiadis. Physics-Informed Neural Networks for AC Optimal Power Flow. *Electric Power Systems Research*, 212:108412, November 2022. ISSN 0378-7796. doi: 10.1016/j.epsr.2022.108412. URL https://www.sciencedirect.com/science/article/pii/S0378779622005636.

[7] Priya L. Donti, David Rolnick, and J. Zico Kolter. DC3: a learning method for optimization with hard constraints. In *International conference on learning representations*, 2021. URL https://arxiv.org/abs/2104.12225.

[8] Peter W. Battaglia, Jessica B. Hamrick, Victor Bapst, Alvaro Sanchez-Gonzalez, Vinicius Zambaldi, Mateusz Malinowski, Andrea Tacchetti, David Raposo, Adam Santoro, Ryan Faulkner, Caglar Gulcehre, Francis Song, Andrew Ballard, Justin Gilmer, George Dahl, Ashish Vaswani, Kelsey Allen, Charles Nash, Victoria Langston, Chris Dyer, Nicolas Heess, Daan Wierstra, Pushmeet Kohli, Matt Botvinick, Oriol Vinyals, Yujia Li, and Razvan Pascanu. Relational inductive biases, deep learning, and graph networks, October 2018. URL http://arxiv.org/abs/1806.01261. arXiv:1806.01261 [cs.LG].

[9] Haggai Maron, Heli Ben-Hamu, Nadav Shamir, and Yaron Lipman. Invariant and Equivariant Graph Networks. In *International Conference on Learning Representations*, 2019. URL https://openreview.net/forum?id=Syx72jC9tm.

[10] Damian Owerko, Fernando Gama, and Alejandro Ribeiro. Optimal Power Flow Using Graph Neural Networks. In *ICASSP 2020 - 2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 5930–5934, May 2020. doi: 10.1109/ICASSP40776.2020.9053140. URL https://ieeexplore.ieee.org/abstract/document/9053140. ISSN: 2379-190X.

[11] Salah Ghamizi, Aoxiang Ma, Jun Cao, and Pedro Rodriguez Cortes. OPF-HGNN: Generalizable Heterogeneous Graph Neural Networks for AC Optimal Power Flow. In *2024 IEEE Power & Energy Society General Meeting (PESGM)*, pages 1–5, July 2024. doi: 10.1109/PESGM51994.2024.10688560. URL https://ieeexplore.ieee.org/abstract/document/10688560. ISSN: 1944-9933.

[12] Ladislav Rampášek, Michael Galkin, Vijay Prakash Dwivedi, Anh Tuan Luu, Guy Wolf, and Dominique Beaini. Recipe for a General, Powerful, Scalable

Graph Transformer. In *Advances in Neural Information Processing Systems*, volume 35, pages 14501–14515. Curran Associates, Inc., 2022. doi: 10.52202/068431-1054. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/5d4834a159f1547b267a05a4e2b7cf5e-Abstract-Conference.html.

[13] Aihui Wen, Bao Wen, Jining Li, and Jin Xu. Heterogeneous Graph Neural Network with Local and Global Message Passing for AC-Optimal Power Flow Solutions. *Applied System Innovation*, 9(1):18, January 2026. ISSN 2571-5577. doi: 10.3390/asi9010018. URL https://www.mdpi.com/2571-5577/9/1/18.

[14] Helge Spieker, Nassim Belmecheri, Arnaud Gotlieb, and Nadjib Lazaar. Evaluating human trajectory prediction with metamorphic testing. In *Proceedings of the 9th ACM international workshop on metamorphic testing*, pages 34–40, 2024. doi: 10.1145/3679006.3685071.

[15] Andreas Venzke, Guannan Qu, Steven Low, and Spyros Chatzivasileiadis. Learning optimal power flow: Worst-case guarantees for neural networks. In *2020 IEEE International Conference on Communications, Control, and Computing Technologies for Smart Grids (SmartGridComm)*, pages 1–7. IEEE, November 2020. doi: 10.1109/SmartGridComm47815.2020.9302963. URL https://doi.org/10.1109/SmartGridComm47815.2020.9302963.

[16] Derek Lim, Joshua Robinson, Lingxiao Zhao, Tess Smidt, Suvrit Sra, Haggai Maron, and Stefanie Jegelka. Sign and basis invariant networks for spectral graph representation learning. In *International Conference on Learning Representations*, 2023. URL https://arxiv.org/abs/2202.13013.

[17] Meng Li, Xiaohua Yang, Jie Liu, and Shiyu Yan. Domain-validity-gated metamorphic testing of scientific ML surrogates, 2026. URL https://arxiv.org/abs/2606.17529.

[18] Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models' sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting. In *The twelfth international conference on learning representations*, 2024. URL https://openreview.net/forum?id=RIu5lyNXjT.

[19] Moran Mizrahi, Guy Kaplan, Dan Malkin, Rotem Dror, Dafna Shahaf, and Gabriel Stanovsky. State of what art? A call for multi-prompt LLM evaluation. *Transactions of the Association for Computational Linguistics*, 12:933–949, 2024. doi: 10.1162/tacl_a_00681.

[20] Oscar Sainz, Jon Ander Campos, Iker García-Ferrero, Julen Etxaniz, Oier Lopez de Lacalle, and Eneko Agirre. NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark. In *Findings of the association for computational linguistics: EMNLP 2023*, pages 10776–10787, 2023. doi: 10.18653/v1/2023.findings-emnlp.722.

[21] Anka Reuel, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J. Kochenderfer. BetterBench: Assessing AI benchmarks, uncovering issues, and establishing best practices. In *Advances in neural information processing systems 37*, pages 21763–21813, 2024. doi: 10.52202/079017-0685.

[22] Maria Eriksson, Erasmo Purificato, Arman Noroozian, João Vinagre, Guillaume Chaslot, Emilia Gómez, and David Fernandez-Llorca. Can we trust AI benchmarks? An interdisciplinary review of current issues in AI evaluation. In *Proceedings of the AAAI/ACM conference on AI, ethics, and society*, volume 8, pages 850–864, 2025. doi: 10.1609/aies.v8i1.36595.

[23] Weiwei Yang, Andrea Britto, Thiago Spina, Spencer Fowers, Baosen Zhang, and Chris M. White. GridSFM: A Foundation Model for AC Optimal Power Flow, May 2026. URL https://www.microsoft.com/en-us/research/wp-content/uploads/2026/05/GridFM_white_paper.pdf.

[24] Microsoft Corporation. GridSFM-Open model card, 2026. URL https://huggingface.co/microsoft/GridSFM_Open/tree/v1.1. Hugging Face, checkpoint gridsfm_open_v1.1.pt.

[25] Microsoft Corporation. GridSFM software repository, 2026. URL https://github.com/microsoft/GridSFM. GitHub, audited at commit 1ca775f.

[26] Yijiang Li, Zeeshan Memon, Hongwei Jin, Stefano Fenu, Keunju Song, Sunash B. Sharma, Parfait Gasana, Hongseok Kim, Liang Zhao, and Kibaek Kim. LUMINA: Foundation models for topology transferable ACOPF. In *ICLR 2026 workshop on foundation models for science: Real-world impact and science-first design*, 2026. doi: 10.48550/arXiv.2603.04300. URL https://doi.org/10.48550/arXiv.2603.04300.

[27] Argonne National Laboratory. LUMINA-2M model card, 2026. URL https://huggingface.co/argonne/LUMINA-2M/tree/v0.1.3. Hugging Face, version v0.1.3.

[28] Sean Lovett, Miha Zgubič, Sofia Liguori, Sephora Madjiheurem, Hamish Tomlinson, Sophie Elster, Chris Apps, Sims Witherspoon, and Luis Piloto. OPFData: Large-scale datasets for AC optimal power flow with topological perturbations, June 2024. URL https://arxiv.org/abs/2406.07234. arXiv:2406.07234.

[29] Sogol Babaeinejadsarookolaee, Adam Birchfield, Richard D. Christie, Carleton Coffrin, Christopher DeMarco, Ruisheng Diao, Michael Ferris, Stephane Fliscounakis, Scott Greene, Renke Huang, Cedric Josz, Roman Korab, Bernard Lesieutre, Jean Maeght, Terrence W. K. Mak, Daniel K. Molzahn, Thomas J. Overbye, Patrick Panciatici, Byungkwon Park, Jonathan Snodgrass, Ahmad Tbaileh, Pascal Van Hentenryck, and Ray Zimmerman. The Power Grid Library for Benchmarking AC Optimal Power Flow Algorithms, January 2021. URL http://arxiv.org/abs/1908.02788. arXiv:1908.02788 [math.OC].

[30] Carleton Coffrin, Russell Bent, Kaarthik Sundar, Yeesian Ng, and Miles Lubin. PowerModels.jl: An Open-Source Framework for Exploring Power Flow Formulations. In *2018 Power Systems Computation Conference (PSCC)*, pages 1–8. IEEE, June 2018. doi: 10.23919/PSCC.2018.8442948. URL https://doi.org/10.23919/pscc.2018.8442948.

[31] Vijay Prakash Dwivedi, Chaitanya K. Joshi, Anh Tuan Luu, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. Benchmarking Graph Neural Networks. *Journal of Machine Learning Research*, 24(43):1–48, 2023. URL http://jmlr.org/papers/v24/22-0567.html.

[32] Vijay Prakash Dwivedi, Anh Tuan Luu, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. Graph Neural Networks with Learnable Structural and Positional Representations. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=wTTjnvGphYj.

[33] Shuxiao Chen, Edgar Dobriban, and Jane H. Lee. A Group-Theoretic Framework for Data Augmentation. *Journal of Machine Learning Research*, 21(245):1–71, 2020. ISSN 1533-7928. URL http://jmlr.org/papers/v21/20-163.html.

[34] Bryn Elesedy and Sheheryar Zaidi. Provably Strict Generalisation Benefit for Equivariant Models. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139, pages 2959–2969. PMLR, July 2021. URL https://proceedings.mlr.press/v139/elesedy21a.html.

[35] Oskar Nordenfors, Fredrik Ohlsson, and Axel Flinth. Optimization Dynamics of Equivariant and Augmented Neural Networks. *Transactions on Machine Learning Research*, January 2025. ISSN 2835-8856. URL https://openreview.net/forum?id=PTTa3U29NR.

[36] Omri Puny, Matan Atzmon, Heli Ben-Hamu, Ishan Misra, Aditya Grover, Edward J. Smith, and Yaron Lipman. Frame averaging for invariant and equivariant network design. In *International conference on learning representations*, 2022. URL https://arxiv.org/abs/2110.03336.

[37] Sékou-Oumar Kaba, Arnab Kumar Mondal, Yan Zhang, Yoshua Bengio, and Siamak Ravanbakhsh. Equivariance with learned canonicalization functions. In *Proceedings of the 40th international conference on machine learning*, volume 202, pages 15546–15566, 2023. URL https://proceedings.mlr.press/v202/kaba23a.html.

[38] Yinan Huang, William Lu, Joshua Robinson, Yu Yang, Muhan Zhang, Stefanie Jegelka, and Pan Li. On the stability of expressive positional encodings for graphs. In *International conference on learning representations*, 2024. URL https://arxiv.org/abs/2310.02579.

[39] Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabás Póczos, Ruslan Salakhutdinov, and Alexander J. Smola. Deep sets. In *Advances in neural information processing systems 30*, 2017. URL https://arxiv.org/abs/1703.06114.

[40] Daniel Bienstock and Abhinav Verma. Strong NP-hardness of AC power flows feasibility. *Operations Research Letters*, 47(6):494–501, November 2019. ISSN 0167-6377. doi: 10.1016/j.orl.2019.08.009. URL https://doi.org/10.1016/j.orl.2019.08.009.

[41] Waqquas A. Bukhsh, Andreas Grothey, Ken I. M. McKinnon, and Paul A. Trodden. Local Solutions of the Optimal Power Flow Problem. *IEEE Transactions on Power Systems*, 28(4): 4780–4788, November 2013. ISSN 1558-0679. doi: 10.1109/TPWRS.2013.2274577. URL https://ieeexplore.ieee.org/abstract/document/6581918.

[42] Clare Lyle, Mark van der Wilk, Marta Kwiatkowska, Yarin Gal, and Benjamin Bloem-Reddy. On the benefits of invariance in neural networks. *arXiv preprint arXiv:2005.00178*, 2020. URL https://arxiv.org/abs/2005.00178.

[43] GridFM Project. gridfm-graphkit Release v0.8.1, 2026. URL https://github.com/gridfm/gridfm-graphkit/releases/tag/v0.8.1. GitHub release, tag v0.8.1, commit b3d663b.

[44] Andreas Wächter and Lorenz T. Biegler. On the implementation of an interior-point filter line-search algorithm for large-scale nonlinear programming. *Mathematical Programming*, 106 (1):25–57, March 2006. ISSN 0025-5610. doi: 10.1007/s10107-004-0559-y. URL https://doi.org/10.1007/s10107-004-0559-y.

A Relation catalog and output maps

Table 4 defines each transformation, states its validity conditions, and gives the output map that returns predictions to the original physical quantities. Figs. 2 and 3 illustrate the two tiers: `REP` relations change representation or angle reference, while `ELC` relations change graph or component bookkeeping and require aggregation or an inverse map.

Table 4: Catalog of the 13 metamorphic relations. For each relation, T_r constructs an equivalent input and P_r aligns or aggregates the transformed prediction with the original physical quantities; the validity condition states when that equivalence holds.

Relation	Tier	Transformation T_r	Validity condition	Output map P_r
component-relabel	`REP`	Reorder the generator array (or load, or shunt) by a permutation σ, moving its bus-link rows with it; buses and branches untouched.	always (every σ)	apply σ^{-1} to the per-component outputs
bus-relabel	`REP`	Reorder the bus array by a permutation π; rewrite every stored bus index through π (branch from/to endpoints and all component-to-bus links).	always (every π)	apply π^{-1} to the per-bus outputs
row-order	`REP`	Reorder the rows of one edge family (`ac_line`, transformer, or a component-link table); every field value and endpoint unchanged (a pure serialisation reshuffle).	always	identity on physical quantities; re-sort through `canonical_link_order` before positional conversion
branch-reversal	`REP`	Reverse a branch: swap its from/to endpoints, swap the shunt halves $b_{fr} \leftrightarrow b_{to}$, and replace $(\text{angmin}, \text{angmax})$ with $(-\text{angmax}, -\text{angmin})$.	per row: $\tau{=}1$, shift $=0$, $g_{fr}{=}g_{to}{=}0$ (all `ac_line` rows; only untapped, unshifted transformers)	reverse that branch's directed-flow channels, $(P_f, Q_f, P_t, Q_t) \to (P_t, Q_t, P_f, Q_f)$
ref-bus-move	`REP`	Move the angle reference: set a chosen generator bus's type to reference and demote the former reference to PV; no other field changes.	the chosen bus is a generator bus in the same island (one reference per island)	identity on components; re-anchor both solutions' angles to a common physical bus
basekv-scale	`REP`	Multiply the model-visible `base_kv` field at every bus by a common factor $\alpha > 0$ after OPFData export.	pipeline-specific: the standard per-unit AC-OPF target is unchanged	identity (solver quantities are unchanged)
tap-side-flip	`REP`	Swap a transformer's from/to endpoints. For complex tap $\tau{=}te^{\mathrm{j}\phi}$, set $\tau'{=}\tau^{-1}$, $z'{=}t^2 z$, $(b'_{fr}, b'_{to}){=}(b_{to}/t^2, b_{fr}/t^2)$, and $(\theta'_{\min}, \theta'_{\max}){=}(-\theta_{\max}, -\theta_{\min})$.	$t > 0$; $Y'{=}\Pi Y \Pi$ for the endpoint-swap matrix Π	reverse directed-flow channels, $(P_f, Q_f, P_t, Q_t) \to (P_t, Q_t, P_f, Q_f)$
island-ref-mix	`REP`	Multi-island only: permute which island serialises first and choose each island's reference independently (bus relabeling on island order composed with per-island ref-bus-move).	multi-island instance; one valid generator reference per island	re-anchor each island's angles, then undo the island/bus permutation
component-split	`ELC`	Split load ℓ into co-located loads ℓ_1, ℓ_2 satisfying $(P^d_{\ell_1}{+}P^d_{\ell_2}, Q^d_{\ell_1}{+}Q^d_{\ell_2}){=}(P^d_\ell, Q^d_\ell)$; or split generator g into two with halved $(P^{\min}_g, P^{\max}_g, Q^{\min}_g, Q^{\max}_g)$ bounds and cost $(c_{g,2}, c_{g,1}, c_{g,0}) \to (2c_{g,2}, c_{g,1}, c_{g,0}/2)$ each.	loads: always; generators: convex quadratic cost and symmetric halved bounds	*sum* the two parts' injections back (aggregation, not a permutation)
parallel-split	`ELC`	Replace one branch by two identical parallel branches: $r, x \to 2r, 2x$ (parallel admittance preserved), $b_{fr}, b_{to} \to b_{fr}/2, b_{to}/2$, rate $\to$ rate$/2$; angle limits kept.	always (admittances recombine; flow splits evenly)	sum the two branches' flows
island-union	`ELC`	Concatenate two solved scenarios into one instance with two disconnected islands, renumbering the second scenario's buses and components after the first.	always (no coupling; optimum and cost are the direct sum)	split the joint prediction into the two island blocks
load-as-gen	`ELC`	Recast a fixed load ℓ with demand (P^d_ℓ, Q^d_ℓ) as a zero-cost generator g' at the same bus with $P^{\min}_{g'}{=}P^{\max}_{g'}{=}{-}P^d_\ell$ and $Q^{\min}_{g'}{=}Q^{\max}_{g'}{=}{-}Q^d_\ell$, then remove ℓ from the load table.	always (a clamped zero-cost generator imposes the identical injection)	read $(P_{g'}, Q_{g'})$ and map it back to the load (must equal $(-P^d_\ell, -Q^d_\ell)$)

Continued on next page

Table 4 (continued)

Relation	Tier	Transformation T_r	Validity condition	Output map P_r
null-pad	ELC	Add an inert element: a load ℓ' with $P^d_{\ell'}{=}Q^d_{\ell'}{=}0$, a shunt with $G^{\mathrm{sh}}{=}B^{\mathrm{sh}}{=}0$, or a zero-cost generator g' with $P^{\min}_{g'}{=}P^{\max}_{g'}{=}Q^{\min}_{g'}{=}Q^{\max}_{g'}{=}0$; bus types are unchanged.	always (the element injects nothing; only counts change)	drop the padded element from the prediction

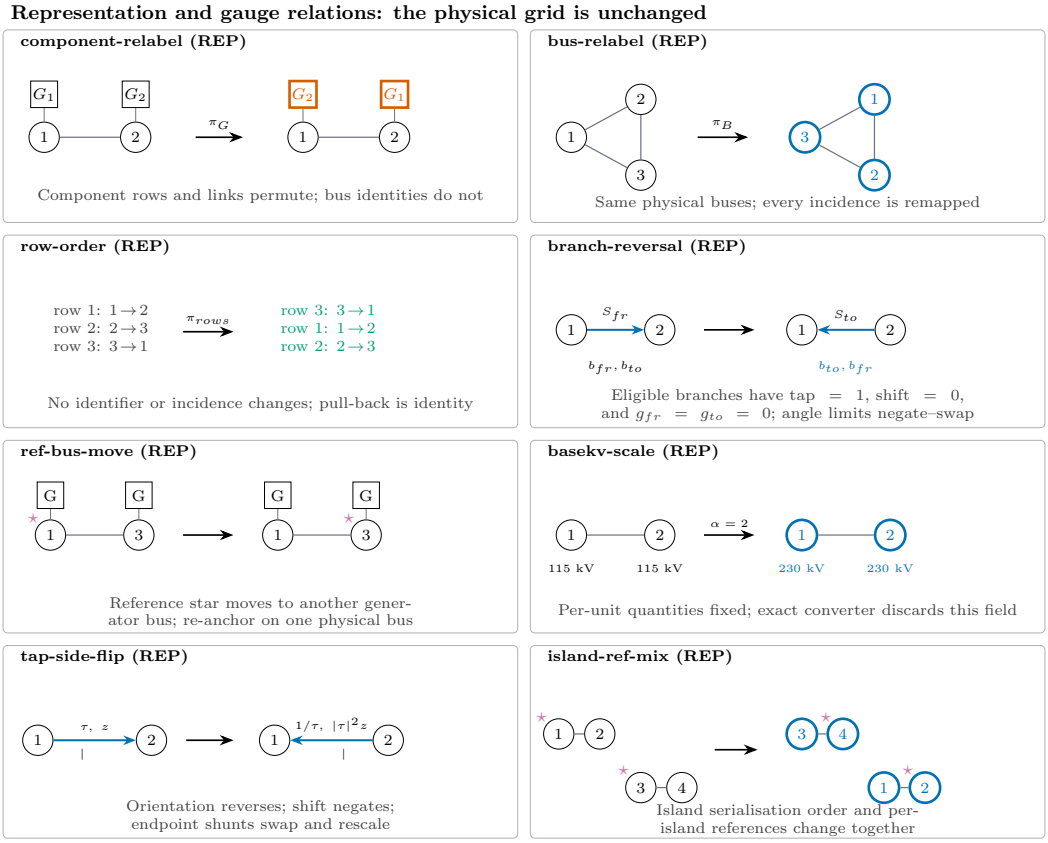

Figure 2: Examples of REP relations. They change labels, row order, branch orientation, or angle reference without changing the underlying AC-OPF problem, and we compare predictions after undoing the declared map. The pipeline-specific `basekv-scale` relation changes an exported model input that is absent from the standard per-unit AC-OPF target.

B Technical foundation

B.1 AC optimal power flow and non-uniqueness

Let $\mathcal{N}$, $\mathcal{E}$, and $\mathcal{G}$ denote buses, branches, and generators. For bus i, the AC-OPF variables include voltage magnitude V_i, voltage angle θ_i, and active and reactive generation (P_g, Q_g); fixed demands are (P^d_i, Q^d_i). With commitment fixed, a common economic-dispatch objective is

$$\min_{V,\theta,P,Q} \quad \sum_{g\in\mathcal{G}} c_{g,2}P_g^2 + c_{g,1}P_g + c_{g,0}. \tag{2}$$

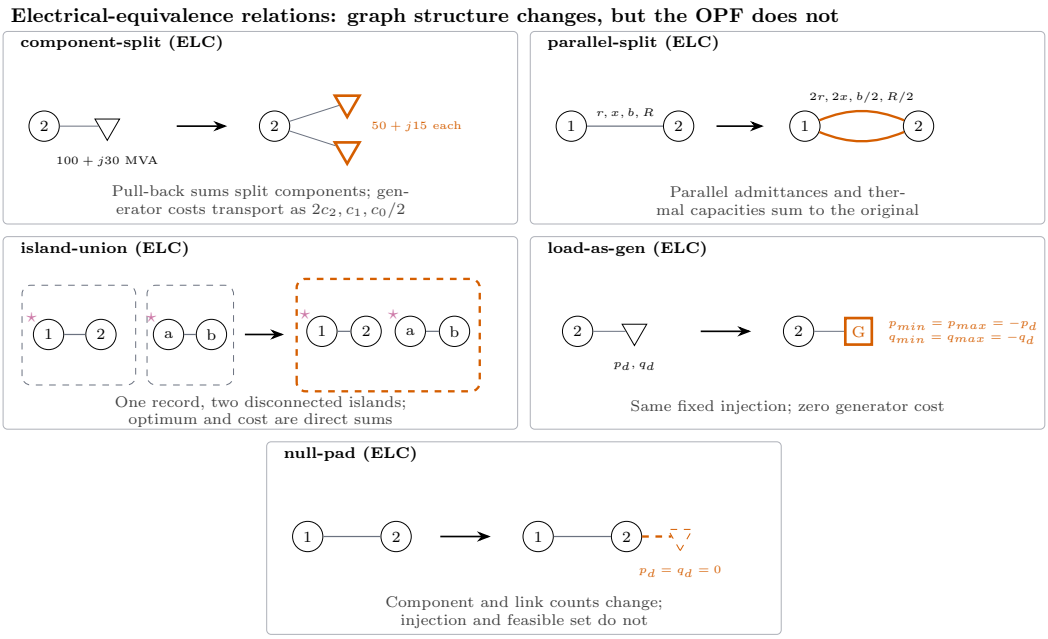

Figure 3: Examples of ELC relations. They change graph size or component-family bookkeeping while preserving the AC-OPF feasible set and objective under the stated validity conditions. Their output maps sum split quantities, separate joined islands, reassign a fixed injection, or remove inert padding, depending on the relation.

Writing $Y = G + \mathrm{j}B$ and $\theta_{ij} = \theta_i - \theta_j$, active and reactive nodal balance require

$$\sum_{g \in \mathcal{G}_i} P_g - P_i^d = V_i \sum_{j \in \mathcal{N}} V_j \left(G_{ij} \cos\theta_{ij} + B_{ij} \sin\theta_{ij} \right), \tag{3}$$

$$\sum_{g \in \mathcal{G}_i} Q_g - Q_i^d = V_i \sum_{j \in \mathcal{N}} V_j \left(G_{ij} \sin\theta_{ij} - B_{ij} \cos\theta_{ij} \right). \tag{4}$$

Generator outputs and voltage magnitudes obey bounds; branch flows obey apparent-power and angle-difference limits. One reference angle per connected island removes the additive angle gauge. Voltage products and trigonometric angle differences make the feasible set nonconvex. AC feasibility is strongly NP-hard in general, and multiple local optima can occur [40, 41]. Production cost may also leave reactive dispatch and voltage magnitude weakly identified. Our test checks whether one deterministic surrogate is consistent across equivalent inputs and does not assume a unique AC-OPF optimum.

B.2 Fibers of the data-to-problem map

Serialized cases include choices that the physical optimization does not depend on, such as bus names, array order, stored branch direction, and angle reference. The problem itself depends on aggregate bus injections, series and shunt admittances, transformer ratios, incidence, variable bounds, and cost coefficients. Let $s : \mathcal{D} \to \mathcal{Z}$ map serialized data to these sufficient quantities. A *fiber* of s contains the encodings with the same sufficient quantities. Valid relations stay within such a fiber, allowing for label permutations or changes of angle reference. After mapping outputs back, a representation-consistent surrogate satisfies $P_r f(T_r d) = f(d)$, or $\Delta_r(d) = 0$ in (1).

Each row of Table 5 identifies a serialization choice that leaves the optimization problem unchanged. The catalog tests one concrete transformation within each listed fiber; Sec. B.3 states the limits of this coverage.

Table 5: Schema distinctions discarded by the data-to-problem map s, with the relations that test each resulting many-to-one collapse.

Collapse in s	Relations
Incidence is label-free	`component-relabel`, `bus-relabel`, `row-order`, `island-ref-mix`
Only the aggregate branch model enters the equations	`branch-reversal`, `tap-side-flip`
Bus balance sees aggregate injection instead of per-device	`component-split`, `load-as-gen`, `null-pad`
Parallel branches collapse to one effective admittance	`parallel-split`
Angle reference is a gauge convention	`ref-bus-move`, `island-ref-mix`
The post-export per-unit target does not use model-visible `base_kv`	`basekv-scale`
Disconnected islands cannot interact	`island-union`

B.3 Fiber coverage and extensions

A sufficiently small relaxation of a strictly inactive bound could extend the catalog: under local solution stability, it preserves the same local solution, although it need not preserve the global OPF problem. Unequal or N-way component splits and arbitrary shunt redistribution would extend the discrete tests to continuous families.

Coverage remains incomplete even for the listed fibers. Most are continuous, so one discrete transformation tests only part of each. For example, `branch-reversal` swaps endpoints without testing arbitrary redistribution of $b_{fr} + b_{to}$, and `component-split` uses equal halves rather than an arbitrary N-way split. We have not shown that compositions of the 13 relations generate every fiber. The catalog is therefore representative but not exhaustive.

B.4 When augmentation theory applies

For squared loss, orbit averaging projects a predictor onto the equivariant function class when the transformations form a compact group, the target is equivariant, and the input distribution is invariant or orbit-balanced. Under these assumptions, the non-equivariant component contributes avoidable orbit-balanced risk [33, 34, 42]. Sampling transformations from the normalized Haar measure gives an unbiased stochastic estimate of this orbit-averaged objective.

Only part of the catalog satisfies these assumptions. Finite relabelings and eligible involutions fit the compact-group setting. By contrast, `basekv-scale` is unbounded and pipeline-specific, reference changes are conditional gauge actions, and `ELC` relations can change graph and output dimensions. The model-specific 0.5 mixture over eligible relations therefore does not sample a single group's Haar measure. The theorem concerns orbit-balanced risk, not error on one canonical serialization, and nonconvex training need not reach the projection [35]. We report canonical accuracy, transformed accuracy, metamorphic discrepancy, and feasibility separately.

C Audit design and model inventory

Table 6 lists every architecture and mechanism in the built-in consistency screen described in Sec. 3. The reimplementations test five design choices: output partitions based on array index versus slack/PV/PQ role; raw Laplacian coordinates; summaries over index-selected landmarks versus all buses; learned per-index versus graph-fixed filters; and directed versus orientation-symmetric branch processing. Where practical, matched pairs use the same surrounding network and random seed and change only the component of interest. They isolate that component's effect but do not reproduce a published model's training or accuracy.

We also include two MPNN rows that connect the architecture screen to the augmentation study. They are randomly initialized instances of the study architectures, each with four message-passing layers, width 128, sum aggregation, orientation-invariant branch features, and masked solution inputs; the second row adds 16 raw Laplacian eigenvectors. The three plain MPNNs provide smaller family controls for mean versus sum aggregation and for branch orientation.

A random-weight failure shows that the tested pipeline does not guarantee invariance to that relation; it does not establish trained accuracy, feasibility, or discrepancy magnitude. The DC3 rows test only the partition-and-completion choice with a stand-in. They do not run DC3's Newton power-flow completion or correction loop.

Table 6: Inventory of the 17 model, reimplementation, and control rows in the built-in consistency screen. "Scope" names the released pipeline or design mechanism behind each row.

Row	Class	Scope
GridSFM-Open v1.1	Open weight	Released 15,148,227-parameter GridSFM graph-transformer checkpoint and inference pipeline [23–25].
LUMINA-2M v0.1.3	Open weight	Released 2,375,548-parameter, six-layer heterogeneous model and inference pipeline [26, 27].
GridFM-GNS v0.8.1	Public architecture	Released heterogeneous preprocessing, model, and decoder at random initialization; we found no compatible public AC-OPF checkpoint as of August 25, 2026 [43].
DC3 (index partition)	Design reimplementation	Random index masks for partial V_m, V_a, P_g outputs followed by covariant mean completion [7].
DC3 (physical partition)	Matched control	Same completion stand-in with masks keyed by slack/PV/PQ roles rather than array position.
Graph transformer (LapPE-4)	Design reimplementation	Global bus self-attention with four raw, uncanonicalized Laplacian eigenvectors [12].
Landmark PE (K=64)	Mechanism control	Four shortest-path-distance moments to the first 64 bus indices, appended to an MPNN.
Landmark PE (K=all)	Matched control	Same moments over all buses, removing the index-selected subset.
Effective resistance (K=64)	Design reimplementation	Five unit-weight effective-resistance moments over the first 64 bus indices, appended to an MPNN [23, 25].
Effective resistance (K=all)	Matched control	Same moments over all buses [3]. Both effective-resistance rows use unit edge weights rather than the susceptance weighting of the cited designs, isolating index selection from electrical weighting.
Index filter (fixed)	Matched control	Three bilinear graph-filter layers using degree-normalized adjacency directly.
Index filter (learned)	Design reimplementation	Same layers with learned per-grid $N \times N$ values masked to graph support [2].
MPNN (matched)	Matched control	Four-layer, width-128 heterogeneous MPNN with sum aggregation, symmetric branch features, and masked solution inputs.
MPNN + LapPE (matched)	Matched control	Same MPNN with 16 raw, uncanonicalized Laplacian eigenvectors.
Plain MPNN (mean)	Family control	Two-layer heterogeneous MPNN with shared encoders and heads, mean aggregation, and symmetric branch encoding; this control is inspired by OPF-HGNN but does not reproduce its exact aggregation or branch encoding [11].
Plain MPNN (sum)	Family control	Same two-layer MPNN with sum aggregation.
Plain MPNN (orientation naive)	Family control	Same sum MPNN with raw endpoint fields and only the stored branch direction.

C.1 Data, sampling, and solver admission

Scenarios come from the N-1 topology-perturbed OPFData release; both released checkpoints were trained on this type of perturbation [24, 27, 28]. LUMINA metadata includes all five audited grids. GridSFM documents a training-size range of 500 to 4,661 buses, placing case14/118 outside that range and the larger grids inside it. These scenarios therefore do not provide a common unseen-topology benchmark for the two checkpoints.

The architecture screen uses case14/500, three source scenarios per grid, three draws per applicable case–relation pair, and three initializations where applicable. Fourteen `component-split` cells retain only case500 evidence because the case14 control was unavailable. Source scenarios are the independent evaluation units; transformed views of one source are repeated measurements.

Five-grid comparison sampling. The separate trained-versus-random audit uses three source scenarios from each of case14, case118, case500, case2000, and case4661, with three sampled transformations per source–relation pair. The same source scenarios and transformation draws are used for each backbone's released checkpoint and its three random initializations (seeds 0, 1, and 2). Each grid and model/seed run therefore contains $3 \times 13 \times 3 = 117$ executed source–relation–draw comparisons before aggregation. The trained-checkpoint runs stored under three seed labels reuse the same checkpoint and draws; these repetitions do not increase the independent sample count. The independent units remain the three source scenarios per grid.

What the solver controls establish. PowerModels.jl and IPOPT check the numerical implementation of the relations [30, 44]. The algebraic conditions still determine validity, and a converged solve need not be globally optimal. The archived expanded-relation configuration uses the polar AC formulation

(ACPPowerModel), a flat start in each view, an IPOPT tolerance of 10^{-8}, and a 300-second limit. The wrapper accepts its recorded converged statuses: LOCALLY_SOLVED, OPTIMAL, and their ALMOST variants. These labels do not independently verify residual bounds or global optimality. The saved control summaries do not pin the linear solver.

Expanded-relation admission. The historical control implementation admits a grid–relation pair only if all its sampled solves converge, the worst relative objective difference is at most 10^{-9}, and the worst aligned primal difference meets its relation-specific limit. Relative objective error is $|C' - C_{\mathrm{ref}}| / \max(|C_{\mathrm{ref}}|, 10^{-12})$; disconnected-union controls use the sum of component reference costs. The primal limits are 0 for basekv-scale, 10^{-9} for ref-bus-move/load-as-gen, 10^{-7} for null-pad, 10^{-6} for component-split/parallel-split, 2.2×10^{-6} for tap-side-flip, and 10^{-5} for island-union/island-ref-mix.

The archived case14/500 table contains 18 grid–relation pairs with nine control records each. All 162 pairs of solves converge, and 17 grid–relation pairs pass admission. Case14 component-split is unavailable because its objective difference, 7.26×10^{-9}, exceeds 10^{-9}, even though IPOPT converged. The model screen reuses these controls and does not run a fresh solve for every model draw. Although the configuration names 10^{-6} equality/inequality targets, the expanded control CSV does not store independently audited residuals, so we do not claim that those targets were separately verified for every admitted view.

Original relabeling and orientation controls. For the four original relations, we check transported labels and residual preservation over 300 records per grid–relation on case14/118/500 and 90 on case2000/4661. We also check bus relabeling against an independent converter. Transported labels match, and residual changes meet 10^{-12} in all 20 grid–relation summaries. A separate re-solve comparison uses 10^{-6} cost and primal tolerances. The case4661 row-order and branch-reversal controls converge but exceed the primal threshold. We retain these outcomes as sensitivity of the solver solution, rather than reclassifying them as surrogate failures. This distinction matters when solutions are non-unique or weakly identified.

Verdict versus descriptive magnitude. The screen takes the largest channel-mean discrepancy over all executed draws and compares it with 10^{-6}; the separate maximum-component check uses 10^{-5}. The five-grid comparison first averages draws within each source, then takes the median across sources and the maximum across grids, retaining a value for each of three random seeds. Its checkmarks mean that this aggregate is below threshold, even if some individual draws exceed it. Neither finite screen passes nor below-threshold aggregates prove universal invariance.

C.2 Threshold sensitivity and component tails

We reaggregate the 221 saved model–relation summaries from the 17-row screen, changing only the reporting threshold and keeping the executed samples and applicability mask fixed. No model or solver is rerun. Every configuration still fails at least six relations at 10^{-7}, 10^{-6}, and 10^{-5} in native channel units. At the primary 10^{-6} channel-mean threshold, all 64 passing cells also have saved maximum-component discrepancies at or below 10^{-5}. The component check therefore reveals no additional failure among the screen passes. This result does not establish an all-draw pass for the separate five-grid aggregate comparison.

Table 7: Sensitivity of the saved architecture-screen verdicts to the channel-mean reporting threshold. Each row uses the same 221 cells and partial-coverage mask.

Threshold	Failed cells / 221	Failed relations per model	Models with a failure
10^{-7}	160	6–13	17/17
10^{-6}	157	6–13	17/17
10^{-5}	154	6–12	17/17
10^{-4}	144	4–12	17/17
10^{-3}	134	4–12	17/17

Excluding basekv-scale leaves 204 model–relation cells. At the primary 10^{-6} threshold, 140 of these cells fail, and every configuration fails between five and twelve of the twelve remaining relations.

At 10^{-4} and 10^{-3}, respectively, 127 and 118 cells fail; every configuration still fails between three and eleven of the twelve remaining relations.

The tolerance is a prespecified reporting threshold, not an empirically calibrated noise floor for every pipeline. The saved results lack a comprehensive set of repeated-forward, round-trip, CPU/GPU, and float32/float64 comparisons. Near-threshold cells therefore cannot identify a particular numerical cause. The table shows that the broad failure-count conclusion survives changes to the reporting threshold; it does not establish robustness to changes in arithmetic precision.

D Detailed audit evidence

D.1 Trained versus random-weight consistency

Table 8 uses the worst-grid summary defined in Sec. C.1. Each trained entry uses the released checkpoint; the random entry gives the range across three seed-specific aggregates. Checkmarks classify the aggregate, not every draw.

Table 8: Trained and random-weight worst-grid discrepancies by relation. Entries are unnormalized worst-channel mean absolute discrepancies; random entries span three weight seeds. A checkmark marks an aggregate value at or below the prespecified 10^{-6} channel tolerance; it does not mean that every draw passes. Because values inherit the worst channel's unit (p.u. for V_m, P_g, Q_g or radians for V_a), they are comparable with the tolerance but not across channels.

	GridSFM-Open v1.1		LUMINA-2M v0.1.3	
Relation	Trained	Random seed range	Trained	Random seed range
`component-relabel`	3.80×10^{-7} ✓	$[3.00 \times 10^{-8}, 3.53 \times 10^{-8}]$ ✓	2.18×10^{-8} ✓	$[4.07 \times 10^{-9}, 1.03 \times 10^{-8}]$ ✓
`bus-relabel`	0.279	[0.00957, 0.0143]	5.01×10^{-9} ✓	$[5.60 \times 10^{-10}, 3.37 \times 10^{-9}]$ ✓
`row-order`	0.780	[0.574, 0.578]	4.45×10^{-8} ✓	$[5.55 \times 10^{-9}, 1.17 \times 10^{-8}]$ ✓
`branch-reversal`	0.326	[0.0133, 0.0365]	0.299	[0.00528, 0.00903]
`ref-bus-move`	0.438	[0.00755, 0.0113]	0.182	$[2.40 \times 10^{-4}, 0.00789]$
`component-split`	0.00940	$[3.45 \times 10^{-4}, 0.00232]$	0.00325	$[2.45 \times 10^{-5}, 7.04 \times 10^{-5}]$
`parallel-split`	0.0253	[0.00171, 0.00206]	1.15×10^{-4}	$[2.86 \times 10^{-5}, 6.86 \times 10^{-5}]$
`island-union`	0.680	[0.0515, 0.160]	3.30×10^{-8} ✓	$[7.51 \times 10^{-9}, 2.23 \times 10^{-8}]$ ✓
`basekv-scale`	0.397	[0.0168, 0.0232]	0.513	[0.0547, 0.0782]
`tap-side-flip`	0.104	[0.00569, 0.0141]	0.0333	$[1.47 \times 10^{-4}, 2.02 \times 10^{-4}]$
`load-as-gen`	0.0322	[0.00782, 0.0175]	0.00157	$[3.12 \times 10^{-5}, 2.06 \times 10^{-4}]$
`null-pad`	0.104	[0.00248, 0.00767]	0.0231	$[8.96 \times 10^{-4}, 0.00226]$
`island-ref-mix`	0.675	[0.244, 0.268]	0.124	$[2.40 \times 10^{-4}, 0.00789]$

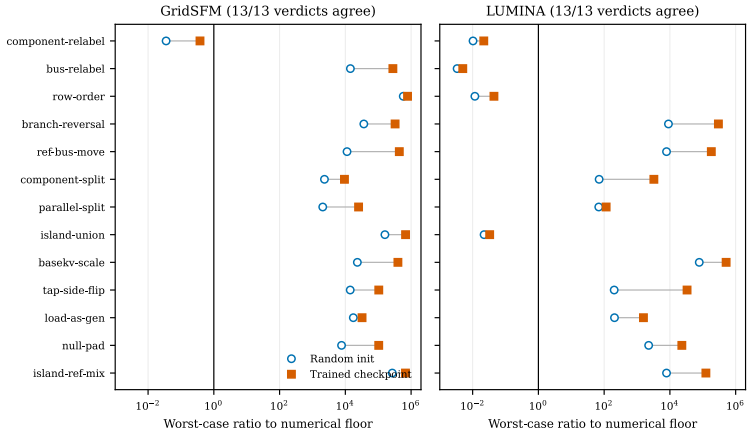

Figure 4: Trained and random-weight models agree on all 13 threshold classifications of the displayed aggregate statistic. For each relation, the horizontal coordinate is the worst-channel discrepancy, maximized over five grids, divided by the 10^{-6} tolerance; random-weight markers also take the worst of three seeds. The vertical line at one marks the aggregate reporting threshold and does not serve as an all-draw pass criterion. Connecting segments show the change in magnitude after training.

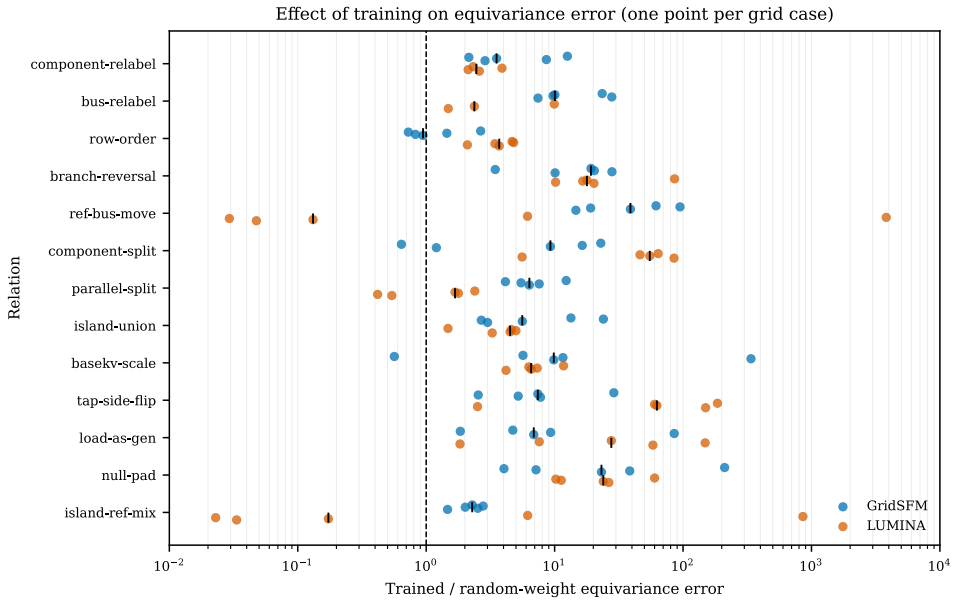

Figure 5: Per-grid effect of training on metamorphic discrepancy. Each point is the trained discrepancy divided by its random-weight control for one grid and relation; values above one indicate a larger mismatch after training. Short black bars mark relation-level medians. Ratios for `component-relabel` compare values that are already below the 10^{-6} tolerance, so they carry little information.

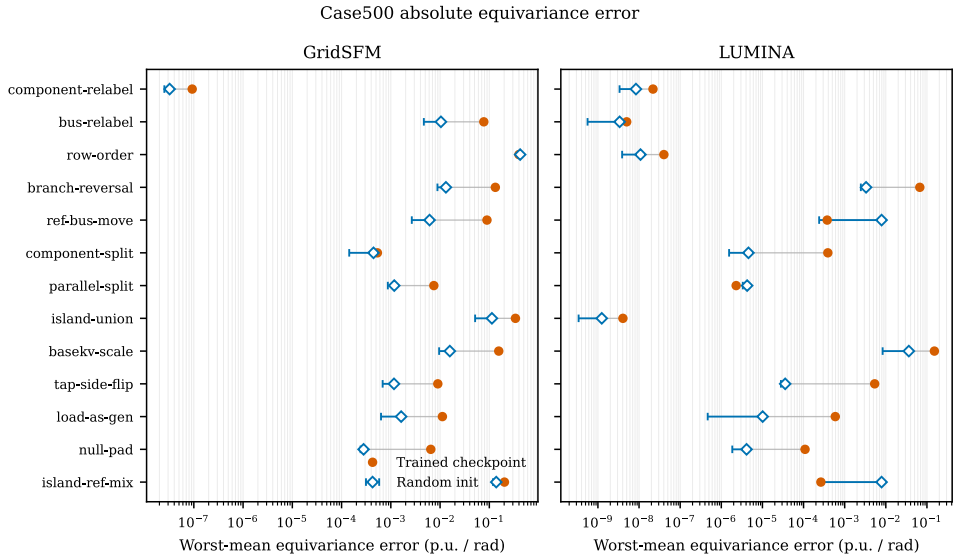

Figure 6: Absolute case500 discrepancies underlying the ratio comparison. Circles show released checkpoints; diamonds and intervals show the central value and range across three random seeds. The plot distinguishes a large relative change near zero from a discrepancy far above 10^{-6}. Each point uses its worst channel, so values should be read against the tolerance and not compared across channel units.

Table 8 and Figs. 4 to 6 show the same aggregate threshold classifications for trained and random weights. GridSFM is below threshold only on `component-relabel`; LUMINA is also below it on `bus-relabel`, `row-order`, and `island-union`. Excluding `component-relabel`, whose trained and random discrepancies are below tolerance, training increases discrepancy in 55 of 60 finite GridSFM case–relation cells and 50 of 58 finite LUMINA cells. The median ratios are 7.50 and 6.23, respectively. Ratios must be read alongside absolute discrepancies: a near-zero random denominator can produce a large ratio, and a ratio below one can still leave a discrepancy above tolerance. Random networks may also be nearly constant. These comparisons measure representation sensitivity, not predictive quality; all trained and random arms fail the AC-feasibility criterion.

D.2 Auxiliary dispatch-cost sensitivity

Let $\widehat{P}_g(d) = [f(d)]_{P_g}$ and $\widehat{P}_{g,r}(d) = [P_r f(T_r d)]_{P_g}$ be the canonical and aligned transformed active-power predictions. With the cost C in Eq. (2), define

$$S_r(d) = \frac{\left|C\left(\widehat{P}_{g,r}(d)\right) - C\left(\widehat{P}_g(d)\right)\right|}{\max\left(\left|C\left(P_g^\star(d)\right)\right|, \varepsilon\right)}, \tag{5}$$

where $P_g^\star(d)$ is the solver dispatch and ε is machine epsilon. Table 9 reports the median numerator in cost units and the median $S_r(d)$. For example, $S_r(d) = 0.513$ means that an encoding change alters the implied cost by 51.3% of the solver's reference cost. The predictions are AC infeasible, so $S_r(d)$ measures representation sensitivity rather than a realized dispatch loss or optimality gap. Random-weight values likewise do not measure economic quality.

Table 9: Median absolute difference between the costs implied by canonical and aligned transformed GridSFM predictions. Absolute values are in cost units; relative values divide that difference by the solver's reference dispatch cost. Because the predictions are infeasible, these values measure sensitivity and are not optimality gaps.

Relation	Trained weights		Random weights	
	Absolute	Relative	Absolute	Relative
`component-relabel`	7.8×10^{-3}	1.57×10^{-8}	2.15×10^{-3}	8.18×10^{-9}
`bus-relabel`	1.25×10^{4}	0.022	1.65×10^{3}	3.94×10^{-3}
`row-order`	6.99×10^{4}	0.234	1.34×10^{3}	4.86×10^{-3}
`branch-reversal`	1.98×10^{4}	0.0404	1.15×10^{3}	4.23×10^{-3}
`ref-bus-move`	2.60×10^{4}	0.0737	281	1.28×10^{-3}
`component-split`	57.5	1.40×10^{-4}	2.04	4.92×10^{-5}
`parallel-split`	180	3.76×10^{-4}	45.2	3.65×10^{-4}
`island-union`	4.63×10^{4}	0.502	1.98×10^{3}	5.27×10^{-3}
`basekv-scale`	3.21×10^{4}	0.0350	1.03×10^{3}	1.81×10^{-3}
`tap-side-flip`	366	7.69×10^{-4}	105	7.16×10^{-4}
`load-as-gen`	2.03×10^{3}	2.46×10^{-3}	66.8	3.32×10^{-4}
`null-pad`	401	1.08×10^{-3}	9.43	9.71×10^{-5}
`island-ref-mix`	9.40×10^{4}	0.513	4.22×10^{3}	0.0120

E Augmentation protocol and extended results

E.1 Training and evaluation protocol

For each of four architectures, we train a baseline on canonical inputs and a second model with transformed inputs. The architectures are a four-layer, width-128 heterogeneous MPNN; that MPNN with 16 raw Laplacian eigenvectors (LapPE+MPNN); that MPNN with a 16-step random-walk structural encoding (RWSE+MPNN); and LUMINA. RWSE uses per-node return probabilities, avoiding eigenvector sign and basis ambiguities [12, 31, 32]. Each architecture uses three seeds (42, 31415, 271828), 2,000 case500 training instances, 89 epochs, 22,250 optimizer steps, batch size 8, and weight decay 10^{-4}. Learning rates are 10^{-3} for the MPNN variants and 3×10^{-4} for LUMINA. Each instance in the augmented arm is transformed by one eligible relation with probability 0.5.

All arms minimize the same equal-channel MSE plus AC constraint penalties,

$$\mathcal{L} = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \mathrm{MSE}_c + \lambda_{\mathrm{eq}} \log\Big(1 + \frac{1}{B} \sum_b (r_{P,b}^2 + r_{Q,b}^2)\Big) + \lambda_{\mathrm{ineq}} \frac{1}{K} \sum_k \max(0, g_k)^2. \tag{6}$$

Here $r_{P,b}, r_{Q,b}$ are the real and reactive parts of $S^{\mathrm{inj}} - V \odot \overline{Y_{\mathrm{bus}} V}$; g_k includes variable bounds, branch angle differences, and thermal limits at both ends. The frozen configuration sets $\lambda_{\mathrm{eq}} = 3.606662 \times 10^{-4}$ and $\lambda_{\mathrm{ineq}} = 5.929293 \times 10^{-3}$ for all 24 arms. These weights were calibrated offline using the median final validation-MSE-to-penalty ratios from nine earlier baseline arms on a 500-scenario validation split. The recovered record reports all nine arms usable and reproduces both weights after rounding. No one-epoch tuning procedure was involved. We select checkpoints using the validation objective. Within this experiment, the objective is fixed and augmentation changes; a comparison with the earlier pure-MSE experiments would change both the loss and the relation sets.

Relation selection. We fixed the training relation sets using discrepancies measured in earlier trained, unaugmented baseline arms and exclusions specific to each architecture. These sets did not come from the random-initialization screen. The seven MPNN and RWSE relations are `ref-bus-move`, `component-split`, `parallel-split`, `tap-side-flip`, `load-as-gen`, `null-pad`, and `island-ref-mix`. LapPE+MPNN adds `bus-relabel` and `island-union`. LUMINA uses the common seven plus `basekv-scale` and `branch-reversal`. The configuration records baseline discrepancies below 10^{-6} for `basekv-scale` in MPNN/LapPE and `island-union` in MPNN. Case500 features saturate, so these small discrepancies do not imply invariance for all inputs. Dashes in Tables 10 and 11 mark relations excluded from training; evaluation still covers all thirteen relations.

Evaluation cohorts. The main evaluation uses 50 held-out case500 source scenarios per seed. Each has 53 views: one canonical input and four draws for each of thirteen relations. The 2,650 views are repeated measurements of 50 independent sources. A separate batched evaluation uses 200 canonical scenarios from each of case14, case118, case500, case2000, and case4661, without further training. None of its case500 scenario IDs occurs in the main cohort. Sec. E.4 explains how the different cohorts and evaluation paths limit the comparison.

Accuracy averages the native-unit MAEs of V_m, V_a, P_g, Q_g with equal weights. This convention combines per unit and radians; Table 10 reports the per-channel values needed to interpret it. For the main tables, improvement is the mean of three paired seed-level percentages, $100(e_{\mathrm{base},s} - e_{\mathrm{aug},s})/e_{\mathrm{base},s}$, so positive values mean lower error. Table 13 uses the opposite sign: negative values mean lower error. The spreads in Sec. E.2 are population standard deviations across the three percentage changes, not confidence intervals.

To summarize metamorphic consistency, we average sampled draws within each source, take the median across sources, then the median across the architecture's trained relations, and finally the mean over seeds. Per-relation entries instead use the seed with the largest source median. Models train on different relation sets, so these aggregate medians support baseline-versus-augmentation comparisons within a model but cannot rank architectures. Older evaluation artifacts used a 10^{-3} reporting scale. We recompute discrepancies from native-unit errors using the paper's 10^{-6} scale.

We assess AC feasibility separately. For each prediction, we first take the largest violation within each constraint family. Within a seed, we then take the worst of each source's 53 views and compute the median, linearly interpolated P95, and maximum over the 50 sources. Tables report the mean and sample standard deviation of the resulting three seed-level statistics. Equality violation is the larger

of the maximum nodal active and reactive balance violations, in p.u. Pooled inequality violation is the maximum over constraint families after dividing each family's violation by its 10^{-6} native-unit tolerance; it is dimensionless. A source is feasible only if every view meets the equality and all inequality thresholds. No evaluated prediction meets these thresholds.

Table 10: Full augmentation results under the AC constraint-penalty objective. Positive improvement (Imp.) denotes the mean seed-paired reduction under augmentation. Metamorphic values report the dimensionless normalized discrepancy $\Delta_r(d)$ defined in (1); 0 denotes exact invariance and dashes mark excluded relations. Constraint rows use the scenario-clustered seed-level P95 defined in the text and report mean ± sample standard deviation across three seeds. Equality is in per unit. The pooled inequality metric is dimensionless and displayed in millions of family-specific 10^{-6} tolerances. Absolute family-level results appear in Table 12.

Outcome	Presentation	Metric	MPNN			LapPE+MPNN			RWSE+MPNN			LUMINA		
			Base	Aug.	Imp. (%)	Base	Aug.	Imp. (%)	Base	Aug.	Imp. (%)	Base	Aug.	Imp. (%)
Mean absolute error (MAE)	Canonical	Mean	0.01295	0.01345	−3.9	0.09983	0.01309	86.8	0.01255	0.01250	0.4	0.01834	0.01889	−3.0
		V_m	0.00425	0.00471	−12.7	0.01453	0.00447	69.1	0.00407	0.00424	−4.2	0.00674	0.00758	−12.4
		V_a	0.01289	0.01290	−0.1	0.12493	0.01320	89.2	0.01274	0.01263	0.9	0.03500	0.03606	−3.0
		P_g	0.01198	0.01268	−5.9	0.15529	0.01176	92.3	0.01103	0.01088	1.4	0.00952	0.00961	−0.9
		Q_g	0.02270	0.02352	−3.6	0.10459	0.02292	77.9	0.02236	0.02226	0.4	0.02211	0.02233	−1.0
	Transformed	Mean	0.01312	0.01353	−3.2	0.10211	0.01318	87.0	0.01271	0.01257	1.1	0.03594	0.02089	41.9
		V_m	0.00427	0.00472	−12.5	0.01468	0.00447	69.5	0.00409	0.00425	−4.0	0.00845	0.00826	2.2
		V_a	0.01308	0.01300	0.6	0.13365	0.01339	89.9	0.01293	0.01271	1.6	0.06418	0.04219	34.3
		P_g	0.01227	0.01279	−4.2	0.15438	0.01187	92.2	0.01133	0.01097	3.1	0.03263	0.00977	70.0
		Q_g	0.02285	0.02361	−3.4	0.10573	0.02299	78.1	0.02251	0.02234	0.7	0.03850	0.02333	39.3
Out-of-distribution (OOD) MAE	case14	Mean	0.29138	0.27818	3.8	0.34591	0.26954	21.8	0.29786	0.28374	4.3	0.15930	0.14110	11.0
	case118	Mean	0.47254	0.53065	−11.3	0.48121	0.46921	2.8	0.44662	0.51278	−13.5	0.29558	0.33423	−14.2
	case2000	Mean	0.25068	0.24089	3.9	0.21860	0.23550	−7.7	0.25316	0.24444	3.5	0.22905	0.21569	5.7
	case4661	Mean	0.49780	0.49983	−0.2	0.41962	0.47545	−13.4	0.47101	0.49675	−6.0	0.48663	0.44790	8.0
Metamorphic	Across trained relations	Median Δ_r	852	595	30.1	1.00×10^3	805	16.2	883	563	36.1	4.63×10^3	1.26×10^3	70.8
	Per relation	`bus-relabel`	0	0	—	1.95×10^5	5.32×10^3	97.3	0	0	—	0	0	—
		`branch-reversal`	—	—	—	—	—	—	—	—	—	2.16×10^5	5.47×10^4	74.7
		`ref-bus-move`	1.47×10^3	923	37.1	1.50×10^3	949	36.9	988	1.02×10^3	−3.3	8.02×10^4	6.29×10^4	21.5
		`component-split`	898	674	24.9	577	668	−15.7	1.00×10^3	703	29.9	190	139	26.5
		`parallel-split`	1.42×10^3	975	31.5	1.05×10^5	3.89×10^3	96.3	1.25×10^3	999	20.1	23.8	17.0	28.6
		`island-union`	0	0	—	1.62×10^5	5.81×10^3	96.4	0	0	—	—	—	—
		`basekv-scale`	—	—	—	—	—	—	—	—	—	2.07×10^5	1.57×10^4	92.5
		`tap-side-flip`	167	90.9	45.4	61.9	31.0	49.9	41.3	50.0	−21.2	6.33×10^3	1.56×10^3	75.3
		`load-as-gen`	1.01×10^3	644	36.1	658	747	−13.4	1.19×10^3	594	50.0	1.98×10^3	478	75.9
		`null-pad`	1.32×10^3	736	44.4	803	1.04×10^3	−29.4	1.20×10^3	1.02×10^3	15.6	797	316	60.4
		`island-ref-mix`	1.10×10^3	695	37.0	1.59×10^5	5.74×10^3	96.4	744	766	−2.9	6.01×10^4	4.70×10^4	21.7
Constraint violations	All views	Equality P95 (p.u.)	58.1±18.7	42.2±2.3	22.1	113.9±42.1	39.5±5.2	61.7	51.8±10.9	33.5±2.4	33.4	871±342	164±11	78.7
		Ineq. P95 (10^6× tol.)	19.3±4.2	11.9±1.6	35.2	42.1±4.6	12.7±3.6	70.1	18.8±2.8	10.2±0.8	45.4	276±91	20.1±2.1	92.3

Panel A of Table 11 takes the median over the relations used to train the models. Per-channel prediction MAEs appear in Table 10.

E.2 Accuracy and relation-level effects

On transformed scenarios, mean MAE improves by $+87.0\%$ for LapPE+MPNN, $+41.9\%$ for LUMINA, and $+1.1\%$ for RWSE+MPNN. MPNN has an improvement of -3.2%, meaning its error increases. Population standard deviations across the seed-level changes are 1.2, 1.5, 0.7, and 3.3 percentage points, respectively. Most of the LapPE+MPNN gain reflects an eigensolver mismatch between training and this audit (Sec. E.4). LapPE+MPNN, RWSE+MPNN, and LUMINA improve in all three seeds; MPNN improves in one of three. On canonical scenarios, improvements are $+86.8\%$ for LapPE+MPNN, $+0.4\%$ for RWSE+MPNN, -3.9% for MPNN, and -3.0% for LUMINA. MPNN and LUMINA worsen in every seed.

Median normalized discrepancy Δ_r falls by 30.1%, 16.2%, 36.1%, and 70.8% for MPNN, LapPE+MPNN, RWSE+MPNN, and LUMINA, respectively. It decreases in 11 of 12 model–seed pairs; one LapPE+MPNN seed is almost unchanged at -0.1%. Results differ by relation. LapPE+MPNN reduces discrepancy by 96 to 97% on its four spectral failures: `bus-relabel`, `island-union`, `island-ref-mix`, and `parallel-split`. These relations expose the arbitrary signs, bases, and ordering of its raw Laplacian eigenvectors.

These gains coexist with regressions. LapPE+MPNN worsens on `null-pad` (-29.4% improvement), `component-split` (-15.7%), and `load-as-gen` (-13.4%). RWSE+MPNN worsens on `tap-side-flip` (-21.2%), `ref-bus-move` (-3.3%), and `island-ref-mix` (-2.9%). Reducing selected spectral discrepancies while worsening other relations does not establish a complete repair. MPNN improves on all seven trained relations and LUMINA on all nine. The aggregate median summarizes the overall direction but cannot establish improvement on every relation.

Table 11: Expanded augmentation results under the AC constraint-penalty objective. Panel A compares three-seed mean accuracy, median Δ_r over each model's trained relations, and scenario-clustered constraint P95; violation entries are mean $\pm$ sample standard deviation across seeds. Equality is in per unit, while pooled inequality is dimensionless and displayed in millions of family-specific 10^{-6} tolerances. Feasible is the percentage of 50 source scenarios per seed whose 53 views all satisfy both 10^{-6} tolerances. Panel B gives the percentage reduction in Δ_r by relation; positive values denote improvement under augmentation, and dashes mark excluded relations.

A. Aggregate outcomes

Model	Arm	Canonical MAE	Transformed MAE	Median Δ_r	Equality P95 (p.u.)	Inequality P95 ($10^6\times$ tol.)	Feasible (%)
MPNN	Base	0.01295	0.01312	852	58.1±18.7	19.3±4.2	0
	Aug.	0.01345	0.01353	595	42.2±2.3	11.9±1.6	0
LapPE+MPNN	Base	0.09983	0.10211	1.00×10^3	113.9±42.1	42.1±4.6	0
	Aug.	0.01309	0.01318	805	39.5±5.2	12.7±3.6	0
RWSE+MPNN	Base	0.01255	0.01271	883	51.8±10.9	18.8±2.8	0
	Aug.	0.01250	0.01257	563	33.5±2.4	10.2±0.8	0
LUMINA	Base	0.01834	0.03594	4.63×10^3	871±342	276±91	0
	Aug.	0.01889	0.02089	1.26×10^3	164±11	20.1±2.1	0

B. Metamorphic-discrepancy change by relation (%)

Relation	MPNN	LapPE+MPNN	RWSE+MPNN	LUMINA
`bus-relabel`	—	97.3	—	—
`branch-reversal`	—	—	—	74.7
`ref-bus-move`	37.1	36.9	−3.3	21.5
`component-split`	24.9	−15.7	29.9	26.5
`parallel-split`	31.5	96.3	20.1	28.6
`island-union`	—	96.4	—	—
`basekv-scale`	—	—	—	92.5
`tap-side-flip`	45.4	49.9	−21.2	75.3
`load-as-gen`	36.1	−13.4	50.0	75.9
`null-pad`	44.4	−29.4	15.6	60.4
`island-ref-mix`	37.0	96.4	−2.9	21.7

E.3 Feasibility and out-of-distribution behavior

No evaluated prediction meets both 10^{-6} tolerances. The feasible-source rate is therefore 0% for every arm: none of the 50 sources per seed has all 53 views feasible. The scenario-clustered equality and pooled inequality P95 values fall with augmentation for all four models, but these reductions do not reach AC feasibility.

Panel A of Table 12 reports absolute violations for each constraint family separately. Within each seed, a source contributes its worst value over all evaluated views. The table gives the mean $\pm$ sample standard deviation across the seed-level medians, P95 values, and maxima. Active and reactive power, generator bounds, and thermal limits are in per unit on the system's 100-MVA base. Voltage is in per unit and angle-difference excess in radians. Keeping the families separate avoids taking an absolute maximum across incompatible units. The LUMINA baseline violates angle-difference limits even though all other arms have zero violations in that family.

Panel B expresses violations relative to physical scales. In Panel B1, the nodal complex-power residual is $\sqrt{r_{P,b}^2 + r_{Q,b}^2}$. Its maximum is divided by the source scenario's total apparent demand, $\sum_b \sqrt{P_{d,b}^2 + Q_{d,b}^2}$, using the same denominator for every transformed view of that source. The canonical metric averages these values within each seed; improvement is calculated within seed before averaging across seeds. In Panel B2, each positive inequality excess is divided by the relevant element's nominal voltage, generator capacity range, branch rating, or violated angle limit. We then

take the worst view per source before computing the P95 and maximum. These percentages give scale to each family's violations; absolute tolerances still determine feasibility. All evaluated denominators were present, finite, nonzero, and active, so no element was excluded from relative aggregation.

Cross-grid generalization results are mixed, and all three seeds agree in direction for only a few model–grid combinations (Table 13). LapPE+MPNN improves on case14 (21.8%) but worsens on case2000 and case4661 (7.7% and 13.4%). RWSE+MPNN improves on case2000 (3.5%). LUMINA improves on case14, case2000, and case4661 (11.0%, 5.7%, and 8.0%) and worsens slightly on the in-distribution case500 (3.4%). LUMINA is the only model with improvements on three unseen grids that each agree in sign across all three seeds. This agreement does not extend to case118, where its mean error increases by 14.2% and seedwise signs are mixed. Every model–grid combination without an asterisk in Table 13 has mixed seedwise signs. Our experiment therefore does not support a general cross-grid benefit from using transformations as data augmentation.

E.4 Cohort and evaluation-path provenance

The main accuracy table and the case500 row of Table 13 use different evaluation cohorts and answer different empirical questions. The main table uses the 50-source canonical subset of the transformed-view audit. The cross-grid table uses a separately selected 200-source clean cohort, with no shared case500 scenario IDs. Reaggregating the saved native-unit MAEs reproduces both tables. The cross-grid run records 24 checkpoint hashes and 24,000 model–scenario rows: four models, two arms, three seeds, five grids, and 200 scenarios. Its MPNN-family evaluation is batched; the transformed-view audit uses a different prediction path.

The LapPE baseline has mean canonical MAE 0.09983 in the main cohort and 0.01253 in the cross-grid case500 cohort. The corresponding augmented MAEs are 0.01309 and 0.01246. A matched follow-up traces the gap mainly to the eigensolver backend. LapPE+MPNN training and cross-grid evaluation use batched single-precision CUDA `torch.linalg.eigh` to compute Laplacian eigenvectors. The transformed-view audit computes them separately for each scenario using double-precision NumPy (LAPACK) `eigh` on CPU. Neither path fixes eigenvector signs or bases within repeated eigenspaces.

We re-evaluated the same six checkpoints, with hashes matching the cross-grid run, on the same 200 cross-grid case500 scenarios after changing only the eigensolver backend to CPU. Baseline mean MAE rises from 0.01253 to 0.08853 ($7.1\times$; seed values 0.0797, 0.0992, and 0.0867), close to the main-cohort value. Augmented MAE changes from 0.01246 to 0.01283 ($+3.0\%$).

Most of the large main-cohort LapPE gain therefore measures augmentation's robustness to CPU/LAPACK spectral coordinates after CUDA training. It does not demonstrate an accuracy gain under the training backend or a general repair of spectral encodings. LapPE+MPNN's main-audit baseline violation P95 and normalized discrepancies use the same CPU path and may also be inflated by the mismatch; this follow-up does not quantify those effects. The remaining cross-grid values describe transfer under their stated cohorts and sign convention.

Table 12: AC constraint violations from the augmentation training experiment. Panel A reports absolute family-level violations as baseline → augmented; each side is the mean ± sample standard deviation across three seed-level statistics after clustering 53 views under each of 50 source scenarios. P95 uses linear interpolation. Panel B reports normalized context using scenario- or element-specific denominators. The AC-feasible source-scenario rate is 0% for every arm at the 10^{-6} tolerance.

A. Absolute violations by constraint family

Model	Constraint family	Unit	Median		P95		Maximum	
			Base	Aug.	Base	Aug.	Base	Aug.
MPNN	Active-power balance	p.u.	21.1 ± 2.39	13.2 ± 1.59	58.0 ± 18.7	42.2 ± 2.26	64.1 ± 16.5	50.8 ± 6.4
	Reactive-power balance	p.u.	8.47 ± 0.476	6.5 ± 0.746	24.2 ± 9.59	14.9 ± 1.68	36.9 ± 10.2	22.3 ± 1.14
	Voltage bounds	p.u.	0.00756 ± 0.00303	0.00617 ± 0.00361	0.0288 ± 0.0149	0.0225 ± 0.0133	0.0406 ± 0.0186	0.0491 ± 0.0351
	Generator active-power bounds	p.u.	0.206 ± 0.116	0.0976 ± 0.0385	2.68 ± 0.482	0.647 ± 0.215	10.2 ± 2.16	1.29 ± 0.468
	Generator reactive-power bounds	p.u.	0.0333 ± 0.00659	0.0272 ± 0.00723	0.138 ± 0.041	0.109 ± 0.0848	0.418 ± 0.173	0.163 ± 0.0549
	Branch thermal limits	p.u.	4.92 ± 1.82	2.5 ± 0.253	19.3 ± 4.16	11.9 ± 1.57	28.8 ± 11.4	20.2 ± 3.97
	Angle-difference limits	rad	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0
LapPE+MPNN	Active-power balance	p.u.	89.2 ± 29.3	12.6 ± 2.86	114 ± 42.1	39.5 ± 5.17	128 ± 45.1	52.0 ± 17.5
	Reactive-power balance	p.u.	38.2 ± 5.89	6.18 ± 1.14	51.3 ± 2.59	12.2 ± 1.76	61.5 ± 6.8	14.7 ± 2
	Voltage bounds	p.u.	0.0447 ± 0.0123	$0.00566 \pm 6.30 \times 10^{-4}$	0.0583 ± 0.0147	0.0162 ± 0.0048	0.0777 ± 0.0302	0.0464 ± 0.0117
	Generator active-power bounds	p.u.	1.32 ± 0.712	0.0828 ± 0.0272	2.29 ± 0.733	0.565 ± 0.0711	9.16 ± 4.15	1.04 ± 0.208
	Generator reactive-power bounds	p.u.	0.23 ± 0.0212	0.0252 ± 0.0166	0.404 ± 0.0964	0.0709 ± 0.014	0.581 ± 0.186	0.103 ± 0.0321
	Branch thermal limits	p.u.	32.9 ± 5.78	2.48 ± 0.685	42.1 ± 4.6	12.7 ± 3.59	53.1 ± 5.37	17.9 ± 7.29
	Angle-difference limits	rad	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0
RWSE+MPNN	Active-power balance	p.u.	19.2 ± 3.05	10.6 ± 2.18	51.8 ± 10.9	33.5 ± 2.38	74.2 ± 13.3	46.6 ± 4.53
	Reactive-power balance	p.u.	8.39 ± 1.49	5.18 ± 0.414	19.3 ± 4.62	12.1 ± 0.743	30.2 ± 5.67	13.5 ± 1.31
	Voltage bounds	p.u.	0.00707 ± 0.00302	0.00629 ± 0.00287	0.0206 ± 0.00408	0.0196 ± 0.00329	0.0347 ± 0.00747	0.0333 ± 0.012
	Generator active-power bounds	p.u.	0.162 ± 0.0675	0.0709 ± 0.0325	2.21 ± 0.666	0.525 ± 0.112	10.6 ± 1.41	1.02 ± 0.211
	Generator reactive-power bounds	p.u.	0.0268 ± 0.00296	0.0219 ± 0.00326	0.186 ± 0.128	0.0868 ± 0.0609	0.289 ± 0.103	0.115 ± 0.0765
	Branch thermal limits	p.u.	4.33 ± 0.889	2.72 ± 0.468	18.8 ± 2.82	10.2 ± 0.843	25.8 ± 3.6	13.6 ± 3.63
	Angle-difference limits	rad	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0	0 ± 0
LUMINA	Active-power balance	p.u.	736 ± 411	91.2 ± 9.83	871 ± 342	164 ± 10.6	989 ± 401	188 ± 8.73
	Reactive-power balance	p.u.	568 ± 512	24.2 ± 2.75	693 ± 405	40.3 ± 5.05	1473 ± 1007	54.5 ± 2.02
	Voltage bounds	p.u.	$2.38 \times 10^{-8} \pm 0$	0 ± 0	$2.38 \times 10^{-8} \pm 0$	0 ± 0	$2.38 \times 10^{-8} \pm 0$	0 ± 0
	Generator active-power bounds	p.u.	$7.12 \times 10^{-8} \pm 4.58 \times 10^{-8}$	$3.21 \times 10^{-8} \pm 1.71 \times 10^{-8}$	$1.06 \times 10^{-7} \pm 3.03 \times 10^{-8}$	$3.21 \times 10^{-8} \pm 1.71 \times 10^{-8}$	$1.24 \times 10^{-7} \pm 0$	$4.20 \times 10^{-8} \pm 0$
	Generator reactive-power bounds	p.u.	$6.44 \times 10^{-9} \pm 7.46 \times 10^{-9}$	0 ± 0	$1.86 \times 10^{-8} \pm 2.83 \times 10^{-8}$	0 ± 0	$2.86 \times 10^{-8} \pm 4.56 \times 10^{-8}$	0 ± 0
	Branch thermal limits	p.u.	238 ± 119	10.8 ± 0.727	276 ± 91.0	20.1 ± 2.13	502 ± 386	32.0 ± 3.52
	Angle-difference limits	rad	10.9 ± 1.85	0 ± 0	11.9 ± 0.445	0 ± 0	13.2 ± 1.9	0 ± 0

Table 12 (continued). AC constraint violations from the augmentation training experiment. Panel B1 gives system-scale equality context and the voltage global maximum. Panel B2 reports relative inequality overruns by family. Each entry is the mean ± sample standard deviation across three seeds.

B. Normalized interpretability metrics

B1. System-scale equality and voltage context

	Global-max nodal complex-power balance residual (% total apparent demand)		Global-max voltage-bound excess (% nominal)		Canonical mean maximum nodal complex-power balance residual (% total apparent demand)		
Model	Base	Aug.	Base	Aug.	Base	Aug.	Imp. (%)
MPNN	36.9 ± 7.29	28.8 ± 2.42	4.06 ± 1.86	4.91 ± 3.51	4.32 ± 0.484	3.94 ± 0.393	8.3
LapPE+MPNN	72.6 ± 25.1	28.6 ± 9.52	7.77 ± 3.02	4.64 ± 1.17	29.6 ± 6.21	3.34 ± 0.385	88.4
RWSE+MPNN	40.4 ± 7.42	25.7 ± 2.71	3.47 ± 0.747	3.33 ± 1.2	3.78 ± 0.28	3.31 ± 0.157	12.2
LUMINA	893 ± 499	103 ± 5.4	$2.38 \times 10^{-6} \pm 0$	0 ± 0	10.0 ± 0.493	8.85 ± 0.481	11.8

B2. Relative inequality overruns (%)

Model	Constraint family	P95		Maximum	
		Base	Aug.	Base	Aug.
MPNN	Voltage bounds	2.88 ± 1.49	2.25 ± 1.33	4.06 ± 1.86	4.91 ± 3.51
	Generator active-power bounds	157 ± 13.4	135 ± 47.5	270 ± 77.9	202 ± 129
	Generator reactive-power bounds	838 ± 434	628 ± 282	3212 ± 2238	2570 ± 1685
	Branch thermal limits	892 ± 197	375 ± 108	1418 ± 675	708 ± 158
	Angle-difference limits	0 ± 0	0 ± 0	0 ± 0	0 ± 0
LapPE+MPNN	Voltage bounds	5.83 ± 1.47	1.62 ± 0.480	7.77 ± 3.02	4.64 ± 1.17
	Generator active-power bounds	616 ± 122	165 ± 104	942 ± 431	409 ± 293
	Generator reactive-power bounds	1203 ± 408	494 ± 249	2687 ± 1645	1934 ± 198
	Branch thermal limits	2098 ± 288	541 ± 222	2309 ± 306	649 ± 185
	Angle-difference limits	0 ± 0	0 ± 0	0 ± 0	0 ± 0
RWSE+MPNN	Voltage bounds	2.06 ± 0.408	1.96 ± 0.329	3.47 ± 0.747	3.33 ± 1.20
	Generator active-power bounds	180 ± 84.6	216 ± 33.6	499 ± 405	679 ± 485
	Generator reactive-power bounds	688 ± 548	438 ± 299	1663 ± 1174	1534 ± 939
	Branch thermal limits	803 ± 73.9	458 ± 79.9	1141 ± 229	601 ± 127
	Angle-difference limits	0 ± 0	0 ± 0	0 ± 0	0 ± 0
LUMINA	Voltage bounds	$2.38 \times 10^{-6} \pm 0$	0 ± 0	$2.38 \times 10^{-6} \pm 0$	0 ± 0
	Generator active-power bounds	$1.19 \times 10^{-5} \pm 1.23 \times 10^{-6}$	$4.76 \times 10^{-6} \pm 2.62 \times 10^{-6}$	$1.19 \times 10^{-5} \pm 1.23 \times 10^{-6}$	$4.85 \times 10^{-6} \pm 2.53 \times 10^{-6}$
	Generator reactive-power bounds	$7.01 \times 10^{-7} \pm 7.38 \times 10^{-7}$	0 ± 0	$8.63 \times 10^{-7} \pm 9.99 \times 10^{-7}$	0 ± 0
	Branch thermal limits	9400 ± 5026	859 ± 27.6	9975 ± 5084	1537 ± 233
	Angle-difference limits	2274 ± 85.0	0 ± 0	2530 ± 363	0 ± 0

Table 13: Effect of transformation augmentation on other grids. Entries are average per-seed percent changes in equal-channel MAE, and **negative** values denote improvement under augmentation. An asterisk marks entries where all three seeds agree in sign. Case500 is a separate 200-source clean cohort with no scenario overlap with the main 50-source audit; evaluation paths also differ (Sec. E.4). The other grids receive no additional training.

Grid	MPNN	LapPE+MPNN	RWSE+MPNN	LUMINA
case14	-3.8	-21.8^{*}	-4.3	-11.0^{*}
case118	$+11.3$	-2.8	$+13.5$	$+14.2$
case500	$+4.8$	-0.4	$+0.4$	$+3.4^{*}$
case2000	-3.9	$+7.7^{*}$	-3.5^{*}	-5.7^{*}
case4661	$+0.2$	$+13.4^{*}$	$+6.0$	-8.0^{*}

F Design and transparency implications

F.1 Designing consistency by construction

An admitted relation links two encodings of the same dispatch problem, with predictions compared after the declared output map. A material difference between those aligned predictions means that the deterministic surrogate depends on the encoding as well as on the physical problem. Low reference-label error on one serialization can miss this dependence.

A relation enforced by construction holds for any learned weights, whether or not it appeared in training. The architecture screen points to positional summaries that preserve permutation equivariance, output partitions based on component type or electrical role, and symmetric transport of orientation-dependent branch attributes. Preprocessing can address some failures by standardizing units, angle references, and row conventions. Relations that change component count or type instead require common component encodings and additive aggregation at physical buses or devices. These changes remain compatible with graph models, provided each proposed remedy addresses the relation that failed.

Where an exact construction is unavailable, training-time augmentation can reduce sensitivity, including for `ELC` relations that change graph size. Its effect is approximate and depends on optimization. In our experiment, aggregate representation sensitivity falls, but some relations worsen, canonical error can increase or decrease, and every prediction remains infeasible. A useful AC-OPF surrogate therefore needs exact consistency for representation and angle-reference choices where practical, measured consistency for the remaining equivalences, and separate tests of per-channel accuracy, AC feasibility, bounds, and dispatch cost.

F.2 Minimum transparency package

To reproduce a consistency verdict, an independent auditor must construct both views of a case and align their outputs. The following information makes those steps possible. A versioned hosted endpoint can substitute for released parameters if it exposes the same documented inputs and outputs.

- Runnable inference requires a versioned checkpoint with checksums and a pinned environment, or a versioned endpoint that accepts documented case files.
- The input conventions cover schema and feature order, physical units and per-unit bases, the angle-reference convention, and normalization statistics.
- The output schema gives the physical quantity, unit, and stable bus or device identity of each output channel.
- The training topologies list the grids or grid families represented in training.
- Each claimed invariance names a relation family that preprocessing, architecture, or training is said to enforce.
- A worked example pairs one case file with its encoded model input and numerical canonical outputs.