arXiv:2611.16662v1 [cs.CL] 12 Nov 2026

MuSeToM: Multi-Sensory Evaluation of Belief Revision with Textual Cues Across Six Senses

Chunkit Chan[♠*] **Kawai Chung**[♠*] **Sze Heng Douglas Kwok**[♢*] **Yuxuan Liu**[♠] **Yauwai Yim**[♠]
Tianshi Zheng[♠] **Haochen Shi**[♠] **Qing Zong**[♠] **Wong Kai Chung**[♠] **Pak Hin Lau**[♠]
YANG Calvin[♠] **Yifan Gao**[♡] **Xiaoman Pan**[♡] **Ginny Wong**[♣]
Simon See[♣] **Yangqiu Song**[♠]
[♠]The Hong Kong University of Science and Technology [♢]Stanford University
[♡]Amazon [♣]NVIDIA AI Technology Center (NVAITC)
{ckchancc, yqsong}@cse.ust.hk

Abstract

Theory of Mind (ToM), the ability to infer mental states in others, is pivotal for human social cognition. Existing evaluations of ToM in large language models (LLMs) mainly focus on static belief attribution and false-belief reasoning, leaving a critical aspect of social reasoning underexplored: *can LLMs decide whether perceptual evidence is sufficient to justify another person's belief revision?* To address this gap, we present MuSeToM, a human-validated, text-based benchmark for evaluating epistemic belief revision in ToM reasoning. MuSeToM contains 2,415 instances adapted from BigToM and OpenToM, with textual cues spanning six sensory categories, and tests whether models can determine when a protagonist should revise a belief and infer the resulting updated belief state. Using MuSeToM, we systematically evaluate state-of-the-art LLMs under direct and chain-of-thought prompting, revealing a pronounced dissonance: while models often detect that a belief may change, they struggle to identify the updated belief and frequently over-update on perceptually salient but epistemically insufficient cues[1].

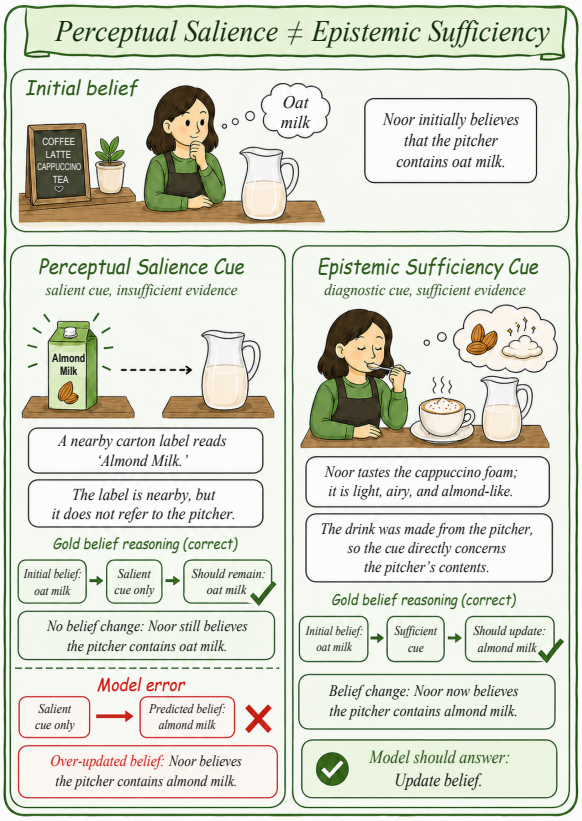

Figure 1: An example in MuSeToM. The model incorrectly treats this cue as sufficient for belief change, whereas the gold reasoning preserves Noor's prior belief.

1 Introduction

Theory of Mind (ToM), the ability to reason and attribute the mental states of others, is increasingly recognized as an important capability for large language models (LLMs) deployed in human-centered environments (Premack and Woodruff, 1978; Ma et al., 2023). Recent studies show that modern LLMs can achieve strong performance on classical ToM benchmarks, particularly on false-belief reasoning tasks (Strachan et al., 2024; Bubeck et al., 2023; Kosinski, 2023). However, most existing evaluations focus primarily on belief attribution: given a narrative context, models are asked to infer what another agent currently believes (Le et al., 2019a; Kim et al., 2023a). In contrast, a fundamental aspect of human social reasoning remains underexplored: *determining whether newly observed evidence is sufficient to justify revising a prior belief.*

In realistic environments, humans continuously update their beliefs based on perceptual evidence. However, belief revision does not occur merely because new information is observed. Instead, humans evaluate whether the available evidence provides sufficient epistemic support for updating an

[*] Equal contribution.
[1]Code and data are available at https://github.com/HKUST-KnowComp/MuSeToM.

existing belief. In Figure 1, Noor tastes the cappuccino foam and finds it light, airy, and almond-like, providing sufficient evidence to revise her belief from oat milk to almond milk. In contrast, a nearby carton labeled "almond milk" may be perceptually salient and attract attention, yet such evidence alone remains insufficient to justify the same conclusion.

Perceptual salience refers to the extent to which a stimulus stands out and captures attention due to its perceptual prominence (Itti, 2007; MacLeod, 2015). Nevertheless, perceptually salient cues are not necessarily evidentially informative. In contrast, we use *epistemic sufficiency* to describe whether available evidence provides adequate support to rationally justify belief revision under uncertainty, consistent with theories of probabilistic and causal inference in human cognition (Tenenbaum et al., 2011; Knill and Pouget, 2004). This distinction between perceptual salience and epistemic sufficiency is fundamental to human reasoning under uncertainty. Prior work in cognitive and developmental psychology further suggests that humans use perceptual evidence not only to form beliefs about the world, but also to infer when other individuals should revise their beliefs (Baron-Cohen, 1991; Gervais, 2013; Raleigh and Vindrola, 2021; Pratt and Bryant, 1990). Consequently, robust ToM reasoning requires more than tracking what another agent can perceive; it also requires evaluating whether the perceived evidence is sufficiently informative to warrant belief revision.

Despite recent progress in ToM evaluation, this capability remains underexplored. Most prior benchmarks primarily evaluate whether models can infer another agent's belief from a narrative context, rather than whether newly encountered perceptual evidence should trigger belief revision (Le et al., 2019a; Kim et al., 2023a; Xiao et al., 2025). Even perception-oriented ToM work has mainly emphasized recovering perceptual access or using a character's perspective to infer current beliefs, rather than testing whether a newly encountered cue warrants revision of a prior belief (Jung et al., 2024; Wilf et al., 2024). Moreover, perceptual information is often simplified into coarse visibility signals (e.g., whether an event was observed), which bypasses the more challenging problem of assessing whether evidence is epistemically sufficient to support belief revision. As a result, current benchmarks provide limited insight into whether LLMs can distinguish genuinely informative evidence from perceptually salient but misleading or irrelevant cues. To address this gap, we introduce MuSeToM, a benchmark for evaluating epistemic belief revision in LLM theory of mind reasoning through multi-sensory perceptual evidence. MuSeToM presents scenarios in which a protagonist initially holds a belief and subsequently encounters perceptual cues arising from events in the environment. These cues span multiple sensory modalities and vary in epistemic strength: some provide sufficient evidence to justify belief revision, while others are perceptually salient yet insufficient, indirect, or entirely irrelevant. All model inputs are text; the six sensory categories refer to the protagonist's experiences described in the cues. The benchmark evaluates two core abilities: (1) determining whether the protagonist should revise their belief, and (2) identifying the resulting updated belief state.

Experiments across a diverse set of state-of-the-art LLMs reveal several consistent findings. First, models perform substantially better at detecting potential belief changes than at identifying the resulting updated belief state. Second, models exhibit a systematic over-updating bias, frequently treating perceptually salient cues as sufficient evidence for belief revision even when the cues are weak or indirect. These findings suggest that current LLMs often conflate perceptual salience with epistemic sufficiency, leading to unreliable belief revision and underscoring the need for more robust and realistic evaluations of theory of mind reasoning. Our contributions are summarized as follows:

- We introduce MuSeToM, the first human-validated benchmark designed to evaluate epistemic belief revision in LLM theory of mind reasoning from multi-sensory perceptual evidence. MuSeToM contains 2,415 human-validated instances adapted from BigToM and OpenToM, spanning six sensory modalities and both clean and noisy reasoning settings.
- We propose evaluation tasks that measure belief-change detection, updated belief-state identification, and epistemic relevance filtering under distractor cues.
- Through extensive experiments across state-of-the-art LLMs, we identify systematic reasoning failures, including over-updating biases and difficulties distinguishing perceptual salience from epistemic sufficiency.

2 Related Work

2.1 Theory of Mind Benchmarks

Early computational ToM benchmarks build on classic false-belief tests such as the Sally-Anne task (Baron-Cohen et al., 1985). ToM-bAbi and ToMi formulate false-belief reasoning as question answering over synthetic narratives (Grant et al., 2017; Nematzadeh et al., 2018; Le et al., 2019a). Subsequent work expands this setting to more realistic and challenging social reasoning scenarios. BigToM uses model-generated stories to evaluate social reasoning (Gandhi et al., 2023a), while FANToM stresses ToM in dialogue interactions (Kim et al., 2023a). ToMBench provides a broader evaluation framework over social cognition abilities (Chen et al., 2024b), and OpenToM introduces longer narratives with richer character and object descriptions (Xu et al., 2024a). Other benchmarks study situated ToM, higher-order belief reasoning, temporal mental-state tracking, and benchmark robustness (Ma et al., 2023; Wu et al., 2023; Street et al., 2026; Xiao et al., 2025; Ullman, 2023; Shapira et al., 2024; Riemer et al., 2025). Recent work further extends ToM evaluation to multilingual, culturally adapted, and multimodal settings (Chan et al., 2025; Sadhu et al., 2024; Jin et al., 2024a; Chen et al., 2024a; Shi et al., 2025b; Villa-Cueva et al., 2025). These benchmarks mainly evaluate belief attribution or state tracking, which do not directly test whether newly perceived evidence is sufficient to revise another agent's prior belief. More related works are provided in Appendix D.

2.2 Perception and Belief Revision

Several recent studies examine perception and perspective-taking in ToM reasoning. PercepToM introduces Percept-ToMi and Percept-FANToM to evaluate perception inference and perception-to-belief inference (Jung et al., 2024). SimToM filters the narrative context to the target character's perspective before answering ToM questions (Wilf et al., 2024). SymbolicToM tracks character belief states with explicit symbolic representations (Sclar et al., 2023). These works are closest to MuSeToM because they connect perceptual access, perspective context, and belief attribution. However, they focus on what an agent perceived or what belief follows from accessible context. MuSeToM instead asks whether a newly encountered perceptual cue is strong enough to revise a prior belief. Belief-R is also related because it evaluates whether language models revise their own conclusions after new evidence is introduced (Wilie et al., 2024). In contrast, MuSeToM evaluates whether a model can infer when another agent should update a belief, and what the updated belief should be, from perceptual evidence available to that agent. Appendix D provides a structured comparison with Belief-R.

3 MuSeToM

We introduce MuSeToM, a multi-sensory benchmark designed to evaluate whether large language models (LLMs) can reason about epistemic belief revision from perceptual evidence and determine when a protagonist should revise a prior belief and infer the resulting updated belief state from perceptual cues in the environment.

3.1 Design Considerations for MuSeToM

A benchmark for perception-driven Theory of Mind like MuSeToM should satisfy three key requirements. First, it should provide realistic narrative contexts with sufficiently rich characters and environments to support perceptual experiences across multiple sensory modalities. Second, it should specify a clear initial belief held by the target agent. Third, it should contain, or allow the construction of, an event that either justifies revising the initial belief or confirms that the original belief should remain unchanged. These requirements are essential for evaluating epistemic belief revision rather than simple belief attribution. In particular, the benchmark must distinguish between perceptually salient observations and evidence that is genuinely sufficient to support belief change.

3.2 Source Data Collection

Based on these criteria, we construct MuSeToM using stories adapted from BigToM (Gandhi et al., 2023a) and OpenToM (Xu et al., 2024a). BigToM contains 200 structured stories that explicitly specify several components necessary for constructing belief-revision scenarios, including an initial belief, a belief-changing causal event, a non-belief-changing random event, and the resulting updated belief. These elements make BigToM a natural source for constructing perception-based belief revision examples. OpenToM complements BigToM by providing 7,000 narrative scenarios with longer contexts, richer character settings, and more detailed object-centric descriptions. It also contains explicit inter-agent belief statements that can be reformulated as initial beliefs. These properties

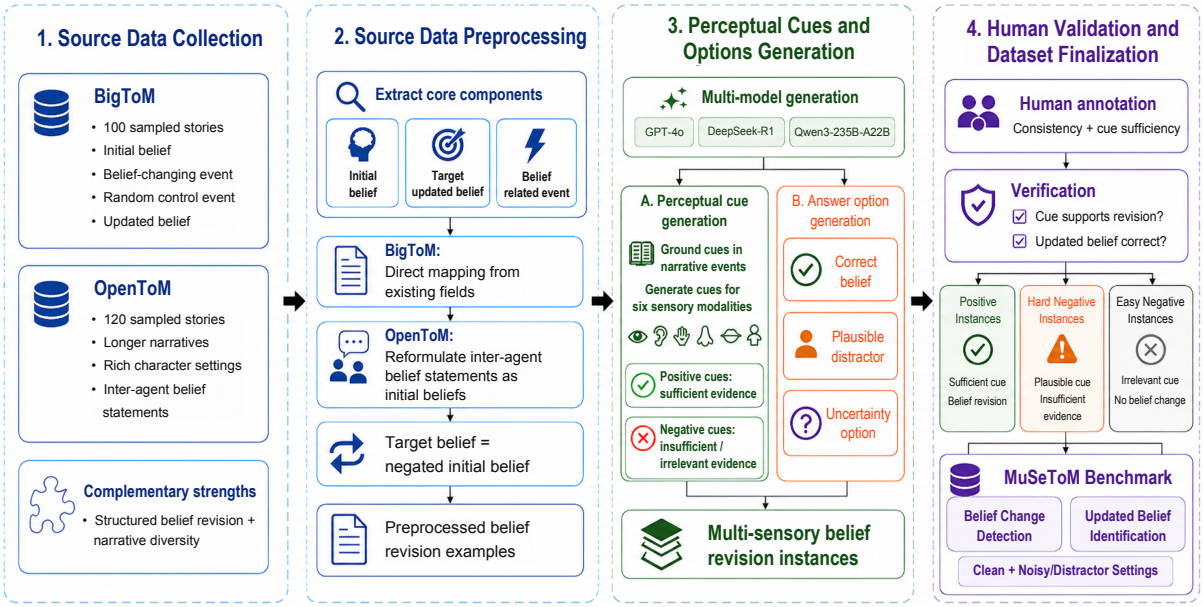

Figure 2: Overview of the MuSeToM dataset construction pipeline.

make OpenToM well suited for constructing scenarios in which perceptual evidence may or may not justify belief revision. The two datasets provide complementary advantages: BigToM offers explicit belief revision structure, while OpenToM contributes broader narrative diversity and richer environmental detail. To assess potential data contamination risk, we report a detailed contamination check in Appendix A.5.

3.3 Task Definition

MuSeToM evaluates epistemic belief revision in Theory-of-Mind reasoning. Given a textual story context, an initial belief, and one or more text-described sensory cues, models are evaluated on four task formats: (1) *clean belief-change detection*, which asks whether a cue provides sufficient evidence to trigger belief revision; (2) *clean updated-belief identification*, which asks models to identify the resulting belief state; (3) *noisy belief-change detection*, which adds a distractor cue; and (4) *noisy updated-belief identification*, which tests belief-state inference under distractor cues. This design evaluates both belief revision and epistemic relevance filtering. We evaluate categorical belief revision under the shared evidence-sufficiency criterion described in Appendix A.2; a no-change label does not rule out reduced confidence in the prior belief. Detailed task questions and prompting formats are provided in Appendix B.2.

3.4 Dataset Construction

Source Data Preprocessing We sample 100 stories from BigToM and 120 stories from OpenToM as the source data for MuSeToM. For each example, we extract three core components: (1) the initial belief, (2) the target updated belief, and (3) the event associated with belief revision. For BigToM, this mapping is straightforward because the dataset already explicitly provides the relevant belief-revision components. For OpenToM, each example involves two agents and an object embedded within a narrative context. We extract explicit inter-agent belief statements (e.g., "[Agent 1] believes [Agent 2] likes [object]") as the initial beliefs. The target updated belief is then constructed as the corresponding contrastive belief, while the belief-changing events are instantiated during the next stage.

Perceptual Cues and Options Generation We use multiple LLMs, including GPT-4o, DeepSeek-R1, and Qwen3-235B-A22B, to generate candidate perceptual cues grounded in the narrative events. Using multiple model families aims to reduce dependence on any individual generator's linguistic style (OpenAI et al., 2024; DeepSeek-AI, 2025a; Qwen Team, 2025a). MuSeToM covers six sensory modalities. In addition to the traditional Aristotelian five senses—visual, auditory, tactile, olfactory, and gustatory—we additionally include proprioception, which refers to awareness of body position and movement (Aristotle, 1984). Propri-

oception plays an important role in how humans interact with their environment and construct beliefs about external events (Gao et al., 2019; Biel and Peske, 2005).

For each story, we generate candidate cues intended to support belief revision, alongside candidates intended to be epistemically insufficient or irrelevant to the protagonist's belief. Final cue labels are assigned after human validation and filtering. For the Updated Belief Identification task, we further generate multiple-choice answer options, including a plausible but incorrect distractor belief and an uncertainty option. The full generation prompts are provided in the Appendix, and additional details on OpenToM contrastive belief construction and noisy distractor construction are provided in Appendix A.1.1.

Human Validation All generated cues are validated, filtered, and labeled by three human annotators who assess logical consistency, plausibility, and epistemic appropriateness. Annotators determine whether each perceptual cue provides sufficient evidence to justify belief revision and verify the corresponding updated belief state. Unreasonable, implausible, or logically inconsistent cues are excluded. Based on the human validation process, the final instances are grouped into three categories: **Positive Instances**, where the perceptual cue provides sufficient evidence for belief revision; **Hard Negative Instances**, where the cue is coherent with the story, available to the protagonist, and relevant to the target belief, but insufficient to support a definite revised belief; **Easy Negative Instances**, where the cue is intentionally irrelevant to the protagonist's belief. Appendix A.2 details the normative cue-label criteria and annotation protocol.

3.5 Benchmark Validation

To evaluate annotation reliability, we measure inter-annotator agreement using Fleiss' κ (Fleiss, 1971). Three annotators annotated each perceptual cue with respect to two questions: whether the cue causes a belief change, and what the protagonist's updated belief is. To improve annotation consistency, we first evaluated annotators on their initial 100 instances and provided detailed feedback on common error patterns before proceeding to the full dataset. After completing the annotation process, we filter out instances with unreasonable or implausible cues. Despite the complexity of inferring belief revision from perceptual cues across diverse sensory modalities, we observe strong inter-annotator agreement on both source datasets. For *Belief Change Detection*, Fleiss' κ is 0.776 on BigToM and 0.742 on OpenToM. For *Updated Belief Identification*, κ is 0.744 on BigToM and 0.742 on OpenToM. Appendix A.2.1 reports annotation robustness and bootstrap intervals. Appendix A.4 shows that removing the story and prior belief reduces performance in both tested cue-only models, while the cue-text probe remains above chance. Context contributes to performance, although these controls do not rule out partial reliance on cue artifacts.

3.6 Dataset Statistics

MuSeToM is constructed from 220 source stories, including 100 stories from BigToM and 120 stories from OpenToM. After human validation and filtering, the final benchmark contains 2,415 perceptual-cue instances spanning six sensory modalities, including 1,169 instances derived from BigToM and 1,246 derived from OpenToM. Detailed statistics are provided in Table 6 in the Appendix.

4 Experimental Setup

We evaluate state-of-the-art LLMs on MuSeToM under two prompting settings: direct prompting and chain-of-thought (CoT) prompting. Each model is tested on both binary belief-change detection and multiple-choice updated-belief identification. The evaluated models cover GPT-3.5-Turbo (OpenAI, n.d.), GPT-4o (OpenAI et al., 2024), GPT-OSS-120B (OpenAI et al., 2025), Llama-3.1 (Llama Team, 2024), Llama-4-Maverick (Meta Llama, 2025), Mistral-7B (Jiang et al., 2023a), Mixtral-8x22B (Mistral AI, 2024b), Qwen3-235B-A22B (Qwen Team, 2025a), QwQ-32B (Qwen Team, 2025b), Qwen3.5-397B-A17B (Qwen Team, 2026), DeepSeek-R1 (DeepSeek-AI, 2025a), DeepSeek-V3.1 (DeepSeek-AI, 2024, 2025b), DeepSeek-V3.2 (DeepSeek-AI, 2025c), Kimi-K2-Thinking (Moonshot AI, 2025), MiniMax-M2.1 (MiniMax, 2025), MiniMax-M2.5 (MiniMax, 2026), and Phi-3-Medium-4K (Microsoft, 2024). We report results in two input settings. In the *clean setting*, each instance contains a single perceptual cue. In the *noisy setting*, an additional perceptual cue is included as a distractor. We use accuracy and macro F1 as aggregate metrics. To analyze belief-revision behavior, we further report true positive rate (TPR) on positive cases, true negative

Model / reference	Belief Change				Belief Change with Noise			
	Task 1 (Binary)		Task 2 (Multiple Choice)		Task 3 (Binary)		Task 4 (Multiple Choice)	
	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)
GPT-3.5-Turbo	44.76	44.55	41.57	41.60	64.73	64.51	62.91	62.79
GPT-4o	70.23	70.09	50.02	49.82	71.73	71.72	60.01	59.54
GPT-OSS-120B	82.53	82.53	61.04	61.08	**80.51**	**80.48**	70.94	70.88
Llama-8B	52.09	51.32	34.20	34.09	62.37	60.82	52.26	51.80
Llama-405B	77.27	77.25	54.66	54.63	71.34	71.13	59.27	59.15
Llama-4-Maverick-17B	67.83	67.51	56.11	55.95	63.86	62.35	64.33	64.19
Mistral-7B	55.82	55.26	41.86	41.79	60.35	59.67	58.53	58.61
Mixtral-8x22B	63.52	63.35	43.60	43.50	65.88	65.67	56.71	56.68
Qwen3-235B-A22B	67.83	67.75	57.10	57.12	68.71	68.71	66.76	66.73
QwQ-32B	82.61	82.61	63.52	63.64	77.07	77.07	67.16	67.33
Qwen3.5-397B-A17B	85.34	85.34	74.58	74.57	67.84	67.81	**71.27**	**71.24**
DeepSeek-R1	77.81	77.54	70.02	69.97	68.54	68.24	67.15	67.10
DeepSeek-V3.1	82.57	82.57	65.13	65.12	55.97	55.70	56.98	56.92
DeepSeek-V3.2	75.65	75.57	70.72	70.68	69.86	69.86	60.62	60.48
Kimi-K2-Thinking	67.41	65.50	**75.40**	**75.39**	77.07	77.05	69.45	69.45
MiniMax-M2.1	78.55	78.54	73.00	73.00	76.67	76.67	71.07	71.07
MiniMax-M2.5	**86.05**	**86.05**	74.12	74.10	79.84	79.83	70.73	70.74
Phi-3-Medium-4K	56.56	55.72	49.77	49.77	63.79	63.37	66.42	66.37
Human reference	94.29	94.28	93.83	93.83	90.76	90.76	90.02	89.95

Table 1: Direct-prompting results on MuSeToM across clean and noisy belief-change settings. Accuracy (%) and Macro-F1 (%) are reported for binary belief-change detection and multiple-choice updated-belief identification. Best and second-best LLM accuracy results in each column are marked in bold and underlined, respectively.

rate on easy negatives (TNR-e), and true negative rate on hard negatives (TNR-h). To mitigate position bias, we randomly shuffle the order of the answer options and cue presentation for every instance (Zheng et al., 2024). We also report a small-sample human reference based on responses from three computer science graduate students who completed the same evaluation tasks. Appendix B.3 reports human evaluation details, modality results, individual scores, and agreement.

5 Main Results

Table 1 reports direct-prompting results across four tasks. All models score below this three-evaluator reference, with especially large gaps in updated-belief identification. Clean binary detection generally yields higher accuracy than updated-belief identification: GPT-OSS-120B scores 82.53% versus 61.04%, and Llama-405B scores 77.27% versus 54.66%. Accuracy leaders also differ: MiniMax-M2.5 and Kimi-K2-Thinking lead the clean binary and multiple-choice tasks, while GPT-OSS-120B and Qwen3.5-397B-A17B lead their noisy counterparts. Clean-to-noisy binary accuracy comparisons vary by model: GPT-3.5-Turbo increases from 44.76% to 64.73%, whereas DeepSeek-V3.1 decreases from 82.57% to 55.97%. These differences motivate the following diagnostics of positive detection, hard-negative rejection, and relevance filtering.

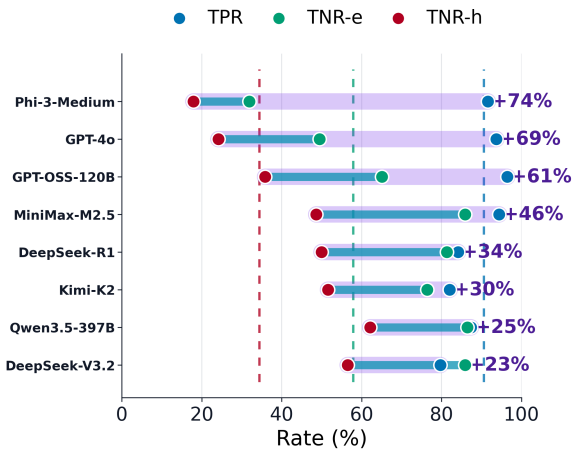

Figure 3: Diagnostic comparison of positive-cue detection and easy- and hard-negative rejection in clean direct-prompting tasks. Dashed vertical lines mark the corresponding all-model means.

5.1 Models Over-Predict Belief Revision on Perceptually Salient Cues

Figure 3 compares positive-cue detection with easy- and hard-negative rejection. Here, high TPR means

that a model correctly predicts belief change when the cue provides sufficient evidence, while high TNR-h means that a model correctly rejects belief change when the cue is plausible but epistemically insufficient. Models generally achieve higher TPR than TNR-h, indicating better acceptance of sufficient cues than rejection of plausible but insufficient ones. The TPR–TNR-h gaps reach 74% for Phi-3-Medium, 69% for GPT-4o, and 61% for GPT-OSS-120B. Even MiniMax-M2.5, the strongest clean binary model in Table 1, shows a 46% gap. These gaps suggest that perceptually salient cues can be treated as sufficient evidence even when they do not justify a revised belief. The resulting asymmetry can be obscured when positive and negative cases are summarized by aggregate accuracy. The all-model means likewise show substantially higher positive-cue detection and easy-negative rejection than hard-negative rejection. Models can therefore predict no belief change, but do so less reliably for plausible, insufficient evidence than for irrelevant cues.

Appendix C.2 reports a separate clean Task 1 analysis of 138 matched groups holding the story, prior belief, event, and sensory category fixed. Direct-prompting joint correctness is 26.7% for positive–hard-negative pairs and 47.4% for positive–easy-negative pairs. Models therefore struggle more to distinguish sufficient from relevant-but-insufficient evidence within matched contexts.

Appendix C.3 reports a separate clean Task 2 analysis: under direct prompting, 49.4% of hard-negative responses adopt the target belief and 25.9% attribute uncertainty to the protagonist. Non-prior responses therefore include both unsupported belief replacement and attributed uncertainty. These percentages describe response categories, not the models' internal confidence; the uncertainty option concerns the protagonist's belief state. This distinction matters when interpreting an incorrect response as an unsupported categorical belief change.

5.2 Distractor Cues and CoT Test Epistemic Relevance Filtering

Figure 4 compares hard-negative overprediction across clean and noisy settings, while Table 2 decomposes CoT effects into changes in positive-cue detection and hard-negative rejection. Distractor cues require models to identify relevant evidence before deciding whether belief revision is warranted, making noisy settings a stress test for epistemic relevance filtering. In particular, aggregate performance alone may obscure whether a model succeeds by detecting sufficient belief-changing cues, by rejecting hard negatives, or by shifting its response bias between these two behaviors. A lower change rate can improve rejection while reducing positive detection, so both components are needed to interpret the effect of CoT.

Model	ΔTPR	ΔTNR-h	Filtering gain
Phi-3-Medium	-13.7	10.2	23.9
GPT-4o	2.3	9.8	7.5
GPT-OSS-120B	-0.6	3.8	4.4
MiniMax-M2.5	-2.4	-1.5	1.0
DeepSeek-R1	9.1	-3.1	-12.2
Kimi-K2-Thinking	-12.8	-1.7	11.1
Qwen3.5-397B	1.2	3.6	2.3
DeepSeek-V3.2	0.4	-1.7	-2.1
Mean	-2.1	2.4	4.5

Table 2: Effect of chain-of-thought prompting on epistemic relevance filtering for selected models under noisy settings. Filtering gain is defined as $\Delta\text{TNR-}h - \Delta\text{TPR}$.

Figure 4 shows that hard-negative overprediction remains high across both clean and noisy conditions, indicating that distractor cues do not eliminate the tendency to over-predict belief revision on plausible but insufficient evidence. Several models, including GPT-3.5-Turbo, Llama-8B, Mistral-7B, Mixtral-8x22B, Phi-3-Medium, and Qwen3-235B-A22B, show hard-negative overprediction rates around or above 70% in at least one direct-prompting setting. At the same time, noisy inputs do not produce a uniform degradation pattern. For some models, such as DeepSeek-V3.2, noisy hard-negative overprediction is lower than clean hard-negative overprediction, whereas other models continue to show high overprediction under noise. This suggests that noisy settings change the balance between accepting sufficient evidence and rejecting plausible but insufficient cues, rather than simply making all examples harder. Table 2 further shows why aggregate performance can mask different effects on positive detection and hard-negative rejection. If CoT mainly improves epistemic relevance filtering, gains in TNR-h should exceed gains in TPR, meaning that the model becomes better at rejecting hard negatives rather than merely predicting more belief changes. On average, CoT slightly decreases positive-cue detection by 2.1% while increasing hard-negative rejection by 2.4%, producing a mean filtering gain of 4.5%. However, this average hides substantial model-level

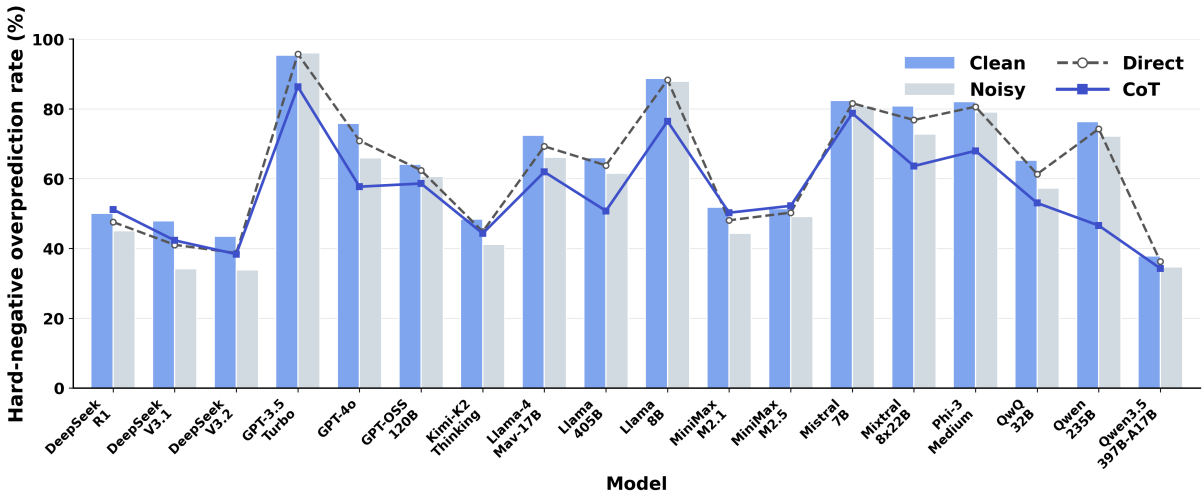

Figure 4: Hard-negative overprediction rates across clean and noisy settings under direct and chain-of-thought prompting. Bars report setting-level hard-negative overprediction rates, defined as $1 - \text{TNR-}h$, while lines compare direct and chain-of-thought prompting.

variation. Phi-3-Medium shows the largest filtering gain, 23.9%, because CoT reduces TPR by 13.7% while increasing TNR-h by 10.2%. GPT-4o and GPT-OSS-120B show more direct filtering improvements, with TNR-h increasing by 9.8% and 3.8%, respectively, while TPR changes only modestly. In contrast, DeepSeek-R1 shows a negative filtering gain of -12.2%, because TPR increases by 9.1% while TNR-h decreases by 3.1%. DeepSeek-V3.2 also shows a negative filtering gain, indicating that CoT does not consistently improve hard-negative rejection. These results show that noisy evaluation can reveal different trade-offs between positive detection and hard-negative rejection that are obscured by aggregate accuracy or F1.

5.3 Aggregate Accuracy Masks Modality-Level Hard-Negative Rejection Failures

Figure 5 compares macro-averaged overall accuracy with macro-averaged TNR-h across sensory modalities. Visual cases achieve the highest mean accuracy of 75.6% but the lowest TNR-h of 48.7%, whereas gustatory cases have lower overall accuracy but the strongest TNR-h of 58.8%. Thus, aggregate accuracy can mask failures to reject plausible but insufficient evidence within a modality. These comparisons characterize reasoning over text-described sensory categories and do not isolate intrinsic sensory capabilities. Appendix C.6.1 shows that category-standardized clean Task 1 accuracy ranks gustatory above visual cues and that hard-negative rejection varies by source. Modality rankings therefore depend partly on item composition. Category standardization holds the proportions of positive, hard-negative, and easy-negative cases constant; it does not equalize source composition. Source-stratified rejection rates therefore complement the standardized comparison.

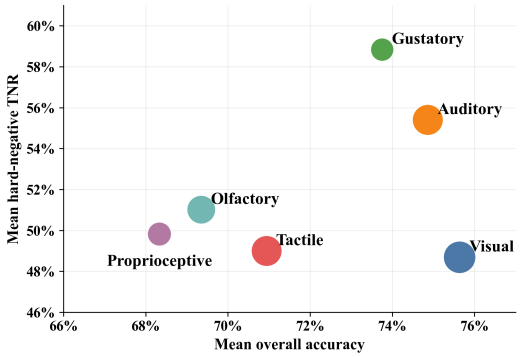

Figure 5: Modality-level comparison between overall accuracy and hard-negative rejection (TNR-h).

6 Conclusion

MuSeToM evaluates LLM Theory-of-Mind belief revision from text-described sensory cues through belief-change detection and updated-belief identification, in both clean and noisy settings. Models often accept sufficient evidence yet systematically over-predict revision when perceptually salient cues are plausible but epistemically insufficient. Noisy evaluation and modality-level diagnostics further show how aggregate performance can obscure failures to reject insufficient evidence. These findings highlight the need to distinguish perceptual salience from evidential sufficiency when evaluating belief revision in LLMs.

Limitations

Generated Perceptual Cues LLM-generated cues may differ stylistically from naturally occurring perceptual descriptions. Using multiple generator families and human validation aims to mitigate these artifacts: annotators assess logical consistency, plausibility, and epistemic appropriateness and remove unsuitable cues. These measures do not establish that generation artifacts are absent. The cue-text probes and cue-only LLM evaluations in Appendix A.4 provide additional diagnostics, while leaving possible partial reliance on cue artifacts unresolved.

Potential Contamination Most existing NLP benchmarks were released before the training of many recent LLMs, meaning that some benchmark data may have been seen during pre-training or post-training (Golchin and Surdeanu, 2024; Li and Flanigan, 2024). Since MuSeToM is constructed from BigToM and OpenToM, potential data contamination cannot be completely ruled out. However, MuSeToM substantially transforms the original benchmark instances into new multi-sensory belief-revision examples through generated perceptual cues, noisy distractors, updated-belief options, and human validation. We further conduct free-generation and partial-reconstruction contamination checks on sampled source instances. Across the evaluated models, none of the generated responses matched the sampled BigToM or OpenToM instances under our matching criteria. These results suggest that the sampled source stories used in MuSeToM are not identified as suffering from direct data contamination. Therefore, although contamination remains a possible limitation, it may not be a primary concern for MuSeToM. Details are provided in Appendix A.5.

Ethics Statement

In this work, we conformed to recognized privacy practices and rigorously followed the data usage policy. This paper introduces MuSeToM, a multi-sensory benchmark for assessing theory of mind belief revision in large language models. MuSeToM is constructed from existing theory of mind benchmarks, including BigToM and OpenToM, and further transforms their original stories into multi-sensory belief-revision instances through newly generated perceptual cues, answer options, noisy distractors, and human validation. We conduct human validation to ensure the logical consistency, plausibility, and epistemic appropriateness of the generated perceptual cues. In particular, annotators verify whether each cue provides sufficient evidence to trigger belief revision and whether the corresponding updated belief state is valid. We also filter out unreasonable, implausible, or potentially inappropriate instances during the validation process. We do not anticipate immediate harmful social consequences from the benchmark itself; however, because the benchmark is derived from existing datasets and generated cues, future work should audit MuSeToM for possible social or cultural bias and for artifacts introduced by cue generation or filtering. Moreover, the licenses and usage conditions of the source datasets allow us to use and adapt them for evaluation and research purposes. We follow the intended research use of these resources and will clearly acknowledge the original datasets.

Resource Copyright This work presents a new resource: MuSeToM, which is a multi-sensory extension of BigToM and OpenToM for evaluating perception-driven belief revision in theory of mind reasoning. Both source datasets are publicly available for research purposes, and we do not add any additional restrictions for accessing MuSeToM beyond the terms required by the original resources. We will highlight the sources of our data and ask users to cite the original BigToM and OpenToM papers when they use MuSeToM for research.

Acknowledgements

The authors of this paper were supported by the ITSP Platform Research Project (ITS/189/23FP) from ITC of Hong Kong, SAR, China, and the AoE (AoE/E-601/24-N), the RIF (R6021-20) and the GRF (16205322) from RGC of Hong Kong, SAR, China. We also thank the support from NVIDIA AI Technology Center (NVAITC) and Amazon.

References

Aristotle. 1984. De anima (on the soul). In Jonathan Barnes, editor, *The Complete Works of Aristotle: The Revised Oxford Translation*, volume 1, pages 641–692. Princeton University Press, Princeton, NJ. Translated by J. A. Smith.

Simon Baron-Cohen. 1991. Precursors to a theory of mind: Understanding attention in others. In Andrew Whiten, editor, *Natural Theories of Mind: Evolution,*

Development and Simulation of Everyday Mindreading, pages 233–251. Basil Blackwell, Cambridge, MA.

Simon Baron-Cohen, Alan M. Leslie, and Uta Frith. 1985. Does the autistic child have a "theory of mind"? *Cognition*, 21(1):37–46.

Lindsey Biel and Nancy K. Peske. 2005. *Raising a Sensory Smart Child: The Definitive Handbook for Helping Your Child with Sensory Integration Issues*. Penguin Books.

Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, Harsha Nori, Hamid Palangi, Marco Tulio Ribeiro, and Yi Zhang. 2023. Sparks of artificial general intelligence: Early experiments with GPT-4. *CoRR*, abs/2303.12712.

Chunkit Chan, Jiayang Cheng, Xin Liu, Yauwai Yim, Yuxin Jiang, Zheye Deng, Haoran Li, Yangqiu Song, Ginny Y. Wong, and Simon See. 2024a. Audience persona knowledge-aligned prompt tuning method for online debate. In *ECAI 2024 - 27th European Conference on Artificial Intelligence, 19-24 October 2024, Santiago de Compostela, Spain - Including 13th Conference on Prestigious Applications of Intelligent Systems (PAIS 2024)*, volume 392 of *Frontiers in Artificial Intelligence and Applications*, pages 3851–3858. IOS Press.

Chunkit Chan, Jiayang Cheng, Weiqi Wang, Yuxin Jiang, Tianqing Fang, Xin Liu, and Yangqiu Song. 2024b. Exploring the potential of chatgpt on sentence level relations: A focus on temporal, causal, and discourse relations. In *Findings of the Association for Computational Linguistics: EACL 2024, St. Julian's, Malta, March 17-22, 2024*, pages 684–721. Association for Computational Linguistics.

Chunkit Chan, Cheng Jiayang, Yauwai Yim, Zheye Deng, Wei Fan, Haoran Li, Xin Liu, Hongming Zhang, Weiqi Wang, and Yangqiu Song. 2024c. Negotiationtom: A benchmark for stress-testing machine theory of mind on negotiation surrounding. In *Findings of the Association for Computational Linguistics: EMNLP 2024, Miami, Florida, USA, November 12-16, 2024*, pages 4211–4241. Association for Computational Linguistics.

Chunkit Chan, Xin Liu, Tsz Ho Chan, Jiayang Cheng, Yangqiu Song, Ginny Y. Wong, and Simon See. 2023a. Self-consistent narrative prompts on abductive natural language inference. In *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics, IJCNLP 2023 -Volume 1: Long Papers, Nusa Dua, Bali, November 1 - 4, 2023*, pages 1040–1057. Association for Computational Linguistics.

Chunkit Chan, Xin Liu, Jiayang Cheng, Zihan Li, Yangqiu Song, Ginny Y. Wong, and Simon See. 2023b. Discoprompt: Path prediction prompt tuning for implicit discourse relation recognition. In *Findings of the Association for Computational Linguistics: ACL 2023, Toronto, Canada, July 9-14, 2023*, pages 35–57. Association for Computational Linguistics.

Chunkit Chan, Yauwai Yim, Hongchuan Zeng, Zhiying Zou, Xinyuan Cheng, Zhifan Sun, Zheye Deng, Kawai Chung, Yuzhuo Ao, Yixiang Fan, Cheng Jiayang, Ercong Nie, Ginny Y. Wong, Helmut Schmid, Hinrich Schütze, Simon See, and Yangqiu Song. 2025. XToM: Exploring the multilingual theory of mind for large language models. *CoRR*, abs/2506.02461.

Chunkit Chan, Yauwai Yim, Hongchuan Zeng, Zhiying Zou, Xinyuan Cheng, Zhifan Sun, Zheye Deng, Kawai Chung, Yuzhuo Ao, Yixiang Fan, Cheng Jiayang, Ercong Nie, Ginny Y. Wong, Helmut Schmid, Hinrich Schütze, Simon See, and Yangqiu Song. 2026. Xtom: Exploring the multilingual theory of mind for large language models. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026*, pages 17681–17716. Association for Computational Linguistics.

Zhanwen Chen, Tianchun Wang, Yizhou Wang, Michal Kosinski, Xiang Zhang, Yun Fu, and Sheng Li. 2024a. Through the theory of mind's eye: Reading minds with multimodal video large language models. *CoRR*, abs/2406.13763.

Zhuang Chen, Jincenzi Wu, Jinfeng Zhou, Bosi Wen, Guanqun Bi, Gongyao Jiang, Yaru Cao, Mengting Hu, Yunghwei Lai, Zexuan Xiong, and Minlie Huang. 2024b. ToMBench: Benchmarking theory of mind in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15959–15983, Bangkok, Thailand. Association for Computational Linguistics.

Jiayang Cheng, Lin Qiu, Tsz Ho Chan, Tianqing Fang, Weiqi Wang, Chunkit Chan, Dongyu Ru, Qipeng Guo, Hongming Zhang, Yangqiu Song, Yue Zhang, and Zheng Zhang. 2023. Storyanalogy: Deriving story-level analogies from large language models to unlock analogical understanding. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 11518–11537. Association for Computational Linguistics.

Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 4299–4307.

Kawai Chung, Chunkit Chan, Yuxuan Liu, Yauwai Yim, Xi Yang, Wenbin Hu, Huihao Jing, Chiu Yeung Tang,

Changxuan Fan, Ki Sen Hung, et al. 2026a. Jev-tom: Evaluating theory of mind in typed decision models. *arXiv preprint arXiv:2610.18610*.

Kawai Chung, Chunkit Chan, Yauwai Yim, Yuxuan Liu, Haochen Shi, Weiqi Wang, Qing Zong, Tianshi Zheng, Yixuan Fu, Kai Chung Wong, Hao Liang, Yifan Gao, Xi Yang, Janet Hui-wen Hsiao, and Yangqiu Song. 2026b. Multivationbench: A benchmark for multimodal sequential motivation reasoning. *CoRR*, abs/2607.26465.

DeepSeek-AI. 2024. Deepseek-v3 technical report. *CoRR*, abs/2412.19437.

DeepSeek-AI. 2025a. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *CoRR*, abs/2501.12948.

DeepSeek-AI. 2025b. DeepSeek-V3.1 model card. Official Hugging Face model card.

DeepSeek-AI. 2025c. Deepseek-v3.2: Pushing the frontier of open large language models. *CoRR*, abs/2512.02556.

Zheye Deng, Chunkit Chan, Weiqi Wang, Yuxi Sun, Wei Fan, Tianshi Zheng, Yauwai Yim, and Yangqiu Song. 2024. Text-tuple-table: Towards information integration in text-to-table generation via global tuple extraction. *CoRR*, abs/2404.14215.

Zheye Deng, Chunkit Chan, Tianshi Zheng, Wei Fan, Weiqi Wang, and Yangqiu Song. 2025. Structuring the unstructured: A systematic review of text-to-structure generation for agentic AI with a universal evaluation framework. *CoRR*, abs/2508.12257.

Joseph L. Fleiss. 1971. Measuring nominal scale agreement among many raters. *Psychological Bulletin*, 76(5):378–382.

Kanishk Gandhi, Jan-Philipp Fränken, Tobias Gerstenberg, and Noah D. Goodman. 2023a. Understanding social reasoning in language models with language models. In *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10–16, 2023, Datasets and Benchmarks Track*, volume 36, pages 13518–13529.

Kanishk Gandhi, Jan-Philipp Franken, Tobias Gerstenberg, and Noah D. Goodman. 2023b. Understanding social reasoning in language models with language models. *ArXiv*, abs/2306.15448.

Qiyang Gao, Xianjie Ping, and Wei Chen. 2019. Body influences on social cognition through interoception. *Frontiers in Psychology*, 10:2066.

Will M. Gervais. 2013. Perceiving minds and gods: How mind perception enables, constrains, and is triggered by belief in gods. *Perspectives on Psychological Science*, 8(4):380–394.

Shahriar Golchin and Mihai Surdeanu. 2024. Time travel in LLMs: Tracing data contamination in large language models. In *The Twelfth International Conference on Learning Representations, ICLR 2024*.

Erin Grant, Aida Nematzadeh, and Thomas L. Griffiths. 2017. How can memory-augmented neural networks pass a false-belief task? In *Proceedings of the 39th Annual Meeting of the Cognitive Science Society*, CogSci 2017 - Proceedings of the 39th Annual Meeting of the Cognitive Science Society: Computational Foundations of Cognition, pages 427–432. The Cognitive Science Society.

Laurent Itti. 2007. Visual salience. *Scholarpedia*, 2(9):3327.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de Las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Lélio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. 2023a. Mistral 7B. *CoRR*, abs/2310.06825.

Yuxin Jiang, Chunkit Chan, Mingyang Chen, and Wei Wang. 2023b. Lion: Adversarial distillation of proprietary large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 3134–3154. Association for Computational Linguistics.

Cheng Jiayang, Chunkit Chan, Qianqian Zhuang, Lin Qiu, Tianhang Zhang, Tengxiao Liu, Yangqiu Song, Yue Zhang, Pengfei Liu, and Zheng Zhang. 2024a. ECON: on the detection and resolution of evidence conflicts. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing, EMNLP 2024, Miami, FL, USA, November 12-16, 2024*, pages 7816–7844. Association for Computational Linguistics.

Cheng Jiayang, Lin Qiu, Chunkit Chan, Xin Liu, Yangqiu Song, and Zheng Zhang. 2024b. Eventground: Narrative reasoning by grounding to eventuality-centric knowledge graphs. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation, LREC/COLING 2024, 20-25 May, 2024, Torino, Italy*, pages 6622–6642. ELRA and ICCL.

Cheng Jiayang, Qianqian Zhuang, Haoran Li, Chunkit Chan, Xin Liu, Lin Qiu, and Yangqiu Song. 2025. Integround: On the evaluation of verification and retrieval planning in integrative grounding. In *Findings of the Association for Computational Linguistics: EMNLP 2025, Suzhou, China, November 4-9, 2025*, pages 13587–13602. Association for Computational Linguistics.

Chuanyang Jin, Yutong Wu, Jing Cao, Jiannan Xiang, Yen-Ling Kuo, Zhiting Hu, Tomer D. Ullman, Antonio Torralba, Joshua B. Tenenbaum, and Tianmin

Shu. 2024a. MMToM-QA: Multimodal theory of mind question answering. *CoRR*, abs/2401.08743.

Chuanyang Jin, Yutong Wu, Jing Cao, Jiannan Xiang, Yen-Ling Kuo, Zhiting Hu, Tomer David Ullman, Antonio Torralba, Joshua B. Tenenbaum, and Tianmin Shu. 2024b. Mmtom-qa: Multimodal theory of mind question answering. *ArXiv*, abs/2401.08743.

Chani Jung, Dongkwan Kim, Jiho Jin, Jiseon Kim, Yeon Seonwoo, Yejin Choi, Alice Oh, and Hyunwoo Kim. 2024. Perceptions to beliefs: Exploring precursory inferences for theory of mind in large language models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 19794–19809, Miami, Florida, USA. Association for Computational Linguistics.

Hyunwoo Kim, Melanie Sclar, Xuhui Zhou, Ronan Le Bras, Gunhee Kim, Yejin Choi, and Maarten Sap. 2023a. FANToM: A benchmark for stress-testing machine theory of mind in interactions. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 14397–14413.

Hyunwoo Kim, Melanie Sclar, Xuhui Zhou, Ronan Le Bras, Gunhee Kim, Yejin Choi, and Maarten Sap. 2023b. Fantom: A benchmark for stress-testing machine theory of mind in interactions. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 14397–14413. Association for Computational Linguistics.

David C. Knill and Alexandre Pouget. 2004. The bayesian brain: the role of uncertainty in neural coding and computation. *Trends in Neurosciences*, 27(12):712–719.

Michal Kosinski. 2023. Theory of mind may have spontaneously emerged in large language models. *CoRR*, abs/2302.02083. Version 1.

Matthew Le, Y-Lan Boureau, and Maximilian Nickel. 2019a. Revisiting the evaluation of theory of mind through question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5872–5877, Hong Kong, China. Association for Computational Linguistics.

Matthew Le, Y-Lan Boureau, and Maximilian Nickel. 2019b. Revisiting the evaluation of theory of mind through question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing, EMNLP-IJCNLP 2019, Hong Kong, China, November 3-7, 2019*, pages 5871–5876. Association for Computational Linguistics.

Changmao Li and Jeffrey Flanigan. 2024. Task contamination: Language models may not be few-shot anymore. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence*, pages 18471–18480.

Haoran Li, Yulin Chen, Jinglong Luo, Yan Kang, Xiaojin Zhang, Qi Hu, Chunkit Chan, and Yangqiu Song. 2023. Privacy in large language models: Attacks, defenses and future directions. *CoRR*, abs/2310.10383.

Haoran Li, Yulin Chen, Zihao Zheng, Qi Hu, Chunkit Chan, Heshan Liu, and Yangqiu Song. 2024a. Backdoor removal for generative large language models. *CoRR*, abs/2405.07667.

Haoran Li, Dadi Guo, Donghao Li, Wei Fan, Qi Hu, Xin Liu, Chunkit Chan, Duanyi Yao, Yuan Yao, and Yangqiu Song. 2024b. Privlm-bench: A multi-level privacy evaluation benchmark for language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2024, Bangkok, Thailand, August 11-16, 2024*, pages 54–73. Association for Computational Linguistics.

Yangning Li, Weizhi Zhang, Yuyao Yang, Wei-Chieh Huang, Yaozu Wu, Junyu Luo, Yuanchen Bei, Henry Peng Zou, Xiao Luo, Yusheng Zhao, Chunkit Chan, Yankai Chen, Zhongfen Deng, Yinghui Li, Hai-Tao Zheng, Dongyuan Li, Renhe Jiang, Ming Zhang, Yangqiu Song, and Philip S. Yu. 2025. Towards agentic RAG with deep reasoning: A survey of rag-reasoning systems in llms. *CoRR*, abs/2507.09477.

Fangzhou Liang, Tianshi Zheng, Chunkit Chan, Yauwai Yim, and Yangqiu Song. 2025. Llm-hanabi: Evaluating multi-agent gameplays with theory-of-mind and rationale inference in imperfect information collaboration game. *CoRR*, abs/2510.04980.

Zizheng Lin, Chunkit Chan, Yangqiu Song, and Xin Liu. 2024. Constrained reasoning chains for enhancing theory-of-mind in large language models. In *PRICAI 2024: Trends in Artificial Intelligence - 21st Pacific Rim International Conference on Artificial Intelligence, PRICAI 2024, Kyoto, Japan, November 18-24, 2024, Proceedings, Part II*, volume 15282 of *Lecture Notes in Computer Science*, pages 354–360. Springer.

Llama Team. 2024. The Llama 3 herd of models. *CoRR*, arXiv:2407.21783.

Ziqiao Ma, Jacob Sansom, Run Peng, and Joyce Chai. 2023. Towards a holistic landscape of situated theory of mind in large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 1011–1031, Singapore. Association for Computational Linguistics.

Bethany MacLeod. 2015. A critical evaluation of two approaches to defining perceptual salience. *Ampersand*, 2:83–92.

Meta Llama. 2025. Llama 4 model card. Official model card.

Microsoft. 2024. Phi-3 technical report: A highly capable language model locally on your phone. *CoRR*, abs/2404.14219.

MiniMax. 2025. MiniMax M2.1: Significantly enhanced multi-language programming, built for real-world complex tasks. Official release page.

MiniMax. 2026. MiniMax M2.5: Built for real-world productivity. Official release page.

Mistral AI. 2024a. Mistral-7B-Instruct-v0.3 model card. Official model card.

Mistral AI. 2024b. Model card for Mixtral-8x22B-Instruct-v0.1. Official model card.

Yunxiang Mo, Tianshi Zheng, Qing Zong, Jiayu Liu, Baixuan Xu, Yauwai Yim, Chunkit Chan, Jiaxin Bai, and Yangqiu Song. 2025. Dixitworld: Evaluating multimodal abductive reasoning in vision-language models with multi-agent dixit gameplay. *CoRR*, abs/2510.10117.

Moonshot AI. 2025. Kimi K2 Thinking model card. Official model card.

Aida Nematzadeh, Kaylee Burns, Erin Grant, Alison Gopnik, and Tom Griffiths. 2018. Evaluating theory of mind in question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2392–2400, Brussels, Belgium. Association for Computational Linguistics.

OpenAI. n.d. GPT-3.5 Turbo. Official model documentation. Accessed 2026-10-13.

OpenAI et al. 2024. Gpt-4o system card. *Preprint*, arXiv:2410.21276.

OpenAI et al. 2025. gpt-oss-120b & gpt-oss-20b model card. *Preprint*, arXiv:2508.10925.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, volume 35.

Chris Pratt and Peter Bryant. 1990. Young children understand that looking leads to knowing (so long as they are looking into a single barrel). *Child Development*, 61(4):973–982.

David Premack and Guy Woodruff. 1978. Does the chimpanzee have a theory of mind? *Behavioral and Brain Sciences*, 1(4):515–526.

Qwen Team. 2025a. Qwen3 technical report. *CoRR*, abs/2505.09388.

Qwen Team. 2025b. QwQ-32B: Embracing the power of reinforcement learning. Official release blog/-model description;.

Qwen Team. 2026. Qwen3.5-397B-A17B model card. Official Hugging Face model card;.

Thomas Raleigh and Filippo Vindrola. 2021. Perceptual experience and degrees of belief. *The Philosophical Quarterly*, 71(2):378–406.

Matthew Riemer, Zahra Ashktorab, Djallel Bouneffouf, Payel Das, Miao Liu, Justin D. Weisz, and Murray Campbell. 2025. Position: Theory of mind benchmarks are broken for large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 82091–82130. PMLR.

Jayanta Sadhu, Ayan Antik Khan, Noshin Nawal, Sanju Basak, Abhik Bhattacharjee, and Rifat Shahriyar. 2024. Multi-ToM: Evaluating multilingual theory of mind capabilities in large language models. *CoRR*, abs/2411.15999.

Melanie Sclar, Sachin Kumar, Peter West, Alane Suhr, Yejin Choi, and Yulia Tsvetkov. 2023. Minding language models' (lack of) theory of mind: A plug-and-play multi-character belief tracker. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13960–13980, Toronto, Canada. Association for Computational Linguistics.

Natalie Shapira, Mosh Levy, Seyed Hossein Alavi, Xuhui Zhou, Yejin Choi, Yoav Goldberg, Maarten Sap, and Vered Shwartz. 2024. Clever hans or neural theory of mind? stress testing social reasoning in large language models. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2257–2273, St. Julian's, Malta. Association for Computational Linguistics.

Haochen Shi, Tianshi Zheng, Weiqi Wang, Baixuan Xu, Chunyang Li, Chunkit Chan, Tao Fan, Yangqiu Song, and Qiang Yang. 2025a. Inferencedynamics: Efficient routing across llms through structured capability and knowledge profiling. *Preprint*, arXiv:2505.16303.

Haojun Shi, Suyu Ye, Xinyu Fang, Chuanyang Jin, Leyla Isik, Yen-Ling Kuo, and Tianmin Shu. 2025b. MuMA-ToM: Multi-modal multi-agent theory of mind. In *Proceedings of the Thirty-Ninth AAAI Conference on Artificial Intelligence*, pages 1510–1519.

STAR Institute for Sensory Processing. n.d. Your 8 senses. Accessed 13 October 2026.

James W. A. Strachan, Dalila Albergo, Giulia Borghini, Oriana Pansardi, Eugenio Scaliti, Saurabh Gupta, Krati Saxena, Alessandro Rufo, Stefano Panzeri, Guido Manzi, Michael S. A. Graziano, and Cristina Becchio. 2024. Testing theory of mind in large language models and humans. *Nature Human Behaviour*, 8(7):1285–1295.

Winnie Street, John Oliver Siy, Geoff Keeling, Adrien Baranes, Benjamin Barnett, Michael McKibben, Tatenda Kanyere, Alison Lentz, Blaise Agüera y Arcas, and Robin I. M. Dunbar. 2026. LLMs achieve adult human performance on higher-order theory of mind tasks. *Frontiers in Human Neuroscience*, 19:1633272.

Joshua B. Tenenbaum, Charles Kemp, Thomas L. Griffiths, and Noah D. Goodman. 2011. How to grow a mind: Statistics, structure, and abstraction. *Science*, 331(6022):1279–1285.

Tomer D. Ullman. 2023. Large language models fail on trivial alterations to theory-of-mind tasks. *CoRR*, abs/2302.08399.

Emilio Villa-Cueva, S. M. Masrur Ahmed, Rendi Chevi, Jan Christian Blaise Cruz, Kareem Elzeky, Fermin Cristobal, Alham Fikri Aji, Skyler Wang, Rada Mihalcea, and Thamar Solorio. 2025. MoMentS: A comprehensive multimodal benchmark for theory of mind. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 22591–22611. Association for Computational Linguistics.

Weiqi Wang, Tianqing Fang, Chunyang Li, Haochen Shi, Wenxuan Ding, Baixuan Xu, Zhaowei Wang, Jiaxin Bai, Xin Liu, Cheng Jiayang, Chunkit Chan, and Yangqiu Song. 2024. CANDLE: iterative conceptualization and instantiation distillation from large language models for commonsense reasoning. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2024, Bangkok, Thailand, August 11-16, 2024*, pages 2351–2374. Association for Computational Linguistics.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed H. Chi, Quoc V. Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, volume 35.

Alex Wilf, Sihyun Lee, Paul Pu Liang, and Louis-Philippe Morency. 2024. Think twice: Perspective-taking improves large language models' theory-of-mind capabilities. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8292–8308, Bangkok, Thailand. Association for Computational Linguistics.

Bryan Wilie, Samuel Cahyawijaya, Etsuko Ishii, Junxian He, and Pascale Fung. 2024. Belief revision: The adaptability of large language models reasoning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 10480–10496, Miami, Florida, USA. Association for Computational Linguistics.

Yufan Wu, Yinghui He, Yilin Jia, Rada Mihalcea, Yulong Chen, and Naihao Deng. 2023. Hi-ToM: A benchmark for evaluating higher-order theory of mind reasoning in large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10691–10706.

Yang Xiao, Jiashuo Wang, Qiancheng Xu, Changhe Song, Chunpu Xu, Yi Cheng, Wenjie Li, and Pengfei Liu. 2025. Towards dynamic theory of mind: Evaluating LLM adaptation to temporal evolution of human states. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24036–24057, Vienna, Austria. Association for Computational Linguistics.

Hainiu Xu, Runcong Zhao, Lixing Zhu, Jinhua Du, and Yulan He. 2024a. OpenToM: A comprehensive benchmark for evaluating theory-of-mind reasoning capabilities of large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8593–8623, Bangkok, Thailand. Association for Computational Linguistics.

Hainiu Xu, Runcong Zhao, Lixing Zhu, Jinhua Du, and Yulan He. 2024b. Opentom: A comprehensive benchmark for evaluating theory-of-mind reasoning capabilities of large language models. *ArXiv*, abs/2402.06044.

Yuqi Yang, Weiqi Wang, Baixuan Xu, Wei Fan, Qing Zong, Chunkit Chan, Zheye Deng, Xin Liu, Yifan Gao, Changlong Yu, Chen Luo, Yang Li, Zheng Li, Qingyu Yin, Bing Yin, and Yangqiu Song. 2025. Sessionintentbench: A multi-task inter-session intention-shift modeling benchmark for e-commerce customer behavior understanding. *CoRR*, abs/2507.20185.

Yauwai Yim, Chunkit Chan, Tianyu Shi, Zheye Deng, Wei Fan, Tianshi Zheng, and Yangqiu Song. 2024. Evaluating and enhancing llms agent based on theory of mind in guandan: A multi-player cooperative game under imperfect information. *CoRR*, abs/2408.02559.

Weizhi Zhang, Yangning Li, Yuanchen Bei, Junyu Luo, Guancheng Wan, Liangwei Yang, Chenxuan Xie, Yuyao Yang, Wei-Chieh Huang, Chunyu Miao, Henry Peng Zou, Xiao Luo, Yusheng Zhao, Yankai Chen, Chunkit Chan, Peilin Zhou, Xinyang Zhang, Chenwei Zhang, Jingbo Shang, Ming Zhang, Yangqiu Song, Irwin King, and Philip S. Yu. 2025. From web search towards agentic deep research: Incentivizing search with reasoning agents. *CoRR*, abs/2506.18959.

Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. 2024. Large language models are not robust multiple choice selectors. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net.

Qing Zong, Jiayu Liu, Tianshi Zheng, Chunyang Li, Baixuan Xu, Haochen Shi, Weiqi Wang, Zhaowei

Wang, Chunkit Chan, and Yangqiu Song. 2025. Critical: Can critique help LLM uncertainty or confidence calibration? *CoRR*, abs/2510.24505.

A Appendix for Data Construction and Validation

The appendix text is organized in Appendices A–F. All supporting tables and figures are collected after the text in Appendix G. Appendix A.1 describes cue construction and sensory categories, while Appendix A.2 covers annotation and disagreement. Appendix A.3 reports dataset statistics; Appendix A.4 presents cue-artifact diagnostics; and Appendix A.5 checks source-story contamination.

A.1 Appendix for Data Construction

Data Construction We construct MuSeToM by augmenting BigToM and OpenToM stories with perceptual cues and events. For BigToM, the generation prompts produce perceptual cues from the original causal-event and random-event fields. For OpenToM, the generation prompts produce belief-changing one-line events and unrelated random one-line events. The generated candidates cover six sensory modalities: visual, auditory, olfactory, gustatory, tactile, and proprioception. By generating perceptual cues from both belief-relevant and unrelated events, we obtain a balanced pool of candidate instances that vary in their relevance to the protagonist's initial belief. After human annotation and filtering, retained instances are labeled as positive, hard-negative, or easy-negative cues. Table 6 reports the final number of source stories, validated cue instances, and modality-specific cue counts after human filtering.

We note that BigToM random-event cue generation uses the same story–belief–event prompt structure as the causal-event setting, but replaces the causal event with the original BigToM random-event field and generates one cue per modality. Since the substituted random event is unrelated to the belief-changing event, the resulting cues are intended to be irrelevant to the belief-changing event, with final labels determined after human annotation and filtering.

Generation Prompt The generation prompt templates used for BigToM and OpenToM are shown in Tables 32 and 33. These prompts are used only for generating candidate cues and events; all retained instances are subsequently validated by human annotators.

A.1.1 Source-Specific and Noisy Instance Construction

OpenToM contrastive belief construction. OpenToM instances provide narrative contexts with two agents, an object, and inter-agent beliefs about object preferences. In MuSeToM, we use these cases as controlled contrastive belief-revision settings. The "negation" used during generation should therefore be understood as a target contrastive belief for preference-style beliefs, not as a claim that belief revision is always simple logical negation. Candidate OpenToM events are retained only when the resulting updated belief is natural in the story, coherent with the observed character's preference, and supported by the perceptual cue. During human validation, annotators verify both whether the cue warrants belief revision and whether the updated belief option is appropriate. Cases that require uncertainty, partial revision, or an unnatural contrastive belief are filtered out.

Noisy distractor construction. The noisy setting is designed to test epistemic relevance filtering rather than introduce arbitrary out-of-context noise. For each noisy example, a task-relevant cue from the validated positive or hard-negative pool is paired with an additional distractor cue sampled from the same story's easy-negative pool. These distractor cues are coherent with the story context but do not provide sufficient evidence for the queried belief revision. Distractor position is randomized across examples so that models cannot rely on cue order. The same answer-label randomization used in the clean setting is preserved in the noisy setting. As the added distractor is chosen from the easy-negative pool, it does not determine the answer by itself. This construction makes the noisy setting a test of whether a model can identify which perceptual evidence is relevant and whether that evidence is epistemically sufficient.

A.1.2 Sensory Category Definitions

Table 3 summarizes the six sensory systems used in MuSeToM, adapted from the STAR Institute's broader eight-sense framework (STAR Institute for Sensory Processing, n.d.).

A.2 Appendix for Human Annotation

We employ three human annotators to validate the generated MuSeToM instances. Annotators judge each perceptual cue with respect to two questions: whether the cue provides sufficient evidence for the

protagonist to revise the initial belief, and what the protagonist's updated belief should be when belief revision is warranted. Before the full annotation round, annotators complete an initial calibration batch of 100 instances. We inspect disagreements in this batch and provide feedback on common error types, especially cases where a cue is perceptually salient but not evidentially sufficient.

Operational criteria for cue labels. These labels provide a normative reference for whether the available evidence justifies definite belief revision under the stated context; they do not describe every individual's response to weak evidence. Table 4 summarizes the operational criteria. Annotator disagreement triggers adjudication when needed, but is neither necessary nor sufficient for a hard-negative label. Both initially unanimous no-change cases and disputed cases adjudicated as insufficient are retained when they meet the same criteria. The annotation and adjudication procedure is described below.

A positive cue supports the annotated revised belief. A hard negative bears on the target belief but remains indirect, ambiguous, or compatible with multiple explanations, so it does not support a definite revision. An easy negative does not bear on the target belief. Incoherent or implausible cues are excluded. Annotator disagreement is not a labeling criterion: retained hard negatives include both unanimous judgments and adjudicated cases, as reported in Appendix A.2.1. The hard-negative label concerns the insufficiency of the evidence, even when that insufficiency is unanimously judged. These rules operationalize evidential support without specifying a numerical probability threshold or requiring deductive entailment.

Rationale and scope of the shared standard. MuSeToM uses a shared sufficiency criterion to make categorical belief-revision judgments comparable across models. For a fixed story and cue, the evaluation asks whether the evidence available to the protagonist warrants the specified revised belief. Allowing unspecified personal thresholds would leave the reference answer dependent on unobserved factors such as prior confidence or risk preferences. Human validation and adjudication therefore establish a common operational reference for distinguishing warranted revision from responses to merely salient or relevant cues. Agreement under this protocol measures consistency in applying the reference; it does not establish a uniquely correct standard for human belief revision.

The tradeoff is that the benchmark does not characterize the full distribution of human updating behavior. For hard negatives, "no change" means that the cue does not sufficiently support categorical replacement of the stated belief; it does not rule out reduced confidence in the retained belief. Modeling heterogeneous thresholds would address the complementary objective of predicting individual responses and would require additional information about those thresholds or a reference distribution of judgments. MuSeToM scores should therefore be interpreted as performance relative to the stated evidence-based annotation criteria, rather than as predictions of how every individual would respond.

During final annotation, annotators apply the same operational criteria to all cues. Initially unanimous no-change cues that remain plausible and belief-relevant but insufficient are retained as hard negatives. Cases without an initial consensus are discussed and adjudicated using the same criteria; a disputed cue is retained as a hard negative when the annotators agree that it is plausible and relevant but insufficient for justified belief revision. Disagreement serves as a quality-control signal for adjudication and does not define the category. Unreasonable, implausible, or logically inconsistent cues, and cases that remain unresolved after discussion, are excluded. Table 7 reports Fleiss' κ for both annotation questions on BigToM and OpenToM. Figure 6 shows the instruction and the interface provided to annotators.

A.2.1 Annotation Disagreement and Sampling Uncertainty

We examine whether hard-negative errors are concentrated on cases that annotators initially disputed. Among the 477 retained hard negatives, 330, or 69.2%, received unanimous no-change judgments, while 147, or 30.8%, received one change vote out of three before adjudication. We compare model predictions across these two annotation strata.

Evaluation and aggregation. We use clean binary belief-change detection (Task 1) under direct and CoT prompting. For each item, we average its valid repetitions, then average items within each model–prompt condition, and finally give the 18 models $\times$ two prompting conditions equal weight. Invalid responses are excluded from the denominators. This Task 1-only, item-weighted analysis is distinct from the clean direct-prompting diagnostic

in Figure 3, which pools valid answer slots across Tasks 1 and 2.

We estimate a percentile 95% confidence interval for the disputed-minus-unanimous gap using 10,000 bootstrap resamples of source-story groups. We first average each item's error rate across model–prompt conditions, then retain all hard-negative cues belonging to each sampled story together. This interval quantifies story-sampling uncertainty for the fixed evaluated model cohort.

Table 5 shows a higher change-prediction rate of 58.9% on disputed hard negatives versus 43.3% on unanimous hard negatives, a difference of 15.6 percentage points, with a story-clustered 95% CI of 10.7–20.6. The difference has the same direction in 32 of 36 model–prompt conditions. Initial disagreement is associated with higher difficulty, but substantial over-updating remains on the unanimously labeled subset. This comparison does not establish a universal human threshold for belief revision.

Qualitative analysis of disputed cases. We examine three retained hard negatives with recorded binary judgments of one change vote and two no-change votes, and final gold labels of no change. The cases illustrate different gaps between a salient cue and the specific revised belief. The available records preserve votes and final labels, but not written adjudication rationales; the competing readings below are retrospective applications of the annotation criteria, not accounts of individual annotators' reasoning.

Linking a cue to its referent (BigToM, 0-0-1; visual). Noor believes that a pitcher contains oat milk. The cue describes an "Almond Milk" label on a nearby carton. One reading links the carton to the pitcher's contents and favors changing the belief to almond milk. Another distinguishes the observed carton from the pitcher: the cue does not establish that milk from this carton entered the pitcher. Under the annotation criteria, the label is relevant but does not sufficiently support the specific replacement belief. The final gold answer retains the oat-milk belief.

Inferring a cause from an effect (BigToM, 2-0-2; olfactory). Carlos believes that a water valve is closed. He smells decaying plant matter and rain-soaked earth from a flooded field. Flooding may suggest an open valve, but the described odor does not distinguish water released through that valve from water accumulating through rainfall or another route. The cue therefore bears on the irrigation situation without identifying the valve's state. The final gold answer retains the closed-valve belief.

Separating past and current preferences (OpenToM, 1-0-4; auditory). Paxton believes that Hadley dislikes a hat. As Hadley is about to throw it away, Paxton hears Hadley say fondly, "This used to be my favorite." The affectionate tone may suggest liking, whereas "used to" locates the preference in the past and is compatible with present dislike. The cue does not establish that Hadley currently likes the hat. The final gold answer retains Paxton's belief that Hadley dislikes it.

These cases illustrate how relevance, causal plausibility, or affective salience can invite a change judgment without sufficiently supporting the proposed revised belief. They show possible sources of disagreement under the same operational criteria, without estimating how frequently each occurs across the disputed subset.

A.3 Appendix for Dataset Statistics and Annotation Agreement

Table 6 summarizes the number of validated MuSeToM instances by source dataset, belief-change category, and sensory modality. Table 7 reports inter-annotator agreement for the two validation questions used to decide belief-change labels and updated belief states.

A.4 Appendix for Cue-Text and Cue-Only Validation

We assess label predictability without the surrounding story using all 1,489 positive and hard-negative cues, comprising 1,012 positive and 477 hard-negative cues grouped into 220 distinct story-text groups. We use five-fold stratified cross-validation grouped by story: all cues from a story remain in the same fold, and preprocessing is fitted on training folds only. A metadata probe uses source dataset, sensory category, word count, and character count. A cue-text probe combines word unigram/bigram and character 3–5-gram TF–IDF features. Both use class-balanced logistic regression and receive no story context, answer options, or gold-answer-derived features. Metrics are computed from pooled out-of-fold predictions.

Table 8 shows that the cue-text probe achieves 67.8% balanced accuracy, compared with 60.7%

for metadata and length. This result indicates that cue wording carries label information, but it does not distinguish unintended generation artifacts from the intended semantic difference in evidential strength. Nor does it establish whether the evaluated LLMs exploit these features. The probe therefore characterizes cue-text predictability without ruling out artifact-based shortcuts.

Cue-only LLM evaluation. We additionally evaluate MiniMax-M2.5 and Kimi-K2-Thinking with the story context and the protagonist's initial belief removed.

Table 9 shows lower cue-only performance for both models. Accuracy drops by 30.12 and 37.05 percentage points on Tasks 1 and 2 for MiniMax-M2.5, and by 19.65 and 37.69 points for Kimi-K2-Thinking. Thus, cue-only inputs do not recover full-context performance for either model. This complements the TF–IDF probe: cue wording carries label information, while the removed context contributes substantially to LLM performance. These results do not rule out partial reliance on cue artifacts.

A.5 Appendix for Verification of Potential Contamination

Since MuSeToM is constructed from existing ToM benchmarks, namely BigToM and OpenToM, it may be exposed to potential data contamination risks. These source datasets were released before the training of many recent LLMs and therefore could have appeared during either pre-training or post-training, including supervised fine-tuning (Ouyang et al., 2022) or RLHF (Christiano et al., 2017). To assess this risk, we follow two established contamination-detection protocols. First, following Li and Flanigan (Li and Flanigan, 2024), we use a free-generation check that asks a model to generate an example and answer from a named source benchmark. Second, following Golchin and Surdeanu (Golchin and Surdeanu, 2024), we use a partial-reconstruction check that asks a model to reconstruct the next sentence from the first sentence of a sampled source instance. For each source dataset, evaluated model, and protocol, we run 100 trials and compare model outputs against the corresponding ground-truth source instances. All responses are screened by automatic matching and then manually reviewed under the strong, likely, and possible match criteria reported in Table 10. The prompt templates are shown in Table 38. Across BigToM and OpenToM, none of the evaluated LLMs' generated responses match the sampled source instances under either protocol. This suggests that the sampled source stories used in MuSeToM are not identified as suffering from direct data contamination. Moreover, MuSeToM transforms the original stories into multi-sensory belief-revision instances through newly generated perceptual cues, noisy distractors, updated-belief options, and human validation, which further reduces the likelihood that models can solve the benchmark by memorizing the original datasets alone. Therefore, although data contamination cannot be completely ruled out, it may not be a primary concern for MuSeToM.

B Appendix for Evaluation Setup and Human Reference

B.1 Appendix for Hyperparameters

We use the following hyperparameters for assessing the large language models mentioned in this paper. For all Task 1–4 evaluations on MuSeToM, we use `temperature=0` and `top_p=1.0`. For Llama-3.1-8B-Instruct (Llama Team, 2024), Llama-3.1-405B-Instruct (Llama Team, 2024), Llama-4-Maverick-17B-128E-Instruct (Meta Llama, 2025), Mistral-7B-Instruct-v0.3 (Mistral AI, 2024a), Mixtral-8x22B-Instruct-v0.1 (Mistral AI, 2024b), QwQ-32B (Qwen Team, 2025b), Qwen3-235B-A22B (Qwen Team, 2025a), GPT-OSS-120B (OpenAI et al., 2025), GPT-4o (OpenAI et al., 2024), and GPT-3.5-Turbo (OpenAI, n.d.), we set `max_tokens=1024`. For DeepSeek-V3.1 (DeepSeek-AI, 2024, 2025b), DeepSeek-V3.2 (DeepSeek-AI, 2025c), DeepSeek-R1 (DeepSeek-AI, 2025a), MiniMax-M2.1 (MiniMax, 2025), MiniMax-M2.5 (MiniMax, 2026), Kimi-K2-Thinking (Moonshot AI, 2025), and Qwen3.5-397B-A17B (Qwen Team, 2026), we set `max_tokens=10240`.

B.2 Appendix for Evaluation Prompting

Appendix F.2 indexes the full evaluation prompt templates. In this study, there are two types of prompting methods. The first one is zero-shot prompting, and we use the prompt templates shown in Tables 34–37. The second one is the chain-of-thought (CoT) prompting method, which follows Wei et al. (2022) and uses the prompt "Let's think step by step." The CoT prompting template is formed by appending "Let's think step by step." to the zero-shot prompt template.

B.3 Appendix for Human Reference Evaluation

To provide a small-sample human reference on MuSeToM, we employ three graduate-student evaluators to complete the same benchmark evaluation tasks as the models. These evaluators are distinct from the three annotators who constructed the dataset and produced the gold labels. Each benchmark instance is evaluated under four task settings: clean belief-change detection, clean updated-belief identification, noisy belief-change detection, and noisy updated-belief identification. The questions and instructions are identical to the templates indexed in Appendix B.2. Evaluators use the annotation interface shown in Figure 6. We compute the final human predictions by majority vote over the three evaluators' responses. This results in 94.29% accuracy and 94.28% macro-F1 on clean belief-change detection, 93.83% accuracy and 93.83% macro-F1 on clean updated-belief identification, 90.76% accuracy and 90.76% macro-F1 on noisy belief-change detection, and 90.02% accuracy and 89.95% macro-F1 on noisy updated-belief identification, as reported in Table 1.

Human reference by sensory modality. Table 11 provides a human reference for interpreting the modality-level model results. We group the 2,415 clean items by the stored sensory label of their perceptual cue. Task 1 pools the three judgments per item using the stored vote counts; Task 2 scores each of the three evaluations separately and averages their accuracy and macro-F1. The response-level scores, individual results, and agreement statistics below all use these same clean-item records. They are reported separately from the aggregate majority-vote reference above. Macro-F1 uses binary labels for Task 1 and canonical belief-state labels for Task 2. Within each modality, we include classes present in the gold labels or any human prediction, using the same class set across evaluations and excluding classes absent from both. Overall scores are recomputed across all items, rather than averaged across modalities.

Individual scores and response agreement. Table 12 uses the same records as Table 11. Panel (a) reports the three recorded evaluations' individual Task 2 scores and sample variance. Task 1 retains vote counts without evaluator identities, so individual Task 1 scores cannot be recovered. Panel (b) reports agreement separately for each sensory category using Fleiss' κ (Fleiss, 1971) and mean pairwise agreement. Each item contributes three response pairs; pairwise agreement is the fraction of matching pairs. We compute $\kappa = (\bar{P} - P_e)/(1 - P_e)$, where $\bar{P}$ is mean pairwise agreement and P_e is chance agreement from the group's response proportions. Task 2 options are mapped to canonical belief states before pooling the two source datasets. Overall agreement is recomputed from all items.

Agreement is lowest for olfactory cues, with $\kappa = 0.569$ for Task 1 and 0.557 for Task 2; overall agreement is $\kappa = 0.758$ and 0.743, respectively. These statistics characterize response agreement on the clean Tasks 1–2 records used in Table 11; the individual scores and sample variance in Table 12 describe the same three recorded evaluations. This small sample does not establish population-level human variability.

Confidence intervals and paired comparisons. Table 13 reports uncertainty for the response-level human reference in Table 11 and comparisons with all 18 direct-prompting models. We average each model's valid repeated answers within each cue item, then weight items equally. Each comparison uses the same items for the human and model scores: 2,388 for DeepSeek-R1 and 2,415 for every other model. These comparisons use the response-level aggregation described here and complement the aggregate results in Table 1.

We compute pointwise 95% percentile confidence intervals using 10,000 bootstrap replicates, resampling source stories separately within BigToM and OpenToM. The 2,415 records contain 221 story clusters, comprising 100 BigToM and 121 OpenToM stories. All cues and responses from a sampled story remain together, and each human–model difference uses identical resampling weights. Human accuracy and macro-F1 follow the scoring in Table 11, with the overall label set fixed across replicates. For significance testing, accuracy is the primary endpoint. We use 99,999 two-sided story-level paired randomizations, swapping the human and model item-accuracy vectors jointly within each story. This test assumes exchangeability of the paired story-level outcomes under the null. We calculate Monte Carlo p-values as $(b+1)/100{,}000$, where b is the number of randomized absolute differences at least as large as the observed difference, and apply Holm correction across all 36 model–task comparisons.

All paired accuracy intervals are above zero. The smallest gaps are 7.72 percentage points against MiniMax-M2.5 on Task 1, with a 95% CI of 6.05–9.37, and 18.29 points against Kimi-K2-Thinking on Task 2, with a 95% CI of 15.94–20.69; all 36 comparisons have Holm-adjusted $p < 0.001$. These results quantify uncertainty over source stories for the observed three human evaluations and model responses. They do not estimate variability across a broader human population, and the small, homogeneous evaluator sample remains a limitation.

C Appendix for Supplementary Diagnostic Analyses

The following analyses examine sampling uncertainty, evidence discrimination within matched contexts, the belief states underlying errors, sensitivity to prompting, qualitative errors, and modality composition. Each subsection specifies its evaluation sample and aggregation. Appendix E indexes the full model and task-level results.

C.1 Appendix for Uncertainty in the Principal Diagnostic Contrasts

Table 14 reports source-story-clustered 95% confidence intervals for the three principal diagnostic contrasts. We use 10,000 bootstrap replicates clustered by 221 source stories. Within each replicate, we recompute each contrast separately for every model–prompt condition, then macro-average across all 36 conditions. These intervals quantify source-story sampling uncertainty for the fixed evaluated model–prompt cohort.

The intervals indicate substantial hard-negative overprediction and positive gaps between easy- and hard-negative rejection, and between positive detection and hard-negative rejection. They support the aggregate within-condition patterns under source-story sampling variation; they do not establish the statistical significance of small differences between individual models or their rankings.

C.2 Appendix for Evidence Discrimination within the Same Context

We group clean Task 1 items by source, story, initial belief, event, and sensory modality, retaining groups with positive (P), hard-negative (HN), and easy-negative (EN) cues for all 18 models under direct and CoT prompting. This yields 138 groups from 114 distinct story texts, comprising 74 BigToM and 64 OpenToM groups. We exclude the 21 items in the full dataset with inconsistent binary and multiple-choice gold labels. Each item's change-response rate averages its valid repetitions; rates are then averaged over cues of the same class within a group. A P–negative pair is correct only if the P rate exceeds 0.5 and the negative rate is below 0.5; ties count as failures. We average group scores within each story, then give stories and models equal weight.

Table 15 shows lower joint correctness for P–HN than P–EN under both prompts, even when the story, initial belief, event, and modality are held fixed. CoT increases P–HN joint correctness by 9.4 percentage points, with improvements in 14 of 18 models. Its positive change-response rate remains nearly unchanged, while both negative rates decrease. Excluding groups tied on either member of a pair preserves the pattern: P–HN correctness is 32.5%/44.0% for direct/CoT, compared with 58.7%/70.1% for P–EN. These descriptive results concern the complete-contrast subset and do not replace full-benchmark item accuracy or identify the mechanism behind the prompting difference.

C.3 Appendix for Distinguishing Belief Replacement from Uncertainty

We map existing Task 2 options to four belief states: retaining the prior belief (Prior), adopting the target updated belief (Target), character uncertainty (Unc.), and another belief (Other). The uncertainty option is identified from its text, accounting for its different canonical position in BigToM and OpenToM. For each model, we retain HN items with valid outputs in all four clean/noisy $\times$ direct/CoT conditions and report the clean Task 2 results. The union contains 477 items, comprising 246 BigToM and 231 OpenToM items. Each model covers all 246 BigToM items; OpenToM coverage is 216 for DeepSeek-R1, 228 for MiniMax-M2.1, and 231 for each other model. We average state frequencies over valid repetitions within each item, then equally over items and the 18 models. Invalid repetitions are excluded.

Table 16 separates concrete belief replacement from uncertainty responses. Under direct prompting, target-belief responses are frequent in both sources, while uncertainty is more common in BigToM at 39.2% than in OpenToM at 11.7%. Across sources, CoT reduces target responses by 6.02 percentage points, alongside increases of 4.45 percentage points in prior retention, 0.94 in uncertainty, and 0.62 in other beliefs. These are changes in

aggregate response shares, not transitions traced through individual generations. The source contrast also includes differences in story content and question type. Thus, HN errors should not all be interpreted as categorical belief reversal: a substantial share instead attributes uncertainty to the protagonist.

C.4 Appendix for Prompt Wording and an Uncertainty Option

We test sensitivity to change-focused wording and binary answers in a 2×2 experiment crossing original versus neutral wording with binary versus ternary response options. We select 30 source stories, each with a positive, hard-negative, and easy-negative cue matched within story and sensory category, yielding 90 items. The sample includes 15 stories from each source and 15 unanimous versus 15 disputed hard negatives. DeepSeek-V3.2, Llama-4-Maverick, MiniMax-M2.5, and Qwen3.5-397B-A17B each answer all 90 items in all four conditions, yielding 1,440 valid responses. Each model–item–condition receives one valid response, with temperature set to zero and answer order shuffled using a fixed seed separately within each condition.

Prompt variants. The original question asks, "Did the event cause the main character [name] to change their belief?" The neutral question asks, "Based only on the information available to [name], which option best describes whether [name]'s belief changed?" Original binary options are Yes/No; neutral options explicitly state changed/did not change. Ternary variants add an option stating that the available information is insufficient to determine whether the character changed their belief. This option permits an indeterminate judgment about change; it does not assign hard negatives a new gold label or represent the character's own uncertainty about the underlying event. The comparison varies question and answer wording together, with answer positions randomized separately across conditions.

Table 17 shows that change predictions depend on wording and response format. Under the original wording, adding the uncertainty option lowers the hard-negative change rate from 41.7% to 35.0%, while also lowering positive-cue change predictions from 76.7% to 59.2%. The additional option therefore affects responses beyond hard negatives.

With neutral wording and ternary options, the hard-negative/easy-negative gap remains 19.2 percentage points, with a 95% CI of 9.2–29.2. All four models retain a positive gap in this condition, ranging from 16.7 to 23.3 points. These results indicate that wording and response options affect the magnitude of over-updating, while the hard-negative/easy-negative difference persists within this sample. The deliberately balanced subset and single response per condition limit generalization to other prompts, models, and repeated generations.

C.5 Appendix for Qualitative Error Analysis

Table 18 summarizes examples of the three cue classes. Appendix A.2.1 separately analyzes recorded human disagreement; the examples here illustrate model predictions.

C.5.1 Appendix for Over-Updating on a Hard Negative

Table 19 shows a representative hard-negative instance from MuSeToM. The perceptual cue is salient and related to the target belief, but it does not establish an evidential link between the nearby almond-milk label and the milk already inside the pitcher. The gold label therefore preserves the protagonist's initial belief. The model instead treats the label as sufficient evidence for belief revision, illustrating the salience-sufficiency confusion observed in Section 5.1.

C.6 Appendix for Modality Diagnostics

Descriptive modality summary. Table 20 provides the descriptive modality-level diagnostics. The composition-controlled analysis below uses its separately specified clean Task 1 cohort and aggregation.

C.6.1 Category-Standardized and Source-Stratified Modality Analysis

Table 21 compares observed and category-standardized accuracy, alongside source-specific hard-negative rejection, using existing clean Task 1 outputs from all 18 models under direct and CoT prompting. This diagnostic includes 2,394 items, excluding 21 BigToM items whose binary gold labels disagree with their cue categories and updated-belief gold labels. We first average valid repetitions within each item, then weight items equally within each model–prompt condition, and finally average the 36 conditions equally. Scores are conditional on available valid responses. This cohort and aggre-

gation define a supplementary analysis; the values do not replace the modality summaries in the main results.

For each modality, standardized accuracy combines its positive, hard-negative, and easy-negative accuracies using the same fixed proportions, $(991, 477, 926)/2{,}394$, from the eligible benchmark. This adjustment controls category composition; it does not equalize source composition. We therefore also report TNR-h separately for BigToM and OpenToM.

Category standardization changes the observed ordering: visual accuracy decreases from 76.26% to 74.09%, while gustatory accuracy increases from 75.75% to 77.48%. Thus, aggregate modality differences partly reflect category composition. Source-specific rates also vary: olfactory TNR-h is 44.71% on BigToM and 57.83% on OpenToM. Small source–modality cells warrant caution; for example, OpenToM has only six proprioceptive hard negatives. Since all cues are expressed as text and story content remains uncontrolled, these comparisons do not isolate intrinsic sensory capabilities.

D Appendix for Related Works

D.1 Related Works for Theory of Mind

Theory of Mind (ToM) evaluation draws on psychological paradigms such as the Sally-Anne false-belief test. Early computational benchmarks operationalized false-belief reasoning as question answering over synthetic narratives (Grant et al., 2017; Nematzadeh et al., 2018; Le et al., 2019b). Subsequent work broadened evaluation to richer social narratives through BigToM and OpenToM (Gandhi et al., 2023b; Xu et al., 2024b), diverse social cognition abilities through ToMBench (Chen et al., 2024b), and higher-order and temporally evolving mental states through HiToM and DynToM (Wu et al., 2023; Xiao et al., 2025). These benchmarks extend the contexts and reasoning demands of belief attribution beyond isolated false-belief questions.

Interactive settings introduce further challenges. FANToM and NegotiationToM assess mental-state reasoning in conversations and negotiations (Kim et al., 2023b; Chan et al., 2024c), while Guandan and Hanabi connect ToM and rationale inference to cooperation under imperfect information (Yim et al., 2024; Liang et al., 2025). Complementary approaches investigate constrained reasoning over related mental-state dimensions (Lin et al., 2024) and typed decision models for mental-state selection (Chung et al., 2026a). Nevertheless, benchmark success does not necessarily imply effective adaptation to other agents, motivating a distinction between predicting mental states and using such predictions in interaction (Riemer et al., 2025).

Evaluation has also expanded across languages and input modalities. XToM and Multi-ToM examine multilingual ToM (Chan et al., 2026; Sadhu et al., 2024), while MMToM-QA, ToM's Eye, MuMA-ToM, and MoMentS investigate mental-state reasoning from multimodal observations and social contexts (Jin et al., 2024b; Chen et al., 2024a; Shi et al., 2025b; Villa-Cueva et al., 2025). Related benchmarks examine evolving motivations and intention shifts (Chung et al., 2026b; Yang et al., 2025). Although these constructs differ from belief, they share an emphasis on interpreting agents from contextual evidence. MuSeToM uses textual descriptions spanning six sensory categories, rather than raw visual, audio, or video inputs, to evaluate reasoning about the implications of perceptual evidence.

Perception-oriented and perspective-taking approaches are particularly close to our setting. Perceptions to Beliefs evaluates perception inference and perception-to-belief inference (Jung et al., 2024), SimToM filters context according to a character's perspective (Wilf et al., 2024), and SymbolicToM explicitly tracks multi-character belief states (Sclar et al., 2023). These approaches address what an agent can access or believe. MuSeToM instead isolates whether newly accessible evidence is sufficient to revise a prior belief, distinguishing the availability of a cue from its epistemic force.

Comparison with Belief-R. Belief-R evaluates whether language models revise their own conclusions after receiving new evidence (Wilie et al., 2024), whereas MuSeToM evaluates third-person belief revision: whether another agent should update a belief and what the updated belief should be. Table 22 summarizes these complementary evaluation targets. MuSeToM's hard and easy negatives distinguish relevant but insufficient evidence from irrelevant evidence, enabling diagnosis of over-updating on plausible, perceptually salient cues. This distinction does not imply that belief revision or resistance to unnecessary updates is absent from Belief-R; it identifies MuSeToM's focus on the suf-

ficiency of evidence available to a protagonist.

D.2 Related Works for LLMs

Research on LLMs has explored capability development through knowledge distillation and model routing (Jiang et al., 2023b; Shi et al., 2025a), alongside evaluations and methods for task-specific reasoning. These include temporal, causal, and discourse understanding (Chan et al., 2024b, 2023b); abductive and analogical reasoning in textual and multimodal settings (Chan et al., 2023a; Cheng et al., 2023; Mo et al., 2025); and knowledge-grounded narrative and commonsense reasoning (Jiayang et al., 2024b; Wang et al., 2024). Together, these directions examine how models interpret relations, construct explanations, and draw on background knowledge.

Reliable information use introduces further challenges, including evidence-conflict resolution (Jiayang et al., 2024a), verification and retrieval-augmented reasoning (Jiayang et al., 2025; Zhang et al., 2025; Li et al., 2025), and uncertainty and confidence calibration (Zong et al., 2025). Related studies investigate structured information integration (Deng et al., 2024, 2025) and audience-aware argument assessment (Chan et al., 2024a). Beyond reasoning, privacy evaluation and defenses against privacy leakage and backdoors address complementary deployment risks (Li et al., 2023, 2024b,a).

MuSeToM focuses on a specific aspect of reasoning reliability: whether textual descriptions of perceptual evidence justify revising another agent's belief. Unlike assessing explanatory plausibility, persuasive impact, or confidence alone, this requires distinguishing perceptual salience from epistemic sufficiency. By separately evaluating belief-change detection and updated-belief identification, MuSeToM examines whether models can both reject insufficient cues and infer the belief warranted by sufficient evidence.

E Appendix for Full Model and Task-Level Results

E.1 Appendix for Overall Results and Full-Model Diagnostic

Table 23 reports the full direct- and CoT-prompting results corresponding to the condensed main table. Figure 7 provides the full-model version of the diagnostic figure used in Section 5.1. The figure uses canonical belief-change decisions for clean direct-prompting examples: for multiple-choice questions, option 0 is treated as no belief change and all other canonical options are treated as belief change.

E.2 Appendix for Detailed Task-Level Results

Tables 24–27 report dataset-specific direct-prompting results for Tasks 1–4, and Tables 28–31 report the corresponding CoT-prompting results. Task 1 and Task 2 include positive, hard-negative, and easy-negative cases, so they report TPR, TNR-h, TNR-e, accuracy, and macro-F1. Task 3 and Task 4 use noisy inputs with positive and hard-negative cases, so they report TPR, TNR-h, accuracy, and macro-F1.

F Appendix for Prompt Templates

F.1 Cue and Event Generation

Tables 32 and 33 contain the candidate-generation templates. Construction and validation are described in Appendix A.

F.2 Benchmark Evaluation

Tables 34–37 contain the four zero-shot evaluation templates. The CoT setting appends "Let's think step by step." to these templates; Appendix C.4 specifies the variants used in the wording and uncertainty diagnostic.

F.3 Source-Story Contamination Checks

Table 38 gives the prompts for the existing source-story checks reported in Appendix A.5.

G Supplementary Tables and Figures

The following tables and figures support the preceding appendix text.

Sensory System	Definition
Visual	Processes visual input, including information about color, shape, orientation, motion, object identity, and object location.
Auditory	Processes sound input, including hearing, sound frequency, changes in volume or pitch, and language-related auditory information.
Olfactory	Processes smell input from the nose, supporting odor detection, odor discrimination, and filtering of background smells.
Gustatory	Processes taste input, helping distinguish food-related qualities such as sweet, salty, sour, and bitter tastes.
Tactile	Processes touch information from the body, including touch, pressure, temperature, and pain.
Proprioceptive	Senses body position and movement through muscles and joints, including the relative position of body parts and the effort used to move them.

Table 3: Definitions of the six sensory systems, adapted from the STAR Institute for Sensory Processing (STAR Institute for Sensory Processing, n.d.).

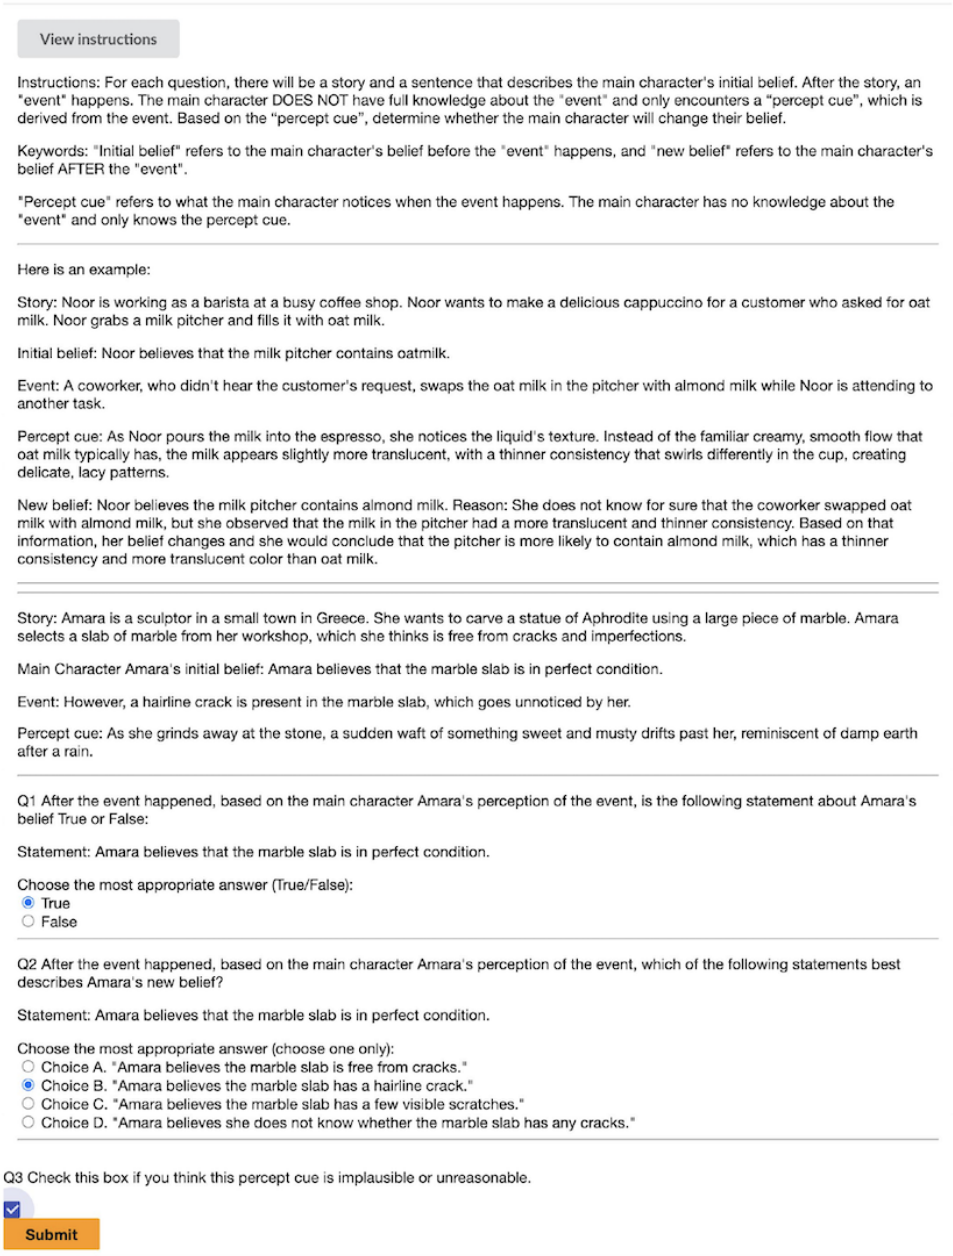

Figure 6: Human annotation interface. The figure shows the interface and instructions for annotators.

Class	Coherent and available	Relevant to target belief	Supports definite revision
Positive	Yes	Yes	Yes
HN	Yes	Yes	No
EN	Yes	No	No

Table 4: Operational cue-label criteria. HN and EN denote hard and easy negatives. Coherence is judged against the story, and availability against the protagonist's perspective.

HN subset	n	Changed (%)
All	477	48.1
Unanimous	330	43.3
Disputed and adjudicated	147	58.9

Table 5: Task 1 change predictions by initial annotation agreement, averaged across 36 model–prompt conditions. The disputed-minus-unanimous gap is 15.6 percentage points, with a story-clustered 95% CI of 10.7–20.6.

Dataset	Stories	Total	Pos.	Hard Neg.	Easy Neg.
BigToM	100	1,169	501	246	422
OpenToM	120	1,246	511	231	504
Total	220	2,415	1,012	477	926

Dataset	Vis.	Aud.	Olf.	Gus.	Tac.	Prop.
BigToM	349	259	135	73	247	106
OpenToM	401	327	174	80	198	66
Total	750	586	309	153	445	172

Table 6: Dataset statistics of MuSeToM after human validation and filtering. The upper block reports source and category statistics, and the lower block reports the distribution of perceptual cues across six sensory modalities.

Annotation Question	BigToM	OpenToM
Q1: Belief Change Detection	0.776	0.742
Q2: Updated Belief Identification	0.744	0.742

Table 7: Inter-annotator agreement measured by Fleiss' κ.

Probe	Balanced acc.	Macro-F1
Majority baseline	50.0	40.5
Stratified-random baseline	50.9	50.9
Metadata and cue length	60.7	59.1
Cue-text TF–IDF	67.8	67.8

Table 8: Positive versus hard-negative label prediction under five-fold source-story-grouped evaluation. Values are percentages over 1,489 held-out predictions.

Input	Task 1		Task 2	
	Acc.	F1	Acc.	F1
MiniMax-M2.5				
Full context	86.05	86.05	74.10	74.12
Cue-only	55.93	55.93	37.05	37.06
Kimi-K2-Thinking				
Full context	65.50	67.41	75.39	75.40
Cue-only	45.85	47.19	37.70	37.70

Table 9: Full-context and cue-only LLM performance (%) on belief-change detection (Task 1) and updated-belief identification (Task 2). Cue-only inputs omit the story context and the protagonist's initial belief.

Dataset	Model	Prompt Type	Trials	Strong	Likely	Possible
BigToM	DeepSeek-V3.2	Free Generation	100	0	0	0
BigToM	DeepSeek-V3.2	Partial Reconstruction	100	0	0	0
BigToM	Kimi-K2-Thinking	Free Generation	100	0	0	0
BigToM	Kimi-K2-Thinking	Partial Reconstruction	100	0	0	0
OpenToM	DeepSeek-V3.2	Free Generation	100	0	0	0
OpenToM	DeepSeek-V3.2	Partial Reconstruction	100	0	0	0
OpenToM	Kimi-K2-Thinking	Free Generation	100	0	0	0
OpenToM	Kimi-K2-Thinking	Partial Reconstruction	100	0	0	0

Table 10: Potential-contamination verification results for 100 trials per dataset, model, and protocol. Free Generation follows Li and Flanigan (Li and Flanigan, 2024), and Partial Reconstruction follows Golchin and Surdeanu (Golchin and Surdeanu, 2024). Strong flags indicate exact or near-exact story reconstruction; likely and possible flags indicate high fuzzy similarity and are manually reviewed.

Modality	Task 1		Task 2	
	Acc.	F1	Acc.	F1
Visual	93.69	93.67	93.47	83.55
Auditory	94.71	94.65	94.31	77.94
Gustatory	96.51	95.25	95.42	80.33
Tactile	94.31	93.58	94.01	79.46
Olfactory	90.29	88.56	90.51	83.31
Proprioceptive	93.22	91.07	95.93	85.07
Overall	**93.76**	**93.53**	**93.69**	**82.87**

Table 11: Human reference scores (%) by sensory modality on clean belief-change detection (Task 1) and updated-belief identification (Task 2). F1 denotes macro-F1.

(a) Individual Task 2 scores

Evaluation	Acc.	Macro-F1
1	93.83	81.63
2	93.58	84.29
3	93.66	82.70
Mean	93.69	82.87
Sample SD	0.13	1.34
Sample variance	0.0160	1.7899

(b) Agreement by sensory category

Modality	n	Task 1		Task 2	
		κ	Agree.	κ	Agree.
Visual	750	0.754	87.73	0.745	86.93
Auditory	586	0.792	89.76	0.773	88.62
Gustatory	153	0.806	93.03	0.750	90.85
Tactile	445	0.758	89.36	0.734	88.01
Olfactory	309	0.569	81.66	0.557	81.01
Proprioceptive	172	0.788	91.86	0.793	91.86
Overall	2,415	0.758	88.38	0.743	87.38

Table 12: Individual scores and agreement for the clean records in Table 11. Panel (a): accuracy and macro-F1 are percentages; sample SD is in percentage points and sample variance in squared percentage points, based on $n = 3$ evaluations. Panel (b): n is the number of items; Agree. is mean pairwise agreement (%).

(a) Human reference with 95% confidence intervals

Task	Accuracy [95% CI]	Macro-F1 [95% CI]
Task 1	93.76 [93.23, 94.29]	93.53 [92.97, 94.08]
Task 2	93.69 [93.23, 94.15]	82.87 [79.45, 84.77]

(b) Paired human–model accuracy comparisons

Model	n	Task 1		Task 2	
		Model acc.	Human − model [95% CI]	Model acc.	Human − model [95% CI]
GPT-3.5-Turbo	2,415	44.74	49.02 [47.18, 50.85]	41.39	52.30 [50.33, 54.31]
GPT-4o	2,415	70.06	23.70 [21.62, 25.80]	49.83	43.87 [41.26, 46.42]
GPT-OSS-120B	2,415	82.88	10.88 [9.32, 12.48]	61.33	32.36 [29.71, 34.96]
Llama-8B	2,415	53.56	40.20 [38.84, 41.59]	34.32	59.37 [57.30, 61.46]
Llama-405B	2,415	77.44	16.32 [14.44, 18.26]	54.26	39.43 [36.51, 42.26]
Llama-4-Maverick-17B	2,415	68.24	25.52 [23.88, 27.17]	55.73	37.96 [35.30, 40.62]
Mistral-7B	2,415	56.69	37.07 [35.57, 38.53]	41.70	51.99 [50.02, 53.92]
Mixtral-8x22B	2,415	64.10	29.67 [27.85, 31.48]	43.44	50.25 [47.81, 52.70]
Qwen3-235B-A22B	2,415	68.04	25.72 [23.63, 27.83]	57.01	36.68 [33.73, 39.61]
QwQ-32B	2,415	82.61	11.15 [9.16, 13.24]	63.52	30.17 [27.54, 32.83]
Qwen3.5-397B-A17B	2,415	85.34	8.42 [6.94, 9.99]	74.58	19.12 [16.59, 21.68]
DeepSeek-R1	2,388	77.61	16.35 [14.95, 17.73]	69.65	24.23 [21.41, 27.11]
DeepSeek-V3.1	2,415	82.57	11.19 [9.44, 12.97]	65.13	28.56 [25.77, 31.38]
DeepSeek-V3.2	2,415	75.65	18.11 [16.27, 19.95]	70.72	22.97 [20.48, 25.57]
Kimi-K2-Thinking	2,415	67.41	26.35 [24.41, 28.29]	75.40	18.29 [15.94, 20.69]
MiniMax-M2.1	2,415	78.55	15.21 [13.07, 17.42]	73.00	20.69 [18.20, 23.25]
MiniMax-M2.5	2,415	86.05	7.72 [6.05, 9.37]	74.12	19.57 [17.40, 21.84]
Phi-3-Medium-4K	2,415	56.45	37.31 [35.66, 38.96]	49.64	44.05 [41.44, 46.62]

Table 13: Uncertainty for the clean response-level human reference and paired accuracy comparisons. Panel (a) reports percentages. Panel (b) reports model accuracy (%), human-minus-model differences in percentage points, and common item counts n; both scores are recomputed on the matched items. Intervals use 10,000 source-stratified story-bootstrap replicates. All 36 two-sided accuracy tests have Holm-adjusted $p < 0.001$ under the story-level paired-randomization null. Intervals and tests condition on the observed evaluators, models, and valid responses.

Principal contrast	Estimate	95% CI
HN overprediction	59.2%	56.0–62.4%
TNR-e − TNR-h	24.5 pp	21.7–27.1 pp
TPR − TNR-h	47.4 pp	43.6–51.2 pp

Table 14: Principal diagnostic contrasts macro-averaged across 36 model–prompt conditions. Intervals use 10,000 source-story-clustered bootstrap replicates. HN denotes hard negative; pp denotes percentage points.

Metric (%)	Direct	CoT
P–HN jointly correct	26.7	36.1
P–EN jointly correct	47.4	58.4
All three classes correct	24.4	32.1
P change-response rate	83.3	83.1
HN change-response rate	53.0	44.5
EN change-response rate	29.4	20.4

Table 15: Task 1 discrimination within matched contexts, averaged equally over stories and models. Joint correctness requires the appropriate decision for both cue classes; the three-class score requires all three decisions to be correct.

Prompt	Prior	Target	Unc.	Other
BigToM				
Direct	15.7	44.4	39.2	0.8
CoT	19.1	39.6	39.6	1.7
OpenToM				
Direct	32.8	54.8	11.7	0.7
CoT	38.4	47.5	13.2	1.0
All sources				
Direct	23.9	49.4	25.9	0.7
CoT	28.4	43.4	26.8	1.4

Table 16: Task 2 belief-state response shares (%) on hard negatives. Prior is the gold state; Unc. denotes the protagonist's uncertainty, not model confidence. Each row sums to 100% before rounding. All-source results pool items within each model before averaging models.

Wording	Answers	P changed	HN changed	EN changed	HN uncertain	HN–EN gap (pp) [95% CI]
Original	Binary	76.7	41.7	17.5	—	24.2 [14.2, 35.0]
Neutral	Binary	70.8	33.3	11.7	—	21.7 [10.8, 33.3]
Original	Ternary	59.2	35.0	8.3	38.3	26.7 [15.0, 39.2]
Neutral	Ternary	66.7	27.5	8.3	41.7	19.2 [9.2, 29.2]

Table 17: Wording and response-format experiment on matched cue triplets. Rates (%) are macro-averages of four models, each evaluated on 30 items per cue class. P, HN, and EN denote positive, hard-negative, and easy-negative cues. Confidence intervals use 5,000 source-stratified story-bootstrap resamples of the 30-story sample, retaining all conditions and the four fixed models for each story.

Case type	Perceptual cue	Gold / prediction	Diagnostic reading
Hard negative error	A bitter, metallic taste appears after seawater enters the character's mouth.	No change / change	The cue is plausible for the belief target, but it does not provide sufficient evidence for belief revision.
Easy negative rejection	A clatter is heard elsewhere in the coffee shop.	No change / no change	The cue is irrelevant to the target belief, and the model rejects it as evidence.
Positive acceptance	A milk frother makes a distinct higher-pitched sound associated with nut-based milk.	Change / change	The cue directly supports an updated belief, and the model predicts belief change.

Table 18: Representative qualitative examples. The examples are shortened for readability; full prompts and model outputs are retained in the analysis files.

Field	Content
Dataset / setting	BigToM, visual hard negative, clean Task 1 with CoT prompting.
Story context	Noor is working as a barista. She fills a milk pitcher with oat milk for a customer who requested oat milk. While Noor is attending to another task, a coworker swaps the oat milk in the pitcher with almond milk, but Noor does not observe the swap.
Initial belief	Noor believes that the milk pitcher contains oat milk.
Perceptual cue	A bright green label on a nearby milk carton catches Noor's eye. The label reads "Almond Milk."
Gold belief state	No belief change: Noor should still believe that the pitcher contains oat milk, because the nearby label is not direct evidence about the contents of the pitcher.
Model prediction	Belief change: GPT-4o with CoT predicts that Noor no longer believes the pitcher contains oat milk.
Failure pattern	The model reasons from the presence of an almond-milk label to possible doubt about the pitcher's contents. This treats a salient, belief-related cue as if it were diagnostic evidence, even though the cue is not linked to the object whose contents Noor is reasoning about.

Table 19: Qualitative hard-negative case study illustrating over-updating on perceptually salient but epistemically insufficient evidence.

Modality	Acc.	TNR-h	HN O.P.	Slot gap
Visual	75.6	48.7	51.3	34.3
Auditory	74.9	55.4	44.6	32.9
Gustatory	73.8	58.8	41.2	23.2
Tactile	70.9	49.0	51.0	31.5
Olfactory	69.3	51.0	49.0	22.8
Proprioceptive	68.3	49.8	50.2	31.7

Table 20: Modality-level diagnostic summary. Acc., TNR-h, and hard-negative overprediction are means over model-prompt observations. Slot gap reports the overall-accuracy minus TNR-h gap computed before table-level aggregation, so it need not equal the difference between the displayed rounded Acc. and TNR-h columns.

Modality	Accuracy		TNR-h	
	Raw	Std.	BigToM	OpenToM
Visual	76.26	74.09	51.31	50.65
Auditory	74.60	73.92	56.37	50.62
Gustatory	75.75	77.48	53.80	67.56
Tactile	72.15	73.88	49.88	46.74
Olfactory	70.22	72.22	44.71	57.83
Proprioceptive	70.59	73.15	48.95	48.96

Table 21: Task 1 modality results (%), averaged equally over 18 models and two prompting conditions. Raw and Std. denote accuracy under observed and common category proportions, respectively. TNR-h is hard-negative rejection within each source.

Aspect	Belief-R	MuSeToM
Reasoning target	The model's own inferred conclusion	Another agent's warranted belief revision
Evidence	Two initial text premises, followed by a new conditional premise	A story, initial belief, and textual descriptions of sensory cues
Evaluation	Conclusions before and after adding a premise	Belief-change detection and updated-belief identification, with and without distractors
Diagnostic classes	Belief Update and Belief Maintain	Positive, hard-negative, and easy-negative cues

Table 22: Comparison of evaluation designs. Both benchmarks provide textual inputs; MuSeToM's sensory categories describe cue content rather than raw sensory channels.

LLMs	Belief Change				Belief Change with Noise			
	Binary		MC		Binary		MC	
	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)	Acc. (%)	F1 (%)
GPT-3.5-Turbo	44.76	44.55	41.57	41.60	64.73	64.51	62.91	62.79
GPT-3.5-Turbo (CoT)	53.25	53.19	49.23	49.23	68.17	65.77	64.53	64.51
GPT-4o	70.23	70.09	50.02	49.82	71.73	71.72	60.01	59.54
GPT-4o (CoT)	76.94	76.89	61.08	61.08	77.95	74.46	64.40	64.30
GPT-OSS-120B	82.53	82.53	61.04	61.08	<u>80.51</u>	<u>80.48</u>	70.94	70.88
GPT-OSS-120B (CoT)	<u>85.47</u>	<u>85.46</u>	63.77	63.78	**81.59**	**81.57**	**72.29**	**72.26**
Llama-8B	52.09	51.32	34.20	34.09	62.37	60.82	52.26	51.80
Llama-8B (CoT)	70.02	70.00	50.97	50.94	70.13	70.04	61.83	61.81
Llama-405B	77.27	77.25	54.66	54.63	71.34	71.13	59.27	59.15
Llama-405B (CoT)	84.51	84.51	59.59	59.53	80.45	80.35	62.58	62.50
Llama-4-Maverick-17B	67.83	67.51	56.11	55.95	63.86	62.35	64.33	64.19
Llama-4-Maverick-17B (CoT)	79.21	79.21	53.37	53.37	78.76	78.70	61.56	61.54
Mistral-7B	55.82	55.26	41.86	41.79	60.35	59.67	58.53	58.61
Mistral-7B (CoT)	66.05	65.99	42.69	42.69	62.44	62.44	59.14	59.12
Mixtral-8x22B	63.52	63.35	43.60	43.50	65.88	65.67	56.71	56.68
Mixtral-8x22B (CoT)	74.20	74.18	53.29	53.28	71.34	71.34	59.34	59.44
Qwen3-235B-A22B	67.83	67.75	57.10	57.12	68.71	68.71	66.76	66.73
Qwen3-235B-A22B (CoT)	84.55	84.55	67.04	67.02	77.88	77.79	63.32	63.34
QwQ-32B	82.61	82.61	63.52	63.64	77.07	77.07	67.16	67.33
QwQ-32B (CoT)	82.19	82.13	61.16	61.33	78.49	78.31	64.46	64.71
Qwen3.5-397B-A17B	85.34	85.34	74.58	74.57	67.84	67.81	<u>71.27</u>	<u>71.24</u>
Qwen3.5-397B-A17B (CoT)	84.84	84.84	<u>75.24</u>	<u>75.27</u>	71.81	71.81	68.78	68.75
DeepSeek-R1	77.81	77.54	70.02	69.97	68.54	68.24	67.15	67.10
DeepSeek-R1 (CoT)	79.17	79.08	70.14	70.19	76.20	72.71	68.85	68.83
DeepSeek-V3.1	82.57	82.57	65.13	65.12	55.97	55.70	56.98	56.92
DeepSeek-V3.1 (CoT)	82.65	82.65	70.93	70.91	73.50	73.48	63.99	63.90
DeepSeek-V3.2	75.65	75.57	70.72	70.68	69.86	69.86	60.62	60.48
DeepSeek-V3.2 (CoT)	84.10	84.10	68.12	68.09	68.91	68.70	61.63	61.59
Kimi-K2-Thinking	67.41	65.50	**75.40**	**75.39**	77.07	77.05	69.45	69.45
Kimi-K2-Thinking (CoT)	68.70	66.90	72.22	72.30	63.05	60.47	62.78	62.82
MiniMax-M2.1	78.55	78.54	73.00	73.00	76.67	76.67	71.07	71.07
MiniMax-M2.1 (CoT)	69.22	68.39	55.93	56.09	69.30	69.00	58.80	58.94
MiniMax-M2.5	**86.05**	**86.05**	74.12	74.10	79.84	79.83	70.73	70.74
MiniMax-M2.5 (CoT)	84.35	84.35	73.13	73.10	77.67	77.65	68.64	68.63
Phi-3-Medium-4K	56.56	55.72	49.77	49.77	63.79	63.37	66.42	66.37
Phi-3-Medium-4K (CoT)	58.18	58.18	49.28	48.89	54.35	54.35	59.20	59.11

Table 23: Full main results on MuSeToM, including both direct and CoT prompting. Each model is shown with two rows. We report Accuracy (%) and Macro-F1 (%) under two settings, Belief Change and Belief Change with Noise, and two question formats, binary and multiple-choice. Best and second-best results in each accuracy column are marked in bold and underlined, respectively.

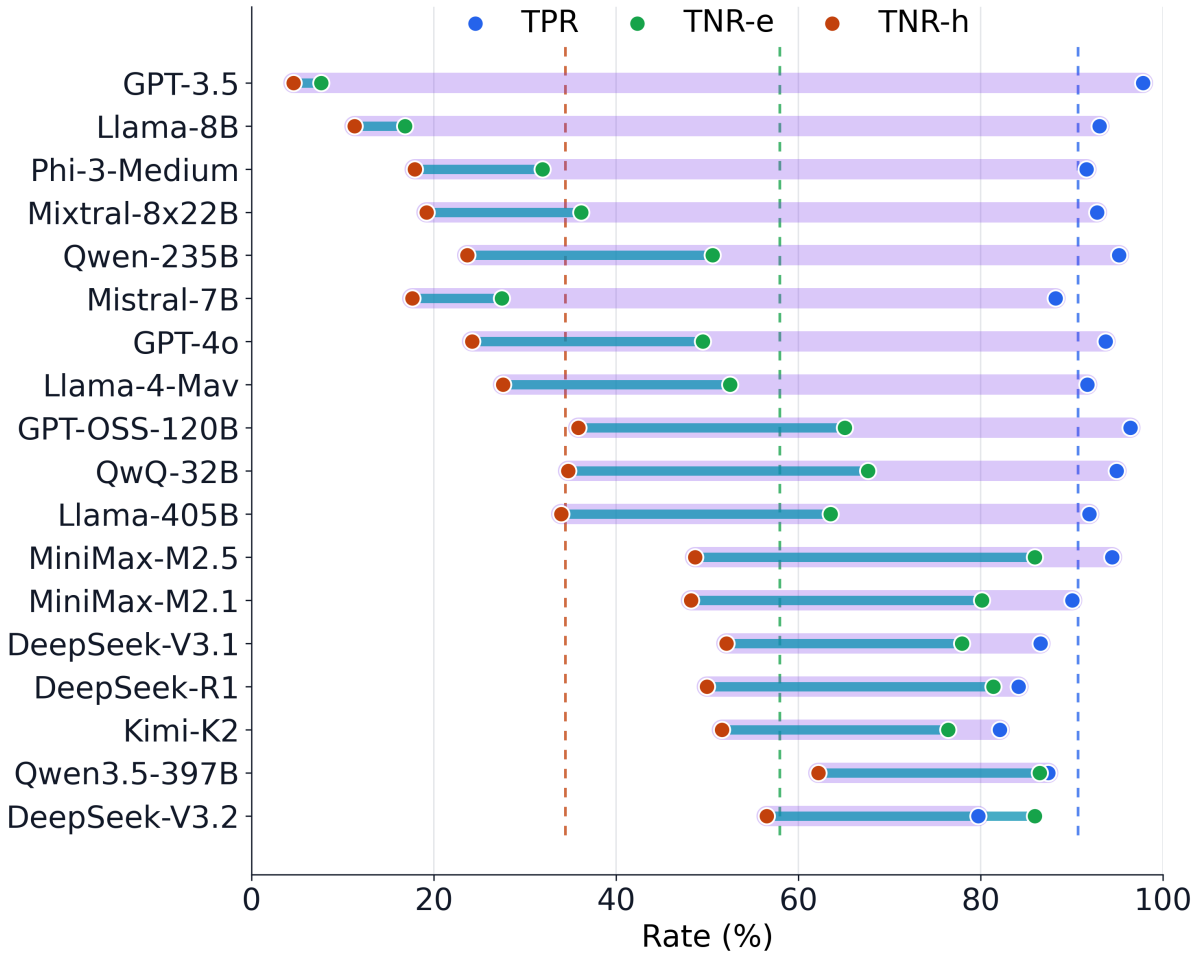

Figure 7: Full-model version of the over-updating diagnostic in Figure 3. Each row reports clean direct-prompting TPR, TNR-e, and TNR-h for one model. The lavender span marks TPR $-$ TNR-h and the teal span marks TNR-e $-$ TNR-h; gap labels are omitted for readability.

Model Name	bigtom					opentom				
	TPR	TNR-h	TNR-e	Acc	F1	TPR	TNR-h	TNR-e	Acc	F1
GPT-4o	90.82	25.20	64.93	67.66	67.48	82.85	50.00	72.30	72.63	72.53
GPT-3.5-Turbo	95.41	0.81	1.42	41.57	41.57	92.20	10.27	19.45	47.75	47.08
GPT-OSS-120B	90.82	40.24	80.57	76.48	76.47	94.52	61.04	94.25	88.20	88.20
Llama-8B	79.64	32.52	33.65	53.12	52.61	90.45	14.29	27.70	51.12	50.05
Llama-405B	85.63	56.10	84.12	78.87	78.85	84.02	46.88	80.16	75.76	75.75
Llama-4-Maverick-17B	89.22	30.89	51.42	63.30	63.00	78.47	48.48	76.39	72.07	71.75
Mistral-7B-Instruct	88.42	14.23	21.56	48.67	48.24	63.01	45.45	69.84	62.52	61.87
Mixtral-8x22B	88.22	25.61	43.84	59.02	58.45	80.12	47.77	64.05	67.74	67.72
Qwen3-235B-A22B	89.22	28.05	48.10	61.51	61.44	92.95	36.36	71.43	73.76	73.68
QwQ-32B	85.23	64.23	92.42	83.40	83.40	84.54	64.94	86.90	81.86	81.85
DeepSeek-R1	81.43	62.20	83.41	78.02	77.69	65.50	72.32	92.14	77.61	77.39
Phi-3-Medium-4K	92.22	10.57	21.56	49.53	49.12	71.43	41.13	64.88	63.16	61.87
DeepSeek-V3.1	83.23	70.33	92.65	83.92	83.91	85.32	60.61	86.71	81.30	81.29
DeepSeek-V3.2	68.86	86.18	98.58	83.23	83.20	57.34	60.61	83.53	68.54	68.38
Kimi-K2-Thinking	70.26	67.07	78.44	72.54	71.39	63.80	56.71	64.09	62.60	59.70
MiniMax-M2.1	85.43	65.04	84.83	80.92	80.90	80.23	58.01	80.75	76.32	76.32
MiniMax-M2.5	89.02	73.17	97.87	88.88	88.88	91.00	59.31	86.71	83.39	83.38
Qwen3.5-397B-A17B	82.44	80.08	97.63	87.43	87.42	80.43	76.62	89.48	83.39	83.35

Table 24: Dataset-specific direct-prompting performance on Task 1 (belief-change detection).

Model Name	bigtom					opentom				
	TPR	TNR-h	TNR-e	Acc	F1	TPR	TNR-h	TNR-e	Acc	F1
GPT-4o	78.44	5.28	23.70	43.28	42.81	83.04	27.68	42.04	56.34	56.19
GPT-3.5-Turbo	88.22	0.00	0.47	37.98	37.99	98.05	8.93	7.27	44.94	45.09
GPT-OSS-120B	90.82	7.32	24.17	49.19	49.20	95.69	37.23	64.29	72.15	72.17
Llama-8B	68.46	1.22	1.66	30.20	30.01	90.25	2.23	0.98	37.96	37.88
Llama-405B	72.26	12.60	43.84	49.44	49.42	86.74	28.12	45.97	59.55	59.40
Llama-4-Maverick-17B	76.65	14.63	44.08	51.84	51.36	95.69	22.08	41.47	60.11	60.10
Mistral-7B-Instruct	87.82	0.81	0.95	38.15	38.21	88.45	16.88	14.68	45.35	45.11
Mixtral-8x22B	76.05	0.00	3.08	33.70	33.70	88.50	22.77	30.26	52.89	52.72
Qwen3-235B-A22B	83.63	9.76	33.18	49.87	49.77	96.67	28.14	47.02	63.88	63.87
QwQ-32B	85.43	20.33	44.31	56.89	56.97	92.76	38.10	60.91	69.74	69.88
DeepSeek-R1	77.85	27.24	76.30	66.37	66.20	87.52	46.88	70.73	73.35	73.22
Phi-3-Medium-4K	92.81	0.00	0.47	39.95	39.95	92.56	29.00	38.69	58.99	58.96
DeepSeek-V3.1	73.85	28.86	61.61	59.97	59.82	79.06	48.92	70.44	69.98	69.97
DeepSeek-V3.2	63.87	32.11	79.62	62.87	62.69	87.48	46.75	82.94	78.09	78.09
Kimi-K2-Thinking	76.05	33.33	81.28	68.95	68.95	94.52	49.35	82.94	81.46	81.44
MiniMax-M2.1	80.24	23.17	75.83	66.64	66.57	93.15	47.19	79.17	78.97	78.96
MiniMax-M2.5	81.04	20.73	74.88	66.12	66.07	97.06	41.56	84.33	81.62	81.60
Qwen3.5-397B-A17B	74.05	42.68	83.89	71.00	70.91	92.56	49.35	76.19	77.93	77.92

Table 25: Dataset-specific direct-prompting performance on Task 2 (updated-belief identification).

Model Name	bigtom				opentom			
	TPR	TNR-h	Acc	F1	TPR	TNR-h	Acc	F1
GPT-4o	86.23	53.25	75.27	75.24	73.68	55.90	68.19	68.19
GPT-3.5-Turbo	95.35	0.41	63.83	63.83	89.47	12.23	65.63	64.88
GPT-OSS-120B	90.10	50.00	76.79	76.78	92.37	66.23	84.23	84.18
Llama-8B	78.59	29.67	62.35	61.19	81.68	19.21	62.40	60.39
Llama-405B	83.23	52.85	73.14	73.02	77.00	52.84	69.54	69.20
Llama-4-Maverick-17B	77.17	40.65	65.05	63.97	65.56	56.28	62.67	60.63
Mistral-7B-Instruct	90.10	11.38	63.97	63.89	56.56	57.14	56.74	54.66
Mixtral-8x22B	86.46	30.08	67.75	67.46	68.23	54.59	64.02	63.88
Qwen3-235B-A22B	87.68	29.27	68.29	68.29	80.82	43.29	69.14	69.14
QwQ-32B	82.63	67.48	77.60	77.59	80.23	68.40	76.55	76.55
DeepSeek-R1	75.80	64.63	71.95	71.76	62.77	71.03	65.20	64.75
Phi-3-Medium-4K	92.73	12.60	66.13	66.05	67.12	48.92	61.46	60.38
DeepSeek-V3.1	59.60	80.08	66.40	66.14	29.55	80.95	45.55	45.28
DeepSeek-V3.2	63.23	91.87	72.74	72.72	64.38	72.73	66.98	66.98
Kimi-K2-Thinking	79.60	73.17	77.46	77.40	81.41	66.23	76.68	76.68
MiniMax-M2.1	84.65	72.36	80.57	80.55	77.50	62.34	72.78	72.78
MiniMax-M2.5	88.08	69.92	82.05	82.03	86.11	58.87	77.63	77.63
Qwen3.5-397B-A17B	68.08	75.61	70.58	70.57	62.04	71.86	65.09	65.05

Table 26: Dataset-specific direct-prompting performance on Task 3 (noisy belief-change detection).

Model Name	bigtom				opentom			
	TPR	TNR-h	Acc	F1	TPR	TNR-h	Acc	F1
GPT-4o	71.72	11.79	51.82	51.12	84.02	32.75	68.19	67.81
GPT-3.5-Turbo	88.08	0.00	58.84	58.90	94.35	5.68	66.98	66.74
GPT-OSS-120B	88.89	11.79	63.29	63.29	95.89	40.26	78.57	78.40
Llama-8B	70.71	0.41	47.37	47.08	82.46	0.44	57.14	56.40
Llama-405B	65.05	15.04	48.45	48.22	86.74	32.75	70.08	69.98
Llama-4-Maverick-17B	70.71	20.73	54.12	53.43	95.50	28.14	74.53	74.52
Mistral-7B-Instruct	90.10	0.41	60.32	60.30	75.34	15.58	56.74	56.97
Mixtral-8x22B	73.33	0.00	48.99	48.95	81.09	27.07	64.42	64.40
Qwen3-235B-A22B	81.01	15.45	59.24	59.28	93.54	31.60	74.26	74.11
QwQ-32B	79.39	23.58	60.86	61.02	86.50	44.59	73.45	73.61
DeepSeek-R1	74.52	32.52	60.03	59.82	84.99	48.13	74.14	74.15
Phi-3-Medium-4K	92.73	0.41	62.08	62.06	87.08	34.63	70.75	70.63
DeepSeek-V3.1	60.81	41.46	54.39	54.22	58.71	61.47	59.57	59.59
DeepSeek-V3.2	55.76	43.09	51.55	51.30	75.54	56.71	69.68	69.64
Kimi-K2-Thinking	67.07	36.59	56.95	56.96	91.98	59.74	81.94	81.93
MiniMax-M2.1	76.36	34.96	62.62	62.58	91.39	53.25	79.51	79.53
MiniMax-M2.5	77.58	30.08	61.81	61.84	95.50	44.59	79.65	79.66
Qwen3.5-397B-A17B	68.28	57.72	64.78	64.68	87.67	55.84	77.76	77.78

Table 27: Dataset-specific direct-prompting performance on Task 4 (noisy updated-belief identification).

Model Name	bigtom					opentom				
	TPR	TNR-h	TNR-e	Acc	F1	TPR	TNR-h	TNR-e	Acc	F1
GPT-4o	89.42	53.66	83.41	79.73	79.69	84.93	51.95	73.81	74.32	74.26
GPT-3.5-Turbo	91.82	18.70	31.04	54.49	54.49	98.83	3.90	26.79	52.09	51.86
GPT-OSS-120B	88.82	54.88	91.00	82.46	82.46	92.95	60.17	96.43	88.28	88.28
Llama-8B	82.83	45.12	57.11	65.61	65.56	94.32	27.71	75.00	74.16	74.15
Llama-405B	79.64	77.24	95.50	84.86	84.86	91.19	58.44	88.89	84.19	84.18
Llama-4-Maverick-17B	90.02	49.19	77.73	76.99	76.97	82.00	66.67	87.30	81.30	81.30
Mistral-7B-Instruct	87.03	27.24	39.34	57.23	57.18	78.86	48.05	81.75	74.32	74.26
Mixtral-8x22B-Instruct	87.62	47.15	70.62	72.97	72.87	85.13	55.84	74.40	75.36	75.36
Qwen3-235B-A22B	83.23	74.80	97.39	86.57	86.57	86.30	61.04	88.89	82.66	82.66
QwQ-32B	83.63	65.85	94.31	83.75	83.67	86.11	59.31	85.12	80.74	80.68
DeepSeek-R1	90.22	50.81	92.42	82.72	82.68	81.41	61.90	76.59	75.84	75.68
Phi-3-Medium-4K-Instruct	89.42	21.95	28.44	53.21	53.20	36.79	69.70	86.11	62.84	62.84
DeepSeek-V3.1	80.84	74.80	94.79	84.60	84.59	84.54	62.34	85.52	80.82	80.82
DeepSeek-V3.2	74.05	87.80	98.82	85.89	85.88	83.37	65.80	89.09	82.42	82.42
Kimi-K2-Thinking	67.66	68.29	72.27	69.46	67.09	67.12	63.20	71.03	67.98	66.65
MiniMax-M2.1	62.48	64.63	76.78	68.09	66.12	70.59	58.01	75.60	70.28	70.05
MiniMax-M2.5	89.02	70.73	95.50	87.51	87.51	91.00	53.68	84.33	81.38	81.38
Qwen3.5-397B-A17B	79.64	85.77	99.29	88.02	88.02	79.45	69.26	90.08	81.86	81.82

Table 28: Dataset-specific CoT-prompting performance on Task 1 (belief-change detection).

Model Name	bigtom					opentom				
	TPR	TNR-h	TNR-e	Acc	F1	TPR	TNR-h	TNR-e	Acc	F1
GPT-4o	81.84	14.63	49.53	56.03	55.80	77.50	43.72	64.09	65.81	65.80
GPT-3.5-Turbo	84.63	6.50	19.19	44.57	44.56	84.15	24.24	36.11	53.61	53.56
GPT-OSS-120B	90.42	8.94	30.33	51.58	51.57	95.89	43.72	68.65	75.20	75.16
Llama-8B	86.63	4.88	20.62	45.59	45.54	90.80	18.18	38.10	56.02	55.90
Llama-405B	62.67	19.11	61.61	53.12	53.04	84.54	38.96	58.73	65.65	65.59
Llama-4-Maverick-17B	74.05	12.60	27.73	44.40	44.39	87.48	25.97	52.18	61.80	61.79
Mistral-7B-Instruct	86.43	0.81	4.27	38.75	38.76	90.02	16.45	15.87	46.39	46.41
Mixtral-8x22B-Instruct	72.26	7.32	25.83	41.83	41.82	83.56	32.90	58.53	64.04	64.02
Qwen3-235B-A22B	73.45	28.46	69.19	62.45	62.39	86.69	39.83	70.24	71.35	71.31
QwQ-32B	81.04	20.73	47.39	56.20	56.25	88.06	36.36	56.75	65.81	65.95
DeepSeek-R1	79.84	27.64	73.46	66.55	66.52	88.45	44.16	71.83	73.52	73.50
Phi-3-Medium-4K-Instruct	91.62	2.44	3.08	40.89	40.63	78.08	40.26	43.65	57.14	56.68
DeepSeek-V3.1	69.06	32.93	78.67	64.93	64.85	87.48	46.32	79.37	76.57	76.57
DeepSeek-V3.2	57.49	32.11	72.27	57.49	57.38	84.74	49.35	84.52	78.09	78.11
Kimi-K2-Thinking	68.46	34.55	79.38	65.27	65.41	86.11	51.52	83.73	78.73	78.78
MiniMax-M2.1	59.28	25.61	54.55	50.47	50.66	66.14	47.39	62.10	61.04	61.16
MiniMax-M2.5	82.44	19.92	68.25	64.16	64.10	96.09	40.69	85.52	81.54	81.52
Qwen3.5-397B-A17B	65.67	45.53	86.49	68.95	69.06	93.74	49.35	82.94	81.14	81.12

Table 29: Dataset-specific CoT-prompting performance on Task 2 (updated-belief identification).

Model Name	bigtom				opentom			
	TPR	TNR	Acc	F1	TPR	TNR	Acc	F1
GPT-4o	86.26	67.48	80.03	77.39	84.15	57.58	75.88	71.35
GPT-3.5-Turbo	91.11	15.45	65.99	65.97	98.63	7.79	70.35	60.25
GPT-OSS-120B	87.07	63.01	79.08	79.08	91.59	67.53	84.10	84.04
Llama-8B	77.17	52.44	68.96	68.88	89.43	31.17	71.29	71.20
Llama-405B	79.80	78.05	79.22	79.15	89.63	64.07	81.67	81.54
Llama-4-Maverick-17B	87.88	60.16	78.68	78.61	83.56	68.40	78.84	78.79
Mistral-7B-Instruct	90.10	17.89	66.13	66.13	55.97	64.94	58.76	58.76
Mixtral-8x22B-Instruct	80.20	58.13	72.87	72.85	71.82	65.37	69.81	69.81
Qwen3-235B-A22B	81.21	76.02	79.49	79.37	85.13	56.71	76.28	76.21
QwQ-32B	85.45	69.51	80.16	79.92	84.34	60.17	76.82	76.70
DeepSeek-R1	87.07	61.79	78.68	75.60	79.65	60.61	73.72	69.81
Phi-3-Medium-4K-Instruct	85.66	21.14	64.24	64.22	36.59	61.90	44.47	44.40
DeepSeek-V3.1	70.10	82.52	74.22	74.20	76.71	64.07	72.78	72.76
DeepSeek-V3.2	64.44	90.65	73.14	73.13	62.43	69.70	64.69	63.95
Kimi-K2-Thinking	63.43	68.29	65.05	62.30	58.12	67.53	61.05	58.64
MiniMax-M2.1	71.52	70.73	71.26	70.51	71.76	57.58	67.34	67.30
MiniMax-M2.5	86.44	71.95	81.62	81.60	80.82	58.01	73.72	73.72
Qwen3.5-397B-A17B	68.89	84.55	74.09	74.09	65.75	77.92	69.54	69.54

Table 30: Dataset-specific CoT-prompting performance on Task 3 (noisy belief-change detection).

Model Name	bigtom				opentom			
	TPR	TNR	Acc	F1	TPR	TNR	Acc	F1
GPT-4o	75.56	23.17	58.16	57.94	81.21	47.19	70.62	70.53
GPT-3.5-Turbo	82.63	10.98	58.84	58.85	92.37	21.21	70.22	70.14
GPT-OSS-120B	89.90	15.85	65.32	65.32	95.50	43.29	79.25	79.12
Llama-8B	79.80	10.16	56.68	56.65	91.19	13.42	66.98	66.75
Llama-405B	66.26	29.27	53.98	53.75	85.91	38.53	71.16	71.17
Llama-4-Maverick-17B	70.91	17.07	53.04	53.06	87.67	31.17	70.08	69.99
Mistral-7B-Instruct	85.45	0.81	57.35	57.34	81.80	14.72	60.92	60.85
Mixtral-8x22B-Instruct	72.73	13.01	52.90	52.95	77.10	40.69	65.77	65.94
Qwen3-235B-A22B	69.29	35.77	58.16	58.14	80.82	41.13	68.46	68.49
QwQ-32B	78.38	20.33	59.11	59.34	83.95	38.53	69.81	69.98
DeepSeek-R1	74.75	33.33	61.00	60.90	87.87	51.95	76.68	76.67
Phi-3-Medium-4K-Instruct	92.73	3.25	63.02	62.82	62.04	40.69	55.39	55.34
DeepSeek-V3.1	61.62	46.75	56.68	56.46	80.23	51.52	71.29	71.28
DeepSeek-V3.2	55.35	41.87	50.88	50.79	80.04	55.41	72.37	72.38
Kimi-K2-Thinking	59.19	39.02	52.50	52.38	81.60	54.11	73.05	73.12
MiniMax-M2.1	59.84	29.27	49.66	49.66	78.04	45.41	67.93	68.13
MiniMax-M2.5	77.58	26.02	60.46	60.44	92.76	41.56	76.82	76.78
Qwen3.5-397B-A17B	61.01	54.07	58.70	58.66	87.87	58.87	78.84	78.86

Table 31: Dataset-specific CoT-prompting performance on Task 4 (noisy updated-belief identification).

BigToM Generation Prompt Template

Belief-changing cue generation:

Story: {Context of Story Scenario from BigToM} {Protagonist's Desires in Story Scenario from BigToM} {Initial Binary Perception Cues in Story Scenario from BigToM}

Main character's belief: {Protagonist's Belief in Story Scenario from BigToM}

Event: {Causal Event in Story Scenario from BigToM}

Instructions: Based on the following story and event, generate perceptual cues that the main character might encounter, which MUST challenge their belief. Each cue should be creative, reasonable, and systematically explore the sensory modalities of: visual, auditory, olfactory, gustatory, tactile, proprioception (sense of body position/weight). Be descriptive and only make vivid descriptions about the perceptual cue, without describing the main character's thoughts. Avoid making explicit comparisons between the event and the main character's belief. Generate 2 perceptual cues for each sensory modality. Be super duper creative! Do not use these words: different, difference, instead of, contrasts.

For each cue, write your output in a format like this example:

Visual Cues:

1. {cue #1}
2. {cue #2}

Random-event cue generation:

The same template is used, replacing {Causal Event in Story Scenario from BigToM} with {Random Event in Story Scenario from BigToM} and only generating one perceptual cue per sensory modality instead of two.

Table 32: The generation prompt template used for BigToM-based MuSeToM instances.

OpenToM Generation Prompt Template

Belief-changing event generation:

Story: {Story Scenario from OpenToM}

Main character's belief: {Protagonist's Belief in Story Scenario from OpenToM}

Instructions: Based on the following story, generate a one-line story consisting an event that MUST NEGATE {protagonist name}'s initial belief and change it to: "{new belief}". Note that this story MUST also be consistent with {observed character}'s preference, i.e., {observed character's preference}. Each story should be creative, reasonable, realistic, and systematically explore the sensory modalities of: visual, auditory, olfactory, gustatory, tactile, proprioception (sense of body position/weight). Generate 2 stories for each sensory modality. Be super duper creative! Show, don't tell. Write the stories in this format:

Visual Stories:

1. {story #1}
2. {story #2}

Random-event generation:

Story: {Story Scenario from OpenToM}

Main character's belief: {Protagonist's Belief in Story Scenario from OpenToM}

Instructions: Based on the following story, generate a one-line story consisting a random event that occurs in the environment, which is unrelated to the object and cannot influence the main character {protagonist name}'s belief. Each story should be creative, reasonable, realistic, and systematically explore the sensory modalities of: visual, auditory, olfactory, gustatory, tactile, proprioception (sense of body position/weight). Generate 1 story for each sensory modality. Be super duper creative! Show, don't tell. Write the stories in this format:

Visual Stories:

{story #1}

Table 33: The generation prompt templates used for OpenToM-based MuSeToM instances.

MuSeToM Task 1 Prompt Template: Belief Change Detection

BigToM Template:
Story: {story}
Event: {event}
Regarding the event, the main character {protag_name} only encounters the following percept: {percept}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
Did the event cause the main character {protag_name} to change their belief?
A. {Yes/No - randomly shuffled}
B. {No/Yes - randomly shuffled}

OpenToM Template:
Story: {story}
Regarding the event, the main character {protag_name} only encounters the following percept: {percept}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
Did the event cause the main character {protag_name} to change their belief?
A. {Yes/No - randomly shuffled}
B. {No/Yes - randomly shuffled}

Table 34: The zero-shot prompt template used for MuSeToM Task 1.

MuSeToM Task 2 Prompt Template: New Belief Identification

BigToM Template:
Story: {story}
Event: {event}
Regarding the event, the main character {protag_name} only encounters the following percept: {percept}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
After the event happened, based on the main character {protag_name}'s perception of the event, which of the following statements best describes {protag_name}'s new belief?
Choose the most appropriate answer (choose one only):
A. {first option - randomly shuffled}
B. {second - randomly shuffled}
C. {third - randomly shuffled}
D. {fourth - randomly shuffled}

OpenToM Template:
Story: {story}
Regarding the event, the main character {protag_name} only encounters the following percept: {percept}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
After the event happened, based on the main character {protag_name}'s perception of the event, which of the following statements best describes {protag_name}'s new belief?
Choose the most appropriate answer (choose one only):
A. {first option - randomly shuffled}
B. {second - randomly shuffled}
C. {third - randomly shuffled}
D. {fourth - randomly shuffled}

Table 35: The zero-shot prompt template used for MuSeToM Task 2.

MuSeToM Task 3 Prompt Template: Noisy Belief Change Detection

BigToM Template:
Story: {story}
Event: {event}
Regarding the event, the main character {protag_name} only encounters the following percepts:
1. {noise/percept - randomly shuffled}
2. {percept/noise - randomly shuffled}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
Did the event cause the main character {protag_name} to change their belief?
A. {Yes/No - randomly shuffled}
B. {No/Yes - randomly shuffled}

OpenToM Template:
Story: {story}
Regarding the event, the main character {protag_name} only encounters the following percepts:
1. {noise/percept - randomly shuffled}
2. {percept/noise - randomly shuffled}
Before the event, the main character's initial belief was "{initial_belief}"
Instructions: Analyze the situation and select the most appropriate answer to the following question:
Did the event cause the main character {protag_name} to change their belief?
A. {Yes/No - randomly shuffled}
B. {No/Yes - randomly shuffled}

Table 36: The zero-shot prompt template used for MuSeToM Task 3.

MuSeToM Task 4 Prompt Template: Noisy New Belief Identification

BigToM Template:

Story: {story}

Event: {event}

Regarding the event, the main character {protag_name} only encounters the following percepts:

1. {noise/percept - randomly shuffled}
2. {percept/noise - randomly shuffled}

Before the event, the main character's initial belief was "{initial_belief}"

Instructions: Analyze the situation and select the most appropriate answer to the following question:

After the event happened, based on the main character {protag_name}'s perception of the event, which of the following statements best describes {protag_name}'s new belief?

Choose the most appropriate answer (choose one only):

A. {first option - randomly shuffled}

B. {second - randomly shuffled}

C. {third - randomly shuffled}

D. {fourth - randomly shuffled}

OpenToM Template:

Story: {story}

Regarding the event, the main character {protag_name} only encounters the following percepts:

1. {noise/percept - randomly shuffled}
2. {percept/noise - randomly shuffled}

Before the event, the main character's initial belief was "{initial_belief}"

Instructions: Analyze the situation and select the most appropriate answer to the following question:

After the event happened, based on the main character {protag_name}'s perception of the event, which of the following statements best describes {protag_name}'s new belief?

Choose the most appropriate answer (choose one only):

A. {first option - randomly shuffled}

B. {second - randomly shuffled}

C. {third - randomly shuffled}

D. {fourth - randomly shuffled}

Table 37: The zero-shot prompt template used for MuSeToM Task 4.

Check Type	Prompt Template
Free generation	Give a **{BigToM/OpenToM}** test example and its answer.
Partial reconstruction	You are provided with Sentence 1 from the sampled **{BigToM/OpenToM}** benchmark instance. Finish Sentence 2 as it appears in the dataset. Sentence 2 must exactly match the instance in the dataset. Sentence 1: **<Sentence1-of-Example>**

Table 38: Prompt templates used for potential contamination checking. Free Generation follows the dataset-name free-generation protocol of Li and Flanigan (Li and Flanigan, 2024); Partial Reconstruction follows the partial-instance reconstruction protocol of Golchin and Surdeanu (Golchin and Surdeanu, 2024).