arXiv:2611.16748v1 [cs.CV] 12 Nov 2026

PhysLoop: Harnessing LLMs' Physics Awareness for Self-Improving Video Generation Models

Yuxuan Chen[1*], **Xiaojie Li**[2*], **Hao Chen**[1], **Juncheng Ma**[1,2], **Kewei Zhang**[1], **Zhongjing Du**[1], **Guoqiang Wei**[2], **Fangfang Wang**[6], **Zhen Dong**[3], **Zhiyang Dou**[4], **Liming Zhao**[2‡], **Yan Zeng**[2], **Jianwei Zhang**[5†§], **Daquan Zhou**[1†§]

[1]Peking University, [2]ByteDance Seed, [3]University of California, Santa Barbara, [4]Massachusetts Institute of Technology, [5]Universität Hamburg, [6]Hangzhou Normal University

[*]Equal contributions, [‡]Industry advisor, [†]Project lead, [§]Corresponding authors

Abstract

Video generation models produce visually realistic videos but still struggle with physical dynamics and causal consistency. Learning from generated videos requires fine-grained physical supervision, yet reliable quality judgments, correction pairs, and controlled variations are difficult to obtain. Large language models (LLMs) possess physical knowledge, but struggle to reliably guide video generation models in producing their training videos. We present **PhysLoop**, a unified harness designed to support physically grounded self-improvement in video generation models. The harness guides and constrains LLMs throughout video generation, selection, and correction, distilling their physical knowledge into physically consistent videos. The framework supports controllable training-data construction and inference-time video repair. By identifying and correcting failures, PhysLoop marks a key milestone toward recursive self-improvement (RSI) in video generation models with limited human intervention. We release 100K videos spanning diverse, realistic indoor and outdoor scenes, along with generation tools and two evaluation datasets: PhysLoop-Gen and PhysLoop-Repair. We evaluate controllable generation and localized repair on these datasets, respectively, and assess physical behavior on the external Physics-IQ Verified benchmark.

Project Page: https://chenyx24.github.io/PhysLoop-site/
GitHub Repo: https://github.com/DAGroup-PKU/PhysLoop
Corresponding Authors: Jianwei Zhang; Daquan Zhou

1 Introduction

Video generation models can produce realistic dynamic scenes with temporal coherence and viewpoint consistency [24, 36]. However, visual realism does not ensure physical or causal consistency: objects may disappear or multiply, and contact may fail to produce a plausible response [3, 4, 29, 33]. Other common failures include unintended slow-motion effects and changes in speed or direction that are inconsistent with gravity, inertia, or contact forces [18]. Such errors disrupt the temporal and causal relationships within a scene, so a video can appear convincing frame by frame while depicting a physically implausible sequence of events.

Improving physical behavior requires training videos that systematically cover how events unfold under

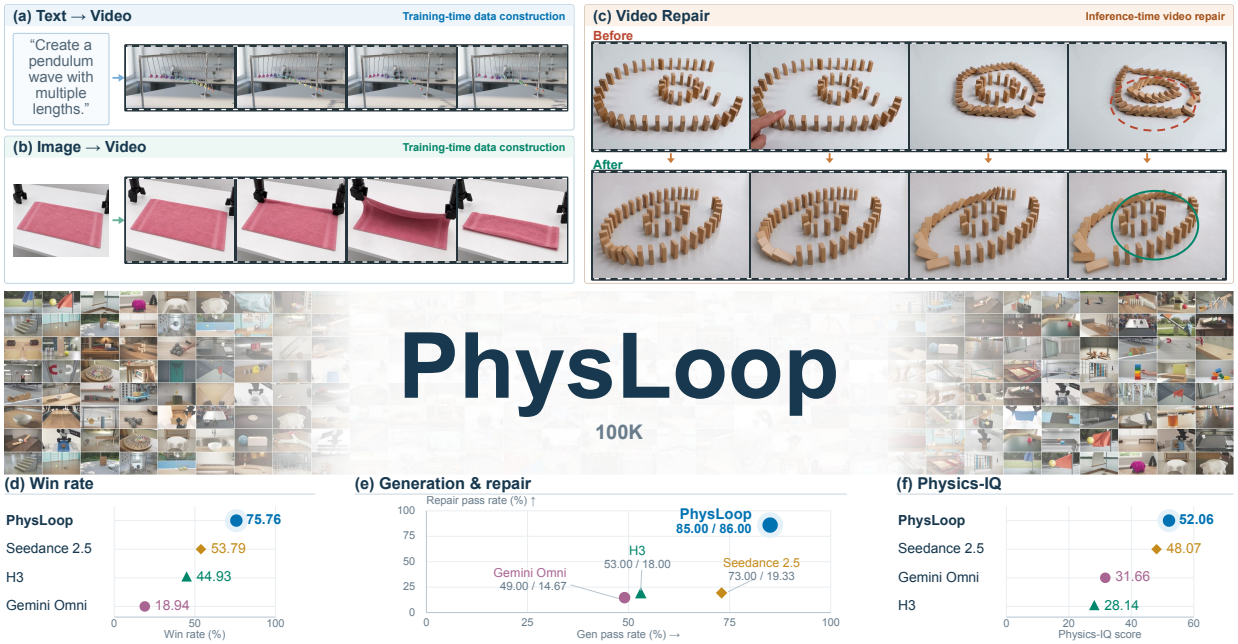

different conditions. Such data should capture physical interactions across diverse objects and scenes [54]. Collecting these videos through real-world experiments is difficult to scale while maintaining controlled conditions and reliable annotations [16]. Consider a glass falling and shattering on impact: varying its release height, initial velocity, or surroundings requires repeated setup, while unintended changes between trials make the effect of any one factor difficult to isolate. This collection burden also makes it difficult to obtain targeted fine-tuning data for physical behaviors that a model handles poorly, since suitable examples require deliberate coverage of the relevant conditions and reliable annotations. The challenge is therefore to obtain diverse physical videos together with dependable signals for identifying and correcting errors.

A long-term direction is recursive self-improvement (RSI), in which video generation models produce and learn from their own training data with minimal human intervention. Recent work such as KBAlign [46] and SEAL [55] adapts language models using self-generated training data, while Absolute Zero [51] learns to propose and solve tasks through verifiable execution feedback. For video generation, such a loop would first autonomously identify weaknesses in the model's physical behavior and generate candidate videos targeting them, then assess and refine the candidates and select suitable data for fine-tuning. The updated model would then generate the next round of candidates, whose quality would again be assessed. However, current video models still produce physical errors, and general-purpose multimodal evaluators can struggle to judge their physical validity reliably [4]. Generating more videos therefore does not by itself ensure that the model receives better training data. A key requirement is to identify what is physically wrong, guide its correction, and turn the resulting videos into dependable learning signals. PhysLoop addresses this need through a unified harness that guides LLMs to generate video data in a sandbox, applies physical constraints through an internal simulator, and iteratively refines the data through physical assessment and correction.

LLMs already possess physical knowledge that supports reasoning about physical processes, but they do not apply it consistently across tasks [37]. Realizing that knowledge in a video requires spatial relationships, physical conditions, and motion to remain consistent throughout an event; choices that appear plausible individually may still produce an incorrect outcome when combined [26]. Our insight is to guide and constrain the application of this existing knowledge within an executable physical environment. The simulator inside PhysLoop supplies physical constraints and observable outcomes, allowing the harness to check whether the generated behavior satisfies the intended physical requirements. Feedback from these outcomes guides the LLM in correcting how its knowledge is applied, while repeated assessment determines which videos are

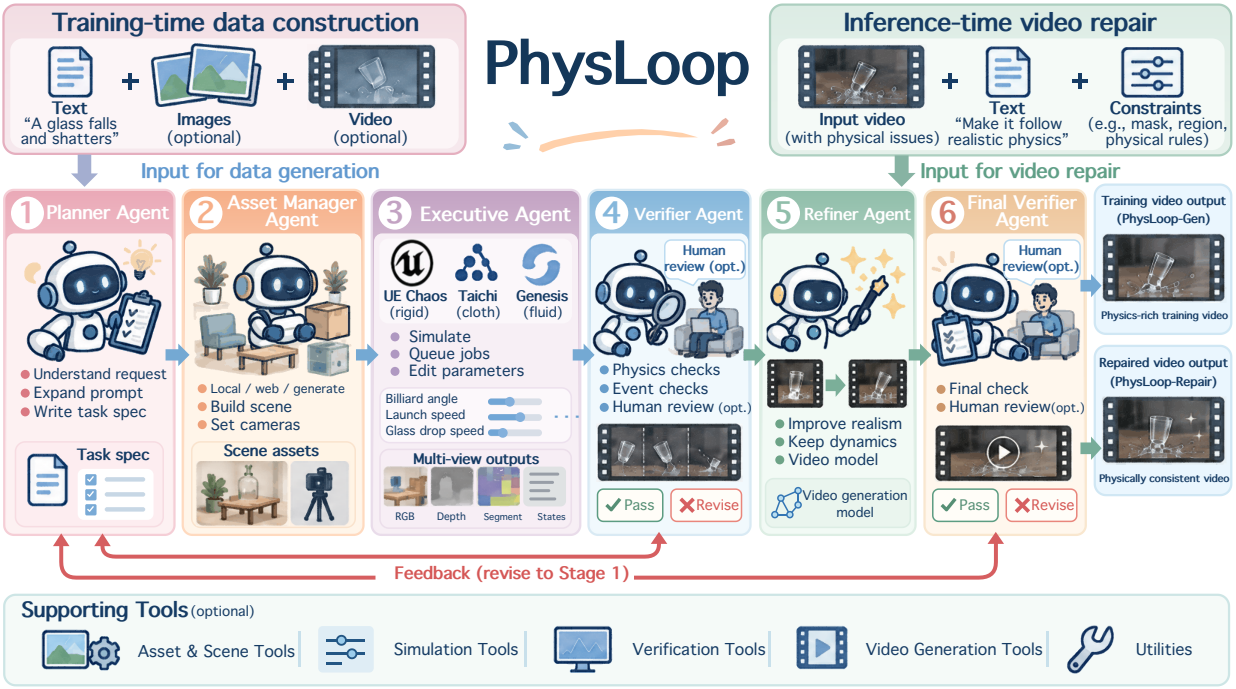

Figure 1 Overview of PhysLoop and its two applications: training-data construction and video repair. The harness guides LLMs to identify physical errors and generate physically consistent videos through simulators and other tools.

suitable to retain. The accepted videos serve as physical supervision from which video generation models can learn, providing a path to distilling the LLM's knowledge into their generated behavior. This shared principle supports controllable training-data construction and the assessment and repair of physical errors in generated videos.

We make the following contributions:

- **A unified harness.** We introduce PhysLoop, a unified harness that guides LLMs to turn existing physical knowledge into video data under simulator constraints, integrating physical assessment and iterative correction to support video-model self-improvement.
- **Datasets and generation tools.** We release 100K videos covering common physical phenomena in diverse, realistic indoor and outdoor scenes, alongside the PhysLoop-Gen and PhysLoop-Repair evaluation datasets and generation tools.
- **Evaluation across generation and repair.** We evaluate controllable generation, localized repair, and physical behavior on PhysLoop-Gen, PhysLoop-Repair, and the external benchmark [38], respectively, using human judgments and automatic metrics.

2 Related Work

Self-improvement through generated data. Learning from self-generated data offers a route to model adaptation with less human annotation [42, 45]. Two relevant directions are constructing training supervision from supplied knowledge or examples and autonomously generating new learning tasks [20]. For adaptation to supplied inputs, self-verification and feedback from downstream performance offer complementary ways to assess and improve the generated supervision [46, 55]. For autonomous task generation, an executable environment can validate proposed tasks and solutions, allowing task proposal and solving to be optimized together [51]. For video models, VLM feedback guides offline fine-tuning to improve object interactions [12], while embodied video planning supports further training on successful trajectories collected from the environment [40]. These

approaches connect generated content to model updates through feedback on its validity or learning utility.

Physical supervision for video generation. Controlled physical data connect known scene conditions with observed behavior, providing structured supervision for learning physical dynamics. Task-oriented datasets cover stability and support, temporal and causal reasoning, and physical prediction across object interactions [5, 17, 43]. A complementary direction varies latent physical properties to test whether models infer dynamics from observed motion and interactions [8, 41]. Reusable simulation platforms provide annotated scenes and multimodal observations, supporting data construction beyond individual tasks [13, 16]. A unified physical representation allows captured 3D Gaussian assets, meshes, and fluids to interact within the same simulation [28]. Integrated suites combine physical domains and visual learning tasks [54], while simulation-based supervision and feedback from vision-language models support learning physical controls for video generation [34].

LLM-guided simulation and physical video generation. Work on language and vision-language agents spans interactive simulator access, programmatic scene construction, and coordinated generation of dynamic physical worlds [23, 39, 44, 47]. Visual feedback and physical constraints help refine language-generated scene layouts while avoiding interpenetration and unstable support [27]. For video generation, simulated trajectories supply motion guidance [11], while agentic planning uses tool execution and verification to revise generation decisions [10]. Stage-specific verification of reconstructed scenes and simulated events supports targeted program repair before video synthesis [26]. When simulated motion guides visual refinement, the final video still needs to preserve the dynamics verified in simulation.

Physical assessment and video correction. Video assessment separates perceptual quality from physical behavior. Perceptual protocols evaluate appearance and temporal coherence [21], while complementary benchmarks assess physical commonsense, rule adherence, and consistency with real-world principles [3, 4, 29, 52]. Reference-based tests compare generated motion with recorded experiments or simulated references [33, 38, 54], and temporally localized diagnoses identify when physical failures occur and explain them [48]. At inference, self-refining sampling revises denoising predictions without external verifiers [22], while VLM feedback supports iterative refinement and selection of video plans for embodied tasks [40]. Motion controllers instead guide generation through user-specified trajectories or optical flow [6, 50]. For existing videos, diagnosis-guided localized refinement corrects text–video misalignment while preserving faithful regions [25]. Inpainting fills missing regions [49, 53], and correspondence-based editing maintains consistency across frames [14]. These methods provide different forms of control, but visual coherence or adherence to a prescribed motion alone does not establish physical correctness. PhysLoop connects diagnosis to simulation-grounded correction and evaluates both the resolution of the target physical error and the preservation of unaffected content.

3 PhysLoop

3.1 Workflow

PhysLoop takes a text request, optionally accompanied by image or video references, and produces a generated or repaired video together with its task specification and execution record. A code agent uses harness tools to specify the task, acquire assets, and construct an executable simulator scene. The harness includes a task manager for converting user requests into task specifications, an asset-management and scene-generation module, a simulator execution module, a verifier, and a video refiner, together with supporting tools. Simulation verification precedes appearance refinement, and final-video verification follows it. The stages of data construction and video repair have different task constraints and output purposes. Figure 1 illustrates the harness components and the workflow shared by data construction and video repair.

3.2 Task Specification and Scene Construction

The agent expands the request into a task specification that records the target event, construction assumptions, and conditions used for verification. Text specifies the intended behavior, while images and videos supply visual references. For a falling-glass request, the agent must choose a release height and an impact surface when

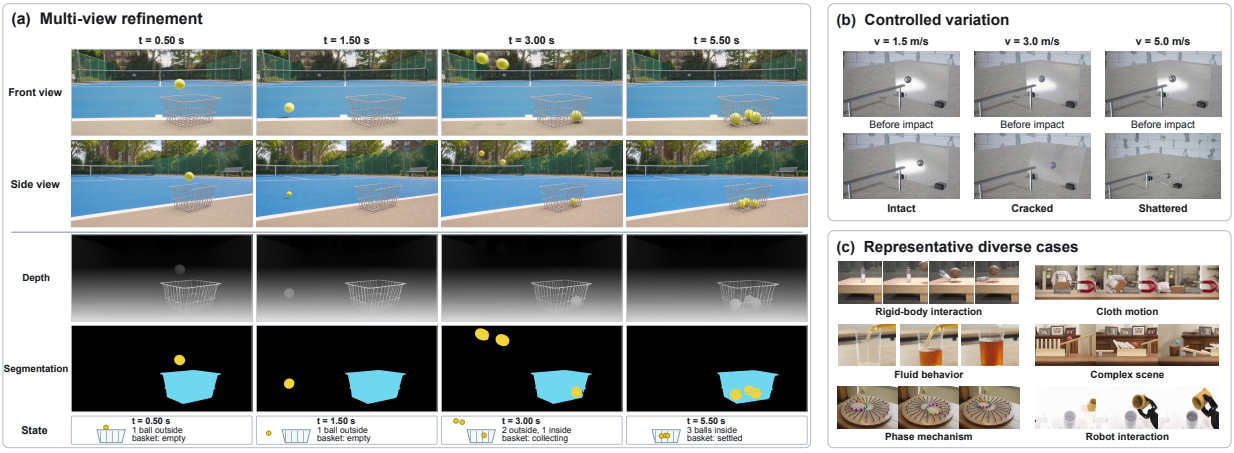

Figure 2 Dataset overview: (a) multi-views, depth, segmentation, and states; (b) controlled variations; (c) representative diverse cases.

these are unspecified. These assumptions are recorded alongside the expected outcome so that construction and verification refer to the same task.

The agent retrieves assets that meet the visual and interaction requirements, using external search or Meshy [30] to obtain or generate 3D assets when suitable objects are unavailable. Before scene construction, the harness checks asset scale and collision geometry against the target interaction. Visual meshes and physical representations are handled separately, so detailed appearance can be paired with an appropriate collision proxy. For example, a glass can retain a detailed rendering mesh while a separate physical representation determines its contact with the ground.

The agent writes a scene program that combines retrieved assets and procedural geometry to specify object placement and relationships. For a falling glass, the program sets the release position and impact surface, while the physics solver computes the subsequent motion. The program can place the event in different settings and compose structures that are not predefined in the asset catalog.

Before full simulation, the harness renders a low-resolution preview frame from each planned camera view to inspect asset placement and visibility of the interaction region. If an object is misplaced or an interaction is occluded, the agent adjusts the layout or camera and repeats the preview before full simulation and video rendering.

3.3 Physics Execution and Video Refinement

Unreal Engine (UE) [9] is the primary simulator integrated into the harness, while Taichi [19] and Genesis [2] provide additional support. Given the initial conditions and permitted actions specified by the agent, the physics solver numerically computes the resulting motion according to the physical model of the scene. Simulator adapters are conversion interfaces that translate scene configurations into each simulator's inputs and transfer the resulting motion to UE with aligned coordinates and timing.

Multiple camera views share the same simulation timeline. Alongside RGB, the harness captures depth and instance segmentation and records object states and backend-provided contact or event signals. Each frame is linked to its simulation time and camera configuration, allowing verification to relate visible behavior to physical states.

After simulation checks pass, appearance refinement adapts the rendered scene to the requested visual appearance. The video refiner takes the verified simulation video as a motion reference, together with the task description and visual references that specify the desired appearance.

3.4 Verification and Revision

Successful execution does not establish that the requested event occurred. The verifier compares simulation states and visual observations with the predefined event requirements and physical conditions. For a collision task, it checks whether contact occurred and whether the subsequent motion satisfies the specification. Conservation checks account for external forces, dissipation, and numerical error.

The final-video check evaluates whether visible behavior preserves the target event and its physical relationships after refinement. Refinement may alter motion that passed simulation checks, so the verifier reassesses the refined result against the same event requirements.

If either simulation or final-video verification fails, the verifier returns a failure description with supporting simulation or visual evidence to the agent. Using this feedback and the recorded history, the agent revises the configuration for the responsible stage. A missed collision may require changes to the scene or initial conditions, while a refinement error may require different refiner inputs.

The task objective and verification criteria remain fixed throughout revision. Execution ends when all required checks pass or the attempt or resource budget is exhausted; unresolved failures remain in the run record. The specification, configurations, observations, and revision history remain linked in a persistent run record. Users can participate mainly during verification by reviewing candidate results, with their judgments recorded together with automated checks and agent revisions.

3.5 Application: Training-Data Construction

The agent exposes task-specific parameters in the generated scene configuration, such as release height and initial velocity for a falling glass. The agent or user varies selected parameters while holding the remaining conditions fixed, producing different instances of the same event. Scene and camera settings can also be edited, and multiple views can be rendered from the same simulation.

For data construction, the task description and visual references specify the appearance to be produced. Videos that pass final verification are paired with their task descriptions as training examples for video generation models. Each sample retains its generating configuration for inspection, subsequent expansion, and regeneration. Downstream training can use these video–text pairs to update the video generation model.

Raw simulation renders retain their synchronized depth, segmentation, and object-state annotations. Because refinement may alter geometry or object positions, these annotations are not treated as pixel-aligned supervision for refined videos. The annotated simulation renders remain a separate resource from the refined video–text pairs.

3.6 Application: Video Diagnosis and Repair

For an existing generated video, the agent first identifies the physical or event-level error and where it occurs, then specifies the event to reconstruct and the segment to replace. Frames before and after the selected segment constrain how the corrected event must connect to the source and which content must be preserved. This diagnosis and the source constraints define the task specification used by the shared scene-construction and simulation stages. During refinement, the source video supplies the appearance reference for the reconstructed objects and environment.

Repair replaces only the selected interval $[a, b)$ of a source video V with T frames. The corrected segment $\hat{V}_{a:b}$ matches the source resolution and frame rate and contains exactly $b - a$ frames. The output is assembled as

$$V' = V_{0:a} \parallel \hat{V}_{a:b} \parallel V_{b:T}, \tag{1}$$

where $\parallel$ denotes concatenation. Blending remains within the replacement interval, leaving source frames outside it unchanged.

The replacement is evaluated within the complete video because a plausible isolated segment may still conflict with the surrounding event. The verifier checks target-error correction, consistency with preceding and following frames, and preservation of unaffected content. If refinement changes the intended physical behavior

Method	Human (Gen)				Model (Gen)				Task		
	Phys.	Real.	Causal	Win	PC	SA	IQ	AQ	Gen	Repair	P-IQ
Video Generation Models											
Gemini Omni	3.11	3.82	3.18	18.94	20.79	**39.17**	71.17	50.40	49.00	14.67	31.66
H3	3.63	4.27	3.55	44.93	**22.98**	38.68	68.47	50.62	53.00	18.00	28.14
Seedance 2.5	3.80	**4.29**	3.80	53.79	21.51	38.69	67.63	**50.63**	73.00	19.33	48.07
Seedance 2.5											
Claude Opus 5	3.82	3.82	3.84	64.18	20.23	36.69	67.66	49.71	71.00	78.67	**53.75**
Kimi K3	3.63	3.69	3.69	54.41	19.90	37.32	68.54	50.42	69.00	68.67	52.04
GPT-6 Astra											
Gemini Omni	3.53	2.96	3.51	45.45	12.71	35.84	**75.11**	46.72	67.00	62.00	12.08
H3	3.86	3.45	3.90	58.21	17.48	37.12	68.77	48.65	73.00	72.67	39.10
w/o Refiner	3.94	2.59	3.95	34.33	12.53	34.88	72.42	45.75	80.00	60.67	–
GPT-6 Astra + Seedance 2.5											
PhysLoop	**4.06**	3.72	**4.12**	**75.76**	15.29	36.28	69.93	48.70	**85.00**	**86.00**	52.06

Table 1 Evaluation of PhysLoop across different model combinations and video generation models. Win is computed from 301 unblinded supplementary pairwise comparisons by one author, with ties counted as half a win.

or the segment fails to connect to the source, the resulting evidence guides further revision under the same stopping conditions.

4 PhysLoop Dataset

Our data resources comprise the PhysLoop-100K training set and two evaluation sets, PhysLoop-Gen and PhysLoop-Repair. PhysLoop-100K provides video–text pairs and simulation annotations, while the evaluation sets assess generation and physical-error repair, respectively.

4.1 PhysLoop-100K

PhysLoop-100K covers broad physical categories, including kinematics, dynamics, rigid-body dynamics, magnetism, fluid mechanics, flexible and elastic materials, and fragmentation and fracture. For each case, the system automatically varies relevant physical parameters, such as mass and initial velocity, together with assets, scenes, lighting, and camera views. The dataset includes all valid combinations of the parameter settings defined for each case, allowing users to select configurations of interest.

PhysLoop-100K contains 100K accepted videos from 2.5K distinct physical events. Each event contains an average of 40 videos, captured from 3-5 perspectives and covering 2-4 key parameters that can be configured to multiple levels. Alongside refined video–text pairs and their generating configurations, the dataset provides the original simulation videos and associated annotations. Depth maps, instance segmentation maps, and object states are provided as additional information and are not pixel-aligned with the refined videos. The split protocol groups each source event with its camera views and parameter variants. Coverage statistics, parameter ranges, and split details are provided in the supplementary material.

4.2 PhysLoop-Gen

PhysLoop-Gen addresses generation requests specifying a target physical event. Each instance contains a text request, with image references provided for some cases. The requests cover rigid-body, cloth, and fluid phenomena, ranging from simple bouncing or sliding to coupled interactions and multistage mechanical chains.

Task categories and difficulty are defined from input properties, including interaction count, event ordering, and observable conditions, rather than from a particular model's performance.

PhysLoop-Gen contains 100 generation tasks. The tasks support comparisons among video generation approaches, including agentic frameworks such as PhysLoop and direct video generation models. Task descriptions and representative examples are provided in the supplementary material.

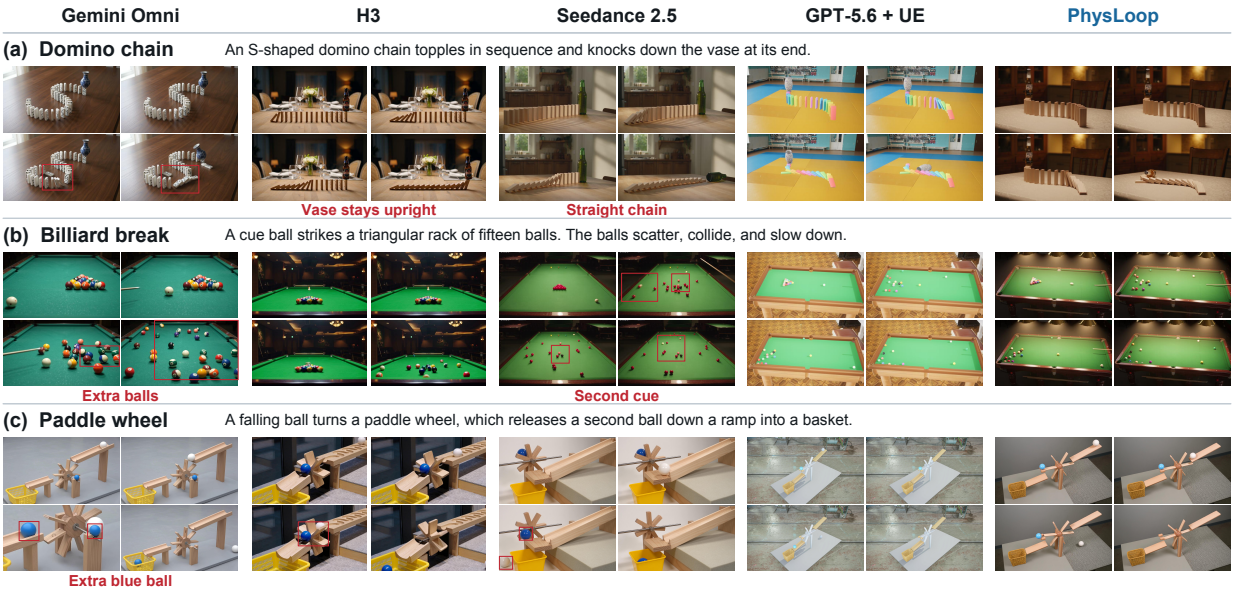

Figure 3 Qualitative comparison of PhysLoop and baseline methods.

4.3 PhysLoop-Repair

PhysLoop-Repair addresses existing videos with identifiable physical or event-level errors. Each instance contains a source video with a physical error, an optional repair request, and annotations describing the error and its temporal interval. Annotations locate the error and specify conditions for a successful correction, such as maintaining a falling object's presence and motion consistent with the event conditions. The primary evaluation protocol supplies only the source video and request; error and interval annotations remain evaluation-only.

Source videos are generated by different video models, including Seedance series, MiniMax H3, and Gemini Omni 1.1 Flash, yielding 50 repair tasks. Inclusion requires an observable error, a specified correction, and source content to preserve. Evaluation checks target-error removal, connection to surrounding events, and preservation of unrelated content.

5 Evaluation

5.1 Models and Experimental Settings

We evaluate video generation on PhysLoop-Gen, video repair on PhysLoop-Repair, and how well generated video continuations match recorded physical events on Physics-IQ Verified. The default configuration uses GPT-6 Astra [35] as the agent backbone and Seedance 2.5 [7] for video refinement. With the refiner fixed, we compare the default backbone against Claude Opus 5 [1] and Kimi K3 [32]. With the backbone fixed, we compare the default refiner against MiniMax H3 [31], Gemini Omni 1.1 Flash [15], and a no-refiner configuration. Results for additional models are provided in the supplementary material.

Across the three benchmarks, both PhysLoop and standalone baselines allow up to five repeat calls to the video generation model per task and return one final result. The supplementary material provides model

versions and settings, prompts, output examples, costs, and runtimes.

5.2 Metrics

Human Evaluation. Human and model-based quality metrics are restricted to PhysLoop-Gen. Human raters score physical plausibility, visual realism, and causal consistency on five-point scales. Overall win rate reflects pairwise preferences for physical and visual quality, with ties counting as half a win. Each video is assigned to multiple raters, with method identities hidden and presentation order randomized.

Automatic Evaluation. VideoPhy's VideoCon-Physics evaluator [3] provides Auto-PC for physical commonsense and Auto-SA for adherence to the shared task description. VBench [21] measures imaging quality (IQ) and aesthetic quality (AQ). For each metric, we average scores over cases valid across compared configurations.

Task-Specific Metrics. Task-specific performance is the primary criterion, while human ratings provide the main assessment of generated-video quality. Compared with standalone Seedance 2.5, PhysLoop scores slightly lower on SA and AQ, lower on PC, and higher on IQ; the disagreement on physical plausibility may partly reflect limitations of automatic evaluation. Following prior work, we treat automatic metrics only as supporting references: they may favor visual styles common to video generators, while physical errors unfolding over time can be difficult to detect from sparsely sampled frames [4, 54].

5.3 Result Analysis

Overall Comparison. Table 1 summarizes results on PhysLoop-Gen, PhysLoop-Repair, and Physics-IQ Verified for standalone video generators and PhysLoop configurations. On PhysLoop-Gen, PhysLoop achieves an 85% task-success rate, compared with 73% for standalone Seedance 2.5. It also receives higher human ratings for physical plausibility (4.06 vs. 3.80) and causal consistency (4.12 vs. 3.80), although its visual realism rating is lower than that of standalone Seedance 2.5 (3.72 vs. 4.29). The gap in perceived realism is most apparent in material appearance and backgrounds.

On PhysLoop-Repair, PhysLoop achieves an 86.00% success rate, compared with 19.33% for the strongest standalone baseline, Seedance 2.5. All three standalone generators remain below 20%, showing that they struggle to satisfy the repair task's combined requirements. A successful repair must correct the physical error while preserving the required source content and connecting coherently to surrounding events. PhysLoop expresses these requirements in an explicit reconstruction task and checks both the simulated event and the refined video against them. These results support the shared workflow of scene construction, simulation, and verification for correcting existing videos, extending its utility beyond controllable generation.

With GPT-6 Astra fixed as the agent backbone, the no-refiner configuration already achieves 80.00% generation success, indicating that scene construction and simulation can realize many of the requested events. Adding Seedance 2.5 refinement raises generation success to 85.00% and repair success from 60.67% to 86.00%. Human-rated visual realism increases from 2.59 to 3.72, accompanied by slightly higher ratings for physical plausibility and causal consistency. In this configuration, refinement improves appearance alongside task completion. Results also vary across refiners: H3 and Gemini Omni yield lower success rates on both tasks under the same backbone. This variation motivates evaluating refiners within the complete workflow, where the final video must satisfy both the requested appearance and the physical requirements of the task.

On Physics-IQ Verified, the default configuration scores 52.06, compared with 48.07 for standalone Seedance 2.5. With Seedance 2.5 fixed as the refiner, Claude Opus 5 achieves the highest score in our comparison (53.75), despite lower success on PhysLoop-Gen and PhysLoop-Repair. The differing rankings suggest that task completion and agreement with an observed physical continuation capture distinct aspects of performance. A plausible event may still differ from the specific dynamics of the reference video. Refiner choice also matters: with GPT-6 Astra fixed, H3 and Gemini Omni score 39.10 and 12.08, respectively, both below the default Seedance 2.5 configuration.

Evaluation Interpretation. Task-specific performance is the primary criterion, while human ratings provide the main assessment of generated-video quality. The full PhysLoop configuration achieves the highest scores across all human-evaluation dimensions, although its automatic scores are slightly lower than those of standalone video generators. Following prior work, we treat automatic metrics only as supporting references: they may favor visual styles common to video generators, while physical errors unfolding over time can be difficult to detect from sparsely sampled frames [4, 54].

5.4 Case Analysis

Figure 3 shows why plausible local motion is insufficient for judging task completion. In the domino example, H3 depicts sequential toppling but leaves the vase upright, so the intended interaction chain remains incomplete. Standalone Seedance 2.5 replaces the requested S-shaped layout with a straight chain, violating the task's spatial specification. In the billiard example, Gemini Omni introduces extra balls and Seedance 2.5 adds an extra cue ball, altering the object set required by the task. These failures concern the event outcome and scene consistency, which visual realism alone cannot establish. The UE-only and refined PhysLoop outputs both depict the rack breaking apart, while differing in material appearance and lighting. This comparison illustrates how event realization and appearance refinement serve complementary purposes in the pipeline.

6 Limitations

PhysLoop depends on the coverage of its simulators, assets, and verification conditions. Reconstructing a video requires assumptions about unobserved geometry, materials, and initial states, so a valid simulation may not recover the original physical process. Appearance refinement can introduce new errors, and local replacement may be insufficient when a mistake has already changed later events. These limitations constrain which videos can be accepted as learning signals and may require human review for ambiguous cases.

Extending PhysLoop to a complete RSI training loop requires connecting weakness discovery, candidate generation, verification and repair, data selection, and model updates across successive rounds. Each updated video model should influence which failures are targeted and which candidates are generated next. Future work must establish whether this process yields sustained gains on held-out events while preserving video quality and data diversity. Independent evaluation, tracking the effects of incorrect quality judgments, and controlling simulation and training costs will be necessary to support reliable self-improvement with decreasing human intervention.

7 Conclusion

We introduced PhysLoop, a unified harness that guides LLMs to realize physical knowledge as videos under simulator constraints. Its shared workflow supports controllable training-data construction and video repair through failure identification, correction, and verification. PhysLoop provides a quality-control component toward recursive self-improvement, enabling video models to generate, assess, and improve learning data with limited human intervention.

References

[1] Anthropic. Introducing Claude Opus 5. https://www.anthropic.com/news/claude-opus-5, 2026. Official announcement; accessed September 17, 2026.

[2] Genesis Authors. Genesis: A generative and universal physics engine for robotics and beyond, December 2024. URL https://github.com/Genesis-Embodied-AI/genesis-world.

[3] Hritik Bansal, Zongyu Lin, Tianyi Xie, Zeshun Zong, Michal Yarom, Yonatan Bitton, Chenfanfu Jiang, Yizhou Sun, Kai-Wei Chang, and Aditya Grover. Videophy: Evaluating physical commonsense for video generation. In *International Conference on Learning Representations*, volume 2025, pages 102075–102121, 2025.

[4] Hritik Bansal, Clark Peng, Yonatan Bitton, Roman Goldenberg, Aditya Grover, and Kai-Wei Chang. Videophy-2: A challenging action-centric physical commonsense evaluation in video generation. In *International Conference on Learning Representations*, volume 2026, pages 118456–118470, 2026.

[5] Daniel M. Bear, Elias Wang, Damian Mrowca, Felix J. Binder, Hsiao-Yu Fish Tung, R. T. Pramod, Cameron Holdaway, Sirui Tao, Kevin Smith, Fan-Yun Sun, Li Fei-Fei, Nancy Kanwisher, Joshua B. Tenenbaum, Daniel L. K. Yamins, and Judith E. Fan. Physion: Evaluating physical prediction from vision in humans and machines. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021. URL https://datasets-benchmarks-proceedings.neurips.cc/paper_files/paper/2021/hash/d09bf41544a3365a46c9077ebb5e35c3-Abstract-round1.html.

[6] Ryan Burgert, Yuancheng Xu, Wenqi Xian, Oliver Pilarski, Pascal Clausen, Mingming He, Li Ma, Yitong Deng, Lingxiao Li, Mohsen Mousavi, Michael Ryoo, Paul Debevec, and Ning Yu. Go-with-the-flow: Motion-controllable video diffusion models using real-time warped noise. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 13–23, June 2025.

[7] ByteDance Seed. Seedance 2.5. https://seed.bytedance.com/en/seedance2_5, 2026. Official model page; accessed September 17, 2026.

[8] Zhenfang Chen, Kexin Yi, Yunzhu Li, Mingyu Ding, Antonio Torralba, Joshua B. Tenenbaum, and Chuang Gan. ComPhy: Compositional physical reasoning of objects and events from videos. In *International Conference on Learning Representations*, 2022. URL https://arxiv.org/abs/2205.01089.

[9] Epic Games. Physics in Unreal Engine. https://dev.epicgames.com/documentation/unreal-engine/physics-in-unreal-engine. Accessed: 2026-09-15.

[10] Yuxiang Feng, Juncheng Wang, Chao Xu, Yijie Qian, Huihan Wang, Wenlong Hou, Yang Liu, Baigui Sun, Yong Liu, and Shujun Wang. Newton: Agentic planning for physically grounded video generation, 2026. URL https://arxiv.org/abs/2605.18396.

[11] Lin Geng Foo, Mark He Huang, Alexandros Lattas, Stylianos Moschoglou, Thabo Beeler, and Christian Theobalt. Physical simulator in-the-loop video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4301–4311, June 2026.

[12] Hiroki Furuta, Heiga Zen, Dale Schuurmans, Aleksandra Faust, Yutaka Matsuo, Percy Liang, and Sherry Yang. Improving dynamic object interactions in text-to-video generation with ai feedback. *arXiv preprint arXiv:2412.02617*, 2024. URL https://arxiv.org/abs/2412.02617v2. Version 2, 17 April 2026.

[13] Chuang Gan, Jeremy Schwartz, Seth Alter, Damian Mrowca, Martin Schrimpf, James Traer, Julian De Freitas, Jonas Kubilius, Abhishek Bhandwaldar, Nick Haber, Megumi Sano, Kuno Kim, Elias Wang, Michael Lingelbach, Aidan Curtis, Kevin Feigelis, Daniel M. Bear, Dan Gutfreund, David Cox, Antonio Torralba, James J. DiCarlo, Joshua B. Tenenbaum, Josh H. McDermott, and Daniel L. K. Yamins. ThreeDWorld: A platform for interactive multi-modal physical simulation. In *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, volume 1, 2021. URL https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/735b90b4568125ed6c3f678819b6e058-Abstract-round1.html.

[14] Michal Geyer, Omer Bar Tal, Shai Bagon, and Tali Dekel. Tokenflow: Consistent diffusion features for consistent video editing. In B. Kim, Y. Yue, S. Chaudhuri, K. Fragkiadaki, M. Khan, and Y. Sun, editors, *International Conference on Learning Representations*, volume 2024, pages 1608–1620, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/file/069274129275ec1c6511abb5f6d96a31-Paper-Conference.pdf.

[15] Google. Gemini Omni Flash. https://ai.google.dev/gemini-api/docs/models/gemini-omni-flash, 2026. Model documentation; accessed September 17, 2026.

[16] Klaus Greff, Francois Belletti, Lucas Beyer, Carl Doersch, Yilun Du, Daniel Duckworth, David J. Fleet, Dan Gnanapragasam, Florian Golemo, Charles Herrmann, Thomas Kipf, Abhijit Kundu, Dmitry Lagun, Issam Laradji, Hsueh-Ti (Derek) Liu, Henning Meyer, Yishu Miao, Derek Nowrouzezahrai, Cengiz Oztireli, Etienne Pot, Noha Radwan, Daniel Rebain, Sara Sabour, Mehdi S. M. Sajjadi, Matan Sela, Vincent Sitzmann, Austin Stone, Deqing Sun, Suhani Vora, Ziyu Wang, Tianhao Wu, Kwang Moo Yi, Fangcheng Zhong, and Andrea Tagliasacchi. Kubric: A scalable dataset generator. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3749–3761, 2022. URL https://openaccess.thecvf.com/content/CVPR2022/html/Greff_Kubric_A_Scalable_Dataset_Generator_CVPR_2022_paper.html.

[17] Oliver Groth, Fabian B. Fuchs, Ingmar Posner, and Andrea Vedaldi. ShapeStacks: Learning vision-based physical intuition for generalised object stacking. In *Proceedings of the European Conference on Computer Vision*, 2018. URL https://openaccess.thecvf.com/content_ECCV_2018/html/Oliver_Groth_ShapeStacks_Learning_Vision-Based_ECCV_2018_paper.html.

[18] Jing Gu, Xian Liu, Yu Zeng, Ashwin Nagarajan, Fangrui Zhu, Daniel Hong, Yue Fan, Qianqi Yan, Kaiwen Zhou, Ming-Yu Liu, and Xin Eric Wang. PhyWorldBench: A comprehensive evaluation of physical realism in text-to-video models. In *The Fourteenth International Conference on Learning Representations*, 2026. URL https://openreview.net/forum?id=rlZeILv3fm.

[19] Yuanming Hu, Tzu-Mao Li, Luke Anderson, Jonathan Ragan-Kelley, and Frédo Durand. Taichi: A language for

high-performance computation on spatially sparse data structures. *ACM Transactions on Graphics (TOG)*, 38(6): 201, 2019.

[20] Chengsong Huang, Wenhao Yu, Xiaoyang Wang, Hongming Zhang, Zongxia Li, Ruosen Li, Jiaxin Huang, Haitao Mi, and Dong Yu. R-Zero: Self-evolving reasoning LLM from zero data. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/d49b9aacebda61051166335af6fd3061-Abstract-Conference.html.

[21] Ziqi Huang, Yinan He, Jiashuo Yu, Fan Zhang, Chenyang Si, Yuming Jiang, Yuanhan Zhang, Tianxing Wu, Qingyang Jin, Nattapol Chanpaisit, Yaohui Wang, Xinyuan Chen, Limin Wang, Dahua Lin, Yu Qiao, and Ziwei Liu. Vbench: Comprehensive benchmark suite for video generative models, 2023. URL https://arxiv.org/abs/2311.17982.

[22] Sangwon Jang, Taekyung Ki, Jaehyeong Jo, Saining Xie, Jaehong Yoon, and Sung Ju Hwang. Self-refining video sampling, 2026. URL https://arxiv.org/abs/2601.18577.

[23] Haoqiang Kang, Xiaokang Ye, Yuhan Liu, Siddhant Hitesh Mantri, Lingjun Mao, James Fleming, Drishti Regmi, and Lianhui Qin. Simworld studio: Automatic environment generation with evolving coding agent for embodied agent learning, 2026. URL https://arxiv.org/abs/2605.09423.

[24] Weijie Kong, Qi Tian, Zijian Zhang, Rox Min, Zuozhuo Dai, Jin Zhou, Jiangfeng Xiong, Xin Li, Bo Wu, Jianwei Zhang, et al. Hunyuanvideo: A systematic framework for large video generative models. *arXiv preprint arXiv:2412.03603*, 2024.

[25] Daeun Lee, Jaehong Yoon, Jaemin Cho, and Mohit Bansal. Self-correcting text-to-video generation with misalignment detection and localized refinement, 2026. URL https://arxiv.org/abs/2411.15115.

[26] Qirui Li, Jinkun Hao, Yibo Li, Ran Yi, Paul L. Rosin, and Yu-Kun Lai. PhysAgent: Reflective agentic physics control for physically plausible video generation. *arXiv preprint arXiv:2607.16355*, 2026. URL https://arxiv.org/abs/2607.16355.

[27] Lu Ling, Chen-Hsuan Lin, Tsung-Yi Lin, Yifan Ding, Yu Zeng, Yichen Sheng, Yunhao Ge, Ming-Yu Liu, Aniket Bera, and Max Li. Scenethesis: A language and vision agentic framework for 3d scene generation. In *International Conference on Learning Representations*, volume 2026, pages 136596–136629, 2026.

[28] Xiaoyang Liu, Shangzhe Wu, and Kai Han. Scene-level heterogeneous physics simulation with 3d gaussian splats. *arXiv preprint arXiv:2606.21753*, 2026.

[29] Fanqing Meng, Jiaqi Liao, Xinyu Tan, Wenqi Shao, Quanfeng Lu, Kaipeng Zhang, Yu Cheng, Dianqi Li, Yu Qiao, and Ping Luo. Towards world simulator: Crafting physical commonsense-based benchmark for video generation. *arXiv preprint arXiv:2410.05363*, 2024.

[30] Meshy. Meshy: Text to 3d api. https://docs.meshy.ai/en/api/text-to-3d. Accessed: 2026-09-15.

[31] MiniMax. MiniMax H3: An Open Model Breaking the Boundaries Between Tasks and Modalities. https://www.minimax.io/blog/minimax-h3, 2026. Official technical blog; accessed September 17, 2026.

[32] Moonshot AI. Kimi K3 Tech Blog: Open Frontier Intelligence. https://www.kimi.com/en/blog/kimi-k3, 2026. Official technical blog; accessed September 17, 2026.

[33] Saman Motamed, Laura Culp, Kevin Swersky, Priyank Jaini, and Robert Geirhos. Do generative video models understand physical principles? In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 948–958, 2026. URL https://openaccess.thecvf.com/content/WACV2026/html/Motamed_Do_Generative_Video_Models_Understand_Physical_Principles_WACV_2026_paper.html.

[34] Sriram Narayanan, Ziyu Jiang, Srinivasa Narasimhan, and Manmohan Chandraker. Phyco: Learning controllable physical priors for generative motion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 41892–41902, June 2026.

[35] OpenAI. GPT-6 Astra. https://developers.openai.com/api/docs/models/gpt-6-astra, 2026. Model documentation; accessed September 17, 2026.

[36] Adam Polyak, Amit Zohar, Andrew Brown, Andros Tjandra, Animesh Sinha, Ann Lee, Apoorv Vyas, Bowen Shi, Chih-Yao Ma, Ching-Yao Chuang, et al. Movie gen: A cast of media foundation models. *arXiv preprint arXiv:2410.13720*, 2024.

[37] Shi Qiu, Shaoyang Guo, Zhuo-Yang Song, Yunbo Sun, Zeyu Cai, Jiashen Wei, Tianyu Luo, Yixuan Yin, Haoxu Zhang, Yi Hu, Chenyang Wang, Chencheng Tang, Haoling Chang, Qi Liu, Ziheng Zhou, Tianyu Zhang, Jingtian Zhang, Zhangyi Liu, Minghao Li, Yuku Zhang, Boxuan Jing, Xianqi Yin, Yutong Ren, Zizhuo Fu, Jiaming Ji, Weike Wang, Xudong Tian, Anqi Lv, Laifu Man, Jianxiang Li, Feiyu Tao, Qihua Sun, Zhou Liang, Yushu Mu, Zhongxuan Li, Jing-Jun Zhang, Shutao Zhang, Xiaotian Li, Xingqi Xia, Jiawei Lin, Zheyu Shen, Jiahang Chen, Qiuhao Xiong, Binran Wang, Fengyuan Wang, Ziyang Ni, Bohan Zhang, Fan Cui, Changkun Shao, Qing-Hong Cao, Ming-xing Luo, Yaodong Yang, Muhan Zhang, and Hua Xing Zhu. PHYBench: Holistic evaluation of

physical perception and reasoning in large language models. In D. Belgrave, C. Zhang, H. Lin, R. Pascanu, P. Koniusz, M. Ghassemi, and N. Chen, editors, *Advances in Neural Information Processing Systems*, volume 38. Curran Associates, Inc., 2025. doi: 10.52202/085713-0044. URL https://proceedings.neurips.cc/paper_files/paper/2025/file/01e793f8cf5689aa0ea46a1b01071bea-Paper-Datasets_and_Benchmarks_Track.pdf.

[38] Tim Rädsch, Yuki M. Asano, Hilde Kuehne, Stefan Bauer, Priyank Jaini, Robert Geirhos, and Carsten T. Lüth. Physics-IQ Verified. *arXiv preprint arXiv:2606.18943*, 2026. URL https://arxiv.org/abs/2606.18943.

[39] Jiawei Ren, Yan Zhuang, Xiaokang Ye, Lingjun Mao, Xuhong He, Jianzhi Shen, Mrinaal Dogra, Yiming Liang, Ruixuan Zhang, Tianai Yue, Yiqing Yang, Eric Liu, Ryan Wu, Kevin Benavente, Rajiv Mandya Nagaraju, Muhammad Faayez, Xiyan Zhang, Dhruv Vivek Sharma, Xianrui Zhong, Ziqiao Ma, Tianmin Shu, Zhiting Hu, and Lianhui Qin. Simworld: An open-ended realistic simulator for autonomous agents in physical and social worlds, 2026. URL https://arxiv.org/abs/2512.01078.

[40] Achint Soni, Sreyas Venkataraman, Abhranil Chandra, Sebastian Fischmeister, Percy Liang, Bo Dai, and Sherry Yang. VideoAgent: Self-improving video generation. *arXiv preprint arXiv:2410.10076*, 2024. URL https://arxiv.org/abs/2410.10076v3. Version 3, 9 February 2025; manuscript title adds "for Embodied Planning".

[41] Hsiao-Yu Tung, Mingyu Ding, Zhenfang Chen, Daniel Bear, Chuang Gan, Joshua B. Tenenbaum, Daniel L. K. Yamins, Judith E. Fan, and Kevin A. Smith. Physion++: Evaluating physical scene understanding that requires online inference of different physical properties. In *Advances in Neural Information Processing Systems*, 2023. URL https://papers.neurips.cc/paper_files/paper/2023/hash/d3e8011c912e651ab2a76e7935a1e464-Abstract-Datasets_and_Benchmarks.html.

[42] Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. Self-Instruct: Aligning language models with self-generated instructions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 13484–13508, 2023. doi: 10.18653/v1/2023.acl-long.754. URL https://aclanthology.org/2023.acl-long.754/.

[43] Kexin Yi, Chuang Gan, Yunzhu Li, Pushmeet Kohli, Jiajun Wu, Antonio Torralba, and Joshua B. Tenenbaum. CLEVRER: Collision events for video representation and reasoning. In *International Conference on Learning Representations*, 2020. URL https://arxiv.org/abs/1910.01442.

[44] Shaofeng Yin, Jiaxin Ge, Zora Zhiruo Wang, Chenyang Wang, Xiuyu Li, Michael J. Black, Trevor Darrell, Angjoo Kanazawa, and Haiwen Feng. Vision-as-inverse-graphics agent via interleaved multimodal reasoning, 2026. URL https://arxiv.org/abs/2601.11109.

[45] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah Goodman. STaR: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, volume 35, pages 15476–15488, 2022. doi: 10.52202/068431-1126. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/639a9a172c044fbb64175b5fad42e9a5-Abstract-Conference.html.

[46] Zheni Zeng, Yuxuan Chen, Shi Yu, Ruobing Wang, Yukun Yan, Zhenghao Liu, Shuo Wang, Xu Han, Zhiyuan Liu, and Maosong Sun. Kbalign: Efficient self adaptation on specific textual knowledge bases. In *EMNLP (Findings)*, pages 13519–13532, 2025.

[47] Hongxin Zhang, Chunru Lin, Junyan Li, Zhou Xian, Tsun-Hsuan Wang, and Chuang Gan. Gs-agent: Creating 4d physical worlds with generative simulation, 2026. URL https://arxiv.org/abs/2607.21522.

[48] Qin Zhang, Peiyu Jing, Hong-Xing Yu, Fangqiang Ding, Fan Nie, Weimin Wang, Yilun Du, James Zou, Jiajun Wu, and Bing Shuai. Physion-eval: Evaluating physical realism in generated video via human reasoning, 2026. URL https://arxiv.org/abs/2603.19607.

[49] Zhixing Zhang, Bichen Wu, Xiaoyan Wang, Yaqiao Luo, Luxin Zhang, Yinan Zhao, Peter Vajda, Dimitris Metaxas, and Licheng Yu. Avid: Any-length video inpainting with diffusion model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 7162–7172, June 2024.

[50] Zhongwei Zhang, Fuchen Long, Zhaofan Qiu, Yingwei Pan, Wu Liu, Ting Yao, and Tao Mei. Motionpro: A precise motion controller for image-to-video generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 27957–27967, June 2025.

[51] Andrew Zhao, Yiran Wu, Tong Wu, Quentin Xu, Yang Yue, Matthieu Lin, Shenzhi Wang, Qingyun Wu, Zilong Zheng, and Gao Huang. Absolute zero: Reinforced self-play reasoning with zero data. In *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pages 105816–105879, 2025. doi: 10.52202/085713-3534. URL https://proceedings.neurips.cc/paper_files/paper/2025/file/9837dc00ff67d176373268ed48042d49-Paper-Conference.pdf.

[52] Dian Zheng, Ziqi Huang, Hongbo Liu, Kai Zou, Yinan He, Fan Zhang, Lulu Gu, Yuanhan Zhang, Jingwen He, Wei-Shi Zheng, Yu Qiao, and Ziwei Liu. Vbench-2.0: Advancing video generation benchmark suite for intrinsic faithfulness, 2025. URL https://arxiv.org/abs/2503.21755.

[53] Shangchen Zhou, Chongyi Li, Kelvin C.K. Chan, and Chen Change Loy. Propainter: Improving propagation and transformer for video inpainting. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 10477–10486, October 2023.

[54] Siyuan Zhou, Hejun Wang, Hu Cheng, Jinxi Li, Dongsheng Wang, Junwei Jiang, Yixiao Jin, Jiayue Huang, Shiwei Mao, Shangjia Liu, Yafei Yang, Hongkang Song, Shenxing Wei, Zihui Zhang, Bing Wang, Zhihua Wang, Chuhang Zou, and Bo Yang. PhysInOne: Visual physics learning and reasoning in one suite. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 33131–33142, 2026. URL https://openaccess.thecvf.com/content/CVPR2026/html/Zhou_PhysInOne_Visual_Physics_Learning_and_Reasoning_in_One_Suite_CVPR_2026_paper.html.

[55] Adam Zweiger, Jyo Pari, Han Guo, Yoon Kim, and Pulkit Agrawal. Self-adapting language models. In *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pages 74084–74115, 2025. doi: 10.52202/085713-2483. URL https://proceedings.neurips.cc/paper_files/paper/2025/file/6b41e04c41726e2a60e456d0a2b961ab-Paper-Conference.pdf.

A Data Resources and Construction

A.1 Training-Data Construction

We construct training data by expanding physical event families into multiple scene configurations. For each family, we vary physical parameters such as mass, initial velocity, and material properties, together with object assets, backgrounds, lighting, and camera views. Each configuration undergoes simulation, verification, appearance refinement, and a final-video check before the output is paired with a text description. The training pair is linked to its generating configuration and native simulation records, including the rendered video, depth maps, instance segmentation, and object states.

Table 2 lists representative event families and the parameters varied to build their variants. A composite scene can combine several families; for example, a falling ball turns a wheel that releases a second ball into a basket.

Category	Representative events	Varied parameters
Rigid-body motion and contact	Projectile motion, bouncing, sliding, rolling, billiard collisions	Mass, initial velocity, friction, restitution
Rigid-body mechanisms	Domino chains, pendulums, Newton's cradles	Release angle, constraint length, object spacing
Elastic and deformable systems	Elastic-body impact, cloth draping, flag motion	Stiffness, damping, Young's modulus, wind velocity
Magnetic interactions	Attraction and repulsion	Field mode, strength, initial separation
Fluids	Container transfer, spilling, swirling flow	Initial volume, release height, initial velocity field
Fracture	Glass-panel and wooden-crate impact	Impactor mass, launch velocity, impact position, damage threshold

Table 2 Representative physical event families and the parameters varied to construct their variants.

Each family defines a parameter grid in its scene configurations; Table 3 shows examples. Because a parameter change can shift both event timing and the final state, each variant keeps its own simulation timeline.

Each sample record links the task, its scene configuration, the native observations, and the refined output. Native RGB, depth, and instance segmentation share timestamps and camera settings; object-state records contain identities, positions, orientations, and the contact and event records exposed by the simulator. Physics step rate and render frame rate are stored separately so that each saved frame maps to a simulation time. Depth and instance labels index the native render, not the refined video.

Multiple cameras observe the same simulation run. Cameras are placed around the objects and contacts that verification must observe, and a low-resolution layout preview checks scale, framing, and occlusion before the full capture. For example, the four-ball basket task is recorded from multiple viewpoints, with native RGB, depth, and instance segmentation.

Family	Parameter	Values
Ramp rolling/sliding	Friction (static, dynamic)	0.02, 0.12, 0.55
Newton's cradle	Release angle	25°, 35°, 45°
Elastic-body impact	Young's modulus	50, 100, 200 kPa
Pinned flag in wind	Wind speed	3.0, 6.5 m/s
Fluid volume release	Initial center height	0.55, 0.65, 0.75 m
Glass-panel impact	Impactor speed	1.5, 3.0, 5.0 m/s

Table 3 Example parameter grids from archived scene configurations.

Assets. Scenes combine assets from three sources: licensed content packages imported into the local catalog, Poly Haven environment maps and materials (CC0), and meshes generated with Meshy when no catalog asset fits. Every scene object records its asset source and license tier. Visual meshes are paired with separate collision proxies; for example, an open funnel uses a continuous conical collision shell matched to the visible inner wall.

A.2 PhysLoop-Gen

PhysLoop-Gen contains 100 tasks: 68 rigid-body tasks, 12 cloth tasks, 9 fluid tasks, and 11 image-conditioned tasks. The image-conditioned group is defined by its input modality and can involve rigid-body or other interactions. Each task contains a text request, optional image references, and success conditions that name the required objects, key interactions, event order, and final outcome.

Task difficulty is defined by the input: the number of interacting objects, required contacts, constraints, and event stages. Multistage tasks are checked stage by stage. In the image-conditioned paddle-wheel task (Fig. 4), a falling ball turns a wheel, whose blade pushes a second ball off its support and down a ramp into a basket.

A.3 PhysLoop-Repair

Each repair task contains a source video, an optional repair request, and evaluation-only annotations of the error, its interval, and the conditions for a successful correction. Error types include motion or contact failures, changes in object count or identity, missing support, wrong event order, and inconsistent material response.

Repairs follow the interval replacement defined in the main paper. A localized repair replaces the interval $[a, b)$ with exactly $b - a$ frames at the source resolution and frame rate, keeps blending inside the interval, and is checked for unchanged decoded frames outside it. When the error has already altered later events, the agent reconstructs the whole clip, which corresponds to $a = 0$ and $b = T$. Each task records which scope was used.

A repair succeeds when the target error is removed, the result connects to the surrounding events, required source content is preserved, and no new severe error appears. Evaluation always inspects the complete output video.

B Harness Execution Details

B.1 Verification Checks

Simulation-stage verification has three layers.

Task assertions. The task specification declares machine-checkable assertions over the native simulation record: existence, count, or order of contacts between named objects; thresholds on an object's state value or state change; trajectory integrity; and completeness of the required media and state artifacts. Contact evidence combines simulator hit events, when available, with oriented-box proximity tests using a 4 cm tolerance. Proximity-based contacts are geometric inferences, not native solver contact impulses. For the paddle-wheel task, the assertions require contact between the driving ball and a wheel blade, contact between another blade and the second ball, a displacement of the second ball of at least 0.45 m, and complete trajectories.

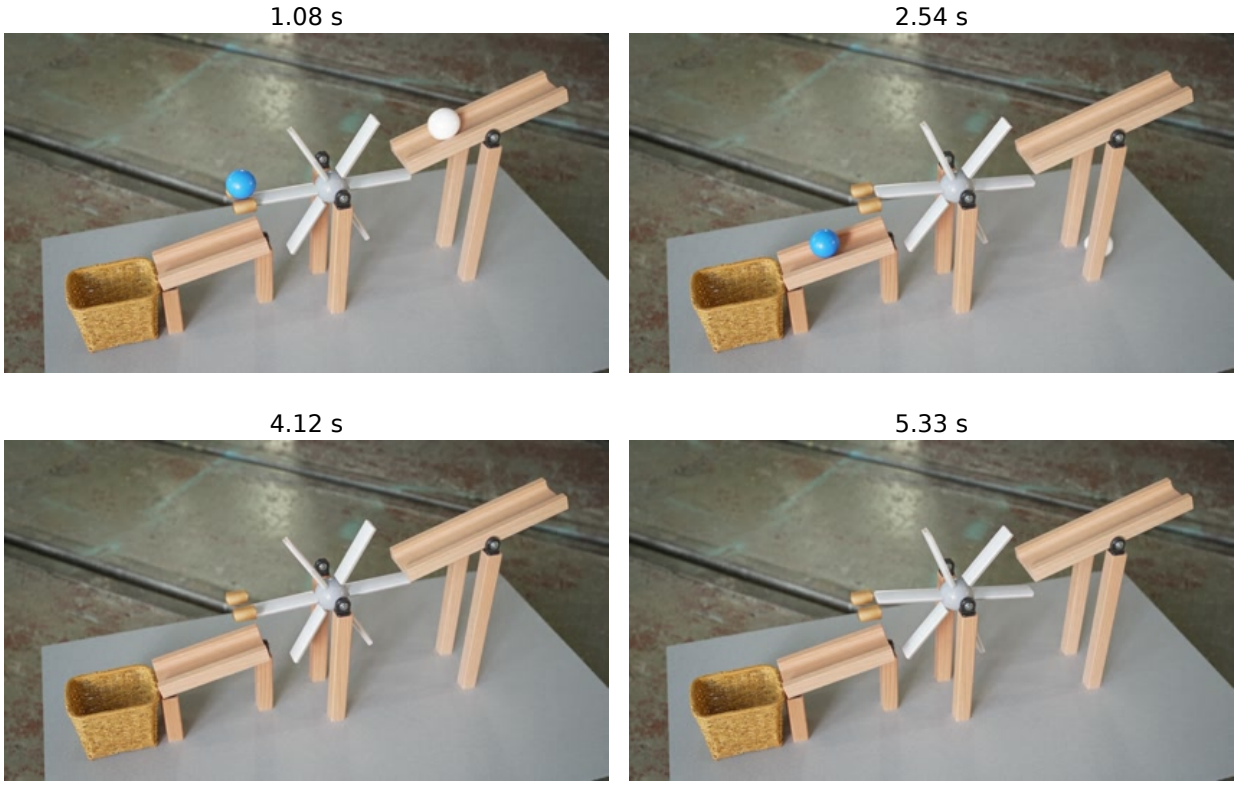

Figure 4 Example frames of the paddle-wheel task.

Technical quality gate. Independently of the task, the gate checks media validity (codec, frame count, frame rate), that each camera moves only as declared, that the initial scene is free of unintended contacts, that compiled object placements match the simulator's pre-simulation state, and, when a depth reference exists, that native depth agrees with the declared reference support surface (mean absolute error $\leq$1 cm, 95th percentile $\leq$2 cm). Joint frames must coincide within 2 mm and 0.1° at the initial pose.

Semantic review. A reviewer checks the evidence bundle against the task conditions. The bundle contains event-selected keyframes from each camera, the contact records, sampled trajectories and constraint states, and the quality-gate results. After refinement, the final-video check applies the same task conditions to the refined video, because refinement can alter motion that passed simulation.

B.2 Budgets, Stopping, and Human Involvement

Each task runs under a controller budget that limits scene-specification revisions, simulator launches, technical retries, and wall-clock time. Configured limits and consumed usage are recorded for each run. A run stops when all checks pass or a limit is reached; unresolved failures stay in the run record. Both PhysLoop and the standalone baselines permit at most five video-generation calls per task.

Users can review candidates during verification. Each human judgment, the change it requested, and the resulting revision are stored in the run record, so automated and human-requested revisions can be counted separately.

C Experimental Settings and Evaluation Protocols

C.1 Compared Configurations

Table 4 lists the nine configurations. Each provides one final output for each of the 100 PhysLoop-Gen tasks, for 900 videos in total.

Agent	Video model
– (text only)	Gemini Omni 1.1 Flash
– (text only)	Seedance 2.5
– (text only)	MiniMax H3
Claude Opus 5	Seedance 2.5
Kimi K3	Seedance 2.5
GPT-6 Astra	Gemini Omni 1.1 Flash
GPT-6 Astra	Seedance 2.5
GPT-6 Astra	MiniMax H3
GPT-6 Astra	– (native render, no refiner)

Table 4 Compared configurations. Each row contributes 100 final outputs.

Configuration records store the provider model identifier, invocation date, input references, resolution, duration, frame rate, and output-selection rule.

C.2 Cost and Runtime Accounting

We record token usage on ten PhysLoop-Gen tasks stratified by category and event stages, using GPT-6 Astra with Seedance 2.5. Each task uses a fresh session, including usage from failed attempts, revisions, and rejected requests. Simulation runs on a Mac mini (Apple M4, 10-core CPU and GPU, 16 GB) with UE 5.7; Seedance produces 1080p, 24 fps videos. Cached catalog assets are reused.

For the price comparison, we apply standard short-context input/cached-input/output rates of \$4/\$0.4/\$20 for GPT-5.6 Sol and \$10/\$1/\$50 for GPT-6 Astra per million tokens to identical usage. The ten-task totals are 199.05, 196.53, and 0.58 million tokens, respectively; output includes reasoning tokens, and cached input is a subset of total input. Figure 5 gives the resulting estimates.[1]

Both estimates use the same video-generation and upload charges, totaling \$42.99 over ten tasks. These are price estimates at fixed usage; local compute, electricity, and two requests with unresolved billing are excluded.

Excluding scene construction, UE simulation and rendering take approximately 2–6 minutes (median 3 min 25 s over 20 records), the low-resolution pre-check takes about 3 minutes, and Seedance refinement takes 2–3 minutes (median 2 min 13 s over 28 records). The combined generation time is estimated at 7–12 minutes (Fig. 5).

C.3 Automatic Evaluation

Physical commonsense (PC) and semantic adherence (SA) are scored by VideoCon-Physics on 32 sampled frames together with the task description. The archived evaluator inputs use the original Chinese task descriptions; no English rerun is mixed into these results. Imaging quality (IQ, MUSIQ-SPAQ) and aesthetic quality (AQ, LAION predictor) are computed by VBench on all frames. Stored scores are normalized to $[0, 1]$ (MUSIQ divided by 100, aesthetic score divided by 10) and multiplied by 100 for display.

For configuration m and metric q, scores are averaged over the cases $\mathcal{C}_q$ that are valid in every configuration:

$$\bar{s}_{m,q} = \frac{100}{|\mathcal{C}_q|} \sum_{c \in \mathcal{C}_q} s_{m,c,q}. \tag{2}$$

[1] https://developers.openai.com/api/docs/pricing, accessed 4 November 2026.

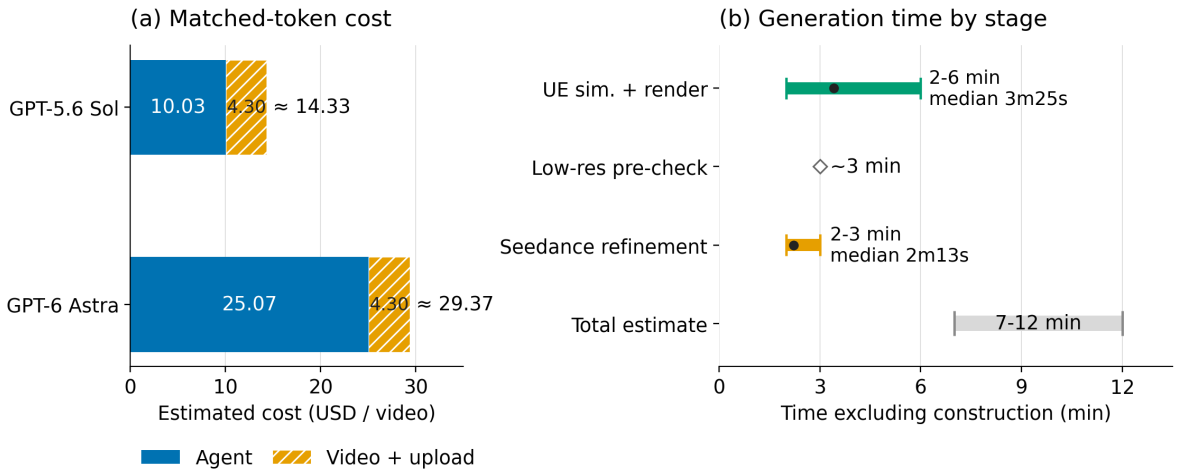

Figure 5 Matched-token cost estimates and generation times excluding scene construction.

The 900 videos give 3,600 metric records, of which 3,546 are valid. Six task descriptions exceed the SA input length in every configuration, so $|\mathcal{C}_{\mathrm{SA}}| = 94$; the other three metrics use all 100 cases.

C.4 Human Evaluation

Raters see the task description and one video, then give separate integer scores from 1 to 5 for physical plausibility, visual realism, and causal consistency. Figures 6 and 7 show the annotation interfaces. Raters judge observable evidence at normal playback speed: attractive appearance should not raise the physical-plausibility score, and physical correctness should not be equated with visual realism. N/A is used when an attribute cannot be judged; it is not a zero score.

Task completion is judged independently as pass, fail, or unjudgeable. An explanation is required for an unjudgeable response; raters can also explain their other judgments. A pass requires completion of the prompt's core event with plausible key physical relationships and causal order. Unjudgeable responses are excluded from the valid-vote denominator. Pairwise items show two outputs for the same task and ask for an overall preference (A, B, tie, or unjudgeable), considering all three attributes and task fulfillment without mechanically adding the scores. Method identities are hidden and presentation order is randomized.

Attribute scores are averaged first over the valid ratings of each video and then over videos, so each video has equal weight; N/A responses are treated as missing:

$$H_{m,a} = \frac{1}{|\mathcal{O}_{m,a}|} \sum_{c \in \mathcal{O}_{m,a}} \frac{1}{|\mathcal{R}_{m,c,a}|} \sum_{r \in \mathcal{R}_{m,c,a}} h_{m,c,r,a}, \tag{3}$$

where $\mathcal{R}_{m,c,a}$ are the valid ratings of attribute a for case c under configuration m, and $\mathcal{O}_{m,a}$ are the cases with at least one valid rating. Pairwise preference counts a tie as half a win and excludes unjudgeable items,

$$P_m = \frac{W_m + 0.5\,T_m}{W_m + T_m + L_m}. \tag{4}$$

Each pairwise contrast compares a UE-guided output with the direct output of the same video model. For task completion, a case passes only if pass votes strictly outnumber fail votes; unjudgeable votes are excluded before this comparison. Ties, cases without valid votes, and missing outputs count as failures. The task-success rate is the number of passing cases divided by the fixed 100-task set, rather than the fraction of individual pass votes. The 19 October pilot records reproduce 85/100 passing cases for the default refined route and 73/100 for text-only Seedance 2.5 under this rule.

The allocation targets three ratings per output across 20 assignment identifiers, each with 135 single-video and 45 pairwise items in six sessions of 30. The audited pilot snapshot dated 19 October 2026 contains 2,565 single-video response records and 855 pairwise records, of which 849 are valid judgments. It contains 19 returned answer sets under 18 original assignment identifiers. All 900 videos have at least one returned response, and 630 have at least three returned answer sets; this is not complete three-rating coverage. The archive retains its pilot-only designation, and assignment identifiers are not a verified count of independent participants.

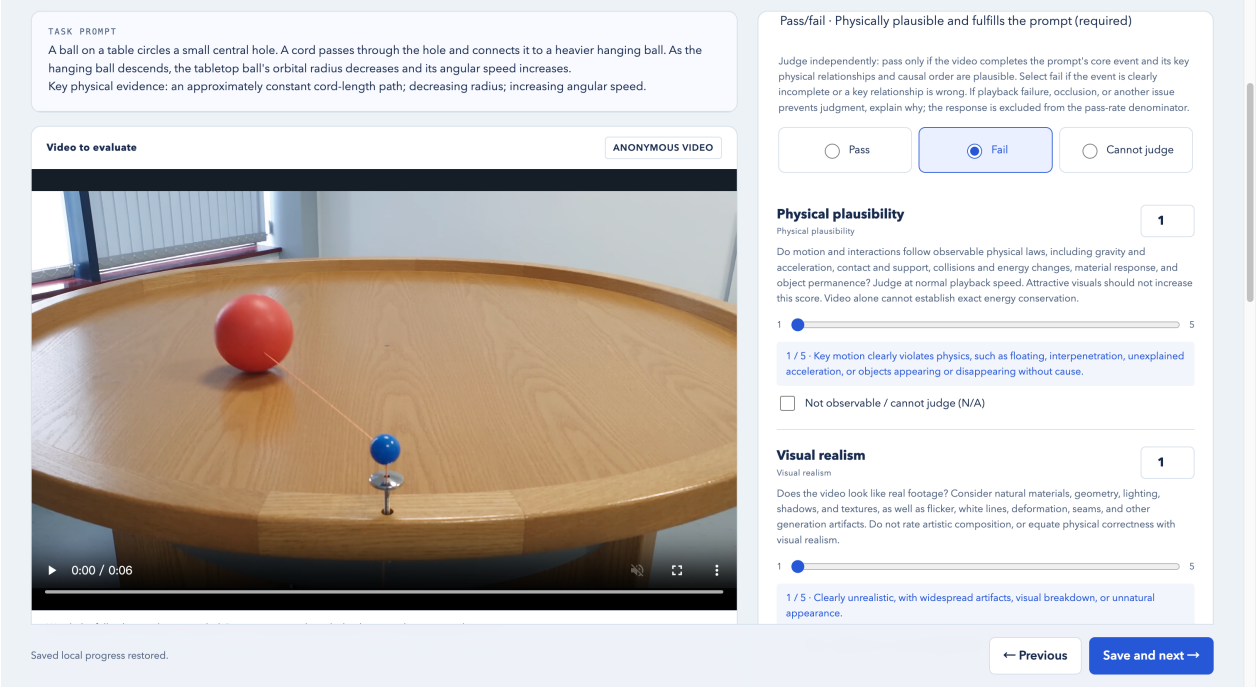

Figure 6 Single-video evaluation interface for task completion and attribute ratings.

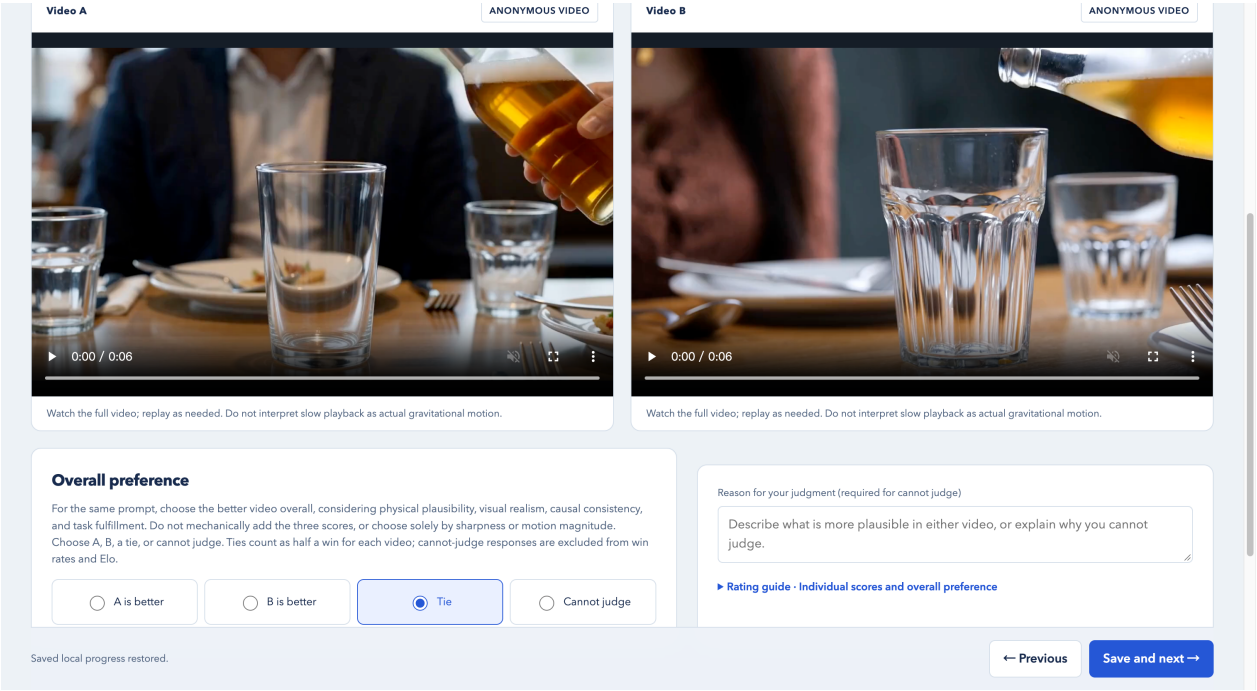

Figure 7 Pairwise evaluation interface with A, B, tie, and cannot-judge options.

C.5 Physics-IQ Verified

We report a single-view diagnostic on 66 Physics-IQ Verified cases using the left view and the first 120 output frames at 24 fps. This covers 66 continuations, rather than the 198 continuations of the three-view protocol. The two route types receive different conditioning: direct continuation receives the recorded three-second past clip, while UE routes receive the switch frame and a scene description from which the agent reconstructs the scene. Scores therefore measure agreement with the recorded continuation under each route's own input.

A primary benchmark result requires a candidate-selection rule fixed before scoring. Archived variants were developed with test feedback and are therefore reported as diagnostics until a final version is frozen. A single-candidate result uses one output per case chosen without benchmark scores. Where several variants were generated, we additionally report the retrospective best-variant score

$$S_{\text{best}} = \frac{1}{N} \sum_{c=1}^{N} \max_{v \in \mathcal{V}_c} s(c, v) \tag{5}$$

as a diagnostic only, together with the number of variants per case.

D Representative Prompts and Skill Rules

We use the paddle-wheel task (Fig. 4) to illustrate the prompts at each stage. Chinese originals are translated into English.

Task request.

> A steel ball falls from a short chute onto one blade of a six-bladed rigid paddle wheel. The wheel rotates about a horizontal axle, and another blade knocks a colored ball off its support. The second ball rolls down a short ramp into a basket.

Scene-construction request.

> Use the reference image for the layout of the wheel axle, blades, both balls, the support, the ramp, and the basket. The first ball must hit the wheel before it rotates; a blade must then hit the second ball, which leaves its support and enters the basket. The video must be at least five seconds long.

Refinement prompt (Seedance 2.5).

> *Scene and assets:* a wooden paddle-wheel chain in a real science museum: the original wooden chute, a six-bladed wheel on a horizontal axle, a white and a blue rubber ball, a short wooden ramp, and a yellow basket. Turn the static supports in the reference into sturdy wooden stands and metal bearing blocks; every raised part must have a visible load-bearing connection. Do not add moving parts or block the ball paths.
>
> *Motion:* the white ball rolls down the upper-right chute and hits the wheel, which turns about its axle; a blade knocks the blue ball down the lower-left ramp into the basket; the white ball falls to the right onto the display base and rolls out of frame. Motion and contact timing follow the reference exactly.
>
> *Shared constraints:* use the reference video as the only source of motion. Keep the number, identity, initial placement, trajectories, speed changes, contact timing, and final state of all objects. Change only the scene, appearance, and materials described above; do not add, duplicate, merge, or remove objects, and do not reorder or repeat events. Keep the reference camera and framing, with no text, cuts, or music. After the reference becomes still, hold the final state until the end.

Verification assertions.

> `event_exists`: contact(`drive_ball`, `blade_3`); `event_exists`: contact(`blade_1`, `second_ball`); `state_delta`: $\Delta x_{\text{second_ball}} \geq 0.45\,\text{m}$; `trajectory_integrity`; `artifact_complete`.

Skill rules. The agent follows a condensed `run-physics-video` skill:

1. Before continuing a task, `inspect` its state, stage, blocker, and allowed actions; keep the original request and references unchanged across revisions.
2. Submit an intent and scene specification that match the current generation context; the controller validates them and records execution.
3. Check placement and framing with `layout-preview`, dynamics with solver-backed smoke runs, and later-time framing with `replay-preview` on the cached trajectory.
4. Revise only the configuration fields implicated by the failure evidence; task conditions, acceptance rules, and usage counters are preserved. Technical retries resume from the last checkpoint.
5. Invoke `review` separately once the evidence bundle exists, and report completion only after the controller reaches `completed`.